

Capability Matters



A CASEBOOK

---

Learning Engineering for Next-Generation Systems

# Capability *Matters*

EDITED BY

William Gray-Roncal, PhD  
James Diamond, PhD

---

JOHNS HOPKINS UNIVERSITY  
LEARNING DESIGN AND TECHNOLOGY

EDITION · AUGUST 2026

An ongoing artifact of the LENS Practice Flywheel — Identify, Activate, Prototype, Analyze, Transition.

DEDICATION

---

*To everyone who believes in a dream and refuses to  
give up.*



## COLOPHON

---

### **Learning Engineering for Next-Generation Systems: Capability Matters**

A Casebook.

First edition, 2026.

Edited by William Gray-Roncal, PhD and James Diamond, PhD.

Compiled for the LENS concentration within the Learning Design and Technology program at the Johns Hopkins University School of Education.

Set in Instrument Serif and DM Sans.

Printed via Lulu at 8 × 10 in, perfect-bound.

Copyright © 2026. All rights reserved.

## HOW THIS VOLUME WAS MADE

This casebook was assembled using AI tools as part of an iterative learning-engineering process — the same Practice Flywheel (Identify → Activate → Prototype → Analyze → Transition) that the volume's case studies argue for. The methodology treats AI as the dual entity the curriculum names it: a creative partner that accelerated drafting, cross-referencing, layout, and citation lookup, and an epistemic risk that had to be checked against the record. Every case has been reviewed by the editors and checked against a seven-column rubric: that the sources cited exist and are findable, that the quoted attributions are fairly represented, and that the case is an accurate account of the incident and the investigation it draws on. Refinement is iterative and ongoing. Items where the source could not be confirmed are marked "Paraphrasing..." so the attribution is honest about what is the author's reconstruction and what is verbatim from the record.

The volume is not finished. It is the first iteration of a reference dataset that LENS will continue to develop. We will revise the book based on reader feedback, reviewer findings, and our own ongoing practice of learning engineering — adding new cases as the operational record grows, correcting our record where it can be improved, and updating the curriculum crosswalk as the program itself iterates. Corrections, new cases, and reader feedback are welcomed.

## SOURCES, ATTRIBUTION & LEGAL NOTICES

---

All cases reference publicly reported incidents and published investigations. Quoted material is attributed to source investigations, court documents, and academic publications. Reflection questions are original to this volume. The audit record of citations checked and changes made is maintained alongside the source repository.

References to convictions, settlements, and findings reflect the public record at the time of writing. Where individuals are named, they are named only as participants in investigations, court proceedings, or public inquiries already in the public domain. The casebook makes no legal characterization beyond what those records establish.

Quoted material is reproduced under fair use (17 U.S.C. § 107) for the purposes of criticism, comment, teaching, scholarship, and research. Trademarks and product names are used nominatively to identify the products and organizations under discussion; no affiliation or endorsement is claimed or implied.

Corrections, requests for clarification, and notice of any inaccuracy can be sent to LDT / LENS at the Johns Hopkins University School of Education. Substantive corrections will be incorporated in subsequent printings and noted in the casebook's public errata.

WELCOME

---

*Learning engineering is a toolbox, not a single tool.*

At its center it is a dialogue between two traditions — the engineering disciplines that design, build, and sustain complex systems, and the learning sciences and education that study how people come to know and do. Around that dialogue gather many other fields: cognitive and human-factors engineering, measurement, implementation science, design, data and computation, and the operational domains themselves. The work is plural by design — many disciplines at once — and it draws on domain knowledge, coding, theory, systems thinking, design, teaching, and analysis.

It welcomes everyone who brings a real tool to it. It is not, for that, a free-for-all: learning engineering stands in an intellectual tradition with its own methods, evidence standards, and lineage, and it holds the work accountable to them. No one carries every tool. This is a team sport — capability gets built at the seams between people who each see part of the system clearly and trust the others to see the rest.

So you don't need to know everything before you start. You need to know what you bring, stay honest about what you don't, and find the people who carry the tools you're missing. What you bring is yours to decide. The cases that follow are an invitation to find your part in the work.



INTRODUCTION

---

# Capability is a system parameter.

*Every complex system depends on what people can do inside it. That dependency is measurable, designable, and too important to leave to chance.*

- 01. Systems don't fail because the technology breaks. They fail because someone couldn't do what the system required.
- 02. We engineer every parameter of a system except the people.
- 03. Capability is not a soft problem. It is a systems engineering problem.
- 04. Decision-grade evidence. Not opinions about training.
- 05. Capability without agency is automation. The goal is operators who can perform, adapt, and lead — not comply.
- 06. Sometimes the constraints make the goal unreachable. Saying so early is a result, not a failure.

The cases in this book are not, individually, hard to understand. A pilot recovered from a stall the wrong way. A nurse could not stop a doctor who skipped a step. A safety operator was watching television on her phone when a pedestrian crossed the road. A captain shot down a civilian airliner because the radar data and the operational framing pointed in opposite directions, and the framing won. The proximate cause is usually obvious. The first investigator on the scene almost always finds it within hours.

What is hard to understand — and what this book exists to make legible — is the system that produced the gap into which each incident fell. In every case the proximate actor was operating inside an architecture of training, procedure, authority, measurement, and incentive that had, for years, been quietly degrading. The pilot had never trained to recover a stall at cruise altitude. The nurse had no procedural path to intervene. The safety operator had been hired into a role nobody had figured out how to make performable. The radar operator had been trained to a tempo and a presumption that the system above him no longer shared with the system around him.

Each of these architectures was designed. Each was funded. Each was reviewed. Each carried, written somewhere in its specifications, the assumption that the people inside it would be able to do what the system required of them when the system required it. That assumption is what this book calls the **capability parameter**. It is a property of the entire sociotechnical system, not of any individual operator. Capability lives at the **interface** between what a system requires of its operators and the impact the system has to deliver — and when that interface is wrong, the gap may originate in

human development, in system design, in organizational performance, or in the interaction among them. Distinguishing which is itself a measurement and engineering problem; the casebook calls it **gap attribution**, and treats it as the diagnostic move that the rest of the discipline depends on. When the interface fails — wherever the source — the consequences are paid in lives, in dollars, in trust, and in time the institution will never recover.

The premise of the Johns Hopkins School of Education's Learning Design & Technology program — and of the Learning Engineering for Next-Generation Systems concentration that has grown out of it — is that the capability parameter is engineerable. Not by accident. Not by declaration. By the same kind of evidence-grounded, methods-grounded, cross-domain discipline that built modern reliability engineering, modern epidemiology, and modern systems safety. The discipline does not yet exist at the scale the evidence demands. LDT and LENS are organized to help build it.

## I · THE COST OF THE GAP

The Institute of Medicine's **To Err Is Human** (1999) estimated that preventable medical error caused between forty-four thousand and ninety-eight thousand deaths a year in the United States — already, at the lower bound, a top-ten cause of death. The report catalyzed a generation of patient-safety work: surgical checklists, central-line bundles, team training, computerized order entry. Some of those interventions are documented in this book's What Works chapters as successes. And yet in 2016 Makary and Daniel returned to the question from the same institution that produces this casebook and estimated that medical error kills more than two hundred fifty thousand people a year in the United States, which would rank it third among causes of death <sup>1</sup>. That estimate is substantively contested on methodological grounds, and Case 8 sets out the dispute; what is not contested is that the count has not moved the way the reform arc promised.

The interval between the IOM report and the Makary update is approximately the canonical gap implementation science has identified between research evidence and clinical practice: about seventeen years <sup>2</sup>, with only about fourteen percent of research findings reaching practice at all <sup>3</sup>. Both are estimates rather than measurements — the first a lag timed across nine mostly primary-care services, the second derived from assumed attrition along the publication-to-guideline pipeline — and Case 13 shows the derivation. They are the right order of magnitude and the wrong kind of number to quote precisely. The system to surface them, vet them, deploy them, sustain them, and measure their effect at the scale at which they would matter has never been built as an institution. It has been built as a series of grants. The difference is the point.

This is not only a healthcare problem. The U.S. Navy lost seventeen sailors in two avoidable destroyer collisions in 2017 after cutting its Surface Warfare Officers School course down to self-study CD-ROMs in 2003 <sup>4</sup>. The world's most expensive aircraft program flew at a 44% mission-capable rate in FY2025, down from 67% in FY2021, because the platform was fielded faster than the maintainer pipeline could be built <sup>5</sup>. One hundred million dollars of educational infrastructure dissolved in fourteen months because the institution building it did not engineer the governance and trust that the technology presupposed <sup>6</sup>. A standardisation algorithm downgraded some 39% of England's teacher-estimated A-level grades in a single afternoon and was withdrawn four days later, because the agency had tested whether the calculation worked and not what it would do to the individuals it was calculating <sup>7</sup>.

None of these failures is a failure of effort. None is a failure of intelligence. Each is a failure of the capability infrastructure to match the system it was operating. In each case, the gap was diagnosable in advance. In most cases, it was actually diagnosed in advance, by someone — and the diagnosis did not produce the remediation. The cost of the unrepaired diagnosis is why **capability** matters.

## II · WHAT AN ENGINEERABLE DISCIPLINE LOOKS LIKE

Crew Resource Management did not exist as a discipline in 1976. In March 1977 two 747s collided on a foggy runway at Tenerife and 583 people died. Within five years United Airlines had operationalized the first CRM curriculum <sup>8</sup> within twenty years the Commercial Aviation Safety Team had built the closed-loop hazard-identification system that lets the FAA know whether CRM is still working <sup>9</sup>. U.S. commercial aviation fatality risk fell 83% over 1998–2008 and fatalities per hundred million passengers fell 95% across the longer arc <sup>10</sup> — a portfolio result across some seventy safety enhancements, of which CRM is one and from which its separate contribution has never been isolated. Case 117 is careful about this, and so should the reader be: the arc is evidence that the discipline paid, not a measurement of what any one part of it bought. CRM works not because it taught individual airmanship — pilots were already excellent — but because it redesigned the **system of interaction** in the cockpit: the authority gradient, the communication protocol, the standard brief and debrief. The discipline treated team coordination as an engineerable property of the system.

The same pattern shows up in every domain where the cost of failure has been high enough to fund the discipline. The Kemeny Commission reported in October 1979 that the industry should set and police its own standards of excellence; by December — eight months after Three Mile Island — the U.S. nuclear industry had stood up the Institute of Nuclear Power Operations, because every utility had recognized that an accident at any single plant would affect every operator's license to operate <sup>11</sup>. INPO and the National Academy for Nuclear Training, founded in 1985, have presided over more than four decades without an INES Level 4-or-above accident at a U.S. commercial reactor. The comparison with the surface Navy, which cut its training during the same era, is the cleanest available test of what happens when capability is engineered versus when it is deferred.

Healthcare has its own version of this arc, but only partially. Peter Pronovost's central-line checklist, paired with the nurses' authority to enforce it, drove bloodstream infections to near zero across 103 Michigan ICUs and saved an estimated fifteen hundred lives in eighteen months <sup>12</sup>. Atul Gawande's nineteen-item surgical safety checklist, paired with three required pauses in the operative timeline, cut the death rate from 1.5% to 0.8% across eight hospitals in eight countries <sup>13</sup> — though a later province-wide Ontario rollout found no significant mortality benefit, which is the transfer problem this book keeps returning to. TeamSTEPPS, jointly developed by AHRQ and DoD, translated decades of high-reliability research from aviation and the military into clinical practice in years rather than decades — because the implementation infrastructure was funded as part of the intervention rather than as an afterthought <sup>14</sup>.

These are not stories about individual hospitals or individual ships. They are stories about the difference between a domain that has organized itself to engineer capability and one that has not. The pattern is consistent across all of them: a catastrophe makes the capability gap visible; an institution is built that treats the gap as the institution's responsibility; the institution funds the measurement architecture that lets it know whether the gap is closing; and twenty years later the death rate is half what it was. The intervention is always paired. A technical artifact — a checklist, a brief, a cord

— combined with a cultural artifact — protected authority to use it, a no-blame protocol, a credible governance body. Neither alone moves the curve. Both together move it dramatically.

### III · THE DISCIPLINE WE HAVE, THE DISCIPLINE WE NEED

The intellectual material to build this discipline already exists. The learning sciences have produced four decades of converging evidence on how skill, knowledge, and judgment develop and decay <sup>15</sup>. Cognitive engineering has produced rigorous methods for analyzing human-machine systems <sup>16</sup>. Human factors and ergonomics have produced an evidence base for interface and workflow design <sup>17</sup>. Implementation science has produced frameworks for taking effective interventions to scale <sup>18</sup>. Systems safety engineering has produced both diagnostic tools — Reason’s swiss-cheese model, Rasmussen’s migration-to-the-boundary, Leveson’s STAMP — and control-theoretic methods for design <sup>19</sup>. High-reliability- organization theory has documented what the institutions that have solved the capability problem actually do <sup>20</sup>. The discipline is not missing its evidence base. It is missing its integration.

Learning engineering is the integration. The term, in its modern usage, traces to Bror Saxberg and was elaborated by Goodell, Kolodner, and the IEEE Industry Connections Industry Consortium on Learning Engineering (ICICLE) <sup>21</sup>. The IEEE Learning Engineering Body of Knowledge (LEBoK) and its associated process model are the most developed available specification of what the discipline does <sup>22</sup>. The premise is that the work of building, deploying, and sustaining capability at scale is an engineering activity: it requires explicit requirements, evidence-based design, instrumented measurement, and feedback-driven iteration. It is not exclusively a research activity. It is not exclusively a managerial activity. It is engineering, and it can be taught and practiced as engineering.

LENS — Learning Engineering for Next-Generation Systems — is a concentration within the Johns Hopkins School of Education’s Learning Design & Technology program that operationalizes this premise for high-consequence operational domains. Its five competencies — systems analysis; iterative development; human-system collaboration; test and evaluation; and navigating sociotechnical constraints — are organized to produce graduates who can do the work the cases in this book required, and who can build the institutions that the book’s success cases had to invent. The curriculum threading commitments — implementation science as a through-line, equity as a design commitment rather than an audit, decision-grade evidence as a deliverable rather than a report, communication and system-of-systems integration as engineerable properties of the system, the disciplined use of learning- technology aids (from XR to adaptive platforms) as instruments whose evaluation is part of the work, and **AI as both a creative partner and an epistemic risk** — leveraging AI’s generative reach while guarding against offloading the thinking — are drawn directly from the patterns the cases reveal.

Three of those commitments deserve a moment of their own — they are the substrate the rest of the curriculum runs on. The first is **communication, translation, and integration within and across disparate systems**. A striking number of cases in this book are, at their proximate cause, translation or integration failures across boundaries — failures of language, convention, unit, role, agency, or discipline to carry meaning reliably from one part of a system to another, or failures of subsystems built to different assumptions to compose into a working whole. Mars Climate Orbiter was a metric-versus- imperial translation failure (Case 98). Tenerife was a translation failure at the takeoff-clearance boundary between two cockpits and a tower under fog and time pressure (Case 117). The Patriot at Dhahran failed because the manufacturer’s assumption about run time and the



operator's actual run time were on opposite sides of a documentation boundary that had not been engineered to be crossed (Case 129). Eagle Claw failed because the rotary-wing, fixed-wing, ground, and command components were assembled across service boundaries that had no shared operating practice (Case 128). 9/11 was a cross-agency translation failure measured in thousands of dead (Case 135). Boeing 737 MAX was, at one level, an integration failure: a single-string sensor, a flight-control law that assumed pilot mental models from prior models, and a certification chain that did not knit the pieces together (Case 97). The successful cases — AlphaFold (Case 36), MICrONS, CRM (Case 117 paired with Case 118), Keystone (Case 19), and the paired-intervention examples that close every part — are, equivalently, cases of disciplined translation and integration: biology into computation, science into operational practice, technical reform into cultural reform, multiple subsystems into a working whole. Engineering capability across these boundaries — language, discipline, agency, subsystem-to-subsystem, system-of-systems — is itself a designable target. The LENS curriculum treats it as such.

The second is **technology aids and learning platforms**. The discipline has at its disposal a fast-evolving toolkit: extended-reality (XR) simulation environments — VR, AR, and mixed-reality systems used for procedural training, spatial-reasoning practice, and high-fidelity rehearsal under conditions that would be unsafe or impossible in the field — together with learning-management systems, adaptive learning platforms (the lineage from Cognitive Tutor and ASSISTments to current ITS work; Cases 67, 139), intelligent tutoring frameworks such as GIFT, learning-analytics infrastructures such as xAPI and the Total Learning Architecture (Case 142), game-based learning environments, LLM-augmented tutors and authoring tools, and high-fidelity simulators in aviation, surgery, defense, and process industries. These tools are not the discipline. They are the discipline's instruments — and the cases in this book show, repeatedly, that powerful instruments do not by themselves engineer capability (Cases 53, 54, 2, 139, 22, 157). The LENS curriculum teaches the practitioner to choose, configure, evaluate, and govern these tools with the same evidentiary discipline applied to any other intervention, and to recognize when the binding constraint is the surrounding architecture rather than the tool itself.

The third — and the design constraint on all the others — is **human agency**. The capability discipline can be misread as optimization: produce operators whose behavior matches a specification. That reading collapses the very property the cases reveal as decisive. Every successful institution in the What Works chapters — INPO operators trained to challenge their reactor team, Keystone ICU nurses authorized to halt a procedure (Case 19), the Nuclear Navy's questioning attitude (Case 137), the first-officer authority CRM built into the cockpit (Case 117) — engineered for capability **and** for the operator's capacity to exercise independent judgment, to adapt to conditions the system was not designed for, and to shape the surrounding architecture rather than only execute inside it. Capability without that capacity is automation. LENS treats the preservation and development of agency — the room the system leaves its operators to think, decide, and lead — as a design constraint on every intervention, including those in which AI is the partner. The paired risk runs the same direction: an AI collaborator can accelerate the flywheel or quietly offload the thinking, and which one it does is a design and governance decision, not a property of the tool.

#### IV · THE METHOD · UNPACKING THE CASE

Learning engineering is not a one-time act of design. Its defining feature, as articulated in the IEEE Industry Connections Industry Consortium on Learning Engineering's process model, is iteration: understand the operational context, specify capability requirements, prototype an intervention, in-

strument it, measure, learn, redesign. The work cycles. The cases in this book that produced sustained outcomes — CRM and CAST, INPO, Keystone ICU, the WHO Surgical Safety Checklist, TeamSTEPPS — all share that loop structure. They were not built once and left alone. They were built, measured, rebuilt, remeasured, and rebuilt again. The institutional capability to sustain the loop **is** what distinguishes them from interventions whose effects decayed.


The loop also has a legitimate negative result, and the discipline has to be honest enough to reach it. Capability engineering always operates inside constraints — budget, schedule, regulation, institutional will, politics, ethics, public trust — and many of the binding ones are not technical. Part of the work is **managing those constraints and the risk they carry**: reading the constraint space accurately enough to recognize when external, non-technical factors have narrowed the solution space to the point where a particular capability goal is no longer realistic to pursue. Reaching that conclusion — and documenting **why** — is itself a valid outcome of a project, not a failure of one. It redirects effort toward goals that are actually reachable, surfaces the non-technical barriers for the stakeholders who can change them, and prevents the far more expensive failure of driving an infeasible target all the way to the operational record before it collapses there. A number of cases in this book are, in part, stories of institutions that could not say “not like this, not yet” in time.

The casebook turns this loop into a reading method. Each case is read, diagnosed, paired with a learning-engineering response, and tested in studio against students’ own operational domains. The cycle — framing the system whose capability is at stake, eliciting requirements from operational analysis, building the evidence architecture that lets an institution know whether its interventions worked, examining the human and AI roles inside the designed system, and running the loop end-to-end on a real problem — is the same discipline the casebook applies to every case it reads. The cases that return as success stories in the What Works chapters are the cases whose institutions completed the loop. The What Fails chapters hold the cases whose institutions did not — or did not loop fast enough. Unpacking the case **is** the method: identify, instrument, and close the next such loop before the next case is written.

## V · THE ANALYTIC LENS · THE FIVE CAPABILITIES THE CASES TURN ON

The cases in this book turn on five capabilities. They are the analytic lens the casebook reads each case through, and the cases that follow are tagged to them: where capability was absent, one or more of these five was the capability that was missing; where it was engineered, one or more of these five was the capability that was built. Named as the five LENS Educational Objectives — the concentration tier of the program’s objectives — they are also what the discipline sets out to teach.

## What LENS graduates can do.



### WHAT LENS GRADUATES CAN DO — FIVE COMPETENCIES

#### 1 · SYSTEMS ANALYSIS — SEE THE WHOLE SYSTEM

##### SYSTEMS THINKING AND ANALYSIS

Decompose system performance requirements into measurable human capability requirements; apply systems-engineering lifecycle models to scope, sequence, and evaluate capability interventions **within and across disparate systems**; analyze and communicate the interdependencies between human performance and system performance to predict operational impact at scale.

#### 2 · ITERATIVE DEVELOPMENT — BUILD, TEST, REFINE

##### LEARNING ENGINEERING DESIGN AND IMPLEMENTATION

Design capability-development interventions that integrate learning-sciences principles with measurable outcomes and system-design constraints, and that function at **speed and scale** in operational environments; construct and communicate implementation plans that address adoption, sustainment, and lifecycle integration across organizational and system boundaries.

#### 3 · HUMAN-SYSTEM COLLABORATION — WORK TOGETHER

##### HUMAN-SYSTEM COLLABORATION AND ADAPTIVE SYSTEMS

Design role architectures, interface and alert systems, mode/state transparency, authority gradients, and recoverability mechanisms that engineer collaborative capability at the human-system boundary — including human-AI partnership as one sub-pattern; evaluate the measured impact of system-mediated work on human performance, naming automation bias, cognitive offloading, and skill atrophy; specify **delegation with revocation** — the disconfirming evidence that would revoke a delegation — and measure collaboration as the unit of analysis.

#### 4 · TEST AND EVALUATION — SHOW WHAT WORKS

## DATA, MEASUREMENT, AND EVALUATION

Design ethical instrumentation strategies that produce measurable evidence in regulated or high-consequence environments; construct **decision-grade evidence** artifacts — dashboards, reports, governance documentation — under irreducible uncertainty, naming what is known, what is assumed, and what would change the decision; differentiate capability gaps from system-design and organizational-performance problems using **gap attribution**; demonstrate that omitting a protected attribute does not establish fairness.

## 5 · NAVIGATING SOCIOTECHNICAL CONSTRAINTS — MAKE IT WORK IN THE REAL WORLD

## CONTEXT AND DOMAIN FLUENCY

Analyze the regulatory, organizational, and cultural constraints that shape capability development in healthcare, defense, or education contexts; apply human-systems-integration frameworks to evaluate the measurable impact of capability approaches on operational environments; synthesize and communicate stakeholder requirements across disciplinary and institutional boundaries into coherent design specifications, including cross-regime / platform-dependency governance where capability is deployed on a platform governed by a different regime than the one operating it.

*See the whole system. Build, test, refine. Work together. Show what works. Make it work in the real world.*

Three phrases recur across those five capabilities and across the cases that follow: **within and across disparate systems** (the scope at which capability is actually engineered), **speed and scale** (the operational tempo the discipline must match), and **gap attribution between learning, system design, and organizational performance problems** (the diagnostic that decides what kind of intervention will actually move the outcome). They are the working language the cases are read in.

One phrase inside those capabilities is worth unpacking, because it is the one most often misread: **decision-grade evidence**. **Evidence**, in LENS, is knowledge about a system's **functional outputs** — the capability its humans and its technology deliver together, not either half alone — and it arrives at very different grades of certainty depending on where it comes from, from a widely validated theory to a single observation self-reported once. **Decision-grade** does not mean the uncertainty is gone; it means the evidence is sufficient to justify an action — enough to frame the choice, commit to it, and know in advance when the decision must be revisited. What the practitioner does with it is **reasoning under uncertainty**: weigh the veracity of the information and the sociotechnical context of the problem, then act — or deliberately decline to. It is a sufficiency judgment under irreducible uncertainty, not a threshold to be cleared. The LENS Companion decomposes the term in full.

## VI · HOW TO READ THIS BOOK

The book is organized topically, in seven parts: Healthcare & Patient Safety; Education, Training & the Learning Workforce; Aviation & Aerospace; Defense & National Security; Industry, Energy & Enterprise Systems; Disaster Prevention & Recovery; and Algorithms, Governance & Public Systems. Each part

splits into two chapters. **What Fails** collects the part's failure cases; **What Works — and the Frontier** collects the interventions that moved the curve and the frontier cases where the discipline is still being asked to specify what good looks like before the catastrophe arrives.

Two analytic threads run across the topical parts. The first is the failure-mode taxonomy carried on every case header: training gap, capability designed out, normalization of deviance, interface failure, governance failure, and knowledge loss. The taxonomy is not a theory. It is a finding: six categories that account for almost every well-documented case in the literature of preventable failure in complex sociotechnical systems, and that recur across every part of this book. The second is the **paired intervention**. Every success case in the dataset shares it: a technical artifact paired with a cultural authority; a measurement instrument paired with a governance body that listens to it; a curriculum paired with the institutional architecture to sustain it. The pattern is that the cases that moved the curve shipped both halves together. What the record does not establish is which half carried the result: these interventions were installed as bundles, and the trials that measured them could not decompose them. That is the strongest empirical pattern in the book, stated at the strength the evidence supports, and it informs the LENS commitment to co-optimization across technical and human design.

**Cases have more than one cause, and the book says so.** Each case names a primary reading — the one its LENS anchor is drawn from, and the one the Lens page teaches. It is rarely the only reading the record supports. Where rival explanations survive the evidence and imply different things to check before the pattern is carried somewhere else, they are set out in a **Competing readings** block at the end of the narrative; where a case demonstrably fails to show something it is often made to show, that boundary is stated as **What this case does not show**. Neither block is a caveat attached to a claim we did not want to defend — an error is corrected in the case, not annotated. They are there because a reader who takes one cause away from a failure with four interlocking ones has learned the wrong lesson, and because a student who meets the alternatives here will not be ambushed by them later. Therac-25, the opening case, carries both blocks as the worked example.

One part departs from the domain logic. Part VI, **Disaster Prevention & Recovery**, reads its cases along the lifecycle of a disaster instead: before the event — the prevention regimes whose quiet thinning is the real story — and after it, the response, the recovery, and the institutions built (or not built) so the same failure cannot recur. Cases whose central lesson belongs to a domain stay in their domain parts; what qualifies a case for Part VI is that its lesson is the pre/post lifecycle itself.

Most cases run to four pages; the shortest run to three and the densest to five. The narrative comes first — what happened, the evidence, the attribution — and closes with a diagram; where the record supports rival explanations, the **Competing readings** and **What this case does not show** blocks follow it, then a pull quote and the numbered sources. The case ends in the **Learning Engineering Lens**: a synthesis of what the case teaches, how the LENS curriculum addresses the pattern, and reflection questions designed for studio discussion. The Lens begins wherever the narrative ends rather than on a page of its own. The casebook is built to be taught from, read straight through, or consulted as a reference. The curriculum crosswalk in the closing matter maps each case to the specific courses in the program that use it as a worked example.

**Reading the figures.** Each figure marks one focal value — the quantity the case turns on. Where a figure carries a glyph on that value, the glyph states the direction of the outcome, and states it twice so that

it survives both the colour and the black-and-white printing: a **star** for a favourable outcome, a **triangle** for an adverse one, in deep teal-green and in burnt orange respectively where colour is available. Figures whose focal value is not an outcome — a mechanism, a subject, a destination — carry it in a single neutral accent instead. Read the glyph where it appears; where none appears, read the caption.

**How these cases were chosen.** The collection is purposive, not systematic. The cases are exemplars — drawn from the documented public record and from the editors' own practice — selected because each makes a capability-engineering pattern legible, not because they constitute a representative or exhaustive sample. The corpus over-represents catastrophic, well-investigated failures, which generate the richest public record, and under-represents the everyday, program-scale work that fills most of a practitioner's career; a deliberate expansion has begun to close that gap. Read the book as a teaching instrument rather than a systematic review: the claim is that these patterns recur and are engineerable — not that this is every case, or the average one.



## HOW TO USE THIS BOOK

# A field manual for capability engineering.

This book collects a growing record of real incidents in which capability — what people could or could not do inside a complex system — determined whether the system worked. Some are failures: lives lost, systems wrecked, billions written off. Some are successes: deliberate interventions that engineered capability into the architecture of the work. Together they form an evidence base for the argument that capability is a designable, measurable property of every complex system.

Most cases run to four pages, beginning on a fresh page rather than a fixed side of the spread. An **In brief** summary opens; the narrative follows in five sourced beats — what happened, what investigators found, and the documentary record — closing with a figure, any **Competing readings** or **What this case does not show** the record requires, a pull quote, and the numbered sources. The case ends in the **Learning Engineering Lens** — what the case teaches about capability as a system parameter, how the LENS curriculum addresses the pattern, and a short set of reflection questions designed for studio discussion.

The Lens also carries a **Who Builds This** note: the mix of expertise and tools a team would bring to engineer the fix, read off the case's failure modes. A Training-Gap case pulls in learning scientists and instructional designers; an Interface case, human-factors and interaction designers; a Governance case, policy, ethics, and measurement. Most cases need several at once, held together by domain experts and a learning engineer. No single discipline carries a case alone — reading the modes as a staffing list is itself a LENS skill, and what **you** bring is yours to decide.

## READING THE TAXONOMY

Each case is tagged with one or more failure modes. Most cases involve several. The primary mode is listed first; contributing modes follow.

- T** Training Gap  
Personnel were never taught what they needed to know.
- D** Designed Out  
Human capability was deliberately removed from the system.
- N** Normalization of Deviance  
Known risks were accepted as routine.
- H** Human-System Interface  
Correct performance was made unreasonably difficult.

## LENS COURSE CODES

Each case is tagged in its metadata to the courses it serves as a worked example for. The tags are not printed on the cases; the mapping is inverted in the course index at the back.

### SHARED LDT FOUNDATIONS · ADDITIONAL CONTEXT

- F1** **Learning Sciences Studio.** How learning, motivation, and skill develop — applied to technology-mediated design.
- F2** **Critical Perspectives on Educational Technology.** How power, equity, and society shape — and are shaped by — learning tech.

### LENS COURSES · † ELECTIVE

- LEN 1** Principles of Learning Engineering for Complex Systems



**G Governance & Trust**

Ethics, accountability, and stakeholder trust were after-thoughts.

**K Knowledge & Institutional Memory**

Organizational knowledge degraded or was never captured.

LEN 2

LEN 3

LEN 4

LEN 5

LEN 6

LEN 7

LEN 8

LEN 9

LEN 10

Human-Machine Teaming and Adaptive Learning Systems  
Learning Engineering Systems  
Evidence, Analytics, and Measurement for High-Consequence Domains  
Human Capability Analysis and Requirements †  
Applied Problem Solving in Learning Engineering  
Bias, Risk, and Governance in Learning System Design †  
Knowledge Transfer and Organizational Learning †  
Computational and AI Methods for Learning Engineers †  
Learning Engineering Project (capstone)

*The strongest cases pair a catastrophic failure with a systematic capability intervention. Together they show both the cost of the gap and the tractability of the solution. The discipline that closes that gap is learning engineering.*

## THE CASE MATRIX

## The cases at a glance

Gold numbers indicate failures, frontier cases, and systemic conditions; teal numbers indicate paired-intervention successes. Numbering is global and a few numbers are absent: cases withdrawn in review keep their numbers rather than have the sequence re-flowed around them. Numbers marked ° appear in the complete digital edition but not the printed main volume; unmarked cases appear in both.

|            |                                                                                                    |                |                  |
|------------|----------------------------------------------------------------------------------------------------|----------------|------------------|
| <b>1</b>   | Therac-25                                                                                          | 1985 – 1987    | <b>H-D-G</b>     |
| <b>2°</b>  | EHR / CPOE Implementation                                                                          | 2005 – present | <b>H-D-G</b>     |
|            | Watson for Oncology — Delegation                                                                   |                |                  |
| <b>3°</b>  | Marketed Ahead of Capability                                                                       | 2013 – 2018    | <b>D-K-N</b>     |
| <b>4°</b>  | Mid Staffordshire NHS Foundation Trust                                                             | 2005 – 2009    | <b>G-N-K</b>     |
| <b>5</b>   | Epic Sepsis Model                                                                                  | 2017 – 2021    | <b>D-G-N</b>     |
| <b>6</b>   | Racial Bias in Pain Assessment — The False-Belief Mechanism                                        | 2016           | <b>T-K-N</b>     |
| <b>7°</b>  | VA Wait-Time Scandal                                                                               | 2014           | <b>G-K-N</b>     |
| <b>8°</b>  | Medical Errors as Systemic Failure                                                                 | 1999 – present | <b>T-H-N-K-G</b> |
| <b>9°</b>  | Vioxx Withdrawal                                                                                   | 1999 – 2004    | <b>G-D</b>       |
| <b>11°</b> | California Nurse Staffing Ratios — Legislating a Capability Requirement                            | 2004 – 2010    | <b>G-K</b>       |
| <b>12°</b> | ACGME 80-Hour Resident Duty-Hour Reform                                                            | 2003–2017      | <b>T-K-N</b>     |
| <b>13°</b> | Implementation Science in Healthcare — The 17-Year Gap                                             | ongoing        | <b>K-G-N</b>     |
|            | Barsuk SBML — Simulation-Based                                                                     |                |                  |
| <b>14°</b> | Mastery Learning Dissemination from Northwestern to the VA                                         | 2009 – 2020s   | <b>T-K</b>       |
| <b>15°</b> | I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer                                    | 2014           | <b>H-K-N</b>     |
| <b>16°</b> | NCSBN National Simulation Study — Licensing the 50% Substitution Rule                              | 2014           | <b>G-K-D</b>     |
| <b>17°</b> | Spaced Education RCTs in Medical Training                                                          | 2007 – 2009    | <b>H-K-D</b>     |
| <b>18°</b> | PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption | 2023           | <b>K-N-D</b>     |
| <b>19</b>  | Keystone ICU / Pronovost Checklist                                                                 | 2004 – present | <b>T-N</b>       |
| <b>20</b>  | TREWS / Sepsis Watch                                                                               | 2018 – 2022    | <b>H-K-G</b>     |
| <b>21°</b> | Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff                   | 2024 – 2025    | <b>H-K-N</b>     |
| <b>22°</b> | ChatGPT in Healthcare — Hallucination Cases                                                        | 2023 – present | <b>H-D</b>       |
| <b>23°</b> | WHO Surgical Safety Checklist                                                                      | 2008 – present | <b>T-N</b>       |

|            |                                                                                               |                    |              |
|------------|-----------------------------------------------------------------------------------------------|--------------------|--------------|
| <b>24°</b> | Bristol Heart Babies Reform                                                                   | 1984 – present     | <b>G·N</b>   |
| <b>25°</b> | Removing the Race Coefficient from eGFR                                                       | 2021               | <b>D·G·N</b> |
| <b>26°</b> | Pulse Oximetry Across Skin Tones                                                              | 1990 – 2020 – 2025 | <b>D·G·N</b> |
| <b>27°</b> | Anesthesia Monitoring Standards and the APSF — The Mortality Decline                          | 1986 – present     | <b>H·K·G</b> |
| <b>28°</b> | Interprofessional Education and the Evidence Gap                                              | 2013 – 2015        | <b>K·N</b>   |
| <b>29°</b> | Bar-Code Medication Administration — Cue at the Point of Care                                 | 2010               | <b>H·K·D</b> |
| <b>30°</b> | Surgical Skill Video Peer-Rating Predicts Complications                                       | 2013               | <b>H·D·K</b> |
| <b>31°</b> | Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units                       | 2006 – 2016        | <b>T·K·H</b> |
| <b>32°</b> | Annual-Screening UI Redesign at University of Missouri Health Care                            | 2020               | <b>H·N</b>   |
| <b>33°</b> | Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal                 | 2019 – 2024        | <b>T·D·N</b> |
| <b>34°</b> | COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case                                | 2022 – 2024        | <b>T·K·D</b> |
| <b>35°</b> | Radiology AI Miscalibration                                                                   | 2018 – present     | <b>H·K·D</b> |
| <b>36°</b> | AlphaFold — Protein Structure Prediction                                                      | 2020 – present     | <b>K·G</b>   |
| <b>37°</b> | DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback                      | 2020               | <b>K·G</b>   |
| <b>38°</b> | iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms                    | 2006 – 2011        | <b>G·D·N</b> |
| <b>39°</b> | TeamSTEPPS                                                                                    | 2006 – present     | <b>T·N</b>   |
| <b>40°</b> | Team Science Training for Clinical and Translational Scientists                               | 2020 – 2022        | <b>K·N</b>   |
| <b>41°</b> | Translational-Workforce Training — Stated Goals Outrunning Named Practice                     | 2020s              | <b>K·N</b>   |
| <b>42°</b> | Australian Hospital-Pharmacy Technician Role Redesign                                         | 2016               | <b>D·N·H</b> |
| <b>43°</b> | Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface           | 2013 – 2018        | <b>H·N</b>   |
| <b>44°</b> | Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD            | 2014 – present     | <b>G·N</b>   |
| <b>45°</b> | Tennessee Voluntary Pre-K Study                                                               | 2009–2018          | <b>G·D</b>   |
| <b>46°</b> | Algorithmic Bias in Educational Predictive Analytics                                          | ongoing            | <b>G·H·D</b> |
| <b>47°</b> | Algorithmic Bias in Automated Exam Proctoring                                                 | 2022               | <b>D·N·K</b> |
| <b>48°</b> | School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism              | 2010s – 2022       | <b>G·K·N</b> |
| <b>49°</b> | UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days | 2020               | <b>D·K·N</b> |

|                                                                                                            |                                            |       |
|------------------------------------------------------------------------------------------------------------|--------------------------------------------|-------|
| Wisconsin DEWS — A Decade of                                                                               |                                            |       |
| 50° Algorithmic Dropout Prediction                                                                         | 2012 – 2023                                | D-K-N |
| Atlanta Public Schools Cheating Scandal                                                                    | 2009 – 2015                                | G-N   |
| Purdue Course Signals — The Reverse-Causality Retention Claim                                              | 2012 – 2013                                | D-K-N |
| 53 inBloom                                                                                                 | 2014                                       | G     |
| 54° Summit Learning / Personalized Learning Rollout                                                        | 2014–2019                                  | G-T-K |
| 55° Enrollment-Algorithm Yield Optimization Across U.S. Higher Education                                   | 2010s – present (Brookings synthesis 2021) | T-K-N |
| 57° GAO Online Program Manager Oversight Gap (GAO-22-104463)                                               | 2022                                       | D-K-N |
| 58° USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC)                            | 2010s – 2024                               | D-K-N |
| 60 Houston EVAAS — Value-Added Teacher Evaluation on Trial                                                 | 2011–2017                                  | G     |
| 61 Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap                        | 1997–2023                                  | K-T   |
| 62° Gates Intensive Partnerships — Measurement Without a Coupled Lever                                     | 2009–2018                                  | G-K   |
| 63° LAUSD iPad Program — The Common Core Technology Project                                                | 2013–2015                                  | D-H-G |
| 65° Training Transfer — The Gap Between Learning and Doing                                                 | 2010                                       | K-N   |
| 66° Growth-Mindset National Experiment — Heterogeneous Effects                                             | 2019                                       | D-N-K |
| 67° Cognitive Tutor / Carnegie Learning                                                                    | 1990s – present                            | T     |
| 68 CIRCUIT — A Scalable, Equity-Centered Research Workforce Model                                          | 2017 – 2023 (six cycles)                   | T-K   |
| 69° Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale                                    | 2016                                       | T-D   |
| 70° High-Impact Learning System — Engineering the Environment, Not Just the Event                          | 2001 – present                             | K-N   |
| 71° Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development | 2025                                       | K-N   |
| 72 ASSISTments — National Replication and Long-Term Follow-Through                                         | 2014 – present                             | T-K-D |
| 73° The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension                    | 2016 – 2025                                | T-K-D |
| 74° Zhang/Scardamalia — Knowledge Building Across Three Cohorts                                            | 2009                                       | T-K-D |
| 75° Chen et al. — Rural China AI Devices and the Equity-Direction Finding                                  | 2025                                       | T-K-D |
| 76° Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account                     | 2023                                       | T-K-N |
| 77 Hybrid Human-AI Tutoring — Augmentation, Not Delegation                                                 | 2024                                       | T-K-D |

|      |                                                                                                      |       |
|------|------------------------------------------------------------------------------------------------------|-------|
| 78°  | CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale<br>2017 – present                   | H·K·N |
| 79°  | The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops<br>1959 – present                      | K·N   |
| 80°  | Georgia State University Predictive Analytics<br>2012 – present                                      | T·K   |
| 81°  | Open University 'Ethical Use of Student Data' and OU Analyse<br>2014 – 2025                          | G·K·N |
| 82°  | MMALA — A Maturity Model for Responsible Learning–Analytics Adoption (Brazil)<br>2024                | K·N   |
| 83°  | Brinkerhoff Success Case Method — Tails as the Evaluation Instrument<br>2003 – present               | K·N   |
| 84°  | Cognitive Tutor Algebra I at Scale — Year–One Null, Year–Two Positive<br>2007 – 2014                 | T·K·D |
| 85°  | OU Analyse — Predictive Learning Analytics and Teacher Use at Scale<br>2019 – 2023                   | T·K·D |
| 86°  | Gándara — Detecting and Mitigating Algorithmic Bias in College Student–Success Prediction<br>2024    | D·K·N |
| 87°  | Yu / Lee / Kizilcec — Protected Attributes in Learning–Analytics Models<br>2021 – 2022               | D·K·N |
| 88°  | LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models<br>2025                       | T·K·N |
| 89°  | Data Privacy and Learning Analytics on the African Continent<br>2022                                 | K·G·N |
| 90°  | Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions<br>2025                 | T·K·N |
| 91°  | LALA — Building Learning–Analytics Governance Capacity Across Latin America<br>2017 – 2020           | K·N   |
| 92°  | Norway's National Expert Commission on Learning Analytics<br>2021 – 2023                             | G·K·N |
| 93°  | Singapore SkillsFuture — National Workforce Capability at Scale<br>2015 – present                    | G·K·D |
| 94°  | Learning Analytics on the African Continent — A Scoping Review as the Present–State Map<br>2022      | K·N   |
| 95°  | PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop<br>1997–2024 | T·G   |
| 96°  | AeroPerú Flight 603<br>1996                                                                          | H·T   |
| 97°  | Boeing 737 MAX / MCAS<br>2018 – 2019                                                                 | D·T·H |
| 98°  | Mars Climate Orbiter — Unit Mismatch<br>1999                                                         | D·K   |
| 100° | Glass–Cockpit Transition in Light Aircraft — Technology Outran Training<br>2010                      | H·N·K |
| 101° | Ariane 5 Flight 501<br>1996                                                                          | D·K·H |
| 102° | Air France Flight 447<br>2009                                                                        | T·H   |
| 103° | Kegworth / British Midland 92<br>1989                                                                | T·H·K |
| 104° | Asiana Airlines Flight 214<br>2013                                                                   | T·H   |
| 105° | Helios Airways Flight 522<br>2005                                                                    | T·H   |

|             |                                                                                             |                                   |              |             |                                                                                   |                                   |              |
|-------------|---------------------------------------------------------------------------------------------|-----------------------------------|--------------|-------------|-----------------------------------------------------------------------------------|-----------------------------------|--------------|
| <b>106°</b> | TransAsia Airways Flight 235                                                                | 2015                              | <b>T-H</b>   | <b>122°</b> | Singapore Airlines Safety Transformation                                          | 1980s – present                   | <b>T-N</b>   |
| <b>107°</b> | Eastern Air Lines Flight 401                                                                | 1972                              | <b>H-T</b>   |             | FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike | 1992 – 2000s                      | <b>G-K-N</b> |
| <b>108°</b> | NASA EVA-23 — Water Intrusion Inside a Spacesuit                                            | 2013                              | <b>H-N-D</b> | <b>123°</b> | USS Fitzgerald & USS John S. McCain                                               | 2017                              | <b>T-K-N</b> |
| <b>109°</b> | Colgan Air Flight 3407                                                                      | 2009                              | <b>T</b>     | <b>124°</b> | F-35 Sustainment & the Maintainer Capability Gap                                  | ongoing                           | <b>T-K-D</b> |
| <b>110°</b> | Atlas Air Flight 3591                                                                       | 2019                              | <b>T-K</b>   | <b>126°</b> | Military Fratricide — Desert Storm to Afghanistan                                 | 1991 – 2014                       | <b>T-H-K</b> |
| <b>111°</b> | Challenger & Columbia                                                                       | 1986 / 2003                       | <b>N-K-G</b> | <b>127°</b> | Operation Eagle Claw                                                              | 1980                              | <b>T-K</b>   |
| <b>112</b>  | NASA Saturn V Documentation                                                                 | 1972 – present                    | <b>K</b>     | <b>128°</b> | Patriot Missile / Dhahran                                                         | 1991                              | <b>D-H-K</b> |
| <b>113°</b> | Boeing Starliner                                                                            | 2019 – 2025                       | <b>K-D</b>   | <b>129</b>  | USS Vincennes & Iran Air Flight 655                                               | 1988                              | <b>H-T</b>   |
| <b>114°</b> | FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule                 | 1988 – 2010s                      | <b>G-D-K</b> | <b>130°</b> | Stanislav Petrov / 1983 False Alert                                               | 1983                              | <b>H-T</b>   |
| <b>115°</b> | FAA NextGen and the ADS-B Out Transition                                                    | 2003 – 2020                       | <b>G-D-K</b> | <b>131</b>  | F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap                      | 2008 – 2012                       | <b>H-K-N</b> |
| <b>116°</b> | Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change | 2005                              | <b>G-K-H</b> | <b>132°</b> | Mark 14 Torpedo Failures                                                          | 1941 – 1943                       | <b>T-K-G</b> |
| <b>117</b>  | Crew Resource Management & CAST                                                             | 1981 – present                    | <b>T-H-N</b> | <b>133</b>  | V-22 Osprey                                                                       | 1991 – present                    | <b>T-H-N</b> |
| <b>118°</b> | Korean Air Safety Transformation                                                            | 2000 – present                    | <b>T-N</b>   | <b>134°</b> | 9/11 Intelligence Sharing Failures                                                | 1996 – 2001                       | <b>G-K</b>   |
| <b>119</b>  | Aviation Safety Reporting System (ASRS)                                                     | 1976 – present                    | <b>T-K-N</b> | <b>135°</b> | U.S. Nuclear Navy / Rickover Training Model                                       | 1954 – present                    | <b>T-K-N</b> |
| <b>120</b>  | EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation                             | 1996 – 2002                       | <b>H-K-G</b> | <b>137</b>  | MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard               | 2020 revision (series since 1968) | <b>G-K</b>   |
| <b>121°</b> | TCAS — Coordinated Collision Avoidance and the Überlingen Lesson                            | 1981 – 2008 (TCAS II Version 7.1) | <b>H-K-G</b> | <b>138°</b> | GIFT and the Adoption Gap                                                         | 2012 – present                    | <b>K-G-N</b> |
|             |                                                                                             |                                   |              | <b>139°</b> |                                                                                   |                                   |              |

|             |                                                                                                |         |             |                                                                                               |         |
|-------------|------------------------------------------------------------------------------------------------|---------|-------------|-----------------------------------------------------------------------------------------------|---------|
| <b>140°</b> | Navy Surface Warfare Readiness Reform 2018 – present                                           | T-K-N   | <b>160°</b> | Davis-Besse Reactor Head Corrosion 2002                                                       | N-K-G   |
| <b>141</b>  | DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks 2009 – 2014              | H-K-D   | <b>161</b>  | Texas City BP Refinery Explosion 2005                                                         | N-T-K-G |
| <b>142°</b> | xAPI / Total Learning Architecture — Interoperability Gap ongoing                              | K-G     | <b>162°</b> | Upper Big Branch Mine Explosion 2010                                                          | N-G-K   |
| <b>143°</b> | Knight Capital Trading Loss 2012                                                               | D-K     | <b>163°</b> | Deepwater Horizon 2010                                                                        | T-N-K   |
| <b>144</b>  | Kodak Digital Camera Stagnation 1975–2012                                                      | D-K-N   | <b>164</b>  | Grenfell Tower 2017                                                                           | G-T-K-N |
| <b>146°</b> | Northeast Blackout 2003                                                                        | H-K     | <b>165°</b> | Camp Fire / PG&E 2018                                                                         | G-N-K   |
| <b>147°</b> | Takata Airbag Inflators 2008 – 2026                                                            | D-G     | <b>166°</b> | Three Mile Island 1979                                                                        | T-H     |
| <b>148</b>  | Wells Fargo Fake Accounts 2011 – 2016                                                          | G-N     | <b>167</b>  | Hurricane Katrina — The FEMA/DHS Response Failure 2005                                        | G-K     |
| <b>149°</b> | Volkswagen Dieselgate 2015                                                                     | D-G     | <b>169°</b> | Fukushima Daiichi 2011                                                                        | N-G-K   |
| <b>151°</b> | GM Ignition Switch 2002 – 2014                                                                 | D-G     | <b>170°</b> | West Africa Ebola — The Delayed International Response 2014 – 2016                            | G-N     |
| <b>152</b>  | TSB Bank IT Migration 2018                                                                     | H-G     |             | Hurricane Maria — Puerto Rico Response and Logistics Failure 2017 – 2018                      | G-N     |
| <b>153°</b> | Equifax Data Breach 2017                                                                       | G-K     | <b>172°</b> | CrowdStrike Falcon Outage 2024                                                                | D-K-G   |
| <b>154°</b> | INL Turbine–Control Upgrade — Verification and Validation as the Cutover Deliverable 2014      | G-K-H   | <b>173</b>  | Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable 1963 – present        | G-K-H   |
| <b>155</b>  | Toyota Production System / Andon Cord 1950s – present                                          | N-G     | <b>174°</b> | Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed 1996 – 2000 | G-D-K   |
| <b>156°</b> | INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet 2010 – present | D-H-K   | <b>175°</b> | INPO and the Nuclear Academy 1979 – present                                                   | T-K-G   |
| <b>157</b>  | AI-Augmented Coding Tools 2021 – present                                                       | T-H     | <b>176</b>  | Tylenol Recall 1982                                                                           | G-N     |
| <b>159°</b> | Bhopal 1984                                                                                    | T-K-N-G | <b>177°</b> | CIRAS — Confidential Incident Reporting for UK Rail 1996 – present                            | G-K-N   |

|             |                                                                                           |                       |                    |
|-------------|-------------------------------------------------------------------------------------------|-----------------------|--------------------|
| <b>178°</b> | The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami         | 2005 – 2020           | <b>G</b>           |
| <b>179</b>  | The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed          | 2020 – 2023           | <b>G-K</b>         |
| <b>180°</b> | Healthcare.gov Launch                                                                     | 2013                  | <b>K-T-G</b>       |
| <b>181°</b> | EU Human Brain Project — Top-Down Vision That Unraveled                                   | 2013 – 2023           | <b>G-N</b>         |
| <b>182°</b> | Amazon Hiring AI — Trained Bias, Deprecated 2017                                          | 2014 – 2018           | <b>D-K-N</b>       |
| <b>183°</b> | Uber ATG / Tempe Fatality                                                                 | 2018                  | <b>T-N-G-H</b>     |
| <b>184</b>  | UK Post Office Horizon Scandal                                                            | 1999 – 2015           | <b>G-H-K</b>       |
| <b>185</b>  | Air Canada Chatbot Liability — Delegation Without Revocation                              | 2022 – 2024           | <b>D-K-N</b>       |
| <b>186°</b> | Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity            | 2009 – 2022           | <b>D-G-N</b>       |
| <b>187</b>  | COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate                           | 2014 – 2018           | <b>D-K-N</b>       |
| <b>189°</b> | Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds                          | 2014 – 2020           | <b>D-G-N</b>       |
| <b>190°</b> | Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment                   | 2023 – 2024           | <b>G-K-N</b>       |
| <b>191</b>  | Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict           | 2016 – 2023           | <b>D-G-N</b>       |
| <b>192°</b> | Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model                 | 2019 – 2021           | <b>D-K-N</b>       |
| <b>193°</b> | Bernard Madoff / SEC Failures                                                             | 1992 – 2008           | <b>G-K-N</b>       |
| <b>194°</b> | Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox) | 2001 – present        | <b>G-K-N</b>       |
| <b>195°</b> | Tesla Autopilot — Recurring Fatalities                                                    | 2016 – present        | <b>T-N-G-H</b>     |
| <b>196</b>  | Fintech Lending Fairness Audit — When Including Race Reduces Disparity                    | 2025                  | <b>D-G-N</b>       |
| <b>197°</b> | Predictive Policing — PredPol                                                             | 2011 – present        | <b>G-H-D</b>       |
| <b>198°</b> | Launching the BRAIN Initiative — Governance of a Grand Challenge                          | 2011 – 2015 – present | <b>G-K-N</b>       |
| <b>199</b>  | Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact       | 2023 – 2026           | <b>G-K-N</b>       |
| <b>200°</b> | CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver          | 2020 – 2026           | <b>G-K-D</b>       |
| <b>201°</b> | Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India      | 2017 – 2025           | <b>G-N-H</b>       |
| <b>203</b>  | NYC Local Law 144 — Bias Audits as Governance Artifact                                    | 2023 – present        | <b>D-G-K</b>       |
| <b>205</b>  | The Discipline We Build Next                                                              | ongoing               | <b>T-K-N-G-H-D</b> |



**Modes.** T Training Gap · D Designed Out · N Normalization of Deviance · H Human–System Interface · G Governance & Trust · K Knowledge & Institutional Memory

**Courses.** The LDT / LENS courses each case serves as a worked example for are indexed separately in **The cases by LENS course** in the appendix.

**Parts.** The book is organized topically in seven parts — Healthcare & Patient Safety; Education, Training & the Learning Workforce; Aviation & Aerospace; Defense & National Security; Industry, Energy & Enterprise Systems; Disaster Prevention & Recovery; Algorithms, Governance & Public Systems. Each part splits into **What Fails** and **What Works — and the Frontier**; the open question closes the volume at Case 205.



Chapter 1

# Healthcare & Patient Safety — What Fails

*When care systems assume capability that was never specified, trained, or measured.*

---

*The harm arrived through the workflow, not around it.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

# Therac-25

**H** Human–System Interface   **D** Designed Out   **G** Governance & Trust

**IMPACT** Six known massive radiation overdoses; at least three deaths; canonical case in software safety engineering

## IN BRIEF

The Therac-25, a radiation-therapy machine, massively overdosed at least six patients between 1985 and 1987, killing at least three. Its predecessors carried independent hardware circuits that monitored the electron beam and stopped it; building the Therac-25, AECL decided not to duplicate all of the earlier machines' hardware safety mechanisms and interlocks, leaving software — adapted from those machines and never engineered for safety — to hold the two beam modes apart. Two defects were established: a data-entry race condition behind the cryptic "Malfunction 54" deaths at Tyler, and a counter overflow that skipped a position check in the fatal January 1987 Yakima accident. Three of the six accidents were never definitively explained. What made a defect lethal rather than annoying is the second half of the chain: the ion chambers saturated and reported an underdose, the fault code was undocumented, and the operator could clear it and fire again. Leveson and Turner's 1993 investigation, the founding case of software-safety engineering, found a systemic failure. The lesson — a safeguard you delete does not remove the hazard; it relocates it to whatever you failed to build.

## BACKGROUND

The Therac-25 was a radiation-therapy linear accelerator. Its predecessors, the Therac-6 and Therac-20, carried independent protective circuits monitoring the electron-beam scanning alongside mechanical interlocks — a backstop that could not be talked out of stopping the beam. Building the Therac-25, AECL took advantage of the computer's ability to control and monitor the hardware and, in Leveson and Turner's words, "decided not to duplicate all the existing hardware safety mechanisms and interlocks," leaving software — adapted from the older machines and never engineered from the ground up for safety-critical use — to keep its two beam modes, the photon setting drawing a hundred times the electron beam current, safely separated. The machine was not interlock-free; what it lacked was an independent circuit covering the failure that killed people. The safety case migrated quietly from steel to code.<sup>9</sup>

## WHAT HAPPENED

Between 1985 and 1987 the machine massively overdosed at least six patients. At Tyler, a race condition the engineers never knew about meant that if an operator entered a prescription, caught a mistake, and corrected it within about eight seconds — the speed an experienced operator naturally reaches — the full-power beam could fire with no target in place, delivering up to a hundred times the intended dose. Then every channel the operator had failed in the same direction. The ion chambers, hit with that dose, saturated and read low: the

console showed six monitor units delivered against the 202 requested — an underdose. The fault read “MALFUNCTION 54,” a code the manual did not explain, and which a technician later testified meant a dose either too high or too low. It was classed as a treatment **pause**, clearable with a single keystroke and repeatable up to five times, rather than a suspend requiring the parameters to be re-entered. The video display was unplugged and the audio monitor broken, so the patient — the one sensor in the room still reporting correctly — could not be seen or heard. The operator pressed the key and fired again. The fatal January 1987 Yakima overdose came from an unrelated defect: a one-byte counter in the setup test wrapped to zero on every 256th pass, and on that pass the routine skipped setting the flag that verifies collimator position. At least three patients died of the radiation burns. Three of the six accidents were never definitively explained: AECL could not reproduce them, and its testing of a suspected microswitch was inconclusive.<sup>1</sup>

#### THE INVESTIGATION


When patients reported searing pain, the manufacturer first insisted an overdose was impossible and treated each report as an isolated complaint rather than a signal. Nancy Leveson and Clark Turner’s 1993 investigation — the founding case study of software-safety engineering — found the fault was systemic rather than a single bug: overconfidence in software, hardware safeguards not carried forward, meaningless error messages, reused code never audited for safety, and an incident-reporting posture that dismissed the early warnings instead of compounding them into evidence. AECL had run a safety analysis in March 1983 — a fault tree that apparently excluded software, on the stated assumption that residual software errors were not included, and that where it did reach the computer assigned a probability of ten to the minus eleven to “computer selects wrong energy” and four times ten to the minus nine to “computer selects wrong mode,” with no justification given for either. The pathology was not a missing analysis but one whose scope had a hole exactly where the safety case had moved. The sharpest evidence came from outside the manufacturer: after the second Tyler accident a physicist at the University of Chicago Joint Center for Radiation Therapy checked whether the same thing could happen on a Therac-20 and found that whenever a new class of students began using it, fuses and breakers tripped. The same defect was there. On that machine it was a nuisance, because the independent hardware circuits were still watching the beam.<sup>2</sup>

#### THE CAPABILITY GAP

The independent circuit had not been redundant; it **was** the safety case, the one thing standing between a typing error and a lethal dose. Nothing took its place — no verified software check, no independent monitor watching the beam. The operator stayed nominally in the loop, and a safety function assigned to a person needs four things to be real: a true reading of the hazardous variable, time to act inside the hazard’s own timescale, an action cheaper to take than to skip, and confirmation it worked. She had none of them. The dose display reported the opposite of the truth; the exposure was irreversible in under a second; clearing the fault took one key and stopping took judgment against a readout saying nothing was wrong. The question capability engineering would have forced — **what function now carries the interlock’s load, and can whatever now holds it actually see?** — was never asked of the redesign.<sup>3</sup>

## AFTERMATH &amp; REFORM

Therac-25 reshaped safety-critical software practice, making the case by counterexample for independent hazard analysis, safeguards that do not all rest on the same software, error messages that say what is wrong so the operator can act, and treating field reports as evidence to be aggregated rather than complaints to be closed. It remains the canonical teaching case across software, medical-device, and systems-safety curricula.<sup>4</sup> No rule then required a site to report such an incident, so each accident was unknown to the others and Tyler's physicist had to reproduce the fault himself; the Safe Medical Devices Act of 1990 changed that. Its lesson is exact and portable: a safeguard you remove does not remove the hazard it guarded — it relocates the hazard to whatever you failed to put in its place, and waits there.



## COMPETING READINGS

- ◆ The removed circuit is the load-bearing absence — set against the reading that the decisive deficit was **feedback**, not shutdown. At Tyler the second dose was delivered because the display said underdose and the fault class permitted a one-key retry; a truthful dosimeter and a suspend-class fault stop the site after the first exposure without any interlock being restored. The corrective action plan deployed hardware and interface changes together, so the record cannot decompose them.
- ◆ Engineering was the missing capability — set against the reading that **reporting** was. Each site's accident was unknown to the others, and the Tyler and January 1987 Yakima accidents were serially preventable: a mandatory reporting regime, with no design change at all, plausibly prevents those deaths. This reading predicts transfer to any system whose hazard is rare per site and repeating across sites.
- ◆ The removal was a design choice — set against the reading that it was a diffuse non-decision. AECL's stated rationale is thin and largely reconstructed after the fact, and the evidence is at least as consistent with a redesign driven by compactness and cost in which no artifact and no role existed to ask the question. The two readings point at different places to intervene: at the decision, or at the absence of a decision point.
- ◆ Better information would have saved the patients — set against the reading that at a sub-second, irreversible timescale no achievable human capability would have helped, and the correct design was to take the human out of the detection loop entirely. On this reading the human-shaped capabilities in the case sit at other timescales: noticing a pattern across sessions, believing the patient, escalating against a manufacturer's assurance, and learning across sites.

What this case does not show — The denominator is unknown. The failure was undetectable from inside the system, so a sub-fatal overdose would have presented as an unexplained treatment reaction and never

been attributed to the machine; “six known” is a lower bound and no exposure-adjusted risk figure drawn from this case can be trusted in either direction. The account itself is a documentary reconstruction from FDA records, correspondence, depositions and lawsuits, with limited manufacturer cooperation and design documentation that in places did not exist — the best available record, and not an audited engineering root-cause report.

*The software error is just a nuisance on the Therac-20 because this machine has independent hardware protective circuits for monitoring the electron-beam scanning.*

LEVESON & TURNER, IEEE COMPUTER, 1993

## REFERENCES

1. N. G. Leveson & C. S. Turner, “An Investigation of the Therac-25 Accidents,” IEEE Computer 26(7): 18–41 (1993). doi:10.1109/MC.1993.274940 — the removed hardware interlocks and the adapted software.
2. Leveson & Turner (1993) — the two software defects (data-entry race condition and counter overflow), the uninformative “Malfunction 54”, overdoses up to 100×, six accidents, and at least three deaths.
3. Leveson & Turner (1993); N. G. Leveson, Safeware: System Safety and Computers (Addison-Wesley, 1995) — the manufacturer’s denial and the systemic findings.
4. N. G. Leveson, Engineering a Safer World: Systems Thinking Applied to Safety (MIT Press, 2011) — why removing a safeguard requires explicitly reassigning its safety function.
5. The case’s role in medical-device software regulation and software-safety practice (FDA software guidance; IEC 62304 lineage); see also Ethics Unwrapped, UT Austin.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Therac-25 is the case in which a safety function was reassigned and nothing was built to receive it. The record does not show anyone weighing the removal; it shows a redesign in which no artifact and no role existed to ask the question, and a 1983 fault tree that left software outside its scope. The operator was kept in the system and stripped of any means of seeing what it was doing wrong. Capability engineering would have asked, before the circuit came out, **what function takes its load, and can whatever now holds it see?**

### LENS APPROACH

LENS frames Therac-25 as a **Design-Out** failure made visible through Interface and Governance pathologies. Studio projects in LEN 5 ask students to produce capability-load diagrams tracing every safety function to a hardware backstop, a software check, or a trained operator action with the information needed to perform it. LEN 7 examines incident reporting as governance.

| DOMAIN     | FAILURE MODES                | LENS COMPETENCY                                     | PROBLEM TYPE | INDUCED | LEO |
|------------|------------------------------|-----------------------------------------------------|--------------|---------|-----|
| HEALTHCARE | → <b>H</b> <b>D</b> <b>G</b> | → <b>D1</b> <b>D2</b> <b>D3</b> <b>D4</b> <b>D5</b> | 3            | 3.1     | 3   |
| TECHNOLOGY |                              | Human-System Collaboration<br>+ D4                  |              |         |     |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- Before deleting any hardware interlock, write down the safety function it performs and explicitly reassign that function to a verified software check, an independent monitor, or an informed operator action.

### IN OPERATION / AFTER

- Treat every field report of unexpected harm as evidence to be aggregated, not a complaint to be closed, with a standing path to halt the device when a pattern emerges.

- ▶ Do not let safety rest entirely on one software path; require an independent channel that does not share the same code, so a single defect cannot defeat the whole safety case.
- ▶ Specify error messages as a safety deliverable — each fault code must say what is wrong and what the operator should do, never offer to proceed past an undiagnosed condition.
- ▶ Audit any safety-critical code reused from a prior system against the new hazard set, since assumptions safe in the old context may be lethal in the new one.
- ▶ Instrument the machine so an independent monitor can detect a beam fired without its target and stop it, restoring the interlock's function even where the operator cannot see the fault.

## REFLECTION QUESTIONS

1. Identify a system in your domain that migrated a safety function from hardware to software. Where did the human-capability load go, and who is accountable for sustaining it?
2. Therac-25 operators saw “MALFUNCTION 54” and continued treatment. Redesign that interface using LEN 2 principles so that the operator's correct action is the easiest action.
3. The Therac-25's safety-critical software was inherited from earlier machines and never re-audited. Identify reused code in a system you build that now carries a load it was never written for, and specify the review it should have received.

**WHO BUILDS THIS** human factors & interaction design · systems & human-factors engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · task analysis · design prototyping · governance frameworks · bias & evidence audits



# EHR / CPOE Implementation

**H** Human–System Interface   **D** Designed Out   **G** Governance & Trust

**IMPACT** ~\$30B federal investment under HITECH; mortality among children transported to one hospital rose 2.80% to 6.57% after CPOE went live; ongoing usability harm at scale

## IN BRIEF

The 2009 HITECH Act poured roughly \$30 billion into accelerating electronic health records, and adoption surged — but the systems had been built to billing and administrative specifications, not clinical workflow. New categories of harm followed. A disputed but canonical 2005 study found that deploying a commercial order–entry (CPOE) system at one pediatric ICU was associated with a more–than–doubling of mortality; later deployments showed mixed or improved results, yet the warning stuck. A 2018 JAMA study of 1.735 million safety reports found 557 whose language explicitly tied an EHR usability problem to possible patient harm. Specific harms recur — a boy discharged on lab results the record showed as complete, a girl killed by a drug the disabled interaction alert would have caught. There is still no regulatory framework monitoring EHR safety. It is the book’s live interface–mismatch case.

## BACKGROUND

The 2009 HITECH Act authorized roughly \$30 billion in incentives to accelerate electronic health record adoption, and adoption surged. But the systems clinicians now had to use had been built to administrative and billing specifications, not to clinical workflow or human–factors specifications — a mismatch baked in before the first order was entered.<sup>0</sup> Because the incentives rewarded adoption rather than usability, hospitals raced to install whatever cleared the bar, and the specification gap propagated to the bedside at national scale before anyone had to demonstrate that the tools supported the work of care rather than merely the work of billing.

## WHAT HAPPENED

New categories of harm appeared alongside the new tools. A 2005 study reported that after a commercial computerized order–entry (CPOE) system went live hospital–wide at one children’s hospital, mortality among children arriving by interfacility transport rose from 2.80% to 6.57% — a single–institution result that provoked debate and that later deployments elsewhere did not reproduce, but that became canonical as a warning that powerful tools without workflow integration can disrupt care at the moment of greatest acuity, when the seconds the interface adds to an order are the seconds a critical patient does not have.<sup>1</sup> Specific harms recur: a twelve–year–old sent home to die of sepsis on lab results the record presented as complete; a twelve–year–old with long QT syndrome killed by a drug whose

interaction alert had been switched off — each a case where the interface’s defaults and toggles, not the clinician’s competence, determined whether the right information reached the decision.<sup>2</sup>

### THE INVESTIGATION


Howe, Adams, Hettinger, and Ratwani read 1.735 million patient-safety reports filed in Pennsylvania and one academic health system between 2013 and 2016. In 557 of them — 0.03 percent — the reporter’s own words tied an EHR usability problem to possible patient harm: 468 reached the patient, 80 potentially caused temporary harm, seven potentially permanent harm, and two could have required intervention to sustain life.<sup>3</sup> Reporting is voluntary and the details are sparse, so the count is a floor, not a measure. What it establishes is that interface failures reach patients at all — which is what turns usability from a satisfaction question into a safety one.

### THE CAPABILITY GAP

EHR/CPOE is the canonical ongoing case of an interface designed for one set of requirements and deployed against another. The systems work for billing and administration because that was the design constraint; they fail for clinical workflow because clinical workflow was not. Decades and tens of billions of dollars later, the harm usability failures do is still counted only in voluntary reports built to catch other things — a cost that persists precisely because it was never a design requirement and so was never engineered out, only worked around by clinicians absorbing the friction shift after shift.<sup>4</sup>

### AFTERMATH & REFORM

Usability and safety have risen on the agenda — AMA/Pew/MedStar usability work, KLAS benchmarking, safety-focused design guidance — but there is still no regulatory framework that monitors deployed EHR safety the way drugs or devices are monitored. The lesson is that buying the tool is not the same as engineering the capability: an interface deployed against the wrong specification keeps extracting a cost no one is formally counting, and without a monitoring regime the harm stays diffuse — spread across millions of encounters, attributable to no single failure, and therefore easy for the system to keep paying indefinitely.<sup>5</sup>



*Despite HIT's promise in improving safety, recent literature has revealed potential safety hazards associated with its use.*

SITTIG & SINGH, DEFINING HEALTH INFORMATION TECHNOLOGY-RELATED ERRORS, ARCHIVES OF INTERNAL MEDICINE, 2011

## REFERENCES

1. Health Information Technology for Economic and Clinical Health (HITECH) Act (2009) — the \$30B EHR incentive program.
2. Y. Han et al., "Unexpected Increased Mortality After Implementation of a Commercially Sold CPOE System," *Pediatrics* 116(6): 1506–1512 (2005) — the disputed pediatric-ICU mortality result.
3. M. Grissinger, "The Absence of a Drug–Disease Interaction Alert Leads to a Child's Death," *P&T* 43(2): 71–72 (2018); F. Schulte & E. Fry, "Death by 1,000 Clicks," *Fortune/KHN* (2019).
4. J. Howe, K. Adams, A. Z. Hettinger & R. Ratwani, "Electronic Health Record Usability Issues and Potential Contribution to Patient Harm," *JAMA* 319(12): 1276–1278 (2018) — 557 of 1.735M safety reports tied usability to possible harm.
5. D. Sittig & H. Singh, "Defining Health Information Technology-related Errors: New Developments Since To Err is Human," *Archives of Internal Medicine* 171(14): 1281–1284 (2011).
6. Pew / MedStar / AMA, *Ways to Improve Electronic Health Record Safety* (2018); R. Wachter, *The Digital Doctor* (2015).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

EHR/CPOE is the canonical ongoing case of an interface designed for one set of requirements and deployed against another. The systems work for billing and administration. They fail for clinical workflow because clinical workflow was not the design constraint. Thirty billion dollars later, the harm usability failures do is still counted only in voluntary reports never built to catch it.

### LENS APPROACH

LENS treats EHR/CPOE in LEN 7 as the live, ongoing example of Design-Out and Interface failure under governance opacity, in LEN 2 as a human–AI teaming problem (alert fatigue, default surfacing), and in LEN 9 as a computational systems problem (the alerting architecture itself is a learnable model). The Han 2005 transported–children finding is the durable warning; the Howe 2018 review of 1.735 million safety reports is the contemporary evidence that usability failures reach patients.

| DOMAIN                     | FAILURE MODES | LENS COMPETENCY            | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|---------------|----------------------------|--------------|---------|-----|
| HEALTHCARE                 | → [H] [D] [G] | → [D1] [D2] [D3] [D4] [D5] | 3            | 3.1     | 3   |
| Human–System Collaboration |               |                            |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Write clinical workflow and human–factors performance into the procurement specification, so a system that serves billing but not bedside care cannot clear the buy.
- ▶ Validate defaults, alerts, and result–surfacing against real clinical scenarios before deployment, since it was the default and the unactivated alert — not clinician skill — that determined the harm.
- ▶ Tie any adoption incentive to demonstrated usability, so the money rewards a tool that supports care rather than merely one that is installed.

#### IN OPERATION / AFTER

- ▶ Stand up a post–deployment safety–monitoring regime for fielded EHRs comparable to how drugs and devices are surveilled, with thresholds and authority to force changes.
- ▶ Read safety reports for the words clinicians use about the interface, since it was reporters' own language that made usability–driven harm countable at all.
- ▶ Audit recurring interface–driven harms — surfaced defaults, disabled alerts — and feed them back into vendor design requirements rather than absorbing them as clinician burden.

## REFLECTION QUESTIONS

1. What is the equivalent system in your domain that was designed for one specification and deployed against another? How would you measure the harm?
2. Design the regulatory architecture that would surface EHR safety harms at scale. Be specific about signal, threshold, and authority.
3. The Han 2005 mortality result was disputed; the Howe 2018 count rests on voluntary reports that catch only a fraction of the harm. What ongoing measurement architecture would have to exist for the equivalent emerging clinical-AI deployment in your domain to be evaluated honestly while in use?

**WHO BUILDS THIS** human factors & interaction design · systems & human-factors engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · task analysis · design prototyping · governance frameworks · bias & evidence audits

# Watson for Oncology — Delegation Marketed Ahead of Capability

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**EVIDENCE TIER** journalism-tier — source confidence flagged; future validation ongoing.

**IMPACT** IBM marketed Watson for Oncology with tumor-board concordance rates as high as 96%; STAT News investigation (Ross & Swetlitz, 2017–2018) and internal IBM documents documented unsafe and incorrect cancer-treatment recommendations and that the system was trained on a small number of synthetic cases rather than real patient outcomes; hospitals worldwide had purchased the tool on the marketed concordance claim

## IN BRIEF

Watson for Oncology was IBM's heavily marketed clinical-decision- support system, sold to hospitals worldwide as a cancer- treatment recommendation engine whose concordance with expert tumor boards was advertised at rates as high as 96%. The capability the system was sold under was AI-grade treatment recommendation across cancer types, delivered with the institutional credibility of IBM and MD Anderson partnerships. The capability the system actually carried — documented in the investigative reporting by Casey Ross and Ike Swetlitz at STAT News through 2017 and 2018, drawing on leaked IBM internal documents — was substantially smaller. Watson was trained on a small number of synthetic cases curated by Memorial Sloan Kettering oncologists rather than on real patient outcomes, generated unsafe and incorrect recommendations in some documented cases, and the internal IBM record showed engineers raising concerns about the gap between marketed concordance claims and what the system actually delivered. Hospitals around the world had purchased Watson on the marketed capability; adoption stalled, deployments were discontinued (MD Anderson shelved its Watson project in 2016 after about \$62 million; Rigshospitalet in Copenhagen dropped it after agreeing with Watson in roughly a third of cases), and IBM wound the product down and, in 2022, sold the Watson Health data and analytics assets to Francisco Partners for roughly \$1 billion, relaunched as Merative — the clinical-recommendation business effectively shuttered before the divestiture. The case is the canonical instance in the AI-delegation typology of capability marketed ahead of capability validated. The evidence-tier flag rendered under the title is binding: STAT News journalism is the primary public source, and the academic record cites it secondarily. "Future validation ongoing" is the standing language; independent peer-reviewed evaluation of Watson for Oncology's operating accuracy at the level the marketing claimed was never produced.

## BACKGROUND

Watson for Oncology was the most publicly visible commercial clinical-AI deployment of the mid-2010s. IBM had positioned Watson — originally the Jeopardy-winning question-answering system — as a healthcare engine after the 2011 Jeopardy performance, and oncology was the flagship domain. The product was sold as a cancer-treatment recommendation system: a clinician would enter the patient's case, Watson would query its knowledge base against the training data provided by Memorial Sloan Kettering oncologists, and the system would return ranked treatment recommendations annotated with evidence citations. The marketed capability was treatment recommendation that concurred with expert tumor boards at rates as high as 96%. Hospitals worldwide purchased the system on that claim.<sup>0</sup>

## WHAT HAPPENED

Casey Ross and Ike Swetlitz at STAT News produced the investigation of record across 2017 and 2018. Drawing on leaked IBM internal documents and interviews with engineers, oncologists, and former IBM staff, the reporting documented a substantially different operating picture. Watson had been trained not on real patient outcomes — the standard a treatment-recommendation engine would have to clear — but on a small number of synthetic cases curated by MSK oncologists to reflect what those oncologists would have done. The training corpus was a model of expert opinion at one institution, not an outcome-anchored learning record. In documented cases the system generated unsafe and incorrect treatment recommendations. Internal IBM documents showed engineers raising concerns about the gap between marketed capability and operating reality.<sup>1</sup>

## THE INVESTIGATION

The structural failure pattern the case carries is delegation of consequential authority — cancer-treatment recommendation — to a tool whose validated capability was substantially smaller than its marketed capability. The procurement mechanism was the carrier: hospitals purchased Watson through capital and IT procurement processes that evaluated the marketed concordance claims and the IBM/MSK institutional credentials, but did not independently validate operating accuracy on the receiving institution's patient population. The capability the procurement was buying — accurate, real-world treatment recommendation at the marketed concordance rate — was not the capability the system carried. Deployments were discontinued as the gap surfaced — MD Anderson shelved its project, Rigshospitalet in Copenhagen dropped Watson outright, and IBM ultimately wound down and divested Watson Health.<sup>2</sup>

## THE CAPABILITY GAP

The evidence-tier flag rendered under the title is binding and survives into the case verbatim. The STAT News investigation by Ross and Swetlitz is the primary public source, and the academic record cites it secondarily. No independent peer-reviewed evaluation of Watson for Oncology's operating accuracy at the level the marketing claimed was ever produced — partly because the system was proprietary and partly because the deployments were wound down before the kind of prospective study that would have established the operating record could be conducted. The case is teachable on the structural failure pattern at journalism-grade evidence; the standing "source confidence flagged; future validation ongoing" language is the honest tier the case rests on.<sup>3</sup>

## AFTERMATH & REFORM

In the AI-delegation typology (Cases 20 TREWS, 5 Epic Sepsis Model, 189 SyRI, and this case), Watson is the marketing-ahead-of-capability failure mode: the procurement decision rode on marketed concordance rates that were not validated at the operating institutions, and the institutional credentials did not substitute for the validation. TREWS shows delegation with strong outcome evidence works; Epic Sepsis shows delegation without validation fails by accuracy disconfirmation; SyRI shows delegation halted on rights grounds; Watson shows delegation driven by procurement and marketing rather than by validated capability. The four together teach the typology of when AI delegation is and is not legitimate. Watson is the case where the marketing finance — the institutional credibility and the headline concordance number — substituted for the evidence the procurement should have required, and the capability the hospitals thought they were buying was not the capability they got.

*The capability the procurement was buying — accurate, real-world treatment recommendation at the marketed concordance rate — was not the capability the system carried.*

EDITORS' SYNTHESIS OF ROSS & SWETLITZ (STAT NEWS 2017–2018).

## REFERENCES

1. Somashekhar et al. and the WFO lung/colorectal early-experience abstracts, J Clin Oncol 35 (2017), 15\_suppl, presented at the ASCO Annual Meeting — the marketed concordance figures: 96% lung, 93% rectal, 81% colon.
2. Ross & Swetlitz, "IBM's Watson supercomputer recommended 'unsafe and incorrect' cancer treatments, internal documents show," STAT (25 July 2018), and "IBM pitched its Watson supercomputer as a revolution in cancer care. It's nowhere close," STAT (5 Sept 2017) — the primary sources; investigative journalism drawing on leaked IBM internal documents.
3. Strickland (2019), "How IBM Watson Overpromised and Underdelivered on AI Health Care," IEEE Spectrum — independent retrospective synthesis of the public record.
4. Topol (2019), Deep Medicine, Basic Books — secondary academic situating of Watson within the broader clinical-AI delegation discourse.
5. v2 paired cases: TREWS (Case 20), Epic Sepsis Model (Case 5), SyRI (Case 189) — the AI-delegation typology.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Watson for Oncology is the canonical instance of clinical-AI delegation marketed ahead of validated capability. The operating capability was substantially smaller than the marketed concordance claims; the procurement mechanism evaluated the marketing rather than the validation; the deployments were wound down as the gap surfaced. The evidence-tier flag is binding: journalism-grade reporting is the source, and future validation ongoing remains the honest qualifier.

### LENS APPROACH

Watson for Oncology is the marketing-ahead-of-capability AI-delegation failure (induced 1.1; LENS D3/PT6) — Domain 3 for **Delegation with revocation** (the discontinued deployments are the revocation half); Domain 4 for requirements-vs- marketing. Pair with Cases 20, 5, 189. Evidence-tier flag binds.

| DOMAIN                                                          | FAILURE MODES | LENS COMPETENCY                              | PROBLEM TYPE | INDUCED | LEO |
|-----------------------------------------------------------------|---------------|----------------------------------------------|--------------|---------|-----|
| HEALTHCARE/<br>ONCOLOGY<br><br>COMMERCIAL AI<br><br>PROCUREMENT | D K N         | D1 D2 D3 D4 D5<br>Human-System Collaboration | 6            | 2.4     | 1-4 |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Separate marketed capability from validated capability in any procurement decision involving consequential delegation; the marketed concordance number is the carrier of the procurement story, not the evidence the procurement should rest on.</li><li>Require independent operating-accuracy validation on the receiving institution's patient population before delegation; the AI-delegation literature now treats this as the procurement floor, and Watson is one of the reasons.</li><li>Anchor the procurement evidence to outcome-anchored training, not to expert-curated synthetic cases; expert opinion at one institution is not outcome data the procurement can rely on.</li></ul> | <ul style="list-style-type: none"><li>Render the evidence tier honestly when the primary public source is journalism; the STAT investigation is the source, and the academic record's secondary citation is the accurate description, not "peer-reviewed evidence."</li><li>Carry the AI-delegation typology — TREWS, Epic Sepsis, SyRI, Watson — as a unit in any curricular discussion of when delegation is legitimate; the four-case set teaches the typology more clearly than any single case.</li><li>Treat the discontinuation of Watson for Oncology deployments as the case's own correction signal; the institutional decisions to discontinue are themselves evidence about what the procurement should have required up front.</li></ul> |

REFLECTION QUESTIONS

1. Identify a procurement decision in your domain that delegated consequential authority to an automated system. Did the procurement evaluate marketed capability claims or validated capability on the receiving institution's data? Where is the gap most likely to have been smoothed by institutional credentials standing in for evidence?
2. Specify what operating-accuracy validation you would require before delegating: which prospective study, on which population, with which comparison and which outcome metric. The AI-delegation typology has established this as the procurement floor; the question is whether your domain's procurement processes meet it.
3. The evidence-tier flag under this case's title is binding because the primary public source is journalism. Identify a case in your domain where the strongest available evidence is journalism-grade. What is the responsible drafting discipline that surfaces the tier rather than implying peer-reviewed validation?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards



# Mid Staffordshire NHS Foundation Trust

**G** Governance & Trust   **N** Normalization of Deviance   **K** Knowledge & Institutional Memory

**IMPACT** Mortality ran above expected at Stafford Hospital for years, which Francis refused to convert into a death count; the Francis Inquiry produced 290 recommendations and restructured UK patient-safety governance

## IN BRIEF

Between roughly 2005 and 2009, Stafford Hospital, run by the Mid Staffordshire NHS Foundation Trust, subjected patients to appalling neglect — left without food, water, or basic care — while mortality ran above expected. The trust had been chasing financial targets that depended on staffing cuts. Robert Francis QC's public inquiry produced a three-volume, 1,782-page report and 290 recommendations, identifying the structural problem as the gap between the trust's reported performance and patients' actual experience: every governance layer above the ward received reports that targets were being met, and none checked those reports against what was happening to patients. Mid Staffs is the dataset's strongest case for the harm done when measurement and reality diverge over years.

## BACKGROUND

The Mid Staffordshire NHS Foundation Trust ran Stafford Hospital in England. Pursuing Foundation Trust status and the financial targets that came with it, the board cut staffing — and the cuts fell on the wards. The institution's reported performance, the numbers that travelled upward, stayed on target.<sup>0</sup> The targets the board chased were financial and procedural, so cutting ward staff improved the very figures the trust was measured on even as it removed the people on whom patient care directly depended.

## WHAT HAPPENED

On the wards the reality was appalling. Patients were left in their own excrement, denied food and water, given the wrong medication or none, for years; mortality ran substantially above expected. The harm was not a single incident but a sustained condition — visible to anyone on the ward and invisible in the reports that left it.<sup>1</sup> Because the suffering was a continuous state rather than a nameable event, it never generated the kind of discrete incident a reporting system is built to catch, and so it accumulated for years beneath numbers that stayed reassuringly on target.

## THE INVESTIGATION

The Healthcare Commission investigated; the layers above did not. Robert Francis QC's public inquiry ran to 1,782 pages across three volumes and 290 recommendations.<sup>2</sup> Its structural finding was the divergence between reported performance and patient experience: every governance layer above the ward had received reports that the hospital was


meeting its targets, and none had checked them against what was happening to patients. “Not just the Trust’s Board but the system as a whole failed in its most essential duty — to protect patients from unacceptable risks of harm and from unacceptable, and in some cases inhumane, treatment that should never be tolerated in any hospital.”<sup>3</sup> The phrase “the system as a whole” located the failure deliberately above any single ward or manager: no one layer was solely at fault, because each had trusted the layer below to be reporting reality rather than targets.

#### THE CAPABILITY GAP

Mid Staffs is the British analog of the VA wait-time scandal (Case 7): measurement and reality diverged over years while every layer of governance acted on the measurement. The capability that was missing was not clinical skill on the ward but the institutional habit of testing whether the numbers corresponded to the patients — a check no layer above the ward performed.<sup>4</sup> Each layer reasonably assumed the check belonged to someone else, so the verification that the report matched the patient fell into the gap between layers — exactly the place a reporting chain built only to pass numbers upward is structurally unequipped to look.

#### AFTERMATH & REFORM

The Francis Inquiry restructured how the NHS treats patient safety, and the Berwick review that followed pressed for a culture of learning over targets. The lesson is the measurement one in its starkest form: a reporting chain can run clean from ward to Department of Health while, underneath it, the thing being reported on quietly fails.<sup>5</sup> Berwick’s framing — learning over targets — named the deeper correction: as long as the target is the thing the institution rewards, the report will describe the target, and only a culture that prizes finding the gap will keep checking the report against the patient.



*each layer was acting on the metric, not the patient*

MID STAFFS — MEASUREMENT AND REALITY DIVERGED ACROSS YEARS

*Not just the Trust’s Board but the system as a whole failed in its most essential duty — to protect patients from unacceptable risks of harm and from unacceptable, and in some cases inhumane, treatment that should never be tolerated in any hospital.*

ROBERT FRANCIS QC, REPORT OF THE MID STAFFORDSHIRE NHS FOUNDATION TRUST PUBLIC INQUIRY, 2013

#### REFERENCES

1. R. Francis QC, Report of the Mid Staffordshire NHS Foundation Trust Public Inquiry (2013) — the staffing cuts and Foundation Trust targets.
2. Healthcare Commission, Investigation into Mid Staffordshire NHS Foundation Trust (2009) — ward conditions and excess mortality.
3. Francis QC (2013) — the three-volume, 1,782-page report and 290 recommendations.
4. Francis QC (2013) — “not just the Trust’s Board but the system as a whole failed in its most essential duty — to protect patients from unacceptable risks of harm and from unacceptable, and in some cases inhumane, treatment” (quoted).
5. D. Berwick, A Promise to Learn — A Commitment to Act (National Advisory Group on the Safety of Patients in England, 2013).
6. K. Walshe & J. Higgins (2002) on NHS safety governance; cf. the VA wait-time scandal (Case 7).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Mid Staffordshire is the British analog of the VA wait-time scandal (Case 7). Measurement and reality diverged over years; every layer of governance above the operating environment was acting on the measurement; patients paid the cost of the divergence. The Francis Inquiry recommendations changed how the NHS thinks about patient safety as an institutional capability.

### LENS APPROACH

LENS uses Mid Staffs in LEN 4 for the divergence-of-measurement problem and in LEN 7 for the governance failure across multiple layers. Studio projects examine the Francis recommendations as a template for institutional reform deliverables.

| DOMAIN              | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------|---------------|------------------|--------------|---------|-----|
| HEALTHCARE          | → G N K       | → D1 D2 D3 D4 D5 | 5            | 2.1     | 4   |
| Test and Evaluation |               |                  |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Design targets so that gaming them (e.g., cutting ward staff) cannot improve the metric while degrading the outcome it stands for.
- ▶ Build a direct patient-experience signal — independent of the financial and procedural targets — into the reporting chain from the start.
- ▶ Assign explicit ownership of the report-versus-reality check at each layer, so verification cannot fall into the gap between them.

#### IN OPERATION / AFTER

- ▶ Audit sustained conditions, not just discrete incidents, since continuous harm never trips an event-based reporting system.
- ▶ Cross-check upward reports against ground truth at the ward periodically, treating a clean report as a hypothesis to be tested.
- ▶ Reward a culture of learning over target attainment, so the institution keeps looking for the gap rather than describing the target.

### REFLECTION QUESTIONS

1. Identify a multi-layer reporting chain in your domain. What would it take for the top layer to know whether the reports correspond to reality?
2. The Francis Inquiry produced 290 recommendations. Pick five that you think were most load-bearing and explain why.
3. The verification fell into the gap between layers because each assumed it belonged to someone else. In your domain, who explicitly owns the check that a report matches reality — and how would you know if no one does?

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards · knowledge bases · data pipelines

# Epic Sepsis Model

**D** Designed Out **G** Governance & Trust **N** Normalization of Deviance

**IMPACT** External validation across 38,455 hospitalizations found AUROC 0.63 versus reported 0.76–0.83, missing ~67% of sepsis at the operational threshold; deployed in hundreds of hospitals without independent validation or FDA clearance

## IN BRIEF

The Epic Sepsis Model was a proprietary machine-learning sepsis prediction tool embedded in the Epic EHR and deployed in hundreds of U.S. hospitals — the most widely operational clinical AI in American medicine. Until Wong et al. (JAMA Internal Medicine 2021) ran an external validation across 38,455 hospitalizations, no independent evaluation had been published. The reported AUROC was 0.63, well below the 0.76–0.83 Epic had reported; at the operational threshold, the model missed roughly two-thirds of sepsis cases, with a 12% positive predictive value and substantial alert burden. The case is not principally about the model's performance. It is about the governance seam that let the deployment happen at scale without independent validation: as an EHR-embedded proprietary feature, the model sat outside the FDA's medical-device oversight regime, so the machinery that would have surfaced its limitations at clearance was never engaged. The post-deployment surveillance pattern (Vioxx, Case 9) is the analog: the harm was the absence of a standing system to check the tool once it was in the population's hands.

## BACKGROUND

By the late 2010s the Epic Sepsis Model was, by deployment count, the most widely operational clinical AI in American medicine — embedded in hundreds of hospitals as a default feature of the dominant inpatient EHR. The tool was presented as a help to clinicians trying to catch sepsis earlier in a workflow already saturated with alerts. The model's design and validation were proprietary and not externally evaluated.<sup>o</sup>

## WHAT HAPPENED

In 2021 Wong et al. published the first large external validation in JAMA Internal Medicine: 38,455 hospitalizations at the University of Michigan, with sepsis diagnosed using two consensus definitions. The model achieved an AUROC of 0.63, materially below Epic's reported 0.76–0.83. At the operational threshold for bedside alerting, the model missed roughly two-thirds of sepsis cases. The 12% positive predictive value meant most alerts were false; the alert burden landed on clinicians who could not distinguish the few real alerts from the many spurious ones.<sup>1</sup>

## THE INVESTIGATION

The governance seam is the structural lesson. Because the Epic Sepsis Model was distributed as a feature of an EHR rather than as a stand-alone clinical-decision-support device, it was never brought to the FDA for clearance — even though FDA’s own 2022 guidance says a risk score for a specific condition does not qualify for the section 520(o)(1)(E) exclusion and is a device. The machinery that would normally require independent validation, post-market surveillance, and demographic stratification of performance was never engaged. The model’s deployment was a regulatory non-event because the regulatory regime treated the EHR layer as out of scope. The clinical AI question and the device-oversight question diverged.<sup>2</sup>

## THE CAPABILITY GAP

What surfaced the failure was post-deployment external validation — the exact discipline that the clearance pathway omits. The Wong et al. paper was disconfirmation in the form the system did not otherwise provide. In October 2022 Epic overhauled the model — changing the sepsis-onset definition, reducing its reliance on antibiotic-order signals, and moving to a gradient-boosted model (ESM v2) that hospitals must train on their own local data before use; many had already turned the alert off, recalibrated, or replaced it. Later multicenter validations of the updated model reported stronger discrimination (AUROC roughly 0.82–0.92) but persistent institution-to-institution variability, low positive predictive value, and heavy alert burden — reinforcing that local validation remains necessary. The corrective action worked through publication, not through governance. That is the gap: a tool can be deployed at hundreds of sites, alert at the bedside for years, and still be disconfirmable only by an academic paper rather than by the surveillance architecture the deployment was supposed to have.<sup>3</sup>

## AFTERMATH & REFORM

The Epic case is the negative pair to TREWS (Case 20) and the governance-seam analog to Radiology AI Miscalibration (Case 35). Together they teach a typology: delegation done well as a paired intervention with engineered interface and outcome-grounded evidence (TREWS); delegation done badly as a proprietary EHR-embedded model deployed outside the device-oversight regime without independent validation (Epic); delegation halted on rights grounds because the system was both ineffective and rights-violating (SyRI); delegation marketed ahead of capability (Watson for Oncology). The four together are the canonical AI delegation typology.

*A deployment is not a validation. Deployment without independent validation is delegation without evidence.*

EDITORS’ SYNTHESIS OF WONG ET AL. (2021).

## REFERENCES

1. Habib et al. (2021), commentary on Wong et al., JAMA Internal Medicine — on the implications for proprietary clinical AI.
2. Wong et al. (2021), “External Validation of a Widely Implemented Proprietary Sepsis Prediction Model in Hospitalized Patients,” JAMA Internal Medicine 181(8):1065–1070, doi:10.1001/jamainternmed.2021.2626.
4. Ross, C. (2022), “Epic overhauls its sepsis algorithm,” STAT News (Oct 2022) — the ESM v2 redesign and the shift to locally-trained models.
5. Adams et al. (2022), Nature Medicine — the paired positive case (20).

3. FDA, Clinical Decision Support Software: Final Guidance (2022)  
— the post-Wong reframing of the EHR-embedded oversight question.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Epic Sepsis Model is the canonical case of consequential clinical-AI delegation at scale without independent validation. The structural lesson is not the model’s poor performance; it is the governance seam that let the deployment proceed without the validation and surveillance machinery that the medical-device pathway would have required, surfaced only by post-deployment external work.

LENS APPROACH

Epic is the Domain 4 + Domain 3 / Problem Type 6 failure that motivates the post-deployment-surveillance discipline LENS teaches. Used in Domain 4 (Test and Evaluation) for measurement architecture under proprietary opacity and the gap-attribution LEO; in Domain 3 (Human-System Collaboration) for the delegation-with-revocation LEO — Epic was delegated without a pre-specified revocation criterion; and in Domain 5 (Navigating Sociotechnical Constraints) for the cross-regime / platform governance seam. Pairs directly against TREWS (Case 20) and sits in the AI-delegation typology with SyRI and Watson.

| DOMAIN      | FAILURE MODES | LENS COMPETENCY          | PROBLEM TYPE | INDUCED | LEO       |
|-------------|---------------|--------------------------|--------------|---------|-----------|
| HEALTHCARE  | D G N         | D1 D2 D3 D4 D5           | 6            | 2.4     | 4 · 5 · 3 |
| CLINICAL AI |               | Test and Evaluation + D3 |              |         |           |
| GOVERNANCE  |               |                          |              |         |           |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Require independent external validation before deployment of any consequential clinical AI, regardless of whether it ships as a stand-alone device or as a feature of a host platform.
- Specify in advance the disconfirming evidence — population, threshold, PPV, alert burden — that would revoke the delegation, and the channel through which that evidence would surface.
- Identify the regulatory regime the tool falls under, and where the seam between regimes is — proprietary EHR features should not be exempt from clinical-AI oversight by virtue of their packaging.

IN OPERATION / AFTER

- Build post-deployment surveillance as a standing institutional capability — outcome metrics, demographic stratification, alert-burden audit — so disconfirmation does not require a single academic paper to surface.
- Close the cross-regime seam: clinical AI embedded in EHRs should be subject to the same independent validation and surveillance as stand-alone clinical-decision-support devices.
- When disconfirming evidence arrives, treat it as a designed input: revise, recalibrate, or remove on a defined timeline, with the corrective action visible to the clinicians who used the tool.

## REFLECTION QUESTIONS

1. Identify a clinical AI tool deployed in your domain. Where in the regulatory architecture would independent validation have been required, and where could it slip the seam? What pre-specified disconfirming evidence would revoke the delegation?
2. Design the post-deployment surveillance deliverable that should accompany every deployment of consequential clinical AI — including embedded-in-EHR features that currently fall outside the device-oversight regime.
3. The disconfirmation in this case came from a single academic paper, not from a standing institutional architecture. What is the minimum surveillance machinery that would have surfaced the model's performance gap at the operational threshold without requiring the Wong et al. paper to exist?

**WHO BUILDS THIS** systems & human-factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Racial Bias in Pain Assessment — The False-Belief Mechanism

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** About half of surveyed white medical students and residents endorsed at least one false belief about biological differences between Black and White people; those who held more false beliefs rated Black patients' pain as less severe and recommended less accurate treatment

## IN BRIEF

Hoffman, Trawalter, Axt, and Oliver (PNAS 2016) surveyed medical students and residents on a battery of false beliefs about biological differences between Black and White people (e.g., "Black people's skin is thicker," "Black people's nerve endings are less sensitive"). About half endorsed at least one such belief. The paper's experimental layer showed that respondents who endorsed more false beliefs rated the pain of mock Black patients as less severe than the same pain in mock White patients, and made less accurate treatment recommendations. The mechanism the case identifies is specific and unusually precise for a bias study: the pain-assessment gap is traceable to a small set of nameable false biological beliefs, not to diffuse implicit bias or structural-only explanation. That precision is what makes Hoffman the human-development case in the race-construct trio. The capability deliverable is not awareness training; it is curriculum that specifically disconfirms the named false beliefs and instrumentation that surfaces when bedside ratings of pain diverge by patient race.

## BACKGROUND

Pain assessment is a clinician's judgment, made repeatedly across a day, on patients whose subjective report of pain has to be translated into a numeric rating and a treatment decision. A documented finding in the medical literature is that Black patients in the United States are systematically under-treated for pain across emergency-department, post-surgical, and end-of-life settings. The bias has been measured at the population level for decades; the mechanism was less precisely named.<sup>o</sup>

## WHAT HAPPENED

Hoffman et al. (2016) administered a battery of statements about biological differences between Black and White people to 222 white medical students and residents — some true, some false (e.g., "Black people's skin is thicker," "Black people's blood coagulates more quickly," "Black people's nerve endings are less sensitive"). About half of respondents endorsed at least one false belief; a smaller subset endorsed several. The experimental layer



of the study presented respondents with two mock patient cases identical except for race, asked them to rate the patients' pain, and asked them to recommend treatment.<sup>1</sup>

#### THE INVESTIGATION

The pattern was that respondents who endorsed more false beliefs rated the pain of the Black mock patient as less severe than the pain of the White mock patient, and recommended less accurate treatment for the Black mock patient. Respondents who endorsed fewer false beliefs showed the opposite pattern, rating the Black mock patient's pain as slightly higher. The case is unusual in identifying a specific cognitive mechanism — a small set of named false biological beliefs — that mediates a documented population-level disparity. Most bias studies leave the mechanism diffuse; Hoffman names it precisely enough that a curriculum or assessment can target it.<sup>2</sup>

#### THE CAPABILITY GAP

What the case teaches at the construct layer is that the capability deliverable in medical education is not generic "implicit bias" awareness — it is curriculum that specifically disconfirms the named false beliefs, with assessment instruments that test whether the beliefs were actually disconfirmed. Operationally, the deliverable is a bedside instrument or surveillance pattern that surfaces when pain ratings diverge by patient race in ways that survive case-mix adjustment. The Hoffman finding makes both deliverables specifiable in a way that more diffuse bias findings did not.<sup>3</sup> The mechanism has since proven portable in a way that extends the deliverable: Omiye et al. (2023, npj Digital Medicine) found that four commercial large language models reproduce false beliefs Hoffman catalogued — about skin thickness and pain threshold — alongside race-corrected kidney and lung outputs, when queried with medical prompts, so the disconfirmation task now reaches past the medical-school curriculum to the AI systems entering the same clinical workflows.<sup>4</sup>

#### AFTERMATH & REFORM

Hoffman pairs with pulse oximetry (Case 26) and eGFR (Case 25) in the race-construct trio. The three cases are the same surface harm — minority patients systematically underserved across a clinical decision — attributable to three distinct layers of the system: the construct definition (eGFR), the validation architecture (pulse oximetry), and the human-development mechanism (Hoffman). The trio is the case-grounded basis for the subobjective **Gap attribution** (4.2): distinguishing capability gaps attributable to human development, system design, and organizational performance, and selecting measurement that isolates the intended cause.

*The mechanism the paper names is precise enough to teach against. Awareness training is not a curriculum; a curriculum has to disconfirm something specific.*

EDITORS' SYNTHESIS OF HOFFMAN ET AL. (2016).

#### REFERENCES

1. Anderson, Green, & Payne (2009), "Racial and ethnic disparities in pain: causes and consequences of unequal care," *Journal of Pain* 10(12):1187–1204 — the population-level disparity.
2. Hoffman, Trawalter, Axt, & Oliver (2016), "Racial bias in pain assessment and treatment recommendations, and false beliefs about biological differences between blacks and whites," *PNAS* 113(16):4296–4301, doi:10.1073/pnas.1516047113.
3. Vyas, Eisenstein, & Jones (2020), *NEJM* — connecting race-in-clinical-algorithms to race-in-clinical-judgment.
4. Omiye, Lester, Spichak, Rotemberg, & Daneshjou (2023), "Large language models propagate race-based medicine," *npj Digital Medicine* 6:195 — commercial LLMs reproduce the Hoffman false beliefs.

3. Sabin & Greenwald (2012), “The influence of implicit bias on treatment recommendations for 4 common pediatric conditions,” American Journal of Public Health — the diffuse-mechanism backdrop the Hoffman precision improves on.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Hoffman et al. is the human–development case in the race–construct trio. The pain–assessment disparity in medical settings is mediated, in measurable part, by a nameable set of false biological–difference beliefs held by clinicians in training. The capability deliverable is curriculum that specifically disconfirms the beliefs and instrumentation that surfaces the bedside rating gap when it persists.

LENS APPROACH

Hoffman is the human–development case in the race–construct trio (Cases 25, 26 and 6). LENS uses it in Domain 4 (Test and Evaluation) for the subobjective **Gap attribution** (4.2) — the gap is in the clinician’s training, not the construct or the device — and in Domain 2 (Learning Engineering Design) for the curriculum–design implication. The trio together is the case–grounded basis for **Gap attribution**: same surface harm, three distinct layers, three distinct remediations. Adjacent to the lending pair (Cases 186–113) at the construct layer.

| DOMAIN            | FAILURE MODES                              | LENS COMPETENCY                                                       | PROBLEM TYPE | INDUCED | LEO |
|-------------------|--------------------------------------------|-----------------------------------------------------------------------|--------------|---------|-----|
| CLINICAL MEDICINE | → <div>T</div> <div>K</div> <div>N</div> → | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 5            | 8.1     | 4   |
| MEDICAL EDUCATION |                                            | Test and Evaluation                                                   |              |         |     |
| HEALTH EQUITY     |                                            |                                                                       |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Build curriculum that specifically disconfirms the named false biological beliefs identified in the Hoffman survey, with assessment items that test whether the disconfirmation took hold.
- ▶ Instrument bedside pain ratings to surface case–mix–adjusted divergence by patient race; the gap is otherwise invisible to the system that produces it.
- ▶ Identify the layer of the gap before designing the remediation: construct, validation, or human–development. A construct fix cannot remediate a clinician–belief fix and vice versa.

IN OPERATION / AFTER

- ▶ Re-administer the Hoffman survey periodically as a curriculum–evaluation instrument; a curriculum that does not move the false–belief endorsement rate is not closing the mechanism the paper identifies.
- ▶ Track whether the bedside pain–rating disparity narrows in cohorts that received the disconfirming curriculum, with reporting at intervals long enough for selection effects to settle.
- ▶ Cross-reference the human–development result against the construct (eGFR) and validation–architecture (pulse oximetry) layers, so the overall equity capability of the clinical system is not assessed only at the layer the institution finds easiest to instrument.

## REFLECTION QUESTIONS

1. Identify a documented disparity in your domain whose mechanism is treated as diffuse. What would a Hoffman-style survey look like — a battery of named false beliefs or assumptions whose endorsement could be measured and whose presence predicts the operational decision?
2. Design the curriculum-evaluation instrument you would use to test whether a curriculum has actually disconfirmed the false beliefs. What endorsement-rate change would you require before claiming the mechanism has been addressed?
3. Hoffman is the human-development case in the race-construct trio. Pulse oximetry (Case 26) is the validation-architecture case; eGFR (Case 25) is the construct-definition case. Which of the three layers does your domain currently address, and which does it leave invisible?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

# VA Wait-Time Scandal

**G** Governance & Trust   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Veterans died waiting for care; ~120,000 waiting 90+ days or never got onto lists; staff falsified records

## IN BRIEF

In 2014 the VA Inspector General found that staff at the Phoenix VA — and then nationwide — had created secret waiting lists and falsified appointment data to hide that veterans were waiting weeks or months for care; some died waiting while the system reported success. A 14-day access target, unrealistic given staffing, was met by hiding reality rather than surfacing it. The warning signs ran back fifteen years: GAO had flagged data-reliability problems since 2000, the IG since 2005, with no systemic fix — and turnover among schedulers, the staff operating the measurement, ran high enough that the institution kept losing the knowledge even to see it was lying to itself. The VA case is the book's canonical example of normalization of deviance applied to measurement itself.

## BACKGROUND

The Veterans Health Administration measured access to care against a 14-day appointment target — a target that, given staffing, was often unrealistic to meet honestly. Turnover among schedulers, the staff who operate that measurement, is high enough that GAO treats scheduler training as a standing risk.<sup>0</sup> The target functioned as the headline number leadership watched, so the pressure to show 14-day compliance bore down hardest on the very front-line role least equipped, through constant churn, to record the access data accurately or to question what the number was leaving out.

## WHAT HAPPENED

In 2014 the VA Office of Inspector General found that staff at the Phoenix VA — and then across the system — had created secret waiting lists and falsified appointment data to hide that veterans were waiting weeks or months for care. Some veterans died while waiting — though the Inspector General could not conclusively tie the deaths to the scheduling delays — inside a system that, by its own metrics, was succeeding.<sup>1</sup> The secret lists were the mechanism by which an impossible target was reconciled with reality: official records showed appointments inside 14 days while the true wait accumulated off the books, so the more the metric was gamed, the more confidently the reporting line above it declared the access problem solved.

## THE INVESTIGATION

The warning signs went back fifteen years: GAO had flagged data-reliability concerns since 2000, and the VA IG had identified problems in 2005 and 2007, with follow-up work in


2008 finding nine agreed recommendations still unimplemented — none of which produced systemic change.<sup>2</sup> GAO found the reported wait times unreliable because VA never made schedulers' definitions of the measurement dates consistent, and called training and oversight vital against high scheduler turnover; five years on it was still concerned about the reliability of the wait times.<sup>3</sup> Each warning had named a real defect in how access was recorded, yet because the schedulers who held the practical knowledge of the system kept turning over, every wave of findings landed on a workforce that had to relearn the instrument from scratch, so the same defect resurfaced report after report.

#### THE CAPABILITY GAP

This is normalization of deviance applied to measurement itself. When the measurement system cannot capture reality — and the people operating it churn before the gaming can be unlearned — the gap between reported and actual performance becomes invisible. Veterans died inside a system whose numbers said it was fine, which is the lethal form of the evidence problem this book treats as a design failure, not a reporting one.<sup>4</sup> An institution that cannot retain the staff who run its measurement loses not only accuracy but the memory that the number was ever wrong, so the gaming stops registering as deviance at all and hardens into the ordinary way the work is done.

#### AFTERMATH & REFORM

The scandal forced resignations, the Veterans Access, Choice, and Accountability Act (2014), and a long, uneven effort to rebuild scheduling and data integrity.<sup>5</sup> Its lesson is that decision-grade evidence is a design requirement: an institution must be able to surface its own failures without relying on the very people incentivized — and too transient — to hide them. The Choice Act bought access by routing care outside the VA, but the underlying capability — a measurement the institution could trust even as its schedulers churned — was the harder thing to rebuild, and the data-reliability concerns that persisted for years afterward show why.



*Given the high turnover among schedulers, it is important that VA remain vigilant about scheduler training, ensuring all who need it receive it.*

DEBRA A. DRAPER (GAO DIRECTOR OF HEALTH CARE), HOUSE VA COMMITTEE TESTIMONY, GAO-19-687T, 24 JULY 2019

#### REFERENCES

1. D. Draper (GAO), House VA Committee testimony, GAO-19-687T (2019) — “given the high turnover among schedulers, it is important that VA remain vigilant about scheduler training” (quoted).
2. VA Office of Inspector General, Report 14-02603-267 (2014) — secret waiting lists and falsified appointment data.
3. VA OIG reports (2005, 2007, 2008) — prior, unactioned findings.
4. GAO Veterans Health Administration reports (2000–2019) — fifteen years of data-reliability warnings.
5. C. Argyris & D. Schön, *Organizational Learning: A Theory of Action Perspective* (1978); G. Bevan & C. Hood, “What’s Measured Is What Matters,” *Public Administration* 84(3): 517–538 (2006).
6. Veterans Access, Choice, and Accountability Act (2014) — the legislative response.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The VA case is the canonical example of measurement failure as a capability failure. The system that should have surfaced the gap instead generated reports that hid it. The turnover among schedulers — the human capability operating the measurement instrument — meant the institution lost the knowledge to even notice it was lying to itself. The case stands as the strongest argument in this book for treating decision-grade evidence as a design requirement, not a reporting requirement.

### LENS APPROACH

LENS treats the VA case in LEN 4 as the canonical evidence-gap case (the measurement system itself was the source of harm), in LEN 7 as a governance failure (multiple warnings unactioned over fifteen years), and in LEN 8 as a knowledge-loss case via turnover.

| DOMAIN              | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------|---------------|------------------|--------------|---------|-----|
| GOVERNMENT          | → G K N       | → D1 D2 D3 D4 D5 | 5            | 2.2     | 4   |
| Test and Evaluation |               |                  |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Set access targets against actual staffing capacity, so the metric cannot be met only by falsifying it.
- ▶ Design the scheduling instrument so correct data entry is the path of least resistance, not a discipline that churning staff must be retrained into.
- ▶ Build an independent read on real wait times — separate from the staff incentivized to report them — into the measurement architecture from the start.

#### IN OPERATION / AFTER

- ▶ Audit reported access against an out-of-band signal (direct veteran survey, third-party booking records) with authority to act when the two diverge.
- ▶ Treat scheduler turnover as a measurement risk: monitor it, and protect the knowledge of how to run the instrument honestly against constant churn.
- ▶ Track whether old data-reliability findings keep recurring — a repeat finding is the signal that the institution is relearning the same gap rather than closing it.

### REFLECTION QUESTIONS

1. Identify a measurement system in your domain that is also operated by a high-turnover role. What is the institutional risk that the system stops measuring reality?
2. Design the evidence pipeline that would have surfaced the Phoenix VA gap without relying on the people who were gaming the metrics.
3. The 14-day target was unrealistic given staffing, so it was met by hiding reality. Identify a target in your organization that the people measured against it cannot honestly meet — and design the correction that surfaces the gap rather than burying it.

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

## CASE 8 HEALTHCARE

1999 – PRESENT

# Medical Errors as Systemic Failure

**T** Training Gap   **H** Human-System Interface   **N** Normalization of Deviance  
**K** Knowledge & Institutional Memory   **G** Governance & Trust

**DISCLOSURE** Institutional overlap: an editor shares an institution (Johns Hopkins) with an author of this work; no editor was personally involved. The disputed mortality estimate is contested in the case itself.

**IMPACT** IOM 1999 estimate of 44,000–98,000 U.S. deaths/year from medical error; Makary & Daniel (2016) estimate of ~250,000 deaths/year — substantively contested on methodological grounds; 2023 NEJM inpatient-harm study confirms persistence; 25-year reform arc with bounded successes and an unmoved population count

## IN BRIEF

Medical error in the United States is not a single incident but a systemic condition that the system's own measurement instruments cannot see. The Institute of Medicine's *To Err Is Human* (1999) raised the alarm with a 44,000–98,000 annual-deaths estimate and reframed harm as a systems problem rather than an individual one; *Crossing the Quality Chasm* (2001) and the 2015 *Improving Diagnosis in Health Care* sequel extended the program. In 2016, Makary and Daniel of Johns Hopkins published a BMJ analysis estimating 250,000 deaths a year — which would make medical error the third leading cause of death behind heart disease and cancer. Death certificates do not record medical error as a cause, so the problem is structurally invisible to the systems meant to track it. The 250,000 figure has been substantively contested on methodological grounds (Shojania & Dixon-Woods, BMJ Quality & Safety, 2017; Carr in Health Affairs), a dispute that is itself a worked example of the gap-attribution problem. The field's 25-year arc shows bounded interventions that work — computerized order entry, handoff protocols, the WHO surgical checklist (Case 23), the Keystone ICU project (Case 19), TeamSTEPPS (Case 39) — alongside a population-scale mortality count that has resisted both intervention and precise estimation.

## BACKGROUND

The Institute of Medicine's *To Err Is Human* (1999) was the field-defining moment: it estimated 44,000–98,000 deaths annually from medical error in U.S. hospitals — at the lower bound, more Americans than die in motor-vehicle accidents — and made the explicit case that the problem was a systems problem, not an individual-clinician problem. The 2001 sequel *Crossing the Quality Chasm* set six aims for the redesign of care, and the 2015 sequel *Improving Diagnosis in Health Care* extended the framing to diagnostic error. The framing of harm as systemic mattered: it created the cultural permission for non-punitive incident reporting, root-cause analysis, and the broader patient-safety movement that followed.<sup>o</sup> Yet modern medicine's core mortality-measurement instrument — the death certificate —

records a proximate physiological cause and has no field for medical error. Because the certificate is keyed to an ICD billing taxonomy built for disease, a death set in motion by a care-process breakdown is recorded under whatever organ ultimately failed, and the causal role of the system disappears into the physiology.

#### WHAT HAPPENED

In 2016, surgeon Martin Makary and Michael Daniel published an analysis in the BMJ estimating that medical errors cause more than 250,000 deaths a year in the United States — which would rank third behind heart disease and cancer. Their core claim was that “people don’t just die from heart attacks and bacteria, they die from system-wide failings and poorly coordinated care.” By relocating the cause from the individual clinician to the coordination of care, the framing recast a ledger of isolated mistakes as a single population-scale failure mode the existing statistics were never built to count. The number itself was computed by extrapolating from four studies published since the IOM estimate, the 2010 OIG Medicare adverse-events study, the Landrigan 2010 NEJM study, and the Classen 2011 Global Trigger Tool study — to the contemporary inpatient population, an extrapolation the authors acknowledged as approximate.<sup>1</sup>

#### THE INVESTIGATION

The 250,000 estimate was substantively contested almost immediately. Shojania and Dixon-Woods, writing in BMJ Quality & Safety in 2017, challenged both the extrapolation and the attribution method, arguing that counting deaths “due to” error is far harder than a single headline number implies. A Health Affairs commentary by Carr and a series of letters in the BMJ pressed the same point: the source studies use different definitions of “preventable death,” the extrapolation rests on prior estimates already known to be unstable, and the ranking against CDC cause-of-death categories is methodologically incompatible (CDC counts use ICD codes; the Makary estimate does not).<sup>2</sup> The objection turned on counterfactual attribution — how confidently one can say a frail, already-dying patient would have survived but for the error — a judgment that resists the clean tallying a headline number demands. The dispute is itself a worked example of the gap-attribution problem: how much of a counted death is the learning system, how much the system design, how much the underlying disease.<sup>3</sup>


#### THE CAPABILITY GAP

The deeper failure is one of measurement: a system that cannot see its own failure modes cannot manage them. With no field on the certificate, no reliable count, and no agreed attribution method, every safety program competes for resources against a harm that the official record renders invisible, so even effective interventions struggle to prove their worth at population scale. The bounded interventions that have worked — computerized physician order entry, structured handoff protocols (I-PASS), the WHO Surgical Checklist (Case 23), the Keystone ICU central-line bundle (Case 19), TeamSTEPPS (Case 39) — each move a specific harm in a specific setting, and each can prove it. What none can prove is movement on the population count, because the population count does not exist in a form precise enough to be moved. The missing capability is an instrument that captures medical error as a tracked cause of harm — and an attribution method robust enough that the resulting number can guide intervention rather than fuel a methodological stalemate.<sup>4</sup>



## AFTERMATH & REFORM

The 25-year arc since *To Err Is Human* is the case's most important teaching artifact. The IOM framing catalyzed a patient-safety movement, an Agency for Healthcare Research and Quality patient-safety program, and the bounded interventions noted above. Yet later work, including the 2023 NEJM study by Bates et al. of inpatient harm across eleven Massachusetts hospitals, confirms that the problem persists at scale: about a quarter of admissions involved an adverse event, and roughly a quarter of those events were preventable. The system has not built the instrument the original report implied it would need — no national active-surveillance system for inpatient harm, no death-certificate field for care-process failure, no agreed attribution method for “deaths due to error.” The interventions exist; the measurement and the implementation still lag, and the headline mortality kept escaping measurement even as bounded harms fell. The case is the canonical worked example of the gap-attribution problem at population scale, and a standing reminder that an instrument the system cannot see through is one it cannot govern.<sup>5</sup>



CONTESTED AS THIRD-RANKING — AND STRUCTURALLY INVISIBLE

*People don't just die from heart attacks and bacteria, they die from system-wide failings and poorly coordinated care.*

MARTIN MAKARY, PRESS STATEMENTS ACCOMPANYING MAKARY & DANIEL, BMJ (2016)

## REFERENCES

1. Institute of Medicine, *To Err Is Human: Building a Safer Health System* (1999); *Crossing the Quality Chasm* (2001); *Improving Diagnosis in Health Care* (2015) — the field-defining trilogy and the 44,000–98,000 estimate; the systems framing.
2. Makary, M. & Daniel, M. (2016), “Medical error — the third leading cause of death in the U.S.,” *BMJ* 353:i2139 — the 250,000 estimate, the quoted framing, the extrapolation from four studies published since the IOM report (*HealthGrades* 2004; *OIG* 2010; *Landrigan* 2010; *Classen* 2011).
3. Shojania, K. & Dixon-Woods, M. (2017), “Estimating deaths due to medical error: the ongoing controversy and why it matters,” *BMJ Quality & Safety* 26(5):423–428; with Bates, D. W. & Singh, H. (2018), *Health Affairs* 37(11):1736–1743 — methodological contestation of the Makary extrapolation and CDC-ranking comparison.
4. Makary & Daniel (2016), *BMJ* — death certificates do not capture medical error as a cause; ICD billing taxonomy as the structural reason.
5. Bates, D. W., Levine, D. M., Salmasian, H., et al. (2023), “The Safety of Inpatient Health Care,” *NEJM* 388(2):142–153 — eleven-hospital Massachusetts cohort of 2,809 admissions; at least one adverse event in 23.6% of admissions, 25% of those preventable; persistence of harm at scale.
6. Agency for Healthcare Research and Quality, *National Healthcare Quality and Disparities Reports* (annual); CDC WONDER ICD-coded mortality data — institutional context for the missing national active-surveillance instrument.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The Makary data is the anchor evidence for the LENS argument because of its scale, its provenance (Johns Hopkins), and its framing — system failure rather than individual error. The seventeen years between **To Err Is Human** and Makary's reassessment is also the implementation gap of Case 13 in another guise — and the cost of leaving it open is measured in lives at population scale.

### LENS APPROACH

Medical error is the central evidence anchor of the curriculum (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation; LEO-4) for the foundational measurement problem: the death certificate cannot record the failure mode it is supposed to govern, and the contested mortality estimates show what happens when an instrument's gap is filled by extrapolation. LENS uses it in Domain 1 (Systems Analysis) as the foundational problem statement of the program and in Domain 5 (Navigating Sociotechnical Constraints) for industry-level institution building. The 25-year arc pairs the case with the WHO Surgical Checklist (Case 23), Keystone ICU (Case 19), and TeamSTEPPS (Case 39) as the bounded interventions that worked; pair with Vioxx (Case 9) at the population-scale-surveillance layer.

| DOMAIN              | FAILURE MODES | LENS COMPETENCY | PROBLEM TYPE | INDUCED | LEO |
|---------------------|---------------|-----------------|--------------|---------|-----|
| HEALTHCARE          | → T H N K G → | D1 D2 D3 D4 D5  | 5            | 2.1     | 4   |
| Test and Evaluation |               |                 |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Give the mortality-recording instrument an explicit field for care-process failure, so a system-caused death is captured, not absorbed into the proximate cause.
- ▶ Specify the attribution method up front — the counterfactual test for “due to” error — so the count guides action rather than collapsing into dispute.
- ▶ Pair each safety intervention with the population measure it should move, gating reforms on demonstrated effect.

#### IN OPERATION / AFTER

- ▶ Audit reported mortality against an independent count of care-related harm.
- ▶ Use active surveillance of inpatient harm rather than waiting for the death certificate.
- ▶ Hold the count and the intervention to the same measurement discipline, closing the implementation gap.

### REFLECTION QUESTIONS

1. Identify a measurement instrument in your domain that systematically fails to capture the failure modes it should be designed to surface. What would it cost to fix?
2. Two hundred fifty thousand deaths a year is the third leading cause of death in the U.S. Design the measurement and intervention regime that would shift the curve over a five-year horizon. Estimate the deliverable and the evidence.
3. Makary and Shojania disagreed not on whether error kills but on how to attribute a death to it. Specify an attribution method robust enough to survive that dispute — and name who would hold the authority to act on the number it produces.

**WHO BUILDS THIS** learning science & instructional design · human factors & interaction design · safety science & organizational analysis · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs · incident analysis · safety dashboards · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

# Vioxx Withdrawal

**G** Governance & Trust   **D** Designed Out

**IMPACT** Tens of thousands of excess cardiovascular events estimated (FDA's Graham, 2004; Lancet 2005); Merck withdrew Vioxx in September 2004; ~\ \$4.85B settlement; FDA Amendments Act of 2007 and the Sentinel Initiative (2008) institutionalized active post-market surveillance

## IN BRIEF

Merck's painkiller Vioxx (rofecoxib) was approved by the FDA in May 1999 and prescribed to an estimated 80 million people worldwide. The 2000 VIGOR trial (Bombardier et al., NEJM) found a roughly five-fold higher rate of myocardial infarction in the Vioxx arm than in the naproxen comparator — a signal Merck and many readers attributed to a hypothesized cardio-protective effect of naproxen rather than a cardiovascular risk of Vioxx, a reading that required the absent placebo arm to be true. Not until the placebo-controlled APPROVe trial (Bresalier et al., NEJM) was halted in September 2004 was the risk established and the drug withdrawn. For nearly four years the disclosure architecture between Merck's internal data, FDA reviewers, and prescribers failed to surface the magnitude of the danger to a decision point. FDA Office of Drug Safety scientist David Graham testified to the Senate Finance Committee in November 2004 that the risk had been visible well before withdrawal, and estimated 88,000–139,000 excess cardiovascular events attributable to Vioxx; the 2005 Graham et al. Lancet analysis put the figure in similar range. Merck litigation discovery produced internal communications later used to argue that guest authorship and Merck-employee authorship patterns had shaped how the cardiovascular signal appeared in the published record. The reforms that followed — Risk Evaluation and Mitigation Strategies (REMS), the FDA Amendments Act of 2007, and the Sentinel Initiative (2008) — built the active post-market surveillance architecture that had not previously existed.

## BACKGROUND

Vioxx, a selective COX-2 inhibitor for arthritis pain, was approved by the FDA in May 1999 and became one of the most widely prescribed drugs of its era — an estimated 80 million people worldwide were exposed before withdrawal. Detecting rare or delayed harms in a drug at that scale depends on post-market surveillance — aggregating adverse-event data after approval — a function that, at the time, was thin relative to the size of the exposed population. The FDA's MedWatch system relied on voluntary submission from clinicians; no active query of the underlying claims and electronic-health-record data was operational. A risk too small to surface in a pre-approval trial becomes a large absolute toll once a drug reaches millions, which is precisely the regime that demands strong after-market monitoring — and precisely the regime in which voluntary reporting fails to detect it.<sup>o</sup>

## WHAT HAPPENED

The VIGOR trial (Bombardier et al., NEJM 2000) randomized 8,076 patients in all, and reported a five-fold higher rate of myocardial infarction among those taking Vioxx than among those taking naproxen. Merck and many readers interpreted the gap as naproxen being cardio-protective rather than Vioxx being harmful. That reading was not absurd — it was the more comfortable of two explanations for the same numbers — but it required the absent placebo arm to be true, and the data could not adjudicate between the readings. The placebo-controlled APPROVe trial (Bresalier et al., NEJM 2005), originally designed to test Vioxx for colorectal-polyp prevention, was terminated early in September 2004 when its data safety monitoring board observed a doubling of cardiovascular events in the Vioxx arm relative to placebo. Merck withdrew the drug worldwide within days.<sup>1</sup>

## THE INVESTIGATION

Senate Finance Committee hearings in November 2004 put on the record that signals of cardiovascular harm had been present in the trial record for years before withdrawal. FDA Office of Drug Safety scientist David Graham testified under oath that the cardiovascular risk had been visible to him by 2000, that he had been pressured by FDA management not to publish his estimate, and that he believed Vioxx had caused 88,000–139,000 excess cases of heart attack or sudden cardiac death in the United States, of which 30–40% were probably fatal. The Graham et al. Lancet 2005 analysis, using Kaiser Permanente data, produced a population-level estimate in similar range. Merck litigation discovery, made public through New Jersey and federal court filings and reported in the NEJM editorial trail, included internal Merck communications and ghost-authorship patterns in published Vioxx papers; a 2008 JAMA analysis by Ross et al. documented the guest-authorship and ghostwriting patterns directly.<sup>2</sup> The harm was not hidden in some unmeasured corner — it sat in the trial record the whole time, waiting for an architecture that would carry it to a decision rather than leave it to interpretation.<sup>3</sup>


## THE CAPABILITY GAP

The capability gap was in post-market surveillance: the disclosure architecture between manufacturer-held data, FDA reviewers, and prescribers did not aggregate the signal to a decision boundary for nearly four years. The FDA's MedWatch adverse-event reporting system relied on voluntary submissions structured to minimize signal, and was not adequate to the volume of the drug's distribution. A system that waits for a clinician to choose to file a report will always lag a harm spread thinly across millions of prescriptions, because no single prescriber sees enough of the pattern to recognize it. The Merck-side gap was structural in a different way: VIGOR was designed to test gastrointestinal safety (the COX-2 selling point), and the cardiovascular signal was a secondary finding that the trial was not powered to adjudicate against the naproxen-protective hypothesis. That trial-design choice compounded with a publication-bias pattern, an FDA adverse-event-reporting architecture insufficient for the population at risk, and a regulator-manufacturer disclosure protocol that had no decision boundary for ambiguous signals.<sup>4</sup>

## AFTERMATH & REFORM

Merck eventually settled U.S. litigation for about \$4.85 billion across approximately 27,000 plaintiffs. The case drove a coordinated regulatory response. The FDA Amendments Act of 2007 gave the FDA explicit authority to require Risk Evaluation and Mitigation Strategies

(REMS) and post-market study commitments as conditions of approval. In 2008 the FDA launched the Sentinel Initiative — an active post-market surveillance system that queries distributed health-data partners covering hundreds of millions of patient-years, rather than waiting for voluntary reports — a direct response to the Vioxx-era detection failure. The NEJM tightened conflict-of-interest disclosure for trial reports and, in December 2005, took the unusual step of issuing an Expression of Concern: three myocardial infarctions had been omitted from the data submitted, and other cardiovascular data had been deleted from the manuscript's table two days before submission — evidence that hardens the story from “the signal aggregated too slowly” toward “the signal was partly, and knowingly, kept out of the record.” The JAMA Ross et al. (2008) analysis became the reference point for diagnosing guest authorship and ghostwriting in drug safety. By going out to the data instead of waiting for it to arrive, the reform inverted the logic that had let the signal sit unaggregated for years. The reform built the surveillance infrastructure that had not previously existed; the Vioxx record is what made the case for it.<sup>5</sup>



VIOXX — POST-MARKET SURVEILLANCE AS A MISSING DELIVERABLE

*The cardiovascular risk was visible in Merck's internal data years before the drug was withdrawn.*

PARAPHRASING DAVID GRAHAM (FDA OFFICE OF DRUG SAFETY), SENATE FINANCE COMMITTEE TESTIMONY, NOVEMBER 2004

## REFERENCES

1. Graham, D. J. et al. (2005), “Risk of acute myocardial infarction and sudden cardiac death in patients treated with COX-2 selective and non-selective NSAIDs: nested case-control study,” *Lancet* 365(9458):475–481 — Kaiser Permanente population-level cardiovascular risk analysis.
2. Bresalier, R. et al. (2005), “Cardiovascular events associated with rofecoxib in a colorectal adenoma chemoprevention trial,” *APPROVe*, *NEJM* 352(11):1092–1102 — placebo-controlled confirmation of cardiovascular risk; early trial termination.
3. Ross, J. S., Hill, K. P., Egilman, D. S., Krumholz, H. M. (2008), “Guest authorship and ghostwriting in publications related to rofecoxib,” *JAMA* 299(15):1800–1812 — guest-authorship and ghostwriting documentation from Merck litigation discovery.
4. U.S. Senate Committee on Finance, hearings on Vioxx and FDA's drug-safety system (November 18, 2004) — Graham testimony; “88,000 to 139,000 Americans” estimate of excess cardiovascular events; described FDA management pressure.
5. Bombardier, C. et al. (2000), “Comparison of upper gastrointestinal toxicity of rofecoxib and naproxen in patients with rheumatoid arthritis,” *VIGOR trial*, *NEJM* 343(21):1520–1528 — the five-fold myocardial-infarction signal and the naproxen-protective hypothesis.
6. FDA Amendments Act of 2007 (P.L. 110–85) and FDA Sentinel Initiative documentation (2008–present) — REMS authority, post-market study requirements, and active distributed-data post-market surveillance.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Vioxx is the canonical pharmaceutical case for post-market surveillance as a capability deliverable. The signal existed; the institutional architecture to aggregate it did not. The reform pattern — Sentinel — built the architecture. The case is the drug-industry analog of the EHR case (Case 2): a measurement architecture too thin for the system that depended on it.

LENS APPROACH

Vioxx is the canonical pharmaceutical post-deployment-surveillance case (induced 2.4; LENS D4+D3/PT5). LENS uses it in Domain 4 (Test and Evaluation; LEO-4) for post-market surveillance as a measurement deliverable — the signal existed; the institutional architecture to aggregate it did not. LENS uses it in Domain 3 (Human-System Collaboration; LEO-3) for the change-control and disclosure architecture between manufacturer, regulator, and prescriber that determines whether ambiguous safety signals reach a decision boundary. Pair with Radiology AI (Case 35) as the post-market-surveillance failure pattern at a new technological boundary; pair with EHR/CPOE (Case 2) at the measurement-architecture-too-thin layer; pair with Medical Errors and IOM (Case 8) at the population-scale-instrument layer.

| DOMAIN                   | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO   |
|--------------------------|---------------|------------------|--------------|---------|-------|
| HEALTHCARE               | → G D         | → D1 D2 D3 D4 D5 | 5            | 2.4     | 4 · 3 |
| Test and Evaluation + D3 |               |                  |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Field an active post-market surveillance architecture sized to the exposed population before scale-up, so a harm spread thinly across millions can still aggregate to a decision boundary.
- ▶ Design trials and data collection to adjudicate between competing readings of a signal, rather than leaving an ambiguous result to the more comfortable interpretation.
- ▶ Treat signal-aggregation as a defined deliverable carrying evidence from manufacturer data through reviewers to prescribers, not a byproduct of voluntary reporting.

IN OPERATION / AFTER

- ▶ Audit the trial record and adverse-event data for signals already present but never carried to a decision, closing the lag between visibility and action.
- ▶ Monitor by querying health-data partners actively rather than waiting for voluntary reports that no single prescriber sees enough of the pattern to file.
- ▶ Sustain post-market surveillance funding against the commercial pressure to defer aggregation, since the cost of waiting compounds across the exposed population.

REFLECTION QUESTIONS

1. Identify a post-deployment surveillance architecture in your domain that is too thin for the scale of the system it monitors. What is the missing deliverable?
2. Design the Sentinel-equivalent post-market surveillance system for a new domain.
3. The VIGOR signal admitted two readings — harmful drug or protective comparator — and the data could not decide. Identify a measurement in your domain whose interpretation is underdetermined, and design the study or instrument that would force the question to a decision.

WHO BUILDS THIS governance, ethics & measurement · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · task analysis · design prototyping

Chapter 2

# Healthcare & Patient Safety — What Works — and the Frontier

*When measurement, teamwork, and design close the loop on patient harm.*

---

*Every durable safety gain shipped with an owner and an instrument.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

# California Nurse Staffing Ratios — Legislating a Capability Requirement

**G** Governance & Trust   **K** Knowledge & Institutional Memory

**IMPACT** Under California's mandated minimum nurse-to-patient ratios its nurses cared for 1–2 fewer patients each than nurses in comparison states; modeling implied 10–14% fewer surgical patient deaths in those states if matched — observational, cross-sectional, no California baseline

## IN BRIEF

California in 2004 became the first U.S. state to mandate minimum nurse-to-patient ratios in acute-care hospitals — unit-level minimums written into law and enforced through inspection. The Aiken et al. (Health Services Research, 2010) study surveyed 22,336 nurses across California, Pennsylvania, and New Jersey (the latter two with no mandate), and found California nurses cared for 1–2 fewer patients each. Modeling implied that if the two comparison states had matched California's medical-surgical ratios, New Jersey would have seen 13.9% and Pennsylvania 10.6% fewer surgical patient deaths. The study is observational and cross-sectional, and the authors are explicit that the **absence of California baseline measures** fueled a contested stakeholder debate over whether the ratios themselves caused the improvement. The case is the canonical recent instance of a capability requirement (adequate staffing) being converted from a stated aspiration into an engineered, enforced specification by law — and the canonical instance of the methodological hedge such a conversion carries.

## BACKGROUND

Nurse staffing is the variable in the capability-of-care literature: a substantial body of evidence links nurse workload to medical errors, mortality, and rescue from complications. The professional consensus on the direction of the effect has been durable since the late 1990s. What had been contested is whether a mandated minimum, written into law and enforced through inspection, would convert the stated requirement (adequate staffing) into an engineered one (a specific ratio actually present at the bedside).<sup>0</sup>

## THE INTERVENTION

California in 2004 became the first U.S. state to do it. The law specifies unit-by-unit minimum nurse-to-patient ratios — 1:6 in medical-surgical at the 2004 start and 1:5 from 2005, with tighter ratios in step-down and intensive care — enforceable through state inspection. The political process took five years from statutory enactment (1999) to regulatory implementation, with industry argument that the ratios would close hospitals and worsen access. The 2004 implementation went ahead substantially as written.<sup>1</sup>



## HOW IT WORKED

Aiken et al. (Health Services Research, 2010) is the strongest published evaluation. The study surveyed 22,336 nurses across California (mandated), Pennsylvania, and New Jersey (no mandate) on workload, work environment, and patient-outcome measures. The headline finding is that California nurses cared for 1–2 fewer patients each across comparable unit types. The modeled mortality implication — computed by applying the workload–mortality relationship from the broader Aiken et al. corpus to the workload difference — was that if Pennsylvania and New Jersey had matched California’s medical–surgical ratios, surgical mortality would have been 10.6% and 13.9% lower respectively.<sup>2</sup>

## THE EVIDENCE

The hedge survives into the case verbatim. The Aiken et al. study is observational and cross-sectional, not a controlled experiment. There was no California baseline measurement to provide a pre/post within-state comparison. The modeled mortality estimates rest on the workload–mortality relationship from prior studies, and debate followed publication over whether the ratios themselves caused the observed improvement, or whether California hospitals differed in other ways the cross-section did not capture. The authors do not claim closure; their conclusion is that the evidence is the strongest available and consistent with the broader literature.<sup>3</sup>

## WHAT TRANSFERRED

What the case teaches is the form of converting a stated capability requirement into an engineered, enforced one by law, and the methodological hedge that conversion has to carry. The political path is part of the deliverable: five years from enactment to implementation, sustained industry opposition, contested evidence, and a study design constrained by the absence of a pre-implementation California baseline. In pair with SUBSAFE (Case 173), the case shows that engineered capability requirements at scale require both the requirements–discipline mechanism and an evidence architecture that survives the political process around the requirement.

*The headline finding is direct; the modeled mortality estimate is the strongest available evidence and not closed proof.*

EDITORS’ SYNTHESIS OF AIKEN ET AL. (2010).

## REFERENCES

1. Aiken, Clarke, Sloane, Sochalski, & Silber (2002), “Hospital nurse staffing and patient mortality, nurse burnout, and job dissatisfaction,” JAMA 288(16):1987–1993 — the workload–mortality relationship the 2010 modeling rests on.
2. California Department of Health Services (1999–2004), AB 394 regulatory implementation documentation.
3. Aiken, Sloane, Cimiotti, Clarke, Flynn, Seago, Spetz, & Smith (2010), “Implications of the California Nurse Staffing Mandate for Other States,” Health Services Research 45(4):904–921, doi:10.1111/j.1475-6773.2010.01114.x.
4. Spetz, Chapman, Herrera, Kaiser, Seago, & Dower (2009), “Assessing the impact of California’s nurse staffing ratios on hospitals and patient care,” California HealthCare Foundation — implementation-period analysis.

## THE LEARNING ENGINEERING LENS

## LE INSIGHT

California’s nurse–ratio mandate is the canonical recent instance of converting a stated capability requirement (adequate staffing) into an engineered, enforced one by law. The workload effect is direct; the modeled mortality estimate is observational and cross–sectional with no California baseline. The hedge is the case.

LENS APPROACH

Nurse ratios is the legislated–requirement case (induced 1.1; LENS D1/PT3). LENS uses it in Domain 1 (Systems Analysis) for the conversion of stated requirement to engineered specification by law, and in Domain 4 (Test and Evaluation) for the subobjective **Judgment under inadequate evidence** (4.4) — the evaluation evidence is observational cross–sectional, the strongest available and not closed proof. Pair with Case 173 SUBSAFE at the requirements–as–deliverable layer.

| DOMAIN           | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO   |
|------------------|-----------------|------------------|--------------|---------|-------|
| HEALTHCARE       | G K             | D1 D2 D3 D4 D5   | 3            | 1.1     | 1 · 4 |
| WORKFORCE POLICY |                 | Systems Analysis |              |         |       |
| PATIENT SAFETY   |                 |                  |              |         |       |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>When converting a stated capability requirement into an engineered specification by law, design the evidence architecture <b>before</b> implementation: pre–registered measurement, a baseline window, and a comparison strategy that survives the political process.</li><li>Specify the unit–level requirement at the granularity the work actually has — for nurse staffing, by unit type and acuity — rather than at an aggregate level that lets the requirement be averaged away.</li><li>Anticipate sustained industry argument as part of the implementation cost; a five–year path from enactment to implementation is the political price of the engineered requirement.</li></ul> | <ul style="list-style-type: none"><li>Report the workload effect (1–2 fewer patients per nurse) separately from the modeled mortality effect (10.6% and 13.9%); the workload finding is direct, the mortality estimate is modeled.</li><li>Preserve the no–baseline hedge in any communication of the evidence; the strongest available evidence is not closed proof of causation.</li><li>Treat the contested stakeholder debate as program documentation, not noise; the debate is part of why the case is teachable.</li></ul> |

REFLECTION QUESTIONS

1. Identify a capability requirement in your domain that is currently stated but not engineered. What would a legislated minimum look like at unit–level granularity, and what evidence architecture would you put in place **before** implementation to make the conversion auditable?
2. Specify the political–process cost you would budget for: California’s path was five years from enactment to implementation under sustained industry argument. What is the analog in your context, and is the engineered requirement worth that cost?
3. The Aiken evidence is observational cross–sectional with no California baseline. What pre/post or controlled comparison would you build into the next state’s implementation to give the next round of evaluation a stronger foundation?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · knowledge bases · data pipelines

# ACGME 80-Hour Resident Duty-Hour Reform

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** ACGME capped U.S. resident physician work hours at 80/week to reduce fatigue-related errors; subsequent RCTs (FIRST, iCOMPARE) found flexible schedules non-inferior on patient outcomes, and the promised safety gain did not appear

## IN BRIEF

After the 1984 death of Libby Zion focused decades of concern on resident-physician fatigue, the ACGME capped resident work at 80 hours a week in 2003 and limited first-year shifts to 16 hours in 2011. The logic was intuitive: tired doctors err, so cut the hours. But hours were one input to a many-variable system; cutting them multiplied error-prone hand-offs and cost residents continuity and procedural repetitions. Two large randomized trials — FIRST (2016) and iCOMPARE (2019) — found flexible schedules non-inferior to the strict caps on patient outcomes, so the promised safety gain never appeared, and in 2017 the ACGME relaxed the intern cap. The case is the book's clearest example of a single-variable intervention into a capability system — and of evidence catching up with a well-meant policy.

## BACKGROUND

After the 1984 death of Libby Zion — blamed on overworked, under-supervised residents — resident fatigue became a decades-long argument. New York's Bell Commission produced the first hours limits in 1989, and pressure for a national standard built from there — the reform's intuition, that exhausted physicians make errors, was strong enough that the simplest lever, capping the hours, became the obvious answer long before anyone tested what the long shift was producing.<sup>o</sup>

## THE INTERVENTION

In 2003 the Accreditation Council for Graduate Medical Education capped resident work at 80 hours a week; in 2011 it limited first-year residents to 16-hour shifts. The logic was clean and intuitive — fatigue causes error, so reduce the hours. It was, in capability terms, a single-variable intervention: change one input and expect the outcome to move — a model that holds only if the rest of the system stays fixed, which in a teaching hospital it never does, since the hours removed had to be absorbed somewhere else.<sup>1</sup>

## HOW IT WORKED

But hours were one input to a system with many. Cutting them redistributed the work and multiplied patient hand-offs — themselves a documented site of error — while reducing residents' continuity and the procedural repetitions that build skill; many reported feeling


less prepared, not better rested. The long shift had quietly been doing work no one accounted for — sustaining a patient's care through one set of hands and accumulating the repetitions that turn a trainee into a clinician — and nothing was put in its place when it was cut.<sup>2</sup>

### THE EVIDENCE

Two large randomized trials tested the policy. FIRST (Bilimoria et al., NEJM 2016), in surgery, found flexible duty hours non-inferior to the strict caps on patient outcomes and no worse for resident well-being — putting a randomized result against an intuition that had driven policy for years.<sup>3</sup> iCOMPARE (Silber et al., NEJM 2019), in internal medicine, reached a parallel result in a second specialty, making the finding harder to dismiss as an artifact of surgery.<sup>4</sup> Neither found the safety gain the cap had promised. In 2017, with FIRST reported and iCOMPARE under way, the ACGME relaxed the 16-hour intern limit. The trials did not show fatigue is harmless — only that cutting one input, without rebuilding supervision and hand-offs, did not produce a safer system.<sup>5</sup>

### WHAT TRANSFERRED

Duty-hour reform is the book's clearest case of a single-variable intervention into a multi-variable system. Read it against the successes here — the Keystone ICU project, crew resource management, the surgical safety checklist — which worked because they engineered supervision, hand-offs, and measurement **together with** the behavioral change, redesigning the surrounding architecture rather than pulling a single lever and hoping the rest would hold.<sup>6</sup> The lesson is not that fatigue does not matter; it is that capability is a property of the whole system, and a reform that moves one variable while leaving the others untouched will be judged, in the end, by what it actually produced rather than by the plausibility of its intuition.



*As compared with standard duty-hour policies, flexible, less-restrictive duty-hour policies for surgical residents were associated with noninferior patient outcomes and no significant difference in residents' satisfaction with overall well-being and education quality.*

FIRST TRIAL CONCLUSIONS (BILIMORIA ET AL., NEJM 2016)

### REFERENCES

1. B. H. Lerner, The Libby Zion Case and the Reform of Medical Education (2006); and the 1989 New York State (Bell Commission) duty-hour regulations — the origin of the reform.
2. Accreditation Council for Graduate Medical Education, Common Program Requirements (2003 and 2011 revisions) — the 80-hour weekly cap and the 16-hour first-year shift limit.
3. D. A. Asch, K. Y. Bilimoria & S. V. Desai, "Resident Duty Hours and Medical Education Policy — Raising the Evidence Bar," NEJM 376: 1704–1706 (2017); Institute of Medicine (Ulmer, Wolman & Johns, eds.), Resident Duty Hours: Enhancing Sleep, Supervision, and Safety (2008) — hand-offs and continuity trade-offs.
4. Bilimoria, Chung, Hedges et al., "National Cluster-Randomized Trial of Duty-Hour Flexibility in Surgical Training" (FIRST), NEJM 374: 713–727 (2016). doi:10.1056/NEJMoa1515724.
5. Silber, Bellini, Shea et al., "Patient Safety Outcomes under Flexible and Standard Resident Duty-Hour Rules" (iCOMPARE), NEJM 380: 905–914 (2019). doi:10.1056/NEJMoa1810642.
6. ACGME, Common Program Requirements (2017 revision) — relaxation of the 16-hour first-year shift limit.
7. P. Pronovost et al., Keystone ICU intervention, NEJM 355: 2725–2732 (2006). doi:10.1056/NEJMoa061115 A. B. Haynes et al., surgical safety checklist, NEJM 360: 491–499 (2009). doi:10.1056/NEJMsa0810119 — integrated interventions that engineered the surrounding architecture.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The clearest healthcare case in the dataset of a single-variable intervention into a multi-variable capability system. Pairs with Case 127 (fratricide) and Case 134 (V-22). The success cases — Keystone (19), CRM (117), Korean Air (118) — engineered supervisory, hand-off, and measurement architecture **together with** the behavioral change.

### LENS APPROACH

LENS uses duty-hour reform in LEN 5 as the foundational capability-system case (what was the long-hour regime **producing** that was lost when the input was capped?), in LEN 4 to discuss measurement architecture (what FIRST and iCOMPARE measured, and what they did not), and in LEN 10 as a studio prompt for the integrated resident-training redesign that the reforms did not deliver.

| DOMAIN                | Modes Addressed | LENS Competency  | Problem Type | Induced | LEO |
|-----------------------|-----------------|------------------|--------------|---------|-----|
| HEALTHCARE            | → T K N         | → D1 D2 D3 D4 D5 | 3            | 2.3     | 2   |
| Iterative Development |                 |                  |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Map the full set of variables the targeted input is coupled to — continuity, supervision, hand-offs, procedural reps — before changing one of them.
- ▶ Design the substitute for whatever the changed input was producing (e.g., structured hand-offs and supervision) into the same reform, not as a follow-on.
- ▶ Build the measurement that will test the reform's promised gain into the rollout, so the policy is falsifiable against the operational record.

#### IN OPERATION / AFTER

- ▶ Audit the intervention against patient and trainee outcomes with a controlled comparison, as FIRST and iCOMPARE did, rather than trusting the intuition.
- ▶ Monitor the variables that absorbed the change (hand-off frequency, procedural exposure) for the harms a single-lever fix can displace.
- ▶ Sustain a willingness to revise the policy when evidence catches up, as the 2017 relaxation did, rather than defending the original lever.

### REFLECTION QUESTIONS

1. What capability is the long-hours / heavy-workload regime in your domain currently producing — supervisory exposure, continuity, procedural reps, tacit-knowledge transfer — that a simple cap would lose?
2. Design the integrated redesign — supervision, hand-off, measurement, exposure — that would substitute for the capability the input cap removes.
3. The reform was intuitive enough to set national policy years before FIRST and iCOMPARE tested it. Design the randomized or quasi-experimental check you would build into a future single-variable reform so its promised gain is measured before it is mandated, not after.

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

# Implementation Science in Healthcare — The 17-Year Gap

**K** Knowledge & Institutional Memory   **G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** The canonical implementation-gap pair — about seventeen years from evidence to practice, about 14 percent of original research arriving — both estimates: Balas and Boren timed nine mostly primary-care services to 50 percent adoption and derived the 14 percent from a chain of assumed attrition points

## IN BRIEF

Implementation science has a canonical pair of figures: about seventeen years for research evidence to reach clinical practice, and only about 14% of original research arriving at all. Both come from Balas and Boren (2000), and both are estimates rather than measurements — the seventeen years is a lag across nine mostly primary-care services, and the 14% is derived from assumed attrition at each stage of the publication-to-guideline pipeline. This is not a single incident but a systemic condition — effective interventions exist; the system to adopt, sustain, adapt, and measure them at scale does not. Frameworks like the Active Implementation Frameworks (Fixsen et al., 2005) and EPIS (Aarons et al., 2011) were built specifically to attack this gap, and LENS threads implementation science throughout its curriculum in direct response. The seventeen-year gap is the meta-case for the whole book: every success case is a closure of this gap in one domain, every failure case the gap left open. The discipline exists to make seventeen years shorter.

## THE SHIFT

Medicine generates more validated knowledge than it can absorb. Implementation science arose to study a stubborn fact: knowing what works and having it practiced are different problems, separated by years.<sup>o</sup> Treating the two as one problem is the error the field formed to correct: a validated finding is not yet a changed practice, and the distance between them is itself a phenomenon to be studied, measured, and engineered rather than waited out.

## WHAT IS EMERGING

The canonical figures are stark, and their provenance matters. Balas and Boren identified nine mostly primary-care services supported by randomized trials and timed each from trial publication to 50 percent adoption: 15.6 years, which becomes the familiar seventeen once publication lag is added. The 14 percent is not a count of anything; it is what remains after applying assumed attrition rates at each stage from submission through systematic review to guideline. Later scrutiny found the utilization baselines wrong in places — the

analysis assumed 20 percent of diabetic patients received foot exams in 1998 when the rate was nearer 60, and dated 50 percent mammography adoption to 1993 when it had been reached by 1990 — and the denominators include patients without regular checkups, so the lag partly measures patient behaviour rather than clinician uptake.<sup>1</sup> Read together, the two figures describe a pipeline that is both slow and leaky: most of what is learned never reaches the bedside at all, and the fraction that does arrives long after the patients who first needed it, so the delay is compounded by sheer attrition.

#### THE CAPABILITY QUESTION


This is not a single case but a systemic condition — the same structural problem the medical-error data (Case 8) describes from the outcome side. Effective interventions exist; the institutional system to adopt, sustain, adapt, and measure them at scale does not.<sup>2</sup> Where the medical-error data counts the harm at the far end of the pipeline, the translation figures name the mechanism that produces it: the same missing adoption-and-measurement system shows up as a delay from one vantage and as a body count from the other.

#### EARLY EVIDENCE

Frameworks built to attack the gap — the Active Implementation Frameworks (Fixsen et al., 2005) and the EPIS framework (Aarons et al., 2011) — show that implementation can be engineered rather than left to chance, and they inform LENS's choice to thread implementation science through every course rather than isolate it in a module.<sup>3</sup> Threading the discipline through every course rather than confining it to a single module is itself a claim these frameworks support: if implementation is an engineerable property of any intervention, it cannot be quarantined as a specialty and must inform how every design is taught.

#### OPEN PROBLEMS

The seventeen-year gap is the meta-case for this book. Every success case in the chapters ahead is a closure of the gap in one domain; every failure case is the gap left open.<sup>4</sup> The open problem is general and unglamorous: building, funding, and owning the adoption-and-measurement pathway that turns a proven intervention into routine practice — which is, in one sentence, what the LENS discipline exists to do. Because the problem is general rather than domain-specific, no single clinical result closes it; what closes it is the repeatable, owned, and funded pathway that any proven finding can be run through, again and again.



*The answer is 17 years, what is the question*



MORRIS, WOODING & GRANT, J ROYAL SOC MED, 2011

REFERENCES

1. E. A. Balas & S. A. Boren (2000), "Managing clinical knowledge for health care improvement," Yearbook of Medical Informatics — the nine-service adoption analysis behind the seventeen-year figure, and the pipeline-attribution derivation behind the 14 percent.

2. Z. S. Morris, S. Wooding & J. Grant, "The answer is 17 years, what is the question: understanding time lags in translational research," J. Royal Society of Medicine 104(12):510–520 (2011) (quoted).

3. D. Fixsen et al., Implementation Research: A Synthesis of the Literature (2005) — the Active Implementation Frameworks.

4. G. Aarons et al. (2011), the EPIS framework; L. Damschroder et al. (2009), CFIR.

5. Cf. medical error as systemic failure (Case 8); Goodell & Kolodner, Learning Engineering Toolkit (2022).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The 17-year gap is the structural problem that LENS exists to address. It is the difference between knowing what works and having it deployed. Every case in this book is a sample from a distribution governed by that gap. The success cases shorten it; the failure cases let it run.

LENS APPROACH

LENS uses this case in LEN 1 as the foundational problem statement of the discipline, in LEN 10 as a studio constraint (designs must consider implementation, not just efficacy), and in LEN 8 as the central knowledge-transfer challenge. The case is referenced at least once in every required course.

| DOMAIN                | MODES AT STAKE | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|-----------------------|----------------|------------------|--------------|---------|-----|
| HEALTHCARE            | → K G N        | → D1 D2 D3 D4 D5 | 4            | 1.4     | 2   |
| Iterative Development |                |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Design every intervention with its adoption-and-measurement pathway attached, so implementation is engineered in rather than left to chance after the evidence is published.
- Use an implementation framework (Active Implementation, EPIS) from the outset to plan adoption, sustainment, and adaptation as deliverables of the project.
- Name an owner and a funding line for the pathway, since a proven finding with no one accountable for fielding it is exactly what the gap is made of.

IN OPERATION / AFTER

- Measure both reach and speed — what fraction of practice has adopted the intervention and how long it took — to see the slow-and-leaky pipeline rather than assume publication equals uptake.
- Sustain and adapt fielded interventions, treating drift back to old practice as a measurable failure mode, not a one-time rollout that holds itself.
- Track the gap as a standing metric across the institution, so closing it in one domain becomes a repeatable pathway rather than a one-off success.

## REFLECTION QUESTIONS

1. Pick an evidence-based intervention in your domain. Estimate the gap between when the evidence became conclusive and when the intervention reached majority of practice. What did the gap cost?
2. Design the deliverable that would shorten that gap by half in your domain. Be specific about who funds it, who owns it, and what evidence demonstrates the reduction.
3. The translation pipeline is both slow and leaky — most findings never reach practice, and those that do arrive late. For your domain, is the binding problem the delay or the attrition, and what would you measure to tell which one to attack first?

**WHO BUILDS THIS** knowledge management & data engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA

**T** Training Gap   **K** Knowledge & Institutional Memory

**IMPACT** Barsuk et al. (Northwestern/Feinberg) demonstrated that simulation-based mastery learning for central venous catheter (CVC) insertion reduced catheter-related bloodstream infection rates and procedural complications, with cost-effectiveness shown in a single-hospital evaluation; the program was subsequently disseminated nationally to Department of Veterans Affairs medical centers

## IN BRIEF

Central venous catheter (CVC) insertion is a high-volume, high-consequence procedure whose complications — pneumothorax, arterial puncture, catheter-related bloodstream infection (CRBSI) — are well-characterized and partly attributable to procedural skill at the bedside. Barsuk and colleagues at Northwestern/Feinberg published a series of studies through the late 2000s and 2010s establishing simulation-based mastery learning (SBML) — practice on a simulator to a defined performance standard, not to a clock — as an effective approach for trainee CVC insertion. The single-center evidence reported reduced CRBSI rates, reduced procedural complications, and favorable cost-effectiveness. The program was subsequently disseminated to Department of Veterans Affairs medical centers as a national rollout — the deliberate cross-institutional transfer that distinguishes the case from a successful single-site intervention. The case is the canonical small-tier instance of cross-domain adaptation of a proven intervention (C6.4), and pairs with the multidisciplinary-translation trio (Cases 40 team science, 28 IPE, 41 implementation-science training) as cross-domain workforce evidence. The honest hedge: the dissemination outcome literature is thinner than the original single-center evidence; the case is teachable on the SBML method and the documented dissemination effort, with measurement at the multi-site level a live evidence frontier.

## BACKGROUND

Central venous catheter insertion is a procedure trainees in medicine, surgery, and critical care perform routinely; the complication profile — pneumothorax, arterial puncture, catheter-related bloodstream infection — is well-characterized and attributable in part to procedural skill at the bedside. The historical training model relied on graduated exposure on live patients, with supervision but without a defined performance standard before unsu-

pervised practice. The Barsuk group at Northwestern's Feinberg School of Medicine, working across critical care and medical education, set out to replace exposure-based training with simulation-based mastery learning.<sup>9</sup>

#### THE INTERVENTION

Simulation-based mastery learning is the specific structural form: trainees practice the procedure on a simulator to a defined performance standard, not to a clock. Practice continues until the trainee demonstrates competence; the standard, not the calendar, governs progression. Barsuk et al. (Archives of Internal Medicine, 2009; Critical Care Medicine, Journal of Hospital Medicine, and a series of subsequent papers) reported that the SBML cohort, compared with historical and concurrent non-SBML cohorts at the same hospital, performed CVC insertions with fewer needle passes, fewer arterial punctures, and lower CRBSI rates on the patients they subsequently catheterized.<sup>1</sup>

#### HOW IT WORKED

The single-center economic evidence was the second pillar. Cohen et al. (Simulation in Healthcare, 2010) reported that the SBML program was cost-saving at the single hospital, driven principally by the reduction in CRBSI cases — each bloodstream infection averted is expensive enough that even a modest reduction recovers the simulation-training investment. The combination of the procedural-outcome evidence and the cost-effectiveness evidence made the program transferable beyond the original site.<sup>2</sup>

#### THE EVIDENCE

The dissemination step is what distinguishes the case from a successful single-site intervention. The SBML CVC program was subsequently disseminated to Department of Veterans Affairs medical centers as a national rollout, supported by the VA's investment in simulation infrastructure and clinical training architecture. The deliberate cross-institutional transfer — from an academic medical center to a federated health system — is the structural feature the C6.4 cross-domain-adaptation competency exists to name. The dissemination outcome literature is thinner than the single-center evidence; multi-site CRBSI tracking is harder to attribute, and the published evidence at the national-VA scale is at the program-report and small-evaluation tier rather than the controlled-comparison tier the single-center papers established.<sup>3</sup>

#### WHAT TRANSFERRED

In pair with the multidisciplinary-translation trio (Cases 40 team science, 28 IPE, 41 implementation-science training), the Barsuk SBML case completes the workforce-evidence picture: an intervention with strong single-center controlled-comparison evidence (this case) sits alongside a structured team-science training program with validated measurement (40), the field-scale enthusiasm-evidence gap in interprofessional education (28), and the operational-practice gap inside implementation-science training programs (41). The four cases together stage the cross-domain workforce-evidence pattern: the mechanism that works at single-program scale is the demonstrator; the field-scale measurement and the multi-site dissemination evidence are what the discipline is still building.<sup>4</sup>

*The standard, not the calendar, governs progression. The trainee practices until they demonstrate competence; the simulator absorbs the cost of the practice.*

EDITORS' SYNTHESIS OF BARSUK ET AL. (2009).

## REFERENCES

1. Barsuk, McGaghie, Cohen, Balachandran, & Wayne (2009), "Use of simulation-based mastery learning to improve the quality of central venous catheter placement in a medical intensive care unit," *Journal of Hospital Medicine* 4(7):397–403.
2. Barsuk, Cohen, Feinglass, McGaghie, & Wayne (2009), "Use of simulation-based education to reduce catheter-related bloodstream infections," *Archives of Internal Medicine* 169(15):1420–1423, doi:10.1001/archinternmed.2009.215.
3. Cohen, Feinglass, Barsuk, Barnard, O'Donnell, McGaghie, & Wayne (2010), "Cost savings from reduced catheter-related bloodstream infection after simulation-based education for residents in a medical intensive care unit," *Simulation in Healthcare* 5(2):98–102, doi:10.1097/SIH.0b013e3181bc8304.
4. Barsuk, Cohen, Nguyen, Mitra, O'Hara, Okuda, Feinglass, Cameron, McGaghie, & Wayne (2016), "Attending physician adherence to a 29-component central venous catheter bundle checklist during simulated procedures," *Critical Care Medicine* 44(10):1871–1881 — the published record of the train-the-trainer dissemination across 58 Veterans Affairs Medical Centers.
5. McGaghie, Issenberg, Cohen, Barsuk, & Wayne (2011), "Does simulation-based medical education with deliberate practice yield better results than traditional clinical medical education? A meta-analytic comparative review of the evidence," *Academic Medicine* 86(6):706–711.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Barsuk SBML for CVC insertion is the canonical small-tier intervention case for cross-institutional dissemination of a proven mechanism. The single-center evidence is controlled-comparison and cost-effective; the VA national dissemination is the cross-institutional transfer step. The multi-site outcome evidence is thinner than the single-center evidence — the live frontier of the dissemination case.

### LENS APPROACH

Barsuk SBML is the cross-institutional dissemination case (induced 6.4; LENS D2/PT4) — Domain 2 for mastery-learning replacing exposure-based progression; Domain 4 for the procedural-outcome + cost-effectiveness pairing. Pair with the translation trio (Cases 40, 28, 41).

| DOMAIN                  | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|-------------------------|-----------------|-----------------------|--------------|---------|-------|
| MEDICAL EDUCATION       | T K             | D1 D2 D3 D4 D5        | 4            | 6.4     | 2 · 4 |
| PATIENT SAFETY          | →               | Iterative Development |              |         |       |
| WORKFORCE DISSEMINATION |                 |                       |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Replace exposure-based training with a defined performance standard on a simulator; the SBML deliverable is competence demonstrated, not hours accumulated.
- ▶ Design the simulator and the standard around the specific complication modes the procedure produces — pneumothorax, arterial puncture, CRBSI for CVC — so the training closes the gaps the outcome data identifies.
- ▶ Pair the procedural-outcome evidence with cost-effectiveness evidence at single-site scale; the combination is what makes the program disseminable beyond the original institution.

### IN OPERATION / AFTER

- ▶ Plan the dissemination as a deliverable of the original program, not as a separate downstream activity; the VA rollout is what makes Barsuk SBML a cross-institutional intervention rather than a successful single-site study.
- ▶ Build multi-site outcome measurement into the dissemination from the start; the published evidence at national-VA scale is thinner than the single-center evidence, and the gap is a live frontier the case names.
- ▶ Carry the cross-domain pairing with team-science training (Case 40), IPE (Case 28), and implementation-science training (Case 41) into the curricular framing — the four cases together stage what cross-domain workforce dissemination looks like with measurement and what it looks like without.

## REFLECTION QUESTIONS

1. Identify a high-volume, high-consequence procedural skill in your domain where training is currently exposure-based rather than performance-standard-based. What is the analog of SBML, and what is the simulator that would absorb the cost of the practice?
2. Specify the procedural-outcome and cost-effectiveness pairing you would design at single-site scale; both halves are what made the Barsuk program disseminable, and the pairing is what distinguishes a teachable single-site case from a national-rollout candidate.
3. The VA dissemination evidence is thinner than the single-center evidence. Design the multi-site outcome-tracking architecture you would build into the next dissemination — what instrument, what comparison, what cadence — so the multi-site evidence catches up to the demonstration evidence.

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines

# I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer

**H** Human-System Interface   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Across nine pediatric residency programs, implementation of the I-PASS handoff bundle (mnemonic + training + faculty development + sustainability campaign) was associated with a 23% relative reduction in medical errors and a 30% reduction in preventable adverse events, without negatively affecting resident workflow — the study design 'precludes definitively establishing a causal link'

## IN BRIEF

The shift-change handoff is the moment in inpatient care where patient state has to be transferred accurately across a human boundary under time pressure. Loss of safety-critical information at handoff is a documented failure mode, and the cognitive demand on the outgoing resident — synthesize, prioritize, and convey — exceeds what unaided improvisation can reliably deliver. I-PASS is a structured handoff bundle: a mnemonic (Illness severity, Patient summary, Action list, Situation awareness and contingency planning, Synthesis by receiver), paired with formal trainee education, faculty development, and a sustainability campaign. Starmer et al. (NEJM, 2014) studied its implementation at nine pediatric residency programs and reported a 23% relative reduction in medical errors and a 30% reduction in preventable adverse events, with no negative effect on resident workflow. The hedge survives verbatim: the authors state plainly that “our study design precludes definitively establishing a causal link.” Published correspondence cautions against implementing the mnemonic alone without the full bundle. The case is the structured-transfer companion to Case 177 (CIRAS) at the cultural-half-of-capability layer, and the small-tier intervention spine for state-transparency under stress.

## BACKGROUND

The shift-change handoff is a structural risk point in inpatient care. The outgoing resident — who has the most detailed mental model of each patient's state and trajectory — has to synthesize, prioritize, and convey enough of that model to the incoming team that the patient's care is not disrupted by the boundary. Loss of safety-critical information at handoff is a documented failure mode in the patient-safety literature; preventable adverse events that trace to communication breakdown are a substantial fraction of inpatient harm. The cognitive demand exceeds what unaided improvisation can reliably deliver, and the institutional tolerance for that improvisation has been falling.<sup>o</sup>

## THE INTERVENTION

I-PASS is an explicitly bundled intervention. The mnemonic — Illness severity, Patient summary, Action list, Situation awareness and contingency planning, Synthesis by receiver — structures the handoff conversation around the information classes the receiver needs in order to take over safely. The mnemonic alone is the visible piece; the bundle around it is what makes the mnemonic land. The published implementation pairs the mnemonic with formal trainee education on its use, faculty development so that supervisors model and reinforce it, and a sustainability campaign so that the discipline does not erode after the rollout. The published correspondence cautions explicitly against implementing the mnemonic alone without the full bundle.<sup>1</sup>

## HOW IT WORKED

Starmer et al. (New England Journal of Medicine, 2014) studied the I-PASS bundle's implementation at nine pediatric residency programs in the United States and Canada. The design was a pre/post evaluation across the participating sites, with the bundle deployed sequentially and outcomes tracked through chart review and observation. The headline finding was a 23% relative reduction in medical errors and a 30% reduction in preventable adverse events from the pre-implementation period to the post-implementation period, with no observed negative effect on resident workflow — the bundle did not impose net friction that displaced other clinical work.<sup>2</sup>

## THE EVIDENCE

The hedge is binding and survives into the case verbatim. The authors state in the published paper that “our study design precludes definitively establishing a causal link” — pre/post evaluation across multiple sites is the strongest practical design in this setting, and it cannot rule out secular trends in patient-safety culture, simultaneous quality-improvement work, or the sites' own selection into the trial. The 23% reduction is the strongest available evidence; it is not closed proof that the bundle alone produced the effect. Treating it as either is a misreading of what the study design supports.<sup>3</sup>

## WHAT TRANSFERRED

What the case carries for the corpus is the state-transparency-under-stress pattern at the human-human boundary (induced 3.3, LENS D2/PT5). The handoff is the moment of cognitive transfer at the team boundary; the bundle is the workflow artifact that makes the patient state legible across that boundary. Paired with Case 177 (CIRAS) at the cultural-half-of-capability layer, the case shows that the mnemonic alone is not the intervention — the institutional commitment to the bundle (trainee education, faculty development, sustainability) is what makes the mnemonic real in operation. The pattern is partly borrowed from aviation discipline, so the case also seeds the cross-domain-adaptation conversation that v1 Cases 117 (CRM) and 89 (ASRS) anchor in aviation evidence.

*Our study design precludes definitively establishing a causal link between implementation of the I-PASS Handoff Bundle and improved patient safety.*

STARMER ET AL., NEJM 2014.

## REFERENCES



1. Starmer, A. J., Spector, N. D., Srivastava, R., West, D. C., Rosenbluth, G., Allen, A. D., Noble, E. L., Tse, L. L., Dalal, A. K., Keohane, C. A., Lipsitz, S. R., Rothschild, J. M., Wien, M. F., Yoon, C. S., Zigmont, K. R., Wilson, K. M., O'Toole, J. K., Solan, L. G., Aylor, M., Bismilla, Z., Coffey, M., Mahant, S., Blankenburg, R. L., Destino, L. A., Everhart, J. L., Patel, S. J., Bale, J. F., Spackman, J. B., Stevenson, A. T., Calaman, S., Cole, F. S., Balmer, D. F., Hepps, J. H., Lopreiato, J. O., Yu, C. E., Sectish, T. C., & Landrigan, C. P. (2014). Changes in medical errors after implementation of a handoff program. *New England Journal of Medicine*, 371(19):1803–1812. doi:10.1056/NEJMs1405556 — the case's primary evaluation.
2. Colligan, L., Brick, D., & Patterson, E. S. (2015). Changes in medical errors with a handoff program. *New England Journal of Medicine*, 372(5):490 — the correspondence cautioning against implementing the mnemonic alone; authors' reply at 372(5):490–491.
3. Sectish, T. C., Starmer, A. J., Landrigan, C. P., & Spector, N. D. (2010). Establishing a multisite education and research project requires leadership, expertise, collaboration, and an important aim. *Pediatrics*, 126(4):619–622 — the I-PASS Study Group methodology and design rationale.
4. Cohen, M. D., & Hilligoss, P. B. (2010). The published literature on handoffs in hospitals: deficiencies identified in an extensive review. *Quality and Safety in Health Care — the broader hand-off-failure-mode literature*.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

I-PASS is the structured-handoff intervention the patient-safety literature anchors on — a 23% relative reduction in medical errors and a 30% reduction in preventable adverse events across nine residency programs, with no negative effect on resident workflow. The hedge that survives verbatim is the authors' own: the study design “precludes definitively establishing a causal link,” and the bundle, not the mnemonic alone, is what the evidence is about.

### LENS APPROACH

I-PASS is the state-transparency-under-stress case at the human-human boundary (induced 3.3; LENS D2/PT5) — Domain 2 for bundle-as-intervention; Domain 5 for the institutional faculty-development commitment. Pair with Case 177 (CIRAS) and v1 Cases 117 (CRM) and 89 (ASRS).

| DOMAIN             | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|--------------------|-----------------|-----------------------|--------------|---------|-------|
| CLINICAL CARE      | H K N           | D1 D2 D3 D4 D5        | 5            | 3.3     | 2 · 5 |
| PATIENT SAFETY     | →               | →                     |              |         |       |
| TEAM COMMUNICATION |                 | Iterative Development |              |         |       |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Specify the information classes the receiver needs in order to take over safely (illness severity, patient summary, action list, situation awareness, synthesis by receiver) and structure the handoff conversation around them, rather than around the outgoing operator's narrative preference.
- ▶ Treat the mnemonic, the trainee education, the faculty development, and the sustainability campaign as one bundle; the published correspondence is explicit that the mnemonic alone does not carry the result.
- ▶ Design the evaluation to track preventable adverse events and resident workflow together, so that the intervention's effect on patient harm and on the work it adds are visible in the same evidence record.

#### IN OPERATION / AFTER

- ▶ Report the 23% relative reduction together with the verbatim hedge that “our study design precludes definitively establishing a causal link”; the strongest available evidence is not closed proof.
- ▶ Carry the bundle-not-mnemonic warning into any adaptation: implementations that drop the faculty-development and sustainability components are not implementations of the intervention the published evidence is about.
- ▶ Build the cross-domain adaptation conversation deliberately — aviation CRM and ASRS evidence (v1 Cases 117, 119) is part of the I-PASS lineage and CIRAS (Case 177) is the non-aviation companion at the cultural-commitment layer.

## REFLECTION QUESTIONS

1. Identify a moment in your domain's workflow where state has to be transferred accurately across a human boundary under time pressure. What are the information classes (the analog of I-PASS's five) the receiver needs in order to take over safely, and what bundle around the structuring artifact would make it land?
2. The 23% relative reduction is the strongest available evidence and the authors are explicit it is not closed proof. What additional evidence — a cluster-randomized rollout, a residual-error analysis at the handoff point — would you require before treating the figure as closure rather than strong signal?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · incident analysis · safety dashboards

# NCSBN National Simulation Study — Licensing the 50% Substitution Rule

**G** Governance & Trust **K** Knowledge & Institutional Memory **D** Designed Out

**IMPACT** A longitudinal RCT randomized students across multiple U.S. nursing programs to control, 25%, or 50% simulation substitution for traditional clinical hours; 660 took the NCLEX with no statistically significant differences in clinical competency, nursing knowledge, or NCLEX pass rates — the number of nursing regulatory bodies permitting up to 50% simulation substitution in prelicensure RN programs rose from 1 in 2014 to 23 in 2022

## IN BRIEF

Pre-licensure nursing education had long rested on a regulatory requirement for traditional clinical hours — time spent caring for patients in a real clinical setting under the supervision of a clinical faculty member. As clinical-placement capacity came under pressure, the question facing nursing regulators was whether high-quality simulation could substitute for some fraction of those hours without degrading the clinical capability of new nurses. The National Council of State Boards of Nursing (NCSBN) ran the study that the field then built on: a longitudinal RCT randomized students across multiple U.S. nursing programs to control, 25%, or 50% simulation substitution. Hayden et al. (Journal of Nursing Regulation, 2014) reported no statistically significant differences in clinical competency, nursing knowledge, or NCLEX pass rates across the three groups; 660 participants took the NCLEX. The follow-up (2023) documents the institutional transfer: the number of nursing regulatory bodies permitting up to 50% substitution in RN programs rose from 1 in 2014 to 23 in 2022 — an unusually clean case of a single evidence base propagating across an entire regulatory field. The hedge survives: the result holds only “under conditions comparable to those described in the study” (high-quality simulation, trained faculty). Pair with Case 40 (Colorado CTSA team-science training) at the cross-domain workforce-intervention layer.

## BACKGROUND

Pre-licensure nursing education in the United States has rested on a regulatory requirement, set by state nursing boards, that students complete a defined number of traditional clinical hours — time spent caring for patients in a real clinical setting under the supervision of a clinical faculty member. As nursing programs grew and clinical-placement capacity in hospitals came under pressure, the field was asking whether high-quality simulation — manikin-based or standardized-patient scenarios run in a controlled environment — could substitute for some fraction of traditional clinical hours without degrading the clinical capability of new nurses.<sup>o</sup>

## THE INTERVENTION

The National Council of State Boards of Nursing — the umbrella body for the state nursing-regulatory boards that together control nursing licensure — commissioned the definitive study. Hayden et al. (Journal of Nursing Regulation, 2014) reported the NCSBN National Simulation Study: a longitudinal randomized controlled trial that assigned pre-licensure nursing students at multiple U.S. nursing programs to a control group, a 25%-substitution group, or a 50%-substitution group. The substitution condition replaced an equivalent fraction of traditional clinical hours with simulation hours run to specified quality standards (validated scenarios, trained simulation faculty, structured debriefing).<sup>1</sup>

## HOW IT WORKED

The headline finding was a null result on the outcome the regulators most cared about: no statistically significant differences in clinical competency, nursing knowledge, or NCLEX pass rates across the three groups. Six hundred and sixty participants took the NCLEX licensure examination, providing the comparison the regulatory decision rested on. The longitudinal design followed cohorts through their first nursing positions, with comparable outcome measures at multiple time points. The study's qualifying language is precise: the result holds only "under conditions comparable to those described in the study," and those conditions are named — high-quality simulation, trained simulation faculty, structured debriefing.<sup>2</sup>

## THE EVIDENCE

What followed is the unusual part. Nursing regulation in the U.S. is genuinely decentralized: each state board sets its own pre-licensure requirements, and there is no federal forcing mechanism. The 2023 follow-up study documents that the number of nursing regulatory bodies permitting up to 50% simulation substitution in prelicensure RN programs rose from 1 in 2014 to 23 in 2022. A single evidence base — one study, run by the regulators' own collaborative — propagated across an entire regulatory field, on the strength of the published design and the null result on the licensure examination the boards control.<sup>3</sup>

## WHAT TRANSFERRED

The case carries the cross-organization knowledge-transfer pattern at the regulatory-institution layer (induced 6.1, LENS D2/PT4). The mechanism that made the propagation work was the credibility of the study's design (longitudinal RCT, multi-site, blinded outcomes), the alignment of its outcome measure with the boards' own licensure mechanism (NCLEX), and the institutional ownership of the evidence (the regulators commissioned it and own it). Pair with Case 40 (Colorado CTSA team-science training) at the cross-domain workforce-intervention layer — both are small-tier interventions with measurable workforce-capability claims, and both depend on the legitimacy of the assessment instrument the institution then has to defend.

*The result holds under conditions comparable to those described in the study — high-quality simulation, trained faculty, structured debriefing.*

EDITORS' SYNTHESIS OF HAYDEN ET AL. (2014).

## REFERENCES

1. Jeffries, P. R. (2012). Simulation in Nursing Education: From Conceptualization to Evaluation (2nd ed.). National League for
3. Hayden, J. K., Smiley, R. A., Alexander, M., Kardong-Edgren, S., & Jeffries, P. R. (2014). The NCSBN National Simulation Study:

Nursing — the simulation–quality framework the NCSBN study’s quality conditions rest on.

2. INACSL Standards Committee (2016). INACSL Standards of Best Practice: Simulation. *Clinical Simulation in Nursing*, 12(S) — the simulation–practice standards downstream programs adopt as the analog of the NCSBN quality conditions.

A longitudinal, randomized, controlled study replacing clinical hours with simulation in prelicensure nursing education. *Journal of Nursing Regulation*, 5(2 Suppl):S1–S64. [https://www.journalofnursingregulation.com/article/s2155-8256\(15\)30062-4/fulltext](https://www.journalofnursingregulation.com/article/s2155-8256(15)30062-4/fulltext) — the case’s primary study.

4. Smiley, R. A., & Martin, B. (2023). Simulation in Nursing Education: Advancements in Regulation, 2014–2022. *Journal of Nursing Regulation*. doi:10.1016/S2155-8256(23)00086-8 — the 2023 follow-up documenting the regulatory propagation.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The NCSBN National Simulation Study is the unusual case of a single, well–designed, regulator–commissioned study propagating a substantial workforce–capability change across an entire decentralized regulatory field. The null result on NCLEX is conditional on the quality conditions the study names — high–quality simulation, trained faculty, structured debriefing — and the qualifying language has to travel with the result.

LENS APPROACH

NCSBN is the regulator–owned cross–organization knowledge–transfer case (induced 6.1; LENS D2/PT4) — Domain 2 for the RCT–design discipline; Domain 5 for the institutional–ownership move that propagated evidence across decentralized regulators. Pair with Case 40 (Colorado CTSA).

| DOMAIN              | Modes Addressed | LENS Competency         | Problem Type | Induced | LEO   |
|---------------------|-----------------|-------------------------|--------------|---------|-------|
| NURSING EDUCATION   | G K D           | D1 D2 D3 D4 D5          | 4            | 6.1     | 2 · 4 |
| CLINICAL SIMULATION | →               | → Iterative Development |              |         |       |
| REGULATION          |                 |                         |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Commission the study from inside the regulatory body that will then have to act on it; the institutional ownership of the evidence is what makes the propagation across decentralized regulators possible.
- ▶ Specify the quality conditions the result is conditional on (high–quality simulation, trained faculty, structured debriefing) at design time, so the qualifying language travels with the result and is not stripped at the citation stage.
- ▶ Align the study’s outcome measure with the licensure mechanism the boards control (NCLEX pass rate); the credibility of the propagation depends on the outcome being the one the regulators already use.

IN OPERATION / AFTER

- ▶ Carry the qualifying language (“under conditions comparable to those described in the study”) into every downstream adoption; the study’s null result is conditional on quality standards, not unconditional.
- ▶ Document the propagation as a transferable institutional pattern: a decentralized regulatory field can move on the strength of a single, well–designed, body–owned study when the outcome measure is the one the regulators already use.
- ▶ Treat the rise from 1 nursing regulatory body to 23 between 2014 and 2022 as evidence of the mechanism of regulatory transfer, not as evidence that 50% substitution is safe outside the study’s quality conditions.

## REFLECTION QUESTIONS

1. Identify a decentralized regulatory field in your domain where a substitution or adoption decision rests on whether the new modality preserves capability. What would a NCSBN-class RCT, commissioned from inside the regulatory body, look like — and what would the licensure-aligned outcome measure be?
2. The result is conditional on quality conditions (high-quality simulation, trained faculty, structured debriefing). What is the analog conditional in your context, and what mechanism would carry the qualifying language into the regulatory adoptions that follow?
3. The rise from 1 nursing regulatory body to 23 is evidence of regulatory transfer, not of safety outside the study's quality conditions. What would you do operationally to ensure the adoptions in your field actually meet the conditions the propagation rests on?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · task analysis · design prototyping

# Spaced Education RCTs in Medical Training

**H** Human–System Interface **K** Knowledge & Institutional Memory **D** Designed Out

**IMPACT** A multi-institution RCT of 537 urology residents across the U.S. and Canada randomized participants to bolus versus spaced-pattern email delivery of validated study questions; spaced education improved acquisition and retention of medical knowledge but did not raise In-Service Examination scores, and a follow-up showed the learning benefit persisting for two years at a modest effect size

## IN BRIEF

Spacing — distributing study sessions over time rather than massing them — is one of the most robust findings in basic learning-science research, with effects across age, population, and content domain. Whether the basic finding transfers into the workplace-L&D context of practicing medical trainees is a separate empirical question. Kerfoot et al. (Journal of Urology, 2007) ran the test that closed the loop. A multi-institution RCT randomized 537 urology residents across U.S. and Canadian programs to receive validated study questions in either a bolus pattern (concentrated delivery) or a spaced pattern (distributed delivery), both via email. The spaced-education condition improved acquisition and retention of medical knowledge against the bolus comparison. The 2009 follow-up (Kerfoot, Journal of Urology, 181:2671) documented that the learning benefit persisted for two years. The case is a methodologically clean small-tier intervention with replication, with standard RCT scope: the result speaks to the delivery-pattern manipulation in the email-delivered validated-question context, not to spacing in general or to other modalities. The strongest randomized strength in the small-tier batch. Pair with Case 69 (Duolingo half-life regression) at the spacing-mechanism-in-deployment layer.

## BACKGROUND

Spacing — distributing study sessions over time rather than massing them in a single concentrated block — is among the most robust findings in basic learning-science research. The spacing effect appears across age groups, populations, content domains, and outcome measures, and the broader retrieval-practice literature complements it. What is less obvious in the basic literature is whether the spacing effect transfers to the workplace-L&D context, where the learner is a practicing trainee, the content is the trainee's own clinical specialty, and the delivery mechanism is one the trainee will tolerate as part of their work week.<sup>o</sup>

## THE INTERVENTION

Kerfoot et al. (Journal of Urology, 2007) ran the test that closed the loop. The study was a multi-institution randomized controlled trial of urology residents at multiple U.S. and Canadian residency programs. Participants were randomized to receive validated study questions drawn from the American Urological Association Self-Assessment Study Program,

delivered by email, in either a bolus pattern (concentrated delivery over a short window) or a spaced pattern (distributed delivery across a longer window). Outcomes were a 32-item online test taken 1 to 14 weeks after the spaced program ended, and the November 2005 Urology In-Service Examination.<sup>1</sup>

#### HOW IT WORKED

The result was direct. The spaced-education condition improved acquisition and retention of medical knowledge against the bolus comparison. The intervention and the measurement formed a tight closed loop on the delivered questions: the spaced cohort scored higher on the staggered online tests, and its scores held steady while the bolus cohort's declined. The gain did not carry over to the In-Service Examination. The 2009 follow-up (Kerfoot, Journal of Urology, 181:2671) reported a separate analysis of the 147 residents still in training who retested and documented that the learning benefit persisted for two years — a duration that is rare in the workplace-L&D RCT literature.<sup>2</sup>

#### THE EVIDENCE

The case's scope is what makes it teachable rather than over-extended. The result speaks to the delivery-pattern manipulation (bolus vs. spaced) in the email-delivered validated-question context, with urology residents at multiple programs. It does not speak to spacing in general, to other modalities (mobile, in-clinic, simulation), or to content classes substantially different from in-service-examination knowledge. The strength of the design — multi-institution RCT, two-year persistence — is the strongest randomized strength in the small-tier batch, and the scope discipline is the qualifying language that makes the strength interpretable. The case is the methodologically clean small-tier closed-loop intervention the LENS Iterative Development domain anchors on.<sup>3</sup>

#### WHAT TRANSFERRED

Pair with Case 69 (Duolingo half-life regression) at the spacing-mechanism-in-deployment layer. Duolingo is the large-scale operational deployment of a spacing-and- retrieval-practice system in a consumer language-learning product; Kerfoot is the small-scale randomized evidence in the workplace-L&D medical-training context. Together the two cases anchor the spacing-mechanism conversation across laboratory, deployed-product, and randomized-workplace evidence. The Iterative Development discipline LENS teaches depends on having both the controlled-randomized closed loop (Kerfoot) and the operational-mechanism evidence (Duolingo) available in the corpus.

*On-line spaced education can generate improvements in learning that are retained 2 years later.*

KERFOOT, JOURNAL OF UROLOGY 2009 FOLLOW-UP.

#### REFERENCES

1. Cepeda, N. J., Pashler, H., Vul, E., Wixted, J. T., & Rohrer, D. (2006). Distributed practice in verbal recall tasks: A review and quantitative synthesis. *Psychological Bulletin*, 132(3):354–380 — the basic-literature spacing-effect review the workplace-L&D RCT sits against.
2. Kerfoot, B. P., Baker, H. E., Koch, M. O., Connelly, D., Joseph, D. B., & Ritchey, M. L. (2007). Randomized, controlled trial of spaced education to urology residents in the United States
3. Kerfoot, B. P. (2009). Learning benefits of on-line spaced education persist for 2 years. *Journal of Urology*, 181(6):2671 — the two-year persistence follow-up.
4. Settles, B., & Meeder, B. (2016). A trainable spaced repetition model for language learning. *Proceedings of ACL 2016*, 1848–1858 — Duolingo half-life regression (Case 69), the operational deployment companion.



and Canada. Journal of Urology, 177(4):1481–1487. doi:10.1016/j.juro.2006.11.074 — the case’s primary RCT.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Kerfoot et al. is the methodologically clean small-tier closed-loop spaced-education RCT in workplace medical training, with two-year persistence in the follow-up. It is the strongest randomized strength in the small-tier batch. The scope discipline is what makes it interpretable: the result speaks to the email-delivered validated-question delivery pattern, not to spacing in general.

LENS APPROACH

Kerfoot is the canonical small-tier closed-loop spacing-in-workplace-L&D case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the intervention-and-measurement-as-single-design discipline, and in Domain 4 (Test and Evaluation) for the persistence- follow-up that makes the result a workplace-L&D headline. Pair with Case 69 (Duolingo half-life regression) at the spacing-mechanism-in-deployment layer — controlled- randomized evidence with deployed-product evidence together.

| DOMAIN            | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|-------------------|-----------------|-----------------------|--------------|---------|-------|
| MEDICAL EDUCATION | H K D           | D1 D2 D3 D4 D5        | 4            | 2.3     | 2 · 4 |
| RETENTION         |                 | Iterative Development |              |         |       |
| WORKFORCE L&D     |                 |                       |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Specify the closed loop in advance: the intervention manipulation, the outcome measure, and the validated-question pool together, so the intervention and the measurement form a single design rather than two separately published artifacts.
- ▶ Choose the delivery modality (email, in this case) that the trainees will tolerate as part of their work week; the workplace-L&D context is unforgiving of modalities that compete with clinical work.
- ▶ Pre-register the persistence follow-up at the design stage; the two-year retention figure in Kerfoot 2009 is what gives the intervention its workplace-L&D strength, and follow-up has to be designed into the original RCT.

IN OPERATION / AFTER

- ▶ Report the result in the scope it speaks to: email-delivered validated-question delivery pattern with urology residents; the spacing-in-general claim sits in the basic literature, and the case’s contribution is the workplace-L&D closed-loop evidence.
- ▶ Treat the two-year persistence figure as the workplace-L&D headline; persistence over years is rare in the workplace-L&D RCT literature, and the figure is what makes the result actionable as a training design.
- ▶ Pair the result with Duolingo (Case 69) at the spacing-mechanism layer in any communication of the broader conversation, so the controlled-randomized evidence and the deployed-product evidence are read together.

## REFLECTION QUESTIONS

1. Identify a workplace-L&D context in your domain where the basic-literature evidence (spacing, retrieval practice, deliberate practice) is robust and the workplace-transfer evidence is thinner. What would a Kerfoot-class closed-loop RCT look like — the intervention manipulation, the outcome measure, the validated-content pool together?
2. The two-year persistence figure is the workplace-L&D headline. What follow-up cadence would you design into your study at the start, so that persistence over years is part of the original RCT rather than a separately commissioned downstream check?
3. The result's scope is the email-delivered validated-question delivery pattern. What is the analog scope discipline in your context — the qualifying language that has to travel with the headline so the result is read as evidence about a specific mechanism rather than as evidence about a general principle?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · task analysis · design prototyping

## CASE 18

GLOBAL HEALTH

HIV CARE

TRAINING-MODALITY COMPARISON

SUB-SAHARAN AFRICA

2023

# PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption

**K** Knowledge & Institutional Memory**N** Normalization of Deviance**D** Designed Out

**IMPACT** Across 16 PEPFAR-supported Sub-Saharan African countries, pre/post knowledge and self-confidence assessments showed mean increases across all training modalities — in-person (pre-COVID), virtual, and blended — when pandemic disruption forced a delivery-mode switch, with the in-person gains significantly larger and the online formats reaching more women; the L1–L2 limitation is explicit and the L3/L4 question remains open

## IN BRIEF

Across 16 PEPFAR-supported Sub-Saharan African countries, a 2023 real-world program evaluation compared in-person (pre-COVID), virtual, and blended HIV-care training delivery. Across all assessed knowledge domains and self-perceived confidence dimensions, pre/post assessments showed mean increases regardless of modality. The honest framing the case carries into print is that the outcome metric is knowledge and self-rated confidence — Kirkpatrick L1–L2 territory (Case 79) — not L3 on-job behavior change or L4 patient outcomes. Even so, the study is one of the few real-world cross-country modality comparisons at meaningful scale, and it is the L&D evaluation pattern playing out in global health at multi-country scope. The case is cross-listed with the corporate / workforce L&D cluster (Cases 79, 65, 83, 70) and with the non-US/UK/EU geographic-coverage gap (Cases 93, and the cases in the later supplemental batches). The study is journal-published (the medRxiv item is its preprint), so the open question is not review but reach: whether L1–L2 knowledge gains translate to L3 behavior change or L4 patient outcomes is untested.

## BACKGROUND

PEPFAR (the U.S. President’s Emergency Plan for AIDS Relief) has been one of the largest sustained global health programs of the past two decades. Workforce capability — clinicians and frontline workers trained on current HIV diagnosis, treatment, and prevention protocols — is one of its core deliverables. COVID disruption halted in-person training in March 2020 and forced a switch to virtual and blended delivery by May 2021, before any of the evidence base on modality equivalence had been built for this workforce.<sup>9</sup>

## THE INTERVENTION

The 2023 study compares pre/post-assessment outcomes across 16 PEPFAR-supported Sub-Saharan African countries and across the three modalities. The headline finding is that across all knowledge domains and self-perceived confidence dimensions assessed, mean increases were observed regardless of modality. But the modalities were not equivalent: the in-person format produced significantly greater knowledge gains (13.6% increase vs. 6.0% for the virtual workshop and 7.6% for the online course) and greater confidence gains, regardless of training level or professional cadre. The countervailing finding is equity of access — the virtual and online formats reached a significantly higher proportion of women — so the policy question under future disruption is a trade-off between L1–L2 learning gains and reach, not a claim of modality equivalence.<sup>1</sup>

## HOW IT WORKED

The honest framing the case must carry into print is the Kirkpatrick limitation. The study's outcomes are knowledge and self-rated confidence — Level 1 (reaction) and Level 2 (learning, which covers confidence) in the Kirkpatrick framework (Case 79). It is not Level 3 (behavior change on the job — whether the clinicians actually changed practice) and it is not Level 4 (results — whether patient outcomes improved). The structural limitation Case 79 names is exactly the limitation this study sits inside.<sup>2</sup>

## THE EVIDENCE

Even so, the study is one of the few real-world cross-country modality comparisons at meaningful scale, and the multi-country scope means the comparison is not trivially confounded by single-country labor-market or health-system factors. It is the L&D evaluation pattern playing out in global health, and the modality-comparison finding is policy-relevant for PEPFAR program design across the deployment region.<sup>3</sup>

## WHAT TRANSFERRED

The sourcing is settled: the medRxiv version is the preprint and the PMC version is the journal-published article, and the case cites the published one. What is not settled is whether the measured gains reach behavior or outcomes. The LENS teaching point is the cross-listed pair with the corporate / workforce L&D cluster (Cases 79, 65, 83, 70) and the non-US/UK/EU cluster (Cases 93). The subobjective **Judgment under inadequate evidence** (4.4) is exercised: the study is the strongest evidence the field has on modality choice in this context, it found the in-person gains larger, and it does not establish L3 / L4 outcomes. Future validation is ongoing.

*The outcomes are knowledge and self-rated confidence. The L3 and L4 questions — did practice change, did patient outcomes improve — remain open.*

EDITORS' SYNTHESIS OF THE PEPFAR 16-COUNTRY MODALITY COMPARISON.

## REFERENCES

1. Kiguli-Malwadde, E., Forster, M., Eliaz, A., et al. (2023). "Comparing in-person, blended and virtual training interventions; a real-world evaluation of HIV capacity building programs in 16 countries in sub-Saharan Africa," PLOS Global Public Health 3(7):e0001654, doi:10.1371/journal.pgph.0001654 — the published version of medRxiv 2023.02.08.23285641.
2. Kirkpatrick & Kirkpatrick (2006), Evaluating Training Programs — the chain-of-evidence framework the L1–L2 limitation references (paired Case 79).
3. Blume et al. (2010), Journal of Management 36(4):1065–1105 — the meta-analytic transfer finding the L3 question references (paired Case 65).

- 2. The evaluated program is STRIPE HIV, led by the African Forum for Research and Education in Health (AFREhealth) and the University of California, San Francisco, and funded by HRSA through PEPFAR — named in the study’s program description and funding statement.

THE LEARNING ENGINEERING LENS

LE INSIGHT

PEPFAR’s 16-country modality comparison is the L&D evaluation pattern in global health: mean L1–L2 gains across all modalities, with the Kirkpatrick limitation (Case 79) explicit. The study is journal-published; the L3 and L4 questions are what remain open.

LENS APPROACH

PEPFAR is the global-health workforce-capability case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the modality-design iteration COVID forced and in Domain 5 (Navigating Sociotechnical Constraints) for the multi-country program scope. The subobjective **Judgment under inadequate evidence** (4.4) is exercised because the study is the strongest available at L1–L2 and does not settle L3 / L4. Pairs with Case 93 (SkillsFuture) as the workforce L&D cross-listing.

| DOMAIN                       | MODES ADDRESSED | LENS COMPETENCY         | PROBLEM TYPE | INDUCED | LEO       |
|------------------------------|-----------------|-------------------------|--------------|---------|-----------|
| GLOBAL HEALTH                | K N D           | D1 D2 D3 D4             | 4            | 2.3     | 2 · 4 · 5 |
| HIV CARE                     |                 | D5                      |              |         |           |
| TRAINING-MODALITY COMPARISON | →               | → Iterative Development |              |         |           |
| SUB-SAHARAN AFRICA           |                 |                         |              |         |           |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the modality comparison around what the assessment instrument can actually establish — knowledge and confidence at L1–L2 are tractable; L3 and L4 require longitudinal data and patient-outcome linkage the study did not have.
- ▶ Build the multi-country sampling so the comparison is not trivially confounded by single-country factors — the 16-country scope is part of the study’s evidential weight.
- ▶ Carry the Kirkpatrick limitation into the program documentation honestly — Case 79’s stop-at-L2 pattern is exactly the limitation this study sits inside, and naming it preserves the case’s teaching value.

IN OPERATION / AFTER

- ▶ Pair with Case 93 (SkillsFuture) as the workforce-capability counterparts at national- and multi-country scale; together they teach what L&D measurement looks like in non-US/UK/EU settings with the evidence-tier honesty intact.
- ▶ Use the subobjective **Judgment under inadequate evidence** (4.4): the study is the strongest evidence the field has on modality choice in this context, and the practitioner has to trade larger in-person gains against wider online reach on L1–L2 evidence while L3/ L4 evidence develops.
- ▶ Carry the measurement limit honestly: the study is journal-published, and what it does not establish is L3 behavior change or L4 patient outcomes; confirmatory replication is the outstanding step.

## REFLECTION QUESTIONS

1. Identify a training program in your context that was forced to switch modality under disruption (COVID, budget constraint, geographic dispersion). What evidence would you have needed at the time to make the modality decision well, and at which Kirkpatrick level was the available evidence?
2. Specify the L3 / L4 measurement architecture you would build to extend the PEPFAR study into behavior-change and patient-outcome evidence — what data sources, what cadence, what linkage between training cohort and clinical outcomes.
3. The case is journal-published and stops at L1–L2. What additional confirmatory evidence — replication in different country sets, L3 / L4 follow-up studies, peer-reviewed meta-analysis — would you require before treating the modality trade-off as settled for program-design decisions?

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards · task analysis · design prototyping

# Keystone ICU / Pronovost Checklist

**T** Training Gap    **N** Normalization of Deviance

**DISCLOSURE** Institutional overlap: an editor shares an institution (Johns Hopkins) with the study's lead author; no editor was personally involved in the Keystone ICU work. Included on the published peer-reviewed evidence (NEJM 2006) and its sustained, independently replicated results.

**IMPACT** Central-line-associated bloodstream infections (CLABSI) reduced to near zero across 103 Michigan ICUs; ~1,500 lives saved in 18 months; ~\$175M saved; sustained at ten years

## IN BRIEF

Peter Pronovost's central-line checklist has five items — wash hands, use full barrier precautions, clean the skin with chlorhexidine, avoid the femoral insertion site, and remove unnecessary catheters — and not one of them was unknown to the physicians it governed. The question was never what to do, but whether it would be done every time. The Keystone project, launched across 103 Michigan ICUs in 2004, paired the checklist with a cultural change: nurses were empowered to stop the procedure, in non-emergency situations, whenever a step was skipped. That authorization is the element the case treats as decisive — though it rode inside a multi-component bundle (clinician education, a dedicated line cart, daily catheter-review goals, and monthly infection-rate feedback) that the trial tested as a whole and could not decompose. Central-line-associated bloodstream infections fell to near zero, an estimated 1,500 lives and \$175 million were saved in eighteen months, and the effect was sustained at ten years.

## BACKGROUND

Central-line-associated bloodstream infections were a common, often fatal complication of intensive care, and the steps to prevent them were well established and uncontroversial. The problem was reliability: in the existing culture, a nurse who saw a physician skip a sterile step had no procedural path to intervene without crossing the hospital's authority gradient. The knowledge was universal and the steps cheap; what was missing was any mechanism that made the right action happen every time rather than most of the time, which is precisely where a fatal infection finds its opening.<sup>o</sup>

## THE INTERVENTION

In 2004, Peter Pronovost's team launched the Keystone ICU project across 103 Michigan units. It combined a simple five-item central-line checklist — hand washing, full-barrier precautions at insertion, chlorhexidine skin prep, avoiding the femoral site where possible, and removing unnecessary catheters — with an explicit authorization: nurses were empowered to stop any non-emergency procedure in which a step was missed. The authorization was deliberate — a nurse who would otherwise have had to challenge a senior physician on

her own account now acted on a standing institutional warrant, which is what lowered the personal risk of intervening.<sup>1</sup>

#### HOW IT WORKED


The checklist was the technical half; the nurses' enforcement authority was the cultural half. The trial could not weigh the two, and the program's own ex post theory credits the networked collaborative and the monthly infection data alongside the stop. Before Keystone, the path to intervene did not exist; after it, the path was institutional and expected. The pairing converted knowledge everyone already had into behavior that happened every time. The checklist gave the nurse an objective, impersonal basis for the stop — a missed item, not a judgment about the physician — which is what made the authority usable in practice rather than merely on paper.<sup>2</sup>

#### THE EVIDENCE

Across the Michigan ICUs, the median quarterly CLABSI rate fell to zero, and the state's units outperformed 90% of ICUs nationwide. The program was estimated to have saved roughly 1,500 lives and \$175 million within eighteen months. Results were published in the New England Journal of Medicine in 2006, and follow-up showed the effect sustained at ten years. The durability mattered as much as the size of the drop: an improvement that survives a decade is evidence the change was built into the system's structure rather than riding on the initial enthusiasm of a single project.<sup>3</sup>

#### WHAT TRANSFERRED

What Keystone established was not which component worked but that the bundle worked and traveled. The model was packaged as the AHRQ CUSP toolkit, adopted in more than forty states, and replicated internationally — the checklist, the nurse-stop authorization, the line cart, the daily goals sheet, and the monthly infection feedback moving together, because that is the unit the evidence covers. Packaging the approach as a reusable toolkit was itself part of what transferred — it turned a single project's success into something other institutions could adopt without rediscovering the load-bearing role of the authority change.<sup>4</sup>



*The checklist gave the nurse an impersonal reason to stop: a missed item rather than a judgment about the physician.*

EDITORS' SYNTHESIS OF PRONOVOST & VOHR, SAFE PATIENTS, SMART HOSPITALS, 2010



## REFERENCES

1. Pronovost & Vohr (2010), Safe Patients, Smart Hospitals — the checklist-plus-nurse-authority pairing (paraphrased).
2. Pronovost, P. et al. (2006), "An Intervention to Decrease Catheter-Related Bloodstream Infections in the ICU," NEJM 355 — the trial and the near-zero result.
3. Bosk, C. et al. (2009), "Reality check for checklists," The Lancet — the critique: the Keystone result cannot be attributed cleanly to the checklist, and the bundle's active ingredient is not established.
4. Pronovost, P. et al. (2016), "Sustaining Reductions in CLABSI in Michigan ICUs: A 10-Year Analysis," Am J Med Qual 31(3):197–202 — the effect sustained through 2013.
5. Agency for Healthcare Research and Quality, CUSP toolkit — dissemination across states.

## THE LEARNING ENGINEERING LENS

## LE INSIGHT

Keystone is the strongest healthcare evidence that a protocol and the authority to enforce it should be designed and shipped together — and the clearest demonstration of what a bundled trial cannot tell you afterwards. The result is large, durable, and widely replicated; it does not decompose. Treating the cultural half as engineering rather than aspiration is the commitment the case supports. Naming which half carried the effect is not.

## LENS APPROACH

LENS uses Keystone in LEN 4 and LEN 10 as the canonical worked example of measurement linked to intervention. Studio projects require students to specify both halves of the pair, name the authority that authorizes the cultural half, and identify the measurement signal that will confirm or refute the effect. The course treats "is the cultural change actually authorized?" as falsifiable, not stated.

| DOMAIN                     | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|-----------------|------------------|--------------|---------|-----|
| HEALTHCARE                 | → T N           | → D1 D2 D3 D4 D5 | 3            | 4.1     | 3   |
| Human-System Collaboration |                 |                  |              |         |     |

## APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- Pair the technical artifact with the authority change from the outset — a checklist plus the explicit, institution-backed right of a junior member to stop the procedure when a step is missed.
- Make the intervention an obligation rather than a permission, so the corrective action does not depend on an individual's willingness to challenge a senior colleague.
- Anchor the stop to an objective, impersonal trigger (a missed checklist item) so the authority is usable without it reading as a judgment of the more senior operator.

## IN OPERATION / AFTER

- Measure the outcome directly (the CLABSI rate) and publish it, so the cultural half can be confirmed as authorized in practice rather than merely declared.
- Track the effect over years, not months, to confirm the change is built into the system's structure rather than riding on a project's initial enthusiasm.
- Package the paired design as a reusable toolkit so other institutions can adopt it without rediscovering the load-bearing role of the authority change.

## REFLECTION QUESTIONS

1. Identify an evidence-based protocol in your domain that is **known** to work but is not consistently performed. What is the authority change required to pair with it?
2. Design the measurement signal that would confirm the cultural half of a Keystone-style intervention is actually being authorized — not merely declared.
3. Keystone made the nurse's stop an obligation the institution stood behind, not a personal risk. Identify a corrective action in your domain that currently costs the person who takes it, and design the institutional backing that would remove that cost.

**WHO BUILDS THIS** learning science & instructional design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards

# TREWS / Sepsis Watch

**H** Human–System Interface   **K** Knowledge & Institutional Memory   **G** Governance & Trust

**DISCLOSURE** Institutional overlap: an editor shares an institution (Johns Hopkins) with authors of this work and its deployment sites (five JHU hospitals); no editor was personally involved. Included on the published peer-reviewed evidence (Nature Medicine 2022).

**IMPACT** Prospective multi-site evidence of reduced mortality, organ failure, and length of stay when clinicians engaged with ML sepsis alerts in deployed care

## IN BRIEF

The Targeted Real-time Early Warning System (TREWS) is a machine-learning sepsis-detection tool deployed at five Johns Hopkins hospitals; Duke's Sepsis Watch follows the same pattern. The Adams et al. prospective multi-site study (Nature Medicine 2022) reported reduced in-hospital mortality, reduced organ failure, and shorter length of stay associated with deployment; the companion Henry et al. paper reported the antibiotic-timing effect — conditional on clinicians acting on alerts within a defined window. The benefit was not the model in isolation. It was the model plus a deliberately engineered alert, workflow, and clinician–confirmation layer that made the alert actionable at the bedside. The honest hedge in the literature is that the evidence is prospective and observational, not RCT; the field notes RCTs are still pending. The case is the positive counter to the Epic Sepsis Model (Case 5): same delegation task, opposite outcome, and the difference is in the engineering of the human–machine boundary and the discipline of the surrounding evidence work.

## BACKGROUND

Sepsis is among the most consequential time-dependent diagnoses in hospital medicine: every hour of delayed antibiotics is associated with increased mortality, and the disease is heterogeneous enough that bedside clinicians frequently miss the earliest signal in a patient already being treated for something else. The promise of machine learning has been to surface that earliest signal from the continuously updated EHR trace — vitals, labs, medications — and route it to a clinician who can act in time.<sup>0</sup>

## THE INTERVENTION

TREWS, deployed across five Johns Hopkins hospitals, and Sepsis Watch at Duke are the two best-documented instances of this approach. The Adams et al. prospective study of 590,000 patient encounters reported that when clinicians confirmed an alert within three hours, in-hospital mortality, organ failure, and length of stay were lower than for patients whose alert was not confirmed within that window, after adjustment for presentation and severity. The evidence is observational rather than randomized, but it is multi-site, pre-registered, and outcome-grounded.<sup>1</sup>

## HOW IT WORKED

The capability the deployment supplied was not “the model.” It was the alert designed to fit into a specific clinical workflow, the confirmation step that put a clinician between the model and the action, and the institutional commitment to measure outcomes — not adoption — as the metric of success. The reported benefit collapsed when clinicians did not engage with the alert: the model on its own did nothing. The deliverable was the interface, the role design, and the surrounding evidence loop, not the prediction.<sup>2</sup>

## THE EVIDENCE

The honest hedge survives into the literature. The Adams et al. paper, and the broader sepsis-AI field, explicitly note that the outcome inference is conditional on the population, the workflow, and the engagement pattern measured at these sites — and that randomized trials remain pending. The benefit is real on the evidence presented, but it is not a closed proof; it is the strongest available evidence that delegation of early detection to ML can be done as a paired intervention with measurable outcome improvement.<sup>3</sup>

## WHAT TRANSFERRED

What TREWS teaches is that the failure pattern of clinical AI (Case 35) is not inevitable. When the model is treated as one component of a deliberately designed human-machine system — with an alert that fits the workflow, a clinician role that retains authority, a deployment that is observed prospectively, and a willingness to report null effects in non-engaged subgroups — the delegation can produce capability rather than alert fatigue. The case is the engineering counter to Watson for Oncology (the model marketed ahead of its capability) and the Epic Sepsis Model (the model deployed ahead of its validation).

*The benefit is not the model. It is the model used as part of a system clinicians can act on.*

EDITORS' SYNTHESIS OF ADAMS ET AL. (2022) AND HENRY ET AL. (2022).

## REFERENCES

1. Sendak et al. (2020), “Real-world integration of a sepsis deep learning technology into routine clinical care: implementation study,” *JMIR Medical Informatics* 8(7):e15182 (Sepsis Watch implementation).
2. Adams et al. (2022), “Prospective, multi-site study of patient outcomes after implementation of the TREWS machine learning-based early warning system for sepsis,” *Nature Medicine* 28(7):1455–1460, doi:10.1038/s41591-022-01894-0.
3. Henry et al. (2022), “Factors driving provider adoption of the TREWS machine learning-based early warning system and its effects on sepsis treatment timing,” *Nature Medicine* 28(7):1447–1454, doi:10.1038/s41591-022-01895-z.
4. Wong et al. (2021), “External Validation of a Widely Implemented Proprietary Sepsis Prediction Model,” *JAMA Internal Medicine* 181(8):1065–1070 — the foil case (102).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

TREWS is the strongest current evidence that delegating early sepsis detection to a machine-learning system can improve patient outcomes — when the delegation is engineered as a paired intervention. The capability deliverable is the alert design, the clinician role, and the outcome-grounded evidence loop, not the model.

### LENS APPROACH

TREWS is the positive Domain 3 / Problem Type 6 case the corpus needed: a documented delegation to AI that worked, with the explanation locatable in the human-machine interface and the governance discipline rather than the model. LENS uses it in Domain 3 (Human-System Collaboration) for the delegation-with-

revocation pattern, in Domain 4 (Test and Evaluation) for outcome-grounded prospective evidence under the judgment-under-inadequate-evidence LEO, and in Domain 5 (Navigating Sociotechnical Constraints) for the workflow-fit discipline. It is the foil drafted directly against the Epic Sepsis Model (Case 5).

| DOMAIN                    | MODES ADDRESSED    | LENS COMPETENCY                                                                                        | PROBLEM TYPE | INDUCED | LEO   |
|---------------------------|--------------------|--------------------------------------------------------------------------------------------------------|--------------|---------|-------|
| HEALTHCARE<br>CLINICAL AI | → <div>HKG</div> → | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div><br>Human-System<br>Collaboration | 6            | 3.1     | 4 · 3 |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the alert to fit a specific workflow, not the average workflow — including the bedside action it should provoke and the timeline within which it must be confirmed.
- ▶ Keep a clinician between the model and the action: model flags, human confirms, the model’s authority is to surface, not to decide.
- ▶ Instrument the deployment for outcomes (mortality, organ failure, antibiotic timeliness) before the first alert fires — adoption is not the metric of success.

IN OPERATION / AFTER

- ▶ Report engagement-stratified outcomes honestly — including the subgroups where the alert was not acted on and the benefit was not observed.
- ▶ Treat the prospective/observational design as a constraint to be replaced by RCT evidence when feasible, not as a result that ends the evidence work.
- ▶ Carry the model’s hedges into the deployment documentation so the next site adopts the model and the discipline that produced it.

REFLECTION QUESTIONS

1. Identify a delegation of detection or screening in your domain that succeeded. What were the components of the human-machine interface that made the model actionable, and what would happen to the outcome metric if those components were removed?
2. Specify the engagement-stratified outcome you would report from a deployment at your site — including the subgroup where the alert was not acted on. What would you need to instrument before the first alert fires?
3. The TREWS evidence is prospective and observational, not RCT. What is the minimum additional evidence you would require before recommending the same model deployment at a new site that differs from the validation sites in population, workflow, or EHR configuration?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

# Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff

**H** Human-System Interface **K** Knowledge & Institutional Memory **N** Normalization of Deviance

**IMPACT** A clinical adaptation of the aviation 'sterile cockpit' rule — no non-essential communication or interruptions during high-workload phases — applied to surgical ward rounds; across 71 control and 70 intervention rounds the published evaluation reports documentation accuracy rising from 63.6% to 77.9%, higher patient satisfaction, and fewer interruptions, at no additional time cost, on a single-ward single-study evidence base whose generalization the authors are explicit about

## IN BRIEF

The “sterile cockpit” rule in commercial aviation (FAR 121.542, codified 1981) prohibits non-essential activity among flight crew during taxi, takeoff, landing, and any other flight operation below 10,000 feet except cruise — the high-workload, high-consequence windows when distraction is most likely to produce error. The principle has been adapted across other high-consequence domains; recent work in clinical care has extended the principle to ward rounds, where interruption-driven information loss is a documented contributor to patient-safety events. Treloar et al. (2025), in *World Journal of Surgery*, describe a structured intervention on a single surgical ward: roles allocated before each patient (leader, assessor, scribe), no interruptions except medical emergencies, one speaker at a time, nursing and allied health invited in, and a checkback in which the scribe reads the documented plan back. Across 71 control and 70 intervention rounds, documentation accuracy rose from 63.6% to 77.9% and patient satisfaction rose, with interruptions and parallel conversations falling. The hedges that survive into the case verbatim: this is a single-unit single-study evidence base, the documentation and satisfaction effects are direct while any inference to adverse events rests on the established link between round quality and downstream harm rather than on direct measurement in this study, and the sustainability of the protected-window discipline beyond the study period is not yet established. The case is the cross-domain adaptation case: an aviation safety-culture rule, adapted with explicit attribution and adapted again for a clinical context, with the evidence at the tier the adaptation has reached.

## BACKGROUND

The aviation sterile-cockpit rule (FAR 121.542) was codified in 1981 in response to a documented pattern of accidents in which flight-crew distraction during high-workload phases — non-operational conversation, cabin coordination, administrative communication — contributed to the error chain. The rule is operationally simple: below 10,000 feet, in taxi, takeoff, approach, and landing, only operationally necessary communication is permitted

in the cockpit, cruise flight excepted. The principle has been extended in safety-culture practice to other high-workload phases and operations.<sup>0</sup>

#### THE INTERVENTION

Clinical care has long carried an analogous problem at the handoff and ward-round boundary. Interruption-driven information loss during handoff is a documented contributor to patient-safety events; I-PASS and SIGNOUT<sup>?</sup> and similar structured-handoff interventions address the information content, but the workflow context — the interruptions that fragment the handoff — has historically been treated as ambient noise rather than as a design variable. The structural argument the clinical adaptation makes is that the workflow context is a design variable, and that the sterile-cockpit principle provides a worked template for engineering it.<sup>1</sup>

#### HOW IT WORKED

Treloar et al. (2025) describe a structured intervention on a surgical ward. Rounds were audio-visually recorded, 71 conducted normally and 70 under the sterile-cockpit rules: roles allocated before each patient, no interruptions except medical emergencies, no parallel conversations, nursing and allied health invited to contribute, and a closing checkback in which the scribe repeats the documented plan verbatim. The primary outcomes were accuracy of documentation, scored by blinded senior surgeons against the recorded speech, and patient satisfaction.<sup>2</sup>

#### THE EVIDENCE

The sterile-cockpit rounds documented case notes more accurately (77.9% vs. 63.6%), drew higher patient satisfaction, brought nurses to the bedside more often (68% vs. 45%), and cut interruptions and parallel conversations — with no additional time taken. The structural learning the case carries is the adaptation discipline: an aviation safety-culture rule, with its specific operational scoping (below 10,000 feet, defined permitted communication), was carried into a clinical context with the operational scoping translated (a defined window, defined permitted interruption classes, defined triage responsibility) rather than the principle imported without translation. Pair the case with I-PASS (a structured-handoff intervention; Case 15 when drafted) at the handoff-as-capability layer, and with CRM (Case 117) at the aviation-safety-culture-to-clinical-care translation layer.<sup>3</sup>

#### WHAT TRANSFERRED

The hedges have to survive verbatim. This is a single-unit single-study evidence base, and the authors name the Hawthorne effect and the impossibility of randomising ward rounds as limits on it; the case should not be generalized beyond a strong structural argument until multi-site replication is in the literature. The documentation-accuracy and satisfaction effects are direct; the inference that better-documented rounds reduce downstream patient-safety events depends on the established link in the round- and handoff- quality literature rather than on adverse-event measurement in this study. The sustainability of the protected-window discipline beyond the study period is not yet established — the operational cost of maintaining the workflow design over the long term is the open question. The case is teachable at the adaptation-discipline level today; the multi-site multi-cycle evidence base is the next deliverable.

*The scoping is the adaptation. Sterile cockpit is “below 10,000 feet”; the clinical analog is a defined window with defined permitted interruption classes.*

EDITORS’ SYNTHESIS OF PMC12515027 AND THE STERILE-COCKPIT ADAPTATION LITERATURE.

REFERENCES

1. Federal Aviation Administration, 14 CFR § 121.542 (codified 1981) — origin of the aviation sterile-cockpit rule.

2. Starmer, A. J. et al. (2014), “Changes in medical errors after implementation of a handoff program,” New England Journal of Medicine 371(19):1803–1812 — I-PASS handoff intervention; structural cousin in the structured-information half of handoff capability.

3. Treloar, E., Herath, M., Altree, M., Potter, S., Penhall, M., Walsh, D., Kennedy, L., Bruening, M., Edwards, S., Ey, J. D., Bradshaw, E. L., & Maddern, G. J. (2025), “A Simple Solution for a Complex Problem: The ‘Sterile Cockpit’ to Improve Ward Rounds,” World Journal of Surgery 49(10):2769–2776, doi:10.1002/wjs.70074, PMID:40930848, PMCID:PMC12515027 — the cited adaptation study.

4. Broom, M. A., Capek, A. L., Carachi, P., Akeroyd, M. A., & Hilditch, G. (2011), “Critical phase distractions in anaesthesia and the sterile cockpit concept,” Anaesthesia 66(3):175–179 — prior anaesthesia-domain sterile-cockpit adaptation establishing the cross-domain transfer pattern.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The sterile-cockpit ward-rounds case is the worked example of cross-domain adaptation discipline at small scale: an aviation safety-culture rule, carried into clinical care with its operational scoping translated and its cultural half preserved. The single-unit evidence is direct on documentation accuracy and patient satisfaction; the adverse-event inference rests on the established round-quality link, and the multi-site replication is the explicit next deliverable.

LENS APPROACH

Sterile-cockpit ward rounds is the cross-domain adaptation case (induced 3.2; LENS D3/PT5) — Domain 3 for workflow- context-as-design-variable; Domain 5 for the cultural- commitment half. Pair with I-PASS (Case 15), CRM (Case 117), and Case 27.

| DOMAIN                                        | MODES ADDRESSED | LENS COMPETENCY                                 | PROBLEM TYPE | INDUCED | LEO   |
|-----------------------------------------------|-----------------|-------------------------------------------------|--------------|---------|-------|
| HEALTHCARE<br>PATIENT SAFETY<br>HUMAN FACTORS | H K N           | D1 D2 D3 D4 D5<br>Human-System<br>Collaboration | 5            | 3.2     | 4 · 5 |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Adapt the principle with its operational scoping translated, not imported. Sterile cockpit is “below 10,000 feet, defined permitted communication”; the clinical analog is a defined window, defined permitted interruption classes, defined triage responsibility — the scoping is the adaptation.
- Design the triage workflow that handles emergent interruptions during the protected window; the rule cannot be “no interruptions” without a path for true emergencies, and the design of that path is part of the deliverable.
- Build the cultural-commitment half alongside the workflow half; the rule depends on the unit’s willingness to honor it under workload pressure, and

IN OPERATION / AFTER

- Carry the single-unit single-study hedge into communication; the case is teachable at the adaptation-discipline level today, with the multi-site replication as the explicit next deliverable.
- Distinguish the direct measure (interruption-frequency and information-transfer-quality reductions during the protected window) from the inferred outcome (downstream patient-safety events), and report them at their respective evidence tiers.
- Plan the long-run sustainability measurement: does the protected-window discipline survive six months, twelve months, leadership turnover? The operational-cost half of the intervention is what the multi-cycle evidence base will eventually answer.



the cultural commitment is what makes the rule operative.

### REFLECTION QUESTIONS

1. Identify a high-workload, high-consequence window in your domain that is currently treated as ambient workflow rather than as a protected period. What would the operational scoping of a sterile-window adaptation look like — permitted communication classes, triage responsibility, signaling?
2. Specify the cultural-commitment half of the rule. The workflow design is necessary but not sufficient; honor under workload pressure is what makes the rule operative, and that depends on leadership and unit culture more than on policy.
3. The case's evidence is single-unit single-study. Design the multi-site replication you would want to see before treating the adaptation as established, and the long-run sustainability measurement you would use to know whether the protected-window discipline survives leadership turnover.

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · incident analysis · safety dashboards

# ChatGPT in Healthcare — Hallucination Cases

**H** Human–System Interface   **D** Designed Out

**IMPACT** Evaluations of LLMs on medical tasks document fabricated citations — 47% of generated references invented — and drug–information answers judged safe and complete in 13 of 50 cases

## IN BRIEF

Since ChatGPT’s public release in late 2022, the clinical and peer-reviewed literatures have documented a recurring pattern: clinicians use large language models to draft patient education, summaries, or treatment guidance, and the output contains fabricated citations, hallucinated drug dosages, or fictitious clinical guidelines. The failures range from cosmetic — invented references in academic submissions — to potentially clinical, such as unsafe dosing or fabricated contraindications. The capability gap is at the human–verification interface: the model presents hallucinated content with the same fluent confidence as accurate content, and early reports suggest clinicians accept LLM output less critically than a colleague’s. The case is the live, foundational case for LLM integration into clinical workflow — the discipline must specify what verification looks like at the moment of use, while deployment is already happening.

## THE SHIFT

Large language models became broadly available with ChatGPT’s release in late 2022, and clinicians began using them almost immediately — to draft patient–education materials, summarize records, and look up guidance. For the first time, a tool that produces fluent, authoritative–sounding medical text is in routine, informal use at the point of care — adopted ahead of any guideline, credential, or institutional sign–off, so the practice spread faster than any structure to govern it could be put in place.<sup>9</sup>

## WHAT IS EMERGING

A recurring failure pattern has been documented across the clinical and peer-reviewed literatures: LLM output containing fabricated citations, hallucinated drug dosages, or fictitious clinical guidelines. The failures range from cosmetic — invented references in academic submissions — to potentially clinical, such as unsafe medication doses or invented contraindications — a span that matters because the same tool, used the same way, can produce a harmless error and a dangerous one with no change in how confident it sounds.<sup>1</sup>

## THE CAPABILITY QUESTION

The capability gap is at the human–verification interface. The model presents hallucinated content with exactly the same fluent confidence as accurate content; the interface does not distinguish the two. The question is whether clinicians can — and will — develop the routine verification practice that the tool’s fluency actively discourages, since the very smoothness

that makes the output easy to accept is what removes the friction a reader would normally use as a warning.<sup>2</sup>

EARLY EVIDENCE

The WHO’s 2024 guidance on large multi-modal models names the mechanism directly: the models “can also encourage ‘automation bias’ by health care professionals and patients, whereby errors are overlooked that would otherwise have been identified.” The gap is not in the documentation of errors but in practice: no established routine for verifying model output stands between the model and the clinician at the moment of use — the concern is not that the model errs but that the practice for catching its errors at the moment of use has not yet been defined or taught. Through 2025 the pattern held: pharmacists comparing ChatGPT against a drug information center measured an overall accuracy rate of 50%, even as clinical adoption ran well ahead of any settled regulatory framework.<sup>3</sup>

OPEN PROBLEMS

This is the live frontier case for human–AI teaming when the AI is fluent across both accurate and hallucinated content. The capability that does not yet exist is a routine clinical verification practice — an analog to the bibliographic discipline of academic writing — specified at the moment of use rather than after harm. The discipline is being asked to define what good looks like while deployment is already underway, so the verification standard has to be built around a tool already in millions of hands rather than gated in front of it.<sup>4</sup>

| TONE                                                   | CONTENT        |
|--------------------------------------------------------|----------------|
| identical                                              | accurate / not |
| the interface does not distinguish; the clinician must |                |

LLMS IN CLINICAL USE — FLUENCY WITHOUT WARRANTY

*LLMs produce hallucinations indistinguishable in tone from accurate information, and clinicians have not yet developed the practice of routine verification.*

EDITORS' SYNTHESIS OF JAMA EDITORIALS ON LLM CLINICAL USE (2023–2024)

REFERENCES

1. Wachter, R. M., & Brynjolfsson, E. (2024). "Will Generative Artificial Intelligence Deliver on Its Promise in Health Care?" JAMA 331(1):65–69, doi:10.1001/jama.2023.25054 — genAI adoption and the productivity paradox in health care.

2. Morath, B., Chiriack, U., Jaskowski, E., et al. (2024). "Performance and risks of ChatGPT used in drug information: an exploratory real-world analysis," European Journal of Hospital Pharmacy 31(6):491–497 — only 13 of 50 answers correct and complete enough to act on without risk of harm.

3. Bhattacharyya, M., Miller, V. M., Bhattacharyya, D., & Miller, L. E. (2023). "High Rates of Fabricated and Inaccurate References in

4. Triplett, S., Ness-Engle, G. L., & Behnen, E. M. (2025). "A comparison of drug information question responses by a drug information center and by ChatGPT," American Journal of Health-System Pharmacy 82(8):448–460, doi:10.1093/ajhp/zxae316 — pharmacist-rated overall accuracy of 50% across the ChatGPT responses.

5. World Health Organization (2024). Ethics and Governance of Artificial Intelligence for Health: Guidance on Large Multi-Modal Models, 18 January 2024 — names automation bias by health care professionals and the risk of false or inaccurate output.

6. Sallam, M. (2023). "ChatGPT Utility in Healthcare Education, Research, and Practice: Systematic Review on the Promising

ChatGPT-Generated Medical Content," Cureus 15(5):e39238 — 47% of 115 generated references fabricated.

Perspectives and Valid Concerns," Healthcare 11(6):887 — 60 records reviewed; benefits cited in 85%, with the accuracy and citation concerns alongside.

THE LEARNING ENGINEERING LENS

LE INSIGHT

LLM use in clinical practice is the live frontier case for human-AI teaming when the AI’s output is fluent across both accurate and hallucinated content. The capability that does not yet exist is the routine verification practice — a clinical analog to the bibliographic discipline of academic writing.

LENS APPROACH

The learning engineer’s deliverable here is a specific artifact, not a posture: a verification-at-point-of-use guardrail that counters cognitive offloading. Concretely, a confirm-before-act gate wrapped around the LLM output — every dosage, contraindication, or guideline citation the clinician would rely on must clear a forced source-attribution step that surfaces the underlying reference (or its absence) and requires an explicit confirmation against it before the output can be acted on. The gate is designed so the model’s fluency cannot satisfy it: a smooth, confident answer with no attributable source fails the check and is held back from patient-facing use. LENS builds and critiques exactly this guardrail in LEN 2 (the human-AI verification interface), stresses it in LEN 7 (governing the gate as deployment policy), and in LEN 10 (capstone) asks the student to design the confirm-before-act artifact for a clinical workflow already in use.

| DOMAIN     | MODES AT STAKE             |                           | LENS COMPETENCY |               |               |               |               | PROBLEM TYPE  | INDUCED | LEO |       |
|------------|----------------------------|---------------------------|-----------------|---------------|---------------|---------------|---------------|---------------|---------|-----|-------|
| HEALTHCARE | →                          | <div>H</div> <div>D</div> | →               | <div>D1</div> | <div>D2</div> | <div>D3</div> | <div>D4</div> | <div>D5</div> | 6       | 5.2 | 4 · 3 |
| TECHNOLOGY | Human-System Collaboration |                           |                 |               |               |               |               |               |         |     |       |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Specify a routine verification practice at the moment of use — what a clinician must independently confirm before acting on LLM output — before the tool enters clinical workflow.</li><li>Design the interface to mark provenance and uncertainty, so fluency alone cannot stand in for warranty of accuracy.</li><li>Restrict the sanctioned uses to those where verification is feasible and cheap, keeping unverifiable high-stakes outputs out of patient-facing work.</li></ul> | <ul style="list-style-type: none"><li>Monitor the documented failure pattern — fabricated citations, dosages, guidelines — across in-use cases to keep the risk profile current as models change.</li><li>Track whether the verification practice is actually being performed, since the tool’s fluency discourages exactly the checking it requires.</li><li>Govern adoption against emerging guidance so the standard for verification keeps pace with a tool already in widespread informal use.</li></ul> |

REFLECTION QUESTIONS

1. Identify a workflow in your domain currently being augmented by LLMs. What is the verification practice — and does it exist at the moment of use, or only after?
2. Design the verification deliverable that should accompany every clinician’s adoption of an LLM tool for patient-facing work.
3. The same fluency that makes LLM output easy to accept is what removes the cues a reader normally uses to doubt it. What interface signal would restore that friction at the moment of use without making the tool unusable?

WHO BUILDS THIS human factors & interaction design · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. TOOLS usability testing · cognitive walkthroughs · task analysis · design prototyping

# WHO Surgical Safety Checklist

**T** Training Gap    **N** Normalization of Deviance

**IMPACT** Death rate 1.5% → 0.8% in eight-site pilot; complications down >33%; adopted, a decade on, by facilities in 70% of the world's countries; Ontario population study (2014) found no significant mortality benefit after a province-wide rollout

## IN BRIEF

In 2008, Atul Gawande's team and the WHO introduced a single-page, nineteen-item surgical checklist applied at three junctures — before anesthesia, before incision, and before the patient leaves the operating room. Piloted across eight hospitals in eight countries, from Toronto to Tanzania, it nearly halved the surgical death rate (1.5% to 0.8%) and cut serious complications by more than a third — results published in the NEJM in 2009. The artifact was the checklist; the intervention was the system of forced pauses it created, three moments when a moving team had to stop and confirm shared state. A later Ontario study found no mortality benefit after a mandated rollout, surfacing the fidelity question: the artifact works only when the institution authorizes its honest use.

## BACKGROUND

Surgical harm — wrong-site operations, retained instruments, missed allergies, post-operative infection — was widespread across health systems rich and poor, and much of it stemmed not from a lack of skill but from teams that never paused to confirm shared understanding before acting. The knowledge to prevent these events existed; the reliable practice did not. A surgical team in motion rarely stopped to verify that everyone held the same picture of the patient and procedure, so a mismatch any member could have caught passed silently into the operation. The cost was global and patterned — the WHO estimated that, of the roughly 234 million operations performed each year, at least seven million produced major complications and about a million ended in death — and the bulk of that harm sat in failure modes that a brief verbal cross-check would have closed.<sup>o</sup>

## THE INTERVENTION

In 2008, the WHO and a Harvard team led by Atul Gawande introduced a single-page, nineteen-item Surgical Safety Checklist applied at three critical junctures: before anesthesia, before skin incision, and before the patient leaves the operating room. The team piloted it across eight hospitals in eight countries spanning very different economic and operational conditions — Toronto, London, Seattle, Auckland, Amman, New Delhi, Manila, and Ifakara — selecting sites deliberately to span the global income gradient. Testing the same artifact from Toronto to Tanzania was deliberate — it had to work in settings with vastly different resources to demonstrate that the gain came from the forced pause itself, not from any

one system's wealth. The artifact was intentionally minimal: items the team had tested against the surgical-safety literature for the smallest set that still spanned the highest-risk junctures.<sup>1</sup>

#### HOW IT WORKED

The checklist was the artifact; the intervention was the system of pauses it imposed. At each juncture a team that would otherwise keep moving had to stop, look at one another, and confirm names, the procedure, allergies, and equipment aloud. The pauses were the load-bearing element — the requirement to halt and speak mattered more than the specific list of items. Saying the confirmations aloud to one another, rather than each member assuming them privately, is what surfaced the mismatches a moving team would otherwise have carried into the incision. The mechanism turned the operating room into a momentarily flat-hierarchy space: a circulating nurse who noticed a missing antibiotic dose at the time-out had institutional cover to say so before the incision, a transaction the pre-checklist culture had not reliably authorized. The artifact's authority — that the team had to stop — was what made the speech act safe.<sup>2</sup>

#### THE EVIDENCE

The 2009 NEJM study reported the surgical death rate falling from 1.5% to 0.8% and major complications dropping by more than a third across all eight sites — confirming Gawande's framing that "gaps in teamwork and safety practices in surgery are substantial in countries both rich and poor." The checklist spread fast: a decade on, the Ariadne Labs / Lifebox review found facilities in 70% of the world's countries attesting to its use. That the pooled death rate roughly halved, and that complication rates fell at all eight sites — significantly at three — despite their differences in wealth and practice, is what made the result so persuasive: the effect tracked the intervention rather than the setting. The published mortality reduction was contested almost from the moment of publication: observational, pre/post, no concurrent control, susceptible to secular trends and Hawthorne-style attention effects. The pilot was always best read as the formation of a strong hypothesis under pragmatic conditions, not as a closed proof of population-scale mortality benefit.<sup>3</sup>

#### WHAT TRANSFERRED

Urbach and colleagues, writing in NEJM in 2014, examined more than 200,000 procedures across 101 Ontario hospitals before and after a province-wide rollout — Ontario's health ministry had required hospitals to publicly report their checklist compliance rate from July 2010, which is to say it mandated the reporting rather than the pause — and found no statistically significant reduction in operative mortality, in-hospital complications, length of stay, or readmission. The null result did not refute the checklist so much as illustrate its dependence on implementation fidelity: where a mandate replaced genuine authorization of the pause, the measured effect attenuated, making the checklist a paired lesson in both minimal-artifact design and the limits of mere compliance. The contrast between the pilot and the mandated rollout is the lasting teaching point: when the pause was genuinely used it worked, and when it became a box to check before proceeding the same paper produced nothing. Subsequent multi-country replications and the WHO's own follow-up work have been mixed — the institutional uptake outran the closed evidence of population-scale mortality benefit, and the contested mortality reduction is the part of the case that does not resolve.<sup>4</sup>



*Gaps in teamwork and safety practices in surgery are substantial in countries both rich and poor.*

ATUL GAWANDE, HARVARD GAZETTE, JANUARY 2009

REFERENCES

1. Haynes, A. et al. (2009), "A Surgical Safety Checklist to Reduce Morbidity and Mortality," NEJM — the 1.5%→0.8% result and major-complication reduction.

2. WHO Safe Surgery Saves Lives campaign documentation — the nineteen-item checklist and the three junctures.

3. Gawande, A. (2009), The Checklist Manifesto — the pause as the active mechanism.

4. Urbach, D. et al. (2014), "Introduction of Surgical Safety Checklists in Ontario, Canada," NEJM — null mortality result after mandate.

5. Bosk, C. et al. (2009), "Reality check for checklists," The Lancet — implementation fidelity.

6. Buissonniere, M. (2020), Checking In on the Checklist, Ariadne Labs and Lifebox — the decade review: facilities in 70% of the world's countries attest to using the checklist, with adoption far lower in low-income settings.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Surgical Safety Checklist is the canonical evidence that a tiny artifact — one page, nineteen items — can produce population-scale effects when paired with the structural change of **requiring a pause**. The checklist alone is paper. The pause alone is anxiety. Together they constitute the smallest effective capability intervention in the dataset. The Ontario follow-up underscores the secondary requirement: the artifact carries the effect only when the institution actually enforces the pause. Where mandate replaced authorization, the measured effect attenuated.

LENS APPROACH

LENS uses the WHO checklist as the canonical pre-committed institutional response paired with closed-loop evidence (induced 4.4; LENS D4/PT5). It runs in LEN 10 (capstone) as the end-to-end design exemplar — a problem identified, a minimal artifact prototyped, a multi-site pilot, a measurement regime, and global scale-out — and in LEN 4 for the measurement architecture that made the pilot effect provable and the Ontario null interpretable. Adjacent to SUBSAFE (Case 173) at the mandatory-mechanism layer, where the artifact's authority is the intervention.

| DOMAIN              | Modes Addressed | LENS Competency  | Problem Type | Induced | LEO   |
|---------------------|-----------------|------------------|--------------|---------|-------|
| HEALTHCARE          | → T N           | → D1 D2 D3 D4 D5 | 5            | 2.3     | 4 · 5 |
| Test and Evaluation |                 |                  |              |         |       |

APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- ▶ Design the minimal artifact — a single-page checklist — around a small number of forced pauses at the highest-risk junctures, keeping the list short enough to be used every time.
- ▶ Make the confirmations spoken aloud and shared across the team, so the pause surfaces mismatches rather than letting each member assume them privately.
- ▶ Pilot across deliberately varied sites so the measured gain can be shown to track the intervention rather than any one setting's resources.

## IN OPERATION / AFTER

- ▶ Measure outcomes (death and complication rates) directly so the effect can be confirmed and the artifact is not assumed to work merely because it is in use.
- ▶ Guard implementation fidelity as the scale grows — a mandate that turns the pause into a box to check reproduces the artifact without the effect.
- ▶ Build a governance signal that distinguishes genuine, authorized use from compliance check-off, so the fidelity gap is visible before a rollout is judged a failure.

## REFLECTION QUESTIONS

1. What is the smallest possible capability artifact in your domain that, paired with a required pause, would shift outcomes?
2. The WHO checklist was studied across eight countries. Design the multi-site evaluation that would establish whether your candidate intervention generalizes.
3. The Ontario mandated rollout produced no measurable mortality reduction. What governance signal would have surfaced the fidelity gap between authorized use and compliance check-off before the trial was declared a failure?

**WHO BUILDS THIS** learning science & instructional design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards



# Bristol Heart Babies Reform

**G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** Foundational UK case in clinical outcomes transparency; produced specialty-level performance reporting in UK cardiac surgery

## IN BRIEF

Through whistleblowing and a public inquiry, the Bristol Royal Infirmary was found to have had substantially worse pediatric cardiac-surgery outcomes than other UK centers for years. The Kennedy Inquiry — one of the most influential UK healthcare inquiries in modern times — located the capability gap in the absence of routine specialty-level outcomes reporting: surgeons did not know how their results compared with peers, patients did not know, and referrals did not respond to actual outcome data. The reform built national specialty-level outcomes registries, first in cardiac surgery and then across other specialties, making the UK one of the few countries that routinely publishes surgeon-level results — a paired intervention of technical registry plus institutional commitment to publish that ended outcomes data as the private property of clinicians.

## BACKGROUND

In pediatric cardiac surgery, small differences in a unit's performance can mean the difference between a child living and dying — yet in the UK of the 1980s and early 1990s, no system routinely compared outcomes between centers. A surgeon, a hospital, or a referring physician had no reliable way to know whether a given unit's results were normal or dangerously poor. Without comparison across centers, a dangerously poor unit and an ordinary one looked alike from inside, so the very gap that cost children's lives was the one no one was positioned to see.<sup>o</sup>

## THE INTERVENTION

After whistleblowing and a public inquiry into deaths at the Bristol Royal Infirmary between 1984 and 1995, the Kennedy Inquiry recommended routine, risk-adjusted, specialty-level outcomes reporting. The reform built national registries — beginning with cardiac surgery and extending to other specialties — together with a commitment to publish results, including at the level of individual surgeons. Starting in cardiac surgery and then extending outward was deliberate sequencing — the specialty where the harm had been exposed proved the model, and the registry then spread to fields that had not had their own Bristol.<sup>1</sup>

## HOW IT WORKED

The intervention was explicitly paired. The technical half — registry design, statistical risk adjustment so that surgeons taking hard cases are not penalized, and a publication architecture — was necessary but not sufficient. The cultural half — surgeons accepting that their comparative results would be visible — was equally load-bearing, and was the harder of

the two to secure. The risk adjustment was what made the cultural half securable: without it, surgeons who took the sickest patients would have been punished by the raw numbers, giving them every reason to resist publication or avoid hard cases.<sup>2</sup>

### THE EVIDENCE

The UK became one of the few countries that routinely publishes surgeon- and unit-level outcomes, and the cardiac-surgery registry is among the most mature specialty-outcomes regimes in any country. Outcomes data ceased to be the private property of clinicians and became a shared resource for patients, referrers, and surgeons themselves. That surgeons themselves became users of the data, not just subjects of it, is part of why the regime endured — comparison that had once felt like exposure became a tool the profession relied on to know where it stood.<sup>3</sup>

### WHAT TRANSFERRED

Bristol is the foundational UK case for outcomes transparency as a paired-intervention deliverable, and its registry model has been extended across NHS specialties and influenced later national-quality reforms. It pairs with Keystone ICU (Case 19) as healthcare-outcomes interventions operating at different layers — the bedside protocol and the system-level measurement regime. The two layers are complementary rather than competing: Keystone changes what happens at the point of care, while Bristol changes what the system can see about results across centers, and a mature regime needs both.<sup>4</sup>



*Patients should be able to gain access to information about the relative performance of a hospital, or a particular service or consultant unit.*

LEARNING FROM BRISTOL, FINAL REPORT SUMMARY, PARAGRAPH 18 (2001)

### REFERENCES

- Kennedy, I. (2001), Learning from Bristol: The Report of the Public Inquiry into Children's Heart Surgery at the Bristol Royal Infirmary 1984–1995 — the inquiry and its recommendations, paraphrased; the pull quote is verbatim from paragraph 18 of the Final Report Summary, and differs in wording from Recommendation 155, which carries the same requirement.
- Society for Cardiothoracic Surgery in GB & Ireland, national outcomes reports — the registry and surgeon-level publication.
- Berwick, D. (2013), A Promise to Learn — A Commitment to Act — the broader NHS-safety reform.
- Sherlaw-Johnson, C., Gallivan, S., Treasure, T. & Nashef, S. A. M. (2004), "Computer tools to assist the monitoring of outcomes in surgery," *European Journal of Cardio-Thoracic Surgery* 26(5):1032–1036 — risk-adjusted outcome monitoring (variable life-adjusted display) as it developed in UK cardiac surgery.
- NHS national clinical audit and registry documentation — extension across specialties.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Bristol is the canonical UK case for outcomes-transparency as a paired intervention. The technical capability — registry design, statistical risk adjustment, publication architecture — was necessary. The cultural capability — surgeons accepting that their results would be visible and comparable — was equally necessary. The pair has produced one of the most mature specialty-outcomes regimes in any country.

### LENS APPROACH

LENS uses Bristol in LEN 4 for outcomes-transparency as a paired- intervention deliverable and in LEN 7 for the institutional reform that made surgeon-level publication acceptable. The case pairs with Keystone ICU (Case 19) as healthcare-outcomes interventions at different layers.

| DOMAIN              | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------|-----------------|------------------|--------------|---------|-----|
| HEALTHCARE          | → G N           | → D1 D2 D3 D4 D5 | 5            | 2.1     | 4   |
| Test and Evaluation |                 |                  |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Pair the technical instrument — a registry with statistical risk adjustment and a publication architecture — with the cultural change of practitioners accepting that comparative results will be visible.
- ▶ Build risk adjustment in from the start so those who take the hardest cases are not punished by raw numbers, which is what makes the cultural half securable.
- ▶ Sequence the rollout to begin where the harm was exposed, proving the model in one specialty before extending it to fields that have not had their own crisis.

#### IN OPERATION / AFTER

- ▶ Make practitioners users of the data, not just subjects of it, so comparison becomes a tool the profession relies on rather than an exposure it resents — which is what sustains the regime.
- ▶ Extend the registry model across specialties over time, turning a single reform into a system-wide measurement regime.
- ▶ Pair the system-level transparency layer with point-of-care interventions (as with Keystone), since a mature outcomes regime needs both what the system can see and what happens at the bedside.

### REFLECTION QUESTIONS

1. What is the equivalent of surgeon-level outcomes transparency in your domain? What cultural change would have to accompany the technical instrument?
2. Design the registry and publication architecture for a specialty in your domain currently operating without outcomes transparency.
3. Bristol's risk adjustment was what let surgeons accept publication, by ensuring those who took the hardest cases were not punished by raw numbers. Identify a transparency measure in your domain that practitioners resist, and design the adjustment that would make the comparison fair enough to accept.

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Removing the Race Coefficient from eGFR

**D** Designed Out **G** Governance & Trust **N** Normalization of Deviance

**IMPACT** A clinical estimating equation that for two decades raised estimated kidney function for Black patients was retired through a governance process; the documented effect on disparities remains unknown

## IN BRIEF

For two decades the standard equation used to estimate glomerular filtration rate from serum creatinine — the kidney-function number on routine lab reports — included a coefficient that raised the estimate for patients reported as Black. The downstream effects were documented: later nephrology referral, later wait-listing for transplant. In 2021 the joint NKF-ASN Task Force, after reviewing over twenty candidate approaches with patient and clinician input, recommended immediate national implementation of the race-free 2021 CKD-EPI creatinine equation (Inker et al., NEJM 2021); clinical laboratories moved to adopt it. The case is a governance intervention — a change-control process that retired a construct after two decades of operational use. The honest hedge, preserved from the Task Force report and from follow-up commentaries, is that the new equation introduces a small bias for all groups, the disparities effect **remains unknown**, and the case is the construct-definition act, not a closed outcome proof. It pairs with pulse oximetry (Case 26) and pain assessment (Case 6) as the trio of “what counts as the same patient across race.”

## BACKGROUND

Glomerular filtration rate (GFR) is the standard measure of kidney function and the threshold variable for nephrology referral, medication dosing, and transplant wait-listing. Because direct GFR measurement is impractical at clinic scale, equations estimate it from serum creatinine and demographic inputs. The dominant equations from the late 1990s onward included a Black-race coefficient that raised the estimate for patients reported as Black — making them appear to have better kidney function than the same creatinine value implied for White patients.<sup>0</sup>

## THE INTERVENTION

The downstream effects of that coefficient are documented across the nephrology literature: later referral to specialty care, later transplant wait-listing, and a population-level shift in who was counted as having advanced chronic kidney disease. The defenders of the coefficient pointed to differences in creatinine generation across self-identified populations; the critics pointed to the logical and ethical problems of embedding a population-level adjustment in an individual-patient decision tool, and to the construct-definition question — what is the equation supposed to be measuring, and across whom does it have to mean the same thing?<sup>1</sup>

## HOW IT WORKED

The NKF-ASN Task Force, formed in 2020, ran the construct-revision process as a governance exercise: an external panel, patient-clinician input, more than twenty candidate alternatives, a sustained review window, and a published report (Delgado et al., *Am J Kidney Dis* 2021). The Task Force recommended immediate national adoption of the 2021 CKD-EPI creatinine equation, which eliminates the race coefficient. The replacement equation was published in parallel (Inker et al., *NEJM* 2021). The large commercial laboratories moved within months; a March 2022 survey of U.S. laboratories found 30.3% had adopted the new equation.<sup>2</sup> In December 2022 the governance response reached past the equation itself: the OPTN mandated that kidney-transplant programs review their waiting lists and backdate the accrued waiting time of Black candidates who had been disadvantaged by the former race-inclusive eGFR (effective January 2023, with the review required by January 2024) — a concrete, population-scale remediation of exactly the transplant-wait-listing disparity this case documents.

## THE EVIDENCE

The hedge in both the Task Force report and the follow-up commentary is binding and is preserved in the case. The new equation introduces a small bias for all groups relative to measured GFR; the **net effect on documented disparities in nephrology care and transplant access** is not yet measured at population scale, and the literature explicitly states that the disparities outcome **remains unknown**. The case is the construct-definition act — the right kind of governance intervention, run with the right kind of process — and the outcome evidence is the continuing work.<sup>3</sup>

## WHAT TRANSFERRED

What the eGFR case teaches is that some equity capability deliverables are construct-definition acts: choosing what gets predicted and what counts as the same patient is the design decision, not a downstream remediation. It pairs with pulse oximetry (Case 26), where the failure was in the validation architecture rather than the equation; and with pain assessment (Case 6), where the failure was in clinician-held false beliefs rather than the instrument. The trio together is the case-grounded basis for the subobjective **Gap attribution** (4.2) — distinguishing capability gaps attributable to human development, system design, and organizational performance — and for the subobjective **Fairness beyond omission** (4.6), of which eGFR is the construct-removal instance.

*What counts as the same patient across race is a construct-definition decision, not a downstream remediation.*

EDITORS' SYNTHESIS OF DELGADO ET AL. (2021) AND VYAS ET AL. (2020).

## REFERENCES

1. Vyas, Eisenstein, & Jones (2020), "Hidden in Plain Sight — Reconsidering the Use of Race Correction in Clinical Algorithms," *NEJM* 383(9):874–882 — broader race-in-algorithms survey.
2. Eneanya, Yang, & Reese (2019), "Reconsidering the Consequences of Using Race to Estimate Kidney Function," *JAMA* 322(2):113–114 — the equity argument that motivated the revision.
3. Genzen, J. R., Souers, R. J., Pearson, L. N., et al. (2022), "Reported Awareness and Adoption of 2021 Estimated Glomeru-
4. Delgado et al. (2021), "A Unifying Approach for GFR Estimation: Recommendations of the NKF-ASN Task Force on Reassessing the Inclusion of Race in Diagnosing Kidney Disease," *American Journal of Kidney Diseases* 79(2):268–288 (published online 2021; in print vol. 79, 2022), doi:10.1053/j.ajkd.2021.08.003. Cited by online-first year, the year of the Task Force recommendation.

lar Filtration Rate Equations Among US Clinical Laboratories, March 2022,” JAMA 328(20):2060–2062, doi:10.1001/jama.2022.15404 — the College of American Pathologists survey: 1,124 of 3,708 responding U.S. laboratories (30.3%) had adopted the 2021 CKD–EPI creatinine equation.

5. Inker et al. (2021), “New Creatinine- and Cystatin C-Based Equations to Estimate GFR without Race,” New England Journal of Medicine 385(19):1737–1749, doi:10.1056/NEJMoa2102953.
6. OPTN/UNOS (2022–2023), “Kidney waiting-time modifications for candidates affected by race-inclusive eGFR calculations” — the mandated backdating of Black candidates’ accrued waiting time (effective January 2023).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

eGFR is the canonical recent instance of construct-definition as an equity capability deliverable. A race coefficient was embedded in a continuous estimating equation for two decades, documented downstream effects on referral and transplant access, and was retired through a governance process. The disparities effect of the change remains unknown; the case is the construct-revision act, not the closed outcome.

### LENS APPROACH

eGFR is the construct-definition equity intervention in the race-construct trio (Cases 25, 26 and 6). LENS uses it in Domain 4 (Test and Evaluation) for the subobjective **Gap attribution** (4.2) — attributing the gap to construct definition, not to clinicians or patients — and in Domain 5 (Navigating Sociotechnical Constraints) for the governance process that revised the standard. Direct trio with Case 26 (pulse oximetry — the validation-architecture mechanism) and Case 6 (pain assessment — the human-development mechanism). Adjacent to the lending pair (Cases 186 and 196) at a different layer.

| DOMAIN            | MODES ADDRESSED | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|-------------------|-----------------|---------------------|--------------|---------|-------|
| CLINICAL MEDICINE | → D G N         | → D1 D2 D3 D4 D5    | 5            | 8.1     | 4 · 5 |
| HEALTH EQUITY     |                 | Test and Evaluation |              |         |       |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Reconsider the construct: ask what the equation should predict and across whom it should mean the same thing, separately from optimizing the residual.
- ▶ Run construct-revision as a governance process: an external panel, patient and clinician input, candidate alternatives examined openly, and a published report.
- ▶ Carry the hedge — the small all-group bias introduced and the unmeasured disparities effect — into the deployment documentation, not the press release.

#### IN OPERATION / AFTER

- ▶ Instrument the downstream outcomes (nephrology referral, transplant wait-listing) demographically so the disparities effect of the construct revision is actually measurable at population scale.
- ▶ Treat construct revision as ongoing: a coefficient was removed; a new validation architecture is the longer work.
- ▶ Make the construct-definition decision auditable — publish the candidate alternatives that were considered and the basis on which the chosen alternative was selected.

### REFLECTION QUESTIONS

1. Identify a clinical or algorithmic construct in your domain that embeds a population-level adjustment in an individual-patient decision. Was the construct revised, examined, or never questioned? What governance process would you run if it had to be revisited?
2. Specify the downstream outcomes you would instrument demographically to make the disparities effect of a construct revision measurable at population scale.
3. The new CKD–EPI equation introduces a small bias for all groups and the disparities effect **remains unknown**. What is the minimum follow-up evidence you would require before concluding the construct revision improved or worsened the equity outcome you care about?

**WHO BUILDS THIS** systems & human-factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Pulse Oximetry Across Skin Tones

**D** Designed Out   **G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** A bedside device validated on a non-representative population systematically under-detected hypoxemia in Black patients for thirty years; the bias persisted because validation asked only for a token darkly-pigmented minimum and reported accuracy as a single pooled number, never broken out by skin tone

## IN BRIEF

Pulse oximetry — the noninvasive bedside measurement of blood oxygen saturation — is one of the most widely used patient-monitoring tools in clinical medicine. Sjoding et al. (NEJM 2020) found that across two large cohorts, Black patients had nearly three times the frequency of **occult hypoxemia** (arterial saturation <88% despite a pulse-ox reading of 92–96%) as White patients. The finding replicated Jubran & Tobin (Chest 1990), published thirty years earlier and overlooked operationally. The bias persisted because validation asked only for two darkly pigmented subjects or fifteen percent of the pool and pooled every measurement into one accuracy number — the aggregate on FDA clearance documentation concealed a group-specific failure mode. The discovery drove FDA review and, in January 2025, a draft guidance recommending that manufacturers evaluate device performance across diverse skin pigmentations during validation. The case is a **failure-to-intervention arc**: the failure sat in the validation architecture for three decades; the intervention is the regulatory change-control on validation, and its measured effect on the disparities outcome is not yet documented.

## BACKGROUND

Pulse oximetry replaced repeated arterial blood draws as the bedside standard for monitoring oxygen saturation in the 1980s and 1990s. The device shines two wavelengths of light through tissue and infers saturation from the absorbance ratio. The inference depends, among other things, on the absorbance properties of the intervening tissue — including melanin pigmentation. The clearance documentation reports an aggregate accuracy number against arterial blood gas measurements.<sup>o</sup>

## THE INTERVENTION

In 1990 Jubran & Tobin reported, in Chest, that pulse oximeters tended to overestimate true oxygenation in patients with darker skin. The paper was published, cited intermittently, and did not drive a change in validation practice. Thirty years later Sjoding et al. (NEJM 2020) revisited the question in two large modern cohorts and reported that Black patients had nearly three times the frequency of **occult hypoxemia** — true arterial saturation below 88 percent despite a pulse-oximeter reading of 92–96 percent — as White patients. The structural form of the finding was the same as Jubran & Tobin's; the population and the salience



were different, and the COVID-19 pandemic, which made pulse oximetry a household-scale triage tool, made it harder to ignore.<sup>1</sup>

#### HOW IT WORKED

Validation at clearance asked for two darkly pigmented subjects or fifteen percent of the pool, and pooled every paired measurement into one accuracy figure. The aggregate accuracy number — clinically acceptable on average — concealed a group-specific failure mode. The capability gap was not in the clinician using the device or in the manufacturer's engineering; it was in the regulatory machinery that approved a measurement device without checking whether its measurement held across the patients it would meet. The Sjoding paper's lasting contribution was not the technical finding alone — Jubran & Tobin had supplied that — but the disconfirmation of the validation architecture.<sup>2</sup>

#### THE EVIDENCE

The FDA's January 2025 draft guidance recommends that pulse oximeter manufacturers evaluate device performance across diverse skin pigmentations during validation, and that the validation protocol explicitly stratify accuracy by skin tone. The guidance is the corrective-action half of the case: an intervention in the validation architecture, not in the device itself. The measured effect on the downstream disparities outcome — under-treated hypoxemia in patients of color — is not yet documented; it is the continuing work the intervention sets up.<sup>3</sup>

#### WHAT TRANSFERRED

What the case teaches is that a measurement-device failure can persist for three decades inside an aggregate accuracy number, and that the capability deliverable is a validation architecture that surfaces group-specific failure modes by stratifying outcome metrics by relevant demographic characteristics. Pulse oximetry pairs with eGFR (Case 25) and pain assessment (Case 6) in the race-construct trio. The mechanisms are distinct — eGFR is construct definition; pulse oximetry is validation architecture; Hoffman is clinician-held false belief — and the case-grounded lesson is that the diagnosis of the same surface harm has to attribute the gap to the right layer of the system before a remediation can land.

*Aggregate accuracy is not group accuracy. A device can be acceptable on average and unsafe for one population.*

EDITORS' SYNTHESIS OF SJODING ET AL. (2020) AND THE FDA 2025 DRAFT GUIDANCE.

#### REFERENCES

1. Jubran & Tobin (1990), "Reliability of Pulse Oximetry in Titrating Supplemental Oxygen Therapy in Ventilator-Dependent Patients," *Chest* 97(6):1420–1425 — original finding, published thirty years earlier.
2. Fawzy et al. (2022), "Racial and Ethnic Discrepancy in Pulse Oximetry and Delayed Identification of Treatment Eligibility Among Patients With COVID-19," *JAMA Internal Medicine* — downstream effect during the pandemic.
3. Sjoding, Dickson, Iwashyna, Gay, & Valley (2020), "Racial Bias in Pulse Oximetry Measurement," *New England Journal of Medicine* 383(25):2477–2478, doi:10.1056/NEJMc2029240.
4. FDA (2025), "Pulse Oximeters for Medical Purposes — Non-Clinical and Clinical Performance Testing, Labeling, and Premarket Submission Recommendations: Draft Guidance for Industry and Food and Drug Administration Staff," issued January 7 2025, Docket No. FDA-2023-N-4976; Federal Register notice 2024-31540, published January 7 2025 — the regulatory corrective-action artifact, language may evolve in final.

#### THE LEARNING ENGINEERING LENS

LE INSIGHT

Pulse oximetry is the validation–architecture instance of the race–construct trio. The bias was published in 1990, persisted for thirty years inside aggregate clearance accuracy, was re–documented in 2020, and reached the regulatory architecture in 2025. The capability deliverable is demographic stratification at validation, not after deployment.

LENS APPROACH

Pulse oximetry is the validation–architecture intervention in the race–construct trio (Cases 25, 26 and 6). LENS uses it in Domain 4 (Test and Evaluation) for the subobjective **Gap attribution** (4.2) — the gap is in the validation architecture, not the device or the clinician — and in Domain 5 (Navigating Sociotechnical Constraints) for the FDA clearance / device oversight regime. Adjacent to Case 35 (radiology AI miscalibration), which is the same diagnosis at a different technological boundary, and to the Epic Sepsis Model (Case 5) for the post–deployment–surveillance pattern.

| DOMAIN          | Modes Addressed | LENS Competency     | Problem Type | Induced | LEO   |
|-----------------|-----------------|---------------------|--------------|---------|-------|
| MEDICAL DEVICES | D G N           | D1 D2 D3 D4 D5      | 5            | 8.2     | 4 · 5 |
| CLINICAL CARE   |                 | Test and Evaluation |              |         |       |
| HEALTH EQUITY   |                 |                     |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Stratify device–validation outcomes by every demographic characteristic that can change the measurement physics, before clearance, not after deployment.
- ▶ Specify the group–specific accuracy that would count as acceptable, separately from the aggregate; do not allow the aggregate to substitute.
- ▶ Treat replication of an earlier finding (Jubran & Tobin → Sjoding) as a verification trigger, not a duplication — the same finding in a different population is itself evidence.

IN OPERATION / AFTER

- ▶ Audit deployed devices on the population that actually uses them, on a schedule that surfaces drift; aggregate accuracy is not a substitute.
- ▶ Tie the regulatory clearance update to a downstream–outcome surveillance plan — under–treated hypoxemia, in this case — so the intervention’s effect on the harm can be measured.
- ▶ When a published finding sits operationally inert for years, ask whether the publication channel reached the operational community; the structural problem may be in how validation evidence diffuses, not in whether it exists.

REFLECTION QUESTIONS

1. Identify a measurement device in your domain whose validation reports an aggregate accuracy number. Across which demographic characteristics could the measurement physics change? What would a stratified validation protocol look like, and who would have to approve it?
2. The Sjoding finding replicated Jubran & Tobin thirty years later. What is the institutional architecture that should have surfaced the original finding to the regulatory regime? Where did the publication–to–operations channel break, and where does it still break in your domain?
3. The FDA 2025 draft guidance corrects the validation architecture. Specify the downstream outcome (under–treated hypoxemia in patients of color) and the surveillance plan you would tie to the guidance so the intervention’s effect on harm is measurable.

**WHO BUILDS THIS** systems & human–factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Anesthesia Monitoring Standards and the APSF — The Mortality Decline

**H** Human–System Interface   **K** Knowledge & Institutional Memory   **G** Governance & Trust

**IMPACT** Continuous pulse oximetry and capnography — encouraged in the 1986 Harvard/ASA minimum-monitoring standards and required by later ASA revisions (pulse oximetry from 1989; expired-CO<sub>2</sub> confirmation of tube placement in the early 1990s; continual capnography in general anesthesia from 1999) — converted silent hypoxemia and esophageal intubation from undetectable failures into monitored, recoverable ones; anesthesia-related mortality fell dramatically over subsequent decades — multifactorial decline

## IN BRIEF

Through the early 1980s, anesthesia in the United States carried a documented patient-safety crisis: silent intraoperative hypoxemia and unrecognized esophageal intubation produced catastrophic outcomes — brain injury and death — that were structurally undetectable until the harm had occurred. A 1982 ABC news special on anesthesia catastrophes converged with a malpractice-insurance crisis to force institutional attention. In 1986 Eichhorn et al. (JAMA, 1986) published the Harvard Medical School minimum monitoring standards — which mandated an oxygen analyzer and disconnection alarm and **encouraged** the newer continuous pulse oximetry and capnography, the additions designed to make hypoxemia and misplaced endotracheal tubes detectable early enough to recover; both were required by later ASA revisions (pulse oximetry from the 1989 amendment; expired-CO<sub>2</sub> confirmation of tube placement in the early 1990s; continual capnography in general anesthesia from the amendment effective July 1999). The ASA adopted essentially the same standards in 1986. The Anesthesia Patient Safety Foundation, founded in 1985, institutionalized the broader change-management effort. Anesthesia-related mortality fell dramatically over subsequent decades — Braz et al.'s systematic review of 33 studies published between 1954 and 2007 puts Brazilian and worldwide anaesthesia-related mortality at fewer than one death per 10,000 anaesthetics over its final two decades — and malpractice premiums declined alongside. The field's own histories preserve the hedge: the decline has multiple co-varying causes (training, device design, pharmacology, team composition), and the device standards themselves still fail in documented edge cases.

## BACKGROUND

Through the early 1980s, anesthesia in the United States carried a documented patient-safety crisis with two structurally similar failure modes at its center. Silent intraoperative hypoxemia — falling oxygenation in a paralyzed, ventilated patient — produced brain injury and death without the clinical signal a conscious patient would have given. And esophageal intubation — endotracheal tube misplaced in the esophagus instead of the trachea —

produced no ventilation at all, often noticed only when the patient was already decompensating. Both failures were detectable in principle and in practice undetectable, because the cue did not reach the anesthetist in time.<sup>9</sup>

#### THE INTERVENTION

The institutional forcing function arrived from outside the field. A 1982 ABC news special documented a string of catastrophic anesthesia outcomes, and the malpractice insurance environment for anesthesiology had reached a crisis level that made the status quo untenable. In 1985 the Anesthesia Patient Safety Foundation was founded — among the first specialty-organized patient-safety foundations in medicine. In 1986 Eichhorn et al. published in JAMA the Harvard Medical School minimum monitoring standards: a defined minimum set of continuous monitors for every anesthetized patient. Pulse oximetry and capnography — then just entering wide use — were **encouraged** rather than required in the 1986 text, and became the additions when later revisions mandated them: pulse oximetry in the 1989 amendment, identification of expired carbon dioxide to verify tube placement in the early 1990s, and continual capnography through a general anesthetic in the amendment effective July 1999.<sup>1</sup>

#### HOW IT WORKED

Pulse oximetry made oxygenation continuously visible to the anesthetist; capnography made the end-tidal carbon-dioxide trace visible, which immediately distinguishes tracheal from esophageal intubation in the first breaths. Together they converted two structurally undetectable failures into monitored, recoverable ones. The American Society of Anesthesiologists adopted essentially the same minimum standards in 1986–87. Eichhorn's follow-up (Anesthesiology, 1989) reported declines in preventable mishaps in the Harvard hospitals consistent with the mechanism the standards targeted.<sup>2</sup>

#### THE EVIDENCE

Anesthesia-related mortality fell dramatically over the subsequent decades. Braz and colleagues, reviewing 33 studies published between 1954 and 2007, put both Brazilian and worldwide anaesthesia-related mortality at fewer than one death per 10,000 anaesthetics over the final two decades of that window. Malpractice premiums for anesthesiology declined alongside — one of the few examples in American medicine of a specialty's malpractice cost falling as the specialty's safety record improved. The case is regularly cited as a canonical example of the cue/alert design as the capability deliverable: when the right signal is made visible at the right point in the workflow, an undetectable failure becomes recoverable, and the institutional outcome record moves.<sup>3</sup>

#### WHAT TRANSFERRED

The hedges the field's own histories preserve are decisive. The mortality decline has multiple co-varying causes — anesthetist training, device design improvements, pharmacological change (newer agents are inherently safer), team composition with broader monitoring presence, and the monitoring standards themselves — and attribution of the entire effect to the standards overstates what the evidence can support. The device standards themselves still fail in documented edge cases: pulse oximetry is unreliable across skin tones (Case 26 in this same v2 batch), and capnography can mislead in specific physiology. The case teaches the cue/alert-design form of capability engineering at its most durable, with the

honest qualification that the institutional decline is consistent with the standards but not isolatable to them alone.<sup>4</sup>

*The deliverable is not the device. It is the standard that makes the device's signal non-waiverable across the specialty.*

EDITORS' SYNTHESIS OF EICHHORN ET AL. (1986) AND APSF FOUNDING DOCUMENTS.

REFERENCES

1. Eichhorn, Cooper, Cullen, Maier, Philip, & Seeman (1986), "Standards for Patient Monitoring During Anesthesia at Harvard Medical School," JAMA 256(8):1017–1020.

2. American Society of Anesthesiologists, "Standards for Basic Anesthetic Monitoring," approved 21 October 1986 and amended since — the original standard and the amendment sequence that made pulse oximetry required (1989), expired-CO2 verification of tube placement required (early 1990s), and continual capnography through a general anesthetic required (effective July 1999).

3. Eichhorn (1989), "Prevention of intraoperative anesthesia accidents and related severe injury through safety monitoring," Anesthesiology 70(4):572–577.

4. Braz, L. G., Módolo, N. S. P., do Nascimento, P. Jr., et al. (2009), "Mortality in anesthesia: a systematic review," Clinics (São Paulo) 64(10):999–1006 — 33 studies, 1954–2007; anaesthesia-related mortality below one death per 10,000 anaesthetics in the final two decades.

5. Anesthesia Patient Safety Foundation (1985 – present), founding documents and the APSF Newsletter — institutional-history record of the broader change effort.

6. Sjoding et al. (2020), "Racial Bias in Pulse Oximetry Measurement," NEJM 383(25):2477–2478 — the racial-bias edge case the standard still carries (cross-reference Case 26).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Harvard / ASA / APSF anesthesia-monitoring intervention is the canonical cue/alert-design success: silent intraoperative hypoxemia and esophageal intubation were converted from undetectable failures into monitored, recoverable ones, and the specialty's mortality and malpractice record moved over subsequent decades. The decline is multifactorial; the device standards still fail in documented edge cases (Case 26).

LENS APPROACH

Anesthesia monitoring is the canonical cue/alert intervention (induced 3.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the cue/alert design as deliverable; in Domain 5 (Navigating Sociotechnical Constraints) for the APSF / ASA institutional discipline that made the standard non-waiverable; and in Domain 3 (Human-System Collaboration) at the human-device boundary. Adjacent to Case 26 (pulse oximetry across skin tones) — same instrument, the canonical edge case the standard still carries. Pair with Case 23 (WHO Surgical Checklist) at the mandatory-standard layer.

| DOMAIN          | MODES ADDRESSED | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|-----------------|-----------------|---------------------|--------------|---------|-------|
| ANESTHESIOLOGY  | H K G           | D1 D2 D3 D4 D5      | 5            | 3.1     | 4 · 3 |
| PATIENT SAFETY  |                 | Test and Evaluation |              |         |       |
| MEDICAL DEVICES |                 |                     |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the cue/alert to surface the failure mode that produces harm — silent hypoxemia, misplaced tube — at the point in the workflow where the operator can still recover, not in a post-hoc record.
- ▶ Pair the device standard with the institutional commitment (APSF, ASA) that makes the standard non-waiverable across the specialty, so adoption is a

IN OPERATION / AFTER

- ▶ Track outcome metrics (mortality, malpractice premiums) alongside the standards, while attributing carefully — multifactorial declines are the norm, not the exception, in long-window safety interventions.
- ▶ Treat the malpractice-cost signal as institutional evidence the standard is recognized to have worked, not as the safety evidence itself; the safety evidence is mortality and harm.

- profession-level deliverable rather than a per-institution choice.
- ▶ Build the standard with edge-case acknowledgment: the device is the strongest available signal, not closed proof; pulse oximetry across skin tones (Case 26) is the canonical edge case the standard has to keep visible.
- ▶ Maintain the edge-case surveillance the standard sets up: pulse oximetry across skin tones, capnography under specific physiology — each documented edge case is a place where the standard's coverage is incomplete.

### REFLECTION QUESTIONS

1. Identify an undetectable failure mode in your domain — one where the cue does not reach the operator in time to recover. What is the analog of pulse oximetry / capnography for that failure: the continuous signal that would convert it from undetectable to monitored?
2. Specify the institutional commitment (specialty foundation, regulatory standard) that would make the new signal non-waiverable. The device alone is not the deliverable; the standard is.
3. The anesthesia mortality decline is multifactorial. What is the minimum decomposition you would publish — training, device, pharmacology, team — to let downstream practitioners learn which components carried the most weight, rather than attributing the outcome to a single intervention?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

## CASE 28

HEALTH PROFESSIONS EDUCATION

INTERPROFESSIONAL COLLABORATION

PATIENT SAFETY

2013 – 2015

# Interprofessional Education and the Evidence Gap

**K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Decades-long well-funded movement to educate health professionals together for collaborative care; Cochrane 2013 found only 15 studies between 1999 and 2011 met inclusion criteria; IOM 2015 made the gap the central finding — the commissioned review in its Appendix A names a 'paucity of high-quality research' measuring the impact of IPE on practice processes and patient outcomes

## IN BRIEF

Interprofessional Education (IPE) is a decades-long, well-funded movement premised on the idea that doctors, nurses, pharmacists, and allied professionals should learn together so they can collaborate better in practice. The Cochrane review (Reeves et al., 2013) found only 15 studies between 1999 and 2011 met its inclusion criteria, and while those studies showed some positive outcomes, the evidence base for linking IPE to measurable changes in practice and patient outcomes was thin. The IOM report (2015) made the gap the central finding: the commissioned review it carries as Appendix A names a paucity of high-quality research measuring IPE's impact on practice processes and patient outcomes, and it proposed a conceptual model for doing the measurement properly. The case is the canonical instance in the corpus of a large, sincere, multidisciplinary translation effort whose core problem is that the field cannot yet measure whether the intervention works. It is the case-grounded basis for the enthusiasm-evidence gap as a curricular concept and motivates the Domain 3 subobjective **Collaboration measurement** (3.4) — the recurring pattern that a field instruments its enthusiasm faster than its outcomes.

## THE SHIFT

The premise of Interprofessional Education is straightforward and credible: health-professions practice is collaborative, so health-professions training should be collaborative. Doctors, nurses, pharmacists, dietitians, social workers, and allied professionals should learn alongside each other, ideally in clinical and simulation contexts, so that the coordination capability is built during training rather than improvised on the wards. The movement is decades old, has institutional support across major accreditation bodies, and is well-funded by government and foundation sources.<sup>o</sup>



## WHAT IS EMERGING

Reeves et al.'s 2013 Cochrane review applied the standard systematic-review machinery to the IPE outcome literature: what studies, between 1999 and 2011, met the inclusion criteria for evaluating IPE interventions against collaborative-practice or patient outcomes? The answer was

15. Across that small set, the studies showed some positive outcomes — on diabetes care, emergency-department culture and satisfaction, and a small number of clinical metrics — but the evidence base for linking IPE to measurable changes in practice and patient outcomes was thin. The reviewers' own conclusion was that the field had not produced the outcome evidence its scale of investment implied it should have.<sup>1</sup>

## THE CAPABILITY QUESTION

The IOM's 2015 report, *Measuring the Impact of Interprofessional Education on Collaborative Practice and Patient Outcomes*, made the evidence gap the central finding. The commissioned review the report carries as Appendix A puts it plainly: there is a paucity of high-quality research measuring the impact of interprofessional education on practice processes and patient outcomes. The report proposed a conceptual model for doing the measurement properly — a chain from IPE intervention through learner outcome, collaborative behavior, organizational practice, and patient outcome — and named the methodological work the field had to do to get from the current evidence state to a defensible causal claim. The report is itself the artifact: a national-academy synthesis that the central problem of the field is the evidence architecture, not the intervention.<sup>2</sup>

## EARLY EVIDENCE

What the case teaches at the LENS framing layer is the structural form of the enthusiasm-evidence gap at the multidisciplinary-translation scale. IPE is a sincere, well-funded, decades-long movement on which a great deal of curricular and operational investment has been made; the strongest synthesis of the outcome literature concludes that the field instrumented its enthusiasm faster than it instrumented outcomes. The case is the case-grounded basis for the Domain 3 subobjective **Collaboration measurement** (3.4) — naming the enthusiasm-evidence gap as a canonical anti-pattern — and for the subobjective **Judgment under inadequate evidence** (4.4): practitioners and program designers in IPE have had to decide for decades on what to build, on evidence that does not establish the causal claim the field's premise rests on.<sup>3</sup>

## OPEN PROBLEMS

In pair with the Colorado CTSA team-science training case (40) and implementation-science training (41), IPE is the frontier endpoint of the multidisciplinary-translation thread. Team science showed the measurement can be done at single-program scale with a validated instrument; IPE shows what happens when the field-scale evidence architecture has not been built. The pair-plus-trio teaches the Domain 3 collaboration-measurement subobjective in both directions: collaboration as a unit of measurement is possible (Case 40), and the field-scale evidence is structurally absent (this case).

*The field instrumented its enthusiasm faster than its outcomes. The strongest synthesis names the gap as the central finding, not the intervention as the failure.*



EDITORS' SYNTHESIS OF REEVES ET AL. (2013) AND IOM (2015).

## REFERENCES

1. Reeves, Perrier, Goldman, Freeth, & Zwarenstein (2013), "Interprofessional education: effects on professional practice and healthcare outcomes (update)," *Cochrane Database of Systematic Reviews*, doi:10.1002/14651858.CD002213.pub3.
2. Institute of Medicine (2015), *Measuring the Impact of Interprofessional Education on Collaborative Practice and Patient Outcomes*, National Academies Press, doi:10.17226/21726, NCBI Bookshelf NBK338360.
3. WHO (2010), *Framework for Action on Interprofessional Education and Collaborative Practice* — the international policy backdrop.
4. v2 paired cases: Team-science training (Case 40), Implementation-science training (Case 41).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Interprofessional Education is the canonical enthusiasm-evidence-gap case in the corpus. A decades-long, well-funded movement; the strongest synthesis of the outcome literature concludes that the evidence base for the field's central claim — that IPE produces measurable changes in practice and patient outcomes — is structurally thin. The IOM 2015 conceptual model is the proposed evidence architecture; the field's task is to build to it.

### LENS APPROACH

IPE is the field-scale enthusiasm-evidence-gap case (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) as the case-grounded basis for the enthusiasm-evidence-gap reading of subobjective 3.4 and for the subobjective **Judgment under inadequate evidence** (4.4) — IPE is the pattern's largest instance. Pair with Case 40 (team-science training, where measurement is possible at program scale) and Case 41 (implementation-science training, where stated goals run ahead of operational practices).

| DOMAIN                          | MODES AT STAKE | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------|----------------|-----------------------|--------------|---------|-----|
| HEALTH PROFESSIONS EDUCATION    | K N            | D1 D2 D3 D4 D5        | 5            | 2.1     | 4   |
| INTERPROFESSIONAL COLLABORATION | →              | → Test and Evaluation |              |         |     |
| PATIENT SAFETY                  |                |                       |              |         |     |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Design the IPE intervention with an outcome-measurement chain in mind from the start: learner outcome → collaborative behavior → organizational practice → patient outcome. The IOM model is the published reference.
- ▶ Specify, before the intervention launches, which links in the chain the evaluation will measure and which it will leave as assumed; the field-scale gap is partly the result of every program leaving the same links unmeasured.
- ▶ Treat enthusiasm-as-evidence as a foreseeable failure mode in any field-scale capability movement; the IPE pattern recurs across other interdisciplinary translation efforts (see Case 41).

### IN OPERATION / AFTER

- ▶ Commission field-level evidence architecture, not only program-level evaluation; the gap the IOM names is at the field level and the program-level studies cannot close it on their own.
- ▶ Publish negative or null findings as a contribution to the evidence base, not as a program failure; the field's evidence gap is partly the result of selective publication on the positive side.
- ▶ Carry the IOM conceptual model into curriculum design conversations: a sub-competency that names the enthusiasm-evidence gap explicitly is the case-grounded curricular response to this pattern.

## REFLECTION QUESTIONS

1. Identify a field-scale capability movement in your domain whose evidence architecture has not kept pace with its operational and curricular investment. What would the IOM-style conceptual model look like for that field's outcome chain?
2. Specify, for the next IPE-style program you would design or evaluate, which links in the outcome chain (learner / collaborative behavior / organizational practice / patient outcome) the evaluation will measure and which it will leave assumed. The field-scale gap is the accumulated result of leaving the same links assumed.

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

# Bar-Code Medication Administration — Cue at the Point of Care

**H** Human-System Interface **K** Knowledge & Institutional Memory **D** Designed Out

**IMPACT** A before-and-after study at an academic medical center associated bar-code electronic medication administration (bar-code eMAR) with a 41.4% reduction in nontiming administration errors and a 50.8% reduction in potential adverse drug events; transcription errors on order transcription were eliminated (6.1% to zero); a later staged single-site rollout reported a 55.4% reduction in harmful medication errors, 0.65 to 0.29 per 100,000 medications

## IN BRIEF

Wrong-drug and wrong-patient administration errors are a persistent failure mode in hospital pharmacy: the cue the human operator needs in order to catch the mismatch is structurally absent at the bedside, because the order, the dispensed medication, and the patient are connected only by paper documentation and clinical memory under time pressure. Bar-code medication administration (BCMA) supplies the cue in hardware: a bedside scan of the medication's bar code against the patient's wristband, gated by the electronic medication administration record. Poon et al. (NEJM, 2010) evaluated bar-code eMAR at an academic medical center using a before-and-after observational design and reported a 41.4% reduction in nontiming administration errors, a 50.8% reduction in potential adverse drug events, and the elimination of transcription errors on order transcription (6.1% to zero). A later staged single-site rollout (Thompson et al., 2018) reported a 55.4% reduction in harmful medication errors, from 0.65 to 0.29 per 100,000 medications. The study is explicit that the design is quasi-experimental — before-and-after / observational — not a randomized trial, and that significant workflow changes were required for the cue to land. The case is the canonical point-of-care cue/alert intervention, paired with Case 27 (anesthesia monitoring / APSF) at the cue-as-deliverable layer and with Case 23 (WHO Surgical Checklist) at the mandatory-mechanism layer.

## BACKGROUND

The wrong-drug and wrong-patient administration error is a structural failure mode of the hospital medication chain. A physician's order, a pharmacist's dispense, a nurse's bedside administration, and a patient's wristband identification are connected only by paper documentation and the operator's clinical memory under time pressure. The cue that would let the bedside nurse catch a mismatch — "this medication's identity does not match this patient's record" — is not present in the workflow unaided. The conventional safeguards (the

five rights, double-checks, naming protocols) are cognitive and procedural; the cue itself is structurally absent.<sup>9</sup>

#### THE INTERVENTION

Bar-code medication administration supplies the cue in hardware. The medication unit-of-use carries a bar code, the patient wears a coded wristband, and the electronic medication administration record (eMAR) gates the administration on a bedside scan that matches medication-to-order-to-patient. The cue lands at the moment of administration, before the harm, and the system refuses to acknowledge the administration if any element does not match. The structural form is the same as the cue/alert interventions in Case 27 (continuous pulse oximetry, capnography under anesthesia): the cue at the right point in the workflow converts an undetectable failure into a monitored, recoverable one.<sup>1</sup>

#### HOW IT WORKED

Poon et al. (New England Journal of Medicine, 2010) evaluated bar-code eMAR at a large academic medical center using a before-and-after observational design. The headline findings were a 41.4% reduction in nontiming administration errors (the failure mode the cue is designed to catch), a 50.8% reduction in potential adverse drug events (the harm the errors would have produced), and the elimination of transcription errors on order transcription (6.1% to zero) upstream of administration. Thompson and colleagues, tracking a staged single-site rollout across fifty inpatient units in an interrupted-time-series analysis, documented a 55.4% reduction in harmful medication errors — 0.65 to 0.29 per 100,000 medications — supporting the transfer of the result across institutional contexts.<sup>2</sup>

#### THE EVIDENCE

The study's design is what the case's hedge rests on. The evaluation is quasi-experimental — before-and-after / observational — not a randomized trial, and the authors are explicit that significant workflow changes were required for the cue to land in operation. The cue is not the deliverable alone; the workflow is also the deliverable, and the attribution of the observed reduction to the cue rather than to the workflow change rests on plausibility and on the mechanism the cue is designed against. The 55.4% transfer number from the later rollout supports the mechanism, and that rollout figure also rests on observational data rather than randomized comparison.<sup>3</sup>

#### WHAT TRANSFERRED

What the case carries for the corpus is the cue-as-deliverable pattern at the point of care, with the explicit before-and-after-design hedge. Paired with Case 27 (anesthesia monitoring / APSF), bar-code eMAR shows the same structural form in a different specialty: the cue at the right point in the workflow converts an undetectable mismatch into a monitored one. Paired with Case 23 (WHO Surgical Checklist), the case sits at the mandatory-mechanism layer: the bar-code scan, like the checklist, is a workflow artifact the operator cannot route around, and the institutional commitment to that non-routability is part of why the cue lands.

*The cue is supplied in hardware at the point of administration; the workflow change is part of the deliverable, and the design is observational rather than randomized.*

EDITORS' SYNTHESIS OF POON ET AL. (2010) AND THE BCMA IMPLEMENTATION LITERATURE.

## REFERENCES

1. Poon, E. G., Keohane, C. A., Yoon, C. S., Ditmore, M., Bane, A., Levitzion-Korach, O., Moniz, T., Rothschild, J. M., Kachalia, A. B., Hayes, J., Churchill, W. W., Lipsitz, S., Whittemore, A. D., Bates, D. W., & Gandhi, T. K. (2010). Effect of bar-code technology on the safety of medication administration. *New England Journal of Medicine*, 362(18):1698–1707. doi:10.1056/NEJMs0907115 — the case's primary evaluation.
2. Bonkowski, J., Carnes, C., Melucci, J., Mirtallo, J., Prier, B., Reichert, E., Moffatt-Bruce, S., & Weber, R. J. (2013). Effect of barcode-assisted medication administration on emergency department medication errors. *Academic Emergency Medicine*, 20(8):801–806 — adjacent transfer evidence.
3. Thompson, K. M., Swanson, K. M., Cox, D. L., Kirchner, R. B., Russell, J. J., Wermers, R. A., Storlie, C. B., Johnson, M. G., & Naessens, J. M. (2018), "Implementation of Bar-Code Medication Administration to Reduce Patient Harm," *Mayo Clinic Proceedings: Innovations, Quality & Outcomes* 2(4):342–351, doi:10.1016/j.mayocpiqo.2018.09.001, PMID:30560236, PMCID:PMC6257885 — later single-site roll-out reporting a 55.4% reduction in harmful medication errors (0.65 to 0.29 per 100,000 medications), by interrupted time series across fifty inpatient units.
4. Institute for Safe Medication Practices (2019), *Guidelines for Safe Electronic Communication of Medication Information — the display and communication standards that bar-code scanning systems and eMARs are built against.*

## THE LEARNING ENGINEERING LENS

## LE INSIGHT

Bar-code medication administration is the canonical small-tier cue/alert intervention at the point of care. The cue lands at the bedside before the harm, the workflow change is part of the deliverable, and the headline results (41.4%, 50.8%, transcription errors eliminated, and a 55.4% transfer figure) rest on a before-and-after / observational design that the case preserves verbatim.

## LENS APPROACH

BCMA is the small-tier point-of-care cue/alert intervention (induced 3.1; LENS D4/PT5) — Domain 4 for cue-as-deliverable; Domain 3 for the workflow-around-the- cue. Pair with Case 27 (anesthesia/APSf) and Case 23 (WHO Surgical Checklist).

| DOMAIN         | MODES ADDRESSED | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|----------------|-----------------|---------------------|--------------|---------|-------|
| CLINICAL CARE  | H K D           | D1 D2 D3 D4 D5      | 5            | 3.1     | 4 · 3 |
| PATIENT SAFETY |                 | Test and Evaluation |              |         |       |
| HEALTH IT      |                 |                     |              |         |       |

## APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- ▶ Design the cue to land at the point in the workflow where the operator can still recover — for bar-code eMAR, the bedside administration scan, before the dose reaches the patient — not at a post-hoc reconciliation point.
- ▶ Treat the workflow change as part of the deliverable, not as friction to minimize; the cue's effect is conditional on the workflow that surrounds it, and the evaluation has to acknowledge both.
- ▶ Pair the device standard with the institutional commitment that makes the scan non-routable around — the equivalent of the APSf / ASA monitoring standard in anesthesia, or the WHO checklist's mandatory-mechanism status.

## IN OPERATION / AFTER

- ▶ Report the headline reductions (41.4%, 50.8%, transcription errors eliminated, 55.4%) together with the design hedge — before-and-after / observational, significant workflow changes required — so the evidence the field cites is the evidence the field can actually use.
- ▶ Track the transfer figure (the 55.4% from the later single-site rollout) as evidence the mechanism survives a new institutional context, while preserving the same observational-design qualification on the new figure.
- ▶ Build the next round of evaluation against the residual error rate (the failures that survive the cue) so the institution learns where the workflow leaks remain, rather than treating the reduction as the closure of the failure mode.

## REFLECTION QUESTIONS

1. Identify a point in your domain's workflow where the cue the operator needs in order to catch a mismatch is structurally absent. What would the hardware-supplied analog of bar-code eMAR look like — and at what point in the workflow does it have to land for the operator to still recover?
2. The headline results (41.4%, 50.8%, transcription errors eliminated, 55.4%) rest on a before-and-after observational design. What additional evidence — randomized comparison, segmented analysis by unit type, residual-error analysis — would you require before treating the figures as a closure rather than a strong signal?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · task analysis · design prototyping

# Surgical Skill Video Peer-Rating Predicts Complications

**H** Human-System Interface **D** Designed Out **K** Knowledge & Institutional Memory

**IMPACT** Twenty bariatric surgeons each submitted a representative gastric-bypass video, rated blind by at least 10 peers; skill ratings ranged 2.6–4.8; the bottom skill quartile had a higher complication rate than the top (14.5% against 5.2%) across a registry of 10,343 patients, and greater skill was associated with fewer reoperations, readmissions, and emergency department visits

## IN BRIEF

Surgical complications are conventionally attributed to patient factors, hospital factors, and case-mix. Birkmeyer et al. (NEJM, 2013) asked a more direct question: can the surgeon's actual technical capability be measured well enough to predict patient outcomes? Twenty bariatric surgeons each submitted a representative video of a laparoscopic gastric bypass; the videos were rated blind by at least 10 peers on a structured skill scale. Skill ratings ranged from 2.6 to 4.8. Linked to a Michigan registry of 10,343 patients, the bottom skill quartile had a higher complication rate than the top — 14.5% against 5.2% — and greater skill was associated with fewer reoperations, readmissions, and emergency department visits. The authors call the findings preliminary and name the skill-versus-volume confound explicitly: low-skill surgeons also did fewer cases and operated more slowly, so the contribution of skill itself versus the contribution of case volume is partly open. The hedge is part of the case. The proposed primary anchor is 2.1 (measuring the failure mode you care about) with C3 and C1 alternates; the editor may move it. Adjacent to JIGSAWS (Case 31) at the surgical-skill-measurement layer.

## BACKGROUND

The conventional accounting for surgical complications attributes outcomes to patient factors (age, comorbidity, severity), hospital factors (volume, ICU support), and case-mix. The surgeon's actual technical capability — what the surgeon does with the instruments in this operation — is conventionally treated as too hard to measure at scale, and so as a hidden variable in the outcome equation. Birkmeyer et al. (New England Journal of Medicine, 2013) asked whether that variable could be brought into the open: can the surgeon's skill be measured well enough, at scale, to predict patient outcomes?<sup>o</sup>

## THE INTERVENTION

The study's design was deliberately simple. Twenty bariatric surgeons in a Michigan collaborative each submitted a representative video of a laparoscopic gastric bypass — a standardized, common procedure. The videos were rated blind by at least 10 peers on a structured skill scale derived from the surgical-education literature: tissue handling, time and motion, instrument knowledge, exposure, completion of the operation. Skill ratings ranged from 2.6 to 4.8, with the inter-rater reliability adequate to support the rank-ordering the study then took into the outcomes analysis.<sup>1</sup>

## HOW IT WORKED

The outcomes evidence linked the peer-rated skill scores to a Michigan registry of 10,343 patients treated by the same twenty surgeons. The headline finding: the bottom skill quartile had a higher complication rate than the top — 14.5% against 5.2% — and greater skill was associated with fewer reoperations, fewer readmissions, and fewer emergency department visits in the months after surgery. The skill score predicted the complication signal at the cohort scale. In the surgical-outcomes literature, this is the first large-registry evidence that operative skill, measured directly from operative video by peer rating, predicts the patient-outcome signal the institution actually cares about.<sup>2</sup>

## THE EVIDENCE

The authors call the findings preliminary and name the decisive qualification in the conclusion itself: "Although these findings are preliminary." The lowest skill quartile did fewer cases — 105 bariatric procedures a year against the top quartile's 240 — and operated more slowly. The skill-versus-volume confound is the central methodological hedge of the case. It is plausible — and consistent with the broader volume-outcome literature — that what the skill rating captured was partly the surgeon's accumulated operative experience, and that the rated skill is downstream of case volume rather than the other axis of the outcome. The published study does not separate the two; the editor may move the primary anchor away from 2.1 toward 1.1 (engineered vs. stated requirements) or 6.2 (operator-to-institution feedback channels) on that basis.<sup>3</sup>

## WHAT TRANSFERRED

What the case carries for the corpus is the evidence-architecture move of measuring the operator's actual technical capability against the outcome the institution cares about, with the volume-confound hedge intact. The case is adjacent to JIGSAWS (Case 31) at the surgical-skill-measurement layer — JIGSAWS provides the controlled-task instrumented evidence and Birkmeyer the naturalistic operative-video evidence, and the two together anchor the small-tier C3 conversation about measuring skill in surgery. The preliminary-findings language is part of what the case teaches.

*Greater skill was associated with fewer reoperations, readmissions, and emergency department visits; the lowest skill quartile also ran longer operations — 137 minutes against 98.*

EDITORS' SYNTHESIS OF BIRKMEYER ET AL. (2013) AND THE ACCOMPANYING EDITORIAL COMMENT.

## REFERENCES

1. Birkmeyer, J. D., Finks, J. F., O'Reilly, A., Oerline, M., Carlin, A. M., Nunn, A. R., Dimick, J., Banerjee, M., & Birkmeyer, N. J. O.
3. Vassiliou, M. C., Feldman, L. S., Andrew, C. G., et al. (2005). A global assessment tool for evaluation of intraoperative laparo-



for the Michigan Bariatric Surgery Collaborative. (2013). Surgical skill and complication rates after bariatric surgery. *New England Journal of Medicine*, 369(15):1434–1442. doi:10.1056/NEJMs1300625 — the case’s primary evaluation.

2. Birkmeyer, J. D., Stukel, T. A., Siewers, A. E., Goodney, P. P., Wennberg, D. E., & Lucas, F. L. (2003). Surgeon volume and operative mortality in the United States. *New England Journal of Medicine*, 349(22):2117–2127 — the volume–outcome literature the confound rests against.

scopic skills. *American Journal of Surgery*, 190(1):107–113 — the laparoscopic skill-rating literature this video assessment sits in.

4. Jacobs, D. O. (2013). Cut well, sew well, do well? *New England Journal of Medicine*, 369(15):1466–1467. doi:10.1056/NEJMe1309785 — the accompanying editorial.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Birkmeyer skill–outcomes study is the first large- registry evidence that operative skill, measured directly from operative video by blind peer rating, predicts the complication signal the institution cares about. The findings are preliminary; the skill-versus-volume confound is explicit; the multi-anchor (2.1 primary, 1.1 and 6.2 alternates) is the editor’s call.

LENS APPROACH

The surgical-skill peer-rating study is the small-tier measure-the-failure-mode-you-care-about case (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the measure-the-capability-against-the-outcome design move, and in Domain 3 (Human-System Collaboration) for the peer-rating workflow that makes the assessment scalable. Adjacent to JIGSAWS (Case 31) at the surgical-skill-measurement layer — JIGSAWS the controlled-task instrumented evidence, Birkmeyer the naturalistic operative-video evidence. The skill-versus-volume confound is the case.

| DOMAIN            | MODES ADDRESSED | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|-------------------|-----------------|---------------------|--------------|---------|-------|
| SURGERY           | H D K           | D1 D2 D3 D4 D5      | 5            | 2.1     | 4 • 3 |
| CLINICAL OUTCOMES |                 | Test and Evaluation |              |         |       |
| ASSESSMENT        |                 |                     |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the capability the assessment is built to measure (operative skill on a standard procedure) and the outcome the institution actually cares about (complication rate, reoperation, readmission), so the evidence architecture connects the measured variable to the institutional signal.
- Use blind peer rating on standardized video to make the assessment scalable and resistant to the gaming patterns of self-report or volume-weighted reputation.
- Pre-register the threats-to-validity table — volume confound, selection of submitted video, inter-rater reliability — so the published headline is reported alongside the qualifications the design carries.

IN OPERATION / AFTER

- Report the 14.5% bottom-quartile complication finding together with the volume-confound hedge; the preliminary-findings language is what makes the result interpretable for the next study.
- Design the follow-up analysis to separate skill from volume — for example, by matching on case volume within skill quartile, or by instrumenting the case-volume time-course — so the next round of evidence resolves the confound the present study leaves open.
- Carry the result into the surgical-education and credentialing conversation as evidence that peer-rated operative-video skill is a measurable, outcome-predictive variable, while preserving the qualifications the authors themselves attach.

## REFLECTION QUESTIONS

1. Identify a capability in your domain where the operator's actual technical capability is conventionally treated as too hard to measure at scale and therefore left as a hidden variable in the outcome equation. What would a Birkmeyer-style scalable, blind peer-rated measure look like?
2. The skill-versus-volume confound is the case. Design the follow-up analysis that would separate skill from volume in your context. What matching, what instrumentation of the volume time-course, what controlled-task companion would be required?
3. Birkmeyer is anchored at 2.1 with C3 and C1 alternates. Which anchor would you choose for your own use, and what does that choice say about which conversation in your domain the case is the evidence for?

**WHO BUILDS THIS** human factors & interaction design · systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · task analysis · design prototyping · knowledge bases · data pipelines

# Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units

**T** Training Gap **K** Knowledge & Institutional Memory **H** Human-System Interface

**DISCLOSURE** Institutional overlap: this case originates at an editor's institution (Johns Hopkins); no editor was personally involved. Anchored to the peer-reviewed PLOS One paper and the public JIGSAWS dataset, not institutional press.

**IMPACT** JHU's Language of Surgery project treated surgical motion as language — decomposing tasks into gesture and sub-gesture motion primitives — and released JIGSAWS, a public da Vinci kinematic/video/gesture/skill-rating dataset that became a standard benchmark; experts used fewer gestures (26.29 vs 31.30) and fewer gesture errors than novices on a suture-throw-and-surgeon's-knot task

## IN BRIEF

The Language of Surgery project at Johns Hopkins, led by Gregory Hager and a cross-departmental team spanning engineering, computer science, and surgery, treated surgical motion as language — decomposing tasks into gestures (the “surgesmes”) and sub-gesture motion primitives (the “dexemes”) fine enough to distinguish expert from novice. The project released JIGSAWS (JHU-ISI Gesture and Skill Assessment Working Set), a publicly available da Vinci surgical-robot dataset with synchronized kinematic, video, gesture-annotation, and skill-rating tracks for suturing, knot-tying, and needle-passing tasks. JIGSAWS became a standard benchmark in surgical skill-assessment and gesture-recognition research. Vedula et al. (PLOS One, 2016) applied the same decomposition to a purpose-collected set of 18 surgeons and showed that experts used fewer gestures (26.29 vs. 31.30 on a suture throw and surgeon's knot) and made fewer gesture errors than novices (1.00 vs. 2.84), with quantifiable sub-task-level differences. The case establishes that surgical skill is decomposable and machine-measurable; the honest open question — preserved here verbatim — is whether automated motion-level feedback accelerates trainee skill acquisition or improves patient outcomes. The dataset enables the question more than it answers it. The case pairs directly with Case 30 (Birkmeyer's video-rated bariatric-surgical-skill outcome study), which establishes that skill matters; together they form the skill-measurement spine the corpus needed.

## BACKGROUND

Surgical training has carried, for decades, a structural measurement gap: skill at the trainee-versus-attending level is universally acknowledged to matter, and is universally measured by rater judgment — case logs, milestone evaluations, OSATS scores — that resolves the construct only at the summary level. What such measures cannot resolve is the unit of

skill: the specific motion sequence, the gesture choice, the sub-gesture smoothness, the corrective re-engagement after a near-error. The Language of Surgery project began from the position that skill could be made measurable at the gesture level if surgical activity were decomposed the way language is decomposed into phonemes and words.<sup>9</sup>

#### THE INTERVENTION

The intervention is dataset-and-method, jointly. The team instrumented the da Vinci surgical robot to capture synchronized kinematic traces of the arms, stereo video of the surgical field, and frame-level human annotation of which gesture was being performed. Onto that base they layered expert skill ratings task by task. The technical decomposition was structured as a hierarchy: a task (e.g., knot-tying) consists of a sequence of gestures or surgemes (e.g., “reach for needle”, “position needle”, “drive needle”), each of which is composed of motion primitives or dexemes at the kinematic-trajectory level. The team released the resulting corpus as JIGSAWS in open form so the broader research community could test methods against a common benchmark, rather than each lab building its own private dataset.<sup>1</sup>

#### HOW IT WORKED

Vedula et al. (PLOS One, 2016) applied the decomposition to a separately collected set — 18 surgeons, four expert attendings and fourteen trainees, performing a single suture throw and a 2–1 surgeon’s knot — and reported that experts used fewer gestures (26.29) than novices (31.30), and that novices made more gesture errors (2.84 against 1.00). The differences were not at the gross-task-outcome level — both groups completed the task — but at the gesture-composition level the decomposition made visible. Subsequent work on JIGSAWS by other groups developed automated gesture-recognition methods, automated skill-rating classifiers, and motion-primitive analyses; the dataset’s role was to make these results comparable across labs.<sup>2</sup>

#### THE EVIDENCE

The honest open question survives into the case verbatim. The project demonstrates that surgical skill is decomposable and machine-measurable at the gesture level. What remains open is whether automated motion-level feedback delivered to trainees actually accelerates skill acquisition, and whether gesture-level skill differences translate into patient-outcome differences for the trainees once they reach the operating room. JIGSAWS enables both questions to be asked rigorously; it does not, by itself, answer either. The case’s learning-engineering content is the construction of the evidence architecture — the measurable unit, the open dataset, the cross-lab benchmark — that makes downstream measurement possible. The home-institution disclosure under the title is the standing safeguard against boosterism.<sup>3</sup>

#### WHAT TRANSFERRED

In pair with Case 30 (Birkmeyer et al.’s video-rated bariatric surgical-skill study), the case completes a skill-measurement pair the corpus had needed: Birkmeyer shows that rated skill predicts patient outcome at scale (skill matters), and this case shows that skill is decomposable into machine-measurable units (skill is engineerable). Drafted together they operationalize the new framing the v2 induced framework proposes — that capability is engineerable when the unit of capability is named and the measurement instrument follows.

JIGSAWS’s continued use as a benchmark, more than a decade after release, is the field-scale evidence that the decomposition was the right resolution for the question.

*Surgical motion is language. The unit of skill is the gesture; the corpus is the dataset; the benchmark is the field’s continued use of it more than a decade after release.*

EDITORS’ SYNTHESIS OF VEDULA ET AL. (2016) AND THE JIGSAWS RELEASE.

REFERENCES

1. Vedula, Malpani, Tao, Chen, Gao, Poddar, Ahmidi, Paxton, Vidal, Khudanpur, Hager, & Chen (2016), “Analysis of the Structure of Surgical Activity for a Suturing and Knot-Tying Task,” PLOS One 11(3):e0149174, doi:10.1371/journal.pone.0149174.

2. Gao, Vedula, Reiley, Ahmidi, Varadarajan, Lin, Tao, Zappella, Bejar, Yuh, Chen, Vidal, Khudanpur, & Hager (2014), “JHU-ISI Gesture and Skill Assessment Working Set (JIGSAWS): A Surgical Activity Dataset for Human Motion Modeling,” MICCAI Workshop — JIGSAWS dataset release paper.

3. Reiley, Lin, Yuh, & Hager (2011), “Review of methods for objective surgical skill evaluation,” Surgical Endoscopy 25(2):356–366 — situates the decomposition within the broader skill-assessment literature.

4. Birkmeyer, Finks, O’Reilly, et al. (2013), NEJM — the paired skill-matters study (Case 30).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Language of Surgery / JIGSAWS is the corpus’s worked example of decomposing a tacit capability — surgical skill — into machine-measurable units, and releasing the measurement infrastructure openly so the field can build on it. The decomposition is established; the downstream question — whether motion-level feedback accelerates skill or improves patient outcomes — is open. The case enables the question rather than answering it.

LENS APPROACH

Language of Surgery is the skill-decomposition case (induced 2.1; LENS D4/PT5) — Domain 4 for construct-resolution; Domain 1 for task-decomposition. Pair with Case 30 (Birkmeyer) and Case 27. Home-institution disclosure under the title.

| DOMAIN                | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|-----------------------|-----------------|-----------------------|--------------|---------|-------|
| SURGERY               | T K H           | D1 D2 D3 D4 D5        | 5            | 2.1     | 2 · 4 |
| SKILL ASSESSMENT      | →               | → Test and Evaluation |              |         |       |
| HUMAN-MOTION ANALYSIS |                 |                       |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Decompose the capability into units fine enough to distinguish expert from novice but coarse enough to be reliably annotated; the surgeme / dexe hierarchy is the worked example of the trade-off.
- ▶ Release the dataset openly with synchronized signal tracks (kinematic, video, annotation, rating) so that downstream methods can be compared on a common benchmark rather than each lab evaluating itself.
- ▶ Pair the dataset with a baseline analysis paper that demonstrates the decomposition resolves real expert/ novice differences; the Vedula et al. analysis is what makes JIGSAWS more than a data dump.

IN OPERATION / AFTER

- ▶ Carry the open question explicitly — whether automated motion-level feedback accelerates skill acquisition is not what the dataset establishes — so downstream researchers and trainees see the gap as a research target, not a settled claim.
- ▶ Preserve the home-institution affiliation disclosure in any communication; the standing language anchors the case to the peer-reviewed paper and the public dataset, not to institutional press.
- ▶ Treat the cross-lab benchmark adoption as the case’s strongest evidence that the decomposition was correct at the resolution chosen; the field’s continued

use of JIGSAWS more than a decade after release is itself the test.

### REFLECTION QUESTIONS

1. Identify a tacit capability in your domain that is currently rated at the summary level. What is the unit of the capability — the gesture-equivalent — at which the decomposition would resolve expert/novice differences, and what signal tracks would the dataset need to synchronize?
2. Specify the open question your decomposition would not answer on its own. JIGSAWS does not establish that motion-level feedback accelerates skill; it enables the question. What is the analog in your context — the question the dataset enables but does not close?
3. The case's evidence of correctness is field-scale adoption of the benchmark a decade after release. What is the publication, release, and open-license strategy that would let your decomposition be tested by labs that have no stake in the original design?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · usability testing · cognitive walkthroughs

# Annual-Screening UI Redesign at University of Missouri Health Care

**H** Human-System Interface **N** Normalization of Deviance

**IMPACT** A site-level redesign of the EHR interface for four ambulatory annual screens (advance directives, depression, fall risk, alcohol/drug misuse) cut task time 8.3 to 6.0 seconds and errors 0.83 to 0.10 per task, raised the System Usability Scale from 80.8 to 96.9, and turned four flat screening-rate trends upward — fall risk 53.5 to 83.2 percent over nine months

## IN BRIEF

A multidisciplinary team at University of Missouri Health Care redesigned the EHR interface nurses and medical assistants use to prompt and perform annual screening — advance directives, depression, fall risk, alcohol and drug misuse. The interface was not the vendor's product but the health system's own configuration of it, and the redesign was a site-level quality-improvement project the vendor had no direct role in. Summative usability testing with twelve staff showed the redesigned interface cut task time for identifying overdue screens from 8.3 to 6.0 seconds and errors from 0.83 to 0.10 per task, and raised the System Usability Scale from 80.8 to 96.9. In an interrupted time series across eight months before and nine months after go-live, all four screening rates — flat or drifting down beforehand — turned upward. It is a small-tier intervention case for cue design as a capability deliverable, with both human-factors and clinical-process outcomes in the same report. The corpus has long needed a small-tier C3 positive example to set against the interface failures already documented at the big tier (Therac-25, CPOE/EHR adoption, the Helios pattern). The evidence base is a single-institution QI study published peer-reviewed in *Applied Clinical Informatics* (2020). Future validation will continue as the downstream clinical-outcome literature on screening-rate gains matures.

## BACKGROUND

Ambulatory annual screening — advance directives, depression, fall risk, alcohol and drug misuse — is the kind of guideline-recommended care that is easy to declare and hard to land. Three of the four are Merit-based Incentive Payment System quality measures and the fourth is defined by the state Medicaid Primary Care Health Home program, so the yardstick was written outside the project. The cue lives in the EHR; the action lives in a time-pressured encounter; and the gap between prompt and completion is where most screening programs lose their numbers. What started this project was not a metric but a complaint: two of the authors kept meeting nurses and medical assistants who called the annual-assessment screen “not user friendly” and said it took “too many clicks.” A heuristic analysis against Nielsen's usability principles found why. The interface required the same date to be entered

twice, did not validate dates, made the nurse calculate in her head whether a screen was due, ignored patient-specific factors such as age and problem list that decide whether it is due at all, and gave no way to record that a screen was due but could not be done. The screen had been in place for several years.<sup>0</sup>

#### THE INTERVENTION

The redesign was multidisciplinary by construction. Health-system physicians and nurse informaticists led it, with input from nurses, medical assistants, administrators and clinical quality-improvement specialists, and with vendor-employed system architects on the team for the site-specific build; the vendor itself had no direct role. Over eight months of iterative, user-centered work the team removed one of the duplicate date fields and autopopulated the other, showed the prior screening date read-only, and — the load-bearing move — resolved the conditional logic the vendor platform could not express by computing it in a separate program at runtime and passing a hidden binary variable back in. Patient age, current problems and date of last screen now produced an unambiguous color-coded indication that a screen was due, and four new controls let staff record why one was being deferred. The design move is the one the induced framework flags as the C3 deliverable: put the judgment the interface was asking a nurse to make in her head into the interface itself. Note what the redesign did **not** add — the HIMSS write-up of this project describes clinical decision support, but the peer-reviewed report describes no decision-support module, no ordering integration, and no change to the screening instruments themselves beyond defaulting their date field.<sup>1</sup>

#### HOW IT WORKED

The reported outcomes cross two layers. At the usability layer, summative testing with twelve nurses and medical assistants found that identifying which screens were overdue took 6.0 seconds on the redesigned interface against 8.3 on the original ( $p = 0.001$ ), with 0.10 errors per task against 0.83 ( $p < 0.0001$ ), and the System Usability Scale rose from 80.8 to 96.9 ( $p = 0.021$ ). One measure moved the wrong way: the redesigned interface took **more** clicks, 1.17 against 0.10 ( $p = 0.007$ ), because the added deferral controls pushed the bottom of the form below the fold and most users had to scroll. The authors report it as an unexpected negative finding and an opening for the next iteration. At the process layer, an interrupted time series over eight months before and nine months after the June 2019 go-live found all four screening rates flat or drifting slightly downward beforehand and rising afterward: the monthly slope changed by +0.44 points for advance directives ( $p = 0.017$ ), +2.97 for depression, +3.52 for fall risk and +3.07 for alcohol and drug misuse (all  $p < 0.0001$ ). Between the last month before go-live and the last month observed, the twelve-month rolling rates ran 75.1 to 80.1 percent for advance directives, 38.0 to 61.5 for depression, 53.5 to 83.2 for fall risk and 49.3 to 74.6 for alcohol and drug misuse. The denominators are not the same population: advance-directive and fall-risk screening are measured on roughly 13,000 eligible patients aged 65 and over, depression on about 63,000 aged 12 and over, alcohol and drug misuse on about 65,000 adults — and the count of eligible patients stayed roughly constant throughout. That second layer is what makes the case a C3 small-tier intervention rather than a usability study: the interface gain translated into the downstream behavior, at least over the reported observation window.<sup>2</sup>



## THE EVIDENCE

Three things bound what the case can claim. The usability test was sized to industry practice for finding defects, not to detect differences: twelve participants, and the authors say plainly that it “was not powered.” Go-live carried an education component — an email to staff and a live demonstration to nursing representatives who then briefed their colleagues — so some of the screening-rate movement may be education rather than interface; the authors argue that a one-off, unrepeated briefing is unlikely to explain a rise sustained over nine months, which is an argument, not a control. And an interrupted time series is quasi-experimental: there is no randomized comparison, the rates are twelve-month rolling averages that lag and damp the underlying change, no seasonal adjustment was made, and the series has fewer points than some authors recommend. What the regulatory frame supplies is the reason this kind of project is rare rather than routine: the Office of the National Coordinator’s certification criteria require user-centered design and summative usability testing of **vendors**, and nothing equivalent applies to the site-specific configuration layer where this interface actually lived — the gap Ratwani and colleagues had in view when they called for evaluation across the whole lifecycle, “from initial product design to local implementation followed by periodic reassessment.” Replication at other institutions and durability across EHR upgrades and personnel rotation have not been independently audited. Future validation will continue as the institution and others publish follow-on screening-rate data and as the relationship between screening completion and downstream patient outcomes is tracked.<sup>3</sup>

## WHAT TRANSFERRED

What the case teaches at the LENS layer is that the C3 interface-failure pattern documented at the big tier is not a cosmic constraint — it is redressable by deliberate design at the small tier, by the site rather than the vendor, on a platform whose own limits had to be worked around. The capability deliverable is the redesigned cue plus the measurement of the downstream process the cue exists to drive. The case is the missing positive example for induced 3.1 at the small tier and a paired teaching companion for the failures already in the corpus.

*The C3 interface failure mode is not a cosmic constraint. It is redressable by deliberate design — and the design authority that mattered sat at the site, not at the vendor.*

EDITORS’ SYNTHESIS.

## REFERENCES

1. Pierce RP, Eskridge BR, Rehard L, Ross B, Day MA, Belden JL (2020), “The Effect of Electronic Health Record Usability Redesign on Annual Screening Rates in an Ambulatory Setting,” *Applied Clinical Informatics* 11(4):580–588, doi:10.1055/s-0040-1715828 — the source of record for every figure in this case. HIMSS carries a secondary write-up of the project; it describes the redesign as adding clinical decision support, which the peer-reviewed report does not, and this case follows the report.
2. Hussain MI, Reynolds TL, Zheng K (2019), “Medication safety alert fatigue may be reduced via interaction design and clinical role tailoring: a systematic review,” *JAMIA* 26(10):1141–1149, doi:10.1093/jamia/ocz095 — adjacent evidence that interaction-design changes move alert behavior. Note its outcome is prescriber acceptance of medication safety alerts, not screening completion; it is background, and no figure in this case rests on it.
3. Ratwani RM, Reider J, Singh H (2019), “A decade of health information technology usability challenges and the path forward,” *JAMA* 321(8):743–744, doi:10.1001/jama.2019.0161 — the call for lifecycle evaluation “from initial product design to local implementation followed by periodic reassessment,” and the framing for the site-configuration gap in federal certification requirements.
4. Middleton et al. (2013), “Enhancing patient safety and quality of care by improving the usability of electronic health record systems,” *JAMIA* 20(e1):e2–e8 — the framing peer-reviewed paper on EHR-usability-as-safety.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The UMHC redesign is the small-tier C3 positive example the corpus needed: cue design as a capability deliverable, with both usability and downstream screening-rate gains in the same project. The evidence is a single-institution quality-improvement study, peer-reviewed in Applied Clinical Informatics (2020): a twelve-participant usability test the authors say was not powered, and a quasi-experimental interrupted time series carrying an education co-intervention the authors argue against but did not control. Magnitudes await independent replication and durability tracking. Future validation ongoing.

LENS APPROACH

UMHC is the C3 small-tier intervention case (induced 3.1; LENS D4/PT5 — the anchor is deliberately the evaluation domain rather than the design domain, because what the case contributes is the outcome evidence for a cue intervention; induced 3.1 carries the design linkage). LENS uses it in Domain 4 (Test and Evaluation) for the LEO-4 deliverable that cue redesign must show its downstream process effect — and show it with its bounds attached, including the measure that moved the wrong way — and in Domain 1 (Systems Analysis) for LEO-1: the heuristic analysis of what the interface was asking the nurse to do was the precondition for the redesign. The case is the paired positive example for the big-tier C3 failures (Therac-25, CPOE/EHR adoption, Helios) the corpus already documents.

| DOMAIN           | Modes Addressed               | LENS Competency                                                       | Problem Type | Induced | LEO   |
|------------------|-------------------------------|-----------------------------------------------------------------------|--------------|---------|-------|
| HEALTHCARE       | → <div>H</div> <div>N</div> → | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 5            | 3.1     | 4 · 1 |
| EHR USABILITY    |                               | Test and Evaluation                                                   |              |         |       |
| NURSING WORKFLOW |                               |                                                                       |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Frame the screening problem as cue-and-alert design from the start; do not separate usability from the clinical-process outcome the cue exists to drive.
- ▶ Seat the people who will use the screen — here nurses and medical assistants, not physicians — on the redesign team alongside the informaticists, and treat the vendor's platform limits as something to engineer around rather than accept.
- ▶ Specify the downstream process metric (screening-rate completion) before the redesign ships, so the usability-layer and process-layer outcomes are measured against the same scope — and prefer a metric someone else defined, as three of these four were.

IN OPERATION / AFTER

- ▶ Treat the QI report honestly: a single-institution peer-reviewed study is a small-tier intervention, not a settled magnitude. Report the measure that moved the wrong way — here the extra click — and name the co-interventions you could not control, as this team named its go-live briefing.
- ▶ Track durability — across EHR upgrades, personnel rotation, and downstream clinical outcomes — as a separate post-deployment commitment, not a footnote to the implementation report.
- ▶ Pair with the big-tier C3 failures already in the corpus when teaching; the failure-and-intervention pair is the teaching artifact, not either alone.

## REFLECTION QUESTIONS

1. Identify a screening or recommended-care prompt in your EHR. What proportion of prompts result in the completed action? Where does the cue-action gap sit, and what is the redesign that would close it?
2. Specify the downstream process metric you would commit to before any usability redesign ships. The UMHC case is teachable because it reported the screening-rate outcome, not only the SUS gain. What would be the equivalent in your context?
3. The case is a single-institution peer-reviewed QI study. What is the minimum independent replication evidence you would require before treating the reported magnitudes as a basis for an institutional investment in your own setting?

**WHO BUILDS THIS** human factors & interaction design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · incident analysis · safety dashboards

## CASE 33

HEALTHCARE

CLINICAL DECISION SUPPORT

EHR DESIGN

NURSING WORKFLOW

2019 – 2024

# Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal

**T** Training Gap   **D** Designed Out   **N** Normalization of Deviance

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** Structured EHR alert redesign — fewer alerts, severity reclassification, interruptive-to-passive conversion, role-tailoring — reduced alert burden in published systematic-review and quality-improvement evidence; the 2024 case studies report alert-rate reduction, and the COVID-precautions project is the one that also reports the clinical outcome rising alongside it

## IN BRIEF

Alert fatigue is the structural failure mode the C3 thread names at the small tier: an EHR that fires so many alerts the actionable ones are lost in the noise. The 2019 JAMIA systematic review by Hussain and colleagues aggregates the evidence that structured redesign — interaction design changes and clinical-role tailoring — can raise the rate at which prescribers accept an alert's advice; a second JAMIA review, Van Dort and colleagues in 2021, found that five of the six hospital optimization programs that evaluated their own changes reported a reduction in alert rate. Two 2024 quality-improvement publications extend the pattern with named maneuvers: a nursing-workflow redesign of four high-firing, low-action alerts using quantitative alert-firing analysis, empathy mapping, and iterative user feedback; and the replacement of an interruptive COVID-precautions alert with passive clinical decision support, targeting both alert burden and the timeliness of precautions orders. The case is the small-tier intervention companion to a C3.2 failure thread that v1 left almost entirely populated by failures (Uber ATG, Robodebt, Northeast Blackout, UK Post Office, Tesla Autopilot). The evidence tier is mixed: the systematic review is peer-reviewed; the per-site QI projects are practice-tier publications. Future validation will continue on whether the redesigns survive EHR upgrades and personnel rotation.

## BACKGROUND

Monitoring under alert burden is the unsupportable role the C3.2 induced sub-competency names — the operator asked to keep attention on a stream of low-signal alerts and to spot the consequential one in real time. In the EHR setting the burden becomes structural: high-firing, low-action alerts train the clinician to dismiss alerts as the default, which is where the actionable alert is missed. Alert fatigue is the failure mode; the question for C3 is whether deliberate redesign can reduce the burden without cutting the safety signal.<sup>o</sup>

## THE INTERVENTION

The 2019 JAMIA systematic review by Hussain and colleagues aggregates the published evidence on EHR alert optimization. The headline finding is narrower than it is often reported to be: across the thirty-nine studies the review included, interruptive modal dialogs were accepted least often, and of the alternatives examined only role tailoring appeared to increase prescriber acceptance. The burden side of the ledger comes from a companion JAMIA review — Van Dort and colleagues, 2021 — which found that of the eight published hospital optimization programs, six evaluated their own changes and five of those reported a reduction in alert rate. The mechanism is not a single intervention but a family of moves — severity reclassification, conversion of interruptive alerts to passive decision support, role-based tailoring so the alert reaches the clinician who can act on it, and removal of alerts whose firing-to-action ratio shows the alert is no longer working.<sup>1</sup>

## HOW IT WORKED

The 2024 quality-improvement publications instantiate the moves. One project redesigned four high-firing, low-action alerts in the nursing workflow using mixed methods — quantitative analysis of firing data, empathy mapping of the nursing experience, and iterative user feedback as the redesign was refined. A second project replaced an interruptive COVID-precautions alert with passive CDS, with dual outcomes: reduce alert burden and improve the timeliness of precautions orders. Together the projects show the redesign-pattern is operable at the per-alert level and that the evaluation can report both halves of the trade-off the redesign exists to manage.<sup>2</sup>

## THE EVIDENCE

The evidence-tier flag matters. The 2019 systematic review is peer-reviewed, and the 2024 Applied Clinical Informatics paper sits in a peer-reviewed informatics journal. The practical maneuver of treating the per-site QI projects as generalizable, though, rests on practice-synthesis logic across the body of work rather than on a single multi-site randomized evaluation. Magnitudes vary by site, EHR vendor, and alert category; the redesigns must be re-verified after EHR upgrades and personnel rotation. Future validation will continue as the optimization-study literature consolidates.<sup>3</sup>

## WHAT TRANSFERRED

The teaching point pairs with Case 32 and with the v1 C3.2 failure thread. C3.2 in v1 is entirely failures; this case is the small-tier intervention that demonstrates the failure mode is redressable by design. The capability deliverable is the redesigned alert architecture itself — severity tiers, interruptive-vs-passive decisions, role-tailoring rules, and an ongoing measurement loop on firing-to-action ratios — not a one-time clean-up. The new LEO around delegation with revocation applies here: when automated decision support oversight is delegated to the bedside, the redesign discipline is part of the delegation.<sup>4</sup>

*Alert fatigue is the failure mode the alert architecture trains. The redesign discipline is part of the delegation.*

EDITORS' SYNTHESIS OF HUSSAIN ET AL. (2019) AND THE 2024 QI LITERATURE.

## REFERENCES

1. Ancker et al. (2017), “Effects of workload, work complexity, and repeated alerts on alert fatigue in a clinical decision support system,” BMC Medical Informatics and Decision Making 17:36 — adjacent measurement evidence.

2. Hussain, M. I., Reynolds, T. L., & Zheng, K. (2019), “Medication safety alert fatigue may be reduced via interaction design and clinical role tailoring: a systematic review,” JAMIA 26(10):1141–1149, doi:10.1093/jamia/ocz095 — thirty-nine studies; interruptive modals accepted least often; role tailoring the only alternative that appeared to increase prescriber acceptance.

3. Fallon A, Haralambides K, Mazzillo J, Gleber C (2024), “Addressing Alert Fatigue by Replacing a Burdensome Interruptive Alert with Passive Clinical Decision Support,” Applied Clinical Informatics 15(1):101–110, doi:10.1055/a-2226-8144 — 80% cut in alert firings; COVID precautions ordered rose from 23% to 61%.

4. Van Dort, B. A., Zheng, W. Y., Sundar, V., & Baysari, M. T. (2021), “Optimizing clinical decision support alerts in electronic medical records: a systematic review of reported strategies adopted by hospitals,” JAMIA 28(1):177–183, doi:10.1093/jamia/ocaa279 — eight published optimization programs; six evaluated their changes and five of those reported a reduction in alert rate.

5. Office of the National Coordinator for Health IT, SAFER Guides on CDS — practitioner-tier guidance the redesigns instantiate.

6. Patterson E (2024), “Navigating Alert Fatigue: A Case Study in Electronic Health Record Alert Design Optimization,” Studies in Health Technology and Informatics 315:447–451, doi:10.3233/SHTI240188 — nursing-workflow QI redesign of four high-firing alerts; all four removed, 877 unactionable disruptive nursing hours released.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The 2019 JAMIA review plus the 2024 QI projects are the small-tier C3.2 intervention companion the corpus needed — the failure thread (Uber ATG, Robodebt, UK Post Office, Tesla) is redressable by deliberate alert redesign. The systematic review is peer-reviewed; the per-site QI publications are practice-tier; magnitudes and durability open.

LENS APPROACH

Alert-fatigue redesign is the C3.2 small-tier intervention case (induced 3.1 and 3.2; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for LEO-4 — the redesign must report the safety signal alongside the alert-rate cut — and in Domain 3 (Human-System Collaboration) for LEO-3 oversight of automated decision support delegated to the bedside, with the LEO on delegation with revocation explicit. Pair with Case 32 for the small-tier C3 thread.

| DOMAIN                    | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|---------------------------|-----------------|-----------------------|--------------|---------|-------|
| HEALTHCARE                | T D N           | D1 D2 D3 D4 D5        | 5            | 3.1     | 4 · 3 |
| CLINICAL DECISION SUPPORT | →               | → Test and Evaluation |              |         |       |
| EHR DESIGN                |                 |                       |              |         |       |
| NURSING WORKFLOW          |                 |                       |              |         |       |

APPROACHES TO CONSIDER

- DURING DEVELOPMENT

▶ Instrument every consequential alert with its firing-to-action ratio; the metric that names the failure mode must be reported alongside the redesign.

▶ Decide the interruptive-vs-passive call deliberately for each alert; the conversion is the C3 move that reduces burden when the signal does not require the interruption to be actionable.

▶ Use mixed methods (firing data + clinician experience + iterative feedback) to design the per-alert redesign; quantitative alone misses why the alert is being dismissed.
- IN OPERATION / AFTER

▶ Track the safety signal the alert exists to protect as a separate post-redesign outcome; an alert-rate cut that lost the signal is a failure, not a win.

▶ Re-audit the redesigned alerts after EHR upgrades and personnel rotation; the redesign is not a one-time clean-up and the v1 thread shows the failure mode returns.

▶ Treat the per-site QI publications as practice-synthesis evidence when generalizing; the systematic review supports the pattern, the magnitudes require local replication.

## REFLECTION QUESTIONS

1. Identify an EHR alert in your context with a high firing-to-action ratio. Which of the redesign moves (severity reclassification, interruptive-to-passive conversion, role-tailoring, removal) would you apply, and what would the measured outcome be on both alert burden and the safety signal?
2. Specify the mixed-methods design (firing data + clinician experience + iterative feedback) you would use to make a per-alert redesign decision. What does the empathy-mapping step add beyond quantitative firing analysis alone?
3. The per-site QI publications are practice-synthesis tier. What is the minimum cross-site or randomized evidence you would require before generalizing the magnitudes from the 2024 reports to your own setting?

**WHO BUILDS THIS** learning science & instructional design · systems & human-factors engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · task analysis · design prototyping · incident analysis · safety dashboards

# COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case

**T** Training Gap   **K** Knowledge & Institutional Memory   **D** Designed Out

**IMPACT** Boussina et al. NPJ Digital Medicine 2024 before-and-after quasi-experimental evaluation of the COMPOSER (CONformal Multidimensional Prediction Of SEpsis Risk) deep-learning model at UC San Diego Health; reported a 1.9 percentage-point absolute decrease (17% relative) in in-hospital sepsis mortality and a 5.0 percentage-point absolute increase in sepsis bundle compliance during the deployment period, evaluated against the pre-deployment baseline period within the same two-emergency-department health system

## IN BRIEF

Aaron Boussina, Shamim Nemati, and colleagues at UC San Diego Health published the COMPOSER (CONformal Multidimensional Prediction Of SEpsis Risk) prospective implementation study in NPJ Digital Medicine in early 2024. The model is a deep-learning sepsis-risk prediction system that uses conformal prediction to handle uncertainty — alerting only when the model's calibrated confidence threshold is met and abstaining when the input is sufficiently out-of-distribution. The deployment evaluation reported a 1.9 percentage-point absolute decrease in in-hospital sepsis mortality and a 5.0 percentage-point absolute increase in sepsis bundle compliance during the deployment period at two UC San Diego emergency departments, evaluated against the pre-deployment baseline period within the same health system. The case completes the AI-delegation typology in sepsis at three deployments — Case 20 (TREWS at Johns Hopkins), Case 5 (Epic Sepsis Score across multiple health systems), and COMPOSER at UCSD. The honest hedges from the source are binding: the study is not RCT-grade; every emergency department in the system ran the model, so no contemporaneous comparison arm exists; the site is a single academic center with a standing interest in sepsis and clinical informatics; and the authors did not evaluate what the alert did to patients who turned out not to have sepsis — antibiotic overuse and its costs are unmeasured. Pair with Case 20, Case 5, and Case 35 (Radiology AI Miscalibration).

## BACKGROUND

The contemporary clinical-AI sepsis-prediction literature has converged on a small number of well-evidenced deployments. Case 20 documents TREWS — the Targeted Real-Time Early Warning System at Johns Hopkins — with its prospective evaluation showing mortality reduction associated with prompt clinician response to alerts. Case 5 documents the Epic Sepsis Score deployed across multiple U.S. health systems, with external validation by Wong



and colleagues finding substantially lower sensitivity and higher false-positive rates than the vendor's validation data implied. COMPOSER is the third deployment in the typology and the second prospective-positive case. The deployment is at two UC San Diego emergency departments, integrated into the electronic health record as a nurse-facing Best Practice Advisory in the emergency-department workflow.<sup>9</sup>

#### THE INTERVENTION

The methodological feature that distinguishes COMPOSER is the conformal-prediction framework. Conformal prediction is a statistical method for producing calibrated uncertainty intervals around individual predictions: rather than producing a single risk score, the model produces a prediction together with a measure of how confident the model is in that prediction given the input. The operational implication is that the model can abstain from alerting when the input is sufficiently out-of-distribution — when the model's calibrated confidence does not support a clinically actionable alert. The abstention structure is what handles one of the failure modes of Case 5 (Epic Sepsis), where the model alerted at high rates against patient populations it had not been adequately calibrated against. COMPOSER's abstention structure is one mechanism for refusing to alert in those circumstances.<sup>1</sup>

#### HOW IT WORKED

The deployment outcomes Boussina and colleagues report are the intervention-evidence. Across the deployment period at the two UC San Diego emergency departments — 6,217 adult septic patients between January 2021 and April 2023, 705 pre-intervention days against 145 post-intervention days — in-hospital sepsis mortality decreased by 1.9 percentage points in absolute terms, and sepsis bundle compliance increased by 5.0 percentage points, both evaluated against the pre-deployment baseline period within the same health system. The evaluation is prospective implementation, not RCT — the comparison is against the same hospital sites' historical baseline rather than against a contemporaneous randomized control arm — and the authors are explicit about the limitation. The authors record no other quality-improvement initiative running at UCSD's sepsis-care pathway over the same period, so the reduction is not confounded by concurrent process change; what limits the inference is the absence of a contemporaneous control arm, not a tangle of co-occurring reforms.<sup>2</sup>

#### THE EVIDENCE

The case pairs with Case 20 (TREWS) for the prospective-positive sepsis-prediction thread: both deployments report mortality benefits, both are integrated into specific health-system workflows, and the comparison between TREWS and COMPOSER turns on the alerting architecture and on the role of the surrounding process. Pair with Case 5 (Epic Sepsis Score) for the external-validation-revealed-gaps thread: COMPOSER's abstention structure is one response to the calibration-gap failure mode that the Epic Sepsis evidence base named. Pair with Case 35 (Radiology AI Miscalibration) for the broader medical-AI miscalibration-in-deployment thread. The three sepsis cases together — Case 20, Case 5, COMPOSER — define the AI-delegation typology in sepsis: prospective-positive with workflow integration (TREWS), externally validated and found wanting against vendor claims (Epic Sepsis), and prospective-positive with conformal-prediction abstention (COMPOSER).<sup>3</sup>

## WHAT TRANSFERRED

The hedges the case carries are binding and preserved in the prose. The deployment is prospective implementation, not an RCT; the mortality reduction and bundle-compliance gain are evaluated against the same health system's historical baseline, and every emergency department in the system ran the model, so no contemporaneous comparison arm exists against which the reduction could be checked. What the authors do not claim is a tangle of co-occurring reforms: they record no other quality-improvement initiative in the same period, which is what leaves the missing control arm, rather than confounding by concurrent process change, as the binding limitation. Two further hedges survive verbatim — the site is a single academic center with a major interest in sepsis and clinical informatics, so external validation in community hospitals and less IT-rich settings is outstanding; and the study did not evaluate the alert's effect on patients who turned out not to have sepsis, leaving inappropriate antibiotic use and its costs unmeasured. The conformal-prediction abstention structure is the methodological contribution the case anchors, and the cue-and-alert design LEO is anchored at the deployment seam where abstention is a clinical-workflow capability: the model that can refuse to alert when its calibrated confidence does not support an alert is the deployment artifact the case names as the decisive one.

*The study is before-and-after, not randomized, and every emergency department in the system ran the model, so there is no contemporaneous comparison arm; the conformal-prediction abstention is the methodological contribution the case anchors.*

EDITORS' SYNTHESIS OF BOUSSINA ET AL. (2024) NPJ DIGITAL MEDICINE AND THE SURROUNDING SEPSIS-PREDICTION DEPLOYMENT LITERATURE.

## REFERENCES

1. Boussina, A., Shashikumar, S. P., Malhotra, A., Owens, R. L., El-Kareh, R., Longhurst, C. A., Quintero, K., et al. (2024), "Impact of a deep learning sepsis prediction model on quality of care and survival," NPJ Digital Medicine 7:14, doi:10.1038/s41746-023-00986-6.
2. Shashikumar, S. P., Wardi, G., Malhotra, A., & Nemati, S. (2021), "Artificial intelligence sepsis prediction algorithm learns to say 'I don't know,'" NPJ Digital Medicine 4:134 — the methodological-foundation paper for the conformal-prediction abstention structure.
3. Wong, A., Otlis, E., Donnelly, J. P., Krumm, A., McCullough, J., De-Troyer-Cooley, O., et al. (2021), "External Validation of a Widely Implemented Proprietary Sepsis Prediction Model in Hospitalized Patients," JAMA Internal Medicine 181(8):1065–1070 — the external-validation paper on Epic Sepsis (Case 5).
4. Adams, R., Henry, K. E., Sridharan, A., Soleimani, H., Zhan, A., Rawat, N., Johnson, L., et al. (2022), "Prospective, multi-site study of patient outcomes after implementation of the TREWS machine learning-based early warning system for sepsis," Nature Medicine 28:1455–1460 — the prospective-positive TREWS deployment paper (Case 20).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

COMPOSER is the third clinical-AI sepsis case in the delegation typology — prospective-positive with conformal-prediction abstention. The deployment was associated with a 1.9 percentage-point fall in in-hospital sepsis mortality and a 5.0 percentage-point rise in sepsis bundle compliance; the decisive hedges are that the evaluation is before-and-after rather than randomized, that every ED in the system ran the model so no comparison arm exists, and that the harms to non-sepsis patients went unmeasured. The abstention structure is the methodological contribution the case anchors.

### LENS APPROACH

COMPOSER is the cue-and-alert-design-as-capability-deliverable case at clinical-deployment scale (induced 3.1; LENS D4+D3/PT6; LEO-4 and LEO-3). LENS uses it in Domain 4 (Test and Evaluation) for the

prospective–implementation– evaluation discipline and in Domain 3 (Human–System Collaboration) for the abstention–as–clinical–workflow– capability anchor. Pair with Case 20 (TREWS), Case 5 (Epic Sepsis), and Case 35 (Radiology AI Miscalibration). The three sepsis cases together — Case 20, Case 5, Case 34 — define the AI–delegation typology in sepsis at the deployment–evidence level.

| DOMAIN                    | MODES ADDRESSED | LENS COMPETENCY          | PROBLEM TYPE | INDUCED | LEO   |
|---------------------------|-----------------|--------------------------|--------------|---------|-------|
| HEALTHCARE                | T K D           | D1 D2 D3 D4 D5           | 6            | 3.1     | 4 · 3 |
| CLINICAL DECISION SUPPORT | →               | →                        |              |         |       |
| MACHINE LEARNING          |                 | Test and Evaluation + D3 |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Build the abstention mechanism as part of the deployment, not as a post–hoc filter; the conformal–prediction framework’s clinical value is that the model can refuse to alert when its calibrated confidence does not support an alert, and the abstention is a deployment capability rather than an analysis artifact.
- ▶ Pre–specify how the attribution question will be answered before evaluation; the case demonstrates that a deployment reaching every unit at once leaves no comparison arm, and that naming the substitute — a confounder–adjusted causal–impact analysis, and a record of what else changed — is the discipline.
- ▶ Integrate the model into specific workflow surfaces — ED admission, inpatient deterioration — rather than as a hospital–wide alerting layer; the deployment’s specificity is part of what makes the conformal–prediction abstention structure operationally meaningful.

### IN OPERATION / AFTER

- ▶ Carry the missing–comparison–arm hedge into print without softening; the case’s pedagogical value depends on the attribution question being explicit alongside the deployment–outcome reporting.
- ▶ Pair in syllabi with Case 20 (TREWS) and Case 5 (Epic Sepsis) so the AI–delegation typology in sepsis is taught as a three–case set at the deployment–evidence level.
- ▶ Use the case to anchor the cue–and–alert design LEO; the curricular target is the discipline of building abstention into the alerting architecture as a capability deliverable, and of evaluating the abstention against the calibration–gap failure mode the prior cases named.

## REFLECTION QUESTIONS

1. Identify a clinical–AI deployment in your setting whose alerting architecture does not include an abstention mechanism. What is the calibration–gap failure mode the deployment would encounter, and what would the conformal–prediction abstention structure look like in your workflow?
2. Specify the multifactorial–attribution acknowledgement you would carry in a prospective–implementation evaluation. What concurrent process improvements would have to be enumerated, and what is the documented decision rule for what the deployment evaluation can and cannot establish?
3. The three sepsis cases — TREWS, Epic Sepsis, COMPOSER — together define an AI–delegation typology. Pick a clinical–AI category in your setting and ask: what is the analogous typology, and which deployments would have to be evidenced to define it at the same level of resolution?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · systems & human–factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

# Radiology AI Miscalibration

**H** Human–System Interface   **K** Knowledge & Institutional Memory   **D** Designed Out

**IMPACT** FDA-cleared radiology AI is cleared on evidence that cannot establish deployment performance — of the 130 devices Wu et al. audited, 126 were evaluated only retrospectively and none of the 54 high-risk devices prospectively, 93 carried no publicly reported multi-site assessment, and 17 reported demographic subgroup performance; the canonical v1 anchor for clinical-AI deployment without surveillance, cross-referenced by the Epic Sepsis (Case 5) and pulse-oximetry (Case 26) deployment-evidence cases

## IN BRIEF

FDA-cleared radiology AI tools — for chest X-ray classification, mammography, CT triage — are cleared on evidence that cannot establish how they will perform in the populations that will meet them. The peer-reviewed record of degradation is built on research classifiers trained on the same public datasets as the cleared tools, not on the cleared tools themselves, which are largely unmeasured after clearance; the degradation in those classifiers concentrates in patient groups under-represented in the training data. Larrazabal et al. (PNAS 2020) showed this structurally for chest-X-ray classifiers across gender imbalance; Seyyed-Kalantari et al. (Nature Medicine 2021) extended the finding to under-served racial and socioeconomic subgroups across three large public chest-X-ray datasets and a multi-source combination of them — a result three published critiques then contested and the authors defended. Obermeyer et al. (Science 2019) showed that bias in the labels used to train clinical AI can under-allocate care to Black patients even when the model looks well-calibrated on its chosen proxy. Wu et al. (Nature Medicine 2021) put numbers on the clearance evidence itself: of 130 devices cleared between January 2015 and December 2020, 126 were evaluated only retrospectively — none of the 54 high-risk devices prospectively — 93 reported no publicly available multi-site assessment, and 17 reported performance by demographic subgroup. The FDA's 510(k) clearance pathway — the route most cleared radiology AI tools have taken — does not routinely require demographic stratification of validation metrics or post-market monitoring of how a tool performs in the population using it; the De Novo pathway used for a small number of novel tools imposes more, but is rarely the chosen route. The FDA's final Predetermined Change Control Plan guidance (December 2024) and its January 2025 draft guidance on the lifecycle management of AI-enabled device software, together with the Good Machine Learning Practice principles, begin to address this gap; the institutional architecture for demographic post-market surveillance is still being built. The capability gap is in the regulatory architecture, not the model: clearance is not the same thing as clinically performable deployment. The case is the v1 anchor for the cross-references in the Epic Sepsis (Case 5) and pulse-oximetry (Case 26) deployment-evidence cases.

## THE SHIFT

Machine-learning tools are moving rapidly into radiology and other diagnostic medicine, cleared for market and integrated into clinical workflows that affect real patients. The FDA's

public list of AI-enabled device authorizations passed a thousand entries during 2024 and has kept growing since; the great majority came through the 510(k) clearance pathway as devices substantially equivalent to a predicate. Unlike a drug, a model can pass its validation study and still behave very differently once it meets a population that differs from its training data — the same model file that scored well on the clearance set can quietly carry a different error profile into every hospital whose patients do not resemble it. Wu and colleagues measured how thin the clearance evidence is: across the 130 AI devices the FDA authorized between January 2015 and December 2020, 126 were evaluated only against data collected before deployment, and not one of the 54 high-risk devices was evaluated prospectively; 93 carried no publicly reported multi-site assessment; 59 did not report the size of the evaluation sample at all, and among the 71 that did, the median was 300 patients; 17 reported any performance breakdown by demographic subgroup. Their own cross-site experiment shows what that costs. A pneumothorax classifier trained at Stanford scored an AUC of 0.903 on Stanford patients, 0.870 at Beth Israel Deaconess and 0.852 on the NIH set; the NIH-trained model scored 0.883 at home and 0.759 at Beth Israel. The 510(k) substantial-equivalence framing was built for an era of physical devices whose behavior was largely determined by their design; it was not built for a class of systems whose behavior depends on the distribution of the data they meet.<sup>o</sup>

#### WHAT IS EMERGING

The peer-reviewed evidence for the divergence comes from research classifiers built on the same public chest-X-ray datasets as the cleared tools, and it is worth being exact about what those studies measure: not deployment, but performance broken out by patient group, and performance carried across from the population a model was trained on to a different one. On both measures the shortfall concentrates in the same place. Larrazabal et al. (PNAS 2020) demonstrated structural performance drops for the under-represented gender in chest-X-ray training data, using the NIH ChestX-ray14 and CheXpert datasets across three network architectures to show that a classifier's AUC for a given group tracks that group's share of the training set — and that even a 25/75 split degrades the minority class relative to a balanced one. Seyyed-Kalantari et al. (Nature Medicine 2021) extended the finding directly: across three large public chest-X-ray datasets and a multi-source combination of them, classifiers under-diagnosed Black, Hispanic, female, and lower-socioeconomic patients at higher rates, with the underdiagnosis rate highest for patients at the intersection of those groups — evidence that the shortfall is not a stray bug but a predictable consequence of which patients the training set did and did not contain. That second result is contested in the literature and the case does not rest on it alone: three Matters Arising in Nature Medicine in 2022 argued that training and test splits drawn from the same biased dataset make the reported disparities hard to interpret, and that confounders were not accounted for; the authors replied and maintained the finding. The dispute is about how much of the gap is model and how much is dataset — not about whether the gap is there.<sup>1</sup>

#### THE CAPABILITY QUESTION

The problem is not confined to imaging. Obermeyer et al. (Science 2019) showed that a widely deployed care-management algorithm — one of a class of commercial risk-prediction tools applied to an estimated 200 million Americans a year — systematically under-allocated resources to Black patients because it was trained on healthcare cost as a

proxy label for need: equally sick Black patients had lower historical costs because they had received less care, so the model rated them as lower-need. Correcting the label — predicting illness rather than cost — would have raised the share of Black patients referred for extra help from 17.7 to 46.5 percent. The question is how a regulator can certify a model as safe without checking how it behaves across the populations and labels it will actually meet — since a model can look well-calibrated on its chosen proxy while the proxy itself encodes the inequity it then propagates. Note what the other side of the argument looks like. Wachter and Brynjolfsson (JAMA 2024) argue that health care is now readier for generative AI than it was for its predecessors, and that the usual lag between deploying a general-purpose technology and seeing benefit from it may be shorter this time. That is a claim about the speed of adoption, not about the adequacy of evaluation — and it is precisely why the evaluation gap matters: the faster a class of tools reaches patients, the more the absence of standing surveillance costs.<sup>2</sup>

#### EARLY EVIDENCE

Similar population-dependence has been reported across modalities. In mammography, the most-publicized system of all — the DeepMind evaluation of Case 37 — cut false positives by 5.7 percentage points on its U.S. set and 1.2 on its UK set. Read that gap carefully: the UK set was assembled to be representative and the U.S. set was cancer-enriched, so the difference is partly a property of how the two datasets were built and not only of the two populations. That is the point, not an aside — an effect size read off a curated evaluation set is not a prediction about a screening program. In sepsis prediction, the Epic Sepsis Model (Case 5) is the corpus's canonical example of a tool deployed at scale whose external validation found it substantially worse than its developer-reported numbers. Pulse-oximetry miscalibration (Sjoding et al., Case 26) sits structurally underneath several downstream AI sepsis and triage tools. And in dermatology, Daneshjou and colleagues built the first pathologically confirmed, expertly curated image set with diverse skin tones and found that state-of-the-art models degrade substantially on dark skin — along with a finding that closes the loop back to Obermeyer: the dermatologists who label these training sets also perform worse on dark skin, so the label carries the gap before the model ever sees it. Yet the FDA's 510(k) clearance pathway does not routinely require demographic stratification of validation metrics, nor post-market monitoring of in-use performance — so the divergence between clearance and deployment is largely invisible while the tool is in use, and a shortfall concentrated in one patient group can persist unmeasured across the entire period the tool is influencing care.<sup>3</sup>

#### OPEN PROBLEMS

The FDA's evolving AI/ML-device guidance begins to address this gap. The Predetermined Change Control Plan — finalized as its own guidance in December 2024 — lets manufacturers pre-specify the model updates and validation procedures that can be applied without a new clearance; the broader device-lifecycle draft guidance that followed in January 2025, together with the Good Machine Learning Practice principles, emphasizes representative training data, transparent performance metrics, and lifecycle monitoring, and remained in draft through mid-2026. The De Novo pathway, used for a small number of novel tools, imposes more — but is rarely the chosen route. What the guidance does not yet require, as of its draft form, is mandatory demographic stratification of validation metrics at clearance and mandatory population-level post-market surveillance of in-use performance for



the cleared model. Radiology AI thus remains the live, recurring case for the gap between regulatory clearance and clinically performable deployment. The capability gap is at the regulatory architecture, not the model: the institutional machinery to require demographic post-market surveillance has not yet been built. It is the medical-AI analog of the Vioxx post-market-surveillance failure (Case 9) at a new technological boundary — a case where the harm comes not from a hidden defect but from the absence of any standing system to watch the tool once it is in the population's hands.<sup>4</sup>



*Performance metrics on a clearance dataset are not the same as performance metrics in deployment populations.*

EDITORS' SYNTHESIS OF THE FDA AI/ML-BASED SAMD DISCUSSION PAPER (2019)

## REFERENCES

1. Wu, E., Wu, K., Daneshjou, R., Ouyang, D., Ho, D. E., & Zou, J. (2021), "How medical AI devices are evaluated: limitations and recommendations from an analysis of FDA approvals," *Nature Medicine* 27(4):582–584, doi:10.1038/s41591-021-01312-x — the 130-device audit (126 retrospective-only, none of the 54 high-risk devices prospective, 93 without publicly reported multi-site assessment, 17 reporting demographic subgroup performance) and the cross-site pneumothorax experiment behind this case's diagram.
2. Seyyed-Kalantari, L., Zhang, H., McDermott, M. B. A., Chen, I. Y., & Ghassemi, M. (2021), "Underdiagnosis bias of artificial intelligence algorithms applied to chest radiographs in under-served patient populations," *Nature Medicine* 27:2176–2182 — disparities across Black, Hispanic, female, and lower-socioeconomic subgroups; underdiagnosis highest for intersectional subpopulations such as Hispanic female patients.
3. Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019), "Dissecting racial bias in an algorithm used to manage the health of populations," *Science* 366(6464):447–453 — proxy-label bias in a care-management algorithm; 200 million Americans/year are subject to tools of this class.
4. Daneshjou, R., Vodrahalli, K., Novoa, R. A., et al. (2022), "Disparities in dermatology AI performance on a diverse, curated clinical image set," *Science Advances* 8(32):eabq6147, doi:10.1126/sciadv.abq6147 — the Diverse Dermatology Images set; models and labeling dermatologists both degrade on dark skin tones.
5. FDA, "Proposed Regulatory Framework for Modifications to AI/ML-Based Software as a Medical Device" (2019); FDA's public "Artificial Intelligence-Enabled Medical Devices" authorization list; final Predetermined Change Control Plan guidance (December 2024) and the draft "Artificial Intelligence-Enabled Device Software Functions: Lifecycle Management and Marketing Submission Recommendations" (January 2025), with Good Machine Learning Practice principles — the regulatory trajectory; absence of mandatory demographic stratification at clearance and post-market monitoring of in-use performance.
6. Larrazabal, A. J., Nieto, N., Peterson, V., Milone, D. H., & Ferrante, E. (2020), "Gender imbalance in medical imaging datasets produces biased classifiers for computer-aided diagnosis," *PNAS* 117(23):12592–12594 — AUC drops for under-represented genders on NIH ChestX-ray14 and CheXpert, across three network architectures.
7. Bernhardt, M., Jones, C., & Glocker, B. (2022), "Potential sources of dataset bias complicate investigation of underdiagnosis by machine learning algorithms," *Nature Medicine* 28(6):1157–1158 — one of three Matters Arising contesting Seyyed-Kalantari et al.; the authors' reply is at 28(6):1161–1162, doi:10.1038/s41591-022-01854-8. Cited here so the contested status travels with the finding.
8. Wachter, R. M. & Brynjolfsson, E. (2024), "Will Generative Artificial Intelligence Deliver on Its Promise in Health Care?" *JAMA* 331(1):65–69, doi:10.1001/jama.2023.25054 — the counterweight argument that health care is readier for generative AI than for its predecessors and the lag to benefit may be shorter. It is not a bias or evaluation paper; do not cite it for the proxy-and-population problem.

LE INSIGHT

Radiology AI is the canonical contemporary case for the gap between regulatory clearance and clinical deployment performance in medical AI. The clearance dataset and the deployment population diverge. The institutional architecture to surface the divergence — demographic post-market surveillance — has not yet been built.

LENS APPROACH

Radiology AI is the canonical contemporary clinical-AI deployment-and-surveillance case (induced 2.4 — post-deployment surveillance and disconfirmation, where `competencies.md` places it; LENS D4+D3/PT5). LENS uses it in Domain 4 (Test and Evaluation; LEO-4) for the clearance-vs-deployment measurement architecture and for demographic stratification of validation as a deliverable. LENS uses it in Domain 3 (Human-System Collaboration; LEO-3) for the human-AI deployment-without-surveillance pattern, and in Domain 5 (Navigating Sociotechnical Constraints) for the FDA AI/ML regulatory trajectory itself. The Obermeyer (2019) finding generalizes the diagnosis: bias enters through the labels and through the population, both of which the 510(k) process currently treats as outside its scope. Pair with Vioxx (Case 9) as the post-market-surveillance-failure pattern at a new technological boundary; cross-references the Epic Sepsis (Case 5) and pulse-oximetry (Case 26) cases at the clinical-AI-without-surveillance layer.

| DOMAIN     |   | MODES AT STAKE |   | LENS COMPETENCY          | PROBLEM TYPE | INDUCED | LEO   |
|------------|---|----------------|---|--------------------------|--------------|---------|-------|
| HEALTHCARE | → | H K D          | → | D1 D2 D3 D4 D5           | 5            | 2.4     | 4 · 3 |
| TECHNOLOGY |   |                |   | Test and Evaluation + D3 |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Require validation metrics stratified by demographic group at clearance, so a tool’s performance is established on the populations it will actually meet, not an aggregate.
- ▶ Specify the intended deployment population and label definition explicitly, and treat divergence from the training distribution as a known risk to be measured, not assumed away.
- ▶ Build the data pipeline for in-use performance capture before deployment, so post-market signals are collectable from the first patient rather than reconstructed after harm.

IN OPERATION / AFTER

- ▶ Mandate demographic post-market surveillance of in-use performance, the institutional machinery the clearance pathway currently lacks.
- ▶ Monitor for the clearance-to-deployment performance gap continuously, since a shortfall concentrated in one group can otherwise persist invisibly.
- ▶ Tie continued authorization to demonstrated in-population performance, so a tool that degrades in deployment can be withdrawn before the divergence compounds.

REFLECTION QUESTIONS

1. Identify a model in your domain whose deployment population diverges from its training population. What is the institutional architecture to surface the divergence?
2. Design the demographic post-market surveillance deliverable that should accompany every FDA clearance of medical AI.
3. FDA 510(k) clearance does not currently require demographic stratification of validation metrics, nor does it require post-market monitoring of how a cleared tool actually performs on the population using it. What is the minimum reporting deliverable a regulator should require so the gap is visible while the tool is in use?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · task analysis · design prototyping



# AlphaFold — Protein Structure Prediction

**K** Knowledge & Institutional Memory **G** Governance & Trust

**IMPACT** Substantially solved a 50-year protein-structure prediction problem — median GDT 92.4 across all CASP14 targets — and released predictions for over 214 million protein sequences; the book's foundational positive AI case, and, after the AlphaFold 3 access reversal, its clearest evidence that open release is a decision that can be taken back

## IN BRIEF

DeepMind's AlphaFold 2 substantially solved the half-century-old protein-structure prediction problem in computational biology — demonstrated at CASP14 in November 2020 with a median GDT of 92.4 across all targets, published in *Nature* in July 2021. The publicly released AlphaFold Protein Structure Database now contains predicted structures for more than 214 million protein sequences — essentially the entire known protein universe — and has been integrated into structural-biology and drug-discovery workflows worldwide. AlphaFold is the strongest positive AI case in this book, and its lesson is in the conditions that made the benefit possible: a benchmark (CASP) the field had agreed on for decades, high-quality training data, an output biologists could verify against experimental structures, and an open release that let the global community adopt it fast. Each was a precondition for the success — and none of them was the model itself. Two of the four are less settled than the headline suggests. Verification reaches the few dozen blind CASP targets far more securely than it reaches the database: 58 percent of human residues carry a confident prediction and 36 percent a very high-confidence one, which is why the per-residue confidence score is part of the release rather than an annotation on it. And AlphaFold 3 shipped in May 2024 with neither code nor weights, reopening the code only in November after more than a thousand researchers objected — and the weights only on application. Open release is a decision that gets made again at every version.

## THE SHIFT

Predicting a protein's three-dimensional structure from its amino-acid sequence was one of biology's grand challenges for half a century — slow, expensive experimental work that bottlenecked drug discovery and basic research. Deep learning offered, for the first time, the prospect of solving it computationally at scale, turning a problem that had been a years-long experimental undertaking per protein into one that could be approached for the whole proteome at once.<sup>o</sup>

## WHAT IS EMERGING

DeepMind's AlphaFold 2 substantially solved the problem at CASP14 in November 2020, predicting structures at accuracies rivaling experiment: a median GDT of 92.4 across all targets, with scores above 90 out of 100 on roughly two-thirds of them. The publicly released AlphaFold Protein Structure Database now holds predicted structures for more than 214

million protein sequences — close to the entire known protein universe, up from the 300,000 structures of the first 2021 release — and has been folded into research workflows worldwide and into the primary data resources of the field (PDB, UniProt, Ensembl, InterPro), so the benefit arrived not as a single laboratory's advantage but as a shared resource the wider community could build on immediately.<sup>1</sup>

#### THE CAPABILITY QUESTION

What made AlphaFold a benefit rather than just a benchmark win? The case poses the question of which conditions allow an AI capability to be safely and widely useful — and the answer turns out to lie around the model, not in it: an agreed benchmark, trustworthy training data, a verifiable output, and a deliberate decision about release. Each of these is an institutional or evidentiary precondition, not an artifact of the architecture, which is why the case reads as a lesson about capability infrastructure rather than about a model.<sup>2</sup>

#### EARLY EVIDENCE

AlphaFold's success rested on four features, each a precondition rather than the model: the CASP benchmark the field had used for decades, high-quality experimental training data, output that biologists could check against known structures, and an open release that let the global community adopt the tool quickly. Where those conditions hold, AI amplifies capability; the technical model alone does not — and the same architecture dropped into a domain missing any one of those four would not have produced a comparable, trusted, widely adopted result.<sup>3</sup> The work was extended by AlphaFold 3 (published in Nature on 8 May 2024), which predicts proteins in complex with DNA, RNA, small molecules, ions and modified residues, and in October 2024 Demis Hassabis and John Jumper shared half the Nobel Prize in Chemistry, "for protein structure prediction"; the other half went to David Baker for computational protein design. AlphaFold 3 also complicated the open-release story this case credits, and the sequence is worth following exactly. The paper appeared with neither code nor model weights — only a rate-limited web server — and Nature published it that way, prompting its own editorial asking why. More than a thousand scientists signed an open letter arguing the release breached the norms of reproducibility. Six months later, on 11 November 2024, Google DeepMind put the code on GitHub under a non-commercial license — but the weights remain available only to academics who apply for them. So the release was not reversed and then restored; it was reversed, contested, and partially restored, on terms the authors set. Open release is a repeated governance decision that can be taken back under commercial pressure, not a settled property of the work — and the demand that forced the partial retreat is the same demand the Haibe-Kains comment made of DeepMind's mammography paper in Case 37.<sup>4</sup>

#### OPEN PROBLEMS

AlphaFold is the strongest positive AI case in the dataset for what supports beneficial deployment in a well-defined scientific domain — which makes it worth being exact about how far the verification actually reaches. CASP verified the model on a few dozen blind targets; the database holds 214 million predictions, and almost none of them will ever meet an experimental structure. Before AlphaFold, experimental structures covered 17 percent of the residues in human protein sequences; the AlphaFold human proteome covers 98.5 percent of human proteins but returns a confident prediction for 58 percent of residues, of which 36 percent of all residues are very high confidence. The verifiable-output precondition holds at

the benchmark and thins out across the database, which is why the per-residue confidence score travels with every structure and is the part of the release a user is obliged to read. The open problem is generalization: most consequential problems lack an agreed benchmark, clean training data, or verifiable output. The frontier question is how much of the AlphaFold pattern can be reconstructed where those preconditions are not given for free — that is, whether a field can deliberately build the benchmark, the data, and the verification path that protein structure happened to have accumulated over decades.<sup>5</sup>

# 214M

PREDICTED STRUCTURES · PUBLICLY RELEASED

*agreed benchmark, training data, verifiable output, open release*

ALPHAFOLD — THE CONDITIONS FOR BENEFICIAL AI DEPLOYMENT

*We have been stuck on this one problem — how do proteins fold up — for nearly 50 years. To see DeepMind produce a solution for this, having worked personally on this problem for so long and after so many stops and starts, wondering if we'd ever get there, is a very special moment.*

JOHN MOULT, CO-FOUNDER AND CHAIR OF CASP, QUOTED IN GOOGLE DEEPMIND'S 30 NOVEMBER 2020 ANNOUNCEMENT OF THE CASP14 RESULT

## REFERENCES

1. Jumper, J., Evans, R., Pritzel, A., et al. (2021), "Highly accurate protein structure prediction with AlphaFold," *Nature* 596(7873):583–589, doi:10.1038/s41586-021-03819-2 — the method, and the CASP14 validation.
2. Varadi, M., Bertoni, D., Magana, P., et al. (2024), "AlphaFold Protein Structure Database in 2024: providing structure coverage for over 214 million protein sequences," *Nucleic Acids Research* 52(D1):D368–D375, doi:10.1093/nar/gkad1011 — the database at 214M, up from 300k in 2021, and its integration into PDB, UniProt, Ensembl and InterPro.
3. CASP14 press release, Protein Structure Prediction Center, 30 November 2020 — GDT above 90 on roughly two-thirds of targets. Note the Moulton quotation this case carries is **not** in the press release: it is from Google DeepMind's post of the same date, "AlphaFold: a solution to a 50-year-old grand challenge in biology," which also reports the 92.4 median GDT.
4. Tunyasuvunakool, K., Adler, J., Wu, Z., et al. (2021), "Highly accurate protein structure prediction for the human proteome," *Nature* 596(7873):590–596, doi:10.1038/s41586-021-03828-1 — 98.5% of human proteins covered; 58% of residues confidently predicted, 36% of all residues at very high confidence, against 17% of residues covered by experimental structure beforehand.
5. AlphaFold 3 access chronology: publication with neither code nor weights (May 2024), the open letter signed by more than a thousand researchers, the *Nature* editorial "AlphaFold3 — why did *Nature* publish it without its code?" (2024, doi:10.1038/d41586-024-01463-0), and the 11 November 2024 GitHub release of the code under a non-commercial license with weights by application — the open-release governance decision as actually written.
6. Abramson, J., Adler, J., Dunger, J., et al. (2024), "Accurate structure prediction of biomolecular interactions with AlphaFold 3," *Nature* 630(8016):493–500, doi:10.1038/s41586-024-07487-w; and the 2024 Nobel Prize in Chemistry, one half to David Baker for computational protein design and one half jointly to Demis Hassabis and John M. Jumper for protein structure prediction.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

AlphaFold is the strongest positive AI case in the dataset. The technical achievement is real. The conditions that made the benefit possible — agreed benchmark, training data, verifiable output, open release — are the capability infrastructure around the model, not the model itself. Two of the four are less settled than the headline suggests: verification reaches the CASP targets far more securely than it reaches the 214 million database entries, and AlphaFold 3 showed that open release is a decision that can be taken back.

LENS APPROACH

AlphaFold is the preconditions-for-beneficial-deployment case (induced 2.1; LENS D4+D5/PT2; LEO-4 and LEO-5). LENS uses it in Domain 4 (Test and Evaluation) for the benchmark-and-verification half — CASP as an instrument the field agreed on and an independent body ran, and the per-residue confidence score as the honest limit on what the database can be trusted for — and in Domain 5 (Navigating Sociotechnical Constraints) for the open-release decision and its partial reversal at AlphaFold 3. In the course sequence: LEN 1 as a problem-framing case for what productive AI deployment looks like, LEN 9 as a technical achievement, LEN 7 for the release governance. Pair with Case 37, where the same demand for released code and methods was made of the same laboratory.

| DOMAIN     | MODES AT STAKE           |     | LENS COMPETENCY |    |    |    |       | PROBLEM TYPE | INDUCED | LEO   |
|------------|--------------------------|-----|-----------------|----|----|----|-------|--------------|---------|-------|
| HEALTHCARE | →                        | K G | →               | D1 | D2 | D3 | D4 D5 | 2            | 2.1     | 4 · 5 |
| TECHNOLOGY | Test and Evaluation + D5 |     |                 |    |    |    |       |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Establish an agreed benchmark and high-quality training data for the target problem before building the model, treating these as preconditions rather than afterthoughts.
- ▶ Engineer the output to be verifiable against an independent ground truth, so users can check predictions rather than having to trust them.
- ▶ Decide the release and access terms deliberately as a governance choice, since open release is what distributed AlphaFold’s benefit globally.

IN OPERATION / AFTER

- ▶ Monitor downstream use to confirm the verifiable-output property holds in practice, and that users are in fact reading the per-residue confidence score where stakes are high; a database entry is a prediction with a stated confidence, not a structure.
- ▶ Track whether the four preconditions still hold as the tool is applied to new protein families or adjacent problems beyond its validated domain.
- ▶ Sustain the open resource and benchmark over time, so the community-wide benefit does not erode as the field and the data move on — and treat the release terms as re-decided at every version, because AlphaFold 3 shipped without code or weights and only partially reopened after a thousand researchers objected.

REFLECTION QUESTIONS

1. Identify a domain in your work where the conditions that supported AlphaFold’s success (benchmark, training data, verifiable output) are present. What is the analogous opportunity?
2. The open release of AlphaFold’s predictions was a governance decision. Design the equivalent decision for an AI capability your institution might develop.
3. AlphaFold inherited a benchmark, clean data, and a verification path that protein structure had accumulated over decades. Pick a problem in your domain that lacks one of those preconditions, and lay out how a field would deliberately build it.

WHO BUILDS THIS knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. TOOLS knowledge bases · data pipelines · governance frameworks · bias & evidence audits

# DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback

**K** Knowledge & Institutional Memory **G** Governance & Trust

**IMPACT** McKinney et al. 2020 Nature paper reported a deep-learning mammography screening system outperforming radiologists on a representative retrospective UK dataset and a cancer-enriched U.S. one, with reductions in false-positives (5.7 percentage points in the U.S. set, 1.2 in the UK set) and false-negatives (9.4 and 2.7 percentage points respectively); Haibe-Kains et al.'s October 2020 Nature comment held that "the lack of detailed methods and computer code undermines its scientific value," and the authors' addendum in the same issue supplied hyperparameters and the augmentation strategy but not the code or the trained models

## IN BRIEF

Scott Mayer McKinney and colleagues at Google Health and DeepMind published "International evaluation of an AI system for breast cancer screening" in Nature on January 1, 2020. The paper reported that a deep-learning system outperformed radiologists on two retrospective datasets — a large representative UK screening set and a large cancer-enriched U.S. one — with reductions in false-positives of 5.7 percentage points (U.S.) and 1.2 percentage points (UK) and reductions in false-negatives of 9.4 and 2.7 percentage points respectively. The paper drew unusual press attention and rapidly entered the policy conversation on AI-assisted screening. Haibe-Kains and colleagues' October 14, 2020 Nature comment titled "Transparency and reproducibility in artificial intelligence" critiqued the McKinney paper for failing to release code, trained models, or sufficient methodological detail to permit independent reproduction. The load-bearing hedge the Haibe-Kains comment delivers is exact: "the lack of detailed methods and computer code undermines its scientific value." In the same issue the authors published an addendum expanding the supplementary methods with optimization hyperparameters and a fuller account of the augmentation strategy — a partial answer that left the code and the trained models unreleased. The screening-comparison framing the original paper offered has since been refined by deployment evidence rather than confirmed at the deployment scale the headline implied — and the randomized evidence that eventually arrived came from a different, commercial system, not this one. The case pairs with Case 35 (Radiology AI Miscalibration), Case 5 (Epic Sepsis), and Case 26 (Pulse oximetry).

## THE SHIFT

The McKinney et al. paper was published on January 1, 2020, in Nature — a top-tier venue and an unusually high-profile publication for a deep-learning medical-imaging study. The

work was a collaboration across Google Health, DeepMind, and clinical partners at Cancer Research UK Imperial Centre, Northwestern Medicine, the Royal Surrey County Hospital, and the University of Cambridge. The retrospective evaluation used two screening datasets that were not built alike — the UK set assembled to be large and representative, the U.S. set large and **enriched** for cancer — and compared the AI system’s outputs against single-reader and double-reader radiologist performance. The headline framing was that the AI system reduced both false-positives and false-negatives relative to radiologists, with the U.S. dataset showing larger absolute reductions than the UK dataset. Two further results shaped how the paper was read. In a separate study of six radiologists the system beat every one of them, with an AUC 11.5 percentage points above the average reader. And in a simulation of the UK double-reading protocol, the system held non-inferior performance while removing 88 percent of the second reader’s workload — a simulated number that set the expectation the randomized trials would later be measured against.<sup>9</sup>

#### WHAT IS EMERGING

The press response was substantial. Mainstream coverage framed the result as “AI outperforms radiologists at breast cancer screening,” and the framing entered the policy conversation on AI-assisted medical imaging quickly. The framing carried more weight than the underlying retrospective comparison was designed to support: a retrospective evaluation against historical reader performance is informative about model output, but prospective deployment against current radiologists operating in their current workflow involves variables — reader fatigue, screen presentation, integration with reading worklists, recall thresholds — that the retrospective study does not measure. The McKinney paper itself was careful in its claims; the gap between the paper’s careful claims and the press’s framing of the headline is part of the case.<sup>1</sup>

#### THE CAPABILITY QUESTION

The October 14, 2020 Nature comment by Benjamin Haibe-Kains, George Adam, Ahmed Hosny, Farnoosh Khodakarami, Massive Analysis Quality Control (MAQC) Society Board of Directors, Levi Waldron, Bo Wang, Chris McIntosh, Anna Goldenberg, Anshul Kundaje, Casey S. Greene, Tamara Broderick, Michael M. Hoffman, Jeffrey T. Leek, Keegan Korthauer, Wolfgang Huber, Alvis Brazma, Joelle Pineau, Robert Tibshirani, Trevor Hastie, John P. A. Ioannidis, John Quackenbush, and Hugo J. W. L. Aerts is the load-bearing reproducibility critique. The comment argued that the McKinney paper had not released code, had not released trained models, and had not provided sufficient methodological detail to permit independent reproduction. The comment was specific: the hyperparameters of the three networks, the augmentation pipeline and the training pipeline were not documented well enough for another group to rebuild the system. The comment did not allege error in the paper; it argued that reproducibility had not been established. What happened next is the part usually left out. In the same 14 October 2020 issue, the McKinney authors published both a formal reply and an addendum. The addendum is the substantive half: it expanded the supplementary methods with the optimization hyperparameters and a fuller description of the augmentation strategy, explicitly to help others replicate the result. The two specific documentation gaps the comment named were therefore closed within one issue. The code and the trained models were not released, and have not been. So the exchange settled at a partial remedy — enough to document the method, not enough to rebuild the system —



which is the honest shape of most reproducibility disputes and the reason the demand has to be made before publication rather than after.<sup>2</sup>

#### EARLY EVIDENCE

The case pairs with Case 35 (Radiology AI Miscalibration) for the medical-imaging-AI-deployment-evidence thread: retrospective evaluation produces one class of evidence; prospective deployment produces another, and the two are not interchangeable. Pair with Case 5 (Epic Sepsis) for the high-profile-result-versus-deployment-evidence thread in healthcare AI; Epic Sepsis is the reference case in the corpus for the gap between vendor or developer claims and external evaluation, and DeepMind Mammography sits in the same conceptual family at a different domain. Pair with Case 26 (Pulse oximetry) for the population-heterogeneity-in-medical-AI thread — with one caution this case has to carry rather than exploit. The McKinney paper's UK-versus-U.S. effect-size difference (5.7 vs 1.2 pp on false-positives) is often read as evidence that the system's performance varies across screening populations. It may; but the U.S. set was enriched for cancer and the UK set was not, so dataset construction and population difference are confounded in that comparison and the paper does not separate them. The teachable point is the confound itself: two numbers from one paper are not comparable unless the sets behind them were built the same way.<sup>3</sup>

#### OPEN PROBLEMS

The honest hedges the case carries are not decoration. The Haibe-Kains comment is not a finding that the McKinney paper was wrong; it is a finding that the paper as published did not establish reproducibility. The arc that followed did not run straight from the paper to vindication. In 2021 Freeman and colleagues reviewed the whole retrospective evidence base for AI in breast screening for the BMJ — twelve studies, 131,822 women, McKinney among them — and found no prospective studies at all, poor methodological quality throughout, and, in the three studies that compared AI against the original radiologists' clinical decisions, 34 of 36 AI systems less accurate than a single radiologist and all 36 less accurate than two radiologists in consensus. Their summary of the pattern is the one this book keeps meeting: "promising results in smaller studies are not replicated in larger studies." The randomized evidence arrived later, from the MASAI trial in Sweden — 105,934 women randomized between April 2021 and December 2022, interim safety analysis in Lancet Oncology 2023, screening performance in Lancet Digital Health 2025, and the primary interval-cancer endpoint in The Lancet in January 2026. It found AI-supported reading non-inferior on interval cancer, higher in sensitivity (80.5 against 73.8 percent), identical in specificity, detecting more cancers (6.4 against 5.0 per 1,000) and cutting screen-reading workload by about 44 percent. Note carefully whose system that was: MASAI ran on Transpara, a commercial product from ScreenPoint Medical, not on the DeepMind system this case is about. The DeepMind system was never taken to a randomized trial. So the arc vindicates the **demand** the Haibe-Kains comment made — prospective, independently runnable evaluation — while leaving the original result exactly where the comment left it, and it was a different developer, working in the open with a regulated commercial product, who supplied the evidence. The case teaches the verification-as-deployment-event pattern: a high-profile retrospective result is the starting point of a verification arc, not its endpoint, and the reproducibility infrastructure the Haibe-Kains comment named is the condition for the arc to be possible.

*In their study, McKinney et al. showed the high potential of artificial intelligence for breast cancer screening. However, the lack of detailed methods and computer code undermines its scientific value.*

HAIBE-KAINS ET AL., "TRANSPARENCY AND REPRODUCIBILITY IN ARTIFICIAL INTELLIGENCE," NATURE 586:E14–E16 (14 OCTOBER 2020)

REFERENCES

1. McKinney, S. M., Sieniek, M., Godbole, V., Godwin, J., Antropova, N., Ashrafian, H., Back, T., et al. (2020), "International evaluation of an AI system for breast cancer screening," Nature 577:89–94, doi:10.1038/s41586-019-1799-6.

2. Haibe-Kains, B., Adam, G. A., Hosny, A., Khodakarami, F., MAQC Society Board of Directors, Waldron, L., Wang, B., et al. (2020), "Transparency and reproducibility in artificial intelligence," Nature 586:E14–E16, doi:10.1038/s41586-020-2766-y.

3. McKinney, S. M., Karthikesalingam, A., Tse, D., Kelly, C. J., Liu, Y., Corrado, G. S., & Shetty, S. (2020), "Reply to: Transparency and reproducibility in artificial intelligence," Nature 586:E17–E18, doi:10.1038/s41586-020-2767-x — the developers' response, published alongside the addendum below.

4. McKinney et al. (2020), "Addendum: International evaluation of an AI system for breast cancer screening," Nature 586:E19, doi:10.1038/s41586-020-2679-9 — published in the same issue as the Haibe-Kains comment; expands the supplementary methods with optimization hyperparameters and the augmentation strategy to assist replication. Code and trained models were not released.

5. Freeman, K., Geppert, J., Stinton, C., Todkill, D., Johnson, S., Clarke, A., & Taylor-Phillips, S. (2021), "Use of artificial intelligence for image analysis in breast cancer screening programmes: systematic review of test accuracy," BMJ 374:n1872, doi:10.1136/bmj.n1872 — twelve studies, 131,822 women, McKinney included; no prospective studies found, studies of poor methodological quality, 34 of 36 AI systems less accurate than a single radiologist and all less accurate than two in consensus.

6. MASAI trial (Lång et al.): clinical safety analysis, Lancet Oncology 2023;24(8):936–944, doi:10.1016/S1470-2045(23)00298-X; Hernström et al., screening performance, Lancet Digital Health 2025;7(3):e175–e183, doi:10.1016/S2589-7500(24)00267-X; Gommers et al., interval cancer, sensitivity and specificity, The Lancet 2026;407(10527):505–514, doi:10.1016/S0140-6736(25)02464-X. 105,934 women randomized; the AI system is Transpara v1.7.0 (ScreenPoint Medical), not the system evaluated by McKinney et al.

THE LEARNING ENGINEERING LENS

LE INSIGHT

DeepMind Mammography is the verification-as-deployment-event case at the high-profile-publication scale. The McKinney paper's retrospective result was the starting point of a verification arc, not its endpoint; the Haibe-Kains comment named the reproducibility infrastructure as the condition for the arc, and the authors' addendum met part of the demand without releasing code or models. The randomized evidence that eventually settled the screening question came six years later, from a different vendor's system.

LENS APPROACH

DeepMind Mammography is the deployment-and-reuse-as- verification-events case at the high-profile-publication seam (induced 7.2; LENS D4+D3/PT6; LEO-4 and LEO-3). LENS uses it in Domain 4 (Test and Evaluation) for the reproducibility-infrastructure-as-verification-condition discipline and in Domain 3 (Human-System Collaboration) for the retrospective-versus-prospective-evidence distinction — including the discipline of noticing that a later randomized result does not validate an earlier system unless it is the same system. Pair with Case 35 (Radiology AI Miscalibration), Case 5 (Epic Sepsis), and Case 26 (pulse oximetry population heterogeneity). The Haibe-Kains comment is a reproducibility finding, not a finding of error; the distinction is the decisive hedge.

| DOMAIN           | MODES AT STAKE                      | LENS COMPETENCY                                                              | PROBLEM TYPE | INDUCED | LEO   |
|------------------|-------------------------------------|------------------------------------------------------------------------------|--------------|---------|-------|
| HEALTHCARE       | <div><div>K</div><div>G</div></div> | <div><div>D1</div><div>D2</div><div>D3</div><div>D4</div><div>D5</div></div> | 6            | 7.2     | 4 · 3 |
| MEDICAL IMAGING  | →                                   | →                                                                            |              |         |       |
| MACHINE LEARNING |                                     | Test and Evaluation + D3                                                     |              |         |       |



## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Release code, trained models, and sufficient methodological detail to permit independent reproduction as a condition of publishing a high-profile retrospective medical-AI result; the Haibe-Kains comment names the reproducibility infrastructure as the condition for the verification arc that the original paper opens.
- ▶ Specify in advance the deployment variables — reader fatigue, recall threshold, worklist integration, screening-population characteristics — that a retrospective comparison does not measure but that a deployment will encounter.
- ▶ Treat the gap between the paper's careful claims and the press's framing as a deployment surface, not a communications problem; the framing the field receives is the framing the deployment will operate under in the policy conversation.

### IN OPERATION / AFTER

- ▶ Carry the Haibe-Kains comment's specific framing into print without softening; the comment is a finding on reproducibility, not a finding of error, and the case's pedagogical value depends on the distinction being preserved.
- ▶ Pair in syllabi with Case 35 and Case 5 so the high-profile-result-versus-deployment-evidence pattern is taught across the medical-imaging-AI deployment seam at multiple instances — and check, every time, that the validating trial ran on the system under discussion. MASAI validated Transpara, not the DeepMind system.
- ▶ Use the case to anchor the verification-as-deployment-event frame; the curricular target is the discipline of treating a high-profile retrospective result as the starting point of a multi-year verification arc rather than as a deployment-ready endpoint.

### REFLECTION QUESTIONS

1. Identify a high-profile retrospective result in your domain whose reproducibility infrastructure — code, trained models, methodological detail — has not been released. What would the verification arc the result opens require to proceed, and what currently blocks it?
2. Specify the deployment variables a retrospective comparison in your setting does not measure but a deployment will encounter. What is the prospective evaluation design that would surface those variables before deployment scale?
3. The press framing of a result often carries deployment implications the paper's careful claims do not. Pick a result in your domain and ask: what is the gap between the careful claim and the framing the field receives, and what would have to be true for the gap to be narrowed in advance of deployment decisions?

**WHO BUILDS THIS** knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · governance frameworks · bias & evidence audits

# iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms

**G** Governance & Trust   **D** Designed Out   **N** Normalization of Deviance

**IMPACT** An FDA-mandated risk-evaluation and mitigation program with pregnancy testing, two-method contraception requirements, and lockout authorization before dispensing — Kaiser Permanente cohort (n=8,344; 9,912 treatment courses) reported 29 fetal exposures and 'no evidence that iPLEDGE significantly decreased the risk of fetal exposure' relative to the prior SMART program

## IN BRIEF

Isotretinoin is a highly effective acne medication and a known teratogen: fetal exposure causes severe birth defects. The FDA has required risk-management programs since the 1980s; in 2006 the agency replaced the prior SMART program with iPLEDGE, later approved as a Risk Evaluation and Mitigation Strategy (REMS), requiring pregnancy testing, two contraception methods (or documented abstinence), and pharmacy lockout authorization before each dispense. The case is the productive counterpoint to SUBSAFE (Case 173) and the WHO Surgical Checklist (Case 23): the structural move is the same — mandatory pre-authorization gating a consequential action — and the measured outcome is very different. The Shin et al. Kaiser Permanente cohort (J Am Acad Dermatol, 2011; n=8,344 patients across 9,912 treatment courses) found 29 fetal exposures and concluded “no evidence that iPLEDGE significantly decreased the risk of fetal exposure” compared with the prior program. Collins et al. (2014) put roughly 150 isotretinoin-exposed pregnancies a year in the U.S. despite the program; FDA adverse-event reports give 218 to 310 a year after 2011, with non-adherence — missed pills, inconsistent condom use — the driver Collins documents. The teaching point is that an authorization mechanism without adherence support does not reliably deliver the capability it is built to enforce.

## THE SHIFT

Isotretinoin (Accutane and successors) is a highly effective systemic treatment for severe acne, and a well-characterized teratogen. Fetal exposure during treatment causes a recognized pattern of severe birth defects. The FDA has therefore required some form of risk-management program around its prescribing since the late 1980s, evolving through pregnancy testing requirements, contraception counseling, and structured enrollment programs. iPLEDGE, introduced in 2006, was the first to require pharmacy lockout — no authorization, no dispense — and was folded into the Risk Evaluation and Mitigation Strategy (REMS) architecture Congress created in 2007.<sup>o</sup>

## WHAT IS EMERGING

The mechanism is what the case turns on. iPLEDGE requires the patient to register, document either two contraception methods or abstinence (with declared method), complete monthly pregnancy tests, and receive a per-cycle authorization code before the pharmacy can fill the prescription. Each prescriber, pharmacy, and wholesaler must also enroll in the program. The structural form is the same one SUBSAFE (Case 173) and the WHO Surgical Checklist (Case 23) both use successfully: mandatory pre-authorization gating a consequential action, with the authorization withheld until pre-conditions are verifiably met.<sup>1</sup>

## THE CAPABILITY QUESTION

The measured outcome diverges sharply. Shin et al. (Journal of the American Academy of Dermatology, 2011) studied 8,344 Kaiser Permanente patients across 9,912 isotretinoin treatment courses from March 2004 to February 2008, spanning both programs, and identified 29 fetal exposures — a rate that fell from 3.11 to 2.67 per 1000 courses after iPLEDGE ( $P = .69$ ). The paper's conclusion is the decisive sentence: "no evidence that iPLEDGE significantly decreased the risk of fetal exposure" relative to the prior SMART program. Collins et al. (2014) report approximately 150 isotretinoin-exposed pregnancies a year in the United States despite the program operating as designed; FDA adverse-event reports analyzed by Tkachenko et al. (2019) settle into 218 to 310 annual pregnancy reports after 2011 — down from a 2006 peak of 768, but persistent enough that the authors call for "investigation into the efficacy of iPLEDGE." The driver Collins et al. document is non-adherence: missed pills, inconsistent condom use, the assumption that one cycle of missed contraception is unlikely to coincide with the teratogenic window. Tkachenko et al. do not adjudicate the cause, listing reporting fatigue and a national decline in teenage pregnancy alongside the program itself.<sup>2</sup>

## EARLY EVIDENCE

The teaching point is precise and load-bearing. iPLEDGE is not a program-design failure in the sense that SUBSAFE succeeded and iPLEDGE was sloppy. The mechanism is elaborate, the enrollment burden on prescribers and pharmacies is real, and the program runs as specified. What the evidence shows is that an authorization mechanism without adherence support does not reliably deliver the capability it is built to enforce — because the capability sits downstream of the authorization in patient behavior, and the program does not reach into that behavior. The structural form (mandatory pre-authorization) is the same as SUBSAFE's; the missing piece is the equivalent of SUBSAFE's lifecycle audit and recurring training that keep the certification real in operation.<sup>3</sup>

## OPEN PROBLEMS

Drafted alongside SUBSAFE (Case 173) and the WHO Surgical Checklist (Case 23), iPLEDGE is the most analytically useful "mixed" case in the v2 sweep. It is not a failure of the form (the form has demonstrated successes); it is evidence that the form alone does not deliver capability when the capability depends on a behavior the program does not instrument. The case is the case-grounded basis for the new subobjective **Judgment under inadequate evidence** (4.4) — the practitioner designing a REMS-class program has to decide on the available evidence what additional adherence-support architecture the mechanism needs in order to deliver, and the iPLEDGE evidence is the worked example of why the question matters.

*The authorization mechanism is operating as designed. The capability sits downstream of the authorization, and the program does not reach that far.*

EDITORS' SYNTHESIS OF SHIN ET AL. (2011).

REFERENCES

1. FDA, iPLEDGE program documentation (2006 – present) — REMS architecture and enrollment requirements.

2. Sullivan et al. (2003), House Science Committee statement on SUBSAFE — the structural counterpoint (Case 173).

3. Tkachenko, E., Singer, S., Sharma, P., Barbieri, J., & Mostaghimi, A. (2019), "US Food and Drug Administration reports of pregnancy and pregnancy-related adverse events associated with isotretinoin," JAMA Dermatology, 155(10):1175–1179 — source of the 218 to 310 annual reports after 2011.

4. Shin, J., Cheetham, T. C., Wong, L., Niu, F., Kass, E., Yoshinaga, M. A., Sorel, M., McCombs, J. S., & Sidney, S. (2011), "The impact of the iPLEDGE program on isotretinoin fetal exposure in an integrated health care system," Journal of the American Academy of Dermatology, PMID:21565419.

5. Collins, M. K., Moreau, J. F., Opel, D., Swan, J., Prevost, N., Hastings, M., Schwarz, E. B., & Ferris, L. K. (2014), "Compliance with pregnancy prevention measures during isotretinoin therapy," Journal of the American Academy of Dermatology, 70(1):55–59, PMID:24157382 — source of the 150 annual exposures figure.

THE LEARNING ENGINEERING LENS

LE INSIGHT

iPLEDGE is the productive counterpoint to SUBSAFE and the WHO Surgical Checklist. The structural form — mandatory pre-authorization gating a consequential action — is the same; the measured outcome is very different because the capability the program exists to enforce depends on patient adherence the program does not instrument. The “no significant decrease” finding is binding and survives into the case verbatim.

LENS APPROACH

iPLEDGE is the analytically useful “mixed” pre-authorization case (induced 4.4; LENS D3/PT5). LENS uses it in Domain 3 (Human-System Collaboration) for the subobjective **Delegation with revocation** (3.3) — what the authorization mechanism delegates and what it does not — and in Domain 4 (Test and Evaluation) for the subobjective **Judgment under inadequate evidence** (4.4) on whether the form will deliver in a specific context. Direct pair with SUBSAFE (Case 173) at the form-vs-context layer; pair with WHO Surgical Checklist (Case 23) at the mandatory-mechanism layer.

| DOMAIN              | MODES AT STAKE | LENS COMPETENCY            | PROBLEM TYPE | INDUCED | LEO   |
|---------------------|----------------|----------------------------|--------------|---------|-------|
| PHARMA SAFETY       | G D N          | D1 D2 D3 D4 D5             | 5            | 4.4     | 4 · 3 |
| REGULATORY PROGRAMS | →              | →                          |              |         |       |
| CLINICAL CARE       |                | Human-System Collaboration |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify in advance what the authorization mechanism reaches and what it does not — for iPLEDGE, the authorization controls dispensing but does not reach into patient adherence — so the gap is named at design time rather than discovered in outcome data.
- Pair the authorization architecture with adherence support that instruments the behavior the capability depends on: in this case, contraception adherence, not just contraception declaration.
- Design the outcome metric around the harm the program exists to prevent (fetal exposure), not around the mechanism’s own throughput (enrollment counts,

IN OPERATION / AFTER

- Treat a “no significant decrease” finding as program evidence about the mechanism’s reach, not as an argument against authorization architectures in general; SUBSAFE and the WHO checklist demonstrate the form can deliver under the right conditions.
- Re-engineer the adherence-support layer based on the documented driver of failure — for iPLEDGE, missed pills and inconsistent condom use — rather than tightening the authorization layer that is already working as specified.
- Carry the exposure figures with their sources attached — roughly 150 a year in Collins et al., 218 to 310 a year in the FDA reports — since iPLEDGE is operating as

authorization codes issued), so the evaluation is of the capability and not of the bureaucracy.

designed and the design's reach is what the evidence is about.

### REFLECTION QUESTIONS

1. Identify a mandatory pre-authorization program in your domain. What does the authorization mechanism actually reach, and what does it not? Where in the patient or operator behavior does the capability the program exists to deliver sit downstream of the authorization?
2. Specify the adherence-support architecture you would add to convert an iPLEDGE-class authorization into a delivered capability. For isotretinoin the gap is documented (missed pills, inconsistent condom use); what is the analog in your context?

**WHO BUILDS THIS** governance, ethics & measurement · systems & human-factors engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · task analysis · design prototyping · incident analysis · safety dashboards

# TeamSTEPPS

**T** Training Gap   **N** Normalization of Deviance

**IMPACT** Improved teamwork, communication, and patient-safety culture across diverse settings; on-time first surgical starts +21% in one urology service (Weld 2016); 1,500+ hospitals, 5,000+ master trainers, 300,000+ staff trained in the first eight years

## IN BRIEF

TeamSTEPPS — developed jointly by the Agency for Healthcare Research and Quality and the Department of Defense and released in 2006 — is the healthcare analog of Crew Resource Management (Case 117): an evidence-based team-training framework distilled from decades of aviation, military, and nuclear research and adapted for clinical settings. It trains four core competencies: communication, leadership, situation monitoring, and mutual support. It is explicitly the translation pathway from high-reliability research into bedside practice — the cross-domain capability transfer LENS is built to teach. Because its implementation infrastructure was funded as part of the program, TeamSTEPPS moved from research to scaled deployment in years rather than decades. In its first eight years more than 1,500 hospitals implemented it, over 5,000 master trainers were certified, and more than 300,000 health professionals were trained.

## BACKGROUND

Decades of research in aviation, the military, and nuclear power had shown that teamwork — not just individual expertise — drives safety in high-consequence work, and the Joint Commission's root-cause data had made communication failure the leading cause of sentinel events. What clinical settings lacked was a structured, evidence-based way to build those team skills. The evidence existed and the diagnosis was clear; what was missing was the translation — a route from the high-reliability research base into a curriculum a hospital could actually run at the bedside.<sup>0</sup>

## THE INTERVENTION

In 2006, the Agency for Healthcare Research and Quality and the Department of Defense jointly released TeamSTEPPS — Team Strategies and Tools to Enhance Performance and Patient Safety — the healthcare analog of Crew Resource Management. It trains four core competencies: communication, leadership, situation monitoring, and mutual support, with a structured curriculum and ready-made implementation materials. The joint AHRQ-DoD authorship mattered: it paired a healthcare research agency's clinical reach with the military's deep team-training experience, so the framework arrived already grounded in both the evidence base and the practicalities of running it.<sup>1</sup>

## HOW IT WORKED


TeamSTEPPS is explicitly the translation pathway from high-reliability-organization research into clinical practice — decades of cross-domain evidence adapted for the bedside. Crucially, its implementation infrastructure (master-trainer programs, toolkits, an institutional support center) was funded as part of the intervention, so adopting organizations had a route from training to sustained practice rather than a binder on a shelf. Funding the implementation alongside the curriculum is precisely what compressed the usual decades-long gap between a proven method and its scaled use, because the path from adoption to sustained practice was paid for in advance.<sup>2</sup>

## THE EVIDENCE

Studies across diverse settings report improved teamwork, communication, and patient-safety culture; on one urology service, Weld et al. measured a 21% improvement in on-time first surgical starts ( $P < .001$ ), alongside a 12.7-minute reduction in mean case time and a fall in patient-safety issues from 16% to 6%. In its first eight years the framework reached more than 1,500 hospitals, certifying over 5,000 master trainers who trained more than 300,000 health professionals; AHRQ has continued to develop it, releasing TeamSTEPPS 2.0 in 2014 and 3.0 in 2023. The continued development into a third version is itself evidence the transfer took hold — a framework that is still being maintained and revised nearly two decades on is one institutions kept using, not one that was issued and abandoned.<sup>3</sup>

## WHAT TRANSFERRED

TeamSTEPPS is the canonical evidence that cross-domain capability transfer is engineerable — and that the long implementation gap can be dramatically shortened when the transfer is funded as part of the intervention rather than as an afterthought. Its four competencies map directly onto the argument that capability engineering is itself a teachable discipline. The deeper lesson is that the bottleneck in cross-domain transfer is rarely the knowledge, which already existed, but the funded path that carries it from research into routine practice — the part programs most often leave unbudgeted.<sup>4</sup>



*TeamSTEPPS represents the translation pathway from high-reliability research into clinical practice.*

EDITORS' SYNTHESIS DRAWING ON AHRQ TEAMSTEPPS 3.0 (2023) AND SALAS ET AL.

REFERENCES

1. Salas, E., DiazGranados, D., Klein, C., Burke, C. S., Stagl, K. C., Goodwin, G. F. & Halpin, S. M. (2008), "Does Team Training Improve Team Performance? A Meta-Analysis," Human Factors 50(6):903–933 — the cross-domain team-training evidence base.

2. Umscheid, C. & Haugstetter, M. (2023), "With TeamSTEPPS 3.0, AHRQ Refreshes a Landmark Patient Safety Training Curriculum," AHRQ Views (12 September) — the DoD partnership and the first-eight-years adoption figures (1,500 hospitals, 5,000 master trainers, 300,000 professionals).

3. Weaver, S. J., Dy, S. M. & Rosen, M. A. (2014), "Team-training in healthcare: a narrative synthesis of the literature," BMJ Quality & Safety 23(5):359–372 — bundled team-training interventions, with tools and organisational change to support transfer, show the largest effects.

4. Weld, L. R. et al. (2016), "TeamSTEPPS Improves Operating Room Efficiency and Patient Safety," American Journal of Medical Quality 31(5):408–414 — the on-time first-start figure.

5. AHRQ, TeamSTEPPS 3.0 Curriculum (2023) — the framework and four competencies.

THE LEARNING ENGINEERING LENS

LE INSIGHT

TeamSTEPPS is the canonical evidence that cross-domain capability transfer is engineerable and that it can shorten the implementation gap dramatically when the transfer is funded as part of the intervention rather than as an afterthought. The four competencies map directly to the LENS curriculum’s argument for what capability engineering competence looks like as a deliverable.

LENS APPROACH

LENS treats TeamSTEPPS in LEN 8 as the canonical cross-domain transfer case and in LEN 1 as evidence that the program’s core proposition — that learning engineering is a discipline — has institutional precedent. Studio projects (LEN 10) reference TeamSTEPPS as the worked example of cross-domain adaptation methodology.

| DOMAIN                                | MODES ADDRESSED             | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|-----------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| HEALTHCARE                            | → <div>T</div> <div>N</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 4            | 6.4     | 5   |
| Navigating Sociotechnical Constraints |                             |                                                                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Distill the cross-domain evidence base into a structured curriculum of a few teachable competencies rather than leaving adopters to translate the research themselves.
- ▶ Pair authorship across the source and target domains (as AHRQ and DoD did) so the framework arrives grounded in both the evidence and the practicalities of running it.
- ▶ Fund the implementation infrastructure — master-trainer programs, toolkits, a support center — as part of the intervention, not as an afterthought.

IN OPERATION / AFTER

- ▶ Maintain and revise the framework over time (as with successive TeamSTEPPS versions) so it stays in use rather than being issued and abandoned.
- ▶ Track concrete operational gains (such as on-time first surgical starts) alongside culture measures, so the transfer’s effect is shown, not assumed.
- ▶ Identify and budget the specific path from adoption to sustained practice when scaling, since that unbudgeted step is the usual bottleneck in cross-domain transfer.



### REFLECTION QUESTIONS

1. What is the cross-domain transfer your domain has not yet executed? What evidence base in another industry should it draw on?
2. TeamSTEPPS funded its implementation infrastructure. Design the implementation-infrastructure budget for an equivalent transfer in your domain.
3. TeamSTEPPS shortened the usual decades-long gap because the path from adoption to sustained practice was paid for in advance. Identify a proven method in your domain that has not scaled, and name the specific unbudgeted step that is blocking it.

**WHO BUILDS THIS** learning science & instructional design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards

# Team Science Training for Clinical and Translational Scientists

**K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Colorado CTSA team-science training (221 registrants; pre/post survey N=117/46) reported statistically significant improvement ( $p < 0.05$ ) across all three TEAMS instrument competencies — Team Planning, Managing a Team, and Interpersonal Relations — with the largest gains in structured/tool-based domains and the smallest in Interpersonal Relations

## IN BRIEF

The Colorado Clinical and Translational Sciences Institute built a structured team-science training program — a prerequisite online module, six interactive workshops, active-learning exercises — and evaluated it across three 2020–2022 cohorts (221 registered participants; 117 pre-program and 46 post-program survey respondents) using a pre/post design. The evaluation introduces and validates a measurement instrument (TEAMS) that resolves team-science capability into three components: Team Planning, Managing a Team, and Interpersonal Relations. Participants showed statistically significant improvement ( $p < 0.05$ ) across all three components — with the largest gains in the structured / tool-based domains (Team Planning, Managing a Team) and the smallest gain in Interpersonal Relations, the only component scored on frequency of use rather than familiarity. The case treats interdisciplinary collaboration as an engineerable, measurable capability — the amended Domain 3 sub-competency (collaboration as a unit of measurement, distinct from any individual operator) — and is honest about the interpersonal half moving least. The honest limit the authors name: outcomes are self-reported and the design lacks longitudinal tracking and integration metrics; this is perceived-competency-gain evidence, not yet demonstrated downstream research- collaboration impact.

## BACKGROUND

Clinical and translational research is the structural example of a discipline whose capability sits in the collaboration itself rather than in any single team member. A protocol moves across investigators, study coordinators, biostatisticians, clinical staff, and regulatory specialists; the science depends on the team's coordination as much as on any individual's expertise. The capability question the case addresses is whether interdisciplinary collaboration can be trained — and whether the resulting capability can be measured.<sup>o</sup>

## THE INTERVENTION

The Colorado Clinical and Translational Sciences Institute built a structured team–science training program: a prerequisite Teaming 101 web module and six interactive workshops. It ran across three 2020–2022 cohorts with 221 registered participants; the pre/post design drew 117 pre–program and 46 post–program survey respondents. The evaluation paper’s central methodological contribution is the validation of a measurement instrument — TEAMS — that resolves team–science capability into three competency components: Team Planning (the structured front–end of collaboration), Managing a Team (the structured operational half), and Interpersonal Relations (the unstructured half, scored on frequency of use because participants already knew it).<sup>1</sup>

## HOW IT WORKED

Pre/post comparison showed statistically significant improvement ( $p < 0.05$ ) across all three components. The pattern of gains is what makes the case teachable: the largest gains were in the structured / tool–based domains — Team Planning and Managing a Team — and the smallest gain was in Interpersonal Relations. The authors do not soften this; the paper states it plainly, and the smaller–gain–on–the–hardest–half finding is the decisive honest result.<sup>2</sup>

## THE EVIDENCE

The honest limit carries into the case. The authors name the response bias and the self–report; the design shows the rest. The TEAMS instrument and the pre/post comparison measure self–reported perceived competency, not downstream collaboration outcome. The design did not track participants longitudinally into their next collaborative projects, did not measure integration metrics on the projects they ran after the training, and did not compare to a non–trained control. The evidence is one of the few published pre/post evaluations of a structured team–science training program; it is not closed proof of downstream collaboration impact, and the case carries the qualification rather than collapsing it.<sup>3</sup>

## WHAT TRANSFERRED

In pair with IPE (Case 28) and Implementation Science Training (Case 41), the Colorado CTSA case is the small–tier intervention companion to two frontier cases. Team–science training is one of the few cases in the corpus that operationalizes “collaboration as the unit of measurement” — the LEO — with a validated instrument and structured evaluation. The case is a worked example of how a discipline can convert collaboration capability from an aspiration to a measurable target, while preserving the honest qualifications about what the measurement does and does not establish.

*The largest gains were in the structured half. The hardest half — interpersonal relations — moved least, and the paper states it plainly.*

EDITORS’ SYNTHESIS OF COLORADO CTSA TEAM–SCIENCE TRAINING EVALUATION.

## REFERENCES

1. Mook, A., Knerich, V., Komaie, G., Cicutto, L., & Cross, J. (2025), “Team science training for clinical and translational scientists: An assessment of effectiveness,” *Journal of Clinical and Translational Science* 9(1):e158, doi:10.1017/cts.2025.10088.
2. Falk–Krzyszinski et al. (2011), “Mapping a research agenda for the science of team science,” *Research Evaluation* 20(2):143–156 — broader team–science literature backdrop.
3. National Research Council (2015), *Enhancing the Effectiveness of Team Science* (National Academies Press) — the National Academies team–science synthesis.
4. Paired cases in this volume: IPE evidence gap (Case 28), Implementation Science Training (Case 41) — the frontier/measurement companions.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Colorado CTSA team-science training program is one of the few cases in the corpus that operationalizes “collaboration as a unit of measurement.” The TEAMS instrument resolves the capability into three components; structured training moved all three with  $p<0.05$  but moved interpersonal relations least. The honest limit — self-reported perceived competency, no longitudinal tracking, no control — survives into the case.

LENS APPROACH

Team-science training is the case-grounded basis for the collaboration-as-measurement LEO (induced 4.3; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the team-coordination redesign and in Domain 4 (Test and Evaluation) for the TEAMS instrument validation and the perceived-competency vs. downstream-impact discipline. Direct pair with Cases 28 (IPE) and 41 (implementation science training) — the intervention with measurement against two frontier cases where the measurement is the gap.

| DOMAIN                                                                                                          | MODES ADDRESSED           | LENS COMPETENCY                                                       | PROBLEM TYPE | INDUCED | LEO   |
|-----------------------------------------------------------------------------------------------------------------|---------------------------|-----------------------------------------------------------------------|--------------|---------|-------|
| <div>CLINICAL/<br/>TRANSLATIONAL<br/>RESEARCH</div> <div>TEAM SCIENCE</div> <div>HEALTHCARE<br/>WORKFORCE</div> | <div>K</div> <div>N</div> | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 4            | 4.3     | 4 · 5 |
|                                                                                                                 | →                         | → Navigating Sociotechnical Constraints                               |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Resolve collaboration capability into measurable components — Team Planning, Managing a Team, Interpersonal Relations — and structure the training so each component has tools, exercises, and an assessment instrument.
- ▶ Validate the measurement instrument before relying on it for evaluation; the TEAMS instrument’s three-component structure is the methodological contribution.
- ▶ Acknowledge in the program design that the interpersonal component moved least and did not reach the intended level of use; the authors’ own recommendation is more role-play and real-situation practice, not more lecture.

IN OPERATION / AFTER

- ▶ Report the pattern of gains, not just the aggregate; the smaller-gain-on-interpersonal finding is the honest signal about what the structured training can and cannot do.
- ▶ Commission longitudinal follow-up that tracks participants into their next collaborative projects and measures integration outcomes — the case names the gap and the next study addresses it.
- ▶ Carry the “collaboration as the unit of measurement” sub-competency into the curriculum: the Domain 3 sub-competency is grounded in evidence here, and the team-science training program is one of the few cases in the corpus that operationalizes it.

## REFLECTION QUESTIONS

1. Identify a collaboration-dependent capability in your domain. What three or four components would you resolve it into for measurement purposes, and what instrument would you validate to test the resolution?
2. Specify the longitudinal follow-up you would build into the next iteration of a team-science training program — what integration metric, on what cadence, against what comparison — to convert perceived-competency-gain evidence into downstream-collaboration-impact evidence.
3. The Colorado CTSA finding is that the structured half moves first and the interpersonal half moves least. What does that imply about the curriculum you would design — and about which components are realistic targets for a single-program training intervention vs. an institutional / cultural intervention?

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

## CASE 41

IMPLEMENTATION SCIENCE

TRANSLATIONAL WORKFORCE

EDUCATION AT SCALE

2020S

# Translational-Workforce Training — Stated Goals Outrunning Named Practice

**K** Knowledge & Institutional Memory    **N** Normalization of Deviance

**IMPACT** Survey of CTSA-funded TL1 training programs (N=50, fielded 2019): 80% name collaboration and/or team science among their stated goals and 46% name multi/inter/cross-disciplinary training, but goals that name a specific practice are far rarer — program evaluation 38%, experiential learning 24%, and competency-based training named by just 6% of the 48 programs describing what distinguishes them

## IN BRIEF

Implementation science — the discipline of moving validated evidence into operational practice — is one of the most consequential cross-domain-transfer competencies in contemporary medicine and adjacent fields, and the CTSA T32 / TL1 program survey (N=50 programs) is the closest national picture of what that training arm says it is for. The headline finding is the same structural pattern as the IPE case (Case 28) at smaller scale: 80% of programs name collaboration and/or team science among their goals and 46% name multi/inter/cross-disciplinary training, but the goals that name a specific practice are rarer — program evaluation 38%, experiential learning 24%, and competency-based training in 6% of the programs saying what sets them apart. The picture is of a field that has converged on what interdisciplinary translation training should aim for, but where the operational practices lag the stated goals — the same enthusiasm-ahead-of-evidence pattern as IPE. The case is the workforce-training counterpart to Case 13 (the “17-year gap” between research evidence and practice adoption); the v2 frame is the gap between what implementation-science training programs **say** they do and what they **operationally** do.

## THE SHIFT

Implementation science exists because the gap between research evidence and operational practice in medicine — the so-called “17-year gap” the corpus documents at Case 13 — is large enough to constitute a discipline-level capability question. Moving validated evidence into clinical practice reliably is itself a competency that has to be designed, trained, and evaluated. The CTSA (Clinical and Translational Science Awards) program is the primary U.S. mechanism for building that workforce, and the TL1 / T32 programs are its training arm.<sup>9</sup>

## WHAT IS EMERGING

The case's evidence base is a survey of CTSA-funded TL1 training programs fielded in 2019 and published in 2021 (N=50 of 51 invited hubs, a 98% response; 48 answered the distinguishing-goals question). The survey asked two open questions: what the program's stated goals are, and which goals distinguish it from other NIH-funded training. General aims are where the programs converge: of the 50 answering the stated-goals question, 80% name collaboration and/or team science and 46% name multi/inter/cross-disciplinary training. Goals that name a specific practice are rarer — program evaluation 38% and experiential learning 24% on the same question, and, on the distinguishing-goals question, competency-based training and/or assessment named by 3 of the 48 respondents (6%), the joint-lowest theme in the table. The survey does not reach past the goal statement: curriculum content and evaluation methods were reported separately.<sup>1</sup>

## THE CAPABILITY QUESTION

The structural form is the same pattern IPE shows at field scale (Case 28): the field has converged on what interdisciplinary translation training should aim for, and the practice-specific goals trail the general ones. At the implementation-science training scale, the lag has specific consequences: programs that intend to produce graduates capable of moving evidence into practice often do not systematically assess whether they have produced that capability, and the field's measurement of its own workforce-development effectiveness is correspondingly sparse.<sup>2</sup>

## EARLY EVIDENCE

What the case teaches is the operational-practice gap as a designable target, not as a curricular failure. The low uptake of the specific competency practices is not an indictment of the programs that fall on the wrong side of it; it is a finding about the operational-infrastructure investment that competency-based training and assessment require, and that field-scale stated-goal convergence does not by itself produce. The pair with Case 13 (the 17-year gap) gives the case its frame: the implementation-science workforce is the recovery mechanism for the research-to-practice gap, and the operational capacity of that workforce is itself the capability the training programs are trying to build.<sup>3</sup>

## OPEN PROBLEMS

In the multidisciplinary-translation trio (Cases 40 + 28 + 41), implementation-science training sits between team science (where the measurement is possible at program scale) and IPE (where the measurement is the field-scale gap). The trio teaches the enthusiasm-evidence-gap sub-competency from three angles — the program-scale success, the field-scale gap, and the operational-practice gap inside training programs themselves — and it teaches the subobjective **Judgment under inadequate evidence** (4.4) by example: practitioners designing implementation-science training programs have to decide what to build on incomplete evidence of what works, while the field is still building the evidence architecture that would let them decide better.

*Stated goals converge on the right targets. Operational practices lag. The gap is designable, not a curricular failure.*

EDITORS' SYNTHESIS OF THE CTSA TL1 PROGRAM-GOALS STUDY (2021).

## REFERENCES

1. Morris, Wooding, & Grant (2011), "The answer is 17 years, what is the question: understanding time lags in translational research," *Journal of the Royal Society of Medicine — the original 17-year-gap source* for Case 13.
2. Sancheznieto, F., Sorkness, C. A., Attia, J., et al. (2021), "Clinical and translational science award T32/TL1 training programs: program goals and mentorship practices," *Journal of Clinical and Translational Science* 6(1):e13, doi:10.1017/cts.2021.884.
3. Brownson, Colditz, & Proctor (2018), *Dissemination and Implementation Research in Health* (2nd ed.) — the broader implementation-science synthesis.
4. Paired cases in this volume: Team-science training (Case 40), IPE evidence gap (Case 28).

## THE LEARNING ENGINEERING LENS

## LE INSIGHT

The CTSA TL1 program survey is the workforce-training instance of the enthusiasm-evidence gap pattern: 80% of programs name collaboration and team science among their goals, while the goals that name a practice are named by far fewer — competency-based training and/or assessment by 3 of the 48 programs answering the distinguishing-goals question (6%), experiential learning by 24%, program evaluation by 38%. The implementation-science workforce is the recovery mechanism for the 17-year research-to-practice gap; the case names the gap inside the recovery mechanism itself.

## LENS APPROACH

Implementation-science training is the frontier workforce-training case (induced 6.4; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the cross-disciplinary translation discipline; in Domain 4 (Test and Evaluation) for the enthusiasm-evidence-gap sub-competency; and in Domain 2 (Iterative Development) for the subobjective **Judgment under inadequate evidence** (4.4) as practitioners design training programs while the field-scale evidence is still being built. Pair with Cases 40, 28 as the multidisciplinary-translation trio; workforce-training counterpart to Case 13.

| DOMAIN                  | MODES AT STAKE    | LENS COMPETENCY                         | PROBLEM TYPE | INDUCED | LEO   |
|-------------------------|-------------------|-----------------------------------------|--------------|---------|-------|
| IMPLEMENTATION SCIENCE  | <b>K</b> <b>N</b> | D1 D2 D3 D4 <b>D5</b>                   | 4            | 6.4     | 4 · 5 |
| TRANSLATIONAL WORKFORCE | →                 | → Navigating Sociotechnical Constraints |              |         |       |
| EDUCATION AT SCALE      |                   |                                         |              |         |       |

## APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- ▶ Specify the competency-based assessment instrument before launching the training program; the gap the survey documents is partly the result of stated competencies that were never operationalized into measurement.
- ▶ Build experiential learning into the operational structure of the program, not as an add-on; the survey found experiential learning named by only 24% of programs — one of the operational dimensions that lags.
- ▶ Plan program evaluation as a deliverable of the training program itself, with the cadence, instrument, and reporting venue named at launch.

## IN OPERATION / AFTER

- ▶ Report the operational-practice gap honestly: stated goals are the convergence point, operational practices are the lag, and the lag is the addressable target.
- ▶ Treat the multidisciplinary-translation trio (Cases 40, 28, 41) as a unit; the three-angle teaching of the enthusiasm-evidence pattern is more useful than any single case can be.
- ▶ Connect the case explicitly to Case 13 (the 17-year gap): the implementation-science workforce is the recovery mechanism for that gap, and the operational capacity of the workforce is the capability that has to be built — which is the case's pedagogical point.



## REFLECTION QUESTIONS

1. Identify a training program in your domain whose stated goals include competency-based assessment, program evaluation, or experiential learning. What proportion of those stated goals are operationalized into specific instruments, cadences, and reporting structures — and which are at the goal-statement layer only?
2. Specify the competency-based assessment instrument you would build into the next iteration of an implementation-science training program. The survey found a competency-based curriculum named by just 3 of the 48 programs answering the distinguishing-goals question — the instrument is what is missing more than the intent; what is the instrument?
3. The implementation-science workforce is the recovery mechanism for the 17-year research-to-practice gap (Case 13). What is the analog in your domain — the workforce whose operational capacity is the recovery for a documented systemic gap — and what is the case for investing in that workforce's training architecture?

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

# Australian Hospital-Pharmacy Technician Role Redesign

**D** Designed Out   **N** Normalization of Deviance   **H** Human-System Interface

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** The 2016 SHPA Redesign Project — a six-month national scoping study — found pharmacy technicians employed in 95% of Australian hospital pharmacies but held to technical rather than clinical work, with no national training or regulatory framework for expanded scope; a September 2018 single-hospital staff survey (61 of 122 responding), published in 2021, then found 82% supported technician final accuracy checking and 81% supported technician drugs-of-addiction register entries, while technician-taken medication histories split pharmacists (38% supportive) from technicians (92%)

## IN BRIEF

Australian hospital pharmacies faced the structural problem most healthcare workforces meet at some point: the pharmacist's capacity was being absorbed by dispensing-accuracy checking, which crowded out the clinical work that requires pharmacist judgment. The 2016 Pharmacy Technician and Assistant Role Redesign project did not expand scope; it mapped the ground for expanding it — a six-month study of the national workforce built from a literature review, a survey, focus groups, interviews and case studies. It found technicians in 95% of hospital pharmacies, roles varying widely between states, no national regulation of the workforce, and board guidance confining technicians to work not requiring professional judgment. A cross-sectional survey fielded in September 2018 at one Victorian metropolitan teaching hospital and published in the *Journal of Pharmacy Practice and Research* in 2021 then measured what staff would accept: 82% supported technician final accuracy checking, 81% supported technician drugs-of-addiction register entries, and technician-taken medication histories split pharmacists (38%) from technicians (92%). The evidence-tier flag renders under the title: the operational figures come from a program-report rather than an independent multi-site audit. Future validation will continue on long-term safety durability and on whether the role-redesign pattern generalizes to other healthcare-workforce roles.

## BACKGROUND

The structural problem the case sits inside is the C3 thread where monitoring or checking duties are loaded onto the role that needs to be doing the clinical work — and the checking work absorbs the capacity. Australian hospital pharmacies in the early 2010s were in this

pattern: the pharmacist's day was consumed by dispensing-accuracy checking, and the clinical-support work that requires pharmacist judgment was crowded out. The induced 3.2 framing is the precise diagnosis: the unsupportable role is the pharmacist as primary accuracy-checker.<sup>0</sup>

#### THE INTERVENTION

The three roles the project put on the table were final accuracy checking, drugs-of-addiction management, and clinical support. The design move is not the addition of a checker but the redistribution of authority: the technician is given the role that the pharmacist was doing, with the training, certification, and supervisory structure to back it. The pharmacist is freed to do the clinical-judgment work that the system needed pharmacists for. The redesign is the C3 role-design intervention the induced framework calls for at the small tier.<sup>1</sup>

#### HOW IT WORKED

What the project produced was a map, not an outcome series: technicians in 95% of hospital pharmacies, roles differing state by state, no national regulation of the workforce, and board guidance holding technicians to work that does not require professional judgment. The Anderson et al. cross-sectional survey, fielded over four weeks in September 2018, carried it into the workforce-attitudes layer — of 122 staff invited at one Victorian teaching hospital, 61 responded; 82% supported technician final accuracy checking, 81% supported technician drugs-of-addiction entries, and fewer than 10% thought technicians were adequately used as things stood.<sup>2</sup>

#### THE EVIDENCE

The evidence-tier flag is binding. The project report is a scoping document and reports no throughput, turnaround or error outcomes; the acceptance survey is peer-reviewed but covers one hospital and 61 respondents. No multi-site outcome evaluation of the expanded role exists in the Australian record. The cross-sectional survey is peer-reviewed but measures attitudes rather than outcomes. The case is included because the pattern is teachable and the practitioner literature is consistent; the magnitudes should travel with the tier flag intact. Future validation will continue.<sup>3</sup>

#### WHAT TRANSFERRED

What the case teaches at the LENS layer is the role-design move as a sociotechnical-constraints intervention. The pharmacist-as-checker pattern was not a property of any individual; it was a property of the role architecture, and the redesign treated the architecture as the design variable. The Domain-5 frame applies: the work was to redraw the authority gradient and team-coordination boundary so that the system's capability — fast, safe dispensing plus pharmacist clinical judgment — emerged from the team rather than depending on the pharmacist's individual heroism. The case is also a Gap-5 echo: a non-US small-tier role-redesign programme — mapped and accepted, but not yet measured in throughput or safety outcomes.

*The pharmacist-as-checker pattern was a property of the role architecture, not the pharmacist. The redesign treated the architecture as the design variable.*

EDITORS' SYNTHESIS OF THE 2016 PROJECT REPORT AND ANDERSON ET AL. (2021).

REFERENCES

1. Pharmacy Board of Australia and SHPA practice statements limiting technician work to activities not requiring professional judgement — the regulatory framing the project sits inside.

2. Society of Hospital Pharmacists of Australia (November 2016), “Exploring the role of hospital pharmacy technicians and assistants to enhance the delivery of patient centred care: A White Paper on the findings and outcomes of the Pharmacy Technician and Assistant Role Redesign within Australian Hospitals (Redesign) Project.”

3. Anderson et al. (2021), “Perceptions of hospital pharmacists and pharmacy technicians towards expanding roles for hospital pharmacy technicians: a cross-sectional survey,” *Journal of Pharmacy Practice and Research* 51(3):216–224, doi:10.1002/jppr.1697 — fielded September 2018; 61 of 122 staff (50%) responded.

4. Boughen, Sutton, Fenn, & Wright (2017), “Defining the Role of the Pharmacy Technician and Identifying Their Future Role in Medicines Optimisation,” *Pharmacy* 5(3):40 — UK companion analysis.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Australian pharmacy-technician redesign is a small-tier C3 role-redesign case at the stage before the outcome data: a national scoping study that named the role architecture, and a single-hospital survey showing 82% acceptance of technician final accuracy checking. The evidence base is a program report plus a peer-reviewed attitudes survey; no throughput or safety outcomes have been published. Future validation ongoing.

LENS APPROACH

Australian pharmacy-technician redesign is the C3 small-tier role-redesign case (induced 3.2 and 4.3; LENS D5/PT3). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for LEO-5 — the work was to redraw the authority gradient — and in Domain 1 (Systems Analysis) for LEO-1, since the role-architecture diagnosis was the precondition for the redesign. Pair with Case 32 (annual-screening UI redesign) for the small-tier C3 thread; Gap-5 echo as a non-US case.

| DOMAIN                                                                                | Modes Addressed                        | LENS Competency                                                       | Problem Type | Induced | LEO   |
|---------------------------------------------------------------------------------------|----------------------------------------|-----------------------------------------------------------------------|--------------|---------|-------|
| <div>HOSPITAL PHARMACY</div> <div>HEALTHCARE WORKFORCE</div> <div>ROLE REDESIGN</div> | <div>D</div> <div>N</div> <div>H</div> | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 3            | 3.2     | 1 · 5 |
| → Navigating Sociotechnical Constraints                                               |                                        |                                                                       |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Diagnose the role-architecture failure before redesigning: which capability does the system need from which role, and where is checking work absorbing capacity that the system needs elsewhere.
- ▶ Redistribute authority, not just tasks: the technician's expanded scope must come with training, certification, and supervisory structure, not just a new line in the procedure manual.
- ▶ Specify both throughput and safety metrics before the redesign ships, so the joint movement (or trade-off) of both is what the report has to defend.

IN OPERATION / AFTER

- ▶ Treat the program-report figures as practice-synthesis evidence; carry the tier flag into any generalization to other sites or other workforce roles.
- ▶ Track the safety metric from the first day of the expanded role, not after it; without a pre-change baseline the redesign cannot answer the only question that matters, which is whether checking accuracy held.
- ▶ Use the workforce-acceptance evidence (2021 survey) to identify where the redesign meets resistance, and stage the rollout against that resistance rather than ignoring it.

## REFLECTION QUESTIONS

1. Identify a healthcare role in your setting where checking or monitoring work is absorbing the capacity the system needs for clinical judgment. What is the technician- or assistant-equivalent role redesign that would redistribute the authority, and what training and supervisory structure would it require?
2. Specify the throughput and safety metrics you would commit to before the redesign ships, so the joint movement of both is what the report has to defend. The Australian project teaches by omission — it mapped the workforce and measured acceptance, and left the outcome question open — so what would you measure that it did not?
3. The acceptance figures come from one hospital and 61 respondents. What independent or multi-site evidence would you require before treating them as generalizable to other workforce roles (nursing, radiology, lab) or other countries?

**WHO BUILDS THIS** systems & human-factors engineering · safety science & organizational analysis · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · incident analysis · safety dashboards · usability testing · cognitive walkthroughs

# Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface

**H** Human-System Interface    **N** Normalization of Deviance

**EVIDENCE TIER** mixed-evidence-tier — source confidence flagged; future validation ongoing.

**IMPACT** A Rwanda mHealth monitoring system equipped community health workers with mobile decision support for maternal and newborn care; the published evaluations were mixed — uptake of maternal and newborn services rose only where the system was paired with added training, supervision, and equipment, while an interrupted time series on national survey data found no significant change in uptake, coverage already sitting above 90% — with the technology designed to extend the CHW's diagnostic-and-referral role rather than replacing it

## IN BRIEF

Between 2013 and 2018 Rwanda's Ministry of Health, working with research partners, deployed an mHealth monitoring system that placed mobile-phone-mediated decision support into the hands of community health workers (CHWs) responsible for maternal and newborn care — pregnancy registration, follow-up across the continuum of care, and emergency alerts to the ambulance system. Peer-reviewed evaluations were mixed: uptake rose only where the system was paired with added support (Ruton et al., 2018), while an interrupted time series on national survey data found no significant change in uptake at all (Hategeka et al., 2019). The case sits at the small (program/study) tier and teaches a precise pattern: capability delivered at the periphery of the formal health system, with the technology designed to extend the CHW's diagnostic-and-referral role rather than replace it. The evidence-tier flag is honest — two peer-reviewed evaluations pointing different ways, with broader impact claims still resting partly on practitioner reporting. As a non-US small-tier case it pairs naturally with the PEPFAR Sub-Saharan training-modality comparison as the African workforce-capability evidence the corpus has been thin on, and it carries the standing "future validation ongoing" language into print rather than overclaiming what one study can settle.

## BACKGROUND

Rwanda's health system, post-genocide, was rebuilt around community-level delivery: every village has elected CHWs who handle a defined scope of maternal, newborn, and child-

health work and refer upward when their scope is exceeded. The capability question for the mHealth program was specific — could a mobile decision-support tool, in CHW hands, increase the uptake of maternal and newborn health services in a way that the prior paper-based workflow had not.<sup>0</sup>

#### THE INTERVENTION

The system delivered structured prompts, reminders, and a referral pathway through the CHW's mobile phone: CHWs registered pregnancies, followed mother-infant pairs, and raised real-time alerts to the ambulance system, the health facility, the district hospital and the center. The design move worth naming is that the technology was scoped to extend the CHW's existing diagnostic-and-referral role, not to substitute for clinical judgment further up the chain. The CHW remained the capability interface; the phone was the cue and the record.<sup>1</sup>

#### HOW IT WORKED

Ruton et al. (2018) report increased uptake of maternal and newborn health services only where the mHealth system was paired with additional training, supervision, and equipment; RapidSMS alone was ineffective. An interrupted time series on national survey data (Hategeka et al., 2019) found no significant change on any indicator studied. Coverage of antenatal care, facility delivery and vaccination already sat above 90%, leaving a ceiling effect — but the authors are explicit that this does not explain the whole null: first-trimester ANC (50.3%) and four-or-more ANC visits (40.6%) had room to move and did not, and “further research is required to understand why RapidSMS had no impact on indicators where there was enough room for improvement.” The published evaluation is the anchor finding; downstream reporting (MIT News 2022 and subsequent AI-augmented maternal-care work) describes the program's continuation and adjacent developments but is journalism-tier, not investigation-grade. The case carries its evidence-tier flag honestly: two peer-reviewed process evaluations do not close the question of long-term outcome durability or generalization to other low-resource settings.<sup>2</sup>

#### THE EVIDENCE

The teaching point is the location of capability. The mHealth tool did not centralize the work; it instrumented the periphery. The decisive variables — household trust in the CHW, the CHW's standing in the village, the referral pathway upward — sit in the sociotechnical layer the tool could support but not constitute. This is the inverse of capability deployments that try to relocate judgment to the center; Rwanda's system kept judgment where it already was and used the technology to make that judgment more reliable and more visible to the formal health system.<sup>3</sup>

#### WHAT TRANSFERRED

What survives the evidence-tier flag is the structural pattern: capability extension at the frontline, with mobile-mediated decision support designed around an existing role, in a low-resource setting where the formal health system cannot reach every birth. Future validation ongoing — both replication in other African and South Asian settings, and longer-term outcome data linking process uptake to maternal and newborn outcomes downstream. The case is included with the flag intact and pairs with the PEPFAR cross-country modality comparison as the Sub-Saharan workforce-capability evidence the v2 corpus needs.

*The mHealth tool did not centralize the work. It instrumented the periphery — and kept judgment where it already was.*

EDITORS' SYNTHESIS OF HATEGEKA ET AL. (2019), RUTON ET AL. (2018), AND THE RWANDA CHW PROGRAM DOCUMENTATION.

REFERENCES

1. Rwanda Ministry of Health, community health program documentation and CHW scope-of-practice guidance, 2013–2018.

2. Ruton, H., Musabyimana, A., Gaju, E., Berhe, A., Grépin, K. A., Ngenzi, J., Nzabonimana, E., & Law, M. R. (2018), "The impact of an mHealth monitoring system on health care utilization by mothers and children: an evaluation using routine health information in Rwanda," *Health Policy and Planning* 33(8):920–927, doi:10.1093/heapol/czy066 — interrupted time series on 461 health centres, 2012–2016; RapidSMS plus additional support raised utilization, RapidSMS alone did not.

3. Hategeka, C., Ruton, H., & Law, M. R. (2019), "Effect of a community health worker mHealth monitoring system on uptake of maternal and newborn health services in Rwanda," *Global Health Research and Policy* 4:8, doi:10.1186/s41256-019-0098-y.

4. MIT News (2022), reporting on subsequent AI-augmented maternal-care work in Rwanda — journalism-tier companion to the peer-reviewed evaluation.

5. Cross-reference: Case 18 (PEPFAR HIV training-modality comparison) for the paired Sub-Saharan workforce-capability evidence.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Rwanda mHealth is a small-tier capability-extension case at the frontline: technology designed around an existing CHW role, with peer-reviewed evaluation showing uptake gains only when the system was paired with added support, and an interrupted national-survey time series showing no significant change at all. The evidence-tier flag is binding — two process studies do not close the durability or generalization question, and the broader impact claims rest partly on practitioner reporting. Future validation ongoing.

LENS APPROACH

Rwanda mHealth is the small-tier non-US frontline-capability case (induced 6.4; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the periphery- extension pattern — the technology supports the CHW's diagnostic-and-referral role without relocating judgment to the center — and in Domain 2 (Iterative Development) for the scope-the-tool-to-the-role design discipline. Direct pair with Case 18 (PEPFAR Sub-Saharan training-modality comparison) as the African workforce-capability evidence the v2 corpus needs.

| DOMAIN                                                                           | MODES ADDRESSED           | LENS COMPETENCY                                                                                                        | PROBLEM TYPE | INDUCED | LEO   |
|----------------------------------------------------------------------------------|---------------------------|------------------------------------------------------------------------------------------------------------------------|--------------|---------|-------|
| <div>GLOBAL HEALTH</div> <div>MHEALTH</div> <div>MATERNAL AND NEWBORN CARE</div> | <div>H</div> <div>N</div> | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> <div>Navigating Sociotechnical Constraints</div> | 4            | 6.4     | 2 · 5 |

APPROACHES TO CONSIDER

- DURING DEVELOPMENT

▶ Scope the tool to extend an existing frontline role rather than substitute for it; the CHW's village standing is the capability the technology can support but cannot create.

▶ Instrument the failure modes that the formal health system cannot see — a pregnancy never registered, a
- IN OPERATION / AFTER

▶ Track uptake of the targeted services as the primary process measure, and link to longer-term maternal and newborn outcomes as the data matures; resist treating short-term process gains as outcome evidence.

▶ Report the CHW workload and tool-acceptance trajectory separately from the headline uptake figures;



- danger sign at home, an ambulance never called — so the tool turns village-level observation into health-system-visible signal.
- ▶ Carry the evidence-tier flag through the deployment documentation: one peer-reviewed evaluation is the anchor finding, not the closure of the question.
- the periphery-extension pattern only holds if the frontline role remains sustainable.
- ▶ When asked whether the pattern generalizes, name the conditions Rwanda's system supplies (village-level CHW network, post-conflict institutional rebuild, single payer) before asserting transfer to other low-resource settings.

## REFLECTION QUESTIONS

1. Identify a frontline capability role in your context (CHW, field technician, ward nurse, line supervisor). What would it mean to scope a mobile decision-support tool to extend that role rather than relocate judgment to the center? Which failure modes only the frontline can see should the tool surface upward?
2. The case rests on two peer-reviewed evaluations that disagree, with broader claims supported by practitioner reporting. What is the minimum additional evidence — durability follow-up, replication in a second setting, linkage to maternal/newborn outcomes — you would require before treating the uptake finding as settled for a transfer decision?
3. Specify the village-level and health-system-level instrumentation you would put in place for a comparable deployment so that the frontline-extension pattern can be evaluated against an alternative deployment that relocates judgment to the clinic level.

**WHO BUILDS THIS** human factors & interaction design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · incident analysis · safety dashboards

# Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD

**G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** Japan's 2019 PMD Act amendment (enforced September 2020) introduced IDATEN (Improvement Design within Approval for Timely Evaluation and Notice): a pre-agreed device-improvement plan approved at initial approval, streamlined review for in-scope changes. The 2025 scoping review of PMDA-approved AI radiology software found 20 AI radiology products among 151 approved SaMD, and no product whose package insert reported complete case demographics — hedge preserved

## IN BRIEF

Japan's Pharmaceuticals and Medical Devices Agency (PMDA) operates one of the regulatory regimes most explicitly designed for the AI/Software-as-a-Medical-Device (SaMD) update problem. A 2019 amendment to the Pharmaceutical and Medical Device Act (enforced September 2020) introduced conditional early approval and IDATEN (Improvement Design within Approval for Timely Evaluation and Notice): a company can submit a proposed device-improvement plan at initial submission; once the plan is approved, subsequent changes within its scope receive streamlined review rather than full re-approval. The DASH for SaMD initiative supports faster reviews and earlier application. A 2025 scoping review on PMDA-approved AI radiology software documents transparency variability across approvals — a decisive hedge the case preserves. The teaching point is structural: the regulator designed change-control as a deliverable rather than defaulting to “approve once, then watch,” which is the governance failure pattern that Epic-Sepsis-class deployments surface. What stays open is both the sourcing and the reach: the 2025 transparency review is a preprint, the regulator's own 2023 Science Board report says IDATEN “has yet to be fully used,” and no source reports outcome durability across the approved devices.

## BACKGROUND

Medical-device regulation evolved for hardware whose design is largely fixed at approval. AI/SaMD breaks that assumption: a learning model can be updated, a training-data distribution can shift, and the device approved on Tuesday can be a meaningfully different device by Friday. The governance question for any regulator is whether the approval architecture

handles the update or defaults to approve-once-then-watch. Japan's PMDA chose the former and built the architecture explicitly.<sup>0</sup>

#### THE INTERVENTION

The 2014 Pharmaceutical and Medical Device Act laid the groundwork; the 2019 amendment introduced conditional early approval and the IDATEN system (Improvement Design within Approval for Timely Evaluation and Notice). Under IDATEN a manufacturer submits, alongside the initial approval package, a structured description of the modifications the manufacturer expects to make over the device's life — with the test and validation plan that will accompany each modification. Once the improvement plan is approved, subsequent changes within its scope receive streamlined regulatory review rather than full re-approval.<sup>1</sup>

#### HOW IT WORKED

The DASH for SaMD initiative complements IDATEN by supporting faster reviews and earlier application of SaMD products, anchoring the change-control architecture in a workflow that can keep pace with iteration. Together IDATEN and DASH are the regulator-side design for delegation-with-revocation: the deployer is delegated the modification authority within a pre-agreed scope, and the regulator retains the authority to review and revoke. The structure is the inverse of static-approval regimes that either block iteration or fail to track it.<sup>2</sup>

#### THE EVIDENCE

The 2025 medRxiv scoping review on PMDA-approved AI radiology software read the package inserts of the 20 AI radiology products among 151 approved SaMD: reported sensitivity ran from 67.7 to 100 percent, and no product reported complete case demographics for the studies behind it. The hedge is binding: IDATEN is a structural improvement in regulatory architecture, but the per-approval transparency the framework enables varies and has been documented as a research finding rather than assumed as a feature. The 2021 medRxiv-then-published systematic review on PMDA AI/ML medical devices through 2020 supplies the prior baseline — 11 approved ML/DL SaMD as of 2020, six of them radiology. The transparency review is itself a preprint and has not been through peer review; the earlier baseline review has. And the PMDA Science Board's own 2023 report on AI-based SaMD says of IDATEN that "the system has yet to be fully used," naming post-market performance deterioration and catastrophic forgetting among the reasons — the architecture exists and is largely unexercised.<sup>3</sup>

#### WHAT TRANSFERRED

The case pairs structurally with the Epic-Sepsis governance gap earlier in the corpus and with the FDA's evolving change-control SaMD policy. PMDA shows what designing change-control as a deliverable looks like at the regulator layer; Epic-Sepsis shows what happens when neither the vendor nor the deploying health system holds the change-control deliverable explicitly. The limits worth stating are both tier and reach: the regulatory framework is documented in program-report sources, the 2025 per-approval transparency analysis is a preprint rather than a peer-reviewed paper, the regulator's own report records that the mechanism is little used, and none of them follows the approved devices far enough to report outcome durability.

*The PMDA designed change-control as a deliverable rather than defaulting to approve once, then watch.*

EDITORS' SYNTHESIS OF THE PMD ACT AMENDMENT AND THE PMDA SAMD SCOPING REVIEWS.

REFERENCES

1. Kikuchi, et al. (2025), "Scoping Review of Regulatory Transparency in AI-based Radiology Software: Analysis of PMDA-approved SaMD Products," medRxiv 2025.10.02.25336333 — preprint, not peer reviewed; 151 approved SaMD as of 31 December 2024, 40 using AI, 20 in radiology.

2. Aisu, N., Miyake, M., Takeshita, K., et al. (2022), "Regulatory-approved deep learning/machine learning-based medical devices in Japan as of 2020: A systematic review," PLOS Digital Health 1(1):e0000001 (medRxiv preprint 2021).

3. Pharmaceuticals and Medical Devices Agency of Japan, PMD Act amendment (2019, enforced stepwise from September 2020) and DASH for SaMD — "DX Action Strategies in Healthcare for SaMD," announced November 2020 — program documentation.

4. PMDA Science Board, Subcommittee on Software as a Medical Device Utilizing AI and Machine Learning (28 August 2023), Report on AI-based Software as a Medical Device (SaMD) — the regulator's own statement that IDATEN "has yet to be fully used."

5. Cross-reference: "A decade of review in global regulation and research of artificial intelligence medical devices (2015–2025)," PMC12310608 — comparative regulatory context.

THE LEARNING ENGINEERING LENS

LE INSIGHT

PMDA's IDATEN and DASH for SaMD are the non-US regulatory architecture for AI/SaMD change control: the modification scope is pre-agreed at initial approval, in-scope changes get streamlined review, and the delegation-with-revocation structure is explicit. The 2025 scoping review — a preprint — documents transparency variability across approvals, and the PMDA Science Board records that IDATEN "has yet to be fully used." Decisive hedges preserved. What no source yet reports is outcome durability across the approved devices.

LENS APPROACH

PMDA is the non-US regulator-designed change-control case (induced 5.4; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the delegation-with-revocation structure and in Domain 3 (Human-System Collaboration) for the subobjective **Delegation with revocation** (3.3) — the regulator delegates in-scope modification authority and retains revocation, rather than defaulting to static approval. Pairs with the Epic-Sepsis governance gap as the structural contrast case.

| DOMAIN                    | MODES ADDRESSED |                                       | LENS COMPETENCY |               |               |               |               | PROBLEM TYPE  | INDUCED | LEO |       |
|---------------------------|-----------------|---------------------------------------|-----------------|---------------|---------------|---------------|---------------|---------------|---------|-----|-------|
| MEDICAL-DEVICE REGULATION | →               | <div>G</div> <div>N</div>             | →               | <div>D1</div> | <div>D2</div> | <div>D3</div> | <div>D4</div> | <div>D5</div> | 6       | 5.4 | 5 · 3 |
| AI/SAMD                   | JAPAN           | Navigating Sociotechnical Constraints |                 |               |               |               |               |               |         |     |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat post-approval change control as a deliverable at the initial-approval submission, not a downstream addendum: the modification scope, test plan, and revocation conditions should be on the approval package from the start.
- ▶ Specify the delegation boundary explicitly — what the manufacturer can modify without re-approval and what triggers a full review — so the delegation-with-revocation structure is operationally clear to both parties.

IN OPERATION / AFTER

- ▶ Audit transparency across approvals — what each package insert publishes about study design, case demographics, and measured performance — as a separate program-evaluation question; the 2025 scoping review found no product reporting complete case demographics.
- ▶ Use IDATEN as the structural contrast with Epic-Sepsis-class deployments: the regulator-designed change-control deliverable is the inverse of the static-approval-plus-watch pattern.

- ▶ Carry the reach of the evidence through any policy decision the scoping review supports: the regulatory framework is documented and the transparency analyses are peer-reviewed, but both are recent and neither follows outcomes.
- ▶ When importing the IDATEN pattern into another jurisdiction, name what the Japanese regulatory institution supplies (centralized agency capacity, established conditional-approval precedent) before asserting transfer.

### REFLECTION QUESTIONS

1. Identify an AI/SaMD-adjacent product in your context whose post-deployment modification is anticipated. What would a IDATEN-style submission look like — the pre-agreed modification scope, the per-modification test plan, the revocation conditions — and which party currently holds each element?
2. The case rests on program-report sources for the framework and, for the per-approval transparency findings, on a 2025 preprint. What is the minimum additional evidence you would require — outcome durability across approved devices, independent transparency audits — before treating IDATEN as a settled best-practice template?
3. Specify the delegation-with-revocation boundary you would write for a deploying organization adopting a SaMD product: which modifications can ship under the pre-agreed scope, which trigger re-review, and what evidence the deployer must publish at each modification event.

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards



### Chapter 3

# Education, Training & the Learning Workforce — What Fails

*When learning systems scale faster than their evidence.*

---

*The pilot results were real. The scale-up assumed they would travel on their own.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

# Tennessee Voluntary Pre-K Study

**G** Governance & Trust   **D** Designed Out

**IMPACT** Vanderbilt longitudinal RCT of a state-funded universal pre-K program: early gains faded by third grade; sixth-grade outcomes were worse than the control group on several measures

## IN BRIEF

Tennessee's Voluntary Pre-K program enrolled some 18,000 four-year-olds a year, and Vanderbilt researchers studied it with a rare large-scale randomized controlled trial — children randomly assigned to an offer of admission against children randomly assigned to the wait list. Through kindergarten the pre-K children showed the expected gains in literacy and vocabulary. By third grade the control children had caught up and surpassed them, and by sixth grade the pre-K group scored lower on state reading, math and science tests and had more disciplinary infractions and special-education placements. The unwelcome result was contested, its methods attacked, and the field largely declined to absorb it. The case is in the book not because pre-K is bad — other programs show durable effects — but because the measurement was unusually rigorous and the discipline's capacity to act on an inconvenient finding was tested and largely failed.

## BACKGROUND

The Tennessee Voluntary Pre-Kindergarten Program served roughly 18,000 four-year-olds a year. Demand exceeded supply, and that scarcity created an ethical lottery — a fair way to ration scarce seats that doubled as a clean randomizer. It let Vanderbilt's Peabody Research Institute run something rare in education: a randomized controlled trial, following children randomly assigned to an offer of admission against children randomly assigned to the wait list, the kind of design rarely available in a live policy setting.<sup>0</sup>

## WHAT HAPPENED

Through kindergarten the pre-K children showed the expected gains — stronger letter knowledge, vocabulary, early literacy — exactly the early result the program had been funded to produce. By third grade the control children had caught up and surpassed them. By sixth grade, the researchers reported, the pre-K children scored **lower** on state reading, math and science tests and had more disciplinary infractions and special-education placements — a reversal that turned an expected success story into an uncomfortable one the longitudinal design had been built to detect.<sup>1</sup>

## THE INVESTIGATION

The result contradicted policy consensus and provoked an unusual response: the study was attacked, its methods contested, and the field largely declined to internalize the findings, defending the policy rather than interrogating it. Other rigorous programs — Perry Preschool,




Abecedarian — remain durable, so the lesson is not that pre-K fails; it is that a discipline met an unwelcome, well-measured answer and mostly looked away, treating the inconvenient evidence as an attack to repel rather than a finding to absorb.<sup>2</sup>

#### THE CAPABILITY GAP

Tennessee Pre-K is the cleanest case in the dataset for what happens when a field has not engineered its own capacity to update on contrary evidence. The measurement instrument worked — a real RCT in a real policy setting, the rare study whose design could not easily be waved off. The institutional architecture for acting on what it found did not, so a strong instrument fed a discipline with no pathway to receive its answer. A measurement that returns an inconvenient result is only as valuable as the discipline's willingness to absorb it, and here the willingness was the part that was missing.<sup>3</sup>

#### AFTERMATH & REFORM

The debate continued for years through follow-up studies and counter-analyses, and the episode became a touchstone in the methodology of early-childhood research — argued over more than acted upon.<sup>4</sup> Its place in this book is as a governance-of-evidence case: the capability that needed engineering was not a better study but an implementation-science pathway that could route an unwelcome finding into program redesign rather than rejection, so the cost of the study bought a course correction instead of a controversy.



*By sixth grade, children who had attended TN-VPK were doing somewhat worse on academic achievement and discipline measures than children in the control group.*

DURKIN, LIPSEY, FARRAN & WIESEN (2022), VANDERBILT PEABODY RESEARCH INSTITUTE

#### REFERENCES

1. National Institute for Early Education Research, State of Preschool yearbooks (2010–2022) — program scale and context.
2. K. Durkin, M. Lipsey, D. Farran & S. Wiesen, “Effects of a Statewide Pre-Kindergarten Program... Through Sixth Grade,” *Developmental Psychology* 58(3): 470–484 (2022) — the sixth-grade reversal (quoted).
3. Responses and counter-analyses to the TN-VPK findings (2018–2022) — the contested reception.
4. M. Lipsey, D. Farran & K. Durkin, “Effects of the Tennessee Pre-Kindergarten Program... Through Third Grade,” *Early Childhood Research Quarterly* 45: 155–176 (2018) — the RCT and fade-out.
5. D. Phillips et al., *Puzzling It Out: The Current State of Scientific Knowledge on Pre-Kindergarten Effects* (2017); J. Heckman on durable early-childhood programs.

#### THE LEARNING ENGINEERING LENS

LE INSIGHT

Tennessee Pre-K is the cleanest case in the dataset for what happens when a discipline has not engineered its own capacity to update on contrary evidence. The measurement instrument worked. The institutional architecture for acting on what it found did not. Compare to Case 8 (Makary methodology debate) and Case 7 (VA Wait-Time): in each, a measurement that returned an unwelcome answer was contested rather than absorbed.

LENS APPROACH

LENS uses Tennessee Pre-K in LEN 4 as a measurement-architecture case (longitudinal RCT in a real policy context), in LEN 7 to discuss the institutional politics of unwelcome findings, and in LEN 10 as a studio prompt for designing the implementation-science pathway that would absorb such findings into program redesign rather than rejection.

| DOMAIN              | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------|---------------|------------------|--------------|---------|-----|
| EDUCATION AT SCALE  | → G D         | → D1 D2 D3 D4 D5 | 5            | 2.4     | 4   |
| Test and Evaluation |               |                  |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>▶ Commit in advance to a decision rule for how an unwelcome but well-measured result will change the program, so the response is designed before the finding arrives.</li><li>▶ Build the longitudinal measurement to detect fade-out and reversal, not just the early gains a program is funded to show.</li><li>▶ Establish an implementation-science pathway that can route a contrary finding into redesign, giving the evidence somewhere to go.</li></ul> | <ul style="list-style-type: none"><li>▶ Audit how the discipline actually receives inconvenient evidence, treating reflexive method-attacks as a symptom of a missing absorption pathway.</li><li>▶ Sustain follow-up measurement through later grades so a faded or reversed effect cannot be obscured by an early success.</li><li>▶ Protect the independence of the measurement function so a strong instrument is not dismantled for returning an answer the field did not want.</li></ul> |

REFLECTION QUESTIONS

1. What measurement instrument in your domain has returned an unwelcome answer, and how did the discipline respond?
2. Design the implementation-science pathway that would absorb a Tennessee-Pre-K-style finding into program redesign rather than rejection.
3. The Tennessee RCT was strong enough that its methods were attacked rather than its conclusion accepted. What decides, in your field, whether a rigorous but inconvenient result is absorbed or contested — and who holds the authority to act on it?

WHO BUILDS THIS governance, ethics & measurement · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · task analysis · design prototyping

# Algorithmic Bias in Educational Predictive Analytics

**G** Governance & Trust   **H** Human-System Interface   **D** Designed Out

**IMPACT** Predictive "at-risk" models misclassify by race in two directions — a community-college degree model would route support to fewer marginal Black students, and national models over-predict failure for Black and Hispanic students; 63 percent of institutions in a 2020 Chronicle survey use predictive analytics for retention

## IN BRIEF

Sixty-three percent of the 589 institutions in a 2020 Chronicle survey use predictive analytics to improve retention and graduation, most often by flagging "at-risk" students for early support. Research finds these models carry racial calibration bias, and it runs two ways: Bird and colleagues find a degree-completion model would send extra support to fewer marginal Black students, and Gándara and colleagues find models predict failure for Black and Hispanic students who in fact succeed. Either way the flag misdirects the support the prediction is meant to trigger. The magnitude depends heavily on how "at-risk" is defined, making this partly a construct-definition problem: the choice of what to predict is itself a capability decision with equity consequences. The models inherit historical patterns of discrimination from their training data, and a "flag" can confirm a biased instructor's low expectations rather than prompt help. It is the educational analog of the UK A-level case — an algorithm built to do good that can allocate help away from the students who need it most.

## BACKGROUND

Colleges increasingly use predictive analytics to identify students at risk of failing or dropping out, so advisors can intervene early — a genuinely well-meant aim, to reach struggling students before they vanish. Sixty-three percent of the 589 institutions surveyed by the Chronicle in 2020 said they use some form of these models, which makes any bias they carry a quiet, sector-wide feature of how support is allocated rather than an isolated experiment.<sup>o</sup>

## WHAT HAPPENED

The intervention is well-intentioned, but research finds the models carry racial calibration bias: they miscalibrate by race in two directions: a Virginia community-college degree model over-predicts success for Black students, so fewer marginal Black students are flagged for support, while national models over-predict failure for Black and Hispanic students who actually succeed. In both readings the flag lands on the wrong students. Crucially, the magnitude of the bias depends on how "at-risk" is defined, which makes it partly a

construct-definition problem rather than a coding bug — a flaw in what the model was asked to predict, not in how it computes the prediction.<sup>1</sup>

#### THE INVESTIGATION


Gándara and colleagues, working with nationally representative Education Longitudinal Study data, find the false-negative rate — predicting failure for a student who in fact earns the degree — is systematically higher for Black and Hispanic students, and that the leading bias-mitigation techniques are “generally ineffective at eliminating disparities in prediction outcomes and accuracy between racialized groups.” The bias is not a residual a fairness constraint cleans up after the fact; it is in what the model was built from, training data that encodes historical patterns of discrimination.<sup>2</sup> Baker and Hawn’s review catalogs which groups algorithmic bias in education is known to harm and which stages and agents in the development and deployment of educational algorithms are implicated — and the harm compounds when an instructor with deficit assumptions reads an “at-risk” flag as confirmation rather than a cue to help, letting the prediction become a self-fulfilling label.<sup>3</sup>

#### THE CAPABILITY GAP

The bias is not in the math; it is in the definition of “at-risk.” Define it one way and support flows to one population; define it another way and it flows to another, so the construct silently sets who the system decides to help. The choice of what the model predicts — the construct — is a capability-engineering decision with measurable equity consequences, and it is the part most often made implicitly, by whoever assembles the training labels, rather than governed deliberately by anyone accountable for where help is sent.<sup>4</sup>

#### AFTERMATH & REFORM

Unlike the discrete failures elsewhere in this chapter, this one is ongoing and quiet — embedded in advising dashboards at hundreds of institutions — which is what makes it dangerous: there is no single collapse to force a reckoning, only a steady misallocation no headline announces.<sup>5</sup> Its lesson, pushed upstream, is the chapter’s: governing an algorithm’s fairness begins not at deployment but at construct definition and label choice, with an equity audit of what the model is asked to predict and for whom the prediction allocates help — catching the bias where it is introduced rather than where it surfaces.



*Algorithmic bias in educational systems poses significant threats to educational equity, potentially amplifying existing social and economic disparities.*

EDITORS' PARAPHRASE OF BAKER & HAWN (2022)

REFERENCES

1. The Evolving State of Predictive Analytics, The Chronicle of Higher Education (2020), survey conducted by Maguire Associates, 589 institutions — 63 percent use predictive analytics for retention and graduation; only 20 percent institution-wide.

2. K. Bird, B. Castleman & Y. Song, "Are Algorithms Biased in Education? Exploring Racial Bias in Predicting Community College Student Success," Journal of Policy Analysis and Management 44(2) (2025), 379–402 — two Virginia Community College System models; bias over 5× higher when "at-risk" is the bottom decile than the bottom half for degree completion, and the reverse for course completion.

3. D. Gándara, H. Anahideh, M. Ison & L. Picchiarini, "Inside the Black Box: Detecting and Mitigating Algorithmic Bias across Racialized Groups in College Student-Success Prediction," AERA Open 10 (2024), doi:10.1177/23328584241258741 — higher false-negative rates for Black and Hispanic students in nationally representative ELS:2002 data; leading bias-mitigation techniques generally ineffective at eliminating the disparities.

4. R. Baker & A. Hawn, "Algorithmic Bias in Education," International Journal of Artificial Intelligence in Education 32 (2022), 1052–1092 — foundational review of algorithmic bias in education.

5. M. Ekowo & I. Palmer, Predictive Analytics in Higher Education: Five Guiding Practices for Ethical Use, New America (2017) — the deficit-versus-asset mindset and the handling of "at-risk," "low-risk," and "high-risk" labels from early-alert systems.

6. Cf. UK A-Level / Ofqual (Case 49); V. Eubanks, Automating Inequality (2018).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Educational predictive analytics is the ongoing live case for algorithmic bias at the construct level. The bias is not in the sigmoid; it is in the definition of "at-risk." Defining at-risk in one way allocates support to one population; defining it in another way allocates support to another. The choice of definition is a capability-engineering decision with measurable equity consequences.

LENS APPROACH

LENS treats Georgia State (Case 80) as the positive counterpart to this case. LEN 4 examines construct definition as the decisive measurement choice. LEN 7 examines the governance architecture that determines whose construct gets adopted. LEN 9 covers the technical bias-mitigation methods.

| DOMAIN              | FAILURE MODES      | LENS COMPETENCY             | PROBLEM TYPE | INDUCED | LEO |
|---------------------|--------------------|-----------------------------|--------------|---------|-----|
| EDUCATION AT SCALE  | → <div>G H D</div> | → <div>D1 D2 D3 D4 D5</div> | 5            | 8.1     | 4   |
| Test and Evaluation |                    |                             |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Govern the construct definition deliberately; have an accountable owner decide and document what "at-risk" means and whom the prediction will route help toward.
- ▶ Audit training data for encoded historical discrimination before fitting, so the model does not learn past inequities as future predictions.
- ▶ Test calibration across racial groups during development and treat disparate misclassification as a design defect, not a tolerable residual.

IN OPERATION / AFTER

- ▶ Monitor deployed dashboards for the quiet, distributed misallocation that has no single failure event to announce it.
- ▶ Train advisors and instructors to read an "at-risk" flag as a cue to help rather than confirmation of a deficit assumption.
- ▶ Sustain a recurring equity audit of what each model predicts and for whom, since a bias embedded across hundreds of institutions persists until someone is tasked to find it.

## REFLECTION QUESTIONS

1. Pick a predictive analytic in your institution. Reconstruct the construct definition behind it. What is the equity consequence of that definition?
2. Design the governance review that a new predictive model should pass before it allocates resources to or away from a population.
3. This failure is ongoing and quiet, with no single collapse to force a reckoning. What would it take to make a slow, distributed misallocation visible enough that someone with authority had to act on it?

**WHO BUILDS THIS** governance, ethics & measurement · human factors & interaction design · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs · task analysis · design prototyping

# Algorithmic Bias in Automated Exam Proctoring

**D** Designed Out **N** Normalization of Deviance **K** Knowledge & Institutional Memory

**IMPACT** The first quantitative study of facial-detection bias in automated exam proctoring software found that students with darker skin tones and Black students were significantly more likely to be flagged for instructor review for potential cheating; at the race-sex intersection, women with the darkest skin tones were 4.36 times more likely to be flagged than medium-skin-tone women and 5.6 times more likely than lighter-skinned women

## IN BRIEF

The COVID-era expansion of remote learning produced a rapid deployment of automated exam-proctoring software across higher education: computer-vision systems that monitor the student via webcam during an exam and flag suspicious behavior for instructor review. Yoder-Himes et al. (Frontiers in Education, 2022) ran the first quantitative study of facial-detection bias in this class of software. Students with darker skin tones and Black students were significantly more likely to be flagged for instructor review for potential cheating than students with lighter skin tones; at the race-sex intersection, women with the darkest skin tones were far more likely to be flagged than other groups. The study examined one major proctoring product and concludes the product “may employ biased AI algorithms that unfairly disadvantage students.” It documents the disparity but not a remediation, so it is a failure / diagnosis case. The harm comes directly from validation that did not stratify across skin tone, surfacing a group-specific failure in a deployed system. Trio with Cases 25 (eGFR), 26 (pulse oximetry), and 6 (Hoffman pain bias) at the race-construct-and-validation layer.

## BACKGROUND

Automated exam proctoring is a recent and rapidly deployed class of ed-tech: a webcam-based monitoring system that observes the student during a remote-administered exam and uses computer vision to flag suspicious behavior — looking off-camera, multiple faces in the frame, the face leaving the frame — for instructor review. The COVID-era pivot to remote instruction expanded the deployment of this software across U.S. higher education at speed. The systems’ face detection and face tracking are the computer-vision components: a failure in face detection translates directly into a false flag for instructor review, and systematic failures in face detection across demographic groups translate into systematic group-level harm.<sup>o</sup>

## WHAT HAPPENED

Yoder-Himes et al. (Frontiers in Education, 2022) ran the first published quantitative study of facial-detection bias in this class of software. The study examined the instructor outputs for about 357 students in four large STEM courses at a single university during Fall 2020, using the proctoring product the institution had adopted — one the paper notes is used by at least 1,500 universities nationally. Skin tone was classified against an expanded Fitzpatrick scale by two trained raters; race and sex were the students' own self-reports (sex assigned at birth; gender was not available from the institution). The headline finding ran in the direction the broader face-recognition-bias literature had documented (Buolamwini and Gebru, Raji, and others) and extended that finding into the assessment context. Students with darker skin tones, and Black students specifically, were significantly more likely to be flagged for instructor review for potential cheating than students with lighter skin tones.<sup>1</sup>

## THE INVESTIGATION

The intersectional analysis sharpened the finding. Women with the darkest skin tones were 4.36 times more likely to be flagged than medium-skin-tone women and 5.6 times more likely than lighter-skinned women — the race-by-sex intersection produced the largest disparity, consistent with the broader face-recognition literature's intersectional findings, and it appears only intersectionally: in aggregate the study found no significant male-female difference. What the study measures is the software's output, not the harm downstream of it. The authors are explicit that they collected no information on whether these students were reported or faced bias from instructors, and that at this institution the software was never the primary means by which reported cheating was identified. The exposure is real — a flag routes a student into an academic-integrity review that no institutional resolution mechanism is designed to handle when the suspicion was generated by a software bias — but the study documents the disparity in the flags, not the consequences that followed them.<sup>2</sup>

## THE CAPABILITY GAP

The study's claim is calibrated and direct: the product "may employ biased AI algorithms that unfairly disadvantage students." The study documents the disparity in a real deployment of a real product and does not document a remediation, so it is a failure / diagnosis case rather than a failure-to-intervention arc. The harm comes directly from validation that did not stratify across skin tone: the computer-vision face-detection model behind the proctoring system was deployed without demographic validation of its detection rates, and the group-specific failure was therefore present in the deployed system from day one and only surfaced post-hoc by external researchers.<sup>3</sup>

## AFTERMATH & REFORM

Drafted alongside the race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman pain bias), the proctoring case extends the validation-must-be- demographically-stratified lesson into the ed-tech and assessment context. The structural form is the same as pulse oximetry: a deployed system measured on an unrepresentative sample, producing aggregate accuracy that conceals a group-specific failure for years, until external researchers stratify the validation post-hoc. The eGFR cross-reference is the construct-definition counterpart; the Hoffman cross-reference is the human-judgment counterpart. All four sit in the small-and-big-tier conversation about validation discipline as an equity design commitment.



*Together, these results suggest that a major automated proctoring software may employ biased AI algorithms that unfairly disadvantage students.*

YODER-HIMES ET AL., FRONTIERS IN EDUCATION 2022.

## REFERENCES

1. Yoder-Himes, D. R., Asif, A., Kinney, K., Brandt, T. J., Cecil, R. E., Himes, P. R., Cashon, C., Hopp, R. M. P., & Ross, E. (2022). Racial, skin tone, and sex disparities in automated proctoring software. *Frontiers in Education*, 7:881449. doi:10.3389/educ.2022.881449 — the case's primary study.
2. Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. *Proceedings of Machine Learning Research*, 81:77–91 — the foundational intersectional-bias finding in face recognition.
3. Raji, I. D., & Buolamwini, J. (2019). Actionable auditing: Investigating the impact of publicly naming biased performance results of commercial AI products. AAAI/ACM Conference on AI, Ethics, and Society — the audit-and-disclosure mechanism the case calls for.
4. Sjoding, M. W., Dickson, R. P., Iwashyna, T. J., Gay, S. E., & Valley, T. S. (2020). Racial bias in pulse oximetry measurement. *New England Journal of Medicine*, 383(25):2477–2478 — the structural analog in the medical-device context (Case 26).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Automated exam proctoring is the assessment-domain counterpart to pulse oximetry's medical-device bias. Yoder-Himes et al. 2022 is the first quantitative study of the disparity in this class of software, with the intersectional finding the broader face-recognition literature predicts. The case documents the disparity in one product; the remediation is not yet documented.

### LENS APPROACH

Proctoring bias is the small-tier ed-tech validation-stratification failure (induced 8.2; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the validation-must-be-demographically-stratified discipline, and in Domain 5 (Navigating Sociotechnical Constraints) for the institutional resolution mechanism that has to handle the case where the flag was generated by software bias. Trio with Cases 25 (eGFR), 26 (pulse oximetry), and 6 (Hoffman pain bias) at the race-construct-and-validation layer.

| DOMAIN        | FAILURE MODES | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|---------------|---------------|---------------------|--------------|---------|-------|
| ED-TECH       | D N K         | D1 D2 D3 D4 D5      | 5            | 8.2     | 4 • 5 |
| ALGO FAIRNESS | →             | →                   |              |         |       |
| ASSESSMENT    |               | Test and Evaluation |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Stratify validation across skin tone (and across the race-sex intersection) before deployment, not after, so the group-specific failure surfaces in the engineering record before it surfaces in the harm record.
- ▶ Specify the harm class the false-positive flag produces — academic-integrity defense, downstream consequence to the student — and design the institutional resolution mechanism to handle the case where the flag was generated by software bias.
- ▶ Require the proctoring vendor to publish demographic stratification of their detection rates, on a standard skin-tone scale, as a deployment condition rather than as a post-hoc disclosure.

### IN OPERATION / AFTER

- ▶ Treat the Yoder-Himes finding as the diagnosis it is: the disparity is documented in one product; the remediation is not.
- ▶ Carry the intersectional reading explicitly (women with the darkest skin tones far more likely to be flagged) in any communication about the case; the broader race-by-sex intersectional finding is consistent with the face-recognition-bias literature and should not be smoothed into a single-dimension finding.
- ▶ Build the cross-case reading with eGFR (Case 25), pulse oximetry (Case 26), and Hoffman pain bias (Case 6): the validation-must-be-demographically-stratified lesson runs across clinical, device, and assessment domains, and the proctoring case is the assessment-domain anchor.

## REFLECTION QUESTIONS

1. Identify a deployed system in your domain whose validation rests on an aggregate accuracy figure rather than a demographically stratified one. What would the stratified validation actually require, and who would have to commission it before deployment rather than after?
2. The Yoder–Himes finding is intersectional: women with the darkest skin tones were 4.36 times more likely to be flagged than medium–skin–tone women. What is the analog intersectional structure of the harm in your domain, and is it visible in the engineering record before the harm record?
3. The case documents the disparity in one product; the remediation is not yet documented. What would a remediation look like — vendor disclosure, regulatory disclosure requirement, institutional resolution–mechanism redesign — and which of those is in your scope?

**WHO BUILDS THIS** systems & human-factors engineering · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · incident analysis · safety dashboards · knowledge bases · data pipelines

## CASE 48

K-12 EDUCATION

SCHOOL SAFETY INFRASTRUCTURE

RACIAL DISPARITIES

2010S – 2022

# School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism

**G** Governance & Trust   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**DISCLOSURE** Institutional overlap: an editor of this volume is appointed in the same school as the study's lead author — the Johns Hopkins University School of Education, where the author holds a Bloomberg Distinguished Professorship jointly with the Bloomberg School of Public Health. No editor was personally involved in the work. Included on the published peer-reviewed evidence (Journal of Criminal Justice, 2022); editorial framing keeps critical distance.

**IMPACT** Johnson and Jabbari (Journal of Criminal Justice, 2022) find that the most heavily surveilled schools suspend more students even with school disorder and student misbehavior controlled, and that surveillance lowers math scores and college enrollment for suspended and non-suspended students alike; the "safety tax" is greatest for Black students, who are far more likely to attend high-surveillance schools — one of the motivating cases for the LEO Gap attribution

## IN BRIEF

Johnson and colleagues, writing in the Journal of Criminal Justice in 2022, analyzed the relationship between school-surveillance infrastructure — security cameras, metal detectors and random metal-detector checks, dog sniffs and contraband sweeps, drug testing, uniform and dress-code requirements, clear book bags, student and faculty ID badges, closed campuses — and outcomes for Black students across U.S. schools. The study's learning-engineering content is in where it locates the mechanism. The analysis holds school social disorder and student misbehavior constant and still finds that the most heavily surveilled schools suspend more students, and that surveillance depresses math scores and college enrollment for suspended and non-suspended students alike; because Black students are far more likely to attend those schools, they carry the largest share of the cost the authors call a safety tax. The mechanism, in other words, is the infrastructure, not the students. The case is one of the motivating cases in the v2 sweep for the subobjective **Gap attribution** (4.2) — the discipline of asking, when a disparity in outcomes is observed, what share of the disparity is attributable to the institutional or technical infrastructure rather than to the population the infrastructure is operating on. The standing COI rendered under the title is binding: an editor of this volume shares an institution with the study's authors, and the case is anchored to the peer-reviewed Journal of Criminal Justice evidence rather than to institutional press. The case extends the race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman) into the K-12 education domain at the infrastructure layer.

## BACKGROUND

U.S. public schools have, over the past two decades, deployed a progressively more elaborate surveillance infrastructure. The measures Johnson and Jabbari index are the ordinary furniture of a fortified school: security cameras; metal detectors and random metal-detector checks; random dog sniffs and contraband sweeps; drug testing; uniform requirements and strict dress codes; clear book bags; student and faculty identification badges; closed campuses at lunch. The distribution of that infrastructure across schools has not been uniform. Schools serving predominantly Black student populations have, on average, carried more of it. The institutional rationale has typically been safety — the surveillance infrastructure is described as keeping students safe — and the question the Johnson study confronts is whether the surveillance is operating as safety infrastructure or as something else, measured by the outcomes the infrastructure produces.<sup>0</sup>

## WHAT HAPPENED

Johnson and colleagues (Journal of Criminal Justice, 2022) analyzed the relationship between school-surveillance infrastructure and outcomes for Black students across the U.S. public-school sector. The study's design uses counterfactual and multi-level modeling on national longitudinal data, comparing students in the highest-surveillance schools with comparable students in the lowest. The headline analytic result the case rests on: with school social disorder and student misbehavior controlled, students in heavily surveilled schools are suspended more, score lower in mathematics, and enroll in college at lower rates — the non-suspended along with the suspended. The infrastructure is acting as a driver of the outcome gap, not only as a response to it.<sup>1</sup>

## THE INVESTIGATION

The mechanism is what makes the case the canonical gap-attribution case. When a disparity in outcomes is observed, the analytical default in many institutional settings is to attribute the disparity to the population the outcomes are measured on — to differences in preparation, behavior, family context. The Johnson study shows that, in the school-surveillance case, that default is wrong in a measurable sense: the institutional infrastructure is itself a driver of the disparity, and attributing the outcome gap to the students rather than to the infrastructure mis-locates the mechanism in a way the data does not support. The case extends the race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman) into K-12 education at the infrastructure layer: in eGFR and pulse oximetry the mechanism was in the device or the formula; in Hoffman the mechanism was in the practitioner's cognitive baseline; here the mechanism is in the institutional architecture itself.<sup>2</sup>

## THE CAPABILITY GAP

The standing COI rendered under the title is binding. An editor of this volume is appointed in the same school as the study's lead author — the Johns Hopkins University School of Education, where the author holds a Bloomberg Distinguished Professorship jointly with the Bloomberg School of Public Health. The overlap is at school level, not merely the university. The case is anchored to the published peer-reviewed evidence in the Journal of Criminal Justice, not to institutional press or to commentary by the editor or the authors outside the paper. The editorial framing has been written to maintain critical distance from the home-institution affiliation; the case's claim is what the published evidence supports, and

the disclosure is the safeguard that makes the institutional adjacency visible at the point of reading.<sup>3</sup>

#### AFTERMATH & REFORM

The case is one of the motivating cases in the v2 sweep for the subobjective **Gap attribution** (4.2) — the discipline of asking, when a disparity in outcomes is observed, what share of the disparity is attributable to the institutional or technical infrastructure rather than to the population the infrastructure is operating on. The race-construct trio established the pattern at the device/formula/cognitive- baseline layers; this case carries the pattern at the institutional-infrastructure layer. The four-case set (Cases 25, 26, 6, 48) is the case-grounded basis for Gap attribution as a designed competency: practitioners and program designers have to be trained to look for the mechanism in the infrastructure they built, not only in the population they are serving, and the evidence the case carries is what makes the proposed LEO defensible.

*The “safety tax” students pay with their average levels of test performance and college going rates is greatest for Black students of both genders given their overrepresentation in high-surveillance schools.*

JOHNSON AND JABBAR, JOURNAL OF CRIMINAL JUSTICE, 2022.

#### REFERENCES

1. O. Johnson & J. Jabbari (2022), “Infrastructure of social control: A multi-level counterfactual analysis of surveillance and Black education,” *Journal of Criminal Justice* 83:101983 — the peer-reviewed source for the case.
2. Hoffman, Trawalter, Axt, & Oliver (2016), “Racial bias in pain assessment and treatment recommendations, and false beliefs about biological differences between blacks and whites,” *PNAS* 113(16):4296–4301 — race-construct trio at the cognitive-baseline layer (Case 6).
3. Inker, Eneanya, Coresh, et al. (2021), “New Creatinine- and Cystatin C–Based Equations to Estimate GFR without Race,” *NEJM* — race-construct trio at the formula layer (Case 25).
4. Sjoding, Dickson, Iwashyna, Gay, & Valley (2020), “Racial Bias in Pulse Oximetry Measurement,” *NEJM* 383:2477–2478 — race-construct trio at the device layer (Case 26).

#### THE LEARNING ENGINEERING LENS

##### LE INSIGHT

The Johnson school-surveillance study locates the mechanism of an outcome disparity in the institutional infrastructure, not in the population the infrastructure operates on. The case extends the race-construct trio (eGFR, pulse oximetry, Hoffman) into K-12 education at the infrastructure layer and is one of the motivating cases for the LEO Gap attribution. COI under the title — shared institution — is binding and rendered openly.

##### LENS APPROACH

Johnson school surveillance is the infrastructure-as-mechanism gap-attribution case (induced 8.1; LENS D4/PT5) — Domain 4 as the case-grounded basis for **Gap attribution**; Domain 5 for the institutional-architecture-as-mechanism framing. Pair with Cases 25, 26, 6. COI binds.

| DOMAIN                       | FAILURE MODES | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO |
|------------------------------|---------------|---------------------|--------------|---------|-----|
| K-12 EDUCATION               | G K N         | D1 D2 D3 D4 D5      | 5            | 8.1     | 4   |
| SCHOOL SAFETY INFRASTRUCTURE | →             | →                   |              |         |     |
|                              |               | Test and Evaluation |              |         |     |
| RACIAL DISPARITIES           |               |                     |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Design the analysis to separate population-level variables from infrastructure-level variables so the share of the disparity attributable to each can be estimated; the Johnson study’s design is the worked example.</li><li>Treat the institutional infrastructure as a candidate mechanism, not as a neutral background; the default attribution-to-population analysis cannot test the infrastructure if the infrastructure is not a variable in the model.</li><li>Anchor the editorial framing to the peer-reviewed evidence rather than to institutional press; the COI is rendered openly under the title and the framing maintains critical distance.</li></ul> | <ul style="list-style-type: none"><li>Carry the Gap-attribution discipline into curriculum design: practitioners have to be trained to look for the mechanism in the infrastructure they built, not only in the population they are serving.</li><li>Pair the case with the race-construct trio (Cases 25, 26, 6) so the gap-attribution pattern is taught at multiple layers — device, formula, cognitive baseline, institutional architecture.</li><li>Preserve the COI render — shared institution, anchored to peer-reviewed evidence, critical editorial distance — as the standing language for home-institution-shared cases across the corpus.</li></ul> |

REFLECTION QUESTIONS

1. Identify a disparity in outcomes in your domain that is currently attributed primarily to the population the outcomes are measured on. What infrastructure-level variables would you have to add to the analysis to test whether the institutional architecture is itself a mechanism of the disparity?
2. Specify the design pattern the Johnson study uses: separating student-level from school-level variables so the share of the disparity attributable to each can be estimated. What is the analog in your context, and where is the default analysis most at risk of mis-locating the mechanism?
3. The case is one of the motivating cases for the LEO Gap attribution. What instance from your work — a device, a formula, a cognitive baseline, an institutional architecture — would you carry as the case-grounded basis for training practitioners in your context to look for the mechanism in the infrastructure they built?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

# UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Ofqual standardisation algorithm applied to summer 2020 A-level grades following examination cancellation downgraded approximately 39.1% of teacher-estimated grades; results released August 13 2020; algorithm withdrawn August 17 2020 after four days of public protest; Centre Assessment Grades (teacher estimates) substituted; the Office for Statistics Regulation's 2021 review found the model failed to command public confidence and that acceptability testing had examined the calculation process rather than the impact on individuals

## IN BRIEF

With summer 2020 examinations cancelled in response to the COVID-19 pandemic, the UK Office of Qualifications and Examinations Regulation (Ofqual) deployed a statistical standardisation model to produce A-level grades from Centre Assessment Grades (teacher estimates) and Centre-level historical performance. Results were released on August 13, 2020. Approximately 39.1 percent of teacher-estimated grades

were downgraded by the algorithm. State-school students in larger cohorts were downgraded at higher rates than independent-school students in smaller cohorts, because the model relied on Centre-level historical performance more heavily where Centre-level cohorts were larger. After four days of public protest, on August 17, 2020, the algorithm was withdrawn and Centre Assessment Grades were substituted. The Ofqual technical report states that its statistical measures are meaningful only for groups of students and would over-interpret the evidence if used to predict an individual result, and that greater reliance on teacher grades at small centres carries an absolute disadvantage for students at larger ones. The case pairs with Case 86 (Gándara / AERA Open community-college fairness), Case 88 (LiveHint AI bias across foundation models), and Case 48 (Johnson school surveillance).

## BACKGROUND

The summer 2020 examination cancellation removed the mechanism by which A-level grades had historically been produced. The Department for Education and Ofqual judged that teacher estimates alone would generate grade inflation incompatible with university

admissions and higher-education capacity planning. The decision was to combine teacher Centre Assessment Grades with a statistical standardisation model that drew on Centre-level historical performance to adjust the distribution of grades each Centre's cohort received. The intent of the standardisation was to preserve year-on-year comparability of the national grade distribution; the seam that surfaced under deployment was that population-level comparability and individual-student fairness are not reconcilable when the standardisation mechanism depends on cohort size.<sup>0</sup>

#### WHAT HAPPENED

The model's mechanics carried the seam. Where a Centre's cohort for a subject was small, the model relied more heavily on the teacher-submitted Centre Assessment Grade and rank order, because small-cohort historical performance was not informative enough to standardise against. Where a Centre's cohort was large, the model relied more heavily on the Centre's historical performance, because the larger cohort gave the standardisation more purchase. The distributional consequence was structural: independent schools and selective settings with small cohorts received grades close to teacher estimates; state schools and large-cohort comprehensive settings received grades pulled downward toward the Centre's historical distribution. The headline result was that approximately 39.1 percent of teacher-estimated grades were downgraded and that the downgrade rate was higher for state-school students in large cohorts.<sup>1</sup>

#### THE INVESTIGATION

The four-day withdrawal arc is the governance record. Grades were released on Thursday, August 13, 2020. Public protest — students gathering with signs reading “the algorithm stole my future,” extensive press coverage of named individual cases, and rapid political pressure — built across the weekend. On Monday, August 17, 2020, Ofqual and the Department for Education announced that A-level and GCSE grades would be reissued at the teacher-submitted Centre Assessment Grade level. The withdrawal was structural — it affected the entire 2020 cohort — and it was rapid in a way that few national-scale algorithmic deployments have been. The Office for Statistics Regulation's March 2021 review adjudicated the governance record and named the engagement failure as the decisive one: acceptability testing had examined the process of calculating grades rather than the impact on individuals. The technical report had been internally honest about the cohort-size dependence, and the failure was that the dependence had not been surfaced to affected schools and students in advance of deployment.<sup>2</sup>

#### THE CAPABILITY GAP

The case pairs with Case 86 (Gándara / community-college predictive equity in AERA Open) at the higher-education scale: both cases turn on the question of whether a standardisation or prediction mechanism that is statistically defensible at the population level can be deployed in a way that is defensible at the individual-student level. Pair with Case 88 (LiveHint AI bias across foundation models) for the bias-surfacing thread in education-deployed algorithms. Pair with Case 48 (Johnson school surveillance) for the algorithmic-administration-in-education parallel at a different scale. The Ofqual case is unusual in the casebook because it is the rare deployment that was withdrawn within days under public pressure; most cases in the corpus document deployments that ran for years before withdrawal or that were never withdrawn.<sup>3</sup>



## AFTERMATH & REFORM

The technical report's hedge is binding. Ofqual's own document states that the statistical measures behind the standardisation are meaningful only for groups of students, that using them to predict an individual result would over-interpret them, and that the leniency granted to small centres imposes an absolute disadvantage on students entering through larger ones. The case teaches the change-control-and-disclosure-as- governance-artifacts pattern: an algorithm that is internally documented as carrying a distributional seam cannot be deployed at population scale without the distributional seam being surfaced to the affected population in advance of deployment. The four-day withdrawal is the governance evidence that the population had not been consulted; the structural argument the case anchors is that the consultation is the governance artifact whose absence the withdrawal made visible.

*Approximately 39.1 percent of teacher-estimated grades were downgraded; the algorithm was withdrawn within four days under public protest; the technical report was internally honest about the cohort-size dependence and the failure was that the honesty did not travel out of the document.*

EDITORS' SYNTHESIS OF THE OFQUAL TECHNICAL REPORT AND THE OFFICE FOR STATISTICS REGULATION REVIEW (2021).

## REFERENCES

1. Ofqual, Awarding GCSE, AS, A level, advanced extension awards and extended project qualifications in summer 2020: interim report, August 2020.
2. Royal Statistical Society, Submission to the Office for Statistics Regulation on the summer 2020 grading process, 2020 — independent statistical review of the standardisation methodology.
3. UK House of Commons Education Committee, Getting the grades they've earned: Covid-19: the cancellation of exams and 'calculated' grades, First Report of Session 2019–21, HC 617, 11 July 2020 — the pre-deployment warning; Office for Statistics Regulation, Ensuring statistical models command public confidence, 2 March 2021 — the post-withdrawal review.
4. Smith, H. (2020), "Algorithmic bias: should students pay the price?" *AI & Society* 35(4):1077–1078 — early academic commentary on the equity dimensions of the withdrawal.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The Ofqual A-level case is the rare national-scale algorithmic deployment withdrawn within days under public pressure. The technical report was internally honest about the cohort-size dependence and the incompatibility of population-level standardisation with individual-level fairness; the four-day withdrawal is the evidence that the internal honesty did not travel out of the document to the affected population in advance of deployment.

### LENS APPROACH

Ofqual A-level 2020 is the change-control-and-disclosure-as- governance-artifacts case at national scale (induced 5.4; LENS D4+D5/PT5; LEO-4 and LEO-5). LENS uses it in Domain 4 (Test and Evaluation) for the consultation-with-affected- stakeholders process as the test surface and in Domain 5 (Navigating Sociotechnical Constraints) for the cohort-size dependence as the distributional seam the deployment carried. Pair with Case 86 (Gándara community-college predictive equity), Case 88 (LiveHint AI bias across foundation models), and Case 48 (Johnson school surveillance). The technical report's acknowledgement of incompatibility is the hedge that matters.

| DOMAIN                 | FAILURE MODES | LENS COMPETENCY          | PROBLEM TYPE | INDUCED | LEO   |
|------------------------|---------------|--------------------------|--------------|---------|-------|
| GOVERNMENT             | D K N         | D1 D2 D3 D4 D5           | 5            | 5.4     | 4 · 5 |
| EDUCATION AT SCALE     | →             | →                        |              |         |       |
| HIGH-STAKES ASSESSMENT |               | Test and Evaluation + D5 |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Surface the distributional seam an internally documented standardisation mechanism carries to the affected population in advance of deployment; the Ofqual technical report was internally honest about the cohort-size dependence, and the governance failure was that the honesty did not travel out of the document.
- ▶ Treat the consultation-with-affected-stakeholders process as the change-control artifact a national-scale algorithmic deployment requires; the four-day withdrawal under public protest is the evidence that the consultation had not occurred.
- ▶ Pre-specify the individual-student fairness criterion against which a standardisation mechanism will be evaluated, and refuse deployment when the criterion is incompatible with the standardisation goal.

### IN OPERATION / AFTER

- ▶ Carry the technical report's hedge — the statistical measures "are only sufficiently meaningful to be informative for groups of students," and small-centre leniency "would lead to an absolute disadvantage to those students entering through larger centres" — into print without softening; the case's pedagogical value depends on the internal documentation of the seam being visible alongside the public withdrawal.
- ▶ Pair in syllabi with Case 86 (Gándara) so the population-level-versus-individual-level fairness tension is taught at both the secondary-to-higher-education transition scale and the community-college transition scale.
- ▶ Use the case as the rare example of an algorithmic deployment withdrawn at national scale within days; the four-day withdrawal arc is the curricular target for governance-response speed under public pressure.

## REFLECTION QUESTIONS

1. Identify a standardisation or prediction mechanism in your domain whose internal documentation flags a distributional seam. What is the consultation process that would surface the seam to the affected population in advance of deployment, and what would deployment without consultation look like?
2. Specify the individual-level fairness criterion against which a population-level standardisation in your setting would be evaluated. Is the criterion compatible with the standardisation goal, and if not, which governs?
3. The Ofqual case is unusual for the speed of withdrawal — four days. Pick a deployment in your domain whose distributional seam has not yet surfaced publicly, and ask: what would have to be true for a four-day withdrawal arc to be possible, and what would have to be true for the seam to have been surfaced in advance instead?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

# Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Wisconsin Department of Public Instruction's Dropout Early Warning System (DEWS) ran from 2012 until DPI withdrew the dashboards on 12 October 2023, scoring every sixth- through ninth-grader in the state — in 2020–21 more than 32,000 of them, 15 percent of the cohort, carried a "high risk" label; Perdomo, Britton, Hardt, & Abebe FAccT 2025 regression-discontinuity analysis on 2013–2021 data put the effect of a higher risk label on on-time graduation at about 5 percentage points, on a 95% interval of -2% to 11.6% that cannot rule out zero; DPI's own 2021 internal equity analysis found false-alarm rates 42 points higher for Black and 18 points higher for Hispanic students, and DPI told no one and changed nothing

## IN BRIEF

The Wisconsin Department of Public Instruction has operated the Dropout Early Warning System (DEWS) from 2012 until it withdrew the dashboards on 12 October 2023, producing dropout-risk scores twice a year for every sixth- through ninth-grader in the state; in 2020–21 more than 32,000 of them — 15 percent of the cohort — carried a "high risk" label. Two evidence streams converge on the deployment. Perdomo, Britton, Hardt, and Abebe's 2025 FAccT paper used a regression-discontinuity design on 2013–2021 DEWS-and-graduation data and estimated that moving a student from the moderate to the high-risk category raises on-time graduation by about 5 percentage points, on a 95 percent confidence interval of -2 to 11.6 percent that cannot rule out zero effect. The Wisconsin DPI's own 2021 internal equity analysis — the slide The Markup obtained is headed "Is DEWS Fair?" — found false-alarm rates 42 percentage points higher for Black students and 18 points higher for Hispanic students than for White students, and the agency neither told districts nor changed the model. The Markup's 2023 investigation by Todd Feathers documented the disparate-impact finding and the agency's response. Both findings are binding. The case pairs with Case 48 (Johnson school surveillance), Case 86 (Gándara community-college predictive equity), and Case 52 (Purdue Course Signals — reverse causality). This case carries both the peer-reviewed and journalism-investigation evidence streams in one entry rather than parallel ones.

## BACKGROUND

The Wisconsin Dropout Early Warning System was built in-house at the Wisconsin Department of Public Instruction, under senior research analyst Jared Knowles, and deployed in 2012. The system ingests student-level administrative data — attendance, course perfor-

mance, behavioral incidents, mobility, and demographics — and produces a dropout-risk score for each student in grades six through nine. The score is surfaced to school administrators, counselors, and teachers, with the operational theory that the early warning enables school-level interventions — additional counseling, tutoring, mentoring, or family contact — that improve the student’s graduation trajectory. The system operated at scale for more than a decade, scoring every sixth- through ninth-grader in the state twice a year; in 2020–21 more than 32,000 of them, 15 percent of the cohort, carried a “high risk” label.<sup>0</sup>

#### WHAT HAPPENED

Perdomo, Britton, Hardt, and Abebe’s 2025 paper at FAcCT (also available as arXiv 2304.06205) is the decisive causal-inference evidence on the deployment. The researchers used a regression-discontinuity design, exploiting the DEWS risk-threshold cutoff to compare students just above the threshold (classified as high-risk and surfaced to the school’s intervention surface) against students just below it (not so surfaced), controlling for the continuous risk score itself. The analysis covered the DEWS scores from 2013 to 2021 — nearly a decade — and the corresponding graduation outcomes. The headline finding is a point estimate of about 5 percentage points on on-time graduation for a student moved from moderate to high risk, on a 95 percent confidence interval of –2 to 11.6 percent. The interval contains zero, so the analysis cannot rule out that the risk category changes nothing; the estimate is directionally correct and not rigorously conclusive. The hedge the paper preserves is binding: the RDD analysis cannot establish a negative finding either (the interval includes small positive effects), but it can and does establish that nearly a decade of statewide data leaves the prediction-triggers-intervention-that-changes-outcome theory unsettled. The paper’s own conclusion goes further than the null: the authors show that a targeting rule using only a student’s school and district — no individual risk score at all — would direct interventions about as efficiently as DEWS does, and argue this holds “even if individual predictions are highly accurate and effective interventions exist,” because dropout is structurally determined. The question the paper leaves on the table is not whether the model is accurate — it sorts students by risk accurately — but whether the individual prediction was ever the thing carrying the value.<sup>1</sup>

#### THE INVESTIGATION

The Wisconsin DPI’s own 2021 internal equity audit, titled “Is DEWS Fair?”, is the decisive disparate-impact evidence on the deployment. The analysis found that DEWS raised false alarms about Black and Hispanic students at far higher rates than about their White classmates — 42 percentage points higher for Black students, 18 for Hispanic students. The agency’s own response to its own analysis was to tell no one and change nothing. The Markup’s 2023 investigation by Todd Feathers documented both the finding and the agency’s response, and the journalism-investigation evidence stream entered the public record in 2023. The case carries both evidence streams — the peer-reviewed causal-inference null and the agency-and-investigation disparate-impact finding — in one entry, rather than as parallel cases, because the deployment is one deployment and the two streams describe complementary structural problems with it.<sup>2</sup>

#### THE CAPABILITY GAP

The case pairs with Case 48 (Johnson school surveillance) for the algorithmic-public-administration-in-education parallel; both cases involve administrative-data predictions

deployed against student populations and both surface disparate-impact concerns at the deployment surface. Pair with Case 86 (Gándara community-college predictive equity) for the predictive-equity-in-education thread at adjacent population scale. Pair with Case 52 (Purdue Course Signals) for the reverse-causality and null-effect thread in education predictive analytics; the Purdue case named the same conceptual problem that the Perdomo et al. analysis evidences at population scale. DEWS is the rare deployment in the corpus that is evidenced from both the peer-reviewed causal-inference direction and the disparate-impact-investigation direction, and the editorial decision to carry both streams in one entry reflects the structural unity of the deployment.<sup>3</sup>

#### AFTERMATH & REFORM

The hedges the case carries are binding and both streams are preserved. The Perdomo et al. RDD analysis cannot rule out zero treatment effect on graduation across 2013–2021 data — eight years of statewide records still cannot settle whether the prediction triggers the intervention the deployment theory requires. The DPI's own equity analysis found DEWS raised false alarms about Black and Hispanic students at far higher rates than about White students, and the agency continued operating the system unchanged for two years afterward. Both findings are binding and travel together. The LEO on designing predictions to trigger support rather than gatekeeping is anchored by the case at the deployment-with-null-causal-effect-and-disparate-accuracy seam: a prediction system operating at population scale for more than a decade, without evidence that the prediction triggers outcome-changing intervention and with evidence that the prediction's accuracy varies across protected-attribute subgroups, is the deployment form whose persistence the case asks the reader to account for.

*The RDD estimate is about 5 points on graduation, on an interval that cannot rule out zero; the agency's own equity analysis found DEWS raised false alarms about Black students 42 points more often than about White students; the agency told no one.*

EDITORS' SYNTHESIS OF PERDOMO ET AL. (2025, FACCT), THE WISCONSIN DPI INTERNAL EQUITY AUDIT (2021), AND THE MARKUP INVESTIGATION (FEATHERS, 2023).

#### REFERENCES

1. Perdomo, J. C., Britton, T., Hardt, M., & Abebe, R. (2025), "Difficult Lessons on Social Prediction from Wisconsin Public Schools," Proceedings of FACCT 2025, doi:10.1145/3715275.3732175 (also arXiv:2304.06205).
2. Wisconsin Department of Public Instruction (2021), internal equity analysis of the Dropout Early Warning System — unpublished; the presentation summarizing it, whose slide is headed "Is DEWS Fair?", was obtained by The Markup under public records.
3. Feathers, T. (2023), "False Alarm: How Wisconsin Uses Race and Income to Label Students 'High Risk,'" The Markup, April 27, 2023 — investigation documenting the disparate-impact finding and the agency response.
4. Wisconsin Department of Public Instruction, "WISEdash for Districts — Dropout Early Warning System (DEWS) Dashboards" — notice that early-warning dashboard data (DEWS/CCREWS) is no longer available as of 12 October 2023 and that DPI is evaluating the future of the systems.
5. Knowles, J. E. (2015), "Of Needles and Haystacks: Building an Accurate Statewide Dropout Early Warning System in Wisconsin," Journal of Educational Data Mining 7(3):18–67 — the original DEWS technical-methodology paper.

#### THE LEARNING ENGINEERING LENS

#### LE INSIGHT

Wisconsin DEWS is the reference case for a prediction system operating at population scale for more than a decade without conclusive evidence that the prediction triggers outcome-changing intervention and with evidence that the prediction's error rates vary across protected-attribute subgroups. Both the peer-reviewed null and the agency-audit disparate-impact finding travel together; the case carries both evidence streams in one entry rather than parallel ones.

### LENS APPROACH

Wisconsin DEWS is the designing-predictions-to-trigger-support-not-gatekeeping case at population scale (induced 8.3; LENS D4+D5/PT6; LEO-4 and LEO-5). LENS uses it in Domain 4 (Test and Evaluation) for the causal-inference-on-multi-year-deployment discipline and in Domain 5 (Navigating Sociotechnical Constraints) for the equity-audit-as-binding-input anchor. Pair with Case 48 (Johnson school surveillance), Case 86 (Gándara community-college predictive equity), and Case 52 (Purdue Course Signals reverse causality). Both the peer-reviewed null and the journalism-and-agency-audit disparate-impact finding are decisive.

| DOMAIN               | FAILURE MODES | LENS COMPETENCY          | PROBLEM TYPE | INDUCED | LEO   |
|----------------------|---------------|--------------------------|--------------|---------|-------|
| EDUCATION AT SCALE   | D K N         | D1 D2 D3 D4 D5           | 6            | 8.3     | 4 · 5 |
| PREDICTIVE ANALYTICS |               | Test and Evaluation + D5 |              |         |       |
| ALGORITHMIC EQUITY   |               |                          |              |         |       |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- Pre-specify the causal-inference design that will evaluate whether the prediction triggers outcome-changing intervention; the Perdomo et al. RDD is the deployment-evidence form that the prediction-triggers-support theory requires for verification at multi-year horizons.
- Treat the agency's own equity audit as a binding governance input, not as an internal document; the DPI's "Is DEWS Fair?" audit was internally honest about the disparate-accuracy finding, and the case's central observation is the agency's continued operation of the system unchanged.
- Design predictions to trigger support, not gatekeeping; the case's central LEO anchor is the discipline of building the intervention surface that the prediction is meant to trigger, and of verifying — not assuming — that the prediction's operational role produces the outcome change the deployment theory requires.

#### IN OPERATION / AFTER

- Carry both evidence streams into print as decisive and complementary; the case's pedagogical value depends on the peer-reviewed null and the agency-audit disparate-impact finding traveling together rather than being separated into parallel entries.
- Pair in syllabi with Case 48 (Johnson) and Case 52 (Purdue) so the algorithmic-public-administration-in-education and reverse-causality threads are taught alongside the multi-year-deployment-with-null-causal-effect finding.
- Use the case as the anchor for the designing-predictions-to-trigger-support LEO; the curricular target is the multi-year deployment record that demonstrates the gap between the prediction's operational theory and the evidence the deployment produces.

## REFLECTION QUESTIONS

1. Identify a prediction system in your domain that has operated at scale for years without RDD-style causal evaluation of whether the prediction triggers outcome-changing intervention. What would the causal-evaluation design look like in your setting, and what would a null finding require of the deployment?
2. Specify the equity-audit-as-binding-input discipline you would apply when an internal audit finds disparate accuracy across protected-attribute subgroups. What is the documented decision rule for modifying or withdrawing the system, and what is the rule when the agency continues operation unchanged?
3. The case carries both the peer-reviewed null and the disparate-impact-investigation finding in one entry. Pick a deployment in your domain that has evidence streams in tension, and ask: what is the editorial discipline that carries the streams together rather than separating them, and what would separating them lose?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

# Atlanta Public Schools Cheating Scandal

**G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** 178 educators named, 38 of them principals; cheating found in 44 of 56 schools examined; 35 indicted; 11 convicted under RICO statute; foundational U.S. education measurement-gaming case

## IN BRIEF

Under a celebrated superintendent, Atlanta Public Schools was held up nationally as a high-performing urban district. The performance was substantially fabricated. A state special investigation, supported by the Georgia Bureau of Investigation, and reporting by the Atlanta Journal-Constitution found that for several years educators across dozens of schools had systematically erased and corrected students' answers on state standardized tests. The cheating was organized — principals pressured teachers, staff held weekend "erasure parties" — and the incentive system rewarded the gaming with bonuses and promotions. The state investigation named 178 educators, 35 indicted, and 11 convicted under Georgia's racketeering statute. The capability gap lay in the measurement architecture: the institution being measured operated the instrument that measured it, with no independent audit.

## BACKGROUND

During the 2000s, Atlanta Public Schools under superintendent Beverly Hall was celebrated nationally for rapid gains on Georgia's high-stakes standardized tests. Bonuses, public recognition, and job security were tied directly to those scores, and the district administered the state tests on which it was being judged and held the answer documents afterward. With the reward attached to the number and the number produced in-house, the people under pressure to improve the score also controlled the answer sheets that determined it.<sup>o</sup>

## WHAT HAPPENED

Over several years, educators at dozens of APS elementary and middle schools systematically changed students' answers after testing. The cheating was organized rather than incidental: 38 of the 178 educators named were principals, who pressured teachers to hit targets, and staff held weekend "erasure parties" to correct answer sheets. Suspiciously high rates of wrong-to-right erasures flagged the pattern — a statistical fingerprint the gaming left behind precisely because correcting a wrong answer to a right one is far rarer, by chance, than the reverse.<sup>1</sup>

## THE INVESTIGATION

A 2008 Atlanta Journal-Constitution analysis of improbable score gains, then a February 2010 erasure analysis by the Governor's Office of Student Achievement, then a Governor-ordered special investigation supported by the Georgia Bureau of Investigation,



documented the scheme in a 2011 report. Investigators found cheating in 44 of the 56 schools they examined on the 2009 CRCT and traced organized misconduct as far back as 2001. The report named 178 educators, 38 of them principals; 35 were indicted and 11 convicted under Georgia’s RICO statute, the report finding the administration had emphasized results and praise “to the exclusion of integrity and ethics.” That an investigation came from outside the district — a newspaper and state investigators, not the schools — is itself the mark of the missing independent check.<sup>2</sup>

#### THE CAPABILITY GAP

The capability gap was at the measurement architecture: the institution being measured also operated the instrument that measured it. High-stakes incentives without an independent audit are a textbook setup for Campbell’s Law — the more a measure is used for decision-making, the more it will be corrupted. The students were not learning; the system was rewarding the appearance of learning, so the score rose while the underlying capability it was meant to certify went unmeasured and, for the children, unmet.<sup>3</sup>

#### AFTERMATH & REFORM

The convictions made APS the most prominent U.S. case of high-stakes-testing fraud and fuelled a broader reassessment of accountability-testing regimes. The structural lesson — that a measurement used for consequential decisions needs an audit independent of the institution being measured — is direct, but is still rarely implemented at scale, because the independent audit costs money and political will that the high score, once reported, makes seem unnecessary.<sup>4</sup>

# 178

EDUCATORS NAMED · 44 OF 56 SCHOOLS EXAMINED

*the institution being measured operated the instrument that measured it*

ATLANTA PUBLIC SCHOOLS — MEASUREMENT GAMING UNDER HIGH-STAKES TESTING

*The superintendent and her administration emphasized test results and public praise to the exclusion of integrity and ethics.*

GEORGIA SPECIAL INVESTIGATORS’ REPORT ON ATLANTA PUBLIC SCHOOLS, 2011

#### REFERENCES

1. Office of the Governor, Special Investigators (Bowers, Wilson, Hyde), Special Investigation into Test Tampering in Atlanta’s School System, Vol. 1 (June 30, 2011; released July 2011) — cheating in 44 of 56 schools examined, 178 educators named including 38 principals. The verified quoted span is “emphasized test results and public praise to the exclusion of integrity and ethics”; the subject clause is the reporting’s, not the report’s — do not re-quote it as the report’s own sentence.
4. Campbell, D. (1976), “Assessing the Impact of Planned Social Change” — Campbell’s Law.
5. Koretz, D. (2017), The Testing Charade — high-stakes-testing distortion.

2. Atlanta Journal-Constitution investigative series (2008–2011) — the December 2008 analysis of statistically improbable CRCT scores and the subsequent erasure-rate reporting.
3. State of Georgia v. Hall et al. (2013–2015) — indictments and RICO convictions.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The APS cheating scandal is the strongest available case for measurement-gaming under high-stakes incentives in education. The capability gap was at the measurement architecture: the institution being measured operated the instrument that measured it. The reform pattern is direct — independent measurement audit — but rarely implemented at scale.

### LENS APPROACH

APS is the protect-the-instrument-from-the-operator-who-controls-it case (induced 2.2): the educators measured on the test also administered and scored it, so the instrument was capturable by the very people it judged. LENS pairs it explicitly with VA wait-times (Case 7), a same-shape instrument-integrity failure where the unit being measured operated the measurement and quietly rewrote the number. The lesson is an instrument-integrity design choice — separate the measured from the measurer, and audit independently of the institution being judged — not the surface reading that people cheated under pressure.

| DOMAIN             | FAILURE MODES       | LENS COMPETENCY                                     | PROBLEM TYPE | INDUCED | LEO |
|--------------------|---------------------|-----------------------------------------------------|--------------|---------|-----|
| EDUCATION AT SCALE | → <b>G</b> <b>N</b> | → <b>D1</b> <b>D2</b> <b>D3</b> <b>D4</b> <b>D5</b> | 3            | 2.2     | 4   |
| GOVERNMENT         | Test and Evaluation |                                                     |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Separate test administration and scoring from the institution being judged, so the party with the incentive to inflate the number never controls the instrument that produces it.
- ▶ Engineer a statistical integrity check — such as wrong-to-right erasure-rate monitoring — into the measurement system before scores carry consequences.
- ▶ Anticipate Campbell's Law when attaching bonuses and job security to a metric, and design the audit as part of the incentive, not as an afterthought.

#### IN OPERATION / AFTER

- ▶ Run independent, ongoing audits of the high-stakes measure rather than relying on a newspaper or state investigators to surface fraud after years.
- ▶ Monitor for the appearance of learning diverging from independent evidence of learning, so the gap is caught while the score still means something.
- ▶ Sustain the audit's funding and authority against the complacency a high reported score creates, since the cost of checking looks unnecessary exactly when it is most needed.

### REFLECTION QUESTIONS

1. Identify a high-stakes measurement in your domain whose audit is operated by the institution being measured. What would change with independent audit?
2. Apply Campbell's Law to a current high-stakes measurement system. What distortion is predicted?
3. The erasure-rate analysis caught APS only after the fact. Design a continuous integrity signal — a statistical fingerprint of gaming — that would flag the pattern as it emerged rather than years later.

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

## CASE 52

HIGHER-ED ANALYTICS

EARLY-WARNING SYSTEMS

EVIDENCE ARCHITECTURE

2012 – 2013

# Purdue Course Signals — The Reverse-Causality Retention Claim

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Purdue's widely cited claim that students taking two or more Course Signals classes were 21% more likely to be retained was challenged by Mike Caulfield in 2013, who identified the dose-response curve as an artifact of selection: students persist and therefore take more Signals courses, not the reverse — Alfred Essa reproduced the apparent retention gain by substituting 'received a piece of chocolate' for 'took a Signals class' in a simulation

## IN BRIEF

Purdue's Course Signals was one of the most-cited early-warning learning-analytics interventions of the early 2010s. The program's headline outcome claim — students who took two or more Signals courses were 21 percent more likely to be retained — was published by Arnold and Pistilli at LAK 2012 and circulated widely in the learning-analytics community and in vendor materials. In 2013 Mike Caulfield, blogging at Hapgood and then at e-Literate, identified the dose-response curve as a reverse-causality artifact: students who persist in college therefore take more Signals courses, so the apparent retention gain reflects selection, not effect. Alfred Essa built a simulation that substituted "received a piece of chocolate" for "took a Signals class" and reproduced the apparent retention gain, demonstrating the methodological flaw. The study was peer-reviewed for LAK and never reanalysed on its own data, yet became one of the most-referenced learning-analytics studies of its era — which is itself the cautionary point about the field's evidence architecture. The case is the small-tier methodological companion to v1 Cases 46 (educational predictive-analytics bias) and 80 (Georgia State predictive analytics).

## BACKGROUND

Course Signals at Purdue was a faculty-facing dashboard that classified students enrolled in a course as green, yellow, or red, based on a predictive model of academic risk built from LMS and grade-book signals. Faculty used the classification to send targeted communications to students flagged at risk. The design and the operational use were not the subject of the later critique; what became contested was the system's headline outcome claim, published by Arnold and Pistilli at Learning Analytics and Knowledge 2012: students who took two or more Signals courses were 21 percent more likely to be retained at the institution than students who took fewer.<sup>o</sup>

## WHAT HAPPENED

The figure circulated. It appeared in conference keynotes, vendor materials, accreditor presentations, and in widely shared accounts of what early-warning analytics could deliver. The claim's status in the literature outran its evidentiary base: the LAK 2012 paper was peer-reviewed but four pages long, and the institutional dataset behind the headline number was not made available for independent reanalysis — Purdue moved the underlying study off its own site during the dispute. The field cited the result anyway, because it was the kind of result the field wanted to be true.<sup>1</sup>

## THE INVESTIGATION

In 2013 Mike Caulfield, writing at Hapgood and then at e-Literate, asked a specific methodological question: was the dose-response curve — more Signals classes, more retention — what it appeared to be? Caulfield argued the relationship was a reverse-causality artifact. Students who persist at the institution have more semesters in which to take Signals courses; students who depart cannot. The “took two or more Signals courses” group was therefore an inadvertent selection on persistence — not a sample exposed to a different treatment intensity. The tell was in the paper's own numbers: taking **one** Signals course correlated with a **lower** retention rate than the non-Signals population, which no dose-response account of a beneficial treatment explains, and which a coin-flip account of course-load explains exactly. Alfred Essa then built a simulation that substituted “received a piece of chocolate” for “took a Signals class,” with chocolate having no causal effect on anything, and reproduced the apparent retention gain. The reverse-causality reading survived the simulation replication.<sup>2</sup>

## THE CAPABILITY GAP

The methodological point is precise: the published analysis did not isolate the treatment from the selection mechanism that determined treatment intensity, and the dose-response curve that looked like the effect was generated by the selection itself. The case is not an argument that Course Signals had no effect on retention. It is an argument that the published evidence could not distinguish effect from selection, and that the institution measured a number which **felt** like the failure mode it cared about (retention) using a design that could not actually answer the causal question. This is the textbook 2.1 failure: measuring the failure mode with a design the institution can deceive itself with.<sup>3</sup>

## AFTERMATH & REFORM

Drafted as a deeper-evidence-of v1 Cases 46 and 80, the Purdue case carries a named methodological failure into the corpus's predictive-analytics conversation. The cautionary thread runs through three places at once: the original study's standing in the field outpaced the four peer-reviewed pages it rested on; the field's citation practice amplified the headline without probing the design; and the correction (Caulfield, Essa) was mounted from outside the original study's institutional network. The case teaches the evidence-architecture failure mode that the LENS Iterative Development domain and the Navigating Sociotechnical Constraints domain both have to protect against — and that v1 Cases 46 and 80 anchor at the bias and the institutional-deployment layers respectively.

*Substituting “received a piece of chocolate” for “took a Signals class” in a simulation reproduces the apparent retention gain.*

EDITORS' SYNTHESIS OF ESSA'S 2013 SIMULATION DEMONSTRATION.

## REFERENCES

1. Arnold, K. E., & Pistilli, M. D. (2012). Course Signals at Purdue: Using learning analytics to increase student success. *Proceedings of LAK 2012*, 267–270. doi:10.1145/2330601.2330666 — the original study at the center of the critique.
2. Caulfield, M. (2013). “Why the Course Signals Math Does Not Add Up,” Hapgood, September 26, 2013, and “Purdue Course Signals Data Issue Explainer,” e-Literate, November 12, 2013 — the reverse-causality critique.
3. Straumsheim, C. (2013). “Mixed Signals,” *Inside Higher Ed*, November 6, 2013 — documents Alfred Essa's placebo simulation, which substituted a randomly distributed piece of chocolate for a Signals class and reproduced the retention curve.
4. Feldstein, M. (2013). “Purdue's Non-Answer on Course Signals,” e-Literate, November 2013 — the field-level critique of promoting research and then declining scrutiny.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Course Signals is the named methodological failure in the predictive-analytics conversation: the institution measured a number that felt like the failure mode it cared about, using a design that could not isolate effect from selection. The published claim was never reanalysed on the institution's own data; the correction came from outside the original network. Both are part of the cautionary reading.

### LENS APPROACH

Course Signals is the small-tier evidence-architecture failure (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for measurement design that cannot deceive the institution, and in Domain 5 (Navigating Sociotechnical Constraints) for the citation-practice failure mode in which a headline outcome outpaces the evidence it rests on. Deeper-evidence-of v1 Cases 46 (predictive-analytics bias) and 80 (Georgia State predictive analytics) — distinct because this is a named methodological-validity failure, not a bias finding.

| DOMAIN                | FAILURE MODES              | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|-----------------------|----------------------------|-----------------------|--------------|---------|-------|
| HIGHER-ED ANALYTICS   | <b>D</b> <b>K</b> <b>N</b> | D1 D2 D3 <b>D4</b> D5 | 5            | 2.1     | 4 · 5 |
| EARLY-WARNING SYSTEMS | →                          | → Test and Evaluation |              |         |       |
| EVIDENCE ARCHITECTURE |                            |                       |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Specify the causal question the analysis is built to answer in advance, and design the comparison group so that the selection mechanism into “treatment intensity” is not itself the variable doing the work.
- ▶ Pre-register the dose-response analysis with the threats-to-validity table — selection, censoring, reverse causality — visible at design time, so the published headline cannot outrun the analysis it rests on.
- ▶ Choose the publication venue that fits the evidentiary claim: conference proceedings for an early result, peer-reviewed journal with independent reanalysis for the figure the field will cite.

### IN OPERATION / AFTER

- ▶ Make the dataset available for independent reanalysis; the Caulfield / Essa correction succeeded because the published claim was simulatable on plausible data, not because the original dataset was inspected.
- ▶ Treat the simulation-replication finding as program evidence about the analysis design, not as a verdict on the intervention; Course Signals may have had an effect, and the published study could not detect or measure it.
- ▶ Carry the cautionary reading into the field's citation practice: a headline outcome circulating widely without peer review outside conference proceedings is itself an evidence-architecture failure, separate from any specific methodological flaw.

## REFLECTION QUESTIONS

1. Identify a predictive or early-warning analytics deployment in your domain whose published outcome claim circulates more widely than the peer-reviewed evidence supports. What threats to validity — selection, reverse causality, censoring — would a Caulfield-style external critic name first?
2. Specify the comparison-group design that would isolate effect from selection in your context. What pre-registered analysis plan, with simulated placebo treatment, would let the field check the claim before it circulates?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

# inBloom

## G Governance & Trust

**IMPACT** \$100M initiative collapsed in ~14 months; all 9 announced partner states pulled out or disavowed participation; data infrastructure for education set back years

### IN BRIEF

inBloom was a \$100-million, Gates- and Carnegie-funded shared data infrastructure for U.S. student records — and the technology worked: no bug, no breach, no performance failure. What killed it in about fourteen months was everything around the technology. It launched without consent frameworks, community engagement, transparency about data use, or a way for parents to participate in decisions about their children's data. Parent-privacy groups organized opposition state by state, and the coalition of nine shrank to three in seven months. Analysts read inBloom as low public tolerance for risk meeting an initiative that never communicated its benefits — the assumption that technically sound infrastructure generates its own legitimacy. It is the purest governance failure in the dataset, and the book's clearest argument that in education at scale, stakeholder trust and governance are not optional features but load-bearing structure.

### BACKGROUND

inBloom was an ambitious shared data infrastructure for U.S. K-12 student records, backed by roughly \$100 million from the Gates Foundation and the Carnegie Corporation of New York, hosted on commercial cloud, and built by enterprise engineers. The premise was that a common store would spare districts from rebuilding the same plumbing and let applications interoperate across systems that had never spoken to one another. Technically it was sound — no bug, no breach, no performance failure ever undid it, and that very soundness is what makes the case instructive.<sup>o</sup>

### WHAT HAPPENED

What undid it was everything around the technology. inBloom launched without adequate consent frameworks, without meaningful community engagement on data governance, without ever explaining what the platform was for, in terms a teacher or parent could use, and without any way for parents to participate in decisions about their children's data. Each omission read, to a worried parent, as a decision made about their child without them in the room. Parent groups organized opposition state by state, and the coalition of nine shrank to three in seven months — Louisiana out first, then Kentucky, Georgia and Delaware, then Colorado's pilot district — as the political cost of staying overtook any promised efficiency. New York, the last and largest partner, was barred from sharing student data with inBloom by a provision in the state budget enacted at the end of March 2014, and on 21 April 2014

inBloom announced it would wind down. Within about fourteen months the \$100-million initiative had collapsed.<sup>1</sup>

#### THE INVESTIGATION


Analysts at Data & Society read the collapse as two forces meeting: a public with low tolerance for risk and uncertainty, and an initiative that never communicated its benefits or won buy-in from the stakeholders it needed. Legitimacy, on that reading, is earned from the people a system acts upon, not conferred by the quality of its engineering. The technology was never the problem; the governance was — the consent, transparency, and trust that had been treated as add-ons rather than as the foundation the whole effort needed before a single record moved.<sup>2</sup>

#### THE CAPABILITY GAP

inBloom is the purest governance failure in this dataset: nothing technical was wrong, and everything sociotechnical was. The missing capability was the design of stakeholder trust — consent, accountability, and a voice for the families whose data was at stake — treated as a precondition for deployment rather than a feature to add later. Once opposition formed, no patch could retrofit the trust that should have been built in from the start, because trust withheld at launch cannot be engineered back in under fire. In education at scale, those are load-bearing elements, not optional ones.<sup>3</sup>

#### AFTERMATH & REFORM

inBloom's collapse set back shared education-data infrastructure for years and became a standard cautionary tale; it also helped drive a wave of state student-data-privacy laws that codified, after the fact, the consent and transparency the project had skipped.<sup>4</sup> The lesson the book takes from it is that ethics-as-design-constraint is not ideology but engineering — and inBloom is the \$100-million empirical test of what happens when you skip it, paid not in downtime but in a dead initiative and a chilled field.



*inBloom was a great solution, but it also highlighted a policy and governance gap that was in K-12, and unfortunately that gap eroded the value proposition of inBloom.*

JAT PANNU, COO OF THE ILLINICLOUD AND K-12 FEDERATION, INTERVIEWED IN BULGER, MCCORMICK & PITCAN, DATA & SOCIETY, 2017

#### REFERENCES



1. M. Bulger, P. McCormick & M. Pitcan, The Legacy of inBloom, Data & Society Research Institute (2017) — inBloom as a failure of technocratic reform.
2. Education Week and Hechinger Report coverage of the state withdrawals (2013–2014) — nine states exiting.
3. Bulger et al. (2017) — the diagnosis: low public tolerance for risk and uncertainty meeting a failure to communicate benefits and win stakeholder buy-in.
4. N. Selwyn, Distrusting Educational Technology (2014); d. boyd & K. Crawford, “Critical Questions for Big Data” (2012).
5. Parent Coalition for Student Privacy archives and the wave of state student-data-privacy legislation that followed inBloom.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

inBloom is the purest governance failure in this dataset. Nothing technical was wrong. Everything sociotechnical was. The case is the strongest argument in the book for treating ethics-as-design- constraint not as ideology but as engineering: a \$100M empirical test of what happens when you do not.

### LENS APPROACH

LENS uses inBloom in LEN 7 as the canonical governance failure and in LEN 1 as a stakeholder-analysis case. The case anchors the LENS threading commitment that equity and accountability are design commitments, not modules. Studio projects (LEN 10) require students to produce a stakeholder-trust deliverable as a precondition for deployment.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY                                     | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|-----------------------------------------------------|--------------|---------|-----|
| EDUCATION AT SCALE                    | → <b>G</b>    | → <b>D1</b> <b>D2</b> <b>D3</b> <b>D4</b> <b>D5</b> | 4            | 5.1     | 5   |
| Navigating Sociotechnical Constraints |               |                                                     |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Treat consent, transparency, and a parent-facing voice as requirements gathered before the data store is built, not features bolted on after launch.
- ▶ Engage the families and districts whose data is at stake as design stakeholders from the outset, so the governance questions surface in requirements rather than in opposition.
- ▶ Make legitimacy an explicit deliverable: document who must agree, on what terms, before any student record moves into shared infrastructure.

#### IN OPERATION / AFTER

- ▶ Audit live deployments for the gap between technical soundness and stakeholder trust, since a clean system can still be losing the political ground it stands on.
- ▶ Monitor state-by-state consent and withdrawal signals as a leading indicator, treating organized parent opposition as data about a governance defect, not noise.
- ▶ Sustain a standing transparency channel so families can see what is collected and why throughout operation, not only at adoption.

### REFLECTION QUESTIONS

1. What is the equivalent unbuilt governance infrastructure in your domain? What would the \$100M empirical test of its absence look like?
2. Design the stakeholder-trust deliverable that a future inBloom-equivalent should have to produce before deployment.
3. inBloom’s engineers were excellent and its technology never failed, yet the project collapsed. Where in your own domain is technical soundness being mistaken for legitimacy — and who would have to consent before a system you build could claim it?

**WHO BUILDS THIS** governance, ethics & measurement — plus domain experts and a learning engineer to integrate.

**TOOLS** governance frameworks · bias & evidence audits

# Summit Learning / Personalized Learning Rollout

**G** Governance & Trust   **T** Training Gap   **K** Knowledge & Institutional Memory

**IMPACT** Personalized-learning platform deployed across ~380 U.S. schools and 72,000 students; parent and student revolts in Cheshire, Brooklyn, McPherson and Wellington; districts suspended, scaled back, or conceded opt-outs between 2017 and 2019

## IN BRIEF

Summit Learning, a personalized-learning platform from Summit Public Schools backed by the Chan Zuckerberg Initiative, was offered free to U.S. districts from 2015 and reached roughly 380 schools and 72,000 students by 2018. Adopters pulled back under parent and student pressure — Cheshire, Connecticut suspended the platform midyear in 2017, a Brooklyn high school cut it for eleventh and twelfth grade after a student walkout in 2018, and McPherson, Kansas conceded an opt-out in 2019 — amid complaints about screen time, disengagement, and data privacy. The pedagogy itself (competency-based progression, projects, mentoring) was defensible and often effective; what failed was deployment governance. There was no evaluation framework districts could read before adopting, no parent-facing data agreement, and no exit path independent of the vendor. Summit is a clean test of the book's claim that the governance architecture must be engineered alongside the tool.

## BACKGROUND

Summit Learning was a personalized-learning platform developed by Summit Public Schools with technical and financial support from the Chan Zuckerberg Initiative, offered free to U.S. districts from 2015 — a price that lowered the bar to adoption while raising no governance questions at the door. Its pedagogy — competency-based progression, self-directed projects, mentor check-ins — was defensible and in many places effective, which is why the eventual revolt could not be blamed on the instructional design.<sup>o</sup>

## WHAT HAPPENED

By 2018 the platform reached roughly 380 schools and an estimated 72,000 students, scaling fast on the strength of a free offer and a well-funded sponsor. By 2019 its most visible adopters were pulling back under parent and student pressure: about a hundred students at Brooklyn's Secondary School for Journalism walked out in November 2018 and the school cut the program for eleventh and twelfth grade; Cheshire, Connecticut had suspended it midyear in December 2017 after a parent petition; McPherson, Kansas conceded an opt-out for up to 225 students after a January 2019 walkout. Walkouts and complaints about

screen time, eye strain, disengagement, and data privacy converged into a revolt that was not about the instructional design at all.<sup>1</sup>

#### THE INVESTIGATION

Press coverage and later analyses located the failure in deployment governance, not pedagogy: there was no academic research on whether Summit's model worked, and Summit had declined the Harvard evaluation whose planning it helped fund while still naming that collaboration in its own materials, no parent-facing data-handling agreement, and no exit pathway that did not depend on the vendor's goodwill — three absences that each became a grievance the moment families looked for them. The implementation never surfaced the governance questions parents would ask, so when those questions arrived they arrived as opposition rather than as design input, and the argument was lost before it started.<sup>2</sup>

#### THE CAPABILITY GAP

Summit is a clean test of the book's central claim: a technology that worked at the pedagogical level still failed because the governance architecture — consent, evidence, measurement, exit — had not been engineered alongside it. A working tool with no accountability contract is a liability waiting for the first organized objection. The pattern recurs across the ed-tech dataset (inBloom, Case 53): a well-intentioned tool, a well-funded rollout, and no institutional contract with the families and teachers operating inside it — the same omission producing the same collapse in a second case.<sup>3</sup>

#### AFTERMATH & REFORM

Several districts withdrew or rebranded their use, CZI and Summit revised their outreach, and the episode became a standard caution in ed-tech adoption — a reputational cost paid for governance work that would have been cheaper to do first. The churn was larger than Summit said: pushing back on the Brooklyn coverage, the organization told reporters that only 10 percent of schools quit the platform each year, but records Chalkbeat obtained from Summit itself put the real figure at 18 percent of schools quitting a year after using it — the 10 percent came from averaging the first three years, when the denominator was 19 schools, then 338.<sup>4</sup> Its lesson for the field is concrete: an adoption decision should have to produce a public evidence summary at parent reading level, a data-handling agreement at the same resolution, and a documented exit path — governance artifacts that make a tool's deployment legitimate, not just its design sound, and that turn the questions parents will ask into inputs gathered before launch rather than weapons raised after it.

*The tools were free. The accountability architecture had not been built.*

EDITORS' SYNTHESIS OF SUMMIT LEARNING ROLLOUT COVERAGE (NEW YORK TIMES, WIRED, EDUCATION WEEK, 2018–2019)

#### REFERENCES

1. N. Singer, "The Silicon Valley Billionaires Remaking America's Schools," New York Times (2017) — the CZI/Summit rollout.
2. L. Genota, "Brooklyn Students Protest Use of Online Learning Platform Designed by Summit Learning," Education Week (November 13, 2018) — the 380 schools and 72,000 students figures and the walkout of nearly 100 students at the Secondary School for Journalism.
3. M. Barnum, "Summit Learning declined to be studied, then cited collaboration with Harvard researchers anyway," Chalkbeat (January 17, 2019) — the missing evaluation framework.
5. M. Barnum, "Summit Learning, the Zuckerberg-backed platform, says 10% of schools quit using it each year. The real figure is higher," Chalkbeat (May 23, 2019) — the 18 percent annual attrition figure from Summit's own records.
6. N. Bowles, "Silicon Valley Came to Kansas Schools. That Started a Rebellion," New York Times (2019) — the parent revolt.
7. S. Schwartz, "Two Districts Roll Back Summit Personalized Learning Program," Education Week (December 22, 2017) — the Cheshire suspension and the Indiana Area rollback.

4. Chan Zuckerberg Initiative & Summit Learning public program documentation (2015–2019); cf. inBloom (Case 53).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Summit Learning is a clean test of the book’s central claim: technology that worked at the pedagogical level still failed because the **governance** architecture (consent, evidence, measurement, exit) had not been engineered alongside it. The pattern — well-intentioned tool, well-funded rollout, no institutional contract with the families and teachers operating inside it — recurs across the educational-technology dataset (Cases 53, 139, 67) and is the educator’s-side analog of the governance failures in Cases 49 and 191.

LENS APPROACH

LENS uses Summit Learning in LEN 7 as the foundational consent-and-evidence case for educational technology, and in LEN 10 as a studio prompt for the governance artifacts that any educational-technology adoption decision should produce: a public evidence summary at parent reading level, a data-handling agreement at the same resolution, and a documented exit pathway that does not depend on the vendor’s goodwill.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|------------------|--------------|---------|-----|
| EDUCATION AT SCALE                    | → G T K       | → D1 D2 D3 D4 D5 | 4            | 5.4     | 5   |
| Navigating Sociotechnical Constraints |               |                  |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>▶ Engineer the governance architecture — consent, evidence, measurement, exit — in lockstep with the pedagogy, so a sound tool ships with a sound accountability contract.</li><li>▶ Produce, before adoption, a public evidence summary and data-handling agreement at parent reading level that answer the questions families will raise.</li><li>▶ Build a documented exit pathway that does not depend on the vendor’s goodwill, so a district can leave without being captured.</li></ul> | <ul style="list-style-type: none"><li>▶ Treat organized parent and student objections as governance signal about a missing contract, not as resistance to the instructional design.</li><li>▶ Monitor adopting districts for the early grievances — screen time, privacy, disengagement — that precede a withdrawal, and route them to a decision-maker.</li><li>▶ Maintain a re-adoption pathway so a withdrawn district can return only after completing the governance work that the first rollout skipped.</li></ul> |

REFLECTION QUESTIONS

1. What is the equivalent of the “free tool, free of governance” pattern in your domain — the offer that bypasses the accountability architecture because it does not yet exist?
2. Design the parent-reading-level governance artifact that a district should require before adopting an educational-technology platform.
3. Summit’s withdrawals in Brooklyn, Cheshire, and McPherson were led by parents, not regulators. What is the equivalent local constituency in your domain that institutional accountability has not yet accommodated, and how would they be heard before deployment rather than after?

**WHO BUILDS THIS** governance, ethics & measurement · learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · scenario simulation · competency models · knowledge bases · data pipelines

## HIGHER EDUCATION

## CASE 55 PREDICTIVE ANALYTICS

2010S – PRESENT (BROOKINGS SYNTHESIS 2021)

## ACCESS PRICING

# Enrollment-Algorithm Yield Optimization Across U.S. Higher Education

**T** Training Gap **K** Knowledge & Institutional Memory **N** Normalization of Deviance

**DISCLOSURE** Institutional overlap: John Katzman collaborates with an editor's institution (Johns Hopkins) but was not involved in this work. He is not associated with the specific vendors named here.

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** An academic study reports a 23.3% out-of-state yield gain (12.1% to 14.8%) at an unnamed large public university; vendor case studies report 33% net tuition increases with a 6-point cut to discount rate (EAB) and 173 additional freshmen without aid-budget increases (Othot); algorithms from at least six named vendors reach some 700 institutions, predicting each accepted applicant's likelihood to enroll, with the second stage — pricing the aid offer against modeled willingness to pay — emerging rather than established, and Othot's optimizer shipping set by default to minimize scholarship aid while maintaining enrollment

## IN BRIEF

Engler's 2021 Brookings paper documents the two-stage architecture of contemporary enrollment-management algorithms: predict each accepted applicant's probability of enrollment, then optimize the financial-aid offer to maximize either net tuition revenue or yield. He names the vendor landscape — Ruffalo Noel Levitz, EAB, Rapid Insight, Capture Higher Ed, Othot, Whiteboard Higher Education, and others — touching at least 700 U.S. institutions. Vendor-reported case studies cite large gains in yield or tuition revenue per matriculant. The structural critique is the inversion of Case 80 (Georgia State, where prediction triggered support): here, the algorithm identifies "willingness to pay" so the institution can offer the minimum scholarship that will still yield enrollment, reducing aid per low-income student. The honest hedges Engler preserves are binding: critical details are obscured by vendors and colleges; algorithmic optimization is hard to separate from manual leveraging; without auditing specific college data and models, fairness impacts on protected classes cannot be confirmed. The evidence-tier flag under the title carries the standing language; future validation ongoing.

## BACKGROUND

The structural seam the case opens is that financial-aid distribution is itself a deployed prediction system, operating across hundreds of U.S. institutions, with measurable consequences for who attends college and how much they pay. Engler's Brookings synthesis is the most thorough public mapping of the deployment surface, drawing on vendor case studies and the higher-education enrollment-management literature. The first stage of the two-stage architecture is prediction: estimate the probability that an accepted applicant will enroll, given observable attributes from the application, the FAFSA, and third-party data. The second stage is optimization: choose the financial-aid offer that maximizes a chosen objective — net tuition revenue per matriculant, total tuition revenue, or class size at a target discount rate.<sup>9</sup>

## WHAT HAPPENED

The vendor landscape Engler names is large and concentrated. Ruffalo Noel Levitz works with roughly 300 institutions; EAB with 150; Rapid Insight with 150; Capture Higher Ed with 100; Othot with around 30; Whiteboard Higher Education rounds out the named tier, with Civitas Learning named separately as a firm exclusively focused on student success. A 2015 Educause survey put analytics for enrollment management at over 75 percent of colleges, up from just over 60 percent in 2012. The reported effect sizes: a 23.3 percent out-of-state yield gain — 12.1 percent to 14.8 percent — at an unnamed large public university in a University of Washington study; EAB's 33 percent net-tuition gain at "Aster University," paired with a six-point cut to the tuition discount rate from around 22 to around 16 percent; Othot's 173 additional freshmen at the New Jersey Institute of Technology, recruited without a corresponding rise in scholarship aid. Two hedges travel with these numbers. The vendor effect sizes are vendor-reported and are not auditable from outside the institution. And the second stage is the emerging one: scholarship optimization has historically been done manually, and Engler writes that it is "not clear that this approach is widely in use" — the 700 institutions are running the prediction step, not necessarily the optimizer.<sup>1</sup>

## THE INVESTIGATION

The structural critique runs through what the algorithm is optimized for. The chosen construct — willingness to pay — is the operational target the prediction system was built around, and Engler's argument is that this construct turns the financial-aid award into a pricing instrument rather than a need-or-merit one. The clearest evidence that the construct is the operational target and not an inference about it is Othot's default: Engler reports that the vendor's optimizer ships set to minimize scholarship aid while maintaining enrollment. The downstream evidence he marshals is the literature linking institutional aid to graduation outcomes: roughly a 0.9 percent gain in graduation odds per additional \$1,000 in merit aid, with larger effects for need-based aid; a more than 5 percent cut in low-income graduation likelihood per additional \$1,000 in unsubsidized loans. If the algorithm reduces aid per student to find the minimum that still yields enrollment, the downstream cost is the completion gap that grant aid was buying down.<sup>2</sup>

## THE CAPABILITY GAP

The case sits as the structural inverse of Case 80 (Georgia State's predictive advising) and pairs with Case 186 (mortgage-lending fairness) and Case 86 (community-college predictive equity). Georgia State used prediction to trigger support; enrollment-management

algorithms use prediction to reduce the help allocated. Bartlett's lending analysis names the same construct-substitution pattern at the pricing layer across a different deployed prediction system; Gándara's community-college work names it at the access layer. The anchor the three cases share is the inversion of the gatekeeping-vs-support frame: the prediction is used to gatekeep the help, and the gatekeeping is invisible at the applicant's end of the transaction.<sup>3</sup>

#### AFTERMATH & REFORM

The hedges Engler names are binding on the case's framing, and the evidence-tier flag's standing language — future validation ongoing — applies in a precise sense. Critical details of the optimization are obscured by vendor non-disclosure and college contracting confidentiality; algorithmic optimization is hard to distinguish from manual "leveraging" of the same logic by human enrollment officers; without auditing specific college data and models, the protected-class fairness questions cannot be answered with the precision Bartlett-class auditing would require. The Brookings synthesis is the strongest public evidence available; it is policy-tier analysis built on vendor case studies and Engler's read of the operational record, and the case rests at that tier, not at the audited-deployment tier the corpus would prefer.

*The algorithms excel at identifying a student's exact willingness to pay. The construct is the decisive fairness decision, and the institution that does not name it has nonetheless made it.*

EDITORS' SYNTHESIS OF ENGLER (2021), BROOKINGS INSTITUTION.

#### REFERENCES

1. Engler, A. (2021), "Enrollment algorithms are contributing to the crises of higher education," Brookings Institution, 14 September 2021 — the two-stage architecture, the vendor tally to at least 700 institutions, and the effect sizes. The phrase "algorithms excel at identifying a student's exact willingness to pay" is Engler's own, from the executive summary.
2. Franke, R. (2012), Towards the Education Nation: Revisiting the Impact of Financial Aid, College Experience, and Institutional Context on Baccalaureate Degree Attainment, UCLA doctoral dissertation — the aid-and-graduation effect sizes Engler cites.
3. Goldrick-Rab, S. (2016), Paying the Price — broader synthesis on net-price, unsubsidized loans, and low-income completion.
4. Vendor case studies cited in Engler (Othot, EAB, Ruffalo Noel Levitz, University of Washington) — vendor-reported and not externally audited; flagged at evidence-tier under the title.

#### THE LEARNING ENGINEERING LENS

##### LE INSIGHT

Enrollment-algorithm yield optimization is the construct-choice case at the pricing layer of higher-education access: prediction is used to reduce aid per applicant, not to trigger support, and the operational target is "willingness to pay." The evidence-tier flag is binding — vendor case studies are not auditable, the algorithmic-vs-manual distinction is bounded, and the protected-class fairness question requires audit access the public synthesis cannot supply. Future validation ongoing.

##### LENS APPROACH

Engler / enrollment-management is the construct-choice case at population scale (induced 8.3; LENS D5/PT5). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the construct-substitution anchor and the disclosure architecture the deployment lacks, and in Domain 4 (Test and Evaluation) for the evidence-tier discipline — practice-synthesis is the strongest available tier, and the case says so. Pair



with Case 80 (Georgia State support-trigger inversion), Case 186 (Bartlett lending fairness), and Case 86 (Gándara community- college equity). coi-light render under the title is binding.

| DOMAIN               | FAILURE MODES | LENS COMPETENCY                         | PROBLEM TYPE | INDUCED | LEO   |
|----------------------|---------------|-----------------------------------------|--------------|---------|-------|
| HIGHER EDUCATION     | T K N         | D1 D2 D3 D4 D5                          | 5            | 8.3     | 4 · 5 |
| PREDICTIVE ANALYTICS | →             | → Navigating Sociotechnical Constraints |              |         |       |
| ACCESS PRICING       |               |                                         |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Name the optimization construct explicitly. “Willingness to pay” is not “need” and is not “merit”; the choice of construct is the decisive fairness decision, and the institution that does not name it has nonetheless made it.
- ▶ Require vendor disclosure of the model’s inputs, training data, and objective function as a condition of contracting; the case’s evidence-tier limit is partly the result of contractual non-disclosure that institutions could refuse to sign.
- ▶ Tie the deployed algorithm’s outputs back to downstream completion data; the literature linking aid to graduation odds is the evidence base the optimization should be tested against, not abstracted from.

### IN OPERATION / AFTER

- ▶ Commission an external audit of the deployed enrollment-management model against protected-class outcomes; the audit Engler says cannot be done from outside the institution is the audit institutions can choose to commission from inside.
- ▶ Publish the discount-rate-and-completion link as a paired metric; institutions that report only net-tuition gains and not the completion consequences are publishing a partial scorecard.
- ▶ Treat the gatekeeping-vs-support inversion as a curricular pattern: pair this case in syllabi with Case 80 so the design choice — which direction the prediction points — is taught as the design choice, not as an institutional default.

## REFLECTION QUESTIONS

1. Identify a deployed prediction system in your domain whose optimization target is named on the institutional side and obscure on the applicant or user side. What is the construct the optimization is built around — and what would change if the construct were named at the point of transaction?
2. Specify a vendor-disclosure clause you would require as a condition of contracting an enrollment-management or analogous optimization system. What inputs, training data, and objective function would the institution insist on auditing, and which would the vendor be willing to disclose under contract?
3. The case sits as the inversion of Case 80 (Georgia State, prediction to trigger support). Pick a prediction system in your domain and ask: in which direction does the prediction point — toward more help or less — and where is that design choice documented?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards



# GAO Online Program Manager Oversight Gap (GAO-22-104463)

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**DISCLOSURE** Institutional overlap: John Katzman collaborates with an editor's institution (Johns Hopkins) but was not involved in this work. Katzman founded 2U — central to the OPM debate, though not named in this GAO report.

**IMPACT** GAO found at least 550 colleges contracted with OPMs to support at least 2,900 academic programs as of July 2021; in a survey GAO cites, of the 25 respondent colleges paying their OPM a share of tuition revenue, 16 paid between 41% and 60% and four paid 61% or more; instructions to auditors and colleges lacked key details on OPM arrangements, leaving a risk that ED would not have the information it needs to detect incentive-compensation violations

## IN BRIEF

The Government Accountability Office's April 2022 audit (GAO-22-104463) names a structural oversight gap in the federal regime governing online program managers. As of July 2021, at least 550 colleges contracted with OPMs to support at least 2,900 academic programs; in survey data GAO cites, of the 25 respondent colleges paying their OPM with a share of tuition revenue, 16 paid between 41 and 60 percent and four paid 61 percent or more. The 1992 Higher Education Act amendments prohibited incentive compensation for student recruiters as a fraud-prevention measure; the Department of Education's 2011 "bundled-services" guidance exempted OPMs from the ban when recruiting was bundled with other services. The GAO found that instructions to auditors and to colleges lacked key details about identifying OPM arrangements, leaving a risk that the Department would not have the information it needs to detect violations. 2U's July 2024 Chapter 11 filing closed one boundary of the policy debate at the commercial level; at the federal level the Department opened a review of the guidance but did not rescind it — the January 2025 Dear Colleague Letter reaffirmed the 2011 bundled-services guidance. The underlying delegation- and-oversight problem persists for successor OPMs and revenue-share structures.

## BACKGROUND

The federal regime against incentive compensation in higher-education recruiting dates to 1992 and was built to address a specific fraud pattern: recruiters paid by enrollment will enroll students the program cannot serve, because the recruiter's compensation is tied to the enrollment rather than to the student's outcome. The 1992 Higher Education Act

amendments banned it for any person or entity engaged in student recruiting or admissions. The 2011 “bundled-services” guidance the Department of Education issued exempted online program managers from the ban when student recruiting was bundled with other services such as program design, platform delivery, and student support, and further safeguards were met. The exemption it created runs to the payment the college makes to the vendor: the statutory ban still prohibits the OPM’s own recruiters from being paid on their success in enrolling students. That is the boundary the construct allowed the revenue-share contracting model to grow across the next decade without anyone auditing.<sup>0</sup>

#### WHAT HAPPENED

The GAO’s audit documents the scale the regime grew to. By July 2021, at least 550 colleges had contracted with OPMs to support at least 2,900 academic programs. In the Eduventures Research survey of 145 college leaders that GAO cites, 49 respondents’ colleges contracted with an OPM, 32 of those contracted for student recruiting, and 25 of the 32 paid the OPM with a share of tuition revenue; of those 25, 16 paid between 41 and 60 percent of tuition and four paid 61 percent or more. The OPM operated under the university brand — recruiting, marketing, and program operations conducted by OPM employees identifying as the university — while receiving compensation tied directly to the enrollments those operations generated. The structural seam the regime had created was an exemption that allowed the prohibited compensation structure as long as the structure was administered by a contracted vendor.<sup>1</sup>

#### THE INVESTIGATION

The oversight gap the GAO documents is operational. Auditor instructions did not reference OPMs or the 2011 guidance, so compliance audits might not assess relevant arrangements; instructions to colleges lacked the detail that would have made them report those arrangements, and GAO found a risk that the Department would not have the information it needs to detect violations. Education’s own data made the gap measurable: 16 completed compliance audits and five completed program reviews had reported incentive-compensation findings since 2012, and not one involved a college’s arrangement with an OPM. The audit’s central finding is not that the OPM regime was designed to fail; it is that the oversight architecture required to police the exemption’s boundaries was never built, and the contracting structure that the exemption allowed grew faster than the monitoring capacity that would have surfaced violations.<sup>2</sup>

#### THE CAPABILITY GAP

2U’s July 2024 Chapter 11 filing closed one boundary of the policy debate at the commercial level; at the federal level the Department opened a review of the bundled-services guidance in 2023–2024 but did not rescind it, and its January 2025 Dear Colleague Letter reaffirmed the 2011 guidance. Neither closed the underlying delegation-and-oversight problem. Successor OPMs and revenue-share structures continue to operate; the universities that delegated student recruitment under the pre-2024 regime retain the operational dependencies and the brand-and-program commitments built during the decade of growth. The pair with Case 58 (USC × 2U Luna class action) shows the consumer-side litigation half of the same delegation; the pair with Case 55 (Engler / enrollment algorithms) shows the pricing-side optimization half.<sup>3</sup>

## AFTERMATH & REFORM

The case is investigation-grade — a GAO audit is the strongest tier of evidence the corpus carries for regulatory-oversight failure — and it is the structural delegation-with-revocation case at population scale. The university delegated student recruitment to a contracted vendor under a regulatory exemption that did not include the monitoring architecture the exemption's boundaries required. The delegation was reversible in principle: the contracts could be terminated, the exemption could be withdrawn, the operations could be brought back inside the university. In practice, the delegation accumulated operational and contractual lock-in across the decade the GAO audit covers, and the revocation when it came in 2024 was forced by commercial collapse rather than by the oversight architecture the audit recommended.

*Without clearer instructions to auditors and colleges about the information on OPM arrangements that must be assessed during compliance audits and program reviews, there is a risk that Education will not have the information it needs to detect incentive compensation violations.*

U.S. GOVERNMENT ACCOUNTABILITY OFFICE, GAO-22-104463 (2022).

## REFERENCES

1. U.S. Government Accountability Office (2022), "Higher Education: Education Needs to Strengthen Its Approach to Monitoring Colleges' Arrangements with Online Program Managers," GAO-22-104463.
2. U.S. Department of Education (2011), "Dear Colleague" guidance on incentive-compensation bundled-services exemption — the regulatory artifact the GAO audits.
3. Higher Education Amendments of 1992, Pub. L. No. 102-325, incentive-compensation prohibition, 20 U.S.C. § 1094(a)(20) — the statutory basis the 2011 guidance interpreted.
4. 2U Inc. (2024), Chapter 11 bankruptcy filing (July 25, 2024); U.S. Department of Education, "Dear Colleague" Letter (Jan. 14, 2025) reaffirming the 2011 bundled-services guidance (GEN-11-05) — the commercial-boundary closure alongside the federal guidance's continued validity.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

GAO-22-104463 is the investigation-grade delegation-with-revocation case at population scale. A regulatory exemption permitted a contracting structure across 550+ colleges and 2,900+ programs; the monitoring architecture required to police the exemption's boundaries was never built. 2U's July 2024 Chapter 11 closed the commercial boundary; ED opened a review but did not rescind the guidance (reaffirmed in the January 2025 DCL); the underlying delegation problem persists for successor OPMs.

### LENS APPROACH

GAO OPM oversight gap is the regulatory-seam case (induced 5.3; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the cross-regime delegation pattern and in Domain 3 (Human-System Collaboration) for the delegation-with-revocation frame — the contract is reversible in principle, locked-in in practice. Pair with Case 58 (USC × 2U Luna, the litigation half) and Case 55 (Engler / enrollment algorithms, the pricing-optimization half). coi-light render under the title is binding.

| DOMAIN               | FAILURE MODES | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO   |
|----------------------|---------------|---------------------------------------|--------------|---------|-------|
| HIGHER EDUCATION     | D K N         | D1 D2 D3 D4 D5                        | 6            | 5.3     | 5 · 3 |
| REGULATORY OVERSIGHT | →             | →                                     |              |         |       |
| ONLINE EDUCATION     |               | Navigating Sociotechnical Constraints |              |         |       |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Build the monitoring architecture as a condition of any regulatory exemption; the bundled-services exemption created a foreseeable contracting structure, and the oversight infrastructure to police its boundaries should have been built with it.</li><li>Require contract-disclosure as a federal reporting obligation, not as an institutional discretion; the GAO's data-collection finding is operationally addressable by mandatory reporting of revenue-share rates and recruiter-compensation arrangements.</li><li>Treat the delegation as reversible from the start: contract terms should preserve termination rights and operational-knowledge transfer; the lock-in the universities experienced was partly contractual and partly operational, and both halves are designable.</li></ul> | <ul style="list-style-type: none"><li>Carry the investigation-grade audit into the curriculum without softening: the GAO's central finding is that the oversight architecture was not built, and that is the teaching point.</li><li>Pair the case with Case 58 (USC × 2U) so the regulator-side audit and the consumer-side litigation are taught together; one half names what the regulator missed, the other names what the delegated marketing actually did.</li><li>Track post-review and post-2U-bankruptcy successor structures as a continuation of the case; the underlying delegation problem persists, and the case's frame is the regime-level oversight gap, not the specific 2U arrangement.</li></ul> |

REFLECTION QUESTIONS

1. Identify a regulatory exemption in your domain that permits a contracting structure without specifying the monitoring architecture required to police its boundaries. What would the audit-detectable artifact be — and is anyone currently collecting it?
2. Specify the contract terms you would require to preserve revocability of a delegated capability. Termination rights, operational-knowledge transfer, data portability — which of these are typically written into your domain's standard contracts, and which are left to negotiation?
3. GAO-22-104463 documents what happened when oversight architecture lagged contracting growth across a decade. What is the analog in your domain — a contracting pattern that grew under an exemption whose monitoring infrastructure was not built — and where would the equivalent investigation-grade audit have to come from?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

## CASE 58

HIGHER EDUCATION

ONLINE PROGRAM MANAGEMENT

PROFESSIONAL LICENSURE

2010S – 2024

# USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC)

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**DISCLOSURE** Institutional overlap: John Katzman collaborates with an editor's institution (Johns Hopkins) but was not involved in this work. Katzman founded 2U and left the company in 2012, before the USC MSW expansion and the 2019-onward conduct documented here.

**EVIDENCE TIER** journalism-tier — source confidence flagged; future validation ongoing.

**IMPACT** USC's MSW program grew from ~300 students per cohort pre-2010 to more than 3,000 students enrolled, almost entirely through online growth under the 2U partnership; 2023 class action alleges USC marketed the online program as the 'same' as the residential program while outsourcing recruiting, advising, and clinical placement to 2U employees; partnership wind-down announced November 2023

## IN BRIEF

USC's Master of Social Work program grew from about 300 students per cohort before 2010 to more than 3,000 students enrolled, almost entirely through online growth under the 2U partnership. The May 2023 Luna v. USC class-action complaint alleges that USC marketed the online program as the "same" as the residential program while outsourcing recruiting, advising, and clinical placement to 2U employees who carried usc.edu email addresses. Plaintiffs further allege that 2U marketers targeted students of color and veterans with aggressive, deceptive tactics, and that the online program's tuition (over \$100,000) reflected on-campus pricing while the delivered student experience did not. USC and 2U announced the partnership's termination on the MSW and other programs in November 2023. The case is still in litigation with no ruling on the merits — USC's demurrers were overruled in substantial part in April 2024 and again in September 2024, and a class-certification motion filed in November 2025 was set for argument in 2026 — and the predatory-targeting reconstruction rests on the complaint and on contemporaneous reporting rather than on a fact-finder's ruling — the evidence-tier flag is binding, and the case carries the journalism-tier framing with the standing language: future validation ongoing.

## BACKGROUND

The structural pattern the case names is the inversion of a stable institutional practice. USC's residential MSW program had been an established, modestly sized clinical training

program with a known student experience. The 2U partnership took the program's name and credential and built a parallel online operation at roughly ten times the cohort scale, with student-facing operations — recruiting, enrollment counseling, clinical-placement coordination — run by 2U employees who identified to applicants and enrollees with usc.edu email addresses. The university's delegation extended beyond program operations to the student-facing surface that defines what the credential means to the student paying for it.<sup>0</sup>

#### WHAT HAPPENED

The Luna class-action complaint, filed in Los Angeles County Superior Court in May 2023, structures its case around two central allegations. The first is misrepresentation: USC marketed the online program as the “same” as the residential program — same faculty, same standards, same credential — while the operational substance of the program had been outsourced. The tuition price (over \$100,000) tracked the residential pricing structure while the delivered student experience tracked the OPM-operated structure. The second is targeting: 2U marketers, the complaint alleges, focused aggressive and deceptive recruitment tactics on students of color and veterans — populations for whom the credential cost-and-mobility calculation is structurally different and for whom the misrepresentation has differential consequences.<sup>1</sup>

#### THE INVESTIGATION

The downstream consequences the complaint identifies extend into the licensure question. The MSW credential is a professional-licensure prerequisite, and clinical placement is the operational core of the training the licensure depends on. The complaint alleges that the delegated clinical-placement coordination — handled by 2U employees under usc.edu cover — did not consistently produce placements at the quality the residential program's reputation implied. The licensure-board half of the consequence chain — what graduates were actually able to do in practice, given clinical-placement quality — has not been independently studied, and the case carries that gap honestly rather than collapsing it.<sup>2</sup>

#### THE CAPABILITY GAP

USC and 2U announced the partnership's termination on the MSW and other programs in November 2023. The termination closed the operational arrangement at the boundary the litigation and the broader 2U commercial collapse forced; it did not adjudicate the legal claim, did not refund tuition paid under the prior arrangement, and did not produce an independent record of what the delegated operations actually delivered. The case sits at the consumer-side counterpart to Case 57 (GAO OPM oversight gap): one half is the regulator-side audit of the regime, the other is the litigation that names what happened to specific applicants and enrollees under the regime. Pair also with Case 55 (Engler / enrollment algorithms) for the pricing- optimization half.<sup>3</sup> That commercial collapse became terminal: 2U filed a prepackaged Chapter 11 in July 2024 and emerged in September 2024 as a privately held company, delisted from Nasdaq — the delegated vendor at the center of the arrangement ceasing to exist in the public-company form the disclosure record had assumed, and USC reabsorbing the programs it had outsourced.

#### AFTERMATH & REFORM

The journalism-tier evidence flag is binding on the case's framing. The complaint is a contested pleading. It has survived two rounds of demurrer — the Los Angeles Superior Court

overruled USC's demurrer to the First Amended Complaint on April 2, 2024 in substantial part, sustaining it only as to an Unruh Act claim with leave to amend, and denied the demurrer to the Second Amended Complaint on September 17, 2024 — but surviving a demurrer decides only that the allegations, taken as true, state a claim; it is not a finding that any of them happened. Plaintiffs filed a Third Amended Complaint on July 30, 2025 and moved for class certification that November. The partnership-termination record is contemporaneous reporting; neither is investigation-grade in the sense a GAO audit is, and the journalism-tier flag governs both. The predatory-targeting reconstruction relies on the complaint's allegations and on contemporaneous reporting — Higher Ed Dive, classaction.org summaries, the Project on Predatory Student Lending's statement — rather than on a fact-finder's ruling or an independent audit of 2U's marketing operations. Future validation ongoing on the litigation's outcome, on the licensure-board half of the consequence chain, and on whether the post-2024 successor-OPM contracts incorporate the disclosure architecture the case names as missing.

*The tuition tracked residential pricing. The delivered student experience tracked the OPM-operated structure. The gap between what was promised and what was delivered is the case.*

EDITORS' SYNTHESIS OF THE LUNA V. USC COMPLAINT AND CONTEMPORANEOUS REPORTING.

## REFERENCES

1. Stephanie Luna v. University of Southern California, class action complaint, Los Angeles County Superior Court, filed May 4, 2023 (First Amended Complaint Sept. 2023; Second Amended May 2, 2024; Third Amended July 30, 2025).
2. Higher Ed Dive reporting on the Luna complaint, May 2023; classaction.org and topclassactions.com summaries.
3. Project on Predatory Student Lending statement on USC-2U partnership termination, November 2023.
4. 2U Inc. and USC public statements on partnership termination, November 2023; broader 2U commercial-collapse reporting referenced through Case 57.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Luna v. USC is the consumer-side journalism-tier counterpart to GAO-22-104463 (Case 57). The complaint alleges USC marketed the online MSW as the "same" as the residential program while outsourcing the student-facing operations to 2U; the licensure-board half of the consequence chain is independently unstudied. Journalism-tier flag is binding — the predatory-targeting reconstruction is allegation-tier; future validation ongoing on litigation outcome and licensure consequences.

### LENS APPROACH

USC × 2U Luna is the disclosure-as-deliverable case at OPM-delegation scale (induced 5.4; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the promised-vs-delivered disclosure gap and in Domain 3 (Human-System Collaboration) for the delegation-with-revocation pattern — the partnership wind-down was announced in November 2023, but the consequences for students who enrolled under the prior arrangement persist. Pair with Case 57 (GAO regulator-side audit), Case 55 (Engler pricing optimization), and Case 5 (Epic Sepsis governance gap). coi-light render under the title is binding.



| DOMAIN                    | FAILURE MODES | LENS COMPETENCY                         | PROBLEM TYPE | INDUCED | LEO   |
|---------------------------|---------------|-----------------------------------------|--------------|---------|-------|
| HIGHER EDUCATION          | D K N         | D1 D2 D3 D4 D5                          | 6            | 5.4     | 5 · 3 |
| ONLINE PROGRAM MANAGEMENT | →             | → Navigating Sociotechnical Constraints |              |         |       |
| PROFESSIONAL LICENSURE    |               |                                         |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Treat the student-facing surface — recruiting communications, advising emails, clinical-placement coordination — as part of the disclosable program substance, not as operational discretion to be delegated under brand cover.
- ▶ Require that delegated personnel identify their actual employer in student-facing communications; the usc.edu email cover the complaint names is operationally addressable by branding policy.
- ▶ Build a paired disclosure that ties the delivered student experience back to the marketed program description; the gap the complaint alleges is between what was promised and what was delivered, and the gap is reportable.

### IN OPERATION / AFTER

- ▶ Carry the journalism-tier framing into print without softening; the litigation is ongoing, the targeting reconstruction is allegation-tier, and the case's pedagogical power rests on naming the evidence tier honestly.
- ▶ Pair the case in syllabi with Case 57 so the regulator-side and consumer-side halves of the OPM regime are taught together; one half names what the audit missed, the other names what the affected students alleged happened.
- ▶ Track the licensure-board half over time; an independent study of clinical-placement quality and post-graduation practice capacity is the audit the case names as the missing evidence.

## REFLECTION QUESTIONS

1. Identify a program in your domain where the student-facing surface — recruiting, advising, placement — has been delegated to a contracted vendor operating under institutional brand. What is the disclosure your institution makes to applicants about that delegation, and at what point in the transaction?
2. Specify the paired disclosure you would build to tie delivered student experience back to the marketed program description. What data — placement outcomes, advising-load ratios, vendor-employee proportion of student-facing communications — would you commit to publishing alongside enrollment marketing?
3. The journalism-tier flag is binding because the targeting reconstruction is allegation-tier. What is the investigation-grade study that would convert the case from allegation-tier to audit-tier — and who could commission it?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards



## CASE 60

## K-12 EDUCATION

## TEACHER EVALUATION

## ALGORITHMIC ACCOUNTABILITY

2011–2017

# Houston EVAAS — Value-Added Teacher Evaluation on Trial

 Governance & Trust

**IMPACT** Houston ISD used the proprietary SAS EVAAS value-added model to attribute student test-score growth to individual teachers and made termination decisions on the scores during the 2011–15 school years; SAS held the model as a trade secret, so teachers could not inspect, replicate, or contest their scores; in May 2017 the federal district court in Hous. Fed'n of Teachers v. Hous. Indep. Sch. Dist., 251 F. Supp. 3d 1168 (S.D. Tex.), denied summary judgment on the procedural due process claim alone — granting it on all others — holding teachers had a protectable property interest that unverifiable scores could deprive them of without due process; HISD settled in October 2017, agreeing not to use unverifiable value-added scores as a basis for termination and paying \$237,000 in fees

**IN BRIEF**

From 2011 the Houston Independent School District used the SAS Education Value-Added Assessment System (EVAAS), a proprietary statistical model, to attribute student test-score growth to individual teachers, and made termination decisions on the resulting scores. SAS held the model as a trade secret, so teachers could not inspect, replicate, or contest the number that ended their careers. In *Houston Federation of Teachers v. Houston ISD*, the federal district court denied summary judgment in May 2017, holding that teachers had a protectable property interest and that unverifiable scores could violate Fourteenth Amendment procedural due process — a “black box” evaluation denies any meaningful opportunity to contest. HISD settled in October 2017, agreeing to stop using unverifiable value-added scores for termination and paying \$237,000 in fees. The research record — year-to-year score instability, the ASA 2014 and AERA 2015 statements, and documented subgroup-composition bias — makes this the casebook’s cleanest adjudicated gap-attribution failure.

**BACKGROUND**

Beginning in 2011, the Houston Independent School District tied high-stakes teacher-appraisal decisions — including bonuses and terminations — to the SAS Education Value-Added Assessment System (EVAAS), a proprietary value-added model (VAM) that estimates an individual teacher’s contribution to student growth on standardized tests. The operational theory was direct: student test-score growth, statistically adjusted, isolates the teacher’s effect, and the isolated effect is decision-grade evidence for employment action.

The statistical profession disagreed on both counts. The American Statistical Association's 2014 statement cautioned that VAMs measure correlation rather than causation, are sensitive to model and data choices, and that most VAM studies find teachers account for about 1% to 14% of the variability in test scores — with the majority of improvement opportunities in system-level conditions. The AERA's 2015 statement enumerated eight technical requirements for valid VAM use and cautioned against giving VAM dispositive weight in high-stakes evaluation. HISD's deployment preceded both statements and continued after them.<sup>0</sup>

#### WHAT HAPPENED

Teachers rated ineffective by EVAAS could be, and were, exited on the strength of the score, but no teacher could verify it. SAS classified the model's source code and implementation as a trade secret and refused to release them to the district or its employees; the score arrived as a number without a derivation. The independent research record gave teachers concrete reasons to want one. Amrein-Beardsley and Collins's 2012 study reconstructed the records of four teachers whose contracts were not renewed in 2011 and found that for three of the four, whether the score showed added or detracted value from one year to the next was roughly the flip of a coin — and teachers across the district reported that their principals were skewing observational ratings to match the EVAAS number. A score that is unstable year-to-year, generated by a model no one outside the vendor can examine, was the sole or decisive input to termination decisions.<sup>1</sup>

#### THE INVESTIGATION

The Houston Federation of Teachers and seven teachers sued in 2014. In May 2017, the federal district court (S.D. Tex., Magistrate Judge Stephen Wm. Smith) denied HISD summary judgment on the procedural due process claim, holding that teachers had a constitutionally protectable property interest in continued employment and that the EVAAS scores' unverifiability could deny them due process: "HISD teachers have no meaningful way to ensure correct calculation of their EVAAS scores, and as a result are unfairly subject to mistaken deprivation of constitutionally protected property interests in their jobs." The opinion noted that the score "might be erroneously calculated for any number of reasons, ranging from data entry mistakes to glitches in the computer code itself," and that algorithms are human creations, subject to error like any other human endeavor. In October 2017 HISD settled: the district agreed not to use unverifiable value-added scores as a basis for termination so long as they remain unverifiable, created a joint instructional consultation panel on the appraisal process, and paid \$237,000 in legal fees.<sup>2</sup>

#### THE CAPABILITY GAP

The capability gap sits at the attribution step. Student test-score growth is a system outcome — shaped by curriculum, student assignment patterns, school effects, peer composition, and prior schooling — and EVAAS attributed it to an individual operator, the teacher, through a statistical coupling that was never validated for the decision it was driving. The ASA's 1%-to-14% figure is the quantitative form of the gap: the model assigned to the individual a residual dominated by the system. The evidence then failed the decision-grade test so completely that a federal court said so — not that the model was inaccurate, which no one could establish either way, but that a score no affected party can inspect, replicate, or contest cannot support a consequential deprivation. The failure is adjudicated gap misattribution: the institution located a system-level shortfall in the individual because

the individual was the addressable unit, and defended the attribution with secrecy rather than validation.<sup>3</sup>

#### AFTERMATH & REFORM

The bias record is documented, not asserted, and its hedges do real work. Amrein-Beardsley and Geiger's 2020 analysis of more than 1,700 HISD teachers' EVAAS results found that teachers in schools serving larger populations of English-language learners, free-or-reduced-lunch students, and special-education students received systematically lower value-added scores — a school-composition correlation, which is evidence of bias against teachers serving those populations but not a teacher-level causal decomposition. Concerns about teachers of highly mobile student populations appear in the Houston survey research and practitioner record but are less precisely quantified in the published literature. The settlement did not invalidate value-added modeling; it established that a proprietary, unverifiable score cannot be the basis for termination, and the due-process holding has since anchored the legal literature on algorithmic accountability in public employment. Houston reads, in retrospect, as the leading edge of a broader retreat: after the federal Every Student Succeeds Act (2015) removed the test-based-evaluation mandate that had driven adoption, high-stakes value-added teacher evaluation was rolled back across most of the roughly forty states that had embraced it — the national policy following the direction the ASA's 2014 warning had pointed. Pair with Case 51 (Atlanta Public Schools) for the adjacent measurement-architecture failure and Case 50 (Wisconsin DEWS) for the education-prediction equity thread.<sup>4</sup>

*When a public agency adopts a policy of making high stakes employment decisions based on secret algorithms incompatible with minimum due process, the proper remedy is to overturn the policy.*

HOUS. FED'N OF TEACHERS V. HOUS. INDEP. SCH. DIST., 251 F. SUPP. 3D 1168 (S.D. TEX. 2017)

#### REFERENCES

1. American Statistical Association (2014), ASA Statement on Using Value-Added Models for Educational Assessment, April 8, 2014.
2. Amrein-Beardsley, A., & Collins, C. (2012), "The SAS Education Value-Added Assessment System (SAS® EVAAS®) in the Houston Independent School District (HISD): Intended and Unintended Consequences," Education Policy Analysis Archives 20(12), doi:10.14507/epaa.v20n12.2012.
3. Education Week (October 2017), "Houston District Settles Lawsuit With Teachers' Union Over Value-Added Scores"; Houston Public Media (October 10, 2017), "Federal Lawsuit Settled Between Houston's Teacher Union and HISD" — settlement terms and the \$237,000 fee figure.
4. Hous. Fed'n of Teachers v. Hous. Indep. Sch. Dist., 251 F. Supp. 3d 1168 (S.D. Tex. 2017) — memorandum opinion denying summary judgment on the procedural due process claim.
5. Amrein-Beardsley, A., & Geiger, T. (2020), "Methodological Concerns About the Education Value-Added Assessment System (EVAAS): Validity, Reliability, and Bias," SAGE Open 10(2), doi:10.1177/2158244020922224.
6. American Educational Research Association Council (2015), "AERA Statement on Use of Value-Added Models (VAM) for the Evaluation of Educators and Educator Preparation Programs," Educational Researcher 44(8):448–452, doi:10.3102/0013189X15618385.

#### THE LEARNING ENGINEERING LENS

##### LE INSIGHT

Houston EVAAS is the casebook's cleanest adjudicated gap-attribution failure: a system outcome — student growth, shaped by curriculum, assignment patterns, and school effects — was attributed to an individual operator through a coupling never validated for the termination decisions it drove, and the evidence failed the decision-grade test so completely that a federal court said so. The fairness finding — lower scores for

teachers serving larger ELL, free-or-reduced-lunch, and special-education populations — is documented, not asserted, and travels as a school-composition correlation.

LENS APPROACH

Houston EVAAS is the unvalidated-attribution-under-secrecy case (induced 2.1; LENS D4/PT4; LEO-4). LENS uses it in Domain 4 (Test and Evaluation) as the anchor for the decision-grade-evidence standard: the score driving the consequential decision was unverifiable by design, unstable year-to-year, and assigned to the individual a variance residual the ASA’s own statement located overwhelmingly in the system. The court’s due-process holding gives the standard its adjudicated form — evidence no affected party can inspect, replicate, or contest cannot support a consequential deprivation. Pair with Case 51 (Atlanta Public Schools) for the measurement-architecture thread and Case 50 (Wisconsin DEWS) for the education-prediction equity thread.

| DOMAIN                     | FAILURE MODES | LENS COMPETENCY | PROBLEM TYPE        | INDUCED | LEO |
|----------------------------|---------------|-----------------|---------------------|---------|-----|
| K-12 EDUCATION             | G             | D1 D2 D3 D4 D5  | 4                   | 2.1     | 4   |
| TEACHER EVALUATION         | →             | →               | Test and Evaluation |         |     |
| ALGORITHMIC ACCOUNTABILITY |               |                 |                     |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Validate the attribution before the deployment: a model that assigns a system outcome (student growth) to an individual operator (the teacher) must be validated for the specific consequential decision it will drive, not just for aggregate statistical fit — and the ASA’s system-dominated variance decomposition is the null hypothesis the validation must beat.
- ▶ Treat contestability as a design requirement for any consequential score: if the affected party cannot inspect, replicate, or meaningfully contest the number, the evidence is not decision-grade regardless of the model’s internal quality, and a trade-secret claim does not waive the requirement.
- ▶ Take the professional bodies’ statements as binding inputs: the ASA 2014 and AERA 2015 statements enumerated the technical conditions for VAM use in high-stakes evaluation, and a deployment that cannot meet them should not carry dispositive weight in employment action.

IN OPERATION / AFTER

- ▶ Use the adjudicated record as the teaching text: the court’s holding — that unverifiable algorithmic scores deny a meaningful opportunity to contest — is the clearest available statement of the decision-grade-evidence standard applied to an operating deployment.
- ▶ Audit the equity record at the school-composition level and preserve its hedges: the documented finding is that teachers serving larger ELL, free-or-reduced-lunch, and special-education populations scored systematically lower, and the finding travels as a composition correlation, not a causal decomposition.
- ▶ Carry the settlement’s structural remedy forward: the joint instructional consultation panel — affected-party representation inside the appraisal-design process — is the governance form that the original deployment lacked.

## REFLECTION QUESTIONS

1. Identify a metric in your domain that attributes a system-shaped outcome to an individual operator for a consequential decision. What validation evidence exists for the attribution step specifically — not for the model's aggregate fit — and what fraction of the outcome's variance does the individual plausibly control?
2. Specify the contestability requirement for a consequential score in your setting: what would an affected party need — data, code, derivation, an appeal surface — to meaningfully verify or contest the number, and does any trade-secret or proprietary claim currently block it?
3. The Houston court found the due-process violation without resolving whether EVAAS was accurate, because no one could establish that either way. Pick a proprietary evaluation system in your domain and ask: if its accuracy can be neither confirmed nor refuted by the people it judges, what decisions is it currently supporting, and what is the institution's documented answer for why that is acceptable?

**WHO BUILDS THIS** governance, ethics & measurement — plus domain experts and a learning engineer to integrate.  
**TOOLS** governance frameworks · bias & evidence audits

# Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap

**K** Knowledge & Institutional Memory   **T** Training Gap

**EVIDENCE TIER** journalism-tier — source confidence flagged; future validation ongoing.

**IMPACT** National Reading Panel (2000) and decades of peer-reviewed cognitive science established systematic phonics as essential to early reading instruction; classrooms across the United States instead deployed balanced-literacy approaches built on three-cueing, institutionalized in popular curricula (Fountas & Pinnell, Calkins Units of Study, Reading Recovery); Emily Hanford's APM Reports investigations — "Hard Words" (2018) and "Sold a Story" (October 2022) — traced the divergence to teacher preparation and curriculum publishing; by late 2024 roughly 40 states plus DC had passed evidence-based reading laws, Columbia Teachers College dissolved the Reading and Writing Project (September 2023), and New York City mandated replacement curricula (May 2023); student-outcome effects of the legislative wave are still emerging

## IN BRIEF

The peer-reviewed evidence base on early reading — decades of cognitive science consolidated by the congressionally mandated National Reading Panel report in 2000, and restated for a practitioner audience by Castles, Rastle, and Nation in 2018 — established systematic, explicit phonics instruction as an essential component of teaching children to read. Across the same decades, a large share of American classrooms deployed balanced-literacy approaches built on three-cueing (meaning, structure, visual — MSV), which directs early readers toward pictures, context, and syntax to identify words, institutionalized in the era's most widely adopted curricula and interventions: Fountas & Pinnell, Lucy Calkins's Units of Study, and Reading Recovery. Emily Hanford's APM Reports investigations — "Hard Words" (2018) and the podcast series "Sold a Story" (launched October 20, 2022) — assembled the reconstruction of how the evidence failed to couple to the operator layer: teacher-preparation programs that did not teach the research, curricula that institutionalized contradicted methods, and an institution that had forgotten that the reading wars had been substantially settled before. The reporting was followed by a wave of state legislation — roughly 40 states plus the District of Columbia with evidence-based reading laws or policies by late 2024 — and by curriculum withdrawals and revisions. The evidence-tier flag is binding: the NRP findings are peer-reviewed, but the causal narrative of why practice diverged is journalistic synthesis, and the student-outcome effects of the legislative wave are still emerging.

## BACKGROUND

The evidence base predates the failure by decades. Jeanne Chall's 1967 synthesis had already concluded that code-emphasis instruction outperformed meaning-emphasis approaches for beginning readers, and the accumulating cognitive science of the 1980s and 1990s converged on the same architecture: skilled reading runs through letter-sound decoding, and beginners need systematic, explicit instruction in the code. The National Reading Panel, convened by Congress and reporting in 2000, reviewed the experimental literature and found systematic phonics instruction significantly more effective than unsystematic or no-phonics alternatives for early readers. Castles, Rastle, and Nation's 2018 review — titled, pointedly, "Ending the Reading Wars" — restated the consensus for a practitioner and policymaker audience. The scientific question of how children learn to read words was, by the field's own account, not open during the years the divergent practice scaled.<sup>0</sup>

## WHAT HAPPENED

What scaled in classrooms instead was balanced literacy built on three-cueing: the instructional theory that early readers identify words by coordinating meaning, structure, and visual cues — pictures, context, syntax, first letters — rather than by decoding letter by letter. The approach descended from Marie Clay's Reading Recovery and Ken Goodman's "psycholinguistic guessing game," and it was institutionalized in the most commercially successful curricula of the era: Fountas & Pinnell's leveled-literacy systems and Lucy Calkins's Units of Study, the latter used before the pandemic by roughly half of New York City elementary schools that reported a curriculum. Hanford's reporting — "Hard Words" in 2018, then the "Sold a Story" podcast series launched October 20, 2022 — documented the cueing practice in classrooms, traced its lineage, and put to its publishers and authors the question the evidence base had already answered. The series' reconstruction is journalistic synthesis, and the case carries it as such; but the classroom practice it documented was not in serious dispute.<sup>1</sup>

## THE INVESTIGATION

The analytical spine of the case is the uncoupled operator layer. The National Reading Panel report was a federal consensus document; it did not, by itself, change what teacher-preparation programs taught. Hanford's reporting, and the National Council on Teacher Quality's program reviews alongside it, found that large numbers of preparation programs did not teach the settled science — graduating teachers who had never encountered the evidence their profession had commissioned. The curricula then institutionalized the contradicted method at the point of use, so that a teacher's daily materials argued against the research even where the teacher had met it. This is implementation science's core question — does the practice hold when it moves into real settings? — inverted: the practice never moved at all. And it is a knowledge-loss failure as much as a training-gap failure: the reading wars had been substantially settled before, in 1967 and again in 2000, and the institution forgot its own evidence each time the settlement failed to reach the operator pipeline.<sup>2</sup>

## THE CAPABILITY GAP

The reporting was followed by institutional movement at a pace the evidence alone had not produced in two decades. By Education Week's tracking, roughly 40 states plus the District of Columbia had passed laws or adopted policies on evidence-based reading instruction by late 2024 — addressing curricula, teacher preparation, coaching, and intervention — with



about 15 states acting in 2024 alone. New York City announced in May 2023 that its elementary schools would be required to adopt one of three approved reading curricula, with the chancellor stating publicly that the balanced-literacy approach “has not worked.” Columbia Teachers College announced in September 2023 that it would dissolve the Teachers College Reading and Writing Project, the consultancy Calkins founded and directed. Calkins had removed cueing from the 2022 revision of Units of Study and continued her work through a successor organization; Heinemann and Fountas & Pinnell disputed that their materials constituted three-cueing while adding decodable texts to subsequent editions. None of the authors or publishers conceded the reporting’s central account.<sup>3</sup>

#### AFTERMATH & REFORM

The hedges the case carries are not decoration. The evidence-tier flag is binding: the NRP findings and the reading science are peer-reviewed, but the causal narrative of why practice diverged — the publishing economics, the teacher-preparation gaps, the role of specific authors — is journalistic synthesis, attributed to the reporting rather than to an adjudicated or systematic record. The measured student-outcome effects of the legislative wave are still emerging: Mississippi’s NAEP fourth-grade reading gains after its 2013 Literacy-Based Promotion Act are suggestive and widely cited, but the attribution is contested — the law bundled curriculum, coaching, retention, and screening, and the contribution of each is not isolated. Whether the 40-state wave produces population-scale reading gains is an open empirical question, and the case must be taught with that question open. What the case establishes without hedge is the structural finding: an evidence base can be settled, federally consolidated, and publicly available for decades while the operator layer — teacher preparation and daily curriculum — carries the contradicted practice at national scale.

*The scientific question of how children learn to read words was, by the field’s own account, settled while the divergent practice scaled; the evidence base failed to reach teacher preparation and daily curriculum for decades, and the causal narrative of why is journalistic synthesis — source confidence flagged, future validation ongoing.*

EDITORS’ SYNTHESIS OF HANFORD (APM REPORTS, 2018, 2022), THE NATIONAL READING PANEL REPORT (2000), AND CASTLES, RASTLE & NATION (2018).

#### REFERENCES

1. National Reading Panel (2000), Teaching Children to Read: An Evidence-Based Assessment of the Scientific Research Literature on Reading and Its Implications for Reading Instruction, National Institute of Child Health and Human Development, NIH Pub. No. 00-4769.
2. Hanford, E. (2022 – 2024), “Sold a Story: How Teaching Kids to Read Went So Wrong,” APM Reports, podcast series launched October 20, 2022.
3. Hanford, E. (2018), “Hard Words: Why Aren’t Our Kids Being Taught to Read?” APM Reports, September 10, 2018.
4. Schwartz, S. (updated through 2024), “Which States Have Passed ‘Science of Reading’ Laws? What’s in Them?” Education Week — legislation tracker; 40 states plus DC by late 2024.
5. Castles, A., Rastle, K., & Nation, K. (2018), “Ending the Reading Wars: Reading Acquisition From Novice to Expert,” *Psychological Science in the Public Interest* 19(1):5 – 51.
6. Seidenberg, M. (2017), *Language at the Speed of Sight: How We Read, Why So Many Can’t, and What Can Be Done About It*, Basic Books.

#### THE LEARNING ENGINEERING LENS

##### LE INSIGHT

Sold a Story is the clearest case of an evidence base that never moved into deployment: decades of peer-reviewed reading science, federally consolidated in 2000, coexisted at national scale with classroom



practice built on the contradicted method, because the operator layer — teacher preparation and curriculum — was never coupled to the evidence. The knowledge-loss reading travels with the training-gap reading: the settlement had been reached before and forgotten, and it took journalism, not the field’s own feedback channels, to surface the gap.

LENS APPROACH

Sold a Story is the research-practice-coupling case at national scale (induced 6.4; LENS D2/PT2; LEO-2). LENS uses it in Domain 2 (Iterative Development) for the discipline of treating evidence-to-deployment coupling as the deliverable — implementation science’s core question inverted, because the practice never moved at all — and for the operator-pipeline-as-deployment-surface framing. Pair with Case 50 (Wisconsin DEWS) and Case 51 (Atlanta Public Schools) for the education-at-scale evidence thread. The evidence-tier flag is binding: the reading science is peer-reviewed, the causal narrative is journalistic synthesis, and the outcome effects of the legislative wave are still emerging.

| DOMAIN                | FAILURE MODES | LENS COMPETENCY         | PROBLEM TYPE | INDUCED | LEO |
|-----------------------|---------------|-------------------------|--------------|---------|-----|
| K-12 EDUCATION        | K T           | D1 D2 D3 D4 D5          | 2            | 6.4     | 2   |
| LITERACY INSTRUCTION  | →             | → Iterative Development |              |         |     |
| RESEARCH-PRACTICE GAP |               |                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the operator pipeline as part of the deployment surface; an evidence consensus that does not reach teacher preparation and daily curriculum has not been deployed, and the coupling — not the consensus — is the deliverable to verify.
- ▶ Audit the installed base against the evidence base on a cycle; the reading wars had been substantially settled twice before the divergent practice scaled, and the institution had no mechanism that compared what its classrooms ran against what its research had concluded.
- ▶ Instrument the point of use, not the policy layer; the contradicted method survived in curricula and classroom routines long after the consensus documents shipped, and only observation at the operator layer — what the materials actually direct teachers to do — surfaced the gap.

IN OPERATION / AFTER

- ▶ Carry the evidence-tier discipline into teaching; the peer-reviewed reading science and the journalistic causal narrative are different confidence classes, and the case’s value depends on students learning to hold both without collapsing either.
- ▶ Track the legislative wave as an open experiment; roughly 40 states changed law faster than the evidence changed practice in twenty years, and the population-scale outcome data now accumulating is the test of whether mandate-led coupling works.
- ▶ Use the case as the anchor for research-practice coupling in any domain with an operator pipeline; the structural finding — settled evidence, uncoupled preparation layer, institutionalized contradicted practice — generalizes beyond literacy.

## REFLECTION QUESTIONS

1. Identify an evidence consensus in your domain that has been settled for a decade or more. Map the operator pipeline — the preparation programs, the daily materials, the point-of-use routines — and ask at which layer the consensus stops. What would the Hanford-style audit of your domain's classrooms find?
2. The reading wars were substantially settled in 1967 and again in 2000, and the settlement was lost each time. Specify the institutional memory mechanism your domain would need for a settled question to stay settled at the operator layer — and name who currently owns it, if anyone.
3. The corrective wave arrived through journalism and legislation rather than through the field's own channels, and its student-outcome effects are still emerging. What is the disciplined position between "the mandate fixed it" and "nothing is known" — and what evidence, on what timeline, would move you off it?

**WHO BUILDS THIS** knowledge management & data engineering · learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · scenario simulation · competency models

## CASE 62

K-12 EDUCATION

TEACHER EFFECTIVENESS REFORM

PROGRAM EVALUATION

2009–2018

# Gates Intensive Partnerships — Measurement Without a Coupled Lever

**G** Governance & Trust   **K** Knowledge & Institutional Memory

**IMPACT** ~\$575M Intensive Partnerships for Effective Teaching initiative (Gates Foundation ~\$212M, 37% of the total; the remainder state and federal funds) across Hillsborough County FL, Memphis/Shelby County TN, Pittsburgh PA, and four California charter management organizations, 2009 – 2016; RAND/AIR six-year evaluation (Stecher et al. 2018, RR-2242) found the initiative did not achieve its goals for student achievement or graduation and produced no evidence of improved outcomes for low-income minority students relative to comparison sites; Hillsborough drew down reserves and dismantled its system in 2015

## IN BRIEF

The Bill & Melinda Gates Foundation's Intensive Partnerships for Effective Teaching initiative (2009 – 2016) spent approximately \$575 million — roughly \$212 million from the foundation, 37 percent of the total, the remainder state and federal funds — across three districts (Hillsborough County FL, Memphis/Shelby County TN, Pittsburgh PA) and four California charter management organizations to redesign teacher evaluation, feedback, and personnel decisions around multi-measure effectiveness ratings: value-added scores, structured classroom observations, and student surveys, informed by the foundation's parallel Measures of Effective Teaching (MET) project (2009 – 2013). The sites built the evaluation architecture. RAND and AIR's six-year evaluation — Stecher et al., Improving Teaching Effectiveness: Final Report (2018, RR-2242) — found the initiative did not achieve its goals for student achievement or graduation, and found no evidence of improvement in low-income minority students' outcomes or access to effective teaching relative to comparison sites. Hillsborough, the flagship implementation, saw its main reserve fall from \$298 million in 2012 to \$152 million in 2015, absorbed costs well beyond the foundation's contribution, and dismantled the system in 2015. The measurement layer was real; the levers it was meant to drive — retention, placement, development — did not move the outcome, and the districts could not carry the system the theory of change required. The RAND evaluation itself is the exemplar of decision-grade evidence in the case; the intervention is the failure.

## BACKGROUND

In 2009 the Gates Foundation launched the Intensive Partnerships for Effective Teaching, selecting three large districts — Hillsborough County Public Schools in Florida, Memphis City Schools in Tennessee (merged into Shelby County Schools in 2013), and Pittsburgh

Public Schools — and four California charter management organizations (Alliance College-Ready, Aspire, Green Dot, and Partnerships to Uplift Communities). Each site agreed to redesign teacher evaluation around a multi-measure effectiveness rating combining growth in student achievement, structured classroom-observation rubrics, and student surveys, and then to couple the rating to the personnel system: recruitment, hiring, placement, tenure, dismissal, individualized professional development, career pathways, and performance-linked compensation. The theory of change was explicit — better measurement of teaching effectiveness would drive personnel and development decisions that would increase students' access to effective teaching, and that access would improve achievement and graduation, particularly for low-income minority students. Total spending across the sites reached approximately \$575 million, of which the foundation contributed roughly \$212 million; the remainder came from site and federal funds — site funds were the largest non-foundation source at four sites, federal funds at the other three.<sup>9</sup>

#### WHAT HAPPENED

The measurement layer was built with unusual rigor. The foundation's parallel Measures of Effective Teaching (MET) project (2009 – 2013) studied roughly three thousand teachers across six districts to validate the multi-measure composite, culminating in Kane, McCaffrey, Miller, and Staiger's January 2013 final report, *Have We Identified Effective Teachers?*, which used random assignment of approximately 800 teachers in grades four through eight to test whether composite effectiveness measures predicted student achievement under reassignment. The reported answer was broadly yes — the composite of value-added, observations, and student surveys carried predictive signal — with a decisive caveat the project itself documented: substantial noncompliance with the random assignment weakened the experimental protocol, and critics noted the validation spoke to prediction within the studied settings, not to what a district's personnel system would do with the score. A second limit ran directly into the initiative's own theory of change: students and teachers could not be randomly assigned across school sites, so the validation says nothing about comparing teachers' measured effectiveness between schools — which is precisely the comparison the plan to move effective teachers into high-need schools depended on. The Intensive Partnerships proceeded on the measurement evidence. What the MET project did not and could not establish was the second half of the theory of change: that a district could operate the rating at scale, couple it to retention, placement, and development decisions, and sustain the machinery long enough for the coupling to move student outcomes.<sup>1</sup>

#### THE INVESTIGATION

RAND and AIR's six-year evaluation — Stecher et al., *Improving Teaching Effectiveness: Final Report: The Intensive Partnerships for Effective Teaching Through 2015 – 2016* (2018, RR-2242) — is the evidence on the deployment. The evaluation found that the sites did implement the measurement architecture: new multi-measure evaluation systems were built and operated, and the effectiveness ratings were generally regarded as an improvement over the prior binary evaluations. But by 2014 – 2015, student achievement, dropout rates, and low-income minority students' access to effective teaching were not dramatically better than at similar sites that did not participate. The evaluation's site-level implementation record is where the decoupling shows: the personnel levers the ratings were meant to drive moved weakly or not at all. About 1 percent of teachers were dismissed for poor perfor-

mance in 2015 – 2016, in part because the new evaluation systems identified very few poor performers. No site consistently took the steps that would have moved low-income minority students into effective teachers' classrooms. Sites struggled to individualize professional development against the weaknesses an evaluation identified, and generally did not require participation, monitor it, or examine whether participants' effectiveness improved. Retention of effective teachers did not rise, though retention of ineffective teachers fell in most sites. RAND is explicit that the evaluation cannot identify the reasons the initiative missed its student-outcome goals — insufficient time for effects to appear, comparison sites adopting similar reforms, and a flawed theory of action all stay live — so the coupling account is the reading the implementation record best supports, not a causal finding the evaluation makes. The evaluation itself — pre-specified outcomes, comparison-site design, six years of implementation tracking, published in full at 500-plus pages — is the model of decision-grade evidence the initiative's own theory of change was never subjected to in advance.<sup>2</sup>

#### THE CAPABILITY GAP

Hillsborough County was the flagship implementation and is the sustainability record. The district's Empowering Effective Teachers program, launched in 2009 with a foundation pledge of \$100 million, built a corps of more than 260 peer evaluators and mentors and a performance-linked pay system. Tampa Bay Times reporting (Sokol, 2015) documented the fiscal arc: more than \$180 million went into the program over six years; the foundation paid about \$80 million against a pledge its agreement worded as "up to" \$100 million, leaving the district to absorb the difference; and the district's main reserve fell from \$298 million in 2012 to \$152 million in 2015, the years spanning the new pay system. In the fall of 2015 superintendent Jeff Eakins announced, in an email to the peer evaluators themselves, the dismantling of the Gates-funded system, and the district wound down the peer-evaluator corps and the evaluation machinery it could no longer afford. The sustainment requirement — what the district would have to carry after the grant, at what recurring cost, against what revenue — was never engineered as a first-class requirement of the design, and when the foundation's share ended and costs ran past projections, the system's political and fiscal support collapsed together.<sup>3</sup>

#### AFTERMATH & REFORM

The hedges the case carries are not decoration. The RAND evaluation is comparison-based, not randomized at the site level: it cannot rule out small effects, and implementation varied enough across sites that no single component is isolated as the point of failure. The MET validation evidenced the measurement, not the coupling, and its own random-assignment protocol was weakened by noncompliance. Hillsborough's dollar figures are drawn from contemporary journalism and district statements, not an audited public accounting. What the record does establish is structural: a rigorous, well-funded reform built a genuinely strong measurement layer and failed at the coupling between measurement and organizational capability — attribution (what the effectiveness rating actually measures and what decisions it can bear) and sustainability (what the district can carry) were both underengineered relative to the measurement itself. The case pairs with Case 51 (Atlanta Public Schools) for the inverse failure — measurement corrupted rather than uncoupled — and with Case 50 (Wisconsin DEWS) for the same prediction-without-evidenced-intervention structure in a student-facing system.

*The sites implemented new multi-measure teacher-evaluation systems, but the initiative did not achieve its goals for student achievement or graduation, and low-income minority students' outcomes and access to effective teaching showed no evidence of improvement relative to comparison sites; Hillsborough drew down reserves and dismantled its system in 2015.*

EDITORS' SYNTHESIS OF THE RAND/AIR FINAL EVALUATION (STECHER ET AL., 2018, RR-2242) AND TAMPA BAY TIMES REPORTING ON THE HILLSBOROUGH IMPLEMENTATION (SOKOL, 2015).

## REFERENCES

1. Stecher, B. M., Holtzman, D. J., Garet, M. S., Hamilton, L. S., et al. (2018), Improving Teaching Effectiveness: Final Report: The Intensive Partnerships for Effective Teaching Through 2015–2016, RAND Corporation, RR-2242-BMGF.
2. Kane, T. J., McCaffrey, D. F., Miller, T., & Staiger, D. O. (2013), Have We Identified Effective Teachers? Validating Measures of Effective Teaching Using Random Assignment, MET Project research paper, Bill & Melinda Gates Foundation, January 2013.
3. Sokol, M. (2015), "Sticker shock: How Hillsborough County's Gates grant became a budget buster," Tampa Bay Times, September 21, 2015 — with companion reporting, "Hillsborough schools to dismantle Gates-funded system that cost millions to develop" (fall 2015).
4. Stecher, B. M., Garet, M. S., Hamilton, L. S., Steiner, E. D., Robyn, A., Poirier, J., Holtzman, D. J., Fulbeck, E. S., Chambers, J., & Brodzia de los Reyes, I. (2016), Improving Teaching Effectiveness: Implementation: The Intensive Partnerships for Effective Teaching Through 2013–2014, RAND Corporation, RR-1295.
5. Will, M. (2018), "An Expensive Experiment: Gates Teacher-Effectiveness Program Shows No Gains for Students," Education Week, June 21, 2018.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The Intensive Partnerships are the reference case for a reform that engineered its measurement and did not engineer its coupling. A validated multi-measure effectiveness rating was built and deployed at scale, and the levers it was meant to drive — retention, placement, development — did not move the outcome, while the sustainment load was never specified as a requirement the district could be verified against. The intervention is the failure; the RAND evaluation that established it is the exemplar of decision-grade evidence.

### LENS APPROACH

The Intensive Partnerships case is the requirements-as-a-sustainment-deliverable anchor at philanthropic-reform scale (induced 1.4; LENS D1/PT1; LEO-1). LENS uses it in Domain 1 (Systems Analysis) for the discipline of engineering the full theory-of-change chain — measurement, coupling, capacity, sustainment — as explicit requirements rather than letting a strong measurement layer stand in for the system. Pair with Case 51 (Atlanta Public Schools) for measurement corrupted rather than uncoupled, and Case 50 (Wisconsin DEWS) for prediction-without-evidenced-intervention. The RAND/AIR final report is taught alongside the failure as the model of the evaluation the design phase never had.

| DOMAIN                       | FAILURE MODES | LENS COMPETENCY    | PROBLEM TYPE | INDUCED | LEO |
|------------------------------|---------------|--------------------|--------------|---------|-----|
| K-12 EDUCATION               | G K           | D1 D2 D3 D4 D5     | 1            | 1.4     | 1   |
| TEACHER EFFECTIVENESS REFORM | →             | → Systems Analysis |              |         |     |
| PROGRAM EVALUATION           |               |                    |              |         |     |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- Engineer the full theory of change as a requirements chain, not a measurement project; every link — rating

### IN OPERATION / AFTER

- Publish the null in full; the 500-plus-page RAND final report, commissioned by the funder whose initiative it

to personnel decision, personnel decision to access to effective teaching, access to outcome — is a stated requirement with its own evidence plan, and a strong first link does not discharge the rest.

- ▶ Treat sustainment as a first-class design requirement; specify what the district must carry after the grant, at what recurring cost, against what revenue and what political tolerance, before the architecture is built — Hillsborough’s collapse was fiscal and political before it was technical.
- ▶ Pre-commit the evaluation design and the decision rule; the RAND/AIR comparison-site evaluation with pre-specified student outcomes is the form the initiative’s own go/no-go and mid-course decisions required, and it arrived as a retrospective rather than a control.
- ▶ Attribute the failure to the coupling, not the measurement; the post-initiative lesson is not that multi-measure teacher evaluation is worthless but that measurement uncoupled from movable organizational levers cannot produce the outcome, and reforms that inherit the instruments must re-engineer the levers.
- ▶ Recover the sustainment accounting; reconstruct what each site actually carried — recurring cost, staffing, political capital — so the next reform’s requirements include the load the last one could not bear.

### REFLECTION QUESTIONS

1. Take a reform or system in your domain whose measurement layer is strong and write out the full theory-of-change chain from measurement to outcome. Which links are engineered requirements with evidence plans, and which are assumptions a strong first link is being allowed to stand in for?
2. Specify the sustainment requirement for a granted or pilot-funded capability you know: what must the institution carry after the external funds end, at what recurring cost, against what revenue and political tolerance? What would a Hillsborough-style drawdown look like in your setting, and when would you detect it?
3. The funder commissioned and published the evaluation that established its own initiative’s failure. What is the standing mechanism in your organization for buying decision-grade evidence against your own program — and what is the pre-committed decision rule if the finding is null?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines

# LAUSD iPad Program — The Common Core Technology Project

**D** Designed Out   **H** Human-System Interface   **G** Governance & Trust

**EVIDENCE TIER** journalism-tier — source confidence flagged; future validation ongoing.

**IMPACT** Los Angeles Unified's plan to put an iPad loaded with Pearson Common Core curriculum in every student's hands — publicly discussed at roughly \$1.3 billion (approximately \$500 million for devices and curriculum plus approximately \$800 million for network infrastructure, drawn from school-construction bonds) — collapsed within two years: students bypassed device security within about a week at initial schools, the district's own evaluator observed the Pearson curriculum in use in one of 245 classrooms, the Apple contract was halted in August 2014, Superintendent John Deasy resigned in October 2014, the FBI seized twenty boxes of documents in December 2014, the SEC examined the bond-funding disclosures, and the district recovered a settlement of \$6.45 million, paid by Pearson, in October 2015

## IN BRIEF

In June 2013 the Los Angeles Unified School District board approved the first phase of the Common Core Technology Project: an iPad, pre-loaded with Pearson's Common Core curriculum, for every student in the nation's second-largest district, publicly discussed at a full-build scale of roughly \$1.3 billion — approximately \$500 million for devices and curriculum and approximately \$800 million for the network infrastructure to support them, funded substantially from long-term school-construction bonds. The fall 2013 rollout at 47 initial schools outran every layer the deployment depended on: teachers were not prepared, the Pearson curriculum was incomplete, students at Roosevelt High and other campuses bypassed the device security filters within about a week, keyboards had not been budgeted, and home-use liability was unresolved. The procurement governance then unraveled in public: emails published by KPCC in August 2014 showed close contact between Superintendent John Deasy, his deputy, and Apple and Pearson executives predating the competitive bid; the district halted the Apple contract days later; Deasy resigned in October 2014; the FBI seized twenty boxes of documents in December 2014 at a federal grand jury's behest; and the SEC examined whether the bond funding was properly disclosed to investors. The district recovered a settlement of \$6.45 million in October 2015, paid by Pearson. The evidence base is largely journalism — KPCC, the Los Angeles Times, Education Week — plus the district's own commissioned evaluation by the American Institutes for Research, so the evidence-tier flag is rendered under the title; future validation will continue.



## BACKGROUND

The Common Core Technology Project was the signature initiative of Los Angeles Unified Superintendent John Deasy: a one-to-one tablet program intended to put an iPad, pre-loaded with Pearson's digital Common Core curriculum, in the hands of every student and teacher in a district of more than 600,000 students. The board approved the initial Apple contract in June 2013, and the project's full build-out was publicly discussed at roughly \$1.3 billion — approximately \$500 million for devices and curriculum and approximately \$800 million set aside for upgrading school networks — funded substantially from voter-approved school-construction bonds, an instrument whose multi-decade repayment horizon critics noted would long outlive the devices it bought. The stated theory was equity and readiness at once: every student would have the device the new Common Core online assessments assumed, and the curriculum would arrive with it.<sup>0</sup>

## WHAT HAPPENED

The fall 2013 rollout at 47 initial schools outran the operator layer the deployment's theory required. Within about a week of receiving the tablets, students at Roosevelt High, Westchester High, and the Valley Academy of Arts and Sciences deleted their device profiles and bypassed the security filters — 185 devices compromised across the three campuses, by the district's own evaluator's count — and the district suspended home use of the devices while it worked out what it had deployed. Keyboards, which the Common Core online assessments would require, had not been budgeted in the original purchase; home-use liability for a district-owned device in a student's backpack was unresolved; and teacher preparation lagged the hardware. The district's own commissioned evaluator, the American Institutes for Research, observed 245 classrooms across nineteen schools — fifteen CCTP and four comparison — in May 2014 and reported that the Pearson curriculum — the pedagogical core of the purchase — was in use in just one of them; school staff reported that the content promised during initial training was incomplete or unavailable during the school year.<sup>1</sup>

## THE INVESTIGATION

The governance layer failed on a separate and compounding track. In August 2014, KPCC published internal emails, obtained by public-records request, showing that Deasy and his deputy had been in close contact with Apple and Pearson executives about the project well before the competitive bid opened, and that the bid requirements tracked curriculum specifications Pearson had suggested in those private exchanges. Three days after publication, the district announced it was halting the Apple contract and reopening the procurement. Deasy resigned in October

2014. In December 2014, at the behest of a federal grand

jury, FBI agents removed twenty boxes of documents from district offices; the SEC separately examined whether the district's use of construction-bond proceeds for tablets was properly disclosed to bondholders. The district sought a refund from Apple and Pearson and in October 2015 reached a settlement of \$6.45 million, paid by Pearson as the curriculum subcontractor. The federal investigation was later reported closed without charges.<sup>2</sup>

## THE CAPABILITY GAP

The evidence-tier flag rendered under the case title is decisive. The public record on the CCTP is substantially journalism — KPCC's email investigation, sustained Los Angeles Times coverage, and Education Week's multi-year reporting and retrospectives — joined by the

district's own commissioned AIR evaluation and its board-committee oversight reports. No court adjudicated the procurement question; the federal inquiry closed without charges, so the pre-bid-contact finding rests on the published emails and the district's own responsive conduct rather than on an adjudicated record. The dollar figures are publicly reported approximations — “roughly \$1.3 billion” is the scale the district and press discussed, not an audited final cost, and the settlement figure is reported at approximately \$6.4 million. The structural pattern is consistent across all sources; the specific magnitudes carry the hedges. Future validation will continue.<sup>3</sup>

#### AFTERMATH & REFORM

What the case teaches is capability designed out at \$1-billion scale. The deployment's implicit theory was that device provision equals learning capability: buy the tablet, license the curriculum, and the classroom capability follows. Every layer that theory silently assumed — trained teachers, complete curriculum, working security, keyboards for the assessments the program existed to serve, a home-use policy — was unbuilt at rollout, and the district discovered each absence in production, at 47 schools simultaneously, rather than in an instrumented pilot honest about readiness. The governance layer was unbuilt in parallel: a procurement whose specifications tracked a vendor's private suggestions could not survive disclosure, and when it surfaced, the program lost its superintendent, its contract, and its public legitimacy in under three months. The two failures are one failure. A deployment that skips the iteration cycle on the human and curricular layers is the same deployment that skips the contestability of its own procurement: both substitute the founding vision for the evidence architecture the vision required.<sup>4</sup>

*The Pearson curriculum, touted as central to the effort to prepare students for the new Common Core standards, was in use in just one of the 245 classrooms the district's own evaluators observed.*

EDITORS' SYNTHESIS OF THE AMERICAN INSTITUTES FOR RESEARCH FINDING, EVALUATION OF THE COMMON CORE TECHNOLOGY PROJECT: INTERIM REPORT (SEPTEMBER 2014), SECTION 3.D, P. 36.

#### REFERENCES

1. Gilbertson, A. (2014), “LA schools cancel iPad contracts after KPCC publishes internal emails,” KPCC 89.3 (Southern California Public Radio), August 25, 2014 — the public-records email investigation and the contract halt; see also KPCC's “LAUSD iPads: Timeline of a troubled program.”
2. Margolin, J., Haynes, E., Heppen, J., Ruedel, K., Meakin, J., Hauser, A., Blum, J., Chavez, S., & Hubbard, A. (2014), Evaluation of the Common Core Technology Project: Interim Report, American Institutes for Research, September 2014 — 245 classroom observations across 19 schools (15 Phase 1 CCTP, four non-CCTP) visited in May 2014; the Pearson curriculum app observed in use in one, at a single elementary school. Followed by AIR's Evaluation of LAUSD's Instructional Technology Initiative: Year 2 Report (August 2015).
3. Herold, B. (2014), “Hard Lessons Learned in Ambitious L.A. iPad Initiative,” Education Week, September 10, 2014 — retrospective on the implementation failures; see also Herold (2014), “FBI Investigation Leaves L.A. iPad Initiative in Further Disarray,” Education Week, December 2014, and Herold (2015), “Pearson Settles With L.A. Over Curriculum Foul-Up,” Education Week, October 20, 2015 — the \$6.45 million settlement, paid by Pearson.
4. Blume, H., Los Angeles Times coverage of the CCTP (2013–2015) — the rollout, the security bypass, the contract suspension, the Deasy resignation, and the FBI document seizure.
5. Education Week Market Brief (2017), “Feds Drop Investigation Into Los Angeles District Over \$1 Billion iPad Purchase,” February 2017 — the reported closure of the federal inquiry without charges.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The LAUSD Common Core Technology Project is the \$1-billion-scale case for technology deployed without capability coupling: the deployment assumed device provision equals learning capability, and every layer that assumption depended on — teacher preparation, curriculum completeness, device security, keyboards, home-use policy — was unbuilt at rollout, while the procurement’s governance layer was unbuilt in parallel. The district’s own evaluator found the purchased curriculum in use in one of 245 observed classrooms; the published pre-bid emails cost the program its superintendent, its contract, and its legitimacy. Evidence base is journalism-tier plus the district’s own commissioned evaluation; the flag is rendered under the title.

LENS APPROACH

LAUSD CCTP is the ed-tech procurement-governance failure case at district scale (induced 5.1; LENS D5/PT6; LEO-5). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the stakeholder-trust-as-precondition anchor — a procurement that cannot survive disclosure forfeits the deployment’s legitimacy regardless of the technology’s merits — and as a Design-mode teaching case for the device-provision-equals-capability fallacy at the human-interface layer. The evidence-tier flag is binding: the AIR interim report has been confirmed against the primary document, but the Los Angeles Times pin-cites have not; the structural pattern survives the flag, the dollar magnitudes do not.

| DOMAIN                                                                             | FAILURE MODES                          | LENS COMPETENCY                                                                                                        | PROBLEM TYPE | INDUCED | LEO |
|------------------------------------------------------------------------------------|----------------------------------------|------------------------------------------------------------------------------------------------------------------------|--------------|---------|-----|
| <div>K-12 EDUCATION</div> <div>ED-TECH PROCUREMENT</div> <div>DEVICE ROLLOUT</div> | <div>D</div> <div>H</div> <div>G</div> | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> <div>Navigating Sociotechnical Constraints</div> | 6            | 5.1     | 5   |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pilot at capability scope, not device scope: an instrumented pilot honest about readiness would have surfaced the security bypass, the incomplete curriculum, the missing keyboards, and the teacher-preparation gap at one school’s cost instead of forty-seven — the pilot’s job is to falsify the device-provision-equals-capability theory before the bond money moves.
- ▶ Treat procurement contestability as a deliverable of the program, not an obstacle to it; a specification that tracks one vendor’s private suggestions is a governance breach whether or not it is ever adjudicated, and the disclosure record — not the intent — is what the deployment’s legitimacy will be judged on.
- ▶ Couple the financing instrument to the asset’s life: multi-decade construction bonds spent on three-to-four-year devices is a mismatch the district’s own investors, regulators, and public will eventually price, and the SEC’s disclosure inquiry is the form that pricing takes.

IN OPERATION / AFTER

- ▶ When the evidence base is journalism plus the institution’s own commissioned evaluation, flag it under the title and carry the standing “future validation ongoing” language into print; the structural pattern survives the flag, the dollar magnitudes do not.
- ▶ Preserve the district’s own evaluation as the decisive internal-evidence stream: the AIR finding that the purchased curriculum was in use in one of 245 observed classrooms is the deployment’s own instrument reporting the gap between the theory and the classroom.
- ▶ Teach the two failure tracks — the unbuilt operator layer and the unbuilt governance layer — as one failure with one root: the substitution of founding vision for evidence architecture, at the human layer and at the procurement layer simultaneously.

## REFLECTION QUESTIONS

1. The CCTP's implicit theory was that device provision equals learning capability. Identify a deployment in your domain whose budget line is the artifact and whose unbudgeted lines are the operator layer — training, workflow, policy, support. What would an instrumented pilot honest about readiness have to measure to falsify the provision-equals-capability theory before full-scale rollout?
2. The procurement failed on disclosure, not adjudication: the federal inquiry closed without charges, but the published pre-bid emails ended the contract, the superintendent's tenure, and the program's legitimacy anyway. What is the disclosure test you would apply to a procurement you are running — and what specification, if published tomorrow, would your program not survive?
3. The case carries journalism-tier evidence plus the institution's own commissioned evaluation. When the deployment's own instrument (the AIR observation of one classroom in 245) agrees with the external journalism, what does that convergence license you to conclude — and which of the case's specific magnitudes still cannot bear weight?

**WHO BUILDS THIS** systems & human-factors engineering · human factors & interaction design · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · usability testing · cognitive walkthroughs · governance frameworks · bias & evidence audits

Chapter 4

# Education, Training & the Learning Workforce — What Works — and the Frontier

*When instruction, analytics, and workforce design run the full evidence loop.*

---

*The effect survived because somebody kept measuring after the launch.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

# Training Transfer — The Gap Between Learning and Doing

**K** Knowledge & Institutional Memory    **N** Normalization of Deviance

**IMPACT** A meta-analysis of 89 empirical studies finds training transfer is positively related to cognitive ability, conscientiousness, motivation, and a supportive work environment, with no clear superiority of individual over situational predictors; the literature carries explicit hedges about inconsistent measurement and significant variability

## IN BRIEF

Blume, Ford, Baldwin, and Huang (Journal of Management 2010) synthesized 89 empirical studies on training transfer — the extent to which what is learned in training produces meaningful change in on-the-job performance. The headline finding is that transfer is positively related to four categories of variable: cognitive ability, conscientiousness, motivation, and the work environment. The authors report no clear superiority of individual over situational predictors; the work environment is the one category an organization can actively manage, which is what makes it the **system** layer's lever. The authors are explicit, and the decisive hedge survives into the case: the literature is characterized by "inconsistent measurement of transfer and significant variability in findings" — same-source, same-context measurement consistently inflated the relationships, and the pooled effect of transfer interventions themselves was only small to moderate. The case is the paired peer-reviewed half of the corporate-L&D pair with Kirkpatrick (Case 79): together they close the gap the v1 corpus had open at the workforce-L&D layer, and they motivate the LENS framing that the human is the biggest variable at the interface — here, the return-to-work interface.

## THE SHIFT

Training transfer is the structural question corporate L&D sits inside: does what learners did in training produce observable change on the job, and what predicts whether it does? The Kirkpatrick chain of evidence (Case 79) frames the question; the Blume et al. meta-analysis is the strongest single peer-reviewed answer the field has consolidated.<sup>o</sup>

## WHAT IS EMERGING

Blume, Ford, Baldwin, and Huang synthesized 89 empirical studies — 61 field studies and 28 lab experiments — and decomposed transfer predictors into four categories: cognitive ability, conscientiousness, motivation, and the work environment (particularly supervisor and peer support). All four categories were positively related to transfer. Cognitive ability carried the largest corrected correlation (.37) and the general work-environment construct a smaller one (.22, rising to .27 for transfer climate and .31 for supervisor support, the last

on small samples). The size ordering is not the whole story: all but two of the studies estimating the cognitive–ability relationship were laboratory studies, most with no interval at all between the end of training and the transfer measure, and in the two field studies of open-skill transfer the relationship ran small and negative. The authors’ own summary judgement is that there is “no clear superiority of individual variables over situational variables, or the reverse.” What distinguishes the work environment is not its size but its tractability: preselecting trainees on individual characteristics runs into logistic and political constraints, whereas situational variables can potentially be actively managed. It is the **system** variable in the set.<sup>1</sup>

#### THE CAPABILITY QUESTION

The hedge that matters survives. The authors are explicit that the literature is “characterized by inconsistent measurement of transfer and significant variability in findings,” and they make the measurement problem concrete: where transfer outcomes were obtained from the same source in the same measurement context as the predictor, the reported relationships were consistently inflated, so the meta-analysis reports only the estimates that exclude that bias. The second hedge is about the lever itself. Pooled across the studies that tried them, the effect of transfer interventions — goal setting, relapse prevention, optimistic program framing — was only small to moderate, which the authors say leads them to “a less optimistic conclusion” than Burke and Hutchins’s 2007 review had drawn. The meta-analysis is the strongest current synthesis, and it also names what the field has not yet measured well.<sup>2</sup>

#### EARLY EVIDENCE

The teaching point at the LENS-framing layer is precise. The capability deliverable is not the training event; it is the transfer. And the variable at the transfer layer the organization can actually move is the work environment — supervisor support, peer support, the opportunity to practice on actual tasks, the absence of countervailing pressures that punish trying the new behavior. The human is the biggest variable at the interface, and the interface is the return-to-work boundary. Training that ignores this boundary is training that delivers Level 2 evidence (Kirkpatrick, Case 79) and does not deliver Level 3 capability.<sup>3</sup>

#### OPEN PROBLEMS

In pair with Kirkpatrick, the case closes the corporate–L&D gap the v1 corpus had open and is a worked example of the revised “decision-grade evidence” framing in the v2 research backbone: the strongest synthesis the field has is a meta-analysis with explicit hedges about measurement inconsistency. The practitioner must decide on what to build — training design, supervisor support, practice opportunity — knowing the evidence is the strongest available but not closed. The subobjective **Judgment under inadequate evidence** (4.4) is exactly the capability this meta-analysis pattern asks practitioners to bring.

*The training event is not the capability deliverable. The transfer is. And the work environment is the variable at transfer an organization can actually manage.*

EDITORS’ SYNTHESIS OF BLUME ET AL. (2010).

#### REFERENCES

1.

Blume, Ford, Baldwin, & Huang (2010), "Transfer of Training: A Meta-Analytic Review," Journal of Management 36(4):1065–1105, doi:10.1177/0149206309352880.

2.

Baldwin & Ford (1988), "Transfer of Training: A Review and Directions for Future Research," Personnel Psychology 41(1):63–105 — the foundational synthesis the 2010 meta-analysis updates.

3.

Burke & Hutchins (2007), "Training Transfer: An Integrative Literature Review," Human Resource Development Review 6(3):263–296 — the integrative-review companion synthesis.

4.

Kirkpatrick & Kirkpatrick (2006), Evaluating Training Programs — the framework the meta-analysis informs (paired Case 79).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Blume et al. is the strongest current peer-reviewed synthesis on training transfer: cognitive ability, conscientiousness, motivation, and a supportive work environment all predict whether training produces on-job behavior change, with no clear superiority of individual over situational predictors. The literature carries explicit decisive hedges about inconsistent measurement and significant variability — including same-source measurement inflating the relationships, and transfer interventions themselves returning only small-to-moderate effects; the case is included with the hedges intact.

LENS APPROACH

Blume is the corporate-L&D transfer case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the design-iteration sub-competency (the work environment is the design variable) and in Domain 4 (Test and Evaluation) for the subobjective **Judgment under inadequate evidence** (4.4) — the meta-analysis is the strongest synthesis the field has, and it explicitly names what it cannot settle. Direct pair with Case 79 (Kirkpatrick); together they close the corporate-L&D gap the v1 corpus had open.

| DOMAIN                  | MODES AT STAKE | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|-------------------------|----------------|-----------------------|--------------|---------|-------|
| CORPORATE L&D           | K N            | D1 D2 D3 D4 D5        | 4            | 2.3     | 2 · 4 |
| TRAINING TRANSFER       | →              | Iterative Development |              |         |       |
| ORGANIZATIONAL BEHAVIOR |                |                       |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the training intervention with the work environment as a design variable — supervisor briefing, peer-support structures, practice opportunity — not only the classroom content.
- ▶ Specify the on-job behavior the training is supposed to enable and the conditions under which it will be observed; the meta-analysis identifies environment as the manageable predictor, so the environment must be instrumented.
- ▶ Carry the meta-analytic hedge into the deployment documentation: the literature is the strongest current synthesis, and it explicitly notes inconsistent measurement and significant variability.

IN OPERATION / AFTER

- ▶ Track the work-environment variables — supervisor support, peer-support structures, practice opportunity — separately from the training event, and report transfer outcomes conditional on environment.
- ▶ Use the meta-analytic finding to argue for the work-environment investment in the program design, not to declare the question settled; future replications and better-measured studies will move the magnitudes.
- ▶ When transfer is low, attribute the gap with the meta-analysis in hand: was the trainee under-prepared, under-motivated, or returning to an environment that did not support application? The remediation depends on the attribution.



## REFLECTION QUESTIONS

1. Identify a training program in your organization that produced strong Level 2 evidence (learning) but unclear Level 3 evidence (behavior change). What does the meta-analysis predict the variable was, and how would you re-design the program with the work environment as a first-class design variable?
2. Specify the work-environment instrumentation you would put in place for a new training deployment — supervisor briefing artifacts, peer-support structures, on-job practice opportunity — so transfer can be measured conditional on environment, not just on training fidelity.
3. The meta-analysis is the strongest current synthesis and carries explicit hedges about measurement inconsistency. What is the minimum additional evidence you would require before treating any specific magnitude from this literature as settled in your decision-making?

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

# Growth-Mindset National Experiment — Heterogeneous Effects

**D** Designed Out   **N** Normalization of Deviance   **K** Knowledge & Institutional Memory

**IMPACT** A nationally representative RCT of U.S. 9th-graders (65 schools,  $n = 12,490$ ) found a less-than-1-hour online growth-mindset intervention raised core-course GPA among lower-achieving students by 0.10 grade points ( $d = 0.11$ ) and, in an exploratory analysis, increased advanced-math enrollment; both pre-registered moderators bit — effects were smaller in higher-achieving schools, and the intervention changed grades when peer norms aligned with the messages of the intervention

## IN BRIEF

Growth-mindset interventions — short, scalable psychological interventions that teach students that intellectual ability is malleable rather than fixed — accumulated a substantial laboratory and small-school evidence base across the 2000s and 2010s. The open question, framed in the broader scalable-interventions literature, was whether the effects survived at population scale and what the structural moderators were. Yeager et al. (Nature, 2019) ran the test that became the field's reference point. A nationally representative RCT of U.S. 9th-graders received a less-than-1-hour online growth-mindset intervention; the trial documented improved grades among lower-achieving students — 0.10 core-course grade points, a standardized effect of 0.11 — and, in an analysis the authors label exploratory because it was not pre-registered, increased advanced-math enrollment relative to an active control. The headline finding for the case is the conditional: both pre-registered moderators bit. Effects were smaller in higher-achieving schools, and the intervention changed grades where the peer norm surrounding the student supported challenge-seeking. The study is notable for treating treatment-effect heterogeneity as the finding rather than as a nuisance. The intervention is not universally effective; naming where it does and does not work is the contribution. The trial's own conditional — “the intervention changed grades when peer norms aligned with the messages of the intervention” — survives verbatim into the case. Pair with Case 86 (Gándara) at the scalable-equity-intervention layer.

## THE SHIFT

Growth-mindset interventions teach students that intellectual ability is malleable — improvable through effort, strategy, and help-seeking — rather than fixed. The pedagogical claim, developed across two decades of research (Dweck and colleagues), is that students who hold the malleable view respond more constructively to academic challenge: they treat difficulty as informative rather than as a signal of fixed ability, and they persist on problems that the fixed-view student would interpret as evidence to stop. The laboratory and small-

school evidence base accumulated substantially across the 2000s and 2010s; the open question, by the late 2010s, was whether the effects survived at population scale and what the structural moderators were.<sup>9</sup>

#### WHAT IS EMERGING

Yeager et al. (Nature, 2019) ran the trial that became the field's reference point. The design was a nationally representative RCT of U.S. 9th-graders — drawn from a sample stratified to reflect the U.S. ninth-grade population — with the intervention delivered as two self-administered online sessions of roughly 25 minutes, about 20 days apart, against an active control that covered brain function without addressing beliefs about intelligence. Sixty-five schools and 12,490 students entered the analysis. Outcomes included grades, course-taking, and contextual measures of the classroom and peer environment. The trial pre-registered two moderators: the school's own achievement level, and the peer-norm environment the student returned to after the sessions.<sup>1</sup>

#### THE CAPABILITY QUESTION

The headline outcome was a positive finding with structure. The growth-mindset intervention raised core-course GPA among lower-achieving students by 0.10 grade points relative to the active control — a standardized effect of 0.11 — and increased enrollment in advanced mathematics, though that second analysis was not pre-registered and the authors report it as exploratory. Both pre-registered moderators bit. Effects were smaller in higher-achieving schools, which the authors read as those schools already having the resources to prevent failure. And the effect ran with peer norms: the intervention produced a greater GPA difference where the behavioural norm surrounding students supported challenge-seeking — where the peers treated the malleable view as legitimate and the help-seeking behavior the intervention encouraged as socially acceptable. Putting the two moderators together, the conditional treatment effect in low- and medium-achieving schools was 0.14 grade points in the third quartile of behavioural norms and 0.18 in the highest.<sup>2</sup>

#### EARLY EVIDENCE

What makes the case methodologically important is the authors' explicit treatment of treatment-effect heterogeneity as the finding rather than as a nuisance to be averaged away. The trial's structural answer is that the intervention is not universally effective, and naming where it does and does not work is the scientific contribution. The trial's conditional — grades changed when peer norms aligned with the intervention's messages — survives verbatim into the case because it is the hedge that matters: a headline-only reading ("growth mindset works at population scale") misses the substance, and an opposite-headline reading ("growth mindset doesn't work") misses it equally. The contribution is the conditional, and the design — a pre-registered moderator analysis with a nationally representative sample — is what makes the conditional defensible.<sup>3</sup>

#### OPEN PROBLEMS

That reading is now contested, and the dispute is itself part of the case. Macnamara and Burgoyne's 2023 Psychological Bulletin meta-analysis found growth-mindset effects near zero among the best-designed, least-biased studies, and large independent trials — the UK Education Endowment Foundation's 101-school RCT (2019) among them — returned no statistically significant effect on attainment. The critics' sharpest charge lands exactly on

this case's method: that "the effect is conditional on X" can be a genuine finding or an unfalsifiable escape hatch, and that moderator-hunting after a null average is how a non-effect is kept alive. The case is best taught with that fault line visible — the heterogeneity-as-finding stance is a real methodological advance and the live target of a serious replication critique, at once.<sup>4</sup>

Drafted alongside Case 86 (Gándara) at the scalable-equity-intervention layer, the case carries the design-predictions-to-trigger-support pattern (induced 8.3, LENS D2/PT5). The intervention targets the students who benefit (lower-achieving students in peer environments where the norm permits the change), rather than applying a blanket treatment that the average effect would credit and the heterogeneity would conceal. The equity-relevant design commitment is to surface the heterogeneity as part of the deliverable, not to launder it through an average-effect headline. The case is the methodologically clean model of how a population-scale RCT can earn the heterogeneity-as-finding stance, and the "conditional on peer norms" qualification is what makes the result useful for the next adaptation.

*The intervention changed grades when peer norms aligned with the messages of the intervention.*

YEAGER ET AL., NATURE 2019.

## REFERENCES

1. Yeager, D. S., Hanselman, P., Walton, G. M., Murray, J. S., Crosnoe, R., Muller, C., Tipton, E., Schneider, B., Hulleman, C. S., Hinojosa, C. P., Paunesku, D., Romero, C., Flint, K., Roberts, A., Trott, J., Iachan, R., Buontempo, J., Yang, S. M., Carvalho, C. M., Hahn, P. R., Gopalan, M., Mhatre, P., Ferguson, R., Duckworth, A. L., & Dweck, C. S. (2019). A national experiment reveals where a growth mindset improves achievement. *Nature*, 573(7774):364–369. doi:10.1038/s41586-019-1466-y — the case's primary trial.
2. National Study of Learning Mindsets, ICPSR 37353 — the trial dataset.
3. Dweck, C. S. (2006). *Mindset: The New Psychology of Success*. Random House — the broader theoretical framework the intervention rests on.
4. Sisk, V. F., Burgoyne, A. P., Sun, J., Butler, J. L., & Macnamara, B. N. (2018). To what extent and under which circumstances are growth mindsets important to academic achievement? Two meta-analyses. *Psychological Science*, 29(4):549–571 — the prior moderator-analysis literature the Yeager trial extends.
5. Macnamara, B. N., & Burgoyne, A. P. (2023). Do growth mindset interventions impact students' academic achievement? A systematic review and meta-analysis. *Psychological Bulletin* 149(3–4):133–173 — near-null effects among the best-designed studies.
6. Education Endowment Foundation (Foliano, F., Rolfe, H., Buzzeo, J., Runge, J., & Wilkinson, D., 2019). *Changing Mindsets: Effectiveness Trial — Evaluation Report — a 101-school, 5,018-pupil RCT with Year 6 pupils; point estimates of about two additional months' progress in English and maths, neither statistically significant.*

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The growth-mindset national experiment is the methodologically clean model of how a population-scale RCT can earn the heterogeneity-as-finding stance. The intervention improved grades among lower-achieving students by 0.10 grade points, and increased advanced-math enrollment in an exploratory analysis — and the effect was conditional on school achievement level and on peer norms. The qualifying language is the hedge that matters; headline-only readings in either direction miss the substance.

### LENS APPROACH

Yeager et al. 2019 is the designing-predictions-to-trigger-support case (induced 8.3; LENS D2/PT5) — Domain 2 for pre-registered moderator analysis; Domain 5 for the equity-relevant commitment to target support to those who benefit. Pair with Case 86 (Gándara).

| DOMAIN                 | MODES AT STAKE | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|------------------------|----------------|-----------------------|--------------|---------|-------|
| K-12 EDUCATION         | D N K          | D1 D2 D3 D4 D5        | 5            | 8.3     | 2 · 5 |
| SCALABLE INTERVENTIONS | →              | →                     |              |         |       |
| EQUITY                 |                | Iterative Development |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Pre-register the moderator analysis at the design stage; the heterogeneity-as-finding stance depends on the moderator being a planned analysis rather than a post-hoc inspection.
- ▶ Sample so that the moderator can be estimated — a nationally representative sample of U.S. 9th-graders, in this case — so the heterogeneity is observed across the population the headline claim would otherwise speak to.
- ▶ Treat the intervention as scalable on the headline finding and conditional on the moderator at the same time; designing predictions to trigger support means targeting the support to those who benefit, in the environments where the support can land.

### IN OPERATION / AFTER

- ▶ Carry the conditional language (“changed grades when peer norms aligned with the messages of the intervention”) verbatim into every downstream communication; a headline-only reading and an opposite-headline reading both miss the substance.
- ▶ Treat the treatment-effect heterogeneity as program documentation, not as a nuisance to be laundered through an average effect; the equity-relevant design commitment is to surface the heterogeneity, not to conceal it.
- ▶ Build the next adaptation against the moderator: where would the intervention work, what peer-norm context would it land in, and what is the institutional path to creating that context where it does not yet exist?

## REFLECTION QUESTIONS

1. Identify a scalable intervention in your domain whose published evidence runs at the average-effect level. What pre-registered moderator analysis — and what sampling design — would let you treat treatment-effect heterogeneity as the finding rather than as a nuisance?
2. The Yeager finding is that the intervention “changed grades when peer norms aligned with the messages of the intervention.” What is the analog conditional in your context — the structural moderator the intervention’s effect depends on — and how would you instrument the moderator at scale?

**WHO BUILDS THIS** systems & human-factors engineering · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · incident analysis · safety dashboards · knowledge bases · data pipelines

# Cognitive Tutor / Carnegie Learning

**T** Training Gap

**IMPACT** Randomized controlled trials showed learning gains; RAND RCT found a significant Algebra I gain in year two (0.21 SD, high school; no year-one effect, none significant in middle schools); the Cognitive Tutor curricula, Algebra I among five, were used by more than 500,000 students in approximately 2,600 school districts as of August 2008 — the Algebra-I-only count is not available

## IN BRIEF

Carnegie Learning's Cognitive Tutor, built from John Anderson's ACT-R cognitive architecture at Carnegie Mellon, carries one of the largest randomized evidence bases of any intelligent tutoring system in education. It uses Bayesian knowledge tracing to model each student's mastery and adapts instruction accordingly, and a RAND Corporation randomized trial found a statistically significant positive effect on Algebra I achievement — about 0.21 SD for high school students in the second year of use, with no significant effect in year one and a middle-school estimate that did not reach significance. The system is a learning-engineering success in the discipline's own terms — learning science to engineered software to randomized-trial evidence to deployment: as of August 2008 the Cognitive Tutor curricula — Algebra I among five — reached more than 500,000 students in approximately 2,600 school districts, though the What Works Clearinghouse notes the Algebra-I-only count is not available. Its limitations are instructive: it works best in well-defined domains like algebra and less well in ill-structured ones, making it the canonical evidence that the pipeline delivers for problems that fit it — leaving open whether the same discipline can deliver where problems do not.

## BACKGROUND

For decades, intelligent tutoring systems promised individualized instruction at scale, but few were grounded in a validated theory of how people learn or rigorously tested for effect. The opportunity was to build a tutor from a real cognitive model and prove its impact with the methods of experimental science. The field's recurring weakness was that promising systems rested on intuition about learning rather than a validated theory, and were rarely subjected to the kind of controlled trial that could separate genuine effect from novelty.<sup>o</sup>

## THE INTERVENTION

Carnegie Learning's Cognitive Tutor, developed from John Anderson's ACT-R cognitive architecture at Carnegie Mellon, models the specific skills underlying a subject and tracks each student's mastery using Bayesian knowledge tracing. It adapts problem selection to the individual learner and provides step-level feedback, embodying a full learning-science theory in software. Grounding the tutor in ACT-R rather than designer intuition is what made

the system testable: a theory that decomposes a subject into specific skills can be turned into a measurement of mastery and an instrument that responds to it.<sup>1</sup>

#### HOW IT WORKED

The tutor's effectiveness rests on the chain from theory to instrument: a decomposable skill model, a measurement method (knowledge tracing) that estimates mastery from student actions, and an instrumentable interface that adapts in response. Instruction is targeted where the model detects weakness, so practice concentrates on skills not yet mastered rather than on a fixed sequence. Each link in that chain has to hold for the next to work — the skill model makes mastery measurable, the measurement makes adaptation possible, and the adaptation is what concentrates practice where it pays off.<sup>2</sup>

#### THE EVIDENCE

The RAND Corporation's multi-site evaluation matched middle and high schools in seven states into similar pairs and randomly assigned each pair to keep its existing algebra curriculum for two years or adopt Cognitive Tutor Algebra I. It found a statistically significant effect for high schools in the second implementation year — 0.21 standard deviations, enough to move the median student about eight percentile points — with none in the first year and none reaching significance in middle schools. The case demonstrates the learning-engineering process working end to end: learning science, to engineered software, to randomized-controlled-trial evidence, to scaled implementation. A multi-site randomized evaluation is the strong form of the claim — it shows the effect survived contact with many real classrooms rather than a single favorable setting, which is what the field's earlier, untested systems had lacked.<sup>3</sup>

#### WHAT TRANSFERRED

The limitations are as instructive as the success. Cognitive Tutor performs best in well-defined domains like algebra and less well in ill-structured ones, making it the canonical evidence that the learning-engineering pipeline works for problems that fit the pipeline. The frontier question — whether the same discipline can deliver in operational, ill-structured domains where capability matters most — remains open. The dependence on a decomposable skill model is the boundary condition: where a subject cannot be cleanly broken into trackable skills, the very chain that made algebra tractable has nothing to attach to.<sup>4</sup> The program's later history sharpens the evidentiary caution. Carnegie Learning re-platformed the Cognitive Tutor software into MATHia (rewritten for the web as "MATHia X" in 2016), which is now its flagship, and the company still cites this same RAND result as MATHia's headline evidence — carrying a 2014 trial of the predecessor blended product forward onto a re-engineered one. In 2023 it went further, launching LiveHint AI, a generative-AI tutor (Case 88), on the authority of "25 years of data" and self-nominated awards but with no independent efficacy evidence of its own — a reminder that the RCT-grade warrant this case celebrates attaches to the specific product and era that earned it, and does not transfer by brand to its successors.<sup>5</sup>



*the LE process working as intended*

THEORY → ENGINEERING → EVIDENCE → SCALE

*Cognitive Tutors demonstrate the LE process working: theory → engineering → evidence → scale.*

EDITORS' SYNTHESIS OF ANDERSON ET AL. (1995) AND KOEDINGER & CORBETT (2006)

## REFERENCES

1. Anderson, J., Corbett, A., Koedinger, K. & Pelletier, R. (1995), "Cognitive Tutors: Lessons Learned," *Journal of the Learning Sciences — the ACT-R basis and design*.
2. Koedinger, K. & Corbett, A. (2006), *Cambridge Handbook of the Learning Sciences — knowledge tracing and adaptive instruction*.
3. Pane, J. F., Griffin, B. A., McCaffrey, D. F. & Karam, R. (2014), "Effectiveness of Cognitive Tutor Algebra I at Scale," *Educational Evaluation and Policy Analysis* 36(2):127–144, doi:10.3102/O162373713507480 — no effect in year one; a statistically significant year-two effect for high schools, not for middle schools.
4. What Works Clearinghouse (2009), *Cognitive Tutor® Algebra I intervention report*, U.S. Department of Education, July 2009 — "more than 500,000 students in approximately 2,600 urban, rural, and suburban school districts" as of August 2008.
5. Lynch, C., Ashley, K., Pinkwart, N. & Aleven, V. (2009), "Concepts, Structures, and Goals: Redefining Ill-Definedness," *International Journal of Artificial Intelligence in Education* 19(3):253–266 — what makes a domain ill-defined and why the tutor pipeline has less to attach to there.
6. Ritter, S. & Fancsali, S. (2016), "MATHia X: The Next Generation Cognitive Tutor," *Proceedings of the 9th International Conference on Educational Data Mining (EDM 2016)*, 624–625 — the web re-platforming of the Cognitive Tutor software into MATHia; and Carnegie Learning's 2023 announcement of Live-Hint AI (cf. Case 88).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Cognitive Tutor is the canonical evidence that the learning– engineering process exists, works, and produces measurable benefits at scale — when the problem is tractable. It is also the canonical case for what tractability looks like: well-defined domain, decomposable skill model, instrumentable interface, rigorous evaluation. The frontier for the discipline is whether the same pipeline can deliver in the operational, ill-structured domains where capability matters most.

### LENS APPROACH

LENS uses Cognitive Tutor in LEN 1 as the foundational LE-process exemplar, in LEN 4 as the canonical case for Bayesian knowledge tracing as a measurement instrument, and in LEN 9 as a technical case for model-based adaptive instruction.

| DOMAIN                | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|-----------------------|-----------------|------------------|--------------|---------|-----|
| EDUCATION AT SCALE →  | T               | → D1 D2 D3 D4 D5 | 5            | 2.3     | 2   |
| Iterative Development |                 |                  |              |         |     |

### APPROACHES TO CONSIDER



## DURING DEVELOPMENT

- ▶ Ground the intervention in a validated theory that decomposes the subject into specific skills, rather than building on designer intuition about how people learn.
- ▶ Build the full chain — a skill model, a measurement method that estimates mastery from learner actions, and an interface that adapts in response — so each link supports the next.
- ▶ Target instruction where the model detects weakness, concentrating practice on unmastered skills rather than marching through a fixed sequence.

## IN OPERATION / AFTER

- ▶ Prove the effect with a multi-site randomized evaluation, so the gain is shown to survive many real settings rather than one favorable classroom.
- ▶ Scale only where the problem fits the pipeline — well-defined, decomposable domains — and treat the decomposability of a skill model as the boundary condition for applicability.
- ▶ Keep the frontier question explicit when extending the method, testing whether the same theory-to-instrument chain can be built in less-structured domains before assuming it transfers.

## REFLECTION QUESTIONS

1. Cognitive Tutor works best in well-defined domains. Identify a problem in your domain that is currently ill-structured. What would have to be true to make the LE pipeline applicable?
2. The case is the strongest evidence the LE process works. What is the equivalent piece of evidence required to demonstrate the same pipeline in a non-cognitive domain — for example, surgical skill, or operational watchstanding?
3. Cognitive Tutor's chain runs from a decomposable skill model through measurement to adaptation. Take a capability in your domain and attempt the decomposition — where it resists being broken into trackable skills, what does that tell you about whether the pipeline can apply?

**WHO BUILDS THIS** learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models

# CIRCUIT — A Scalable, Equity-Centered Research Workforce Model

**T** Training Gap **K** Knowledge & Institutional Memory

**DISCLOSURE** Authorship: an editor of this volume is the senior author of the underlying study. Included on the published peer-reviewed evidence (ASEE 2023); editorial framing keeps critical distance from the program's self-presentation, and the open question about external/comparative evaluation is preserved in the text.

**IMPACT** An eight-pillar cohort model that in 2022 supported over 100 undergraduate, graduate, and ROTC students from 'trailblazing' backgrounds (first-generation, low-income, limited prior research access); peer-reviewed program description with longitudinal student outcomes across six cycles

## IN BRIEF

CIRCUIT is a research workforce-development program at Johns Hopkins APL built on eight explicit pillars — holistic recruiting, mission engagement, targeted training, leadership development, integrated assessment, diverse mentorship, academic partnerships, and career empowerment. Cervantes, Florzania, Sharp, Gray-Roncal, and Johnson ("Empowering Trailblazers toward Scalable, Systematized, Research-Based Workforce Development," ASEE Annual Conference 2023) presents the model and reports longitudinal student outcomes gathered over six program cycles. In 2022 the program supported over 100 undergraduate, graduate, and ROTC students, positioning CIRCUIT as a replicable model for STEM recruitment and retention of underrepresented students. The strongest honest framing — preserved in the case — is that this is a self-authored multi-cycle program evaluation at a single program; an external comparative evaluation would strengthen the causal claim, and the case says so rather than overstate. The COI render under the title (editor is the senior author) is binding. The case is the paired peer-reviewed companion to CIRCUIT proofreading (Case 78) — that case is about deploying capability against automation failure; this case is about building the capability in the first place at the edge of the trainees' prior preparation.

## BACKGROUND

The recurring story in STEM workforce-development at the undergraduate level is that the pipeline narrows at every stage, and that the narrowing falls disproportionately on students whose prior preparation has not included the kind of access — to research labs, to mentorship networks, to technical-skills training that meet a research project's immediate needs — that converts undergraduate interest into research capability. CIRCUIT was built to engineer that conversion at the edge of trainees' prior preparation, across a cohort drawn from "trail-

blazing” backgrounds: first-generation college students, low-income students, and students with limited prior research access.<sup>9</sup>

#### THE INTERVENTION

The program model is the case’s first contribution. Eight pillars are named and operationalized: holistic recruiting that does not screen on credentials a trailblazer’s background would not have generated; mission-driven research that lets trainees see why the technical skill they are acquiring matters; targeted technical training built around the project’s immediate needs; leadership development; integrated assessment, which runs learning and confidence through feedback loops per module rather than a summative pass/fail; diverse mentorship; university partnerships that route the cohort across institutions; and career empowerment that sustains the capability beyond the program. The model is published as a model in the ASEE paper — a peer-reviewed full paper with DOI, not an institutional press release.<sup>1</sup>

#### HOW IT WORKED

The longitudinal outcome evidence is the case’s second contribution. The ASEE paper reports outcomes gathered over the six program cycles from 2017 through 2023 — roughly 220 undergraduates, 15 graduate students and 23 high schoolers served through the 2022 cycle, with over 100 students supported in 2022 alone. The program presents itself as a replicable model, with the documentation, assessment instruments, and pillar-level operationalization that a replicating institution would need.<sup>2</sup>

#### THE EVIDENCE

The honest framing preserved in the case is the one the editorial discipline demands. This is a self-authored multi-cycle program evaluation at a single program. The ASEE paper clears the peer-review bar that the build list O2 requires; the model and the operational evidence are auditable. What an external comparative evaluation — by a researcher unaffiliated with the program — would add is the causal-claim half of the evidence: did CIRCUIT produce these outcomes, or did the cohort selection produce them? The case names the open question rather than answering it.<sup>3</sup>

#### WHAT TRANSFERRED

In pair with Case 78 (CIRCUIT proofreading + MICrONS), the case completes the CIRCUIT picture: building the capability (this case, peer-reviewed) and deploying it against automation failure (Case 78, frontier with evidence-tier flag). The pair also exercises the corpus’s COI discipline — both cases carry editor-related COI, both are rendered with the standing gold-bordered “Disclosure” block under the title, and both are anchored to the strongest peer-reviewed evidence available with the institutional and program-evaluation gap honestly named. The book’s credibility on these cases rests on plain disclosure, not on hiding the affiliation.

*The deliverable is the replicable model, not the program brand. The pillars are operationalized for independent evaluation.*

EDITORS’ SYNTHESIS OF CERVANTES ET AL. (2023).

#### REFERENCES

1.

National Academies / NSF research on undergraduate research experience effects on STEM persistence (broader literature against which the case sits).

2.

Cervantes, Floryanzia, Sharp, Gray-Roncal, & Johnson (2023), "Empowering Trailblazers toward Scalable, Systematized, Research-Based Workforce Development," ASEE Annual Conference, doi:10.18260/1-2-43271.

3.

CIRCUIT program documentation (JHU Hub, 2017 – present) – institutional program description.

4.

Paired case (78) – CIRCUIT proofreading + MICrONS – the deployed-capability companion.

THE LEARNING ENGINEERING LENS

LE INSIGHT

CIRCUIT is a peer-reviewed eight-pillar workforce-development program for STEM trainees from trailblazing backgrounds, with longitudinal outcomes over six program cycles. The strongest honest framing is self-authored multi-cycle program evaluation; an external comparative would add the causal half. COI under the title – editor is senior author – is binding.

LENS APPROACH

CIRCUIT workforce model is the equity-engineered workforce-development case (induced 1.2; LENS D2/PT4; with 8.3 alternate). LENS uses it in Domain 2 (Iterative Development) for the operationalized pillar model and the multi-cycle iteration evidence; in Domain 4 (Test and Evaluation) for the high-resolution assessment; and in Domain 5 (Navigating Sociotechnical Constraints) for the holistic-recruiting design choice that converts equity from rhetoric to enrolled cohort. Pair with Case 78 (CIRCUIT proofreading) – building capability vs. deploying it against automation failure. COI render is binding.



APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Operationalize the program pillars at the same level of detail an external replicator would need; the deliverable is not the program brand, it is the replicable model.
- Build holistic recruiting to actively de-weight credentials a trailblazer background would not have generated; the design choice converts the equity commitment from rhetoric into engineered enrollment.
- Pair high-resolution assessment with the targeted training so feedback can move at the cadence the work demands, not the semester boundary.

IN OPERATION / AFTER

- Commission external comparative evaluation alongside the self-evaluation; the causal-claim half of the evidence is what the program cannot produce on its own, and the case is honest about the gap.
- Publish the model documentation, instruments, and pillar operationalization openly so the model can be replicated and evaluated by other institutions on independent cohorts.
- Render the COI disclosure under the title – the standing language – and preserve the editorial critical distance from the program’s own self-presentation in any drafting.

## REFLECTION QUESTIONS

1. Identify a STEM workforce-development program in your domain. Which of CIRCUIT's eight pillars are operationalized at the level of detail an external replicator would need, and which are at the level of brand? Where would replication actually be testable?
2. Specify the external comparative evaluation you would commission alongside your own program evaluation. What independent cohort, what instrument, on what cadence — and who is unaffiliated enough to evaluate?
3. The COI render under this case's title is the standing language for editor-author cases. Identify a case in your domain where the program's evaluation is conducted by people with a stake in the program. What is the disclosure architecture that keeps the evaluation honest without hiding the affiliation?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines

# Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale

**T** Training Gap **D** Designed Out

**IMPACT** Settles & Meeder (ACL 2016) introduced Half-Life Regression (HLR), a trainable spaced-repetition model that learns per-item forgetting rates from large-scale learner data; HLR was deployed in Duolingo's production review-scheduling system after cutting prediction error against the Leitner scheduler by at least 45% on 12.9 million practice traces; in the six-week live trial against Leitner the daily-retention changes were not significant and practice sessions fell 7.3%

## IN BRIEF

Settles and Meeder (Association for Computational Linguistics, 2016) introduced Half-Life Regression (HLR), a trainable statistical model for spaced repetition in language learning. HLR combines the psychological theory of memory half-life (Ebbinghaus's forgetting curve and its descendants) with a learned regression that estimates each item's per-learner half-life from observed practice history and item features. The model was deployed in Duolingo's production review-scheduling system; the published evaluation compares HLR against heuristic schedulers (Leitner spacing, Pimsleur intervals) on Duolingo's predictive recall metric and on a daily (next-day) learner-retention metric. HLR cut prediction error against Leitner by at least 45%; on retention it did not beat Leitner. The variant that survived in production was not the headline model: after student complaints traced to overfitted lexeme-tag features, Duolingo shipped HLR without them. The case is one of the few published instances of a spaced-repetition algorithm being deployed and evaluated against meaningful behavioral outcomes at consumer scale (more than 150 million students had enrolled in a Duolingo course by the time of the study). The hedges that survive into the case verbatim: the evaluation is a single-vendor study, the "learning" outcome is measured by Duolingo's predictive metric and next-day retention rather than independent language-proficiency assessment, and the generalization to other content domains rests on the structural argument rather than a multi-vendor evidence base.

## BACKGROUND

The psychological theory of memory has carried, since Ebbinghaus's 1885 forgetting-curve work, the finding that review timed near the point of forgetting produces stronger long-term retention than either massed practice or uniform-interval review. The applied descendants — Leitner boxes (1972), Pimsleur graduated-interval recall (1967), SuperMemo's SM-2 algorithm (1987, and successors) — are all attempts to operationalize the theory as a schedule. Each uses heuristics: items advance through fixed intervals on success and reset on failure.

The heuristics work at the population level; they are coarse at the per-learner, per-item level where the actual half-life varies.<sup>9</sup>

#### THE INTERVENTION

Settles and Meeder's contribution was to treat the half-life as a quantity that can be learned from data. Half-Life Regression (HLR) models the probability that a learner recalls an item at a given delay as an exponential decay whose half-life is a regression on item features (language, part of speech, difficulty proxies) and per-learner practice history. The model is fit on aggregated Duolingo practice data — 12.9 million learning traces — and produces per-item half-life estimates that drive the scheduling: the estimated recall probability sets the skill-tree strength meters and ranks which words a practice session serves. The item features and the practice-history interaction features are separable, and which of the two the production model kept turned out to matter.<sup>1</sup>

#### HOW IT WORKED

The published evaluation compares HLR against several baselines on two outcomes. The first is Duolingo's predictive metric — how well the scheduler's model of recall matches observed recall on held-out practice data. The second is a downstream behavioral outcome: daily (next-day) learner retention, which Duolingo treats as a proxy for engagement-driven learning continuation. HLR cuts mean absolute error by at least 45% against Leitner, which itself kept the highest AUC of any algorithm tried; the six-week retention trial against Leitner returned no significant gain in activity or lessons and a 7.3% drop in practice sessions, which the authors read as a sign of better scheduling — students had been practicing “just to keep the tree gold” — and deployed on. That deployment did not hold. Within months students reported that particular words and skills decayed rapidly no matter how often they practiced, and the complaints traced to lexeme-tag features carrying strongly negative weights: feature-based overfitting the model's regularization had not prevented. A second two-week experiment on 3.3 million students pitted HLR stripped of lexeme features (HLR-lex) against full HLR and found all three retention metrics significantly higher for the stripped variant — any activity +12.0%, lessons +1.7%, practice +9.5%. HLR-lex is what Duolingo shipped to all students. The evaluation includes the production-deployment outcomes, not just offline validation — including the outcome that sent the first deployed model back.<sup>2</sup>

#### THE EVIDENCE

The case sits in the corpus because it is one of the few published, peer-reviewed deployments of a spaced-repetition algorithm at consumer scale with reported behavioral outcomes. Duolingo's scale at the time of publication — over 150 million enrolled students, and just under a million in the retention trial alone — meant the study had statistical resolution that the academic spaced-repetition literature rarely achieves. The contribution to LENS is the worked example of iterative design in the Domain 2 sense: a learning-engineering model deployed against a measurable behavioral outcome, with the outcome instrumented in the production system rather than inferred from offline validation.<sup>3</sup>

#### WHAT TRANSFERRED

The hedges have to survive verbatim. This is a single-vendor study published by Duolingo researchers about a Duolingo product; independent replication at comparable scale is not in the literature. The behavioral outcome is Duolingo's daily next-day retention, which is a

sensible proxy for engagement but is not a language-proficiency assessment — the case does not say HLR makes learners more proficient than Leitner does; it says HLR does better on the predictive metric, and that the retention trial against Leitner showed no significant gain. Generalization to other content domains (clinical knowledge, technical skill, vocabulary outside language-learning) rests on the structural argument that per-item half-life estimation should outperform heuristic scheduling, not on a multi-domain replication base. The case teaches the trainable- scheduling form, with the qualification that the strongest available evidence for the form is single-vendor.

*HLR does better on the predictive metric. Against Leitner it did not do better on next-day retention, and the case does not say it makes learners more proficient.*

EDITORS' SYNTHESIS OF SETTLES & MEEDER (2016).

REFERENCES

1. Ebbinghaus, H. (1885), Über das Gedächtnis — the empirical forgetting-curve foundation HLR formalizes.

2. Settles, B., & Meeder, B. (2016), "A Trainable Spaced Repetition Model for Language Learning," Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics, pp. 1848–1858, doi:10.18653/v1/P16-1174.

3. Pimsleur, P. (1967), "A memory schedule," Modern Language Journal 51(2):73–75 — graduated-interval recall heuristic.

4. Leitner, S. (1972), So lernt man lernen — Leitner-box spacing heuristic.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Half-Life Regression is one of the few peer-reviewed spaced-repetition deployments at consumer scale with reported behavioral outcomes. The structural argument — learn the half-life rather than fix it by heuristic — is strong; the evidence is single-vendor, the outcome is an engagement proxy not a proficiency assessment, and the cross-domain generalization rests on the structural argument rather than replication.

LENS APPROACH

Duolingo HLR is the trainable-scheduling consumer-scale case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the production-deployment evaluation form, and in Domain 4 (Test and Evaluation) for the engagement-proxy-vs-proficiency-assessment distinction. Pair with the spaced-education clinical RCTs (Case 17) for the cross-domain replication base, and with the machine-teaming AI-delegation typology where the delegation is to a scheduler optimizing on a proxy metric.

| DOMAIN             | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|--------------------|-----------------|-----------------------|--------------|---------|-------|
| EDUCATION AT SCALE | T D             | D1 D2 D3 D4 D5        | 4            | 2.3     | 2 · 4 |
| MACHINE LEARNING   |                 | Iterative Development |              |         |       |
| MEMORY AND SKILL   |                 |                       |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the scheduling parameter (item half-life) as a quantity to learn from data rather than to fix by heuristic. The cost of learning it is the data infrastructure the production system already has; the heuristic was a substitute for missing data, not a principled choice.

IN OPERATION / AFTER

- ▶ Carry the single-vendor hedge into communication about the case. HLR is the strongest available evidence for trainable scheduling at consumer scale; it is not multi-vendor evidence and it is not a language-proficiency study.



- ▶ Instrument the production system to record the practice events that the model needs (item, learner, delay, outcome) so the model can be fit and re-fit on observed behavior at consumer scale.
- ▶ Choose the deployment outcome metric carefully. Duolingo's choice — predictive recall + next-day retention — is defensible for an engagement-driven product; whether it is the right metric for a proficiency-driven application is a different decision.
- ▶ Distinguish the structural argument (per-item half-life estimation should outperform heuristic scheduling) from the domain transfer claim (HLR specifically generalizes to other learning content), which the published evidence does not yet support.
- ▶ Pair the case with spaced-education clinical RCTs (Case 17) when those are drafted, so the corpus has both a consumer-scale single-vendor deployment and a clinical-domain replication base for the same underlying mechanism.

### REFLECTION QUESTIONS

1. Identify a scheduling, dosing, or pacing parameter in your domain that is currently set by heuristic. What data does your production system already record that would let the parameter be learned per-unit rather than fixed?
2. Specify the outcome metric you would use for evaluating a learned scheduler. Duolingo chose engagement proxy + predictive recall; a proficiency-driven application would choose differently, and the choice is part of the evaluation design.
3. The single-vendor hedge is structural to the case. What would a multi-vendor replication base look like in your domain, and what minimum independent evidence would you require before treating a single-vendor production result as the basis for a curriculum-wide claim?

**WHO BUILDS THIS** learning science & instructional design · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · task analysis · design prototyping

## CASE 70

CORPORATE L&amp;D

PERFORMANCE CONSULTING

LEARNING TRANSFER

2001 – PRESENT

# High-Impact Learning System — Engineering the Environment, Not Just the Event

**K** Knowledge & Institutional Memory    **N** Normalization of Deviance

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** Brinkerhoff & Apking's HILS reframes corporate L&D as a system spanning pre-training (line-manager alignment, work-context preparation), the event itself, and post-training (supervisor support, on-the-job practice) — a deployable program model that anticipates by nine years Blume's meta-analytic finding that work-environment support is the transfer lever an organization can actually manage

## IN BRIEF

The High-Impact Learning System (HILS), introduced by Brinkerhoff and Apking in 2001, reframes L&D as a system spanning pre-training, the event itself, and post-training. The design principle, in Brinkerhoff's own words, is that "evaluation should focus on the entire training and performance improvement process, not solely on training events" — and that the improvement leverage sits in "the performance management system factors in the larger organization outside the boundaries of the training department or function." Blume's 2010 meta-analysis (Case 65) arrived nine years after HILS: it is the later evidence the model anticipated, not its source. HILS deployments include pre-training line-manager alignment and work-context preparation; the event itself; and post-training supervisor support and on-the-job practice opportunities. Corporate deployments report large gains in transfer, but the per-firm numbers live in vendor whitepapers and conference talks rather than peer-reviewed audits, and no baseline rate is established: Blume et al. find the transfer literature "characterized by inconsistent measurement of transfer and significant variability in findings." The case is the deployed-program counterpart to Case 83 (SCM as evaluation instrument): SCM measures whether the program worked at the tails; HILS designs the program so that the conditions for transfer are engineered. Evidence-tier flag is practice-synthesis: the model is documented in Brinkerhoff's own books and journal articles rather than in independent evaluations, the corporate effect sizes are self-reported, and future validation will continue.

## BACKGROUND

Blume et al.'s 2010 meta-analysis (Case 65) puts the work environment — transfer climate, supervisor and peer support, organizational constraints — at a corrected correlation with transfer of about .22, with supervisor support the strongest component at .31, and argues

its appeal is that situational variables "can potentially be actively managed" while trainee characteristics can only be selected for. The L&D practitioner question that follows is what to **deploy** in response to the finding: if the training event is not the capability deliverable, what is the surrounding architecture that has to be built so the transfer actually happens?<sup>0</sup>

#### THE INTERVENTION

Brinkerhoff and Apking's **High Impact Learning** (2001) is one of the answers in practitioner circulation. HILS reframes the L&D pipeline around three phases. **Before the event:** line-manager alignment, work-context preparation, learner readiness. **The event itself:** content delivery designed against the on-job task it is intended to enable. **After the event:** supervisor support, peer-support structures, deliberately engineered practice opportunities on the actual job tasks. The training event is one component, not the whole program.<sup>1</sup>

#### HOW IT WORKED

The corporate adoption story — across firms identified in Brinkerhoff's own books, articles and consulting materials — reports large gains in transfer, against no established baseline: Blume et al. find the transfer literature carries inconsistent measurement and wide variability in findings. The per-firm numbers live in vendor whitepapers and conference talks, not in peer-reviewed evaluations. The pattern is consistent across the practitioner sources; the magnitudes are self-reported and have not been independently audited.<sup>2</sup>

#### THE EVIDENCE

HILS is the deployed-program counterpart to Case 83 (SCM as evaluation instrument). Where SCM measures whether the program worked at the tails of the outcome distribution, HILS designs the program so that the surrounding conditions for transfer are engineered. The pair is the operational answer to the Kirkpatrick chain-of-evidence problem (Case 79): design the environment so transfer can happen, and measure the tails to confirm it did. Neither component on its own crosses the Level-2 / Level-3 seam — together they structure the cross.<sup>3</sup>

#### WHAT TRANSFERRED

The LENS teaching point is the framework-level claim that capability lives at the operator-system interface, not in the individual. HILS is a documented practice pattern that names the environment-as-intervention move; Blume's later meta-analytic finding (Case 65) is the evidence that caught up with it. The evidence-tier flag is practice-synthesis — the model and the deployment pattern are durable in practitioner literature, the per-firm effect magnitudes are self-reported, future validation will continue. The subobjective **Judgment under inadequate evidence** (4.4) is exactly the capability HILS asks the practitioner to bring: act on the strongest synthesis the field has produced while naming what is and is not independently audited.

*The work environment is the transfer variable an organization can actually manage. That is the argument for making it a design variable, and HILS makes it one.*

EDITORS' SYNTHESIS OF BRINKERHOFF & APKING (2001) AND THE BLUME META-ANALYTIC FINDING.

#### REFERENCES

1. Blume et al. (2010), *Journal of Management* 36(4):1065–1105 — the meta-analytic finding HILS operationalizes (paired Case 65).
2. Brinkerhoff, R. O., & Apking, A. M. (2001), *High Impact Learning: Strategies for Leveraging Performance and Business Results from Training Investments*, Perseus Publishing.
3. Brinkerhoff, R. O. (2006), "Increasing impact of training investments: an evaluation strategy for building organizational learning capability," *Industrial and Commercial Training* 38(6):302–307, doi:10.1108/00197850610685824 — the author's own statement that evaluation should address the whole performance-improvement process, not the training event.
4. Brinkerhoff, R. O. (2005), "The Success Case Method: A Strategic Evaluation Approach to Increasing the Value and Effect of Training," *Advances in Developing Human Resources* 7(1):86–101 — SCM as the paired evaluation instrument (Case 83).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

HILS is the deployed-program form of the environment-as-design-variable claim that Blume's 2010 meta-analysis (Case 65) later put at  $r \approx .22$ , and the design-side counterpart of SCM (Case 83). Evidence is practice-synthesis: the model is durable in practitioner literature, per-firm effect sizes are self-reported, future validation continues. The subobjective **Judgment under inadequate evidence** (4.4) is the capability the case asks for.

### LENS APPROACH

HILS is the L&D environment-as-design-variable case (induced 2.4; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development / Learning Engineering Design) for the amended sub-competency that narrates the design iteration explicitly across pre / event / post phases, and exercises the subobjective **Judgment under inadequate evidence** (4.4) because the practitioner must decide on practice-synthesis-tier evidence. Pairs with Case 83 (SCM) and Cases 65 and 79 as the corporate-L&D cluster.

| DOMAIN                 | MODES ADDRESSED | LENS COMPETENCY         | PROBLEM TYPE | INDUCED | LEO   |
|------------------------|-----------------|-------------------------|--------------|---------|-------|
| CORPORATE L&D          | K N             | D1 D2 D3 D4 D5          | 4            | 2.4     | 2 · 4 |
| PERFORMANCE CONSULTING | →               | → Iterative Development |              |         |       |
| LEARNING TRANSFER      |                 |                         |              |         |       |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Plan the pre-training phase as a first-class design deliverable: line-manager briefings, work-context preparation, learner readiness assessment, instrumented before the event runs.
- ▶ Design the post-training phase before delivery: supervisor support artifacts, peer-support structures, scheduled on-job practice opportunities on the specific tasks the training targets.
- ▶ Treat the training event as one component of a three-phase architecture, not as the whole program — Blume's finding (Case 65) is the decisive reason the architecture has to span the boundary.

#### IN OPERATION / AFTER

- ▶ Pair with Case 83 (SCM) to measure whether the deployed program worked at the tails — HILS designs the environment, SCM samples the outcomes; together they structure the chain-of-evidence cross (Case 79).
- ▶ Carry the practice-synthesis flag honestly: the model is documented and durable, the per-firm effect sizes are self-reported, and any specific magnitude cited from vendor whitepapers should be flagged as such in program documentation.
- ▶ Use the subobjective **Judgment under inadequate evidence** (4.4): the synthesis is the strongest the field has, and the practitioner has to decide whether to deploy HILS at organizational scale on practitioner-tier evidence while naming the validation that remains open.

## REFLECTION QUESTIONS

1. Identify a training deployment in your context that currently invests heavily in the event and lightly in the surrounding environment. What pre-training and post-training artifacts would you build to convert the deployment from a single-component program into a three-phase HILS-style architecture?
2. Specify the line-manager and supervisor briefing materials, the peer-support structures, and the on-job practice opportunities you would design — and identify who in your organization would be accountable for each. The accountability map is part of the design.
3. The HILS corporate effect sizes are self-reported. What independent evidence — third-party audit, peer-reviewed evaluation, cross-firm comparison — would you require before treating any specific magnitude as decision-grade, and how would you act on practitioner-tier evidence in the meantime?

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

# Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development

**K** Knowledge & Institutional Memory    **N** Normalization of Deviance

**IMPACT** A four-year longitudinal study of reflective-practice development on a work-based software-engineering degree, published at IEEE CSEE&T 2025 — one cohort's reflective reports (n = 36–38 per year) coded against two established models, Boud et al.'s and Bain et al.'s 5R; sophisticated forms of reflection rise markedly in the later years

## IN BRIEF

Barr, Nabi and Andrei's "The Development of Reflective Practice on a Work-Based Software Engineering Program: A Longitudinal Study" (IEEE CSEE&T 2025; preprint arXiv:2504.20956) reports a longitudinal study of how reflective-practice capability itself develops across a multi-year work-based software-engineering program. The signal the v2 corpus needs from this case is precise: it applies two established coding schemes to four successive years of one cohort's reflective reports, so the unit of analysis is the trajectory rather than a single point. That is the LENS-revised LEO-2's evaluation problem in miniature — if the program asks the learner to narrate and defend the design iteration in first person, the program also has to be able to evidence that the capability to do so is developing. The work cleared peer review, so what remains open is not review but reach: one cohort at one university, unreplicated. It is the v2 corpus's reference instance of established reflection coding applied longitudinally to evidence **development** rather than presence at a single point — the prior art the editor-commissioned first-person Practice Flywheel accounts will sit alongside.

## THE SHIFT

The amended LEO-2 in v2 asks the learner to narrate and defend the design iteration in first person. The evaluation problem this creates is real: most existing reflective-practice instruments measure whether a learner is reflecting at a given moment, not whether the capability to reflect well is developing across a program. A multi-year work-based software-engineering program is a useful site for this question because the work itself provides successive design iterations the learner can reflect on.<sup>9</sup>

## WHAT IS EMERGING

Barr, Nabi and Andrei report a longitudinal study of reflective–practice development across such a program. The methodological move worth naming is the focus on development rather than snapshot: two established coding schemes are applied to four successive years of assignments from one cohort — between 36 and 38 reports a year — and the study reports the trajectory. The authors are careful about what that buys — the study is not experimental, the reports were never standardized across years, and they make no claims about causation.<sup>1</sup>

## THE CAPABILITY QUESTION

The teaching point is the construct boundary. Measuring the **presence** of reflection is well-trodden ground (rubrics, codings, reflective–essay scoring schemes). Measuring the **development** of reflective capability is a different construct: the comparison runs across time, and this study builds it out of presence coding rather than around it — each stage of the 5R model is scored present, minimal, or absent in every assignment, and the trajectory is the change in that prevalence from Year 1 to Year 4.<sup>2</sup>

## EARLY EVIDENCE

Why the case sits in the v2 supplemental tier rather than in the verified primary list: the evidence is thin, not unreviewed. The arXiv version was published at IEEE CSEE&T 2025, and the case is included on the strength of that review, with the limit stated plainly in the prose rather than as a tier flag — one cohort, one university, no replication. The signal the corpus extracts is structural — the instrument–design move and the construct distinction — rather than the specific magnitudes the preprint reports.<sup>3</sup>

## OPEN PROBLEMS

What the case supplies the editor–commissioned first–person Practice Flywheel accounts (CIRCUIT, ERKS–class) is the evaluation pathway: evidence that reflective–practice capability can be measured as it develops, not only observed once. That is the prior art a Flywheel account’s evaluation design can build on. What is still outstanding is replication across other work–based engineering programs.

*The focus is the development of reflective practice over time. The study is not experimental, and the authors make no claims about causation.*

EDITORS’ SYNTHESIS OF BARR, NABI & ANDREI (2025).

## REFERENCES

1. v2\_research/01\_LENS\_revised\_competencies\_and\_CL0s.md — the proposed LENS revisions, including the amended LEO–2 (first–person narration of design iteration) whose evaluation problem this case addresses.
2. Barr, M., Nabi, S. W., & Andrei, O. (2025), “The Development of Reflective Practice on a Work–Based Software Engineering Program: A Longitudinal Study,” Proc. 37th IEEE Conference on Software Engineering Education and Training (CSEE&T), pp. 53–61, doi:10.1109/CSEET66350.2025.00012 (preprint: arXiv:2504.20956).
3. Boud, Keogh & Walker (eds.), Reflection: Turning Experience into Learning (1985) — reflection as a learning process, and the measurement problem it raises.
4. D. Schön, The Reflective Practitioner (1983) — the foundational account of reflection–in–action the genre rests on.

## THE LEARNING ENGINEERING LENS

## LE INSIGHT

Barr, Nabi & Andrei apply two established reflection coding schemes — Boud et al.'s and Bain et al.'s 5R, each stage scored present, minimal, or absent — across four years of one cohort ( $n = 36, 38, 38, 37$ ) on a work-based program, so the reported result is a trajectory rather than a snapshot. Peer-reviewed at CSEE&T 2025, but one cohort at one university and unreplicated; the case is included on the structural contribution (construct distinction, instrument-design move) rather than specific magnitudes. Replication is the outstanding step.

### LENS APPROACH

The longitudinal SE-program reflective-practice study is the evaluation-pathway case for first-person practice accounts (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the amended LEO-2 and in Domain 4 (Test and Evaluation) for the subobjective **Judgment under inadequate evidence** (4.4) — the study supplies the construct distinction (development vs. presence) from a single unreplicated cohort, so the distinction travels and the magnitudes do not.

| DOMAIN                               | MODES AT STAKE    | LENS COMPETENCY         | PROBLEM TYPE | INDUCED | LEO   |
|--------------------------------------|-------------------|-------------------------|--------------|---------|-------|
| SOFTWARE<br>ENGINEERING<br>EDUCATION | <b>K</b> <b>N</b> | D1 <b>D2</b> D3 D4 D5   | 4            | 2.3     | 2 · 4 |
| WORK-BASED<br>LEARNING               | →                 | → Iterative Development |              |         |       |
| REFLECTIVE PRACTICE                  |                   |                         |              |         |       |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ When designing a program-level evaluation of reflective-practice capability, distinguish the development construct from the presence construct at the instrument-design stage; the preprint is the worked example of why the distinction matters.
- ▶ Specify the temporal cadence of the instrument — when in the program reflections are collected, against what successive iterations of design work — so the development trajectory can be evidenced rather than inferred.
- ▶ Carry the single-cohort limit through any pedagogical decision the case informs; the portable contribution is the construct distinction, and the specific magnitudes await replication.

#### IN OPERATION / AFTER

- ▶ Track intra-learner depth change in reflective output across the program as the primary measurement target; cross-learner presence comparisons are a secondary construct and should be reported as such.
- ▶ Use the preprint as the prior-art reference for any first-person Flywheel evaluation pathway you propose for CIRCUIT, an ERKS-class effort, or a LENS-graduate program; the construct distinction is portable even where the specific instrument is not.
- ▶ When a replication on another work-based program appears, re-evaluate the magnitudes; the case is included on the structural contribution, not on this cohort's specific figures.

### REFLECTION QUESTIONS

1. Identify a multi-year program in your context where reflective-practice capability is supposed to develop. What instrument would you build to measure development (intra-learner depth change over time) rather than presence (snapshot)? At what temporal cadence?
2. The case is peer-reviewed but rests on one cohort at one university. What is the minimum additional evidence — replication across other work-based programs, comparison with snapshot-based instruments — you would require before treating any specific magnitude from this work as settled in your program design?
3. Specify the prior art you would assemble around a first-person Flywheel evaluation pathway for an engineering-practice account; which construct distinctions (development vs. presence) would you carry forward, and which would you supplement with locally produced evidence?



**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

## CASE 72

K-12 MATHEMATICS

HOMEWORK SUPPORT

FORMATIVE ASSESSMENT

2014 – PRESENT

# ASSISTments — National Replication and Long-Term Follow-Through

**T** Training Gap   **K** Knowledge & Institutional Memory   **D** Designed Out

**IMPACT** Cluster RCT across 43 schools and 2,850 students (Maine): 7th-graders assigned to ASSISTments outperformed controls on end-of-year math; largest gains for lower-performing students; the North Carolina replication (Feng et al. 2023) found larger gains for students of color ( $g = 0.14$ ) and Hispanic students ( $g = 0.13$ ) in a sample chosen for national demographic representativeness; the North Carolina replication found a sustained effect at the end of 8th grade, one year after the intervention ended (Hedges'  $g = 0.10$ )

## IN BRIEF

The Roschelle, Feng, Murphy, and Mason cluster RCT (AERA Open 2016), conducted across 43 schools and 2,850 students in Maine, found that 7th-graders assigned to ASSISTments outperformed controls on end-of-year mathematics, with the largest gains for lower-performing students. The North Carolina replication (Feng et al. 2023) reported larger gains for students of color and for Hispanic students — a finding the nearly all-white (93%) Maine trial could not have produced — and found the 7th-grade effect persisted into 8th-grade outcomes. A subsequent Arnold Ventures-funded extension tested a lower-cost virtual-training adaptation in predominantly rural areas, with an 8th-grade follow-through extension that Arnold Ventures and WestEd ended when the study design departed from the planned RCT. The case is one of the few EdTech tools in the corpus with replicated multi-state RCT evidence at meaningful effect sizes and with deliberate attention to the heterogeneity that matters most for equity outcomes. Pair with Case 17 (spaced education RCTs) for the replication-discipline thread. Open questions the authors carry: whether the virtual-training adaptation matches the in-person-training arm; whether the effect persists post-grade-8.

## BACKGROUND

ASSISTments is structurally different from the intelligent tutoring systems that dominate the K-12 EdTech evidence base. It augments homework rather than replacing curriculum; it does not require the institutional commitment to a new instructional system that Cognitive Tutor and its peers require; and the research team behind it (Heffernan and collaborators) has deliberately designed the platform around an evidence-generation loop with classroom

teachers. The cluster RCT Roschelle et al. published in 2016 is the first statewide evaluation of the platform, and it was designed to support claims a single-site trial could not support.<sup>9</sup>

#### THE INTERVENTION

The trial cluster-randomized 43 schools in Maine and assigned 2,850 7th-grade students to either an ASSISTments condition or a business-as-usual homework condition. The outcome instrument was end-of-year mathematics achievement. The headline finding is positive: ASSISTments-assigned students outperformed controls. The effect is meaningful in size, and the cluster randomization supports an inference at the school level rather than only at the student level. The teacher-side change required for the intervention was deliberately minimized — the platform is built around teacher-assigned homework problems, with the formative-assessment and feedback loops automated — so the trial estimates an effect achievable under realistic adoption conditions.<sup>1</sup>

#### HOW IT WORKED

The heterogeneity finding is what makes the case equity-relevant. Effect-size estimates were largest for lower-performing students. The North Carolina replication (Feng et al. 2023) reported that students of color benefited more than the average effect would suggest (Hedges'  $g = 0.14$ , and  $0.13$  for Hispanic students) — a finding the nearly all-white (93%) Maine sample could not have produced, and not an outcome the 2016 trial was powered to estimate. The pattern is the one Case 55 (Engler / enrollment algorithms) names as the inversion target: prediction and adaptive feedback used to trigger support rather than to gatekeep aid.<sup>2</sup>

#### THE EVIDENCE

The North Carolina replication (Feng et al. 2023) extended the evaluation into a longitudinal follow-through, finding the effect sustained at the end of 8th grade, one year after the intervention ended — a persistence finding that the EdTech evidence base does not consistently report, and that converts the case from a single-year effect-size study into a multi-year follow-through case. A subsequent Arnold Ventures-funded extension tested a lower-cost virtual-training adaptation in predominantly rural areas with an 8th-grade follow-through extension that Arnold Ventures and WestEd ended when the study design departed from the planned RCT. The replication structure — trial, replication, longitudinal follow-through, adaptation tested under different deployment conditions — is the closed-loop evidence architecture the corpus's EdTech cases mostly aspire to and rarely report.<sup>3</sup>

#### WHAT TRANSFERRED

The honest open questions the case carries are the ones the research team itself names. Whether the lower-cost virtual-training adaptation matches the in-person-training arm's effect size is still under analysis. Whether the effect persists past grade 8 is the longer-horizon question that the corpus's evaluation-horizon discipline (Case 84) directly applies to. Pair the case with Case 127 (Cognitive Tutor at-scale evaluation) for the evaluation-horizon thread, with Case 17 (spaced education RCTs) for the replication-discipline thread, and with Case 55 (Engler enrollment algorithms) for the prediction-triggers-support inversion — the equity-relevant heterogeneity finding here is the structural complement to Engler's gatekeeping critique.

*The largest gains are for lower-performing students. The North Carolina replication reports larger gains for students of color — a result the nearly all-white Maine sample could not have produced, and not what the 2016 trial was powered to estimate.*

EDITORS' SYNTHESIS OF ROSCHELLE ET AL. (2016) AND FENG, HUANG & COLLINS (2023).

## REFERENCES

1. Heffernan, N. T., & Heffernan, C. L. (2014), "The ASSISTments ecosystem," *International Journal of AI in Education* 24:470–497 — platform design and research-loop description.
2. Roschelle, J., Feng, M., Murphy, R. F., & Mason, C. A. (2016), "On-line Mathematics Homework Increases Student Achievement," *AERA Open* 2(4):1–12, doi:10.1177/2332858416673968.
3. Feng, M., Huang, C., & Collins, K. (2023), Technology-Based Support Shows Promising Long-Term Impact on Math Learning, WestEd — the North Carolina replication, the equity heterogeneity finding, and the 8th-grade persistence result.
4. Arnold Ventures grant documentation of the virtual-training-adaptation arm — including the 8th-grade extension that WestEd and Arnold Ventures ended when the design departed from the planned RCT.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

ASSISTments is the case in the corpus with the cleanest closed-loop evidence architecture for EdTech: cluster RCT, longitudinal follow-through into the next grade, and an adaptation arm under different deployment conditions. The equity-relevant heterogeneity finding is the exception and is treated as one: it surfaced in the North Carolina replication, was not pre-specified, and had no chance of surfacing in the nearly all-white Maine trial. The case grounds the closed-loop evaluation anchor in EdTech the same way Case 40 grounds it in team-science training.

### LENS APPROACH

ASSISTments is the closed-loop EdTech evaluation case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the teacher-side minimum-change design discipline and in Domain 4 (Test and Evaluation) for the heterogeneity-pre-specification and longitudinal follow-through structure. Pair with Case 84 (Cognitive Tutor at-scale evaluation horizon), Case 17 (spaced education RCTs), and Case 55 (Engler enrollment algorithms inversion — prediction-triggers-support).

| DOMAIN               | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|----------------------|-----------------|-----------------------|--------------|---------|-------|
| K-12 MATHEMATICS     | T K D           | D1 D2 D3 D4 D5        | 4            | 2.3     | 2 · 4 |
| HOMEWORK SUPPORT     | →               | Iterative Development |              |         |       |
| FORMATIVE ASSESSMENT |                 |                       |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Pre-specify equity-relevant heterogeneity outcomes before running the trial rather than mining them afterward; the ASSISTments minority-benefit finding surfaced only in a later moderator analysis and is fragile in a nearly all-white sample — a caution about trusting subgroup effects a trial was not powered to estimate.
- ▶ Design the teacher-side change to the minimum required for the intervention to operate; the case's external-validity strength depends on its having estimated an effect achievable under realistic adoption conditions.

### IN OPERATION / AFTER

- ▶ Publish the heterogeneity result with the aggregate result; the case's equity-relevant pedagogical value depends on the heterogeneity finding being on the same page as the average effect.
- ▶ Track the adaptation arm — the lower-cost virtual-training condition — as a separate replication; the closed-loop evidence architecture the case demonstrates includes adaptation-under-different-conditions as a distinct evidence layer.
- ▶ Carry the case in syllabi alongside Case 84 so the evaluation-horizon discipline and the heterogeneity-pre-specification discipline are taught together; the

- ▶ Build the longitudinal follow-through into the trial's data infrastructure from the start; the 7th-to-8th-grade persistence finding required data structures that single-year trials do not necessarily provide.

two methodological lessons are independent and both matter for EdTech-evaluation design.

### REFLECTION QUESTIONS

1. Identify an EdTech intervention in your domain whose equity-relevant heterogeneity outcome was not pre-specified in the trial design. What pre-specification would the next replication require, and what is the data infrastructure that would support it?
2. Specify the longitudinal-follow-through design you would build into the next at-scale EdTech evaluation. What grade-to-grade or year-to-year outcome would you track, and what data infrastructure does the tracking require?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

# The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension

**T** Training Gap   **K** Knowledge & Institutional Memory   **D** Designed Out

**IMPACT** Van Campenhout et al. (LAK 2023) replicated the doer-effect causal claim across seven courses with 15.2 million data points; L@S 2025 replication held with AI-generated practice questions (correlational, two sections of one course); LAK 2025 non-WEIRD radio/phone extension found weaker effect for learners with higher prior educational attainment — the heterogeneity finding

## IN BRIEF

The original “doer effect” causal claim — Koedinger and colleagues at LAK 2016 — held that students who interact with embedded practice activities learn more than students who only read, even after controlling for prior achievement and engagement, and that the effect appears causal. Van Campenhout et al.’s LAK 2023 paper replicated the claim across seven courses with 15.2 million data points; the L@S 2025 follow-up reported the effect held with AI-generated practice questions, though on two sections of one course and as a correlational rather than causal analysis; the Butler et al. LAK 2025 non-WEIRD extension tested the effect for learners receiving lecture content via community radio and practice via basic mobile phones, and reported that the doer-effect relationship was weaker for learners with higher prior educational attainment — the decisive heterogeneity finding the corpus most needs. The case sits with Case 69 (Duolingo half-life) and Case 17 (spaced education RCTs) as the replication-arc thread. The closed loop is closed not by a single trial but by replication; the effect generalizes but not uniformly, and the heterogeneity is itself the teachable result.

## BACKGROUND

The doer-effect causal claim is one of the cleanest published claims in the online-learning literature. Koedinger and colleagues, at LAK 2016, drew on student-level data from large online courses and isolated the effect of doing — interacting with embedded practice activities — from the effect of reading, controlling for prior achievement and engagement. The conclusion: doing causally improves learning more than reading does, by an effect size that has stood up across the field’s subsequent replications. The case Van Campenhout et al. open at LAK 2023 is what happens when the claim is treated not as a single result but as a replication target — what does the doer effect look like at scale, across courses, content domains, and delivery modalities?<sup>o</sup>

## THE INTERVENTION

The LAK 2023 replication is the broadest, by data volume. Van Campenhout et al. assembled data from seven courses and 15.2 million practice and reading interactions, and reported the doer-effect relationship held in the direction and approximate magnitude the original claim established. The 2025 L@S follow-up addressed a question that did not exist when the original claim was published: what happens to the doer effect when the practice questions are AI-generated rather than human-authored? The L@S 2025 replication reported the effect held with auto-generated content — a meaningful generalization given the rising deployment of LLM-generated practice across educational platforms, and a narrow one. It ran on two sections of a single Cognitive Psychology course and followed the established **correlational** doer-effect analysis; it does not repeat the causal test the seven-course LAK 2023 paper carries.<sup>1</sup>

## HOW IT WORKED

The 2025 LAK non-WEIRD extension by Butler and collaborators is the case's decisive methodological contribution. The extension tested the effect for learners receiving lecture content via community radio and practice via basic mobile phones — a delivery substrate structurally different from the platform-mediated online courses the original claim and the LAK 2023 replication operated on. The doer-effect relationship was weaker for learners with higher prior educational attainment. The heterogeneity is not noise the analysis should adjust away; it is the result. What generalizes across delivery modalities is the association, not the causal test: 234 Ugandan students, linear regression, a conceptual rather than direct replication. Within that limit the relationship is conditioned on prior attainment in a way the WEIRD-population evidence base did not surface.<sup>2</sup>

## THE EVIDENCE

The case is closed-loop in a sense the corpus most often aspires to. The loop is closed not by a single trial but by replication: the original claim, the seven-course large-N replication, the AI-generated-content replication, and the non-WEIRD-modality extension are four converging pieces of evidence — though only the first two carry a causal test, and only the first is large-N. The closed loop is also honest about what it closes and what it does not. Long-term retention across multi-year intervals is not yet in the replication record; transfer beyond the included content domains is not yet in the replication record; the non-WEIRD-modality heterogeneity is documented but not yet decomposed into the components — content familiarity, modality affordances, attention conditions — that would let designers act on it. The case carries those open questions rather than collapsing them.<sup>3</sup>

## WHAT TRANSFERRED

The case anchors with Case 69 (Duolingo half-life) and Case 17 (spaced education RCTs) as the replication-arc thread. All three demonstrate the closed-loop discipline at field scale: a design principle established as a single result, replicated across contexts, and surfaced as conditional on population and modality. The LEO on judgment under inadequate evidence is operative in a productive sense: the original claim was adequate evidence for the WEIRD-platform-online-course context, and the non-WEIRD-modality evidence the LAK 2025 extension adds is what extends the principle's actionable scope. The case completes the replication-arc thread the corpus needs to teach design-principle generalization honestly.

*The doer-effect association holds on radio and phone, but it is weaker for learners with more prior education – a condition the WEIRD-population evidence base did not surface. The heterogeneity is the result.*

EDITORS' SYNTHESIS OF VAN CAMPENHOUT ET AL. (2023, 2025) AND BUTLER ET AL. (2025).

REFERENCES

1. Van Campenhout, R., Jerome, B., Dittel, J. S., & Johnson, B. G. (2023), "The Doer Effect at Scale: Investigating Correlation and Causation Across Seven Courses," LAK23, pp. 357–365, doi:10.1145/3576050.3576103.

2. Van Campenhout, R., Autry, K. S., Clark, M. W., & Johnson, B. G. (2025), "Scaling the Doer Effect: A Replication Analysis Using AI-Generated Questions," L@S '25, pp. 24–34, doi:10.1145/3698205.3729545 — two sections of a Cognitive Psychology course, following the established correlational doer-effect analysis.

3. Butler, D. et al. (2025), "Does the Doer Effect Generalize To Non-WEIRD Populations? Toward Analytics in Radio and Phone-Based Learning," LAK '25, doi:10.1145/3706468.3706505 (also arXiv 2412.20923).

4. Koedinger, K. R., McLaughlin, E. A., Jia, J. Z., & Bier, N. L. (2016), "Is the doer effect a causal relationship? How can we tell and why it's important," LAK '16, pp. 388–397, doi:10.1145/2883851.2883957 — the replication target the present case builds on.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The doer-effect replication arc is the closed-loop-via- replication case in the corpus. Original claim, seven-course large-N replication, AI-generated-content replication, non- WEIRD-modality extension — four converging pieces of evidence with the prior-attainment heterogeneity finding as the load- bearing result. The case completes the replication-arc thread alongside Cases 69 and 17.

LENS APPROACH

Doer-effect replication arc is the closed-loop-by- replication case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the replication- arc discipline and in Domain 4 (Test and Evaluation) for the cross-population generalization-with-heterogeneity structure. Pair with Case 69 (Duolingo half-life) and Case 17 (spaced education RCTs) — the replication-arc thread teaches the closed-loop discipline at field scale.

| DOMAIN                   | Modes Addressed | LENS Competency       | Problem Type | Induced | LEO   |
|--------------------------|-----------------|-----------------------|--------------|---------|-------|
| LEARNING ANALYTICS       | T K D           | D1 D2 D3 D4 D5        | 4            | 2.3     | 2 · 4 |
| ONLINE LEARNING          | →               | Iterative Development |              |         |       |
| MOBILE/DISTANCE LEARNING |                 |                       |              |         |       |

APPROACHES TO CONSIDER

- DURING DEVELOPMENT

  - ▶ Treat the design principle as the replication target, not the original trial; the doer-effect case's structure is built on four converging pieces of evidence rather than on the original claim alone.
  - ▶ Pre-specify the cross-context replication conditions before launching the replication: course scope, content domain, delivery modality, and population characteristics are all conditions the original effect was estimated under and that the replication should vary.
  - ▶ Test design-principle generalization at the modality boundary before claiming generalization across
- IN OPERATION / AFTER

  - ▶ Publish the heterogeneity as the result, not the adjusted-away noise; the non-WEIRD prior-attainment finding is what extends the principle's actionable scope and is the decisive pedagogical content of the case.
  - ▶ Build the long-term-retention and cross-content-transfer studies the replication arc names as the next open questions; the closed loop the case completes is honest about what it does and does not close.
  - ▶ Carry the case in syllabi alongside Cases 69 and 17 so the replication-arc thread is taught as a thread, not as three independent examples; the closed-loop



modalities; the LAK 2025 non-WEIRD extension is the methodological model for that boundary test.

discipline at field scale is more visible across the three cases than within any one.

### REFLECTION QUESTIONS

1. Identify a design principle in your domain that has been replicated within WEIRD platform conditions but not across delivery modalities. What is the modality-boundary test you would design, and what would the analog of the radio-lecture-plus-mobile-phone substrate look like in your domain?
2. Specify the pre-specified heterogeneity outcomes you would build into the next replication of a design principle in your domain. What population characteristic do you expect would condition the effect size, and what would constitute disconfirming evidence?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

# Zhang/Scardamalia – Knowledge Building Across Three Cohorts

**T** Training Gap **K** Knowledge & Institutional Memory **D** Designed Out

**IMPACT** Three-year design study in a single Grade 4 classroom (Institute of Child Study, Toronto) using Knowledge Forum across an optics curriculum; documented progression from fixed small-group to opportunistic collaboration across cohorts, with associated gains in the depth and distribution of scientific explanations across the class community

## IN BRIEF

Zhang, Scardamalia, Reeve, and Messina's 2009 Journal of the Learning Sciences paper reports a three-year design study in a single Grade 4 classroom at the Institute of Child Study in Toronto, working through an optics curriculum supported by the Knowledge Forum platform. Across three successive cohorts, the classroom's collaborative structure progressed from a fixed small-group organization to an opportunistic one in which students convened around emerging questions and dispersed when those questions were addressed. Outcomes on the depth and distribution of scientific explanations across the class community improved in step with the structural progression. The hedges that travel with the case from the source are binding: single teacher, single school, design study not causal in the trial sense; transferability to non-ICS contexts and to teachers without knowledge-building expertise is open. The case is the classroom-scale longitudinal counterpart in the small-tier evidence base; the LE Lens uses it as the JLS- anchored small-tier complement to v1 Case 67.

## BACKGROUND

Knowledge Building is the pedagogical program Marlene Scardamalia and Carl Bereiter developed across the 1980s and 1990s, organized around the idea that classrooms can operate as knowledge-creating communities rather than as receivers of established content. The Knowledge Forum platform is the supporting infrastructure: a shared database in which students post conjectures, build on one another's notes, mark rise-above syntheses, and trace the evolution of community ideas. The 2009 Journal of the Learning Sciences paper is the systematic three-year study of how the collaborative structure inside a single Grade 4 classroom developed across cohorts as the teacher and students worked the program through.<sup>o</sup>

## THE INTERVENTION

The setting is deliberately constrained. A single Grade 4 classroom at the Institute of Child Study in Toronto, a laboratory school with a Knowledge Building tradition; a single teacher

across the three years; an optics curriculum delivered to three successive cohorts; Knowledge Forum as the supporting platform. The design study's interventions across cohorts were not randomized treatment arms but iterated redesigns of how collaboration was organized in the classroom — what kind of grouping the students worked in, how questions were surfaced and addressed, how the classroom community decided what counted as a productive direction. The design-based research method treats this iterative redesign as the unit of inquiry.<sup>1</sup>

#### HOW IT WORKED

The structural progression Zhang and colleagues document runs through three organizations, one per year: fixed small groups, then interacting small groups with substantial cross-group knowledge sharing, then opportunistic collaboration in small teams. In the first cohort, students worked in stable small groups that took on assigned subtopics within optics. By the third cohort, the classroom had moved to an opportunistic structure: students convened around emerging questions as those questions surfaced from the community's shared work in Knowledge Forum, and dispersed when the questions had been addressed. Zhang and colleagues read those organizations off a social network analysis of the Knowledge Forum record together with qualitative analysis of the notes themselves. The depth of scientific explanations in the community's shared notes improved in step with the structural change, and the distribution of explanatory contribution across students became less concentrated — more students were authoring substantive contributions, and the substantive contributions were reaching further into the underlying physics.<sup>2</sup>

#### THE EVIDENCE

The hedges that travel with the case from the source are binding and decisive. The study covers a single teacher, a single school, and a single curriculum domain; the design-based research method does not produce a causal estimate in the trial sense, and the authors are explicit that what they report is a developmental account of how the classroom's collaborative practice changed across cohorts, not a controlled comparison. Transferability to non-ICS contexts is open. Transferability to teachers without Knowledge Building expertise is open. The case is not the evidence that the Knowledge Building program works at scale across heterogeneous classrooms; it is the evidence that in a specific well-supported setting, a classroom's collaborative structure can be iteratively redesigned across cohorts in a direction that improves the depth and distribution of scientific explanation.<sup>3</sup>

#### WHAT TRANSFERRED

The case sits as the classroom-scale longitudinal counterpart in the small-tier evidence base. Pair with Case 40 (Colorado CTSA team science) for the collaboration-measurement thread at a different scale and domain; with Case 77 (Hybrid Human-AI Tutoring) for the small-tier deployment-success counterpart; with Case 81 and Case 85 (OU Analyse) for the distance-higher-education governance frame at a different population. The LE Lens uses the case as the JLS-anchored small-tier complement to v1 Case 67 (Cognitive Tutor's multi-site at-scale arc); the two cases together teach that the longitudinal classroom record is the substrate that learning-engineering iterations operate on, and that the iteration unit is the cohort, not the lesson.

*The cohort is the iteration unit. The collaborative structure of a single Grade 4 classroom moved from fixed small-group to opportunistic across three years, and the depth and distribution of scientific explanation moved with it.*

EDITORS' SYNTHESIS OF ZHANG, SCARDAMALIA, REEVE, & MESSINA (2009).

## REFERENCES

1. Zhang, J., Scardamalia, M., Reeve, R., & Messina, R. (2009), "Designs for Collective Cognitive Responsibility in Knowledge-Building Communities," *Journal of the Learning Sciences* 18(1):7–44, doi:10.1080/1058400802581676.
2. Scardamalia, M., & Bereiter, C. (2006), "Knowledge Building: Theory, pedagogy, and technology," in K. Sawyer (ed.), *Cambridge Handbook of the Learning Sciences — programmatic backdrop*.
3. Bereiter, C., & Scardamalia, M. (2014), "Knowledge Building and Knowledge Creation: One Concept, Two Hills to Climb," in S. C. Tan, H. J. So, & J. Yeo (eds.), *Knowledge Creation in Education — extension into the post-2009 program*.
4. Design-based research methodology references — Cobb, Confrey, diSessa, Lehrer, & Schauble (2003), "Design Experiments in Educational Research," *Educational Researcher* 32(1):9–13 — the methodological frame the 2009 study operates inside.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Zhang and colleagues' three-year design study is the classroom-scale longitudinal record of how a collaborative structure can be iteratively redesigned across cohorts in a direction that improves the depth and distribution of scientific explanation. The hedges are binding — single teacher, single school, design study not causal — and the case's pedagogical value depends on the hedges being preserved. The cohort is the iteration unit.

### LENS APPROACH

Zhang/Scardamalia is the cohort-as-iteration-unit case at classroom scale (induced 2.2; LENS D2/PT4; LEO-2 and LEO-4). LENS uses it in Domain 2 (Iterative Development) for the multi-cohort design-based-research discipline and in Domain 4 (Test and Evaluation) for the depth-and-distribution paired outcome measure. Pair with Case 40 (Colorado CTSA team science collaboration measurement), Case 77 (Hybrid Human-AI Tutoring small-tier deployment), and Cases 81 and 85 (OU Analyse — distance higher-education governance at a different scale). The LE Lens uses it as the small-tier complement to v1 Case 67 (Cognitive Tutor's multi-site at-scale arc).

| DOMAIN                 | MODES ADDRESSED | LENS COMPETENCY         | PROBLEM TYPE | INDUCED | LEO   |
|------------------------|-----------------|-------------------------|--------------|---------|-------|
| K-12 SCIENCE           | T K D           | D1 D2 D3 D4 D5          | 4            | 2.2     | 2 · 4 |
| COLLABORATIVE LEARNING | →               | → Iterative Development |              |         |       |
| DESIGN-BASED RESEARCH  |                 |                         |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Treat the cohort as the iteration unit when the classroom collaborative structure is the design target; the case demonstrates that a year is the right horizon for a meaningful redesign cycle and a multi-year horizon is required to surface the progression.
- ▶ Design the infrastructure — Knowledge Forum's shared note database — to make the community's evolving ideas inspectable; the redesign across cohorts

### IN OPERATION / AFTER

- ▶ Carry the binding hedges into print without softening; the case's value to the corpus rests on its specificity, and the transferability questions are open and disclosable as open.
- ▶ Pair with Case 40 (collaboration measurement at team-science scale) so the collaboration-as-design-target thread is taught at both the classroom and team scales.

- depended on the teacher and students having a record of how prior cohorts' ideas had moved.
- ▶ Track depth and distribution of contribution as paired outcome measures; reporting depth alone hides the equity-relevant distribution change, and reporting distribution alone hides the substantive-progress signal.
- ▶ Use the case as the JLS-anchored small-tier complement to v1 Case 67 (Cognitive Tutor's multi-site at-scale arc); the cohort-as-iteration-unit lesson is the bridge between the classroom-design literature and the at-scale evaluation literature.

### REFLECTION QUESTIONS

1. Identify a classroom or team in your domain whose collaborative structure has not been redesigned across cohorts. What would a cohort-as-iteration-unit redesign cycle look like, and what shared-infrastructure artifact would make the prior cohort's evolution inspectable for the next?
2. Specify the paired outcome measures you would track when collaborative structure is the design target. Depth and distribution travel together in the Zhang case; what is the analogous pair in your setting, and what would reporting one without the other obscure?
3. The case's transferability hedges are open — non-ICS contexts, non-Knowledge-Building-expert teachers. Pick a feature of the ICS setting that the case depends on, and ask: what would have to be true in your setting for the redesign discipline to travel, and what wouldn't?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

# Chen et al. — Rural China AI Devices and the Equity-Direction Finding

**T** Training Gap   **K** Knowledge & Institutional Memory   **D** Designed Out

**IMPACT** Two-part study in 12 Chinese schools (4 urban, 8 rural) — a 268-teacher acceptance survey in January–February 2024 and a three-month classroom experiment from September to November 2024; rural experimental classes gained 17.93% on mathematics and 13.46% on history, while urban experimental classes gained 10.96% on mathematics and 9.55% on history — the rural gain exceeded the urban gain across both subjects

## IN BRIEF

Chen, Wu, Chen, and Zhou's 2025 paper in *Frontiers in Psychology* reports a three-month quasi-experimental study of AI-integrated Mixed Reality devices across 12 schools — 4 urban and 8 rural — in China, run from September to November 2024 alongside a 268-teacher acceptance survey fielded that January. The headline result is the equity-direction finding: rural experimental classes gained 17.93% on mathematics and 13.46% on history, while urban experimental classes gained 10.96% on mathematics and 9.55% on history. The rural gain exceeded the urban gain across both subjects. The honest hedges that travel with the case from the source are decisive and preserved in the prose: the horizon is three months; the study covers 12 schools; assignment is non-randomized, with classes matched only on baseline attainment; and the authors name the small sample and the survey's unvalidated self-report basis. All four authors are at Chinese institutions — Hangzhou City University, Lingnan University, and Huazhong University of Science and Technology — and the case stands as the first-author-from-deployment-country evidence the corpus's non-WEIRD thread most needs.

## BACKGROUND

The corpus's non-WEIRD deployment record is structurally thin. The published learning-engineering evidence base is built largely on WEIRD-population studies — Western, educated, industrialized, rich, democratic — and the deployments that operate at meaningful scale outside that population are rarely written up in peer-reviewed journals by authors based in the deployment country. Chen and colleagues' 2025 *Frontiers in Psychology* paper is one of the strongest documented exceptions: a deployment of AI-supported instructional devices across rural and urban Chinese schools, evaluated quantitatively, reported by a four-author team based at Chinese institutions, and published in a peer-reviewed journal with the decisive methodological hedges named.<sup>o</sup>

## THE INTERVENTION

The deployment covers 12 schools — 4 urban and 8 rural — over the September to November 2024 window; a separate acceptance survey of 268 teachers was fielded the previous January. The intervention is AI-integrated Mixed Reality devices used in mathematics and history classroom instruction. The comparison structure is between experimental and standard classes within each school setting across both subjects; post-experiment scripts were anonymized and graded by independent evaluators blind to group assignment. The headline outcome the paper reports is the equity-direction finding: rural experimental classes gained 17.93% on mathematics and 13.46% on history, while urban experimental classes gained 10.96% on mathematics and 9.55% on history. The rural-over-urban gain pattern is the result the case teaches; the equity direction is unusual in the published EdTech literature, where deployments more often surface as widening the gap between higher- and lower-resourced settings rather than narrowing it.<sup>1</sup>

## HOW IT WORKED

The honest hedges the case carries are binding and explicit in the source. The horizon is three months — a September to November 2024 window — and the case's pattern is conditioned on that horizon, not on a school-year or multi-year follow-through. The sample is 12 schools, which is right for the within-subject within-school comparison structure the study runs but is not right for prevalence claims about Chinese rural schooling at population scale. Assignment to experimental and control classes is non-randomized: classes of comparable baseline attainment were selected, with pre-experiment scores differing by less than 5% within each school, and teacher- or school-level selection effects are not ruled out. The authors' own limitations name the small sample and the survey's self-report basis — teachers' stated intentions were never correlated with usage logs or observed teaching behavior — along with two the case has to carry because they bear directly on its equity reading: there is no long-term follow-up, and the analysis does not segment students by characteristic, so the rural-over-urban pattern is a setting-level result and not a claim about which students within a setting gained. Those hedges travel into the prose without softening; the case is published in a peer-reviewed mid-tier journal, and the evidence-tier flag rendering is intentionally not set — the hedges are carried in the case's own argument rather than in the title-bar render.<sup>2</sup>

## THE EVIDENCE

The case's structural placement in the corpus is the equity-direction-finding case at the non-WEIRD deployment seam. Pair with Case 86 (Gándara / student-success predictive equity in AERA Open) for the equity-direction thread at the higher-education scale; with Case 88 (LiveHint AI bias across foundation models) for the bias-surfacing thread in AI-supported instruction; with Case 73 (Doer Effect non-WEIRD radio-and-phone extension) for the non-WEIRD methodological discipline at the heterogeneity-finding axis. The first-author-from-deployment-country structure of Chen et al. is itself the methodological contribution the corpus's non-WEIRD thread needs: the strongest reading of a deployment is the one written by authors with access to the deployment context and published in the journal architecture the field can read.<sup>3</sup>

## WHAT TRANSFERRED

The pedagogical seam the case opens is the equity-direction reading of an AI-supported instructional intervention. The conventional reading — that deployments widen the gap between higher- and lower-resourced settings — is not what Chen and colleagues report on the three-month horizon they studied. The rural gain exceeded the urban gain across both subjects, and the result is the teaching point the case anchors. The case does not claim that the pattern generalizes past three months; it does not claim that the pattern holds outside the 12 schools the study covers; it does not claim causal identification past the bounds the non-randomized assignment supports. What it reports is the equity-direction finding in the deployment window the study describes, and the case uses that finding to anchor the LEO on fairness beyond omission — the finding is the rare published example where an AI deployment in education narrowed rather than widened a between-setting gap on its evaluation horizon, and the methodological structure of the study supports treating the result as the case's central observation rather than as a discounted outlier.

*Rural experimental classes gained 17.93 percent on mathematics and 13.46 percent on history; urban experimental classes gained 10.96 percent on mathematics and 9.55 percent on history. The rural gain exceeded the urban gain across both subjects — and the hedges that matter are three-month horizon, twelve schools, non-randomized assignment matched only on baseline attainment, and a small sample with an unvalidated self-report survey acknowledged by the authors.*

EDITORS' SYNTHESIS OF CHEN, WU, CHEN, & ZHOU (2025).

## REFERENCES

1. Chen, R., Wu, Y., Chen, Z., & Zhou, P. (2025), "Advancing educational equity in rural China: the impact of AI devices on teaching quality and learning outcomes for sustainable development," *Frontiers in Psychology* 16:1588047, doi:10.3389/fpsyg.2025.1588047.
2. Paired Case 86 (Gándara et al., AERA Open) — bias in student-success prediction at the higher-education scale.
3. Paired Case 88 (LiveHint AI bias across foundation models, AIED 2025) — bias-surfacing in AI-supported instruction.
4. Paired Case 73 (Doer Effect non-WEIRD LAK 2025 radio-and-phone extension) — non-WEIRD methodological discipline at the heterogeneity-finding axis.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Chen and colleagues' 12-school three-month deployment of AI-supported instructional devices across rural and urban Chinese classrooms reported an equity-direction finding — rural gain exceeded urban gain across both mathematics and history. The hedges that matter are explicit and carried in the prose: three-month horizon, twelve schools, non-randomized assignment matched only on baseline attainment, and a small sample with an unvalidated self-report survey acknowledged by the authors. The case anchors the LEO on fairness beyond omission with a rare published equity-narrowing result.

### LENS APPROACH

Chen et al. is the equity-direction-finding case at the non-WEIRD deployment seam (induced 8.3; LENS D2/PT4; LEO-4 and fairness beyond omission). LENS uses it in Domain 2 (Iterative Development) for the deployment-on-a-defined-horizon discipline and in Domain 5 (Navigating Sociotechnical Constraints) for the equity-direction reading that anchors the fairness-beyond-omission LEO. Pair with Case 86 (Gándara on bias in student-success prediction), Case 88 (LiveHint AI bias across foundation models), and Case 73



(Doer Effect non-WEIRD radio-and-phone extension). The case’s pedagogical value depends on the four binding hedges traveling with the result into print.

| DOMAIN                   | MODES ADDRESSED | LENS COMPETENCY | PROBLEM TYPE          | INDUCED | LEO |
|--------------------------|-----------------|-----------------|-----------------------|---------|-----|
| K-12 EDUCATION           | T K D           | D1 D2 D3 D4 D5  | 4                     | 8.3     | 4   |
| AI-SUPPORTED INSTRUCTION | →               | →               | Iterative Development |         |     |
| EDUCATIONAL EQUITY       |                 |                 |                       |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pre-specify the equity-direction outcome as the decisive reading; the case demonstrates that a deployment study can credibly report a rural-over-urban gain on a defined horizon when the methodological hedges are named.
- ▶ Carry the four binding hedges — horizon, sample, assignment structure, measurement bias — in the deployment write-up itself; the case’s value to the corpus depends on the hedges being internal to the argument rather than buried in a limitations footnote.
- ▶ Treat first-author-from-deployment-country authorship as a methodological feature; the non-WEIRD evidence base is structurally improved by deployments written up by teams with access to the deployment context.

IN OPERATION / AFTER

- ▶ Commission the longer-horizon and larger-sample replication the three-month 12-school study cannot perform; the equity-direction finding’s persistence past three months and across more schools is the validation question the case names.
- ▶ Pair the case in syllabi with Case 86 (Gándara on bias in student-success prediction) and Case 73 (Doer Effect non-WEIRD extension) so the equity-direction reading is taught across deployment scales and population settings.
- ▶ Use the case to anchor the LEO on fairness beyond omission with a rare published equity-direction finding; the corpus needs equity-narrowing results as visible as equity-widening ones, and the methodological discipline that lets the result be reported credibly is the curricular target.

REFLECTION QUESTIONS

1. Identify a deployment in your domain whose equity-direction finding — narrowing or widening a between-setting gap — has not been pre-specified as a primary outcome. What would the pre-specification require, and what data infrastructure would support it?
2. Specify the four decisive hedges you would carry in the prose of a non-WEIRD deployment write-up — horizon, sample, assignment structure, measurement bias. Which of the four is hardest to name credibly in your domain, and what would naming it cost?
3. The case’s structural feature is first-author-from-deployment-country authorship. Pick a deployment in your domain that has been written up primarily by external authors, and ask: what would change in the methodological reading if the write-up were led by authors with access to the deployment context?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

## CASE 76

MULTIMODAL LEARNING ANALYTICS

CLASSROOM DEPLOYMENT

DESIGN-BASED RESEARCH

2023

# Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Deployment team's lessons-learned account of a two-year in-the-wild multimodal learning analytics study with 399 nursing students and 17 educators — indoor positioning, headset audio, physiological wristbands and video in a healthcare simulation classroom; documents what worked, what failed, what the team would have done differently

## IN BRIEF

Martinez-Maldonado et al.'s "Lessons Learnt from a Multimodal Learning Analytics Deployment In-the-wild," in ACM Transactions on Computer-Human Interaction, is the deployment team's own account of a two-year in-the-wild MMLA study with 399 nursing students and 17 educators — indoor positioning, headset audio, physiological wristbands and video in a healthcare simulation classroom. The paper documents what worked, what failed, and what the team would have done differently. The case is offered not as a deployment- results case (the deployment outcomes live in adjacent peer-reviewed papers) but as a published-first-person Practice Flywheel exemplar — the genre the front-matter "unpacking is the method" reframing calls for. The case pairs structurally with the reflective-practice cases elsewhere in the v2 supplemental tier and grounds the practitioner-reflection-as-evidence discipline at the LE-specific layer. The consolidated lessons-learned synthesis is peer-reviewed — TOCHI 31(1), Article 8 — and arXiv:2303.09099 is that article's author preprint. What remains open is the broader question of whether the genre takes hold across the LE community at the scale the Practice Flywheel exemplar requires.

## THE SHIFT

Multimodal learning analytics is the strand of the learning-analytics field where the data substrate moves beyond click-stream and assessment-response logs to include eye-tracking, audio capture, video, and spatial positioning. The lab-deployment record for MMLA is substantial; the in-the-wild classroom-deployment record is structurally different and structurally sparse. What happens when the sensor stack moves from the controlled lab environment into the classroom — where lighting varies, students move, audio overlaps, and the consent architecture has to accommodate the school's operational norms — is a question the published deployment-results papers can answer only partially.<sup>9</sup>

## WHAT IS EMERGING

Martinez–Maldonado et al.’s 2023 TOCHI paper is structured as the deployment team’s own reflection on lessons from two iterations of a single in–the–wild MMLA study. The paper’s content is what the deployment–results papers do not contain: what worked, what failed, what the team would have done differently. The structure is reflective rather than hypothesis–testing — the team narrates the deployments, identifies the recurring patterns, and surfaces the operational knowledge that does not fit into a methods section of a results paper. The lessons are synthesized under five headings — the technological and physical aspects of the deployment; the multimodal data and the interfaces built on it; the design process; participation, ethics and privacy; and the sustainability of the deployment — and the sensing stack behind them was a positioning locator worn in a belly bag, a wireless headset microphone, an Empatica E4 physiological wristband, and video. The honest framing the paper preserves is that the reflective genre is the right vehicle for the kind of operational knowledge the case contains, and that the peer–review structures of the LE field have not consistently supported the genre.<sup>1</sup>

## THE CAPABILITY QUESTION

The case is offered in the corpus not as a deployment– results case but as a published–first–person Practice Flywheel exemplar. The genre the editor’s memo (B1) anticipates — first–person practitioner accounts of deployment iterations, intended to be paired with the front–matter “unpacking is the method” reframing — has structural analogs in adjacent fields (Lutz on reflective journaling, CBE–LSE on reflective–practice primers, SE work–based reflective–practice longitudinals) but has historically been under–published in LE. The Martinez–Maldonado paper is the LE–specific instance of the genre at deployment scale; the case carries it on that basis.<sup>2</sup>

## EARLY EVIDENCE

The structural anchor the case grounds is the sustaining– tacit–capability–across–generations anchor. Practitioner knowledge — what to do when a positioning tag loses its fix mid–session, how to design the consent architecture for headset audio capture across a cohort of hundreds, what physiological wristband placement looks like when the simulation–room layout shifts — walks out the door if it is not narrated. The paper’s first–person genre is the narration vehicle, and the case grounds the curriculum’s response to the question of how operational knowledge accumulates and propagates across practitioner cohorts. The pair with the CIRCUIT cases (78, 68) is the workforce–and–capability layer; the pair with Case 77 (hybrid human–AI tutoring) is the design–iteration layer.<sup>3</sup>

## OPEN PROBLEMS

The sourcing needs stating precisely. The consolidated lessons–learnt synthesis the case carries is peer–reviewed — it is TOCHI 31(1), Article 8 — and arXiv:2303.09099 is that article’s author preprint, not a separate unreviewed document. The open question is therefore one and not two: whether the first–person practitioner–reflection genre takes hold in the LE community at sufficient scale to function as the Practice Flywheel exemplar the framework names. The case is included not despite the single–account limit but with it; that limit is part of what the case teaches.

*Practitioner knowledge walks out the door if it is not narrated. The first-person reflective genre is the narration vehicle, and the field has not consistently supported it.*

EDITORS' SYNTHESIS OF MARTINEZ-MALDONADO ET AL. (2023).

REFERENCES

1. Blikstein, P., & Worsley, M. (2016), "Multimodal Learning Analytics and Education Data Mining: Using computational technologies to measure complex learning tasks," Journal of Learning Analytics 3(2):220–238 — broader MMLA literature backdrop.

2. Martinez-Maldonado, R., Echeverria, V., Fernandez-Nieto, G., et al. (2023), "Lessons Learnt from a Multimodal Learning Analytics Deployment In-the-wild," ACM Transactions on Computer-Human Interaction 31(1), Article 8, pp. 1–41 (November 2023), doi:10.1145/3622784; author preprint at arXiv:2303.09099.

3. v2\_research/03\_editor\_request\_memo.md, item B1 — the Practice Flywheel commissioning structure this case is offered as a published-first-person exemplar within.

4. Schön, D. (1983), The Reflective Practitioner — the genre's theoretical underpinning, referenced across the reflective-practice case tier.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Martinez-Maldonado et al. is the LE-specific published- first-person Practice Flywheel exemplar at MMLA in-the-wild deployment scale. The case is offered not as a deployment- results case but as a genre exemplar — the reflective- practice account, peer-reviewed in TOCHI 31(1) Article 8. What is still open is whether the genre takes hold across the LE community at the scale the Practice Flywheel exemplar requires.

LENS APPROACH

MMLA in-the-wild is the sustaining-tacit-capability case in the LE-conferences tier (induced 6.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the reflective-narration-of-design-iteration discipline and in Domain 4 (Test and Evaluation) for the discipline that binds the single-account limit to the genre's pedagogical role. Pair with Cases 78 and 68 (CIRCUIT workforce-and-capability layer) and Case 77 (hybrid human- AI tutoring design-iteration layer).

| DOMAIN                        | MODES AT STAKE          | LENS COMPETENCY | PROBLEM TYPE | INDUCED | LEO   |
|-------------------------------|-------------------------|-----------------|--------------|---------|-------|
| MULTIMODAL LEARNING ANALYTICS | T K N                   | D1 D2 D3 D4 D5  | 4            | 6.3     | 2 · 4 |
| CLASSROOM DEPLOYMENT          | → Iterative Development |                 |              |         |       |
| DESIGN-BASED RESEARCH         |                         |                 |              |         |       |

APPROACHES TO CONSIDER

- DURING DEVELOPMENT

  - ▶ Narrate the deployment in first person while it is still operating; the operational knowledge the case names is contemporaneous, and the post-hoc reconstruction loses the texture the first-person genre preserves.
  - ▶ Treat the reflective paper as a deliverable on par with the results paper; the case demonstrates that the operational knowledge has the same evidentiary status as the methods-section content of a results paper, and the publication structure should support it.
  - ▶ Build the consent and ethics architecture around the in-the-wild deployment substrate from the start; the lessons-learned content includes consent-
- IN OPERATION / AFTER

  - ▶ Take the peer-reviewed lessons-learnt article as proof the genre can clear review at a top venue; the validation step still outstanding is not this paper's publication but the field's — whether enough teams publish comparable accounts for the genre to function as evidence.
  - ▶ Carry the Practice Flywheel exemplar designation into the curriculum's first-person-account commissioning structure; the case is the LE-specific anchor for the genre the editor's memo (B1) calls for, and the curriculum is the vehicle that institutionalizes the genre.

architecture failures that the lab-deployment record does not surface.

- Pair the case with the front-matter “unpacking is the method” reframing so the genre’s role in the casebook’s pedagogical architecture is visible; the case is offered as an exemplar of the genre, and the genre is offered as the curriculum’s response to the sustaining-tacit-capability question.

### REFLECTION QUESTIONS

1. Identify a deployment in your domain whose operational knowledge — what worked, what failed, what the team would have done differently — has not been narrated outside the team. What would the first-person reflective account look like, and what publication venue would carry it?
2. Specify the consent and ethics architecture you would build into an in-the-wild deployment in your domain that the lab-deployment record would not have surfaced. The case’s lessons-learned content includes consent-architecture failures; what would the deployment-substrate-specific architecture look like?
3. The case is offered as a published-first-person Practice Flywheel exemplar. Identify a practitioner in your domain whose operational knowledge you would commission a first-person account from. What would the commissioning structure look like, and what would the account contain that the published-results papers do not?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

# Hybrid Human-AI Tutoring — Augmentation, Not Delegation

**T** Training Gap   **K** Knowledge & Institutional Memory   **D** Designed Out

**IMPACT** Three quasi-experimental studies of hybrid human-AI tutoring deployments reported improvements in student learning relative to comparison conditions; the AI is positioned as augmentation, not delegation; the human tutor retains authorization to override and re-direct

## IN BRIEF

Thomas et al.'s LAK 2024 paper, "Improving Student Learning with Hybrid Human-AI Tutoring: A Three-Study Quasi-Experimental Investigation," reports three quasi-experimental studies of hybrid deployments where AI augmentation is added to human tutoring rather than used to replace it. The headline finding is that learning outcomes improved in each of the three studies, and only Site 1 carried a control condition, while Sites 2 and 3 compared students against their own pre-treatment baseline. The contribution the case carries for the LENS framework is the design positioning: the AI is augmentation, the human tutor retains the authorization to override and re-direct, and the measured outcome is student learning rather than AI-system fidelity. The case is the small-tier intervention-side counterpart to Case 20 (TREWS, the clinician-AI teaming case that worked) translated into education. Pair also with Cases 78 and 68 (CIRCUIT human-AI workforce) and Case 5 (Epic Sepsis, the delegation case that did not work). Open questions: longitudinal persistence; transfer to lower-resource tutoring contexts where human-tutor availability is the binding constraint.

## BACKGROUND

The deployment record for AI in tutoring has been pulled in two directions. The fully-automated tutoring track — from Cognitive Tutor through LLM-based tutoring (Case 88) — has tested whether AI alone can replace or substantially reduce the human-tutor role. The augmentation track has tested whether AI can extend the reach and effectiveness of human tutors, with the AI positioned as a tool the tutor uses rather than as a substitute for the tutor. Thomas et al.'s LAK 2024 paper is the strongest published evaluation of the augmentation track to date, and it reports three quasi-experimental studies that converge on a positive finding.<sup>o</sup>

## THE INTERVENTION

The three studies test variants of the augmentation pattern in K-12 tutoring deployments. The AI supports the human tutor with information surfacing, problem recommendation, and student-progress visibility; the human tutor retains the conversational and pedagogical

lead. Only the first site carried a math–software–only control condition; the second and third compared students against their own pre–treatment usage, and the third compared two tutoring intensities against each other. Across the three studies, the hybrid condition produced measurable improvements in student learning. The replication structure across the three studies — same authorship team, varying institutional context, converging direction of effect — is the methodological backbone of the case.<sup>1</sup>

#### HOW IT WORKED

The design positioning the case carries for the LENS framework is the augmentation–not–delegation frame. The AI is positioned as augmentation, the human tutor retains the authorization to override and re–direct, and the measured outcome is student progress inside the learning software, with impact on external assessments still untested. This is the design pattern that worked in clinical–decision–support at Case 20 (TREWS) and that did not work at Case 5 (Epic Sepsis) — where TREWS preserved clinician authorization and built the explanation structure that supported it, the Epic Sepsis deployment pattern collapsed clinician judgment into alert compliance. Hybrid human–AI tutoring is the educational analog of the TREWS pattern, and the LAK 2024 paper is the evidence base that grounds the analog.<sup>2</sup>

#### THE EVIDENCE

The case anchors with the CIRCUIT pair (Cases 78 and 68) at the workforce–augmentation layer. CIRCUIT proofreading positions human capability as the recovery mechanism for automation failure at petabyte scale; the CIRCUIT workforce model builds the capability in the first place; hybrid human–AI tutoring positions AI as augmentation of an already–capable human tutor. The three cases together teach the augmentation–and–correction pattern across three deployment substrates — connectomics proofreading, neuroscience–workforce development, and tutoring — and they ground the curriculum’s machine–teaming and delegation–with–revocation anchors at each substrate.<sup>3</sup>

#### WHAT TRANSFERRED

The open questions the case carries are the ones the authors name. Longitudinal effect persistence is not yet in the evidence base — the three quasi–experimental studies report end–of–intervention outcomes, not multi–year follow–through. Whether the design transfers to lower–resource tutoring contexts, where human–tutor availability is the binding constraint and the augmentation–of–a–tutor frame may not apply, is the open generalization question. The quasi–experimental design is honest about its causal–inference limits — randomization is not at the level a cluster RCT would provide — and the case carries the qualification. Persistence, transfer, and the tutor–scarce–context generalization stay open.

*The AI is positioned as augmentation, not delegation. The human tutor retains the authorization to override and re–direct. The measured outcome is student learning, not AI–system fidelity.*

EDITORS’ SYNTHESIS OF THOMAS ET AL. (2024).

#### REFERENCES

1. Thomas, D. R. et al. (2024), “Improving Student Learning with Hybrid Human–AI Tutoring: A Three–Study Quasi–Experimental Investigation,” LAK ’24, doi:10.1145/3636555.3636896.
2. Case 20 (TREWS) reference set — Henry et al. (2022), *Nature Medicine* — clinician–AI teaming analog.
3. Case 5 (Epic Sepsis) reference set — Wong et al. (2021), *JAMA Internal Medicine* — delegation–collapse analog.

2. Koedinger, K. R. et al. — Cognitive Tutor literature as the fully-automated track the augmentation track contrasts with.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Hybrid human-AI tutoring is the educational analog of the TREWS clinician-AI teaming pattern. AI is positioned as augmentation; human tutor retains override authorization; student-learning outcome is the measure. Three quasi- experimental studies converge on positive learning effects. The case pairs with Cases 20 (TREWS), 5 (Epic Sepsis), and 68 / 78 (CIRCUIT) in the human-AI teaming thread and grounds the delegation-with- revocation LEO at the educational deployment.

LENS APPROACH

Hybrid human-AI tutoring is the augmentation-not-delegation case in education (induced 6.4; LENS D3/ PT6). LENS uses it in Domain 3 (Human-System Collaboration) for the augmentation pattern and the override-authorization frame, and in Domain 2 (Iterative Development) for the three-study converging-design replication. Pair with Cases 20 (TREWS) and 5 (Epic Sepsis) at the clinical analog, and with Cases 78 and 68 (CIRCUIT) at the workforce-augmentation analog.

| DOMAIN                                                                           | MODES ADDRESSED                        | LENS COMPETENCY                                                                                             | PROBLEM TYPE | INDUCED | LEO   |
|----------------------------------------------------------------------------------|----------------------------------------|-------------------------------------------------------------------------------------------------------------|--------------|---------|-------|
| <div>TUTORING</div> <div>HYBRID HUMAN-AI SYSTEMS</div> <div>K-12 EDUCATION</div> | <div>T</div> <div>K</div> <div>D</div> | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> <div>Human-System Collaboration</div> | 6            | 6.4     | 2 · 3 |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Position the AI as augmentation explicitly in the design documentation, not implicitly in the deployment pattern; the augmentation-vs-delegation distinction is the design choice that decides the rest and should be named as such.

► Preserve human-tutor authorization to override and re-direct as a system-design requirement, not as a discretionary affordance; the comparison with Case 5 (Epic Sepsis) is that override authorization collapses when the system pattern does not actively preserve it.

► Measure the student-learning outcome, not the AI-system-fidelity outcome; the case’s pedagogical framing depends on the outcome instrument being the educationally relevant one, not the AI-development-internal one.

IN OPERATION / AFTER

- Commission longitudinal follow-through that extends the evidence base past the end-of-intervention horizon; the open persistence question is testable against the same deployment with additional data infrastructure.

► Test the augmentation design in tutor-scarce contexts; the open generalization question is whether the pattern transfers to settings where the binding constraint is human-tutor availability rather than human-tutor effectiveness.

► Pair the case with Case 20 (TREWS) in the curriculum so the augmentation-and-override pattern is taught across clinical and educational substrates; the two cases together ground the delegation-with-revocation LEO with two converging instances.

REFLECTION QUESTIONS

1. Identify an AI deployment in your domain where the design choice between augmentation and delegation has been implicit rather than explicit. What would change in the system design if the choice were named explicitly, and what comparison condition would you build to test the difference?
2. Specify the override-authorization preservation mechanism in your domain’s analog deployment. Is the human operator’s authority to override and re-direct a system-design requirement, a discretionary affordance, or an implicit assumption? Which of the three is honest about what the system currently supports?



**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

## CASE 78

NEUROSCIENCE/CONNECTOMICS

HUMAN-AI TEAMING

WORKFORCE TRAINING

2017 – PRESENT

# CIRCUIT / MICrONS – The Human Correction Layer at Petabyte Scale

**H** Human-System Interface   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**DISCLOSURE** Contribution: an editor has research adjacency to connectomics programs here, and the case originates at an editor's institution (Johns Hopkins / APL). The method is anchored to peer-reviewed work; program training-outcome evidence is institutional documentation, not independent evaluation — the gap rendered as the evidence-tier flag above.

**EVIDENCE TIER** internal-pm — source confidence flagged; future validation ongoing.

**IMPACT** Under IARPA's MICrONS program, automated electron-microscopy segmentation produces brain maps too large and too error-prone to deploy without human correction; CIRCUIT trains undergraduate cohorts as the proofreading workforce that is the recovery mechanism for automation failure at petabyte scale

## IN BRIEF

Under IARPA's Machine Intelligence from Cortical Networks (MICrONS) program, automated electron-microscopy segmentation produces petabyte-scale brain maps that are too large and too error-prone to be deployed without human verification. APL's CIRCUIT program trains cohorts of undergraduates to proofread and correct these maps; APL's BossDB stores and serves the petabyte-scale data. The learning-engineering content the case carries is the human-in-the-loop correction layer: where automation fails at the petabyte boundary, a trained human capability is the recovery mechanism that makes the data scientifically usable. The honest evidentiary state — preserved as the evidence-tier flag rendered under the title — is that the connectomics method and infrastructure are documented in the peer-reviewed literature, but CIRCUIT's program training outcomes (did the cohorts reliably produce proofreading capability, at what error rates, with what retention, with what transfer to other tasks) sit in institutional documentation rather than peer-reviewed program evaluation. The case is included as a frontier — the structural pattern is the increasingly central question of how to design a human correction layer for generative and automated systems at scale, and future validation will continue as the program-evaluation literature builds out. The COI render under the title is binding: home institution, research adjacency, evidence-tier flag, all visible at the point of reading.

## THE SHIFT

Connectomics — the reconstruction of detailed wiring diagrams of brain tissue from electron-microscopy imaging — is the modern incarnation of a problem with the same structural

shape as the early radiology-AI cases (Case 35): automation produces a primary output at scale that no human team could produce manually, and the output is too important and too error-prone to be deployed without verification. The MICrONS program, funded by IARPA, set out to close the performance gap between human analysts and automated pattern recognition by reverse-engineering the algorithms of the brain. The reconstructed volume is about a cubic millimetre of mouse visual cortex — some 200,000 cells, roughly 120,000 of them neurons, and over 500 million detected synapses — imaged at petabyte scale.<sup>0</sup>

#### WHAT IS EMERGING

The automation produces petabytes of data. Automated segmentation labels each voxel with the neuron it belongs to, but at petabyte scale the per-voxel error rate, however low, compounds into a reconstruction that contains many incorrect neuron boundaries — exactly the boundaries the downstream science depends on. The structural form is familiar: the model is good enough to make the project tractable, and not good enough to deploy without a verification layer. The verification layer in this case is human: trained proofreaders work through the segmentation and correct it, neuron by neuron, with tools designed for the task.<sup>1</sup>

#### THE CAPABILITY QUESTION

The CIRCUIT program at APL trained cohorts of undergraduates to be that proofreading workforce. The program is documented across institutional news, program description, and the MICrONS program literature; CIRCUIT trainees and APL infrastructure (BossDB) together produced the human and computational substrate of the verification layer. The learning-engineering content is the human capability as the recovery mechanism for automation failure at the petabyte boundary: where the model fails, the trained human is the design decision that makes the system's output usable.<sup>2</sup>

#### EARLY EVIDENCE

The honest evidentiary state is what the evidence-tier flag under the title encodes. The connectomics method, the automated-segmentation literature, and the BossDB infrastructure are documented in peer-reviewed connectomics publications. CIRCUIT's program **training outcomes** — did the cohorts reliably produce proofreading capability at the error rates and retention required, did the trainees transfer to other tasks, what is the program's measured effect on the downstream science — are documented in institutional news and program description rather than in peer-reviewed program evaluation. The case is teachable on the structural pattern; the operating outcomes are at the institutional-program-management tier and should be read at that tier. The COI under the title — research adjacency, home institution — makes the institutional-tier evidence claim auditable.<sup>3</sup>

#### OPEN PROBLEMS

The frontier note the case carries is the most forward-looking idea from the pass-4 sweep. The human correction layer for generative and automated systems at scale is going to be a defining capability across an expanding number of domains — automated transcription, AI-generated code, model-extracted structured data, document segmentation, scientific imaging reconstruction. CIRCUIT is one instance of the pattern in neuroscience; the structural question — how to design, staff, train, and govern a human correction layer for the gap between what models produce and what is operationally usable — recurs and is not

well-named in the existing curriculum. It is also the shape Wachter and Brynjolfsson give the productivity paradox in medicine: a general-purpose technology pays off only after the complementary organizational work — the workflow redesign, the roles, the training — has been done, and that complementary work is exactly what a correction layer is. The frontier flag on this case is the case-grounded basis for proposing a sub-competency in this area, pairing with the broader v2 AI-delegation typology and the Domain 3 **Delegation with revocation** LEO.<sup>4</sup>

*The recovery mechanism for automation failure is a designed human capability, not an emergent workforce.*

EDITORS' SYNTHESIS OF THE CIRCUIT / MICRONS PROGRAM RECORD.

REFERENCES

1. MICrONS program literature (IARPA) — connectomics method and automated segmentation evidence base.

2. APL BossDB documentation — petabyte-scale connectomics data infrastructure.

3. Cervantes, Floryanzia, Sharp, Gray-Roncal, & Johnson (2023), "Empowering Trailblazers toward Scalable, Systematized, Research-Based Workforce Development," ASEE Annual Conference, doi:10.18260/1-2-43271 — the paired peer-reviewed CIRCUIT case (Case 68).

4. CIRCUIT program documentation (JHU Hub, 2017 – present) — institutional/program description; the training-outcome evidence is at this tier and the evidence-tier flag is binding.

5. Wachter, R. M., & Brynjolfsson, E. (2024), "Will Generative Artificial Intelligence Deliver on Its Promise in Health Care?" JAMA 331(1):65–69 — the productivity-paradox framing: general purpose technologies pay off only after the complementary organizational work is done.

THE LEARNING ENGINEERING LENS

LE INSIGHT

CIRCUIT / MICrONS is the frontier instance of designing the human correction layer for automated systems at scale. The connectomics method is peer-reviewed; the program training outcomes are institutional documentation, and the gap is rendered as the evidence-tier flag under the title. The forward-looking question the case names — how to design the human correction layer for generative and automated systems at scale — recurs across domains and is not well-named in the existing curriculum.

LENS APPROACH

CIRCUIT proofreading is the human-correction-layer frontier case (induced 3.4; LENS D3/PT6) — Domain 3 for **Delegation with revocation**; Domain 4 for the evidence- tier split. Pair with Case 68 and Cases 20, 5, 189, 3. COI binds.

| DOMAIN                        | MODES AT STAKE | LENS COMPETENCY              | PROBLEM TYPE | INDUCED | LEO   |
|-------------------------------|----------------|------------------------------|--------------|---------|-------|
| NEUROSCIENCE/<br>CONNECTOMICS | H K N          | D1 D2 D3 D4 D5               | 6            | 3.4     | 4 · 3 |
| HUMAN-AI TEAMING              | →              | → Human-System Collaboration |              |         |       |
| WORKFORCE<br>TRAINING         |                |                              |              |         |       |

APPROACHES TO CONSIDER

- DURING DEVELOPMENT

▶ Treat the human correction layer as a designed capability deliverable, not an emergent workforce:
- IN OPERATION / AFTER

▶ Commission peer-reviewed program-evaluation work on the training-outcome side so the institutional

- define the error rate to achieve, the verification protocol, the trainee qualification, and the tooling.
- ▶ Identify in advance which errors the automation will make at scale and design the correction tooling around those error classes — automation failure is structured, and the correction layer should be too.
- ▶ Carry the evidence-tier honesty: peer-reviewed sourcing for the method and infrastructure, institutional-program documentation for the program training outcomes, with the gap visible in the case rather than smoothed.
- record can converge toward independent evidence over time, while the program is operating.
- ▶ Track the human correction layer across automation upgrades: as the model improves, what changes about the kind of correction the trainees do, the error classes, the qualification requirements?
- ▶ Carry the frontier note into the curriculum conversation: the design of the human correction layer for generative and automated systems at scale is a sub-competency the existing framework does not yet name, and CIRCUIT is one of several emerging instances.

### REFLECTION QUESTIONS

1. Identify a domain in your work where automation produces a primary output at a scale that exceeds manual review. What is the human correction layer's design — error classes targeted, verification protocol, trainee qualification, tooling — and which of these decisions are designed vs. emergent?
2. Specify the peer-reviewed vs. institutional-evidence split for a human-correction-layer program you would propose. The connectomics method can be cited from the literature; the program training outcomes will not be. What evidence tier is honest for each layer?
3. The case's frontier note — designing the human correction layer for generative/automated systems at scale — is a sub-competency the existing curriculum does not yet name. What instance from your domain (automated transcription, AI code, structured-data extraction, document segmentation) would be a paired case to anchor this sub-competency against?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · incident analysis · safety dashboards

# The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops

**K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** Across a U.S. corporate-training market sized at more than \$125B/year in the ASTD figure Blume et al. cite (Paradise 2007), the dominant evaluation framework thins as it climbs and rarely reaches the top: in ATD's 2016 survey of 199 talent-development professionals, 83 percent evaluated Level 2 (learning), 60 percent reported some Level 3 (behavior), 35 percent evaluated Level 4 (business results) to any extent, and 15 percent measured return on investment — the levels that flatter the program are the ones that get measured

## IN BRIEF

Donald Kirkpatrick's four-level model — Reaction, Learning, Behavior, Results — has been the dominant corporate-training evaluation framework for sixty years, increasingly described as a “chain of evidence” from satisfaction through workplace behavior change to business outcomes. The documented systemic pattern is that the chain thins as it climbs: reported evaluation falls from near-universal at Levels 1 and 2 to 60 percent at Level 3, 35 percent at Level 4 and 15 percent for ROI. Level 3 (behavior change on the job) and Level 4 (business results) require data that the training organization typically cannot access: longitudinal performance traces, business-unit outcome metrics, line-manager observation. Satisfaction scores do not predict on-job performance, and knowledge retention does not guarantee workplace application. The case is the corporate-scale instance of the enthusiasm-evidence gap and a direct illustration of the revised “decision-grade evidence” framing in the v2 research backbone: the evidence most L&D decisions ride on is structurally sub-decision-grade. The evidence base is practice-synthesis: Devlin Peck, D2L, Valamis, and related evaluation-practice guides documenting where evaluation stops, ATD's 2016 association survey for the level-by-level figures, plus the Blume meta-analysis (Case 65) for the transfer half of the chain. The evidence-tier flag is rendered under the title; future validation will continue as the corporate-L&D evaluation literature consolidates.

## THE SHIFT

In 1959 Donald Kirkpatrick proposed four levels at which a training program could be evaluated: Reaction (did learners like it), Learning (did they acquire the content), Behavior (did they apply it on the job), and Results (did business outcomes improve). The four levels

became the dominant corporate–training evaluation framework and were later described as a “chain of evidence” — each level meant to provide the evidence for the next.<sup>0</sup>

#### WHAT IS EMERGING

The documented systemic pattern in corporate L&D evaluation practice is that the chain thins as it climbs and the top of it is rarely reached. Practice–literature synthesis across major evaluation–guidance sources (Devlin Peck, D2L, Valamis, and corporate L&D benchmarking studies) describes programs stopping at Levels 1 and 2; the best single survey figure puts the sharp break a level higher. ATD’s 2016 survey of 199 talent–development professionals found roughly 90 percent evaluating Level 1 and 83 percent Level 2, 60 percent reporting some Level 3 evaluation, 35 percent evaluating business results to any extent, and 15 percent measuring return on investment — and those upper figures count any program evaluated at that level to any extent, not routine completion of the chain. The structural reason is that Level 3 (behavior change) requires longitudinal observation that the training organization cannot conduct, and Level 4 (results) requires business outcome data that often sits outside L&D’s reach.<sup>1</sup>

#### THE CAPABILITY QUESTION

The consequence is a field that spends heavily on training — U.S. organizations reported more than \$125B/year on workplace training in the ASTD figure Blume et al. (Case 65) cite, drawn from ASTD’s 2007 State of the Industry review — while measuring mostly the variable that flatters the program. Satisfaction scores do not predict on–job performance; knowledge retention measured immediately after training does not guarantee workplace application. The chain of evidence is cited as the framework; in practice it thins above Level 2 and breaks decisively at the link into Level 4, and the decisions made on the available data are not decisions about whether training is producing capability change.<sup>2</sup>

#### EARLY EVIDENCE

The evidence–tier flag rendered under the case title is decisive. The case is documented through evaluation–practice guides synthesizing where evaluation stops across many organizations rather than through a single peer–reviewed study of the phenomenon; the strongest single datum in it is an association survey (ATD 2016), not peer–reviewed research. The pattern is consistent across the sources and is the practitioner consensus, but the magnitudes vary by sector and the sectoral breakdown is in flux as the field consolidates its evaluation evidence. Future validation will continue as the empirical synthesis improves; the case is included with the flag because the pattern itself is teachable and the practitioner consensus is durable.<sup>3</sup>

#### OPEN PROBLEMS

What the case teaches in pair with Blume (Case 65) is the structural form of the enthusiasm–evidence gap at corporate scale, and it is the case–grounded basis for the revised “decision–grade evidence” framing proposed in the v2 research backbone. The capability deliverable is an evaluation architecture that crosses the Level 2 / Level 3 seam — by partnering with line management for behavior observation, by instrumenting the workplace tasks the training targets, and by reporting honestly what evidence is and is not available — rather than declaring the chain satisfied at the Level the training organization can control. Until that crossing happens, most corporate L&D decisions are made on structurally sub–

decision-grade evidence, and the subobjective **Judgment under inadequate evidence** (4.4) is exactly the capability the practitioner needs.<sup>4</sup>

*The chain of evidence is cited as a framework. In practice it thins above Level 2 and breaks at the link into Results, and the decisions made on the available data are not decisions about capability.*

EDITORS' SYNTHESIS OF THE KIRKPATRICK MODEL AND CORPORATE L&D EVALUATION PRACTICE.

REFERENCES

1. Kirkpatrick (1959–1960), original four-level evaluation series in the Journal of the American Society of Training Directors (Nov 1959 – Feb 1960); updated as Kirkpatrick & Kirkpatrick, Evaluating Training Programs (3rd ed., 2006).

2. Devlin Peck, “The Kirkpatrick Model: 4 Levels, Examples, and Free Template” — synthesis of the stop-at-L2 pattern in corporate practice.

3. D2L, “Evaluating Training Programs: Models, Metrics and Business Impact,” practitioner guide documenting the same pattern.

4. Valamis, “Kirkpatrick model: how to evaluate whether training worked” — practitioner guide on evaluation-practice gaps.

5. Association for Talent Development (2016), Evaluating Learning: Getting to Measurements That Matter — survey of 199 talent-development professionals; source of the level-by-level figures ( 90% L1, 83% L2, 60% L3, 35% L4, 15% ROI). Association research, not peer-reviewed.

6. Blume, Ford, Baldwin, & Huang (2010), “Transfer of Training: A Meta-Analytic Review,” Journal of Management 36(4):1065–1105 — the paired peer-reviewed case (65); its opening \$125B figure is ASTD’s, via Paradise (2007).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Kirkpatrick chain’s fall-off above Level 2 — about a third of surveyed organizations reaching business results and a seventh reaching ROI — is the corporate-scale instance of the enthusiasm-evidence gap and the direct illustration of the revised “decision-grade evidence” point: the evidence most L&D decisions ride on is structurally sub-decision-grade. Evidence base is practice-synthesis tier; the pattern is consistent across sources, the magnitudes still consolidating; future validation will continue.

LENS APPROACH

Kirkpatrick is the corporate-L&D evaluation case (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the subobjective **Judgment under inadequate evidence** (4.4) — the practitioner must decide well on structurally sub-decision-grade evidence — and in Domain 5 (Navigating Sociotechnical Constraints) for the seam above Level 2 — on-job behavior and business results — that the training organization cannot cross alone. Direct pair with Case 65 (Blume transfer meta-analysis) for the workplace-environment half of the chain.

| DOMAIN              | MODES AT STAKE     | LENS COMPETENCY           | PROBLEM TYPE | INDUCED | LEO |
|---------------------|--------------------|---------------------------|--------------|---------|-----|
| CORPORATE L&D       | → <div>K N</div> → | <div>D1 D2 D3 D4 D5</div> | 5            | 2.1     | 4   |
| TRAINING EVALUATION |                    | Test and Evaluation       |              |         |     |
| WORKFORCE DEV       |                    |                           |              |         |     |

APPROACHES TO CONSIDER

- DURING DEVELOPMENT

▶ Design the evaluation across the seam above Level 2 before the training is delivered: partner with line management for behavior observation, identify the workplace tasks the training targets, and instrument them.

▶ Specify in advance the on-job behavior and business outcome that would count as evidence the training
- IN OPERATION / AFTER

▶ When Level 3 and Level 4 data are not available, say so plainly in the evaluation report rather than substituting Level 2 metrics; the subobjective **Judgment under inadequate evidence** (4.4) is what the practitioner needs to do well here.

▶ Track the work environment as a separate variable — Blume’s meta-analysis (Case 65) names it as decisive



worked, with a reporting cadence long enough for transfer effects to surface.

- ▶ Treat the practice-synthesis evidence base honestly: name the pattern, flag the tier, and do not present a stop-at-L2 outcome as if it were a Level-3 or Level-4 finding.

— and report the training outcome conditional on environment, not as a property of the training alone.

- ▶ Use the structural pattern to argue for the evaluation architecture investment, not to abandon evaluation. The gap is a capability gap, not an argument against measurement.

### REFLECTION QUESTIONS

1. Identify a recent corporate training program in your organization. At which Kirkpatrick level did evaluation stop? What would the Level 3 and Level 4 measurement have required, and who would have had to provide the data?
2. Specify the evaluation architecture you would build to cross the seam above Level 2: which on-job behavior, which business outcome, on what cadence, in partnership with which line-management role.
3. The case is practice-synthesis tier. What is the minimum peer-reviewed or program-evaluation evidence you would require before relying on the stop-at-L2 pattern to justify an evaluation-architecture investment in your context?

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

# Georgia State University Predictive Analytics

**T** Training Gap **K** Knowledge & Institutional Memory

**IMPACT** Six-year graduation rate 32% → 54%; equity gaps in graduation eliminated; 2,000+ more graduates per year

## IN BRIEF

Georgia State University built a predictive-analytics advising system that tracks some 800 risk factors per student and triggers proactive outreach when early warning signs appear. It was designed with equity as a primary constraint — the explicit goal was to eliminate, not reproduce, graduation gaps — and the alerts prompt human advisors rather than making automated decisions. The six-year graduation rate rose from 32% to 54%; Black and Pell-eligible students now graduate at the same rate as their peers, and GSU produces roughly 2,000 more graduates a year. The difference from the algorithmic-bias cases (46, 49, 191) is design: GSU built equity in from the start, used predictions to trigger more human support rather than gatekeeping, and tracked outcomes by demographic group as a primary metric.

## BACKGROUND

Georgia State was a regional commuter university where only about a third of students finished in six years, with large gaps by race and income. Like many institutions it had predictive data but such systems, where they existed elsewhere, were typically used for triage or gatekeeping — risking the reproduction of existing inequities rather than their repair. A model that flags at-risk students can just as easily steer them away as toward help; the same prediction serves opposite ends depending on what the institution decides to do with it.<sup>0</sup>

## THE INTERVENTION

Beginning in 2012, GSU deployed a predictive-analytics advising system that monitors roughly 800 behavioral and academic risk factors per student daily and fires an alert to an advisor when warning signs — a missed assignment, a poor grade in a gateway course — appear. The system was built with equity as a primary design constraint, with the explicit aim of closing graduation gaps. Daily monitoring meant the alert fired while there was still time to act — a slipping student was caught at the missed assignment rather than at the failed semester, when intervention could still change the outcome.<sup>1</sup>

## HOW IT WORKED

The load-bearing design choice was the human-loop architecture: alerts trigger proactive advising — a phone call, a meeting, a financial-aid check — rather than automated decisions. Predictions are used to deliver more support to at-risk students, not to gatekeep them out. Human judgment stays in the loop, and the model functions as decision support rather than


decision-maker. Routing every alert through an advisor rather than an automated action is what kept the prediction in service of the student: the model identified who needed attention, and a person decided what that attention should be.<sup>2</sup>

#### THE EVIDENCE

GSU's six-year graduation rate rose from 32% to 54%, and the institution now produces some 2,000 additional graduates a year. The graduation rate for Black students rose to match the overall rate, and Pell-eligible students graduate at the same rate as non-Pell students — the equity gap was eliminated rather than merely narrowed. Eliminating the gap rather than narrowing it is the decisive result: the overall rate rose while the disparities by race and income closed, so the gain did not come at the expense of the students the system was most at risk of leaving behind. The attribution carries a hedge that has to survive: most of the 32%-to-54% rise accumulated across a decade in which GPS Advising (launched 2012) was one component of a broader bundle of reforms — meta-majors, Panther Retention Grants, freshman learning communities, block scheduling — so the isolated causal contribution of predictive advising cannot be cleanly separated from the co-occurring changes.<sup>3</sup>

#### WHAT TRANSFERRED

GSU is the positive counterpart to the algorithmic-harm failures earlier in this part and in Part VII (the A-Level algorithm, Case 49; Robodebt, Case 191; educational predictive bias, Case 46). The same technical capability — a predictive model — produced an equity gain rather than an equity harm because of how it was framed and governed. The case is the strongest evidence that construct definition and human-loop architecture, not the model itself, determine whether prediction helps or harms. Holding the technology constant and varying only the design constraint and governance is what isolates the lesson: the same predictive capability that harmed in those cases helped here, so the framing and the human loop, not the model, are where intent lives.<sup>4</sup> GSU has since moved to export the model: in October 2020 it founded the National Institute for Student Success, which by the end of 2024 had engaged more than 130 campuses serving some 1.5 million students, while GSU's own six-year graduation rate held steady near 53 percent — the test now shifting from whether the bundle worked at one institution to whether the advising-and-governance recipe transfers to others.<sup>5</sup>



*Predictions trigger support, not gatekeeping.*

EDITORS' SYNTHESIS OF THE GSU ADVISING MODEL, DRAWN FROM RENICK & STROM (2020) AND FAUSSET'S NEW YORK TIMES FEATURE (15 MAY 2018)

REFERENCES

1. Georgia State University institutional research and Strategic Plan reports — graduation-rate and equity data.

2. Fausset, R., "Georgia State, Leading U.S. in Black Graduates, Is Engine of Social Mobility," New York Times, 15 May 2018 — the 800-factor advising model.

3. EDUCAUSE Review on GSU predictive advising — the human-loop architecture.

4. Renick, T. & Strom, A. (2020) on GSU's advising transformation — the system design and outcomes.

5. Complete College America, Game Changers documentation — dissemination of the model.

6. Georgia State University National Institute for Student Success (founded 2021) — reported engagement with 130-plus campuses and 1.5 million students by 2024.

THE LEARNING ENGINEERING LENS

LE INSIGHT

GSU is the positive counterpart to A-Level (49), Robodebt (191), and educational algorithmic bias (46). The same technical capability — a predictive model — was deployed under a different design constraint and produced an equity outcome rather than an equity harm. The case is the strongest available evidence that the **construct definition** and the **human-loop architecture** determine whether a predictive model produces good or harm, not the model itself.

LENS APPROACH

LENS treats GSU as the canonical positive case for predictive analytics in education. LEN 4 examines the measurement architecture that made equity a primary outcome. LEN 7 examines the governance structure that kept the system as decision support rather than automated decision. LEN 1 uses it as a problem-framing exemplar.

| DOMAIN              | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------|-----------------|------------------|--------------|---------|-----|
| EDUCATION AT SCALE  | → T K           | → D1 D2 D3 D4 D5 | 5            | 8.3     | 4   |
| Test and Evaluation |                 |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Set equity as a primary design constraint from the start — the explicit aim of closing gaps — rather than discovering disparities after deployment.
- ▶ Build a human-loop architecture so alerts trigger proactive support routed through an advisor, with the model as decision support and a person deciding the action.
- ▶ Tune the monitoring to fire early — at the missed assignment, not the failed semester — so the intervention reaches the student while it can still change the outcome.

IN OPERATION / AFTER

- ▶ Track outcomes by demographic group as a primary metric, so the system is judged on whether it closes gaps rather than merely raises the average.
- ▶ Confirm the overall gain does not come at the expense of the most at-risk students — eliminating the gap, not just narrowing it, is the test that the design held.
- ▶ Keep human judgment in the loop as the system scales, so prediction continues to deliver more support rather than drifting into automated gatekeeping.

## REFLECTION QUESTIONS

1. What is the difference between GSU's predictive analytics and the algorithmic-bias failure cases in Parts II and VII (e.g., Cases 46, 49, 187, 191)? Be specific about what makes the GSU implementation work.
2. Design the equity-as-primary-constraint version of a predictive system in your domain. What would you measure first?
3. GSU used predictions to deliver more support rather than to gatekeep, with an advisor between the alert and the action. Identify a predictive system in your domain and specify the human-loop architecture that would keep it serving the people it flags rather than screening them out.

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines

# Open University 'Ethical Use of Student Data' and OU Analyse

**G** Governance & Trust**K** Knowledge & Institutional Memory**N** Normalization of Deviance

**IMPACT** The first institutional 'Ethical Use of Student Data' policy in higher education (2014); an eight-principle consent architecture co-designed with students that unblocked predictive analytics across tens of thousands of learners and supported documented intervention work

## IN BRIEF

Predictive learning analytics at the Open University faced a real governance objection: large-scale processing of student data for intervention, with pre-GDPR scrutiny and credible surveillance concerns. The OU's response in 2014 was to author the first institutional "Ethical Use of Student Data for Learning Analytics" policy in higher education — eight principles built through wide stakeholder consultation including students, framing students as participants rather than data subjects. The consent-and- transparency architecture was the enabling engineering, not a compliance afterthought; the deployment followed because trust was established first. OU Analyse — a weekly machine-learning at-risk prediction system — operated on top of that architecture, and the Analytics4Action framework (Rienties et al., JIME 2016) paired predictions with tutor judgment and documented interventions across 18 large-scale modules carrying 42,848 registered students. A 2019 evaluation (Herodotou et al., BJET) of 559 teachers — 189 of them given access to OU Analyse — covering more than 14,000 students in 15 undergraduate courses examined how degree of system usage related to outcomes. The honest open question, raised by the OU's own researchers, is whether predictive analytics genuinely serves students versus surveils them — a tension that remains contested, and the policy has since been superseded twice — by a broader Data Ethics Policy, and that in turn by a Responsible AI Policy. Governance as a living artifact, not a solved problem.

## BACKGROUND

The Open University runs higher education at distance scale — cohorts of tens of thousands across a single module, with no classroom signal that a student is falling behind. Predictive learning analytics on the engagement and assessment trace can, in principle, surface that signal in time for a tutor to act. The governance objection in the early 2010s was not abstract: the data volumes were large, the scrutiny pre-GDPR but tightening, and the surveillance concern — that institutional analytics treat students as objects rather than partners — was credible enough that Slade and Prinsloo set it out in 2013 as an unresolved dilemma for the field.<sup>o</sup>

## THE INTERVENTION

The OU's intervention in 2014 was to author the "Ethical Use of Student Data for Learning Analytics" policy, the first such policy in higher education. The development process was the teaching point as much as the document was: wide stakeholder consultation including students, eight principles that framed students as participants in the analytics rather than its subjects, and a published artifact the institution could be held to. The policy was the **enabling engineering** for analytics deployment, not a compliance overlay applied to a deployment that was happening anyway.<sup>1</sup>

## HOW IT WORKED

OU Analyse — a weekly machine-learning at-risk prediction system — operated on top of that architecture. The Analytics4Action framework (Rienties et al., JIME 2016) paired the prediction with tutor judgment rather than treating the model output as a decision: tutors reviewed flagged students, made the intervention call, and documented what action they took. The framework was tested across 18 introductory modules carrying 42,848 registered students in 2014/15, illustrated by a health and social care module of 3,000-plus. A 2019 evaluation (Herodotou et al., BJET) of 559 teachers, 189 of whom were given access to OU Analyse, covering more than 14,000 students in 15 undergraduate courses, examined how degree of teacher engagement with the system related to outcomes; the benefit concentrated in teachers who made average use — the dashboard opened in 10 to 40 percent of a course's weeks — and who intervened with flagged students.<sup>2</sup>

## THE EVIDENCE

The honest open question, raised by the OU's own researchers, is whether predictive learning analytics genuinely serves students versus surveils them. The OU's policy did not foreclose that question; it made the deployment legible enough to argue about, and it superseded the 2014 policy with a broader Data Ethics Policy — itself since replaced by a Responsible AI Policy — as practice and the scrutiny regime evolved. Governance as a living artifact is the case. The capability deliverable is not a final answer to the surveillance question — there isn't one — but a consent-and-transparency architecture that lets the institution decide and the students participate in the deciding.<sup>3</sup>

## WHAT TRANSFERRED

What OU teaches in the pair (Cases 81 + 80) is the governance-objection diagnostic: when the objection is about trust and accountability — as it was at the OU — good design can dissolve it, and the deployment can proceed under a credibly co-authored consent architecture. The pair's contrast case is the Dutch SyRI (Case 189), where the governance objection was correct: the system was both rights-violating and operationally ineffective, and the District Court of The Hague stopped it on Article 8 ECHR grounds. The diagnostic capability is to tell those two situations apart before deployment, not after.

*The consent architecture is the enabling engineering. The deployment is what follows once trust is established.*

EDITORS' SYNTHESIS OF THE OU 2014 POLICY AND SLADE & PRINSLOO (2013).

## REFERENCES

1. Slade & Prinsloo (2013), "Learning Analytics: Ethical Issues and Dilemmas," *American Behavioral Scientist* 57(10):1510–1529, doi:10.1177/0002764213479366.
2. Open University (2014), "Ethical Use of Student Data for Learning Analytics" — first institutional policy of its kind in higher education.
3. Herodotou, Hlosta, Boroowa, Rienties, Zdrahal, & Mangafa (2019), "Empowering online teachers through predictive learning analytics," *British Journal of Educational Technology* 50(6):3064–3079, doi:10.1111/bjet.12853 — OU Analyse evaluation across 559 teachers (189 with OUA access) and 14,000+ students in 15 undergraduate courses; average-use teachers benefited students most.
4. Rienties, Boroowa, Cross, Kubiak, Mayles, & Murphy (2016), "Analytics4Action Evaluation Framework: A Review of Evidence-Based Learning Analytics Interventions at the Open University UK," *Journal of Interactive Media in Education* 2016(1):2, doi:10.5334/jime.394.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The Open University case is the cleanest instance in the sweep of a governance objection dissolved by design. The 2014 ethical-use policy was the enabling engineering; OU Analyse and Analytics4Action operated on top of it; the 2019 evaluation showed average-use teachers who intervened benefited their students most. The open question — serve vs. surveil — remains contested, and the policy is a living artifact, not a solved problem.

### LENS APPROACH

OU is the positive Domain 4 / Problem Type 5 stakeholder governance case (induced 5.1; LENS D4/PT5). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the governance-objection diagnostic — the dissolvable-objection side — and in Domain 4 (Test and Evaluation) for the average-use-teacher outcome evidence. Direct pair with Case 189 (Dutch SyRI), where the governance objection was correct and design could not have dissolved it. Adjacent to Case 5 (Epic Sepsis Model) as the inverse pattern: OU built consent before deployment; Epic deployed without validation, and the objection that should have been raised wasn't.

| DOMAIN             | MODES ADDRESSED | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|--------------------|-----------------|---------------------|--------------|---------|-------|
| HIGHER ED (UK)     | G K N           | D1 D2 D3 D4 D5      | 5            | 5.1     | 4 • 5 |
| LEARNING ANALYTICS | →               | Test and Evaluation |              |         |       |
| DATA GOVERNANCE    |                 |                     |              |         |       |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Build the consent-and-transparency architecture as enabling engineering, not as compliance afterthought; the document is the artifact the deployment can be held to.
- ▶ Co-author the policy with the people the analytics will be applied to — students or operators — including their voice in the principles before the system is built.
- ▶ Pair predictions with human judgment by design: the prediction surfaces a candidate; the tutor or operator makes the decision and documents what they did.

#### IN OPERATION / AFTER

- ▶ Evaluate the system on degree of use and intervention quality, not adoption alone; at the OU it was average use plus intervening that tracked with student outcome.
- ▶ Treat the governance policy as a living artifact: when the scrutiny regime or practice evolves, supersede the policy openly rather than letting it fall into disuse.
- ▶ Keep the open question — does this genuinely serve students or surveil them? — visible in the institutional record; the legitimacy of the program depends on the question staying askable.



## REFLECTION QUESTIONS

1. Identify a deployment in your domain that faced a governance objection about trust and accountability. Was the objection dissolved by design — a consent or transparency architecture co-authored with the affected parties — or was it managed as compliance overlay? What did the difference cost?
2. Specify the engagement metric you would track, separately from adoption, to know whether your deployment is producing intervention quality. At the OU it was teacher engagement; what is the analog in your context?
3. The OU's open question — does predictive analytics genuinely serve students or surveil them — remains contested, and the policy has been superseded. What is the institutional discipline that keeps the question visible and the policy a living artifact rather than a one-time document?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

## CASE 82

HIGHER EDUCATION (BRAZIL)

LEARNING ANALYTICS

INSTITUTIONAL GOVERNANCE

2024

# MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil)

**K** Knowledge & Institutional Memory    **N** Normalization of Deviance

**IMPACT** MMALA is a maturity model spanning five categories — data management, administration and training, pedagogical support, data analysis, and legislation/privacy/ethics (16 process areas, four maturity levels — ad hoc, initial, structured, systematic); 90 percent of 13 expert evaluators judged it comprehensive, consistent, and adequate; a pilot at three Brazilian universities self-assessed their maturity across all 16 areas on a 0-to-4 scale, where 0 means the activity is not performed at all — instrument for responsible adoption, downstream learning outcome open

## IN BRIEF

Freitas, Fonseca, Garcia, Pontual Falcão, Marques, Gasevic, and Mello (Journal of Learning Analytics, 2024) developed and validated MMALA — a Maturity Model for Adopting Learning Analytics — designed to let an institution self-assess its readiness across the five categories responsible adoption actually depends on: data management, administration and training, pedagogical support, data analysis, and legislation, privacy and ethics — resolved into 16 process areas and four maturity levels: ad hoc, initial, structured, systematic. Thirteen learning-analytics experts evaluated the model, and about 90 percent judged it comprehensive, consistent, and adequate to support adoption; a pilot at three Brazilian universities then used MMALA to self-assess their maturity across all 16 process areas. The case is one of the corpus's clearest worked examples of governance-as-instrument: a structured artifact an institution can use to convert the abstract goal "we should adopt learning analytics responsibly" into specific assessments of where its current capability sits and what it has to build next. The honest limit preserved verbatim: the instrument is validated by expert opinion and a three-institution exercise — one senior respondent per institution, self-reporting — not by longitudinal outcomes of institutions that used it to adopt LA; it is an instrument for responsible adoption, with the downstream effect on student learning still to be measured. The case pairs with the OU policy (Case 81) and the LALA CANVAS (Case 91) as the institutional-instrument layer of the non-US LA governance set.

## BACKGROUND

Institutional adoption of learning analytics typically collapses into one of two patterns: a top-down deployment that meets resistance because the institution is not yet ready (governance is thin, staff are not analytically trained, the pedagogical integration was not designed), or a stalled aspiration where leadership wants responsible adoption but does

not have a structured way to know what “ready” means. MMALA was built to address the second pattern: to give an institution a structured self-assessment instrument across the dimensions readiness actually has, so the adoption decision becomes specific rather than aspirational.<sup>0</sup>

#### THE INTERVENTION

The instrument is dimensional. Data management asks whether collection, storage and stewardship are in place at the level the planned adoption requires. Data analysis asks whether the institution has the analytical capability and staffing the tools will need. Pedagogical support asks whether the analytics are tied to instructional decisions rather than reporting, and whether they integrate with the institution’s learning design or would be bolted on as a separate track. Administration and training asks whether leadership, funding, communication and staff training are in place. Legislation, privacy and ethics is a single category — the model merges them, because compliance with law does not resolve into separate levels — and asks whether the governance and consent architecture has been built. Each process area is resolved into four maturity levels — ad hoc, initial, structured, systematic — so the self-assessment yields a structured picture of which dimensions are ready and which are not, rather than a single overall readiness score.<sup>1</sup>

#### HOW IT WORKED

The Freitas et al. validation has two halves. The first is expert evaluation: thirteen learning-analytics experts rated the model’s process areas, and about 90 percent judged it comprehensive, consistent, and adequate to support adoption. The second is a pilot at three universities in the Brazilian state of Pernambuco, where one senior respondent at each used MMALA to self-assess maturity across all 16 process areas — scored 0 (activity not performed) to 4 — and returned three sharply different profiles. The instrument was usable in practice and produced actionable structure for the institutions that piloted it.<sup>2</sup>

#### THE EVIDENCE

The honest limit survives verbatim. MMALA is validated by expert opinion and a three-institution exercise, not by longitudinal outcomes of institutions that used it to adopt LA. The authors hold the pilot to its size themselves: three institutions cannot characterize Brazilian higher education, and a single respondent per institution may not have known of every local LA initiative — they instructed respondents to consult colleagues in other units to reduce that risk. The instrument’s claim is to be a suitable structure for responsible adoption — what the published evidence supports — and the downstream effect on student learning is the next study, not this one. The case is teachable on the instrument and the validation method; it is honest that the causal chain from “institution adopted via MMALA” to “students learned more” is not yet closed by evidence.<sup>3</sup>

#### WHAT TRANSFERRED

In the non-US LA governance set — OU (Case 81, policy and operating system), LALA (Case 91, participatory adoption framework), African data privacy (Case 89, cross-regime seam), Norway (Case 92, national commission), and MMALA (this case, institutional maturity model) — the cases together teach the level structure of governance for the same underlying capability. MMALA’s specific contribution is the institutional-instrument layer: when a single institution is the unit of adoption and the governance artifact has to be usable by the

institution itself, the maturity model is the form. The pentad demonstrates that governance is producible at multiple levels for the same capability, and the choice of level is itself a governance decision.

*The instrument converts “we should adopt responsibly” into a structured per-dimension self-assessment. The instrument is validated; the downstream learning outcome is the next study.*

EDITORS’ SYNTHESIS OF FREITAS ET AL. (2024).

REFERENCES

1. Freitas, Fonseca, Garcia, Pontual Falcão, Marques, Gašević, & Ferreira Mello (2024), “MMALA: Developing and Evaluating a Maturity Model for Adopting Learning Analytics,” *Journal of Learning Analytics* 11(1):67–86, doi:10.18608/jla.2024.8099. The first version had 18 process areas across the same five categories; consolidation plus one added area (“result-based intervention”) produced the 16 of record.

2. Open University Ethical Use of Student Data policy (2014) — institutional-policy companion (Case 81).

3. Hilliger et al. (2020), Internet and Higher Education — LALA participatory adoption companion (Case 91).

4. Norwegian Expert Commission on Learning Analytics, final NOU (2023) — national-scale companion (Case 92).

THE LEARNING ENGINEERING LENS

LE INSIGHT

MMALA is the institutional-instrument instance of governance for learning-analytics adoption: a structured maturity model across five categories — data management, administration and training, pedagogical support, data analysis, and legislation, privacy and ethics. The validation evidence — 13-expert evaluation and a three-institution pilot — is what the case claims; the downstream effect on student learning is the next study, and the case is honest about that gap.

LENS APPROACH

MMALA is the institutional-maturity-model case (induced 5.1; LENS D4/PT4). LENS uses it in Domain 4 (Test and Evaluation) for the structured self-assessment instrument and the two-layer validation discipline; in Domain 5 (Navigating Sociotechnical Constraints) for the governance-as-instrument framing; and on the subobjective **Judgment under inadequate evidence** (4.4) for the instrument-validation-vs-outcome-validation tier distinction. Pair with Cases 81, 91, 89, 92 as the non-US LA governance pentad; MMALA is the institutional-instrument layer.

| DOMAIN                    | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|---------------------------|-----------------|-----------------------|--------------|---------|-------|
| HIGHER EDUCATION (BRAZIL) | K N             | D1 D2 D3 D4 D5        | 4            | 5.1     | 4 · 5 |
| LEARNING ANALYTICS        | →               | → Test and Evaluation |              |         |       |
| INSTITUTIONAL GOVERNANCE  |                 |                       |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Resolve readiness into the categories adoption actually depends on — data management, administration and training, pedagogical support, data analysis, and legislation, privacy and ethics — rather than a single overall score; the per-category structure is what makes the instrument actionable.

IN OPERATION / AFTER

- Report the evidence at its tier: expert opinion and three-institution validation establish the instrument; longitudinal outcome evidence on institutions that used it is the next study, and the case names the gap rather than glossing it.

- ▶ Build maturity levels within each process area — MMALA uses four, ad hoc through systematic — so the self-assessment yields a structured next-step picture rather than a binary ready / not-ready.
- ▶ Validate the instrument in two layers: expert evaluation for comprehensiveness and suitability, and institutional pilot for usability — both are necessary, and the case carries both.
- ▶ Commission follow-up studies on the institutions that pilot MMALA to convert the instrument-validation evidence into adoption-outcome evidence over time.
- ▶ Carry the level-of-governance lesson into the broader v2 governance set: institutional-instrument, multi-country participatory, national commission, and judicial review are different governance modes for the same capability; MMALA is the institutional-instrument instance.

### REFLECTION QUESTIONS

1. Identify a capability your institution is considering adopting where readiness is currently discussed as a single overall question. What are the dimensions readiness actually has, and what maturity levels within each dimension would yield a structured next-step picture rather than a binary judgment?
2. Specify the two-layer validation you would build for the instrument: expert evaluation for comprehensiveness and suitability, and institutional pilot for usability. Where would each layer be at risk of being skipped, and what would that skip cost the instrument's credibility?

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

# Brinkerhoff Success Case Method — Tails as the Evaluation Instrument

**K** Knowledge & Institutional Memory    **N** Normalization of Deviance

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** When ROI-style evaluation of corporate training is intractable, Brinkerhoff's Success Case Method samples the tails of the outcome distribution — the highest- and lowest-impact participants — and reconstructs the system conditions that made the program work for some and fail for others; deployed at Cargill, Ford, Merck, the World Bank, and the Red Cross and Red Crescent

## IN BRIEF

The Success Case Method (SCM), introduced by Robert Brinkerhoff in his 2003 book and formalized in a 2005 peer-reviewed paper, deliberately samples the tails of a training program's outcome distribution. Instead of attempting to derive an average effect that flatters most programs and gives L&D nothing actionable, SCM identifies the highest- and lowest-impact participants, studies them in detail, and reconstructs the system conditions that made the program work for some and fail for others. The method's argument — and the one that places it inside the LENS framework — is that extreme cases reveal whether the program ever produces meaningful work-performance change and what the surrounding system has to provide for transfer to happen. The method is peer-reviewed in **Advances in Developing Human Resources**; the corporate deployments at Cargill, Ford, Merck, World Bank, and the Red Cross and Red Crescent are documented in practitioner channels — case-study writeups, conference talks, vendor whitepapers — rather than in peer-reviewed evaluation literature. SCM is the operational answer to the chain-of-evidence problem named in Case 79 (Kirkpatrick): the practitioner instrument that crosses the Level-2 / Level-3 seam by sampling where the evidence is most informative. Evidence-tier flag is practice-synthesis; future validation will continue as more firms publish their SCM outcome data.

## BACKGROUND

Corporate L&D evaluation faces a structural problem the Kirkpatrick chain-of-evidence pattern names (Case 79): Level 3 (behavior change on the job) and Level 4 (business results) require longitudinal data the training organization typically cannot access. Average-effect statistics across whole programs flatter most programs, on Brinkerhoff's premise that the distribution of training outcomes is heavily skewed — a minority of participants produce most of the work-performance change, and the average buries that signal under the par-

ticipants for whom the program produced nothing. The premise is the method's starting argument rather than an independently established distributional fact; SCM is built to test it program by program.<sup>9</sup>

#### THE INTERVENTION

Robert Brinkerhoff's 2005 **Advances in Developing Human Resources** paper proposes a different sampling logic: do not estimate the average; identify the participants in the tails of the outcome distribution and study them in detail. The Success Case Method identifies the highest-impact participants (where did the program work, and what conditions made the work-performance change possible), the lowest-impact participants (where did the program fail, and what conditions were missing), and reconstructs the system around each.<sup>1</sup>

#### HOW IT WORKED

The method's argument is that the tails carry the decision-grade information. A program that produces meaningful work-performance change for any participants proves it **can** work; the conditions that distinguish the success cases from the failure cases name what the surrounding system has to provide for transfer. The argument is the practical complement of the Blume meta-analytic finding (Case 65) that work environment is the decisive transfer variable: SCM operationalizes the finding by sampling where the variable's effects are most visible.<sup>2</sup>

#### THE EVIDENCE

The Brinkerhoff Evaluation Institute lists deployments at Cargill, Ford, Merck, the World Bank, and the Red Cross and Red Crescent. These deployments are documented in case-study writeups, conference talks, and vendor whitepapers, not in peer-reviewed program evaluations. The evidence-tier flag is therefore practice-synthesis: the method itself is peer-reviewed, but the per-firm impact data on which the corporate adoption story rests live in practitioner channels. The honest framing in print is that the methodological pattern is teachable and durable, but the per-firm effect sizes are not independently audited.<sup>3</sup>

#### WHAT TRANSFERRED

The LENS teaching point pairs SCM with the chain-of-evidence problem from Case 79 and the meta-analytic transfer finding from Case 65. SCM is the operational instrument that crosses the Level-2 / Level-3 seam by sampling where the evidence is most informative — and it exercises the subobjective **Judgment under inadequate evidence** (4.4) directly, because the method asks practitioners to act on detailed case reconstructions rather than wait for population-scale causal estimates that corporate L&D often cannot produce. Future validation will continue as more firms publish their SCM outcome data.

*The tails carry the decision-grade information. The average flatters the program.*

EDITORS' SYNTHESIS OF BRINKERHOFF (2005) AND THE CORPORATE SCM DEPLOYMENTS.

#### REFERENCES

1. Kirkpatrick & Kirkpatrick (2006), Evaluating Training Programs — the chain-of-evidence framework SCM operationalizes (paired Case 79).
2. Brinkerhoff, R. O. (2005), "The Success Case Method: A Strategic Evaluation Approach to Increasing the Value and Effect of Training," *Advances in Developing Human Resources* 7(1):86–101, doi:10.177/1523422304272172 — the peer-reviewed state-
3. Blume, Ford, Baldwin, & Huang (2010), *Journal of Management* 36(4):1065–1105 — the meta-analytic environment-as-finding SCM samples around (paired Case 65).
4. Brinkerhoff Evaluation Institute deployment list — International Red Cross and Red Crescent, World Bank, Cargill, Merck, Ford Motor — practitioner channel.

ment of the method; the book-length treatment is Brinkerhoff (2003), *The Success Case Method: Find Out Quickly What's Working and What's Not* (Berrett-Koehler).

THE LEARNING ENGINEERING LENS

LE INSIGHT

SCM is the practitioner instrument that operationalizes Blume’s environment-as-decisive finding (Case 65) by sampling the tails of the outcome distribution. The method is peer-reviewed; the per-firm impact data at Cargill, Ford, Merck, World Bank and the Red Cross and Red Crescent live in practitioner channels. Evidence-tier flag is practice-synthesis; future validation will continue as more firms publish.

LENS APPROACH

SCM is the corporate-L&D tail-sampling case (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the subobjective **Judgment under inadequate evidence** (4.4) — the method asks practitioners to act on case-reconstruction evidence rather than population estimates — and in Domain 2 (Iterative Development) by way of the evaluation-feedback loop SCM enables. Pairs with Case 70 (HILS) for the redesign of the surrounding work environment.

| DOMAIN              | Modes Addressed | LENS Competency | Problem Type        | Induced | LEO   |
|---------------------|-----------------|-----------------|---------------------|---------|-------|
| CORPORATE L&D       | K N             | D1 D2 D3 D4 D5  | 5                   | 2.1     | 4 · 2 |
| TRAINING EVALUATION | →               | →               | Test and Evaluation |         |       |
| WORKFORCE DEV       |                 |                 |                     |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Before the program is delivered, plan the SCM sampling: define the outcome metric (specific work-performance change), identify how the success and failure cases will be surfaced (manager nominations, performance data, structured interviews), and pre-commit to the sampling logic.
- ▶ Structure the success-case and failure-case interviews around the system conditions Blume’s meta-analysis (Case 65) names as decisive — supervisor support, peer support, practice opportunity, environment — so the reconstruction is theory-grounded, not anecdote-grounded.
- ▶ Report the tails honestly: the success cases are existence proofs the program **can** work; the failure cases are evidence of what the surrounding system did not provide. Neither is a population effect size.

IN OPERATION / AFTER

- ▶ Translate the tail evidence into a redesign of the surrounding system — supervisor briefings, peer-support structures, on-job practice opportunity — rather than into a redesign of the training event alone. Pair with Case 70 (HILS).
- ▶ Use the subobjective **Judgment under inadequate evidence** (4.4): act on the case-reconstruction evidence the method produces while being explicit that it is not a population causal estimate; document the uncertainty.
- ▶ Carry the practice-synthesis flag honestly in any program documentation citing SCM corporate deployments — the method is peer-reviewed, but the per-firm effect sizes are not.



## REFLECTION QUESTIONS

1. Identify a recent training program in your organization. Define how you would identify the highest- and lowest-impact participants and what structured interview you would conduct with each to reconstruct the system conditions around them.
2. Map the SCM tails to the Blume meta-analytic variables (Case 65): supervisor support, peer support, practice opportunity, environment. What pattern would the tails have to show for you to redesign the surrounding system rather than the training event itself?
3. The corporate SCM deployments at Cargill, Ford, Merck, the World Bank and the Red Cross and Red Crescent are documented in practitioner channels. What additional evidence — independent program audit, peer-reviewed evaluation, cross-firm comparison — would you require before treating any specific per-firm effect size as decision-grade?

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

# Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive

**T** Training Gap   **K** Knowledge & Institutional Memory   **D** Designed Out

**IMPACT** Cluster-randomized 147 schools (middle and high schools) across seven states; year-one posttest scores showed no significant difference between Cognitive Tutor and control conditions; year-two posttest scores showed CTAI high schools significantly outperforming controls (the middle-school effect was not statistically significant); a one-year evaluation would have published the wrong answer

## IN BRIEF

Pane, Griffin, McCaffrey, and Karam's 2014 paper in *Educational Evaluation and Policy Analysis* reports the central at-scale evaluation of Cognitive Tutor Algebra I (CTAI), the canonical learning-sciences-to-classroom translation. The RAND team cluster-randomized 147 middle and high schools across seven states to continue their current Algebra I curriculum or to adopt CTAI for two years. Year-one posttest scores: no significant difference between conditions. Year-two posttest scores: CTAI high schools significantly outperformed controls (the middle-school effect was not statistically significant). The honest reading the authors press is that a one-year evaluation would have published a null result and the two-year evaluation surfaced a real effect — and both findings were in the same trial. The case is the deeper-evidence-of update on Case 67 (Cognitive Tutor), translating that case from a system-description case into a methodological-discipline case about evaluation horizons. The timeline of the evaluation is itself a falsifiable design choice; the case grounds the subobjective **Judgment under inadequate evidence** (4.4) where the inadequacy is the evaluation horizon, not the sample size.

## BACKGROUND

Cognitive Tutor is the case the learning-sciences-to-classroom translation thread cannot avoid teaching at v1 Case 67. The system instantiates a model of student cognition, drives a problem-by-problem adaptive curriculum, and is the canonical published instance of an intelligent tutoring system in K-12 mathematics. The question the v1 case left open — and the question the Pane et al. 2014 paper answers — is what happens when the system is deployed at scale, in the classroom and curricular conditions that classroom adoption actually generates, rather than in the developer-supported conditions of an early efficacy trial.<sup>o</sup>

## THE INTERVENTION

The RAND team built the evaluation to support that question. 147 schools — 73 high schools and 74 middle schools, in 51 districts across seven states — were cluster-randomized in matched pairs to either continue their current Algebra I curriculum or to adopt CTAI for two years. The sample was deliberately wide — multiple states, multiple district contexts, classroom adoption with the support structures the publisher could realistically supply. The outcome instrument was an Algebra I posttest applied at end of school year. Each school held its assigned condition for two years and taught two successive cohorts of students under it, so the design supports an effect-size estimate for the year-one cohort and a separate estimate for the year-two cohort in the same schools — not a two-year follow-up of the same students.<sup>1</sup>

## HOW IT WORKED

The year-one result was a null. Posttest scores in CTAI schools did not differ significantly from posttest scores in control schools, across the cluster-randomized sample. The year-two result was different. Posttest scores in CTAI high schools were significantly higher than in control schools (the middle-school estimate was of similar size but not statistically significant) — about 0.20 SD, or roughly eight percentile points for the median student. The two findings were generated by the same trial, run in the same schools, with the same instrument, on two successive cohorts of students. The methodological discipline the case teaches turns on what the field would have learned if the evaluation had been designed to a one-year horizon: it would have published a null result against an intervention that the same trial, on its planned horizon, established as positive.<sup>2</sup>

## THE EVIDENCE

The authors press the honest reading without softening it. The timeline of the evaluation is itself a falsifiable design choice, not a methodological default. The reason the analysis was run separately by cohort is on the record: the researchers hypothesized that implementation would be better in the second year as teachers gained experience with the curriculum. What the implementation record does not show is a matching rise in fidelity — the year-two gain tracked teachers using fewer of the activities the guidelines prescribe for non-computer class time, not more. The two-year horizon was the right horizon because the effect appeared only in the second year; what changed between the two years is still contested in the trial's own implementation record.<sup>3</sup>

## WHAT TRANSFERRED

The case is the deeper-evidence-of update on Case 67. Case 67 carries Cognitive Tutor as the system description; this case carries the at-scale evaluation and the timeline-of-evidence lesson. The subobjective **Judgment under inadequate evidence** (4.4) is directly motivated: the one-year evaluation would have been inadequate evidence, and a field that publishes the one-year null without noting the planned second-year horizon publishes a wrong answer on the available data. Pair with Case 72 (ASSISTments) for the multi-year follow-through paired case, and with Case 5 (Epic Sepsis) for the evaluation-horizon discipline in clinical AI. The closed loop the case completes is the two-year-horizon-was-the-right-horizon record that lets a field design the next at-scale evaluation honestly.

A decade on, the record carries a second, sharper lesson the original evaluation-horizon reading did not anticipate: that single year-two effect has never been independently reproduced by a comparably rigorous trial, yet it now bears almost the entire evidentiary weight the product carries. Every “Strong”/“Positive” badge traces back to this trial and two pre-2015 studies, and only one registry issues one: Evidence for ESSA rates the program Strong, while the What Works Clearinghouse, reviewing the high-school arm in June 2016, adjusted the year-two difference for multiple comparisons, found it no longer statistically significant, and rated Cognitive Tutor Algebra I mixed effects on algebra; Evidence for ESSA’s own pooled estimate across the qualifying studies is a mere +0.04 SD, far below the 0.20 SD headline. Kraft’s benchmarks put the median of 747 education RCTs at 0.10 SD and call 0.20 SD large rather than exceptional; the sharper comparison is his median of 0.03 SD for trials enrolling more than 2,000 students. And the trial’s own implementation record, as the NBER ed-tech review reports it (Escueta et al., 2017), ties the year-two gain not to higher fidelity but to teachers reducing — though not eliminating — their use of the activities the CTAI guidelines prescribe for non-computer class time. The case’s warning thus deepens: the danger is not only the one-year evaluation that would have published the wrong null, but the single positive replication-year result scaled into a durable “evidence-based” claim without the fresh independent confirmation that never came.

*A one-year evaluation would have published a null. The two-year evaluation surfaced a real effect. Both findings were in the same trial.*

EDITORS’ SYNTHESIS OF PANE ET AL. (2014).

## REFERENCES

1. Evidence for ESSA, program record for Carnegie Learning High School Math Solution Algebra I — ESSA rating “Strong,” weighted mean effect size +0.04 across three qualifying studies (Cabalo & Vu 2007; Campuzano et al. 2009; Pane et al. 2014).
2. Pane, J. F., Griffin, B. A., McCaffrey, D. F., & Karam, R. (2014), “Effectiveness of Cognitive Tutor Algebra I at Scale,” *Educational Evaluation and Policy Analysis* 36(2):127–144, doi:10.3102/O162373713507480.
3. Pane, Griffin, McCaffrey, & Karam (2014), Addendum to Effectiveness of Cognitive Tutor Algebra I at Scale, RAND Working Paper WR-1050-DEIES — the addendum the WWC reviewed alongside the published paper.
4. Koedinger, K. R., Anderson, J. R., Hadley, W. H., & Mark, M. A. (1997), “Intelligent tutoring goes to school in the big city,” *IJAIED* — the Case 67 system description Cognitive Tutor builds from.
5. What Works Clearinghouse (2016), Cognitive Tutor® intervention report, Secondary Mathematics, June 2016 — rates Cognitive Tutor Algebra I “mixed effects” on algebra; records the Pane et al. year-two difference as no longer significant after adjustment for multiple comparisons.
6. Escueta, Quan, Nickow, & Oreopoulos (2017), “Education Technology: An Evidence-Based Review,” NBER Working Paper 23744 — reports that the Pane year-two gain tracked reduced use of the prescribed non-computer activities, not higher fidelity; published as Escueta, Nickow, Oreopoulos, & Quan (2020), “Upgrading Education with Technology: Insights from Experimental Research,” *Journal of Economic Literature* 58(4):897–996.
7. Kraft, M. A. (2020), “Interpreting Effect Sizes of Education Interventions,” *Educational Researcher* 49(4):241–253 — 1,942 effect sizes from 747 RCTs, median 0.10 SD, 0.03 SD for samples above 2,000 students, and the benchmark that treats 0.20 SD as large.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Pane et al. is the at-scale evaluation case for Cognitive Tutor and the worked example of evaluation horizon as a falsifiable design choice. Year one: null. Year two: significantly positive. The case is the deeper-evidence-of update on Case 67 and the curriculum’s primary anchor for the subobjective **Judgment under inadequate evidence** (4.4) where the inadequacy is the horizon, not the sample size.

### LENS APPROACH

Pane / CTAI at scale is the closed-loop evaluation-horizon case (induced 2.3; LENS D4/PT4). LENS uses it in Domain 4 (Test and Evaluation) for the horizon-as-design-choice discipline and in Domain 2 (Iterative Development) for the deployment-substrate-stabilization frame — teacher fluency, scheduling, workflow familiarity are designable parts of the intervention substrate, not noise. Pair with Case 72 (ASSISTments follow-through), Case 67 (Cognitive Tutor system), and Case 5 (Epic Sepsis horizon discipline).

| DOMAIN               | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO   |
|----------------------|-----------------|-----------------------|--------------|---------|-------|
| K-12 MATHEMATICS     | T K D           | D1 D2 D3 D4 D5        | 4            | 2.3     | 2 · 4 |
| INTELLIGENT TUTORING | →               | → Test and Evaluation |              |         |       |
| EDUCATION POLICY     |                 |                       |              |         |       |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Plan the evaluation horizon as a design decision, not a budget default; the case demonstrates that the horizon choice is causally consequential and that a shorter horizon can produce a wrong-direction finding.
- ▶ Document the deployment-substrate stabilization that the intervention depends on: teacher fluency, scheduling adjustment, student familiarity with the workflow are all year-over-year stabilization mechanisms the trial should track.
- ▶ Design the trial so the year-one and year-two estimates are both interpretable; the case's pedagogical value depends on having both estimates from the same trial, not on having only the longer-horizon estimate.

#### IN OPERATION / AFTER

- ▶ Publish the year-by-year result pattern, not the aggregate effect; the case's teaching power depends on the year-one null being on the same page as the year-two positive.
- ▶ Carry the timeline-of-evidence lesson into the curriculum's evaluation-design teaching; the case grounds the subobjective **Judgment under inadequate evidence** (4.4) with a rare worked example where the inadequacy is the horizon, not the sample size.
- ▶ Pair the case with Case 72 (ASSISTments multi-year follow-through) so the field-level discipline is taught with two converging examples; the methodological lesson is more useful as a pattern than as a single instance.

### REFLECTION QUESTIONS

1. Identify an intervention in your domain whose effect depends on a deployment-substrate stabilization that takes more than a year (teacher fluency, workflow familiarity, scheduling alignment). What evaluation horizon is honest for that intervention, and what would the year-one estimate, taken alone, falsely conclude?
2. Specify the year-by-year result pattern your next evaluation will report. The case's pedagogical value depends on having both the year-one and year-two estimates from the same trial; what reporting structure would you commit to that supports the pattern view, not only the aggregate?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

# OU Analyse — Predictive Learning Analytics and Teacher Use at Scale

**T** Training Gap   **K** Knowledge & Institutional Memory   **D** Designed Out

**IMPACT** Across 15 undergraduate courses, 559 teachers (189 with OU Analyse access) and more than 14,000 students, teachers who accessed the dashboard in 10–40 percent of course weeks and intervened with flagged students had students who significantly outperformed their own students of the previous year; a survey of 366 teachers three years post-implementation finds only 42 percent still active users, and maps what drives adoption and what blocks it

## IN BRIEF

Herodotou et al.'s 2019 British Journal of Educational Technology paper reports the evaluation of OU Analyse, the Open University UK's predictive-learning-analytics dashboard for online tutors, across 15 courses and 559 teachers (189 of whom had access to OU Analyse) with more than 14,000 students. Teachers' engagement with the predictive learning analytics was associated with measurable improvements in student performance. The LAK 2023 study surveys 366 teachers across four faculties three years after the dashboard became business as usual, and models adoption with UTAUT: use rises with performance expectancy, self-efficacy, positive attitudes and low anxiety, and falls where facilitating conditions are absent. The case is distinct from the OU consent-and-ethical-use frame Case 81 covers; this case carries the post-deployment teacher-use evaluation at multi-cohort scale. The authors' hedges are binding: causal attribution to OU Analyse use specifically — versus teacher selection effects — is bounded; the 2019 study compares a teacher's students against that same teacher's students in the prior presentation, controlling for academic performance, rather than randomizing. Future validation ongoing on multi-institution transfer: the system runs inside the Open University, and no evaluation of it at another institution is in the peer-reviewed literature.

## BACKGROUND

Predictive learning analytics dashboards have been deployed across higher-education institutions for more than a decade. The deployment record is largely a collection of pilots and single-cohort studies, with Course Signals (Case 52, the discontinued Purdue deployment) as the structural cautionary case. OU Analyse is structurally different: deployed across an entire distance-learning institution, used by hundreds of tutors across multiple cohorts and multiple courses, and evaluated at the scale the institutional deployment supports. The 2019 Herodotou et al. paper is the central peer-reviewed evaluation of the teacher-use side of the deployment.<sup>o</sup>

## THE INTERVENTION

The evaluation design covers 15 courses, 559 teachers, and more than 14,000 students. 189 of the 559 teachers had access to OU Analyse; the comparison structure sets a teacher's students against the students that same teacher taught in the previous year's presentation of the course, when no predictions were available. The outcome is student performance on course assessments. Engagement was operationalized as how many of a course's weeks the teacher accessed the dashboard, together with whether the teacher intervened with flagged students. The gain is not monotonic in usage: teachers who accessed OU Analyse in between 10 and 40 percent of the weeks a course ran, and intervened, benefited their students the most.<sup>1</sup>

## HOW IT WORKED

The 2023 LAK three-year-post-implementation follow-up extends the picture into a question the single-cohort pilots cannot address: what happens to teacher use of the predictions across multiple cohorts and across years? The 2023 paper answers it with a survey of 366 teachers across four faculties, three years after the dashboard became business as usual, read through the Unified Theory of Acceptance and Use of Technology. Adoption rises with performance expectancy, self-efficacy, positive attitudes, low anxiety and less teaching experience, and falls where facilitating conditions are absent and the tool is felt to be effortful. The usage split is the sharper finding, and it is not a success story: three years into business as usual, 42 percent of the surveyed teachers were active users, 20 percent had tried the dashboards and stopped, 18 percent had heard of them and never used them, and 19 percent had not heard of them at all. The delegation-with-revocation structure is operative: each teacher chose, in each cohort, whether to act on the prediction, and the choice itself is the decisive capability the dashboard supports.<sup>2</sup>

## THE EVIDENCE

The case is distinct from the OU consent-and-ethical-use frame Case 81 covers. Case 81 carries the governance story — the Policy on Ethical Use of Student Data, the institutional review of analytic deployments, the consent architecture. This case carries the post-deployment teacher-use evaluation at multi-cohort scale. The editor's decision (memo A6) anchors LENS primary, induced secondary, LEO carried; the two cases together teach the deployment across both governance and post-deployment-use frames without collapsing one into the other. The frame the present case adds to the corpus is the rare successful learning-engineering intervention with both deployment scale and longitudinal teacher-use evidence at the journal-graded evidence tier the corpus needs.<sup>3</sup>

## WHAT TRANSFERRED

The authors' hedges are binding. The causal attribution to OU Analyse use specifically — as distinct from teacher selection effects — is bounded. The 2019 study compares a teacher's students with that same teacher's students in the previous presentation, controlling for prior academic performance, and does not randomize at the teacher level; the 2023 follow-up extends the evidence base on adoption and perception but does not close the teacher-selection question. Future validation ongoing on multi-institution transfer — the Open University's own project page describes OU Analyse as available to its staff, and no evaluation of the system at another institution is in the peer-reviewed literature. Pair with Case 81 (OU consent governance), Case 84 (Cognitive Tutor at-scale evaluation), and Case 52



(Purdue Course Signals) — the OU Analyse evaluation is the methodologically more careful descendant of the Course Signals lineage.

*The delegation-with-revocation structure is operative: each teacher chose, in each cohort, whether to act on the prediction. The choice itself is the capability the dashboard supports.*

EDITORS' SYNTHESIS OF HERODOTOU ET AL. (2019, 2023).

## REFERENCES

1. Herodotou, C., Hlosta, M., Boroowa, A., Rienties, B., Zdrahal, Z., & Mangafa, C. (2019), "Empowering online teachers through predictive learning analytics," *British Journal of Educational Technology* 50(6):3064–3079, doi:10.1111/bjet.12853.
2. Herodotou, C., Maguire, C., Hlosta, M., & Mulholland, P. (2023), "Predictive Learning Analytics and University Teachers: Usage and perceptions three years post implementation," *LAK '23*, pp. 68–78, doi:10.1145/3576050.3576061 — survey of 366 teachers across four faculties; 42% active users, 20% tried and stopped, 18% aware but unused, 19% unaware; UTAUT model of adoption.
3. Herodotou, Rienties, Boroowa, Zdrahal, & Hlosta (2019), "A large-scale implementation of predictive learning analytics in higher education: the teachers' role and perspective," *Educational Technology Research and Development* 67(5):1273–1306, doi:10.1007/s11423-019-09685-0 — the complementary teacher-perspective study (59 teachers, nine courses, 1,325 students).
4. Arnold, K. E., & Pistilli, M. D. (2012), "Course Signals at Purdue: Using Learning Analytics to Increase Student Success," *LAK '12* — the structural precursor the corpus carries as a cautionary case at Case 52.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

OU Analyse is the rare successful learning-engineering intervention with both deployment scale and longitudinal teacher-use evidence at journal tier. The 2019 BJET paper establishes the effect on student performance at the institutional deployment; the 2023 LAK survey of 366 teachers finds 42 percent still active users three years on, with a fifth having tried the dashboards and stopped. The delegation-with-revocation structure is operative and teachable. Hedges binding on causal attribution and on multi-institution transfer.

### LENS APPROACH

OU Analyse teacher-use is the human-in-the-loop predictive- analytics case at institutional scale (induced 5.2; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the multi-cohort longitudinal-follow-through design and in Domain 3 (Human-System Collaboration) for the delegation-with-revocation structure — the teacher chooses, each time, whether to act on the prediction. Pair with Case 81 (OU consent governance), Case 84 (Cognitive Tutor at-scale), and Case 72 (ASSISTments multi-year).

| DOMAIN             | MODES ADDRESSED | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|--------------------|-----------------|---------------------|--------------|---------|-------|
| HIGHER EDUCATION   | T K D           | D1 D2 D3 D4 D5      | 5            | 5.2     | 4 • 3 |
| LEARNING ANALYTICS | →               | →                   |              |         |       |
| DISTANCE LEARNING  |                 | Test and Evaluation |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Treat teacher use of the prediction as the capability the dashboard supports; the delegation-with-revocation structure — each teacher chooses, in each cohort, whether to act — is the design point, not an emergent behavior.
- ▶ Build the deployment at multi-cohort and multi-course scale from the start; the case's evidence-tier strength depends on the deployment having been

### IN OPERATION / AFTER

- ▶ Commission the next-tier evaluation that addresses the teacher-selection-effect question; an RCT at the teacher level, or an instrumental-variable design exploiting institutional dashboard-rollout variation, would close the residual causal-attribution gap.
- ▶ Publish multi-institution-transfer evaluations if the system moves beyond the OU; the case's value as a



- institutional, not pilot, and the institutional deployment supports the longitudinal-follow-through evidence the case provides.
- ▶ Design the dashboard's prediction surface around the action the teacher would take — a flagged-student list with an intervention prompt, not a probability score the teacher has to interpret unaided.
  - ▶ Pair the case in the curriculum with Case 81 (OU consent governance) so the deployment is taught across both governance and post-deployment-use frames; the two cases together teach the institutional anchor without collapsing one frame into the other.

### REFLECTION QUESTIONS

1. Identify a predictive-analytics dashboard in your domain whose operator-use stratification — active, background reference, non-engagement — has not been longitudinally tracked. What would the multi-cohort follow-up look like, and what data infrastructure would it require?
2. Specify the design choice for the prediction surface in your domain: a flagged-list-with-intervention-prompt vs. a probability-score-the-operator-interprets-unaided. Which design is in place currently, and what would change if the design switched?
3. The case's hedge on teacher selection effects is the open methodological question. What randomized or quasi-experimental design would close the residual attribution gap in your domain's analog deployment — and what institutional support would the design require?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

# Gándara — Detecting and Mitigating Algorithmic Bias in College Student–Success Prediction

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Gándara, Anahideh, Ison, and Picchiarini (AERA Open, 2024) audited predictive models of college student success and showed that models which look acceptable on overall accuracy are systematically less accurate for Black and Hispanic students and overestimate success for white and Asian students — small-tier frontier evidence that overall accuracy hides the disparity, that the four leading bias-mitigation techniques tested are generally ineffective at eliminating it, and that the evaluation regime and end-user training, not the model alone, determine whether an equity-oriented prediction is fair to the groups the equity commitment is meant to protect

## IN BRIEF

Predictive models of college student success — models that score students on predicted graduation, retention, or course completion to drive advising, outreach, and support decisions — have become a routine tool across community colleges and four-year institutions. Gándara, Anahideh, Ison, and Picchiarini, publishing in AERA Open (2024), audited such models across racialized groups and found that a model which looks acceptable on overall accuracy is systematically less accurate for Black and Hispanic students — making more prediction errors for them — while overestimating success for white and Asian students. The models run on the nationally representative Education Longitudinal Study of 2002 and predict bachelor’s-degree attainment; the disparity surfaces only when the evaluation is stratified by racialized group, and the four bias-mitigation techniques the authors test are, in their own words, “generally ineffective at eliminating disparities in prediction outcomes” — in-processing techniques do better than preprocessing ones, and reweighting the successful students from under-represented groups barely moves the disparity at all. What the prediction then does to a student depends on the decision context it is read in — whether end users are trained to contextualize a flagged prediction or treat it as a verdict. The paper’s contribution is the frontier-shaped finding that stratified evaluation, not aggregate accuracy, is what makes an equity-oriented prediction auditable at all. The case pairs explicitly with the v2 race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman) and with the broader equity-construct competency C8.2: demographic stratification of validation and outcomes as a design commitment.

## THE SHIFT

Predictive modeling of student success in higher education is now routine. Community-college and four-year institutions operate models that score students on predicted graduation likelihood, predicted retention, predicted benefit-from-intervention, or predicted financial need. The scores feed downstream decisions: which students get outreach emails, which get advising appointments, which get need-based aid or enrollment incentives. The structural capability question — whose interest does the model maximize, by what construct, and against what stratification — is the layer where the case operates.<sup>0</sup>

### WHAT IS EMERGING

Gándara, Anahideh, Ison, and Picchiarini, publishing in AERA Open, audited predictive models of college student success across racialized groups. The central finding is that models which look acceptable on overall accuracy are systematically less accurate for Black and Hispanic students — making more prediction errors for them — while overestimating success for white and Asian students. The methodological move is to hold the prediction task fixed — bachelor's-degree attainment in the Education Longitudinal Study of 2002 — and vary the evaluation: four fairness notions across several model classes, scored by racialized group rather than in aggregate. Among students who did not complete a degree, the models predict success for 65 percent of White and 73 percent of Asian students against 28 percent of Hispanic and 33 percent of Black students. Which evaluation an institution runs is a normative decision the model architecture cannot make on its own.<sup>1</sup>

### THE CAPABILITY QUESTION

The second methodological move concerns evaluation stratification. Overall accuracy, the standard summary metric — around 77 percent here — can mask substantial disparity by racialized group, populations whose base rates and observed outcomes differ from the modal student the training data over-represents. The paper's stratified evaluation shows the canonical equity-construct finding: a model that looks fair under overall accuracy can be substantially less accurate, or substantially biased in its prediction direction, for the subgroups the equity commitment is supposed to protect. The induced framework's C8.2 sub-competency — demographic stratification of validation and outcomes as a design commitment — is the analytic anchor.<sup>2</sup>

### EARLY EVIDENCE

The third move concerns the decision context the prediction is consumed in and the mitigations available. A model is not deployed in isolation; its predictions are read by advisors and administrators who decide what to do with a flagged student. Gándara and colleagues stress training end users on the potential for algorithmic bias — so a prediction is contextualized for the individual student rather than treated as a verdict — and evaluate four bias-mitigation techniques, which they find generally ineffective at eliminating the disparity: in-processing methods outperform preprocessing ones, and reweighting under-represented successful students is close to no help — evidence that their under-representation in the training data is not the principal source of the bias. The system's apparent fairness therefore depends not only on the model but on the evaluation regime and the decision practice the institution actually deploys.<sup>3</sup>

OPEN PROBLEMS

What the case teaches at the LENS framing layer is the frontier-shaped finding: fairness in equity-oriented prediction is a construct-definition problem before it is a model-bias problem. The induced framework’s C8.2 sub- competency and the equity-construct sub-objective proposed in the v2 research backbone — fairness beyond omission (4.6) — find their case-grounded basis here. The case explicitly cross-references the v2 race-construct trio — eGFR (Case 25), pulse oximetry (Case 26), and Hoffman pain bias (Case 6) — at the construct-definition layer: in those cases the construct (race correction in eGFR, single-sensor calibration in pulse oximetry, the pain-perception assumption in Hoffman) was the design decision that produced the disparate outcome; in this case the design decision is the evaluation itself — whether accuracy is scored in aggregate or by racialized group, and whether the advisor reading a flagged prediction has been told which way the model errs.

*A model that looks fair on overall accuracy can be substantially less accurate for the students the equity commitment is meant to protect. Stratified evaluation is what makes the difference visible.*

EDITORS’ SYNTHESIS OF GÁNDARA, ANAHIDEH, ISON, & PICCHIARINI, \_AERA OPEN\_ (2024).

REFERENCES

1. Gándara, Anahideh, Ison, & Picchiarini (2024), “Inside the Black Box: Detecting and Mitigating Algorithmic Bias across Racialized Groups in College Student-Success Prediction,” AERA Open 10, doi:10.1177/23328584241258741 — primary case source on cross-group accuracy disparity, stratified evaluation, and bias mitigation in student-success prediction.

2. Barocas, Hardt, & Narayanan, Fairness and Machine Learning (fairmlbook.org) — methodological backdrop on construct definition and stratified evaluation.

3. Friedler, Scheidegger, & Venkatasubramanian (2021), “The (im)possibility of fairness: different value systems require different mechanisms for fair decision making,” Communications of the ACM — the construct-definition argument at field level.

4. v2 cross-referenced cases: 25 (eGFR race correction), 26 (pulse oximetry across skin tones), 6 (Hoffman pain bias) — the race-construct trio at the construct-definition layer.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Gándara’s student-success-prediction audit is the small-tier frontier instance of fairness-as-stratified-evaluation. Which stratification is used in validation, which mitigation technique is applied, and whether the advisor reading the prediction knows how it errs each determine the fairness properties of the deployed system; the audit found models systematically less accurate for Black and Hispanic students. Cross-references the v2 race-construct trio at the construct-definition layer.

LENS APPROACH

Gándara student-success-prediction fairness is the stratified-evaluation case (induced 8.2; LENS D4/PT5) — Domain 4 for stratified evaluation; Domain 5 for the decision-context and mitigation question. Cross-reference Cases 25, 26, 6 — case-grounded basis for the equity- construct subobjective, fairness beyond omission (4.6).

| DOMAIN           | MODES AT STAKE | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|------------------|----------------|---------------------|--------------|---------|-------|
| HIGHER EDUCATION | D K N          | D1 D2 D3 D4 D5      | 5            | 8.2     | 4 · 5 |
| ALGO FAIRNESS    |                | Test and Evaluation |              |         |       |
| POLICY ANALYSIS  |                |                     |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Make the construct choice explicit at design time: predicted enrollment, predicted benefit, predicted need are different constructs with different policy implications, and the institution's normative position on which to optimize for has to be on the record.
- ▶ Build stratified evaluation by income, race/ethnicity, first-generation status, and other equity-relevant axes into the model's validation; the C8.2 sub-competency operationalized.
- ▶ Pair the targeting model with the specific intervention it will feed; the fairness properties of the system depend on the pairing, and evaluating the model in isolation underestimates the variation a deployed system actually exhibits.

### IN OPERATION / AFTER

- ▶ Report which construct the model maximizes, on which stratifications it was evaluated, and which intervention it is paired with — three pieces of information any deployed targeting system should publish together.
- ▶ Treat the fairness conversation as a construct-definition conversation first, and as a model-bias conversation second; the latter is operationally important, the former determines what the model is fair **about**.
- ▶ Carry the race-construct trio (Cases 25, 26, 6) as the cross-domain pair in the curriculum: in each, the construct definition is the upstream design decision; the framework's C8.2 sub-competency and the equity-construct subobjective (4.6) are anchored here.

## REFLECTION QUESTIONS

1. Identify a predictive-targeting model in your domain. Which construct does it maximize — observed outcome, predicted benefit, predicted need — and is the choice on the record as a normative institutional decision, or absorbed into the model's training objective without disclosure?
2. Specify the stratifications you would build into the model's validation: which equity-relevant axes (income, race/ethnicity, first-generation, geography, disability) and what comparison structure (per-axis accuracy, per-axis calibration, per-axis intervention effectiveness)?
3. The case cross-references the v2 race-construct trio (Cases 25, 26, 6). In each, the construct definition is the upstream design decision that produced the disparate outcome. Identify a construct decision in your domain that is currently absorbed into the design rather than on the record — and what would it take to make the choice explicit?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

# Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Yu, Lee, and Kizilcec, publishing at ACM Learning @ Scale in 2021, tested whether four protected attributes (gender, underrepresented-minority status, first-generation status, high financial need) should be included as features in dropout prediction across 93,457 students at one U.S. research university, and found that including them left overall prediction performance unchanged and improved algorithmic fairness only marginally — small-tier frontier evidence that the include-or-exclude question is settled locally, by audit, rather than by principle

## IN BRIEF

A long-running practical question in learning analytics is whether protected attributes — race/ethnicity, gender, socioeconomic status, first-generation status — should be included as features in predictive models. The intuitive “fairness through unawareness” answer is to exclude them. The technical-fairness literature has shown the unawareness answer is incomplete: omitted protected attributes are typically reconstructable from correlated features (zip code, course history, prior achievement), so excluding the attribute does not exclude its predictive footprint and can make discrimination harder to detect and audit. Yu, Lee, and Kizilcec, publishing at ACM Learning @ Scale, examined the include-or-exclude question empirically on one prediction task at one institution. The headline finding is the frontier-shaped one: including the four protected attributes left overall prediction performance unchanged and improved algorithmic fairness only marginally, which the authors read as reason to include them and to evaluate the effect in each local context. The case is the small-tier frontier instance of “surfacing bias through governance, not just technique” (C8.4 in the induced framework). It cross-references the v2 race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman) at the model-fairness layer: in those cases the construct was the design decision; in this case the question is whether the attribute is allowed into the model that operationalizes the construct.

## THE SHIFT

A practical question that recurs across learning-analytics deployments is whether protected attributes — race/ethnicity, gender, socioeconomic status, first-generation status, and similar — should be included as features in predictive models. The intuitive policy answer (“fairness through unawareness”) is to exclude them: a model that does not see race cannot discriminate on race. The intuition is operationally appealing and statutorily aligned in some domains where regulators have read the omission requirement strictly.<sup>o</sup>

## WHAT IS EMERGING

The technical-fairness literature has shown the unawareness answer is incomplete. Protected attributes that are omitted from a model are typically reconstructable from correlated features the model does see — zip code, school assignment, prior achievement, course history, language at home, family income proxies. Omitting the attribute removes the label, not the predictive footprint, and can in some cases make discrimination harder to detect and audit precisely because the auditable record no longer carries the attribute that would let an evaluator stratify the model's output.<sup>1</sup>

## THE CAPABILITY QUESTION

Yu, Lee, and Kizilcec, publishing at the ACM Conference on Learning @ Scale, examined the include-or-exclude question empirically on institutional records for the whole student body of a large U.S. research university — 93,457 students, residential and fully online — predicting dropout after one academic year with logistic regression and gradient-boosted trees, with and without four protected attributes: gender, underrepresented-minority status, first-generation status, and high financial need. The headline finding is the frontier-shaped one: including the attributes left overall prediction performance unchanged and improved algorithmic fairness only marginally. The result favors inclusion, and it favors it weakly enough that the answer has to be re-established in each local context.<sup>2</sup>

## EARLY EVIDENCE

The right framing the paper develops is governance and audit. The decision to include or exclude a protected attribute is one of several decisions a learning-analytics deployment makes that determine its fairness properties; the decision has to be made deliberately, recorded with reasoning, and paired with stratified evaluation and audit cadence that catches the consequences of the decision in operation. The case is the induced framework's C8.4 instance at small scale: "surfacing bias through governance, not just technique." The technical-fairness machinery alone does not answer the question; the governance architecture around the model is the carrier of the answer in any specific deployment.<sup>3</sup>

## OPEN PROBLEMS

In pair with Case 86 (Gándara on student-success prediction) and with the v2 race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman), the case completes the v2 equity-construct frontier picture. Case 86 names the construct-definition layer; this case names the protected-attribute-in-the-model layer; the race-construct trio names the construct-encoded-in-the-instrument layer (race correction in eGFR, sensor calibration in pulse oximetry, pain perception in Hoffman). Together, the five cases stage the equity-construct competency across the construct-definition / attribute-handling / instrument-construct axes — the case-grounded basis for the equity-construct sub-competencies proposed in the v2 research backbone.<sup>4</sup>

*Including the protected attributes changed prediction performance not at all and fairness only marginally. What settles the question is not principle but a local audit, recorded with its reasoning.*

EDITORS' SYNTHESIS OF YU, LEE, & KIZILCEC (L@S 2021).

## REFERENCES

1. Yu, R., Lee, H., & Kizilcec, R. F. (2021), "Should College Dropout Prediction Models Include Protected Attributes?" in Proceedings of the Eighth ACM Conference on Learning @ Scale (L@S '21), doi:10.1145/3430895.3460139 — primary paper on the include-or-exclude empirical analysis.
2. Kizilcec, R. F., & Lee, H. (2022), "Algorithmic Fairness in Education," in Holmes & Porayska-Pomsta (eds.), The Ethics of Artificial Intelligence in Education, Routledge, doi:10.4324/9780429329067-10 — broader synthesis of the fairness-in-learning-analytics frontier.
3. Dwork, Hardt, Pitassi, Reingold, & Zemel (2012), "Fairness through awareness," Proceedings of ITCS — foundational paper on the inadequacy of fairness-through-unawareness.
4. Barocas, Hardt, & Narayanan, Fairness and Machine Learning (fairmlbook.org) — technical-fairness backdrop.
5. v2 cross-referenced cases: 86 (Gándara), 25 (eGFR), 26 (pulse oximetry), 6 (Hoffman) — equity-construct five-case set.

## THE LEARNING ENGINEERING LENS

## LE INSIGHT

Yu, Lee, and Kizilcec's protected-attributes work is the frontier instance of surfacing bias through governance, not just technique. Including four protected attributes in a dropout model built on 93,457 student records left overall performance unchanged and improved fairness only marginally; the authors read that as reason to include them. The governance architecture around the model is the carrier of the answer in any specific deployment.

## LENS APPROACH

Yu/Lee/Kizilcec protected attributes is the model-fairness- governance case (induced 8.4; LENS D4/PT5) — Domain 4 for stratified-evaluation-by-attribute; Domain 5 for the decision-plus-reasoning-plus-audit architecture. Cross- reference Case 86 and the race-construct trio (25, 26, 6).

| DOMAIN             | MODES AT STAKE | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|--------------------|----------------|---------------------|--------------|---------|-------|
| LEARNING ANALYTICS | D K N          | D1 D2 D3 D4 D5      | 5            | 8.4     | 4 · 5 |
| ALGO FAIRNESS      |                | Test and Evaluation |              |         |       |
| ED-TECH RESEARCH   |                |                     |              |         |       |

## APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- ▶ Treat the include-or-exclude decision for each protected attribute as a deliberate design choice, recorded with reasoning, rather than absorbed into the data-engineering pipeline.
- ▶ Build stratified evaluation by the protected attribute regardless of whether the model itself uses it as a feature; auditing the model's output by attribute is independent of whether the attribute is an input.
- ▶ Pair the decision with the downstream intervention context; the include-or-exclude answer that produces fairer outcomes in one intervention context may produce less fair outcomes in another.

## IN OPERATION / AFTER

- ▶ Publish the protected-attribute handling decisions for any deployed learning-analytics model — included, excluded, and with what reasoning — as part of the model's governance documentation.
- ▶ Operate stratified audit on a regular cadence; the include-or-exclude consequences in operation are what the audit catches, and audit absence makes the decision functionally invisible to the institution.
- ▶ Carry the five-case equity-construct set into the curriculum: Cases 25 (eGFR), 26 (pulse oximetry), 6 (Hoffman) on the instrument-construct layer; Case 86 (Gándara) on the construct-definition layer; this case on the protected-attribute-in-the-model layer. The set is the case-grounded basis for the equity-construct sub-competencies.



## REFLECTION QUESTIONS

1. Identify a learning-analytics or analogous predictive model in your domain. Which protected attributes are inputs to the model, which are not, and is the include-or-exclude decision on the record with reasoning, or absorbed into the data-engineering pipeline?
2. Specify the stratified-audit cadence you would operate for the model regardless of whether protected attributes are inputs; the audit catches the consequences of the include-or-exclude decision in operation, and its absence makes the decision functionally invisible.

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

# LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Repeated probing of LiveHint AI (an LLM-based tutor in development at Carnegie Learning, released to a small number of U.S. school districts) found no evidence of harmful bias and no dialect effect, but response differences in tone and level of detail tied to stated nationality/ethnicity and to prompt formality; choice of foundation model materially affected the level of differentiation

## IN BRIEF

The AIED 2025 paper “Evaluating an AI Tutor for Bias Across Different Foundation Models” tests LiveHint AI, an LLM-based tutor under development at Carnegie Learning, against a structured probing protocol. Realistic student queries are modified to include explicit or implicit statements of identity — nationality, dialect markers, demographic cues — and the tutor’s responses are assessed for tone and level of detail. The choice of foundation model materially affects the level of differentiation in responses. The authors describe the work as a collaborative audit run with Carnegie Learning’s cooperation, offered as a template for future audits of LLM-based tutors; LiveHint is in research and development and released to a small number of U.S. school districts. The case extends the race-construct trio (Cases 25, 26, 6, 48) into the LLM-tutoring layer where the structurally new variable is the foundation model. Open questions: whether lab-style probing matches deployed-conversation patterns; whether vendor selection across foundation models becomes a routine fairness deliverable.

## THE SHIFT

The deployment surface for LLM-based tutoring expanded faster than the fairness-evaluation literature for it. Most pre-deployment fairness work on educational AI was built for classifiers — proctoring systems, recommendation engines, automated essay scoring — where the output is a score or a flag and the fairness question can be formulated as a disparate-impact comparison across groups. LLM tutoring shifts the output to a conversational response, and the fairness question shifts with it: it is no longer whether two equivalently-prepared students receive equivalent flags, but whether they receive equivalent pedagogical engagement.<sup>o</sup>

## WHAT IS EMERGING

The AIED 2025 paper builds the evaluation method against that shifted target. LiveHint AI, an LLM-based tutor under development at Carnegie Learning, is the subject system.

The probing protocol generates realistic student queries — calibrated to the kinds of mathematics questions a tutor would receive — and produces matched variants of each query that differ only in identity-marker content: explicit statements of nationality, dialect markers, demographic cues. The tutor’s responses to the matched pairs are then assessed on two dimensions, using lexical metrics: tone (is the response respectful and appropriately framed for a tutoring context) and level of detail (does the response give the same quality of pedagogical content). Whether the resulting differentiation reflects sound pedagogy or undesirable behavior is the interpretive question the paper leaves open.<sup>1</sup>

#### THE CAPABILITY QUESTION

The headline finding is that the choice of foundation model materially affects the level of differentiation in responses. Three versions of the tutor were probed — one each on GPT-4, GPT-4o and Claude 3.5 Sonnet — and they produce different patterns of identity-conditioned variation: Claude 3.5 differentiated most, GPT-4 next, GPT-4o least. The variation is not uniform across the nine lexical metrics, nor across identity-marker types; and the comparison that carries the case is that responses to the unmodified baseline prompts differ more across the three foundation models than any version’s responses differ across identity markers. The concrete instances are specific: all three versions sometimes answer an English-language prompt from a self-identified Spanish or Italian student entirely in Spanish or Italian. The vendor-selection decision — which foundation model the tutoring product is built on — is itself a fairness-relevant design choice, and the case names it as such. The structurally new variable in the case, by comparison with the race-construct trio in the corpus, is the foundation-model layer, which did not exist as a deployment-side decision when the proctoring-bias and pulse-oximetry cases were studied.<sup>2</sup>

#### EARLY EVIDENCE

The authors’ framing is explicit and binding on the case. LiveHint AI is in research and development and released to a small number of U.S. school districts; the paper is a collaborative audit offered as a template for auditing LLM-based tutors. The case is not the deployment-bias-audit case the corpus carries at Cases 6 (Hoffman pain assessment), 26 (pulse oximetry), 25 (eGFR), and 48 (Johnson school surveillance). It is the structurally new methods-development case at the layer above those deployments — the foundation-model layer — and it grounds the curriculum’s demographic-stratification anchor at that layer.<sup>3</sup>

#### OPEN PROBLEMS

The open questions the authors preserve are the case’s decisive hedges. The authors also report that they found no evidence of harmful biases — the differentiation they measured may reflect sound pedagogy or undesirable behavior, an interpretive question they leave open. Whether the differentiation patterns documented in lab-style probing match what students encounter in deployed conversations — where session length, follow-up turns, and student adaptation affect the response trajectory — is a question the present study cannot answer. Whether vendor selection across foundation models becomes a routine fairness deliverable for the educational-AI procurement pipeline is a market-evolution question the present study can name but cannot resolve. The case pairs with the race-construct trio for the stratified-validation discipline and with Case 77 (hybrid human-AI tutoring) for the deployment-side complement — the augmentation pattern Case 77 documents depends on the foundation-model-level evaluation Case 88 is methodologically grounding.

*The choice of foundation model is itself a fairness-relevant design decision. The evaluation should be a pre-deployment deliverable, not a post-deployment audit.*

EDITORS' SYNTHESIS OF THE AIED 2025 LIVEHINT AI BIAS EVALUATION.

REFERENCES

1. Bommasani, R. et al. (2021), "On the Opportunities and Risks of Foundation Models," Stanford CRFM — the foundation-model framing the case builds on.

2. Vinodh, A., Harvey, E., Almoubayyed, H., Yu, R., Brooks, C., Koe-neck, A., & Kizilcec, R. F. (2025), "Evaluating an AI Tutor for Bias Across Different Foundation Models," AIED 2025, LNAI 15882, pp. 341–348, doi:10.1007/978-3-031-98465-5\_43; author copy at renzheyu.com/papers/AIED2025\_Tutor.pdf.

3. Race-construct trio reference set: Hoffman et al. (2016), Sjöding et al. (2020) pulse oximetry, Inker et al. (2021) eGFR-without-race — paired with Cases 25, 26, 6.

4. Carnegie Learning LiveHint product documentation — the subject system; case framing is binding on LiveHint being in development, not deployment.

THE LEARNING ENGINEERING LENS

LE INSIGHT

LiveHint AI is the methods-development case at the foundation- model layer of LLM-tutoring fairness evaluation. The matched- pair identity-probing protocol surfaces tone and level-of-detail differences across identities; the foundation-model choice materially affects the differentiation level. The case extends the race-construct trio into the LLM- tutoring layer with a structurally new variable — the foundation model — that did not exist at the deployment-audit layer.

LENS APPROACH

LiveHint AI is the demographic-stratified validation case at the foundation-model layer (induced 8.2; LENS D4/PT6). LENS uses it in Domain 4 (Test and Evaluation) for the matched- pair probing methodology and in Domain 3 (Human-System Collaboration) for the foundation-model-selection-as- fairness-decision frame. Pair with the race-construct trio (Cases 25, 26, 6, 48) at the deployment-audit layer and with Case 77 (hybrid human-AI tutoring) as the augmentation-pattern complement.



APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Build the probing protocol to vary identity markers within matched query pairs; the methods-development contribution depends on the matched-pair design that isolates the identity-conditioned response variation.
- ▶ Treat the choice of foundation model as a fairness-relevant design decision, not as an upstream procurement decision; the case demonstrates that the foundation-model layer materially affects the tutoring-layer's identity-conditioned response patterns.
- ▶ Conduct the fairness evaluation before deployment, not after; the methods-development framing is that the evaluation should be a pre-deployment deliverable, parallel to the race-construct-trio cases at the deployment-audit layer.

IN OPERATION / AFTER

- ▶ Commission the deployed-conversation evaluation that the lab-style probing cannot perform; the open question on whether probing patterns match deployment patterns is testable against deployment logs as the system moves toward release.
- ▶ Publish the foundation-model-level fairness findings as part of the procurement record; the case argues that vendor selection across foundation models is a fairness deliverable, and the publication discipline is what would make that deliverable operational.
- ▶ Pair the case in the curriculum with the race-construct trio (Cases 25, 26, 6, 48) so the demographic-stratification anchor is taught across both the deployed-system layer and the foundation-model layer above it.

## REFLECTION QUESTIONS

1. Identify an LLM-based system in your domain whose foundation-model choice was treated as an upstream procurement decision rather than as a fairness-relevant design decision. What would the matched-pair probing protocol look like for that system, and what response dimensions would you assess?
2. Specify the pre-deployment fairness deliverable you would build into the procurement record for an LLM-based product in your domain. What would the published artifact contain — probe set, response-dimension assessment, foundation-model comparison — and what would the procurement decision turn on?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

# Data Privacy and Learning Analytics on the African Continent

**K** Knowledge & Institutional Memory

**G** Governance & Trust

**N** Normalization of Deviance

**IMPACT** Mapped the legal and regulatory privacy landscape across African jurisdictions and surfaced the structural governance seam — African higher education frequently uses learning platforms and analytics hosted by external providers, creating cross-regime gaps where student data crosses jurisdictions with inconsistent protection — and recommended common cross-border data-sharing frameworks

## IN BRIEF

Prinsloo and Kaliisa (British Journal of Educational Technology, 2022) mapped the legal and regulatory privacy landscape across African jurisdictions — a scoping review of privacy regulations in 32 countries, complemented by 15 scholarly papers — and found numerous national and regional instruments already giving clear pointers on student data privacy. The core governance seam the paper surfaces — and what makes the case a frontier rather than a settled instance — is what the v2 sweep names **extraterritorial platform governance**: African higher education frequently uses learning platforms, learning-management systems, and analytics services hosted by external (typically Global-North or Asian) providers, creating a structural gap where student data is generated under one regulatory regime, processed under another, and the institution's governance authority does not reach the operating regime. The paper recommends common cross-border data-sharing frameworks as the architectural response. The case is included as a frontier — it documents the governance architecture needed for responsible adoption rather than a completed success — and it is valuable because it surfaces a seam (extraterritorial hosting) the U.S.-centric canon rarely confronts, and that is increasingly universal as institutions everywhere build on cloud and AI services they do not control.

## THE SHIFT

The standard framing of learning-analytics governance assumes that the operating institution, the regulatory regime, and the hosting infrastructure are co-located — a U.S. university under U.S. law on U.S.-hosted infrastructure, a UK university under UK law on UK or EU infrastructure. African higher education sits structurally outside this assumption. Most universities across the continent rely on learning-management systems, predictive analytics tools, and adjacent platforms hosted by providers based in the Global North or in Asia, and the regulatory regimes governing those providers are not the same as the regimes governing the institutions or the students.<sup>9</sup>

## WHAT IS EMERGING

Prinsloo and Kaliisa's 2022 paper in the *British Journal of Educational Technology* maps the legal and regulatory privacy landscape across African jurisdictions. The picture is one of a continent already legislated: a scoping review of privacy regulations in 32 countries, complemented by 15 scholarly papers, found numerous national and regional instruments giving clear pointers on student data privacy, with substantial variation across countries in scope, enforcement capacity, and cross-border data-transfer policy. Some countries have comprehensive frameworks; some have sectoral protections; some have constitutional privacy provisions without operational data-protection legislation; and the patchwork makes the cross-border picture itself uneven.<sup>1</sup>

## THE CAPABILITY QUESTION

The structural seam the paper surfaces is what makes the case teachable. When student data is generated under one jurisdiction's regulatory regime, processed under another's because the platform is hosted abroad, and the institution's governance authority does not reach the operating regime, the governance question becomes architectural rather than institutional. The authors put the cross-border data-sharing frameworks to African governments and the platform question to the institutions, which should weigh student data privacy when entering service level agreements with educational-technology and platform providers — the kind of inter-regime instrument that GDPR's adequacy decisions and successor cross-border-transfer mechanisms began to construct between the EU and other regimes, but that no comparable continent-wide African framework yet provides. The case names the gap and the architectural response; what it does not yet document is the build-out.<sup>2</sup>

## EARLY EVIDENCE

The honest evidentiary state is what the frontier tag carries. The paper is peer-reviewed and the mapping is grounded in legal documentation across jurisdictions; what the case does not establish — because the architecture is not yet built — is that any specific cross-border framework has been adopted and that the seam has been closed in practice. The case is teachable on the structural pattern: the governance question is the architectural one across regimes, not the policy one inside any single regime. The outcome is open in the sense that no continent-scale resolution yet exists. Future validation will continue as African data-protection law consolidates.<sup>3</sup>

## OPEN PROBLEMS

The frontier note the case carries names what the v2 sweep calls extraterritorial platform governance — capability deployed on platforms governed by a different regime than the one operating them. The pattern is not African-specific. It is increasingly universal as institutions everywhere build on cloud, AI, and platform services they do not control: a U.S. university on a U.S. cloud is the historical case; a U.S. university running analytics through a model-hosted-elsewhere service, a Latin American university on a U.S. cloud, an African university on a European platform — all share the structural seam, and the African case surfaces it most starkly because the regime asymmetry is largest there. The case is the case-grounded instance of a sub-competency both frameworks already name — LENS 5.6, cross-regime and platform-dependency governance, and induced 5.3, which is this case's own anchor.

*The governance question is architectural across regimes, not policy inside any single regime. The seam is at the joints, and the joints have to be drawn before the architecture can be designed.*

EDITORS' SYNTHESIS OF PRINSLOO & KALIISA (2022).

REFERENCES

1. Prinsloo, P., & Kaliisa, R. (2022), "Data privacy on the African continent: Opportunities, challenges and implications for learning analytics," British Journal of Educational Technology 53(4):894–913, doi:10.1111/bjet.13226.

2. African Union Convention on Cyber Security and Personal Data Protection (Malabo Convention, adopted 2014; in force 8 June 2023) — continental policy backdrop, ratified by a minority of AU member states.

3. Slade & Prinsloo (2013), American Behavioral Scientist — earlier framing on which the 2022 paper builds.

4. Open University Ethical Use of Student Data policy (2014) — single-regime governance-by-design companion.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The African data-privacy case names a governance seam the existing curriculum underweights: extraterritorial-platform governance, where the operating, regulating, and hosting regimes diverge. The case is a frontier — peer-reviewed mapping, architectural recommendation, no continent-scale resolution yet built. The pattern is increasingly universal, and the case-grounded basis is the case for naming a sub-competency the framework does not yet have.

LENS APPROACH

African data privacy is the extraterritorial-platform governance case (induced 5.3; LENS D5/PT6) — Domain 5 for the cross-regime seam, on subobjective 5.6 (cross-regime / platform-dependency governance); Domain 4 for **Judgment under inadequate evidence** (4.4) on a frontier where the architectural response is recommended but not built. Pair with Case 81 and Case 91.

| DOMAIN                                                              | Modes at Stake | LENS COMPETENCY                                         | PROBLEM TYPE | INDUCED | LEO   |
|---------------------------------------------------------------------|----------------|---------------------------------------------------------|--------------|---------|-------|
| HIGHER EDUCATION (AFRICA)<br>DATA PRIVACY LAW<br>LEARNING ANALYTICS | K G N          | D1 D2 D3 D4 D5<br>Navigating Sociotechnical Constraints | 6            | 5.3     | 5 · 4 |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Map the regime structure explicitly: which jurisdiction generates the data, which processes it on the platform, which the institution operates under, which the student lives under. The seam is at the joints, and the joints have to be drawn before the architecture can be designed.
- Choose the platform stack with the regime structure in view; a procurement decision is also a governance decision when the platform is in a different jurisdiction from the operating institution.
- Treat the architectural response as common cross-border frameworks rather than institution-by-institution contract negotiation; the inter-regime instrument is what scales.

IN OPERATION / AFTER

- Surface the extraterritorial-platform-governance pattern in any v2 case that involves cloud/AI infrastructure; the African case is the starkest instance, not the only one.
- Carry the outcome-open honesty: the architecture is not yet built, and the case names the gap rather than a completed success.
- Build the case-grounded basis for proposing extraterritorial-platform-governance as a sub-competency the existing curriculum does not yet name; pair with the OU and LALA cases at the governance-by-design layer.



## REFLECTION QUESTIONS

1. Identify a platform or service your institution depends on whose operating regime is different from your own. Which jurisdiction generates the data, which processes it, which regulates the institution, which regulates the user — and where is the seam between regimes?
2. Specify the architectural response your context needs: an institution-by-institution contractual approach, a sectoral framework, a cross-border data-sharing instrument. Which level resolves the seam, and which level cannot?
3. The case names extraterritorial platform governance as an under-named sub-competency. What instance from your domain would anchor it — a cloud-hosted clinical analytics tool, a model-as-a-service inference layer, a learning-platform vendor in a different jurisdiction — and how is the curriculum response different from a single-regime governance case?

**WHO BUILDS THIS** knowledge management & data engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

## CASE 90

HIGHER EDUCATION

ALGORITHMIC DECISION-MAKING

HUMAN-COMPUTER INTERACTION

2025

# Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**DISCLOSURE** Institutional overlap: John Katzman collaborates with an editor's institution (Johns Hopkins) but was not involved in this work. Pyle and Andalibi's study is independent of Katzman; the case sits on the applicant side of enrollment management.

**IMPACT** A content analysis of 52 admissions-technology vendor websites paired with eighteen semi-structured interviews using speculative verbal probes surfaced systematic distances between vendors' claims (efficiency and ease, transparency, bias mitigation and DEI, accuracy) and applicants' perceptions (increased applicant labour, opacity and privacy violation, algorithmic bias and conformity, new inaccuracy and gaming)

**IN BRIEF**

Pyle and Andalibi (CSCW 2025) report an interview study with 18 recent U.S. university applicants, using speculative-design probes to surface how applicants perceive algorithmic systems operating in college admissions and enrollment. The study's central contribution is naming the systematic distance across four axes between vendors' claims — efficiency and ease, transparency, bias mitigation and DEI, accuracy — and what applicants named against each: increased applicant labour, opacity and privacy violation, algorithmic bias and conformity, and new forms of inaccuracy and gaming. The case is the peer-reviewed consent-side companion to Case 55 (Engler / enrollment algorithms) and Case 57 (GAO OPM oversight gap). The authors' own hedge is explicit: recruitment stopped at thematic saturation after 18 interviews, and the study attempted to capture varying experiences and perceptions, not representative or generalizable findings; "future validation ongoing" applies to whether the patterns hold across applicant demographics and institution types.

**THE SHIFT**

The deployment surface for algorithmic systems in college admissions and enrollment has been studied largely from the institutional side — what vendors promise, what colleges buy, what optimization outcomes the institution reports. The applicants whose lives the deployment most directly affects have been the structurally absent voice in the deployment record. Pyle and Andalibi's CSCW 2025 paper is the systematic peer-reviewed correction at the consent-side seam: a content analysis of 52 admissions-technology vendor websites

paired with interviews of 18 recent U.S. university applicants, using speculative verbal probes built from public documentation of deployed systems.<sup>9</sup>

#### WHAT IS EMERGING

The speculative–design method is the right instrument for the question. Applicants in most cases do not know which specific algorithmic systems acted on their applications; the institutional side does not disclose the vendor stack or the optimization objective at the point of application. Speculative–design probes — scenario sketches that concretize plausible algorithmic interventions and ask applicants to react — let the study elicit perceptions that are not contingent on the applicant having direct knowledge of the deployed system. The 18-interview sample is calibrated to the method’s depth rather than to prevalence–claim breadth, and the authors are explicit about that calibration as a methodological choice.<sup>1</sup>

#### THE CAPABILITY QUESTION

The findings name systematic distance between vendor marketing and applicant perception across four axes. Vendors pitch efficiency and ease for admissions officers — the system processes more applications faster — while applicants name the labour the system moves onto them. Vendors pitch transparency for admissions officers, while applicants name opacity and privacy violation: they do not know which decisions are made algorithmically and cannot interrogate the logic. Vendors pitch bias mitigation and DEI, while applicants name algorithmic bias and pressure to conform. Vendors pitch accuracy, while applicants name new forms of inaccuracy and openings to game the system.<sup>2</sup>

#### EARLY EVIDENCE

The case sits as the consent–side counterpart to Case 55 (Engler / Brookings, the deployment–side mapping) and Case 57 (GAO–22–104463, the regulator–side audit). Engler documents the deployed algorithmic optimization; the GAO audit documents the regulatory oversight gap; Pyle and Andalibi document what the affected applicants understand about the deployment. The three cases together name the structural seam: a deployed system about which the institutional side, the regulator, and the affected population each hold partial and non–overlapping information. The case also pairs with Cases 186 (Bartlett lending fairness) and 86 (Gándara on bias in student–success prediction) as the applicant–perception strand of the equity–in– deployed–prediction thread.<sup>3</sup>

#### OPEN PROBLEMS

The authors’ hedges are binding on the case’s framing. Eighteen semi–structured interviews reached thematic saturation and surfaced the dimensions of perception that matter, but the authors state the study did not aim at representative or generalizable findings, so it carries no claim about prevalence — what fraction of applicants hold each perception, how the perceptions distribute across demographic groups, how they correlate with admissions outcomes. The case is the strongest peer–reviewed evidence currently available on the consent–side question; future validation ongoing on whether the perception patterns generalize across applicant demographics, institution types, and the rapidly evolving algorithmic–deployment landscape.

*Vendors pitch efficiency, transparency, bias mitigation, and accuracy. Applicants name labour, opacity, algorithmic bias, and new inaccuracy. The distance between the two is the case.*

EDITORS' SYNTHESIS OF PYLE & ANDALIBI (2025).

REFERENCES

1. Pyle, C., & Andalibi, N. (2025), "Algorithmic College Admissions in the U.S.: Distances Between Vendors' Claims and Applicants' Perceptions," Proceedings of the ACM on Human-Computer Interaction 9(7), CSCW369, doi:10.1145/3757550.

2. Engler (2021), Brookings — paired deployment-side mapping (Case 55).

3. GAO-22-104463 (2022) — paired regulator-side audit (Case 57).

4. Dunne, A., & Raby, F. (2013), Speculative Everything — methodological backdrop for speculative-design probes.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Pyle and Andalibi is the peer-reviewed consent-side case at the applicant end of the enrollment-management deployment. A 52-vendor website analysis and eighteen interviews surface systematic distance between vendor claims and applicant perception across four axes; the authors' methodological hedge is binding on prevalence claims. The case completes the partial-information triangle with Cases 55 (deployment-side) and 57 (regulator-side).

LENS APPROACH

Pyle & Andalibi is the governance-rather-than-technique case at the consent boundary (induced 8.4; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the affected-voice inclusion lesson and in Domain 3 (Human-System Collaboration) for the disclosure-architecture design at the applicant interface. Pair with Cases 55 (Engler deployment), 57 (GAO oversight), 186 (Bartlett), and 86 (Gándara). coi-light render under the title is binding for cluster placement.

| DOMAIN                         | MODES AT STAKE | LENS COMPETENCY                         | PROBLEM TYPE | INDUCED | LEO   |
|--------------------------------|----------------|-----------------------------------------|--------------|---------|-------|
| HIGHER EDUCATION               | T K N          | D1 D2 D3 D4 D5                          | 6            | 8.4     | 5 · 3 |
| ALGORITHMIC<br>DECISION-MAKING | →              | → Navigating Sociotechnical Constraints |              |         |       |
| HUMAN-COMPUTER<br>INTERACTION  |                |                                         |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Include the affected-population voice in the deployment-decision record from the start; the systematic distance the study documents is partly the result of decision processes that did not include the affected voice.
- ▶ Use speculative-design probes when the affected population cannot be expected to know what is deployed; the method is the right instrument for the consent-side question and the institutional side can commission it.
- ▶ Treat the perceived harms — applicant labour, opacity and privacy violation, algorithmic bias, new inaccuracy — as designable targets, not as misperceptions to

IN OPERATION / AFTER

- ▶ Commission the prevalence study the speculative-design study cannot perform; the perception patterns the 18-interview study surfaces are testable at survey scale, and the deployment-side institutions are best positioned to commission the survey.
- ▶ Pair the consent-side study with the regulator-side audit (Case 57) and the deployment-side mapping (Case 55) in the curriculum; the three cases together name the partial-information structure of the deployment.
- ▶ Track the perception findings over time as the algorithmic-deployment landscape evolves; the case's value as a longitudinal baseline depends on the comparison studies that come next.

correct; the perceptions are responses to the actual disclosure architecture.

### REFLECTION QUESTIONS

1. Identify a deployed prediction system in your domain whose affected population was not consulted in the deployment–decision record. What would a speculative–design study look like for surfacing the affected voice — and who is positioned to commission it?
2. Specify the perceived harms — applicant labour, opacity and privacy violation, algorithmic bias, new inaccuracy — for the system you identified. Which is a disclosure–architecture target, which is an objective–choice target, and which is an audit–evidence target?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

## CASE 91

HIGHER EDUCATION (LATIN AMERICA)

LEARNING ANALYTICS

CROSS-NATIONAL GOVERNANCE

2017 – 2020

# LALA — Building Learning-Analytics Governance Capacity Across Latin America

**K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** An EU-funded multi-country project (Chile, Ecuador) that adapted the European SHEILA protocols rather than importing Global-North tools unchanged; 37 manager interviews, 16 focus groups with 51 teaching staff and 45 students, and surveys of 1,884 students and 368 staff produced a stakeholder needs map, and the project published the four-dimension LALA Framework for institutional adoption

## IN BRIEF

The LALA (Building Capacity to Use Learning Analytics to Improve Higher Education in Latin America) project, funded under EU grant 586120-EPP-1-2017-1-ES, ran from 2017 to 2020 across Chilean and Ecuadorian universities. The project began by adapting rather than importing: the European SHEILA project's interview, focus-group and survey protocols were translated and reworked for four Latin American universities, on the grounds that regional adoption had lagged and the needs had never been measured locally. Through structured interviews with administrators and focus groups with students and teachers, the Hilliger et al. team (Internet and Higher Education, 2020) surfaced what stakeholders actually needed for adoption to be locally legitimate, and the project published the LALA Framework — four dimension manuals, institutional, technological, ethical, and community, that an institution works through before it deploys. The honest limit preserved verbatim: this is adoption-readiness and capacity-building evidence, not yet long-run outcome evidence that the deployed systems improved student retention or learning. The contribution is the participatory governance method that made adoption locally legitimate — and the case is the non-US companion to OU Analyse (Case 81), where governance-by-design unblocked deployment in a different cross-cultural context.

## BACKGROUND

Learning analytics as a field is heavily Global-North-shaped. The reference tools (predictive at-risk classifiers, dashboard analytics for instructors, early-warning systems) were developed at U.S. and European universities with U.S. and European governance assumptions about data, consent, and institutional decision-making. The LALA project began from a thinner and more awkward observation: regional attempts at learning analytics remained small in scale, and the region lacked a community to exchange ideas, methods and tools because experienced researchers in the field were few. There was no local needs evidence to build on, so the project measured it before it built anything.<sup>o</sup>

## THE INTERVENTION

The project ran across Chilean and Ecuadorian universities under an EU Erasmus+ grant from 2017 to 2020. The method's first commitment was participatory: semi-structured interviews with 37 senior managers surfacing what their decisions actually needed evidence about; 16 focus groups with 51 teaching staff and 45 students; and surveys answered by 1,884 students and 368 teaching staff, triangulated against the qualitative coding. Hilliger et al. (Internet and Higher Education, 2020) is the peer-reviewed mixed-methods report that documents the method and the findings. The headline result is a needs map: students named quality feedback and timely support, teaching staff named timely alerts and meaningful evaluation of their own practice, and managers named quality information to evaluate support interventions. Ethical and privacy considerations the authors list as the next study, not as a finding of this one.<sup>1</sup>

## HOW IT WORKED

The deliverable is the LALA Framework — four dimension manuals, institutional, technological, ethical and community, that walk an institution through what it has to settle before a learning-analytics system is deployed. Its opening instrument is the LALA Canvas, a moderated group template whose six dimensions are adapted from the Overseas Development Institute's ROMA approach: desired behaviours, change strategy, internal capabilities, political context, influential actors, and a measurement and evaluation plan. The framework's contribution is procedural: it converts the governance question from a yes/no gate at deployment time into a structured set of decisions taken openly during the adoption process. The participatory method made the framework locally legitimate across two regimes whose own data governance is differently mature.<sup>2</sup>

## THE EVIDENCE

The honest limit survives into the case. The published evidence is needs and capacity-building evidence: the framework was developed through the participatory process, and the project produced trained local teams with the capacity to lead adoption in their own institutions. The authors put adoption and evaluation of implemented tools in future work. What the evidence does **not** yet establish is that the deployed systems improved long-run student retention or learning outcomes. The case is teachable on the governance method — the participatory route to local legitimacy — and the outcome evidence is the next study, not this one. Drafting that softens this hedge over-states the claim.<sup>3</sup>

## WHAT TRANSFERRED

In pair with the Open University case (Case 81, the UK consent-by-design intervention), LALA shows that the governance-by-design pattern is transferable across regimes: OU built consent for a single-institution intervention under pre-GDPR UK scrutiny; LALA built participatory adoption for multi-country capacity-building under two different Latin American regulatory regimes. The pair plus SyRI (Case 189) teaches the non-US LA governance triple — design that unblocked deployment (OU), participatory governance that built adoption capacity (LALA), and rights-grounded halt (SyRI) — three honest results for the same structural question of when delegation to analytic infrastructure is legitimate.

*Ask the stakeholders what they need before you buy the tool. Students named feedback, staff named alerts they could use, managers named information they could act on — and the ethics manual came from the literature, not from the room.*

EDITORS' SYNTHESIS OF HILLIGER ET AL. (2020).

## REFERENCES

1. Hilliger, Ortiz-Rojas, Pesántez-Cabrera, Scheihing, Tsai, Muñoz-Merino, Broos, Whitelock-Wainwright, & Pérez-Sanagustín (2020), "Identifying needs for learning analytics adoption in Latin American universities: A mixed-methods approach," *The Internet and Higher Education* 45:100726, doi:10.1016/j.iheduc.2020.100726.
2. LALA project — EU Erasmus+ grant 586120-EPP-1-2017-1-ES (2017–2020) — program documentation and the LALA CANVAS artifact.
3. Open University Ethical Use of Student Data policy (2014) and OU Analyse — UK companion governance-by-design case (Case 81).
4. Slade & Prinsloo (2013), "Learning Analytics: Ethical Issues and Dilemmas," *American Behavioral Scientist* 57(10):1510–1529 — the broader field-scale ethics framing.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

LALA converted learning-analytics adoption from a Global-North template-lift into a participatory process that built local legitimacy across two Latin American regimes. The LALA Framework is the artifact; the participatory needs study is the deliverable. The honest limit — adoption-readiness evidence, not yet long-run outcome evidence — is what the case carries, and the outcome study is the next one.

### LENS APPROACH

LALA is the non-US participatory-governance case (induced 5.1; LENS D5/PT4) — Domain 5 for the cross-regime participatory method; Domain 4 on **Judgment under inadequate evidence** (4.4) — adoption-readiness evidence, not closed outcome proof. Pair with Case 81 (OU) and Case 189 (SyRI).

| DOMAIN                              | MODES ADDRESSED   | LENS COMPETENCY                                   | PROBLEM TYPE | INDUCED | LEO   |
|-------------------------------------|-------------------|---------------------------------------------------|--------------|---------|-------|
| HIGHER EDUCATION<br>(LATIN AMERICA) | <b>K</b> <b>N</b> | <b>D1</b> <b>D2</b> <b>D3</b> <b>D4</b> <b>D5</b> | 4            | 5.1     | 5 · 4 |
| LEARNING<br>ANALYTICS               | →                 | → Navigating Sociotechnical Constraints           |              |         |       |
| CROSS-NATIONAL<br>GOVERNANCE        |                   |                                                   |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Begin from the documented failure mode — Global-North tools that did not transfer — so the adoption project is designed against the actual local barrier, not against an imagined universal one.
- ▶ Run the participatory method with administrators, teachers, and students separately; the decisions, the daily use, and the consent each sit with a different stakeholder.
- ▶ Convert the governance question from a yes/no deployment gate into a structured set of decisions surfaced during adoption; the CANVAS is the worked artifact of that conversion.

### IN OPERATION / AFTER

- ▶ Report the evidence at its tier: adoption-readiness and capacity-building is what the published study establishes; long-run student-outcome evidence is the next study, and the case names that gap rather than glossing it.
- ▶ Commission longitudinal follow-up at the institutions that adopted via the CANVAS to convert capacity-building evidence into deployed-outcome evidence over time.
- ▶ Carry the cross-regime lesson into the broader v2 governance set: governance-by-design is transferable, but the local participatory process is the half that does not transfer; each regime has to do it for itself.



## REFLECTION QUESTIONS

1. Identify a tool or framework in your domain that has been lifted from one regime to another without local adaptation. What participatory method would surface what the destination stakeholders actually need, and what would convert the governance question into a structured set of decisions during adoption rather than a gate at deployment?
2. Specify the evidence tier honestly for an adoption-readiness case in your context: capacity-building and process evidence is what you can publish now; long-run outcome evidence is the next study. Where in your communication is the tier most at risk of being smoothed away?

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

# Norway's National Expert Commission on Learning Analytics

**G** Governance & Trust

**K** Knowledge & Institutional Memory

**N** Normalization of Deviance

**IMPACT** Rather than let learning analytics diffuse unregulated or block it, Norway's Ministry of Education convened a national Expert Commission to investigate the pedagogical, ethical, legal, and privacy issues and establish a regulatory foundation before sector-wide deployment; interim report June 2022, final report 2023 (NOU), with central dilemmas explicitly framed

## IN BRIEF

Norway's Ministry of Education convened a national Expert Commission on Learning Analytics in September 2021 — a national-government response to a capability deployment question at sector scale. Rather than let learning analytics diffuse unregulated across Norwegian education, or block it on precautionary grounds, the ministry chose to construct the governance architecture first. The commission delivered an interim report to the Minister in June 2022 and a final report in 2023 (the NOU, Norges offentlige utredninger, the canonical form of Norwegian public commission reports) identifying central dilemmas across the pedagogical, ethical, legal, and privacy dimensions. The case is governance-as-deliberate-artifact at national scale — a country treating change-control and disclosure as the precondition for adoption rather than the consequence of it. The honest limit preserved verbatim: the commission's recommendations were guidance to a ministry, and downstream sector outcomes — whether deployed Norwegian LA systems actually improved learning outcomes or preserved trust under operation — are not yet documented. This is process-level evidence (a national governance artifact exists, the dilemmas are named), not yet deployment-outcome evidence. The case pairs with the OU (Case 81) and SyRI (Case 189) cases as the national Nordic complement to the institutional-UK and judicial-Dutch governance modes.

## BACKGROUND

Learning analytics at sector scale presents a national government with a structural choice. Let the technology diffuse across schools and universities under existing patchwork governance, or block it on precautionary grounds until questions are settled (a path some jurisdictions have pursued for specific applications), or construct the governance architecture first and let deployment follow under it. Norway in 2021 chose the third path. The Ministry of Education convened a national Expert Commission on Learning Analytics to investigate the pedagogical, ethical, legal, and privacy issues across the whole education sector and to establish a regulatory foundation for what sector-scale deployment should look like. The commission was appointed in September 2021 and worked for two years.<sup>o</sup>

## THE INTERVENTION

The commission's mandate covered the full chain. Pedagogically, what kinds of learning-analytics-driven interventions are defensible at primary, secondary, and tertiary levels; what claims about outcome the evidence supports; what teacher-student relationship the analytics should and should not be allowed to alter. Ethically, what is owed to students whose data drives the analytics, what consent architecture is defensible across age groups, and what disclosure structure the analytics-driven decisions should carry. Legally, how the Norwegian data-protection regime under GDPR interacts with the educational context, and what additional sectoral instruments are needed. Privacy, where the line is between pedagogically useful data and surveillance overreach.<sup>1</sup>

## HOW IT WORKED

The commission delivered an interim report to the Minister of Education and Research in June 2022 and a final report in 2023 in the canonical NOU form. The reports name central dilemmas the field has to live with rather than resolve once: the need for information against data protection; learning as an individual process against learning as a social process; centralization against autonomy; and competence needs against competence reality. The reports' framing is that governance for learning analytics is the kind of artifact that has to be revisited as the technology and the evidence base change, not a one-time document.<sup>2</sup>

## THE EVIDENCE

The honest evidentiary state is process-level. The case establishes that a national government can produce a structured governance artifact, identify the dilemmas sector-scale deployment will face, and deliver the architecture as guidance to a ministry. The case does **not** yet establish that downstream Norwegian sector outcomes — improved learning, preserved trust, defensible interventions — have been delivered, because deployment under the new architecture is too recent for outcome evidence. This is a governance-process success, not yet a measured deployment-outcome success, and the case carries the qualification rather than collapsing it.<sup>3</sup>

## WHAT TRANSFERRED

In pair with the Open University (Case 81, institutional governance-by-design, UK) and SyRI (Case 189, judicial rights-grounded halt, Netherlands), Norway's commission is the national-scale governance-architecture mode. The three cases together teach that learning-analytics governance can be produced at the institutional level (OU), constrained by the courts after the fact (SyRI), or constructed by national deliberation before sector deployment (Norway). The structural lesson is that governance is producible as a deliverable at whichever level matches the deployment scope, and the choice of level is itself a governance decision the case literature names.

*The commission did not resolve learning-analytics governance for Norway. It named the dilemmas the sector will live with and produced the artifact deployment can be governed under.*

EDITORS' SYNTHESIS OF THE NORWEGIAN EXPERT COMMISSION INTERIM AND FINAL REPORTS.

## REFERENCES

1. Ekspertgruppen for digital læringsanalyse (2022), Læringsanalyse — noen sentrale dilemmaer, interim report delivered to the Minister of Education and Research, June 2022.

2. NOU 2023: 19, Læring, hvor ble det av deg i alt mylderet? Bruk av elev- og studentdata for å fremme læring (June 2023) — the commission's final report.

3. Wasson, Giannakos, Blikstad-Balas, Uppstad, Langford & Bøhn (2024), "Implementing Learning Analytics in Norway: Four Central Dilemmas," Journal of Learning Analytics 11(2):268–280, doi:10.18608/jla.2024.8241 — written by commission members; the interim report and its four dilemmas (data, learning, governance, competence).

4. Hilliger et al. (2020), Internet and Higher Education — the LALA companion at multi-country participatory scale (Case 91).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Norway's national Expert Commission is the governance-architecture-at-national-scale instance: a country constructing the regulatory and pedagogical foundation before sector-scale learning-analytics deployment, rather than after diffusion or via judicial halt. The artifact exists and the dilemmas are named; downstream sector outcomes are not yet documented. Process-level success; deployment-outcome evidence is the next decade's work.

LENS APPROACH

Norway is the national-scale governance-architecture case (induced 5.4; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the change- control-and-disclosure-as-artifact discipline and in Domain 4 (Test and Evaluation) for the subobjective **Judgment under inadequate evidence** (4.4) on the process-vs-outcome tier. Pair with Case 81 (OU, institutional), Case 189 (SyRI, judicial), Case 91 (LALA, multi-country participatory), and Case 89 (African data privacy, frontier) — the non-US LA governance pentad teaching the level-of-governance decision.

| DOMAIN             | MODES ADDRESSED | LENS COMPETENCY | PROBLEM TYPE | INDUCED | LEO   |
|--------------------|-----------------|-----------------|--------------|---------|-------|
| EDUCATION (NORWAY) | G K N           | D1 D2 D3 D4 D5  | 4            | 5.4     | 5 · 4 |
| NATIONAL POLICY    |                 |                 |              |         |       |
| LEARNING ANALYTICS |                 |                 |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Match the governance artifact to the deployment scope: institutional governance for a single university, national commission for sector-scale deployment, judicial review when neither operates. The choice of level is itself a governance decision.
- Convene the commission with the dimensional breadth the deployment actually has: pedagogical, ethical, legal, privacy — not only the dimension the most-visible objection sits on.
- Frame the artifact as living: governance for a moving technology has to be revisited as the technology and the evidence base change, and the artifact should say so.

IN OPERATION / AFTER

- Report the evidence at its tier: governance-process success is what the commission and NOU establish; deployment-outcome success is the next decade's work, and the case is honest about the gap.
- Carry the central dilemmas into the implementation conversation rather than treating them as resolved by publication of the NOU; the dilemmas are what implementation will live with.
- Use the case as the national-scale instance in the OU / LALA / SyRI / Norway non-US LA governance set; the joint teaching point is that governance is producible at whichever level matches the deployment scope.

## REFLECTION QUESTIONS

1. Identify a sector-scale capability deployment question in your domain. Is the right governance level institutional, sectoral, national, or judicial — and what determines the match between deployment scope and governance level?
2. Specify the dimensional breadth your commission or governance artifact would have to cover. Norway's mandate was pedagogical, ethical, legal, and privacy. What is the analog in your context, and which dimension is most at risk of being narrowed away under stakeholder pressure?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

# Singapore SkillsFuture – National Workforce Capability at Scale

**G** Governance & Trust   **K** Knowledge & Institutional Memory   **D** Designed Out

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** Singapore's SkillsFuture pairs individual training credits with employer subsidies, a cross-sector skills framework, and a two-wave outcome survey (TRAQOM, at end-of-course and at three months — six months for courses ending before June 2024) — an MTI study of 2011–2014 cohorts found a 5.8% real wage premium for WSQ full-qualification trainees and 0.8% for Statement-of-Attainment trainees, with 87% of Work-Study Programme graduates employed full-time within six months

## IN BRIEF

Singapore's SkillsFuture Movement, launched in 2015, pairs individual training credits with employer subsidies, a cross-sector skills framework, and one of the most ambitious national L&D measurement instruments deployed at scale: the two-wave TRAQOM survey, administered at end-of-course and at six months post-training, paired with labor-market data. Self-reported figures from the 2024 Year-in-Review are strong: 98% of trainees report being able to perform better at work; 93% report the course played a pivotal role. A separate Work-Study Programme Outcomes Survey found 87% of graduates employed full-time within six months (2019 cohort; the 2024 review reports more than 9 in 10). A 2018 MTI study found a 5.8% real wage premium for WSQ-trained workers. The honest reading the case carries into print: self-report dominates the headline numbers, and the program has not been subjected to a rigorous quasi-experimental external evaluation that would isolate the program's causal effect from underlying labor-market trends. SkillsFuture is the non-US national-scale L&D case the corpus needs for both the corporate / workforce L&D gap and the non-US/UK/EU gap. The evidence-tier flag is practice-synthesis: the program design and the TRAQOM instrument are documented in SSG annual reports and in ILO and Springer analyses, the headline outcomes are self-report, and future validation — particularly quasi-experimental causal evaluation — is ongoing.

## BACKGROUND

SkillsFuture was launched in 2015 as a Singapore-wide workforce-capability program at the seam between individual upskilling, employer demand, and state coordination. The program design pairs individual training credits (SkillsFuture Credit), employer subsidies for workforce training, a cross-sector skills framework that defines competencies and progression

paths across industries, and a Work-Study Programme that embeds learners in employer contexts during training.<sup>0</sup>

#### THE INTERVENTION

The measurement instrument is unusually ambitious for a national L&D program. The Training Quality and Outcomes Measurement framework (TRAQOM) is a two-wave outcome survey administered at end-of-course and at six months post-training. Employment and wage outcomes are measured separately — by the Work-Study Programme Outcomes Survey and by MTI's administrative-data studies — not inside TRAQOM itself. The design crosses the Kirkpatrick Level-2 / Level-3 seam (Case 79) at policy level, not only program level.<sup>1</sup>

#### HOW IT WORKED

The 2024 Year-in-Review reports headline figures: 98% of trainees report being able to perform better at work; 93% report the course played a pivotal role. A separate Work-Study Programme Outcomes Survey found 87% of graduates employed full-time within six months (2019 cohort; the 2024 review reports more than 9 in 10 employed). A Ministry of Trade and Industry study of the 2011–2014 cohorts found a 5.8% real wage premium for WSQ full-qualification holders and 0.8% for holders of a Statement of Attainment. The labor-market figures are the strongest available external corroboration of the self-report data.<sup>2</sup> Since 2024 the program has broadened well beyond credits and surveys: the SkillsFuture Level-Up Programme gave Singaporeans aged 40 and over a S\$4,000 credit top-up (May 2024) and, from 2025, a training allowance of up to S\$3,000 a month for full-time reskilling, while the SkillsFuture Jobseeker Support scheme (launched April 2025) added up to S\$6,000 over six months for the involuntarily unemployed — extending the model from upskilling into income-supported labor-market transition.

#### THE EVIDENCE

The honest reading is the teaching point. Self-report dominates the headline outcomes. The program has not been subjected to a rigorous quasi-experimental external evaluation that would isolate the program's causal effect from underlying labor-market trends — and Singapore's labor market has been strong across the program's deployment period. The TRAQOM design is one of the strongest national L&D instruments deployed, and what it cannot yet do is what no national L&D instrument yet does well: produce decision-grade causal evidence at population scale. Future validation is ongoing.<sup>3</sup>

#### WHAT TRANSFERRED

The LENS teaching point is that the program is a non-US national-scale case for the corporate / workforce L&D cluster (Cases 79, 65, 83, 70) and a non-US/UK/EU case for the geographic-coverage gap. The amended LEO on collaboration measurement is directly exercised: TRAQOM measures across employer-employee-state, not only across the training organization. Pair with Case 18 (PEPFAR) for the global-health workforce-capability counterpart, and with Case 70 (HILS) for the design-side practitioner pattern that the SSG program operationalizes at policy scale. Evidence-tier flag is practice-synthesis; the design is documented, the causal magnitudes are open.<sup>4</sup>

*The instrument crosses the Level-2/Level-3 seam at policy level. What it cannot yet do is what no national L&D instrument yet does well.*

EDITORS' SYNTHESIS OF THE SKILLSFUTURE MOVEMENT AND THE TRAQOM MEASUREMENT FRAMEWORK.

REFERENCES

1. SkillsFuture Singapore (SSG), Year-in-Review 2024 — program metrics and outcome reporting.

2. Ministry of Education (MOE), Singapore, parliamentary replies on TRAQOM, 2020.

3. Ministry of Trade and Industry (MTI), Singapore, 2018 — WSQ wage-premium study.

4. "Future-Skilling the Workforce: SkillsFuture Movement in Singapore," Springer, 2024 — peer-reviewed program analysis.

5. International Labour Organization (ILO), "Investigating an Upskilling Programme in Singapore" — international comparative analysis.

THE LEARNING ENGINEERING LENS

LE INSIGHT

SkillsFuture is the non-US national-scale L&D case the corpus needs. The TRAQOM instrument is among the most ambitious national L&D measurement architectures deployed; the headline outcomes are self-report dominant; no rigorous quasi-experimental external evaluation yet exists. Evidence-tier flag is practice-synthesis; future validation is ongoing.

LENS APPROACH

SkillsFuture is the national workforce-capability case (induced 2.4; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the amended LEO on collaboration measurement — TRAQOM measures across employer-employee-state — and in Domain 2 as the policy-scale operationalization of the HILS-style environment-and-event integration (Case 70). Pairs with Case 18 (PEPFAR) for the global-health workforce-capability counterpart.

| DOMAIN              | MODES ADDRESSED | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO       |
|---------------------|-----------------|---------------------------------------|--------------|---------|-----------|
| WORKFORCE DEV       | <div>GKD</div>  | <div>D1D2D3D4</div>                   | 4            | 2.4     | 2 · 5 · 3 |
| NATIONAL L&D POLICY | →               | → <div>D5</div>                       |              |         |           |
| ASIA-PACIFIC        |                 | Navigating Sociotechnical Constraints |              |         |           |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the measurement instrument across the training-and-employment seam, not within the training organization alone — TRAQOM's two-wave + labor-market cross-check is the architecture the LENS Domain 4 module should teach.
- ▶ Make the cross-sector skills framework a first-class artifact — without it the credits, the subsidies, and the Work-Study Programme do not cohere as a single workforce-capability deliverable.
- ▶ Treat the self-report dominance honestly: name what TRAQOM can and cannot establish at the design stage, so the program documentation does not have to retrofit the hedge.

IN OPERATION / AFTER

- ▶ Pair with Case 18 (PEPFAR) for the global-health workforce-capability counterpart at multi-country scale; together they teach what national- and program-scale L&D measurement at evidence-flagged tier looks like.
- ▶ Use the amended LEO on collaboration measurement: TRAQOM is a worked example of measurement across employer-employee-state, and the program documentation can teach the architecture in LENS Domain 5 (Sociotechnical Constraints).
- ▶ Carry the practice-synthesis flag honestly: the program design and the TRAQOM instrument are documented, the headline magnitudes are self-report, and future validation requires independent quasi-experimental causal evaluation.



## REFLECTION QUESTIONS

1. Identify a workforce-capability program in your context that currently measures at the training-organization boundary. What would the analogue of TRAQOM — a two-wave outcome survey paired with employment-and-wage data at population scale — require of your measurement infrastructure?
2. Specify the cross-sector skills framework that would coordinate individual credits, employer subsidies, and a Work-Study-style placement program in your context. Without the framework, do the components cohere as a single capability deliverable?
3. SkillsFuture's headline magnitudes are self-report dominant. What independent quasi-experimental evidence — comparison-cohort design, regression discontinuity, instrumented variation — would you require before treating any specific outcome magnitude as decision-grade for a program-scale investment in your context?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · task analysis · design prototyping

# Learning Analytics on the African Continent — A Scoping Review as the Present-State Map

**K** Knowledge & Institutional Memory    **N** Normalization of Deviance

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** A 2022 scoping review found only 15 learning-analytics studies on the entire African continent, overwhelmingly from South Africa (10 of 15 studies), with the remainder from Nigeria, Tanzania, Zambia, and Kenya; the structural finding — limited LMS access, limited institutional resourcing, limited African-scholar visibility at SoLAR — is itself the evidence the field requires before construct-travel claims can be made

## IN BRIEF

Prinsloo and Kaliisa (2022) published a scoping review of learning-analytics research on the African continent for the Journal of Learning Analytics. The review found only 15 studies meeting inclusion criteria, with publication output overwhelmingly from South Africa (10 of 15) and the remainder from Nigeria, Tanzania, Zambia, and Kenya. The structural findings — limited LMS access in many African higher-education institutions, limited institutional resourcing for learning-analytics infrastructure, and limited African-scholar visibility at the Society for Learning Analytics Research (SoLAR) conferences — are the present-state map the field needs before importing U.S./UK/EU learning-analytics constructs into African contexts. The case is included as a frontier scoping case at the practice-synthesis tier: a review of an early-stage research base where the absence of dense primary studies is itself the finding. It pairs with Case 89, the African data-privacy governance case, to articulate the construct-travel problem in both research-base and governance terms. Future validation ongoing as the African learning-analytics literature matures.

## THE SHIFT

Learning analytics as a field consolidated around 2011 with the formation of SoLAR and a research agenda anchored heavily in U.S., Australian, UK, and EU higher-education data infrastructures. The capability question for the African higher-education community is whether the analytic constructs travel — and who is positioned to validate them when they arrive. Prinsloo and Kaliisa set out to map what learning-analytics research had actually been published from African institutions at the time of writing.<sup>o</sup>

## WHAT IS EMERGING

The headline finding is sparse. The scoping review located only 15 learning-analytics studies on the African continent, with publication output overwhelmingly from South Africa (10 of 15), the rest from Nigeria, Tanzania, Zambia, and Kenya. Two of the fifteen are the South African higher-education studies most often cited as the local literature — Lemmens and Henn (2016) and Janse van Vuuren's 2020 contextualised data-analytics framework — so citing them does not enlarge the corpus. The corpus exists; it is small, geographically concentrated, and dwarfed by the parent field's publication output from U.S./UK/EU institutions.<sup>1</sup>

## THE CAPABILITY QUESTION

The structural findings underneath the count are the decisive teaching. African higher-education institutions face limited LMS access at scale (many institutions still rely on heterogenous and intermittent digital infrastructure); limited institutional resourcing for the data-engineering, ethics-review, and analyst capacity that learning analytics presupposes; and limited African-scholar visibility at SoLAR and related venues — of the fifteen studies, one appeared at LAK and none in the Journal of Learning Analytics, and ten are empirical against five conceptual. The construct-travel problem is not abstract — it is the mismatch between what the analytic frameworks assume about data availability and what the institutions can actually generate.<sup>2</sup>

## EARLY EVIDENCE

The case is positioned as frontier rather than failure or intervention because it does not document a deployed capability outcome. It documents the state of the evidence base, which is itself the relevant capability question for anyone considering whether to deploy a learning-analytics program in an African higher-education institution. The review surfaces the construct-validation work that has not yet been done; the absence of dense primary studies is the finding, and the implication is that any deployer must either undertake the validation work themselves or import constructs with the validation gap acknowledged.<sup>3</sup>

## OPEN PROBLEMS

The practice-synthesis-tier flag is binding. A scoping review is the field's snapshot at a moment in time; the individual primary studies it draws on vary in rigor, and the review is the strongest single map of the present state but not the final account. Future validation ongoing as the literature matures, as institutional infrastructure improves, and as African scholars author more of the learning-analytics research that travels out of the continent rather than into it. The case carries the standing tier language into print.

*The absence of dense primary studies is itself the finding.*

EDITORS' SYNTHESIS OF PRINSLOO AND KALIISA (2022).

## REFERENCES

1. Prinsloo, P., & Kaliisa, R. (2022), "Learning Analytics on the African Continent: An Emerging Research Focus and Practice," Journal of Learning Analytics 9(2):218–235, doi:10.18608/jla.2022.7539 — 15 studies (10 empirical, 5 conceptual); ten involve South Africa, two Nigeria, one each Kenya, Tanzania, Zambia.
2. Lemmens, J. C., & Henn, M. (2016), "Learning analytics: A South African higher education perspective," in Botha & Muller (eds.),
3. Janse van Vuuren, E. C. (2020), "Development of a contextualised data analytics framework in South African higher education: involvement of teacher (teaching) analytics as an indispensable component," South African Journal of Higher Education 34(1) — one of the fifteen included studies.
4. Cross-reference: Case 89 (Prinsloo & Kaliisa's companion African data-privacy paper, BJET 2022), for the construct-travel problem stated in governance terms.

Institutional Research in South African Higher Education, 231–253, SUN PReSS — one of the fifteen included studies.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The scoping review is the present-state map of learning-analytics research on the African continent: 15 studies, geographically concentrated, with the structural barriers (LMS access, institutional resourcing, African-scholar SoLAR visibility) underneath the count. Practice-synthesis-tier — a snapshot of an early-stage literature; future validation ongoing as the field matures.

LENS APPROACH

African learning-analytics scoping is the non-US frontier-evidence case (induced 8.4; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the construct-travel problem and in Domain 4 (Test and Evaluation) for the subobjective **Judgment under inadequate evidence** (4.4) — the deployer must decide on what to import and what to validate locally when the validation literature is thin. Pairs with Case 89 (African data privacy) for the construct-travel problem stated in governance terms.

| DOMAIN                                           | MODES AT STAKE | LENS COMPETENCY                                         | PROBLEM TYPE | INDUCED | LEO   |
|--------------------------------------------------|----------------|---------------------------------------------------------|--------------|---------|-------|
| LEARNING ANALYTICS<br>HIGHER EDUCATION<br>AFRICA | K N            | D1 D2 D3 D4 D5<br>Navigating Sociotechnical Constraints | 4            | 8.4     | 4 · 5 |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ When importing learning-analytics constructs into an African higher-education context, treat the scoping-review map as the precondition: do not assume the construct's validation literature applies and identify which validation work has to be repeated locally.
- ▶ Specify the institutional preconditions the analytic framework presupposes — LMS coverage, data-engineering capacity, ethics-review infrastructure — and audit them honestly against the deploying institution's actual conditions.
- ▶ Carry the practice-synthesis-tier flag through any decision document the scoping review supports; the review is the strongest current map but not the final account.

IN OPERATION / AFTER

- ▶ Treat each deployed learning-analytics program in an African institution as a contribution to the construct-validation literature, not only as a service intervention; publish the validation work where SoLAR and related venues can absorb it.
- ▶ Track institutional infrastructure conditions as a separate variable from analytic-construct performance; a construct that travels under one set of conditions may not travel under another.
- ▶ When the literature matures, retire the scoping review as the present-state map and replace it with denser primary-study evidence; the case is included because the present-state matters now, not because the present-state should persist.

## REFLECTION QUESTIONS

1. Identify a learning-analytics construct your program imports from U.S./UK/EU literature. What does the validation literature for that construct assume about LMS coverage, data-engineering capacity, and ethics review? Audit those assumptions against your institution's actual conditions.
2. The case rests on a scoping review of an early-stage literature. What would the minimum additional evidence — replication primary studies, multi-institution validation, longitudinal outcome data — look like before you would treat any specific African learning-analytics finding as settled?
3. Specify the institutional preconditions you would identify and the validation work you would commit to before deploying an analytic framework in a low-resource higher-education setting; treat your deployment as a contribution to the literature, not only as a service intervention.

**WHO BUILDS THIS** knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

# PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop

**T** Training Gap   **G** Governance & Trust

**DISCLOSURE** Institutional overlap: an editor shares an institution (Johns Hopkins School of Education) with leading PBIS researchers; no editor was personally involved. Framed as learning from that peer-reviewed literature.

**IMPACT** A multi-tiered behavior-support framework implemented in more than 25,000 U.S. schools; group-randomized effectiveness trials in 37 Maryland elementary schools found significant reductions in suspensions and office discipline referrals and improvements in child behavior problems when the model was implemented with fidelity; the PBIS Maryland state-university-district partnership (since 1999) trained staff at over 1,000 schools around a standing fidelity-measurement and coaching loop

## IN BRIEF

Positive Behavioral Interventions and Supports (PBIS) is one of the most widely scaled evidence-based frameworks in U.S. education — implemented, by the Center on PBIS's own count, in more than 25,000 schools. The evidence base is unusually strong for education at scale: Bradshaw, Mitchell, and Leaf's five-year group-randomized effectiveness trial in 37 Maryland elementary schools (2010) found that schools trained in school-wide PBIS reached high implementation fidelity and experienced significant reductions in suspensions and office discipline referrals, with companion trials reporting improvements in child behavior problems (Bradshaw, Waasdorp, & Leaf, Pediatrics 2012) and reductions in bullying and peer rejection (Waasdorp et al. 2012). The central finding the implementation-science literature itself reports — and the reason the case is in this book — is that the effect does not live in the framework document. It lives at the interface between the intervention design and the school's implementation capability: fidelity measurement (the School-wide Evaluation Tool, the Tiered Fidelity Inventory), coaching infrastructure, district capacity, and sustained administrative support carry the effect, and schools that adopt the framework without that layer show attenuated or null results. The Maryland scale-up — PBIS Maryland, a state-university-district partnership running since 1999 — is the worked example of scaling **with** the fidelity loop attached. The hedges are the literature's own: fidelity varies widely, sustainability is fragile, and abandonment of effectively implemented practice is commonplace.

## BACKGROUND

The recurring story in school-improvement research is that an intervention validated in trials arrives at a new school as a document — a framework, a binder, a training day — and produces nothing. PBIS is the instructive counter-example because its research community confronted that problem directly and instrumented it. The framework itself, built from applied behavior analysis and organized as a multi-tiered system of support (universal expectations and reinforcement at Tier 1, targeted group supports at Tier 2, individualized supports at Tier 3), was codified in the late 1990s under an OSEP technical-assistance center and has since scaled to more than 25,000 U.S. schools. The question the field then spent two decades answering is the one this casebook cares about: why does the same framework work in some schools and not others?<sup>0</sup>

## THE INTERVENTION

The causal evidence comes principally from the Maryland group-randomized effectiveness trials. Bradshaw, Mitchell, and Leaf (2010) followed 37 Maryland elementary schools over five years; schools randomized to school-wide PBIS training implemented the model with high fidelity and showed significant reductions in student suspensions and office discipline referrals relative to comparison schools. Companion analyses from the same trial infrastructure reported significant improvements in teacher-rated child behavior problems across 12,344 children (Bradshaw, Waasdorp, & Leaf 2012) and reductions in bullying and peer rejection (Waasdorp, Bradshaw, & Leaf 2012). The qualifier that matters is in the trial design itself: the randomized schools did not merely receive the framework — they received structured training, on-site coaching, and annual fidelity assessment, and the published effects are effects of that package implemented with fidelity.<sup>1</sup>

## HOW IT WORKED

The fidelity instruments are the part of the story this book reads as capability measurement. The School-wide Evaluation Tool (SET; Horner et al. 2004) and its successor, the Tiered Fidelity Inventory (TFI; Algozzine et al. 2014, with technical-adequacy validation in McIntosh et al. 2017), turn “is the school actually doing PBIS?” into a scored, repeatable observation — expectations defined and taught, reinforcement systems operating, data-based decision cycles running, administrative leadership engaged. The instruments are not compliance paperwork; they are the closed loop that links the intervention to its evidence. Fidelity scores feed the coaching conversation, coaching moves the scores, and the outcome literature consistently reports that schools implementing below fidelity thresholds — typically those adopting the framework without the coaching and measurement layer — show attenuated or null effects relative to high-fidelity implementers.<sup>2</sup>

## THE EVIDENCE

The Maryland scale-up is the worked example of carrying that loop to state scale. PBIS Maryland, formed in 1999 as a partnership among the Maryland State Department of Education, Sheppard Pratt Health System, and Johns Hopkins University, with all 24 local school systems as partners, has trained staff at over 1,000 Maryland schools. The design decision that distinguishes it from framework-by-memo scale-ups is structural: each district provides local coaching capacity, the state provides training and annual fidelity measurement, and the university partners run the prevention research on the resulting data — including the randomized trials above and a subsequent statewide quasi-experimental study of the

scale-up itself. The scale-up did not distribute a document; it replicated the implementation infrastructure that the trials had shown to carry the effect.<sup>3</sup>

### WHAT TRANSFERRED

The honest framing is the one the literature itself supplies, and the case preserves it as learning rather than critique. Sustainability is the field's own named open problem: McIntosh and colleagues, tracking fidelity across 5,331 schools over five years, sorted them into four patterns — Sustainers (29%), Slow Starters (13%), Late Abandoners (24%), and Rapid Abandoners (34%) — and found middle and high schools markedly likelier than elementary schools to abandon the practice, whether late or fast (McIntosh, Mercer, Nese, & Ghemraoui 2016). Sugai and Horner's 2020 synthesis frames sustaining and scaling as a distinct engineering problem from demonstrating efficacy. That the PBIS research community measured its own implementation variance, published its own null-attenuation conditions, and built the instruments that make the variance visible is precisely why the case sits in the intervention corpus: the field solved — in the open, with peer review — the coupling problem that most deployments in this book never instrumented at all.<sup>4</sup>

*The framework document does not carry the effect. The fidelity loop — measurement, coaching, district capacity, sustained administrative support — is the intervention.*

EDITORS' SYNTHESIS OF THE SWPBIS IMPLEMENTATION LITERATURE.

### REFERENCES

1. Bradshaw, Mitchell, & Leaf (2010), "Examining the Effects of Schoolwide Positive Behavioral Interventions and Supports on Student Outcomes: Results From a Randomized Controlled Effectiveness Trial in Elementary Schools," *Journal of Positive Behavior Interventions*, 12(3), 133 – 148.
2. Bradshaw, Waasdorp, & Leaf (2012), "Effects of School-Wide Positive Behavioral Interventions and Supports on Child Behavior Problems," *Pediatrics*, 130(5), e1136 – e1145.
3. Horner, Todd, Lewis-Palmer, Irvin, Sugai, & Boland (2004), "The School-Wide Evaluation Tool (SET): A Research Instrument for Assessing School-Wide Positive Behavior Support," *Journal of Positive Behavior Interventions*, 6(1), 3 – 12.
4. Bradshaw, Pas, et al. (2012), "A State-Wide Partnership to Promote Safe and Supportive Schools: The PBIS Maryland Initiative," *Administration and Policy in Mental Health*, 39(4), 225 – 237.
5. McIntosh, Mercer, Nese, Strickland-Cohen, & Hoselton (2016), "Predictors of Sustained Implementation of School-Wide Positive Behavioral Interventions and Supports," *Journal of Positive Behavior Interventions*, 18(4), 209 – 218.
6. McIntosh, Massar, Algozzine, George, Horner, Lewis, & Swain-Bradway (2017), "Technical Adequacy of the SWPBIS Tiered Fidelity Inventory," *Journal of Positive Behavior Interventions*, 19(1), 3 – 13.

### THE LEARNING ENGINEERING LENS

#### LE INSIGHT

PBIS is the rare education framework whose research community instrumented its own implementation variance: randomized trials showed real effects on suspensions, referrals, and behavior — conditional on fidelity — and the field built the measurement instruments (SET, TFI), the coaching infrastructure, and the sustainability literature that explain why the same framework works in one school and not another. The effect lives at the coupling between design and implementation capability, and the fidelity instruments are the capability measurement layer.

#### LENS APPROACH

PBIS implementation fidelity is the coupling-and-fidelity intervention case (induced 2.3; LENS D2/PT3). LENS uses it in Domain 2 (Iterative Development) for the fidelity-measurement-to-coaching loop as the mechanism that carries an evidence-based design into practice; in Domain 4 (Test and Evaluation) for the SET/TFI instruments as validated measures of implementation rather than outcome; and in Domain 5 (Navigating Sociotechnical Constraints) for the Maryland state-university-district partnership as the governance architecture that made a 1,000-school scale-up sustain its measurement discipline. The framing is learning



from the implementation–science literature — the field documented its own attenuation conditions — and the COI render under the title is binding.

| DOMAIN                  | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO |
|-------------------------|-----------------|-----------------------|--------------|---------|-----|
| K-12 EDUCATION          | T G             | D1 D2 D3 D4 D5        | 3            | 2.3     | 2   |
| SCHOOL CLIMATE          | →               | →                     |              |         |     |
| IMPLEMENTATION FIDELITY |                 | Iterative Development |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Treat the fidelity instrument as part of the intervention, not as evaluation overhead; a framework shipped without its measurement loop is a different — and weaker — intervention than the one the trials validated.
- ▶ Build the coaching layer at the district, not just the school; the Maryland design routes coaching capacity through every partner district so a principal transition does not orphan the implementation.
- ▶ Score fidelity on a public, repeatable rubric (SET/TFI-style) and feed the scores into the coaching cadence, so the implementation conversation runs on observations rather than impressions.

#### IN OPERATION / AFTER

- ▶ When an evidence-based practice underperforms at a new site, attribute the gap with fidelity data before revising the practice; the PBIS literature shows most of the variance lives in implementation, not design.
- ▶ Plan for sustainability as its own engineering problem: instrument speed-to-fidelity, watch the elevated-risk tiers (middle and high schools, in this literature), and treat abandonment as a measurable failure mode rather than an anomaly.
- ▶ Publish the implementation variance, including the null and attenuated conditions; the PBIS field's credibility rests on having documented where its own framework does not work.

### REFLECTION QUESTIONS

1. Pick an evidence-based practice in your domain that underperformed at a new site. What would its SET or TFI look like — the scored, repeatable observation that distinguishes “the practice failed” from “the practice was never implemented”? Who would administer it, and on what cadence?
2. The Maryland scale-up replicated infrastructure — district coaching capacity, annual fidelity measurement, a university research partner — rather than distributing a framework document. Map your own scale-up plan against those three layers. Which one is missing, and what does the PBIS sustainability literature predict happens without it?
3. McIntosh and colleagues found that most schools fell into an abandoning pattern within five years, and that middle and high schools abandoned fastest. What early signal would flag an abandonment risk in your setting, and is anyone measuring it?

**WHO BUILDS THIS** learning science & instructional design · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · governance frameworks · bias & evidence audits



Chapter 5

# Aviation & Aerospace — What Fails

*When automation, transition, and assumption meet an unforgiving envelope.*

---

*The aircraft did what it was designed to do. So did the crew. The designs disagreed.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

# AeroPerú Flight 603

**H** Human–System Interface   **T** Training Gap

**IMPACT** 70 killed; Boeing 757 crashed into the Pacific after maintenance tape was left over static ports

## IN BRIEF

While polishing the fuselage of a Boeing 757 in Lima in October 1996, maintenance staff taped over the static ports and failed to remove the tape before flight. With the static system blocked, the aircraft's altimeters, airspeed indicators, and altitude-alert and overspeed warnings all read inconsistently. The crew received a cascade of contradictory warnings — overspeed, stick shaker, ground proximity — and could not determine their true altitude or speed. Believing they were higher than they were, they flew into the Pacific, killing all 70 aboard. The investigation named the maintenance error as the primary cause but stressed that every air-data instrument depended on the same blocked sensor: the apparent redundancy was an illusion. One channel stayed correct — the radio altimeters, independent of the air data computers, drove the terrain warnings — and the crew's training had no procedure for "everything you see is wrong."

## BACKGROUND

AeroPerú 603 was a Boeing 757 night departure from Lima to Santiago. Before the flight, ground crew had covered the aircraft's static ports with adhesive tape during cleaning and polishing of the fuselage, and the tape was not removed. The static ports feed the instruments that tell the crew their altitude and airspeed — a single physical source upstream of nearly every primary cockpit display, so blocking it corrupted not one instrument but the whole air-data picture the crew would rely on in the dark.<sup>0</sup>

## WHAT HAPPENED

On 2 October 1996, the aircraft took off into darkness over the ocean with its static system blocked. Altimeters, airspeed indicators, and the altitude-alert and overspeed systems all produced false and contradictory readings. The crew received simultaneous overspeed and stall warnings and ground-proximity alerts they could not reconcile. Believing they were higher than they were, they descended and struck the Pacific, killing all 70 aboard — the contradictory warnings offering no way to tell which, if any, instrument to trust, because every air-data instrument drew on the same blocked source and so failed in concert rather than disagreeing usefully. The true reading arrived as one more alarm in a cockpit where every alarm had already proved false: "too low — terrain" sounded twenty-two times in forty-five seconds about eight minutes before impact, and again through the final minutes to the water.<sup>1</sup>

## THE INVESTIGATION


The Peruvian accident investigation named as the principal cause the error of the maintenance staff, including the crew: the tape was left on, and neither the release inspections nor the pilot-in-command's walk-around caught it. One channel stayed correct. The radio altimeters read height above the sea independently of the air data computers, and it was the radio altimeter that fired the ground-proximity warnings the crew heard for the rest of the flight. The board named the crew's failure to observe the radio altimeter and to fly the GPWS escape procedure as a contributing cause: the radio altimeter was the one altitude source the blockage never reached. The crew's training had no procedure for the case in which all primary instruments are simultaneously wrong — that failure had been assumed away rather than planned for, leaving the crew to improvise against a problem the system design had quietly made possible and the curriculum never named.<sup>2</sup>

## THE CAPABILITY GAP

AeroPerú 603 is the case for redundancy that is not redundancy. The cockpit's apparent instrument redundancy was an illusion at the source — a common-cause failure that defeated every air-data display at once, while the one channel it did not reach went unheeded. The missing capability was both a maintenance control that made the blocked port impossible to miss and a procedure for flying when the instruments themselves cannot be trusted — the first preventing the common-cause failure at its source, the second giving the crew somewhere to stand when the apparent redundancy in front of them collapsed all at once.<sup>3</sup>

## AFTERMATH & REFORM

The accident reinforced industry maintenance practices for conspicuous covering and mandatory removal-verification of static ports and pitot tubes, and entered the human-factors literature as a worked example of common-cause failure and the limits of crew training under total instrument corruption. It is paired with other pitot-static events in discussions of air-data integrity — its lesson being that redundancy counted at the display tells you nothing if the redundant paths share a single upstream point that can fail them all together.<sup>4</sup>



*The crew received contradictory indications they could neither reconcile nor override.*

PARAPHRASING THE PERUVIAN DIRECTORATE GENERAL OF AIR TRANSPORT ACCIDENT INVESTIGATION BOARD REPORT ON AEROPERÚ 603, 2 OCTOBER 1997

REFERENCES

1. Peru Dirección General de Transporte Aéreo, Comisión de Investigación de Accidentes, final report on AeroPerú 603 (2 October 1997; with NTSB/FAA/Boeing participation) — principal cause given as error of the maintenance staff including the crew, and the instrument cascade (paraphrased).

2. Peru DGTA Comisión de Investigación de Accidentes report (2 October 1997) — the masking tape over the three left-side static ports and the contradictory warnings; the tape was not caught at release, at transfer to the apron, or in the pilot-in-command’s walk-around.

3. Peru DGTA Comisión de Investigación de Accidentes report (2 October 1997) — the radio altimeter, independent of the air data computers, working normally and driving the GPWS terrain warnings; crew failure to follow the GPWS procedure and to read the radio altimeter named as contributing. CVR: “too low — terrain” 22 times in 45 seconds beginning about eight minutes before impact.

4. Leveson, N. (2011), Engineering a Safer World (STAMP) — common-cause failure.

5. Dismukes, Berman & Loukopoulos (2007), The Limits of Expertise — crew performance under instrument failure.

THE LEARNING ENGINEERING LENS

LE INSIGHT

AeroPerú 603 is the case for redundancy that is not redundancy. Every air-data indicator depended on the same physical sensor, and the apparent redundancy was an illusion at the source. One independent channel — the radio altimeter driving the terrain warnings — stayed correct and was disbelieved inside a flood of false alarms. The training did not include the failure case because the failure case had been assumed not to occur.

LENS APPROACH

LENS uses AeroPerú in LEN 5 to teach the difference between apparent and actual redundancy and in LEN 2 for the human role when all interface data is unreliable. Studio projects examine the “trust nothing” procedure design.

| DOMAIN           | FAILURE MODES               | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|------------------|-----------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| AVIATION         | → <div>H</div> <div>T</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 1            | 3.1     | 1   |
| Systems Analysis |                             |                                                                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Map redundant displays back to their upstream sources, and eliminate or instrument any common point whose failure would corrupt them all at once.
- ▶ Make critical sensor covers conspicuous and require verified removal, so a blocked static port cannot reach takeoff unnoticed.
- ▶ Provide an independent reference (e.g., a source not fed by the common component) so the crew can detect inconsistency at the source.

IN OPERATION / AFTER

- ▶ Audit maintenance controls for tasks that cover or block critical sensors, confirming each has a conspicuous marker and a removal-verification step.
- ▶ Train and rehearse a “trust nothing” procedure for total instrument corruption, the case the curriculum had assumed away.
- ▶ Monitor for common-cause failure modes whose probability was assumed negligible, and revisit that assumption when the consequence is catastrophic.

### REFLECTION QUESTIONS

1. Identify a redundant interface in your domain whose redundancy depends on a common upstream component. What is the operator procedure when the upstream component fails?
2. Design the “trust nothing” procedure for the AeroPerú crew. What information remains reliable when all instruments are corrupted?
3. A taped-over static port survived to takeoff because nothing made it conspicuous or forced its removal to be verified. Design the maintenance control in your domain that would make a blocked or covered critical component impossible to overlook before operation.

**WHO BUILDS THIS** human factors & interaction design · learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · scenario simulation · competency models

# Boeing 737 MAX / MCAS

**D** Designed Out   **T** Training Gap   **H** Human–System Interface

**IMPACT** 346 killed across two crashes — Lion Air 610 and Ethiopian Airlines 302; 20-month worldwide grounding; ~\$20B direct cost; FAA delegated-authority reform under the Aircraft Certification, Safety, and Accountability Act (2020)

## IN BRIEF

The Boeing 737 MAX was a re-engined 737 meant to fly without new pilot training — the commercial promise that sold it against the Airbus A320neo. Its bigger, forward-mounted engines changed the jet's pitch behavior, so Boeing added software (MCAS) to mask the difference, then kept it out of the manuals and training and let it fire repeatedly on a single angle-of-attack sensor. When that sensor failed on Lion Air 610 (October 2018), MCAS forced the nose down against a crew never told it existed; four months later, on Ethiopian 302, a crew warned of the symptom but not the system ran Boeing's cutout procedure and still could not recover; 346 people died and the fleet was grounded for twenty months. Five major investigations — the NTSB-supported KNKT (Lion Air) and EAIB (Ethiopian) reports, the U.S. House Transportation Committee final report, the DOT Inspector General review, and the multinational Joint Authorities Technical Review (JATR) — converged on the same mechanism: an MCAS that could fire the nose down on a single failing sensor, waved through an FAA delegated-authority regime in which Boeing reviewed much of its own safety judgment. The House Transportation Committee's report went further, documenting the training omission as a cost-driven commercial choice — Boeing had promised Southwest a \$1-million-per-plane rebate if simulator training proved necessary. The MAX is the book's inverse case: human capability engineered out of a system to save the price of sustaining it, with the elision underwritten by a certification process that did not have the independence to catch it.

## BACKGROUND

The 737 MAX was Boeing's answer to the Airbus A320neo: a re-engined version of its best-seller that airlines could fly without retraining pilots on a new type. But the MAX's larger, forward-mounted engines changed how it pitched at high angles of attack, so Boeing added software — the Maneuvering Characteristics Augmentation System, MCAS — to make it handle like the older 737 NG, papering over an aerodynamic change with a control law so the airframe would feel identical from the cockpit. The whole commercial logic depended on that software staying invisible: the less of a "new function" MCAS seemed, the less the MAX would trigger costly new training, and the lower the airline's switching cost away from the A320neo. Boeing reportedly promised Southwest a rebate of about a million dollars per jet if simulator training proved necessary — a clause that turned the training requirement into a direct line on the program's ledger, and one the House investigation later cited as



a structural reason the “no new training” promise behaved as a binding constraint on the engineering rather than as an aspiration.<sup>0</sup>

#### WHAT HAPPENED

To keep MCAS in the background, it was left out of the manual and the type-rating training, and allowed to fire on a single angle-of-attack sensor with no second source to cross-check it. The system’s authority was also expanded late in development — the trim it could command grew from 0.6° to 2.5° per cycle, and its firing was made repeatable rather than one-shot — without the corresponding reassessment of the failure modes that change implied. When the sensor failed on Lion Air Flight 610 in October 2018, MCAS repeatedly forced the nose down; the crew, never told the system existed, fought it cycle after cycle until the jet dove into the Java Sea, killing all 189. Four months later Ethiopian Airlines Flight 302 met the same sequence — sensor failure, repeated trim commands, unrecoverable nose-down attitude — with the difference that this crew had Boeing’s post-Lion-Air bulletin and the FAA emergency directive, ran the runaway-stabilizer cutout they prescribed, and still lost the aircraft, killing 157, for 346 dead across the two crashes, and the entire MAX fleet grounded worldwide for what would become twenty months.<sup>1</sup>

#### THE INVESTIGATION

The investigations made the decisions legible. The Indonesian KNKT and Ethiopian EAIB accident reports, with NTSB participation, established the accident sequences and the single-sensor design as the proximate cause. The multinational Joint Authorities Technical Review (JATR), convened by the FAA in 2019 with international regulators, concluded that MCAS had not been evaluated as the novel flight-control function it actually was. Internal 2013 Boeing meeting minutes showed employees noting that calling MCAS a new function would bring “greater certification and training impact” — the cost the program was built to avoid, written down years before either crash; a 2016 internal survey found 39% of certification staff felt undue management pressure. The House Transportation Committee’s 2020 final report concluded Boeing’s assumption that simulator training was unnecessary “diminished safety, minimized the value of pilot training, and inhibited technical design improvements,” naming the omission as a choice rather than an oversight.<sup>2</sup> The DOT Inspector General traced how the single-sensor design and the training omission passed through a certification process in which the FAA had delegated much of the safety judgment back to Boeing itself under the Organization Designation Authorization (ODA) program — so the company effectively reviewed its own most consequential trade-offs, and the regulator lacked both the technical depth and the institutional independence to challenge the assessment.<sup>3</sup>


#### THE CAPABILITY GAP

The MAX inverts this book’s usual case. Human capability was not overlooked by accident; it was deliberately engineered out to avoid the cost of the training that would have created it. The pilots were not undertrained by oversight but by design — the absence of training was, in effect, a contract deliverable promised to customers and protected by a rebate. The single-sensor architecture was a second engineered-out capability: a redundant cross-check on the angle-of-attack signal would have triggered the kind of design and certification work the program was built to avoid, and was left out for the same reason. Seen together, the crashes are not a training failure that befell a good airplane but exactly what the commercial and engineering decisions specified: a flight-control system that depended on pilots reacting

correctly, in seconds, to a failure whose cause was withheld from them — unnamed in the manuals before Lion Air, and unnamed in the emergency directive after it — with no sensor cross-check that could keep the failure from arriving in the first place.<sup>4</sup>

#### AFTERMATH & REFORM

The MAX was grounded for some twenty months — the longest grounding of a U.S.-certified airliner in the jet era. MCAS was redesigned to use both AOA sensors, to fire once rather than repeatedly, and with limited authority bounded by airspeed — restoring the margins the original had stripped away. The FAA's 2020 return-to-service directive required the simulator training the airplane had been built to avoid. The Aircraft Certification, Safety, and Accountability Act of 2020 then tightened FAA oversight of the ODA delegation regime, required disclosure to pilots of systems that move the flight controls without pilot input, and funded FAA engineering capacity that the delegation system had let atrophy. Boeing entered a Deferred Prosecution Agreement with the DOJ on a fraud charge tied to its disclosures to the FAA's Aircraft Evaluation Group.<sup>5</sup> That agreement's oversight was tested when a door plug blew off a separate 737 MAX 9 (Alaska Airlines 1282) in January 2024, days before the DPA's probation period lapsed; the DOJ found Boeing had breached the agreement, a 2024 guilty plea was rejected by the court, and in May 2025 the parties settled on a non-prosecution agreement worth more than \$1.1 billion — no criminal conviction, over the objection of the crash victims' families. The reform conceded the point the program had spent years resisting: the training was a real requirement all along, the sensor cross-check was a real requirement all along, and removing them from the paperwork only deferred the cost — and then moved it onto two airplanes full of people.



*Boeing's assumption that simulator training was unnecessary diminished safety, minimized the value of pilot training, and inhibited technical design improvements.*

HOUSE TRANSPORTATION COMMITTEE REPORT, 2020

#### REFERENCES

1. U.S. House Committee on Transportation and Infrastructure, *The Boeing 737 MAX Aircraft: Costs, Consequences, and Lessons from Its Design, Development, and Certification* (Final Committee Report, Sept. 2020) — the MCAS omission, the 2013 minutes, the 2016 survey, and the “diminished safety” conclusion.
2. U.S. House Committee on Transportation and Infrastructure, *Boeing 737 MAX 8: A Review of the FAA’s Oversight of the Aircraft’s Design, Development, and Certification* (Final Committee Report, Sept. 2020) — the MCAS omission, the 2013 minutes, the 2016 survey, and the “diminished safety” conclusion.
3. U.S. House Committee on Transportation and Infrastructure, *Boeing 737 MAX 8: A Review of the FAA’s Oversight of the Aircraft’s Design, Development, and Certification* (Final Committee Report, Sept. 2020) — the MCAS omission, the 2013 minutes, the 2016 survey, and the “diminished safety” conclusion.
4. U.S. Department of Transportation, Office of Inspector General, *correspondence and review of FAA oversight of MCAS and the angle-of-attack systems (2019–2020) — the delegated-certification process*.
5. J. Herkert, J. Borenstein & K. Miller, “The Boeing 737 MAX: Lessons for Engineering Ethics,” *Science and Engineering Ethics* 26 (2020) — certification-by-omission as an engineering-ethics failure.

2.

U.S. House report (2020) and MCAS design summary — MCAS firing on a single angle-of-attack sensor; the Lion Air 610 (189 killed) and Ethiopian 302 (157 killed) accident sequences.
3.

U.S. House Transportation Committee report (2020) — Boeing internal communications and the certification-and-training-impact reasoning (quoted).
6.

Federal Aviation Administration, Summary of the FAA's Review of the Boeing 737 MAX (Return-to-Service, 2020) — the MCAS redesign and the simulator-training requirement; and the Aircraft Certification, Safety, and Accountability Act (2020).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The 737 MAX is the documentary record of a design choice to remove human capability to avoid the regulatory and commercial cost of sustaining it. The pilots were not undertrained by accident — they were undertrained as a **requirement**. The omission of training was a contract deliverable. Once that is understood, the accident becomes legible as an engineering decision rather than a training failure.

LENS APPROACH

The 737 MAX is the canonical engineered-out-capability failure (induced 1.1; LENS D1+D3/PT4). LENS uses it in Domain 1 (Systems Analysis) for the LEO-1 work of decomposing system performance requirements into measurable human-capability requirements: the elided pilot-training requirement is the traceable artifact a capability-requirements analysis would have surfaced before the omission could become a contract deliverable. LENS uses it in Domain 5 (Navigating Sociotechnical Constraints; LEO-5) for the delegated-authority regulatory regime in which Boeing reviewed its own most consequential safety judgments — the case is the governance counterpart to Therac-25 (Case 1) at the safeguard-removed-with-nothing-in-its-place layer, and pairs with Patriot/Dhahran (Case 129) at the layer of design assumptions that do not travel with a change of context.

| DOMAIN                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO   |
|-----------------------|---------------|------------------|--------------|---------|-------|
| AVIATION              | → D T H       | → D1 D2 D3 D4 D5 | 4            | 1.1     | 1 · 5 |
| Systems Analysis + D5 |               |                  |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat any control law that masks a handling change as a new function in its own right — document it, surface it to pilots, and size its training as part of the design, not after it.
- ▶ Forbid safety-critical authority from resting on a single sensor; require an independent source or a cross-check before software can command the flight surfaces.
- ▶ Quarantine commercial commitments — rebates, “no new training” promises — from the engineering judgment about what capability the airplane actually requires.

IN OPERATION / AFTER

- ▶ Audit delegated certification so the party making a trade-off is never the same party that signs off on its safety, restoring independent review of the riskiest decisions.
- ▶ Monitor in-service AOA-disagree and uncommanded-trim events as leading indicators, with a route that reaches engineering before a second airframe is lost.
- ▶ Re-validate the “no new training needed” claim against real fleet incidents, and reopen the training requirement the moment operational data contradicts the assumption.

## REFLECTION QUESTIONS

1. If a customer contract required removing a training requirement you judged necessary, what artifact would you produce, who would you route it to, and what would you do if the routing failed?
2. Reconstruct the MAX accident as a **capability-requirements** failure: what was the elided requirement, at what life-cycle stage was it elided, and who had the authority to insert it?
3. MCAS was permitted to act on a single sensor with no cross-check. Identify a system you work on that takes a safety-critical action from one unverified input, and specify the redundancy or human confirmation it lacks.

**WHO BUILDS THIS** systems & human-factors engineering · learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · scenario simulation · competency models · usability testing · cognitive walkthroughs

# Mars Climate Orbiter — Unit Mismatch

**D** Designed Out   **K** Knowledge & Institutional Memory

**IMPACT** Orbiter lost on arrival at Mars, one of the two Mars '98 spacecraft whose development cost \$193.1M; ground software produced thrust output in pound-force while navigation expected newtons

## IN BRIEF

The Mars Climate Orbiter, a NASA spacecraft built by Lockheed Martin and navigated by JPL, was lost on arrival at Mars in September 1999. Lockheed's ground software reported thruster impulse in pound-force seconds; JPL's navigation expected newton-seconds. The conversion — a factor of about 4.45 — was never applied at the boundary between the two teams, so every trajectory correction was mis-modeled and the error accumulated over the cruise. The orbiter arrived about 170 kilometres lower than planned and was either destroyed in the atmosphere or thrown back into heliocentric space; either way it was lost to a unit mismatch. The investigation did not stop at the coding error: the board's finding was that the project's processes never caught it. What failed was the interface between them: the specification existed and required newton-seconds, and nothing verified that the delivered software obeyed it. It is the canonical case of interface-as-requirement.

## BACKGROUND

The Mars Climate Orbiter, launched in December 1998 to study the Martian atmosphere, was one of NASA's "faster, better, cheaper" missions, built on a compressed budget that trimmed margins across the program. Lockheed Martin built the spacecraft and its ground software; JPL navigated it from Earth. Across the months of interplanetary cruise the two teams exchanged data so JPL could command the small trajectory corrections that keep a spacecraft on course — an exchange that crossed a software interface between the two organizations.<sup>0</sup>

## WHAT HAPPENED

Lockheed's ground software reported thruster impulse in pound-force seconds; JPL's navigation expected newton-seconds. The conversion — a factor of about 4.45 — was never applied at the boundary where the two systems met. Each firing was mis-modeled by that factor, and the error accumulated steadily over the long cruise until the predicted and actual trajectories had quietly diverged. When the orbiter reached Mars on 23 September 1999 it arrived about 170 kilometres too low — a closest approach of roughly 57 kilometres against the planned 226, and well below the 80 kilometres the spacecraft could have survived — into the atmosphere it was built to study from orbit, and was either destroyed there or thrown back into heliocentric space. The orbiter was lost to a unit conversion no one made.<sup>1</sup>

## THE INVESTIGATION


The Mishap Investigation Board put the proximate cause exactly there — the failed English-to-metric translation in ground software — but was careful to name the deeper one rather than stop at the bug. The ground software did not follow the interface specification, and the processes meant to catch that never did. What failed was the boundary between them: the specification fixed the units at newton-seconds and nothing checked the delivered software against it, and there was no complete end-to-end verification of the navigation software and its models. Navigators had even noticed odd trajectory behavior in cruise, but the concern was never run fully to ground before arrival, and the chance to catch it passed.<sup>2</sup>

## THE CAPABILITY GAP

The missing capability was ownership of the interface itself. Where a system is split across two organizations, the place they meet is not a documentation footnote but an engineering deliverable in its own right — with an owner, a specification, and a verification step. Here the specification existed and called for metric units; what was missing was any check that the delivered software complied with it, and any escalation path for the navigators who saw the trajectory drifting for months. Each team treated its own side as complete and the boundary as someone else's concern, so the one assumption that had to be shared and checked — what units are we speaking in? — was precisely the one no one verified. The spacecraft did not fail; it performed as built; the seam between the two halves of the organization did.<sup>3</sup>

## AFTERMATH & REFORM

NASA tightened interface management and end-to-end verification and treated the loss as a cautionary tale about how far “faster, better, cheaper” could be pushed before the corners being cut turned out to be decisive. The orbiter became the canonical systems-engineering example of interface-as-requirement — the civilian, software parallel to the Patriot's untraveled assumption (Case 129): two competent halves of a system, a boundary nobody owned or verified, and a small unspecified thing at that boundary that quietly destroyed the whole.<sup>4</sup>



MARS CLIMATE ORBITER — THE UNOWNED INTERFACE

*The “root cause” of the loss of the spacecraft was the failed translation of English units into metric units in a segment of ground-based, navigation-related mission software.*

ARTHUR STEPHENSON, CHAIRMAN, MARS CLIMATE ORBITER MISSION FAILURE INVESTIGATION BOARD, NASA RELEASE 99-134, 10 NOVEMBER 1999

## REFERENCES

1. NASA, Mars Climate Orbiter Mishap Investigation Board: Phase I Report (Nov. 1999) — mission, the Lockheed/JPL split, and the program context.
2. NASA MCO MIB Report (1999) — the pound-force-second vs newton-second mismatch (factor 4.45) in the “Small Forces” ground–software file, the accumulated navigation error, the 57 km closest approach against a planned 226 km periapsis and an 80 km survival limit, and the loss on 23 Sept. 1999 (root-cause statement quoted).
3. NASA MCO MIB Report (1999) — the missing verified interface specification, the absent end-to-end validation, and the unresolved cruise–trajectory anomalies.
4. N. G. Leveson, Engineering a Safer World (MIT Press, 2011) — interfaces as engineering deliverables requiring an owner and a verification step.
5. Sauser, B. J., Reilly, R. R., & Shenhar, A. J. (2009), “Why projects fail? How contingency theory can provide new insights — A comparative analysis of NASA’s Mars Climate Orbiter loss,” International Journal of Project Management 27(7), 665–679.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Mars Climate Orbiter is the textbook case for interface boundaries as engineering deliverables. Each half flew as built. The specification at the boundary was not enforced. The capability that was missing was the interface–specification verification step.

### LENS APPROACH

LENS uses Mars Climate Orbiter in LEN 5 to teach interface-as- requirement and in LEN 8 to discuss cross-contractor capability boundaries. The case is the foundational software-engineering parallel to Patriot (Case 129).

| DOMAIN            | FAILURE MODES       | LENS COMPETENCY                                     | PROBLEM TYPE | INDUCED | LEO |
|-------------------|---------------------|-----------------------------------------------------|--------------|---------|-----|
| SPACE / AEROSPACE | → <b>D</b> <b>K</b> | → <b>D1</b> <b>D2</b> <b>D3</b> <b>D4</b> <b>D5</b> | 1            | 1.3     | 1   |
| Systems Analysis  |                     |                                                     |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Make the interface between two organizations an owned deliverable with a written specification of units and formats, not a footnote each side assumes the other handles.
- ▶ Fix and verify the unit convention at the boundary explicitly, requiring both sides to confirm agreement before data crosses the seam.
- ▶ Run end-to-end validation across the combined system, since each half can be correct within its own assumptions while their composition is wrong.

#### IN OPERATION / AFTER

- ▶ Treat in-cruise anomalies, like the noticed trajectory behavior, as signals to run to ground before a point of no return rather than concerns to revisit after arrival.
- ▶ Monitor predicted-versus-actual trajectory continuously so an accumulating boundary error surfaces while there is still time to correct it.
- ▶ Audit cross-contractor seams for unowned assumptions whenever a program trims margins under “faster, better, cheaper,” confirming the corners cut are not decisive.

### REFLECTION QUESTIONS

1. Identify a contractor-to-contractor interface in your domain whose specification ownership is unclear. What would the equivalent unit-mismatch look like there?
2. Design the interface-verification deliverable that would have caught the pound-force / newton boundary before launch.
3. Navigators noticed odd trajectory behavior in cruise but the concern was never run to ground before arrival. Specify the threshold and process by which an in-flight anomaly in your domain must be resolved before a point of no return, rather than carried past it.

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines

# Glass-Cockpit Transition in Light Aircraft — Technology Outran Training

**H** Human-System Interface   **N** Normalization of Deviance   **K** Knowledge & Institutional Memory

**IMPACT** An NTSB safety study of ~8,000 piston aircraft manufactured 2002–2006, with accidents analyzed for 2002–2008, found that glass-cockpit aircraft had lower total accident rates but higher fatal accident rates than comparable conventional-instrument aircraft, and no significant overall improvement in safety; the Board did not resolve why the fatal rate diverged, but found no FAA equipment-specific training requirement for glass-cockpit pilots and issued recommendations A-10-36 through A-10-41

## IN BRIEF

Glass cockpits — integrated digital primary flight displays and multifunction displays replacing the inherited six-pack of analog instruments — were introduced into light piston aircraft over the 2000s as a fleet-wide modernization. The NTSB safety study NTSB/SS-10/01 examined approximately 8,000 small piston aircraft manufactured 2002–2006 and reached a finding the technology's own advocates did not expect: glass-cockpit aircraft had lower total accident rates than comparable conventional-instrument aircraft over the period studied but higher fatal accident rates, and no significant overall improvement in safety. The NTSB did not resolve why the fatal rate diverged — it recorded that additional research is warranted — and identified the absence of equipment-specific training requirements as the safety issue in reach. The Board issued recommendations A-10-36 through A-10-41 on knowledge-testing standards, simulator availability, and training requirements. The study is explicit that advanced avionics “can increase the safety potential” of light aircraft but that, “for the aircraft and time period studied,” the transition “has not yet resulted in the anticipated improvement in safety” — an open, not closed, verdict. The case is the canonical 7.1 failure of an inherited capability regime (pilot proficiency) not being re-verified against the new envelope the technology transition introduced. Pair with the aging-system transition cases (Cases 114–116, 154, 156, 174, and 194, drafted in parallel).

## BACKGROUND

Glass-cockpit avionics — integrated digital primary flight displays and multifunction displays replacing the inherited six-pack of analog round-dial instruments — were introduced into light piston aircraft over the 2000s. The fleet conversion was largely industry-driven; new-build aircraft came with glass as standard, retrofits became increasingly common, and the FAA's certification framework treated the transition as a positive safety move on the strength of the capability the displays delivered to pilots: integrated situational awareness, moving-map navigation, terrain and traffic display, system-status integration. The opera-



tionally critical assumption was that pilots transitioning from analog to glass would carry their proficiency across the change.<sup>0</sup>

#### WHAT HAPPENED

The NTSB safety study NTSB/SS-10/O1 examined that assumption directly. The Board studied approximately 8,000 small piston airplanes manufactured in the United States between 2002 and 2006 — 2,848 in the conventional-display cohort and 5,116 in the glass-cockpit cohort — and compared their accident records over 2002 through 2008. The headline finding ran against expectation. Glass-cockpit aircraft had lower total accident rates than comparable conventional-instrument aircraft over the period, but a higher fatal accident rate, and no significant improvement in safety overall. The expected fleet-wide safety gain from the displays had not yet materialized in the accident record.<sup>1</sup>

#### THE INVESTIGATION

The NTSB's attribution is careful. It did not name a cause for the fatal-rate gap, and said further research is warranted. What it did establish was that the two cohorts were not flying the same missions: the glass-cockpit accident flights differed significantly in purpose, planned length, and instrument-meteorological exposure, and their pilots differed in age, certification, and total experience — leaving open whether owners bought glass for the flying they already intended or the new capability encouraged longer flights in worse weather. The transitioning pilot population was certificated and current under the inherited training regime that assumed analog instruments, and the glass displays — for all their situational-awareness advantages — introduced new failure modes (mode confusion, automation surprises, attention capture by integrated displays, degraded scan in degraded modes) that the inherited training did not address. The Board issued recommendations A-10-36 through A-10-41 to the FAA: equipment-specific knowledge-testing standards, simulator and training-device availability for transitioning pilots, and structured training requirements before operating glass-cockpit aircraft.<sup>2</sup>

#### THE CAPABILITY GAP

The study's language is open rather than closed. Advanced avionics "can increase the safety potential" of light aircraft, but "for the aircraft and time period studied" the transition "has not yet resulted in the anticipated improvement in safety." The verdict is not that glass is unsafe; it is that the inherited certification of pilot proficiency, built around analog instruments, did not transfer to the new envelope without re-engineered training. The case is the canonical 7.1 failure: a capability regime (pilot proficiency) was not re-verified when the cockpit's interface envelope changed, and the inherited training assumptions silently became inadequate to the new operational reality.<sup>3</sup>

#### AFTERMATH & REFORM

Drafted alongside the aging-system transitions in Cases 114–116, 154, 156, 174, and 194, the glass-cockpit GA study carries the capability-under-system-change pattern at the consumer-aviation scale where the analog cases run at the defense and commercial-aviation scale. The structural lesson is the same: when a system's interface envelope changes — and even when the change is a capability-enhancing one — the inherited certification of operator proficiency has to be re-verified against the new envelope, or the transition outruns the training and the accident record moves in the wrong direction. The NTSB recommendations

are the engineering response to that pattern, and their implementation is what closes (or fails to close) the open verdict the study left.

*Advanced avionics and electronic displays can increase the safety potential of general aviation aircraft operations by providing pilots with more operational and safety-related information and functionality, but more effort is needed to ensure that pilots are prepared to realize that potential.*

NTSB SAFETY STUDY NTSB/SS-10/01 (2010).

REFERENCES

1. Wiener, E. L., & Curry, R. E. (1980). Flight-deck automation: Promises and problems. *Ergonomics*, 23(10):995–1011 — the foundational literature on automation-induced failure modes that the glass-cockpit transition re-introduced at the GA scale.

2. National Transportation Safety Board (2010). Introduction of Glass Cockpit Avionics into Light Aircraft, Safety Study NTSB/SS-10/01. <https://www.nts.gov/safety/safety-studies/Documents/SS1001.pdf> — the case’s primary investigation document.

3. NTSB Safety Recommendations A-10-36 through A-10-41 (2010), issued to the FAA — the engineering response the open verdict points to.

4. Sarter, N. B., Woods, D. D., & Billings, C. E. (1997). Automation surprises. In *Handbook of Human Factors and Ergonomics* (2nd ed.) — the mode-confusion / automation-surprise literature the NTSB findings cross-reference.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The NTSB’s glass-cockpit GA study is the canonical capability-under-system-change failure at the consumer-aviation scale: a positive technology transition that nevertheless outran the inherited certification of operator proficiency. The Board named no cause for the fatal-rate gap and asked for more research, and the verdict is open — the safety potential is there, and the transition has not yet realized it.

LENS APPROACH

Glass-cockpit GA is the small-fleet capability-under- transition failure (induced 7.1; LENS D1+D3/PT1). LENS uses it in Domain 1 (Systems Analysis) for the inherited certification of operator proficiency that has to be re-verified against the new envelope, and in Domain 3 (Human-System Collaboration) for the new automation-induced failure modes (mode confusion, automation surprises, attention capture) the transition introduced. Pair with the aging-system transition cases (Cases 114–116, 154, 156, 174, 194) at the cross-scale capability-under-change layer.

| DOMAIN           | FAILURE MODES | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO |
|------------------|---------------|-----------------------|--------------|---------|-----|
| GENERAL AVIATION | H N K         | D1 D2 D3 D4 D5        | 1            | 7.1     | 1-2 |
| AVIONICS         |               | Systems Analysis + D3 |              |         |     |
| PILOT TRAINING   |               |                       |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the inherited certification of operator proficiency in advance of a technology transition, and design the re-verification against the new envelope as part of the transition’s deliverable, not as a downstream training catch-up.
- Treat capability-enhancing transitions with the same re-verification discipline as capability-degrading ones; the glass-cockpit transition was a positive technology

IN OPERATION / AFTER

- Preserve the NTSB’s open-verdict language (“can increase the safety potential ... has not yet resulted in the anticipated improvement in safety”) in any report on the transition’s outcome; the study did not say glass is unsafe, and the precise language is what makes the recommendation set actionable.
- Track implementation of the NTSB recommendation set (A-10-36 through A-10-41) as the engineering

- move that nevertheless required training re-verification the field skipped.
- ▶ Name the new failure modes the transition introduces (mode confusion, automation surprises, attention capture) at the design stage of the equipment-specific training, rather than waiting for the accident record to surface them.
- ▶ response to the open verdict; the verdict closes when the recommendations are implemented and the next round of evidence is collected.
- ▶ Carry the structural lesson into adjacent transitions — the aging-system transition cases (Cases 114–116, 154, 156, 174, 194) — as evidence that the inherited certification of operator proficiency has to be re-verified against the new envelope across consumer, commercial, and defense scales.

### REFLECTION QUESTIONS

1. Identify a technology transition in your domain that is positive on its capability claim and that nevertheless puts the inherited certification of operator proficiency under question. What re-verification of operator proficiency would the transition require — and is it currently part of the deliverable, or assumed away?
2. The NTSB declined to name a cause for the fatal-rate gap and pointed instead at the absent equipment-specific training requirement. What is the analog distinction in your context — between the engineered artifact and the operator-proficiency regime that has to track it?

**WHO BUILDS THIS** human factors & interaction design · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · incident analysis · safety dashboards · knowledge bases · data pipelines

# Ariane 5 Flight 501

**D** Designed Out   **K** Knowledge & Institutional Memory   **H** Human-System Interface

**IMPACT** Maiden flight destroyed itself 39 seconds after main-engine ignition — about 32 seconds after lift-off; four Cluster science satellites lost, replaced at 214 MECU; reused Ariane 4 code never re-verified for Ariane 5

## IN BRIEF

On its 1996 maiden flight, the Ariane 5 rocket destroyed itself 39 seconds after main-engine ignition — about 32 seconds after lift-off — losing the four Cluster science satellites aboard. The cause was reused software: the inertial reference system inherited code from Ariane 4, where a horizontal-velocity value never exceeded a 16-bit integer's range. Ariane 5's steeper, faster trajectory pushed it past that range; the integer overflow crashed both redundant reference systems almost simultaneously, the vehicle lost guidance and broke up. The inquiry found the offending code path had been irrelevant on Ariane 4 and was never removed or re-verified for Ariane 5 — software reuse had been treated as risk-free. Ariane 5 is the foundational safety-critical- software case for the hazard of reusing code without re-verifying it against the new system's operating envelope.

## BACKGROUND

The Ariane 5 heavy-lift rocket reused proven components from its successful predecessor, Ariane 4 — including, in its inertial reference system, software that had flown reliably for years. Reuse of flight-proven code was regarded as a way to reduce risk and cost.<sup>0</sup> The reasoning was seductive precisely because the code's flight record was genuine: software that had worked on Ariane 4 looked like the safest possible choice, and that confidence is what allowed it to cross onto a different vehicle without being re-examined against the new vehicle's conditions.

## WHAT HAPPENED

On 4 June 1996, Ariane 5 Flight 501 — its maiden flight — veered off course and began to disintegrate under aerodynamic loads about 39 seconds after the main-engine ignition command, roughly 32 seconds into the flight; the on-board destruct system fired as designed, and the four Cluster science satellites were lost.<sup>1</sup> The vehicle was destroyed in well under a minute of flight, before it had done anything but climb — the whole loss flowing from a fault that triggered about thirty seconds after liftoff, when the new trajectory first pushed the inherited code outside the range it had been built to handle.

## THE INVESTIGATION

The inquiry board traced the cause to the reused code. A horizontal-velocity value that never exceeded a 16-bit signed integer on Ariane 4's trajectory did exceed it on Ariane

5's faster one; the resulting overflow shut down both redundant inertial reference systems almost simultaneously, and the vehicle lost guidance.<sup>2</sup> Redundancy gave no protection here because both units ran the identical inherited code and met the identical out-of-range value at the same instant, so the backup failed a data cycle ahead of the primary, leaving nothing to switch to — duplicating a system defends against a part breaking, not against a shared assumption being wrong. The offending code path had become irrelevant to Ariane 5's flight after liftoff, yet had neither been removed nor re-verified against the new vehicle's envelope — reuse had been treated as risk-free.<sup>3</sup>

#### THE CAPABILITY GAP

Ariane 5 is the canonical software-engineering case for the fallacy of risk-free reuse. Code is fit for the envelope it was written and verified against; reusing it in a different envelope is a new design decision and must be re-verified as one. The institutional capability that was missing was the requirement that reused safety-critical code be re-verified, not merely trusted because it had worked before.<sup>4</sup> The flaw was never in the Ariane 4 code, which was correct for the trajectory it was written against; the flaw was the institutional assumption that a record of working in one envelope carried over to another, an assumption no test was required to challenge.

#### AFTERMATH & REFORM

The fix and the lesson reshaped safety-critical software practice: reuse became a verification event, with explicit checking of every inherited assumption against the new system's operating conditions.<sup>5</sup> The case rhymes with the Patriot (Case 129) and the Mars Climate Orbiter (Case 98) — small, unexamined assumptions inherited across a boundary, fatal when the boundary's conditions changed. What unites the three is that each inherited a quantity sound in its original setting and lethal in the new one, so the reform is the same in every case: make the boundary crossing the moment the assumption is re-checked, rather than the moment it is silently trusted.

# 39s

AFTER IGNITION · 16-BIT INTEGER OVERFLOW

*code path disabled by the previous vehicle's profile; re-enabled by the new one*

ARIANE 5 — THE FALLACY OF RISK-FREE CODE REUSE

*The internal SRI software exception was caused during execution of a data conversion from 64-bit floating point to 16-bit signed integer value.*

ARIANE 5 FLIGHT 501 FAILURE INQUIRY BOARD, 1996

#### REFERENCES

1. J. L. Lions (chair), Ariane 5 Flight 501 Failure Inquiry Board Report (1996) — the data-conversion overflow (quoted).
2. Lions Report (1996) — SRI failure at HO+36.7 s (about 30 s after lift-off), disintegration under aerodynamic loads at about HO+39 s, and the automatic destruct.
3. Lions Report (1996) — the 64-bit-to-16-bit conversion overflow and the simultaneous shutdown of both inertial reference systems.
4. Lions Report (1996) — the reused code path neither removed nor re-verified for Ariane 5's envelope.
5. Cf. Patriot (Case 129) and Mars Climate Orbiter (Case 98).
6. N. Leveson, Safeware (1995) — software-reuse hazards; G. Le Lann, "An Analysis of the Ariane 5 Flight 501 Failure" (1997).
7. ESA Science & Technology, "Cluster — Summary" — the loss of the four Cluster spacecraft and the 214 MECU recovery mission approved by the SPC.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Ariane 5 is the canonical software-engineering case for the fallacy of risk-free code reuse. Code is fit for its envelope of operation. Reusing it in a different envelope is a new design decision and must be re-verified as one. The institutional capability that was missing was the requirement to re-verify.

### LENS APPROACH

LENS uses Ariane 5 in LEN 5 for software-engineering capability deliverables and in LEN 8 for the institutional discipline that treats reuse as a verification event rather than as a savings.

| DOMAIN            | FAILURE MODES                | LENS COMPETENCY                                     | PROBLEM TYPE | INDUCED | LEO |
|-------------------|------------------------------|-----------------------------------------------------|--------------|---------|-----|
| SPACE / AEROSPACE | → <b>D</b> <b>K</b> <b>H</b> | → <b>D1</b> <b>D2</b> <b>D3</b> <b>D4</b> <b>D5</b> | 1            | 7.2     | 1   |
| Systems Analysis  |                              |                                                     |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Treat every reuse of safety-critical code as a new design decision: re-verify each inherited assumption against the new system's operating envelope before flight.
- ▶ Document the original operating envelope a component was verified against, so a later reuse can be checked against it rather than blindly trusted.
- ▶ Design redundancy to fail independently — diverse implementations or inputs — so a shared inherited assumption cannot take out primary and backup at once.

#### IN OPERATION / AFTER

- ▶ Audit reused components for code paths that are irrelevant in the new system but still active, and remove or re-verify them rather than leaving them dormant.
- ▶ Make boundary crossings — code moving onto a new vehicle, a new envelope — a mandatory re-verification event in the institution's process, not a savings.
- ▶ Review past reuse decisions for the Ariane / Patriot / Mars Climate Orbiter pattern: a quantity sound in its original setting and unchecked in the new one.

### REFLECTION QUESTIONS

1. Identify a piece of reused infrastructure in your domain whose original operating envelope is not documented. What is the silent assumption?
2. Design the verification deliverable that should accompany every reuse of safety-critical software.
3. Both redundant reference systems failed at the same instant because they shared the same inherited assumption. Identify a place in your domain where redundancy gives false comfort because the duplicated units share a common flaw rather than fail independently.

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · usability testing · cognitive walkthroughs

# Air France Flight 447

**T** Training Gap    **H** Human-System Interface

**IMPACT** 228 killed; the deadliest accident in Air France's history

## IN BRIEF

Air France Flight 447, an Airbus A330, fell into the South Atlantic on 1 June 2009, killing all 228 aboard — the deadliest accident in the airline's history. At cruise altitude the pitot probes iced over, the airspeed readings failed, and the autopilot handed the jet to a crew that had never trained to fly it by hand in that regime. The pilot flying held the nose up into a full aerodynamic stall and never recognized it; the upset ran four minutes and twenty-three seconds from autopilot disconnection to impact. The BEA traced the loss to the airspeed failure, the crew's inappropriate inputs, and a training system that taught stall prevention at low altitude but never stall recovery at altitude — a gap between the trained envelope and the operational one that reshaped global pilot-training rules.

## BACKGROUND

AF447 left Rio for Paris on 31 May 2009 with 228 aboard an Airbus A330 — a highly automated jet with an excellent safety record. Its route crossed the equatorial storm band, and it carried a known vulnerability: Thales AA pitot probes prone to brief icing at high altitude in heavy precipitation, a pattern documented across the A330/A340 fleet in the years before the accident. A replacement program with newer probes was underway, but the accident aircraft had not yet been modified, so the vulnerability the program existed to close was still live on the very jet that crossed the storm band — the retrofit recognized as necessary but not yet fitted where it mattered.<sup>9</sup> The A330's fly-by-wire architecture also carried a design choice that would matter later: the two sidesticks were not mechanically linked, so an input on one was not felt on the other, a philosophy that traded the visual cue of yoke movement for sidestick lightness and independence.

## WHAT HAPPENED

At 35,000 feet the probes iced, the airspeed readings went invalid, and the autopilot and autothrust disconnected into Alternate Law — a degraded control regime in which stall protection no longer held. The pilot flying responded with sustained nose-up input; the jet climbed, stalled, and never recovered, falling some 38,000 feet into the ocean in about three and a half minutes. The stall warning sounded, then cut out at extreme angle of attack and resumed when the nose dropped — warning against the one input that would have begun a recovery, so that the cue meant to guide the crew instead punished the correct action and rewarded the fatal one, an inversion no amount of hand-flying instinct could resolve in the time available.<sup>1</sup> The other pilot's correcting inputs on his own sidestick went unfelt on

the flying pilot's: the design's clean independence became, under stress, an inability to see what the other crew member was doing.

#### THE INVESTIGATION

The BEA's three-year investigation, concluded in the 2012 final report, named a chain: air-speed loss from pitot icing, inappropriate crew inputs, and the crew's failure to recognize and recover from the stall.<sup>2</sup> The pilots were not incompetent — they were outside anything their training had prepared them for, and the BEA said so: "the conditions in which airline pilots are trained and exposed to stalls... does not generate the expected behaviour in any acceptable reliable way."<sup>3</sup> The finding reframed the loss from a question of individual airmanship to one of training design: the crew had been drilled in a regime that never produced the responses the emergency demanded. The BEA report also surfaced the longer industry debate over fly-by-wire philosophies — the Airbus convention of independent sidesticks and protected envelopes versus the Boeing convention of linked yokes and trim feedback — not to adjudicate which was safer in steady state but to argue that whichever a manufacturer chose, the training had to make the degraded-mode behavior of that choice reflexive rather than novel.

#### THE CAPABILITY GAP

The gap was precise. Airlines trained stall recovery at low altitude — the regime certification required and the one simulators could reproduce — while the simulators of the era could not faithfully reproduce a high-altitude stall, so crews never practiced the situation that arrived. Years of reliable automation had also let hand-flying skills atrophy, so when the autopilot handed back control the crew met an unfamiliar regime with rusty manual skills at the worst possible moment. The trained envelope and the operational envelope had quietly diverged, and AF447 fell into the gap between them — a gap no one had been positioned to see widen, because the training record showed full compliance and the operational record showed an excellent safety performance, right up until the regime that neither had covered arrived together with degraded control law.<sup>4</sup>

#### AFTERMATH & REFORM

The BEA's recommendations reached far beyond one airline: the pitot probes were replaced fleet-wide, and the report pressed for training in manual high-altitude flight, stall recovery from cruise altitude, unreliable-airspeed handling, and angle-of-attack indication.<sup>5</sup> Regulators then made Upset Prevention and Recovery Training mandatory for airline pilots — FAA Part 121 adopted UPRT-aligned stall-recovery training in 2014; EASA phased UPRT into ATPL training by 2019; ICAO codified it in Doc 10011 — closing at the regulatory level the gap that had been invisible at the airline level, and moving the fix from a single carrier's discretion to a binding standard every airline had to meet.<sup>6</sup> Simulator manufacturers were pushed to extend their aerodynamic models into the post-stall regime so the training could actually take place. The crew performed exactly as trained; the training was the wrong training, and only a system-wide mandate could keep that mismatch from recurring elsewhere.





system, and with the Boeing 737 MAX MCAS sequence as the near-current echo. LEO mapping: LEO-2 (Iterative Development) primary; LEO-4 (Test and Evaluation) for the BEA investigation framing; LEO-3 (Human-System Collaboration) for the automation-handoff dimension.

| DOMAIN                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO   |
|-----------------------|---------------|------------------|--------------|---------|-------|
| AVIATION              | → T H         | → D1 D2 D3 D4 D5 | 4            | 1.2     | 2 · 4 |
| Iterative Development |               |                  |              |         |       |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>▸ Define the operational capability envelope explicitly — including degraded-automation and high-altitude regimes — and train to its edges, not only to the routine center.</li><li>▸ Engineer warning logic so that cues never punish the correct recovery input; validate stall-warning behavior across the full angle-of-attack range before fielding.</li><li>▸ Treat manual-flying proficiency in degraded modes as a measured deliverable, not an assumed residual of automated operation.</li></ul> | <ul style="list-style-type: none"><li>▸ Audit recurrent training against the actual operational record so the trained envelope cannot silently diverge from the regimes crews encounter.</li><li>▸ Monitor for skill atrophy where reliable automation reduces hands-on exposure, and refresh manual competence before it erodes.</li><li>▸ Sustain the reform at the regulatory level (mandatory upset recovery) so a fix proven in one carrier propagates to all rather than lapsing.</li></ul> |

REFLECTION QUESTIONS

1. The simulators of 2009 could not produce high-altitude stall behavior. What is the equivalent gap in your domain — the operational regime your training environment cannot reproduce?
2. Design the recurrent-training curriculum that would have caught the AF447 gap. Be specific about cost, evidence, and what makes the curriculum falsifiable against the operational record.
3. The autopilot handed control to a crew at the one moment it was least prepared to take it. Identify an automation-to-human handoff in your domain that occurs precisely when the human is least ready, and design the trigger or warning that would change that.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · usability testing · cognitive walkthroughs

# Kegworth / British Midland 92

**T** Training Gap   **H** Human-System Interface   **K** Knowledge & Institutional Memory

**IMPACT** 47 killed and 74 seriously injured when the crew shut down the wrong engine

## IN BRIEF

On 8 January 1989 British Midland Flight 92, a nearly new Boeing 737-400, crashed on the M1 motorway embankment near Kegworth, killing 47 and seriously injuring 74. After a fan blade fractured in the left engine, the crew shut down the right engine — the one still working. They acted before reading the engine displays: the first officer named the right engine and could not afterwards say what he had seen, and the commander accepted it. The new, harder-to-read electronic displays did not correct them, and the cabin crew who saw flames never told the flight deck. The commander later explained his suspicion of the right engine by an air-conditioning layout that held on the earlier 737s he had flown but not on the -400. The AAIB issued 31 recommendations; Kegworth became the textbook case of capability degrading under system change.

## BACKGROUND

On 8 January 1989 British Midland Flight 92, a nearly new Boeing 737-400 (G-OBME), left Heathrow for Belfast with 126 aboard. The -400 was a recent variant — bigger engines, a partly redesigned cockpit with new electronic engine instruments. The crew were experienced 737 pilots, but their conversion onto the variant was a one-day ground course on the differences, taken with no 737-400 simulator in existence anywhere in the United Kingdom — so the display they would have to read under pressure had never been flown, only shown.<sup>0</sup>

## WHAT HAPPENED

Climbing through 28,300 feet, a fan blade fractured in the left engine, filling the cabin with vibration and smoke. Without waiting to read the engine displays, the first officer named the right engine and the commander throttled it back some nineteen seconds after the fan blade fractured; the order to shut it down came about forty-three seconds after the vibration began. The symptoms eased (because disconnecting the autothrottle let the left engine stop surging), seeming to confirm it, and they shut the right engine down. On final approach the damaged left engine failed completely; with the good engine off, the aircraft struck the M1 embankment short of the runway — 47 killed, 74 seriously injured. The brief calm after the throttle-back was the trap: it appeared to confirm a diagnosis that was wrong, removing the doubt that might have prompted a reassessment.<sup>1</sup>

## THE INVESTIGATION

The Air Accidents Investigation Branch traced the disaster to the shutdown of the serviceable engine, and among its contributory factors to the crew reacting prematurely and contrary to their training, and not assimilating the engine instrument display before throttling

back the No 2 engine. The commander's later account — that his reading of the air-conditioning layout pointed him at the right engine — the Board recorded as reasoning that would have held on the earlier 737 variants he had flown but not on the -400, whose flight deck is supplied from the left engine.<sup>2</sup> The new electronic displays did not help: the vibration indicator that would have pointed at the failing engine was harder to read at a glance than the dials it replaced, so the one instrument that could have corrected the diagnosis was the least likely to be consulted under pressure.<sup>3</sup> And though cabin crew and passengers could see flames from the left engine, that never reached the flight deck — a crew-resource-management gap aviation would later train hard against, because those with the decisive evidence had no path to those making the decision.<sup>4</sup>

#### THE CAPABILITY GAP

The crew were skilled, and their mental model was correct — for the airplane they had flown the week before. The -400 was treated, for training, as an incremental change to a familiar type; under emergency pressure it behaved as a categorical one. Reading a manual note about revised bleed air is not the same as drilling a new reflex that fires before the old one, and the cockpit redesign had quietly removed perceptual cues a startled crew relies on. The change was filed as incremental precisely because nothing tested it until an emergency did. Capability had degraded under system change with no one noticing, because nothing failed until the day it did.<sup>5</sup>

#### AFTERMATH & REFORM

The AAIB issued thirty-one safety recommendations spanning conversion and difference training, engine-instrument and vibration-display design, cabin-to-flight-deck communication, and crashworthiness.<sup>6</sup> Kegworth became a standard human-factors teaching case and a reference for how a transition program should be built — so the differences between an old system and a new one are made hard to overlook, and the people closest to the evidence have a path to those making the decision. Each recommendation traced back to a failure the accident exposed — the unforced mental model, the unreadable instrument, the silent cabin — treated together rather than as isolated faults.

737-400 (1989)



▲ 47 dead — the instruments were never read

CREW MENTAL MODEL FROM OLDER VARIANT; CONFIGURATION CHANGED

*They did not assimilate the indications on the engine instrument display before they throttled back the No. 2 engine.*

AAIB AIRCRAFT ACCIDENT REPORT 4/1990, CAUSE AND CONTRIBUTORY FACTORS, 1990

REFERENCES

1. Air Accidents Investigation Branch, Report on the accident to Boeing 737-400 G-OBME near Kegworth, Leicestershire, on 8 January 1989, AAIB Aircraft Accident Report 4/90 (1990) — aircraft, route, and the limited conversion training; see also U.S. FAA, Lessons Learned: G-OBME.

2. AAIB Report 4/90 (1990) — accident sequence: left-engine fan-blade failure, shutdown of the serviceable right engine, and total power loss on approach; 47 killed, 74 seriously injured.

3. AAIB Report 4/90 (1990) — the older-737 bleed-air mental model misapplied to the 737-400's changed configuration.

4. AAIB Report 4/90 (1990) — the new electronic engine instrumentation and the readability of the vibration indicator.

5. AAIB Report 4/90 (1990) — the cabin-to-flight-deck communication gap (crew resource management).

6. AAIB Report 4/90 (1990) — the 31 safety recommendations on difference training, instrument design, CRM, and crashworthiness.

7. J. Reason, Human Error (Cambridge Univ. Press, 1990) — Kegworth as a case in the misapplication of a correct mental model to a changed system.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Kegworth is the textbook example of Capability Degradation Under System Change. The pilots were competent and current. What the AAIB found was that they acted before reading the instruments that would have told them which engine was failing, on a variant whose engine display they had met in one day of ground school and never in a simulator. The new airframe was treated as an incremental change. The crew's response revealed it was a categorical one.

LENS APPROACH

LENS treats Kegworth in LEN 5 as the canonical **system-change** capability problem and in LEN 8 as the institutional version: how knowledge about what has changed travels from engineering to operations and what gets lost in transit. The case sits alongside Patriot (Case 129) in the canonical problem-type pair for system transition.

| DOMAIN                     | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|---------------|------------------|--------------|---------|-----|
| AVIATION                   | → T H K       | → D1 D2 D3 D4 D5 | 6            | 3.3     | 3   |
| Human-System Collaboration |               |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Classify any change that alters operator mental models — even a “minor” variant — as categorical, and gate it on drilled retraining rather than reading.
- ▶ Test new interface designs against startled-crew performance before fielding, so a redesign cannot quietly remove the perceptual cues operators rely on.
- ▶ Engineer a communication path from those closest to the evidence (cabin crew, passengers, sensors) to those making the decision, so decisive observations reach the flight deck.

IN OPERATION / AFTER

- ▶ Audit transition programs for capability that degraded silently — where nothing failed until an emergency tested the change — using simulator scenarios that force the new reflex.
- ▶ Monitor whether difference training actually overwrites the prior model, treating un-disturbed mental models as a measurable latent risk.
- ▶ Track that interface and communication recommendations are implemented together, not as isolated fixes, so the integrated failure mode does not recur.

## REFLECTION QUESTIONS

1. The Kegworth crew's mental model was right for the previous variant. What change in your domain currently risks an analogous misapplication?
2. Difference training is a generic deliverable in transition programs. Design the artifact that would make differences hard to overlook rather than easy.
3. The cabin crew and passengers saw flames the flight deck never learned of. Identify a path in your domain by which the people closest to the evidence cannot reach the people making the decision, and design the channel that would carry it under emergency conditions.

**WHO BUILDS THIS** learning science & instructional design · human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs · knowledge bases · data pipelines

# Asiana Airlines Flight 214

**T** Training Gap   **H** Human-System Interface

**IMPACT** 3 killed, 187 injured; Boeing 777 struck the seawall short of SFO runway 28L

## IN BRIEF

On a clear afternoon at San Francisco in July 2013, the crew of Asiana 214 allowed their Boeing 777 to slow far below approach speed without recognizing that the autothrottle was not maintaining it. The aircraft struck the seawall short of runway 28L, killing three of the 307 aboard and injuring nearly two hundred. The NTSB found the crew had mismanaged the approach and inadequately monitored airspeed — but also that the complexity of Boeing’s autoflight systems contributed: after the pilots disconnected the autopilot and pulled thrust to idle, the autothrottle reverted to a HOLD mode in which it would not wake to hold the selected speed. The captain believed it would. Asiana 214 is the most prominent recent case of automation surprise on a Western wide-body airliner.

## BACKGROUND

Asiana 214 was a scheduled Boeing 777 flight from Seoul to San Francisco. The captain was experienced on other types but new to the 777 and flying it under the supervision of a training captain; the approach was a visual one, with the airport’s instrument glideslope out of service for construction — so a pilot still learning the type’s automation was hand-flying an approach stripped of the vertical guidance that would normally have anchored it, exactly the conditions under which an unfamiliar mode is most likely to bite.<sup>o</sup>

## WHAT HAPPENED

On 6 July 2013, on final approach to runway 28L, the aircraft descended below the proper path and slowed. The crew had set an autopilot mode that left the autothrottle in HOLD, where it would not advance thrust to hold the selected speed. Airspeed decayed to about 103 knots against a target of 137; the 777 struck the seawall short of the runway, killing three and injuring nearly two hundred of the 307 aboard. The thirty-four-knot shortfall built up silently while the crew believed the automation was holding the speed, so the first unambiguous signal of trouble was the airframe itself running out of energy on short final.<sup>1</sup>

## THE INVESTIGATION

The NTSB found the probable cause to be the crew’s mismanagement of the descent, the pilot flying’s unintended deactivation of automatic airspeed control, inadequate monitoring of airspeed, and a delayed go-around — the deactivation traced, in the Board’s analysis, to the pilot flying’s faulty mental model of the automation, and the automation’s complexity inadequately described in Boeing’s documentation and Asiana’s training. The captain told investigators he had assumed the autothrottle would maintain speed; in that configuration, it does not — a gap between the system’s documented behavior and the mental model the

training had left him with, the kind of mismatch that surfaces only when the automation is asked to do what the operator wrongly believes it is already doing.<sup>2</sup>

#### THE CAPABILITY GAP

Asiana 214 is the textbook automation surprise on a Western wide-body: the mode the autothrottle was actually in and the mode the crew believed it was in diverged silently. The missing capability was not raw pilot skill but transparency — a system that made its own state, and the fact that the crew now owned the airspeed, salient at the moment it mattered, rather than leaving the handoff of responsibility to be inferred from a silent mode change the interface did nothing to announce.<sup>3</sup>

#### AFTERMATH & REFORM

The NTSB issued recommendations on autoflight training, low-energy alerting, and the intuitiveness of flight-crew/autoflight interfaces, and emphasized training that preserves manual-flying proficiency to counter automation dependence. The accident became a standard reference in airline automation-policy reviews and in the push toward “automation airmanship” curricula — shifting the lesson from one crew’s error to a recurring design problem about how modes are annunciated and how monitoring is taught when the automation is doing most of the flying.<sup>4</sup>



*The complexities of the autothrottle and autopilot flight director systems that were inadequately described in Boeing’s documentation and Asiana’s pilot training, which increased the likelihood of mode error.*

NTSB AIRCRAFT ACCIDENT REPORT NTSB/AAR-14/01, PROBABLE CAUSE STATEMENT, CONTRIBUTING FACTORS, 2014

#### REFERENCES

1. NTSB, Aircraft Accident Report: Asiana Airlines Flight 214, NTSB/AAR-14/01 (2014) — probable cause and contributing factors (quoted).
2. NTSB/AAR-14/01 (2014) — the autothrottle HOLD reversion and airspeed decay to 103 kt.
3. NTSB/AAR-14/01 (2014) — the captain’s mental model of the autothrottle.
4. Sarter, N. & Woods, D. (1995), “How in the world did we ever get into that mode?,” Human Factors — automation surprise.
5. NTSB/AAR-14/01 (2014), safety recommendations A-14-37 to A-14-43 and A-14-55 on autoflight training, low-energy alerting, interface intuitiveness, and manual flight.

#### THE LEARNING ENGINEERING LENS



LE INSIGHT

Asiana 214 is the aviation case for the LENS Human-System Collaboration proposition that automation transparency is a capability deliverable. The mode the autothrottle was actually in and the mode the crew believed it was in diverged silently — a textbook automation surprise. The training, the documentation, and the interface together did not surface the divergence.

LENS APPROACH

LENS treats Asiana 214 in LEN 2 as a worked case for automation transparency requirements. Students design the interface and training intervention that would have made the mode disconnect salient to the crew at the moment the aircraft slowed.

| DOMAIN                     | FAILURE MODES                   | LENS COMPETENCY                                                                   | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|---------------------------------|-----------------------------------------------------------------------------------|--------------|---------|-----|
| AVIATION                   | → <span>T</span> <span>H</span> | → <span>D1</span> <span>D2</span> <span>D3</span> <span>D4</span> <span>D5</span> | 6            | 3.3     | 3   |
| Human-System Collaboration |                                 |                                                                                   |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Make automation transparency a design requirement: the system must announce its active mode and any silent reversion, not leave it to be inferred.
- ▶ Annunciate transfers of responsibility (e.g., when the crew now owns airspeed) explicitly at the moment they occur, so no handoff is implicit.
- ▶ Design and validate mode logic against operator mental models during development, so configurations that “do not do what pilots assume” are caught before fielding.

IN OPERATION / AFTER

- ▶ Train and audit energy-state monitoring so a slow, silent airspeed decay is caught by the crew before the airframe signals it.
- ▶ Preserve manual-flying proficiency through recurrent practice to counter the automation dependence that hid the decay.
- ▶ Monitor operations for mode-confusion events and feed them into automation-policy and annunciation-design reviews.

REFLECTION QUESTIONS

1. Identify a mode in an automated system you work with where the operator’s mental model and the system’s actual behavior can silently diverge. How would the operator know?
2. Design the interface change that would have alerted the Asiana crew to the mode mismatch before the airspeed decayed.
3. The autothrottle’s HOLD reversion handed airspeed back to a crew that did not know it now owned it. Design the annunciation that would make a transfer of responsibility between automation and operator unmissable at the instant it occurs.

**WHO BUILDS THIS** learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs

# Helios Airways Flight 522

**T** Training Gap   **H** Human-System Interface

**IMPACT** 121 killed; cabin failed to pressurize after a maintenance setting was left on "manual" and the crew misread the warning

## IN BRIEF

Shortly after takeoff from Larnaca in August 2005, the crew of Helios 522 misread a warning horn. The same intermittent horn that signals an incorrect takeoff configuration on the ground sounds, in flight, when cabin altitude climbs past about 10,000 feet — and the pressurization selector had been left in "manual" after a leak check the night before. The cabin never pressurized. The crew, troubleshooting the wrong fault, lost consciousness to hypoxia; the Boeing 737 flew on autopilot until it ran out of fuel and crashed into a hillside near Athens, killing all 121 aboard. The investigation found a single cue carried two meanings with no differentiation, and a pilot hears the horn, across a career, only in the takeoff case.

## BACKGROUND

Helios 522 was a Boeing 737 departing Larnaca, Cyprus, for Athens. The night before, engineers had run a pressurization leak check and left the cabin-pressurization selector in the "manual" position. The pre-flight checks did not catch it, and the aircraft departed with its pressurization system not set to maintain cabin altitude automatically — a maintenance setting left in an abnormal position the following morning, the kind of latent configuration error that a pre-flight verification exists precisely to catch and here did not.<sup>0</sup>

## WHAT HAPPENED

As the aircraft climbed on 14 August 2005, the cabin failed to pressurize and a warning horn sounded. On the 737 the same horn indicates an incorrect takeoff configuration on the ground and excessive cabin altitude in flight. The crew interpreted it as a configuration warning and spent minutes troubleshooting the wrong problem while hypoxia set in. By the time the aircraft neared Athens the flight crew was incapacitated; it circled on autopilot until fuel exhaustion and crashed, killing all 121 aboard. The crew spent their few useful minutes diagnosing a configuration fault that did not exist, the ambiguous horn having steered them away from the one problem — climbing cabin altitude — that was quietly disabling them.<sup>1</sup>

## THE INVESTIGATION

The Hellenic Air Accident Investigation Board listed among its direct causes the non-identification of the warnings and of the reasons for their activation — not only the cabin-altitude warning horn but the passenger oxygen-mask deployment indication and the Master Caution that followed it — alongside the failure to recognize the manual pressur-

ization setting and the crew's incapacitation by hypoxia. Its latent causes reached past the cockpit: deficiencies in the operator's organization, quality management and safety culture; the regulator's long-standing inadequate oversight; inadequate application of crew resource management; and the ineffectiveness of the measures the manufacturer had taken after earlier pressurization incidents on the type. The single horn carried two meanings with no differentiation between them, and a pilot's career exposure to it was the ground meaning; the in-flight case was an edge condition the crew had not been prepared to recognize — so an ambiguous cue and a one-sided training regime combined to make the wrong interpretation the natural one, with hypoxia closing the window before the error could be caught.<sup>2</sup>

#### THE CAPABILITY GAP

Helios 522 is the textbook example of a single cue carrying two meanings without differentiation — an interface that lied by ambiguity while the training filled in only one of the two meanings. The missing capability was a cue, or a training regime, that distinguished the deadly in-flight meaning from the routine ground one before the operators had to guess under hypoxia — a way to resolve the ambiguity at the interface or in training, rather than leaving it to be resolved by a crew already losing the cognition the decision required.<sup>3</sup>

#### AFTERMATH & REFORM

The investigation drove changes to Boeing's warning logic and to pressurization-related training and checklists across operators; the FAA in 2011 required warning lights that separate cabin altitude from the takeoff-configuration warning. The accident is a standard case study in cue ambiguity, pre-flight verification, and the physiology of hypoxia in human-factors curricula — its remedies attacking the failure on both fronts at once, differentiating the warning at the interface and reinforcing the in-flight meaning in training so neither gap could be papered over by the other.<sup>4</sup>

| ACTUAL                                                            | INTERPRETED |
|-------------------------------------------------------------------|-------------|
| cabin alt.                                                        | T/O config. |
| <i>one horn, two meanings · the training had drilled only one</i> |             |
| HELIOS 522 — A CUE CARRYING TWO MEANINGS WITHOUT DIFFERENTIATION  |             |

*The Board concluded that the flight crew confused the two meanings of the warning horn. Similar occurrences had been reported by flight crews worldwide in the past.*

HELLENIC AAIASB AIRCRAFT ACCIDENT REPORT 11/2006 ON HELIOS 522, §2.2.3

#### REFERENCES

1. FAA Airworthiness Directive 2011-03-14, Amendment 39-16598, 76 FR 6529 (7 February 2011) — the mandated installation of separate CABIN ALTITUDE and TAKEOFF CONFIG warning lights on the 737 Classic fleet.
2. AAIASB Report 11/2006 — the dual-meaning warning horn and the training emphasis.
3. Boeing 737 Flight Crew Operations Manual — configuration and cabin-altitude warnings.

2. Hellenic Air Accident Investigation & Aviation Safety Board, Aircraft Accident Report 11/2006 — probable cause and the warning misinterpretation (paraphrased).

3. AAIASB Report 11/2006 — the manual pressurization selector and failure to pressurize.

6. Reason, J. (1990), Human Error — the cue-ambiguity / human-error framing the case draws on.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Helios 522 is the textbook example of a single cue carrying two meanings without differentiation. The interface lied by ambiguity. The training filled in only one meaning. The combination produced a window of a few minutes — the crew’s last useful ones — in which they were solving the wrong problem.

LENS APPROACH

LENS uses Helios in LEN 5 as a case in cue ambiguity and in LEN 2 as a Human-AI Teaming case where the interface failed to communicate which of two possible system states it was warning about. Students redesign the cue.

| DOMAIN                     | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|---------------|------------------|--------------|---------|-----|
| AVIATION                   | → T H         | → D1 D2 D3 D4 D5 | 3            | 3.1     | 3   |
| Human-System Collaboration |               |                  |              |         |     |

APPROACHES TO CONSIDER

- DURING DEVELOPMENT

▶ Differentiate cues that carry distinct meanings in distinct contexts, so a single signal cannot be read as the wrong one under stress.

▶ Train both meanings of any dual-purpose cue, not just the routine one, so the rare-but-deadly interpretation is recognized first.

▶ Design pre-flight verification to catch abnormal maintenance configurations, so a setting left in “manual” cannot survive into operation.
- IN OPERATION / AFTER

▶ Audit checklists and warning logic for cues whose meaning depends on phase or context, and resolve the ambiguity before an operator must guess.

▶ Monitor maintenance-to-operations handoffs for latent configuration errors that pre-flight checks have historically missed.

▶ Sustain training on the physiology that degrades the operator (e.g., hypoxia), so the window for self-correction is understood and guarded.

REFLECTION QUESTIONS

1. Identify a cue in your domain that carries two meanings in different operational contexts. How would the operator know which?

2. Redesign the Helios configuration warning to distinguish ground- and in-flight meaning without adding cognitive load.

3. A maintenance setting left in “manual” survived the pre-flight checks and reached the air. Design the pre-flight verification step in your domain that would make an abnormal maintenance configuration impossible to depart with unnoticed.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · usability testing · cognitive walkthroughs

# TransAsia Airways Flight 235

**T** Training Gap   **H** Human-System Interface

**IMPACT** 43 killed in Taiwan; crew shut down the only working engine after the other failed

## IN BRIEF

About thirty-seven seconds after takeoff from Taipei Songshan in February 2015, the right engine of a TransAsia ATR 72 auto-feathered following an intermittent fault in its auto-feather unit. Working from memory under acute time pressure, the crew shut down the left engine — the one still producing thrust — leaving the aircraft with no power and too little altitude to recover. It clipped a viaduct, struck a taxi, and crashed into the Keelung River, killing 43 of the 58 aboard. The Aviation Safety Council found the captain had failed his command-upgrade simulator check in May 2014, and that the crew skipped required items of the published engine-flameout drill and dropped the cross-check callouts that precede a shutdown. The wrong-engine error mirrors Kegworth (Case 103) twenty-six years earlier.

## BACKGROUND

TransAsia 235 was an ATR 72-600 twin-turboprop departing Taipei Songshan. The captain had failed his command-upgrade simulator check on 31 May 2014 — the unsatisfactory items included insufficient knowledge of the engine-flame-out-at-takeoff drill — passed a re-check a month later, and had been noted in line training for weak handling of abnormal procedures. The published response to an engine failure was a memory-item drill, and the drill carried mandatory cross-check callouts before any lever moved — so a captain already flagged for weak abnormal-procedure handling was exactly the operator most likely to skip the confirming callout precisely when it mattered, a known weakness meeting a known failure mode.<sup>0</sup>

## WHAT HAPPENED

On 4 February 2015, about thirty-seven seconds after takeoff, the right engine's autofeather system activated after an intermittent break in the signal between the number 2 auto-feather unit and its torque sensor, traced afterwards to compromised solder joints inside the unit — the engine itself was capable of producing power. The crew, identifying an engine problem, shut down the left engine, which was operating normally. With neither engine now producing useful thrust and the aircraft low and slow, it stalled, clipped a viaduct and a taxi, and crashed into the Keelung River, killing 43 of the 58 aboard. A recoverable single-engine event became unrecoverable the moment the working engine was shut down, leaving an aircraft with no thrust and no altitude in the seconds just after takeoff when both are least forgiving.<sup>1</sup>

## THE INVESTIGATION

The Taiwan Aviation Safety Council found the crew failed to identify which engine had actually failed and shut down the wrong one. They did not perform the documented failure-identification procedure before acting and skipped required items of the flameout drill; the pilot flying retarded the wrong power lever about four seconds after the master warning, before the pilot monitoring had finished identifying the failure, and that engine was shut off 46 seconds after it. The Board also found that the crew had not rejected the takeoff when the automatic take-off power control system failed to arm on the roll, and that TransAsia had no clear documented policy telling them to. The captain's documented proficiency deficiencies and the airline's training and crew-resource-management shortfalls were contributing factors — the individual lapse sitting atop an organizational one, the carrier having left the very weaknesses its own proficiency checks had recorded uncorrected before the flight.<sup>2</sup>

## THE CAPABILITY GAP

TransAsia 235 is the modern recurrence of the wrong-engine pattern that destroyed British Midland 92 at Kegworth in 1989 (Case 103). Crews under acute time pressure default to memory; if the memory is incomplete, the action follows the flawed memory rather than the checklist. The missing capability is a procedure design and training regime that keeps the verification step in the loop precisely when stress pushes operators to skip it — engineering the confirm-which-engine step so it fires under startle, rather than trusting it to a memory that time pressure is actively eroding.<sup>3</sup>

## AFTERMATH & REFORM

The ASC issued recommendations on TransAsia's pilot training, proficiency-check standards, and crew-resource management; the airline ceased operations in November 2016; GE235 was its second fatal crash in seven months, after Flight GE222 at Magong in July 2014. The case is used to argue that the persistence of the wrong-engine pattern across a quarter-century reflects an un-engineered intervention, not merely individual error — the same failure recurring across crews, airframes, and decades points to a design and training gap that no amount of blaming the individual pilot has ever closed.<sup>4</sup>



TRANSASIA 235 — THE SAME WRONG-ENGINE PATTERN AS KEGWORTH (CASE 103)

*The flight crew was unable to identify which engine had failed and shut down the operating engine in error.*

PARAPHRASING THE AVIATION SAFETY COUNCIL (TAIWAN) FINAL REPORT ON TRANSASIA 235, 2016

REFERENCES

1. Aviation Safety Council (Taiwan), Aviation Occurrence Report: TransAsia Airways Flight GE235, final report (2016) — findings and the wrong-engine shutdown (paraphrased).

2. ASC final report (2016) — the autofeather event and the crash sequence.

3. ASC final report (2016) — non-use of the checklist and the captain's proficiency record.

4. AALB (UK), British Midland Flight 92 (Kegworth) report (1990) — the antecedent wrong-engine case.

5. Dismukes, Berman & Loukopoulos (2007), The Limits of Expertise — memory-driven response under time pressure.

THE LEARNING ENGINEERING LENS

LE INSIGHT

TransAsia 235 is the modern recurrence of the Kegworth pattern under stress. Crews under acute time pressure default to memory; if the memory is incomplete or wrong, the action follows the memory. The capability gap was the same as in 1989. The intervention pattern needed is also the same.

LENS APPROACH

LENS uses TransAsia 235 in LEN 5 as a recurrence case for well-known capability problems (wrong-engine shutdown) and in LEN 2 to teach the difference between checklist-driven response and memory-driven response under stress.

| DOMAIN                     | FAILURE MODES               | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|-----------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| AVIATION                   | → <div>T</div> <div>H</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 3            | 3.1     | 3   |
| Human-System Collaboration |                             |                                                                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the abnormal-procedure response so the confirm-which-unit verification step is forced under startle, not left to memory.
- ▶ Build proficiency checks whose red signals have the authority to ground an operator before the predicted failure, not merely to document a weakness.
- ▶ Train the checklist-driven response under realistic time pressure, so the trained reflex survives the stress that pushes crews to memory.

IN OPERATION / AFTER

- ▶ Audit whether documented proficiency deficiencies are actually acted on, treating an uncorrected red check as an organizational, not individual, failure.
- ▶ Monitor for recurrence of well-known patterns (wrong-engine shutdown) across crews and airframes as evidence the intervention is un-engineered, not merely human error.
- ▶ Sustain crew-resource-management and checklist discipline so the verification step does not erode back to memory under operational pressure.

REFLECTION QUESTIONS

1. Identify a procedure in your domain that is supposed to be checklist-driven but is actually memory-driven under stress. What is the gap?
2. Why has the wrong-engine shutdown pattern persisted across 25 years of aviation reform? What intervention has not yet been engineered?
3. The captain failed his upgrade simulator check in May 2014 and passed a re-check, and the airline never addressed the weakness the check had recorded. Design the proficiency-check regime in your domain whose red signal would actually remove an operator from duty before, not after, the failure it predicts.

**WHO BUILDS THIS** learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs

# Eastern Air Lines Flight 401

**H** Human-System Interface   **T** Training Gap

**IMPACT** 101 killed in the Everglades; the entire flight crew fixated on a landing-gear indicator bulb while the autopilot silently disengaged

## IN BRIEF

On 29 December 1972 an Eastern Air Lines L-1011 descended into the Florida Everglades after breaking off its night approach to Miami, killing 101 of the 176 aboard. The cause was attention, not mechanics: all three flight-deck crew became absorbed in a landing-gear indicator bulb that had failed to light, and while they worked it, the autopilot's altitude hold was inadvertently disengaged. The aircraft sank slowly; an altitude alert chimed once and no crewmember remarked on it, and no one was monitoring the flight path. The NTSB findings launched decades of research into attention, monitoring, and cockpit-alert design, and the accident is a formative event behind Crew Resource Management. Eastern 401 is the canonical case of a low-priority task crowding out a life-critical one.

## BACKGROUND

Eastern Air Lines Flight 401, a wide-body Lockheed L-1011, broke off its night approach to Miami on 29 December 1972 when the nose-gear light failed to come on, and held west of the airport at 2,000 feet with a captain, first officer, flight engineer, and an Eastern maintenance specialist on the jumpseat.<sup>0</sup> Three crew was the era's safeguard against any one person being overloaded, the assumption being that more eyes meant more coverage; the night would show that without a designed division of those eyes, three people could converge on a single trivial problem as easily as one.

## WHAT HAPPENED

The crew noticed that the landing-gear indicator light had not illuminated, and all three focused on the bulb. While they worked the problem, one of them inadvertently nudged the controls and disengaged the autopilot's altitude hold. The TriStar began a slow descent into the Everglades; a single half-second chime sounded at 250 feet off the selected altitude, and no crewmember remarked on it. The aircraft hit the swamp, killing 101 of the 176 aboard.<sup>1</sup> The descent was gentle enough to go unfelt, the flight engineer was below deck in the electronics bay hunting for the gear indices, and the amber altitude-alert light was inhibited below 2,500 feet on Eastern's L-1011s — so a half-second tone was the whole of the warning the cockpit got.

## THE INVESTIGATION

The NTSB found the crew had become so engrossed in the landing-gear difficulty that they failed to monitor the flight path.<sup>2</sup> A burned-out indicator bulb had absorbed the attention




of an entire qualified crew while a wide-body airliner descended unwatched — the disproportion between the trigger and the outcome is exactly what made the case so durable a teaching example. The findings inaugurated decades of research into attention, monitoring, and the design of cockpit alerts — the accident is cited in nearly every introductory cognitive-engineering course as the example of how a low-priority task can crowd out a life-critical one when attention is undivided among the available channels.<sup>3</sup>

#### THE CAPABILITY GAP

The missing capability was a designed division of attention across the crew, and an alert that carried the priority its message deserved. Three competent people attended to a burned-out bulb while no one watched the altitude; the chime existed but did not cut through. The flight deck had the people and the information — it lacked the design that would have kept one channel of attention on the thing that could kill them. Attention had been treated as something the crew would naturally allocate well, rather than as a resource the cockpit had to be built to protect, and that unexamined assumption was the gap.<sup>4</sup>

#### AFTERMATH & REFORM

Eastern 401 was one of the formative events behind Crew Resource Management (Case 117) — the explicit allocation of monitoring and cross-checking roles — and behind modern standards for prioritized cockpit alerting.<sup>5</sup> Both reforms encode the same correction: CRM assigns someone to keep flying the aircraft while others troubleshoot, and prioritized alerting ensures the most dangerous condition is the loudest one — turning the night's two failures into permanent structural defenses. Its lesson is that attention is a designable parameter, not a personal virtue: a system that lets all eyes converge on one task must guarantee the critical channel is still watched.



*Preoccupation with a malfunction of the nose landing gear position indicating system distracted the crew's attention from the instruments and allowed the descent to go unnoticed.*

NTSB AIRCRAFT ACCIDENT REPORT NTSB-AAR-73-14 ON EASTERN 401, 1973, PROBABLE CAUSE

#### REFERENCES

1. NTSB, Aircraft Accident Report: Eastern Air Lines Flight 401, NTSB-AAR-73-14 (1973) — the night approach and full flight deck.
4. C. D. Wickens et al., Engineering Psychology and Human Performance — attention, monitoring, and alert design.

2. NTSB-AAR-73-14 (1973) — the landing-gear-bulb fixation, the inadvertent autopilot disengagement, and 101 deaths of 176 aboard.

3. NTSB-AAR-73-14 (1973) — the crew's preoccupation with the landing-gear indication that distracted them from maintaining altitude (paraphrased).

5. R. Helmreich & H. Foushee (1993), aircrew-coordination research — the Crew Resource Management origins.

6. Cooper, G. E., White, M. D. & Lauber, J. K. (1980), Resource Management on the Flightdeck, NASA CP-2120 — proceedings of the 1979 NASA/industry workshop that named Crew Resource Management.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Eastern 401 is the canonical attention-failure case. The capability that was missing was a designed division of attention across the crew. Eastern 401 is one of the accidents behind CRM, and the discipline of cockpit alert design exists because the altitude warning chime did not carry the priority weight it needed to.

LENS APPROACH

LENS uses Eastern 401 in LEN 5 to teach attention as a designable parameter and in LEN 2 to introduce alert prioritization as a capability deliverable. The case anchors the CRM origin story.

| DOMAIN                     | FAILURE MODES               | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|-----------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| AVIATION                   | → <div>H</div> <div>T</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 3            | 3.1     | 3   |
| Human-System Collaboration |                             |                                                                         |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                   |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div><div>▶ Build an explicit division of attention into the operating procedure, so at least one role is always assigned to the life-critical channel while others troubleshoot.</div><div>▶ Design alerts with a priority hierarchy, ensuring the most dangerous condition is also the most salient rather than one indistinguishable chime among many.</div><div>▶ Treat attention as a resource the system must protect by design, not a virtue the operators are assumed to supply on their own.</div></div> | <div><div>▶ Drill crews on scenarios where a trivial problem competes with a critical one, and verify the monitoring role actually holds under that pressure.</div><div>▶ Audit whether alerts in service are heard and acted on, retiring or redesigning any that get lost in routine clutter.</div><div>▶ Track distraction-related near-misses so attention-displacement failures are visible before they cause a loss.</div></div> |

REFLECTION QUESTIONS

1. Identify a low-priority task in your domain that could plausibly absorb all of an operator’s attention. What is the life-critical task that would be displaced?
2. Redesign the altitude warning chime of an L-1011 so that it cuts through a focused troubleshooting conversation.
3. Eastern 401 had three qualified people and still left the flight path unwatched. How would you assign and verify a “someone is always watching the critical channel” role so it cannot collapse when the whole team is drawn to one problem?

**WHO BUILDS THIS** human factors & interaction design · learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · scenario simulation · competency models

# NASA EVA-23 — Water Intrusion Inside a Spacesuit

**H** Human-System Interface   **N** Normalization of Deviance   **D** Designed Out

**IMPACT** During the second of two ISS spacewalks, ESA astronaut Luca Parmitano's helmet filled with up to 1.5 liters of water from the suit's cooling-and-ventilation loop, blinding him, fouling his communications, and threatening drowning in vacuum; the EVA was terminated and Parmitano recovered to the airlock guided by tether feel — the NASA Mishap Investigation Board classified it a High Visibility Close Call and called the condition life threatening

## IN BRIEF

On July 16, 2013, during EVA-23 outside the International Space Station, water began collecting inside the helmet of European Space Agency astronaut Luca Parmitano. The contamination grew until the water covered his eyes, ears, and nose, fouled his communications microphone, and would have drowned him in vacuum had the spacewalk continued. The crew terminated the EVA and Parmitano made his way back to the airlock partly by feel along the tether. The NASA Mishap Investigation Board (MIB) report issued later that year identified the proximate cause as blocked separator holes in the suit's Fan/Pump/Separator assembly that allowed cooling-loop water into the ventilation loop, and the deeper cause as a missed opportunity at the previous EVA-22 a week earlier when a similar (but smaller) water event had been mis-attributed to in-suit drink-bag leakage. The case is the canonical instance in modern human spaceflight of a failure mode the suit was not instrumented to detect, surfacing first as a near-fatal event because the earlier weak signal was rationalized away. The MIB's recommendations span hardware redesign, in-suit water-detection instrumentation, and the safety-culture response to anomalous EVA telemetry.

## BACKGROUND

EVA-23 was the second of two scheduled spacewalks for ISS Expedition 36. Parmitano and NASA astronaut Chris Cassidy exited the Quest airlock to continue tasks begun during EVA-22 a week earlier. Roughly 44 minutes into the EVA, Parmitano reported water at the back of his head inside the helmet. The EVA ran on for roughly 23 more minutes while ground controllers assessed; the water continued accumulating, eventually reaching the front of the helmet where it pooled across Parmitano's eyes, blocked his nose, and fouled the communications microphone. The EVA was terminated and Parmitano made the return traverse to the airlock with degraded vision and communication, partly guided by his sense of the tether.<sup>o</sup>

## WHAT HAPPENED

The proximate hardware fault, identified by the NASA Mishap Investigation Board, was in the suit's Fan/Pump/Separator assembly — the component that separates cooling water from the breathing-gas ventilation loop. Contamination had blocked drum-holes in the water separator. Once the separator could not perform its function, water from the cooling loop migrated into the ventilation loop and was pushed up into the helmet by the suit's fan. The hardware had no in-suit instrumentation to detect free water in the ventilation path; the crew detected the failure only when Parmitano felt liquid on his skin.<sup>1</sup>

## THE INVESTIGATION

The MIB's deeper finding was about EVA-22, the prior spacewalk one week earlier. Half a litre to a litre of water was found in the same helmet when it was removed after that EVA, in the same suit later worn on EVA-23. The crew reasoned the water must have entered during repressurization, when Parmitano was leaning forward and could have pinched the drink-bag bite valve open with his chin; the ground team accepted that assessment and the water was not investigated further. No suit teardown or anomaly investigation was completed before EVA-23 was approved. Among the MIB's five root causes, all organizational, were the EVA community's unexamined belief that drink bags leak — a determination no one in operations, engineering, safety or the crew office challenged — and the flight control team's perception that invoking the anomaly-report process was so resource-intensive it would not be worth what it turned up. The EVA-22 water was the weak signal of the same failure mode that nearly killed Parmitano a week later, and an unverified benign explanation was treated as adequate clearance to proceed.<sup>2</sup>

## THE CAPABILITY GAP

The corrective actions span three layers. At the hardware layer the suit's water-separator and ventilation-loop hygiene were re-engineered, and a Helmet Absorption Pad and a snorkel were added so a future water event would not immediately threaten the airway. At the operations layer water-quality specifications and monitoring were tightened, but no in-suit free-water sensor was added to the ventilation loop: after a March 2022 spacewalk a crewmember again found water pooled on his visor — noticed by the astronaut, not by instrumentation — and NASA paused routine U.S. spacewalks for seven months, concluding afterwards that the cause was condensation from integrated system performance rather than a separator failure. At the safety-culture layer the MIB pressed the documented-and-investigated discipline: an unexplained EVA event requires investigation rather than the most convenient explanation, particularly when the next planned EVA reuses the same hardware.<sup>3</sup>

## AFTERMATH & REFORM

The hedge that survives into the case is about attribution depth. The MIB was clear that the in-helmet and hardware fixes address the specific failure mode that produced EVA-23. The harder, lifecycle question — how a suit certified for decades of use developed a contamination pathway not represented in its anomaly catalog, and how future suits will avoid the analogous category of uninstrumented failure — is left as an open sustainment-engineering problem. The case is the canonical modern human-spaceflight instance of a recovery from a near-fatal event whose proximate cause was hardware and whose deeper cause was the institutional decision to accept the most convenient explanation for an anomalous prior event.

*The previous occurrence is the strongest available signal of the next one; the institutional decision to accept the convenient explanation is what made the second event near-fatal.*

EDITORS' SYNTHESIS OF THE NASA MIB EVA-23 REPORT.

## REFERENCES

1. NASA International Space Station Program (2013), "International Space Station EVA Suit Water Intrusion High Visibility Close Call" — Mishap Investigation Board final report, December 2013.
2. NASA Aerospace Safety Advisory Panel (2014), Annual Report — EVA-23 follow-up actions including helmet absorption pad and snorkel.
3. Parmitano, L. (2013), "EVA 23: exploring the frontier," ESA astronaut blog, 20 August 2013 — the crewmember's account of the water event and the return traverse.
4. Chappell, Norcross, Abercromby et al. (2017), "Risk of Injury and Compromised Performance Due to EVA Operations," NASA Human Research Program Evidence Report — broader EVA risk context.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

EVA-23 is the canonical modern human-spaceflight case of a failure mode the suit was not instrumented to detect, surfacing first as a near-fatality because the prior weak signal was rationalized away. The hardware fix and the new helmet pad and snorkel are necessary; the safety-culture half — every instance of free water gets investigated — is what generalizes.

### LENS APPROACH

EVA-23 is the human-spaceflight uninstrumented-failure case (induced 3.4; LENS D3/PT5). LENS uses it in Domain 3 (Human-System Collaboration) for the operator's detection-by-skin-contact at the edge of the engineered envelope, and in Domain 4 (Test and Evaluation) for the anomaly-investigation discipline the MIB pressed. Pair with Case 27 (anesthesia monitoring) at the cue/alert-as- capability layer and with Case 132 (F-22 OBOGS) as the aerospace companion in instrumentation-gap failure.

| DOMAIN             | FAILURE MODES                  | LENS COMPETENCY                                   | PROBLEM TYPE | INDUCED | LEO   |
|--------------------|--------------------------------|---------------------------------------------------|--------------|---------|-------|
| SPACE / AEROSPACE  | → <b>H</b> <b>N</b> <b>D</b> → | <b>D1</b> <b>D2</b> <b>D3</b> <b>D4</b> <b>D5</b> | 5            | 3.4     | 4 · 3 |
| HUMAN SPACEFLIGHT  |                                | Human-System Collaboration                        |              |         |       |
| SAFETY ENGINEERING |                                |                                                   |              |         |       |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Instrument the failure modes the system was not designed for. EVA-23 was detectable only by skin contact because the ventilation loop had no free-water sensor; instrumentation is the precondition for recovery.
- ▶ Treat the previous occurrence of an anomaly as the strongest available signal of the next one. EVA-22's smaller water event was the warning EVA-23 ignored; the institutional decision to accept the convenient explanation is what made the second event near-fatal.
- ▶ Engineer the recovery margin into the suit itself — the post-EVA-23 Helmet Absorption Pad and snorkel — so a contamination event has a survivable buffer before it becomes an airway emergency.

### IN OPERATION / AFTER

- ▶ Require an anomaly investigation, not a hypothesis, before reusing the same hardware for a subsequent high-consequence operation; the cost of a teardown is small against the cost of the failure mode.
- ▶ Build the in-suit instrumentation expectation into the certification process for future suits so the lesson generalizes beyond the specific Fan/Pump/Separator failure.
- ▶ Carry the sustainment-engineering hedge into program communication: hardware certified for decades can still develop pathways not represented in its anomaly catalog, and the certification process has to keep pace.

## REFLECTION QUESTIONS

1. Identify an operational system in your domain where a failure mode is detectable only by direct human sensing (skin contact, smell, the operator noticing). What instrumentation would convert that failure mode from sensed-late to sensed-early?
2. Specify the anomaly-investigation threshold you would set in your program: an unexplained prior event of the same class requires what level of teardown before the next operation is approved? EVA-22 to EVA-23 is the worked example of the cost of getting that threshold wrong.
3. The post-EVA-23 fixes addressed the specific failure mode. What is the lifecycle question your program owes for decades-certified hardware: how does the certification process keep pace with contamination pathways the original anomaly catalog did not represent?

**WHO BUILDS THIS** human factors & interaction design · safety science & organizational analysis · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · incident analysis · safety dashboards · task analysis · design prototyping

# Colgan Air Flight 3407

**T** Training Gap

**IMPACT** 50 killed near Buffalo; precipitated the FAA's 1,500-hour rule and the Pilot Records Database

## IN BRIEF

On approach to Buffalo in February 2009, the captain of Colgan Air 3407 responded to a stall warning by pulling back on the controls instead of lowering the nose; the Bombardier Q400 stalled and crashed into a house in Clarence Center, killing all 49 aboard and one person on the ground. The NTSB found the captain had a documented history of training failures, and the first officer was ill, fatigued, and paid about \$16,000 a year. The regulator held the record; the airline never retrieved it — PRIA obliged Colgan to collect prior-employer files, not the FAA's notices of disapproval, so the hiring decision ran on the captain's own account, which disclosed one failure of three. Victims' families mounted one of the most effective aviation-safety campaigns in a generation, producing the 2010 law that raised first-officer experience requirements to 1,500 hours and created the Pilot Records Database. The capability gap was visible to everyone except the system that licensed the pilots.

## BACKGROUND

Colgan Air 3407, a Bombardier Q400 flying a regional route to Buffalo, was crewed by a captain with a documented history of training failures and a first officer who was ill, fatigued, and paid roughly \$16,000 a year. The FAA held the captain's record of checkride disapprovals; Colgan did not obtain it. PRIA required a hiring carrier to gather five years of prior-employer records, not the regulator's own certification file, and the FAA route was optional and consent-gated — so the facts that should have shaped who sat in that seat stayed with the licensing authority and never crossed into the hiring decision.<sup>9</sup>

## WHAT HAPPENED

On 12 February 2009, on approach to Buffalo, the aircraft slowed and its stall-warning stick shaker activated. The captain responded by pulling back on the control column — the opposite of stall recovery — and the Q400 stalled and crashed into a house in Clarence Center, killing all 49 aboard and one person on the ground. The reflex he reached for was exactly wrong for the situation, the kind of error a pattern of failed training events should have predicted long before he was handed the controls.<sup>1</sup>

## THE INVESTIGATION

The NTSB found the probable cause to be the captain's inappropriate response to the stick shaker, against a backdrop of fatigue, weak airline training, and a hiring system blind to the captain's history. The data that should have flagged him existed — multiple failed training events — but did not flow to the decision that put him in the seat, so the hiring choice was

made on a record that omitted the very facts that mattered most, an absence that looked like a clean slate.<sup>2</sup>

### THE CAPABILITY GAP

Colgan is the canonical case for the gap between the information an institution holds about its operators and the information that reaches the decisions about those operators. The capability that was missing was not pilot skill in the abstract but a data flow: a way to make a pilot's documented training history actionable at the hiring and licensing decision, rather than buried in files no one consulted — a record can exist in full and still fail completely if it never reaches the moment and the person whose choice it was meant to inform.<sup>3</sup>

### AFTERMATH & REFORM

The victims' families organized one of the most effective aviation-safety lobbying efforts in a generation, producing the Airline Safety and FAA Extension Act of 2010 — which raised the minimum experience for airline first officers to 1,500 hours and created the Pilot Records Database to make pilot history flow between carriers.<sup>4</sup> The 1,500-hour rule remains debated; the records database fundamentally restructured how the information that was missing in 2009 now moves, converting a buried file into a record a hiring carrier is required to retrieve — the difference between a fact that exists and one that arrives where the decision is made.

50

KILLED, CLARENCE CENTER NY

*the data was in the system; the data flow was not*

COLGAN 3407 — AND THE REFORM THAT FOLLOWED

*The captain's inappropriate response to the activation of the stick shaker, which led to an aerodynamic stall from which the airplane did not recover.*

NTSB AIRCRAFT ACCIDENT REPORT AAR-10/01, PROBABLE CAUSE, 2010

### REFERENCES

1. NTSB/AAR-10/01 (2010) — the captain's training history, and the first officer's fatigue and pay.
2. NTSB, Aircraft Accident Report: Colgan Air Flight 3407, NTSB/AAR-10/01 (2010) — probable cause (quoted).
3. NTSB/AAR-10/01 (2010), §2.7.3 — PRIA and FAA guidance did not require operators to obtain notices of disapproval; recommendation A-05-1 reiterated as "Open—Unacceptable Response".
4. DOT Office of Inspector General, FAA Delays in Establishing a Pilot Records Database Limit Air Carriers' Access to Background Information, AV-2015-079 (2015); the Families of Continental Flight 3407 campaign.
5. Airline Safety and Federal Aviation Administration Extension Act of 2010, Pub. L. 111-216 — the 1,500-hour rule and the Pilot Records Database.

### THE LEARNING ENGINEERING LENS

### LE INSIGHT



Colgan is the canonical case for the gap between the information an institution holds about its operators and the information that reaches the decisions about whether those operators should be in that seat. The data was already in the system — multiple failed training events. The data flow that would have made that data actionable at the hiring decision did not exist.

#### LENS APPROACH

LENS uses Colgan in LEN 4 as a case for evidence-system design (the PRD as a designed information flow) and in LEN 8 as a case for advocacy-driven institutional change. Studio projects examine how families of victims became the load-bearing element of a reform the industry had resisted for years.

| DOMAIN              | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------|---------------|------------------|--------------|---------|-----|
| AVIATION            | → T           | → D1 D2 D3 D4 D5 | 2            | 2.1     | 4   |
| Test and Evaluation |               |                  |              |         |     |

#### APPROACHES TO CONSIDER

##### DURING DEVELOPMENT

- ▶ Make a pilot's documented training history a required, structured input to the hiring and licensing decision, not a record held passively in carrier files.
- ▶ Design the data flow so that adverse signals (failed training events) are pushed to the decision point rather than waiting to be requested.
- ▶ Define record retention and transfer obligations across carriers up front, so a clean-looking application cannot conceal a known history.

##### IN OPERATION / AFTER

- ▶ Audit whether the records actually reach hiring decisions, treating a complete file that never arrives as a live failure, not a solved one.
- ▶ Monitor the pipeline for operators whose adverse history is not surfacing, and act before the gap is paid for in an accident.
- ▶ Sustain the records database against drift and gaps in coverage, since a data-flow remedy decays as fast as its weakest link.

#### REFLECTION QUESTIONS

1. What information about operators in your domain exists somewhere in the system but does not flow to the decisions that depend on it?
2. Design the data-flow architecture that would make a Colgan-equivalent visible **before** the accident rather than after.
3. The 2010 reform was driven by victims' families, not the industry that had resisted it. What decisive constituency in your domain could force a stalled data-flow fix into existence, and what evidence would mobilize them?

**WHO BUILDS THIS** learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models

# Atlas Air Flight 3591

**T** Training Gap **K** Knowledge & Institutional Memory

**IMPACT** 3 killed; Boeing 767 cargo flight crashed in Texas after the first officer mistakenly activated go-around mode

## IN BRIEF

On approach to Houston in February 2019, the first officer of Atlas Air 3591 — an Amazon Prime Air cargo flight — inadvertently triggered the go-around mode during a turbulent descent. Experiencing a somatogravic illusion that the aircraft was pitching up, he pushed the nose down sharply and dove the Boeing 767 freighter into Trinity Bay, killing all three aboard. The NTSB found the first officer had a long, undisclosed history of training failures across multiple carriers — and that Atlas could not see it. The Pilot Records Database mandated by Congress after Colgan Air 3407 (Case 109) was not yet operational, so Atlas was working under PRIA, which retrieves records only from the employers an applicant names — and he named neither of the two airlines he had left without completing initial training, one of them inside PRIA's five-year window. Atlas also did not fully evaluate the adverse training records it did receive. The case is the live recurrence of Colgan: the information existed; the data-flow system was not yet complete.

## BACKGROUND

Atlas Air 3591 was a Boeing 767 freighter operating for Amazon Prime Air from Miami to Houston. The first officer had failed training events at several previous employers and had not disclosed them on his Atlas application. At the time of hiring, the Pilot Records Database directed by Congress after Colgan Air 3407 was not yet operational; Atlas used the older PRIA system, under which a hiring operator requests records from each former employer the applicant discloses. The first officer disclosed neither CommutAir (2011) nor Air Wisconsin (2012), the two airlines he had left after failing to complete initial training — the Air Wisconsin employment falling inside PRIA's five-year window, the CommutAir employment outside it. A retrieval process keyed to the applicant's own list is structurally incapable of showing what the applicant leaves off.<sup>9</sup>

## WHAT HAPPENED

On 23 February 2019, during a turbulent descent toward Houston, the first officer inadvertently activated the aircraft's go-around mode, which commanded a pitch-up. Experiencing a somatogravic illusion that the nose was rising into a stall, he pushed forward hard. The 767 entered a steep dive from which the crew did not recover and crashed into Trinity Bay, killing all three aboard. The startle and spatial disorientation that drove the fatal push were the same kind of breakdown his prior training failures had repeatedly recorded, had that history been visible to the carrier that hired him.<sup>1</sup>

## THE INVESTIGATION


The NTSB found the probable cause to be the first officer's inappropriate response to the inadvertent go-around activation, driven by a startle and spatial-disorientation response, and identified his pattern of training deficiencies and Atlas's inability to access his full record as contributing factors. The records that would have informed the hiring decision existed but did not reach it — the same information-flow failure named after Colgan a decade earlier, recurring here because the remedy built to fix it was authorized but not yet carrying the full history. The Board also found that Atlas did not fully evaluate the training-performance information it did receive from a previous employer, so the gap was not only in what the system delivered but in what the carrier did with what arrived.<sup>2</sup>

## THE CAPABILITY GAP

Atlas Air 3591 is the iteration test of the Colgan reform. The reform partially worked — the Pilot Records Database had been authorized — but it was not yet operational and its coverage was incomplete, so the same information-flow gap that killed fifty people in 2009 was still partly open in 2019. Partial implementation of a remedy leaves a measurable harm inside the aperture — the years between authorization and full coverage are not a neutral transition but a window in which the original failure can recur on the cases the incomplete remedy does not yet reach.<sup>3</sup>

## AFTERMATH & REFORM

The PRD final rule was published in 2021 and became effective for Part 121 carriers in 2022, with full historical coverage phasing in through 2024 — closing the gap the case exposed. Atlas 3591 is cited in implementation-fidelity discussions as evidence that a remedy is only as good as its completed coverage — that the date a rule is authorized and the date it actually reaches every carrier are different dates, and the distance between them is where the avoidable harm accumulates.<sup>4</sup>



THE COST OF PARTIAL IMPLEMENTATION BETWEEN COLGAN AND THE PRD'S FINAL COVERAGE

*Thus, if a pilot intentionally omits a previous employer, as the FO did, the hiring operator may never know the pilot's complete background.*

NTSB AIRCRAFT ACCIDENT REPORT AAR-20/02, SECTION 2.5.1, 2020

## REFERENCES

1. NTSB/AAR-20/02 (2020) — the first officer's training history and Atlas's record-access limits.
4. DOT Office of Inspector General, FAA Delays in Establishing a Pilot Records Database Limit Air Carriers' Access to Back-

2. NTSB/AAR-20/02 (2020) — the inadvertent go-around activation and the dive into Trinity Bay.

3. NTSB, Aircraft Accident Report: Atlas Air Flight 3591, NTSB/AAR-20/02 (2020) — probable cause and contributing factors (paraphrased).
5. FAA Pilot Records Database final rule (Federal Register, 10 June 2021; subpart B/C compliance from June 2022, full historical coverage by September 2024); Pilot Records Improvement Act of 1996.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Atlas Air 3591 is the iteration test of the Colgan reform. The reform partially worked: the PRD exists. The reform was incomplete: not all carriers were covered. The case is a cautionary worked example of why partial implementation of a regulatory remedy leaves the original capability gap partially open.

LENS APPROACH

LENS uses Atlas Air 3591 in LEN 4 as a case for measurement-system coverage (the PRD had been built but its mandatory coverage was incomplete), and in LEN 8 for the iteration cycle of reform: an intervention that leaves an aperture creates a measurable harm inside the aperture.

| DOMAIN              | FAILURE MODES               | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|---------------------|-----------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| AVIATION            | → <div>T</div> <div>K</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 2            | 2.4     | 4   |
| Test and Evaluation |                             |                                                                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the records system to require a full operator history, not a fixed-length window that can hide a multi-employer pattern.
- ▶ Plan the coverage phase-in explicitly, with interim controls protecting the cases the remedy does not yet reach.
- ▶ Specify the data the hiring decision needs and make its delivery mandatory before the decision, not optional or self-reported.

IN OPERATION / AFTER

- ▶ Audit coverage completeness, not just authorization, and treat the gap between the two as an active source of harm until closed.
- ▶ Monitor for adverse events occurring inside the aperture of a partially implemented remedy, and use them to accelerate full coverage.
- ▶ Sustain attention through the full phase-in, since a remedy authorized but not yet universal leaves the original gap measurably open.

REFLECTION QUESTIONS

1. Identify a regulatory remedy in your domain whose coverage is partial. What harm is occurring inside the aperture?
2. Design the deliverable that would close the PRD coverage gap in advance of a future Atlas Air 3591.
3. The PRD was authorized in 2010 but not fully effective until 2024, and Atlas 3591 fell inside that gap. For a remedy in your domain, design the interim control that would protect the cases the remedy does not yet reach while its coverage phases in.

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · knowledge bases · data pipelines

# Challenger & Columbia

**N** Normalization of Deviance   **K** Knowledge & Institutional Memory   **G** Governance & Trust

**IMPACT** 14 astronauts killed (7 per accident); 17 years between identical organizational pathologies

## IN BRIEF

NASA lost two Space Shuttle crews to the same organizational pathology seventeen years apart: Challenger in 1986, when O-ring seals failed in cold weather, and Columbia in 2003, when foam debris breached the wing's thermal protection — fourteen astronauts in all. Both flaws had been seen repeatedly and accepted as routine; the Board found foam loss on more than 80 percent of the 79 missions for which launch imagery existed — including seven left-bipod-ramp sheddings among 72 imaged missions, some of which NASA had not known about — and O-ring erosion in the field joints had been on the engineering record since the early flights. Sociologist Diane Vaughan named the mechanism “normalization of deviance” from the Challenger investigation; the Columbia Accident Investigation Board found the same culture intact seventeen years later and concluded NASA's organizational culture had as much to do with the accident as the foam. The pair is the strongest single-institution evidence that culture is an engineerable property of a system — and that a pathology diagnosed but left unrepaired will recur, at the same cost.

## BACKGROUND

The Space Shuttle flew with two flaws its engineers knew about. The solid-rocket booster joints sealed with O-rings that stiffened in cold; the external tank shed foam insulation that struck the orbiter on ascent. Both had appeared on flight after flight without catastrophe, and both were progressively reclassified from hazards to accepted, “routine” features of flying.<sup>0</sup> Each survived flight became evidence that the flaw was tolerable, so the very absence of disaster fed the reclassification — the safety margin redefined downward not by decision but by the accumulating habit of getting away with it.

## WHAT HAPPENED

On 28 January 1986 Challenger launched on an unusually cold morning; an O-ring failed to seal, and the vehicle broke apart 73 seconds after liftoff, killing all seven aboard. Seventeen years later, on 1 February 2003, foam that had struck Columbia's wing on ascent had opened a breach in its thermal protection; the orbiter disintegrated on reentry, killing its seven.<sup>1</sup> Both flaws were the ones engineers had already flagged and the institution had already filed as routine, so each accident was less a surprise than the arrival of a bill the system had decided, repeatedly, not to pay.

## THE INVESTIGATION


The Rogers Commission traced Challenger to the O-ring and to a launch decision that overrode engineers' cold-weather warnings — the Thiokol teleconference on the eve of launch, at which the engineering position ("do not launch below 53°F") was reversed under management pressure and the launch went ahead at 36°F.<sup>2</sup> Diane Vaughan's *The Challenger Launch Decision* (1996) re-examined the same teleconference with a decade's documentary access and reframed the decision: not deviant individuals overriding good engineering, but a working group whose rules of evidence had drifted, one accepted anomaly at a time, until a launch at 36°F was inside what the group's own decision rules counted as acceptable risk. Seventeen years later the Columbia Accident Investigation Board found the same patterns intact — flying with flaws defined as routine, foam strikes filed under "in-family" across the great majority of imaged missions, and a structure that suppressed the upward flow of safety concerns — and concluded that "the NASA organizational culture had as much to do with this accident as the foam."<sup>3</sup> That the same board found the same structure seventeen years on is the case's sharpest point: the first reform had fixed the hardware and the schedule pressure but not the mechanism by which a flaw becomes redefined as acceptable.

## THE CAPABILITY GAP

Diane Vaughan's "normalization of deviance," developed from Challenger, named the mechanism: deviations from the safety baseline become acceptable through production pressure, weak communication, and habit, one small step at a time. Columbia validated the concept against its author's intent — Vaughan herself was retained by the CAIB and the Board's final report cites the mechanism by name. The pathology had been diagnosed in 1986 and never engineered away — which is the point: culture is a system property, and a diagnosis without a remediation decays.<sup>4</sup> Naming the mechanism did not arrest it, because a name lives in a report while the production pressure and the suppressed warnings live in the daily flow of work, where they kept doing what they had always done. The institutional-learning gap is itself the decisive finding: between 1986 and 2003 NASA had reorganized twice, lost an administrator, and adopted multiple safety initiatives, and the structural pathway from a dissenting engineer to a launch decision was substantially unchanged.

## AFTERMATH & REFORM

Each accident produced reform — the Rogers Commission's redesign of the booster joint, the CAIB's call to treat culture as a safety variable, the establishment of the NASA Engineering and Safety Center, and the Board's instruction to rebuild independent technical authority — and the Shuttle was retired in 2011 after completing the International Space Station assembly. The pair stands as the book's strongest evidence that organizational culture is engineerable, and that leaving it unengineered is a choice with a recurring, lethal cost.<sup>5</sup> The CAIB's insistence on independent technical authority conceded the deeper lesson the booster-joint redesign alone had missed in 1986: the upward path for a dissenting engineer is itself a piece of safety hardware, and one that has to be rebuilt and defended rather than assumed. The hedge that survives into the case: Vaughan's "normalization of deviance" is the analytic claim across both accidents and is never to be smoothed; the seventeen-year institutional-learning gap is the empirical claim, and the CAIB's cross-referencing of Rogers Commission language ("the same decision-making structures") is the documentary anchor.



*These repeating patterns mean that flawed practices embedded in NASA's organizational system continued for 20 years and made substantial contributions to both accidents.*

COLUMBIA ACCIDENT INVESTIGATION BOARD, 2003

## REFERENCES

1. W. Starbuck & M. Farjoun (eds.), *Organization at the Limit: Lessons from the Columbia Disaster* (2005) — independent academic re-analyses of the institutional-learning gap.
2. Columbia Accident Investigation Board, Report Vol. I (2003) — the foam strike; finding F3.2-5 (left bipod ramp shedding on 7 of the 72 missions for which launch or ET-separation imagery exists, including instances NASA was unaware of) and finding F3.2-7 (foam loss on more than 80 percent of the 79 missions with confirming imagery); the filing of strikes as “in-family,” the recurrence of the cultural pattern, and the call for an independent technical authority.
3. Rogers Commission, Report of the Presidential Commission on the Space Shuttle Challenger Accident (1986) — the O-ring failure, the Thiokol teleconference, and the cold-weather launch decision.
4. CAIB (2003) — “the NASA organizational culture had as much to do with this accident as the foam” (quoted); Rogers Commission (1986) and CAIB (2003) jointly on the suppressed upward flow of safety information across both accidents.
5. D. Vaughan, *The Challenger Launch Decision: Risky Technology, Culture, and Deviance at NASA* (Univ. of Chicago Press, 1996) — normalization of deviance; the working group's drift of decision rules.
6. NASA Engineering and Safety Center founding documents (2003 – present) — institutional response to the CAIB's call for independent technical authority; Shuttle retirement (STS-135, July 2011).

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Challenger/Columbia is the strongest documentary evidence that organizational culture is an engineerable property of a system, and that diagnoses without engineered remediations decay. The same pathology, twice, seventeen years apart, in the same institution, with the same human cost — and with the diagnosis already on the record. Vaughan's “normalization of deviance” names the mechanism and is the analytic claim across both accidents. Capability engineering treats culture as a deliverable.

### LENS APPROACH

Challenger/Columbia anchors the multi-layer-drift-and-cumulative- inadequacy competency (induced 7.4; LENS D5/PT4): a cascade of marginal-but-tolerable conditions across decision rules, schedule pressure, and communication structure aligned twice, seventeen years apart. LENS uses the pair in Domain 1 (Systems Analysis) for normalization of deviance as a systems concept and the institutional-learning gap as a measurable property; in Domain 5 (Navigating Sociotechnical Constraints) for the governance failure that allowed a diagnosed pathology to persist and for the upward-channel design the CAIB called a piece of safety hardware. Pair with Deepwater Horizon (Case 163) on the multi-layer-drift form, and with Bhopal (Case 159) and Fukushima (Case 169) on cumulative inadequacy in catastrophic-system operations.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|-----------------|--------------|---------|-----|
| SPACE / AEROSPACE                     | NKG           | D1D2D3D4D5      | 4            | 7.4     | 1·5 |
| Navigating Sociotechnical Constraints |               |                 |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                            | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                              |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Hold known flaws as open hazards with explicit owners, so a clean flight cannot silently reclassify a deviation as routine.</li><li>Build independent technical authority into the launch decision, giving dissenting engineers a path that production pressure cannot override.</li><li>Define the safety baseline quantitatively at design time, so any drift below it is a flagged change rather than an unremarked habit.</li></ul> | <ul style="list-style-type: none"><li>Re-audit every diagnosed-but-unrepaired pathology on a fixed cadence, verifying remediation in practice rather than on paper.</li><li>Track the upward flow of safety concerns as a measured signal — count and resolve dissents — so suppression becomes visible.</li><li>Treat the cultural mechanism, not just the hardware fix, as the deliverable that must persist between major incidents.</li></ul> |

REFLECTION QUESTIONS

1. What is the equivalent “diagnosed but not repaired” pathology in your domain? What evidence would close the loop?
2. The CAIB called culture engineerable. Sketch the engineering deliverable for the cultural intervention you would propose in your domain — including its measurement signal.
3. Seventeen years separated two accidents with the same root pathology. What mechanism in your organization would verify that a past diagnosis has actually been remediated rather than merely documented?
4. Vaughan reframed the Thiokol teleconference not as deviant override but as a working group whose decision rules had drifted. Identify a recurring decision in your domain whose rules of evidence may have drifted incrementally, and specify the audit that would surface the drift before the next high-consequence call.

WHO BUILDS THIS safety science & organizational analysis · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. TOOLS incident analysis · safety dashboards · knowledge bases · data pipelines · governance frameworks · bias & evidence audits



# NASA Saturn V Documentation

**K** Knowledge & Institutional Memory

**IMPACT** Foundational U.S. aerospace case for the cost of losing the institutional capability to build a system

## IN BRIEF

When NASA returned to heavy-lift rockets in the 2000s, it confronted an uncomfortable fact: the institutional capability to build a Saturn V — the rocket that sent Apollo to the Moon — had been lost. The drawings and documents survived; the practical knowledge of how the components were manufactured, tested, and assembled, and why each choice was made, had walked out with the workforce that retired by the 1990s. The vehicle could be redesigned but not reproduced. Apollo was documented to an unprecedented degree, yet the documentation was not the capability: when the engineers who built the system left, the system left with them. Saturn V is the book's strongest evidence that institutional capability lives in people and sustaining practices, not in the artifacts an institution leaves behind.

## BACKGROUND

The Saturn V that launched Apollo to the Moon was, by any measure, one of the best-documented engineering programs in history — exhaustive drawings, specifications, and test records.<sup>0</sup> That thoroughness is what makes the case bite: the program left behind close to everything a successor could ask for on paper, so any later difficulty in rebuilding it cannot be blamed on a thin archive but must lie in what the archive could not hold.

## WHAT HAPPENED

Yet by the mid-2000s, as NASA worked the Constellation program, it had become apparent that the institutional capability to build a Saturn V had been lost. The drawings existed; the practical knowledge of how the components were manufactured, tested, and assembled — and why particular choices had been made — had walked out with the workforce that retired by the 1990s. The vehicle could be redesigned. It could not be reproduced.<sup>1</sup> The distinction is exact: redesign starts from requirements and rebuilds the reasoning afresh, while reproduction needs the original making-knowledge, and it was that knowledge — not the paperwork describing it — that had left with the people who held it.

## THE INVESTIGATION

The case is canonical for the difference between documentation and institutional capability. Apollo's documentation was exhaustive, but it captured the **what**, not the tacit **how** — the judgment, the workarounds, the undocumented reasons — that lived in the people who did the work. When they left, that knowledge was in no archive to recover.<sup>2</sup> A drawing records the decision but not the deliberation behind it; the workaround that made a part manufacturable


and the reason a tolerance was set where it was lived in the doing, and once the doers retired there was no document from which to reconstruct it.

### THE CAPABILITY GAP

Saturn V is the strongest available evidence that institutional capability is not the same as the artifacts an institution produces. Capability lives in people and in the practices that sustain them; a library of drawings is a necessary record but not a transferable ability. An institution that treats documentation as preservation of capability is, quietly, letting the capability expire.<sup>3</sup> The danger is that the archive looks like insurance: its very completeness can reassure managers that the capability is safe, so the people and practices that actually carry it are allowed to disperse precisely because the paperwork seems to stand in for them.

### AFTERMATH & REFORM

The lesson reshaped how serious programs think about knowledge retention — apprenticeship, continuity of teams, deliberate capture of tacit reasoning, and not letting a critical capability rest in a handful of soon-to-retire heads.<sup>4</sup> It pairs with the chapter's other memory cases: knowledge, unlike a document, has to be actively kept alive, or it is gone in a generation. Each of the retention practices addresses the same root cause the Saturn V exposed — that tacit making- knowledge transfers only person to person — so apprenticeship and team continuity are not nice-to-haves but the actual mechanism by which a capability outlives the people who first held it.



*The Saturn V drawings exist. The Saturn V does not.*

PARAPHRASING THE NASA CONSTELLATION-ERA CAPABILITY DISCUSSION, C. 2005

### REFERENCES

1. R. E. Bilstein, Stages to Saturn: A Technological History of the Apollo/Saturn Launch Vehicles (NASA SP-4206, 1980) — the program and its workforce.
2. NASA Constellation Program documentation and reviews (2005–2010) — the difficulty of reproducing Saturn V capability.
3. The documentation-vs-capability distinction (editors' synthesis of the Saturn V record).
4. M. Polanyi, The Tacit Dimension (1966); I. Nonaka & H. Takeuchi, The Knowledge-Creating Company (1995).
5. J. S. Brown & P. Duguid, The Social Life of Information (2000) — tacit and institutional knowledge.

### THE LEARNING ENGINEERING LENS

### LE INSIGHT

The Saturn V case is the canonical evidence that institutional knowledge is not equivalent to documentation. The documents persist; the capability does not. Capability engineering must treat the people who hold tacit knowledge as a primary institutional asset.

LENS APPROACH

LENS uses the Saturn V case in LEN 8 to teach that retaining and overlapping the expert personnel who hold tacit making-knowledge is itself an engineered deliverable (induced 6.3) — not a by-product of writing better documents. The capability to rebuild a Saturn V was lost because workforce overlap and retention were never treated as a designed retention deliverable, and documentation alone cannot carry tacit process knowledge across a generation. The teaching point in LEN 1 is therefore not “document more thoroughly” but “engineer the personnel continuity that the archive can never substitute for.”

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|------------------|--------------|---------|-----|
| AVIATION                              | → K           | → D1 D2 D3 D4 D5 | 2            | 6.3     | 2   |
| Navigating Sociotechnical Constraints |               |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Capture tacit reasoning as the work happens — the why behind a choice, the workaround that made a part buildable — not just the resulting drawing.
- ▶ Build capability into teams and apprenticeship, so the making-knowledge has a living carrier rather than resting in a handful of soon-to- retire heads.
- ▶ Treat documentation as a record, never as a substitute for the people and practices that hold the capability.

IN OPERATION / AFTER

- ▶ Audit critical capabilities for single points of human failure — knowledge held by a few near-retirement practitioners — and transfer it before they leave.
- ▶ Periodically test reproducibility, not just whether the archive is complete, since a full archive can mask a capability that has already expired.
- ▶ Sustain continuity of teams and practice deliberately, so a capability is kept alive across generations rather than rediscovered at need.

REFLECTION QUESTIONS

1. Identify a capability in your domain that is currently held in the tacit knowledge of a small number of senior practitioners. What is your institution’s capacity to reproduce it after they retire?
2. Design the institutional practice that would preserve a capability against the retirement of its holders.
3. Saturn V was exhaustively documented yet could not be reproduced, because the archive captured the **what** and not the tacit **how**. Identify a capability in your institution whose documentation might be giving false reassurance — and describe what the paperwork is failing to capture.

WHO BUILDS THIS knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS knowledge bases · data pipelines

# Boeing Starliner

**K** Knowledge & Institutional Memory **D** Designed Out

**IMPACT** Multiple delays; the 2024 crewed flight left two NASA astronauts at the ISS for months; contemporary case for capability erosion at a legacy contractor

## IN BRIEF

Boeing's Starliner, meant to be the second U.S. commercial crew vehicle alongside SpaceX's Crew Dragon, accumulated failures across a decade: software errors on its 2019 uncrewed flight, valve corrosion scrubbing a 2021 launch, and propulsion-system trouble on its 2024 crewed test that left two NASA astronauts on the ISS rather than returning them as contracted — NASA brought them home on SpaceX months later. GAO and NASA reviews found capability erosion across multiple engineering disciplines at a contractor whose human-spaceflight record had once been definitive. The erosion looks partly generational (Apollo- and Shuttle-era engineers retired) and partly institutional (cost and schedule pressure, and the thinning of NASA's in-house depth to challenge the contractor). Starliner is the contemporary case for capability erosion at a legacy contractor.

## BACKGROUND

Boeing was awarded a NASA Commercial Crew contract to build Starliner as the second U.S. vehicle to carry astronauts to the ISS, alongside SpaceX's Crew Dragon — drawing on a human-spaceflight heritage that had once been definitive in American aerospace.<sup>0</sup> The award rested in part on that heritage: a contractor with a definitive human-spaceflight record was a presumed safe choice, and that presumption is exactly what the program would go on to test, since reputation was standing in for a current measurement of capability.

## WHAT HAPPENED

Instead Starliner stumbled across a decade: software errors marred its 2019 uncrewed test flight, valve corrosion scrubbed a 2021 launch, and propulsion-system problems on the 2024 crewed test flight left the two NASA astronauts aboard the ISS rather than returning on the contracted spacecraft — NASA brought them home on a SpaceX vehicle months later.<sup>1</sup> The three failures fell in different subsystems across three separate years, which is itself telling: a single bad part is bad luck, but software, valves, and propulsion failing in succession points to something broader than any one component — a decline running across the engineering organization rather than within one of its parts.

## THE INVESTIGATION

GAO and NASA reviews identified the program as exhibiting capability erosion across multiple engineering disciplines — software, valves, propulsion, integration testing — at a contractor whose track record had previously been definitive.<sup>2</sup> The erosion looked partly generational, as Apollo- and Shuttle-era engineers retired, and partly institutional: cost


pressure, schedule-driven decisions, and the loss of NASA's own in-house engineering depth to challenge the contractor. NASA's inspector general later named the drivers directly: overconfidence in Boeing's use of heritage systems — which is why integrated testing was never required under the contract — an unachievable schedule, and limited flight-simulation data, compounded by NASA's underuse of the data rights its own contract gave it, which narrowed what the agency could analyse when the flight failures came.<sup>3</sup> The two causes compound: as the contractor's senior depth thinned through retirement, the buyer's own engineering depth had also eroded, so the very expertise NASA would have needed to catch the slipping contractor was the expertise it had let go of.

#### THE CAPABILITY GAP

Starliner is the contemporary case for capability erosion at scale at a legacy contractor. The decline happened over decades and was visible only in retrospect, because the institutional architecture for catching a supplier's slow capability decay — and updating the buyer's confidence to match — did not exist. Reputation outran reality, and no instrument was measuring the gap.<sup>4</sup> Slow decay is the hard case to catch: no single year's results look alarming, the brand keeps the buyer's confidence high, and without an instrument that tracks the supplier's actual current capability the divergence is only legible once a crewed flight forces the reckoning.

#### AFTERMATH & REFORM

NASA leaned harder on independent reviews and on SpaceX as the reliable alternative, and the episode sharpened questions about how to sustain — and verify — capability at sole-source and legacy suppliers.<sup>5</sup> Starliner returned uncrewed in September 2024, and its two astronauts finally came home aboard a SpaceX Crew Dragon in March 2025 after roughly nine months aboard the station. NASA then downgraded Boeing's next flight, Starliner-1, to an uncrewed cargo run to prioritise crew safety; by mid-2026 no launch date was scheduled, the helium leaks and propulsion failures were still unresolved, and the agency's own inspector general judged human-rating certification likely to slip to 2027 — with NASA unable to say when the remaining testing would be finished or certification obtained. It pairs with Saturn V (Case 112): where that case lost a capability to retirement, Starliner shows the same erosion in slow motion at a living institution still carrying the brand of the capability it had let thin. Having a second supplier to fall back on is what let NASA absorb the failure, which is itself the lesson: where a capability can erode unseen, a verified alternative is the difference between an embarrassment and a crew with no way home.



*erosion visible across multiple engineering disciplines*

STARLINER — CAPABILITY DECAY AT A LEGACY CONTRACTOR

*NASA was overconfident in Boeing's design and potential success based on the provider's use of heritage systems and its long-standing space flight experience.*

NASA OFFICE OF INSPECTOR GENERAL, IG-26-011, 30 JUNE 2026

REFERENCES

1. U.S. GAO, NASA Commercial Crew Program: Significant Work Remains to Begin Operational Missions to the Space Station, GAO-20-121 (Jan. 2020); Schedule Uncertainty Persists for Start of Operational Missions, GAO-19-504 (June 2019).

2. NASA and Boeing joint Independent Review Team, Orbital Flight Test findings (7 July 2020) — 80 recommendations across software requirements, testing and simulation, process, and hardware.

3. NASA OIG, NASA's Management of Its Commercial Crew Program, IG-26-011 (30 June 2026); NASA, Starliner Propulsion System Anomalies during the Crewed Flight Test (redacted, February 2026).

4. NASA Aerospace Safety Advisory Panel, annual reports (2020–2025) — Commercial Crew oversight and the Panel's criticism of NASA's mishap-classification requirements.

5. NASA, Starliner's uncrewed landing at White Sands Space Harbor (6 September 2024) and the Crew-9 return of Wilmore and Williams aboard a SpaceX Dragon (18 March 2025).

6. Cf. Saturn V (Case 112); N. Augustine, Augustine's Laws (1986).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Starliner is the case for sustained capability erosion at a legacy contractor whose track record had previously been definitive. The erosion happened over decades and was visible in retrospect; the institutional architecture for catching it did not exist.

LENS APPROACH

LENS uses Starliner in LEN 5 to teach one clean capability gap: the contractor–NASA interface capability that thinned over a decade. As NASA shifted to fixed-price commercial contracting, its own oversight and insight capacity eroded, so the problems an integrated review would have caught instead surfaced in flight (against the cross-subsystem troubles as the visible symptom). The teaching point is that the buyer's review capability is an engineered deliverable in its own right — let it thin and a slipping supplier becomes legible only when a crewed flight forces the reckoning.

| DOMAIN                                | FAILURE MODES   | LENS COMPETENCY         | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|-----------------|-------------------------|--------------|---------|-----|
| SPACE / AEROSPACE                     | → <div>KD</div> | → <div>D1D2D3D4D5</div> | 2            | 6.3     | 1   |
| Navigating Sociotechnical Constraints |                 |                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Verify a supplier's current capability against present evidence rather than awarding on reputation and a once-definitive track record.
- ▶ Preserve enough in-house engineering depth to genuinely challenge a contractor, since a buyer who has let its own expertise erode cannot catch a slipping supplier.
- ▶ Maintain a verified second supplier where the stakes are crewed, so a capability that erodes unseen does not become a single point of failure.

IN OPERATION / AFTER

- ▶ Instrument supplier capability so slow, multi-year decay is visible before a high-stakes flight forces the reckoning.
- ▶ Watch for cross-subsystem patterns — failures in software, valves, and propulsion in succession — as a signal of organization-wide erosion, not isolated bad luck.
- ▶ Update institutional confidence to match measured reality, so a contractor's brand cannot keep outrunning its current performance.

## REFLECTION QUESTIONS

1. Identify a legacy supplier in your domain whose capability track record may have eroded faster than your institutional confidence in them has updated. What is the signal?
2. Design the contractor–capability review that would have caught the Starliner gaps before the 2024 crewed flight.
3. Starliner’s failures spanned software, valves, and propulsion across three separate years — a pattern, not a single bad part. In your domain, what would distinguish a one-off component failure from organization-wide capability erosion, and how soon could you tell them apart?

**WHO BUILDS THIS** knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · task analysis · design prototyping





Chapter 6

# Aviation & Aerospace — What Works — and the Frontier

*When an industry engineers its own error rate down and keeps the books open.*

---

*When reporting became cheaper than concealment, aviation could finally see what it was fixing.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

## CASE 114

AVIATION SAFETY

STRUCTURAL ENGINEERING

REGULATORY RULEMAKING

1988 – 2010S

# FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule

**G** Governance & Trust   **D** Designed Out   **K** Knowledge & Institutional Memory

**IMPACT** After Aloha Airlines 243 (1988) exposed widespread fatigue cracking in an aging Boeing 737-200, the FAA's Aging Aircraft Program and the Airworthiness Assurance Working Group produced two decades of structural inspection programs culminating in 14 CFR Part 26 subpart C, published November 15, 2010 and effective January 14, 2011, which requires design approval holders to establish a Limit of Validity for each affected model and embed widespread fatigue damage protection into the maintenance program

## IN BRIEF

On April 28, 1988, Aloha Airlines Flight 243 lost an 18-foot section of upper fuselage in flight; the aircraft, a Boeing 737-200, had accumulated 89,680 flight cycles and was operating well past the design assumptions of the original certification. The accident — one flight attendant killed, 65 injured — exposed the structural-engineering category of widespread fatigue damage (WFD): multiple small cracks across many similar structural details, simultaneous enough that established single-crack inspection assumptions did not apply. The FAA stood up the Aging Aircraft Program almost immediately, and the Airworthiness Assurance Working Group (AAWG) operated through the 1990s and 2000s producing structural inspection programs across the transport-airplane fleet. The work culminated in 14 CFR Part 26 (Continued Airworthiness and Safety Improvements; WFD provisions effective 2011), which requires manufacturers to establish a Limit of Validity (LOV) for each type — the operational service goal below which the maintenance program protects against WFD — and to embed WFD prevention into the structural maintenance program itself. Part 26 is one of the most concrete recent examples of legacy assets aging past their original oversight regime being pulled back under structured airworthiness governance, and closes a long-standing zero-state in the induced framework's C7 competency.

## BACKGROUND

Aloha Airlines Flight 243 on April 28, 1988, was the sentinel event. The Boeing 737-200 had accumulated 89,680 flight cycles — well above the design service goal — and was operating in a high-cycle short-haul environment that stressed the fuselage skin and lap joints at a rate the original certification analysis had not anticipated. In cruise at 24,000 feet, an 18-foot section of upper fuselage separated. The crew recovered the aircraft and landed it; one flight attendant was lost overboard and 65 occupants were injured. The NTSB investi-

gation identified multiple small fatigue cracks across many lap-joint rivet holes, linking up catastrophically — the textbook presentation of widespread fatigue damage rather than a single-crack failure.<sup>9</sup>

#### THE INTERVENTION

The structural-engineering category WFD names is specifically the regime where established inspection assumptions break down. A single-crack model assumes one crack initiates, propagates, and is caught by scheduled inspection before reaching critical length. WFD assumes many small cracks initiate at similar structural details across the fleet at similar times, and that link-up between adjacent cracks becomes the dominant failure mode. The inspection cadence, detection threshold, and replacement program a designer would build under the single-crack assumption do not adequately cover the WFD regime, and the airworthiness oversight regime Aloha exposed had not made the distinction.<sup>1</sup>

#### HOW IT WORKED

The FAA's response was to stand up the Aging Aircraft Program and convene the Airworthiness Assurance Working Group, which operated through the 1990s and 2000s. The AAWG produced mandatory structural inspection documents and supplemental structural inspection programs for transport-category aircraft model-by-model, working through the older fleet systematically. The work was technically substantial — fleet-survey data, fatigue-test campaigns, inspection-program revisions for in-service aircraft — and institutionally sustained across more than two decades. The intermediate deliverables were per-model inspection programs; the capstone was rulemaking.<sup>2</sup>

#### THE EVIDENCE

14 CFR Part 26 — Continued Airworthiness and Safety Improvements — was finalized in 2007 with the WFD provisions published November 15, 2010 and effective January 14, 2011. Two structural elements are decisive. First, every transport-category type must have a Limit of Validity established by the manufacturer: the operational service goal, in flight cycles or flight hours, below which the maintenance program is shown to protect against WFD. Operation beyond the LOV requires either the manufacturer establishing an extended LOV with supporting analysis, or operator-specific evidence approved by the regulator. Second, WFD prevention is embedded in the structural maintenance program itself, not handled as a separate one-time inspection event. The maintenance program becomes the carrier of WFD protection across the type's service life.<sup>3</sup>

#### WHAT TRANSFERRED

What the case teaches at the LENS framing layer is the structural form of pulling legacy assets back under structured airworthiness governance. The original 737-200 certification did not anticipate the cycle count and the WFD regime that high-cycle short-haul operation produced; the regulatory response was not to retire the type but to engineer the oversight regime forward — fleet survey, AAWG, model-by-model inspection programs, and finally Part 26 codifying the LOV concept across the transport-category fleet. The case is one of the v2 sweep's clearest closes of the C7 (Capability under system change, transition, and aging assumptions) zero-state in the induced framework: a sustained two-decade regulatory program that turned an aging-fleet structural surprise into a governed sustainment discipline.<sup>4</sup>

*The deliverable was not a single rule. It was the two-decade per-model technical record that made the rule a codification of established practice rather than top-down regulation.*

EDITORS' SYNTHESIS OF THE FAA AGING AIRCRAFT PROGRAM AND THE PART 26 RULEMAKING RECORD.

REFERENCES

1. NTSB (1989), Aircraft Accident Report AAR-89/03, Aloha Airlines Flight 243, Boeing 737-200, N73711.

2. Swift, T. (1993), "Widespread Fatigue Damage Monitoring — Issues and Concerns," International Conference on Structural Airworthiness of New and Aging Aircraft, Hamburg, June 16–18, 1993 — technical synthesis of the WFD inspection regime.

3. Airworthiness Assurance Working Group (1999), Recommendations for Regulatory Action to Prevent Widespread Fatigue Damage in the Commercial Airplane Fleet, Final Report, March 11, 1999 (Rev. A, June 29, 1999) — the ARAC record behind the rule.

4. FAA, "Aging Airplane Program: Widespread Fatigue Damage," Final Rule, 75 FR 69746 (November 15, 2010, effective January 14, 2011) — the subpart C limit-of-validity provisions.

5. FAA, "Enhanced Airworthiness Program for Airplane Systems/ Fuel Tank Safety (EAPAS/FTS)," Final Rule, 72 FR 63409 (November 8, 2007), Docket FAA-2004-18379 — the rule that created 14 CFR Part 26, "Continued Airworthiness and Safety Improvements for Transport Category Airplanes."

THE LEARNING ENGINEERING LENS

LE INSIGHT

The FAA Aging Aircraft Program and Part 26 are one of the v2 sweep's clearest closes of the C7 zero-state. Aloha 243 exposed a regime the original certification did not cover; the AAWG operated for two decades; Part 26 codified Limit of Validity and embedded WFD prevention into the maintenance program itself. The sustained two-decade rulemaking is the deliverable, not the rule alone.

LENS APPROACH

FAA aging aircraft is the canonical aging-asset-governance case (induced 7.3; LENS D1/PT3) — Domain 1 for LOV-as- operational-service-goal; Domain 5 for the AAWG institutional discipline. Closes the C7 zero-state. Pair with Cases 115, 174, 156 as the v2 aging-system quartet.

| DOMAIN                 | MODES ADDRESSED | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO |
|------------------------|-----------------|-----------------------|--------------|---------|-----|
| AVIATION SAFETY        | <div>GDK</div>  | <div>D1D2D3D4D5</div> | 3            | 7.3     | 1-5 |
| STRUCTURAL ENGINEERING | →               | → Systems Analysis    |              |         |     |
| REGULATORY RULEMAKING  |                 |                       |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Name the structural regime the original certification analysis did not cover — WFD vs. single-crack here — explicitly, so the inspection program can be designed around the actual failure mode rather than retrofitted to the original assumption.
- ▶ Operate the working group across a long horizon (AAWG: two decades): the per-model deliverables build the technical record that supports the eventual rulemaking, and short-cycle deliverables alone do not.
- ▶ Make the operational service goal (LOV) the artifact the regulation rests on; the goal is auditable per type, and operation beyond it requires additional analysis rather than implicit extension.

IN OPERATION / AFTER

- ▶ Embed the new protection in the maintenance program rather than as a one-time inspection: the maintenance program is the carrier across the service life, and one-time events are not.
- ▶ Treat the model-by-model technical work as the legitimacy basis for the rule; without the AAWG's fleet-survey and inspection-program record, Part 26 would have looked like top-down regulation rather than codification of established practice.
- ▶ Carry the case as the C7 instance the induced framework needed: legacy assets aging past their original oversight regime, pulled back under structured airworthiness governance by sustained two-decade rulemaking.

## REFLECTION QUESTIONS

1. Identify a legacy asset class in your domain whose original certification or design analysis no longer covers the regime the assets are now operating in. What is the analog of the WFD distinction — the failure mode the original analysis did not anticipate?
2. Specify the analog of the LOV: an operational service goal, auditable per asset, beyond which additional analysis is required. What instrument would you use, what cadence, and what would “operation beyond LOV” trigger in your domain?
3. The AAWG operated for two decades before Part 26 codified the practice. What is the institutional discipline you would build for a similar two-decade horizon — the working-group cadence, the per-asset technical record, the eventual rulemaking — and is there a body in your domain that could plausibly carry that horizon?

**WHO BUILDS THIS** governance, ethics & measurement · systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · task analysis · design prototyping · knowledge bases · data pipelines

# FAA NextGen and the ADS-B Out Transition

**G** Governance & Trust   **D** Designed Out   **K** Knowledge & Institutional Memory

**IMPACT** The FAA's Next Generation Air Transportation System (NextGen) shifted U.S. air-traffic management from radar-based surveillance to a satellite-based architecture; the ADS-B Out mandate effective January 1, 2020 required equipage across the controlled-airspace fleet, achieving substantial compliance — with documented schedule slippage and benefit-realization gaps preserved as decisive hedges

## IN BRIEF

The FAA's Next Generation Air Transportation System (NextGen) is the multi-decade transition of U.S. air-traffic management from a radar-based surveillance architecture to a satellite-based architecture built on Automatic Dependent Surveillance – Broadcast (ADS-B). The ADS-B Out final rule, published in 2010, set the equipage mandate for January 1, 2020 — aircraft operating in most controlled airspace must broadcast their GPS-derived position and identity, replacing the radar-only secondary-surveillance model that defined the era prior. The mandate was substantially met at the deadline; equipage compliance across the affected fleet was high, and ADS-B-based surveillance is now the operational backbone in much of U.S. airspace. The case is one of the largest planned aging- infrastructure transitions in the recent regulatory record — closing a long-standing C7 zero-state in the induced framework — and it carries the load-bearing hedges that GAO and DOT Inspector General reviews have repeatedly documented: significant schedule slippage across the program, benefit-realization gaps relative to original projections, and contested cost-benefit accounting. The transition happened; the transition was harder, slower, and more partial than the original NextGen plan implied.

## BACKGROUND

The pre-NextGen U.S. air-traffic management architecture rested on ground-based primary and secondary radar surveillance, with voice communication, paper or quasi-paper flight progress strips at many facilities, and traffic-flow management built around scheduled equipage upgrades that lagged the broader civil-aviation hardware base. By the early 2000s the structural problem was named clearly in the policy literature: a 20th-century surveillance architecture was being asked to handle 21st-century traffic, and incremental modernization inside the radar-paradigm was approaching its limits.<sup>o</sup>

## THE INTERVENTION

NextGen, formally launched in 2003 under the Vision 100 – Century of Aviation Reauthorization Act, was the FAA's multi-decade response. The technical core was the transition from radar-based surveillance to ADS-B: aircraft broadcast their GPS-derived position once

per second, ground stations and other aircraft receive the broadcast, and the resulting surveillance picture is more precise, lower-latency, and less expensive to operate at scale than secondary radar. The broader NextGen program included additional elements (Data Communications, NAS Voice System, System-Wide Information Management) but ADS-B Out was the decisive equipage-mandate piece.<sup>1</sup>

#### HOW IT WORKED

The ADS-B Out final rule was published in 2010 with the January 1, 2020 compliance deadline — a decade of lead time for operators to equip. At the deadline, substantial compliance was reported across the affected fleet; the FAA’s surveillance picture in controlled airspace moved substantially onto the ADS-B architecture, and the operational transition was completed in the sense the rule intended. The transition is one of the largest infrastructure-replacement programs in the recent FAA record, and it executed.<sup>2</sup>

#### THE EVIDENCE

The hedges have to survive. GAO and DOT Inspector General reports across the 2010s repeatedly documented significant schedule slippage across NextGen — multiple elements missing original delivery dates, the broader program’s realized benefits running below original projections, and contested cost-benefit accounting between FAA program reporting and external review. The benefit categories NextGen’s business case rested on — fuel-burn reduction, delay reduction, increased airspace capacity — have materialized in some respects and not in others, and the attribution to NextGen specifically (vs. broader operational and weather variation) is contested in the published audit literature. Equipage happened; benefit realization is mixed and ongoing.<sup>3</sup>

#### WHAT TRANSFERRED

What the case teaches at the LENS framing layer is the sustainment-engineering form of a planned aging- infrastructure transition at continental scale. The transition happened: the regulatory architecture, the decade of lead time, the equipage mandate at the deadline, and the operational migration to ADS-B-based surveillance are real and large. The transition was harder, slower, and more partial than the original NextGen plan implied; the schedule slippage and benefit-realization gaps are not anomalies in the case but part of what infrastructure transitions at this scale look like in practice. Together with the FAA Aging Aircraft program (Case 114), Y2K (Case 174), and LWRS (Case 156), NextGen is part of the v2 quartet that closes the C7 zero-state in the induced framework — and is the instance where the hedges are largest.<sup>4</sup>

*The transition happened. The transition was harder, slower, and more partial than the original plan implied. Both are part of the case.*

EDITORS’ SYNTHESIS OF THE GAO NEXTGEN AUDIT SERIES AND FAA PROGRAM REPORTING.

#### REFERENCES

1. FAA, “Automatic Dependent Surveillance-Broadcast (ADS-B) Out Performance Requirements To Support Air Traffic Control (ATC) Service,” Final Rule, 75 FR 30160 (May 28, 2010) — 14 CFR §§91.225 and 91.227.
2. Vision 100 — Century of Aviation Reauthorization Act, Public Law 108-176 (2003) — NextGen program statutory basis.
3. Department of Transportation Office of Inspector General (2021), NextGen Benefits Have Not Kept Pace With Initial Projections, Report No. AV2021023, March 30, 2021 — FAA’s 2017 benefits projection \$113 billion below the JPDO estimate; \$6 billion in benefits realized 2010–2018.
4. FAA, NextGen Annual Reports (2010 – present) — program-self-report tier; read against the GAO and DOT IG reviews.

3. Government Accountability Office, “NextGen Air Transportation System” report series (2010s) — sustained external audit record on schedule slippage and benefit–realization gaps.

THE LEARNING ENGINEERING LENS

LE INSIGHT

NextGen / ADS-B is one of the largest planned aging– infrastructure transitions in the recent regulatory record. The equipage mandate executed at the January 2020 deadline. The broader NextGen program has documented significant schedule slippage and benefit–realization gaps in sustained external audit. The case closes a C7 zero–state with the largest acknowledged hedges in the v2 aging–system quartet.

LENS APPROACH

NextGen / ADS-B is the planned infrastructure–transition case (induced 7.1; LENS D1/PT4). LENS uses it in Domain 1 (Systems Analysis) for the separation of the auditable equipage mandate from the multifactorial benefit– realization business case, and in Domain 5 (Navigating Sociotechnical Constraints) for the decade of lead time and the sustained external–audit discipline. Closes the C7 zero–state alongside Cases 114 (FAA aging aircraft), 174 (Y2K), and 156 (LWRS) as the v2 aging–system quartet.

| DOMAIN                     | MODES ADDRESSED | LENS COMPETENCY    | PROBLEM TYPE | INDUCED | LEO   |
|----------------------------|-----------------|--------------------|--------------|---------|-------|
| AVIATION<br>INFRASTRUCTURE | G D K           | D1 D2 D3 D4 D5     | 4            | 7.1     | 1 · 5 |
| AIR TRAFFIC<br>MANAGEMENT  | →               | → Systems Analysis |              |         |       |
| REGULATORY<br>TRANSITION   |                 |                    |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the equipage mandate with the lead time the transition actually needs — a decade for ADS-B Out — so the deadline arrives with realistic compliance pathways rather than as a forcing function operators cannot meet.
- Separate the equipage–mandate piece (auditable, has a deadline) from the broader benefit–realization business case (multifactorial, harder to attribute); the equipage piece will deliver, the benefit piece will deliver partially, and conflating them sets the program up to look like a failure where it succeeded.
- Plan for sustained external audit (GAO, DOT IG) as part of the program’s operating environment — the schedule slippage and benefit–realization gap reporting is institutional discipline, not a sign the program is unworkable.

IN OPERATION / AFTER

- Report the equipage transition as the deliverable that executed and the benefit realization as the deliverable that is ongoing and mixed; both are real, and the hedge that matters is that the original NextGen plan understated the difficulty of the latter.
- Sustain the audit–and–rebaseline cadence after the headline equipage mandate; the benefit–realization picture continues to develop, and the institutional record of slippage and partial realization is part of how transitions at this scale are documented honestly.
- Treat the NextGen case as the instance where the C7 transition delivered with the largest acknowledged hedge — useful precisely because the transition is real and the hedges are also real, and conflating either with the other misreads the lesson.



## REFLECTION QUESTIONS

1. Identify an aging-infrastructure transition in your domain whose original business case rested on benefit categories that are multifactorial and contested in attribution. What is the equivalent of the auditable equipage mandate — the piece of the transition that delivers cleanly — and how would you separate its reporting from the broader benefit case?
2. Specify the lead time the transition actually needs from mandate to deadline; ADS-B Out used a decade. What is the analog in your context, and what does the operator-equipage trajectory look like across that horizon?

**WHO BUILDS THIS** governance, ethics & measurement · systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · task analysis · design prototyping · knowledge bases · data pipelines

# Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change

**G** Governance & Trust   **K** Knowledge & Institutional Memory   **H** Human-System Interface

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** A 17,000-line pilot modernization of the Eurocat air-traffic-management system, run in 2005 to assess transformation validity and settle the target language, opened a program whose safety evidence — that the automated transformation was non-distortive of original functionality — convinced Thales customers to accept a system-wide architecture-driven modernization with 100% automated code transformation

## IN BRIEF

The Eurocat Air Traffic Management System was the kind of safety-critical legacy software whose customers cannot accept a big-bang rewrite: the operational system in production cannot tolerate the discontinuity. The 2005 Thales-led pilot modernization was scoped narrowly — 17,000 lines of code — with the deliberate goal of generating safety-equivalence evidence that an automated transformation did not distort the original functionality. The pilot succeeded in producing that evidence, and the evidence was used to convince customers to accept a full architecture-driven modernization with 100% automated code transformation. The teaching pattern is the small-as-gateway-to-big move: the small-tier verification artifact dissolved the customer objection to the large-tier change. The case is documented in a vendor-authored Elsevier technical chapter rather than an independent academic evaluation; the evidence-tier flag is rendered under the title. The case is the small-tier C7 transition-verification companion to the big-tier failures (Patriot/Dhahran, Ariane 5, Knight Capital, CrowdStrike, TSB) already in the corpus. Future validation will continue on the long-run operational behavior of the modernized system.

## BACKGROUND

Air traffic management software is a paradigmatic capability-system-misalignment-at-transition problem: the legacy system carries decades of accreted operational logic, the safety case the regulator and the airspace user expect rests on the legacy's documented behavior, and a big-bang rewrite is unacceptable because the operational discontinuity is

itself the safety hazard. The Eurocat problem in 2005 was the standard one — modernize without breaking the safety case the legacy already carries.<sup>0</sup>

#### THE INTERVENTION

The design move was to scope a pilot narrowly enough that the verification could be exhaustive. Seventeen thousand lines of code is small as a fraction of the full system; it is large enough that a working automated transformation pipeline can be exercised end-to-end and the safety-equivalence evidence can be generated against a realistic subset. The chapter gives the pilot two jobs — assess the validity of the transformation, and settle the target language. It was a deliberately designed evidence artifact: prove, at small scale and on the actual legacy, that the transformation does not distort the function the customer cares about.<sup>1</sup>

#### HOW IT WORKED

The pilot succeeded in producing the evidence. Customers were then willing to accept a system-wide architecture-driven modernization with 100% automated code transformation — as the chapter puts it, fully automated at the code layer and accompanied by iterative semi-automated reengineering of the design and architecture. The modernization that followed ran from 2005 through 2008 across three Ada 83 variants of the Eurocat product line, roughly 1.6 million lines in total; first site acceptance of the Java flight-data-processing subsystem was approved in March 2010. The causal claim the case rests on is not “the modernization worked” — that is a long-run operational question — but “the customer objection to the large-tier change was dissolved by the small-tier verification artifact.” That is the C7 teaching point at the small tier the corpus has not had: the same governance-objection-dissolver move that Case 199 / Case 200 show in the AV regulatory regime, transposed into a legacy-software transition.<sup>2</sup>

#### THE EVIDENCE

The evidence-tier flag matters. The case is documented in a practitioner-authored chapter in an Elsevier book on information-systems transformation. The authors are from the vendor side; there is no independent academic evaluation of the pilot’s safety-equivalence claims or of the post-modernization operational behavior. The pattern the case teaches — verification at the small tier as the gateway to the large-tier change — is robust across the practitioner literature on legacy modernization, but the Eurocat-specific magnitudes and the long-run system outcomes have not been audited in a peer-reviewed source. Future validation will continue.<sup>3</sup>

#### WHAT TRANSFERRED

What the case adds at the LENS layer is the small-tier C7 intervention that the big-tier failures already in the corpus do not teach: when the inherited design has to move to a new envelope, the verification work itself can be the capability deliverable that earns the change its legitimacy. The Domain-1 frame applies: systems analysis of the legacy’s safety case as an artifact, and design of the transformation pipeline as another artifact, are the capability-engineering moves the program made. The case is the missing positive example for induced 7.1 and 7.2 at the small tier.

*The verification at the small tier is the artifact that earns the large-tier change its legitimacy. The pilot is not a feasibility check. It is the evidence.*

EDITORS' SYNTHESIS OF THE EUROCAT MODERNIZATION CHAPTER.

REFERENCES

1.

DelaPeyronnie, Newcomb, Morillo, Trimech, Nguyen, & Purtill (2010), "Modernization of the Eurocat Air Traffic Management System (EATMS)," in Information Systems Transformation: Architecture-Driven Modernization Case Studies (Elsevier / Morgan Kaufmann), Chapter 5.

2.

Ulrich & Newcomb (eds., 2010), Information Systems Transformation — the host volume on architecture-driven modernization patterns.

3.

Brodie & Stonebraker (1995), Migrating Legacy Systems — the framing reference on small-step legacy modernization.

4.

Bisbal et al. (1999), "Legacy Information Systems: Issues and Directions," IEEE Software 16(5):103–111 — peer-reviewed framing companion.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Eurocat is the small-tier C7 transition-verification intervention the corpus needed: a narrowly scoped pilot designed to produce the safety-equivalence evidence that dissolves customer objection to a system-wide modernization. Evidence is vendor-authored practitioner literature; the long-run operational record of the modernized system is the open question. Future validation ongoing.

LENS APPROACH

Eurocat is the C7 small-tier transition case (induced 7.1 and 7.2; LENS D1/PT1). LENS uses it in Domain 1 (Systems Analysis) for LEO-1 — the safety case and the transformation pipeline are both engineered artifacts — and in Domain 2 (Iterative Development) for LEO-2, since the pilot is the iteration that earns the rollout. Pair with the big-tier C7 failures (Patriot/Dhahran, Ariane 5, Knight Capital, CrowdStrike, TSB) for the failure-and- intervention contrast.

| DOMAIN                  | MODES ADDRESSED | LENS COMPETENCY    | PROBLEM TYPE | INDUCED | LEO   |
|-------------------------|-----------------|--------------------|--------------|---------|-------|
| AIR TRAFFIC MANAGEMENT  | G K H           | D1 D2 D3 D4 D5     | 1            | 7.1     | 1 · 2 |
| LEGACY MODERNIZATION    | →               | → Systems Analysis |              |         |       |
| SOFTWARE TRANSFORMATION |                 |                    |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Scope the pilot to be small enough that the verification can be exhaustive against the actual legacy, not just a feasibility check on a synthetic subset.
- ▶ Design the safety-equivalence evidence as a deliberate deliverable from the pilot, not a byproduct: specify in advance what the customer needs to see to accept the large-tier change.
- ▶ Document the transformation pipeline itself as a verifiable artifact, so the customer can audit the transformation, not just inspect the transformed code.

IN OPERATION / AFTER

- ▶ Carry the practitioner-tier flag into any downstream citation; the case is a vendor account and the long-run operational behavior of the modernized system is the open question.
- ▶ Treat the small-as-gateway-to-big pattern as the teaching point; the magnitudes and the system-wide rollout are open and depend on the long-run operational record.
- ▶ Pair with the big-tier C7 failures (Patriot/Dhahran, Ariane 5, Knight Capital, TSB) when teaching; the pair shows the verification-as-deliverable principle is what separates the failure and intervention threads.

## REFLECTION QUESTIONS

1. Identify a legacy system in your context whose customer or operator will not accept a big-bang rewrite. What is the small-tier pilot whose evidence would dissolve the objection to the large-tier change, and what specifically would the customer need to see?
2. Specify the safety-equivalence evidence the pilot would have to generate as a deliberate deliverable. The Eurocat pattern teaches because the pilot was scoped to produce the artifact the customer needed, not just to demonstrate feasibility.
3. The case is vendor-authored practitioner literature. What independent evidence (academic evaluation, multi-customer replication, long-run operational behavior) would you require before treating the small-as-gateway-to-big pattern as a settled engineering practice in your domain?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · usability testing · cognitive walkthroughs

# Crew Resource Management & CAST

**T** Training Gap   **H** Human-System Interface   **N** Normalization of Deviance

**IMPACT** CAST adopted the White House Commission's 80% fatality-risk reduction target and reported 83% for 1998–2008 — a portfolio result across 70 selected safety enhancements, credited by CAST and the FAA to CAST's work "along with new aircraft, regulations and other activities"; CRM's separate contribution is not isolated in any published estimate, and CRM predates the measurement window by seventeen years

## IN BRIEF

In March 1977, two 747s collided in fog at Tenerife and 583 people died — in part because a KLM flight engineer's twice-voiced doubt about the runway was overridden by his captain. The failure was not of skill but of the system by which skill in one seat reached another. Crew Resource Management, formalized by United Airlines in 1981, re-engineered the cockpit as a coordinated team: explicit communication protocols, named authority gradients, structured briefings. Sixteen years later the Commercial Aviation Safety Team (CAST) added closed-loop analysis of operational data to find and fix hazards before they caused accidents. The pairing — cultural redesign plus continuous evidence — helped drive the 83% reduction in U.S. commercial-aviation fatality risk between 1998 and 2008 that CAST achieved alongside new aircraft and regulation, work that earned CAST the 2008 Collier Trophy.

## BACKGROUND

By the 1970s, accident investigations were repeatedly finding that crashes stemmed not from a lack of individual flying skill but from breakdowns in how a crew worked together. The 1977 Tenerife disaster — 583 dead — was the starkest example: a KLM flight engineer questioned, twice and indirectly, whether the runway was clear, and the senior captain dismissed him and continued the takeoff. The pattern that emerged was consistent across investigations — skilled crews failing not because anyone lacked competence, but because the crew's most senior voice could close off information held by a more junior one before it reached the decision.<sup>9</sup>

## THE INTERVENTION

Crew Resource Management, formalized by United Airlines in 1981, treated crew coordination as an engineerable property of the system rather than a matter of personality. It introduced explicit communication protocols, named and trained against authority gradients, and instituted structured briefings and debriefings — rebuilding the cockpit as a team in which information from any seat could reach the decision. Rather than exhorting captains to listen better, it built coordination into the standard procedure itself, so the behavior that Tenerife had lacked became the trained default rather than a matter of individual temperament.<sup>1</sup>

## HOW IT WORKED


CRM did not teach individual airmanship; it engineered the system of interaction in which airmanship operates. By making it legitimate and expected for a junior officer to challenge a captain, and by giving crews a shared protocol for doing so, it closed the path by which Tenerife-style deference had absorbed safety-critical information instead of transmitting it. The change was structural rather than attitudinal: a challenge that had once depended on a junior officer's nerve now had a named procedure behind it, so the same doubt that went unheard at Tenerife had an authorized route to the decision.<sup>2</sup>

## THE EVIDENCE

Sixteen years on, the Commercial Aviation Safety Team added the missing second half: closed-loop hazard identification on operational data, prioritized enhancements, tracked implementation, and measured outcomes. CAST set an 80% fatality-reduction target and exceeded it, reaching 83% across 1998–2008 — work recognized with the 2008 Collier Trophy. Over twenty years, fatalities per 100 million passengers fell roughly 95%. The loop closed on itself: data surfaced the next hazard, enhancements were prioritized against it, and the measured outcome fed back into the priorities, so improvement continued rather than plateauing once the cultural change had taken hold.<sup>3</sup>

## WHAT TRANSFERRED

CRM and CAST together define what a mature capability-engineering apparatus looks like: a cultural redesign, and sixteen years later a continuous-evidence loop to steer it. What the record does not show is what either contributed on its own. CRM ran without CAST from 1981; the 83 percent figure covers 1998–2008; and CAST and the FAA credit it to a portfolio of seventy selected enhancements alongside new aircraft and regulation. The model has been exported to surgery, firefighting, and other high-consequence domains, and is now the template for redesigning human roles in AI-augmented systems. What transferred was not the specific protocols but the design logic itself — that coordination is an engineerable system property and that the cultural change needs a measurement loop to keep it honest over time.<sup>4</sup> The 83% is a portfolio figure, and its CFIT share belongs to the safety enhancements treated in Case 120 (the TAWS mandate) and Case 123 (the FSF task forces); CRM's own share is not separately measured in that data.



*CRM succeeded because it treated crew coordination as an engineerable property of the system.*

PARAPHRASING HELMREICH & FOUSHEE, IN KANKI ET AL., CREW RESOURCE MANAGEMENT (2010)

REFERENCES

1. Spanish CIAIAC / ALPA reports on the 1977 Tenerife collision — the overridden crew challenge.

2. Helmreich, Merritt & Wilhelm (1999), “The Evolution of Crew Resource Management Training in Commercial Aviation,” International Journal of Aviation Psychology.

3. FAA Advisory Circular 120-51E, Crew Resource Management Training — CRM protocols and authority-gradient training.

4. CAST/FAA Safety Enhancement reports (2016, 2018) — the closed-loop data process and the 83% reduction.

5. Collier Trophy citation (2008); Kanki, Helmreich & Anca (2010), Crew Resource Management.

THE LEARNING ENGINEERING LENS

LE INSIGHT

CRM is the canonical evidence that capability is engineerable at the system level, not just the individual. Tenerife was not solvable by hiring better pilots — only by changing the authority structure inside the cockpit. CRM paired a procedural change with a cultural change in how the procedure was authorized, and CAST added the evidence loop sixteen years later. The pairing is the design lesson; which part carried the fatality-rate fall is not something a portfolio result can separate.

LENS APPROACH

LENS treats CRM/CAST as the anchor success case across the curriculum. LEN 1 uses it as the problem-framing exemplar; LEN 4 uses CAST as the model closed-loop evidence system; LEN 2 uses CRM as the template for redesigning human roles in automated environments — the logic now being applied to AI-augmented systems.

| DOMAIN                     | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|-----------------|------------------|--------------|---------|-----|
| AVIATION                   | → T H N         | → D1 D2 D3 D4 D5 | 3            | 4.3     | 3   |
| Human-System Collaboration |                 |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Build the coordination behavior into standard procedure — explicit communication protocols, named authority gradients, and structured briefings — so it is the trained default, not a matter of temperament.
- ▶ Authorize the junior-to-senior challenge explicitly and drill it, so safety-critical information has a procedural route to the decision rather than depending on individual nerve.
- ▶ Stand up the evidence half from the start: instrument operational data so hazards can be found and prioritized before they cause an accident.

IN OPERATION / AFTER

- ▶ Run the closed loop continuously — surface hazards from data, prioritize enhancements, track implementation, and measure outcomes — so improvement does not plateau once the culture shifts.
- ▶ Set and exceed an explicit reduction target (the CAST model) so the intervention is judged against a measured outcome, not against its own activity.
- ▶ Export the design logic, not just the protocols, when transferring to new domains, adapting the paired cultural-plus-evidence structure to each setting.



### REFLECTION QUESTIONS

1. Identify a high-consequence domain whose authority gradient absorbs information rather than transmitting it. What would the CRM equivalent intervention look like there?
2. CRM is paired with CAST. What is the closed-loop evidence system in your domain — and if there is not one, what would it cost to build?
3. CRM made a junior officer's challenge a named procedure rather than an act of nerve. Identify a moment in your domain where the right information depends on someone's courage, and design the protocol that would make raising it the expected default instead.

**WHO BUILDS THIS** learning science & instructional design · human factors & interaction design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs · incident analysis · safety dashboards

# Korean Air Safety Transformation

**T** Training Gap    **N** Normalization of Deviance

**IMPACT** From industry pariah (a 1988–1998 loss rate of 4.79 per million departures against United's 0.27, and 228 dead at Guam in 1997) to spotless passenger safety record since 1999

## IN BRIEF

Between 1988 and 1998, Korean Air had one of the worst safety records in commercial aviation — a loss rate of 4.79 hull losses per million departures against United's 0.27 — and the 1997 Guam crash of Flight 801 killed 228. The NTSB's probable cause was the captain's failure to adequately brief and execute the nonprecision approach and the first officer's and flight engineer's failure to monitor and cross-check it, with his fatigue and Korean Air's inadequate training as contributing factors — and, contributing to the accident, the FAA's intentional inhibition of the minimum safe altitude warning system at Guam. The first officer did call for a missed approach six seconds before impact; on why no challenge came earlier in the approach, the Board wrote that it was unable to identify the reason. The cultural reading — that a steep cockpit authority gradient silenced the right seat — came from the industry and from the airline itself, and it is what Korean Air then acted on. In 2000, Korean Air brought in Delta's David Greenberg to rebuild flight operations — mandating English as the cockpit language, adapting CRM for a high-power-distance culture, and bringing in outside consulting from Boeing and Delta. The airline has had no fatal passenger accident since, winning the 2006 Phoenix Award. Cultural legacy is not destiny when it is deliberately redesigned.

## BACKGROUND

From 1988 to 1998, Korean Air suffered repeated fatal crashes, at a loss rate of 4.79 per million departures against United Airlines' 0.27. The NTSB's report on the 1997 Guam crash of Flight 801, which killed 228, found the captain had failed to adequately brief and execute the nonprecision approach and the first officer and flight engineer had failed to effectively monitor and cross-check it, contributing his fatigue and the airline's inadequate training; a further contributing cause sat outside the airline entirely — the FAA's intentional inhibition of the minimum safe altitude warning system at Guam and its failure to manage that system. It did not reach for culture: the words power distance, hierarchy and authority gradient appear nowhere in it. On why the right seat raised no challenge earlier in the approach — the first officer did call for a missed approach, but only six seconds before impact — the Board said it was unable to identify whether the silence came from inadequate briefing and monitoring, from failing to see the captain's errors, or from unwillingness to confront errors they did perceive. That gap is where the cultural reading entered — from the industry, from the press, and from the airline, which acted on it.<sup>o</sup>

## THE INTERVENTION

After foreign carriers suspended code-shares over its safety record, Korean Air in 2000 hired David Greenberg from Delta to overhaul flight operations; the FAA's August 2001 downgrade of South Korea to Category 2 came a year later, and the rating was restored that December. The interventions were deliberate and cultural: mandated English fluency for all pilots, CRM training adapted for a high-power-distance setting, external consulting from Boeing and Delta, and fleet modernization. The external commercial pressure — suspended code-shares and a downgraded rating — supplied the leverage to act, turning a cultural problem the airline had long lived with into one it could no longer afford to ignore.<sup>1</sup>

## HOW IT WORKED


The decisive move operated on language. Making English the cockpit language stripped out the Korean honorific hierarchy that had silenced first officers, because English has no honorifics to enforce rank. The CRM adaptation then gave crews an explicit, culturally workable protocol for raising concerns — converting deference into communication. The choice was elegant precisely because it changed the medium rather than asking pilots to override their own culture: with no honorifics to encode, the language itself flattened the gradient that had absorbed the first officer's warnings.<sup>2</sup>

## THE EVIDENCE

Korean Air has recorded no fatal passenger accident since the reforms took hold. Air Transport World recognized the turnaround with its 2006 Phoenix Award, and the carrier moved from international pariah to a safety record indistinguishable from the best global operators. The categorical nature of the shift — from a loss rate many times the industry's to none — is what makes the case persuasive: the same crews and the same national culture produced opposite outcomes once the cockpit's communication structure was redesigned.<sup>3</sup>

## WHAT TRANSFERRED

Korean Air is the strongest aviation evidence that cultural legacy is not destiny: a specific cultural feature — high power distance in the cockpit — was the binding constraint, and once it was redesigned, the safety record changed categorically. The case also shows that the interface for cultural redesign can itself be engineered, in this instance through language. The transferable lesson is to locate the single cultural feature that actually binds the capability rather than attempting to remake a culture wholesale, then change the medium through which that feature operates.<sup>4</sup>



THE AIRLINE'S OWN DIAGNOSIS: FLATTEN THE COCKPIT GRADIENT

*Korean Air’s record before 2000 was changed by an external intervention into cockpit culture, not by criticism of it.*

EDITORS’ SYNTHESIS OF NTSB FINDINGS ON KAL 801 AND THE KOREAN AIR TRANSFORMATION

REFERENCES

1. NTSB (2000), AAR-00/01, Controlled Flight Into Terrain, Korean Air Flight 801, Boeing 747-300, HL7468, Nimitz Hill, Guam, August 6, 1997 — probable cause at §3.2: the captain’s failure to adequately brief and execute the nonprecision approach and the first officer’s and flight engineer’s failure to effectively monitor and cross-check it; contributing to those failures, the captain’s fatigue and Korean Air’s inadequate flight crew training; contributing to the accident, the FAA’s intentional inhibition of the MSAW system at Guam and its failure to adequately manage it.

2. Helmreich, Wilhelm, Klinec & Merritt (2001), “Culture, Error, and Crew Resource Management,” in Salas, Bowers & Edens (eds.), Improving Teamwork in Organizations (Erlbaum) — national culture and CRM adaptation.

3. Gladwell, M. (2008), Outliers — the Korean Air chapter on power distance and the English-language change.

4. Air Transport World Phoenix Award documentation (2006) — recognition of the turnaround.

5. Korean Air corporate safety reports — the post-2000 accident-free record.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Korean Air is the strongest aviation evidence that cultural legacy is not destiny. A specific cultural feature — high power distance in the cockpit — was the binding capability constraint. Once it was redesigned, the safety record changed categorically. The intervention is also methodologically interesting because it operated on language: English became the cockpit language because English has no honorifics to enforce. The interface for cultural redesign was a linguistic one.

LENS APPROACH

LENS uses Korean Air in LEN 8 as the canonical organizational- learning case under cultural constraint and in LEN 2 as a CRM- extension case for high-power-distance contexts. The case is paired in this book with the Toyota Andon Cord (Case 155) as cultural intervention success stories.

| DOMAIN                     | MODES ADDRESSED             | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|-----------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| AVIATION                   | → <div>T</div> <div>N</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 3            | 4.3     | 3   |
| Human-System Collaboration |                             |                                                                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Locate the single cultural feature that actually binds the capability — here, high power distance expressed through honorifics — rather than attempting to remake the whole culture.
- ▶ Change the medium through which that feature operates (making English the cockpit language) so the constraint is removed structurally rather than by asking people to override their own culture.
- ▶ Pair the structural change with an adapted protocol — CRM tuned for a high-power-distance setting — that gives crews a culturally workable way to raise concerns.

IN OPERATION / AFTER

- ▶ Use external pressure (suspended code-shares, a downgraded rating) as durable leverage to hold the reform in place against any drift back to the prior norm.
- ▶ Sustain the change with continued outside benchmarking and consulting so the new communication structure is reinforced rather than quietly eroding.
- ▶ Track the safety record over years to confirm the categorical shift holds, treating a maintained accident-free record as the evidence the redesign took.

## REFLECTION QUESTIONS

1. Identify a cultural feature in an institution you work with that constrains capability. Is it engineerable? What would the redesign look like?
2. Korean Air's intervention operated on language. What surface of culture is engineerable in your domain that you have not yet considered?
3. Korean Air acted only once suspended code-shares and a downgraded rating made the status quo unaffordable. What external pressure exists in your domain that could supply the leverage to redesign a long-tolerated cultural constraint — and how would you use it?

**WHO BUILDS THIS** learning science & instructional design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards

# Aviation Safety Reporting System (ASRS)

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** NASA-administered confidential reporting system; more than 2M reports received; the architecture other evidence-driven safety systems copied

## IN BRIEF

The Aviation Safety Reporting System, run by NASA on behalf of the FAA since 1976, accepts confidential reports from pilots, controllers, mechanics, and cabin crew about incidents, near-misses, and safety concerns. Its decisive feature is institutional, not technical: reporting an event to ASRS confers immunity from FAA enforcement for the conduct reported, within specified limits, making honest reporting the rational choice. Over nearly fifty years and more than two million reports, ASRS has become one of the world's largest sources of aviation-safety and human-factors information, surfacing patterns — automation surprise, runway incursions, fatigue effects — before they reached formal investigation thresholds. The architecture has been emulated across domains, and is the canonical success case for an evidence system paired with a credible commitment to non-punitive use.

## BACKGROUND

The most valuable safety information in any high-consequence domain lives with front-line operators — the near-misses and quiet errors that never reach an accident report. But operators will not surrender that information to an institution that can punish them for it, so the data that could prevent the next accident stays locked in the people who hold it. The incentives run backward: the person best placed to report a near-miss is the same person a punitive system gives the strongest reason to stay silent, so the data that matters most is the data least likely to surface.<sup>o</sup>

## THE INTERVENTION

In 1976 the FAA and NASA created the Aviation Safety Reporting System, a confidential channel for pilots, controllers, mechanics, and cabin crew to report incidents, near-misses, and concerns. NASA — not the regulator — administers it, and reporting an event confers immunity from FAA enforcement for the conduct reported, subject to specified limits. Putting a neutral party between the reporter and the enforcer directly addressed the backward incentive — an operator now had a positive reason to report, because doing so converted potential jeopardy into protection.<sup>1</sup>

## HOW IT WORKED

The system pairs a technical artifact (the reporting form and a searchable database) with a cultural commitment (protected, non-punitive use). The immunity provision makes reporting the rational choice for the operator, and NASA's role as a neutral third party makes the protection credible. The two were designed together and have never been separated: the

FAA established the program itself in April 1975 and within four months signed the agreement handing collection to NASA, having judged that its own enforcement role would suppress the reporting the program depended on. That is design reasoning on the record rather than a test of either half alone. The credibility of the protection is what does the work — a promise of non-punishment from the regulator itself would be doubted, so routing it through NASA is what makes operators trust it enough to report.<sup>2</sup>

#### THE EVIDENCE

Over nearly fifty years ASRS has accumulated more than two million reports — the largest single repository of aviation-safety information in the world. Patterns such as automation surprise, runway incursions, and fatigue effects were first identified at scale through ASRS data before they crossed formal investigation thresholds. Surfacing a pattern before it reaches an accident is the whole point — the value of the system is the hazards it lets the industry act on while they are still near-misses, not the reports it collects after the fact.<sup>3</sup>

#### WHAT TRANSFERRED

ASRS has been studied and emulated across domains — patient-safety reporting systems, the maritime and aviation CHIRP scheme, and similar systems in rail and nuclear power. It is the canonical positive case for evidence architecture paired with an institutional commitment to non-punitive learning, the defining design pattern of a “just culture.” The breadth of emulation shows that the load-bearing element travels — wherever the most valuable safety data sits with operators who fear punishment, the same protected-reporting design recurs as the way to unlock it.<sup>4</sup>



*ASRS is the model for confidential, voluntary, non-punitive incident reporting in any high-consequence domain.*

PARAPHRASING JAMES REASON, MANAGING THE RISKS OF ORGANIZATIONAL ACCIDENTS, 1997

#### REFERENCES

1. Reason, J. (1997), Managing the Risks of Organizational Accidents — non-punitive reporting as a model (paraphrased).
2. NASA ASRS Program documentation and annual reports — system design, immunity provision, and report volume.
3. FAA Advisory Circular OO-46F, “Aviation Safety Reporting Program” — the program’s establishment in April 1975, the FAA–NASA Memorandum of Agreement of 15 August 1975, and the immunity provisions.
4. NASA ASRS technical reports (Connell et al.) — patterns first surfaced via ASRS data.
5. Dekker, S. (2012), Just Culture — the cultural commitment to non-punitive use.
6. CHIRP and patient-safety reporting-system documentation — cross-domain emulation.

THE LEARNING ENGINEERING LENS

LE INSIGHT

ASRS is the canonical positive case for paired-intervention evidence architecture. The technical artifact (the reporting form and the database) shipped together with the institutional commitment to non-punitive use and a neutral holder, and has never run without it — so the case teaches how the pair was designed, not what either half would do alone. Together they have produced the operational-safety dataset other domains have tried to copy.

LENS APPROACH

LENS uses ASRS in LEN 4 as the primary positive case for evidence architecture and in LEN 8 for institutional commitment to non-punitive learning. Studio projects design ASRS-equivalents for new domains.

| DOMAIN              | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------|-----------------|------------------|--------------|---------|-----|
| AVIATION            | → T K N         | → D1 D2 D3 D4 D5 | 2            | 4.2     | 4   |
| Test and Evaluation |                 |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Pair a simple reporting artifact (a form and searchable database) with a credible commitment to non-punitive use, since the channel without the protection collects nothing of value.
- Confer immunity for reported conduct so that reporting becomes the rational choice, directly reversing the incentive that otherwise keeps the most valuable data hidden.
- Route the protection through a neutral third party rather than the enforcer, so operators trust the non-punitive promise enough to report against their own interest.

IN OPERATION / AFTER

- Analyze the accumulated reports to surface patterns — automation surprise, runway incursions, fatigue — and act on them while they are still near-misses rather than accidents.
- Protect the immunity provision over time, since a single high-profile punishment of a reporter would collapse the trust the whole system depends on.
- Export the protected-reporting design, not just the database, to new domains, adapting the neutral-party structure wherever valuable safety data sits with operators who fear punishment.

REFLECTION QUESTIONS

1. Identify a domain in your institution that would benefit from an ASRS-equivalent. What cultural commitment would be required for it to function?
2. Design the institutional commitment that makes an ASRS-equivalent operational rather than merely declared.
3. ASRS made the protection credible by routing it through NASA rather than the regulator. Identify a reporting channel in your domain that operators distrust, and specify the neutral party or structural separation that would make its non-punitive promise believable.

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards



# EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation



Human-System Interface



Knowledge &amp; Institutional Memory



Governance &amp; Trust

**IMPACT** AlliedSignal's Enhanced Ground Proximity Warning System (EGPWS, 1996; now Honeywell's), mandated by the FAA as Terrain Awareness and Warning System (TAWS) for U.S.-registered turbine aircraft beginning in 2000 and broadly worldwide by 2002, converted controlled flight into terrain (CFIT) — historically one of the largest categories of commercial-aviation fatalities — into a category whose rate in equipped fleets has fallen sharply; CFIT events on properly equipped and operating airliners are now rare

## IN BRIEF

Controlled flight into terrain (CFIT) — a serviceable aircraft under the pilot's control flown unintentionally into the ground, water, or an obstacle — was for decades one of the largest categories of commercial-aviation fatalities. The 1979 Air New Zealand Mt Erebus crash (257 dead) and the 1995 American Airlines 965 Cali crash (159 dead) are canonical examples. The first-generation Ground Proximity Warning System (GPWS) developed by C. Donald Bateman at Sundstrand / Honeywell in the 1970s used radio altimeter and rate-of-descent inputs to warn of imminent terrain contact; it reduced CFIT but produced late warnings and was blind to terrain ahead of the aircraft. Enhanced GPWS (EGPWS), introduced commercially in 1996, added a digital terrain database and aircraft position to the input set, permitting forward-looking terrain-avoidance alerting. The FAA mandated EGPWS-class equipment (formally TAWS) on U.S.-registered turbine aircraft beginning March 2000, with full compliance required by 2005; ICAO and most national regulators followed. The published outcome record is that CFIT in EGPWS-equipped commercial fleets has become rare — NTSB, FAA, and Flight Safety Foundation analyses consistently report a sharp decline. The hedge that survives: residual CFIT events still occur, typically involving disabled or inhibited equipment, deviation from procedure, or terrain outside the database, and the case has to honor the system- in-operation discipline rather than the system-as-installed claim.

## BACKGROUND

Through the 1960s and 70s, controlled flight into terrain was one of the highest-fatality categories in commercial aviation. The pattern was structurally consistent: serviceable aircraft, qualified crew, often in IMC (instrument meteorological conditions) or at night, flown into rising terrain the crew had not visualized correctly. The Air New Zealand Mt Erebus crash (1979, 257 dead) and the American Airlines 965 Cali crash (1995, 159 dead) are the canonical

examples — competent crews who lost situational awareness about their position relative to terrain in conditions that prevented visual recovery.<sup>0</sup>

### THE INTERVENTION

The first engineered intervention was the Ground Proximity Warning System (GPWS), developed in the early 1970s by C. Donald Bateman at Sundstrand (later Honeywell). GPWS used radio altimeter readings and rate-of-descent to generate “pull up” and similar warnings when the aircraft was descending toward terrain directly below it. GPWS reduced CFIT meaningfully through the 1970s and 80s, but had two structural limits: it produced late warnings (the aircraft was already close to terrain when the alert fired), and it was blind to terrain ahead of the aircraft — at Cali the GPWS fired only about twelve seconds before impact, because the rising terrain was ahead of the flight path, not below.<sup>1</sup>

### HOW IT WORKED

Enhanced GPWS (EGPWS), introduced by Honeywell in 1996, addressed both limits by adding a digital terrain database and aircraft position (GPS / IRS) to the input set. The system can now compute a forward-looking terrain surface relative to the aircraft’s projected flight path and provide alerts well before terrain contact is imminent. The FAA codified the capability in the Terrain Awareness and Warning System (TAWS) regulation, requiring TAWS-class equipment on U.S.-registered turbine aircraft with six or more passenger seats under a rule published 29 March 2000 and effective 29 March 2001, binding new-build aircraft from March 2002 and the whole fleet by March 2005. ICAO and most national regulators followed with parallel mandates.<sup>2</sup>

### THE EVIDENCE

The published outcome record across NTSB accident statistics, FAA reporting, and Flight Safety Foundation analyses is that CFIT in EGPWS-equipped commercial fleets has fallen sharply. The category that once dominated airliner-fatality statistics is now an uncommon-event category in equipped fleets. The structural claim the case makes is the cue/alert-design one: a failure mode in which the operator’s perception of terrain was the limiting variable was converted into a monitored, recoverable mode by surfacing the forward terrain picture as an actionable alert. Pair with anesthesia monitoring (Case 27) at the cue/alert-as-capability layer, and with TCAS (Case 121) at the automated-advisory-system layer.<sup>3</sup>

### WHAT TRANSFERRED

The hedge has to survive into the case. CFIT events still occur, typically involving one or more of: EGPWS disabled or inhibited (crew action, MEL release, maintenance), deviation from procedure (e.g., descent below minimum sector altitude under pressure), or terrain or obstacles not represented in the database (rapidly changing wind-turbine and structure environments are a known frontier). The system-in-operation has to be flying and the crew has to act on the alert; a rule-of-thumb in the safety community is that EGPWS is most useful when its warnings are taken seriously enough that they are rare in operation. The case teaches the cue/alert intervention at its most durable, with the qualification that the capability depends on the standard being honored in operation, not on the equipment being installed.<sup>4</sup> TAWS is CAST Safety Enhancement 1, so the CFIT decline reported here is a component of the 83% portfolio result quoted in Case 117, arrived at through the industry coordination described in Case 123.

*The capability depends on the standard being honored in operation, not on the equipment being installed.*

EDITORS' SYNTHESIS OF FAA TAWS RULE HISTORY AND FSF ALAR ANALYSES.

REFERENCES

1. Royal Commission to Inquire into the Crash on Mount Erebus, Antarctica of a DC10 Aircraft Operated by Air New Zealand Limited (1981), final report (Mahon report).

2. Aeronáutica Civil of the Republic of Colombia (1996), Aircraft Accident Report: Controlled Flight Into Terrain, American Airlines Flight 965, Boeing 757-223, N651AA, near Cali, Colombia, December 20, 1995 (NTSB participating, DCA96RA020).

3. Federal Aviation Administration (2000), 14 CFR §§ 91.223, 121.354, 135.154 — Terrain Awareness and Warning System (TAWS) equipage requirement.

4. Flight Safety Foundation (1998 – 2000), CFIT / ALAR Task Force reports — operational and outcome analyses motivating mandate.

5. Bateman, C. D. (1999), “The Introduction of Enhanced Ground Proximity Warning Systems (EGPWS) into Civil Aviation Operations Around the World,” in Proceedings of the 11th Annual European Aviation Safety Seminar, Flight Safety Foundation, pp. 259–274 — developer history.

THE LEARNING ENGINEERING LENS

LE INSIGHT

EGPWS / TAWS is the canonical cue/alert intervention at fleet scale. The forward-looking terrain alert converted a failure mode in which the crew’s terrain perception was the limiting variable into a monitored, recoverable mode. CFIT in equipped fleets has become rare; residual events typically involve inhibited equipment or procedure deviation, and the hedge is the case.

LENS APPROACH

EGPWS / TAWS is the aviation cue/alert intervention case (induced 3.1; LENS D4/PT5) — Domain 4 for cue-design-as- deliverable; Domain 3 for the operator-cue boundary. Pair with TCAS (Case 121) and Case 27 (anesthesia monitoring).

| DOMAIN             | Modes Addressed | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|--------------------|-----------------|---------------------|--------------|---------|-------|
| AVIATION           | H K G           | D1 D2 D3 D4 D5      | 5            | 3.1     | 4 · 3 |
| SAFETY ENGINEERING |                 | Test and Evaluation |              |         |       |
| HUMAN FACTORS      |                 |                     |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Identify the operator-perception variable that is the limiting variable in a failure mode (here: the crew’s awareness of terrain ahead of the aircraft) and engineer the system that surfaces that variable as an actionable alert with enough lead time to recover.
- ▶ Pair the alert design with a regulatory mandate that makes the equipment non-waiverable across the fleet, so adoption is fleet-level capability rather than per-operator choice. The 1996-to-2001 gap between commercial availability and mandate is the political-process cost.
- ▶ Build the cue’s lead time around the time the operator needs to act, not the time the equipment can produce the alert; a too-late alert is the GPWS limitation EGPWS was specifically built to address.

IN OPERATION / AFTER

- ▶ Carry the system-in-operation hedge into communication: the capability depends on EGPWS being operational, not inhibited, and on the crew acting on the alert. Inhibition discipline is part of the deliverable.
- ▶ Maintain the terrain database as a continuously updated artifact; the equipment as installed is only as good as the database it queries, and rapidly changing obstacle environments (wind turbines, structures) are a known frontier.
- ▶ Treat residual CFIT events as evidence about the operational discipline, not as evidence against the intervention; the system that has reduced a fatality category to rare uncommon events is doing the work the case claims.

## REFLECTION QUESTIONS

1. Identify a failure mode in your domain where operator perception of an external variable is the limiting factor. What is the analog of the digital terrain database — the engineered representation of the variable — and what lead time would the cue need to be actionable?
2. Specify the regulatory or institutional mandate path you would expect: EGPWS reached the market in 1996, was mandated in 2001, and was fully equipaged by 2005. Five years from commercial availability to full equipage is a useful planning datum for a fleet-scale capability mandate.
3. The system-in-operation hedge is binding. What inhibition discipline would your program require so that the engineered recovery layer is operating when the failure mode appears, and how would you instrument that the discipline is being honored?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

## CASE 121

## AVIATION

## AUTOMATION

## HUMAN FACTORS

1981 – 2008 (TCAS II VERSION 7.1)

# TCAS — Coordinated Collision Avoidance and the Überlingen Lesson



Human-System Interface



Knowledge &amp; Institutional Memory



Governance &amp; Trust

**IMPACT** TCAS II — the Traffic Alert and Collision Avoidance System — provides cockpit traffic display and coordinated Resolution Advisories (RAs) between aircraft on conflicting trajectories; mandated on U.S. air-carrier and on most international turbine aircraft, TCAS converted mid-air collision in commercial aviation from a recurring fatality category to a rare event; the 2002 Überlingen mid-air (71 dead) exposed a specific coordination failure mode and drove the 2008 release of TCAS II Version 7.1, which added RA reversal logic and replaced the ambiguous 'adjust vertical speed' advisory with 'level off, level off'

## IN BRIEF

TCAS — the Traffic Alert and Collision Avoidance System, standardized in RTCA DO-185 and successors — is the cockpit automation that monitors transponder returns from nearby aircraft, computes potential conflicts, and issues Traffic Advisories and Resolution Advisories (RAs) to the crew. Operational TCAS II was mandated on U.S. air-carrier aircraft by FAA rule in the early 1990s and on most international turbine aircraft by ICAO. RAs are coordinated: when two TCAS-equipped aircraft are in conflict, one is instructed to climb and the other to descend by negotiated inversion of the data-link. The intervention converted mid-air collision in commercial aviation from a recurring fatality category to a rare event. The case's decisive edge case is the 2002 Überlingen mid-air collision (71 dead), in which one crew followed the TCAS RA and the other followed an ATC instruction in the opposite direction. The BFU investigation identified the human-TCAS-ATC coordination failure mode and drove two separate responses: an ICAO procedural change — Amendment 12 to PANS-OPS, in force from November 2003, making the RA take precedence over a conflicting ATC instruction — and the 2008 release of TCAS II Version 7.1, which added RA reversal logic for the aircraft that does not comply and replaced the ambiguous "adjust vertical speed" advisory with "level off, level off". The hedge survives into the case: TCAS is among the strongest aviation automation interventions in the outcome record, and the Überlingen failure mode and its V7.1 correction are part of the case rather than smoothed away.

## BACKGROUND

Mid-air collision in commercial aviation has been a recurring fatality category since the 1950s. The 1956 Grand Canyon mid-air (128 dead) prompted the modern U.S. air-traffic-control system, but ATC alone cannot always prevent collision when traffic densities or coordination errors exceed the controller's reach. The FAA and RTCA developed the Traffic

Alert and Collision Avoidance System (TCAS) through the 1980s as a cockpit collision-avoidance automation independent of ATC. TCAS II — the operational version with Resolution Advisories — was mandated on U.S. air-carrier aircraft beginning in the early 1990s, and on most international turbine aircraft by ICAO over the following decade.<sup>9</sup>

#### THE INTERVENTION

The TCAS II architecture is what the case rests on. Each equipped aircraft interrogates the transponders of nearby aircraft and computes a closest-point-of-approach projection from range, altitude, and rate data. If the projection enters the conflict envelope, TCAS issues a Traffic Advisory (TA) — a cue to the crew to acquire the other aircraft visually if possible. If the conflict persists, TCAS issues a Resolution Advisory (RA): a specific vertical-rate command (“Climb, climb” or “Descend, descend”). When two TCAS-equipped aircraft are in conflict, the two RAs are coordinated via the Mode S data link so the aircraft are instructed to diverge — one climbing, one descending — rather than both maneuvering in the same direction.<sup>1</sup>

#### HOW IT WORKED

The deployed outcome record across the 1990s and 2000s was strong. Mid-air collision in TCAS-equipped fleets became rare. The case nevertheless turns on Überlingen. On July 1, 2002, a Russian Tu-154 and a DHL Boeing 757 approached on conflicting trajectories at FL360 over southern Germany. Both aircraft received coordinated TCAS RAs — the 757 to descend, the Tu-154 to climb. A Skyguide controller working alone on a degraded console configuration instructed the Tu-154 to expedite descent at 21:34:49, seven seconds after the crew’s own TCAS had called traffic; seven seconds after that instruction, at 21:34:56, the crew pushed the control column forward — and in the same second their TCAS called “climb, climb.” The copilot said the TCAS was calling for a climb; the captain answered that ATC was guiding them down. The Tu-154 crew followed the ATC instruction; the 757 followed the TCAS RA; both aircraft descended into each other. 71 people died, most of them Russian schoolchildren on a charter from Moscow to Barcelona.<sup>2</sup>

#### THE EVIDENCE

The BFU named two immediate causes — ATC did not notice the separation infringement in time, and the instruction to the Tu-154 to descend came when separation could no longer be ensured; and the Tu-154 crew followed that instruction and kept descending after TCAS advised a climb, contrary to the RA — and four systemic causes: the integration of ACAS/TCAS II into the aviation system was insufficient and did not correspond in all points with the system philosophy; the ACAS regulations published by ICAO, and the national, manufacturer and operator instructions that followed from them, were “not standardised, incomplete and partially contradictory”; and the air navigation service company’s management and quality assurance neither ensured that all open workstations were staffed at night nor stopped a single-controller night practice it had tolerated for years. The data-link coordination between the two TCAS units performed as designed and could not enforce the result on the crew. The aviation response was operational and technical: ICAO’s Amendment 12 to PANS-OPS (Doc 8168), in force from November 2003, set out that a TCAS RA takes precedence over a conflicting ATC instruction and that crews must report RAs to ATC; ATC procedures were tightened around coordination of conflict-resolution between controllers and TCAS; and RTCA SC-147 with EUROCAE WG-75 developed TCAS II Version 7.1 (released 2008),

which made two separate changes: it added reversal logic for the case where an aircraft manoeuvres opposite to its resolution advisory, and it replaced the ambiguous “adjust vertical speed, adjust” advisory with the unambiguous “level off, level off”, and clarified RA wording and behavior at the boundary.<sup>3</sup>

WHAT TRANSFERRED

The hedge survives into the case. TCAS is among the strongest aviation automation interventions in the outcome record; the Überlingen failure mode is not a refutation of the system but a documented coordination limit that drove the V7.1 correction. The case teaches the coordinated-automation form at its most durable, with the discipline that the human-automation-human- operator triangle (crew, TCAS, ATC) has to be designed coherently. A single-controller / degraded-console situation and an unclear precedence rule converted a working automation into a fatal outcome; both were addressed in the post-Überlingen response, and both remain part of the case rather than smoothed away.<sup>4</sup>

*The intervention is not less effective for having a documented coordination edge case; it is more credible because the edge case is named and addressed.*

EDITORS' SYNTHESIS OF THE BFU ÜBERLINGEN REPORT AND THE V7.1 DEVELOPMENT RECORD.

REFERENCES

1. Federal Aviation Administration, 14 CFR § 121.356 — TCAS II equipage requirement.

2. RTCA (2008), DO-185B “Minimum Operational Performance Standards for Traffic Alert and Collision Avoidance System II (TCAS II)” — Version 7.1 with reversal logic.

3. Bundesstelle für Flugunfalluntersuchung (BFU) (2004), AX001-1-2/02 — Investigation Report on the mid-air collision on 1 July 2002 near Überlingen.

4. ICAO, PANS-OPS (Doc 8168), Amendment 12 — in force November 2003; the RA-takes-precedence rule and the requirement to report RAs to ATC.

5. Espindle, Griffith & Kuchar, “Safety Analysis of Upgrading to TCAS Version 7.1 Using the 2008 U.S. Correlated Encounter Model,” MIT Lincoln Laboratory Project Report ATC-349 (2009) — the quantitative basis for the V7.1 change.

6. Commission Regulation (EU) No 1332/2011 — ACAS II version 7.1 required in EU airspace from 1 March 2012 for new aeroplanes and from 1 December 2015 for the existing fleet.

THE LEARNING ENGINEERING LENS

LE INSIGHT

TCAS is among the strongest aviation automation interventions in the outcome record; the Überlingen coordination failure mode is part of the case rather than smoothed away. The human-automation-ATC triangle has to be designed coherently; V7.1’s reversal logic and the clarified precedence rule are the iterative-design response to the documented edge case.

LENS APPROACH

TCAS is the coordinated-cockpit-automation case (induced 3.1; LENS D4/PT5) — Domain 4 for the Überlingen-driven V7.1 iteration; Domain 3 for the crew-TCAS-ATC precedence rule. Pair with EGPWS (Case 120) and Case 27.

| DOMAIN        | MODES ADDRESSED | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|---------------|-----------------|---------------------|--------------|---------|-------|
| AVIATION      | H K G           | D1 D2 D3 D4 D5      | 5            | 3.1     | 4 · 3 |
| AUTOMATION    |                 | Test and Evaluation |              |         |       |
| HUMAN FACTORS |                 |                     |              |         |       |

APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- ▶ Design the coordinated-automation logic for the case where one of the coordinated agents does not comply — the Überlingen failure mode — not just for the case where both comply. The V7.1 reversal logic is the worked example of that design discipline.
- ▶ Specify precedence at the human-automation boundary unambiguously and in advance: TCAS RAs take precedence over conflicting ATC instructions. Leaving precedence to crew judgment under time pressure is the design choice that produced Überlingen.
- ▶ Treat the ATC context (single-controller, degraded console) as part of the system the automation operates in, not as a precondition the automation can assume away. Systemic failures at the boundary determine the boundary behavior.

## IN OPERATION / AFTER

- ▶ Carry the Überlingen failure mode into the case rather than smoothing it away. The intervention is not less effective for having a documented coordination edge case; it is more credible because the edge case is named and addressed.
- ▶ Track the post-V7.1 outcome record as evidence about the correction, not just about the original intervention; the lesson is the iterative-design discipline that the human-automation triangle requires after a failure mode is exposed.
- ▶ Use the case as the canonical pair to EGPWS (Case 120): two cockpit automations, two outcome categories closed, one with a coordination edge case that drove a version revision and one without.

## REFLECTION QUESTIONS

1. Identify a coordinated-automation system in your domain where two agents must comply for the resolution to work. What is the analog of the V7.1 reversal logic — the design for the case where one agent does not comply?
2. Specify the precedence rule at the human-automation boundary in your system. Überlingen turned on an ambiguous precedence rule; the post-2002 clarification is the worked example of why precedence has to be unambiguous in advance.
3. The systemic context (Skyguide single-controller / degraded console) was part of the failure. What contextual preconditions does your automation assume that, if they fail, would convert a working automation into a failure mode?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · governance frameworks · bias & evidence audits



# Singapore Airlines Safety Transformation

**T** Training Gap    **N** Normalization of Deviance

**IMPACT** Sustained safety record over decades despite challenging operating conditions; among the most safety-invested carriers in commercial aviation

## IN BRIEF

Singapore Airlines has invested in safety capability across decades in a way that sets it apart from carriers operating under comparable conditions — early adoption of CRM, heavy simulator investment, a young-fleet policy, and a strong reporting culture, sustained even through rapid growth. The 2000 crash of Flight SQ006, which began its takeoff roll at Taipei on a runway open only for taxi, a section of it closed for construction, during a typhoon and killed 83, prompted an institutional response — early public acknowledgement, compensation offered ahead of findings, and operational review — that the crisis-communication literature examines, while the Singapore Ministry of Transport's investigation team formally contested the Taiwan report's apportionment of cause. The airline is the operational successor to Korean Air (Case 118): an Asian carrier that engineered its safety capability deliberately and sustained the investment as a competitive differentiator, not only in response to crisis.

## BACKGROUND

Commercial aviation runs on thin margins, and safety investment — simulators, training hours, fleet renewal — is a cost that competitive pressure constantly pushes downward. The question for any carrier is whether to treat capability as a floor set by regulation or as a deliberate, sustained investment ahead of the minimum. The pressure is structural rather than occasional — every budget cycle invites trimming the margin between regulatory minimum and actual capability, so sustaining the investment requires deciding the question deliberately rather than by default.<sup>0</sup>

## THE INTERVENTION

Singapore Airlines chose sustained investment. From the 1980s it was an early adopter of Crew Resource Management and CRM-style culture work tuned to its operating context, and it committed to heavy simulator investment, a deliberately young fleet, and a strong internal reporting culture — maintaining these even during periods of rapid expansion. Holding the investment through rapid growth is the hard test — expansion is precisely when the temptation to let capability lag the fleet is strongest, and maintaining it then is what separates a sustained commitment from a fair-weather one.<sup>1</sup>

## HOW IT WORKED

The carrier treated safety capability as a competitive differentiator rather than a regulatory burden, pairing technical investment — training systems, modern aircraft — with a culture of transparency and reporting. Investing ahead of regulatory minimums made the capability

a managed system parameter, not a residual of cost-cutting decisions made elsewhere. Framing safety as a differentiator rather than a burden is what made the investment defensible against cost pressure — it tied capability to the brand and the premium the airline charged, giving the spend a commercial rationale, not just a safety one.<sup>2</sup>

#### THE EVIDENCE

The 2000 crash of Flight SQ006 — a takeoff roll begun at Taipei on Runway 05R, which was open only for taxi and had a section closed for construction, during Typhoon Xangsane, killing 83 of the 179 aboard — tested the institution. The Taiwan Aviation Safety Council's final report listed eight findings related to probable causes: the first was the typhoon weather, the other seven the flight crew. The Singapore Ministry of Transport's investigation team filed comments, appended to the report under ICAO Annex 13, arguing that the CKS Airport deficiencies belonged among the probable causes rather than among the "findings related to risk." The airline's own conduct — early acknowledgement, senior management on site, compensation offered ahead of findings — is what the crisis-management literature examines, and it is a narrower claim than a clean institutional-learning record. The test of a safety culture is how it behaves after its own accident, and choosing transparency over defensiveness is what turned SQ006 from a refutation of the airline's reputation into evidence the reporting culture extended to its own failures.<sup>3</sup>

#### WHAT TRANSFERRED

Singapore Airlines is the case for sustained capability investment under competitive pressure, and the operational successor to Korean Air (Case 118): where Korean Air is a transformation forced by crisis, Singapore Airlines is deliberate investment sustained without one. Together they show two routes — crisis-driven and voluntary — to the same engineered safety capability. The voluntary route is the harder one to hold, because it has no catastrophe to point to as justification, which is why framing the investment as a competitive differentiator matters: it supplies the rationale that a crisis would otherwise provide.<sup>4</sup>



*sustained capability investment ahead of regulatory minimums*

SINGAPORE AIRLINES — CAPABILITY INVESTMENT AS COMPETITIVE DIFFERENTIATOR

*Singapore Airlines has consistently invested in safety capability ahead of regulatory minimums.*

EDITORS' SYNTHESIS ON SINGAPORE AIRLINES' SUSTAINED SAFETY INVESTMENT

#### REFERENCES

1. Aviation Safety Council (Taiwan), Crashed on a Partially Closed Runway during Takeoff, Singapore Airlines Flight 006, Boeing 747-400, 9V-SPK, CKS Airport, Taoyuan, October 31, 2000 (2002) — the accident, the eight findings related to probable causes (§3.1: one weather, seven flight crew), and the Singapore MOT investigation team's comments appended at §7.

2. Henderson, J. C. (2003), "Communicating in a crisis: flight SQ 006," Tourism Management 24(3):279–287 — the analysis of Singapore Airlines' early crisis response.

3. Singapore Airlines, "Singapore Airlines keeps fleet young" (SilverKris/company account) and SIA annual reports — the young-fleet policy in the airline's own words; in-service average fleet age about 8.3 years as of 2026.

4. Helmreich, Wilhelm, Klinec & Merritt (2001), "Culture, error, and crew resource management," in Salas, Bowers & Edens (eds.), Improving Teamwork in Organizations — national culture and CRM adaptation.

5. Weick & Sutcliffe (2007), Managing the Unexpected — sustained high-reliability investment.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Singapore Airlines is the case for sustained capability investment in a competitive industry. The carrier has chosen safety capability investment as a primary differentiator. The result over decades is a safety record that distinguishes it from peers operating under comparable conditions.

LENS APPROACH

LENS uses Singapore Airlines in LEN 8 for sustained institutional capability investment under competitive pressure. The case pairs with Korean Air (Case 118) as Asian-carrier capability stories of different shapes — one transformation under crisis, the other sustained investment without crisis.

| DOMAIN                                |   | MODES ADDRESSED |   | LENS COMPETENCY |  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---|-----------------|---|-----------------|--|--------------|---------|-----|
| AVIATION                              | → | T N             | → | D1 D2 D3 D4 D5  |  | 4            | 1.4     | 5   |
| Navigating Sociotechnical Constraints |   |                 |   |                 |  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Decide deliberately to invest ahead of regulatory minimums — simulators, a young fleet, training hours — rather than letting capability settle at the floor by default under cost pressure.
- ▶ Pair the technical investment with a transparency and reporting culture, so capability is a managed system parameter rather than a residual of cost-cutting decisions elsewhere.
- ▶ Frame safety capability as a competitive differentiator tied to the brand, giving the spend a commercial rationale that can survive scrutiny, not just a safety one.

IN OPERATION / AFTER

- ▶ Hold the investment through periods of rapid growth — the moment the temptation to let capability lag the fleet is strongest — since maintaining it then is what makes the commitment sustained rather than fair-weather.
- ▶ Respond to the institution's own accidents by publishing what changed operationally, not only by contesting the apportionment of cause; a reporting culture is credible only when it extends to the institution's own failures.
- ▶ Anchor the investment to the brand and premium so it survives leadership and budget cycles, supplying the durable justification a voluntary commitment lacks without a crisis to point to.

## REFLECTION QUESTIONS

1. Identify an institution in your domain that has chosen capability investment as a primary differentiator. What pattern has it sustained?
2. Design the institutional architecture that makes sustained capability investment defensible against competitive cost pressure.
3. Singapore Airlines sustained its investment voluntarily, without a crisis to point to, by framing capability as a competitive differentiator. Identify a safety or capability investment in your domain that lacks a catastrophe to justify it, and construct the commercial rationale that would defend it against the next budget cut.

**WHO BUILDS THIS** learning science & instructional design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards

## CASE 123

## AVIATION SAFETY

## INDUSTRY COORDINATION

## TRAINING

1992 – 2000S

# FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike

**G** Governance & Trust**K** Knowledge & Institutional Memory**N** Normalization of Deviance

**IMPACT** After Controlled Flight Into Terrain emerged as the leading cause of commercial-jet fatalities through the late 1980s, the Flight Safety Foundation convened industry-wide task forces that produced the CFIT Checklist, the ALAR Tool Kit, and the institutional momentum behind Terrain Awareness and Warning System (TAWS) mandates; CFIT and ALAR accident rates fell sharply over the subsequent decade

## IN BRIEF

Through the late 1980s and into the early 1990s, Controlled Flight Into Terrain — a serviceable aircraft flown under control into the ground, water, or an obstacle — was the leading cause of commercial-jet fatalities worldwide. The Flight Safety Foundation (FSF), an independent industry body, convened the CFIT Task Force in 1992 and, in parallel with the broader ICAO response, produced the CFIT Checklist — a structured tool for operators to assess their own exposure. The Approach-and-Landing Accident Reduction (ALAR) Task Force followed in 1996, producing the ALAR Tool Kit (released 2000) covering the approach phase where roughly half of fatal accidents then occurred. The institutional momentum from those task forces sat behind the eventual Terrain Awareness and Warning System (TAWS) mandates in the U.S. (2002) and ICAO (2007). CFIT and approach-and-landing accident rates fell sharply through the subsequent decade. The case teaches industry-level institution building after a catastrophe-class spike: the deliverable is the cross-operator tool, the diagnostic structure, and the coordinated path to mandate. The hedge survives — the rate decline is multifactorial (TAWS hardware, training, procedural change, fleet turnover) — and the institutional form is what the case is teachable on.

## BACKGROUND

Through the late 1980s and into the early 1990s, Controlled Flight Into Terrain dominated commercial-jet fatality statistics. A serviceable aircraft, flown under control, was finding terrain or water — the crew's mental model of position and trajectory diverged from reality at the worst possible moment, and the existing Ground Proximity Warning Systems (GPWS) generation reached the crew too late to recover in many terrain encounters. The pattern was recognized across operators and regulators, and the response that emerged was industry coordination led by an independent body — the Flight Safety Foundation — rather than regulator-first action alone.<sup>o</sup>

## THE INTERVENTION

The CFIT Task Force convened in 1992 and operated as a working group across operators, manufacturers, regulators, and pilot organizations. Its central deliverable was the CFIT Checklist — a structured self-assessment instrument that let an operator score its own CFIT-risk exposure across route, equipment, training, and procedure dimensions, and identify where the gaps sat. The Checklist was distributed without restriction. The institutional theory was straightforward: a cross-operator diagnostic, owned by no single competitor, would let the industry move together on a problem that no single operator's own data could fully characterize.<sup>1</sup>

## HOW IT WORKED

The Approach-and-Landing Accident Reduction Task Force followed in 1996, scoped to the approach phase, which accounted for roughly half of fatal accidents in the analysis window. The ALAR Tool Kit, released in 2000, was the most substantial deliverable of the entire industry-coordination effort: a multi-element package of briefing notes, training aids, video content, and risk-assessment instruments covering stabilized approach criteria, runway excursion, monitored approach practice, and crew procedural design. The Tool Kit was adopted across operators of all sizes precisely because the FSF had no competitive stake in any one airline's adoption.<sup>2</sup>

## THE EVIDENCE

The eventual regulatory action — the FAA's TAWS mandate (final rule 2000: new-build compliance 2002, retrofit of earlier GPWS-equipped aircraft with the newer terrain-database-driven warning system) and the ICAO Annex 6 requirement — a forward-looking terrain-avoidance function on turbine aeroplanes above 5,700 kg or carrying more than nine passengers, for new certificates of airworthiness from 1 January 2004 and for the rest of that fleet by 1 January 2007 — sat downstream of the task-force momentum, not upstream of it. CFIT accident rates fell sharply through the subsequent decade across both commercial transport and corporate aviation; approach-and-landing accidents declined alongside. The task forces set an explicit target — a 50 percent reduction in the approach-and-landing accident rate in five years — and the decline that followed ran well past it, though no published analysis isolates the task forces' share of it.<sup>3</sup>

## WHAT TRANSFERRED

The hedge that matters survives into the case. The accident-rate decline is multifactorial: TAWS hardware in the cockpit, training changes the task forces catalyzed, stabilized-approach criteria adopted at the operator level, and steady fleet turnover into airframes with more capable equipment all moved together. Attributing the entire decline to the FSF task forces alone overstates the evidence; what the evidence supports is that the industry-coordination form — independent convening body, cross-operator diagnostic tools released without competitive restriction, momentum sustained to a regulatory mandate — was the institutional mechanism that organized the response, and the response worked. The case is the canonical C6.1 instance of industry-level institution building after a catastrophe-class spike, paired with v1 ASRS (Case 119) and CAST (Case 117) at the industry-coordination layer.<sup>4</sup> The task forces, the TAWS mandate (Case 120), and the CRM/CAST apparatus (Case 117) are three components of one decline; the multifactorial hedge governs all three, and no case's

share of the 83% reduction in U.S. commercial-aviation fatality risk that CAST reports for 1998–2008 (Case 117) is separately measured.

*The deliverable is the cross-operator diagnostic, owned by no single competitor, released without restriction. The mandate consolidates what the field already does.*

EDITORS' SYNTHESIS OF THE FSF CFIT AND ALAR TASK-FORCE LITERATURE.

REFERENCES

1. Flight Safety Foundation, "Killers in Aviation: FSF Task Force Presents Facts About Approach-and-landing and Controlled-flight-into-terrain Accidents," Flight Safety Digest (1998–1999).

2. Flight Safety Foundation, ALAR Tool Kit (2000; updated 2010) — distributed multi-element package built on the 34 ALAR Briefing Notes published in Flight Safety Digest, August–November 2000.

3. Khatwa & Helmreich, "Analysis of critical factors during approach and landing in accidents and normal flight," Flight Safety Digest (1998) — the analytical basis of the ALAR Task Force scope.

4. FAA, Terrain Awareness and Warning System (TAWS) Final Rule, 65 FR 16736 (29 March 2000), 14 CFR §§ 91.223, 121.354, 135.154 — new-build compliance 29 March 2002, retrofit by 29 March 2005.

5. ICAO Annex 6, ground-proximity warning with a forward-looking terrain-avoidance function — new certificates of airworthiness from 1 January 2004, the existing turbine fleet above 5,700 kg or over nine passengers by 1 January 2007; the international consolidation.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The FSF CFIT and ALAR task forces are the canonical case of industry-level institution building after a catastrophe-class spike. An independent foundation convened the cross-operator working bodies, released structured diagnostic tools without competitive restriction, and sustained momentum to a regulatory mandate. The accident-rate decline is multifactorial; the institutional form is what the case is teachable on.

LENS APPROACH

FSF CFIT/ALAR is the canonical industry-coordination case (induced 6.1; LENS D5/PT4) — Domain 5 for the independent- convening-body form; Domain 2 for the structured diagnostic instruments as iterating deliverables. Pair with Cases 117, 119, and 13.

| DOMAIN                | Modes Addressed | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO |
|-----------------------|-----------------|---------------------------------------|--------------|---------|-----|
| AVIATION SAFETY       | G K N           | D1 D2 D3 D4 D5                        | 4            | 6.1     | 5   |
| INDUSTRY COORDINATION | →               | →                                     |              |         |     |
| TRAINING              |                 | Navigating Sociotechnical Constraints |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Convene the cross-operator working body before the regulator does — an independent industry foundation, no competitive stake in any one operator's adoption — so the diagnostic tool can be released without restriction and adopted across the field.
- ▶ Design the diagnostic as a structured self-assessment, not a benchmark league table; operators will use what they can apply privately, and the field-level signal aggregates from voluntary adoption.
- ▶ Build the tool kit around the phase of flight that accident analysis says dominates the harm — approach-and-landing in this case — rather than

IN OPERATION / AFTER

- ▶ Sustain the institutional momentum into the regulatory mandate; the task forces did the technical and political work that made the TAWS rule adoptable, and the rule consolidates what the field already does.
- ▶ Report the accident-rate decline with the multifactorial hedge intact; TAWS hardware, training, procedure, and fleet turnover all moved together, and isolating the task-force contribution overstates the available evidence.
- ▶ Treat the institutional form as the transferable artifact: an independent convening body, cross-operator diagnostic tools without competitive restriction,

around evenly distributed coverage that no single operator can adopt fully.

coordinated path to regulatory mandate. The form pairs with CAST (Case 117) and ASRS (Case 119) at the industry-coordination layer.

### REFLECTION QUESTIONS

1. Identify a catastrophe-class failure pattern in your domain whose response has been operator-by-operator rather than industry-coordinated. What would the analog of an independent convening body look like, and which body could plausibly play that role without competitive stake?
2. Specify the cross-operator diagnostic — checklist, tool kit, structured self-assessment — that you would design as the first deliverable of an FSF-style task force. The deliverable has to be applicable privately by each operator, and aggregate into field-level signal.

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards



Chapter 7

# Defense & National Security — What Fails

*When the capability the mission requires goes unspecified at the point of contact.*

---

*The system was fielded. The operator's task was not.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

# USS Fitzgerald & USS John S. McCain

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** 17 sailors killed in two avoidable destroyer collisions in the Western Pacific within nine weeks

## IN BRIEF

In the summer of 2017 two U.S. Navy destroyers of the forward-deployed Seventh Fleet collided with merchant ships nine weeks apart, killing seventeen sailors — seven on the Fitzgerald, ten on the John S. McCain. Their crews' seamanship, navigation, and console skills had eroded under a decade of relentless operational tempo and the self-study "training" that replaced the in-person school the Navy cut in 2003. Investigations by the Navy and the NTSB judged both collisions avoidable and traced them to insufficient training and weak oversight; on the McCain, a confusing touch-screen helm let a watch team believe it had lost steering it never actually lost. The case is this book's canonical training gap: stated readiness and real readiness diverged for years, and the system could no longer see the difference.

## BACKGROUND

The destroyers belonged to the forward-deployed Seventh Fleet in Japan, kept at the fleet's highest tempo — where training was the first thing spent. Unlike home-ported ships, which rotate through a dedicated work-up cycle before deploying, the forward-deployed force was expected to be continuously available, so there was rarely a quiet stretch in which to recover a lapsed qualification. In 2003 the Navy replaced its in-person Surface Warfare Officers School course with self-study CD-ROMs — "SWOS in a Box" — sending newly commissioned officers directly to their ships to learn the trade as time and the ship's qualified watchstanders allowed.<sup>0</sup> The change was framed as a cost-saving modernization; in practice it transferred junior-officer education from the schoolhouse to the bridge at the moment the operational tempo on those bridges left the least slack to absorb it. By 2017 the GAO found 37% of the Japan-based ships' warfare certifications — including basic seamanship — expired, a fivefold rise since 2015, with the lapses routinely waived rather than fixed.<sup>1</sup>

## WHAT HAPPENED

On 17 June 2017 the Fitzgerald, the give-way vessel in a busy shipping channel, took no early action and was struck by the container ship ACX Crystal off the coast of Japan; the sea poured into a berthing compartment where the crew slept, and seven sailors drowned before they could escape.<sup>2</sup> The NTSB later described a bridge team that had lost track of converging traffic on a clear night, and an officer of the deck who made poor navigational decisions and never called the combat information center for the support the situation demanded. Nine weeks later the John S. McCain collided with the tanker Alnic MC near

Singapore, killing ten: while shifting throttle control between consoles, a watchstander unknowingly handed off steering to another station, the ship turned across the strait's traffic, and no one on the bridge understood the touch-screen helm well enough to recognize what had happened.<sup>3</sup> For about three minutes the bridge team believed the ship had lost steering it had never lost — a misdiagnosis the interface invited and the training had not equipped anyone to overturn.

#### THE INVESTIGATION

The Navy's command investigation reports (November 2017) judged both collisions avoidable — Fitzgerald from an accumulation of smaller errors and lapsed navigational practice, McCain from complacency, over-confidence, and lack of procedural compliance.<sup>4</sup> The NTSB found the McCain's probable cause to be a lack of Navy oversight that produced insufficient training and inadequate bridge procedures,<sup>5</sup> and faulted the design of the touch-screen helm, installed to modernize the bridge, whose blending of steering and throttle made an unintended transfer of control easy to trigger and hard to notice — a trap waiting for an under-trained crew. The companion Strategic Readiness Review, commissioned by the Secretary of the Navy and led by retired Admiral Gary Roughead and Michael Bayer, reached further into the institution: a decade of "can-do" responses to mounting demand had eroded the manning, certification, and maintenance margins the surface force was built on, and the readiness reports senior leaders relied on had stopped reflecting the conditions on the ships.<sup>6</sup> Navy rest records showed the McCain's fourteen bridge watchstanders averaging just over 4.9 hours in the 24 hours before the collision — a finding the NTSB folded into its causal chain.

#### THE CAPABILITY GAP

The gap was invisible from inside. The Comprehensive Review described risks that "accumulated over time and did so insidiously," the system no longer able to see that the processes meant to surface shortfalls had themselves failed.<sup>7</sup> Each individual waiver was locally defensible — a deadline met, a deployment kept — but in aggregate they hollowed out the force while every readiness dashboard still reported green. Stated and actual readiness had diverged for a decade; the cost arrived as seventeen lives, not a failed inspection. Two collisions, nine weeks apart, ruled out the comforting reading that one was an outlier: Fitzgerald taught that an under-trained bridge team could fail at the most basic give-way rules, and McCain taught that an unfamiliar interface could be the precipitating mechanism through which the same gap reached the hull. The pair, not either case alone, made the diagnosis structural rather than individual.

#### AFTERMATH & REFORM

The reforms were the deepest in a generation: the in-person officer pipeline was rebuilt around a 2018 training continuum that extended the Basic Division Officer Course and added officer-of-the-deck phases, restoring classroom and simulator instruction the 2003 CD-ROM model had displaced; a Ready-for-Sea Assessment and a Naval Surface Group Western Pacific stood up to give forward-deployed units the independent certification cycle home-ported ships already had; circadian watchbills were adopted fleet-wide to fight fatigue; and the touch-screen helm was slated for replacement by a conventional wheel and throttle across the destroyer fleet.<sup>8</sup> Each change conceded that the training and the

interface had been real requirements all along — and that trading them for tempo had only moved the cost onto two ships full of people.<sup>9</sup>



*The risks that were taken in the Western Pacific accumulated over time and did so insidiously.*

U.S. NAVY STRATEGIC READINESS REVIEW, 2017

## REFERENCES

1. T. C. Miller, M. Faturechi & R. Rotella, “Years of Warnings, Then Death and Disaster,” ProPublica (2019) — the 2003 SWOS-in-a-Box shift.
2. GAO, Navy Readiness, GAO-17-809T (Sept. 2017) — 37% of Japan-based warfare certifications expired by June 2017.
3. NTSB, Collision between USS Fitzgerald and MV ACX Crystal, NTSB/MAR-20/02 (2020), DCA17PM018.
4. NTSB, Collision between USS John S. McCain and Tanker Alnic MC, NTSB/MAR-19/01 (2019), DCA17PM029.
5. U.S. Navy, Comprehensive Review of Recent Surface Force Incidents (Nov. 2017) — both collisions judged avoidable.
6. NTSB/MAR-19/01 (2019) — probable cause: insufficient training and inadequate bridge procedures from a lack of Navy oversight.
7. R. Rotella et al., “The Navy Installed Touch-Screen Steering Systems to Save Money,” ProPublica (2019).
8. U.S. Navy, Strategic Readiness Review (Dec. 2017) — risks that “accumulated over time and did so insidiously.”
9. U.S. Navy corrective actions (2017); NTSB/MAR-19/01 recommendations; USNI News (2017–2020).
10. Surface Warfare Officers School Command, Basic Division Officer Course curriculum and the post-2017 return to in-person instruction; Naval Surface Group Western Pacific stand-up (2018) as the forward-readiness certification authority.
11. NTSB/MAR-19/01 (2019) — watchstander fatigue findings, including average sleep hours and the touch-screen helm misdiagnosis sequence.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Fitzgerald/McCain is the canonical Training Gap case because the gap was invisible from inside the system. Operational tempo and self-study “training” combined to produce a fleet whose stated readiness and actual readiness diverged for more than a decade. The two collisions nine weeks apart converted what could have been read as an outlier into a structural diagnosis: an under-trained bridge team failing the most basic give-way rules, and an unfamiliar interface that made the same gap reach the hull. Seven sailors died on Fitzgerald; ten on McCain. The cost of the divergence was paid in lives long after it could have been measured in dollars or inspections.

### LENS APPROACH

Fitzgerald/McCain is the worked example of induced sub-competency 1.1 (engineered vs. stated requirements) and the LENS D1/PT3 pairing — Systems Analysis applied to the capability-engineering problem of underway watchstanding. Students reconstruct the capability requirements from operational analysis, then design the evidence architecture (LENS D3) that would have surfaced the gap before the collisions and the sociotechnical reforms (LENS D4) that would keep waivers from quietly hollowing the requirement out. The case pairs with Three Mile Island (Case 166) as the failure that engineered durable industry reform through INPO. LEO mapping: LEO-1 (Systems Analysis) primary, LEO-5 (Sociotechnical Constraints) for the waiver-and-reporting institutional dynamics.

| DOMAIN           | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO   |
|------------------|---------------|------------------|--------------|---------|-------|
| DEFENSE          | → T K N       | → D1 D2 D3 D4 D5 | 3            | 1.1     | 1 · 5 |
| Systems Analysis |               |                  |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify watchstanding proficiency as a measured deliverable — define the competency, instrument it, and gate qualification on demonstrated skill rather than sea time.
- Keep human-factors review in the procurement loop: validate that any interface change (e.g., the bridge console) is tested against operator performance before it is fielded.
- Design the readiness signal to report demonstrated capability, not self-attested course completion.

IN OPERATION / AFTER

- Audit the gap between reported and actual readiness with independent assessment (the Ready-for-Sea model) — with authority to pull a unit offline.
- Protect training time against operational tempo so the capability cannot erode silently when schedules tighten.
- Track leading indicators — qualification currency, near-miss reports — so divergence is visible before it is paid for in lives.

REFLECTION QUESTIONS

1. The Navy replaced classroom and simulator instruction with CD-ROM self-study in 2003. What measurement would have detected the capability shortfall before 2017?
2. The Strategic Readiness Review found the readiness-reporting system had itself stopped working. Design a capability-evidence pipeline that would not normalize its own gaps.
3. Identify a capability in your organization that is certified by completion (a checked box) rather than demonstrated performance. What measure would convert it to evidence — and who would hold the authority to act on a red signal?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

# F-35 Sustainment & the Maintainer Capability Gap

**T** Training Gap   **K** Knowledge & Institutional Memory   **D** Designed Out

**IMPACT** Fleet mission-capable rate 44% in FY2025, down from 67% in FY2021 (GAO-26-108113); sustainment cost estimate \$1.58T, total acquisition and sustainment above \$2T (GAO-24-106703); technical-data, training, and depot shortfalls are the binding readiness constraint (GAO-23-105341)

## IN BRIEF

The F-35 is the most expensive weapons program in history — roughly 2,500 jets planned and, on GAO's 2023 figures, a lifecycle cost above \$1.7 trillion, about \$1.3 trillion of it in operating and support; GAO's 2024 review put the sustainment estimate alone at \$1.58 trillion and the acquisition-plus-sustainment total above \$2 trillion. The hard part was never the airplane but keeping a global fleet ready, and that half has lagged from the start. As of March 2023 the fleet's mission-capable rate was about 55%, far short of goal, with more than 10,000 components in the repair queue and depots behind schedule. GAO traced the shortfall to gaps in maintainer training and support equipment, the military's lack of access to technical data, and contractor dependency, and urged a full reassessment of the sustainment strategy. It is the book's cleanest case of a platform fielded faster than the capability infrastructure to sustain it.

## BACKGROUND

The F-35 is the most expensive weapons program in history: the Pentagon plans to buy nearly 2,500, at a total acquisition and sustainment cost GAO put above \$2 trillion in 2024 — \$1.58 trillion of it not the aircraft but the decades of operating and sustaining them.<sup>o</sup> The flyaway jet is the finite part; keeping a global fleet ready — maintainers, technical data, depots — is the open-ended part, and the part that lagged, because the cost that dominates the program is not buying the aircraft but the decades of sustaining them, the very work that received the least attention as the jets rolled off the line.

## WHAT HAPPENED

The failure is a standing condition. As of March 2023 the fleet's mission-capable rate was about 55%, far short of goal; more than 10,000 components waited in the repair queue, average repair cycle time ran to 141 days against a 60-day goal, and the depots were still behind schedule in standing up the capacity to do the work at all — a backlog and a turnaround time that compound, since parts stuck in the queue keep jets grounded and the under-built depots cannot clear the queue fast enough to recover the mission-capable rate.<sup>1</sup>

## THE INVESTIGATION


GAO's September 2023 review was bluntly titled: the Department and services **need to reassess the future sustainment strategy**. It traced the shortfall to gaps in maintainer training and support equipment, the military's lack of access to the technical data needed to do its own repairs, and the resulting dependence on the prime contractor.<sup>2</sup> None of it was new — GAO has repeated the same diagnosis year after year, through a troubled logistics-software backbone and slow progress, even as procurement continued and readiness stayed flat, so the program kept buying more jets it could not fully sustain while the same three shortfalls were named in review after review without being closed.<sup>3</sup>

## THE CAPABILITY GAP

The F-35 is the cleanest modern case of a platform fielded faster than the capability infrastructure to sustain it. Aircraft arrived on schedule; the maintainers, technical data, and depots were treated as follow-on costs rather than deliverables that had to field with the jets. The hardest part is the data: much of what is needed to repair the aircraft stayed controlled by the contractor, so the services cannot freely write procedures, qualify depots, or compete the work, which locks the fleet into a single source for the knowledge needed to keep it flying — merely expensive in peacetime, dangerous in war, when a contractor-dependent sustainment chain is exactly the kind of bottleneck an adversary would seek to exploit.<sup>4</sup>

## AFTERMATH & REFORM

GAO urged a full reassessment of the sustainment strategy — government access to technical data, depot capacity, and a maintainer pipeline — rather than more patching.<sup>5</sup> Those recommendations remain a work in progress, and readiness has since deteriorated rather than recovered: GAO's June 2026 review found the fleet mission-capable rate had fallen from 67 percent in FY2021 to 44 percent in FY2025, and the full-mission-capable rate from 38 to just 25 percent — the fleet able to fly all its missions only a quarter of the time. Air Force officials attributed that year's full-mission-capable decline to the program accepting new aircraft that were non-mission capable because of software delays, alongside scarce parts and corrosion inspections, while the Joint Program Office's remediation ("Global Support Solution Reset") was projected to cost an estimated \$13.7 billion more through FY2031, with readiness expected to worsen before it improves. The F-35 sits in this book as the live argument for treating capability infrastructure — people, data, the means to sustain them — as a fielding gate, not an afterthought. The bill for skipping that gate does not disappear; it compounds, arriving as grounded jets, a swelling repair queue, and a sustainment chain the services cannot control rather than as a line item caught early enough to fix cheaply.



MISSION-CAPABLE RATE VS. FLEET TARGET

*Organizational-level maintenance has been affected by a number of issues, including a lack of technical data and training.*

GAO-23-105341, F-35 AIRCRAFT, 2023 — WHAT GAO FOUND

REFERENCES

1. U.S. Government Accountability Office, F-35 Aircraft: DOD and the Military Services Need to Reassess the Future Sustainment Strategy, GAO-23-105341 (Sept. 2023) — 2,500 planned aircraft and a lifecycle cost exceeding \$1.7 trillion, \$1.3 trillion of it in operating and support.

2. GAO-23-105341 (2023) — 55% fleet mission-capable rate (March 2023); over 10,000 components awaiting repair; 141-day average depot repair cycle time against a 60-to-90-day goal; depot stand-up behind schedule.

3. GAO-23-105341 (2023) — sustainment shortfall traced to insufficient technical data, training and support equipment for maintainers, depot stand-up delays, and contractor dependency.

4. GAO, F-35 Sustainment: DOD Faces Several Uncertainties..., GAO-22-105995 (2022), and the broader GAO F-35 series — the recurring, year-over-year diagnosis.

5. GAO-23-105341 (2023) — recommendation that DOD reassess the future sustainment strategy.

6. GAO, F-35 Sustainment: Costs Continue to Rise While Planned Use and Availability Have Decreased, GAO-24-106703 (2024) — costs rising while readiness stays below goal; \$1.58 trillion sustainment estimate and the above-\$2 trillion total.

7. GAO, F-35 Sustainment: Actions Needed to Ensure Updated Strategy Improves Persistent Readiness Challenges, GAO-26-108113 (11 June 2026) — mission-capable rate 67% (FY2021) to 44% (FY2025), full-mission-capable 38% to 25%, and the Global Support Solution Reset's \$13.7 billion of additional funding through FY2031.

THE LEARNING ENGINEERING LENS

LE INSIGHT

F-35 is the live, current example of fielding a platform faster than its capability infrastructure can be built. The aircraft is the easy part; the maintainers are the hard part. A decade of program schedules treated maintainer training as a follow-on cost, not a fielding gate. The fleet is now operating at half of its design readiness, and the program of record now carries a life-cycle estimate above \$2 trillion, of which about \$1.58 trillion is sustainment.

LENS APPROACH

LENS treats the F-35 in LEN 5 as the canonical case of **Capability- System Misalignment at Transition**, the program's first canonical problem type. Students design the capability infrastructure that should have accompanied each lot of aircraft delivered, including the training pipeline, technical-data deliverables, and depot establishment.

| DOMAIN           | FAILURE MODES | LENS COMPETENCY | PROBLEM TYPE | INDUCED | LEO |
|------------------|---------------|-----------------|--------------|---------|-----|
| DEFENSE →        | T K D →       | D1 D2 D3 D4 D5  | 1            | 1.4     | 1   |
| Systems Analysis |               |                 |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Make sustainment infrastructure — maintainers, technical data, depot capacity — a contractual deliverable that fields with each lot of aircraft, not a follow-on cost.
- Secure government rights to the technical data needed to write procedures, qualify depots, and compete the work, so the fleet is not locked to a single source.

IN OPERATION / AFTER

- Audit the mission-capable rate against repair-queue depth and depot turnaround, and act when the backlog signals a structural, not transient, shortfall.
- Track recurring GAO-style diagnoses to closure, refusing to let procurement outpace the sustainment fixes named year after year.



- Gate fielding on demonstrated sustainment capacity, not just aircraft delivery, so readiness is engineered alongside the platform.
- Monitor contractor dependency as a wartime risk, and sustain a government maintainer pipeline and depot capacity that can hold under pressure.

### REFLECTION QUESTIONS

1. Pick a current technology platform in your domain. Estimate the capability infrastructure that must field with it. What happens if half of that infrastructure is years behind the hardware?
2. The F-35 program treated maintainer training as follow-on cost. Design a fielding gate that would prevent that decision being available to a future program manager.
3. Much of the data needed to repair the F-35 stayed controlled by the contractor, foreclosing the government's ability to compete or qualify the work. For a platform in your domain, what data rights would you make a contractual deliverable up front, and how would you verify they were actually delivered?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

# Military Fratricide — Desert Storm to Afghanistan

**T** Training Gap   **H** Human-System Interface   **K** Knowledge & Institutional Memory

**IMPACT** 24% of U.S. battle deaths in Desert Storm from friendly fire (35 of the 146 killed in action; 35 of 148 fatalities in the OTA's tally — 24% either way) — a rate later reviews found consistent with past wars, not exceptional

## IN BRIEF

Friendly fire killed an unusual share of coalition forces in the 1991 Gulf War: 35 of 146 U.S. combat deaths (24%) and 72 of 467 wounded (15%). (The often-cited “2% historical baseline” from Shrader’s 1982 study is contested — Shrader himself cautioned that the evidence was inadequate for quantitative conclusions, the Office of Technology Assessment’s 1993 review of World War II and Vietnam casualty surveys concluded that “15 to 20 percent may be the norm, not the exception,” and a 1994 Naval Institute analysis treated the 2% norm as a myth the press had seized on.) Post-war reviews blamed the chaos of combat, weak situational awareness and fire-control discipline, and combat-identification failures — and noted the military lacked a shared record to even study its own pattern. Fratricide is the failure of several systems to integrate; despite a generation of combat-ID investment, recurrences in Afghanistan and after show the rate never fell to a confidently low level, because integration is the hardest property to engineer by program.

## BACKGROUND

Friendly fire is as old as war, but its true rate is hard to pin down. The most-cited estimate, from Charles Shrader’s 1982 study *Amicicide*, put it under 2% of battle casualties — a figure Shrader himself hedged, warning that the available evidence was inadequate to support quantitative conclusions, and one later analysts challenged as far too low. The Office of Technology Assessment’s 1993 review of long-extant World War II and Vietnam casualty surveys concluded that estimates of 2 percent are unrealistic and that “15 to 20 percent may be the norm, not the exception”; a 1994 U.S. Naval Institute Proceedings analysis by Charles Hawkins treated the 2 percent “historical norm” as a misperception the press had seized on. The disputed baseline mattered because it became the yardstick against which a modern war would be measured — and a yardstick set too low makes any later rate look like a catastrophe.<sup>o</sup>

## WHAT HAPPENED

The 1991 Gulf War made the question grim and concrete: of 146 U.S. service members killed in action, 35 — about 24% — died by friendly fire, and 72 of 467 wounded (15%) were hit by their

own side. (The OTA counted 615 U.S. battle casualties and 148 fatalities rather than 146; the fratricide share is 24 percent on either denominator.)<sup>1</sup> An A-10 strike on U.S. LAV-25s near Khafji killed seven Marines; an A-10 attack on British Warrior vehicles killed nine — each an aircraft firing on friendly vehicles it had failed to identify, the recurring shape of the problem. Whatever the true baseline, a quarter of American combat deaths from one's own forces could not be waved off as the ordinary friction of battle.<sup>2</sup>

#### THE INVESTIGATION


Post-war reviews converged on familiar causes: the chaos of combat, inadequate situational awareness, weak adherence to fire-control measures, and combat-identification failures.<sup>3</sup> One finding cut deeper — broad-based data on fratricide rates did not exist. The military had no comprehensive, shared record of fratricide incidents, so it could not study its own pattern, separate training failures from doctrine or equipment, or tell whether a fix was working; the Army began requiring Fratricide Incident Reports at its training centers only after the Gulf War. Without a common database, every incident was investigated in isolation and the aggregate signal that might have driven reform never formed. The capability to **learn** was itself missing, a second-order gap beneath the first.<sup>4</sup>

#### THE CAPABILITY GAP

Fratricide is not one problem but the failure of several systems to integrate — situational awareness, fire-control discipline, combat identification, joint coordination — each the subject of dedicated programs, yet the rate never fell to anything negligible. In Afghanistan a satellite-guided strike went wrong after a controller's device reset its coordinates, and in 2014 a B-1's targeting pod could not detect the infrared strobes marking U.S. troops, killing five. Each contributor was worked on; the integration **across** them — which is what actually keeps friendly forces from killing each other — is the hardest thing to engineer by program, because no single office owns it and no single procurement can deliver it — and it is where the capability kept failing.<sup>5</sup>

#### AFTERMATH & REFORM

The response was a generation of combat-identification investment: better IFF systems, blue-force-tracking networks, and changes to fire-control doctrine and training. The improvements were real, yet fratricide never dropped to a confidently low, stable rate, and even measuring it remained contested.<sup>6</sup> That is the lesson: where capability is an emergent property of many systems working together, no single program closes the gap — improving each contributor in isolation still leaves the integration between them unaddressed — and progress has to be measured against an honest baseline rather than a convenient one.



FRIENDLY-FIRE RATE, U.S. CASUALTIES, DESERT STORM VS. HISTORICAL

*Broad-based data on fratricide rates are not available.*

U.S. CONGRESS, OFFICE OF TECHNOLOGY ASSESSMENT, WHO GOES THERE: FRIEND OR FOE? OTA-ISC-537, JUNE 1993

REFERENCES

1. C. R. Shrader, Amicide: The Problem of Friendly Fire in Modern War (U.S. Army Combat Studies Institute, 1982) — the under-2% estimate from 269 incidents, a baseline later challenged as too low (OTA 1993: 15 to 20% may be the norm).

2. “Friendly Fire: Facts, Myths and Misperceptions,” U.S. Naval Institute Proceedings (June 1994) — Desert Storm: 35 of 146 KIA (24%) and 72 of 467 wounded (15%) by friendly fire; critique of the 2% “historical norm.”

3. Khafji and Warrior fratricide incidents (Feb. 1991) — see USNI Proceedings (1994) and Time, “They Didn’t Have to Die”. (Per-incident casualty figures vary across sources; see AUDIT.)

4. U.S. GAO, Operation Desert Storm fratricide investigations — Apache incident (OSI-93-4) and Army fratricide investigation (OSI-95-10) — causes and the Army’s reviews.

5. U.S. Congress, Office of Technology Assessment, Who Goes There: Friend or Foe?, OTA-ISC-537 (June 1993) — the Desert Storm counts (615 battle casualties, 148 fatalities, 35 by friendly fire), the “15 to 20 percent may be the norm” reassessment, the absence of broad-based fratricide data, and the NTC observer distribution of causes (identification failure, command and control, planning, communication and navigation).

6. Later-conflict recurrences (2001 coordinate-reset strike; 2014 B-1 infrared-strobe strike) and S. Snook, Friendly Fire (Princeton Univ. Press, 2000) on the 1994 Black Hawk shootdown — the systems-integration archetype.

7. “IFF Update: Stalled Again,” U.S. Naval Institute Proceedings (June 1994) — the combat-identification and identification-friend-or-foe programs pursued after Desert Storm, and how slowly they matured.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Fratricide is a multi-decade capability problem that resists single- intervention solutions. Each of the contributing causes — situational awareness, fire-control discipline, combat identification, joint coordination — has been the subject of dedicated programs. The rate persists because the integration across the contributors is itself the capability, and integration is the hardest property to engineer by program.

LENS APPROACH

LENS uses fratricide in LEN 5 to teach systems-of-systems capability analysis and in LEN 2 to introduce combat identification as a Human-AI Teaming problem (operators rely on automated IFF systems whose limitations are not in their training). LEN 8 examines why decades of awareness have not produced a sustained reduction.

| DOMAIN           | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|------------------|---------------|------------------|--------------|---------|-----|
| DEFENSE          | → T H K       | → D1 D2 D3 D4 D5 | 1            | 3.3     | 1   |
| Systems Analysis |               |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat combat identification as a systems-of-systems requirement with a single owner for the integration, not as separate IFF, situational-awareness, and fire-control programs.
- ▶ Build a comprehensive, shared incident database from the outset so the pattern across events can be studied and fixes tested against it.
- ▶ Design fire-control discipline and identification into joint coordination, so the integration that prevents fratricide is engineered rather than left emergent.

IN OPERATION / AFTER

- ▶ Audit the friendly-fire rate against an honest baseline, distinguishing a temporary dip from a structural reduction before crediting any program.
- ▶ Monitor each contributor and the integration across them, since improving components in isolation can leave the joint failure mode untouched.
- ▶ Sustain the learning channel so later-conflict recurrences feed back into doctrine and training rather than being investigated in isolation.

## REFLECTION QUESTIONS

1. Identify a capability problem in your domain that has been the subject of repeated programs without sustained improvement. What is the integration gap that the programs do not address?
2. Design a measurement system that would distinguish a temporary improvement in fratricide rate from a structural one.
3. The military lacked a shared incident database, so it could not study its own pattern. Identify a recurring harm in your domain that is investigated case-by-case but never aggregated, and design the shared record that would let the pattern become visible.

**WHO BUILDS THIS** learning science & instructional design · human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs · knowledge bases · data pipelines

# Operation Eagle Claw

**T** Training Gap **K** Knowledge & Institutional Memory

**IMPACT** 8 servicemembers killed in Iran; mission to rescue 53 American hostages aborted; catalyst for the founding of U.S. Special Operations Command

## IN BRIEF

The April 1980 mission to rescue fifty-three American hostages held at the U.S. embassy in Tehran was aborted at a desert staging point when three of eight helicopters were lost to mechanical and weather problems. During the withdrawal a helicopter collided with a refueling C-130, killing eight servicemembers; the hostages remained captive for another nine months. The Holloway Commission found the operation had been assembled ad hoc: each service contributed its own units, equipment, and command relationships, the aircrews had not trained together, and there was no standing joint special-operations command to own the mission. The capability had to be improvised, and the improvisation failed. Eagle Claw catalyzed the creation of JSOC, the Goldwater-Nichols reforms, and ultimately U.S. Special Operations Command.

## BACKGROUND

In November 1979, Iranian revolutionaries seized the U.S. embassy in Tehran and took more than sixty Americans hostage; fifty-three remained by April 1980. After months of failed diplomacy, the Carter administration authorized a military rescue. No standing joint special-operations command existed to plan or execute it; the force had to be assembled from units drawn separately from the Army, Navy, Marines, and Air Force — each bringing its own equipment, procedures, and chain of command to a mission that demanded they act as one, with no institution whose job it was to make them cohere before the night of the raid.<sup>o</sup>

## WHAT HAPPENED

On 24 April 1980, the mission staged at a remote site code-named Desert One. Three of the eight RH-53 helicopters were disabled by a dust storm and mechanical failure, dropping the force below the minimum needed; the commander aborted. During the withdrawal a helicopter collided with a C-130 tanker, and the resulting fire killed eight servicemembers. The mission failed before reaching Tehran — undone not by enemy action but by attrition and a chaotic withdrawal among aircraft and crews that had never operated together, the predictable cost of integration improvised under pressure rather than built and rehearsed in advance.<sup>1</sup>

## THE INVESTIGATION

The Holloway Special Operations Review Group named “the ad hoc nature of the organization and planning” as one of two fundamental concerns: each service contributed its own units, equipment, command relationships, and communications; the aircrews had not


trained together as a unit; the helicopters were RH-53Ds drawn from the Navy's mine-countermeasures force, their first crews Navy minesweeping pilots, later replaced by Marine crews who flew the mission in an airframe and a profile — long-range night penetration — neither unit had been built for. There was no standing organization to own the mission end to end — no single authority responsible for the force's training, equipment fit, and command architecture as a whole, so each gap was someone's problem in part and no one's in full.<sup>2</sup>

#### THE CAPABILITY GAP

Each service was competent inside its own boundary. The capability that did not exist was the integration across them — a joint command, a common communications architecture, and a force that had trained together. That integration had to be improvised for a single high-stakes mission, and the improvisation could not hold — because cross-service cohesion is not summoned on demand but accrued through standing structure and repeated joint training, neither of which existed when the order came.<sup>3</sup>

#### AFTERMATH & REFORM

Eagle Claw produced the Holloway Commission, the creation of the Joint Special Operations Command in 1980, and — alongside the 1983 Beirut bombing and the Grenada invasion — the impetus for the Goldwater-Nichols Act of 1986 and the Nunn-Cohen Amendment establishing U.S. Special Operations Command in 1987. The reform built the institution the mission had needed and not had — converting a capability that had been improvised once and failed into a standing command with its own forces, training, and authority, so the next mission would inherit cohesion rather than assemble it from scratch.<sup>4</sup>



*five organizations · no joint command · no shared training*

#### THE ABSENCE OF AN INSTITUTION THAT COULD OWN THE MISSION

*By not utilizing an existing JTF organization, the Joint Chiefs of Staff had to start literally from the beginning to establish a JTF, create an organization, provide a staff, develop a plan, select the units and train the force.*

HOLLOWAY SPECIAL OPERATIONS REVIEW GROUP, RESCUE MISSION REPORT, 1980

#### REFERENCES

1. Locher, J. (2002), Victory on the Potomac — the reform arc from Desert One to Goldwater-Nichols.
2. Holloway Commission (1980) — the Desert One sequence and the helicopter-C-130 collision.
3. Goldwater-Nichols Department of Defense Reorganization Act of 1986, Pub. L. 99-433.
4. Nunn-Cohen Amendment (1986), establishing U.S. Special Operations Command (1987).

3. Holloway Special Operations Review Group, Rescue Mission Report (1980) — 23 issues affecting the mission, of which command and control and the ad hoc nature of the organization and planning are the two the Review Group treated as fundamental; the helicopter and aircrew selection is among them.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Eagle Claw is the canonical case in U.S. defense for the absence of institutionalized cross-service capability. Each service was competent inside its own boundary; the integration across them did not exist as a deliverable — until the reform that followed built the institution the mission had needed and not had.

LENS APPROACH

LENS uses Eagle Claw in LEN 5 as a worked case for cross- organizational capability requirements and in LEN 8 for the institutional reform that followed; it pairs with INPO (Case 175) as the defense analog of failure-driven institution building. The seven-year arc to Goldwater-Nichols (1986) and USSOCOM (1987) is itself a measurement: the institutional response time when the fix requires statutory action.

| DOMAIN           | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|------------------|---------------|------------------|--------------|---------|-----|
| DEFENSE          | → T K         | → D1 D2 D3 D4 D5 | 1            | 1.3     | 1   |
| Systems Analysis |               |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Stand up a single command that owns the cross-organizational mission end to end — its force, equipment fit, communications, and training — rather than assembling it per mission.

▶ Require the contributing units to train together as one force on a common communications architecture before they are committed.

▶ Select the unit and the equipment for fitness to the actual mission, not for whichever squadron is most expedient to task.

IN OPERATION / AFTER

- ▶ Audit whether a cross-organizational capability that exists on paper has ever actually operated as one force, and treat the absence of joint reps as an unfilled gap.

▶ Sustain the standing institution and its training pipeline so the next mission inherits cohesion rather than improvising it.

▶ Monitor for the recurrence of ad-hoc assembly, since the conditions that produced Desert One reappear whenever no single authority owns the joint mission.

REFLECTION QUESTIONS

1. Where in your domain does a cross-organizational capability exist on paper but not in practice? What would force its institutionalization?
2. Eagle Claw produced USSOCOM and Goldwater-Nichols six years later. Sketch the institutional design that an equivalent failure in your domain would force into existence.
3. The Holloway Commission named the ad hoc nature of the organization and planning as the diagnosis. What standing capability — institution, command, training pipeline — does your domain currently lack that an Eagle-Claw-class failure would force into existence?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · knowledge bases · data pipelines



# Patriot Missile / Dhahran

**D** Designed Out   **H** Human–System Interface   **K** Knowledge & Institutional Memory

**IMPACT** 28 U.S. soldiers killed in their barracks; roughly 100 wounded

## IN BRIEF

On 25 February 1991 a Scud missile struck a U.S. barracks in Dhahran, Saudi Arabia, killing 28 soldiers and wounding about a hundred; the Patriot battery defending the area never engaged it. Designed for short engagements against Soviet aircraft in Europe, the Patriot tracked time in a register whose tiny rounding error grew with every hour of uptime. After about a hundred hours of continuous operation the Gulf War demanded, the radar's range gate was off by a third of a second — enough to look in the wrong place. Israeli operators had flagged the drift two weeks earlier, and a patch reached Dhahran the day after the strike. The capability to hold accuracy under sustained use was assumed away at design time, and no one carried that assumption forward when the mission changed.

## BACKGROUND

The Patriot was built to defend Western Europe against Soviet aircraft in engagements of minutes — switched on, fired, switched off, then moved before it could be targeted. In the 1991 Gulf War it was doing something else: defending fixed sites in Saudi Arabia against ballistic missiles, in batteries left running continuously for more than a hundred hours because the threat came without warning and could not be scheduled. The mismatch between the original concept of operations and the new use was real, and invisible to the operators, because no one had told them the machine had been built around an assumption they were now violating.<sup>0</sup>

## WHAT HAPPENED

The system tracked time in a 24-bit fixed-point register, and a tiny rounding error — invisible in any short engagement — accumulated with every hour of uptime. After about a hundred hours the timing was off by roughly a third of a second, which seems negligible until it is multiplied by a ballistic missile's speed: enough for the radar's range gate to look in the wrong place and reject the real track as noise. On 25 February 1991 an incoming Scud arrived where the Patriot was no longer searching, passed unengaged, and struck a barracks in Dhahran, killing twenty-eight soldiers — thirteen of them from the 14th Quartermaster Detachment — and wounding ninety-nine.<sup>1</sup>

## THE INVESTIGATION

The drift was not unknown: Israeli operators had flagged it two weeks earlier from their own sustained use, and engineers had a patch already in hand — which reached Dhahran the day after the strike, too late to matter.<sup>2</sup> The only field warning said “very long” run times could shift the range gate, never defining “very long” and never naming a reboot interval, so a crew


could heed it to the letter and still drift into the danger band. The General Accounting Office found the Army had simply presumed no one would run a battery continuously for so long, and so never treated the accumulating error as a hazard worth specifying or warning against.<sup>3</sup>

#### THE CAPABILITY GAP

The capability to hold accuracy under sustained operation was never built in, because the original concept of operations did not require it — a defensible omission for the system as first imagined, fighting in short bursts and moving on. The failure came when the concept changed and the assumption did not travel with it. Nothing carried the design's hidden premise forward to the soldiers in Dhahran; no artifact, briefing, or warning told them “this system was built assuming you would turn it off.” The system did not malfunction — it did exactly what it was built to do; the transition between operating contexts did, with nothing in place to catch the broken premise.<sup>4</sup>

#### AFTERMATH & REFORM

The patch was distributed and concrete reboot intervals defined in place of the vague advisory, and the defect — a fixed-point truncation error growing without bound as uptime increased — became a staple of numerical-analysis and software-engineering teaching.<sup>5</sup> The harder lesson is the one this chapter turns on: a capability quietly assumed away at design time does not announce itself when the context shifts. It waits, intact and invisible, until the day the assumption is wrong — and by then the people relying on the system have no way to know the ground has moved beneath them, because the premise was never written where they could read it.



*Army officials presumed that the users would not continuously run the batteries for such extended periods.*

GAO/IMTEC-92-26, 1992

#### REFERENCES

1. U.S. General Accounting Office, Patriot Missile Defense: Software Problem Led to System Failure at Dhahran, Saudi Arabia, GAO/IMTEC-92-26 (1992) — the design-context mismatch and continuous-operation use.
2. R. Skeel, “Roundoff Error and the Patriot Missile,” SIAM News 25(4): 11 (1992) — the 24-bit fixed-point time truncation and the 0.34-second range-gate drift after 100 hours.
3. GAO/IMTEC-92-26 (1992) — the vague “very long run time” advisory and the Army’s presumption that batteries would not be run continuously for such periods (quoted).
4. M. Barr / Barr Group, “An Update on the Patriot Missile Software Problem” — an engineering post-mortem of the accumulated-truncation defect.
5. GAO/IMTEC-92-26 (1992) and Skeel (1992) — the design assumption that failed to travel across the change in operational context, and the case’s teaching legacy.

3. GAO/IMTEC-92-26 (1992) — the prior Israeli warning, the software patch arriving in Dhahran on 26 February (the day after the strike), and the 28 killed.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Patriot is the canonical Designed-Out case from defense. The capability to maintain accuracy under sustained operation was omitted because the original concept of operations did not anticipate it. When the concept changed, the assumption did not travel. The system did not fail; the transition from one operational context to another did, and there was no capability infrastructure to catch it.

LENS APPROACH

LENS treats Patriot as the textbook example of **Capability Degradation Under System Change**. LEN 5 methods require operators of any transitioning system to have current capability-relevant documentation specifying what assumptions of the original design have changed. LEN 8 examines how organizational knowledge about design constraints travels from engineering to operations — and why, here, it arrived a day late.

| DOMAIN           | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|------------------|---------------|------------------|--------------|---------|-----|
| DEFENSE          | → D H K       | → D1 D2 D3 D4 D5 | 1            | 1.2     | 1   |
| Systems Analysis |               |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Make every design assumption the safety case rests on — here, that the system would be cycled off — an explicit, recorded requirement so a later change of use cannot silently violate it.
- ▶ Bound numerical error over the full intended operating envelope, including run times far beyond the original concept, rather than only the durations first imagined.
- ▶ Design the system to detect and bound its own accumulating drift, degrading or warning before the error reaches a magnitude that defeats the mission.

IN OPERATION / AFTER

- ▶ Issue concrete, testable operating limits — exact reboot intervals, not “very long” — derived from the design data and propagated to every crew.
- ▶ Establish a path for field reports of anomalous behavior, like the Israeli warning, to reach engineering and trigger fleet-wide action faster than a patch can lose a race with the threat.
- ▶ Require a capability-transition review whenever a system is redeployed to a new concept of operations, surfacing which original assumptions the new mission breaks.

REFLECTION QUESTIONS

1. The Patriot’s design assumed short engagements. What assumption in a current system you work with would become lethal if the operational context shifted, and how would operators learn of the shift?
2. Construct the capability-transition artifact that should have accompanied the redeployment of Patriot batteries from Europe to Saudi Arabia. What would it have said, and who would have signed it?
3. The field advisory said to reboot after “very long” run times without defining the term. Rewrite a vague operational instruction in your domain into a specific, testable limit, and name the design data the limit must be derived from.

**WHO BUILDS THIS** systems & human-factors engineering · human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · usability testing · cognitive walkthroughs · knowledge bases · data pipelines

# USS Vincennes & Iran Air Flight 655

**H** Human-System Interface **T** Training Gap

**IMPACT** 290 civilians killed — one of the deadliest shootdowns of a commercial airliner by a military force; precipitating case for Aegis CIC display and decision-aid doctrine reform, and a central reference in the human-AI teaming literature

## IN BRIEF

On 3 July 1988, during a surface skirmish with Iranian gunboats in the Persian Gulf, the USS Vincennes shot down Iran Air Flight 655 — a civilian Airbus A300 climbing on a scheduled route — killing all 290 aboard, one of the deadliest airliner shootdowns by a military force. The Aegis combat system worked: its radar correctly showed the aircraft ascending on the published civilian-traffic corridor. Yet operators in the Combat Information Center, primed by a simultaneous surface fight to expect a hostile inbound, tagged the contact as a military F-14 — from an IFF Mode II squawk operators reported on their consoles but the ship's own system never held, the airliner squawking only civilian Mode III — and, on reports of a descent its own tapes contradicted, fired. The Fogarty Report (DoD, 1988) found the Aegis system functioned as designed and judged that stress, task fixation, and unconscious distortion of data — “scenario fulfillment” — may have played a major role. The 1992 Newsweek “Sea of Lies” investigation and a 2018 U.S. Naval Institute Proceedings retrospective reopened both the operational record and the interface-design lessons. Vincennes is the book's foundational human-AI-teaming case: the most advanced surface combatant afloat failed because its CIC interface and the operational framing made the correct reading possible in principle and unsustainable in practice under combat stress.

## BACKGROUND

In the closing weeks of the Iran-Iraq War, U.S. warships patrolled a tense Persian Gulf where civilian airliners and hostile military aircraft shared the same crowded sky. The USS Vincennes, a Ticonderoga-class cruiser with the Navy's most advanced Aegis combat system, was in a surface fight with Iranian gunboats — itself a contested engagement that the later Newsweek “Sea of Lies” investigation argued had taken the ship into Iranian territorial waters — when an aircraft took off from Bandar Abbas, an airfield used by both civil and military traffic.<sup>o</sup> The dual-use field meant a single track could be either threat or scheduled flight, so the crew's reading of the contact carried the whole burden of identification — and the Combat Information Center was already managing divided attention against the surface engagement and against an interface whose track display did not directly show altitude trend, the kind of conditions under which a wrong call becomes far easier to make than a right one.

## WHAT HAPPENED

The aircraft was Iran Air Flight 655, a civilian Airbus A300 climbing through 12,000 feet on its scheduled route to Dubai. The Vincennes' crew identified it as a descending, hostile F-14 and fired two SM-2 surface-to-air missiles; all 290 people aboard, including 66 children, were killed — one of the deadliest shootdowns of a commercial airliner by a military force.<sup>1</sup> An ascending airliner and a diving attack jet are opposite behaviors, yet the crew converged on the second while the Aegis radar reported the first. The engagement consumed only seven minutes from radar detection to launch, collapsing identification, IFF interrogation, and decision into a window too narrow for anyone to slow down and reconcile the contradiction.

## THE INVESTIGATION

The Aegis SPY-1A radar had not malfunctioned: it correctly showed the aircraft ascending. The IFF interrogation system added a structural confusion of its own: the Vincennes' recorded data held only the airliner's Mode III civilian transponder, while operators reported a momentary Mode II-1100 on their consoles that the ship's system never held — a source the investigation never established.<sup>2</sup> The Fogarty Report (Rear Adm. William Fogarty, DoD, August 1988) attributed the tragedy to human error under extreme stress — confirmation bias and the "stress, task fixation, and unconscious distortion of data" — naming "scenario fulfillment" as the psychological mechanism by which operators read every indication through the lens of a presumed hostile attack. Reports of "descending altitude" reached the CO while the ship's own command-and-decision system showed a continuous ascent, and no one in the CIC reconciled the two,<sup>3</sup> a shared error that is more damning than a single misread because it shows the interface offered no cross-check strong enough to break a wrong reading once the team had settled into it. The framing arrived before the data did, so each new return was fitted to the expected threat rather than weighed against it.<sup>4</sup>


## THE CAPABILITY GAP

The system did not lie, and the operators did not act in bad faith. The interface and the operational framing combined to make a particular misreading not merely possible but likely. Correct performance was possible in principle and unsustainable in practice — and the gap between those two is the engineering problem. An Aegis large screen display that could not show altitude at all, that gave equal weight to civilian and military IFF returns under mode confusion, and that offered no mandatory disconfirmation step before weapons release, had not been designed for the stress it would actually meet. The burden of overriding a presumed hostile attack was left entirely to the operator's discipline, exactly when combat had stripped that discipline of the time and calm it needed to work.<sup>5</sup>

## AFTERMATH & REFORM

The Navy's post-incident response included revised tactical doctrine for civilian-traffic de-confliction in the Gulf, a fleet-wide identification matrix for dense air-traffic environments, a review of Aegis IFF operator training on the selectable RCI modes, and a redesign of the large screen display so it could show altitude directly. The case has remained a standing reference in the human-AI-teaming literature: a 2018 US Naval Proceedings retrospective placed it at the heart of how operators should be teamed with automated decision aids under stress, and the situation-awareness, naturalistic-decision-making, and automation-trust literatures (Klein, Endsley, Cummings) treat it as the canonical worked example.<sup>6</sup> Its lesson is that interface design is a capability deliverable, not an aesthetic one — and that

the most advanced system afloat is only as good as the human reading it under fire. The retrospective reframed the loss not as a one-off error but as a predictable outcome of teaming a person with a decision aid that displayed truth without defending it, a pattern that recurs wherever automation is fast and the human is the last check.



*In combat, technology should create time for the decision maker, not simply provide information.*

ANTHONY TINGLE, "THE HUMAN-MACHINE TEAM FAILED VINCENNES," U.S. NAVAL INSTITUTE PROCEEDINGS, JULY 2018

## REFERENCES

1. J. Barry & R. Charles, "Sea of Lies," Newsweek (July 13, 1992) — contemporaneous reinvestigation of the operational record, including the disputed account of Vincennes' position relative to Iranian territorial waters.
2. Rear Adm. W. Fogarty, Formal Investigation into the Circumstances Surrounding the Downing of Iran Air Flight 655 (U.S. Navy, August 1988) — the engagement, the IFF mode-confusion findings, and the 290 deaths.
3. Fogarty report (1988) — the Aegis SPY-1A radar functioned; the aircraft was ascending while the crew perceived a descent; "scenario fulfillment" as the psychological mechanism.
4. Fogarty report (1988) — "human error under extreme stress," confirmation bias, and "unconscious distortion of data" (quoted); shared-error finding across CIC operators.
5. LTC A. Tingle, USA, "The Human-Machine Team Failed Vincennes", U.S. Naval Institute Proceedings 144/7 (July 2018) — the human-AI-teaming reframing and the standing case-study role.
6. M. L. Cummings, "Automation Bias in Intelligent Time Critical Decision Support Systems," AIAA 1st Intelligent Systems Technical Conference (2004), doi:10.2514/6.2004-6313 — interface design and automation under time pressure as the engineering frame for the case.
7. G. Klein, Sources of Power (1998); M. Endsley, "Toward a Theory of Situation Awareness" (1995) — the naturalistic-decision-making and situation-awareness literatures that treat Vincennes as the canonical worked example.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Vincennes is the canonical case for human-AI teaming under stress. The system did not lie. The operators did not act in bad faith. The interface and the operational framing combined to make a particular misreading not just possible but likely. Correct performance was possible in principle and unsustainable in practice — and that gap is the engineering problem.

### LENS APPROACH

Vincennes is the foundational mode-and-state-transparency case under combat stress (induced 3.3; LENS D3/PT6 human-AI teaming). LENS uses it in Domain 2 (Iterative Development; LEO-2) for interface-as-deliverable: a CIC display tested against the worst case it would actually meet — confirmation bias during a simultaneous surface engagement — not against the calm of acceptance testing. LENS uses it in Domain 3 (Human-System Collaboration; LEO-3) for the human-AI teaming problem itself: a decision aid that displays truth without defending it predictably fails wherever automation is fast and the human is the last check.

The case sits at the heart of the program’s argument that interface design is a capability deliverable, not an aesthetic one. Pair with EHR/CPOE (Case 2) at the interface-built-to-wrong-specification layer.

| DOMAIN                     | FAILURE MODES                   | LENS COMPETENCY                                                                   | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|---------------------------------|-----------------------------------------------------------------------------------|--------------|---------|-----|
| DEFENSE                    | → <span>H</span> <span>T</span> | → <span>D1</span> <span>D2</span> <span>D3</span> <span>D4</span> <span>D5</span> | 6            | 3.3     | 3   |
| Human-System Collaboration |                                 |                                                                                   |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>▶ Design the contact display to make ascent-versus-descent unmistakable at a glance, so the cue that contradicts a presumed threat cannot be passed over under time pressure.</li><li>▶ Test the interface against the worst case it will actually meet — confirmation bias during a simultaneous surface engagement — not against the calm conditions of acceptance testing.</li><li>▶ Build a mandatory disconfirmation step into the engagement sequence so identifying a hostile track requires actively ruling out the civilian one.</li></ul> | <ul style="list-style-type: none"><li>▶ Audit live engagements and exercises for cases where the crew’s reading diverged from sensor data, treating each as a near-miss that exposes an interface gap.</li><li>▶ Train and drill operators specifically on the stress regime — divided attention, presumed-hostile framing — rather than only on nominal track identification.</li><li>▶ Track whether decision aids are trusted past their evidence, and feed that signal back into interface and procedure revisions.</li></ul> |

REFLECTION QUESTIONS

1. Identify a high-stress interface in your domain. What framing arrives before the data, and how does it shape what operators see?
2. Vincennes’ operators acted under tunnel vision. Design the procedural intervention that would have forced one of the five to call out the contradiction.
3. The Aegis radar reported the truth and the crew still read its opposite. What would an interface have to make impossible — not merely visible — to keep a pre-formed expectation from overriding correct data under stress?

**WHO BUILDS THIS** human factors & interaction design · learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · scenario simulation · competency models

# Stanislav Petrov / 1983 False Alert

**H** Human–System Interface **T** Training Gap

**IMPACT** Soviet early-warning system reported five incoming U.S. ICBMs; Lt. Col. Petrov correctly assessed the signal as false; retaliation averted

## IN BRIEF

On the night of 26 September 1983 the Soviet Oko early-warning system reported five U.S. intercontinental missiles inbound. The duty officer, Lt. Col. Stanislav Petrov, judged it a false alarm — a real first strike would involve hundreds of missiles, not five — and reported it as such rather than up the launch–decision chain. He was right: sunlight glinting off high-altitude clouds had fooled the satellite’s infrared sensors. Petrov’s judgment is widely credited with averting nuclear war. The case is the book’s strongest positive evidence for human-in-the-loop capability: the automation failed in a mode its designers never anticipated, and a person with contextual judgment and the latitude to override it was the recoverability the system had. Keeping humans in some loops is not nostalgia — it is capability engineering.

## BACKGROUND

Soviet nuclear command rested partly on Oko, a satellite early-warning system meant to detect U.S. missile launches and feed a launch-on-warning posture. The duty officer at the Serpukhov-15 bunker was responsible for verifying any alert and passing it up the chain.<sup>o</sup> A launch-on-warning posture compresses the time between detection and decision to almost nothing, so the duty officer’s verification was not a formality but the single human checkpoint standing between a satellite reading and the machinery of retaliation.

## WHAT HAPPENED

On the night of 26 September 1983 Oko reported a U.S. intercontinental ballistic missile launch — then, minutes later, four more, for five in all. Lt. Col. Stanislav Petrov, the duty officer, assessed the signal as a false alarm — a genuine first strike, he reasoned, would involve hundreds of missiles, not five — and reported it as such. He was correct.<sup>1</sup> The reasoning that saved the situation came from outside the system entirely: Oko could report what it saw, but it could not weigh five launches against the doctrine of a real first strike, and that mismatch between the alert and the strategic picture was exactly the judgment the machine had no way to make.

## THE INVESTIGATION

The cause was an unanticipated automation failure: sunlight reflecting off high-altitude clouds at a particular geometry had fooled the Oko satellite’s infrared sensors into reading launches that were not there.<sup>2</sup> It was a failure of the rarest kind — a benign natural phenomenon the sensors had never been designed to discount — which is precisely why no



automated check existed to catch it. Later investigation identified the satellite-geometry failure mode and modified the algorithm; the scenario is now permanently archived in early-warning training as the canonical false positive, a permanent reminder that the system's worst error was one it could not have flagged for itself.<sup>3</sup>

#### THE CAPABILITY GAP

The automation had failed in a mode its designers had not imagined, and no automated check could catch it. What caught it was a human with the contextual judgment to weigh the alert against what a real attack would look like, and the institutional latitude to override the system rather than simply relay it. Petrov's "funny feeling" was a career's worth of judgment doing the work the automation could not — and the chain of command had left him room to act on it rather than forcing him to pass the alert upward unfiltered, so the recoverability lived as much in the authority structure as in the man.<sup>4</sup>

#### AFTERMATH & REFORM

Petrov's decision is widely credited with averting nuclear war, and the case is permanently studied in command-and-control training.<sup>5</sup> Preserving the episode in training is itself a design choice: it keeps the failure mode and the human override alive as institutional memory, so the lesson does not erode the way the original sensor's blind spot had. It is the book's strongest argument that retaining a human in some loops is capability engineering, not nostalgia: the human role is the recoverability of an automation failure the designers did not anticipate — which is exactly why fully automating a strategic decision removes the one element that can catch the unimagined error.



*I had a funny feeling in my gut.*

STANISLAV PETROV, QUOTED IN THE WASHINGTON POST, 1999

#### REFERENCES

1. D. Hoffman, *The Dead Hand: The Cold War Arms Race and Its Dangerous Legacy* (2009) — the Oco system and Soviet launch-on-warning posture.
2. Accounts of the 26 Sept. 1983 incident (Hoffman; Petrov interviews) — the five-missile alert and Petrov's assessment.
3. Investigations of the Oco incident (incl. V. Aksenov, 2004) — sunlight-on-clouds satellite geometry as the false-positive cause.
4. The subsequent Oco algorithm modification and the incident's use in early-warning training.
5. S. Petrov, interview, *The Washington Post* (1999) — "a funny feeling in my gut" (quoted).
6. S. Sagan, *The Limits of Safety* (1993); M. L. Cummings (2017) — automation in critical decisions and human-in-the-loop.

#### THE LEARNING ENGINEERING LENS

LE INSIGHT

Petrov is the canonical positive case for human-in-the-loop nuclear command-and-control. The automation failed in a mode its designers had not anticipated. The human’s contextual judgment was the backstop that allowed the failure to be recovered before it became catastrophic. The case is the strongest argument in this book that the design choice to retain humans in some loops is not nostalgia but capability engineering.

LENS APPROACH

LENS uses Petrov in LEN 2 as the positive Human-AI Teaming case at the highest possible stakes: the human role is the recoverability of an automation failure that the designers did not anticipate. The case anchors arguments against full-automation of strategic decisions.

| DOMAIN                     | FAILURE MODES | LENS COMPETENCY        | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|---------------|------------------------|--------------|---------|-----|
| DEFENSE                    | → [H][T]      | → [D1][D2][D3][D4][D5] | 6            | 3.4     | 3   |
| Human-System Collaboration |               |                        |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Where automation feeds an irreversible decision, design the human role to be a genuine override — with the context, judgment cues, and authority to refuse a faulty alert, not merely relay it.</li><li>Give operators a strategic picture against which to weigh an alert, so an anomalous signal can be tested against what a real event would look like.</li><li>Treat unimagined failure modes as a design assumption: keep a human in the loop precisely where no automated check can cover the unanticipated.</li></ul> | <ul style="list-style-type: none"><li>When a failure mode surfaces, modify the algorithm and archive the scenario in training so the lesson persists as institutional memory.</li><li>Audit whether human-in-the-loop roles still carry real override authority over time, or have decayed into ceremonial relays.</li><li>Preserve and refresh the contextual judgment override depends on through deliberate training, so the capability does not erode between rare events.</li></ul> |

REFLECTION QUESTIONS

1. Identify an automated system in your domain where retaining a human-in-the-loop is genuinely capability-engineering rather than ceremonial. How would you tell the difference?
2. Petrov’s “funny feeling” was contextual judgment built across a career. Design the training that produces it deliberately.
3. Petrov’s recoverability lived as much in his latitude to override as in his judgment. Map an automated decision in your domain where the operator has the knowledge to catch a failure but lacks the authority to act on it — and redesign the authority structure.

**WHO BUILDS THIS** human factors & interaction design · learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · scenario simulation · competency models

# F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap



Human-System Interface



Knowledge &amp; Institutional Memory



Normalization of Deviance

**IMPACT** Between 2008 and 2012 USAF F-22 Raptor pilots reported a cluster of physiological events consistent with hypoxia; one fatal accident (Capt. Jeffrey Haney, Nov 2010) was attributed by the Air Force accident board in part to operator response to a bleed-air shutoff, a conclusion the DoD Inspector General later found unsupported by the facts in the board's own report; the fleet was stood down in 2011, and the USAF Scientific Advisory Board had not determined a root cause and found OBOGS performance exceeded pilot physiological needs; the Air Force later traced the events to a valve inflating the pilots' pressure vests and restricting breathing

## IN BRIEF

Between 2008 and 2012 the USAF accumulated a cluster of F-22 Raptor “physiological events” — pilot reports consistent with hypoxia, often with persistent post-flight symptoms (chronic cough, disorientation). One fatal class-A accident in November 2010 (Capt. Jeffrey Haney) was attributed in part to operator response after a bleed-air shutoff. In May 2011 the USAF stood the fleet down for several months. The USAF Scientific Advisory Board's investigation, with NASA contribution, reported it had **not determined the root cause** — the finding of the case — while clearing OBOGS itself: system performance exceeded pilot physiological needs. What it found was a lack of robustness in the life-support configuration, the Breathing Regulator Anti-G valve and pressure-garment vest, aircrew flight equipment fit, and the absence of in-cockpit physiological monitoring that would have permitted earlier detection. The case is the canonical recent instance of a sustainment-era failure whose diagnosis was protracted because the system as fielded had no instrumentation of the variable that mattered — pilot blood oxygenation — and the failure mode therefore had to be inferred from pilot self-report and incident clustering.

## BACKGROUND

The F-22 Raptor entered operational service in 2005. Like most modern fighters it uses an On-Board Oxygen Generation System (OBOGS) rather than carrying liquid oxygen — a molecular sieve separates oxygen from engine bleed air to supply the pilot's breathing mask. Beginning in 2008, USAF Raptor pilots reported a cluster of “physiological events”: symptoms consistent with hypoxia during flight, often with persistent post-flight effects — chronic cough, disorientation, exercise intolerance. The cluster did not conform to a single

platform or sortie profile, and the reports accumulated faster than the diagnostic process could resolve them.<sup>0</sup>

#### WHAT HAPPENED

The public forcing event was the loss of Capt. Jeffrey Haney and his aircraft in November 2010 over Alaska, though the Scientific Advisory Board later set it aside. The Air Force Accident Investigation Board found that a bleed-air leak had triggered an automatic shutoff that cut oxygen flow to the pilot, but ruled that hypoxia did not contribute and laid the cause to the pilot's failure to initiate a timely dive recovery. The DoD Inspector General later found the board's conclusions unsupported by the facts; what the accident crystallized was that the platform's response to an OBOGS cutoff placed an unreasonable burden on pilot recognition and reaction in a window the platform did not instrument. In May 2011 the USAF stood the entire F-22 fleet down for several months pending investigation.<sup>1</sup>

#### THE INVESTIGATION

The USAF Scientific Advisory Board investigation, with participation from NASA's review of the hypoxia incidents, reported the finding the case turns on: it had **not determined the root cause**, and a Note in Proof recorded that the Life Support Systems Task Force believed it had. The findings ran the other way on OBOGS — system performance exceeded pilot physiological needs — and settled on a lack of robustness in the life-support configuration, the Breathing Regulator Anti-G valve and upper-pressure-garment vest that under some conditions restricted normal breathing, aircrew flight-equipment fit issues, and the absence of in-cockpit physiological monitoring. The investigation made several corrective recommendations — including modifications to the upper-pressure garment, changes to OBOGS components, and the addition of physiological monitoring during flight — and the symptom reporting decreased substantially in the period after implementation.<sup>2</sup>

#### THE CAPABILITY GAP

What makes this a sustainment-era case rather than a design-era one is the instrumentation gap. The F-22 as fielded had no in-cockpit measurement of the pilot's oxygenation or of the actual gas composition reaching the mask. The variable that mattered most for the failure mode — was the pilot getting enough oxygen? — had to be inferred from pilot self-report after the fact. The diagnostic process took years partly because the population of events had to be characterized statistically from a noisy reporting channel, and partly because the multi-cause structure of the failure made any single intervention only partially testable against the population.<sup>3</sup>

#### AFTERMATH & REFORM

The hedge survives into the case verbatim. The USAF SAB reported it had not yet determined the root cause, and the corrective actions were a bundle: garment modification, OBOGS component changes, training emphasis on hypoxia recognition, and the long-overdue addition of physiological monitoring. Symptom reports decreased after the bundle was implemented, but the attribution of the decline to any single component within the bundle is not possible from the available evidence. The case teaches the sustainment-instrumentation-gap form: a high-performance platform fielded without measurement of the operator-physiology variable that determines whether the mission can be completed, and

a diagnostic process forced to work from self-report and clustering because the platform did not surface the data.

*The variable that mattered most – was the pilot breathing freely? – was the variable the platform did not measure.*

EDITORS' SYNTHESIS OF THE USAF SAB AND NASA REVIEWS.

REFERENCES

1. USAF Scientific Advisory Board, Report on Aircraft Oxygen Generation, SAB-TR-11-04 (1 February 2012) — the 2011 Quick-look Study on the F-22 hypoxia-like incidents.

2. NASA Engineering and Safety Center, F-22 Life Support System Independent Analysis Team (2012) — NESC support to the Air Force life-support investigation.

3. USAF Accident Investigation Board (2011), Capt. Jeffrey Haney F-22 accident report (Nov 2010).

4. U.S. House Armed Services Subcommittee on Tactical Air and Land Forces, F-22 Pilot Physiological Issues, H.A.S.C. No. 112-154 (13 September 2012) — congressional record.

THE LEARNING ENGINEERING LENS

LE INSIGHT

F-22 OBOGS is the canonical recent sustainment-instrumentation-gap case. A high-performance platform was fielded without in-cockpit measurement of pilot oxygenation, the diagnostic process took years against a noisy self-report channel, and the corrective action was an irreducibly multi-cause bundle. The “root cause not determined” finding is the case.

LENS APPROACH

F-22 OBOGS is the aerospace instrumentation-gap case (induced 2.4; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the no-single-root-cause finding and the bundle-of-corrections evaluation problem, and in Domain 3 (Human-System Collaboration) for the operator burden when an automatic action cuts a life-critical supply. Pair with Case 108 (EVA-23) as the human- spaceflight instrumentation-gap companion and with Case 27 (anesthesia monitoring) at the cue/alert-as- capability layer.

| DOMAIN           | FAILURE MODES | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|------------------|---------------|---------------------|--------------|---------|-------|
| AVIATION DEFENSE | H K N         | D1 D2 D3 D4 D5      | 5            | 2.4     | 4 · 3 |
| HUMAN FACTORS    |               | Test and Evaluation |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Instrument the operator-physiology variable that determines mission completion before fielding a high-performance platform. The F-22 OBOGS case is the worked example of the cost of fielding without it: years of self-report-based diagnosis and one fatality.
- ▶ When a failure-event cluster does not conform to a single platform or sortie profile, treat the absence of a single root cause as a finding, not as a failure of the investigation. The bundle of contributing factors is the diagnosis.
- ▶ Design the platform’s automatic responses (e.g., bleed-air shutoff) with the pilot’s recognition window in mind; an automatic action that cuts a life-critical supply has to be paired with cueing the pilot can act on in time.

IN OPERATION / AFTER

- ▶ Implement the corrective actions as a bundle — garment modification, OBOGS components, training, physiological monitoring — and report the population-level outcome decline as bundle evidence, not as evidence about any single component.
- ▶ Carry the “root cause not determined” hedge into program communication; the multi-cause structure is the case’s teaching point and the basis for the sustainment-instrumentation argument.
- ▶ Treat the F-22 OBOGS experience as a fielded-platform sustainment lesson for subsequent platforms (e.g., F-35) — the analogous instrumentation question should be answered before the cluster appears, not after.

## REFLECTION QUESTIONS

1. Identify a fielded high-performance platform in your domain where the variable that determines mission completion is not directly measured. What is the analog of in-cockpit physiological monitoring, and what would it cost to add at sustainment vs. having designed it in?
2. Specify what your program would treat as “no single root cause” — the irreducibly multi-factor finding — and how you would report bundle-of-corrections evidence rather than overstating attribution to any single change.
3. The F-22 cluster appeared in 2008 and stand-down came in 2011. What signal-detection threshold would your program use on a noisy self-report channel to act earlier on a population-level pattern?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · incident analysis · safety dashboards

# Mark 14 Torpedo Failures

**T** Training Gap   **K** Knowledge & Institutional Memory   **G** Governance & Trust

**IMPACT** Persistent torpedo failures across the first ~20 months of the Pacific War; resolved only after fleet-level testing forced acknowledgment of the defects

## IN BRIEF

The U.S. Navy entered World War II with the Mark 14 torpedo, so guarded and so costly to expend that the Bureau of Ordnance had not destructively tested it since 1926. Through 1942 submarine crews reported torpedoes running deep, failing to detonate, or exploding prematurely; the Bureau insisted the weapon was sound and blamed the operators. It took about twenty months and repeated fleet-level forcing — Lockwood's net tests, the fleet-ordered net firings of summer 1942, and Lockwood's Pearl Harbor trials of July 1943 — to confirm three separate defects: the torpedo ran about ten feet too deep, its magnetic exploder fired erratically, and its contact pin crushed on a square hit. The fixes were simple once the defects were acknowledged. The binding constraint was institutional: a bureau insulated from the operators who knew the weapon did not work.

## BACKGROUND

The Mark 14 was the U.S. Navy's standard submarine torpedo at the start of the Pacific War. No destructive test had been fired with its magnetic exploder since 1926 — economy and secrecy together closing off realistic trials; the weapon went to war essentially unproven against realistic conditions, so the very decision meant to conserve a scarce and costly weapon guaranteed that its defects would first be discovered in combat, by the crews who could least afford them to surface there.<sup>9</sup>

## WHAT HAPPENED

From early 1942, submarine crews reported a litany of failures: torpedoes that ran far below their set depth, magnetic exploders that detonated prematurely or not at all, and contact firing pins that crushed on a direct hit. The Bureau of Ordnance insisted the weapon worked and attributed the misses to crew error; captains who reported failures risked their careers — an arrangement that turned every field report into a self-accusation and so suppressed the very evidence that would have isolated the defects the Bureau refused to acknowledge.<sup>1</sup>

## THE INVESTIGATION

It took about twenty months and three separate forcings: Lockwood's net test off Australia in June 1942, the fleet-ordered net firings that confirmed the depth error that summer, and Lockwood's Pearl Harbor trials after the USS Tinosa's July 1943 attack on the Tonan Maru, in which nine torpedoes struck the stopped ship squarely and failed to detonate — forced the issue. Each in turn confirmed that the torpedo ran about ten feet too deep, that the Mark 6 magnetic exploder failed routinely, and that the contact pin buckled on perpendicular


impact — three independent defects that had been masking one another at sea, which is why it took a fleet commander's controlled trials, not further combat patrols, to force each defect in turn onto the record.<sup>2</sup>

#### THE CAPABILITY GAP

The defects were real and separate, but the binding constraint was institutional. The capability that was missing was not at the front; it was a channel by which the bureau that owned the weapon could be made to hear what the boats already knew. The fixes — recalibrating depth, deactivating the magnetic exploder, redesigning the contact pin — were small. The refusal to believe the operators was the expensive part, measured not in engineering hours but in the months of patrols and the targets that escaped while a bureau defended a verdict the boats had already disproven.<sup>3</sup>

#### AFTERMATH & REFORM

By late 1943 the three defects were corrected and the Mark 14 became an effective weapon for the rest of the war. The episode entered U.S. Navy institutional history as the canonical case of a procurement bureau insulated from operator feedback, and is cited in modern organizational-learning literature on the cost of suppressing field reports — a reminder that the structure that decides whose evidence counts is itself a capability, and that its absence can cost more than any single technical defect it conceals.<sup>4</sup>



*The Bureau certified the weapon; field reports of failure were treated as evidence of operator error.*

PARAPHRASING BLAIR, SILENT VICTORY, 1975

#### REFERENCES

1. Rowland, B. & Boyd, W. (1953), U.S. Navy Bureau of Ordnance in World War II — testing policy and the three defects.
2. Blair, C. (1975), Silent Victory: The U.S. Submarine War Against Japan — operator reports and the Bureau's resistance (paraphrased).
3. Blair (1975) — the USS Tinosa's July 1943 attack (eleven consecutive duds on a stopped target) and Lockwood's fleet-level testing.
4. Naval History and Heritage Command, Mark 14 torpedo files — depth error, Mark 6 magnetic exploder, contact-pin failure.
5. Edmondson, A. (2018), The Fearless Organization — suppression of field reports as an organizational-learning failure.

#### THE LEARNING ENGINEERING LENS



## LE INSIGHT

The Mark 14 is the canonical Navy case for the institutional refusal to believe operator feedback. The capability that was missing was not at the front. It was at the bureau that owned the weapon. The cost of the refusal was paid by the submariners forced to use it until the bureau yielded. The eventual fixes — re-setting depth calibration, replacing the magnetic exploder, redesigning the contact pin — were technical and small. The reform was institutional and slow. The capability that had to be built was the channel by which the bureau could be made to hear what the boats already knew.

## LENS APPROACH

LENS uses the Mark 14 in LEN 7 as a governance failure case and in LEN 8 as an organizational-learning case in which the operators had the diagnosis and the institution lacked the structure to receive it. Studio projects (LEN 10) examine what the equivalent operator-to-institution feedback channel should look like.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|------------------|--------------|---------|-----|
| DEFENSE                               | → T K G       | → D1 D2 D3 D4 D5 | 4            | 6.2     | 5   |
| Navigating Sociotechnical Constraints |               |                  |              |         |     |

## APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- ▶ Require realistic live testing before fielding, even for scarce and costly items, so defects surface in trials rather than in combat.
- ▶ Build an operator-feedback channel into the weapon's ownership from the start, with a path that does not put the reporter's career at risk.
- ▶ Separate the authority that certifies a system from the authority that investigates its field failures, so a bureau cannot judge its own product.

## IN OPERATION / AFTER

- ▶ Audit field reports for suppressed or career-risking signals, and treat a pattern of "operator error" verdicts as a warning to test the system, not the crew.
- ▶ Empower an operational commander to order controlled trials when field reports conflict with the owner's certification.
- ▶ Sustain the feedback channel so multiple masking defects can be separated and proven before the cost compounds.

## REFLECTION QUESTIONS

1. Identify a system in your domain whose owners are institutionally insulated from the operators using it. What feedback would they refuse to hear, and what would it cost?
2. Design the operator-feedback channel that the U.S. Navy Bureau of Ordnance should have had in 1941. Who signs, who receives, what triggers action?
3. The USS Tinosa fired a string of torpedoes that struck a stopped ship and failed to detonate before the bureau accepted the diagnosis. What is the operator-evidence threshold in your domain that would force the equivalent institutional acknowledgment — and how would you make sure it is reached before the cost is paid?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

# V-22 Osprey

**T** Training Gap   **H** Human-System Interface   **N** Normalization of Deviance

**IMPACT** 20 service members killed in four crashes since 2022 alone, against a compiled record of 62 deaths across 16 hull losses since 1991; serious-mishap rate above comparable fleets (GAO-26-107285, 2025); some fixes stretch to the 2030s

## IN BRIEF

The V-22 Osprey — the tiltrotor flown by the Marines, Air Force, and Navy — has had about 16 hull-loss accidents and 62 fatalities since 1991, including 19 Marines in a single 2000 test crash and 8 airmen off Yakushima, Japan, in 2023. The Yakushima crash traced to cracks in a transmission gear (a flaw in the X-53 steel alloy) and a pilot's decision to keep flying through warnings. In December 2025, GAO and a NAVAIR review found the V-22's joint program office had failed for years to assess and address mounting safety risks, that serious-mishap rates exceeded comparable fleets, and that some fixes stretch toward

2034. The Osprey is the steady-state form of normalization of deviance:

a documented, reviewed shortfall accepted as the cost of flying the airframe.

## BACKGROUND

The V-22 Osprey tilts its rotors to take off like a helicopter and cruise like a turboprop — an ambitious capability shared awkwardly across three services and a joint program office. Since development began in the 1980s it has suffered 16 hull-loss accidents and 62 fatalities, including 19 Marines in a single crash during 2000 testing.<sup>0</sup> The same configuration that makes the tiltrotor uniquely useful makes it uniquely demanding to sustain, and the joint arrangement diffused ownership of that burden across three services that each carried the airframe but none of which fully owned its safety trajectory.

## WHAT HAPPENED

On 29 November 2023 a CV-22B broke up off Yakushima, Japan, killing all eight airmen aboard. The Air Force traced it to cracks in a transmission gear — rooted in metallic inclusions in the X-53 steel alloy used for the gears — and to the pilot's decision to keep flying despite warnings to land. The crash grounded Osprey fleets worldwide for months.<sup>1</sup> The materiel flaw and the human decision to press on through warnings were the same paired factors earlier reviews had named, so Yakushima read less as a novel failure than as one more draw from a pattern the program had already characterized but not closed.

## THE INVESTIGATION

In December 2025 the Government Accountability Office and a NAVAIR review reported together that the joint program office had failed for years to adequately assess and address mounting safety risks, even as service members died.<sup>2</sup> Serious-mishap rates generally exceeded those of comparable Navy and Air Force fixed- and rotary-wing fleets from FY2015 to FY2024 and spiked in 2023–2024, across a stretch in which four crashes killed twenty service members; the median age of the program's 28 unresolved system safety risk assessments was about nine years, with 17 of the 28 carried between six and fourteen years, and full proprotor-gearbox fixes are not expected until 2034. The NAVAIR review found that material risks the program office had itself identified were not sufficiently mitigated or resolved in time, "which resulted in catastrophic outcomes in 5 of the 12 mishaps."<sup>3</sup> Risks carried six to fourteen years are the timescale of normalization made literal: the deviation persisted long enough to become the airframe's accepted background condition rather than an open defect demanding action.

## THE CAPABILITY GAP


The V-22 is the steady-state form of normalization of deviance. The shortfall is not unknown — it has been documented across successive reviews — but the system built to remediate it has accepted its own incompleteness as the cost of operating the airframe. Each crash carries a familiar set of factors (materiel failure, human error, weak coordination across the three services), and each is followed by adjustments that do not converge.<sup>4</sup> Convergence is the missing property: three services adjusting in parallel, none holding final authority over the whole, produce motion without resolution, so the mishap rate stays elevated while every individual response looks reasonable on its own terms.

## AFTERMATH & REFORM

Groundings, gear-inspection regimes, and a redesign program have followed, with full fixes for some gearbox issues stretching toward

2034. The case is instructive precisely because the problem is

recognized: a known, reviewed capability gap can persist for decades when each incident is treated as an isolated event rather than a sample from a distribution the program keeps drawing from.<sup>5</sup> Fixes that stretch toward 2034 mean the airframe will keep flying for years inside a margin already judged worse than comparable fleets — the program in effect electing to operate the distribution it has been told it is drawing fatalities from.



*Material risks were identified by the Program Office, but were not sufficiently mitigated or resolved in a timely manner, which resulted in catastrophic outcomes in 5 of the 12 mishaps.*

NAVAIR INDEPENDENT REVIEW OF THE V-22, DECEMBER 2025, AS QUOTED IN DEFENSE ONE

REFERENCES

1. Compiled V-22 accident record — 16 hull losses and 62 fatalities since 1991, including the 2000 Marana test crash (19 Marines). (A compiled tally, not an official count; the four-crash, twenty-death figure for 2022–2025 is GAO’s.)

2. U.S. Air Force Accident Investigation Board findings, via USNI News (Aug. 2024) — the 29 Nov. 2023 Yakushima CV-22B crash: transmission-gear cracks (X-53 inclusions) and continued flight despite warnings.

3. U.S. GAO, Osprey Aircraft: Additional Oversight and Information Sharing Would Improve Safety Efforts, GAO-26-107285 (Dec. 2025) — the joint program office’s failure to assess and address risks.

4. GAO-26-107285 (2025) — serious-mishap rates exceeding comparable Navy/Air Force fleets (FY2015–FY2024); 28 unresolved system safety risks with a median age of about nine years, over half carried six to fourteen years.

5. NAVAIR independent review of the V-22 (Dec. 2025), reported in Defense One, “New GAO, Navy reports warn of serious V-22 Osprey safety risks” (Dec. 2025) — material and cross-service-coordination factors, unresolved catastrophic parts issues, the 5-of-12 mishap finding, and proprotor-gearbox fixes running to 2034.

6. USNI News, V-22 program coverage (2024–2025).

THE LEARNING ENGINEERING LENS

LE INSIGHT

V-22 demonstrates the steady-state version of normalization of deviance: a platform whose shortfall has been documented, reviewed, and acted on across multiple administrations without producing sustained improvement. The capability gap has itself been normalized. Each incident is treated as an event rather than as a sample from a distribution the program continues to draw from.

LENS APPROACH

LENS treats V-22 in LEN 5 as a multi-service capability-coordination problem and in LEN 8 as a study in long-cycle organizational learning failure. Students design the evidence system that would distinguish a true reduction in mishap rate from the natural variation of a chronically marginal platform.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|------------------|--------------|---------|-----|
| DEFENSE                               | → T H N       | → D1 D2 D3 D4 D5 | 4            | 7.4     | 5   |
| Navigating Sociotechnical Constraints |               |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Assign one accountable owner for the airframe’s safety trajectory rather than diffusing it across three services with no final authority.
- ▶ Treat each known materiel flaw as an open defect on a clock, so a problem like the gearbox cannot quietly become accepted background condition.
- ▶ Define an absolute mishap-rate threshold against comparable fleets that triggers mandatory action, not merely review.

IN OPERATION / AFTER

- ▶ Aggregate incidents as samples from one distribution, not isolated events, so the elevated rate itself becomes the thing being managed.
- ▶ Audit whether parallel service-level adjustments are actually converging on a lower rate, and escalate when they are not.
- ▶ Gate continued operation on remediation progress, so fixes stretching toward 2034 do not silently license years of degraded margin.

### REFLECTION QUESTIONS

1. Identify a platform or process in your domain that has been operating in a documented shortfall for years. What measurement would have to change for the shortfall to become unacceptable?
2. The V-22's three services do not converge on remediation. Design the governance structure that would force convergence.
3. Osprey safety risks sat unresolved for a median of nine years, and some for fourteen. What mechanism in your domain converts a long-tolerated defect back into an open item that demands action?

**WHO BUILDS THIS** learning science & instructional design · human factors & interaction design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs · incident analysis · safety dashboards

# 9/11 Intelligence Sharing Failures

**G** Governance & Trust   **K** Knowledge & Institutional Memory

**IMPACT** 2,977 killed; the 9/11 Commission found systemic intelligence-sharing failures across the U.S. government

## IN BRIEF

The September 11, 2001 attacks, which killed 2,977 people, were enabled in part by intelligence the U.S. government already held but never integrated. The CIA tracked two future hijackers to a January 2000 meeting in Kuala Lumpur and lost them in Bangkok; it did not watch-list them, and it did not tell the FBI about one man's U.S. visa or the other's travel to the United States. The FBI separately flagged suspicious flight-training activity in Phoenix and Minneapolis in 2001, but that information was never aggregated. Visa issuance, immigration tracking, and watch-listing were each run by a different agency, and the handoffs between them depended on individual initiative that no institution required. The 9/11 Commission called it a "failure of imagination" — a framing critics say understates the structural gap. The reform that followed built the cross-agency architecture, the ODNI and the NCTC, that had not existed.

## BACKGROUND

By 2001, U.S. counterterrorism depended on many agencies — CIA, FBI, NSA, State, INS, FAA — each holding a piece of the picture. No institution was responsible for integrating those pieces; cross-agency information sharing depended on individuals choosing to pass information along rather than on any architecture that required it. Each agency's incentives, classification rules, and turf reinforced the boundary, so the natural tendency was to hold information rather than to push it across a line no one was charged with bridging.<sup>o</sup>

## WHAT HAPPENED

On 11 September 2001, nineteen hijackers seized four aircraft and killed 2,977 people. In the months and years before, the warning signs had been distributed across the government: the CIA tracked two future hijackers from a January 2000 meeting in Kuala Lumpur but did not watch-list them or notify the FBI; FBI field offices flagged suspicious flight-training activity in Phoenix and Minneapolis; none of it was aggregated into a single picture. Any one fragment looked minor in isolation; only assembled would they have shown the shape of the plot, and assembly was exactly the function no one performed.<sup>1</sup>

## THE INVESTIGATION

The 9/11 Commission and the earlier Congressional Joint Inquiry documented specific failures of information sharing across the FBI, CIA, and NSA. The Commission famously concluded that "the most important failure was one of imagination" — a framing later criticized as understating the structural nature of the gap, which was less a lack of imagination than


an absence of any institution responsible for integration. The distinction matters for the remedy: a failure of imagination invites exhortation to think harder, while a structural gap demands an institution be built to close it.<sup>2</sup>

#### THE CAPABILITY GAP

The intelligence sharing did not happen because no institution owned it as a deliverable. Each agency was competent inside its own boundary; the integration across boundaries existed nowhere as a required function. The missing capability was an architecture — shared watch-lists, mandated handoffs, a body responsible for fusing the picture — rather than more raw collection. The government did not lack data; it lacked the connective tissue to turn data held in many places into a single picture anyone was accountable for assembling.<sup>3</sup>

#### AFTERMATH & REFORM

The Intelligence Reform and Terrorism Prevention Act of 2004 created the Office of the Director of National Intelligence and the National Counterterrorism Center, and fusion centers followed — the institutional architecture for cross-agency information sharing that had not previously existed. By creating bodies whose explicit mandate was integration, the reform converted information sharing from an act of individual initiative into a required function someone owned. The case is foundational in U.S. national-security policy for treating cross-agency capability as an engineerable institutional deliverable.<sup>4</sup>



*no institution responsible for the integration as a deliverable*

9/11 — CROSS-AGENCY INFORMATION SHARING AS MISSING ARCHITECTURE

*The most important failure was one of imagination.*

THE 9/11 COMMISSION REPORT, 2004

#### REFERENCES

1. National Commission on Terrorist Attacks Upon the United States, The 9/11 Commission Report (2004) — the quoted “failure of imagination” and the specific sharing failures.
2. 9/11 Commission Report (2004) — the Kuala Lumpur tracking and the Phoenix/Minneapolis flagging.
3. Joint Inquiry into Intelligence Community Activities Before and After the Terrorist Attacks of September 11, 2001, S. Rept. 107-351 (2002) — cross-agency information-sharing failures.
4. Zegart, A. (2007), Spying Blind — structural-organizational analysis of the failures.
5. Intelligence Reform and Terrorism Prevention Act of 2004 — creation of the ODNI and NCTC.

#### THE LEARNING ENGINEERING LENS

LE INSIGHT

9/11 is the foundational U.S. case for cross-agency information- sharing as an engineering deliverable. The architecture had not been built. The reform built it. The cost of the missing architecture was paid in 2001. The discipline LENS represents is the kind of work that, applied across the U.S. intelligence community in the 1990s, would have produced the architecture earlier.

LENS APPROACH

LENS uses 9/11 in LEN 8 as the foundational case for cross- organizational capability and in LEN 7 for the governance architecture of multi-agency systems. Studio projects compare 9/11 with Eagle Claw (Case 128) as institutional-architecture failures of different kinds.

| DOMAIN     | FAILURE MODES | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO |
|------------|---------------|---------------------------------------|--------------|---------|-----|
| DEFENSE    | → G K         | → D1 D2 D3 D4 D5                      | 1            | 1.3     | 5   |
| GOVERNMENT |               | Navigating Sociotechnical Constraints |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Assign integration as an explicit deliverable owned by a named body, so the function does not depend on individuals choosing to pass information across a boundary.</li><li>Engineer mandated handoffs and shared watch-lists into the architecture, making cross-agency sharing a required function rather than an act of initiative.</li><li>Design for fusion of fragments that look minor in isolation, so the system is built to assemble the picture no single holder can see.</li></ul> | <ul style="list-style-type: none"><li>Audit cross-boundary information flows for reliance on individual initiative, and treat any flow that depends on goodwill as an unowned failure mode.</li><li>Monitor whether the integrating body actually receives and fuses what the component agencies hold, rather than assuming the mandate guarantees the practice.</li><li>Sustain the integration architecture against the agency incentives, classification rules, and turf that pull information back behind boundaries over time.</li></ul> |

REFLECTION QUESTIONS

1. Identify a cross-organizational information flow in your domain that depends on individual initiative rather than institutional architecture. What is the foreseeable failure mode?
2. Design the institutional deliverable that would have produced ODNI-level information sharing across U.S. intelligence agencies in the 1990s.
3. The Commission called it a “failure of imagination”; critics called it a structural gap. Take a near-miss in your domain and argue which framing fits — and show how the remedy differs depending on which you choose.

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · knowledge bases · data pipelines



Chapter 8

# Defense & National Security — What Works — and the Frontier

*When training, standards, and sustainment are engineered as one deliverable.*

---

*The readiness that survived contact was the readiness that had been measured.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

# U.S. Nuclear Navy / Rickover Training Model

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Zero reactor accidents in 60+ years of U.S. Naval nuclear operations; the most demanding nuclear operator training program in the world

## IN BRIEF

Admiral Hyman Rickover built a training and qualification culture for the Naval Nuclear Propulsion Program that has produced zero reactor accidents across more than 60 years and thousands of reactor-years of operation, often in extreme conditions. Every nuclear-trained sailor must pass rigorous qualification to zero-defect standards and demonstrate competence in oral examination by senior nuclear-qualified officers; the culture demands personal accountability, deep technical mastery, and a questioning attitude that obliges operators to challenge assumptions, including superiors'. The sharpest contrast in this book is internal: the same Navy that ran the nuclear program to this standard let surface-warfare training decay to CD-ROMs and paid for it at Fitzgerald and McCain (Case 124). Same institution, two philosophies — and the choice shows up in the casualty columns.

## BACKGROUND

In the early 1950s the U.S. Navy set out to put nuclear reactors aboard ships and submarines — machines whose failure could be catastrophic and irreversible, operated by young sailors far from any help. The capability problem was absolute: there was no acceptable accident rate, so the human operating system had to be engineered to an extraordinary standard from inception. Because a reactor at sea could not be evacuated or handed to outside experts, the operators themselves had to be the entire margin of safety — a constraint that forced training to a level no ordinary program would justify.<sup>0</sup>

## THE INTERVENTION

Admiral Hyman Rickover established a training and qualification regime requiring every nuclear-trained officer and enlisted sailor to pass demanding programs held to zero-defect standards. Operators must demonstrate competence through oral examination by senior nuclear-qualified officers, and the program embeds continuous re-qualification rather than one-time certification. The oral board tested understanding rather than recall, and the continuous re-qualification meant competence had to be sustained rather than banked once — so the standard governed the operator's whole career, not just entry to it.<sup>1</sup>

## HOW IT WORKED

The culture pairs technical mastery with a deliberately engineered attitude: personal accountability and a mandatory questioning posture in which operators are trained to challenge assumptions, including those of superiors. Rickover's premise — that people, not organizations or management systems, get things done — made the qualification ladder, not

paperwork, the load-bearing element of safety. The questioning posture is the cultural half of the pair: deep technical mastery alone could still defer to a mistaken superior, so the obligation to challenge assumptions is what keeps competence from being silenced by rank.<sup>2</sup>

#### THE EVIDENCE

The result is the longest-running continuous capability-engineering record in any high-consequence domain: zero reactor accidents across more than six decades and thousands of reactor-years. The cost — the qualification ladder, the zero-defect oral boards, the continuous re-qualification — is the visible budget-line price of that record. The duration is the key evidence: a record sustained across six decades and many generations of sailors shows the safety came from the engineered system rather than from any one cohort's talent or luck.<sup>3</sup>

#### WHAT TRANSFERRED

The internal comparison is the instructive one. The same Navy that engineered the nuclear program to this standard let surface-warfare training decay to CD-ROM self-study and paid the price at Fitzgerald and McCain (Cases 124 and 140). Same institution, same era, opposite philosophies, radically different outcomes — the strongest available demonstration of capability treated as a system parameter versus deferred as a cost. The internal comparison controls for nearly everything that usually confounds such claims — one service, one manpower system, one budget process — leaving the training philosophy itself as the variable that diverged.<sup>4</sup>



*Human experience shows that people, not organizations or management systems, get things done.*

ADMIRAL HYMAN G. RICKOVER, "DOING A JOB," COLUMBIA UNIVERSITY EGLESTON MEDAL ADDRESS, 5 NOVEMBER 1982

#### REFERENCES

1. Polmar, N. & Allen, T. (2007), Rickover: Father of the Nuclear Navy — the program and Rickover's philosophy (paraphrased).
2. Naval Nuclear Propulsion Program documentation (NRC/DOE) — qualification standards and the accident record.
3. Admiral Hyman G. Rickover, "Doing a Job" (Columbia University, Egleston Medal address, 5 November 1982) — "people, not organizations... get things done" (quoted).
4. Duncan, F. (1990), Rickover and the Nuclear Navy — the qualification culture.
5. GAO-21-168 (2021), Navy Readiness: Actions Needed to Evaluate and Improve Surface Warfare Officer Career Path — the surface community's career path set against the submarine, aviation and EOD communities.

#### THE LEARNING ENGINEERING LENS

LE INSIGHT

The Nuclear Navy is the longest-running continuous capability- engineering program in any high-consequence domain. The choice to treat training as a system parameter rather than as a cost center has produced sixty-plus years of zero reactor accidents. The contrast with the Surface Navy shows what happens when capability is engineered versus when it is deferred — and the price of that engineering (the qualification ladder, the zero-defect oral boards, the continuous re-qualification) is visible on the budget line.

LENS APPROACH

LENS treats the Rickover model in LEN 8 as the anchor organizational-learning case and in LEN 5 as a worked example of capability requirements traceable from operational analysis through qualification standards. The internal Navy comparison anchors the program’s core argument about capability as a system parameter.

| DOMAIN           | MODES ADDRESSED |   |   | LENS COMPETENCY |   |    |    |    | PROBLEM TYPE | INDUCED | LEO |     |   |
|------------------|-----------------|---|---|-----------------|---|----|----|----|--------------|---------|-----|-----|---|
| DEFENSE          | →               | T | K | N               | → | D1 | D2 | D3 | D4           | D5      | 1   | 1.4 | 1 |
| Systems Analysis |                 |   |   |                 |   |    |    |    |              |         |     |     |   |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Engineer the human operating system to the standard the consequences demand from inception — where there is no acceptable failure rate, make the operators themselves the margin of safety.
- ▶ Gate qualification on demonstrated understanding through oral examination by senior qualified people, testing comprehension rather than recall.
- ▶ Pair technical mastery with a mandatory questioning posture so competence is obliged to challenge assumptions, including superiors’, rather than be silenced by rank.

IN OPERATION / AFTER

- ▶ Embed continuous re-qualification rather than one-time certification, so competence must be sustained across a career and cannot be banked once and assumed.
- ▶ Protect the training as a durable, visible budget line, since the qualification ladder and oral boards are the price of the safety record and the first thing tempo will erode.
- ▶ Sustain the standard across generations of operators, treating a multi-decade record as the evidence the safety comes from the engineered system rather than any one cohort.

REFLECTION QUESTIONS

1. Identify an institution that operates two divisions under different capability philosophies. What does the comparison reveal that neither division could see alone?
2. Rickover’s standard was zero-defect oral examination. What is the equivalent in your domain — and would you survive it?
3. The Nuclear Navy’s safety record is paid for by a durable, visible training budget. What is the equivalent line-item investment in your domain that a comparable safety claim would require — and is it being made?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

# MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard

**G** Governance & Trust   **K** Knowledge & Institutional Memory

**IMPACT** DoD Design Criteria Standard MIL-STD-1472H, the 2020 revision of a series dating to 1968, converts human-factors and human-engineering findings into binding design criteria across DoD acquisition — controls, displays, anthropometry, workspace, environment, hazards — invoked by acquisition programs as a mandatory or tailored design specification

## IN BRIEF

MIL-STD-1472 is the U.S. Department of Defense's design-criteria standard for human engineering — the discipline of making equipment usable, safe, and effective for the human operator. The series originated in 1968 and has been revised through versions A, B, C, D, E, F, G, and most recently H, released September 15, 2020. The standard codifies established human-factors and ergonomic findings into specific binding design criteria: control and display design, anthropometric accommodation ranges, workspace dimensions and access, labeling and signage, environmental limits (noise, vibration, illumination, thermal), and hazard mitigation. Acquisition programs invoke MIL-STD-1472H either as a mandatory standard or with tailored deviation, making human-factors evidence binding rather than advisory. The 2020 revision integrated accumulated findings since 2012 (the prior G revision), including updates to anthropometric data, environmental criteria, and human-system-integration practices. The case is the structural archetype of converting a body of human-factors research into engineered design requirements; it works at the requirement-discipline layer rather than the per-platform layer, and is the standard that programs referenced when they specify human-engineering deliverables.

## BACKGROUND

Human factors engineering as a defense discipline emerged from World War II — the documented mismatch between cockpit control layouts and pilot perception (Fitts & Jones's "designed errors" work) showed that platforms could be made un-flyable by avoidable design choices. Through the 1950s and 60s the discipline accumulated a body of findings on control / display design, anthropometry, environmental tolerance, and workspace layout that lived across academic publications, military handbooks, and service-specific guidance. What was missing was a single binding design-criteria standard that an acquisition program could specify in a contract.<sup>o</sup>

## THE INTERVENTION

The DoD response was MIL-STD-1472, first issued in 1968, and revised through letter suffixes A (1970), B (1974), C (1981), D (1989), E (1996), F (1999), G (2012), and H (September

2020). Each revision incorporates accumulated findings since the prior version, updated anthropometric data (the surveyed populations evolve), and refinements to specific criteria informed by service experience. The standard is not a list of recommendations; it is a Design Criteria Standard, structured so that an acquisition program can invoke it as a mandatory specification, or tailor specific paragraphs with documented deviation.<sup>1</sup>

#### HOW IT WORKED

The structural mechanism is the conversion of human-factors findings into engineered requirements. Where the research literature might find that controls of a given size and force are operable across a defined population percentile, MIL-STD-1472H carries that finding as a binding design criterion that the program's controls must meet. Where the literature finds that illumination below a certain level degrades reading performance for given character sizes, the standard carries that as an environmental limit. Anthropometric accommodation — the body-size ranges the equipment must fit — is specified to defined population percentiles. Hazard analysis is required as part of the design process, with mitigation criteria for identified hazards.<sup>2</sup>

#### THE EVIDENCE

The case's value to LENS is the requirement-discipline form. MIL-STD-1472H is the human-engineering analog of what SUBSAFE (Case 173) does for submarine watertight-integrity: a binding, recurring, auditable requirement set that the program-management chain cannot trade away without a documented deviation. The acquisition contract invokes the standard, the acquirer's human-factors discipline verifies compliance against it, and the engineered design carries the research findings as binding criteria rather than as advice the program might or might not adopt. The 2020 H revision is notable for the cadence — the eight-year gap from G to H — and for the coverage it added: habitability, cybersecurity, information presentation, and handheld devices, reflecting the rise of complex software-intensive systems.<sup>3</sup>

#### WHAT TRANSFERRED

The honest framing the case requires is that the standard is necessary but not sufficient. MIL-STD-1472H is the requirements-discipline mechanism; whether the requirements are met in any specific program depends on the program's human-factors engineering capacity, the acquirer's verification rigor, and the tailoring decisions made up front. The standard does not by itself guarantee usable equipment; it makes usability a contractable, auditable deliverable. Cases of operator-platform mismatch in fielded DoD systems (e.g., the F-22 OBOGS instrumentation gap in Case 132) are not failures of the standard's existence but of the tailoring and verification around it. The case teaches the requirements-as-deliverable form at the human-engineering scale, with the qualification that the standard is the mechanism, not the outcome.

*The criterion that survives into a binding standard is the criterion the program is held to.*

EDITORS' SYNTHESIS OF MIL-STD-1472 REVISION HISTORY.

#### REFERENCES

1. Fitts, P. M., & Jones, R. E. (1947), "Analysis of factors contributing to 460 'pilot error' experiences in operating aircraft controls" — origin of designed-error analysis.
3. Chapanis, A. (1965), "Man-Machine Engineering" — foundational text for the discipline the standard codifies.

2. Department of Defense (2012), MIL-STD-1472G — the prior revision; revision history documents the 1968 origin and intermediate letters.

4. Department of Defense (2020), MIL-STD-1472H “Department of Defense Design Criteria Standard: Human Engineering,” 15 September 2020 — replaces MIL-STD-1472G (2012).


THE LEARNING ENGINEERING LENS

LE INSIGHT

MIL-STD-1472H is the structural archetype of converting a body of human-factors research into binding engineered requirements an acquisition contract can specify. The standard is the mechanism by which usability becomes a contractable, auditable deliverable rather than design advice. It is necessary but not sufficient; tailoring and verification determine whether the contract is honored.

LENS APPROACH

MIL-STD-1472H is the binding-standard requirements case (induced 1.1; LENS D1/PT3). LENS uses it in Domain 1 (Systems Analysis) for the conversion of research findings into engineered requirements, and in Domain 5 (Navigating Sociotechnical Constraints) for the tailoring and verification disciplines around the standard. Direct pair with Case 173 (SUBSAFE) at the requirements-as-deliverable layer and with Case 132 (F-22 OBOGS) at the standard-versus-tailoring layer.



APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the human-factors literature in your domain as a source of engineered requirements, not as design advice. The criterion that survives into a binding standard is the criterion the program is held to.
- ▶ Specify anthropometric and environmental criteria to defined population percentiles, not to “typical users.” The percentile framing is what makes accommodation auditable.
- ▶ Design the tailoring discipline around the standard so that any deviation from a binding criterion is documented with rationale, and the program retains a record of which criteria it chose not to meet and why.

IN OPERATION / AFTER

- ▶ Maintain the revision cadence the field requires — anthropometric data ages, environmental tolerances drift, human-system-integration practices evolve — so the standard does not become an obsolete reference.
- ▶ Verify compliance as a discipline distinct from the standard itself; the standard is the requirement, the verification process is the assurance that the requirement is met.
- ▶ Carry the necessary-but-not-sufficient framing: a binding human-engineering standard is the mechanism by which research findings become contractable; whether the contract is honored is the program’s responsibility, not the standard’s.

REFLECTION QUESTIONS

1. Identify a body of research findings in your domain that programs treat as advisory. What would it take to convert those findings into a binding design-criteria standard an acquirer could specify in a contract?
2. Specify the tailoring discipline you would put around a binding standard: which paragraphs may be tailored, with what documented rationale, and what is the cadence at which the standard itself is revised against accumulated experience?
3. The standard is necessary but not sufficient. What verification capacity does your program need to know that the binding criteria are actually met in the delivered system, and where is the gap between standard and verification visible today?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines



# GIFT and the Adoption Gap

**K** Knowledge & Institutional Memory

**G** Governance & Trust

**N** Normalization of Deviance

**IMPACT** Active open-source ITS framework with demonstrated learning effectiveness; ubiquitous fielded adoption in routine military training has not been achieved

## IN BRIEF

The Generalized Intelligent Framework for Tutoring (GIFT) is an open-source framework, originated at the U.S. Army Research Laboratory, for authoring and delivering intelligent tutoring systems. Computer-based tutoring has been shown to be nearly as effective as human tutoring in well-defined domains, and GIFT exists to lower the barrier to building it; the framework is actively developed, with regular releases and a peer-reviewed annual symposium. The puzzle is not that GIFT failed — it did not — but that ubiquitous fielded adoption across routine military training remains limited despite decades of supporting research. That gap is the canonical learning-engineering problem: the science is settled, the platform exists, the studies are positive — and the institutional pathway to scaled adoption is still being built.

## THE SHIFT

Five decades of research established that computer-based tutoring can be nearly as effective as human tutoring in well-defined domains —  $d = 0.76$  against  $0.79$  — and well above no tutoring at all. The open question shifted from “does adaptive tutoring work?” to “why isn’t it everywhere?”<sup>o</sup> That shift matters because it moves the problem out of the laboratory: once the efficacy question is answered, every remaining obstacle to scaled use is institutional rather than scientific, and the field’s research strength stops being the binding constraint.

## WHAT IS EMERGING

The Generalized Intelligent Framework for Tutoring (GIFT), originated under the U.S. Army Research Laboratory and now developed at DEVCOM, is an open-source framework for authoring, delivering, and evaluating intelligent tutoring systems — an effort to lower the authoring barrier that has historically made ITS expensive to build. It is actively maintained, with regular releases and a peer-reviewed annual symposium.<sup>1</sup> Open-sourcing the framework and sustaining a research community around it directly attacks the cost-to-build problem, since the expense of authoring a tutor from scratch had long been the practical reason adaptive tutoring stayed confined to well-funded demonstrations.

## THE CAPABILITY QUESTION

So the puzzle is not failure — GIFT was not discontinued. It is that ubiquitous fielded adoption across routine military training remains limited despite a working framework and decades


of positive research. The science is settled and the platform exists; what is missing is the institutional pathway to scaled use.<sup>2</sup> This is the diagnostic feature of the case: the gap is not between idea and prototype but between a maintained, evidence-backed framework and the routine training pipelines that would have to adopt it, and that latter distance is the one no amount of further research closes.

### EARLY EVIDENCE

The effectiveness evidence is strong — the tutoring-effectiveness literature is among the more robust in education — and GIFT-based studies continue to demonstrate learning gains. The bottleneck is not efficacy.<sup>3</sup> Because the supporting literature is among the more robust in the field, a decision-maker hesitating to field adaptive tutoring cannot honestly point to weak evidence as the reason; the hesitation traces instead to the missing pathway that would let the proven approach be bought, integrated, and made routine.

### OPEN PROBLEMS

What has not been built is the adoption pathway: procurement that can buy adaptive tutoring, integration into existing training pipelines, instructor-workflow redesign, and the authority structure to make adaptive tutoring a default rather than an experiment.<sup>4</sup> GIFT is the case in this book closest to the LENS discipline itself — proof that a working technology and settled science do not adopt themselves, and that the adoption pathway is an engineering deliverable in its own right. Each of those pieces — a contracting vehicle, a pipeline integration, a redesigned instructor workflow, an owner with authority — is a concrete artifact someone must build, and their absence, not any technical shortfall, is what keeps the framework experimental.



*The technology works. The institutional pathway to ubiquitous fielded use does not yet.*

EDITORS' SYNTHESIS OF THE GIFT ADOPTION RECORD

### REFERENCES

1. K. VanLehn, "The Relative Effectiveness of Human Tutoring, Intelligent Tutoring Systems, and Other Tutoring Systems," *Educational Psychologist* 46(4) (2011) — tutoring effectiveness.
2. GIFT Project, [gifttutoring.org](http://gifttutoring.org) — the framework, releases, and development under ARL / DEVCOM.
3. Editors' synthesis of the GIFT adoption record — active development but limited ubiquitous fielding (quoted).
4. R. Sottolare, A. Graesser, X. Hu & H. Holden (eds.), *Design Recommendations for Intelligent Tutoring Systems* (U.S. Army Research Laboratory, 2013– ); IJAIED Special Issue on GIFT (2017).
5. J. Goodell & J. Kolodner, *Learning Engineering Toolkit* (2022) — adoption as an engineering problem.

## THE LEARNING ENGINEERING LENS

## LE INSIGHT

GIFT is the most directly relevant case in this book to the learning-engineering discipline itself. The technology works, the learning science works, the framework is active and supported. What has not been built is the institutional adoption pathway — procurement, training-pipeline integration, instructor workflow redesign, the authority structure — that would make adaptive tutoring a default rather than an experiment.

## LENS APPROACH

LENS treats GIFT in LEN 1 as the problem-framing case for the discipline, in LEN 8 as the foundational adoption-pathway case, and in LEN 10 (capstone) as a prompt for designing the institutional deliverables that would close an adoption gap of this shape.

| DOMAIN             | MODES AT STAKE        | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|--------------------|-----------------------|------------------|--------------|---------|-----|
| DEFENSE            | → K G N               | → D1 D2 D3 D4 D5 | 4            | 1.1     | 2   |
| EDUCATION AT SCALE | Iterative Development |                  |              |         |     |

## APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- ▶ Treat the adoption pathway — procurement, pipeline integration, instructor-workflow redesign, authority — as a named deliverable of the program, not a follow-on hope.
- ▶ Build the tutoring capability against an existing training pipeline so integration is designed in, rather than fielding a framework and expecting pipelines to bend to it.
- ▶ Specify who holds the authority to make adaptive tutoring a default, so the decision to scale is owned rather than left to volunteers.

## IN OPERATION / AFTER

- ▶ Measure fielded routine use, not just study-level learning gains, so the institution can see whether adoption is actually happening.
- ▶ Sustain the open-source framework and its community so the authoring-cost barrier it lowered does not quietly rise again as releases age.
- ▶ Audit each stalled pipeline to find which adoption artifact is missing — contracting vehicle, integration, workflow, owner — and treat that as the engineering gap to close.

## REFLECTION QUESTIONS

1. GIFT exists, is supported, and works. Adoption at scale does not. What is the equivalent in your domain — an effective intervention whose adoption pathway has not been engineered?
2. Design the institutional adoption deliverable that would move adaptive tutoring from “available framework” to “default routine practice” in one operational training pipeline.
3. GIFT’s bottleneck is procurement, pipeline integration, instructor-workflow redesign, and authority — not efficacy. For an effective tool in your domain, which of those four is the binding constraint, and who would have to own it for adoption to become the default?

**WHO BUILDS THIS** knowledge management & data engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Navy Surface Warfare Readiness Reform

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Threefold increase in initial ship-driving training hours planned by 2021; 10 pass-or-fail career assessments, 4 of them go/no-go; Ready-for-Sea Assessments — 3 of 18 forward-deployed ships immediately sidelined

## IN BRIEF

After the fatal 2017 Fitzgerald and McCain collisions (Case 124) exposed gutted seamanship training, the U.S. Navy overhauled how it builds surface-warfare competence. It restored the Surface Warfare Officers School from CD-ROM self-study to classroom and simulator instruction, stood up Mariner Skills Training Centers on both coasts, created ten pass-or-fail career assessments — four of them no-go gates — and adopted aviation-style debriefing. New Ready-for-Sea Assessments evaluated forward-deployed ships against a deliverable standard; three of the first eighteen were immediately sidelined. The structural change is real and the investment substantial. What is missing is the third half: GAO found in 2019 that the Navy had not put key processes and assessments in place to evaluate the effectiveness of its changes to ship-driving training — a live success in structure, with evidence of effect still outstanding.

## BACKGROUND

In 2017, the destroyers USS Fitzgerald and USS John S. McCain suffered fatal collisions (Case 124) that investigations traced in part to seamanship and navigation training degraded to CD-ROM self-study. The Navy faced a clear capability gap: officers were going to sea without the hands-on ship-driving competence the job required. The diagnosis was unusually unambiguous — two avoidable collisions in one year pointing at the same eroded fundamentals — which is what converted a long-tolerated shortfall into a mandate for structural change.<sup>o</sup>

## THE INTERVENTION

Beginning in 2018, the Navy restored the Surface Warfare Officers School from self-study to classroom-plus-simulator instruction, established Mariner Skills Training Centers on both coasts, set out a threefold increase in initial ship-driving training hours by 2021, and created ten pass-or-fail assessments across an officer's career path — four of them no-go gates, at the division-officer, department-head, command and major-command milestones, that can halt advancement. The no-go gates were the structural teeth: by tying advancement to demonstrated competence rather than time served, they made the qualification something the system would stop on, not merely something it recorded.<sup>1</sup>

## HOW IT WORKED

The reform paired technical investment — simulators, restored curricula, qualification gates — with a cultural change borrowed from aviation: structured debriefing and explicit gate


ownership. New Ready-for-Sea Assessments evaluated forward-deployed ships against a deliverable standard rather than a paper one; three of the first eighteen ships assessed were immediately sidelined as not ready. That a sixth of the first cohort failed against the new standard showed the assessment had teeth — it measured demonstrated readiness rather than accepting the paper certifications that had masked the pre-collision gap.<sup>2</sup>

#### THE EVIDENCE

Here the case turns instructive. The Government Accountability Office found in November 2019 that “the Navy has not put key processes and assessments in place to evaluate comprehensively the effectiveness of its changes to ship-driving training.” The structural intervention happened and the investment was large, but decision-grade evidence on outcomes is incomplete — the measurement to confirm the effect has lagged the change itself. Without an outcome time-series, the Navy cannot yet distinguish a real capability gain from activity that looks like one, which is the same blindness, in milder form, that let the pre-collision gap go unseen.<sup>3</sup>

#### WHAT TRANSFERRED

The reform is a live, in-progress success: structural change real, evidence of effect outstanding. As a teaching case it argues that mature capability engineering must build the measurement infrastructure from the start — the time-series that lets an institution know whether the capability it bought is materializing — rather than treating evaluation as an afterthought. The two halves the reform did deliver, technical investment and a cultural change, are necessary but not sufficient; the missing third half is the evidence regime that would let the Navy prove, not assume, the gap has closed.<sup>4</sup>



*The Navy has not put key processes and assessments in place to evaluate comprehensively the effectiveness of its changes to ship-driving training.*

GAO-20-154, NAVY READINESS: ACTIONS NEEDED TO EVALUATE THE EFFECTIVENESS OF CHANGES TO SURFACE WARFARE OFFICER TRAINING, NOVEMBER 2019

#### REFERENCES

1. Navy and NTSB reports on the Fitzgerald and McCain collisions (2017–2019) — the training-degradation antecedent.
2. Readiness Reform Oversight Council, One-Year Report (2019) — restored training, assessments, and gates.
3. USNI News reform coverage (2020, 2022) — Ready-for-Sea Assessments (15 of 18 FDNF–Japan ships assessed ready for sea,
4. GAO-20-154, Navy Readiness: Actions Needed to Evaluate the Effectiveness of Changes to Surface Warfare Officer Training (November 2019) — the absent evaluation processes and the planned threefold increase in ship-driving training hours.
5. SWOS Norfolk and San Diego Mariner Skills Training Center documentation — simulators and curriculum.

the remaining three sidelined) and the surface community's accounting of ten milestone assessments with four go/no-go gates.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Navy reform is a paired intervention in progress: technical (training restored, simulators procured, assessments created) plus cultural (debriefing, gate ownership). What it lacks is the third half: evidence that the intervention has worked. LENS treats this as a teachable case for what mature capability engineering should include from the start — the measurement infrastructure that lets the institution know whether its investment is producing the capability it bought.

LENS APPROACH

LENS uses Navy SWO reform in LEN 10 (capstone) as a worked exercise in capability intervention at scale, and in LEN 4 as a case where the measurement infrastructure has lagged the intervention. Students design the evidence regime that should accompany the reform.

| DOMAIN                     | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|-----------------|------------------|--------------|---------|-----|
| DEFENSE                    | → T K N         | → D1 D2 D3 D4 D5 | 4            | 2.3     | 4   |
| Iterative Development + D4 |                 |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pair the technical investment — restored curricula, simulators, training hours — with a cultural change such as structured debriefing and explicit gate ownership, rather than buying tools alone.
- ▶ Install no-go gates that tie advancement to demonstrated competence, so qualification is something the system will stop on, not merely something it records.
- ▶ Assess units against a deliverable standard from the first cohort, accepting that a real test will sideline some — that failure rate is the evidence the gate has teeth.

IN OPERATION / AFTER

- ▶ Build the measurement infrastructure from the start — the outcome time-series that lets the institution distinguish a real capability gain from activity that merely looks like one.
- ▶ Give the senior commander a readiness dashboard that reports demonstrated capability over time, so the same blindness that masked the pre-collision gap cannot recur.
- ▶ Treat evidence of effect as the third half of the intervention, not an afterthought, so the reform can be proven rather than assumed to have closed the gap.

REFLECTION QUESTIONS

1. Navy SWO reform happened without an evidence regime to confirm it worked. Identify a current reform in your domain and the evidence regime that should accompany it.
2. Design the dashboard the Chief of Naval Operations should have to know whether SWO capability is improving in time-series.
3. The Ready-for-Sea Assessment sidelined three of the first eighteen ships because it measured readiness against a deliverable standard. Identify a certification in your domain that ratifies rather than tests, and design the gate that would be willing to fail a sixth of its first cohort.

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

# DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks

**H** Human–System Interface **K** Knowledge & Institutional Memory **D** Designed Out

**IMPACT** An IDA independent assessment found that, after 16 weeks of Digital Tutor instruction, U.S. Navy IT graduates with no prior IT experience outscored fleet Information Systems Technicians with an average 9.1 years of experience on a knowledge test, with an effect size of 4.30, and outperformed them on most troubleshooting and design tasks

## IN BRIEF

DARPA's Digital Tutor program asked whether a one-on-one intelligent tutoring system, modelled on expert human tutoring, could compress years of operational IT expertise into a 16-week pipeline. The independent evaluation by the Institute for Defense Analyses (Morrison and Fletcher, IDA Document D-4686, 2012) compared Digital Tutor graduates — U.S. Navy enlistees with no prior IT experience — against fleet Information Systems Technicians with an average 9.1 years of experience. The Digital Tutor cohort outscored fleet ITs on a knowledge test with an effect size of 4.30 and outperformed them on most troubleshooting and design tasks; only the Security exercise produced a fleet advantage. The IDA report concludes the program “appears to have achieved its goals.” Two hedges survive into the case: knowledge “accounts for about 40 percent of practical-exercise performance variance” and is “an enabler of performance rather than a direct measure of performance itself,” and the system-architecture detail in the available documentation is too scant to fully reproduce. The case is the canonical small-tier instance of compressing the capability envelope at the edge of training, paired with CIRCUIT (Cases 78 and 68) on the workforce-capability-at-the-edge axis.

## BACKGROUND

The U.S. Navy's Information Systems Technician rating has a conventional pipeline: an A-school of several months, followed by years of fleet experience that turn the rated sailor into an operational troubleshooter. The capability that matters at the operational end — diagnose a networking failure under time pressure, recover a degraded system, design a workable configuration for an unfamiliar requirement — is conventionally treated as a thing seat time produces. DARPA's Digital Tutor program asked whether an intelligent tutoring system, modelled on the discipline of expert one-on-one human tutoring, could compress that pipeline into 16 weeks of structured instruction.<sup>o</sup>

## THE INTERVENTION

The program's design choice was to model the system on what expert human tutors actually do: a continuous dialogue around authentic problems, with the tutor pulling the trainee

toward the conceptual move that resolves the situation. The instructional sequence was built around troubleshooting and design problems drawn from the operational domain, not around a syllabus reconstructed from the existing course. The working hypothesis was that the tutorial discipline — not the content coverage — was what produced operational expertise, and that a sufficiently capable system could deliver enough of that discipline at scale to be useful as a training pipeline.<sup>1</sup>

#### HOW IT WORKED

The Institute for Defense Analyses (Morrison and Fletcher, IDA Document D-4686, 2012) ran the sponsor-commissioned evaluation that the case rests on. Digital Tutor graduates — Navy enlistees with no prior IT background, 16 weeks in — were compared against a sample of fleet Information Systems Technicians with an average 9.1 years of operational experience. On a knowledge test, the Digital Tutor cohort outscored fleet ITs with an effect size of 4.30 — a magnitude that is unusual in workforce L&D research and that the report treats as the headline finding. On troubleshooting and design tasks the Digital Tutor cohort outperformed the fleet sample on most exercises; the Security exercise was the exception where fleet experience showed.<sup>2</sup>

#### THE EVIDENCE

The IDA report concludes the effort “appears to have achieved its goals,” and the language is deliberate. Two hedges survive into the case verbatim. First, knowledge “accounts for about 40 percent of practical-exercise performance variance” and is “an enabler of performance rather than a direct measure of performance itself” — so the knowledge-test effect size, as large as it is, is not the same as the operational capability the Navy actually buys. Second, the available program documentation is too scant in system-architecture detail to fully reproduce. The result is teachable; the engineering recipe is not yet open.<sup>3</sup>

#### WHAT TRANSFERRED

What the case carries for the corpus is the capability-envelope argument at the edge of training (induced I.2, LENS D2/PT4). The conventional pipeline assumes operational expertise is a function of seat time and exposure. The Digital Tutor evaluation is evidence that the envelope is reachable substantially faster than the inherited course assumed — under one rating, one program, one evaluation — and the hedges name what the evidence does and does not close. Paired with CIRCUIT (Cases 78, 68), the case anchors the workforce-capability-at-the-edge-of-training axis that connectomics proofreading and submarine-system troubleshooting share at the structural level.

*The Digital Tutor cohort outscored fleet ITs with 9.1 years' experience on the knowledge test at an effect size of 4.30; the hedge is that knowledge accounts for about 40 percent of practical-exercise variance.*

EDITORS' SYNTHESIS OF MORRISON & FLETCHER (2012), IDA DOCUMENT D-4686.

#### REFERENCES

1. Defense Advanced Research Projects Agency, Digital Tutor program documentation — program description and design rationale.
2. Fletcher, J. D. (2009). From behaviorism to constructivism: a philosophical journey from drill and practice to situated learn-
3. Morrison, J. E., & Fletcher, J. D. (September 2012). DARPA Digital Tutor: Assessment Data. IDA Document D-4686 (prepared for DARPA under contract DASW01-O4-C-0003) — the sponsor-commissioned evaluation the case rests on.



ing. — methodological grounding for the Digital Tutor’s tutorial discipline.

4. Anderson, J. R., Corbett, A. T., Koedinger, K. R., & Pelletier, R. (1995). Cognitive tutors: Lessons learned. *Journal of the Learning Sciences*, 4(2):167–207. doi:10.1207/s15327809jls0402\_2 — the broader intelligent-tutoring evidence base the Digital Tutor program sits within.

THE LEARNING ENGINEERING LENS

LE INSIGHT

DARPA’s Digital Tutor is the strongest available evidence that the capability envelope of a training pipeline can be re-specified — from years of seat time to 16 weeks of tutorial-discipline instruction — against an operational comparison the program is built to compete with. The hedges (knowledge as enabler, architecture detail scant) are part of what makes the result interpretable.

LENS APPROACH

Digital Tutor is the canonical workforce-capability-at-the- edge-of-training case (induced 1.2; LENS D2/ PT4). LENS uses it in Domain 2 (Iterative Development) for the tutorial-discipline-as-instructional-artifact design move, and in Domain 4 (Test and Evaluation) for the operational-comparison evaluation against fleet ITs with 9.1 years of experience. Pair with CIRCUIT (Cases 78, 68) at the workforce-capability-at-the-edge axis — connectomics proofreading and Navy IT troubleshooting share the structural pattern of compressing operational expertise through tutorial discipline.

| DOMAIN               | Modes Addressed | LENS Competency       | Problem Type | Induced | LEO   |
|----------------------|-----------------|-----------------------|--------------|---------|-------|
| DEFENSE              | H K D           | D1 D2 D3 D4 D5        | 4            | 1.2     | 2 · 4 |
| WORKFORCE L&D        | →               | Iterative Development |              |         |       |
| INTELLIGENT TUTORING |                 |                       |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the operational capability the pipeline must produce in the language of the work — troubleshoot under time pressure, design a workable configuration — not in the language of the existing course’s content coverage.
- Treat the tutorial discipline (continuous dialogue around authentic problems, pulling toward the resolving conceptual move) as the instructional artifact that did the work, rather than the content sequence the legacy course inherited.
- Design the evaluation against the operational comparison the program is built to compete with — for Digital Tutor, fleet ITs with 9.1 years of experience — so the result speaks to the capability envelope, not to a within-program improvement.

IN OPERATION / AFTER

- Report the knowledge-test effect (4.30) and the practical-exercise variance hedge (knowledge accounts for 40%) together; the headline number is real, and the qualification that knowledge is an enabler rather than a direct measure is part of what makes the result interpretable.
- Treat the absence of reproducible architecture detail as program documentation, not as polish: future builders need the engineering recipe, and the next iteration of evidence is conditional on that recipe being available.
- Carry the result into adjacent capability-envelope debates — CIRCUIT proofreading, submarine-system troubleshooting — as evidence that the envelope is reachable faster than the inherited training assumption, under one program and one evaluation.

## REFLECTION QUESTIONS

1. Identify a capability in your domain where the inherited training assumption is that operational expertise is a function of seat time. What would a Digital-Tutor-class re-specification — tutorial discipline around authentic problems — look like, and what is the operational comparison your evaluation would have to beat?
2. The Digital Tutor knowledge-test effect size is 4.30 and the report still hedges that knowledge accounts for only 40% of practical-exercise variance. What is the analog hedge in your context: which part of the capability does your evaluation measure directly, and which part is enabler rather than direct measure?
3. The architecture detail is too scant for outside builders to reproduce the system. What is the minimum engineering recipe you would publish alongside a similar result, so that the next iteration of evidence rests on a reproducible base rather than a one-off program?

**WHO BUILDS THIS** human factors & interaction design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · task analysis · design prototyping

# xAPI / Total Learning Architecture — Interoperability Gap

**K** Knowledge & Institutional Memory   **G** Governance & Trust

**IMPACT** Despite xAPI adoption, most implementations remain siloed; cross-organizational learning-data interoperability largely unrealized

## IN BRIEF

The Experience API (xAPI) was created to track learning experiences across platforms, and the Advanced Distributed Learning Initiative's Total Learning Architecture envisioned learning records, competencies, and credentials flowing across organizational boundaries to support continuous capability development — the kind of evidence infrastructure LENS depends on. In practice, xAPI implementations remain largely siloed inside individual learning-management systems; the cross-organizational data sharing that matters most for high-consequence domains has not materialized at scale. The technical standard exists and reference implementations exist; what lags is governance — who owns the data, what consent applies, how quality is assured. xAPI mirrors inBloom (Case 53) in structure — technology ahead of governance — across the whole learning-technology ecosystem. It is the book's case for treating data governance as a capability deliverable.

## THE SHIFT

Learning increasingly happens across many systems — courses, simulators, on-the-job tools — and the field recognized it needed a common way to record those experiences, so capability could be tracked over a career rather than a single course.<sup>o</sup> The career-long view is the point: a single course's records say little about whether a person can do the job, whereas experiences stitched across courses, simulators, and on-the-job tools are what let an institution reason about real capability rather than completed seat-time.

## WHAT IS EMERGING

The Experience API (xAPI) was built to do exactly that, and the Advanced Distributed Learning Initiative's Total Learning Architecture envisioned learning records, competency frameworks, and credentials flowing across organizational boundaries — the evidence infrastructure a discipline like LENS depends on.<sup>1</sup> The vision was explicitly cross-boundary: records, competencies, and credentials that move with the learner between organizations are precisely the evidence base on which a capability discipline must stand, which is why the standard's promise mattered well beyond any single training shop.

## THE CAPABILITY QUESTION


In practice, xAPI implementations remain largely siloed inside individual LMS platforms. The cross-organizational data sharing — the capability most relevant to high-consequence operational domains — has not materialized at scale. The standard exists; the ecosystem it promised does not.<sup>2</sup> The siloing is the diagnostic detail: the very cross-boundary flow that made the architecture worth building is the part that did not arrive, so the standard delivers tracking within each LMS while the career-long, cross-organizational record it envisioned stays out of reach.

### EARLY EVIDENCE

The data model is sound and reference implementations exist. What has lagged is semantic conformance — ADL's own acquisition guidance warns that xAPI data not following published Profiles “will have interoperability issues outside of the implementing organization” — and governance: who owns the data, what consent frameworks apply, and how data quality is assured across organizations.<sup>3</sup> With the data model proven and reference implementations in hand, the remaining obstacles are not engineering questions an organization can solve alone but agreements between organizations — ownership, consent, and assured quality — that no technical specification can settle on their behalf.

### OPEN PROBLEMS

xAPI mirrors inBloom (Case 53) in essential structure — technology in advance of governance — but across the whole learning-technology ecosystem rather than one initiative.<sup>4</sup> It is the book's clearest case that an interoperability standard is necessary but not sufficient: without the governance to make organizations willing and able to share, the data stays in its silos, and the evidence infrastructure remains a diagram rather than a system. The parallel to inBloom is instructive because it shows the pattern is not particular to one failed project: a sound standard arriving ahead of the ownership, consent, and quality-assurance arrangements will stall the same way at any scale.



*The standard exists. The governance does not.*

EDITORS' SYNTHESIS OF ICICLE / ADL TLA DISCUSSION

## REFERENCES

1. Advanced Distributed Learning Initiative, Total Learning Architecture documentation — the cross-boundary vision.
4. ADL Initiative, Total Learning Architecture Standards: Digital Learning Acquisition Techniques (December 2023) — xAPI data

2.

IEEE Std 9274.1.1-2023, Standard for Learning Technology — JSON Data Model Format and RESTful Web Service for Learner Experience Data Tracking and Access (xAPI 2.0) — the technical standard of record; the IEEE LTSC project was authorized in May 2018, board-approved 30 March 2023, and published 6 October 2023.
3.

B. Smith and L. Milham, Total Learning Architecture (TLA) Data Pillars and Their Applicability to Adaptive Instructional Systems (ADL Initiative, 2021; ERIC ED628093) — ADL’s own observation that while most LRS solutions offer dashboards, “they are commonly used to view xAPI statements that have been generated within a single learning activity,” against a design intent of data shared “vertically within an organization, horizontally across other organizations, or longitudinally throughout a career.”
5.

that does not follow xAPI Profiles “will have interoperability issues outside of the implementing organization.”  
Cf. inBloom (Case 53) — technology in advance of governance.

THE LEARNING ENGINEERING LENS

LE INSIGHT

xAPI/TLA is the technical-standard analog of the implementation gap — the discipline has the data model, the spec, and the reference implementations. What it does not have is the governance and institutional architecture to make cross-organizational learning data flow. This is the canonical case in this book for treating data governance as a capability-engineering deliverable.

LENS APPROACH

LENS treats xAPI/TLA in LEN 4 as a data-governance and interoperability case and in LEN 8 as an example of organizational- learning infrastructure that has not scaled. The case is the technical underlay to the larger argument about evidence systems that decision-makers can trust.

| DOMAIN             | MODES AT STAKE  | LENS COMPETENCY                         | PROBLEM TYPE | INDUCED | LEO |
|--------------------|-----------------|-----------------------------------------|--------------|---------|-----|
| EDUCATION AT SCALE | → <div>KG</div> | → <div>D1D2D3D4D5</div>                 | 4            | 5.3     | 5   |
| DEFENSE            |                 | → Navigating Sociotechnical Constraints |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the governance — data ownership, consent, and cross-organization quality assurance — as a deliverable alongside the technical standard, not after it.
- Design for the cross-boundary flow from the start, so the standard is not merely implemented inside each LMS but engineered to move records between organizations.
- Secure the inter-organizational agreements that no specification can settle, treating willingness-to-share as a thing to be built rather than assumed.

IN OPERATION / AFTER

- Audit whether learning data is actually crossing organizational boundaries, not just whether xAPI is nominally adopted within each silo.
- Monitor data quality across organizations continuously, since shared records are only trustworthy evidence if their quality is assured at the seams.
- Watch for the inBloom pattern — a sound standard outrunning its governance — and treat any stall as a missing-governance signal, not a technical defect.

## REFLECTION QUESTIONS

1. Why has the xAPI standard not produced cross-organizational interoperability at scale? What governance condition is missing?
2. Design the minimum governance architecture under which xAPI data could flow across two organizations in your domain.
3. The xAPI data model is sound; what stalled was ownership, consent, and quality assurance across organizations. For a data-sharing effort in your domain, which of those three is the unresolved question — and who would have to agree for the data to actually flow?

**WHO BUILDS THIS** knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · governance frameworks · bias & evidence audits

Chapter 9

# Industry, Energy & Enterprise Systems — What Fails

*When incentive, inertia, and migration risk go unengineered.*

---

*The signal arrived years early. The iteration never ran.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

## EDITOR'S NOTE · THE ITERATION GAP

Iterative development is what learning engineering does most often. It is also the capability the casebook documents least directly. Several cases in this chapter — and their counterparts across the book's failure halves — share a shape: an organization received an unambiguous signal, held the engineering or product position to act on it, and failed to run the iteration cycle through product, business model, and operations until the opportunity had passed. The signal was sometimes a field report (Therac-25, the adoption and sustainment cases), sometimes an external launch (BlackBerry meeting the iPhone), sometimes an internal invention the organization could not build a business around (Kodak's 1975 digital camera). In each case the organization had the engineering capability and the institutional position. What it lacked was the discipline of running the iteration through.

The casebook's broader corpus shows what iteration produced — the interventions collected in every part's What Works chapter, the failures whose recovery iterations followed. It does not yet show iteration **being done** in detail: the design moves, the dead ends, the reframes. That gap is exactly what the v2.1 amendment to D2 names — “narrate and defend the design iteration in first person” — and what the editor-commissioned first-person Practice Flywheel accounts will fill in the next edition. This note names the gap honestly so that the reader knows what to expect from the volume, and what the discipline is still working to make legible.



# Knight Capital Trading Loss

**D** Designed Out   **K** Knowledge & Institutional Memory

**IMPACT** \$440M loss in 45 minutes; firm forced to seek emergency capital and was effectively acquired within months

## IN BRIEF

On 1 August 2012 Knight Capital, a major U.S. market maker, deployed new order-routing software to seven of its eight servers and missed the eighth. The new code reused a flag that, on the eighth server's old software, reactivated long-dead "Power Peg" code never removed from the repository. At the opening bell it fired millions of unintended orders; in about 45 minutes Knight lost roughly \$440 million — the SEC order puts the loss on the positions at over \$460 million — and the firm, left with what the order calls "significant net capital problems," survived only on a \$400 million rescue financing and was effectively acquired within months. The SEC found Knight had no procedure to verify the deployment across all servers and no controls to halt the runaway orders. The capability designed out was deployment verification; the dead code was technical debt that exercised its option at the worst possible moment.

## BACKGROUND

Knight Capital was one of the largest market makers in U.S. equities, routing an enormous share of retail order flow through automated systems wired directly into the market. On 1 August 2012 the New York Stock Exchange was launching a new Retail Liquidity Program, and Knight had updated its order-routing software to participate in it from the opening bell. The update went out to production the way countless updates had gone before — a routine deployment, on an ordinary morning, handled as unremarkable.<sup>0</sup>

## WHAT HAPPENED

A technician deployed the new code to seven of the eight routing servers and missed the eighth, leaving one node running stale software. The new code reused a configuration flag that, on the old software still on that eighth server, had once activated the long-dormant "Power Peg" routing function — discontinued in 2003 but left present and callable. At the open the dead code woke and began firing millions of unintended orders into the market. In about forty-five minutes Knight amassed a vast unwanted position — long about \$3.5 billion in 80 stocks and short about \$3.15 billion in 74 — and realized a pre-tax loss of about \$440 million. Its shares fell from \$10.33 to \$2.58 in two sessions; it survived only on a \$400 million rescue financing convertible at \$1.50 a share, which left the existing shareholders a minority, and it was effectively acquired within months.<sup>1</sup>

## THE INVESTIGATION

The Securities and Exchange Commission's October 2013 enforcement order — In re Knight Capital Americas LLC — found Knight had no written procedure requiring a second technician to verify the deployment across all hosts, no written code-development and deployment procedures for the router, and no written protocol governing access to unused code — nor controls able to recognize and halt the flood of erroneous orders once it began. The order catalogued specific violations of the SEC's Market Access Rule (Rule 15c3-5), and Knight settled for a \$12 million penalty.<sup>2</sup> The dead Power Peg code was the proximate trigger and the reused flag the match that lit it; but the underlying cause was the absence of the verification and risk controls that should surround any change to a system wired directly into the live market. The order finds that Knight "did not have technology governance controls and supervisory procedures sufficient to ensure the orderly deployment of new code or to prevent the activation of code no longer intended for use," and that its firm-wide capital thresholds were never linked to controls that could stop the orders being entered.<sup>3</sup>

## THE CAPABILITY GAP

The capability designed out was deployment verification — confirming, every time, that what runs in production is exactly what was intended, on every node, before it touches live money. The dead code was technical debt in its most literal form: a retired function left in the repository is a standing option on a future failure, and reusing its flag exercised that option at the worst possible moment. As at the Mars Climate Orbiter's interface, the boundary that mattered — between "deployed" and "verified as deployed everywhere" — had no owner and no automated check, and a large institution kept walking across it routinely until the day the floor gave way.<sup>4</sup>

## AFTERMATH & REFORM

Knight became the canonical case in modern software-operations practice for why deployment is itself an engineering deliverable, not a clerical step: automated verification that every host runs the intended build, disciplined removal of dead code from the repository, pre-trade risk limits, and kill switches that can stop a runaway process in seconds. Regulators sharpened their attention to automated market-access controls in response. Regulation Systems Compliance and Integrity (Regulation SCI), adopted 19 November 2014, extended formal system-development, testing, change-management, and incident-reporting obligations to registered exchanges, clearing agencies, FINRA, plan processors, and the larger alternative trading systems — the institutional response that wrote the deployment discipline Knight had lacked into a rule. It did not, however, reach Knight's own class: a broker-dealer market maker is not an SCI entity, and remains governed on this point by the Market Access Rule Knight was charged under. The lesson rhymes with the orbiter's across a forty-year, civilian-to-financial gap: a small, unowned boundary inside a large automated system is precisely where the institution is most exposed, and least watching — and as algorithmic trading continued to grow after 2012, the same class of fragility kept producing new incidents that the Knight template explained.<sup>5</sup>

# \$440M

LOSS IN 45 MINUTES · 7 OF 8 SERVERS UPDATED

*dead code on the eighth server; repurposed flag, no deployment verification*

KNIGHT CAPITAL — THE COST OF AN UNVERIFIED DEPLOYMENT

*Knight did not have written procedures requiring a second technician to verify the deployment.*

PARAPHRASING THE SEC ORDER AGAINST KNIGHT CAPITAL, 2013

## REFERENCES

1. U.S. Securities and Exchange Commission, Order Instituting Administrative and Cease-and-Desist Proceedings, In re Knight Capital Americas LLC (2013) — the firm, the Retail Liquidity Program launch, and the deployment.
2. SEC Order (2013) and Knight Capital Group 8-K filings of 2 and 6 August 2012 — the missed eighth server, the reactivation of the dormant “Power Peg” code, the 45-minute event, and the \$400 million rescue financing (preferred convertible at \$1.50 a share); the SEC puts the loss on the positions at over \$460 million, Knight’s own filing at a realized pre-tax loss of \$440 million.
3. SEC Order (2013), Rel. 34-70694, 16 October 2013 — the absence of a second-technician deployment review, the missing written code-development, deployment and unused-code procedures (§126), the capital thresholds unlinked to automated controls, and the \$12 million penalty (quoted).
4. B. Beyer, C. Jones, J. Petoff & N. R. Murphy (eds.), Site Reliability Engineering (O’Reilly, 2016) — deployment verification, dead-code removal, and automated safeguards as engineering deliverables.
5. SEC Market Access Rule (Rule 15c3-5) and subsequent automated-controls guidance — the regulatory response on pre-trade risk and market-access controls.
6. D. Seven, “Knightmare: A DevOps Cautionary Tale” — a widely cited engineering post-mortem on the deployment process, the orphaned eighth server, and the reused feature flag.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Knight Capital is the financial-industry version of Mars Climate Orbiter (Case 98): a small, unspecified boundary inside a large system that took the institution down. The capability that was missing was deployment verification. The dead code was the proximate trigger; the absent procedure was the cause.

### LENS APPROACH

Knight Capital is the canonical inherited-dead-code case (induced 7.1, re-verifying inherited designs in new envelopes, with change control and disclosure — 5.4 — as the post-incident Regulation SCI lens; LENS D1/PT3). LENS uses it in LEN 5 to teach deployment-as-capability — students design the deployment deliverable that would have caught the eighth server — and in LEN 9 for the technical-debt argument: every line of dead code carries an option on a future failure. Adjacent to Mars Climate Orbiter (Case 98) at the small-boundary-no-owner layer and to Regulation SCI as the institutional response that codified the missing controls.

| DOMAIN           | FAILURE MODES   | LENS COMPETENCY         | PROBLEM TYPE | INDUCED | LEO |
|------------------|-----------------|-------------------------|--------------|---------|-----|
| TECHNOLOGY       | → <div>DK</div> | → <div>D1D2D3D4D5</div> | 3            | 7.1     | 1·4 |
| Systems Analysis |                 |                         |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div><div>▶ Remove retired code and dormant flags from the repository as a matter of discipline, since every dead function left behind is a standing option on a future failure.</div><div>▶ Never reuse a configuration flag whose old meaning still exists somewhere in the fleet; treat flag semantics as a versioned, owned interface.</div><div>▶ Design pre-trade risk limits and a kill switch into any system wired to the live market, so a runaway process can be stopped in seconds rather than minutes.</div></div> | <div><div>▶ Require automated verification that every host runs the intended build before a change touches live money, with a second technician confirming the deployment across all nodes.</div><div>▶ Monitor for the runaway condition — an abnormal order rate — and halt automatically, rather than relying on humans to notice and intervene mid-flood.</div><div>▶ Audit deployments for the orphaned node: confirm no host is left on stale code, closing the gap between “deployed” and “verified as deployed everywhere.”</div></div> |

REFLECTION QUESTIONS

1.

Identify a deployment procedure in your domain whose verification step depends on convention rather than on a designed check. What is the eighth server?
2.

Design the deployment deliverable that would prevent a Knight Capital-equivalent loss in your organization.
3.

A retired “Power Peg” function left in the repository became the trigger years later when its flag was reused. Identify dead code or a dormant feature flag in a system you run that is still an option on a future failure, and specify the policy that should have removed it.

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS task analysis · design prototyping · knowledge bases · data pipelines

# Kodak Digital Camera Stagnation

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Inventor of the digital camera (1975) and holder of foundational digital-imaging patents; shipped digital cameras from 1991 and led the consumer market at points in the 2000s, and still filed for Chapter 11 bankruptcy in January 2012, unable to rebuild the business around margins the film franchise had set

## IN BRIEF

In 1975 a young Kodak engineer, Steven Sasson, built the first self-contained digital camera in a Rochester lab: an 8-pound, 0.01-megapixel device that wrote a black-and-white image to cassette tape in 23 seconds. By the 1990s Kodak held foundational digital-imaging patents and had a working line of digital products. Yet the company that invented the category never ran the iteration cycle through to a digital-first business: film and the high-margin consumables it sold remained the protected core, and digital was treated as a threat to be managed rather than a product to be developed against itself. By 2003 digital camera sales overtook film in the United States; in January 2004 Kodak announced the elimination of 12,000 to 15,000 positions; in January 2012 the company filed for Chapter 11. The capability gap was not invention. It was the organization's failure to iterate the business around the technology it already owned.

## BACKGROUND

Kodak's twentieth-century position was anchored on a single durable insight: the camera was a loss-leader and the film, paper, and processing chemistry were the recurring high-margin revenue. The business model — what later analysts would call the "razor-and-blades" of imaging — funded a research apparatus deep and well-resourced enough that in December 1975 a 24-year-old Eastman Kodak engineer, Steven Sasson, was able to assemble the first all-electronic still camera from a Fairchild CCD, a cassette-tape recorder, sixteen nickel-cadmium batteries, and a lens cannibalized from a Super 8 movie camera. The prototype wrote a 0.01-megapixel black-and-white image to tape in 23 seconds and played it back on a television.<sup>o</sup>

## WHAT HAPPENED

Sasson's later published recollection on Kodak's own corporate blog (2007) is that internal reception was polite but unmoved: the demonstration raised the question of when the company would have to compete against its own film business, and the answer settled into a long institutional hedge. Through the 1980s and 1990s Kodak filed foundational digital-imaging patents, built digital products including the 1991 DCS-100 professional digital SLR, and partnered with Apple on the early QuickTake consumer line; the technology was under-

stood and the engineering position was real. What was missing was an iteration cycle that ran the digital signal all the way through product, distribution, pricing, and the corporate identity, rather than holding digital at the perimeter of a film company.<sup>1</sup>

#### THE INVESTIGATION

Paraphrasing Lucas & Goh (Journal of Strategic Information Systems, 2009): Kodak's leadership read digital primarily as a cannibalization threat to film rather than as an opportunity to be engineered. The retrospective record describes successive strategies — digital as a complement to film, digital kiosks at retail, Ofoto / Kodak Gallery for online printing, the EasyShare consumer line — each of which was a move toward digital but none of which committed the operating company to a digital-first product roadmap with the cadence and capital allocation the transition required. The cycle that a healthy iteration would have run — ship, measure, learn, ship again at higher ambition — was repeatedly truncated by the gravitational pull of the film P&L it would have had to displace.<sup>2</sup>

#### THE CAPABILITY GAP

By 2003 digital camera sales overtook film camera sales in the United States; by the mid-2000s smartphone cameras began collapsing the standalone-camera category Kodak had partially entered. Kodak's revenues fell from roughly USD 16 billion in 1996 to under USD 6 billion by 2011. The company announced the elimination of 12,000 to 15,000 positions in January 2004, sold its Health Imaging division in 2007, and on 19 January 2012 filed for Chapter 11 bankruptcy protection in the Southern District of New York. In the bankruptcy proceedings Kodak monetized roughly 1,100 of its digital-imaging patents — many traceable to the engineering lineage Sasson's 1975 prototype had started — for approximately USD 525 million, a sale that drew explicit press commentary on what it meant for a company to monetize patents on a category it had failed to commercialize.<sup>3</sup>

#### AFTERMATH & REFORM

Kodak is the canonical case of an organization that held the technological position and lost on the iteration discipline. The capability gap was not in the lab and was not in the patent portfolio; it was in the leadership system that could not bring itself to run the iteration cycle all the way through the business model it would have replaced. The instructive contrast is Fujifilm, which over the same period diversified into chemicals, materials, and pharmaceuticals on the back of its film-coating expertise — a deliberate organizational iteration against the same signal. The lesson the v2.1 framing names is that "iterative development" at the organizational level is not a product practice but a leadership practice, and the cycle has to run on the business model and the operations as well as on the artifact.<sup>4</sup>

# 1975

FIRST DIGITAL CAMERA PROTOTYPE · 0.01 MEGAPIXEL

*37 years from prototype to Chapter 11*

KODAK — A SIGNAL ACTED ON PARTLY, NEVER ITERATED THROUGH

*You don't get to keep the future just because you invented it.*

EDITORS' SYNTHESIS OF THE SASSON (IEEE SPECTRUM 2007) AND LUCAS & GOH (JSIS 2009) RECORD

REFERENCES

1.

S. Sasson, "We Had No Idea," Kodak corporate blog, Kodak: Plugged In (2007) — first-hand account of the 1975 prototype, the demonstration, and the internal reception.

2.

M. Spector & D. Mattioli, "Kodak Teeters on the Brink," The Wall Street Journal (5 January 2012); follow-up coverage of the patent-sale process in WSJ and Reuters (December 2012).

3.

H. C. Lucas Jr. & J. M. Goh, "Disruptive technology: How Kodak missed the digital photography revolution," Journal of Strategic Information Systems 18(1):46–55 (March 2009).

4.

Eastman Kodak Company, Voluntary Petition for Chapter 11 Reorganization, U.S. Bankruptcy Court, Southern District of New York, Case No. 12-10202 (19 January 2012).

5.

A. R. Sorkin & M. J. de la Merced, "Eastman Kodak Files for Bankruptcy," The New York Times DealBook (19 January 2012).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Kodak is the cleanest available evidence that invention is not iteration. The company held the prototype, the patents, and the engineering depth and still lost the category. The gap was not technical; it was the leadership system's inability to run the iteration cycle through the business model it would have had to displace. v2.1 D2 names this directly: iteration-against-opportunity is an organizational discipline, not a lab discipline, and the cycle has to close on the P&L as well as on the artifact.

LENS APPROACH

Kodak is the v2.1 D2 stagnation exemplar (induced 2.3 transfer to high-consequence settings; LENS D2/PT4 adoption and sustainment; LEO-2 Iterative Development). LENS uses it to make the organizational-iteration claim concrete: 2.2 (run the cycle), 2.3 (transfer the cycle to the high-stakes business decision), and 2.4 (sustain adoption against the gravitational pull of the legacy P&L) all sat un-owned for a generation. Pair with Fujifilm as the in-category counter-iteration that ran the same signal differently.

| DOMAIN                | FAILURE MODES                            | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|-----------------------|------------------------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| INDUSTRIAL            | → <div>D</div> <div>K</div> <div>N</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 4            | 2.3     | 2   |
| Iterative Development |                                          |                                                                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat an internally demonstrated category-changing prototype as a signal to run the full iteration cycle on the business, not only on the artifact — including pricing, distribution, sales force, and identity.
- ▶ Name explicitly the existing high-margin P&L the new category will have to displace, and assign an executive sponsor whose job is to iterate against that displacement on a schedule rather than around it.
- ▶ Set a measured cadence — ship, instrument, learn, ship again at higher ambition — and protect the cadence against the organizational pull of the legacy business at each review cycle.

IN OPERATION / AFTER

- ▶ Audit, periodically and without sentimentality, whether the iteration is running through to the operating company or has been parked at the perimeter as research or partnership; treat parking as the failure mode it is.
- ▶ Compare against an in-category peer running the same iteration with a different leadership posture (Fujifilm), and use the divergence as evidence rather than letting the legacy P&L be the only reference point.
- ▶ Read the late monetization of patents on a category the organization never commercialized as the indictment it is — a signal that the engineering position was held and the iteration discipline was not.

## REFLECTION QUESTIONS

1. Identify a category-changing prototype or signal that your organization has internally demonstrated but parked at the perimeter. What would the full iteration cycle through the operating company actually require, and which P&L is the gravitational pull?
2. Kodak iterated on the artifact (digital products shipped) but not on the business model. Distinguish the two in a current initiative in your domain and specify which one is being held still.
3. Fujifilm ran the same signal differently. Identify an in-category peer running an iteration your organization is not, and construct the honest comparison the leadership review would have to engage with.

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards



# Northeast Blackout

**H** Human–System Interface    **K** Knowledge & Institutional Memory

**IMPACT** 50 million people without power across eight U.S. states and Ontario; \$4–10B U.S. economic cost; FERC Order 693 followed

## IN BRIEF

On 14 August 2003 a high-voltage transmission line in Ohio sagged into a tree and tripped — an event the grid should have absorbed. But FirstEnergy's control-room alarm system had failed silently just after 14:14, so operators did not know the line was gone. Further lines tripped, and a cascade swept across the Eastern Interconnection; within minutes 50 million people across eight U.S. states and Ontario lost power, at a U.S. cost the task force put at \$4 to \$10 billion. It found FirstEnergy lacked situational awareness, its alarm system had failed without notice, vegetation management was poor, and the reliability organizations gave no effective real-time diagnostics. The reforms made reliability standards mandatory and enforceable (FERC Order 693). The gap sat at the automation-operator boundary: a silent failure left operators blind.

## BACKGROUND

The Eastern Interconnection is built to ride through the loss of a single transmission line, and control rooms watch the grid through software and alarms. FirstEnergy, the Ohio utility at the center of the story, ran a control room whose alarm system had — unknown to anyone — failed just after 14:14 and was never restored.<sup>0</sup> The interconnection's single-line resilience assumes the operators can see which line is gone; a silent alarm broke that assumption at its root, leaving a grid designed to tolerate one fault blind to the fault it was tolerating.

## WHAT HAPPENED

On 14 August 2003 a 345 kV line in Ohio sagged into a tree and tripped. With the alarms silent, operators did not know the line was gone and so did not take the corrective steps that would have contained it. Further lines tripped and a cascade swept across the interconnection; within minutes 50 million people across eight states and Ontario lost power, at a U.S. cost estimated at \$4 to \$10 billion.<sup>1</sup> The first trip was the routine single-line loss the grid was built to absorb; what turned it into a cascade was not the tree but the hour in which the operators acted on a picture of the grid that no longer matched the grid.

## THE INVESTIGATION

The U.S.–Canada Power System Outage Task Force named “inadequate situational awareness at FirstEnergy” as a cause — “FE did not recognize or understand the deteriorating condition of its system” — along with FE's failure to manage tree growth and the reliability organizations' failure to provide effective real-time diagnostic support.<sup>2</sup> The reforms produced FERC Order 693, issued 16 March 2007, which approved 83 of the 107 reliability


standards NERC had proposed as mandatory and enforceable under the authority Congress created in section 215 of the Energy Policy Act of 2005.<sup>3</sup> Making the standards mandatory addressed the deeper finding that a voluntary regime had let vegetation management and operator awareness drift: when compliance is optional, the practices that prevent a cascade are exactly the ones a cost-pressured utility lets slip.

#### THE CAPABILITY GAP

The gap sat at the boundary between automation and operator. When the alarm system stopped doing its job, it did so silently — and the operators had no independent signal that they were losing control of their grid segment. The missing capability was the meta-monitor: the watch on the watchman, the check that tells you when the thing that is supposed to tell you has itself stopped.<sup>4</sup> Silence is the most dangerous failure mode because it is indistinguishable from a quiet, healthy grid; the operators were not ignoring a warning but trusting an absence of warning that the broken system could no longer guarantee.

#### AFTERMATH & REFORM

FERC Order 693 and the new mandatory reliability regime — enforced by an Electric Reliability Organization with audits and penalties — treated grid reliability as a deliverable rather than a best practice. The 2003 blackout endures as the canonical case of silent automation failure: a system that quits without announcing it, leaving the humans nominally in charge with nothing to act on.<sup>5</sup> Backing the standards with audits and penalties conceded that reliability could not be left to good intentions: the practices that would have kept the operators sighted had to be enforceable obligations, not best practices a utility could quietly defer.



*Inadequate situational awareness at FirstEnergy. FE did not recognize or understand the deteriorating condition of its system.*

U.S.-CANADA POWER SYSTEM OUTAGE TASK FORCE, FINAL REPORT ON THE AUGUST 14 2003 BLACKOUT, APRIL 2004

#### REFERENCES

1. U.S.-Canada Power System Outage Task Force, Final Report on the August 14, 2003 Blackout in the United States and Canada (2004) — the tree contact, the silent alarm, and the cascade.
2. Task Force (2004) — 50 million people affected across eight states and Ontario; the minute-by-minute sequence.
3. Task Force (2004), Causes Groups 1–4 — “inadequate situational awareness at FirstEnergy,” FE’s failure to manage tree
4. FERC Order No. 693, Mandatory Reliability Standards for the Bulk-Power System, 118 FERC ¶ 61,218 (16 March 2007) — 83 of 107 proposed standards approved as mandatory and enforceable, effective 18 June 2007.
5. M. R. Endsley (1995), situation-awareness theory — the human-factors frame for silent-automation failure.
6. North American Electric Reliability Council reports (2004) and the creation of the Electric Reliability Organization.

growth, and the reliability organizations’ failure to provide effective real-time diagnostic support (quoted).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The 2003 blackout is the canonical case for silent automation failure: a system that stops doing its job without telling its operators. The capability that was missing was the meta-monitor — the system that watches the monitor. The reform that followed treated grid-reliability compliance as a deliverable rather than as a best practice.

LENS APPROACH

LENS uses the 2003 blackout in LEN 2 as a Human-AI Teaming case for silent-automation-failure handling and in LEN 8 for the legislative-reform arc that produced enforceable reliability standards.

| DOMAIN                     | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|---------------|------------------|--------------|---------|-----|
| ENERGY →                   | H K           | → D1 D2 D3 D4 D5 | 6            | 3.3     | 3   |
| Human-System Collaboration |               |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design monitoring with a positive heartbeat so a silent system is distinguishable from a healthy one, never assuming no alarm means no problem.
- ▶ Build a meta-monitor that watches the alarm system itself and announces when the watchman has stopped.
- ▶ Give the regional reliability coordinator the authority and information to intervene across utility boundaries before a cascade forms.

IN OPERATION / AFTER

- ▶ Audit vegetation management and operator situational awareness against enforceable standards, since voluntary practice drifts under cost pressure.
- ▶ Test alarm and monitoring systems for detectable failure, verifying that an outage of the monitor itself raises an alert.
- ▶ Back reliability obligations with audits and penalties so the practices that keep operators sighted cannot be quietly deferred.

REFLECTION QUESTIONS

1. Identify an automated monitoring system in your domain whose silent failure would not be detected by its operators. How would they know?
2. Design the meta-monitor that should have been watching FirstEnergy’s alarm system in 2003.
3. Silence was indistinguishable from a healthy grid. What positive heartbeat signal in your domain would let operators tell a quiet system from a dead one?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS usability testing · cognitive walkthroughs · knowledge bases · data pipelines

# Takata Airbag Inflators

**D** Designed Out    **G** Governance & Trust

**IMPACT** 28 U.S. deaths confirmed by NHTSA and at least 400 U.S. injuries alleged from inflator ruptures; largest automotive recall in history

## IN BRIEF

Takata built its airbag inflators around ammonium nitrate — cheap and energetic but unstable, and used by no other major inflator maker — rivals built on guanidine nitrate. Takata's inflators carried no desiccant. Over years of heat and humidity the propellant degraded and could rupture the metal housing on deployment, spraying shrapnel at the driver; more than two dozen deaths and hundreds of injuries followed. Takata's own tests had shown the ruptures, but it reported them as isolated anomalies and, in places, falsified data; in 2017 it pleaded guilty to wire fraud, paid roughly \$1 billion, and went bankrupt. The recall — some 67 million inflators across 19 automakers in the U.S. alone — is the largest in automotive history. Two capabilities were designed out: the desiccant, and the independent verification regulators never had.

## BACKGROUND

An airbag inflates by burning a propellant inches from a person's face, so the propellant must be energetic **and** stable across the years and climates a car will see. To cut cost, Takata built its inflators around ammonium nitrate — cheap and energetic but notoriously unstable, shifting crystalline structure and absorbing moisture as temperature cycles day after day. Its competitors did not use ammonium nitrate at all, building on guanidine nitrate instead; Takata added a drying desiccant only to later inflator versions, and for years shipped with none.<sup>9</sup>

## WHAT HAPPENED

Over years of heat and humidity the propellant degraded, and a degraded charge can burn too fast, generating more pressure than the housing was built to contain. A firing inflator could then rupture its own metal housing and spray shrapnel into the cabin — turning the device meant to save a life into a fragmentation hazard aimed at the driver. NHTSA has confirmed 28 deaths in the United States, and at least 400 injuries are alleged. The recall grew to roughly 67 million inflators in the United States across some nineteen automakers, and more than 100 million worldwide — the largest in automotive history — and Takata went bankrupt; the toll keeps rising as unrepaired vehicles stay on the road and their inflators keep aging.<sup>1</sup>

## THE INVESTIGATION

Takata's own testing had shown the ruptures, but internal documents revealed engineers raising the alarm internally while the company reported the failures to automakers and

regulators as isolated anomalies rather than a systemic propellant problem — and, in places, manipulated the test data outright.<sup>2</sup> In 2017 Takata pleaded guilty to wire fraud and paid roughly a billion dollars in fine, restitution, and a victims' fund. The legal finding was pointed: not that a part had failed, but that the company had spent years misrepresenting what its own engineers knew.<sup>3</sup>

#### THE CAPABILITY GAP

Two capabilities were designed out, and the second matters as much as the first. The product capability was the desiccant — the stabilizing component Takata itself later added, omitted to save cost, in plain view of anyone comparing the two generations side by side. The system capability was independent verification: the recall regime treated the manufacturer's representations about its own safety data as authoritative, with no independent pipeline to test inflator behavior across the fleet as it aged in the field — which let an obvious failure mode hide for years inside a process that kept receiving the evidence and reporting it back as noise.<sup>4</sup>

#### AFTERMATH & REFORM

The recall is still being worked off, vehicle by vehicle, long after the company that built the inflators ceased to exist — concrete proof that a designed-out capability can outlive the firm that removed it and become someone else's burden. A decade on, the burden persists: the oldest, most dangerous "Alpha" inflators have been escalated to urgent "Do Not Drive" warnings, and in February 2026 FCA extended one to all 225,000 of its unrepaired vehicles. The unrepaired remainder in mid-2026 is put at about 4.8 million by CARFAX and at around 6 million by NHTSA's own completion reporting — the estimates differ, the order of magnitude does not — even as completion of the original "Alpha" inflators passed 99 percent. The episode pushed regulators toward more aggressive, coordinated recall management that does not leave pace to each manufacturer. Its central lesson is the pairing: Takata is the modern Pinto in its product failure, and more in its system failure — the evidence pipeline a regulator relies on is itself a safety capability, and omitting it is as consequential as omitting the desiccant from the inflator.<sup>5</sup>

# 67M

INFLATORS RECALLED IN THE U.S. · 19 AUTOMAKERS

*the desiccant Takata added only later was the designed-out capability*

TAKATA — THE LARGEST AUTOMOTIVE RECALL ON RECORD

*Takata engaged in a sustained pattern of misrepresenting inflator safety data to its automaker customers and to regulators.*

PARAPHRASING THE U.S. DEPARTMENT OF JUSTICE SETTLEMENT WITH TAKATA, 2017

REFERENCES

1. U.S. National Highway Traffic Safety Administration, Takata air-bag inflator recall coordination materials — the ammonium-nitrate-without-desiccant design and the propellant-degradation rupture mechanism.

2. NHTSA recall record and investigative reporting (Reuters, New York Times) — roughly 67 million inflators recalled in the United States across 19 automakers (over 100 million worldwide); the largest automotive recall in history; 28 U.S. deaths confirmed by NHTSA and at least 400 U.S. injuries alleged.

3. U.S. DOJ (2017) and Takata internal documents released in litigation — the sustained misrepresentation of inflator test data to automakers and regulators.

4. U.S. Department of Justice, settlement and guilty plea with Takata Corporation (2017) — wire fraud, \$1B in fine and restitution, and Takata's subsequent bankruptcy.

5. NHTSA consumer alert on FCA's "Do Not Drive" warning for all vehicles with unrepaired Takata recalls (February 2026, 225,000 vehicles); CARFAX estimate of 4.8 million unrepaired Takata recalls as of July 2026; Honda's report that over 99% of the original "Alpha" inflators have been replaced or accounted for — the years-long completion of replacements and the continuing risk in unrepaired vehicles.

6. NHTSA Coordinated Remedy Program for the Takata recalls — the regulator's move to actively prioritize and manage replacement across nineteen automakers rather than leave pace to each manufacturer.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Takata is the modern Pinto: an engineering choice to omit a stabilizing component, internal data showing the consequence, and a regulatory architecture that treated the manufacturer's representation as authoritative rather than as one input. The capability gap was at the regulator's evidence pipeline as much as at the manufacturer's bench.

LENS APPROACH

The lesson is post-deployment surveillance. Field rupture reports arrived over years but were never aggregated and interpreted as a single class-level safety signal, so an obvious failure mode stayed hidden inside a process that kept receiving the evidence and reporting it back as noise. The ammonium-nitrate-without-desiccant chemistry is the substrate of the hazard, but the capability that should have caught it is the surveillance-aggregation function that turns scattered field reports into one population-level signal with authority to act. LENS uses Takata in LEN 5 to build that aggregate-the-field-signal capability and in LEN 7 as an industrial-governance failure spanning manufacturer, regulator, and customer auditors.

| DOMAIN              | FAILURE MODES               | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|---------------------|-----------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| INDUSTRIAL          | → <div>D</div> <div>G</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 3            | 2.4     | 4   |
| Test and Evaluation |                             |                                                                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Qualify a safety-critical material against the full service life and climate range it will see, not just its as-delivered state, so slow degradation cannot turn a safe part lethal years later.
- ▶ When choosing a cheaper, less stable material, require the stabilizing measure the material needs — here a desiccant — rather than omitting it to save cost.
- ▶ Design inflator testing to detect the rupture failure mode as a class across temperature and humidity cycling, so it cannot be dismissed as an isolated anomaly.

IN OPERATION / AFTER

- ▶ Stand up an independent verification pipeline that tests fielded inflators across the aging fleet, so a regulator does not depend on the manufacturer's own representation of its safety data.
- ▶ Aggregate field rupture reports as a population signal rather than closing them one by one, with authority to trigger recall when a pattern appears.
- ▶ Manage the recall as a coordinated program that prioritizes replacement across automakers, since a designed-out hazard outlives the firm and keeps aging in unrepaired cars.

## REFLECTION QUESTIONS

1. Where in your domain does a regulator receive manufacturer test data without an independent verification pipeline? What is the decisive trust assumption?
2. Design the verification regime that should have surrounded ammonium-nitrate inflator testing. Who funds it, who runs it, and what does it produce?
3. Takata's propellant degraded slowly with heat and humidity, so a part safe at delivery became lethal years later. Identify a component in your domain whose qualification testing does not cover its full service life, and specify the aging test that would close the gap.

**WHO BUILDS THIS** systems & human-factors engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits

# Wells Fargo Fake Accounts

**G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** ~3.5 million potentially unauthorized accounts identified; ~\$3B in penalties; CEO resigned; the Federal Reserve capped the bank's assets

## IN BRIEF

To meet aggressive sales quotas, Wells Fargo employees opened accounts customers had not asked for; a third-party review put the total at roughly 3.5 million potentially unauthorized accounts from January 2009 to September 2016. The behavior was widespread and visible to internal risk and compliance functions, but the bank's response was to fire individual "bad apples" while leaving the incentive structure intact — and that structure was the actual cause: sales targets most employees could not meet by ethical means. The 2016 CFPB, OCC, and Los Angeles actions brought \$185 million in penalties — a 2020 DOJ/SEC settlement later reached roughly \$3 billion — the CEO resigned, and the Federal Reserve took the unprecedented step of capping the bank's assets. Wells Fargo is the canonical case of an incentive architecture that manufactured the misconduct the institution then prosecuted at the front line, while insisting the misconduct was individual.

## BACKGROUND

Wells Fargo built its retail strategy on "cross-selling" — opening many products per customer — and drove it with aggressive sales quotas pushed down to branch employees, whose pay and jobs depended on hitting them. The bank's signature metric, "Gr-eight" (an average of eight products per household), and a daily "solution" scorecard reported up the management chain made cross-sell the single most-watched proxy for branch performance. The metric the bank chose to manage by thus became the thing every front-line employee was structurally compelled to maximize, whatever it took. The controls function carried no comparably visible counter-measure — no widely reported figure for the share of those accounts that customers had actually authorized — so the incentive ran without a designed brake.<sup>o</sup>

## WHAT HAPPENED

Unable to meet the quotas honestly, employees opened roughly 3.5 million potentially unauthorized accounts in customers' names — the rational response to a target most could not reach by legitimate means. Practices documented by investigators included opening checking, savings, and credit-card accounts without customer consent, forging signatures, moving funds between accounts to manufacture activity, and enrolling customers in online banking they had not requested. The behavior was widespread and long-standing, and visible to internal risk and compliance functions for years before the 2013 Los Angeles Times reporting made it public; the institutional response was to discipline individual employees



as bad apples — Wells Fargo terminated roughly 5,300 employees between 2011 and 2016 — while leaving the incentive structure intact, fixing the symptom and preserving the cause.<sup>1</sup>

#### THE INVESTIGATION

The 2016 CFPB, OCC, and Los Angeles actions brought \$185 million in penalties (a 2020 DOJ/SEC settlement later reached roughly \$3 billion), and the bank's own independent-directors investigation (2017) documented how the sales-target architecture drove the misconduct — locating the cause in the design, not the people executing it.<sup>2</sup> Investigators tied the misconduct directly to the bank's incentive-compensation structure, which had made such sales practices a foreseeable result rather than an aberration. The CEO resigned, and the Federal Reserve capped the bank's assets, reaching past individual penalties to constrain the institution itself.<sup>3</sup>

#### THE CAPABILITY GAP

Wells Fargo is the strongest evidence in the dataset that incentive architecture is a capability-engineering deliverable. The measurement system used to manage employees produced exactly the behavior the institution then prosecuted — the same structure that demanded the result punished the people who delivered it. The gap was not at the front line but at the governance layer that designed the incentives and then, for years, treated the predictable result as individual moral failure, mistaking a designed outcome for a character flaw.<sup>4</sup>

#### AFTERMATH & REFORM

Stumpf's successor Tim Sloan, promoted from chief operating officer in the immediate aftermath, was himself forced out in 2019 after the Federal Reserve's asset cap proved more durable than the bank had assumed. The cap — imposed in February 2018 at roughly \$1.95 trillion — restricted Wells Fargo's growth pending evidence of governance and risk-management remediation, and it held for roughly seven years — the longest-running enforcement action of its kind against a major U.S. bank — until the Federal Reserve lifted it on 3 June 2025, concluding that the bank had completed the required governance and risk-management remediation. The case became the standard teaching example of measurement-gaming and incentive design.<sup>5</sup> Its lesson is that any quota becomes a target to be gamed, and that an institution is accountable for the behavior its measurement system makes rational, not just the behavior endorsed in its values statement — because employees respond to the incentive they are paid on, not the value they are told to honor.

# 3.5M

POTENTIALLY UNAUTHORIZED ACCOUNTS

*the incentive architecture made misconduct rational for the front line*

WELLS FARGO — THE MEASUREMENT SYSTEM PRODUCED THE BEHAVIOR THE INSTITUTION THEN PROSECUTED

*Wells Fargo’s sales practices were a foreseeable consequence of its incentive compensation structure.*

PARAPHRASING THE 2016 REGULATORY AND INDEPENDENT-DIRECTORS FINDINGS ON WELLS FARGO

REFERENCES

1. Consumer Financial Protection Bureau, Consent Order against Wells Fargo (2016) — the unauthorized accounts and penalties.

2. Office of the Comptroller of the Currency, Consent Order AA-EC-2016-66 (2016) — unsafe or unsound sales practices tied to the incentive-compensation structure (paraphrased).

3. Independent Directors of Wells Fargo, Sales Practices Investigation Report (2017) — how the sales-target architecture drove the conduct.

4. Enforcement record: 3.5 million accounts, \$3 billion in penalties, and the CEO’s resignation.

5. A. C. Edmondson, The Fearless Organization (2018); incentive-design and corporate-governance literature.

6. Federal Reserve asset cap on Wells Fargo (imposed February 2018; lifted 3 June 2025) — the structural growth constraint and its termination after remediation.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Wells Fargo is the strongest available evidence that incentive architecture is a capability-engineering deliverable. The measurement system used to manage employees produced the behavior the institution then prosecuted. The gap was not at the front line. It was at the governance layer that had designed the incentives.

LENS APPROACH

Wells Fargo is the canonical “protecting the measurement from gaming” case (induced 2.2; LENS D4/PT5), with measuring-the- failure-mode-you-care-about (induced 2.1) as the alternate anchor. LENS uses it in LEN 4 as the measurement-gaming case and in LEN 7 for the corporate-governance dynamics that allow such incentives to persist for years. Studio projects redesign the incentive architecture and the countervailing audit measure that would have detected the gaming as a structural pattern rather than as individual misconduct. Adjacent to Texas City (Case 161) at the wrong-measurement-reported-as-a-win layer.

| DOMAIN     | FAILURE MODES | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|------------|---------------|---------------------|--------------|---------|-------|
| TECHNOLOGY | → G N         | → D1 D2 D3 D4 D5    | 5            | 2.1     | 4 · 5 |
| GOVERNMENT |               | Test and Evaluation |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design incentive and quota structures against the behavior they will rationally produce, treating the measurement system as a capability deliverable in its own right.
- ▶ Set targets that are achievable by legitimate means, since a quota most employees cannot reach honestly is an instruction to cheat.
- ▶ Pair every performance metric with a countervailing measure that detects gaming before it becomes systemic.

IN OPERATION / AFTER

- ▶ Audit for the gap between the behavior the measurement system rewards and the conduct the institution claims to value, and treat divergence as a design fault.
- ▶ Empower risk and compliance to escalate a pattern of gaming as a structural cause, not to absorb it as isolated misconduct.
- ▶ Hold the governance layer that designed the incentives accountable for predictable outcomes, rather than disciplining only the front line.

## REFLECTION QUESTIONS

1. Identify a measurement system in your organization that is currently producing the behavior the organization claims to deplore. What is the gap between the measurement and the goal?
2. Design the incentive architecture for a bank's branch sales force that does not produce a Wells-Fargo-equivalent failure.
3. Wells Fargo's misconduct was visible to risk and compliance for years before it was addressed at the structural level. What lets an organization see a pattern of gaming and still treat it as individual bad apples — and who would have to be empowered to name the incentive as the cause?

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Volkswagen Dieselgate

**D** Designed Out    **G** Governance & Trust

**IMPACT** ~11 million vehicles equipped with defeat-device software; USD 4.3B U.S. DOJ criminal-and-civil plea (Jan 2017); total global cost above USD 30B across penalties, vehicle buybacks, and settlements; multiple criminal convictions

## IN BRIEF

Volkswagen engineered a “defeat device” into its diesel emissions software: code that detected when a car was on a regulatory test bench and switched on full emissions controls only then. On the road the controls were largely disabled, producing nitrogen-oxide emissions up to forty times the U.S. limit on the cars EPA tested, across roughly 11 million vehicles worldwide for years. A small West Virginia university team measuring real-world emissions found the anomaly; fifteen months of CARB and EPA pressure turned it into VW’s admission. Internal documents showed the defeat device was an institutional decision — a deliberate response to a standard VW’s engineers did not believe they could meet within cost — not a rogue act. VW pleaded guilty to a USD 4.3 billion criminal and civil settlement with the U.S. Department of Justice (January 2017); total Dieselgate cost across penalties, vehicle buybacks, and settlements has been estimated above USD 30 billion globally, with multiple criminal convictions. Dieselgate is the book’s case for organized evasion of a measurement system.

## BACKGROUND

Volkswagen had staked its U.S. growth on “clean diesel,” promising cars that met strict nitrogen-oxide limits without sacrificing performance or cost — a marketing position that left no room to admit the engineering could not deliver all three at once. Its engineers did not believe they could actually meet the standard within the platform’s cost constraints, putting the company’s public promise on a collision course with its own technical reality.<sup>o</sup>

## WHAT HAPPENED

So they cheated by design. VW’s emissions software detected when a vehicle was on a regulatory test bench — by its steering, speed, and duration patterns, the telltale signature of a lab rather than a road — and switched on full emissions controls only during the test. On the road the controls were largely disabled, producing emissions up to forty times the U.S. limit on the vehicles EPA tested, across roughly 11 million vehicles worldwide for years, so the pollution the standard existed to prevent flowed freely everywhere except where it was measured.<sup>1</sup>

## THE INVESTIGATION

A small university research team in West Virginia, testing for the ICCT, found road emissions far above lab results and reported the gap to CARB and EPA in May 2014; VW admitted the defeat device only in September 2015, after fifteen months of regulator pressure.<sup>2</sup> Internal


documents then showed the defeat device had been authorized inside VW's engineering hierarchy as a deliberate institutional response to a standard the team could not meet — not the work of a rogue engineer, but a decision made up the chain of command. VW pleaded guilty to a USD 4.3 billion U.S. DOJ criminal and civil settlement (January 2017) — total global Dieselgate cost across penalties, vehicle buybacks, and settlements has been estimated above USD 30 billion — and multiple executives were convicted.<sup>3</sup>

#### THE CAPABILITY GAP

The exploitable gap was at the regulator's instrument: the emissions test ran in a regime the vehicle could detect, so a manufacturer determined to evade could tune its behavior to the test rather than to the road, optimizing for the measurement instead of the goal it stood for. Dieselgate is the canonical case of institutionally engineered evasion of a measurement system — and a reminder that any test the measured party can recognize is a test it can defeat, because a detectable test invites the very gaming it is meant to catch.<sup>4</sup>

#### AFTERMATH & REFORM

The reform closed the gap at the instrument: the EU introduced real-world driving-emissions testing, moving the measurement off the predictable bench and onto the road the standard actually cared about, making it far harder to game.<sup>5</sup> The pattern parallels Takata (Case 147) — a manufacturer's fraud meeting a regulator whose evidence pipeline trusted the manufacturer's test conditions, and a fix that had to upgrade the measurement itself, not just punish the cheat, because punishing the cheat leaves the exploitable instrument in place for the next one. Criminal accountability in Germany ran on for years: on 26 May 2025 the Braunschweig regional court convicted four former VW managers of fraud — four and a half years in prison for the former diesel engine development chief, two years and seven months for a former engine-electronics manager, suspended sentences for the other two — while the separate trial of former CEO Martin Winterkorn, opened in September 2024, was provisionally terminated in July 2025: the court cited a health impediment it described as temporary and reviewable, so the charges stand undecided rather than dismissed.<sup>6</sup>



*The defeat device was the product of a long-standing institutional decision to evade emissions standards.*

PARAPHRASING THE U.S. DEPARTMENT OF JUSTICE PLEA AGREEMENT WITH VOLKSWAGEN AG, 2017

#### REFERENCES

1. Volkswagen internal documents released through litigation — authorization of the defeat device within the engineering hierarchy.

2. U.S. EPA, Notice of Violation to Volkswagen (2015) — the defeat device and emissions exceedances.

3. West Virginia University / ICCT real-world diesel-emissions study (2014) — the discovery comparing road to lab.

4. U.S. Department of Justice Plea Agreement with Volkswagen AG (January 2017) — the institutional decision, USD 4.3 billion criminal-and-civil settlement, and convictions (quoted).

5. J. Ewing, *Faster, Higher, Farther: The Volkswagen Scandal* (2017); cf. Takata (Case 147).

6. EU Real Driving Emissions (RDE) testing regulation — the post-Dieselgate measurement reform.

7. Associated Press reporting on the German criminal proceedings — the Braunschweig regional court’s conviction of four former VW managers on 26 May 2025 and its provisional termination of the Winterkorn trial in July 2025 on health grounds.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Dieselgate is the canonical case for institutionally engineered evasion of a measurement system. The capability gap was at the regulator’s test protocol: it operated in a regime the vehicle could detect. The reform — real-world emissions testing — was a capability deliverable upgrade at the regulator’s instrument boundary.

LENS APPROACH

LENS reads Dieselgate at the independent-instrument and operator-inspector boundary (induced 2.2): the defeat device worked by detecting the test condition itself, so the measurement was one the operator could see coming and game. The lesson is designing an emissions-measurement and inspection regime the operator cannot detect and defeat — real-driving emissions, surprise and outside-the-loop testing — rather than the bare observation that fraud occurred. The reform pattern parallels Takata (Case 147).

| DOMAIN     | FAILURE MODES       |                           | LENS COMPETENCY |               |               |               |               | PROBLEM TYPE  | INDUCED | LEO |   |
|------------|---------------------|---------------------------|-----------------|---------------|---------------|---------------|---------------|---------------|---------|-----|---|
| INDUSTRIAL | →                   | <div>D</div> <div>G</div> | →               | <div>D1</div> | <div>D2</div> | <div>D3</div> | <div>D4</div> | <div>D5</div> | 3       | 2.2 | 4 |
| GOVERNMENT | Test and Evaluation |                           |                 |               |               |               |               |               |         |     |   |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div><div>▶ Design measurement protocols the measured party cannot detect and tune to, so behavior is optimized for the goal rather than the test.</div><div>▶ Surface the gap between a public commitment and the engineering reality early, so an unmeetable standard is renegotiated rather than secretly evaded.</div><div>▶ Build a governance check between an impossible target and the team asked to meet it, so the response is escalation, not a defeat device.</div></div> | <div><div>▶ Audit real-world behavior against test-condition behavior, treating a divergence between road and bench as the signature of engineered evasion.</div><div>▶ Upgrade the measurement instrument itself after a discovered cheat, not just penalize the cheat, since the exploitable test invites the next one.</div><div>▶ Sustain independent, outside-the-loop verification — as the university team provided — to catch fraud the regulator’s own protocol cannot see.</div></div> |

REFLECTION QUESTIONS

1. Identify a measurement protocol in your domain that operates in a regime detectable by the entity being measured. What is the evasion potential?

2. Design the upgrade to a regulatory test protocol that makes Dieselgate-style evasion structurally infeasible.

3. VW’s defeat device was authorized up the engineering hierarchy, not the act of a rogue engineer. What allows an institution to sanction fraud as a deliberate response to an unmeetable standard — and what governance check should sit between an impossible target and the team asked to meet it?

**WHO BUILDS THIS** systems & human-factors engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits

# GM Ignition Switch

**D** Designed Out   **G** Governance & Trust

**IMPACT** 124 deaths attributed; ~2.6M vehicles recalled for the defective switch; \$900M federal forfeiture (DOJ, 2015); the fix existed for eight years before the recall

## IN BRIEF

A faulty ignition switch in several GM compact cars (the Chevrolet Cobalt, Saturn Ion) could slip from “run” to “accessory” while driving, cutting power steering and brake assist and — fatally — disarming the airbags. GM approved the switch for production in 2002 below its own torque spec. In 2006 an engineer approved a redesigned switch but did not change its part number, so the fix propagated as neither a revision nor a recall, and defective cars kept selling. The recall came only in 2014 — about 2.6 million vehicles, with 124 deaths attributed through GM’s compensation fund. The Valukas report found a culture that absorbed safety concerns, and GM paid a \$900 million federal penalty. What was designed out was not a part but the pathway by which a known safety problem reaches the decision to recall.

## BACKGROUND

The ignition switches in several GM compact cars — the Chevrolet Cobalt and Saturn Ion among them — had too little detent torque to hold their position, so a jostle, a bump in the road, or a heavy keychain could rotate the switch out of “run” while driving. That cut power steering and braking and, fatally, disarmed the airbags so they would not deploy in the very crash that loss of control often produced. GM approved the switch for production in 2002 below its own torque spec, and when an engineer stalled a pre-production Cobalt in 2004 the company classified the fault as a convenience issue, not a safety one.<sup>0</sup>

## WHAT HAPPENED

In 2006 a GM engineer approved a redesigned switch but did not change its part number — and in an engineering organization the part number is the very mechanism by which the system knows something changed. The fix propagated as neither a revision nor a recall; defective cars kept selling, and crashes with non-deploying airbags were investigated piecemeal over years, no one connecting them to a switch the records said had “never changed.” The recall came only in 2014 — about 2.6 million vehicles, with 124 deaths attributed through the compensation fund. What finally broke the silence was a wrongful-death lawsuit, in which a family’s expert found the part had been quietly changed under the same number — the smoking gun GM’s own records had been built to miss.<sup>1</sup>



## THE INVESTIGATION


GM's commissioned investigation by Anton Valukas (2014) found not a single villain but a culture that absorbed safety concerns until they dissipated — the “GM nod,” in which a room agrees a thing should happen and then no one acts, and the “GM salute,” arms crossed and each person pointing elsewhere — and a fundamental failure to use the escalation processes the company already had on the books.<sup>2</sup> In 2015 GM paid a \$900 million federal forfeiture for concealing the defect; the compensation fund alone offered \$594.5 million on 399 approved claims.<sup>3</sup>

## THE CAPABILITY GAP

What was designed out of GM was not a part but a pathway — the institutional route by which a known safety problem reaches the decision to recall. The fix existed for eight years; the path from fix to recall did not, so the knowledge sat inert. The mechanism is mundane and all the more instructive for it: the part-number convention was the company's own way of seeing what had changed, and quietly breaking it blinded the organization to its own action. Suppressing a signal need not be a conspiracy; here it was one engineer taking the path of least resistance through a process no one was assigned to guard.<sup>4</sup>

## AFTERMATH & REFORM

Mary Barra, who became CEO as the recall broke, used the Valukas report to restructure GM's safety decision-making — a global-safety leadership role, consolidated escalation channels, and a “Speak Up for Safety” program to give concerns a route upward that did not depend on one person's persistence to survive. The case became a standard teaching example in governance and psychological safety: an organization can hold the fix for the better part of a decade and still fail to act if the channel that carries bad news upward has been allowed to fail quietly, with no one accountable for keeping it open.<sup>5</sup>



GM IGNITION SWITCH — EIGHT YEARS BETWEEN THE FIX AND THE RECALL

*There was a fundamental failure to use the formal escalation processes that GM had established.*

PARAPHRASING THE VALUKAS REPORT TO THE GM BOARD OF DIRECTORS, 2014

## REFERENCES

1. A. R. Valukas, Report to the Board of Directors of General Motors Company Regarding Ignition Switch Recalls (Jenner & Block, 2014) — the 2002 identification of the defect and the switch-torque mechanism.
4. U.S. Department of Justice, deferred-prosecution agreement with General Motors (2015) — \$900 million forfeiture for concealing the defect; total penalties/settlements exceeding \$2.6 billion.

2. U.S. NHTSA investigation of the GM ignition switch (2014) and the GM recall record — 2.6 million vehicles; 124 deaths via the GM compensation fund.

3. Valukas Report (2014) — the unchanged part number, the “GM nod” and “GM salute,” and the failure to use established escalation processes (quoted).

5. Valukas Report (2014) and A. C. Edmondson, *The Fearless Organization* (Wiley, 2018) — the part number as an organizational signal, and the suppression of upward safety information.

6. U.S. NHTSA consent order with General Motors (2014, \$35M civil penalty) and the GM ignition-switch victims’ compensation program (K. Feinberg, administrator) — the regulatory penalty and the basis for the 124-death figure.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The GM ignition switch case is the canonical example of a corporate organizational structure that suppressed the upward flow of safety information by procedural design. The fix existed for eight years before the recall. The institutional path between the fix and the recall did not.

LENS APPROACH

The lesson is change control. Engineers changed the ignition switch but kept the same part number, defeating configuration traceability so the field failures could never be tied back to a design change and the records insisted the part had “never changed.” The capability to build is configuration management that forbids altering a part without renumbering it, so the system can always see what changed and connect a field failure to the revision that caused it. LENS uses GM in LEN 4 for that change-control and part-number-as-signal discipline, and secondarily in LEN 7 and LEN 8 for the corporate-cover-up and Valukas-style retrospective accountability framing.

| DOMAIN                                | FAILURE MODES               | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|-----------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| INDUSTRIAL                            | → <div>D</div> <div>G</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 5            | 5.4     | 4   |
| Navigating Sociotechnical Constraints |                             |                                                                         |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div><div>▶ Treat the part number as a safety-critical signal: make changing a part without changing its number impossible by process, so the system cannot be blinded to its own revision.</div><div>▶ Specify the ignition switch’s detent torque as a hold-position requirement and verify it, so a known under-torque defect cannot pass into production as a tolerable quirk.</div><div>▶ Build the route from “fix approved” to “recall decision” as an explicit deliverable, not a path that depends on one person choosing to escalate.</div></div> | <div><div>▶ Correlate field crashes with non-deploying airbags as a population, so piecemeal incidents are connected to a common cause rather than investigated and closed one at a time.</div><div>▶ Run an escalation channel that carries bad news upward independent of any single person’s persistence, with a named owner accountable for keeping it open.</div><div>▶ Audit whether approved fixes have actually propagated to a recall or revision, closing the gap in which a known defect keeps shipping under an unchanged record.</div></div> |

REFLECTION QUESTIONS

1. What information channel in your organization carries the same load that GM’s part-number system did? Could it be silently bypassed?

2. Design the escalation deliverable that would have moved the GM ignition switch fix to a recall in 2006.

3. A single engineer changed the switch without changing its part number, and no one was assigned to guard that convention. Identify a process in your organization whose integrity depends on an unenforced convention, and name who should own enforcing it.

**WHO BUILDS THIS** systems & human-factors engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits

# TSB Bank IT Migration

**H** Human-System Interface   **G** Governance & Trust

**IMPACT** 1.9 million UK customers locked out of accounts; £330M+ in compensation and remediation; CEO resigned

## IN BRIEF

In April 2018 TSB Bank tried to migrate some five million customer accounts off its former parent Lloyds' systems onto a new platform from its current owner, Sabadell, over a single weekend. When services came back online, nearly every component failed: 1.9 million customers were locked out, 440 saw a linked relative's accounts, mortgages vanished, payments bounced. Recovery took months and cost over £330 million; the CEO resigned and the regulator fined the bank. The independent review found the platform had been tested under conditions that did not approximate real load, certified ready by a process that did not challenge the certification, and pushed live on readiness assurances that were statements of expectation rather than of proven fact. TSB is the financial-sector analog of Healthcare.gov: a migration shipped without adequate testing because the assurance process confirmed readiness instead of testing it.

## BACKGROUND

TSB Bank, spun out of Lloyds and acquired by Spain's Sabadell, needed to move some five million customer accounts off Lloyds' systems onto a new Sabadell-built platform. The cutover was scheduled for a single weekend.<sup>0</sup> Compressing a five-million-account migration into one weekend left no room for partial failure: the schedule itself became a forcing function, framing readiness as a date to be hit rather than a condition to be proven, and that framing would later prove decisive when the technical signal said the platform was not ready.

## WHAT HAPPENED

When customer-facing services came back online that Sunday evening in April 2018, nearly every component of the new platform had problems. About 1.9 million customers were locked out; 440 saw the accounts of a closely linked relative, mortgages disappeared, and card payments failed. The recovery took months, cost more than £330 million in compensation and remediation, and the chief executive resigned.<sup>1</sup> That nearly every component failed at once points away from a single defect and toward a platform that had never been exercised under real load — the kind of systemic breakdown that follows when a system is proven only in conditions it will never actually meet.

## THE INVESTIGATION

The Slaughter and May independent review found the migration had been tested under conditions that did not approximate real customer load, and that the platform had been certified ready by a process that did not adequately challenge the certification — a certi-


fication that confirmed readiness rather than interrogating it, which is how a system that would fail under real conditions could be signed off as fit.<sup>2</sup> Decisively, the assurances the board acted on were forward-looking statements of intention rather than statements of completed fact, and no independent adviser was asked to test them; the Financial Conduct Authority and the Prudential Regulation Authority later fined TSB £48.65 million, treating the unchallenged assurance as a failure of governance and not merely of engineering.<sup>3</sup>

#### THE CAPABILITY GAP

The signals existed — testing had slipped, three of the program’s own guiding principles had been abandoned, and attestations were still incomplete on the morning of the decision. But they reached the board as assurances rather than as objections, and the institutional architecture gave no one the standing to convert a slipped test into a halt. The missing capability was not testing knowledge but a governance structure in which a “not ready” could stop a scheduled go-live. Knowing a system is unready is worthless if the knowledge cannot reach the decision with its force intact and the authority to act on it; here the truth was present but powerless.<sup>4</sup>

#### AFTERMATH & REFORM

TSB rebuilt its testing and migration governance, paid out and was penalized, and the case entered the literature as a study in technical-decision authority.<sup>5</sup> Rebuilding governance rather than merely the platform was the right diagnosis: the failure had been one of who could say “stop” and be heeded, so the durable fix had to live in the decision structure rather than the code. It is the financial-sector analog of Healthcare.gov (Case 180): a large migration shipped without the testing the institution knew it needed, because schedule pressure overrode a technical signal that had no authority to win.



*Inadequacies in the safeguards meant TSB went live with the platform before it was ready to do so.*

PARAPHRASING THE FCA FINAL NOTICE TO TSB BANK PLC, 20 DECEMBER 2022, PARAGRAPH 4.154

#### REFERENCES

1. Slaughter and May, Independent Review of the TSB Migration (2019) — the single-weekend cutover and the testing failures.
4. Financial Conduct Authority, Final Notice on TSB Bank plc (20 December 2022) — the £29.75m penalty (with the PRA's

2. Slaughter and May (2019) and FCA materials — 1.9 million customers locked out, £330M+ in costs, and the CEO's resignation.
3. Slaughter and May (2019) — inadequate load testing and an unchallenged readiness certification.
4. Slaughter and May (2019) — inadequate load testing and an unchallenged readiness certification.
5. House of Commons Treasury Committee, report on the TSB IT migration (2018).
6. Cf. Healthcare.gov (Case 180) and the migration-safety literature.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

TSB is the canonical case for schedule pressure overriding technical signal in a regulated industry. The technical shortfalls existed. They reached the decision layer converted into attestations of readiness, and nothing in the governance made the conversion visible. The institutional architecture did not allow the technical layer to halt the migration.

### LENS APPROACH

LENS uses TSB in LEN 7 as a corporate-governance case and in LEN 8 for the institutional structure of technical-decision authority. Studio projects compare TSB and Healthcare.gov.

| DOMAIN                                     | FAILURE MODES                                                                     | LENS COMPETENCY | PROBLEM TYPE | INDUCED | LEO |
|--------------------------------------------|-----------------------------------------------------------------------------------|-----------------|--------------|---------|-----|
| TECHNOLOGY → <span>H</span> <span>G</span> | → <span>D1</span> <span>D2</span> <span>D3</span> <span>D4</span> <span>D5</span> | 4               | 4.1          | 5       |     |
| Navigating Sociotechnical Constraints      |                                                                                   |                 |              |         |     |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- ▶ Test the platform under conditions that approximate real production load, since a system proven only in unrepresentative conditions will fail when it meets the real ones.
- ▶ Make certification adversarial — a process that tries to disprove readiness — rather than a sign-off that confirms it.
- ▶ Avoid single-weekend big-bang cutovers where feasible; stage the migration so partial failure is survivable rather than catastrophic.

#### IN OPERATION / AFTER

- ▶ Build a governance structure in which a technical “not ready” can halt a scheduled go-live, so the signal reaches the decision with its force and authority intact.
- ▶ Ensure the decision authority hears the technical signal directly rather than through intermediate layers that attenuate it.
- ▶ After any failure, fix the decision structure that allowed a known objection to be overridden, not just the system that broke.

### REFLECTION QUESTIONS

1. Where in your organization does a technical signal arrive at the executive layer attenuated by intermediate layers? What is the cost of the attenuation?
2. Design the institutional structure that would allow a technical lead to halt a migration like TSB's without resigning.
3. TSB's certification confirmed readiness rather than interrogating it. Examine a sign-off process in your domain that rubber-stamps rather than challenges — what would it take to make the certification adversarial enough to catch an unready system?

**WHO BUILDS THIS** human factors & interaction design · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · governance frameworks · bias & evidence audits

# Equifax Data Breach

**G** Governance & Trust   **K** Knowledge & Institutional Memory

**IMPACT** 147 million Americans' personal data exposed; CEO resigned; settlement of at least \$575M, up to \$700M; foundational U.S. data-breach case

## IN BRIEF

In 2017 attackers exploited a known, two-month-old vulnerability in Apache Struts on an Equifax web portal — a patch had been available, and Equifax's own security team had told IT to apply it; no one did. Over the next 2.5 months the attackers exfiltrated the personal data of 147 million Americans, and Equifax did not disclose the breach until September. A Senate investigation found systematically inadequate patching, an incomplete asset inventory (so the company did not know which systems needed the fix), and a 2015 internal audit whose warnings on all of it went unremediated for two years. The CEO resigned and Equifax settled for at least \$575 million. No single failure caused the breach; cumulative inadequacy across routine cybersecurity work did.

## BACKGROUND

Equifax, one of the three U.S. credit bureaus, held the most sensitive financial data on virtually every American adult — a concentration of identity information that made any breach catastrophic by definition. One of its public web portals ran a version of Apache Struts with a known critical vulnerability for which a patch had been available for two months — and which Equifax's own security team had flagged for IT to apply, so the warning and the fix both existed inside the company and simply went unacted-upon.<sup>o</sup>

## WHAT HAPPENED

The patch was not applied. Attackers exploited the vulnerability beginning in May 2017 and quietly exfiltrated personally identifying information on 147 million Americans over the next two and a half months, a window long enough that the theft was less a break-in than a sustained occupation. Equifax did not publicly disclose the breach until September, so the people whose identities were taken learned of it only well after the fact.<sup>1</sup>

## THE INVESTIGATION


The U.S. Senate Permanent Subcommittee on Investigations found Equifax's patching practices systematically inadequate and the company lacking "a comprehensive IT asset inventory" — so it could not reliably know which systems needed the patch, leaving the security team's warning with no map to act on.<sup>2</sup> Monitoring was weak — an SSL certificate on the dispute portal had been expired since November 2016, so the traffic went uninspected — and a 2015 audit had named every one of these gaps; the CEO resigned and Equifax settled for at least \$575 million, rising to \$700 million if claims required it.<sup>3</sup>

## THE CAPABILITY GAP

Equifax is the institutional-cybersecurity case for cumulative inadequacy in routine work: patching, asset inventory, monitoring, response. Each function was below standard; none alone produced the breach; the combination did, the marginal weaknesses compounding into a single open door. The capability gap was the management of unglamorous, universally-agreed-necessary maintenance — exactly the work easy to defer because deferring it usually costs nothing, until the one time it costs everything, and a company that has already been told where the gaps are still has to close them.<sup>4</sup>

### AFTERMATH & REFORM

The settlement funded consumer compensation and credit monitoring, and the breach pushed patching discipline, asset inventory, and breach-disclosure timelines up the corporate agenda — elevating, after the loss, the unglamorous work that had been deferred before it.<sup>5</sup> It was a cumulative-inadequacy failure: no dramatic single cause, just several routine defenses each left marginally inadequate, failing together on the day a determined attacker arrived to test all of them at once.



*Equifax lacked a comprehensive IT asset inventory.*

U.S. SENATE PERMANENT SUBCOMMITTEE ON INVESTIGATIONS, HOW EQUIFAX NEGLECTED CYBERSECURITY, MARCH 2019

## REFERENCES

1. U.S. Senate Permanent Subcommittee on Investigations, How Equifax Neglected Cybersecurity and Suffered a Devastating Data Breach (2019) — the unpatched Apache Struts vulnerability.
2. The breach record — 147 million affected, exploitation from May 2017, public disclosure in September 2017.
3. Senate PSI (2019) — “Equifax lacked a comprehensive IT asset inventory” (quoted).
4. U.S. FTC / CFPB / state settlement (at least \$575 million, up to \$700 million, 2019) and the CEO’s resignation.
5. U.S. GAO, Actions Taken by Equifax and Federal Agencies in Response to the 2017 Breach, GAO-18-559 (2018).
6. Apache Struts CVE-2017-5638 advisory.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT



Equifax is the canonical institutional-cybersecurity case for cumulative inadequacy in routine work: patching, inventory, monitoring, response. Each function was below industry standard. None alone produced the breach. The combination was the breach.

LENS APPROACH

LENS uses Equifax in LEN 6 as a feedback-channel failure: the security team identified the Apache Struts patch and told IT to apply it — the operator-to-institution escalation existed, was correct, and died before reaching anyone with the asset map and the authority to act. The teaching point is not “patch your systems” but the engineered escalation deliverable: a flagged, agreed-upon fix has to be wired to an owner who can map it to every affected asset and is funded to close it before the gap costs everything. LEN 7 carries the institutional half — an audit that names the gap does nothing unless someone is made to close it.

| DOMAIN     | FAILURE MODES | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO |
|------------|---------------|---------------------------------------|--------------|---------|-----|
| TECHNOLOGY | → G K         | → D1 D2 D3 D4 D5                      | 3            | 6.2     | 2   |
| GOVERNMENT |               | Navigating Sociotechnical Constraints |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Close audit findings on routine security maintenance — patching, monitoring, response — on a clock with a named owner, so a known gap cannot sit open for two years.
- ▶ Build and maintain a comprehensive asset inventory so a known vulnerability can be mapped to every system that runs it.
- ▶ Wire the security team’s warnings to an owner with the authority and resources to apply a flagged patch promptly.

IN OPERATION / AFTER

- ▶ Audit the routine defenses together — patching, inventory, monitoring, response — since each marginally inadequate layer raises the odds the combination fails.
- ▶ Monitor for sustained, low-noise exfiltration, treating a months-long quiet window as the failure mode to detect, not just a single break-in.
- ▶ Hold breach-disclosure timelines short, so affected people learn of a theft of their identity without months of delay.

REFLECTION QUESTIONS

1. Identify a piece of routine work in your domain that is chronically deferred. What is the cumulative-inadequacy threshold?
2. Equifax did not know which assets ran Struts. Design the asset-inventory deliverable that an organization the size of Equifax should be able to produce on demand.
3. Equifax’s security team had flagged the patch, but the warning had no asset map to act on and no funded function to carry it out. What turns a known, agreed-upon fix into deferred work in your organization — and what would force it to be done before the one time it costs everything?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · knowledge bases · data pipelines



Chapter 10

# Industry, Energy & Enterprise Systems — What Works — and the Frontier

*When operations make capability visible on the floor.*

---

*The mechanism was easy to copy. What made it work sat on a different budget line.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

# INL Turbine–Control Upgrade — Verification and Validation as the Cutover Deliverable

**G** Governance & Trust   **K** Knowledge & Institutional Memory   **H** Human–System Interface

**IMPACT** An INL paper set out the NUREG–0711 human–factors process for a paired plant–process–computer and turbine–control–system upgrade — iterative usability testing, integrated system validation in a full–scope training simulator, a stand–alone operator training program built with the training organization, and validation scenarios reviewed by trainers and procedure writers to confirm the operating procedures need no modification

## IN BRIEF

Nuclear–plant control–room modernization is one of the canonical C7 problems — legacy analog systems that have to be replaced with digital equivalents inside a regulatory regime that demands the safety case survive the transition. The Idaho National Laboratory’s Light Water Reactor Sustainability (LWRS) program has produced a body of technical reports and conference papers documenting the verification–and–validation work behind specific upgrades. The case here is a specific method: for a paired plant–process–computer and turbine–control–system upgrade, the operators are trained on the new interfaces through a stand–alone program built with the training organization, and the validation scenarios are reviewed by trainers and procedure writers to confirm the operating procedures need no modification. The retraining is planned into the cutover, not avoided by it. That specification is the small–tier complement to the big–tier LWRS program case the corpus also tracks: the program produces specific design–V&V case studies, not only program–level claims. The decisive hedge is not the sourcing — a reviewed conference paper in the Human Factors and Ergonomics Society’s annual–meeting proceedings, alongside the program’s own technical reporting — but what it reports. It reports a process, prospectively and in the future tense, and no outcome: the process was, at publication, “currently being implemented at a U.S. nuclear utility with four plants slated for control room modernization,” with details and lessons learned to be published in later papers.

## BACKGROUND

The nuclear–plant control room is the C7 transition problem with the safety–case constraint at the front: the regulator and the licensee both need to know that the upgraded digital control behaves equivalently to the legacy analog under the operating envelopes the safety case covers, and that the operator does not lose capability in the cutover. The LWRS program at INL exists to produce that evidence across the U.S. fleet’s modernization needs. The specific study here is one of the program’s small–tier deliverables: a human–factors verification of a turbine–control–system upgrade.<sup>o</sup>

## THE INTERVENTION

The relevant specification the case rests on is that operators are trained on the new digital interfaces through a stand-alone program developed with the training organization, and that the validation scenarios are reviewed by trainers and procedure writers to ensure the operating procedures do not require modification. That is the substantive human-factors-V&V claim: retraining is a named deliverable of the cutover, and the procedure set is what the validation is asked to hold constant. The paper is a per-subsystem process account rather than a whole-control-room finding, and the LWRS program treats these small-tier evaluations as the building blocks of the larger fleet-modernization safety case.<sup>1</sup>

## HOW IT WORKED

The case is the small-tier complement to the big-tier LWRS program case the corpus also tracks. The big-tier case states the program-level claim: that nuclear plants can be modernized into digital control rooms with the safety case maintained. The small-tier case is the per-subsystem evidence that the program-level claim has to rest on — the individual design-V&V deliverables that, in aggregate, give the regulator and the licensee a reason to accept the modernization. Without the small-tier evidence the program-level claim is unfounded; without the program-level coordination the small-tier evidence does not roll up.<sup>2</sup>

## THE EVIDENCE

What has to stay visible is what the paper does not contain. The sourcing is sound — a reviewed HFES conference paper, alongside the program's own technical reporting — but it reports a planned process rather than an outcome, conditional on the specific subsystems, plant, and operator population; whether the finding generalizes to other subsystems in the same plant, to other plants in the fleet, or to other digital control products is the open question that the LWRS program itself acknowledges and continues to study, and it stays open until the program publishes the follow-on results.<sup>3</sup>

## WHAT TRANSFERRED

What the case adds at the LENS layer is the small-tier example of verification-as-deliverable inside a regulated transition. The Domain-1 frame applies: the human-factors analysis of the operator's task under the digital cutover is a systems-analysis artifact that the regulatory regime consumes as part of its acceptance decision. The Domain-4 frame applies because the V&V evidence is itself the measurement instrument the program rests on. The case is a paired small-tier companion to Case 116 (Eurocat ATM) and sits inside the C7 thread alongside the big-tier failures.

*The program-level claim cannot rest on un-rolled-up small-tier evidence. The per-subsystem V&V is the building block of the modernization safety case.*

EDITORS' SYNTHESIS OF INL LWRS REPORTING.

## REFERENCES

1. Ronald L. Boring (2014), "Human Factors Design, Verification, and Validation for Two Types of Control Room Upgrades at a Nuclear Power Plant," Proceedings of the Human Factors and Ergonomics Society Annual Meeting 58(1), pp. 2295–2299, doi:10.1177/1541931214581478 — the NUREG-0711 process for the paired plant-process-computer and turbine-control-system upgrade. Step 4 of the paper's implementation sequence is the
3. Nuclear Regulatory Commission (NUREG-0711), "Human Factors Engineering Program Review Model" — the regulatory framework the V&V deliverables are produced against.
4. O'Hara et al. (2008), "Human Factors Considerations with Respect to Emerging Technology in Nuclear Power Plants," NUREG/CR-6947 — peer-adjacent framing.

- source of the training claim verbatim: “A stand-alone training program will be developed in cooperation with the training organization. In addition, the scenarios will be reviewed by trainers and procedure writers to ensure that the operating procedures do not require modifications.” The future tense is the paper’s own.
2. Idaho National Laboratory, Light Water Reactor Sustainability Program reports on control-room modernization — series available via OSTI.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The INL turbine-control paper is a small-tier C7 verification-as-deliverable case inside the LWRS nuclear-modernization program: the human-factors V&V process carries the cutover, training included. The paper reports the process and not its outcome, and the generalization beyond the studied subsystem is the open question the program itself keeps open.

LENS APPROACH

INL turbine-control upgrade is the C7 small-tier transition case (induced 7.1; LENS D1/PT1). LENS uses it in Domain 1 (Systems Analysis) for LEO-1 — the operator-task analysis under the digital cutover is the systems-analysis artifact the regulator consumes — and in Domain 4 (Test and Evaluation) for LEO-4, since the V&V evidence is itself the measurement instrument the program rests on. Pair with Case 116 (Eurocat) for the small-tier C7 thread and with the LWRS big-tier program case already in the corpus.

| DOMAIN                     | MODES ADDRESSED | LENS COMPETENCY    | PROBLEM TYPE | INDUCED | LEO   |
|----------------------------|-----------------|--------------------|--------------|---------|-------|
| NUCLEAR POWER              | G K H           | D1 D2 D3 D4 D5     | 1            | 7.1     | 1 · 4 |
| CONTROL-ROOM MODERNIZATION | →               | → Systems Analysis |              |         |       |
| HUMAN FACTORS              |                 |                    |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat each subsystem upgrade as a per-subsystem human-factors V&V deliverable, not as a sub-task of the larger control-room modernization claim.
- ▶ Specify the operator-task analysis the cutover has to satisfy before the upgrade ships, so the V&V evidence has a target the regulator can audit.
- ▶ Document the conditioning factors — plant, operator population, subsystem boundary — explicitly, so the generalization question is honestly framed in the report.

IN OPERATION / AFTER

- ▶ Roll the per-subsystem V&V studies into the program-level fleet-modernization safety case deliberately; the program-level claim cannot rest on un-rolled-up small-tier evidence.
- ▶ Carry the process-not-outcome limit into any downstream citation; what the paper specifies is a method, conditional on the studied subsystem and population.
- ▶ Track durability of the low-retraining finding under personnel rotation, procedure updates, and subsequent upgrades; the cutover is not finished when the digital system is in service.

### REFLECTION QUESTIONS

1. Identify a control-system upgrade in your context whose cutover burden on the operator is the decisive constraint. What is the per-subsystem V&V deliverable that would settle the burden question, and what is its scope?
2. Specify the rollup logic from per-subsystem V&V evidence to a program-level modernization claim. The INL pattern teaches because the small-tier studies are the building blocks of a fleet-level claim — what would the equivalent be in your context?
3. The paper specifies a process and reports no outcome from it. What evidence (measured retraining burden, multi-plant replication, durability tracking under personnel rotation) would you require before treating the low-retraining expectation as a finding at all?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · usability testing · cognitive walkthroughs

# Toyota Production System / Andon Cord

**N** Normalization of Deviance    **G** Governance & Trust

**IMPACT** Front-line signalling authority surfaces defects at the source within the work cycle; most activations never stop the line; the system exported globally with mixed fidelity

## IN BRIEF

The andon cord lets any assembly-line worker signal a defect. Under the fixed-position stop system the line keeps moving: the pull lights the station's andon and calls the team leader, and only a problem still unresolved when the vehicle reaches the end of the work position stops the line. Most activations stop nothing. The cord is trivially cheap; the guaranteed response behind it is the design. What American automakers imported in the 1980s and 1990s was the cord. What they did not import was the response behind it: team leaders staffed on the floor at a ratio that leaves them free to answer inside a cycle, buffers thin enough that a defect cannot be parked downstream, and standardized work to judge the anomaly against. Toyota's own accounts add a cultural layer — no-blame root-cause analysis and the codified "Five Whys" — and the two are entangled in every plant where the system works. No study separates them. The Andon Cord is the manufacturing counterpart of the Keystone nurse-authority intervention (Case 19), and carries the same evidentiary limit.

## BACKGROUND

In high-volume manufacturing, a defect that passes undetected propagates downstream, multiplying the cost of every later correction. Catching problems at the source requires the person who sees them — usually the lowest-ranking worker on the line — to be able to act, in an environment where stopping a line running at millions of dollars per hour is otherwise unthinkable. The economics and the authority structure point in opposite directions: the cheapest moment to fix a defect is the one at which the person who sees it has the least standing to halt production.<sup>9</sup>

## THE INTERVENTION

Toyota's andon is not a switch that halts a factory. Under the fixed-position stop system, pulling the cord or pressing the button lights the station's andon and calls the team leader; the line keeps moving. If the leader resolves the problem before the vehicle reaches the end of the work position — one takt cycle, on the order of a minute — the signal resets and nothing stops. Only an unresolved problem stops the line, and it stops at the fixed position, so no task is interrupted mid-cycle. The design is deliberate: it removes the two things that suppress reporting — the dread of halting production over a small problem, and the confusion of restarting a half-finished job. What the front-line worker holds is the authority to **summon**, backed by a guaranteed response. The stop is the escalation, not the act.<sup>1</sup>



## HOW IT WORKED


What sits behind the cord is a set of conditions a plant has to be funded to hold. Team leaders are staffed on the floor at a ratio that leaves them free to answer within a takt cycle; a fully utilized responder turns the cord into a doorbell. Buffers between stations are kept thin, so a defect cannot be parked downstream and the signal has to be acted on. Standardized work gives the anomaly something to be measured against. Toyota's accounts add a cultural layer on top of these — no-blame root-cause analysis and the codified "Five Whys" — and practitioners generally credit it. In every plant where the system works the structural conditions and the cultural ones are present together, and the record does not separate them.<sup>2</sup>

## THE EVIDENCE

The evidence is thinner than the story. Toyota's Georgetown plant is widely reported to log andon activations in the thousands per day against a handful per week at comparable U.S. plants — figures that trace to a 2007 BBC report and a since-removed plant-tour page, not to audited data. NUMMI is the usual proof: GM's Fremont plant closed in 1982; the joint venture reopened it in 1984 with approximately 85 percent of the same laid-off UAW workforce and came within reach of Takaoka's quality. But NUMMI changed everything at once — Toyota management, team structure, worker-authored standardized work, a no-layoff commitment, and training trips to Takaoka. No study isolates the cord's marginal contribution, and much of the mechanism narrative traces to Toyota's own account of itself. The claim the record supports is narrow: a reporting channel with a guaranteed, non-punitive response gets used, and one without such a response does not.<sup>3</sup>

## WHAT TRANSFERRED

What travels is not the cord. It is a short list of things a plant can be funded to hold or refused: a responder ratio with deliberate slack, buffers thin enough that a signal cannot be absorbed, standardized work to compare against, and employment terms under which surfacing a problem is not self-harm. Each is a line item somebody approves or cuts, which is what makes the transfer question answerable before a cord is hung. Keystone (Case 19) is the medical case with the same shape and the same limit — a reporting channel paired with a guaranteed response, and no trial able to say which half carried the result.<sup>4</sup>



*most calls resolved within the work cycle*

SIGNAL, GUARANTEED RESPONSE, AND A STOP ONLY AT THE FIXED POSITION

## COMPETING READINGS

- ◆ The cord is the causal agent — set against the reading that the cord is an interface to a staffed response chain, with team leaders on the floor and resourced to answer within the cycle. On this reading the expensive half is the staffing, and imitation fails on a budget line rather than on values.
- ◆ Protected authority is the variable — set against the reading that fixed-position stop is an engineering choice that lowers the cost of speaking up, so the design does the work the empowerment is credited with.
- ◆ American imitation failed on culture — set against the reading that it failed on structure, because the copied plants lacked the team-leader ratio and the thin buffers that make a guaranteed response possible. Each reading predicts a different thing to check before a transfer.
- ◆ A high pull count signals health — set against the reading that pull counts are an input measure Toyota controls and publishes, with no external audit distinguishing real defect signalling from routine minor calls.

What this case does not show — The andon is a local, operational control at the work station; it did not and could not surface the design- and organization-level failures of the 2009–10 recall crisis, which the NASA/NHTSA review traced to pedal and floor-mat defects and which Toyota's own president attributed to growth outrunning capability development.

*The cord costs almost nothing. Everything that makes it work — a responder with time to answer, a buffer thin enough that the signal cannot be absorbed — is the part that costs.*

EDITORS' SYNTHESIS

## REFERENCES

1. Liker, J. (2020), *The Toyota Way* (2nd ed.) — the andon cord, the cultural pairing, and the American imitation (paraphrased).
2. Shingo, S. (1989), *A Study of the Toyota Production System — the technical mechanism*.
3. Spear, S. & Bowen, H. (1999), "Decoding the DNA of the Toyota Production System," *HBR* — the response protocol and embedded learning.
4. Adler, P. (1993), "Time-and-Motion Regained," *HBR* 71(1) — NUMMI observed from the floor; the rehired workforce, the team structure, and the standardized-work practice that changed alongside the signalling system.
5. Womack & Jones (1996), *Lean Thinking* — diffusion and the limits of surface imitation.
6. Rother, M. (2009), *Toyota Kata* — the routines that sustain the practice.
7. NASA Engineering and Safety Center / NHTSA (2011), *Technical Assessment of Toyota Electronic Throttle Control Systems — the unintended-acceleration review underlying the scope limit*.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The andon is a ladder, not a switch: a signal engineered to cost almost nothing to send, a responder staffed to answer inside the cycle, and a bounded consequence if the problem outlives it. Each rung is funded or cut separately, and a channel whose response is unfunded is a doorbell. That the signal and the response have to be designed together is the conclusion the record supports. Which of them carried Toyota's outcome is not something any study here settles.

### LENS APPROACH

LENS uses the Andon Cord in LEN 2 as the worked example of a signalling channel designed together with its response, in LEN 8 for cross-domain transfer — where the student's task is to price the responder staffing, the buffer policy, and the employment terms before judging whether the mechanism can move — and in LEN 10 as a studio brief in which the artifact is the easy half: the design work is the response and its cost curve, with the falsification test specified before anything is built.

| DOMAIN                          | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------|-----------------|------------------|--------------|---------|-----|
| INDUSTRIAL                      | → N G           | → D1 D2 D3 D4 D5 | 3            | 4.1     | 3   |
| Human-System Collaboration + D4 |                 |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Put the signal where the problem is first visible and engineer its cost down to near zero — under fixed-position stop the first rung summons help rather than halting production, which is what makes it usable by the person with the least standing.
- ▶ Budget the response before installing the signal. The responder ratio with deliberate slack is the expensive half and the one that decides whether the channel works.
- ▶ Keep buffers thin enough that a defect cannot be parked downstream, so the signal has to be acted on rather than absorbed.

IN OPERATION / AFTER

- ▶ Sustain the pairing with no-blame root-cause analysis and a codified method (the Five Whys) so each stop becomes embedded learning rather than a one-off interruption.
- ▶ Treat the activation rate as a health indicator in which low is bad, and watch it under schedule pressure — a rate that falls when the quarter tightens is the channel closing.
- ▶ Before transferring the model, price the responder staffing, the buffer policy, and the employment terms. Those are the line items a copy drops silently.

REFLECTION QUESTIONS

1. Identify a signalling channel in your domain — a hotline, a flag, a stop-work right. What does using it cost the person who uses it, in seconds and in social exposure? Cut that cost by an order of magnitude and name what you had to change to do it.
2. Name who is structurally free to answer that signal inside the window in which an answer still matters, and what their utilization rate is. If nobody is, the channel is a doorbell; specify the staffing that would make it a channel.
3. Toyota publishes its own pull counts and no external audit distinguishes real defect signals from routine minor calls. Design the falsification test for your channel: what observation would show it is being used ritually rather than substantively, and what would you do on seeing it?

WHO BUILDS THIS safety science & organizational analysis · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. TOOLS incident analysis · safety dashboards · governance frameworks · bias & evidence audits

## CASE 156

NUCLEAR ENGINEERING

CONTROL-ROOM HUMAN FACTORS

RESEARCH AND DEVELOPMENT

2010 – PRESENT

# INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet

**D** Designed Out   **H** Human-System Interface   **K** Knowledge & Institutional Memory

**IMPACT** The U.S. Department of Energy's Light Water Reactor Sustainability (LWRS) program, executed at Idaho National Laboratory in partnership with utilities, has produced a structured research-and-pilot record for modernizing analog control-room instrumentation in the existing U.S. nuclear fleet — pilot-scale evidence covering hybrid digital/analog operator interfaces, human-factors validation, and qualification pathways for digital instrumentation

## IN BRIEF

The U.S. commercial nuclear fleet is dominated by reactors originally licensed in the 1970s and 1980s with analog instrumentation and control-room layouts of that vintage. Aging analog components, vendor obsolescence, and the operational case for digital instrumentation make modernization a foreseeable sustainment requirement; the regulatory environment (Nuclear Regulatory Commission Regulatory Guide 1.180, Branch Technical Position 7-19, and related guidance) makes the qualification path for digital instrumentation deliberately stringent. The U.S. Department of Energy's Light Water Reactor Sustainability (LWRS) program, executed at Idaho National Laboratory in partnership with U.S. utilities, runs the federally-funded research-and-pilot work that supports the fleet's modernization decisions. The program has produced a structured record covering hybrid digital/ analog operator interfaces, human-factors validation studies in full-scope simulators, and qualification-pathway research for safety-related digital instrumentation. The honest hedge survives clearly: LWRS observations are pilot-scale, the sample of plants that have implemented major modernizations is small, and the operational-fleet evidence at scale is forward-looking. The case is teachable on the structured sustainment-research form — a federally-funded research program operating across decades to support an industry's aging-fleet decisions — and closes a C7 zero-state in the induced framework, paired with Cases 114, 115, 174 as the v2 aging-system quartet.

## BACKGROUND

The U.S. commercial nuclear fleet operates roughly 90 reactors, most of them originally licensed in the 1970s and 1980s — the handful of later units (Watts Bar 2, the Vogtle AP1000s) are the exception that the fleet's sustainment problem is not about. The control rooms of that era were built around analog instrumentation — strip-chart recorders, hardwired indicators, control-panel meters — designed for the operational envelope and the human-

factors assumptions of the period. Aging analog components, vendor obsolescence, and the operational case for digital instrumentation (better diagnostics, lower maintenance burden, integration with plant computer systems) make modernization a foreseeable sustainment requirement across the fleet.<sup>0</sup>

#### THE INTERVENTION

The regulatory environment is deliberately stringent. Nuclear Regulatory Commission guidance — Regulatory Guide 1.180 on electromagnetic and radio-frequency compatibility, Branch Technical Position 7-19 on common-cause failure analysis for digital instrumentation and control, and the broader Standard Review Plan Chapter 7 — sets the qualification path for safety-related digital instrumentation. The stringency is intentional: digital systems can introduce failure modes (common-cause software faults, cyber exposure) that analog systems do not, and the regulatory regime is designed to keep those failure modes bounded as plants modernize.<sup>1</sup>

#### HOW IT WORKED

The U.S. Department of Energy's Light Water Reactor Sustainability (LWRS) program is the federally-funded research-and-pilot work that supports the fleet's modernization decisions. Executed at Idaho National Laboratory in partnership with U.S. utilities and EPRI, LWRS operates across a multi-decade horizon with annual research-report deliverables. The control-room modernization research line — including the Human Systems Simulation Laboratory at INL with full-scope plant simulators — covers hybrid digital/analog operator-interface designs, human-factors validation studies, and qualification-pathway research that supports utility-level licensing submissions.<sup>2</sup>

#### THE EVIDENCE

The honest hedge has to be visible. LWRS observations are pilot-scale: full-scope simulator studies, single-plant pilot implementations, and structured human-factors experiments with operator participants from utility partners. The sample of U.S. plants that have implemented major control-room modernizations is small relative to the fleet, and the operational-fleet evidence at scale — fleet-wide reliability, fleet-wide human-error rate, fleet-wide maintenance burden under hybrid digital/analog control rooms — is forward-looking rather than retrospective. The program's own reports characterize the research at this evidence tier, and the case carries the same honesty.<sup>3</sup>

#### WHAT TRANSFERRED

What the case teaches at the LENS framing layer is the structured sustainment-research form — a federally-funded research program operating across decades in partnership with industry and a regulator, producing the research-and-pilot record that supports licensing decisions on aging-fleet modernization. The form is the sustainment-engineering analog of the FAA Aging Aircraft Program (Case 114): a long-horizon institutional discipline that produces the technical record that aging-fleet decisions can rest on. With NextGen/ADS-B (Case 115) and Y2K (Case 174), LWRS completes the v2 aging-system quartet that closes the C7 zero-state in the induced framework. The LWRS instance is where the evidence is most explicitly pilot-scale, and the case carries that as the program's tier acknowledgement rather than as a weakness to smooth.<sup>4</sup>

*The institutional form is what transfers: federal research program plus industry partnership plus regulator engagement, operating across decades to produce a technical record aging-fleet decisions can rest on.*

EDITORS' SYNTHESIS OF THE LWRS ANNUAL REPORT SERIES.

REFERENCES

1. O'Hara, Higgins, Brown, Fink, Persensky, Lewis, Kramer, & Szabo (2008), "Human Factors Considerations with Respect to Emerging Technology in Nuclear Power Plants," NUREG/CR-6947 — foundational human-factors backdrop.

2. Nuclear Regulatory Commission, Branch Technical Position 7-19, "Guidance for Evaluation of Diversity and Defense-in-Depth in Digital Computer-Based Instrumentation and Control Systems."

3. Nuclear Regulatory Commission, Regulatory Guide 1.180, "Guidelines for Evaluating Electromagnetic and Radio-Frequency Interference in Safety-Related Instrumentation and Control Systems."

4. Idaho National Laboratory, Light Water Reactor Sustainability (LWRS) program annual reports (2010 – present) — primary research-and-pilot record.

5. Electric Power Research Institute, control-room modernization technical reports — industry-side sustainment record.

THE LEARNING ENGINEERING LENS

LE INSIGHT

LWRS is the structured sustainment-research case — federally- funded research at INL in partnership with utilities and the regulator, producing the research-and-pilot record that supports aging-fleet modernization decisions across a multi- decade horizon. The observations are pilot-scale; the operational-fleet evidence at scale is forward-looking. The hedge is part of the case.

LENS APPROACH

LWRS control-room modernization is the structured sustainment-research case (induced 7.1; LENS D3/PT4) — Domain 1 for the long-horizon partnership; Domain 5 for the federal-research + industry + regulator triple structure; Domain 3 for the hybrid digital/analog research line. Closes C7 with Cases 114, 115, 174.

| DOMAIN                     | MODES ADDRESSED | LENS COMPETENCY              | PROBLEM TYPE | INDUCED | LEO   |
|----------------------------|-----------------|------------------------------|--------------|---------|-------|
| NUCLEAR ENGINEERING        | D H K           | D1 D2 D3 D4 D5               | 4            | 7.1     | 1 · 3 |
| CONTROL-ROOM HUMAN FACTORS | →               | → Human-System Collaboration |              |         |       |
| RESEARCH AND DEVELOPMENT   |                 |                              |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pair the research program with full-scope simulator infrastructure (the INL Human Systems Simulation Laboratory) so the human-factors validation work runs in the operational envelope rather than in abstracted laboratory conditions.
- ▶ Structure the partnership across federal research (DOE/INL), industry (utilities, EPRI), and regulator (NRC) so the research record is recognized by all three parties — the licensing submissions ultimately rest on that recognition.

IN OPERATION / AFTER

- ▶ Track fleet-scale outcomes as utility implementations proceed: fleet-wide reliability, human-error rate, maintenance burden under modernized control rooms. The forward-looking evidence will accumulate over the next two decades; the program is the institutional carrier of that accumulation.
- ▶ Carry the LWRS case in pair with FAA Aging Aircraft (Case 114) as the sustainment-engineering analog at multi-decade horizon, and with Y2K (Case 174) and NextGen (Case 115) as the broader aging-system quartet.

- ▶ Carry the pilot-scale evidence tier honestly in the program's reports; LWRS characterizes its research at that tier, and downstream practitioners reading the work should read it the same way.
- ▶ Treat the structured-sustainment-research institutional form as the transferable artifact: federal research program plus industry partnership plus regulator engagement, operating across decades to produce the technical record aging-fleet decisions can rest on.

### REFLECTION QUESTIONS

1. Identify an aging fleet of long-lived assets in your domain whose modernization decisions rest on research that does not yet exist at fleet scale. What is the analog of a federally-funded research program — institutional partner, industry partner, regulator partner — that could operate across the horizon the decisions actually need?
2. Specify the full-scope simulator (or equivalent in-envelope test infrastructure) the program would need so its human-factors validation work runs in operational conditions rather than abstracted laboratory ones. What is the cost of building it, and what does its absence mean for the evidence the research produces?
3. LWRS evidence is pilot-scale and the program reports it as such. Identify a program in your domain whose evidence tier should be acknowledged in its reporting but is not. What institutional discipline keeps the evidence-tier honesty visible — to the program, to industry partners, and to the regulator?

**WHO BUILDS THIS** systems & human-factors engineering · human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · usability testing · cognitive walkthroughs · knowledge bases · data pipelines

# AI-Augmented Coding Tools

**T** Training Gap   **H** Human-System Interface

**IMPACT** Twenty million GitHub Copilot users by mid-2025, plus Cursor and peers; productivity gains documented; security and correctness implications still being characterized

## IN BRIEF

AI-augmented coding tools — GitHub Copilot, Cursor, Codeium, and peers — represent the largest real-time experiment in human-AI collaboration in this book, with twenty million Copilot users alone by mid-2025. Published studies (Peng et al. 2023) document real short-term productivity gains; other work (Pearce et al. 2022) finds a substantial share of AI-generated completions in security-sensitive settings contain vulnerabilities, though a controlled study (Sandoval et al. 2023) found AI assistance did not significantly raise the rate of critical security bugs. The capability question is open: are developers becoming more capable, or more dependent? The short-term gains are real; the long-term consequences — especially for those who learn the craft with these tools always available — are not yet known. The discipline is being asked to define good before the longitudinal evidence is in.

## THE SHIFT

AI coding assistants moved from novelty to infrastructure in a few years. GitHub Copilot, Cursor, Codeium, and similar tools now suggest, complete, and generate code for twenty million Copilot users alone — the largest real-time experiment in human-AI collaboration in professional knowledge work to date, conducted not in a study design but in the live practice of an entire profession, with no control group and no agreed measure of what it is doing to the underlying craft.<sup>0</sup>

## WHAT IS EMERGING

Two findings are accumulating in parallel. Controlled studies (Peng et al. 2023) document real short-term productivity gains. At the same time, the security picture is unsettled: Pearce et al. (2022) found a substantial fraction of Copilot completions in security-relevant scenarios contained vulnerabilities, while a controlled study by Sandoval et al. (2023) found AI assistance did not significantly increase the rate of critical security bugs. The two results do not cancel so much as mark how unsettled the picture is — output clearly rises, but the quality and safety of that output resist a single verdict.<sup>1</sup>

## THE CAPABILITY QUESTION

The capability question is open and consequential: are developers using these tools becoming more capable, or more dependent? Short-term output rises, but whether the underlying skill grows or erodes — especially for those who learn the craft with the tools always present




— is precisely what the productivity metrics cannot tell us, because a measure of how much code ships says nothing about whether the person shipping it could still produce or judge it without the assistant.<sup>2</sup>

#### EARLY EVIDENCE

The longitudinal evidence is not yet sufficient to answer. Dell’Acqua et al. (2023) found a “jagged frontier” in professional LLM use: performance improves on tasks inside the tool’s competence and degrades on tasks just outside it, where users over-trust the output. The short-term gains are real; the long-term capability consequences remain uncharacterized — and the jagged frontier is hard to navigate precisely because its edge is invisible from inside the task, so the user cannot tell when they have crossed from where the tool helps to where it misleads.<sup>3</sup>

#### OPEN PROBLEMS

AI-augmented coding is the live frontier for human–AI teaming in professional knowledge work, and the discipline LENS represents is being asked to specify what good looks like before the long-term evidence is in. The open problem is the longitudinal study that could distinguish capability growth from capability erosion — and a training practice that keeps the human’s skill on the growing side, built and adopted while a generation is already learning the craft with the tools always within reach. A July 2025 randomized controlled trial by METR sharpened the warning: experienced open-source developers allowed to use early-2025 AI tools took about 19% longer on real tasks, yet believed the tools had sped them up by roughly 20% — direct evidence that self-reported productivity can invert the true effect, exactly the measurement gap this case argues the discipline must instrument.<sup>4</sup>



*AI assistance changes what developers can do; it may also change what they need to know.*

EDITORS' SYNTHESIS

#### REFERENCES

- Peng et al. (2023), “The Impact of AI on Developer Productivity” — short-term productivity gains.
- Pearce et al. (2022), “Asleep at the Keyboard? Assessing the Security of GitHub Copilot’s Code Contributions.”
- L. Bainbridge, “Ironies of Automation,” *Automatica* 19(6) (1983) — the classic deskilling and over-reliance problem, applied to AI-augmented work.
- Dell’Acqua et al. (2023), “Navigating the Jagged Technological Frontier” (HBS / BCG) — professional LLM use.
- METR (2025), “Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity,” arXiv:2507.09089 — the randomized trial finding a 19% slowdown against a self-reported speedup.
- Sandoval et al. (2023), “Lost at C,” *USENIX Security* — AI assistance did not significantly increase critical security-bug rates.

## THE LEARNING ENGINEERING LENS

## LE INSIGHT

AI-augmented coding is the live test case for human-AI teaming in professional knowledge work. The short-term gains are real. The long-term capability question — does the tool make the operator more capable, or more dependent — is the question the discipline must learn to ask and answer.

## LENS APPROACH

The teaching point is a measurement-design problem, and it is the one that decides the rest: productivity metrics count output and so cannot distinguish a developer whose skill is growing from one whose skill is quietly eroding under augmentation. The learning engineer's task is to build the instrument that separates those two. The design is a longitudinal, tool-removed probe — periodically measure each developer on representative tasks with the assistant withheld, scoring unaided correctness, debugging, and the ability to judge generated code, and track that aided-minus-unaided gap over time. A widening gap (rising aided output, flat-or-falling unaided competence) reads as skill-atrophy; a narrowing one reads as skill-growth. LENS uses this case in LEN 2 (human-AI teaming), in LEN 8 to build the capability-development instrument itself, and in LEN 10 (capstone) to have the student design the atrophy-versus-growth measurement for an augmented practice in their own domain.

| DOMAIN                     | MODES AT STAKE | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|----------------|------------------|--------------|---------|-----|
| TECHNOLOGY                 | → T H          | → D1 D2 D3 D4 D5 | 6            | 5.2     | 3   |
| Human-System Collaboration |                |                  |              |         |     |

## APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- ▶ Define what competence the human must retain independently of the tool, and design the workflow so that skill is exercised rather than quietly handed off.
- ▶ Engineer the assistant to surface the jagged-frontier edge — flagging where a task sits outside its reliable competence — so users do not over-trust output just beyond it.
- ▶ Keep verification of generated code, especially in security-relevant settings, a required step rather than an optional one, given the unsettled quality picture.

## IN OPERATION / AFTER

- ▶ Run the longitudinal study that productivity metrics cannot substitute for, measuring whether underlying skill is growing or eroding over years of use.
- ▶ Monitor for over-reliance at the competence boundary, where the evidence shows performance degrading as users trust the tool past its reliable range.
- ▶ Track outcomes for practitioners who learned the craft with the tools always present, the cohort whose long-term capability is most uncertain.

## REFLECTION QUESTIONS

1. In your domain, identify a class of practitioners whose work is currently being augmented by AI tools. What evidence would tell you whether their capability is growing or eroding?
2. Design the longitudinal study that would distinguish capability growth from capability erosion in an AI-augmented professional practice.
3. The “jagged frontier” is hard to navigate because its edge is invisible from inside the task. Design a signal or practice that would tell a practitioner in your domain when they have crossed from where the tool helps to where it misleads.

**WHO BUILDS THIS** learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs

Chapter II

# Disaster Prevention & Recovery — What Fails

*Before the event, prevention regimes that thinned out; after it, response plans that met reality.*

---

*Every barrier had been inspected. What failed was the system deciding what counted as a barrier.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

#### EDITOR'S NOTE · BEFORE AND AFTER

This part reads the corpus along a different axis than the domain parts that precede it: the lifecycle of a disaster. Every case here has two clocks. Before the event, there is a prevention regime — a certification, an inspection cadence, a defense-in-depth argument — whose quiet thinning is the real story. After the event, there is a response and a recovery, and the institutions that are (or are not) built so the same failure cannot recur. The chapter is therefore ordered in two arcs, marked below: prevention failures first, then response-and-recovery failures; and in the What Works chapter, prevention engineered as a deliverable, then recovery run as a discipline. Cases whose primary lessons belong to a domain — an aviation accident, a clinical harm — remain in their domain parts; what qualifies a case for this part is that its load-bearing lesson is the pre/post lifecycle itself.

#### BEFORE THE EVENT — PREVENTION

---

# Bhopal

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance   **G** Governance & Trust

**IMPACT** ≈ 15,000–20,000 killed; ≈ 500,000 injured; worst industrial disaster in history

## IN BRIEF

On the night of 2–3 December 1984, about forty tons of methyl isocyanate gas escaped from a Union Carbide pesticide plant in Bhopal, India — the worst industrial disaster in history. Thousands died within hours; estimates of total deaths run to 15,000–20,000, with roughly half a million exposed or injured. Safety systems had been off-line for months, the plant was understaffed, and workers were inadequately trained to recognize or handle the emergency. A seven-week New York Times inquiry traced the catastrophe to operating errors, design flaws, maintenance failures, training deficiencies, and cost-cutting that endangered safety. Bhopal catalyzed the creation of the U.S. Chemical Safety Board and reshaped industrial-safety regulation for decades. It is the largest-magnitude capability-and-governance failure on record.

## BACKGROUND

Union Carbide's pesticide plant in Bhopal, India, stored methyl isocyanate (MIC) — an extraordinarily toxic intermediate — in bulk, holding a lethal hazard in tanks beside a populated city. By 1984 the plant was running under heavy cost pressure: understaffed, with several key safety systems out of service for months, and workers inadequately trained to handle an MIC emergency or read its warning signs. Each economy was individually defensible on a ledger; together they thinned every layer of defense the process depended on.<sup>0</sup>

## WHAT HAPPENED

On the night of 2–3 December 1984, water entered an MIC storage tank and triggered a runaway reaction; the safety systems that should have contained it were non-operational, so the one event the plant existed to prevent met no working barrier on its way out. About forty tons of gas vented over the sleeping city. Thousands died within hours; estimates of total deaths run to 15,000–20,000, and roughly half a million people were exposed or injured — the worst industrial disaster in history, its toll set by who happened to be downwind.<sup>1</sup>

## THE INVESTIGATION

A seven-week New York Times inquiry found the catastrophe “resulted from operating errors, design flaws, maintenance failures, training deficiencies and economy measures that endangered safety” — a list with no single villain, which is precisely what made it hard to govern.<sup>2</sup> Human-factors analysis placed Bhopal alongside Three Mile Island in its neglect of the human element, and the U.S. Chemical Safety Board would later find ineffective employee training an underlying cause in nine of its first twenty-three chemical-incident

investigations — a pattern that traces to Bhopal and shows the same gap recurring long after the lesson was available.<sup>3</sup>

#### THE CAPABILITY GAP

Bhopal is the largest-magnitude capability-and-governance failure on record, and a multi-layer one: training, maintenance, design, staffing, and oversight had all degraded together, and no layer above the plant was accountable for the whole. Because the erosion was spread across layers, no single inspection or metric saw it, and each degraded layer made the next one matter more. The capability to operate an extraordinarily hazardous process safely had been hollowed out by cost-cutting, and the governance that should have caught the hollowing did not exist to ask whether the whole was still safe.<sup>4</sup>

#### AFTERMATH & REFORM

The disaster reshaped industrial safety worldwide and, in the United States, catalyzed the creation of the Chemical Safety Board — an INPO-equivalent for industrial chemistry, authorized by the Clean Air Act Amendments of 1990 and not operating until January 1998, fourteen years after the release — and the process-safety regime that followed, an oversight layer built to own exactly the whole that no one had owned at Bhopal.<sup>5</sup> The book's recurring arc runs through Bhopal in its starkest form: a catastrophe forces into being the institution the industry should have built before it, at a price no institution-building exercise should ever have to cost.



▲ ~ 15,000-20,000 dead · ~ 500,000 injured

SAFETY SYSTEMS OFF-LINE FOR MONTHS BEFORE THE RELEASE

*Operating errors, design flaws, maintenance failures, training deficiencies and economy measures that endangered safety.*

NEW YORK TIMES INVESTIGATION, 1985

#### REFERENCES

1. Union Carbide and government investigation reports (1985) — MIC storage, the disabled safety systems, and plant understaffing.
2. Accounts of the 2–3 Dec. 1984 release — the contested toll (thousands of immediate deaths; 15,000–20,000 total estimates; 500,000 exposed). (Figures vary widely across sources; see AUDIT.)
3. New York Times investigation (1985) — “operating errors, design flaws, maintenance failures, training deficiencies and economy measures that endangered safety” (quoted).
4. N. Meshkati, “Human Factors in Large-Scale Technological Systems’ Accidents,” *Industrial Crisis Quarterly* (1991) — Bhopal alongside Three Mile Island on neglect of the human element. G. Joseph, M. Kaszniak & L. Long, “Lessons after Bhopal: CSB a catalyst for change,” *Journal of Loss Prevention in the Process Industries* 18(4–6):537–548 (2005) — ineffective employee training an underlying cause in nine of the CSB’s first twenty-three completed investigations.
5. P. Shrivastava, *Bhopal: Anatomy of a Crisis* (1992); C. Perrow, *Normal Accidents* (1984).
6. The U.S. Chemical Safety and Hazard Investigation Board — authorized by the Clean Air Act Amendments of 1990, operational

from January 1998 — and the post-Bhopal process-safety regime (EPCRA 1986; OSHA Process Safety Management, 1992).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Bhopal is the largest-magnitude capability-and-governance failure on record. It is also the catalyst for the creation of the U.S. Chemical Safety Board, an INPO-equivalent for industrial chemistry. The pattern that follows nearly every case in this book — diagnose, then engineer remediation at scale — runs through Bhopal in canonical form.

LENS APPROACH

LENS names the specific capability gap, not “governance matters” in the abstract: at a transferred and cost-pressured plant, the maintenance and operating capability had been hollowed out — refrigeration and scrubber safety systems off-line, maintenance deferred, staff and training cut, deviance quietly normalized. Each eroded layer was individually tolerated, so the failure is a multi-layer drift (induced 7.4) that no single inspection or metric could see. The teaching point is the multi-layer-drift analysis and the treatment of maintenance capability as a deliverable that must be sustained, not a ledger line available for trimming.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|-----------------|--------------|---------|-----|
| INDUSTRIAL                            | → T K N G →   | D1 D2 D3 D4 D5  | 4            | 7.4     | 1·5 |
| Navigating Sociotechnical Constraints |               |                 |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat each safety system, staffing level, and training requirement as a non-negotiable layer of defense, not a cost line available for trimming under pressure.
- ▶ Assign a single accountable owner for the integrated hazard, so degradation spread across layers cannot fall between everyone’s responsibilities.
- ▶ Require that operators be trained to recognize and respond to the specific emergency the process can produce before the process is run with that hazard in bulk.

IN OPERATION / AFTER

- ▶ Audit the whole defense-in-depth posture together rather than layer by layer, since each degraded barrier silently raises the stakes on the next.
- ▶ Track how long any safety system stays off-line and gate continued operation on its restoration, treating extended downtime as an unacceptable condition.
- ▶ Sustain an oversight layer — an INPO- or CSB-equivalent — with authority and reach above the individual plant to own the integrated risk.

REFLECTION QUESTIONS

1. Bhopal produced the CSB. What is the institution your domain has not yet built that a comparable disaster would force into existence?
2. Identify a plant or facility in your domain that has had safety systems off-line for an extended period. What is the measurement gap that allowed it?
3. At Bhopal training, maintenance, design, staffing, and oversight all degraded together while each looked acceptable alone. Whose job in your organization is to judge whether the whole is still safe — and what would they have to see to declare that it is not?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards · governance frameworks · bias & evidence audits

# Davis–Besse Reactor Head Corrosion

**N** Normalization of Deviance   **K** Knowledge & Institutional Memory   **G** Governance & Trust

**IMPACT** Pineapple-sized cavity — roughly 5 inches by 4 to 5 inches — eaten clean through the carbon steel of the reactor pressure-vessel head; near-miss; plant shut down for repair and fleet-wide head-inspection requirements tightened

## IN BRIEF

During a 2002 refueling outage at the Davis–Besse nuclear plant in Ohio, workers found a pineapple-sized cavity eaten through six inches of the carbon-steel reactor-vessel head by leaking boric acid, leaving only a thin layer of stainless cladding between the cavity and reactor coolant at 2,200 psi — a serious near-miss for a loss-of-coolant accident. The leakage had been visible for years and inspections deferred; FirstEnergy had won an NRC deferral of a mandatory inspection to fit its refueling schedule, and the NRC’s Inspector General later found the decision had weighted economics over safety. Davis–Besse is the canonical post-Three-Mile-Island near-miss — evidence that even an industry that built INPO to engineer safety can let regulator-operator dynamics erode it.

## BACKGROUND

After Three Mile Island, U.S. commercial nuclear power had built — through INPO and a strengthened NRC — a reputation for engineering safety as a system. Davis–Besse, a pressurized-water reactor in Ohio, operated inside that regime, with a known industry-wide hazard: boric acid leaking from cracked nozzles could corrode the carbon-steel reactor-vessel head.<sup>0</sup> The hazard was generic and understood across the fleet, which is what makes the case sharp: the danger was not a surprise mechanism but a recognized one the post-TMI regime was precisely supposed to keep inspected and contained.

## WHAT HAPPENED

During a refueling outage in early 2002, workers found a cavity about 5 inches long and 4 to 5 inches wide eaten clean through six inches of the carbon-steel head, leaving only a thin layer of stainless-steel cladding between the corrosion and reactor coolant at 2,200 psi. Had the cladding given way, the result would have been a medium-break loss-of-coolant accident with a badly degraded safety margin.<sup>1</sup> The thin remaining cladding was the entire distance between routine operation and the kind of accident the entire regulatory regime existed to prevent — a margin measured in a fraction of an inch of unintended material rather than in any engineered barrier.



## THE INVESTIGATION

The boric-acid leakage had been observable for years, and inspections had been deferred. FirstEnergy had requested — and the NRC had granted — a deferral of the inspection Bulletin 2001-01 called for by 31 December 2001. The utility asked to inspect in April, after a 31 March 2002 refuelling shutdown; NRC drafted an order to shut the plant in December, then shelved it and settled on mid-February.<sup>2</sup> The NRC's Office of Inspector General later found the decision had been driven in large part by a desire to lessen the financial impact on the utility — a finding the NRC disputed, and which GAO, reviewing the agency's records in 2004, could neither support nor refute because the key meetings went undocumented — and that the agency had set itself an unreasonably high bar — absolute proof of a safety problem — before it would order a plant shut down.<sup>3</sup> The deferral was not a covert lapse but a documented decision both parties signed off on, which is the unsettling part: the erosion happened through the regulator's normal process, in the open, agreed to be safe to wait.

## THE CAPABILITY GAP

Davis-Besse is a near-miss in the dimension above operations: the regulator-operator relationship. The same industry that had shown, through INPO, that it knew how to engineer safety let a known corrosion mechanism go uninspected because inspecting it was inconvenient and expensive. Regulatory capture — the regulator adopting the operator's economic frame — is a capability failure at the institutional layer above the plant.<sup>4</sup> The plant's own engineering competence was never the missing piece; what failed was the independence of the layer meant to overrule a utility's schedule when safety required it, and that layer had quietly adopted the schedule as its own.

## AFTERMATH & REFORM

INPO and the NRC restructured significant parts of their inspection regimes, and reactor-head inspection requirements were tightened across the fleet. The lesson pairs with the book's TMI / INPO arc: institutional capability is not built once. It erodes if it is not re-engineered — and the erosion is quietest where the regulator and the regulated agree it is safe to wait.<sup>5</sup> Tightening the head-inspection requirement across the fleet conceded that a mandatory inspection had been a real barrier all along — one the deferral process had been allowed to treat as negotiable against a refueling schedule.

3/8"

OF STAINLESS CLADDING REMAINED

*between a 6-inch corrosion cavity and reactor coolant at 2,200 psi*

DAVIS-BESSE — THE POST-TMI NEAR-MISS

*NRC's decision was driven in large part by a desire to lessen the financial impact on FirstEnergy that would result from an early shutdown.*

NRC OFFICE OF THE INSPECTOR GENERAL FINDING, QUOTED IN GAO-04-415 (2004)

REFERENCES

1.

U.S. NRC Office of Inspector General, NRC's Regulation of Davis-Besse Regarding Damage to the Reactor Vessel Head (Case No. 02-03S, December 2002) — the post-TMI regulatory regime and oversight failures.

2.

NRC event records and OIG (2002) — the 6-inch corrosion cavity, the remaining cladding, and the 2,200-psi coolant margin.

3.

NRC Bulletin 2001-01, "Circumferential Cracking of Reactor Pressure Vessel Head Penetration Nozzles" (3 August 2001), and the compromise that let Davis-Besse run to a 16 February 2002 outage.

4.

NRC OIG (2002), as quoted in U.S. GAO, GAO-04-415 (2004) — the finding that NRC's decision "was driven in large part by a desire to lessen the financial impact on FirstEnergy," and NRC's dispute of it.

5.

D. Lochbaum / Union of Concerned Scientists analysis (2002); U.S. GAO, GAO-04-415 (2004).

6.

J. V. Rees, Hostages of Each Other (1994) — INPO and the fragility of institutional safety capability.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Davis-Besse is the case for how regulator-operator dynamics can erode the capability of an industry that had previously demonstrated, through INPO, that it knew how to engineer safety. Regulatory capture is a capability failure at the institutional layer above operations.

LENS APPROACH

LENS uses Davis-Besse in LEN 7 to study regulatory-capture dynamics and in LEN 8 to compare with INPO's earlier success. The case is a reminder that institutional capability requires sustained re-engineering.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|------------------|--------------|---------|-----|
| ENERGY →                              | → N K G       | → D1 D2 D3 D4 D5 | 5            | 5.4     | 5   |
| Navigating Sociotechnical Constraints |               |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Make safety-critical inspections of known generic hazards non-deferrable, so a mandatory check cannot be traded against a refueling schedule.
- ▶ Require any deferral decision to be argued on safety margin alone, with the utility's economic case explicitly excluded from the record.
- ▶ Preserve the independence of the oversight layer so it can overrule an operator's timeline rather than adopt it.

IN OPERATION / AFTER

- ▶ Audit granted deferrals for whether the regulator reasoned from safety or from the operator's economics, and flag drift toward the latter.
- ▶ Track observable degradation (like boric-acid leakage) against inspection currency, so a known mechanism cannot run uninspected for years.
- ▶ Re-engineer institutional safety capability on a cadence, treating post-incident regimes like INPO as maintained, not permanent.

## REFLECTION QUESTIONS

1. Identify a regulator–operator relationship in your domain in which the regulator may be at risk of accepting the operator’s economic argument over its own safety judgment. What signal would surface it?
2. Design the institutional control that would prevent a Davis–Besse–style deferral from being granted.
3. The erosion here happened openly, through the regulator’s normal process. What audit would distinguish a defensible deferral from one in which the regulator has quietly adopted the operator’s schedule as its own?

**WHO BUILDS THIS** safety science & organizational analysis · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** incident analysis · safety dashboards · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

# Texas City BP Refinery Explosion

**N** Normalization of Deviance **T** Training Gap **K** Knowledge & Institutional Memory **G** Governance & Trust

**IMPACT** 15 killed, 180 injured at BP's Texas City refinery; CSB found systemic safety-culture decay

## IN BRIEF

On 23 March 2005 a startup at BP's Texas City refinery overfilled a distillation tower — operators were working from a miscalibrated level transmitter and a redundant high-level alarm that never sounded — and venting hydrocarbon vapor found an ignition source and exploded, killing 15 workers in trailers parked beside the unit and injuring 180. The Chemical Safety Board found accumulated safety-culture decay, deferred maintenance, and a celebrated cost-cutting program. Its central finding reshaped U.S. industrial safety: BP's **personal**-safety metrics were among the best in the industry, while its **process**-safety state — the integrity of the hazardous process itself — had been drifting, unmeasured. Texas City is the canonical evidence that the wrong measurement, reported as a win, is worse than no measurement at all.

## BACKGROUND

BP's Texas City refinery, one of the largest in the U.S., had absorbed years of deferred maintenance and corporate cost-cutting, with instruments and alarms tolerated in a degraded state. On the surface it looked safe: its personal-injury rates were among the best in the industry.<sup>0</sup> The strong injury numbers actively reassured the organization, because the metric the company trusted most was the one that carried no information about the degrading instruments and alarms on which a safe startup actually depended.

## WHAT HAPPENED

On 23 March 2005, during a unit startup, operators overfilled a distillation tower far past its safe level, working from a miscalibrated level transmitter that falsely showed the level declining while the tower overfilled, a redundant high-level alarm that never sounded, and an unreadable sight glass. Hydrocarbon vapor vented, drifted across the site, found an ignition source — an idling truck — and exploded. Fifteen workers in temporary trailers parked beside the unit were killed and about 180 injured.<sup>1</sup> Every contributor had been tolerated as routine — the faulty tower instruments, a high-level alarm reported broken repeatedly in the prior two years, the trailers parked beside a hazardous unit — so the startup was run blind to a danger the site had long since stopped seeing.

## THE INVESTIGATION

The U.S. Chemical Safety Board's investigation became the most influential in the agency's history.<sup>2</sup> It found accumulated safety-culture decay, deferred maintenance, broken instruments tolerated as routine, trailers sited dangerously close to a hazardous unit, and cost

cutting — corporate 25% budget-reduction targets after the 1999 Amoco merger and again for 2005, and a 2003 site “1,000 Day Goals” program the CSB faulted for measuring personal safety and cost over process safety — celebrated as a success while it consumed the process-safety margin. The CSB drew the distinction that would reshape the field: personal safety metrics “are not a good indicator of process safety performance.”<sup>3</sup> The Board’s force came from naming the deeper error as one of measurement: the company had not failed to measure but had measured the wrong dimension and then trusted the reassuring number it produced. BP’s recordable personal-injury rate was excellent at the moment of the disaster, the front-page metric on which management-incentive plans paid out; the process-safety state of the unit that killed fifteen workers had no comparable reporting line at all.

#### THE CAPABILITY GAP

Texas City is the clearest evidence that an organization can measure the wrong thing and report excellent performance while its real capability gap widens. Personal-safety metrics — slips, trips, recordable injuries — carried no information about the integrity of the hazardous process, so the signal regime was blind in the very dimension that killed people. The wrong measurement, trusted, is worse than none.<sup>4</sup> A blank dashboard at least invites suspicion; a confident green reading on the wrong axis manufactures false assurance, which is why the excellent injury record made the process-safety drift harder to see rather than easier.

#### AFTERMATH & REFORM

The Baker Panel — chaired by former Secretary of State James Baker, commissioned by BP at the CSB’s recommendation — reported in 2007 that the failure was a corporate one, not a Texas City one: a safety-culture decay extending across BP’s five U.S. refineries, driven by cost pressure and management-of-change failures that the personal-safety metric system was structurally incapable of surfacing. BP invested heavily in process safety afterward, and the process-safety/personal-safety distinction — developed in the CCPS literature and codified in OSHA’s 1992 process-safety-management standard — moved into mainstream U.S. industrial regulation after 2005, with API Recommended Practice 754 (2010) establishing process-safety leading and lagging indicators as an industry standard. The case’s lasting contribution is a measurement lesson: count the thing that can kill you, not the thing that is easy to count.<sup>5</sup> That the distinction had been available in the CCPS literature and the OSHA standard before the explosion underscores the point: the knowledge existed, but the refinery’s reporting had not been built to carry it upward where the hazard actually lived.



*Personal safety metrics are not a good indicator of process safety performance.*

U.S. CHEMICAL SAFETY BOARD, BP TEXAS CITY FINAL INVESTIGATION REPORT, 2007

REFERENCES

1. CSB (2007) — accumulated safety-culture decay, deferred maintenance, and the siting of occupied trailers beside a hazardous unit.

2. U.S. Chemical Safety and Hazard Investigation Board, Refinery Explosion and Fire, Investigation Report 2005-04-I-TX (2007) — the startup, malfunctioning instruments, and 15 killed / 180 injured.

3. A. Hopkins, Failure to Learn: The BP Texas City Refinery Disaster (CCH, 2008).

4. CSB (2007) — “indicators of personal safety are not indicators of process safety” (quoted).

5. OSHA Process Safety Management standard (29 CFR 1910.119, 1992) and the CCPS process-safety literature — the personal-vs-process-safety distinction.

6. Report of the BP U.S. Refineries Independent Safety Review Panel (the Baker Panel, 2007).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Texas City is the clearest evidence that organizations can be measuring the wrong thing and reporting excellent performance while their actual capability gap widens. Personal-safety metrics had no information about process-safety state. The signal regime was blind in the dimension that killed people.

LENS APPROACH

Texas City is the canonical “measuring the wrong failure mode” case (induced 2.1; LENS D4/PT5), with cue/alert design (induced 3.1) and change-control (induced 5.4) as alternate anchors. LENS uses it in LEN 4 as the wrong-measurement case and in LEN 7 for the governance failure that allowed years of cost-cutting to be reported as wins. Studio projects design the process-safety measurement deliverable BP’s executives should have been receiving in 2003. Adjacent to Wells Fargo (Case 148) at the measurement-system-blind-to-the-real-failure-mode layer.

| DOMAIN     | FAILURE MODES | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO |
|------------|---------------|---------------------------------------|--------------|---------|-----|
| ENERGY     | → N T K G →   | D1 D2 D3 D4 D5                        | 4            | 5.4     | 5   |
| INDUSTRIAL |               | Navigating Sociotechnical Constraints |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Instrument process-safety state directly — barrier integrity, alarm health, instrument validity — rather than inferring safety from personal-injury rates.
- ▶ Treat degraded instruments and silent alarms as startup-blocking conditions, not routine items to defer past the next run.
- ▶ Set facility-siting rules that keep occupied trailers away from hazardous units as a design constraint, not a tolerated exception.

IN OPERATION / AFTER

- ▶ Audit whether the headline metric actually carries information about the hazard that can kill, and retire reassuring numbers that do not.
- ▶ Track deferred maintenance and tolerated-defect counts as process-safety leading indicators reported to executives.
- ▶ Verify that the process-vs-personal-safety distinction is wired into the reporting chain so it reaches the layer that funds maintenance.

REFLECTION QUESTIONS

1. Identify a “personal safety vs. process safety” equivalent in your domain. What capability gap is invisible to the metric your institution currently reports?
2. Design the process-safety dashboard that BP Texas City’s executives should have been receiving in 2003.

**WHO BUILDS THIS** safety science & organizational analysis · learning science & instructional design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** incident analysis · safety dashboards · scenario simulation · competency models · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

# Upper Big Branch Mine Explosion

**N** Normalization of Deviance   **G** Governance & Trust   **K** Knowledge & Institutional Memory

**IMPACT** 29 killed in West Virginia coal mine; MSHA found two sets of books and required examinations not performed; first U.S. mining-industry CEO convicted of a federal mine-safety charge (misdemeanor)

## IN BRIEF

On 5 April 2010 coal dust and methane ignited at Massey Energy's Upper Big Branch mine in West Virginia, killing 29 miners — the worst U.S. mine disaster in forty years. Investigators found Massey had kept two sets of records: an internal log of actual conditions and a separate, clean log for federal inspectors. Required pre-shift, on-shift, and weekly examinations went unperformed, no methane examinations were made on the longwall tailgate in the weeks before the blast, and the mine's superintendent later admitted ordering a methane monitor disabled. CEO Don Blankenship was convicted of a misdemeanor conspiracy to violate mine-safety standards — the first U.S. mining CEO criminally convicted on such a charge. Upper Big Branch is the dataset's clearest case of a measurement system engineered deliberately to defeat the regulator.

## BACKGROUND

Massey Energy's Upper Big Branch mine in West Virginia operated under federal safety rules enforced through inspections and the records the mine kept. Massey kept two: an internal log of actual conditions, and a separate, clean log maintained for the inspectors.<sup>o</sup> The enforcement regime depended entirely on the records reflecting reality, so a second, sanitized set of books did not merely break a rule — it disabled the very mechanism by which the regulator was supposed to see the mine at all.

## WHAT HAPPENED

On 5 April 2010 accumulated coal dust and methane ignited and tore through the mine, killing twenty-nine miners — the worst U.S. coal disaster in forty years. The conditions that fed the blast — high methane, inadequate ventilation, dust not rendered inert — were the very ones the clean, inspector-facing records had been built to hide.<sup>1</sup> The records had done their intended work right up to the blast: they kept the conditions that killed twenty-nine men out of the regulator's view precisely while those conditions were building toward ignition.

## THE INVESTIGATION

MSHA found required pre-shift, on-shift, and weekly examinations unperformed and no methane examinations at all on the longwall tailgate in the weeks before the blast; a parallel U.S. Department of Justice probe established that a methane monitor had been ordered disabled.<sup>2</sup> Massey CEO Don Blankenship was eventually convicted of a misdemeanor count



of conspiring to willfully violate mine-safety standards — the first U.S. mining-industry CEO criminally convicted on a mine-safety charge — though acquitted on the felony counts.<sup>3</sup> The conviction did not reach the deception: the jury convicted Blankenship of conspiring to violate mine-safety standards and acquitted him of conspiring to defraud the regulator, so the records architecture is proved at the mine and never proved above it.

#### THE CAPABILITY GAP

The dual-records architecture was not the work of a few bad foremen; it was a stable institutional practice sustained across years and multiple management layers. The capability gap was not in the miners but in the executive ranks that designed and operated a measurement system specifically to defeat the regulator — which makes it the book's clearest case of measurement engineered as deception.<sup>4</sup> Where other cases show measurement drifting or pointed at the wrong dimension, here it was deliberately constructed to mislead, so no improvement in the regulator's reading of the official records could ever have helped — the records were the lie.

#### AFTERMATH & REFORM

The case anchors the argument that decision-grade evidence needs structural protection from the institution that produces it, when that institution has a stake in what the evidence says. Blankenship's conviction set a marker for corporate-officer accountability — and left open the regulator-side question: what architecture would have detected two sets of books before twenty-nine people died?<sup>5</sup> Holding a CEO criminally accountable raised the cost of designing such a system, but accountability after the fact is not detection; the unanswered question is what independent signal could have exposed the divergence between the two logs while the mine was still running.

| INSPECTOR<br>LOG                                                         | INTERNAL<br>LOG |
|--------------------------------------------------------------------------|-----------------|
| compliant                                                                | actual          |
| <i>two sets of records sustained for years as institutional practice</i> |                 |
| UPPER BIG BRANCH — MEASUREMENT ENGINEERED AGAINST THE REGULATOR          |                 |

*Massey kept two sets of books — one for inspectors, one for itself.*

PARAPHRASING FEDERAL INVESTIGATORS, MSHA UPPER BIG BRANCH REVIEWS (2011 – 2012)

#### REFERENCES

1. U.S. MSHA, Internal Review of MSHA's Actions at Upper Big Branch and the accident investigation (2011–2012) — the dual records and the 29 deaths.
2. Governor's Independent Investigation Panel (J. Davitt McAteer, 2011) — mine conditions: methane, ventilation, and coal-dust inerting.
3. MSHA and U.S. DOJ findings — suppressed methane readings, disabled monitors, and falsified pre-shift examinations.
4. United States v. Blankenship (S.D.W. Va., 2015–2016) — the misdemeanor conviction and felony acquittals.
5. H. Berkes / NPR investigative reporting on Massey Energy.
6. A. Hopkins, *Disastrous Decisions: The Human and Organizational Causes of the Gulf of Mexico Blowout* (2012) — corporate decision-making and safety (comparative).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Upper Big Branch is the case for deliberately engineered measurement deception. The dual-records practice was not error. It was capability design — by management, against the regulator. The case anchors the LENS argument that decision-grade evidence requires structural protection against the institution that produces it having a stake in what it reports.

LENS APPROACH

LENS uses Upper Big Branch in LEN 4 as the most explicit case for measurement-system protection and in LEN 7 as a corporate-criminal accountability case. Studio projects examine what regulator-side architecture would have detected the dual-records system earlier.

| DOMAIN              | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------|---------------|------------------|--------------|---------|-----|
| ENERGYINDUSTRIAL    | → N G K       | → D1 D2 D3 D4 D5 | 2            | 2.2     | 4   |
| Test and Evaluation |               |                  |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                            | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Source safety-critical data from monitors the operator cannot disable or edit, so a clean official log cannot be constructed by hand.</li><li>Design measurement so the producing institution has no unilateral control over the record the regulator relies on.</li><li>Build cross-checks that compare independent streams, making a single sanitized set of books detectably inconsistent.</li></ul> | <ul style="list-style-type: none"><li>Establish whistleblower and anomaly-detection channels that can surface a divergence between actual and reported conditions while operations continue.</li><li>Audit for the structural conditions that make dual records possible — operator-controlled monitors, no independent verification — not just for violations.</li><li>Pair corporate-officer accountability with detection architecture, since raising the cost of deception does not by itself reveal it in time.</li></ul> |

REFLECTION QUESTIONS

- Where in your domain could two sets of records plausibly be kept? What architectural change would make that impossible?
- What does it mean for a measurement system to be “structurally protected” from the institution that produces it?
- Accountability after the fact is not detection. What independent signal could have exposed the divergence between the two logs while the mine was still running?

WHO BUILDS THIS safety science & organizational analysis · governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS incident analysis · safety dashboards · governance frameworks · bias & evidence audits · knowledge bases · data pipelines

# Deepwater Horizon

**T** Training Gap   **N** Normalization of Deviance   **K** Knowledge & Institutional Memory

**IMPACT** 11 killed; largest marine oil spill in U.S. history; \$65B+ in damages

## IN BRIEF

On 20 April 2010 the Macondo well blew out beneath the Deepwater Horizon rig in the Gulf of Mexico, killing 11 workers and releasing roughly 4.9 million barrels of oil on the federal estimate — the trial court later found 4.0 million — the largest marine spill in U.S. history, and more than \$65 billion in eventual costs to BP. The crew had misread the negative-pressure test, the primary check of well integrity, and kept working a well that was no longer sealed. The National Commission traced the blowout to a series of identifiable mistakes by BP, Halliburton, and Transocean — systematic failures in risk management that cast doubt on the safety culture of the whole industry. Every defense — procedure, training, supervision, contractor coordination, equipment — had drifted independently until none caught the others. It is the book's canonical multi-layer normalization failure. Four major investigations — the National Commission, the joint Coast Guard / BOEMRE inquiry, the Chemical Safety Board's process-safety review, and the Deepwater Horizon Study Group at Berkeley — each surfaced a distinct facet of the same drift; the disagreements among them are themselves decisive for the case.

## BACKGROUND

The Deepwater Horizon was drilling BP's Macondo exploratory well in the Gulf of Mexico when, on 20 April 2010, the well blew out. Eleven workers were killed, the rig burned and sank, and roughly 4.9 million barrels of oil escaped the well over 87 days on the federal Flow Rate Technical Group's estimate — the trial court, weighing the same evidence in 2015, put the release at 4.0 million — of which about 810,000 barrels were captured at the wellhead — the largest marine oil spill in U.S. history, eventually costing BP more than \$65 billion.<sup>9</sup> The well sat at the end of a long chain of contractors and schedules, each operating to its own tolerance, so that the 87 days of uncontrolled flow stand as a measure not of one mistake but of how far the accumulated drift had to be unwound before the Gulf was sealed again.

## WHAT HAPPENED

The proximate trigger was a misread safety check. The crew ran a negative-pressure test — the primary test of whether the well was sealed — got anomalous readings, accepted an incorrect explanation for them, and proceeded to displace the heavy drilling mud. The well was not sealed; hydrocarbons surged up the riser to the rig and ignited.<sup>1</sup> The anomalous pressure was the well speaking plainly that it remained live, but the explanation the crew accepted let the displacement proceed on schedule, and once the heavy mud was gone nothing stood between the reservoir and the rig but a barrier that had already failed.

## THE INVESTIGATION

Four major investigations reached the wreck and pulled on different threads. The National Commission called a series of identifiable mistakes across BP, Halliburton, and Transocean the blowout's immediate causes and a systemic failure of industry management its root cause, and found that the mistakes revealed "such systematic failures in risk management that they place in doubt the safety culture of the entire industry."<sup>2</sup> The joint Coast Guard / BOEMRE inquiry traced the blowout-preventer failure to a buckled drill pipe and a maintenance backlog that the leasing-and-safety regulator had not caught. The Chemical Safety Board's separately published process-safety review made the decisive distinction explicit: BP and Transocean both reported improving personal-injury rates in the year before the blowout, yet personal-safety metrics had no purchase on the catastrophic process-safety risks that produced it — the same inversion the CSB had already documented at Texas City. Government and academic reviews found training that had not covered the well-control situation the crew faced, an unclear chain of command, and a cascade of failed defenses that BP's own accident investigation report summed up as "a complex and interlinked series of mechanical failures, human judgments, engineering design, operational implementation and team interfaces."<sup>3</sup> That the Commission reached past the rig to the whole industry's safety culture marked the failure as structural rather than local: the same drift, it judged, was latent wherever the same pressures and the same defenses operated unexamined.


## THE CAPABILITY GAP

Deepwater Horizon is the canonical multi-layer failure: procedure, training, supervision, contractor coordination, and equipment had each drifted independently until none caught the others. The negative-pressure test was the trigger, but the system had been carrying accumulated procedural debt for years, and the capability to recognize an unsafe well-state was simply not present at the moment it was needed.<sup>4</sup> Because each defense had degraded inside its own tolerance, no single review of any one layer would have flagged a crisis; the danger lived in the correlation between layers, which is precisely the property no individual procedure was designed to watch. The cement job, the blowout preventer's design and maintenance, the well-control culture before the incident, and the regulator's split mandate were each, on their own, a partial story; the failure was the alignment among them, and no defense layer had been instrumented to read that alignment in time.

## AFTERMATH & REFORM

The blowout drove a restructuring of offshore regulation — the Minerals Management Service was broken up and replaced by the Bureau of Safety and Environmental Enforcement (BSEE) and the Bureau of Ocean Energy Management (BOEM), separating the leasing function from the safety function so revenue pressure could no longer erode well-integrity enforcement — and drilling-safety and well-integrity rules were tightened, while BP paid tens of billions in penalties and settlements. The deeper lesson is the normalization one: no single layer failed catastrophically on its own; each had quietly drifted within tolerance until the day the tolerances lined up.<sup>5</sup> Splitting the regulator conceded that an agency both leasing acreage and policing safety carried a built-in conflict, and tightening the well-integrity rules conceded that the negative-pressure test had been a real barrier all along — one the system had been quietly permitted to misread. The hedges survive into the case: the process-safety / personal-safety distinction the CSB names is decisive and easy to lose; the OSHA-vs-EPA

enforcement gap on offshore facilities was structural, not incidental; and the four investigations did not converge on a single cause precisely because the catastrophe had several, all of which had been quietly tolerated.



*The immediate causes of the Macondo well blowout can be traced to a series of identifiable mistakes... that reveal such systematic failures in risk management that they place in doubt the safety culture of the entire industry.*

NATIONAL COMMISSION, DEEP WATER (REPORT TO THE PRESIDENT), 2011

## REFERENCES

1. Spill-volume estimates — the Flow Rate Technical Group's 4.9 million barrels against the 4.0 million barrels found by the district court in *In re Oil Spill by the Oil Rig "Deepwater Horizon"* (E.D. La., January 2015), less 810,000 barrels collected — and BP cost disclosures (>\$65 billion); the reorganization of the Minerals Management Service into BSEE/BOEM (Secretarial Order 3299, 2010); A. Lustgarten, *Run to Failure: BP and the Making of the Deepwater Horizon Disaster* (W.W. Norton, 2012) — long-arc account of accumulated procedural debt.
2. National Commission (2011) — the well-control sequence and mud-displacement decision.
3. National Commission (2011) — "systematic failures in risk management... place in doubt the safety culture of the entire industry" (quoted); U.S. Chemical Safety Board, *Drilling Rig Explosion and Fire at the Macondo Well* (final volumes, 2014–2016) — process-safety vs. personal-safety distinction and the improving personal-injury rates BP and Transocean reported as misleading indicators.
4. BOEMRE / U.S. Coast Guard Joint Investigation (2011) — blowout-preventer maintenance backlog and chain-of-command findings; National Academies, *Macondo Well Deepwater Horizon Blowout: Lessons for Improving Offshore Drilling Safety* (2012) — training gaps and the interlinked cascade of failed defenses; BP, *Deepwater Horizon Accident Investigation Report* (the Bly report, 8 September 2010) — BP's own summation, "a complex and interlinked series of mechanical failures, human judgments, engineering design, operational implementation and team interfaces" (quoted).
5. Deepwater Horizon Study Group (UC Berkeley, 2011) final report; N. Leveson, systems-theoretic analysis of Deepwater Horizon — the multi-layer drift and the limits of single-cause framings.
6. National Commission on the BP Deepwater Horizon Oil Spill, *Deep Water: The Gulf Oil Disaster and the Future of Offshore Drilling* (Report to the President, 2011) — systematic management failures as root cause; the misread negative-pressure test.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Deepwater Horizon is the canonical multi-layer failure: training, procedure, supervision, contractor-coordination, and equipment all drifted independently until none caught the others. The negative-pressure test was the proximate trigger, but the system as a whole had been operating with accumulated procedural debt for years. The capability to recognize an unsafe well-state simply was not present at the moment it was needed. The CSB's process-safety / personal-safety distinction is decisive: the improving personal-injury rates BP and Transocean reported were a measurement of the wrong construct, and reading them as safety was itself a normalization.

### LENS APPROACH

Deepwater Horizon anchors the cue-and-alert-design competency (induced 3.1; LENS D4/PT5): the negative-pressure test was the cue, and the cue’s ambiguity at the moment of decision was the capability gap. LENS uses it in Domain 4 (Test and Evaluation) for the cue-design failure and the wrong-construct measurement; in Domain 1 (Systems Analysis) for multi-layer drift across procedure, training, supervision, contractor coordination, and equipment; and in Domain 5 (Navigating Sociotechnical Constraints) for the OSHA-vs-EPA enforcement gap and the leasing-vs-safety regulator conflict the BSEE/BOEM split conceded. Pair with Texas City (Case 161) on the process-safety-vs-personal-safety inversion, and with Challenger / Columbia (Case 111) on the multi-layer-drift form.

| DOMAIN              | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO   |
|---------------------|---------------|------------------|--------------|---------|-------|
| ENERGY              | → T N K       | → D1 D2 D3 D4 D5 | 5            | 3.1     | 4 · 5 |
| Test and Evaluation |               |                  |              |         |       |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>▶ Make the negative-pressure test a hard, instrumented gate: define the pass criteria so an anomalous reading stops work rather than inviting a rationalizing explanation.</li><li>▶ Specify a single, unambiguous well-control chain of command across operator and contractors before drilling begins, so no decision falls between organizations.</li><li>▶ Build well-control training around the ambiguous small-anomaly case the crew actually faced, not only the textbook blowout.</li></ul> | <ul style="list-style-type: none"><li>▶ Audit each defense layer for drift independently — and audit the correlation between them, since the danger lived where tolerances aligned.</li><li>▶ Track accumulated procedural debt as a standing metric so silent degradation surfaces before tolerances coincide.</li><li>▶ Structurally separate the regulator’s leasing and safety roles so revenue pressure cannot erode well-integrity enforcement.</li></ul> |

REFLECTION QUESTIONS

1. Identify a multi-defense process in your domain. For each layer, what is the procedural debt that has accumulated since it was designed?
2. The negative-pressure test was the proximate trigger. Design a capability check that would have caught the misinterpretation in real time — and specify what makes the cue unambiguous enough that an anomalous reading stops work rather than inviting a rationalizing explanation.
3. The Commission judged the drift industry-wide, not rig-specific. What measure in your domain would reveal whether a single failure is a local lapse or a sample from a systemic distribution?
4. The Deepwater Horizon had gone seven years without a single lost-time incident. Name the measurement in your domain that is most at risk of being the wrong construct — a personal-safety analog where the catastrophic process-safety risk lives elsewhere — and design the second instrument that would surface the harm the first one cannot see.

WHO BUILDS THIS learning science & instructional design · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · incident analysis · safety dashboards · knowledge bases · data pipelines

# Grenfell Tower

**G** Governance & Trust   **T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** 72 killed in a residential tower fire in London; decades of regulatory failure

## IN BRIEF

On 14 June 2017 a fire spread up the exterior of Grenfell Tower, a London public-housing block, killing 72 people. It raced because the tower had been wrapped in combustible aluminium-composite cladding during a refurbishment. The public inquiry found the disaster the culmination of decades of failure: cladding firms engaged in “systematic dishonesty,” marketing combustible materials as safe; the council’s building control surveyor was overworked, untrained in cladding, and visited less than he meant to; and the London Fire Brigade, whose “stay put” advice proved fatal, was unprepared for a cladding fire whose risks earlier incidents had already shown. The failure spanned manufacturers, regulators, inspectors, and responders — each contributing a piece, none owning the whole. Grenfell is the book’s case for capability failure distributed across many hands.

## BACKGROUND

Grenfell Tower was a 1970s public-housing block in West London, refurbished between 2012 and 2016 with new exterior cladding intended to improve the building’s appearance and efficiency. The cladding chosen used a combustible aluminium-composite material — installed while the project’s fire engineer never finished the fire safety strategy and never warned the design team, and while published guidance against combustible high-rise cladding sat unread by the people specifying the panels.<sup>0</sup>

## WHAT HAPPENED

On 14 June 2017 a kitchen fire broke out and, instead of staying contained as a tower’s compartment design assumes, climbed the building’s exterior on the combustible cladding, wrapping the tower in flame within minutes and defeating the very principle the building was meant to rely on. Residents, following long-standing “stay put” advice premised on that containment, remained in their flats as the route to safety closed around them; 72 people died.<sup>1</sup>

## THE INVESTIGATION

The Grenfell Tower Inquiry found the fire the culmination of decades of failure by central government and every body responsible. Cladding companies had engaged in “systematic dishonesty,” marketing combustible products as safe and corrupting the very test data buyers relied on; the council’s building control surveyor visited far less often than the monthly schedule he intended, and never recognized that the panels and insulation in use were combustible.<sup>2</sup> The London Fire Brigade was unprepared: the risks of rapidly developing cladding fires were known from prior incidents — Knowsley Heights, Garnock

Court, Shepherd's Court — but the commanders who attended “had received no training in the particular dangers associated with combustible cladding,” so each near-miss taught no one.<sup>3</sup>

### THE CAPABILITY GAP

Grenfell is the book's strongest evidence that capability failure can be distributed across many actors, each contributing a small piece and none accountable for the whole. Manufacturer fraud, regulatory capture, inspection incompetence, training gaps, and lost institutional memory all converged on one building, and because each actor saw only its fragment, each could regard its own part as tolerable. It was what one might call a “grey elephant” — a danger known but ignored, hiding in plain sight — and the missing capability was anyone owning the integrated risk that everyone could see in part but no one held in full.<sup>4</sup>

### AFTERMATH & REFORM

The inquiry's Phase 2 report (2024) and the government response (2025) drove an overhaul of building-safety regulation, cladding remediation, and fire-service doctrine — reforms that, between them, tried to assign the ownership the original system had left vacant.<sup>5</sup> Grenfell's lesson is the governance one this chapter turns on: when responsibility for a known risk is split across dozens of actors, the risk has, in effect, no owner — and a system with no owner for its gravest hazard will eventually pay for it, in the currency of the people living inside it.

manufacturer fraud

regulatory capture

building control

LFB unprepared

▲ 72 dead — “a grey elephant, known but ignored”

LEVELS OF FAILURE: MANUFACTURE, REGULATION, INSPECTION, RESPONSE

*The LFB's preparation and planning for a fire such as that at Grenfell Tower was gravely inadequate.*

GRENFELL TOWER INQUIRY, PHASE 1, 2019

### REFERENCES

1. Grenfell Tower Inquiry, Phase 2 Report (2024) — decades of failure and the combustible-cladding decision.
2. Grenfell Tower Inquiry, Phase 1 Report (2019) — the fire's spread up the cladding and the failure of “stay put.”
3. Phase 2 Report (2024) — cladding firms' “systematic dishonesty” and the RBKC building control failures (ch. 62).
4. Phase 1 Report (2019) — London Fire Brigade unpreparedness; preparation and planning “was gravely inadequate” (quoted).
5. B. Hutter & M. Power (eds.), *Organizational Encounters with Risk* (2005) — distributed risk ownership.
6. UK Government response to the Grenfell Phase 2 report (2025) — building-safety and fire-service reform.

### THE LEARNING ENGINEERING LENS

#### LE INSIGHT



Grenfell is the strongest evidence in the dataset that capability failure can be distributed across many actors, each of whom contributes a small piece, none of whom is accountable for the whole. The “grey elephant” framing — a known danger no one owns — describes a pattern that LENS treats as a primary governance problem in any high-consequence domain.

#### LENS APPROACH

LENS reads Grenfell through the institutional-memory-of-warnings channel (induced 7.4, with a 6.2 secondary): prior cladding-fire warnings existed — Lakanal House, the London Fire Brigade’s own knowledge from earlier incidents — but none ever reached the refurbishment and cladding decision that could have acted on them. The capability deliverable is the channel that carries a known warning to the decision empowered to stop the work; the cladding firms’ systematic dishonesty was a real aggravator but is the secondary thread, not the lesson.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY | PROBLEM TYPE | INDUCED | LEO   |
|---------------------------------------|---------------|-----------------|--------------|---------|-------|
| INDUSTRIAL                            | → G T K N →   | D1 D2 D3 D4 D5  | 4            | 7.4     | 2 · 5 |
| Navigating Sociotechnical Constraints |               |                 |              |         |       |

#### APPROACHES TO CONSIDER

##### DURING DEVELOPMENT

- ▶ Assign one accountable actor to own the building’s integrated fire risk end to end, so no hazard can fall into the gaps between manufacturers, inspectors, and responders.
- ▶ Verify combustible-material claims against independent evidence rather than trusting vendor marketing, treating “systematic dishonesty” as a threat the process must defeat.
- ▶ Preserve the containment principle the building relies on: gate any cladding choice on whether it keeps a compartment fire from climbing the exterior.

##### IN OPERATION / AFTER

- ▶ Route every site inspection and near-miss into a shared record so repeated visits cannot each miss the same banned material in isolation.
- ▶ Feed prior-incident knowledge into firefighting policy, training, and “stay put” doctrine, so lessons from earlier cladding fires actually change practice.
- ▶ Sustain a single line of accountability for the integrated hazard through the building’s life, not only at refurbishment.

#### REFLECTION QUESTIONS

1. What is the “grey elephant” — the well-known risk that nobody owns — in your domain?
2. Design the deliverable that forces a single actor to own the integration risk that Grenfell distributed across dozens.
3. Site inspections, prior cladding fires, and published guidance each touched a fragment of the Grenfell hazard, yet none assembled it. What mechanism in your domain could gather scattered partial warnings into one picture in front of someone empowered to stop the work?

**WHO BUILDS THIS** governance, ethics & measurement · learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

# Camp Fire / PG&E

**G** Governance & Trust   **N** Normalization of Deviance   **K** Knowledge & Institutional Memory

**IMPACT** 85 killed in Butte County, California; deadliest U.S. wildfire in a century; PG&E pleaded guilty to 84 counts of involuntary manslaughter

## IN BRIEF

On 8 November 2018 a worn hook on a nearly century-old PG&E transmission line failed in high winds and drought, igniting the Camp Fire, which swept into the town of Paradise faster than people could evacuate. Eighty-five died — the deadliest U.S. wildfire in a century — and PG&E later pleaded guilty to 84 counts of involuntary manslaughter. Investigators found PG&E had known for years that its transmission infrastructure across high-fire-risk areas was deteriorating, and had run it under a strategy its own documents called Run to Failure while cutting the time and cost budgeted for each inspection, so their thoroughness fell even as their nominal frequency held. Infrastructure built for one climate was operating in another. Camp Fire is the book's foundational climate-era case for utility capability under changed conditions, and it restructured how California regulates utility wildfire risk.

## BACKGROUND

PG&E operated aging transmission lines across the wildfire-prone foothills of Northern California — some hardware approaching a century old — in a climate growing hotter and drier than the one the grid had been built for, so the operating environment had drifted away from the assumptions the infrastructure was designed against.<sup>0</sup>

## WHAT HAPPENED

On 8 November 2018 a worn C-hook on a PG&E transmission tower failed in high winds, dropping a live line and igniting a fire under severe drought conditions — a single piece of aged hardware setting off a catastrophe the dry, windy conditions stood ready to amplify. The fire moved into the town of Paradise faster than its evacuation routes could clear it; 85 people died — the deadliest U.S. wildfire in a century. PG&E later pleaded guilty to 84 counts of involuntary manslaughter, accepting criminal responsibility at a scale rare for a utility.<sup>1</sup>

## THE INVESTIGATION

CalFire's investigation and the Butte County District Attorney's report found that PG&E had known for years about the deteriorating condition of its transmission infrastructure in high-fire-risk areas, and had run the line under a maintenance strategy its own asset-management documents named "Run to Failure" — the District Attorney found that PG&E personnel tried to distance the company from the phrase and that "the evidence clearly establishes quite the opposite." Between the 2005 and 2018 procedure manuals the frequency of inspections and patrols did not change; what PG&E cut was the unit cost — the time allotted to each structure — and with it, the District Attorney found, the thoroughness

of the inspection. Nobody in PG&E was inspecting the C hooks at all, so the hazard was not unknown but a recognized risk left unlooked-at.<sup>2</sup> The gap was simultaneously at the utility's asset-maintenance decisions and at the regulatory architecture that had allowed the deferrals to continue, neither side holding a line that would have forced the work.<sup>3</sup>

#### THE CAPABILITY GAP

Camp Fire is the climate-era case for utility-capability failure under changing risk. The infrastructure had been designed and maintained for one set of conditions and was operating in another, more dangerous one, so the safety margins the original design assumed had quietly eroded as the climate shifted beneath them. The capability to update operations — inspection cadence, vegetation management, de-energization, replacement — to match the actual risk did not exist as an institutional deliverable, on either the utility's side or the regulator's, leaving no one tasked with closing the widening gap.<sup>4</sup>

#### AFTERMATH & REFORM

PG&E entered bankruptcy under tens of billions in wildfire liability, and California restructured how it regulates utility wildfire-risk planning — mandatory mitigation plans, inspections, and public-safety power shutoffs, turning the deferred work into requirements with teeth behind them.<sup>5</sup> Paired with the Northeast Blackout (Case 146), Camp Fire shows utility capability failing in a second way: not a silent control-room failure but a slow, known erosion of physical infrastructure against a rising hazard the institution declined to fund against — a failure measured in years of deferral rather than seconds of cascade.

## 85

### KILLED IN THE CAMP FIRE · SINGLE TRANSMISSION HOOK

*infrastructure designed for one climate, operating in another*

#### CAMP FIRE — CAPABILITY MISMATCH UNDER CHANGED CONDITIONS

*Catastrophic failure of the Caribou-Palermo line was not an “if” question; it was a “when” question.*

BUTTE COUNTY DISTRICT ATTORNEY, THE CAMP FIRE PUBLIC REPORT, 16 JUNE 2020

#### REFERENCES

1. CalFire, Camp Fire Investigation Report (2019) — the worn transmission-line hardware as ignition source.
2. The Camp Fire record — 85 killed; PG&E's guilty plea to 84 counts of involuntary manslaughter (2020).
3. Butte County District Attorney, The Camp Fire Public Report (2020) — PG&E's knowledge and deferred maintenance (quoted).
4. California Public Utilities Commission, Order Instituting Investigation into PG&E (2019) — the regulatory dimension.
5. PG&E bankruptcy and California wildfire-mitigation reforms (2019–).
6. Cf. the Northeast Blackout (Case 146); climate-and-infrastructure literature.

#### THE LEARNING ENGINEERING LENS

LE INSIGHT

The Camp Fire is the canonical climate-era case for utility- capability failure under changing risk conditions. The infrastructure was designed for one set of conditions and operated in another. The capability to update operations to match actual conditions did not exist as an institutional deliverable.

LENS APPROACH

LENS uses the Camp Fire in LEN 7 as a legacy-asset case (induced 7.3): infrastructure designed for one operating regime and run unchanged into another as the climate shifted the fire risk underneath it. The teaching point is the assumption-migration review as a recurring engineered obligation — an asset whose safety case rests on conditions that no longer hold needs a standing, funded re-verification that the original design assumptions still match the world, not a one-time sign-off. It is the single taught example of the legacy-asset-aged-past-its-regime pattern; the case pairs with the Northeast Blackout (Case 146) as utility-capability failures of different kinds.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|------------------|--------------|---------|-----|
| ENERGY                                | → G N K       | → D1 D2 D3 D4 D5 | 1            | 7.3     | 1-5 |
| Navigating Sociotechnical Constraints |               |                  |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Re-derive the infrastructure’s safety margins against current conditions, not the historical climate it was designed for, and treat the gap as a hazard to close.</li><li>Make updating operations to match rising risk — inspection cadence, vegetation management, replacement — an explicit institutional deliverable with an accountable owner.</li><li>Build de-energization and equipment-replacement triggers tied to known high-risk hardware in high-fire-risk areas.</li></ul> | <ul style="list-style-type: none"><li>Audit deferred maintenance against the hazard it guards against, so a recognized, deteriorating risk cannot be postponed year after year.</li><li>Hold the regulator’s line with mandatory mitigation plans, inspections, and penalties, so the cost of resilience is forced before the disaster rather than after.</li><li>Monitor aging hardware in the highest-risk corridors as a standing priority, treating a worn critical component as an active threat, not a backlog item.</li></ul> |

REFLECTION QUESTIONS

1. Identify infrastructure in your domain that was designed for one set of conditions and is now operating in another. What is the capability deliverable to bridge the gap?
2. Design the regulatory architecture that would prevent a utility from deferring critical wildfire-risk maintenance.
3. PG&E knew its hardware was deteriorating in high-fire-risk areas yet deferred the work for years, and the regulator let it. What in your domain lets a recognized, rising hazard be postponed indefinitely — and who would have to hold the line that forces the spending before the catastrophe?

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · incident analysis · safety dashboards · knowledge bases · data pipelines

AFTER THE EVENT — RESPONSE AND RECOVERY

# Three Mile Island

**T** Training Gap    **H** Human-System Interface

**IMPACT** Partial meltdown of a Babcock & Wilcox PWR; minimal off-site dose; most serious accident in U.S. commercial nuclear history; catalyst for industry-wide reform

## IN BRIEF

On 28 March 1979 a small cooling fault at Three Mile Island's Unit 2, a Babcock & Wilcox reactor near Harrisburg, escalated into a partial core meltdown — the most serious accident in U.S. commercial nuclear history. A relief valve stuck open while a control-room light reported it closed, and operators trained for dramatic design-basis ruptures misread the slow, ambiguous cascade and cut the cooling the core needed. A nearly identical near-miss at Davis-Besse eighteen months earlier had never been propagated to the fleet. The Kemeny Commission concluded the fundamental problems were people-related, not equipment. Off-site radiation was minimal, but the accident produced a system of reform — most enduringly INPO — making it the book's paired example of a failure that engineered lasting capability.

## BACKGROUND

Three Mile Island's Unit 2, a Babcock & Wilcox reactor near Harrisburg, was run by operators trained for large, fast, design-basis ruptures — not the slow, ambiguous small-break sequence that actually came.<sup>0</sup> It should not have surprised anyone: in September 1977 a nearly identical stuck-open relief valve had occurred at Davis-Besse, but neither the utility, the vendor, nor the NRC grasped its significance or pushed the lesson out to the fleet — so the precise sequence that would later threaten the core had already been demonstrated and then quietly filed away rather than turned into a drill or a warning.<sup>1</sup> The industry's posture rested on a probabilistic argument that catastrophic accidents in light-water reactors were vanishingly unlikely; that argument had become institutional doctrine, shaping what was trained for, what was instrumented, and what near-misses were treated as.

## WHAT HAPPENED

On 28 March 1979 a minor secondary-cooling upset tripped the reactor; a power-operated relief valve (PORV) opened and then stuck open, draining coolant from the primary loop. The control-room light reported the **command** to close it, not its actual position, so the panel read "closed" while the valve stayed open. Misreading the rising pressurizer level, the operators throttled back the high-pressure injection the starving core depended on — taking the one action that turned a recoverable upset into a meltdown, precisely because the instrument they trusted was telling them a state that was not true. About half the core melted; off-site radiation, as it turned out, was minimal, though the public-communication

failure that accompanied the accident — the contradictory updates from utility, state, and federal officials — would shape every nuclear emergency-response standard built after.<sup>2</sup>

#### THE INVESTIGATION


The Kemeny Commission (October 1979) inverted the industry's assumptions: "the fundamental problems are people-related problems and not equipment problems." The equipment was good enough that, but for the human failures, the accident would have been minor. The criticism reached past the operators to management, the utility, and the NRC — an institution-wide belief that serious accidents were effectively impossible, a complacency that had shaped training priorities, staffing, and oversight long before the relief valve ever stuck, so that the operators inherited a posture they had not chosen.<sup>3</sup> The companion Rogovin report for the NRC documented the Davis-Besse precursor and the institutional failure to disseminate it, and the decades-long literature that followed — Walker's historical synthesis, Rees's institutional study — treated TMI as the moment U.S. nuclear stopped being able to assume away its own human factors.<sup>4</sup>

#### THE CAPABILITY GAP

The gap was not intelligence but the right capability for the event that arrived. Operators drilled on design-basis ruptures had no model for an ambiguous cascade, and a control room reporting commands rather than states made correct diagnosis nearly impossible — the interface and the training were mismatched to the failure that actually arrived. Beneath that lay a second failure: the capacity to **learn** — to turn Davis-Besse into fleet knowledge — had itself broken down, so the industry kept making the same diagnosis blind, which is why the case resists any single-cause reading.<sup>5</sup> The third layer was institutional: there was no organization whose job was to force operating experience from one utility into the training and procedures of every other, and no body the NRC trusted to do it on the industry's behalf.

#### AFTERMATH & REFORM

TMI produced not another accident but a system of reform. Within nine months the industry created the Institute of Nuclear Power Operations (INPO, December 1979) to set standards, run plant evaluations, and force the sharing of operating experience the Davis-Besse failure had lacked.<sup>6</sup> The NRC overhauled licensing and inspection: it required plant-referenced simulators, expanded the resident inspector program to every commercial reactor, tied operator training and re-qualification to a national standard, and in Regulatory Guide 1.97 specified the post-accident instrumentation needed to give crews state — not command — information across the regimes accidents actually produce, with subsequent control-room human-factors standards (NUREG-0700) carrying the principle into every new and modernized plant.<sup>7</sup> Each measure attacked one of the failures the accident had exposed: the training gap, the broken learning channel, the gulf between what indicators displayed and what crews inferred, and the institutional absence of a body responsible for cross-fleet learning. TMI is paired later with INPO (Case 175) as the book's strongest argument that failure can engineer durable industry-level reform.



*training matched worst case, not this case*

A MINOR COOLING FAULT, A VALVE REPORTED SHUT, A MELTDOWN

*The fundamental problems are people-related problems and not equipment problems.*

KEMENY COMMISSION REPORT ON THREE MILE ISLAND, 1979

## REFERENCES

1. Kemeny Commission, Report of the President's Commission on the Accident at Three Mile Island (Oct. 1979) — operator training oriented to large design-basis accidents.
2. M. Rogovin & G. T. Frampton, Three Mile Island: A Report to the Commissioners and to the Public, NUREG/CR-1250 (U.S. NRC, 1980) — the September 1977 Davis-Besse stuck-PORV precursor and the failure to disseminate its lessons.
3. U.S. Nuclear Regulatory Commission, Backgrounder on the Three Mile Island Accident — accident sequence, misleading PORV indication, throttled high-pressure injection, and the off-site dose — about 1 millirem average, under 100 millirem at the site boundary.
4. Kemeny Commission (1979) — central conclusion that the fundamental problems were people-related, not equipment-related (quoted).
5. J. V. Rees, Hostages of Each Other: The Transformation of Nuclear Safety since Three Mile Island (Univ. of Chicago Press, 1994) — TMI as the origin of systematic capability reform.
6. C. Perrow, Normal Accidents (1984); J. Reason, Human Error (1990) — why ambiguous cascades defeat design-basis training and command-not-state interfaces.
7. Institute of Nuclear Power Operations, "Our History" — INPO established December 1979 to set standards and force sharing of operating experience; see also NEI, "Lessons from the 1979 Accident at Three Mile Island."
8. U.S. Nuclear Regulatory Commission, Regulatory Guide 1.97, Rev. 2, Instrumentation for Light-Water-Cooled Nuclear Power Plants to Assess Plant and Environs Conditions During and Following an Accident (Dec. 1980), and NUREG-0700, Human-System Interface Design Review Guidelines — the control-room and post-accident-instrumentation standards that grew out of TMI; see also J. S. Walker, Three Mile Island: A Nuclear Crisis in Historical Perspective (Univ. of California Press, 2004) on the operator-training-and-licensing overhaul.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

TMI is the textbook moment when an industry discovered that its capability assumptions did not survive contact with operational reality. Training to design-basis events left the operators blind to the genuinely ambiguous failures that complex systems actually produce; a control room that reported commands rather than states made correct diagnosis structurally hard; and the channel that should have carried Davis-Besse out to every plant did not exist. The reform — INPO, plant-referenced simulators, resident inspectors, RG 1.97 instrumentation, NUREG-0700 human-factors standards — built the missing infrastructure together rather than one piece at a time, which is the case's decisive teaching.

### LENS APPROACH

TMI is the worked example of induced sub-competency 6.1 (industry-level institution building after catastrophe) and the LENS D5/PT4 pairing — Navigating Sociotechnical Constraints applied to the institutional architecture a catastrophic-system industry needs to learn at scale. Students reconstruct the capability requirements that the design-basis training framework missed (LENS D1), examine the control-room interface and post-accident instrumentation as evidence-architecture problems (LENS D3), and design the cross-fleet learning channel whose absence at Davis-Besse let TMI happen. The case pairs with INPO (Case 175) as the institution-building counterpart to the failure that produced it, and with Fitzgerald/McCain

(Case 124) as the contrasting failure where the missing learning channel was internal to one service rather than across a civilian industry. LEO mapping: LEO-5 (Sociotechnical Constraints) primary for the INPO/NRC institutional architecture; LEO-1 (Systems Analysis) for the interface-and-training requirements analysis the accident exposed.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO   |
|---------------------------------------|---------------|------------------|--------------|---------|-------|
| ENERGY                                | → T H         | → D1 D2 D3 D4 D5 | 4            | 6.1     | 5 · 1 |
| Navigating Sociotechnical Constraints |               |                  |              |         |       |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>▶ Design control-room indicators to report actual system state, not the command issued, so operators diagnose from truth rather than intent.</li><li>▶ Build training around ambiguous, slow-onset cascades — not only dramatic design-basis ruptures — using simulators that mirror the operators’ own plant.</li><li>▶ Engineer a learning channel that propagates precursor events across the fleet as drills, so a near-miss at one site becomes practiced knowledge everywhere.</li></ul> | <ul style="list-style-type: none"><li>▶ Audit whether near-misses are actually reaching operators as changed procedure and training, and force-share operating experience through an independent standards body (the INPO model).</li><li>▶ Sustain plant-referenced simulators and resident oversight so diagnostic skill on ambiguous transients does not decay back to textbook drills.</li><li>▶ Monitor the institutional assumption that serious accidents are impossible, and treat its persistence as a measurable readiness risk.</li></ul> |

REFLECTION QUESTIONS

1. TMI operators were trained for worst-case scenarios but failed in an ambiguous one. What is the equivalent training gap in your domain between the trained case and the messy case?
2. The Kemeny Commission called the human element the dominant risk. What evidence would you need to demonstrate the same conclusion in your own domain?
3. The Davis-Besse precursor occurred eighteen months earlier but never reached the operators who needed it. Design the channel in your domain that would turn a near-miss at one site into a drill at every other before the second event arrives.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · usability testing · cognitive walkthroughs



# Hurricane Katrina — The FEMA/DHS Response Failure

**G** Governance & Trust    **K** Knowledge & Institutional Memory

**IMPACT** Approximately 1,400 deaths under revised counts (earlier official estimates near 1,800; attribution varies by study) and roughly \$125 billion in damage after Katrina's August 29, 2005 landfall; the most-investigated disaster response failure in U.S. history — House Select Bipartisan Committee "A Failure of Initiative" (February 2006), Senate HSGAC "Hurricane Katrina: A Nation Still Unprepared" (May 2006), and the White House Townsend report (February 2006) converged on a response that failed against a scenario the July 2004 Hurricane Pam exercise had simulated in detail thirteen months earlier; reform arrived as the Post-Katrina Emergency Management Reform Act of 2006

## IN BRIEF

Hurricane Katrina made landfall on the Gulf Coast on 29 August 2005; the levee system protecting New Orleans failed, and the federal response failed with it. The casebook's spine for the case is that the missing capability was known before the event: in July 2004, the FEMA-funded Hurricane Pam exercise gathered roughly 300 local, state, and federal officials around a simulated slow-moving Category 3 hurricane whose storm surge overtopped the New Orleans levees, flooded the city, and stranded a population the plans could not move — almost exactly the scenario that arrived thirteen months later. Pam's action items were unfunded and unfinished when Katrina hit. In the event, FEMA's logistics system could not track or deliver commodities; the National Response Plan's Incident of National Significance machinery was invoked only after landfall; unified command across local, state, and federal levels never formed; communications interoperability collapsed; and the Superdome and Convention Center were known to television audiences before they were known to the federal operations picture. Death-toll estimates range from roughly 1,400 under revised counts to earlier official figures near 1,800, with attribution varying by study. Three major investigations followed, and the Post-Katrina Emergency Management Reform Act of 2006 rebuilt FEMA's authorities around a National Preparedness System.

## BACKGROUND

New Orleans's vulnerability was not a discovery. The city sits largely below sea level behind a federal levee system, and the catastrophic-flood scenario had been anticipated for years in emergency-management planning. In July 2004 — thirteen months before Katrina — FEMA funded the Hurricane Pam exercise in Baton Rouge, bringing together roughly 300 local,

state, and federal emergency-response officials around a simulated slow-moving Category 3 hurricane with 120-mile-per-hour sustained winds, up to twenty inches of rain, and a storm surge that overtopped the levees in the New Orleans area. The scenario projected ten to twenty feet of flooding across most of the city and surrounding parishes, the destruction of hundreds of thousands of buildings, 300,000 residents who would not evacuate in advance, and fatalities in the tens of thousands. Pam produced draft functional plans — search-and-rescue, sheltering, temporary medical care, debris — and a schedule of follow-on workshops to turn the drafts into executable, resourced capability. When Katrina made landfall, the follow-on work was partially unfunded and unfinished: the exercise had specified the capability requirement, and no institution had converted the specification into built capability.<sup>9</sup>

#### WHAT HAPPENED

Katrina made landfall on 29 August 2005; the levee system was breached rather than merely overtopped, and roughly eighty percent of New Orleans flooded. The response then discovered, in production, each capability the exercise had specified on paper. FEMA's logistics system could not track commodities in transit or confirm delivery — trucks of ice, water, and food moved without visibility into where they were or whether they arrived. The National Response Plan's machinery for a catastrophic event — the Incident of National Significance designation and the catastrophic-incident annex intended to push resources without waiting for state requests — was invoked only after landfall, in its first operational use, rather than driving pre-landfall mobilization. Unified command across city, state, and federal authorities never formed, and senior officials worked around the designed coordination structures. Communications interoperability collapsed as towers, switches, and emergency-operations centers flooded. Tens of thousands of residents concentrated at the Superdome and the Convention Center in deteriorating conditions; the Convention Center's situation was visible to television audiences before it entered the federal operations picture, and senior officials publicly disputed conditions that broadcast footage had already established. Estimates of the death toll range from roughly 1,400 under revised counts to earlier official figures near 1,800 — attribution of individual deaths to the storm varies by study — with damage of roughly \$125 billion.<sup>1</sup>

#### THE INVESTIGATION

Three major investigations followed within a year, and their convergence is the case's evidentiary backbone. The House Select Bipartisan Committee's A Failure of Initiative (15 February 2006) found failures at every level of government, documented the Hurricane Pam exercise and its unconverted action items in detail, and framed the response in deliberate contrast to the 9/11 Commission's "failure of imagination": the imagination had been exercised, funded, and written down, and what failed was the initiative to act on it. The Senate Homeland Security and Governmental Affairs Committee's Hurricane Katrina: A Nation Still Unprepared (May 2006), built on more than 325 interviews and over 800,000 pages of documents, charged FEMA with failures across personnel deployment, communications, catastrophic planning, commodity pre-staging, medical and search-and-rescue deployment, and evacuation — and recommended dissolving FEMA and rebuilding its function inside DHS. The White House's The Federal Response to Hurricane Katrina: Lessons Learned (the Townsend report, 23 February 2006) issued 125 recommendations and 17 lessons, and found that officials from national leaders to responders on the ground lacked the situational

awareness a prompt response required — the finding that formalized the television-before-operations problem.<sup>2</sup>

### THE CAPABILITY GAP

Reform arrived through statute. The Post-Katrina Emergency Management Reform Act of 2006 (PKEMRA, Public Law 109-295, signed 4 October 2006) rebuilt FEMA's authorities inside DHS rather than dissolving it: it restored preparedness functions to the agency, set qualification requirements for the FEMA Administrator and gave the office a direct line to the President during catastrophes, authorized accelerated federal assistance in advance of a state request for precisely the catastrophic case Katrina represented, and mandated the national preparedness architecture — a national preparedness goal and system under which exercise findings are tracked as corrective actions rather than left as drafts. PKEMRA is the institutional codification of the case's lesson: it converted the discipline that Hurricane Pam lacked — a mechanism obligating someone to fund and finish what an exercise proves is missing — into statutory machinery. Whether the machinery holds is a separate question the casebook leaves open; what the statute establishes is that Congress located the failure where the investigations did, in the seam between demonstrated requirement and built capability.<sup>3</sup>

### AFTERMATH & REFORM

The hedges the case carries are not decoration. The death toll is contested: revised counts place it near 1,400, earlier official figures near 1,800, and the attribution of individual deaths — direct drowning versus indirect cardiovascular and displacement deaths — varies by study and jurisdiction; the casebook reports the range, not a number. The analytical spine is the Hurricane Pam seam: an exercised, documented capability requirement — levee overtopping, a flooded city, 300,000 residents who would not evacuate in advance — that no institution converted into funded, staffed, rehearsed capability in the thirteen months before the event. This is the disaster-preparedness form of stated versus engineered requirements: the stated requirement existed in exercise documentation at a level of specificity the investigations found almost eerie, and the engineered capability did not exist at all. Response is the phase in which the gap is discovered in production, at the cost structure production imposes. The case asks the reader to account not for why no one foresaw the scenario — they did, in writing — but for why foresight without a funded owner is indistinguishable, on the day, from no foresight at all.

*It remains difficult to understand how government could respond so ineffectively to a disaster that was anticipated for years, and for which specific dire warnings had been issued for days.*

A FAILURE OF INITIATIVE, SELECT BIPARTISAN COMMITTEE, FEBRUARY 15, 2006.

### REFERENCES

1. Select Bipartisan Committee to Investigate the Preparation for and Response to Hurricane Katrina (2006), A Failure of Initiative, H. Rept. 109-377, U.S. House of Representatives, February 15, 2006 — including the dedicated Hurricane Pam findings chapter.
2. U.S. Senate Committee on Homeland Security and Governmental Affairs (2006), Hurricane Katrina: A Nation Still Unprepared, S. Rept. 109-322, May 2006 — 325+ interviews, 838,000+ pages of documents reviewed.
3. Post-Katrina Emergency Management Reform Act of 2006, Title VI of Public Law 109-295, signed October 4, 2006 — rebuilt FEMA authorities, accelerated pre-request federal assistance, national preparedness goal and system.
4. U.S. Senate Committee on Homeland Security and Governmental Affairs (2006), Preparing for a Catastrophe: The Hurricane Pam Exercise, hearing, S. Hrg. 109-403, January 24, 2006 — testimony on the July 2004 exercise scope and unfinished follow-on work.

3. The White House (2006), The Federal Response to Hurricane Katrina: Lessons Learned (the Townsend report), February 23, 2006 — 17 lessons, 125 recommendations, including the situational-awareness finding.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Katrina is the load-bearing case for a capability gap that was specified, exercised, and documented before the event and converted into built capability by no one. The Hurricane Pam exercise stated the requirement thirteen months in advance at near-exact fidelity; the response phase then discovered the unbuilt capability in production, at the cost structure production imposes. Foresight without a funded owner is indistinguishable, on the day, from no foresight at all.

LENS APPROACH

Katrina is the stated-versus-engineered-requirements case at national scale (induced 1.1; LENS D1/PT1; LEO-1). The induced anchor is 1.1 — engineered versus stated requirements — because the case’s spine is precisely that seam: Hurricane Pam produced a stated requirement of unusual fidelity, and nothing engineered it into funded, staffed, rehearsed capability before the event; the response failures are downstream expressions of the unconverted requirement, not independent causes. LENS uses it in Domain 1 (Systems Analysis) for the discipline of tracing a documented requirement to the built capability that does or does not exist behind it, with the pre/post disaster lifecycle of this part as the organizing frame. Pair with the prevention-thinning cases earlier in the chapter and with the PKEMRA reform arc for the institution-building sequel.

| DOMAIN                                                                             | FAILURE MODES             | LENS COMPETENCY                                                       | PROBLEM TYPE | INDUCED | LEO |
|------------------------------------------------------------------------------------|---------------------------|-----------------------------------------------------------------------|--------------|---------|-----|
| <div>EMERGENCY MANAGEMENT</div> <div>DISASTER RESPONSE</div> <div>GOVERNMENT</div> | <div>G</div> <div>K</div> | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 1            | 1.1     | 1   |
| Systems Analysis                                                                   |                           |                                                                       |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat an exercise finding as an open capability requirement with a named owner, a funding line, and a completion date; Hurricane Pam’s draft plans and follow-on workshop schedule specified the requirement, and the case’s central observation is that specification without a funded owner produced no capability in thirteen months.
- ▶ Verify the catastrophic-response machinery by operating it before the catastrophe; the Incident of National Significance designation and the push-resources-without-request annex had never been exercised end-to-end in anger, and their first operational use came after landfall of the event they were designed for.
- ▶ Build the common operating picture as an engineered deliverable, not an aspiration; the test is concrete — the operations center should know what broadcast television knows, and a response organization that fails that test has a situational-awareness capability gap, not a communications inconvenience.

IN OPERATION / AFTER

- ▶ Convert investigation findings into statutory obligation where institutional memory alone will not hold; PKEMRA’s corrective-action and national-preparedness mandates are the codified form of the discipline Pam lacked, and the case teaches the statute as the institutional response, not as the proof the gap is closed.
- ▶ Preserve the contested death-toll attribution in any teaching use; the range and its methodological instability are themselves evidence about how poorly instrumented the response was, and smoothing to a single number discards that signal.
- ▶ Pair the case in syllabi with the prevention-side cases in this part so the pre/post lifecycle reads as one arc: a prevention regime that thinned, a rehearsed requirement that went unbuilt, and a response phase that discovered the unbuilt capability in production.

## REFLECTION QUESTIONS

1. Identify an exercise, red-team, or audit finding in your organization that specified a missing capability and was closed as “documented” rather than “built.” What would a funded owner, a completion date, and a re-verification event have looked like, and what converts a finding into an obligation in your setting?
2. The Incident of National Significance machinery saw its first operational use during the event it was designed for. Pick a rarely-invoked escalation mechanism in your domain and specify the rehearsal that would verify it end-to-end — including the authority handoffs — before it is needed in production.
3. The Convention Center was visible to television audiences before it entered the federal operations picture. What is the equivalent external signal in your domain that would outrun your operations picture, and what engineered channel would make your institution aware of what the public already knows?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines

# Fukushima Daiichi

**N** Normalization of Deviance   **G** Governance & Trust   **K** Knowledge & Institutional Memory

**IMPACT** Three reactor meltdowns following the Tōhoku earthquake and tsunami; approximately 150,000 people evacuated (NAIIC); cleanup and compensation put at ¥21.5 trillion (about \$190B) by the Japanese government in 2016, with independent estimates running higher

## IN BRIEF

On 11 March 2011 the Tōhoku earthquake and the tsunami it spawned overwhelmed TEPCO's Fukushima Daiichi plant: the quake severed the off-site transmission lines, the wave rose above the site's ground level and flooded the emergency diesel generators, and cooling failed. Three of the six cores melted down and hydrogen explosions spread radioactive material. NAIIC records that approximately 150,000 people were evacuated in response to the accident; the Japanese government put cleanup and compensation at ¥21.5 trillion — about \$190 billion — in 2016, and independent estimates run higher. The independent Diet commission (NAIIC), chaired by Kiyoshi Kurokawa, called it a disaster “made in Japan” — the product of regulatory capture and a culture reluctant to challenge utility assumptions; evidence of large historical tsunamis had been discussed internally but never forced a change to the seawall. Other reviews stressed under-estimated external hazards. Fukushima is the post-TMI evidence that safety institutions like INPO must be deliberately built, not assumed.

## BACKGROUND

Fukushima Daiichi sat on Japan's northeast coast, its reactors and their backup diesel generators protected by the site's ground elevation and a design basis set by the Japan Society of Civil Engineers method. Evidence of much larger historical waves — back to the ninth-century Jōgan event — had been discussed in TEPCO's own internal assessments and shared with the regulator from 2006, without ever becoming a required change.<sup>9</sup> The gap was not in the data but in the conversion: NAIIC found TEPCO and NISA both knew a wave above site level meant total loss of power, and that NISA gave no compulsory instruction and TEPCO took no action.

## WHAT HAPPENED

On 11 March 2011 the Tōhoku earthquake struck and the tsunami that followed rose above the site's ground level. The earthquake had already severed the off-site transmission lines; the diesel generators meant to keep cooling running were inundated. Cooling failed, three of the six reactor cores melted down, and hydrogen explosions spread radioactive material across the region. Approximately 150,000 people were evacuated in response to the accident; the Japanese government's 2016 estimate put cleanup and compensation at ¥21.5 trillion, about \$190 billion, and independent estimates run higher.<sup>1</sup> Siting the backup generators where a

wave above the site's ground level would reach them tied the entire cooling chain to that single design assumption, so once the water was on the site the loss of cooling followed almost mechanically from the layout itself.

#### THE INVESTIGATION

The independent investigation chaired by Kiyoshi Kurokawa for the National Diet (NAIIC) called the accident "made in Japan" — the product of regulatory capture, deference to authority, and an institutional reluctance to challenge utility assumptions.<sup>2</sup> Other major reviews — the Hatamura government commission (2012) and the IAEA Director General's report (2015) — emphasized the under-estimation of external hazards and defense-in-depth assumptions over the cultural critique; the Kurokawa framing is the most-cited but not the only consensus reading.<sup>3</sup> NAIIC also refused to accept the tsunami as the sole direct cause, holding that the earthquake may have damaged safety equipment at Unit 1; the IAEA report found the on-site destruction tsunami-driven. The book treats Kurokawa's as the most-cited reading, not the settled one.

#### THE CAPABILITY GAP

Fukushima is the post-TMI case showing that the INPO pattern (Case 175) is not self-executing. The U.S. industry built INPO to force operating discipline and shared learning; the Japanese industry did not build an equivalent with the independence to override a utility's optimistic assumptions. The internal tsunami evidence existed; the institutional capability to act on it did not.<sup>4</sup> Evidence without an independent body empowered to act on it is inert: the larger-wave assessments could be acknowledged and shelved indefinitely because the regulator that did stand outside the utility, NISA, knew the risk and issued no compulsory instruction to act on it.

#### AFTERMATH & REFORM

The accident drove a restructuring of Japanese nuclear regulation — a new, more independent Nuclear Regulation Authority — and a global re-examination of external-hazard margins and station-blackout protection. Paired with INPO and Davis-Besse, it triangulates the book's claim: sustained nuclear-safety capability is an institution that must be deliberately built and rebuilt, not a property a competent industry simply has.<sup>5</sup> Creating a more independent regulator after the fact conceded the structural diagnosis directly: the missing piece had been an authority sitting outside the utility, and the reform was precisely to build the independence that the pre-2011 arrangement had lacked.

## 3 of 6

### REACTORS MELTED DOWN

*tsunami exceeded design basis the institutional evidence had already questioned*

FUKUSHIMA DAIICHI — "MADE IN JAPAN," PER THE DIET INQUIRY

*What must be admitted – very painfully – is that this was a disaster "Made in Japan."*

NATIONAL DIET OF JAPAN FUKUSHIMA NUCLEAR ACCIDENT INDEPENDENT INVESTIGATION COMMISSION, 2012

REFERENCES

1. National Diet of Japan Fukushima Nuclear Accident Independent Investigation Commission (NAIIC; K. Kurokawa, chair), Report (2012) – the internal tsunami evidence and the regulatory-capture finding.

2. NAIIC (2012) – the accident sequence: seawall overtopping, generator inundation, three core meltdowns, and “approximately 150,000 people were evacuated in response to the accident.” The widely quoted 160,000 is the 2012 peak evacuee count in prefectural records, not NAIIC’s figure; do not “correct” the text against it. Cleanup and compensation costs are from the Japanese government’s December 2016 estimate of ¥21.5 trillion, not from NAIIC.

3. NAIIC (2012) – the “made in Japan” cultural and regulatory-capture conclusion (quoted).

4. Investigation Committee on the Accident (Hatamura government commission, 2012); IAEA Director General, The Fukushima Daiichi Accident (2015) – external-hazard under-estimation.

5. D. Lochbaum, E. Lyman & S. Stranahan, Fukushima: The Story of a Nuclear Disaster (2014).

6. Y. Funabashi & K. Kitazawa, Fukushima in Review (2012); cf. INPO (Case 175) and Davis-Besse (Case 160).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Fukushima is the post-TMI case that establishes that the INPO pattern (Case 175) is not self-executing. The U.S. industry built INPO; the Japanese industry did not. The cost of the difference, paid in 2011, is the strongest available evidence that capability institutions must be deliberately built, not assumed.

LENS APPROACH

LENS uses Fukushima in LEN 8 as the cross-cultural comparison to INPO and in LEN 7 for the regulator-utility dynamics study. The case is paired with INPO and Davis-Besse to triangulate what sustained nuclear-safety capability requires.

| DOMAIN                                | FAILURE MODES | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|---------------|------------------|--------------|---------|-----|
| ENERGY                                | → N G K       | → D1 D2 D3 D4 D5 | 4            | 7.4     | 5   |
| Navigating Sociotechnical Constraints |               |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Build an institutional path that converts an acknowledged hazard (like the larger-wave evidence) into a binding design requirement, not a discussable assessment.
- ▶ Site critical backups so a single overtopped barrier cannot disable the whole cooling chain at once.
- ▶ Establish an oversight body independent of the utility, with standing to override optimistic in-house assumptions about external hazards.

IN OPERATION / AFTER

- ▶ Audit whether internally held hazard evidence has actually driven design changes, treating shelved assessments as an open finding.
- ▶ Re-examine external-hazard margins on a cadence rather than freezing them at a design-basis figure set once.
- ▶ Verify that the regulator retains the independence to act, since institutional safety capability erodes if it is assumed rather than rebuilt.



## REFLECTION QUESTIONS

1. Identify a regulatory regime in your domain whose effectiveness depends on a cultural willingness to challenge authority. What if the culture changes?
2. INPO is U.S.-specific. Design the structural features that would have to be in place for a comparable institution to function in a different national context.
3. The tsunami evidence was discussed but never converted into a binding requirement. What institutional path in your domain turns an acknowledged hazard into a mandated change rather than a shelved assessment?

**WHO BUILDS THIS** safety science & organizational analysis · governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** incident analysis · safety dashboards · governance frameworks · bias & evidence audits · knowledge bases · data pipelines

# West Africa Ebola – The Delayed International Response

**G** Governance & Trust    **N** Normalization of Deviance

**IMPACT** The 2014–2016 West Africa Ebola epidemic reached approximately 28,600 reported cases and 11,300 reported deaths (WHO situation-report figures, undercount acknowledged) across Guinea, Liberia, and Sierra Leone; first cases December 2013, outbreak confirmed March 2014, MSF called it "out of control" on 21 June 2014 while WHO staff in Guinea and Sierra Leone accused it of causing unnecessary panic, and WHO declared the Public Health Emergency of International Concern only on 8 August 2014 — months into exponential growth

## IN BRIEF

The 2014–2016 West Africa Ebola epidemic is the load-bearing case for the gap between declared international readiness and deployable international capability. The first cases appeared in Guinea in December 2013; the outbreak was confirmed as Ebola in March 2014. Médecins Sans Frontières declared the epidemic "out of control" in June 2014, while WHO staff in Guinea and Sierra Leone called that alarmist. WHO declared the Public Health Emergency of International Concern only on 8 August 2014 — months into exponential growth, and by the third PHEIC ever issued. The response then scaled: UNMEER, the first UN emergency health mission, and U.S. and UK military logistics; and community-led change — safe burials and community care-center networks — bent the curve. The reported toll reached approximately 28,600 cases and 11,300 deaths; the undercount is acknowledged and the figures are reported figures, not a census. Three converging post-mortems — the WHO Ebola Interim Assessment Panel (Stocking Report, July 2015), the Harvard-LSHTM Independent Panel (Moon et al., The Lancet, November 2015), and the UN High-Level Panel (2016) — found the same structural failures: WHO's outbreak-response capacity had been cut after the 2011 budget reductions; most member states had simply not built the IHR (2005) core capacities, and the Panel called those results "not acceptable" and the reliance on voluntary self-assessment to measure them "unacceptable"; and early-warning signals were discounted for political and economic reasons. The reform — the WHO Health Emergencies Programme (2016) — carries the panels' own caveat that post-crisis reform pledges have historically decayed. COVID-19 later tested it.

## BACKGROUND

The epidemiological timeline is the spine of the case. The first cases trace to Guinea in December 2013; the outbreak was laboratory-confirmed as Ebola virus disease in March 2014. In April, Guinea's health ministry and WHO reporting

characterized the outbreak as slowing toward containment. By June, the case count in Guinea, Liberia, and Sierra Leone was climbing, and Médecins Sans Frontières stated publicly that the epidemic was “out of control” and beyond the capacity of the responding organizations — a framing that WHO staff in the affected countries called alarmist. WHO convened its Emergency Committee on 6–7 August 2014 and declared a Public Health Emergency of International Concern on 8 August 2014, by which point roughly 1,700 cases and 930 deaths had been reported. The declaration came months into exponential growth, not ahead of it. It was only the third PHEIC WHO had ever issued.<sup>9</sup>

#### WHAT HAPPENED

The analytical spine is declared capability versus engineered capability. Under the International Health Regulations (2005), member states were required to build and certify a set of core capacities for detecting, assessing, and responding to public-health events. On the certification record, most had not built them. As at November 2014, 64 States Parties told the Secretariat they had achieved the core capacities, 81 requested extensions, and 48 did not communicate their status or intentions. The Stocking Panel called those results “not acceptable,” and said it considered “it unacceptable that only voluntary self-assessment is required for measurement of core capacities” — information that, “when the health of all is at stake,” must be validated by peer review or other external assessment. A legally binding obligation was being measured by the obligated party’s own unverified account of itself. The Stocking Report named the IHR implementation gap directly as one of its three review areas, alongside WHO’s own emergency-response capacity and its role within the wider humanitarian system. The gap itself — obligated readiness minus built readiness — is the artifact the case asks the reader to own.<sup>1</sup>

#### THE INVESTIGATION

The escalation machinery failed as governance, not merely as forecasting. The early signals were present and were discounted: warnings of uncontrolled spread were treated as alarmist, and calls for a larger response were slowed by political and economic considerations, including the fear that a PHEIC declaration would damage trade and travel in the affected economies. This is normalization of deviance in the escalation loop — each week that the exponential signal was reframed as manageable lowered the threshold for reframing it again the next week. The Harvard-LSHTM Independent Panel (Moon et al., *The Lancet*, November 2015) found the Director-General had ample reason to declare the emergency earlier, and its ten reform proposals targeted the governance of outbreak escalation — the machinery by which a discounted signal is supposed to become a declared emergency — rather than the science of detection. WHO’s outbreak-response capacity had itself been reduced after the 2011 budget cuts, so the institution charged with escalating had been hollowed at the point of need.<sup>2</sup>

#### THE CAPABILITY GAP

The response, once it scaled, showed what deployable capability looks like. The UN established UNMEER, its first dedicated emergency health mission; the United States and United Kingdom supplied military logistics, engineering, and treatment-unit construction. But the curve did not bend on logistics alone. Community-led change carried decisive weight: safe and dignified burial practices that interrupted a major transmission route, and community care-center networks that moved isolation and care closer to where transmission was

occurring. The reported toll settled at approximately 28,600 cases and 11,300 deaths across the three countries to mid-2016. The figures are reported figures and the undercount is acknowledged in the source record — the true toll is higher by an amount that the surveillance system, the same system whose inadequacy is the case's subject, could not measure. That hedge is load-bearing and travels with every use of the numbers.<sup>3</sup>

#### AFTERMATH & REFORM

The reform response and its own caveat close the case. WHO established the Health Emergencies Programme in 2016, restructuring its outbreak-and-emergency function; the UN High-Level Panel on the Global Response to Health Crises (2016) issued twenty-seven recommendations for national, regional, and international action, including political-leadership structures above the technical agencies; and all three panels pressed for genuine — not certified — IHR core-capacity building. Each panel attached the same caveat: that reform pledges made in the wake of a health crisis have historically decayed as the crisis recedes and attention and funding move on. This is the case's binding hedge on its own resolution — the reform is real but its durability was, at the time of writing, unestablished. COVID-19 later tested the reformed system, and the question the panels raised — whether the escalation machinery had been engineered rather than re-certified — is the question the reader is asked to carry forward, not the question the case closes.

*The Director-General had ample reason to raise international attention by convening the Emergency Committee or declaring a public health emergency of international concern earlier.*

HARVARD-LSHTM INDEPENDENT PANEL ON THE GLOBAL RESPONSE TO EBOLA (MOON ET AL., THE LANCET, 28 NOVEMBER 2015).

#### REFERENCES

1. Médecins Sans Frontières (2015), Ebola: Pushed to the Limit and Beyond — MSF account of the international outbreak response, March 2015: the "out of control" declaration of 21 June 2014 and the reaction of WHO staff in Guinea and Sierra Leone. WHO declared the PHEIC on 8 August 2014, the Emergency Committee having met 6–7 August; as of 4 August, 1,711 cases and 932 deaths had been reported.
2. Ebola Interim Assessment Panel (2015), Report of the Ebola Interim Assessment Panel (the "Stocking Report," chair Dame Barbara Stocking), World Health Organization, 7 July 2015 — the November 2014 IHR core-capacity figures (64 / 81 / 48), the "not acceptable" and "unacceptable" findings (quoted), and the Panel's three review headings.
3. Moon, S., Sridhar, D., Pate, M. A., Jha, A. K., Clinton, C., et al. (2015), "Will Ebola change the game? Ten essential reforms before the next pandemic. The report of the Harvard-LSHTM Independent Panel on the Global Response to Ebola," *The Lancet* 386(10009):2204–2221, 28 November 2015, doi:10.1016/S0140-6736(15)00946-0 — the "ample reason" finding (quoted) and the ten reform proposals.
4. World Health Organization (2016), Ebola Situation Reports — West Africa: approximately 28,600 cases and 11,300 deaths reported to 10 June 2016 (reported figures; undercount acknowledged).
5. High-level Panel on the Global Response to Health Crises (2016), Protecting Humanity from Future Health Crises (chair Jakaya M. Kikwete), UN doc. A/70/723, 9 February 2016 — twenty-seven recommendations, including political-leadership structures above the technical agencies.

#### THE LEARNING ENGINEERING LENS

##### LE INSIGHT

West Africa Ebola is the reference case for declared international readiness that was certified but not deployable. Most member states had not built the IHR (2005) core capacities and said so, and the

only measurement of that obligation was voluntary self-assessment, which the Stocking Panel called unacceptable; WHO’s own outbreak-response capacity had been cut after 2011; and the escalation loop discounted an exponential signal for months before the PHEIC of 8 August 2014. The reported toll carries an acknowledged undercount, and the reform carries the panels’ own caveat that post-crisis pledges decay.

LENS APPROACH

West Africa Ebola is the declared-versus-engineered- capability case at the scale of the international system (induced 1.1; LENS D5/PT6; LEO-5). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the IHR obligation- versus-built-capability seam and for the governance failure in the escalation machinery — the discounting of the escalating signal as normalization of deviance in the declaration loop. The three converging panels (Stocking, Moon et al., UN High-Level) are carried as the investigation evidence base; the acknowledged undercount and the reform- decay caveat are decisive hedges that travel with the case.

| DOMAIN                   | FAILURE MODES | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO |
|--------------------------|---------------|---------------------------------------|--------------|---------|-----|
| GLOBAL HEALTH            | G N           | D1 D2 D3 D4 D5                        | 6            | 1.1     | 5   |
| OUTBREAK RESPONSE        | →             | Navigating Sociotechnical Constraints |              |         |     |
| INTERNATIONAL GOVERNANCE |               |                                       |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the gap between obligated and built capability as the primary operational artifact: most states had not built the IHR core capacities and reported as much, and the analyst’s job during escalation is to measure the readiness the system can produce, not the readiness it is obliged to hold.
- ▶ Build the escalation loop to resist discounting an exponential signal for political or economic reasons; the recurring failure was that each week’s reframing of the signal as manageable lowered the threshold for reframing it again, and the declaration lagged the growth it existed to trigger.
- ▶ Do not hollow the institution charged with escalation ahead of the emergency; WHO’s outbreak-response capacity had been reduced after the 2011 budget cuts, so the body responsible for sounding the alarm had been degraded at the point of need — a sustainment failure in the readiness deliverable itself.

IN OPERATION / AFTER

- ▶ Preserve the acknowledged-undercount hedge with the reported figures; the 28,600 cases and 11,300 deaths are reported figures produced by the same surveillance system whose inadequacy is the case’s subject, and the true toll is higher by an unmeasured amount.
- ▶ Carry the panels’ own reform-decay caveat into any teaching of the resolution; the WHO Health Emergencies Programme (2016) is real, but the panels flagged that post-crisis reform pledges historically decay, and the case’s honesty depends on holding the reform and its caveat together.
- ▶ Use COVID-19 as the deferred test rather than the closing chapter; the case asks whether the escalation machinery was engineered or merely re-certified, and the later pandemic is the evidence event the reader is directed toward, not a resolution the case supplies.

## REFLECTION QUESTIONS

1. Identify a capability in your domain that is certified present but may not be deployable — a readiness that exists on the compliance record but has never been exercised at full load. What would it take to measure the gap between the declared readiness and the readiness the system could actually produce, and who owns that gap today?
2. The escalation loop discounted an exponential signal for months because raising the alarm carried political and economic cost. Specify the decision rule you would build so that the threshold for reframing a growing signal as “manageable” does not fall each time it is reframed. What makes an escalation trigger resistant to the incentive not to pull it?
3. The reform response carries the panels’ own caveat that post-crisis pledges historically decay. Pick a reform your institution adopted after a crisis, and ask: was the underlying capability re-engineered, or was it re-certified? What evidence — short of the next crisis — would tell you which?

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Hurricane Maria — Puerto Rico Response and Logistics Failure

**G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** Hurricane Maria struck Puerto Rico September 20, 2017, the third major hurricane to strike U.S. territory in under a month after Harvey and Irma; FEMA's own 2017 Hurricane Season After-Action Report (July 12, 2018) found the agency entered Maria with force strength below target, its Caribbean Distribution Center inventory largely depleted by the Irma response, and commodity-tracking visibility it could not sustain to the ports; GAO-18-472 found 54% of FEMA staff serving in a capacity in which they did not hold the "Qualified" title at the October 2017 peak; the grid collapse became the longest blackout in U.S. history (~11 months to full restoration); the commissioned GWU Milken Institute study estimated ~2,975 excess deaths (Sept 2017 – Feb 2018), adopted as the official toll in place of the initial 64

## IN BRIEF

Hurricane Maria made landfall on Puerto Rico on September 20, 2017, as a high-end Category 4 storm, the third major hurricane to strike U.S. territory in under a month. The decisive evidence on the federal response is the responding institution's own self-assessment: FEMA's 2017 Hurricane Season After-Action Report, released July 12, 2018, found that the agency entered Maria with its commodity stockpiles largely drained by Harvey and Irma — the Caribbean Distribution Center in Puerto Rico had distributed more than 80 percent of its inventory of some commodities for Irma and had not been replenished — with force strength below its staffing target, and with logistics visibility so poor it could not comprehensively track commodities moving to Puerto Rico and the U.S. Virgin Islands. GAO-18-472 documented that at the October 2017 deployment peak, 54 percent of deployed staff were serving in a capacity in which they did not hold the "Qualified" title. The island's grid collapse became the longest blackout in U.S. history, with full restoration taking approximately eleven months. Excess-mortality estimates diverge by method and are carried here as estimates: the commissioned George Washington University Milken Institute study estimated approximately 2,975 excess deaths (September 2017 – February 2018), adopted as the official toll in place of the initial count of 64; the Kishore et al. household survey in the New England Journal of Medicine estimated 4,645 with a 95 percent confidence interval of 793 to 8,498 (September 20 – December 31, 2017). The capability gap the case anchors is a logistics and workforce capability designed to a mainland-hurricane envelope and never re-verified for an island scenario with a destroyed grid and ports.

## BACKGROUND

Hurricane Maria struck Puerto Rico on September 20, 2017, two weeks after Hurricane Irma had passed north of the island and less than a month after Hurricane Harvey had drawn down federal response capacity on the Gulf Coast. Maria destroyed the island's electrical grid outright: the resulting outage became the longest blackout in U.S. history, with power not restored to all customers until August 2018 — approximately eleven months — and an estimated 3.4 billion lost customer-hours of electricity service, roughly three times the next-largest blackout in the U.S. record. The operational context was categorically different from the mainland-hurricane scenario the federal response apparatus was sized against: an island 1,000 miles from the mainland, reachable only by sea and air, with ports damaged, the grid destroyed, and no adjacent states from which mutual-aid convoys could drive in.<sup>9</sup>

## WHAT HAPPENED

FEMA's own 2017 Hurricane Season After-Action Report, released July 12, 2018, is the decisive self-assessment and the reason the case carries the program-report evidence tag. The report found that the agency entered the season with force strength below its staffing target, and that the sequential demands of Harvey and Irma had largely drained the national commodity posture before Maria arrived: when Irma struck, FEMA opened the Caribbean Distribution Center — its warehouse of disaster supplies in Puerto Rico — and within days had distributed more than 80 percent of its inventory of some commodities; Maria struck before the supplies were replenished. The report further found that FEMA could not comprehensively track commodities moving across multiple modes of transportation to Puerto Rico and the U.S. Virgin Islands, due to staffing shortages and business-process shortfalls — the agency lost visibility of its own relief supplies at the ports. A 2011 exercise after-action report for Puerto Rico had anticipated that the territory would require extensive federal support in moving commodities from the mainland and to distribution points across the island; the anticipated requirement had not been converted into verified capability by 2017.<sup>1</sup>

## THE INVESTIGATION

GAO-18-472 (September 2018) documented the workforce side of the capability gap. At the height of the 2017 deployments in October — close to 14,000 federal employees deployed across Harvey, Irma, Maria, and the California wildfires — 54 percent of FEMA staff were serving in a capacity in which they did not hold the title of "Qualified" under FEMA's own qualification system standards. The workforce was overwhelmed by the concurrent demand, and the shortage of trained personnel with program expertise complicated the response, particularly after Maria, where the staffing shortfall landed on the disaster with the greatest logistics difficulty and the least pre-existing territorial response capacity. The GAO finding and the FEMA after-action "Force Strength" finding describe the same structural condition from outside and inside the agency: a workforce sized and qualified for the routine disaster-response envelope, deployed into a season three standard deviations outside it.<sup>2</sup>

## THE CAPABILITY GAP

The mortality evidence is carried as estimates, with their methods, and the divergence is not smoothed. The government of Puerto Rico's initial official toll was 64. The Kishore et al. study, published in the *New England Journal of Medicine* in May 2018, surveyed households across the island and estimated 4,645 excess deaths from September 20 through December 31,



2017, with a 95 percent confidence interval of 793 to 8,498 — a 62 percent increase in the mortality rate over the same period in 2016, and an interval whose width is itself a finding about what a survey design can resolve.<sup>3</sup> The commissioned George Washington University Milken Institute study, released in August 2018, used death-registry data and demographic modeling that accounted for post-storm population displacement, and estimated approximately 2,975 excess deaths from September 2017 through February 2018. The government of Puerto Rico adopted 2,975 as the official toll, revising the count from 64. The two studies used different methods, different windows, and different uncertainty structures; they agree that the initial official count understated the mortality by more than an order of magnitude, and they do not agree on a single number.<sup>4</sup>

#### AFTERMATH & REFORM

The hedges the case carries are not decoration. The excess-mortality figures are model-dependent estimates, not counts; the GWU and NEJM designs are not directly comparable and neither supports an attribution chain from any individual logistics failure to any individual death. What the record does support is the capability-envelope finding, and it is the responding institution's own: the logistics and workforce capability had been specified and verified against a mainland-hurricane scenario — adjacent states, intact ports, a grid that fails partially — and was never re-verified for the island scenario that the 2011 exercise report had already described. The after-action report itself is the intervention-shaped element inside the failure: an institution measuring its own capability gap honestly, in public, with its own name on the findings. The systems-analysis discipline the case anchors is the one FEMA's report performs after the fact and the one the deployment record shows was not performed before it — deriving the requirement from the operational reality of the deployed context rather than from the envelope the capability was built and staffed to.

*FEMA entered Maria with force strength below target and its Caribbean Distribution Center largely emptied by the Irma response; it could not comprehensively track commodities to the ports; the grid collapse became the longest blackout in U.S. history; the official toll was revised from 64 to 2,975 on the strength of the commissioned GWU excess-mortality study.*

EDITORS' SYNTHESIS OF THE FEMA 2017 HURRICANE SEASON AFTER-ACTION REPORT (2018), GAO-18-472 (2018), THE GWU MILKEN INSTITUTE STUDY (2018), AND KISHORE ET AL. (2018, NEJM).

#### REFERENCES

1. Houser, T., & Marsters, P. (2018), "The World's Second Largest Blackout," Rhodium Group, April 12, 2018 — approximately 3.4 billion lost customer-hours, "three times larger than the next biggest blackout in US history"; the largest blackout in U.S. history by that metric.
2. Federal Emergency Management Agency (2018), 2017 Hurricane Season FEMA After-Action Report, July 12, 2018 — the agency's own self-assessment: force strength below target, Caribbean Distribution Center depletion, commodity-tracking shortfalls.
3. U.S. Government Accountability Office (2018), 2017 Hurricanes and Wildfires: Initial Observations on the Federal Response and Key Recovery Challenges, GAO-18-472, September 4, 2018 — "at the height of FEMA workforce deployments in October 2017, 54 percent of staff were serving in a capacity in which they did not hold the title of 'Qualified'; close to 14,000 federal employees deployed.
4. Kishore, N., Marqués, D., Mahmud, A., et al. (2018), "Mortality in Puerto Rico after Hurricane Maria," *New England Journal of Medicine* 379(2):162 – 170, doi:10.1056/NEJMs1803972 — household-survey estimate of 4,645 excess deaths (95% CI 793 – 8,498), September 20 – December 31, 2017.
5. Milken Institute School of Public Health, George Washington University (2018), Ascertainment of the Estimated Excess Mortality from Hurricane Maria in Puerto Rico, August 2018 — commissioned by the Government of Puerto Rico; 2,975 excess deaths, September 2017 – February 2018; basis for the official toll revision from 64.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Hurricane Maria is the reference case for a capability specified and verified against one operational envelope — the mainland hurricane, with adjacent states, intact ports, and a partially failing grid — and deployed into another. The stockpile depletion, the workforce qualification gap, and the lost logistics visibility were not separate failures; they were one requirements failure surfacing in three subsystems. The after-action report is the case's second lesson: the institution measured its own capability gap honestly, and that honesty is what makes the gap analyzable.

LENS APPROACH

Hurricane Maria is the engineered-versus-stated-readiness case at national-response scale (induced 1.1; LENS D1/PT1; LEO-1). LENS uses it in Domain 1 (Systems Analysis) for the discipline of deriving requirements from the deployed operational context — island logistics, destroyed grid and ports, sequential-storm depletion — rather than from the envelope the capability was built to, and for the after-action report as the requirements artifact produced after the fact. Pair with Case 124 (USS Fitzgerald and McCain) for the stated-versus-real-readiness thread and Case 126 (F-35 sustainment) for capability requirements as a sustainment deliverable. The excess-mortality estimates are carried as estimates with their methods; the divergence between the GWU and NEJM figures is decisive and is not smoothed.

| DOMAIN                                                                                 | FAILURE MODES             | LENS COMPETENCY                                                                                   | PROBLEM TYPE | INDUCED | LEO |
|----------------------------------------------------------------------------------------|---------------------------|---------------------------------------------------------------------------------------------------|--------------|---------|-----|
| <div>EMERGENCY MANAGEMENT</div> <div>DISASTER LOGISTICS</div> <div>PUBLIC HEALTH</div> | <div>G</div> <div>N</div> | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> <div>Systems Analysis</div> | 1            | 1.1     | 1   |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the operational envelope from the deployed context, not the prior platform: an island response with destroyed ports and grid is a different requirement than a mainland response with adjacent-state mutual aid, and the 2011 exercise report that anticipated the difference was the requirements artifact that needed an owner.
- Treat stockpile posture and force strength as live capability measurements under sequential-event depletion; the Caribbean Distribution Center's inventory state after Irma was a known quantity before Maria's landfall, and the replenishment gap was a decision surface, not a surprise.
- Instrument logistics visibility as a first-class capability, not a reporting afterthought; an agency that cannot track its own commodities at the ports cannot re-plan distribution under the conditions where re-planning matters most.

IN OPERATION / AFTER

- Write the after-action report the way FEMA wrote this one — the institution's own name on the capability-gap findings, in public; the report's honesty is what makes the case teachable and is the standard the casebook holds after-action self-assessment to.
- Carry the excess-mortality estimates with their methods and their divergence; the GWU registry-based 2,975 and the NEJM survey-based 4,645 (95% CI 793 – 8,498) answer differently designed questions, and revising the official toll from 64 to 2,975 is the decisive correction, not the reconciliation of the two studies.
- Convert the envelope finding into verification events: re-verify inherited response designs against the island scenario explicitly — pre-positioned stock floors that survive a sequential season, workforce qualification targets audited against deployment reality, and port-condition logistics exercised rather than assumed.

## REFLECTION QUESTIONS

1. Identify a capability in your domain that was specified and verified against one operational envelope and is routinely deployed near or beyond it. What is the equivalent of the island scenario — the context that breaks the envelope's quiet assumptions — and what verification event would surface the gap before deployment does?
2. FEMA's after-action report named its own force-strength, stockpile, and visibility gaps in public. Specify what an equivalently honest self-assessment would look like for your institution: who signs it, what it is permitted to find, and what pre-committed response follows a finding that the capability was below requirement.
3. The two excess-mortality studies estimate different quantities with different methods and do not agree on a number, and both were necessary to displace the official count of 64. When evidence streams in your domain diverge, what is the discipline for carrying both — with methods attached — rather than picking one or averaging them?

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# CrowdStrike Falcon Outage

**D** Designed Out   **K** Knowledge & Institutional Memory   **G** Governance & Trust

**IMPACT** 8.5 million Windows machines crashed; airlines, hospitals, broadcasters, and banks affected simultaneously; largest IT outage on record

## IN BRIEF

On 19 July 2024 CrowdStrike pushed a content update to its Falcon endpoint sensor that contained a logic error in a small configuration file. On every Windows machine running the affected sensor, the file drove the kernel driver to crash, looping the blue screen of death and requiring hands-on recovery of each device. Roughly 8.5 million machines went down at once — across hospitals, airlines, banks, broadcasters, and governments worldwide — the largest IT outage on record. CrowdStrike's post-incident review found the content update had not been put through the same automated testing or staged rollout as code: the pipeline treated "content" as a lesser category than "code," though the operating system did not. CrowdStrike is the cybersecurity-vendor analog of Knight Capital, six orders of magnitude larger.

## BACKGROUND

CrowdStrike's Falcon sensor runs deep inside Windows — in the kernel — to detect threats, and it updates constantly: not only its code, but "content," the rapid-response detection configuration pushed to customers continuously. The deployment pipeline treated that content as a lighter category than code, with less testing and no staged rollout.<sup>0</sup> The distinction had a logic: content shipped fast and often, precisely so the sensor could keep pace with new threats, and slowing it down with full code-grade testing seemed to defeat its purpose. That speed was the very reason the safety gate was lowered on the artifact that could still crash a kernel.

## WHAT HAPPENED

On 19 July 2024 a content update contained a logic error in a small configuration file. On every Windows machine running the affected sensor, the file drove the kernel driver to read out of bounds and crash, looping the blue screen of death and requiring manual recovery of each device. Roughly 8.5 million machines failed at once — hospitals, airlines, banks, broadcasters, and governments worldwide — the largest IT outage on record.<sup>1</sup> Because the file went to every affected sensor simultaneously with no staged rollout, there was no first wave to catch the fault and no blast radius short of the whole install base; the requirement for hands-on recovery of each device turned a single bad push into weeks of physical labor across the world.

## THE INVESTIGATION

CrowdStrike's own post-incident review found the content update had not been put through the same depth of testing and staged rollout as its code releases. The fault was not exotic: a category boundary in the deployment pipeline — content treated as safer than code — that did not match the operational reality, in which a bad content file could crash the kernel exactly as bad code could. The boundary was an organizational convenience rather than a technical truth, and the kernel, which executes whatever reaches it, recognized no such distinction at all.<sup>2</sup>

## THE CAPABILITY GAP

The missing capability was the recognition that, for deployment safety, content **is** code: anything that can crash the kernel must clear the same testing and staged-rollout gates. CrowdStrike's customers had trusted the vendor's deployment safety, and that trust turned out to be decisive for the operation of a large slice of the global economy on a single morning. Each customer had implicitly outsourced a safety gate to the vendor's pipeline, so the one missing gate inside that pipeline was multiplied across every institution that ran the sensor, with no independent check standing between a bad push and their kernels.<sup>3</sup>

## AFTERMATH & REFORM

CrowdStrike moved content updates onto staged rollouts and stronger validation, Microsoft revisited kernel-level access for security vendors, and the episode prompted scrutiny of concentration risk in endpoint security.<sup>4</sup> The fallout also moved into court: Delta Air Lines, which put its losses near \$500 million, and CrowdStrike filed competing suits on the same day in October 2024, and on 16 May 2025 Judge Kelly Lee Ellerbe of the Fulton County Superior Court let Delta's gross-negligence and computer-trespass claims proceed while dismissing the fraud claims resting on statements made before June 2022 — a narrower fraud claim, that CrowdStrike had promised not to install an unauthorised back door, survived. A separate passenger class action against Delta over the cancellations was allowed to advance in part the same month.<sup>5</sup> Each engineering response targets a different layer of the same failure: staged rollout limits the blast radius of any one push, reconsidering kernel access limits how much a vendor fault can break, and the concentration-risk scrutiny acknowledges that a single vendor had become a shared point of failure for much of the economy. It is the cybersecurity-vendor analog of Knight Capital (Case 143) — an unverified deployment to a system wired into critical operations — at a scale six orders of magnitude larger.

# 8.5M

MACHINES · SINGLE CONFIGURATION FILE

*content treated differently from code in the deployment pipeline*

*Implement a staggered deployment strategy for Rapid Response Content in which updates are gradually deployed to larger portions of the sensor base, starting with a canary deployment.*

CROWDSTRIKE, FALCON CONTENT UPDATE PRELIMINARY POST INCIDENT REPORT, 24 JULY 2024

REFERENCES

1. CrowdStrike, Falcon Content Update Preliminary Post Incident Report (24 July 2024) — Rapid Response Content deploys as configuration data through Channel Files without the automated testing, manual validation and staged rollout that Sensor Content clears.

2. D. Weston, “Helping our customers through the CrowdStrike outage,” The Official Microsoft Blog (20 July 2024) — the 8.5 million affected Windows devices, “less than one percent of all Windows machines”; cf. the CrowdStrike PIR for the Channel File 291 out-of-bounds read and the kernel crash.

3. CrowdStrike PIR (24 July 2024) — the finding that the update did not clear the testing depth or staged rollout applied to sensor code; the staggered-deployment and canary commitment is quoted from this report.

4. U.S. Government Accountability Office, Cyber Resiliency: CrowdStrike Outage Highlights Challenges, GAO-24-107733 (23 September 2024); House Committee on Homeland Security, Subcommittee on Cybersecurity and Infrastructure Protection hearing, 24 September 2024 — supply-chain and concentration risk in endpoint security.

5. CrowdStrike PIR remediation commitments (2024); Microsoft, “Taking steps that drive resiliency and security for Windows customers,” Windows Experience Blog (12 September 2024), following the Windows Endpoint Security Ecosystem Summit of 10 September 2024 — new platform capabilities for security vendors rather than the elimination of kernel access.

6. Delta Air Lines v. CrowdStrike, Fulton County Superior Court, order on the motion to dismiss, 16 May 2025 (Judge Kelly Lee Ellerbe); the passenger class action against Delta, N.D. Ga., May 2025 (Judge Mark Cohen), allowed to proceed in part.

7. B. Beyer et al. (eds.), Site Reliability Engineering (2016) — staged rollout and canarying; cf. Knight Capital (Case 143).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The CrowdStrike outage is the live case for what happens when a deployment safety architecture treats categories of artifact differently than the operational system actually does. Content looked operationally identical to code; it was treated as different in the deployment pipeline. The mismatch became the largest IT outage in history.

LENS APPROACH

LENS uses CrowdStrike in LEN 5 as a categories-and-boundaries capability case and in LEN 2 for the vendor-customer trust architecture: customers trusted CrowdStrike’s deployment safety; that trust was decisive for the operation of the global economy on a single day.

| DOMAIN                                | FAILURE MODES                            | LENS COMPETENCY                                                         | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|------------------------------------------|-------------------------------------------------------------------------|--------------|---------|-----|
| TECHNOLOGY                            | → <div>D</div> <div>K</div> <div>G</div> | → <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 4            | 5.4     | 5   |
| Navigating Sociotechnical Constraints |                                          |                                                                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Define deployment-safety gates by operational impact, not artifact category: anything that can crash the kernel clears the same testing and staged rollout as code.
- Make staged rollout mandatory for every update, so a faulty push is caught by a first wave instead of reaching the entire install base at once.
- Design for recovery, not just prevention — assume a bad update will ship and ensure it does not require hands-on intervention at every device.

IN OPERATION / AFTER

- Audit vendor deployment-safety practices your institution depends on, rather than treating the outsourced safety gate as trustworthy by default.
- Map and reduce concentration risk so a single vendor fault cannot take down a large slice of critical operations simultaneously.
- Run post-incident reviews that interrogate category boundaries in the pipeline and feed the findings back into validation gates.

## REFLECTION QUESTIONS

1. Identify a vendor relationship in your domain whose deployment-safety practice your institution does not audit. What would the audit reveal?
2. Design the cross-vendor staged-rollout protocol that should be standard for endpoint security software.
3. CrowdStrike's pipeline treated content as safer than code, but the kernel did not. What category boundary in one of your systems is an organizational convenience that the operational reality ignores — and what would it cost if it broke?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · governance frameworks · bias & evidence audits





Chapter 12

# Disaster Prevention & Recovery — What Works — and the Frontier

*Prevention as a deliverable before the event; recovery as a discipline after it.*

---

*The disaster that did not happen had a budget, an owner, and a deadline.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

BEFORE THE EVENT — PREVENTION

---

# Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable

**G** Governance & Trust

**K** Knowledge & Institutional Memory

**H** Human-System Interface

**IMPACT** From 1915 to 1963 the U.S. Navy lost 16 submarines to non-combat causes; since 1963 it has lost one — USS Scorpion — which was not SUBSAFE-certified. The Columbia Accident Investigation Board cited SUBSAFE as a model NASA could emulate

## IN BRIEF

USS Thresher was lost with 129 aboard on April 10, 1963. Within fifty-four days the U.S. Navy created SUBSAFE — a program that certifies design, material, fabrication, and testing for every component inside the submarine's watertight-integrity boundary and its safe-recovery systems. The requirements were issued by December 20 of that same year. The program demands what it calls "Objective Quality Evidence" for every step — verifiable fact, not probabilistic assessment — and pairs that with annual training and recurring audits across the entire fleet lifecycle. The documented result is a step-change in non-combat submarine loss rates: 16 losses across the 48 years before SUBSAFE; one loss (USS Scorpion, not SUBSAFE-certified) across the 63 years since. The Columbia Accident Investigation Board cited SUBSAFE in 2003 as one of three programs that could be models for NASA. The honest hedge survives: the zero-loss record is correlational across decades with many co-varying factors — submarine design, reactor maturity, operating procedures, intelligence environment — and SUBSAFE's own program literature notes that the requirements can look "excessive." The case is the archetype of treating capability requirements as a recurring, auditable, non-waiverable deliverable across the entire system lifecycle, with the hedges that decades-long capability-engineering claims have to carry.

## BACKGROUND

USS Thresher, the lead boat of a new class of U.S. nuclear attack submarines, was lost on April 10, 1963 during post-overhaul sea trials off Cape Cod. All 129 aboard died. The investigation identified a likely sequence — silver-brazed piping joint failure, flooding, electrical fault that scrambled the reactor, inability to blow ballast tanks fast enough to recover — that pointed not to a single defective component but to a gap in the whole way the fleet certified the watertight-integrity boundary and the systems for recovering from a flooding casualty.<sup>o</sup>

## THE INTERVENTION

The Navy's response was institutional speed of an unusual order. SUBSAFE was created within fifty-four days of the loss; formal requirements were issued by December 20, 1963. The program scopes itself to two things: the watertight-integrity boundary (every component that holds back seawater at depth), and the safe-recovery systems (the ballast-blow

and propulsion systems that get a boat to the surface if flooding starts). Inside that scope it demands certification of design, material, fabrication, testing, and configuration control — with what the program calls “Objective Quality Evidence” attached at every step. That phrase is the cultural artifact the program rests on: verifiable fact, not probabilistic assessment, is what the certification rests on.<sup>1</sup>

#### HOW IT WORKED

What makes SUBSAFE a sustainment intervention rather than a design intervention is the lifecycle discipline. Annual training across the fleet, recurring audits at shipyards and tenders, change-control on every modification, and a culture that treats requirements as non-waiverable artifacts the program-management chain cannot trade away under schedule pressure. The certification is not done at launch; it is the condition of being allowed to dive. Each overhaul re-engages the certification process. The cost is real — the program’s own histories concede the requirements can look “excessive” — and the program treats that cost as the price of the capability the certification produces.<sup>2</sup>

#### THE EVIDENCE

The documented outcome is one of the cleanest before/after records in the safety-engineering literature. From 1915 to 1963 the Navy lost 16 submarines to non-combat causes (collision, flooding, equipment failure, fire). Since 1963 it has lost one — USS Scorpion in 1968, the only post-1963 loss that was not SUBSAFE-certified. In 2003 the Columbia Accident Investigation Board, examining the loss of the Space Shuttle Columbia, cited SUBSAFE in its final report as one of three independent safety programs whose practices yielded lessons NASA should weigh in redesigning its own organization. The endorsement comes from a board examining a different catastrophic-system domain — and one chaired by a Navy admiral, with the commander of the Naval Safety Center among its members.<sup>3</sup>

#### WHAT TRANSFERRED

The hedge has to survive into the case. The zero-loss record since 1963 is correlational across more than six decades and many co-varying factors: submarine design generations, reactor maturity, training systems, operating procedures, intelligence environment, the absence of certain operational stressors that the Cold War sometimes produced. Attributing the entire outcome to SUBSAFE alone overstates what the evidence can support. What the evidence does support is that the **program** — the requirements discipline, the Objective Quality Evidence standard, the lifecycle audit cycle, the non-waiverable culture — has been a defining feature of the capability since 1963, and has survived endorsement by an independent investigation in a different domain. The case teaches the requirements-as-sustainment-deliverable form at its strongest, with the honest hedge that decades-long capability claims have to carry.<sup>4</sup>

*Objective Quality Evidence — verifiable fact, not probabilistic assessment — is the cultural artifact the program is built on.*

EDITORS’ SYNTHESIS OF THE SUBSAFE PROGRAM LITERATURE.

#### REFERENCES

1. NASA SMA (2006), “USS Thresher Lessons Learned” briefing — safety-message archive.
4. Columbia Accident Investigation Board (2003), final report — Volume I, endorsement of SUBSAFE as a model for NASA.

2. Rear Admiral Paul E. Sullivan, statement to House Science Committee (2003), NASA/Columbia archive — primary congressional record on SUBSAFE.

3. U.S. Navy NAVSEA, SUBSAFE program documentation — operating program publications.

5. MIT Press, “SUBSAFE: An Example of a Successful Safety Program” — book chapter (open access).

THE LEARNING ENGINEERING LENS

LE INSIGHT

SUBSAFE is the archetype of treating capability requirements as a recurring, auditable, non-waiverable deliverable across the entire system lifecycle. The before/after non-combat-loss record is one of the cleanest in safety engineering — and correlational across many co-varying factors over six decades. The hedge is part of the case.

LENS APPROACH

SUBSAFE is the canonical sustainment-engineering case (induced 1.4; LENS D1/PT3). LENS uses it in Domain 1 (Systems Analysis) for the requirements-as-deliverable discipline; in Domain 4 (Test and Evaluation) for the Objective-Quality-Evidence standard and the recurring-audit cycle; and in Domain 5 (Navigating Sociotechnical Constraints) for the non-waiverable culture that resists schedule pressure. Adjacent to the nurse-ratios case (Case 11) at the requirements-becomes- engineered layer, and to the WHO Surgical Checklist (Case 23) at the mandatory-mechanism layer.

| DOMAIN            | MODES ADDRESSED    | LENS COMPETENCY                                                       | PROBLEM TYPE | INDUCED | LEO |
|-------------------|--------------------|-----------------------------------------------------------------------|--------------|---------|-----|
| NAVAL ENGINEERING | → <div>GKH</div> → | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 3            | 1.4     | 1·4 |
| DEFENSE           |                    | Systems Analysis                                                      |              |         |     |
| SAFETY CERT       |                    |                                                                       |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Scope the certification boundary tightly — for SUBSAFE, the watertight-integrity boundary and the safe-recovery systems — so the discipline is enforceable, not aspirational across an undifferentiated whole.
- ▶ Make “Objective Quality Evidence” the cultural standard: verifiable fact at every certification step, not probabilistic assessment, and not signature-without-evidence.
- ▶ Treat the requirements as non-waiverable artifacts the program-management chain cannot trade away under schedule pressure; the program’s resistance to waivers is the program.

IN OPERATION / AFTER

- ▶ Operate the lifecycle discipline as the program: annual training, recurring audits, change-control on every modification, certification re-engaged at each overhaul.
- ▶ Carry the correlational hedge in any communication of the outcome record; a decades-long zero-loss record across co-varying factors is the strongest available evidence the program works, not closed proof.
- ▶ Treat external endorsement (Columbia Accident Investigation Board) as a teaching artifact about the **form** of the program, transferable to other catastrophic-system domains under their own scope discipline.

## REFLECTION QUESTIONS

1. Identify a capability in your domain where the certification is done at launch and not re-engaged across the lifecycle. What is the analog of SUBSAFE's annual training and recurring audit — and what is the resistance to it?
2. Specify what "Objective Quality Evidence" would mean in your context: verifiable fact at every step rather than signature-without-evidence. Which steps in your current process would not survive that standard?
3. SUBSAFE's outcome record is correlational across many co-varying factors. What is the minimum additional evidence you would require before treating a similar long-run record in your domain as evidence the program is what produced the outcome?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · usability testing · cognitive walkthroughs

## CASE 174

SOFTWARE SUSTAINMENT

CRITICAL INFRASTRUCTURE

FEDERAL PROGRAM MANAGEMENT

1996 – 2000

# Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed

**G** Governance & Trust   **D** Designed Out   **K** Knowledge & Institutional Memory

**IMPACT** The U.S. federal government and the broader public and private sectors invested an estimated 100 billion dollars (U.S.) between 1995 and 2001 remediating two-digit-year date handling in legacy systems; the January 1, 2000 rollover passed with minimal disruption to critical infrastructure — the success contributed to the durable counterfactual debate about whether the threat justified the spending

## IN BRIEF

The “Year 2000 problem” — Y2K — was the structural consequence of decades of legacy software representing year fields as two digits, ambiguous between 1900 and 2000 at the rollover. The affected code ran banking systems, embedded controllers in industrial infrastructure, federal benefit-payment systems, air traffic and rail systems, and the broader public and private software base. From 1996 through December 1999 the U.S. federal government, under sustained Office of Management and Budget reporting and GAO audit, drove an inventory-and-remediation program across mission-critical federal systems, while the private sector executed a parallel multi-year effort. The Commerce Department put total U.S. Y2K spending, covering repairs started in 1995 through 2001, at about \$100 billion. The January 1, 2000 rollover passed with minimal disruption to critical infrastructure. The case is the canonical instance of a believed-and-treated aging-system transition in the recent regulatory record: the threat was specific, the program was inventoried at line-item granularity, the deadline was immovable, and the institutional discipline was sustained over four years. The hedge survives in the durable counterfactual debate the success itself produced — would the rollover have passed similarly with less spending? — and the published GAO record characterizes the program as a major federal management success without claiming the counterfactual is closed.

## BACKGROUND

The Y2K problem was an artifact of decades of software development in which year fields had been stored as two digits to save memory and database space when memory and database space were expensive. By the mid-1990s the code base had aged into a regime where the two-digit representation was a foreseeable failure: date arithmetic across the

January 1, 2000 boundary could return wildly incorrect results — interest calculations, scheduling logic, embedded controller timestamps, federal benefit eligibility checks. The structural form is the C7 case par excellence: legacy code aging past the assumptions of its original design, with a hard, foreseeable deadline.<sup>0</sup>

#### THE INTERVENTION

The federal-program-management response began in earnest in 1996 and accelerated through 1997–1998 with the creation of the President’s Council on Year 2000 Conversion (established February 1998) and the OMB quarterly reporting cycle. Federal mission-critical systems were inventoried at line-item granularity — OMB directed the major departments and agencies to submit quarterly reports beginning 15 May 1997, agency by agency, on the proportion of their mission-critical systems that were Y2K-compliant. In that first report the figure across the 24 major departments and agencies was 21 percent; by December 1999 OMB reported 99.9 percent. The discipline of the program rested on three institutional features: line-item inventory at the level of the actual systems, OMB-enforced quarterly status reporting, and GAO-sustained external audit that named agencies whose remediation lagged.<sup>1</sup>

#### HOW IT WORKED

The parallel private-sector effort was as large and is harder to characterize precisely. Total U.S. investment estimates cluster around the \$100 billion figure; global estimates are larger. Major financial institutions, utilities, telecommunications providers, and industrial operators ran their own multi-year inventory-and-remediation programs. Industry-coordination bodies (banking, electric utility, telecom) operated parallel structures to the federal program’s coordination role. The deadline forced parallel execution across the public and private sectors at a scale and pace that legacy-software remediation does not usually see.<sup>2</sup>

#### THE EVIDENCE

January 1, 2000 passed with minimal disruption to critical infrastructure. A small number of localized incidents occurred and were absorbed; the catastrophic-scenario possibilities the program had been built to prevent — major financial system failure, power-grid cascade, air-traffic disruption, federal benefit payment failure — did not materialize. The GAO published a sustained record of the federal program through and past the rollover; the broader retrospective literature characterizes the four-year effort as one of the major federal program-management successes of the era.<sup>3</sup>

#### WHAT TRANSFERRED

The hedge that survives the case is the counterfactual debate the success itself produced. The rollover passed; the counterfactual — whether the same outcome would have obtained with substantially less spending — is structurally unobservable, because the program ran. The post-2000 retrospective literature includes serious arguments on both sides: that the threat was real and the investment was the reason the rollover passed quietly, and that the investment was substantially overestimated relative to the actual latent failure population. The case is teachable on the institutional discipline — line-item inventory, OMB-enforced reporting, GAO audit, sustained four-year cadence, immovable deadline — and on the structural feature that the C7 transition succeeded because it was believed and treated. The counterfactual debate is preserved as part of the case rather than smoothed away.<sup>4</sup>



*The C7 transition succeeded because it was believed and treated. The counterfactual – whether the threat was as large as the response implies – the program structurally cannot answer.*

EDITORS' SYNTHESIS OF THE GAO Y2K REPORT SERIES AND THE POST-2000 RETROSPECTIVE LITERATURE.

REFERENCES

1. General Accounting Office, Year 2000 Computing Challenge report series (1997–2000), particularly GAO/AIMD-00-290, Lessons Learned Can Be Applied to Other Management Challenges (September 2000) — the federal-program-management record and GAO's own retrospective; GAO designated Y2K a high-risk area in February 1997.

2. Office of Management and Budget, quarterly reports on federal Y2K remediation status (1997–1999) — program-self-report tier; OMB “directed the major departments and agencies to submit quarterly reports beginning May 15, 1997,” and reported 21 percent of their mission-critical systems compliant in May 1997 against 99.9 percent in December 1999 (figures as given in GAO/AIMD-00-290).

3. U.S. Department of Commerce, Economics and Statistics Administration, The Economics of Y2K and the Impact on the United States (November 1999) — total U.S. Y2K spending of about \$100 billion covering repairs started 1995 through 2001, of which the federal share was about \$8 billion.

4. President's Council on Year 2000 Conversion, The Journey to Y2K final report (2000) — institutional retrospective of the federal coordination effort; the Council was established by Executive Order 13073 in February 1998.

5. Manion & Evan (2000), “The Y2K Problem and Professional Responsibility: A Retrospective Analysis,” Technology in Society 22(3), 361–387, doi:10.1016/S0160-791X(00)00015-4 — retrospective literature on the counterfactual debate.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Y2K remediation is the canonical case of a believed-and- treated aging-system transition. Line-item inventory, OMB quarterly reporting, GAO sustained audit, an immovable deadline, and a sustained four-year program-management cadence converted a foreseeable legacy-software failure into a transition the rollover passed quietly. The counterfactual debate the success produced is part of the case.

LENS APPROACH

Y2K is the legacy-software-sustainment case (induced 7.1; LENS D1/PT4) — Domain 1 for line-item inventory + immovable deadline; Domain 5 for OMB-reporting + GAO-audit. Closes C7 with Cases 114, 115, 156. Counterfactual-debate hedge preserved.

| DOMAIN                     | MODES ADDRESSED | LENS COMPETENCY    | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|-----------------|--------------------|--------------|---------|-----|
| SOFTWARE SUSTAINMENT       | GDK             | D1D2D3D4D5         | 4            | 7.1     | 1·5 |
| CRITICAL INFRASTRUCTURE    | →               | → Systems Analysis |              |         |     |
| FEDERAL PROGRAM MANAGEMENT |                 |                    |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Inventory the legacy at line-item granularity — actual systems, not categories — so the remediation status can be reported and audited against a denominator the program can defend.
- ▶ Build the immovable deadline into the program's operating discipline; the Y2K deadline could not be

IN OPERATION / AFTER

- ▶ Preserve the counterfactual hedge: a transition that succeeds because it was treated cannot prove the threat was as large as the response implies. The case is teachable on the institutional discipline, not on the closed answer to “was it worth it.”

- rebaselined, and the program's discipline came from the deadline's hardness rather than from management exhortation.
- Pair internal OMB reporting (the program's own status discipline) with external GAO audit (independent verification); the combination is what produced the institutional record the retrospective rests on.
- Carry the public/private parallel-execution lesson: the federal program and the broader private-sector effort moved together because the deadline was external to both, and the coordination mechanisms operated alongside each other rather than depending on each other.
- Treat the Y2K case as part of the v2 aging-system quartet (Cases 114, 115, 156) — the instance where the transition was a software-sustainment problem with the largest counterfactual-debate hedge.

## REFLECTION QUESTIONS

1. Identify a legacy-software or aging-system transition in your domain whose deadline is foreseeable but rebaselinable. What would the program look like if the deadline were treated as immovable — line-item inventory, quarterly reporting, sustained external audit — and what is the institutional cost of treating it that way?
2. Specify the public/private parallel-execution structure you would build if the transition reached beyond a single organization. The Y2K coordination did not depend on one body controlling the others; it depended on the external deadline being equally hard for everyone.
3. The Y2K success produces a counterfactual debate. Identify a similar success in your domain whose justification rests on the absence of the failure it was built to prevent. What institutional discipline keeps the historical record honest about the counterfactual without diminishing the discipline that produced the outcome?

**WHO BUILDS THIS** governance, ethics & measurement · systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · task analysis · design prototyping · knowledge bases · data pipelines

## AFTER THE EVENT — RESPONSE AND RECOVERY

---

# INPO and the Nuclear Academy

**T** Training Gap   **K** Knowledge & Institutional Memory   **G** Governance & Trust

**IMPACT** No INES Level 4-or-above accident at a U.S. commercial reactor since Three Mile Island; sustained improvement in INPO/WANO performance indicators across the industry

## IN BRIEF

Three Mile Island did not produce a reactor accident at the next plant over — it produced an institution. The Kemeny Commission reported on 30 October 1979 that the industry should “set and police its own standards of excellence”; in December 1979, eight months after the accident, the U.S. commercial nuclear industry established the Institute of Nuclear Power Operations on a stark premise: an accident at any single plant would threaten every operator’s license, and no utility could engineer its safety capability alone. Funded by the utilities it evaluated and operating without statutory authority, INPO set training and certification standards, accredited every plant’s programs through the National Academy for Nuclear Training, and ran peer evaluations in which operators from one utility scrutinized another’s control rooms and records. The pre-TMI culture of complacency gave way to mandated vigilance. No U.S. commercial reactor has had a INES Level 4-or-above accident since.

## BACKGROUND

The 1979 partial meltdown at Three Mile Island exposed not just a plant-level failure but an industry with no shared mechanism for learning. The Kemeny Commission found that warnings had been disregarded: a September 1977 event at Davis-Besse foreshadowed TMI, was investigated by Babcock & Wilcox and the NRC, and produced no notice to utilities before the accident. The structural problem sat above any single control room: a lesson learned at one plant had no path to the others, so the same latent failure could surface repeatedly across an industry that never compared notes.<sup>0</sup>

## THE INTERVENTION

The Kemeny Commission’s recommendation to the industry was that it “set and police its own standards of excellence to ensure the effective management and safe operation of nuclear power plants,” and the utilities established the Institute of Nuclear Power Operations in December 1979 in response — eight months after the accident and two months after the report. Its premise was that an accident anywhere threatened everyone’s license to operate. INPO set training and certification standards for operators and supervisors, and in 1985 the National Academy for Nuclear Training began accrediting each facility’s programs. The shared-exposure premise was what gave a body with no statutory power its teeth — every utility had a direct stake in every other utility’s competence, because one failure could end the whole industry’s license to operate.<sup>1</sup>

## HOW IT WORKED

INPO's decisive mechanism was honest peer review: teams of operators from one utility examined another's procedures, control rooms, and incident records, reporting candidly because every utility was, in the title of one history, a hostage of the others. Funded by the utilities it evaluated and holding no statutory authority, INPO depended on shared catastrophic exposure to make its findings stick. Peer review by working operators rather than distant regulators also meant the examiners knew what to look for and the findings carried the weight of professional judgment, not just rule compliance.<sup>2</sup>

## THE EVIDENCE

The post-TMI culture shifted from smugness to mandated vigilance, and no U.S. commercial reactor has had an INES Level 4-or-above accident since INPO's founding — though Davis-Besse's 2002 reactor-head corrosion (Case 160) was rated a Level 3 serious incident, and it was a regulator-mandated inspection rather than peer review that found it. INPO operated alongside NRC regulation and the post-Three-Mile-Island design and procedural changes, and no published evaluation isolates its separate contribution. Industry performance indicators tracked by INPO and its international counterpart WANO improved steadily and broadly across the fleet. The broad, steady improvement across the whole fleet — not just the strongest plants — is the signature of a learning mechanism working as designed: the laggards were pulled up by the same peer-review architecture that held the leaders to standard.<sup>3</sup>

## WHAT TRANSFERRED

INPO is the strongest evidence in any domain that capability engineering can be undertaken at the level of an entire industry rather than a single organization. Its enabling conditions — shared catastrophic exposure, regulatory legitimacy, and an honest peer-review architecture — recur wherever one operator's failure can damage every operator, and it informed the founding of WANO after Chernobyl. That the model crossed national borders to WANO is itself evidence that the design is portable: the enabling conditions, not the particular American institution, are what make the mechanism work.<sup>4</sup>



INDUSTRY-WIDE PEER REVIEW: A LEARNING RING AFTER TMI

*Every utility recognized that an accident at any single plant would affect every operator's license to operate.*

PARAPHRASING THE INSTITUTIONAL ANALYSIS IN REES, HOSTAGES OF EACH OTHER, 1994

REFERENCES

1. Report of the President’s Commission on the Accident at Three Mile Island (Kemeny Commission, 30 October 1979) — the pre-TMI culture, the disregarded September 1977 Davis-Besse precursor, and the recommendation that the industry “set and police its own standards of excellence.”

2. Willard, R. F. (2019), The Role of the Institute of Nuclear Power Operations in Supporting the United States Commercial Nuclear Power Industry’s Focus on Nuclear Safety, testimony for the record, U.S. Senate Committee on Environment and Public Works, 13 November 2019 — INPO’s design, funding, peer evaluations, and the formation of the National Academy for Nuclear Training in 1985. INPO’s own history dates its formal establishment to December 1979, in response to the Commission’s recommendations; the Kemeny report does not name INPO.

3. Rees, J. (1994), Hostages of Each Other: The Transformation of Nuclear Safety since Three Mile Island — INPO’s design, the peer-evaluation mechanism, and the “hostages” premise (paraphrased).

4. World Nuclear Association — Three Mile Island Accident; INPO/ WANO performance indicators.

5. Nuclear Energy Institute, “Lessons from the 1979 Accident at Three Mile Island”; National Academy for Nuclear Training — accreditation, and the founding of WANO after Chernobyl.

THE LEARNING ENGINEERING LENS

LE INSIGHT

INPO is the strongest evidence in any domain that capability engineering can be undertaken at the level of an **industry**, not just an organization. The conditions that made it possible — shared catastrophic exposure, regulatory legitimacy, an honest peer-review architecture — appear wherever a single failure can damage every operator.

LENS APPROACH

LENS uses INPO in LEN 8 as the canonical example of industry-level learning: students identify the structural conditions in their own domain that would permit (or block) an INPO-equivalent and design the peer-review architecture required. LEN 1 uses the founding moment — eight months after TMI — to discuss the **speed** a credible response to catastrophe demands.

| DOMAIN                                | MODES ADDRESSED | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|---------------------------------------|-----------------|------------------|--------------|---------|-----|
| ENERGY                                | → T K G         | → D1 D2 D3 D4 D5 | 4            | 6.1     | 5   |
| Navigating Sociotechnical Constraints |                 |                  |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Found the body on shared catastrophic exposure — make explicit that one operator’s failure threatens every operator’s license — so members have a direct stake in each other’s competence.
- ▶ Set common training, certification, and accreditation standards across the industry rather than leaving each organization to learn alone.
- ▶ Staff peer review with working operators, not distant regulators, so examiners know what to look for and findings carry professional weight beyond rule compliance.

IN OPERATION / AFTER

- ▶ Track fleet-wide performance indicators and confirm the laggards are being pulled up, not just the leaders held to standard — the signature of a learning mechanism working.
- ▶ Sustain candid peer review by keeping the body funded by and accountable to its members, so the honest examination that makes it effective does not erode into formality.
- ▶ Export the enabling conditions rather than the institution when scaling (as INPO informed WANO), adapting the shared-exposure-plus-peer-review design to each new context.

## REFLECTION QUESTIONS

1. What is the equivalent of “an accident at any single plant affects every operator” in your domain? If the answer is “nothing,” what does that tell you?
2. INPO was stood up in eight months. Pick a current cross-organizational capability problem and write the eight-month deliverable that would constitute a credible response.
3. INPO held no statutory authority yet made its findings stick through shared catastrophic exposure and peer review. Design the non-statutory mechanism that could enforce a standard in your domain, and name the shared stake that would give it teeth.

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

# Tylenol Recall

**G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** The founding U.S. corporate crisis-management case; produced tamper-evident packaging regulation and modern recall practice

## IN BRIEF

In 1982, seven people in the Chicago area died after taking Tylenol capsules laced with potassium cyanide. Not knowing who was responsible or how widespread the tampering was, Johnson & Johnson recalled every bottle of Tylenol capsules nationwide — 31 million bottles, at a cost of roughly \$100 million — suspended advertising, and engaged openly with the FBI and FDA. The response was unprecedented in U.S. corporate practice, and it was a direct application of the J&J Credo, written in 1943, which had pre-specified that the company's first responsibility was to its customers. The reform that followed reshaped consumer packaging worldwide — tamper-evident seals and blister packs — and the FDA issued tamper-resistant-packaging rules within months. Tylenol climbed back to a 30 percent share by mid-1983 and to its pre-crisis 35 percent by year's end.

## BACKGROUND

Johnson & Johnson's corporate Credo, written in 1943, pre-specified that the company's first responsibility was to the patients and consumers who used its products, ahead of shareholders. For four decades it was a statement of values; in 1982 it became an operational decision rule under extreme pressure. The ordering was explicit — customers ahead of shareholders — which is precisely the ranking that crisis pressure inverts, so committing to it in advance pre-decided the hardest trade-off before it had to be faced.<sup>0</sup>

## THE INTERVENTION

After seven people in the Chicago area died from Tylenol capsules laced with potassium cyanide, and with the source and scope of the tampering unknown, Johnson & Johnson recalled every bottle of Tylenol capsules nationwide — about 31 million bottles, at a cost near \$100 million — suspended all advertising, and engaged publicly with the FBI and FDA rather than minimizing exposure. Recalling nationwide despite the tampering being known only in Chicago was the decisive choice — it treated the unknown scope as a reason to protect every customer rather than as room to limit the company's own exposure.<sup>1</sup>

## HOW IT WORKED

The decisive element was a commitment made in writing beforehand. Because the Credo had already decided, decades earlier, that the customer came first, the 1982 leadership did not have to improvise an ethical calculus under crisis — it executed a pre-made decision. CEO James Burke later credited the Credo with making clear “exactly what we were all about” the moment the deaths occurred. Pre-commitment worked because it moved the decision

out of the moment of maximum pressure — when fear and legal caution push hardest toward minimization — and into a calmer time when the right ordering could be set down without that distortion.<sup>2</sup>

### THE EVIDENCE

The response, without precedent in U.S. corporate practice, preserved public trust: Tylenol climbed from an 8 percent share back to 30 percent by mid-1983 and to its pre-crisis 35 percent by year's end, despite the enormous short-term cost. The case became the canonical positive example in business education of crisis response driven by capability commitment rather than legal minimization. The market recovery is what makes the case persuasive rather than merely admirable — the \$100 million spent protecting customers was repaid in the trust that brought them back, so the pre-committed choice proved sound on its own terms.<sup>3</sup>

### WHAT TRANSFERRED

The reform reshaped consumer-product packaging worldwide — tamper-evident seals, blister packs, and caplet forms became standard — and the FDA promulgated tamper-resistant-packaging regulations within months. The deeper transfer is the principle that values must be pre-committed in writing to be operational under stress, not invented in the moment. The packaging reform and the decision-rule principle are the two layers of the transfer — one a physical safeguard against the specific threat, the other an institutional safeguard against the improvisation that crisis invites.<sup>4</sup>

# 31M

BOTTLES RECALLED · ~\$100M COST

*the pre-committed institutional credo became operational under stress*

TYLENOL — VALUES PRE-COMMITTED IN WRITING, EXECUTED UNDER CRISIS

*The credo made it very clear at that point exactly what we were all about. It gave me the ammunition I needed to persuade shareholders and others to spend the \$100 million on the recall.*

JAMES BURKE (JOHNSON & JOHNSON CEO) — CLOSING SENTENCE QUOTED IN LASTING LEADERSHIP (WHARTON); OPENING LEAD IS THE EDITORS' PARAPHRASE

### REFERENCES

1. Edmondson, A. (2018), The Fearless Organization — pre-committed values under stress.
2. Kaplan, T. (2014), The Tylenol Crisis — the recall and corporate response.
3. James Burke (J&J CEO), in Lasting Leadership (Wharton) — the Credo quote.
4. Greyser, S., Johnson & Johnson: The Tylenol Tragedy (HBS case, 1992) — market recovery and crisis management.
5. FDA Final Rule on Tamper-Resistant Packaging (1982) — the packaging reform.



THE LEARNING ENGINEERING LENS

LE INSIGHT

Tylenol is the canonical positive case for institutional response to crisis. The capability that mattered was the pre-specified institutional commitment in the Credo. The crisis decision had been made decades earlier; in 1982 it was executed. The case is the strongest evidence in the business-ethics dataset that values must be pre-committed in writing to be operational under stress.

LENS APPROACH

LENS uses Tylenol in LEN 7 as the primary positive case for institutional governance under crisis and in LEN 10 (capstone) as a worked example of pre-committed capability that executed under operational pressure.

| DOMAIN     | MODES ADDRESSED | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO |
|------------|-----------------|---------------------------------------|--------------|---------|-----|
| HEALTHCARE | → G N           | → D1 D2 D3 D4 D5                      | 3            | 4.4     | 5   |
| INDUSTRIAL |                 | Navigating Sociotechnical Constraints |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pre-commit the hardest trade-off in writing before the crisis — rank customer safety ahead of shareholder exposure — so leadership executes a pre-made decision rather than improvising under pressure.
- ▶ Set the rule in a calm period when fear and legal caution cannot distort the ordering, since those forces push hardest exactly when the decision must be made.
- ▶ Make the commitment concrete enough to act on — a nationwide recall, open engagement with regulators — so unknown scope becomes a reason to protect everyone rather than room to limit exposure.

IN OPERATION / AFTER

- ▶ Pair the institutional decision rule with a physical safeguard against the specific threat (tamper-evident packaging) so the response addresses both the improvisation problem and the vulnerability.
- ▶ Treat the preserved trust and market recovery as the measure that the pre-committed choice was sound, not merely admirable, and document it to defend the principle internally.
- ▶ Embed the commitment durably enough that it survives leadership turnover, so the next crisis meets the same pre-decided rule rather than a fresh improvisation.

REFLECTION QUESTIONS

1. What is your institution’s equivalent of the J&J Credo, and is it operational under crisis or aspirational?
2. Pre-commitment is hard to enforce later. Design the institutional architecture that makes a Tylenol-style response the only available option in the worst case.
3. J&J recalled nationwide while the tampering was known only in Chicago, treating unknown scope as a reason to protect everyone. Identify a decision in your domain where uncertainty currently licenses minimizing exposure, and write the pre-committed rule that would flip it toward protection instead.

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# CIRAS — Confidential Incident Reporting for UK Rail

**G** Governance & Trust

**K** Knowledge & Institutional Memory

**N** Normalization of Deviance

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** Between 2008 and 2012 the UK rail Confidential Incident Reporting and Analysis System received 2,228 reports — 45% led to tangible safety improvements and about 33% contained important safety information (program self-report); directly influenced a confidential reporting system in the U.S.

## IN BRIEF

CIRAS began as a 1995–1997 ScotRail pilot — a structured channel for rail workers to report hazards and near-misses confidentially, insulated from the employer’s disciplinary process. After the Ladbroke Grove crash in 1999, the program was mandated across UK mainline rail in 2000; in 2008 it became an independent unit within the Rail Safety and Standards Board. The published record states that between 2008 and 2012 CIRAS received 2,228 reports, of which the operating program reports 45% led to tangible safety improvements and about 33% contained important information about safety that was new to the company concerned. The program directly influenced the design of the Federal Railroad Administration’s Confidential Close Call Reporting System (C3RS) in the United States. The interview-based method surfaces motive and intent — the **why** behind a near-miss — that company databases miss precisely because those databases are tied to discipline. The honest hedge that survives into the case: the 45%-led-to-improvement figure is self-reported by the operating program, not independently audited. The case is the non-aviation companion to ASRS / CRM in v1 (Cases 117 + 119), keeping the non-punitive-reporting-with-credible-commitment competency from resting entirely on aviation evidence.

## BACKGROUND

Confidential incident reporting as a safety-culture intervention rests on a specific structural argument: the richest information about a near-miss — the operator’s intent, the team’s local pressures, the workaround that almost worked — is exactly the information that an employer-controlled reporting channel cannot collect, because reporters with anything to lose will not put it there. The aviation experience (ASRS, CRM; v1 Cases 117 and 119) establishes the pattern at the canonical safety-culture scale. CIRAS is the same pattern, in a different industry, with the credible commitment supplied differently — and the corpus needs the non-aviation evidence.<sup>o</sup>

## THE INTERVENTION

CIRAS began in 1995–1997 as a ScotRail pilot. The structural design was the interview, not the form: trained interviewers took reports from rail workers, anonymized them, and surfaced patterns to the operating companies. The credible commitment that made the channel safe to use was not just confidentiality — it was institutional independence from the employer's disciplinary process. After the Ladbroke Grove crash in October 1999 (31 deaths, the accident that defined the decade's UK rail-safety reform), CIRAS was mandated across UK mainline rail in 2000. In 2008 it became an independent unit within the Rail Safety and Standards Board (RSSB), with the independence from employer discipline written into its operating structure.<sup>1</sup>

## HOW IT WORKED

The published record on outcomes is what the case rests on. Davies et al. (Cognition, Technology & Work) describe the method and its yield. The operating program reports that between 2008 and 2012, CIRAS received 2,228 reports, of which 45% led to tangible safety improvements and approximately 33% contained important information about safety that was new to the company concerned. The program's design has directly influenced the construction of a confidential reporting system in the United States — the Federal Railroad Administration's Confidential Close Call Reporting System (C3RS), which the operating record says used lessons learned from CIRAS. The interview-based method surfaces motive and intent in a way that incident-database schemas tied to discipline cannot.<sup>2</sup>

## THE EVIDENCE

The honest hedge has to survive. The 45%-led-to-improvement figure is self-reported by the operating program — CIRAS itself characterizes its outputs, and there is not yet an independent audit of that classification. The peer-reviewed literature on CIRAS describes the method and the institutional design, and reports the program-supplied figures rather than independently validating them. The case is teachable on the evidence of the institutional design and the published method; the operating outcome statistics are the strongest current claim, and the editor and downstream readers should treat them as program self-report rather than audited finding.<sup>3</sup>

## WHAT TRANSFERRED

What CIRAS adds to the corpus is non-aviation depth in the non-punitive-reporting-with-credible-commitment competency. The mechanism (a structured confidential channel) paired with the credible commitment (independence from the employer's disciplinary process, written into operating structure) is the cultural half of capability that the aviation cases also teach. Drafted together, ASRS (Case 119), CRM/CAST (Case 117), and CIRAS (this case) show that the structural form is transferable across high-consequence operational industries — and that the credible commitment, not the channel alone, is what makes the reporting safe to use.

*The 45%-led-to-improvement figure is the operating program's self-report. The institutional design is the audited finding.*

EDITORS' SYNTHESIS OF DAVIES ET AL. AND THE STRATHCLYDE IMPACT CASE STUDY.

REFERENCES

1. Davies, Wright, Courtney, & Reid (2000), "Confidential Incident Reporting on the UK Railways: The 'CIRAS' System," Cognition, Technology & Work 2, 117–125, doi:10.1007/PL00011494 — the interview-based method and the structural argument for a channel independent of the employer.

2. University of Strathclyde, CIRAS impact case study (REF) — the institutional history: the ScotRail pilot of 1995–97, the national mandate of 2000 after Ladbroke Grove, and "in 2008, CIRAS became an independently operated unit within the Rail Safety and Standards Board." Secondary accounts date the transfer to 2009, when the CIRAS Charitable Trust wound up its funding role; do not "correct" the 2008 date against them.

3. University of Strathclyde, CIRAS impact case study (REF) — the operating-program self-report: "between 2008 and 2012 CIRAS received 2228 reports; 45% of these resulted in tangible safety improvements and approximately 33% contained important information about safety that was new to the company concerned"; and the influence on the FRA's Confidential Close Call Reporting System (C3RS).

4. Rail Safety and Standards Board (RSSB), CIRAS program documentation 2008 – present — operating-program publications; the peer-reviewed literature describes the method and reports the program-supplied figures rather than independently validating them.

5. Ladbroke Grove Rail Inquiry (Cullen, 2001), final report — the public inquiry into the crash that preceded the national mandate.

THE LEARNING ENGINEERING LENS

LE INSIGHT

CIRAS is the non-aviation companion to ASRS and CRM in the corpus. The pattern — confidential channel paired with credible commitment (institutional independence from discipline) — works in rail as it does in aviation. The operating outcome statistics are program self-report and deserve their tier acknowledgement; the institutional design is the audited finding.

LENS APPROACH

CIRAS is the non-aviation pairing-mechanism case (induced 4.2; LENS D5/PT2). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the credible-commitment design — independence written into operating structure — and in Domain 4 (Test and Evaluation) for the subobjective **Judgment under inadequate evidence** (4.4) on the operating self-report tier. Pair with ASRS (Case 119) and CRM/CAST (Case 117) as the cross-industry C4 set, and with the WHO Surgical Checklist (Case 23) at the mandatory-mechanism layer.

| DOMAIN           | MODES ADDRESSED | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO |
|------------------|-----------------|---------------------------------------|--------------|---------|-----|
| RAIL TRANSPORT   | G K N           | D1 D2 D3 D4 D5                        | 2            | 4.2     | 5   |
| SAFETY REPORTING |                 | Navigating Sociotechnical Constraints |              |         |     |
| SAFETY CULTURE   |                 |                                       |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Design the credible commitment first: institutional independence from the employer's disciplinary process, written into operating structure, not relying on goodwill that can be withdrawn.
- Choose the reporting method (interview, not just form) for the information class you want to surface — motive and intent — that form-based incident databases tied to discipline cannot collect.
- Anchor the program's authority in a sector-wide forcing event (Ladbroke Grove for rail; in other domains, the analog) so the mandatory-adoption decision rides on the safety crisis the program is built to address.

IN OPERATION / AFTER

- Report operating outcome statistics with their evidence tier: program self-report vs. independently audited. Treating the 45% figure as audited overstates the evidence; treating it as untested under-states the program's institutional credibility.
- Commission an independent audit of the program's classification of its own reports at a sustainable cadence — five years is a reasonable interval — so the operating record builds toward audited evidence over time.
- Carry the structural lesson across domains: ASRS, CRM, CIRAS — the cultural half of capability is the credible commitment, not the channel alone.

## REFLECTION QUESTIONS

1. Identify a high-consequence operational domain in your context where confidential incident reporting is absent or weak. What is the credible-commitment design that would make the channel safe to use — and is the institutional independence written into operating structure, or relying on goodwill?
2. Specify the information class your reporting program is built to surface: form-based incident counts, or interview-based motive and intent. The choice of method follows the information class, and they answer different questions.
3. The 45%-led-to-improvement figure is program self-report. Design the independent-audit cadence that would convert the operating record into audited evidence over time, without compromising the confidentiality the channel depends on.

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

# The UN Cluster Approach – Engineering Humanitarian Coordination After the Tsunami

**G** Governance & Trust

**IMPACT** IASC cluster approach adopted after the 2005 Humanitarian Response Review found weak sectoral capacity and no locus of accountability across the international system, conceived amid the Darfur response and the failures of the collaborative approach for IDPs; first applied in the October 2005 Pakistan earthquake response; eleven global clusters, each with a designated lead agency accountable for coverage and carrying provider-of-last-resort responsibility; two formal system-wide evaluations (Stoddard et al. 2007; Steets et al. 2010) found improved gap identification and coverage at real transaction cost, with national and local actors often marginalized; revised under the 2011 IASC Transformative Agenda after the Haiti earthquake exposed scale limits

### IN BRIEF

The response to the December 2004 Indian Ocean tsunami — hundreds of NGOs converging on the same coastlines, duplicated relief in accessible areas, unassigned gaps in hard ones — and the parallel critiques of the Darfur response exposed a structural fact about the international humanitarian system: no one was accountable for whole sectors of a response. The 2005 Humanitarian Response Review, commissioned by UN Emergency Relief Coordinator Jan Egeland, named the diagnosis, and the Inter-Agency Standing Committee's answer, rolled out from late 2005, was the cluster approach: each sector — logistics, shelter, health, water and sanitation, protection, camp management, and the rest — assigned a designated lead agency accountable for coverage and carrying provider-of-last-resort responsibility. Coordination was redesigned as a role architecture with named ownership rather than as goodwill. First applied in the October 2005 Pakistan earthquake response, the approach was evaluated twice at system scale: Phase 1 (Stoddard et al., 2007) found early gains in gap identification at real transaction cost; Phase 2 (Steets et al., GPPI/Groupe URD, 2010) found mostly positive coverage and gap effects with persistent weaknesses in local-actor inclusion and inter-cluster coordination. The 2010 Haiti earthquake exposed the mechanism's scale limits, and the IASC's 2011 Transformative Agenda revised it. The evaluations' own hedges are binding: benefits were uneven across clusters and countries, the costs were real, and national and local actors were often marginalized.

### BACKGROUND

The 2004 Indian Ocean tsunami produced the largest humanitarian mobilization to that date and, with it, the clearest exhibit of the system's chronic coordination failure. Hundreds of NGOs, UN agencies, and bilateral responders arrived with independent mandates, independent funding, and no mechanism that made any of them answerable for a sector as a whole. Relief duplicated where access was easy and cameras were present; gaps went unassigned where it was not; and when a sector was underserved, there was no one whose job it was to notice. The Darfur response drew the same critique in the same period, in a conflict setting: sectoral coverage depended on which agencies happened to show up, and the persistent gaps — protection above all — belonged to no one. The failure was not of effort or resources but of architecture: the system had participants, not owners.<sup>0</sup>

### THE INTERVENTION

The 2005 Humanitarian Response Review, commissioned by Emergency Relief Coordinator Jan Egeland from an independent consultant team, made the structural diagnosis explicit: capacity was unevenly distributed across sectors, and accountability for sectoral outcomes was absent by design, because no design existed. The Inter-Agency Standing Committee's response was the cluster approach, agreed in September 2005 as one pillar of a broader humanitarian reform. Each sector of response — logistics, emergency shelter, health, water/sanitation/hygiene, protection, camp coordination and camp management, nutrition, education, and others, eventually eleven global clusters — was assigned a designated lead agency, globally and in each activated country response. The lead is accountable to the Humanitarian Coordinator for the adequacy of coverage in its sector and carries the provider-of-last-resort commitment: where critical gaps persist and no other actor steps in, the lead must do its utmost to fill them. The redesign is the case's analytic core — coordination converted from a voluntary good into a role architecture with named ownership.<sup>1</sup>

### HOW IT WORKED

The mechanism was paired with an authorization, and the pairing is what distinguishes the cluster approach from the coordination exhortations that preceded it. The lead agency's

sectoral accountability was mandated by the IASC and the Emergency Relief Coordinator — the closest thing the voluntary humanitarian system has to an institutional spine — rather than left to inter-agency goodwill, and it was tested almost immediately: the October 2005 Pakistan earthquake, two months after the approach was agreed, became its first field application, with the initial set of clusters stood up mid-response. The Pakistan experience showed both faces at once — sectoral gap analysis that had never existed before, and meeting burdens heavy enough that NGO observers questioned whether the coordination was worth its transaction cost.<sup>2</sup>

### THE EVIDENCE

The formal evaluations are the reason the case can be told with discipline. The Phase 1 evaluation (Stoddard, Harmer, Haver, Salomons, and Wheeler, 2007) found early gains — better identification of gaps and stronger, more predictable sectoral leadership, with no observable increase in ultimate accountability — purchased at real transaction cost in meetings, reporting, and staff time, with benefits uneven across clusters and countries.<sup>3</sup> The Phase 2 evaluation (Steets, Grünewald, Binder, and colleagues, GPPi/Groupe URD, 2010), commissioned by the IASC across six country studies, found mostly positive effects on coverage and gap identification and judged the benefits to already slightly outweigh the costs, while documenting persistent weaknesses that the evaluators declined to smooth: national and local actors were often marginalized or excluded by an architecture built around international agencies, inter-cluster coordination remained weak, and cluster performance varied widely by sector and country. Those hedges are carried here as findings, not caveats — the outcome evidence is genuinely mixed at the margins even where the architectural gain is real.<sup>4</sup>

### WHAT TRANSFERRED

The 2010 Haiti earthquake stress-tested the mechanism at a scale it had not been engineered for — hundreds of organizations in a single cluster's meetings, a coordination apparatus that became process-driven enough to be accused of impeding the delivery it existed to enable — and the IASC Principals' December 2011 Transformative Agenda revised the approach in response, restating the clusters' original purpose around strategic gap analysis and adding system-wide activation protocols for the largest emergencies.<sup>5</sup> That arc is why the case sits on the works side of this book: an institution diagnosed its own coordination failure after catastrophe, redesigned accountability as a property of the multi-organization system rather than of any single agency, evaluated the redesign formally and in public, and revised it when the next catastrophe found its limits. The cluster approach did not solve humanitarian coordination; it made sectoral coverage someone's named job, made the residual gaps visible, and left an evaluation record honest enough to build on.

*Coordination redesigned as a role architecture with named ownership rather than as goodwill — and evaluated, in public, at system scale.*

EDITORS' SYNTHESIS OF THE HUMANITARIAN RESPONSE REVIEW (2005) AND THE TWO CLUSTER EVALUATIONS (2007, 2010).

### REFERENCES

1. Adinolfi, Bassiouni, Lauritzen, & Williams (2005), Humanitarian Response Review, commissioned by the UN Emergency Relief
4. Stoddard, Harmer, Haver, Salomons, & Wheeler (2007), Cluster Approach Evaluation, Final Report, OCHA Evaluation and Stud-



Coordinator, New York/Geneva: OCHA, August 2005 — the diagnosis of uneven sectoral capacity and absent accountability across the international system.

- Adinolfi et al. (2005), Humanitarian Response Review; Inter-Agency Standing Committee Principals' decision of 12 September 2005 — the cluster approach as one pillar of humanitarian reform: designated lead agencies, globally and per activated country response, accountable to the Humanitarian Coordinator and carrying provider-of-last-resort responsibility. Nine clusters were designated initially; the set later grew to eleven.
- Street, A., & Parihar, G. (2007), "The UN cluster approach in the Pakistan earthquake response: an NGO perspective," Humanitarian Exchange 37, ODI Humanitarian Practice Network — the first field application: information-sharing and named sectoral accountability set against "cluster creep," meeting burdens and English-only coordination.
- ies Section, with NYU Center on International Cooperation, ODI Humanitarian Policy Group, and the Praxis Group — Phase 1: "stronger and more predictable leadership over sectors. However, there has been no observable increase in ultimate accountability."
- Steets, Grünewald, Binder, de Geoffroy, Kauffmann, Krüger, Meier, & Sokpoh (2010), Cluster Approach Evaluation 2: Synthesis Report, GPPi / Groupe URD, commissioned by the IASC, Geneva: OCHA — Phase 2, six country studies: the benefits "already slightly outweigh its costs and shortcomings," against persistent marginalisation of national and local actors and weak inter-cluster coordination.
- Inter-Agency Standing Committee (2011 – 2012), IASC Transformative Agenda — Principals' decisions of December 2011 and associated protocols; Grünewald & Binder (2010), Inter-agency Real-Time Evaluation in Haiti: 3 Months After the Earthquake, Groupe URD / GPPi.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The cluster approach is the paired-intervention pattern at global scale: a mechanism — sector lead agencies with coverage accountability and provider-of-last-resort responsibility — paired with an authorization, the ERC/IASC mandate. Accountability became a designed property of a multi-organization system rather than an aspiration. The evaluation record keeps the case honest: gains in gap identification and coverage were real, and so were the transaction costs, the uneven benefits across clusters and countries, and the marginalization of national and local actors. The 2011 revision after Haiti is part of the intervention, not evidence against it.

LENS APPROACH

The cluster approach is the institution-building-after- catastrophe case at the largest organizational scale in the corpus (induced 6.1; LENS D5/PT6; with 4.4 alternate for the mechanism-plus-authorization pairing). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) as the anchor example of governance architecture for a multi-organization system — named sectoral ownership, provider of last resort, and a mandate that binds a voluntary system; in Domain 1 (Systems Analysis) for the Humanitarian Response Review's structural diagnosis of unowned sectors; and in Domain 4 (Test and Evaluation) for the two formal system-wide evaluations whose hedged findings — uneven benefits, real costs, marginalized local actors — drove the 2011 Transformative Agenda revision. Pair with the industry-level institution-building cases in Part VI: the discipline is the same, the scale is the test.



APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Assign every sector of a multi-organization response a named owner with explicit coverage accountability; a gap with no owner is a design decision, not an accident.
- Pair the ownership mechanism with an authorization that can bind — the lead-agency role works because the ERC/IASC mandate stands behind it, not because agencies agree to coordinate in principle.

IN OPERATION / AFTER

- Commission formal, published, system-wide evaluations of the coordination architecture itself — not only of individual operations — and let the findings, including the unflattering ones, drive revision.
- Treat local and national actors' marginalization as an architectural defect to be engineered out, not a cultural footnote; an inclusion weakness documented in 2010 that persists is a standing design debt.

- ▶ Make provider-of-last-resort responsibility explicit at designation time, so the question “who acts when no one volunteers?” has an answer before the gap appears.
- ▶ Budget the transaction cost of coordination honestly and re-verify the mechanism at each new scale; Haiti showed that a role architecture sized for one magnitude of response can fail at the next.

### REFLECTION QUESTIONS

1. Map a multi-organization response system in your domain — a hospital network’s surge plan, a district’s crisis-response stack, a consortium research program. For each sector of the work, who is the named owner accountable for coverage, and who is the provider of last resort? If the answer is “everyone” or “no one,” you have found the tsunami-response architecture.
2. The cluster evaluations found real gains purchased at real transaction cost, with benefits uneven across sectors and countries. Design the evaluation you would commission for a coordination mechanism in your domain: what would count as a gap-identification gain, how would you price the meeting and reporting burden against it, and who is independent enough to say the mechanism is not working?
3. The Phase 2 evaluation documented national and local actors marginalized by an architecture built around international agencies — a weakness named in 2010 and still central to humanitarian-reform debates a decade later. Identify the actors your coordination design treats as participants rather than owners. What would it take, structurally rather than rhetorically, to move them into the ownership architecture?

**WHO BUILDS THIS** governance, ethics & measurement — plus domain experts and a learning engineer to integrate.

**TOOLS** governance frameworks · bias & evidence audits

# The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed

**G** Governance & Trust   **K** Knowledge & Institutional Memory

**DISCLOSURE** Contribution: an editor leads the Johns Hopkins APL group that contributed engineering support to this system; the editor was not personally involved in the build. Included on the published peer-reviewed record and independent recognition.

**IMPACT** Launched 22 January 2020 by two researchers and became the world's de facto epidemiological reference for the COVID-19 pandemic: peak single-day usage above 4.6 billion data requests (29 March 2020), over 226 billion feature-layer requests and 3.6 billion page views by June 2022 per the peer-reviewed record; governments, newsrooms, and researchers built directly on its openly published data; 2022 Lasker-Bloomberg Public Service Award; ceased data collection 10 March 2023

## IN BRIEF

On 22 January 2020 — one day after the United States reported its first COVID-19 case — Lauren Gardner's Center for Systems Science and Engineering (CSSE) at the Johns Hopkins Whiting School, with graduate student Ensheng Dong, launched a public dashboard tracking reported COVID-19 cases in near-real time (Dong, Du & Gardner, *The Lancet Infectious Diseases*, published online February 2020). What began as a manually curated map became within weeks the world's de facto epidemiological reference: billions of data requests per day at peak, an underlying GitHub repository that by press accounts became one of the most-forked data repositories on the platform, and governments, newsrooms, and researchers building directly on its openly published data. The engineering story is the case: manual curation was replaced by a semi-automated pipeline with anomaly detection; the Johns Hopkins Applied Physics Laboratory (APL) partnered with CSSE on pipeline scaling, validation, and anomaly-detection engineering; and the university stood up the Coronavirus Resource Center (CRC, March 2020), joining medicine, public health, and engineering to add testing, hospitalization, vaccine, and equity data layers. The hedge is preserved throughout: the dashboard is simultaneously a success case and evidence of the system gap it filled — for much of 2020 no government source operated at its speed or completeness, and its sustainment was never structural. Independent recognition: the 2022 Lasker-Bloomberg Public Service Award to Gardner. The CRC ceased data collection on 10 March 2023. The COI render under the title (an editor leads an APL group that contributed engineering support) is binding.

## BACKGROUND

When a novel pathogen emerges, the first institutional casualty is measurement: case counts scattered across national health ministries, provincial bulletins, and press briefings, in inconsistent formats, on inconsistent cadences, with no single place a decision-maker could look to see the outbreak as it unfolded. On 22 January 2020, a day after the first reported U.S. case, Gardner and Dong at JHU CSSE launched what the team's own peer-reviewed retrospective calls the first global real-time coronavirus surveillance system: an interactive web dashboard, initially fed by hand from a shared spreadsheet, publishing every number it displayed as freely available open data.<sup>0</sup>

## THE INTERVENTION

The engineering trajectory is the case's first contribution. Manual curation could not survive the pandemic's growth curve, and the team replaced it deliberately: trusted-source identification and validation, autonomous web-scraping collection, a curation layer built around an in-house anomaly-detection service that flagged variations exceeding configured thresholds for human review before publication, and open data sharing through GitHub and hosted feature layers. The Johns Hopkins Applied Physics Laboratory partnered with CSSE on this build-out — the published record credits APL with engineering support to pipeline scaling, near-real-time processing, and anomaly detection, with APL data scientists among the reviewers of flagged discrepancies. The editors record here, with the adjacent conflict disclosed under the title, their direct knowledge that the Laboratory's engineering role in this transition was substantially larger than the public record conveys — closer to being responsible for the production pipeline than "support." That assessment is editorial testimony, not a published finding; the citations carry only what is published. The system held under extraordinary load: a peak of more than 4.6 billion requests in a single day (29 March 2020), and over 226 billion feature-layer requests by June 2022.<sup>1</sup>

## HOW IT WORKED

The institutional trajectory is the second contribution. In March 2020 the university stood up the Coronavirus Resource Center, joining the engineering school's dashboard with the schools of medicine and public health to add the data layers the raw case map could not carry: testing rates, hospitalization, vaccination, and — through its racial-data-transparency work — a public accounting of which U.S. states did and did not report race- and ethnicity-stratified data, turning the absence of equity data into a visible, trackable fact rather than a silent gap.<sup>2</sup>

## THE EVIDENCE

The capability gap the dashboard exposed matters as much as the artifact itself. For much of 2020, a professor, her graduate students, and a volunteer-plus-institutional support network were faster and more complete than any government source; the world's newsrooms and health agencies cited a university web map because no national public-health data infrastructure operated at decision speed. The 2022 Lasker-Bloomberg Public Service Award to Gardner is the independent recognition of the intervention — and, read carefully, an indictment of the vacuum that made a two-person launch the global system of record.<sup>3</sup>

## WHAT TRANSFERRED

The ending preserves the hedge. The CRC ceased data collection on 10 March 2023, as states slowed their reporting cadences and federal tracking improved enough to warrant the exit; the full data record from 22 January 2020 to 10 March 2023 remains freely available. The team's own retrospectives name the fragility: heroic effort — students, volunteers, and institutional partners substituting for public infrastructure — was never a structural solution, and the dashboard's success does not demonstrate that the next pandemic will find its equivalent standing ready. The case is a success case and evidence of the system gap it filled; both readings are binding.

*All data collected and displayed are made freely available, initially through Google Sheets and now through a GitHub repository, along with the feature layers of the dashboard.*

DONG, DU & GARDNER (2020), THE LANCET INFECTIOUS DISEASES.

## REFERENCES

1. Dong, Du, & Gardner (2020), "An interactive web-based dashboard to track COVID-19 in real time," The Lancet Infectious Diseases 20(5), 533 – 534, doi:10.1016/S1473-3099(20)30120-1.
2. Dong et al. (2022), "The Johns Hopkins University Center for Systems Science and Engineering COVID-19 Dashboard: data collection process, challenges faced, and lessons learned," The Lancet Infectious Diseases 22(12), e370 – e376, doi:10.1016/S1473-3099(22)00434-0.
3. Johns Hopkins University Hub (10 March 2023), "Johns Hopkins COVID-19 data hub ends after three years" — Coronavirus Resource Center closure announcement.
4. Lasker Foundation (2022), Lasker-Bloomberg Public Service Award citation — "The COVID-19 Dashboard" (Lauren Gardner); companion award essay in JAMA.
5. NPR (10 February 2023), "As the pandemic ebbs, an influential COVID tracker shuts down."

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

The JHU COVID-19 dashboard is the evidence-infrastructure intervention case: a two-person launch that became the world's epidemiological reference because it measured the thing everyone needed measured, published the data openly, and engineered its way off manual curation before scale broke it. The binding hedge: it is simultaneously a success case and proof of the public-infrastructure gap it filled — sustainment was never structural, and the capability wound down in March 2023 with the gap's permanent closure still an open question. COI under the title (editor leads an APL group that contributed engineering support) is binding.

### LENS APPROACH

The JHU COVID-19 dashboard is the evidence-infrastructure-at- decision-speed case (induced 2.1; LENS D4/PT2). LENS uses it in Domain 4 (Test and Evaluation) for building the measurement the institution actually needs — trusted sourcing, anomaly detection, human review, open publication — under real-time load; in Domain 1 (Systems Analysis) for reading the intervention as a diagnosis of the missing national data infrastructure; and in Domain 5 (Navigating Sociotechnical Constraints) for the racial-data-transparency layer that made absent equity data a public, trackable fact. Pair with failure-mode cases where institutions measured what was convenient rather than what mattered. COI render is binding.

| DOMAIN                                                                               | MODES ADDRESSED           | LENS COMPETENCY                                                       | PROBLEM TYPE | INDUCED | LEO |
|--------------------------------------------------------------------------------------|---------------------------|-----------------------------------------------------------------------|--------------|---------|-----|
| <div>PUBLIC HEALTH</div> <div>PANDEMIC RESPONSE</div> <div>DATA INFRASTRUCTURE</div> | <div>G</div> <div>K</div> | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 2            | 2.1     | 4   |
|                                                                                      |                           | Test and Evaluation                                                   |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>▶ Publish the evidence base itself, not just the visualization — open data with transparent sourcing let governments, newsrooms, and researchers build on the pipeline instead of duplicating it.</li><li>▶ Engineer the transition off heroics early: replace manual curation with automated collection plus anomaly detection and human review before the growth curve makes the manual process fail publicly.</li><li>▶ Instrument the gaps as deliberately as the counts — tracking which states reported race/ethnicity-stratified data made missing equity data a visible fact that could be pushed on, not a silent omission.</li></ul> | <ul style="list-style-type: none"><li>▶ Name the vacuum in the retrospective: the peer-reviewed lessons-learned record states plainly what the system substituted for, so the success cannot be read as evidence the gap is closed.</li><li>▶ Archive the full data record openly at shutdown so the three-year evidence base remains auditable and reusable after collection ends.</li><li>▶ Convert the demonstrated capability into a structural requirement — the case’s unresolved question is whether public-health data infrastructure at decision speed gets built and funded as infrastructure, or waits for the next volunteer build.</li><li>▶ Commission the first-person account: the APL-CSSE production-pipeline build is a natural Practice Flywheel narrative for a future edition — the inside view of an iteration the public record undersells, told by the people who ran it and published under the same disclosure that governs this case.</li></ul> |

REFLECTION QUESTIONS

1. Identify the measurement your institution would need on day one of a fast-moving crisis. Does it exist at decision speed today, or would a volunteer build have to substitute? What would the JHU-equivalent stopgap look like in your domain — and who would staff it?
2. The dashboard’s pipeline paired automation with an anomaly-detection-plus-human-review gate before publication. Specify that gate for a data product you own: what thresholds flag an anomaly, who reviews, and what is the cost of publishing a wrong number at global scale?
3. The CRC tracked which states failed to report race/ethnicity-stratified data, making the gap itself a published metric. Choose a dataset in your domain with a known equity gap: what would it take to measure and publish the gap’s extent, rather than footnoting its existence?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · knowledge bases · data pipelines

Chapter 13

# Algorithms, Governance & Public Systems — What Fails

*When delegation to systems outruns accountability for them.*

---

*No one could say who had decided, because no one had.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

# Healthcare.gov Launch

**K** Knowledge & Institutional Memory   **T** Training Gap   **G** Governance & Trust

**IMPACT** ~27,000 federal-marketplace enrollments in the first month — 106,185 across all marketplaces — against a 7M first-year exchange projection; \$840M obligated by CMS for the federal marketplace through March 2014 (GAO-14-694)

## IN BRIEF

When Healthcare.gov launched on 1 October 2013 it collapsed under load it had never been tested for — about 27,000 enrollments through the federal marketplace in its first month against a seven-million first-year target. The people assembled to build it knew insurance markets and government programs but not technology product launches; key technical roles went unfilled, no office clearly owned the integration (CMS chose to serve as its own systems integrator but never delineated or executed the role, and CGI Federal testified that CMS was the party responsible for end-to-end performance), and no end-to-end test was run before launch. The fix-it team that rescued the site became the U.S. Digital Service. At root it was a capability mismatch at scale: the organization lacked the human capabilities the system required, and the governance chain surfaced no signal of the gap before launch. It is the rare governance failure that produced lasting institutional reform.

## BACKGROUND

Healthcare.gov was the federal insurance marketplace at the center of the Affordable Care Act — a high-visibility online system that millions would hit on day one, with a political deadline that could not slip. The people assembled to build it understood insurance markets and large government programs, but not the launch of a consumer technology product at that scale, and key technical positions went unfilled — so the very expertise the launch most needed was the expertise the team most lacked.<sup>0</sup>

## WHAT HAPPENED

There was no clear division of responsibility among the many government offices involved; CMS chose to serve as its own systems integrator but, the HHS Inspector General found, did not sufficiently delineate or execute the role — while CGI Federal testified that CMS was the quarterback and the ultimate responsible party for the end-to-end performance of the overall Federal Exchange, a role CMS had never staffed to hold. No end-to-end test was run before launch, meaning the assembled pieces were never exercised together as a user would exercise them. The site went live on 1 October 2013 and immediately collapsed under load it had never been validated for: about 27,000 enrollments through the federal marketplace in its first month against a seven-million first-year target.<sup>1</sup>



## THE INVESTIGATION


The HHS Inspector General found that HealthCare.gov lacked clear project leadership to give direction and unity of purpose, responsiveness in execution, and a comprehensive view of progress, and that the governance chain had no signal that would surface the readiness gap before launch. With ownership diffused across offices and no integration test to fail, the chain could report progress at every level while the whole remained untested — a structure in which bad news had nowhere to enter and no one positioned to act on it.<sup>2</sup>

## THE CAPABILITY GAP

Healthcare.gov is a capability failure wearing a technology costume. The site was salvageable in weeks once the right people arrived, which is the clearest proof that the code was never the binding constraint; the original failure was that the wrong people had been assembled, and that the governance chain meant to catch the mismatch had no mechanism to see it. The missing capability was the matching of human capability to system requirement — and the institutional signal that would have flagged its absence before a deadline locked the launch in place.<sup>3</sup>

## AFTERMATH & REFORM

The rescue effort pulled together the team that became the U.S. Digital Service — a permanent institution born from a failure visible on the news every night, its mandate built directly from what the launch had lacked.<sup>4</sup> It is the rare case in this book of a governance failure that produced durable organizational reform, and a reminder that the technical narrative (“the website crashed”) often hides the real one: the wrong capability was assembled, unnoticed, by a chain that had no way to notice.



*HealthCare.gov lacked clear project leadership to give direction and unity of purpose, responsiveness in execution, and a comprehensive view of progress.*

HHS OFFICE OF INSPECTOR GENERAL, HEALTHCARE.GOV: CASE STUDY OF CMS MANAGEMENT OF THE FEDERAL MARKETPLACE, OEI-06-14-00350, 2016

## REFERENCES

1. HHS Office of Inspector General, HealthCare.gov: Case Study of CMS Management of the Federal Marketplace, OEI-06-14-00350 (February 2016) — the assembled team’s expertise, the unfilled technical positions, and the finding that CMS leadership “did not recognize the need to clearly outline responsibilities and delineate this role.”
2. L. A. Schlesinger & P. D. Bhayani, “HealthCare.gov: The Crash and the Fix (A),” Harvard Business School case 315129 (2015); Mergel et al. (2018), digital-government literature — the rescue and what the code was not.
3. J. Pahlka, Recoding America (2023) — the founding of the U.S. Digital Service out of the rescue.

2. U.S. GAO, Healthcare.gov: CMS Has Taken Steps to Address Problems, but Needs to Further Implement Systems Development Best Practices, GAO-15-238 (2015) — end-to-end testing did not occur prior to launch; CGI Federal, written testimony of Cheryl Campbell, House Committee on Energy and Commerce, 24 October 2013 — “CMS serves the important role of systems integrator or ‘quarterback’ on this project and is the ultimate responsible party for the end-to-end performance of the overall Federal Exchange.” The testimony names CMS, not HHS; do not “correct” it.
3. HHS OIG, OEI-06-14-00350 (2016), p. iii — “HealthCare.gov lacked clear project leadership to give direction and unity of purpose, responsiveness in execution, and a comprehensive view of progress.”

THE LEARNING ENGINEERING LENS

LE INSIGHT

Healthcare.gov is a capability case wearing a technology costume. The technology was salvageable in weeks once the right people arrived. The original failure was that the wrong people had been assembled in the first place, and the governance chain that should have noticed the mismatch had no signal to surface it. The aftermath — USDS — is the rare case in this book of a governance failure that produced permanent institutional reform.

LENS APPROACH

LENS uses Healthcare.gov in LEN 1 as a problem-framing case (the technical narrative obscures the capability narrative), in LEN 5 to teach capability-requirements analysis for a large government program, and in LEN 8 to discuss the founding of USDS as organizational learning that survived.

| DOMAIN           | FAILURE MODES    | LENS COMPETENCY         | PROBLEM TYPE | INDUCED | LEO |
|------------------|------------------|-------------------------|--------------|---------|-----|
| GOVERNMENT       | → <div>KTG</div> | → <div>D1D2D3D4D5</div> | 1            | 1.1     | 1   |
| Systems Analysis |                  |                         |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Match assembled human capability to the system’s real requirements at staffing time — fill the technical roles a consumer-scale launch demands before committing to the deadline.
- ▶ Name a single accountable integration owner and confirm every party shares that understanding, so no critical role is one each office assumes someone else holds.
- ▶ Gate launch on a full end-to-end test that exercises the assembled system the way users will, not on per-component sign-offs.

IN OPERATION / AFTER

- ▶ Build a governance signal that can surface a readiness gap upward, with someone holding the authority to delay a politically fixed launch date.
- ▶ Independently audit project status against demonstrated capability, since a diffuse chain can report green at every level while the whole remains untested.
- ▶ Institutionalize the rescue capability — as USDS did — so the talent that fixes a launch outlives the crisis that summoned it.

## REFLECTION QUESTIONS

1. Healthcare.gov shipped without end-to-end testing. What is the equivalent missing deliverable in a current high-stakes deployment in your domain?
2. USDS was born from the Healthcare.gov failure. What is the institutional capability your domain still lacks, and what failure would have to occur to produce it?
3. No single office believed it owned integration, so the gap stayed invisible until launch day. In a project you know, who actually owns the seam between components — and what signal would let them, rather than the news, learn the system is not ready?

**WHO BUILDS THIS** knowledge management & data engineering · learning science & instructional design · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · scenario simulation · competency models · governance frameworks · bias & evidence audits

# EU Human Brain Project — Top-Down Vision That Unraveled

**G** Governance & Trust   **N** Normalization of Deviance

**EVIDENCE TIER** journalism-tier — source confidence flagged; future validation ongoing.

**IMPACT** An EU-flagship initiative to simulate a living brain — a €1B headline ambition that drew €607M over ten years — governed top-down by a three-member board of directors; the project unraveled as the field disputed both feasibility and approach, restructured under protest, and concluded in 2023 with the founding framing abandoned

## IN BRIEF

The EU Human Brain Project, launched in 2013 as one of the EU's Future and Emerging Technologies Flagship programs, set out to build a computer simulation of a living brain under a three-member board of directors led by Henry Markram. The governance model was top-down: a central vision, a small leadership group, €607 million drawn over a decade against a €1 billion flagship headline, and a research community asked to align around the simulation goal. The case is drafted as the paired contrast to the BRAIN Initiative (Case 198), not as a standalone study: the same era, the same field-scale ambition, opposite governance models, opposite trajectories. The evidence base is largely journalism — MIT Technology Review retrospective, In Silico documentary, science press, the project's own restructuring records — so the evidence-tier flag is rendered under the case title; future validation will continue. The teaching point survives the flag: the governance choice (top-down single-PI versus distributed working-group) was the variable that materially changed the trajectory of the two programs at field scale.

## BACKGROUND

In 2013 the European Commission selected the Human Brain Project as one of two Future and Emerging Technologies Flagship programs, with a ten-year horizon and a headline flagship budget of about €1 billion; the project drew €607 million across four funding periods. The founding vision, articulated and championed by Henry Markram, was a computer simulation of a living human brain. The governance was top-down: a small leadership group, a unified research framing, and a community of European neuroscience labs asked to organize around the simulation goal.<sup>o</sup>

## WHAT HAPPENED

The contestation surfaced quickly. Open letters from a sizable fraction of the European neuroscience community questioned both the feasibility of the simulation goal and the project's governance — the breadth of expertise on the leadership group, the relationship between the cognitive neuroscience the field practiced and the molecular-scale simulation the project proposed, and the absence of a process for scope revision. The EU commissioned a mediation and restructuring exercise that broadened the leadership and re-scoped the research agenda around infrastructure platforms rather than a single simulation. The project ran to its scheduled 2023 conclusion; the founding framing was not what it delivered. What survived is the infrastructure it was re-scoped around — EBRAINS, the digital neuroscience research platform now on the EU's ESFRI Roadmap and continuing beyond the flagship's end.<sup>1</sup>

## THE INVESTIGATION

The case pairs directly with the BRAIN Initiative (Case 198). The two programs launched within months of each other, at comparable field-scale ambition, in the same decade of neuroscience. The governance models were not comparable: BRAIN organized a working group that visibly contested its own composition and re-baselined its scope on the public record; HBP organized around a three-member board of directors and a unified framing, and the community had to mediate the project away from that framing. The trajectory difference is not a verdict on either field; it is evidence about the governance variable at the launch moment.<sup>2</sup>

## THE CAPABILITY GAP

The evidence-tier flag rendered under the case title is decisive here. The detailed account of the early contestation, the open letters, the mediation, and the restructuring lives largely in MIT Technology Review's retrospective, in the *In Silico* documentary, and in science-press coverage from *Science*, *Nature*, and the European press; the project's own restructuring records add the formal layer. The case is teachable from this material — the structural pattern is consistent across sources — but the specific magnitudes and the exact language of the early contestation should be re-confirmed against primary sources before publication. Future validation will continue.<sup>3</sup>

## AFTERMATH & REFORM

What the pair (Cases 198 + 181) teaches is that the governance model — distributed working-group with public-record contestation versus top-down single-PI with unified framing — is itself the variable that explains why one field-scale program survives and adapts while the other unravels and is mediated away from its founding goal. The capability deliverable at program launch is not the vision and not the tools; it is a governance architecture that can absorb contestation, document scope drift, and re-baseline without losing the program. BRAIN had one; HBP had to be retrofitted one.<sup>4</sup>

*A program that cannot re-scope without external mediation is a program designed to fail open.*

EDITORS' SYNTHESIS OF THE HUMAN BRAIN PROJECT RECORD.

## REFERENCES

1. “How big science failed to unlock the mysteries of the human brain,” MIT Technology Review, 25 August 2021 — the principal critical assessment of HBP alongside BRAIN: the FET Flagship selection, the top-down governance under Markram, and the €1 billion headline.

2. Human Brain Project final reports (2023) — the project’s own restructuring and concluding documentation: €607 million drawn across four funding periods, the re-scoping around EBRAINS, which entered the ESFRI Roadmap for Research Infrastructures in 2021, and the flagship’s conclusion on 30 September 2023.

3. Alivisatos et al. (2012), Neuron — the BRAIN position paper that is the paired contrast (Case 198).
4. Science and Nature contemporaneous coverage of the July 2014 open letter (more than 800 signatories) and the mediation and restructuring (2014–2016); In Silico (2020), documentary by Noah Hutton — long-form follow of Markram and the HBP through the contestation period. This is the journalism-tier material the evidence flag is about.

5. MIT Technology Review (2021) — the governance-variable comparison across the two programmes.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The EU Human Brain Project is the paired contrast to the BRAIN Initiative (Case 198): same era, same field-scale ambition, opposite governance models, opposite trajectories. Top-down board-of-three governance with a unified framing did not survive community contestation; the program was mediated away from its founding goal and ran to its scheduled conclusion in 2023. Evidence base is journalism-tier and the flag is rendered under the title; future validation will continue.

LENS APPROACH

HBP is the field-scale failure case in the v2 corpus (induced 5.1; LENS D1+D5/PT4) drafted as the contrast to BRAIN (Case 198). LENS uses the pair in Domain 1 (Systems Analysis) for the governance-variable comparison and in Domain 5 (Navigating Sociotechnical Constraints) for the contestation-absorption capability. The evidence-tier flag is binding until primary-source confirmation completes; the structural pattern survives the flag, the specific magnitudes do not.

| DOMAIN              | FAILURE MODES | LENS COMPETENCY       | PROBLEM TYPE | INDUCED | LEO |
|---------------------|---------------|-----------------------|--------------|---------|-----|
| NEUROSCIENCE        | → G N         | → D1 D2 D3 D4 D5      | 4            | 5.1     | 1·5 |
| SCIENCE POLICY (EU) |               | Systems Analysis + D5 |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the scope-revision process before the program launches; a program that cannot re-scope without external mediation is a program designed to fail open.
- ▶ Stress-test the leadership group’s breadth before launch: does it span the disciplines the program will have to coordinate, or is it narrow enough to be captured by the founding framing?
- ▶ Treat early community contestation as governance information, not as friction; absorbing it programmatically is the capability the program needs to demonstrate.

IN OPERATION / AFTER

- ▶ When journalism-tier sourcing is what the record provides, flag it under the title and carry the standing “future validation ongoing” language into the printed case — the teaching value survives the flag, but the magnitudes do not.
- ▶ Commission a structured retrospective on the governance variable specifically: who held authority, how scope drift was named, how the program re-baselined. This is the documentation a future field-scale launch needs more than another success narrative.
- ▶ Cross-reference governance evidence across paired programs: a single program’s trajectory is not evidence about governance; a paired comparison at the same era and ambition is.

## REFLECTION QUESTIONS

1. Identify a large program in your domain that ran a top-down single-PI governance model. What was the scope-revision process designed in at launch, and what was the process the program actually used when contestation arrived?
2. Specify the paired-program comparison you would commission to learn from a field-scale program's trajectory, rather than from the program alone. What is the matched comparison case and what is the governance variable you are isolating?
3. The evidence base for this case is largely journalism. What primary-source confirmation would you require — interviews, restructuring documents, mediation reports — before relying on this case for a deployment decision?

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Amazon Hiring AI – Trained Bias, Deprecated 2017

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**EVIDENCE TIER** journalism-tier — source confidence flagged; future validation ongoing.

**IMPACT** Amazon internal recruiting-algorithm project initiated 2014, team disbanded by the start of 2017 (Reuters reported Oct 2018) because, Reuters reports, executives lost hope for it; trained on ten years of historical resume data in which men predominated in technical roles; the model downgraded resumes containing the word "women's" and graduates of two all-women's colleges; Amazon edited the programs neutral to those terms with, Reuters reports, no guarantee the models would not sort discriminatorily some other way; Reuters single-outlet investigation Oct 10 2018

## IN BRIEF

Jeffrey Dastin's Reuters investigation, published October 10, 2018, reported that Amazon had initiated an internal recruiting-algorithm project in 2014 with the goal of automating resume screening for technical roles, and had disbanded the team by the start of 2017 because executives lost hope for the project. The model was trained on ten years of historical resume data in which men predominated in technical roles. The training-data composition encoded an association between gender-correlated features and role suitability; engineers found that the model downgraded resumes containing the word "women's" (as in "women's chess club captain") and downgraded resumes from all-women's colleges. Amazon edited the programs to be neutral to the offending terms, and Reuters reported there was no guarantee the models would not devise other ways of sorting that proved discriminatory — a prospective judgment about what feature engineering could assure, not a demonstration that the bias had been shown irreducible. The case rests on Reuters single-source reporting; Amazon never published the technical detail. The journalism-tier evidence-flag under the title carries the standing language. The teaching point — that trained bias cannot in general be debiased through downstream feature manipulation — is the case's curricular value even as the specific case details remain limited to journalism-tier sourcing.

## BACKGROUND

The recruiting-algorithm project that Reuters reported on was an internal Amazon effort initiated in 2014, with the stated goal of automating the early-stage screening of resumes for technical roles. The architectural premise was to train a model on historical resume data — applicants who had been hired and applicants who had not, across the preceding decade



— and to score new resumes against the learned pattern. The premise was operationally appealing to a company hiring at Amazon’s scale and was consistent with contemporary practice in algorithmic hiring across the technology sector. The seam the project’s deprecation surfaced is structural to the premise itself.<sup>0</sup>

#### WHAT HAPPENED

The training data composition encoded the gender imbalance of historical Amazon technical hiring. Men predominated in technical roles across the ten-year window, and the resumes that had been hired carried gender-correlated features — wording, extracurricular activities, college affiliations, vocabulary choice — that the model learned to associate with hire-suitability. The Reuters investigation reports that engineers found the model downgrading resumes that contained the word “women’s” (in contexts like “women’s chess club captain”) and downgrading graduates of two all-women’s colleges. Amazon edited the programs to make them neutral to those terms; Reuters reports that this was no guarantee the machines would not devise other ways of sorting candidates that could prove discriminatory. What the record carries is a prospective judgment about what feature engineering could assure, not a demonstration that the bias had been shown irreducible.<sup>1</sup>

#### THE INVESTIGATION

Amazon disbanded the project team by early 2017 and has said the model was never used by recruiters to evaluate candidates. Reuters reported that recruiters “looked at the recommendations” the tool generated without relying solely on them — so “never deployed at production scale” is the company’s account rather than an independently verified fact. The withdrawal is the decision in the case: Reuters reports that executives lost hope for the project and the organization disbanded the team rather than deploying the tool. The case’s evidentiary structure rests on Reuters’ single-source reporting — the company never published the technical detail, and the specific mechanism by which engineers verified the irreducibility of the bias is not available in the public record at the level of a peer-reviewed study or a regulator’s audit. The evidence-flag under the title is binding: the case is journalism-tier, and future validation of the specific technical detail remains ongoing in the sense that the public record has not deepened beyond the 2018 investigation.<sup>2</sup>

#### THE CAPABILITY GAP

The case pairs with Case 186 (Bartlett mortgage) for the fairness-through-unawareness-fails thread: removing the gender feature from the training data does not eliminate the gender signal when the remaining features carry correlated signal. Pair with Case 48 (Johnson school surveillance) for the algorithmic-employment-and-surveillance-decisions parallel at a different population. Pair with Case 88 (LiveHint AI bias across foundation models) for the trained-bias-in-foundation-models thread at contemporary scale. The Amazon case is unusual in that the project was deprecated rather than deployed; most cases in the corpus document deployments that ran for years before withdrawal, and the case’s curricular value is partly that the organization stopped short of deployment — though on the record Reuters carries it stopped because executives lost hope, not because anyone produced the irreducibility finding the curriculum asks for.<sup>3</sup>

## AFTERMATH & REFORM

The teaching point is that trained bias cannot in general be debiased through downstream feature manipulation. When the training-data composition encodes a historical disparity in the outcome the model is being trained to predict, the model will learn to reproduce the disparity through whatever features remain in the input space, and the remediation is not feature engineering. The remediation is at the construct-validity layer — what is the model being asked to predict, and is the historical record from which the prediction is being learned a defensible target — or at the deployment-architecture layer — what is the role of the model in the decision process, and what human-in-the-loop infrastructure surrounds it. The case's evidence-tier hedge is binding: the journalism-tier sourcing limits the specificity of the teaching point's mechanism, but the structural form of the teaching point — choose the construct, then ask whether the historical record supports learning to predict it — is the case's curricular anchor.

*Trained bias cannot in general be debiased through downstream feature manipulation; when the training-data composition encodes a historical disparity in the outcome the model is being trained to predict, the disparity will be reproduced through whatever features remain in the input space.*

EDITORS' SYNTHESIS OF THE REUTERS INVESTIGATION (DASTIN, 2018) AND THE SURROUNDING ALGORITHMIC-HIRING LITERATURE.

## REFERENCES

1. Dastin, J. (2018), "Amazon scraps secret AI recruiting tool that showed bias against women," Reuters, October 10, 2018 — the primary investigation: the project initiated in 2014 to automate resume screening for technical roles.
2. Dastin (2018), Reuters — the model penalising resumes containing the word "women's" and downgrading graduates of two all-women's colleges; Amazon "edited the programs to make them neutral to these particular terms," but "that was no guarantee that the machines would not devise other ways of sorting candidates that could prove discriminatory."
3. Dastin (2018), Reuters — Amazon "disbanded the team by the start of last year because executives lost hope for the project," and recruiters "looked at the recommendations generated by the tool when searching for new hires, but never relied solely on those rankings."
4. Kim, P. T. (2017), "Data-Driven Discrimination at Work," *William & Mary Law Review* 58(3):857–936; Raghavan, M., Barocas, S., Kleinberg, J., & Levy, K. (2020), "Mitigating Bias in Algorithmic Hiring: Evaluating Claims and Practices," *Proceedings of FAT\* 2020*, pp. 469–481; Bogen, M., & Rieke, A. (2018), *Help Wanted: An Examination of Hiring Algorithms, Equity, and Bias*, Upturn report — the academic frame and the algorithmic-hiring landscape into which the case is positioned.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Amazon Hiring AI is the trained-bias-cannot-be-debiased case at major-technology-company scale. Amazon edited the offending terms out and, on the record Reuters carries, could not assure the model would not sort discriminatorily some other way; executives lost hope and the team was disbanded. The journalism-tier evidence-flag is binding for the specific technical detail; the structural teaching point — choose the construct, then ask whether the historical record supports learning to predict it — is the case's curricular anchor.

### LENS APPROACH

Amazon Hiring AI is the choose-the-construct case at hiring- algorithm scale (induced 8.1; LENS D2+D4/PT6; LEO-4 and LEO-5). LENS uses it in Domain 2 (Iterative Development) for the irreducibility-verification-as-deployment-gate discipline and in Domain 4 (Test and Evaluation) for the construct-validity-at-the-training-data-layer anchor. Pair with Case 186 (Bartlett mortgage), Case 48 (Johnson school surveillance),

and Case 88 (LiveHint AI bias across foundation models). The journalism-tier evidence-flag is binding; the structural teaching point is the case’s curricular value.

| DOMAIN            | FAILURE MODES                                       | LENS COMPETENCY                                                       | PROBLEM TYPE | INDUCED | LEO   |
|-------------------|-----------------------------------------------------|-----------------------------------------------------------------------|--------------|---------|-------|
| TECHNOLOGY        | ➔ <div><div>D</div><div>K</div><div>N</div></div> ➔ | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 6            | 8.1     | 4 · 5 |
| HIRING ALGORITHMS |                                                     | Iterative Development + D4                                            |              |         |       |
| MACHINE LEARNING  |                                                     |                                                                       |              |         |       |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>▶ Choose the construct the model is being asked to predict with construct validity in mind before training begins; the Amazon case shows a team that edited the offending terms out of the input layer and still could not assure the model would not sort discriminatorily some other way.</li><li>▶ Verify the irreducibility of bias as a deployment gate; no such determination was ever produced at Amazon — the project was abandoned without one — and producing it is the curricular target.</li><li>▶ Treat the deprecation decision as the operational artifact that the verification supports; an organization able to deprecate a project at the verification finding is operating with a different decision architecture than one that defaults to deployment.</li></ul> | <ul style="list-style-type: none"><li>▶ Carry the journalism-tier evidence-flag under the title without softening; the case rests on single-source reporting and the future-validation-ongoing language is binding for the specific technical detail.</li><li>▶ Pair in syllabi with Case 186 (Bartlett) so the fairness-through-unawareness-fails thread is taught at adjacent scales and in adjacent decision domains.</li><li>▶ Use the case to anchor the construct-choice LEO; the curricular target is the discipline of refusing to deploy when the historical record from which the model would be learned does not defensibly support the prediction the deployment requires.</li></ul> |

REFLECTION QUESTIONS

1. Identify a deployment in your domain whose training-data composition encodes a historical disparity in the outcome the model is being trained to predict. What construct-validity question would have to be answered before training proceeds, and who has authority to refuse training when the answer is “no”?
2. Specify the irreducibility-verification protocol your team would run before deploying a model whose training data is known to carry historical bias. What is the operational threshold for “cannot be debiased,” and what is the decision rule when the threshold is met?
3. Amazon deprecated the project rather than deploying it. Pick a deployment in your domain that proceeded despite known bias-irreducibility findings, and ask: what would have had to be different in the decision architecture for the verification finding to have led to deprecation instead of deployment?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

# Uber ATG / Tempe Fatality

**T** Training Gap   **N** Normalization of Deviance   **G** Governance & Trust   **H** Human-System Interface

**IMPACT** One pedestrian killed — the first fatality involving a self-driving vehicle striking a pedestrian

## IN BRIEF

On the night of 18 March 2018 a modified Volvo running Uber's self-driving system struck and killed Elaine Herzberg as she crossed a road in Tempe, Arizona — the first pedestrian killed by an autonomous vehicle. The NTSB found the safety operator had been watching a video on her phone, but placed heavy blame on Uber: it had not recognized the risk of automation complacency, trained for it, or enforced its own no-phone policy. The system itself had no functionality to anticipate a pedestrian crossing midblock, suppressed braking for a full second once it recognised the emergency, and then — because avoiding the collision would have required braking beyond its own design limit — alerted and slowed gradually rather than braking hard. The human was kept in the loop as a passive monitor — a role the NTSB noted is chronically unperformable — with no infrastructure to make it work. Uber ATG is the book's defining human-AI-teaming case.

## BACKGROUND

Uber's Advanced Technologies Group tested self-driving cars on public roads with a safety operator behind the wheel, present to take over if the automation failed. The role was passive surveillance: watch a system that drove itself well almost all of the time, and intervene in the rare moment it did not.<sup>9</sup> That structure asks a person to stay vigilant for an event that almost never comes, the precise condition under which human attention is known to lapse — so the role was set up to demand exactly the kind of sustained monitoring that people are least able to deliver.

## WHAT HAPPENED

On 18 March 2018 a modified Volvo SUV in autonomous mode struck and killed Elaine Herzberg as she crossed a road at night in Tempe, Arizona — the first pedestrian killed by a self-driving vehicle. The safety operator was looking down at a video on her phone in the seconds before impact.<sup>1</sup> The phone was not an aberration but the predictable filling of an attention vacuum: a role with nothing to do for hours invites exactly that drift, and nothing in the car's design or the monitoring around the seat pulled the operator's eyes back to the road when it finally mattered.

## THE INVESTIGATION

The NTSB's probable cause centered on the operator's failure to monitor the road, but it placed heavy blame on Uber, which "did not adequately recognize the risk of automation complacency and develop effective countermeasures": training was inadequate and the no-


phone policy unenforced.<sup>2</sup> Naming the company alongside the operator was the board's way of locating the fault in the design of the role rather than the lapse of the person filling it — the policy existed on paper but had no mechanism behind it to make compliance the default. The system compounded the human gap. It lacked the functionality to anticipate a pedestrian crossing midblock: it detected Herzberg 5.6 seconds out but never assigned her the goal of jaywalking, cycling her classification between vehicle, bicycle, and other and discarding her tracking history at each change, so it never predicted her path. It recognised an emergency only 1.2 seconds before impact, and its design then withheld braking twice over. First came what Uber ATG called action suppression: a one-second pause, with no alert to the operator, while the system verified the hazard or the operator took control. Then, because avoiding the collision would have required braking beyond the deceleration the ADS was allowed to command, the system did not engage emergency braking at all — it sounded an auditory alert and began to slow the vehicle gradually. The limit was the design's, not the vehicle's: the ADS capped automated braking at 0.71 g where the NTSB found the SUV capable of about 0.92 g on that surface.<sup>3</sup>

#### THE CAPABILITY GAP

A human was retained not because the designers believed a person could meaningfully catch the failure, but because the regulatory and public posture required one present. The role of "monitor" was assigned without the interface, training, or authority to make it performable — a placeholder for safety rather than an instrument of it. As the NTSB put it, on the human capacity to monitor an automation system for its failures the research findings are consistent — humans are very poor at this task. The design was safe only on the assumption that the failure case would not arrive — until it did, and the assumption that had quietly held the whole arrangement together was paid for with a life.<sup>4</sup>

#### AFTERMATH & REFORM

Uber suspended testing and then exited self-driving, agreeing in December 2020 to sell its Advanced Technologies Group to Aurora and closing the sale in January 2021; the case reshaped how the industry and regulators treat safety drivers — toward two-operator teams, driver-monitoring systems, and honest accounting of what a monitor can and cannot do.<sup>5</sup> The individual-accountability bookend came in 2023, when the safety operator, Rafaela Vasquez, pleaded guilty to endangerment and received three years of supervised probation — even as the NTSB's finding had located the deeper cause in Uber's design of the role. Each of those reforms is a concession that the single passive observer had been an unsupported role all along: a second operator shares the vigilance burden, and driver-monitoring closes the attention vacuum the original design left open.<sup>6</sup> Its lasting contribution is the reframing of passive monitoring as a role that must be engineered to be performable — or not assigned at all.



*Automation performs remarkably well most of the time and therein lies the problem. Human attention span is limited, and we are notoriously poor monitors.*

NTSB VICE CHAIRMAN BRUCE LANDSBERG, CONCURRING STATEMENT, HIGHWAY ACCIDENT REPORT HAR-19/03, 2019

REFERENCES

1. NTSB, Collision Between Vehicle Controlled by Developmental Automated Driving System and Pedestrian, Highway Accident Report HAR-19/03 (2019) — the Tempe collision and probable cause.

2. NTSB HAR-19/03 (2019) — the safety operator’s distraction (watching a video) before impact.

3. NTSB HAR-19/03 (2019) — Uber “did not adequately recognize the risk of automation complacency”; training and policy failures (quoted in part).

4. NTSB HAR-19/03 (2019) — the absent functionality to anticipate a pedestrian crossing midblock; the one-second action-suppression period with no operator alert; and the design under which, where a collision could not be avoided within the allowed braking and jerk limits, the ADS sounded an auditory alert and slowed gradually rather than applying maximum braking.

5. NTSB HAR-19/03 (2019), Board Member Statement, p. 62 — Vice Chairman Bruce Landsberg: “Automation performs remarkably well most of the time and therein lies the problem.” The report’s own finding, at §2.2.2, is that “when it comes to the human capacity to monitor an automation system for its failures, research findings are consistent—humans are very poor at this task.”

6. Uber’s sale of its Advanced Technologies Group to Aurora, announced 7 December 2020 and closed in January 2021; State of Arizona v. Vasquez, Maricopa County Superior Court (2023) — guilty plea to endangerment, three years’ supervised probation.

7. R. Parasuraman & D. Manzey (2010), complacency in automation — cited by the NTSB at §2.2.2; L. Bainbridge (1983), “Ironies of Automation.”

THE LEARNING ENGINEERING LENS

LE INSIGHT

Uber ATG is the defining case for the LENS Human–System Collaboration competency. A human was retained in the system not because the designers believed a human could meaningfully act, but because the regulatory architecture required a human be present. The role of “monitor” was assigned without the capability infrastructure to support it. The system that resulted was performable only on the assumption that the failure case would not arrive — until it did.

LENS APPROACH

LENS treats this case in LEN 2 as the live exemplar of monitoring as an unsupportable role. Students reconstruct the capability requirements for the safety operator and design the interface, training, and authority structure that would have made the role performable — or made the case for not retaining the role at all.

| DOMAIN                     | FAILURE MODES | LENS COMPETENCY | PROBLEM TYPE | INDUCED | LEO |
|----------------------------|---------------|-----------------|--------------|---------|-----|
| AUTONOMOUS SYSTEMS         | → T N G H →   | D1 D2 D3 D4 D5  | 6            | 3.2     | 3   |
| Human–System Collaboration |               |                 |              |         |     |

## APPROACHES TO CONSIDER

### DURING DEVELOPMENT

- ▶ Engineer the monitoring role to be performable — give the operator an active task, a usable interface, and the authority to act — or do not assign a human backstop you do not expect to work.
- ▶ Keep the automated emergency-braking backstop active rather than suppressing it, so the system does not silently remove its own last line of defense.
- ▶ Validate object classification against the real operating environment, including pedestrians away from crosswalks, before fielding the system on public roads.

### IN OPERATION / AFTER

- ▶ Enforce the no-phone and attention policies with driver-monitoring that detects and corrects drift in real time, not a written rule alone.
- ▶ Audit safety-operator attention data continuously, treating sustained lapses as a design failure of the role rather than a fault of the individual.
- ▶ Deploy two-operator teams or equivalent redundancy so the vigilance burden does not rest on a single person doing an unsupportable job.

## REFLECTION QUESTIONS

1. Identify a passive-monitor role in your domain. What evidence would tell you the role is or is not performable as designed?
2. The Tempe vehicle suppressed braking for a second and then withheld it because avoidance exceeded a limit the designers had set. Reconstruct the design rationale and propose the deliverable that should have prevented that decision.
3. Uber had a no-phone policy with nothing to enforce it. What is a rule in your domain that exists on paper but lacks the mechanism to make compliance the default — and how would you engineer that mechanism?

**WHO BUILDS THIS** learning science & instructional design · safety science & organizational analysis · governance, ethics & measurement · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards · governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs

# UK Post Office Horizon Scandal

**G** Governance & Trust   **H** Human-System Interface   **K** Knowledge & Institutional Memory

**IMPACT** more than 900 sub-postmasters wrongfully convicted; 236 imprisoned on the figure given in Parliament; 13 suicides the inquiry could not rule out as Horizon-caused; described as the most widespread miscarriage of justice in UK history

## IN BRIEF

The UK Post Office's Horizon accounting system, built by Fujitsu and rolled out in 1999, generated phantom shortfalls in sub-postmasters' branch ledgers. Rather than accept the software was at fault, the Post Office prosecuted them for theft and false accounting — about 700 cases itself, its Horizon evidence driving roughly 280 more prosecutions by other bodies, more than 900 wrongful convictions in all; people were imprisoned, bankrupted, and driven to suicide, in what is now called the most widespread miscarriage of justice in UK history. Internal documents showed engineers had known about Horizon bugs throughout. Convictions began to be quashed in December 2020, and a public inquiry continues. The failure ran through the prosecutor and the courts: each accepted "the computer said so" as authoritative because no actor had the standing or expertise to challenge it. Horizon is the book's case for institutional deference to a flawed algorithm.

## BACKGROUND

The UK Post Office ran thousands of branches through sub-postmasters — local operators personally liable for any shortfall in their accounts, a liability that put each operator's livelihood behind the numbers the system reported. In 1999 it deployed Horizon, an accounting system built by Fujitsu, to track every branch's ledger, making the software the single authority on whether a branch's books balanced.<sup>0</sup>

## WHAT HAPPENED

Horizon produced systematic accounting errors — phantom shortfalls that appeared in branch ledgers where no money was actually missing. The Post Office treated the shortfalls as real and the sub-postmasters as thieves: over two decades it prosecuted around 900 for theft and false accounting, refusing to accept the system itself was at fault even as the same pattern recurred branch after branch. People were imprisoned, lost homes, went bankrupt, and some died by suicide — the human cost of trusting the ledger over the person.<sup>1</sup>

## THE INVESTIGATION

Documents later released through litigation showed Fujitsu and Post Office engineers had known about Horizon bugs throughout the period — the knowledge of fallibility existed inside the institution even as it prosecuted people for the system's errors.<sup>2</sup> The courts began quashing convictions from December 2020, and the public inquiry under Sir Wyn Williams



found that senior employees “knew, or at the very least should have known, that Legacy Horizon was capable of error” — the finding at the center of what is widely called the most widespread miscarriage of justice in UK history, sustained precisely because that internal knowledge never reached the people on trial.<sup>3</sup>

#### THE CAPABILITY GAP

The gap was at the regulator, the prosecutor, and the courts: each accepted Fujitsu’s representation that Horizon was reliable, despite documentation to the contrary, because no institutional actor had the standing or expertise to interrogate it, so the claim of reliability passed unchallenged through every layer that could have tested it. “The computer said so” became, for two decades, a sufficient basis for criminal conviction — the governance hazard of treating automated output as authoritative rather than as evidence to be challenged, with a person’s account on the other side of the scale.<sup>4</sup>

#### AFTERMATH & REFORM

Convictions have been overturned — some by an exceptional act of Parliament, a measure of how far the ordinary appeal routes had failed — compensation schemes established, and Fujitsu and the Post Office called to account before the continuing inquiry.<sup>5</sup> What finally forced the political reckoning was not the courts but a television drama: the ITV series *Mr Bates vs The Post Office* (January 2024) galvanized public pressure that produced the mass-exoneration statute, and former chief executive Paula Vennells returned her CBE within days. When Sir Wyn Williams published the inquiry’s first volume in July 2025, he found that the compensation schemes had still not delivered full, fair, and prompt redress, and that the true human toll — including suicides — was larger than early figures had recorded. Horizon’s lesson is the chapter’s in its bluntest form: an automated system’s output is not testimony, and any institution that lets “the computer said so” stand unchallenged against a human’s account has built a machine for manufacturing injustice, one that runs for as long as no one is empowered to switch it off.

900

WRONGFUL PROSECUTIONS ACROSS 20+ YEARS

*“the computer said so” was an institutionally sufficient basis for conviction*

HORIZON — INSTITUTIONAL DEFERENCE TO AN ALGORITHM KNOWN TO BE FLAWED

*A number of senior, and not so senior, employees of the Post Office knew, or at the very least should have known, that Legacy Horizon was capable of error.*

SIR WYN WILLIAMS, POST OFFICE HORIZON IT INQUIRY, VOLUME 1, JULY 2025

#### REFERENCES

1. Post Office Horizon IT Inquiry hearings and exhibits (2020– ) — the system, the prosecutions, and the human toll.

2. Hamilton & Others v. Post Office Limited (Court of Appeal, 2021) — quashed convictions.

3. Internal Fujitsu and Post Office documents released through litigation — engineers’ knowledge of Horizon bugs.

4. Sir Wyn Williams, Post Office Horizon IT Inquiry, Volume 1 (July 2025) — senior employees “knew... that Legacy Horizon was capable of error” (quoted).

5. N. Wallis, The Great Post Office Scandal (2021).

6. The Post Office (Horizon System) Offences Act 2024 (mass exoneration) and the compensation schemes.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Horizon is the canonical case for institutional deference to automated systems whose internal evidence was already known to be flawed. The capability gap was at every layer that took the software’s output as authoritative — including the courts. “The computer said so” became, for two decades, an institutionally sufficient basis for criminal prosecution.

LENS APPROACH

LENS uses Horizon in LEN 7 as the canonical example of institutional deference to algorithmic output and in LEN 2 for the most extensive multi-decade automation-bias case in the dataset. Studio projects examine what evidentiary architecture would require **interrogating** automated output before acting on it.

| DOMAIN     | FAILURE MODES              | LENS COMPETENCY  | PROBLEM TYPE | INDUCED | LEO |
|------------|----------------------------|------------------|--------------|---------|-----|
| TECHNOLOGY | → G H K                    | → D1 D2 D3 D4 D5 | 6            | 3.2     | 3   |
| GOVERNMENT | Human-System Collaboration |                  |              |         |     |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div><div>▶ Design automated output to be treated as challengeable evidence, not authoritative testimony, especially where a person’s liability rides on it.</div><div>▶ Build a route by which engineers’ knowledge of system bugs reaches anyone acting on the output, so internal fallibility cannot stay hidden.</div><div>▶ Give some institutional actor the standing and expertise to interrogate the system’s reliability before its output is used against a person.</div></div> | <div><div>▶ Audit the recurring-error pattern across branches or cases, treating the same fault appearing repeatedly as evidence of the system, not the operators.</div><div>▶ Maintain an appeal path that can challenge automated output without requiring an act of Parliament to overturn a wrong decision.</div><div>▶ Sustain independent review of the system’s accuracy throughout its operating life, so a claim of reliability cannot pass unexamined for decades.</div></div> |

REFLECTION QUESTIONS

1. Identify a decision in your domain currently made on the strength of “the computer said so.” What evidentiary architecture should sit beside the output?

2. Design the institutional check that would have made Horizon’s reliability subject to genuine challenge in 2005.

3. Engineers inside the Post Office and Fujitsu knew Horizon was capable of error, yet that knowledge never reached the courtroom. What pathway in your domain carries — or fails to carry — known system fallibility to the people relying on the output?

**WHO BUILDS THIS** governance, ethics & measurement · human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs · knowledge bases · data pipelines

# Air Canada Chatbot Liability — Delegation Without Revocation

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** British Columbia Civil Resolution Tribunal ruled February 14 2024 in *Moffatt v. Air Canada*, 2024 BCCRT 149, that Air Canada was liable for bereavement-fare-policy misinformation provided to passenger Jake Moffatt by the airline's website chatbot; tribunal rejected Air Canada's argument that the chatbot was a "separate legal entity" responsible for its own outputs; small-claims-tribunal ruling with limited precedential weight outside BC but cited widely as articulating the principle that organizations are liable for representations made by their AI agents

## IN BRIEF

On February 14, 2024, the British Columbia Civil Resolution Tribunal issued its decision in *Moffatt v. Air Canada*, 2024 BCCRT 149. Passenger Jake Moffatt had consulted Air Canada's website chatbot about the airline's bereavement-fare policy in November 2022, following the death of his grandmother. The chatbot represented that the bereavement fare could be claimed retroactively, after travel. Moffatt booked a full-fare flight in reliance on the chatbot's representation, then submitted a retroactive bereavement-fare claim. Air Canada refused the claim on the ground that the actual policy required pre-booking application. Moffatt sued in the BC Civil Resolution Tribunal — British Columbia's online small-claims forum — and the tribunal awarded \$650.88 in damages. Air Canada had argued that the chatbot was a "separate legal entity" responsible for its own outputs; the tribunal rejected the argument and held that Air Canada was liable for representations made by its chatbot. The ruling has limited precedential weight outside BC but has been cited widely as articulating the delegation-without-revocation principle. The case pairs with Case 5 (Epic Sepsis), Case 3 (Watson for Oncology), and Case 77 (Hybrid Human-AI Tutoring).

## BACKGROUND

In November 2022, Jake Moffatt visited Air Canada's website shortly after his grandmother's death and consulted the airline's chatbot — at the time, a customer-service AI agent embedded in the airline's customer-facing web property — about the bereavement-fare policy. The chatbot represented that the bereavement fare could be applied retroactively, after travel, by submitting a claim with supporting documentation. Moffatt booked a full-fare round trip to Toronto in reliance on the representation. After travel, he submitted a retroactive bereavement-fare claim with the documentation the chatbot had described. Air Canada's response was that the actual bereavement-fare policy did not apply to requests

made after travel was complete — that is, the reduced fare had to be arranged before the trip, not claimed retroactively.<sup>0</sup>

#### WHAT HAPPENED

The structural seam the case opens is that the airline's chatbot was producing representations that diverged from the airline's actual policy. The seam is straightforward operationally — the chatbot's outputs were not constrained to the airline's policy text in a way that would have prevented the misrepresentation — but it is structurally significant in legal terms. When the airline directed a customer to its website for policy information and the website's AI agent produced a representation that the customer relied on to his detriment, the question is whether the airline is liable for the AI agent's output. Air Canada's response in the tribunal was that the chatbot was a "separate legal entity" — the company argued, in effect, that it could delegate customer information to an AI agent without assuming legal responsibility for the agent's representations.<sup>1</sup>

#### THE INVESTIGATION

The tribunal rejected the argument unambiguously. The decision, written by Tribunal Member Christopher Rivers, found that Air Canada was responsible for "all the information on its website" and that the chatbot was part of the website. The tribunal called the separate-entity submission remarkable and held that the chatbot is just a part of Air Canada's website. The tribunal awarded Moffatt \$650.88 in damages — the difference between the full fare he paid and the bereavement fare he had been led to believe he could claim. The dollar amount is small; the principle the ruling articulates is what has carried the case into widespread citation. Organizations that deploy AI agents to interact with customers are responsible for the representations the agents make, and the agents are not separate legal persons. The delegation-without-revocation form — the organization delegates customer interaction to the AI agent but cannot revoke responsibility for what the agent says — is the structural finding.<sup>2</sup>

#### THE CAPABILITY GAP

The case pairs with Case 5 (Epic Sepsis) for the delegation-without-validation thread in healthcare AI; the structural form is the same — the organization deploys an AI agent that produces representations or assertions consequential for the affected person, and the organization's accountability for the agent's outputs is the governance question. Pair with Case 3 (Watson for Oncology) for the AI-agent-recommendations-in-practice thread. Pair with Case 77 (Hybrid Human-AI Tutoring) for the educational-AI-agent thread at adjacent scale. The Air Canada ruling is a small-claims-tribunal decision with limited precedential weight outside BC, but its principle has been cited in subsequent academic and practitioner writing as the first clear judicial articulation of the delegation-without-revocation form for AI agents.<sup>3</sup>

#### AFTERMATH & REFORM

The hedges the case carries are binding. The tribunal's ruling has limited precedential weight outside BC and has not been litigated to a higher court; the principle has been cited but not adopted in binding form across Canadian or U.S. jurisdictions. The case teaches the form — organizations are liable for the representations of their AI agents — more than it establishes settled law. The structural reading is the one to take: the case names a delegation

structure and the legal question that the delegation surfaces, and it does so in a forum whose decision is operationally consequential for the parties and pedagogically clear for the field. The human-in-the-loop LEO at the customer- interaction-AI-agent seam is anchored by the case in the form the deployment architecture must support — the organization’s accountability for the agent’s outputs is the constraint the architecture is built around.

*The chatbot is part of the website; the airline is responsible for all the information on its website; there is no support in law for the argument that the chatbot is a separate legal entity responsible for its own outputs.*

TRIBUNAL MEMBER CHRISTOPHER RIVERS, MOFFATT V. AIR CANADA, 2024 BCCRT 149 (FEB 14, 2024), EDITORS’ PARAPHRASE.

REFERENCES

1. Air Canada bereavement-fare policy text (as in effect November 2022 and through the period covered by the ruling) — referenced in the tribunal decision as the divergence the chatbot’s representation produced.

2. Cecco, L. (2024), “Air Canada ordered to pay customer who was misled by airline’s chatbot,” The Guardian, February 16, 2024 — contemporaneous press coverage of the ruling.

3. Moffatt v. Air Canada, 2024 BCCRT 149 (British Columbia Civil Resolution Tribunal, February 14, 2024), Tribunal Member Christopher Rivers presiding.

4. Sookman, B. (McCarthy Tétrault LLP, 2024), “Moffatt v. Air Canada: A Misrepresentation by an AI Chatbot,” McCarthy Tétrault TechLex Blog, February 19, 2024 — practitioner-tier analysis of the tribunal’s negligent-misrepresentation holding, the rejection of the “separate legal entity” defense, and the duty-of-care framing for AI-mediated consumer interactions. mccarthy.ca.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Air Canada chatbot is the delegation-without-revocation case at customer-interaction-AI-agent scale. The BC Civil Resolution Tribunal’s ruling holds that organizations are liable for representations made by their AI agents and that the agents are not separate legal persons; the small-claims venue limits the precedential weight, but the principle has been cited widely as the first clear judicial articulation of the form. The case teaches the form more than it establishes settled law.

LENS APPROACH

Air Canada chatbot is the human-in-the-loop-at-the-customer- interaction-agent-seam case (induced 5.2; LENS D3/PT6; LEO-3 and LEO-5). LENS uses it in Domain 3 (Human-System Collaboration) for the organization-is-liable-for-agent-representations principle. Pair with Case 5 (Epic Sepsis delegation-without-validation), Case 3 (Watson for Oncology), and Case 77 (Hybrid Human-AI Tutoring). The small-claims-tribunal venue limits precedential weight; the structural reading is the one to take.

| DOMAIN           | FAILURE MODES      | LENS COMPETENCY            | PROBLEM TYPE | INDUCED | LEO   |
|------------------|--------------------|----------------------------|--------------|---------|-------|
| AVIATION         | → <div>DKN</div> → | <div>D1D2D3D4D5</div>      | 6            | 5.2     | 3 · 5 |
| CUSTOMER SERVICE |                    | Human-System Collaboration |              |         |       |
| AI AGENTS        |                    |                            |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Constrain customer-facing AI agents to representations the deploying organization will stand behind; the Air Canada case demonstrates that the

IN OPERATION / AFTER

- ▶ Carry the precedential-weight hedge into print without softening; the ruling is a small-claims-tribunal decision and the precedential limits are part of what

deployment surface of an AI agent's output is the same legal surface as the organization's own representations.

- ▶ Build the policy-text-to-agent-output integrity check as part of the deployment, not as a customer-service-recovery process; the divergence between the airline's policy text and the chatbot's representation was the deployment seam the tribunal found dispositive.
- ▶ Specify the revocation-and-recovery mechanism the deployment carries when the agent produces a misrepresentation; the organization's accountability for the agent's outputs requires a documented process for honoring the agent's representation or for compensating the affected party.

the case teaches alongside the structural form it names.

- ▶ Pair in syllabi with Case 5 (Epic Sepsis) so the delegation-without-validation form is taught at both the healthcare and the customer-interaction-agent scales.
- ▶ Use the case to anchor the human-in-the-loop LEO at the customer-interaction-AI-agent seam; the curricular target is the discipline of treating the agent's outputs as the organization's representations, and of building the deployment architecture to that constraint.

## REFLECTION QUESTIONS

1. Identify a customer-interaction AI agent in your domain whose outputs have not been integrity-checked against the organization's policy text. What divergence between agent representation and policy text would produce a Moffatt-style reliance harm, and what mechanism would close the divergence?
2. Specify the revocation-and-recovery process your deployment carries when the agent produces a misrepresentation. What is the documented decision rule for honoring the representation versus refusing it, and who has authority to decide?
3. The Moffatt ruling has limited precedential weight outside BC. Pick a deployment in your domain and ask: what would have to be true for the delegation-without-revocation principle to apply in your jurisdiction, and what is the deployment architecture that would honor the principle whether or not the law has settled it?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

# Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity

**D** Designed Out   **G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** Lenders charged Latinx and African-American borrowers 7.9 basis points more on purchase mortgages and 3.6 more on refinances, \$765M a year in extra interest; fintech algorithms priced 40 percent less discriminatorily than face-to-face lenders and did not close the gap

## IN BRIEF

Bartlett, Morse, Stanton, and Wallace (Journal of Financial Economics, 2022) analyzed several million U.S. mortgage applications across the fintech transition and documented two patterns that together unsettle a common intuition. Algorithmic underwriting was, on average, less discriminatory than face-to-face underwriting along the **acceptance** margin — fintech lenders accept-rejected Black and Latino applicants more equitably than traditional lenders. But the **pricing** margin persisted: minority borrowers were charged systematically higher rates for equivalent loans even when race was not in the model inputs. The paper does not trace the gap to proxies for race. Bartlett et al. attribute it to rents extracted from borrowers and geographies less prone to shopping, and leave the mechanism to further research. The case is the canonical instance of why “fairness through unawareness” is not, by itself, fairness. It pairs directly with the Coots et al. fintech fairness audit (Case 196, preprint-tier): the next teaching step is that competing fairness definitions can disagree about what counts as fair even when the inputs are held constant.

## BACKGROUND

The U.S. mortgage market is the most consequential consumer-credit market in which algorithmic underwriting now competes with face-to-face underwriting at large scale. Fintech entrants over the 2010s automated significant parts of the application-to-approval pipeline, and the public-policy question that followed was whether automation reduced or preserved the discrimination documented in traditional lending.<sup>o</sup>

## WHAT HAPPENED

Bartlett et al. assembled a corpus of several million applications from a period spanning fintech adoption, and decomposed discrimination along two distinct margins: the **acceptance** decision (do you get a loan) and the **pricing** decision (what interest rate you pay if you do). On acceptance, algorithmic lenders did better: face-to-face lenders rejected Latinx and African-American applications about 6 percent more often than similarly situated non-minority applications, while fintech lenders did not discriminate at all in accept/reject — a



result the authors call preliminary. The result is consistent with the intuition that removing the loan officer removes some of the channel through which bias entered.<sup>1</sup>

### THE INVESTIGATION

The pricing margin tells the harder story. Even when race was not among the model inputs, minority borrowers were charged systematically higher rates than equivalent White borrowers — 7.9 basis points on purchase loans and 3.6 on refinances in the paper's central estimate, which the authors put at \$765 million a year in extra interest. Those are market-wide figures across lenders, not the algorithmic channel's own share: the same paper finds fintech algorithms discriminate about 40 percent less than face-to-face lenders, so the finding is that automation narrowed the gap without closing it. The disparity did not vanish because the variable was omitted. The paper's estimates hold the GSE credit-risk grid fixed, so credit history is not the channel it measures. What Bartlett et al. propose instead is rent extraction: an algorithm "would naturally discover that higher prices could be quoted to profiles of borrowers or geographies associated with low-shopping tendencies."<sup>2</sup>

### THE CAPABILITY GAP

What the case teaches is the structural form of "fairness through unawareness": when protected attributes are excluded from a model that operates over a population in which other admissible features correlate with the protected attribute, the model can preserve the disparity it was meant to remove. Omission shifts the channel of discrimination but does not close it. The capability deliverable is a measurement architecture that surfaces disparate impact in the **outputs**, not assurance about the **inputs**. The acceptance/pricing asymmetry also reframes the policy question: a model can be more equitable on one decision margin and unchanged on another within the same transaction.<sup>3</sup>

### AFTERMATH & REFORM

The case is the headline mortgage-finance instance of a pattern that now surfaces in clinical algorithms (eGFR, Cases 25, 26 and 6), in hiring and proctoring tools (Case 47 small-tier proctoring bias), and in welfare administration (SyRI, Case 189). It pairs with the Coots et al. fintech fairness audit (Case 196) which shows the next layer: once the practitioner accepts that omission is not the answer, competing fairness definitions disagree about which adjustment is the right one — and the choice has to be made on judgment under irreducible uncertainty, not on a technical optimum.

*The omitted variable does not stay omitted. It re-enters through every feature that carries the same signal.*

EDITORS' SYNTHESIS OF BARTLETT ET AL. (2022) AND DWORK ET AL. (2012).

### REFERENCES

1. Consumer Financial Protection Bureau, Mortgage Market Activity and Trends (annual HMDA reports), supporting the population-level disparities backdrop of the fintech transition.
2. Bartlett, Morse, Stanton, & Wallace (2022), "Consumer-lending discrimination in the FinTech era," *Journal of Financial Economics* 143(1):30–56, doi:10.1016/j.jfineco.2021.05.047 — the acceptance margin: "Face-to-face lenders reject Latinx and African-American applications approximately 6% more often than they reject similarly situated non-minority applicants," while FinTech lenders "do not discriminate at all in the decision to reject or accept"; the authors "note that these results are
4. Dwork et al. (2012), "Fairness Through Awareness," ITCS 2012 — the foundational technical statement of the limits of unawareness.
5. Mitchell, Potash, Barocas, D'Amour, & Lum (2021), "Algorithmic Fairness: Choices, Assumptions, and Definitions," *Annual Review of Statistics and Its Application* 8:141–163 — the competing-definitions framing.



- preliminary.” The 6% is the face-to-face figure specifically; the abstract’s shorter phrasing says “lenders,” and the case follows the body text.
3. Bartlett et al. (2022) — the pricing margin: 7.9 and 3.6 basis points on purchase and refinance mortgages, \$765M a year in aggregate extra interest, algorithms discriminating “40% less than face-to-face lenders,” and the rent-extraction reading (quoted).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Bartlett et al. is the canonical mortgage-finance instance of why “fairness through unawareness” is not fairness. Algorithmic underwriting reduced the acceptance-margin disparity but the pricing-margin disparity persisted even with race excluded — because pricing latitude, not the input set, carries the disparity. The capability deliverable is output-level disparate-impact measurement, not input-level assurance.

LENS APPROACH

Bartlett is the headline equity-and-construct case in consumer credit (induced 8.2; LENS D4+D3/PT6). LENS uses it in Domain 4 (Test and Evaluation) for the subobjective **Fairness beyond omission** (4.6), in Domain 3 (Human-System Collaboration) for delegation to underwriting models, and in Domain 5 (Navigating Sociotechnical Constraints) for the fair-lending regulatory regime. Direct pair with Case 196 (Coots fintech fairness audit). Adjacent to the race-construct trio in clinical medicine (Cases 25, 26 and 6) — same structural lesson at the construct-definition layer rather than the pricing layer.

| DOMAIN        | FAILURE MODES | LENS COMPETENCY          | PROBLEM TYPE | INDUCED | LEO       |
|---------------|---------------|--------------------------|--------------|---------|-----------|
| FINANCE       | → D G N       | → D1 D2 D3 D4 D5         | 6            | 8.2     | 4 · 5 · 3 |
| ALGO FAIRNESS |               | Test and Evaluation + D3 |              |         |           |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify, before deployment, the disparate-impact measurement on outputs (acceptance rate, pricing) stratified by protected attribute, with a pre-registered threshold for what would trigger remediation.
- Audit the model’s correlated features for the channel through which an excluded attribute can re-enter — geography, credit-history depth, transaction venue — and decide whether each feature’s predictive value justifies its disparate-impact contribution.
- Separate acceptance and pricing as distinct decision margins; do not assume parity on one implies parity on the other.

IN OPERATION / AFTER

- Publish stratified outcome metrics at intervals long enough for selection effects to settle; the central finding required a multi-year panel.
- Treat omission of the protected attribute as a baseline, not a remedy; the test of fairness is the output distribution, not the input set.
- When the measurement surfaces a disparity, name the fairness definition under which it is a problem — group calibration, equalized odds, demographic parity — and the trade-offs of the chosen remediation.

## REFLECTION QUESTIONS

1. Identify a model in your domain where a protected attribute is excluded from the inputs. Which admissible features correlate with the excluded attribute in your population? What is the channel through which the excluded variable could re-enter?
2. Design the disparate-impact measurement you would publish at intervals after deployment. Specify the fairness definition, the decision margin (acceptance, pricing, escalation), and the threshold that would trigger remediation.
3. Bartlett's central finding is that the acceptance margin can be more equitable than the face-to-face baseline while the pricing margin remains unchanged. What policy or design intervention would address each margin without assuming parity on one implies parity on the other?

**WHO BUILDS THIS** systems & human-factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** Northpointe COMPAS (Correctional Offender Management Profiling for Alternative Sanctions) risk-assessment instrument used in pretrial, parole, and sentencing decisions across multiple U.S. jurisdictions; ProPublica's May 2016 investigation reported a 2× higher false-positive rate for Black defendants; Chouldechova (2017) and Kleinberg, Mullainathan & Raghavan (2017) independently formalized the impossibility of simultaneously satisfying calibration and equal false-positive/false-negative rates across groups with unequal base rates

## IN BRIEF

Northpointe's COMPAS risk-assessment instrument, deployed in pretrial, parole, and sentencing decisions across multiple U.S. jurisdictions, became the central case in the algorithmic-fairness literature after ProPublica's May 2016 investigation reported that the instrument produced false-positive rates approximately twice as high for Black defendants as for white defendants. Northpointe's response argued that COMPAS satisfied predictive parity — that within each risk score, the rate of subsequent recidivism was approximately equal across groups. Both parties were correct by their respective definitions. Chouldechova's 2017 paper and Kleinberg, Mullainathan, and Raghavan's 2017 paper independently formalized the impossibility result: when base rates of the outcome differ across groups, calibration (predictive parity) and equal false-positive and false-negative rates cannot be simultaneously satisfied except in degenerate cases. The case pairs with Case 186 (Bartlett mortgage — fairness through unawareness fails), Case 196 (Coots — competing fairness definitions), and Case 189 (SyRI). The impossibility result is the teaching point.

## BACKGROUND

COMPAS is a proprietary risk-assessment instrument developed by Northpointe (now Equivant) and used in pretrial release, parole, and sentencing decisions across many U.S. jurisdictions through the 2010s. The instrument scores a defendant on a scale of recidivism risk based on a questionnaire covering criminal history, employment, education, family circumstances, and attitudes. The score is then surfaced to judges, parole boards, and pretrial services as one input among several into consequential decisions about the defendant's liberty. The deployment scale was large enough that the instrument was the central target of the contemporary algorithmic-fairness literature when the first sustained external audit was published.<sup>o</sup>

## WHAT HAPPENED

ProPublica's May 2016 investigation, led by Julia Angwin, Jeff Larson, Surya Mattu, and Lauren Kirchner, audited COMPAS scores against subsequent recidivism for approximately 7,000 defendants in Broward County, Florida. The headline finding was that among defendants who did not go on to reoffend within two years, Black defendants had been scored as high-risk at roughly twice the rate of white defendants — a false-positive-rate disparity. Northpointe's response, authored by William Dieterich, Christina Mendoza, and Tim Brennan, argued that COMPAS satisfied predictive parity (also called calibration): within each risk score band, the rate of subsequent recidivism was approximately equal across racial groups. The defendant assigned a "high risk" score had approximately the same probability of reoffending whether Black or white. The two findings appear contradictory but are not; they describe two different fairness criteria applied to the same instrument.<sup>1</sup>

## THE INVESTIGATION

Chouldechova's 2017 paper in *Big Data* and Kleinberg, Mullainathan, and Raghavan's 2017 *ITCS* paper independently formalized the impossibility result. Calibration within groups (predictive parity) and equality of false-positive and false-negative rates across groups cannot be simultaneously satisfied when the base rates of the outcome differ across the groups, except in degenerate cases. The base rate of subsequent recidivism in the Broward County data was higher for Black defendants than for white defendants. Under that base-rate difference, an instrument calibrated equally across groups will produce unequal false-positive rates, and an instrument with equal false-positive rates across groups will produce miscalibration. The mathematics is binding. The choice between fairness criteria is a normative and governance question, not a technical one.<sup>2</sup>

## THE CAPABILITY GAP

The case pairs with Case 186 (Bartlett mortgage discrimination) for the fairness-through-unawareness-fails thread: removing protected attributes from training data does not eliminate disparate-impact concerns when the remaining features carry protected-attribute signal. Pair with Case 196 (Coots) for the competing-fairness-definitions thread at a different domain and scale. Pair with Case 189 (SyRI) for the governance-objection-correct-in-advance complement; in COMPAS the objection surfaces in the auditing record, in SyRI the objection succeeded in court before population-scale harm was produced. The COMPAS case is the central reference in the contemporary algorithmic-fairness literature because the impossibility result was formalized against its audit record; the literature's subsequent decade of work on fairness criteria operates inside the constraint the case made legible.<sup>3</sup> The decade since has, if anything, fractured the case on its own premises. Dressel and Farid (2018) showed COMPAS's 137-feature model is no more accurate than untrained online workers or a two-variable model, undercutting the accuracy claim on which its deployment rested; Rudin, Wang, and Coker (2020) argued that ProPublica's racial-disparity finding is partly an artifact of COMPAS's nonlinearity in age rather than race; and Bao et al. (2021) argued the COMPAS dataset, resting on re-arrest as a proxy for re-offense, should not be used as a fairness benchmark at all. The impossibility theorem remains the durable teaching point — but the benchmark it was proved against is now itself among the most distrusted in the field, which is part of what the case should teach.<sup>4</sup>

## AFTERMATH & REFORM

The hedges the case carries are load-bearing. Both Northpointe and ProPublica are correct by their respective definitions, and the impossibility result formalizes the tension rather than resolving it. The case does not teach that COMPAS is fair or that COMPAS is unfair; it teaches that the choice between fairness criteria is governance and normative work that the deployment did not surface to the affected jurisdictions or to the defendants whose liberty depended on the score. The LEO on fairness beyond omission is anchored by the case in its mature form — the impossibility result requires the deploying institution to choose, document, and disclose which fairness criterion the instrument optimizes, and to make the trade-off legible to the people the criterion does not protect.

*Both Northpointe and ProPublica were correct by their respective definitions of fairness; the impossibility result is that calibration and equal false-positive rates cannot be simultaneously satisfied when base rates differ across groups, except in degenerate cases.*

EDITORS' SYNTHESIS OF THE COMPAS AUDIT RECORD AND THE 2017 IMPOSSIBILITY-RESULT PAPERS.

## REFERENCES

1. Angwin, J., Larson, J., Mattu, S., & Kirchner, L. (2016), "Machine Bias," ProPublica, May 23, 2016 — the audit investigation of COMPAS scores in Broward County, Florida.
2. Dieterich, W., Mendoza, C., & Brennan, T. (2016), COMPAS Risk Scales: Demonstrating Accuracy Equity and Predictive Parity, Northpointe Inc. response document.
3. Chouldechova, A. (2017), "Fair Prediction with Disparate Impact: A Study of Bias in Recidivism Prediction Instruments," Big Data 5(2):153–163, doi:10.1089/big.2016.0047.
4. Kleinberg, J., Mullainathan, S., & Raghavan, M. (2017), "Inherent Trade-Offs in the Fair Determination of Risk Scores," Proceedings of ITCS 2017, doi:10.4230/LIPIcs.ITCS.2017.43.
5. Dressel, J., & Farid, H. (2018), "The accuracy, fairness, and limits of predicting recidivism," Science Advances 4(1):eaao5580 — COMPAS no more accurate than untrained humans or a two-variable model.
6. Rudin, C., Wang, C., & Coker, B. (2020), "The age of secrecy and unfairness in recidivism prediction," Harvard Data Science Review 2(1) — the age-nonlinearity reanalysis of ProPublica's disparity claim.
7. Bao, M., et al. (2021), "It's COMPASlicated: The Messy Relationship between RAI Datasets and Algorithmic Fairness Benchmarks," NeurIPS Datasets & Benchmarks — the case against COMPAS as a fairness benchmark.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

COMPAS is the central reference in the contemporary algorithmic-fairness literature because the impossibility result was formalized against its audit record. Both the ProPublica finding (unequal false-positive rates) and the Northpointe response (predictive parity within risk scores) are correct by their respective definitions; the Chouldechova and Kleinberg/Mullainathan/Raghavan results show that the two cannot be simultaneously satisfied when base rates differ. The choice between fairness criteria is governance work, not technique.

### LENS APPROACH

COMPAS is the impossibility-result case at consequential- decision scale (induced 8.4; LENS D4+D5/PT6; LEO-4 and LEO-5). LENS uses it in Domain 4 (Test and Evaluation) for the multi-criterion-audit discipline and in Domain 5 (Navigating Sociotechnical Constraints) for the surfacing-bias-through-governance-not-just-technique anchor. Pair with Case 186 (Bartlett mortgage — fairness through unawareness fails), Case 196 (Coots — competing fairness definitions), and Case 189 (SyRI governance-objection- correct precedent). The impossibility result is the load- bearing teaching point; both Northpointe and ProPublica are correct by their respective definitions.

| DOMAIN               | FAILURE MODES                                       | LENS COMPETENCY                                                       | PROBLEM TYPE | INDUCED | LEO   |  |  |  |
|----------------------|-----------------------------------------------------|-----------------------------------------------------------------------|--------------|---------|-------|--|--|--|
| CRIMINAL JUSTICE     | ➔ <div><div>D</div><div>K</div><div>N</div></div> ➔ | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 6            | 8.4     | 4 · 5 |  |  |  |
| PREDICTIVE ANALYTICS |                                                     | Test and Evaluation + D5                                              |              |         |       |  |  |  |
| ALGO FAIRNESS        |                                                     |                                                                       |              |         |       |  |  |  |

APPROACHES TO CONSIDER

| DURING DEVELOPMENT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | IN OPERATION / AFTER                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <ul style="list-style-type: none"><li>Choose, document, and disclose the fairness criterion the instrument optimizes for in advance of deployment; the impossibility result requires the deploying institution to make the choice and to make the trade-off legible to the people the criterion does not protect.</li><li>Audit the deployed instrument against multiple fairness criteria simultaneously; the COMPAS record demonstrates that an instrument can satisfy predictive parity and fail equality of false-positive rates at the same time, and the audit must surface both.</li><li>Treat the base-rate difference across groups as a governance fact, not a technical artifact; the difference is what makes the impossibility binding, and pretending it can be eliminated is the rhetorical move the case teaches to refuse.</li></ul> | <ul style="list-style-type: none"><li>Carry the impossibility result into print as the teaching point; the case does not teach that COMPAS is fair or that COMPAS is unfair, and the editorial framing must preserve the formal constraint that both audit findings instantiate.</li><li>Pair in syllabi with Case 186 (Bartlett) so the fairness-through-unawareness-fails thread and the impossibility-of-multiple-criteria thread are taught together as complementary structural arguments about disparate impact.</li><li>Use the case to anchor the fairness-beyond-omission LEO; the curricular target is the discipline of choosing and disclosing the fairness criterion when the impossibility result rules out satisfying all of them simultaneously.</li></ul> |

REFLECTION QUESTIONS

1. Identify a risk-assessment instrument in your domain whose fairness criterion has not been chosen and disclosed in advance of deployment. What are the candidate criteria, and which of them are jointly satisfiable given the base-rate differences across affected groups?
2. Specify the multi-criterion audit you would run against a deployed instrument. What is the format in which the audit surfaces incompatibility findings to the deploying institution, and what is the documented decision rule when incompatibility is found?
3. The impossibility result is mathematical; the choice between criteria is governance and normative work. Pick a setting in your domain and ask: who has authority to make the choice, who is accountable for documenting and disclosing the trade-off, and to whom is the trade-off disclosable?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

## CASE 189

GOVERNMENT/WELFARE (NETHERLANDS)

ALGORITHMIC DECISION-MAKING

PUBLIC-SECTOR AI

2014 – 2020

# Dutch SyRI — Welfare–Fraud Risk Scoring Halted on Rights Grounds

**D** Designed Out   **G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** An opaque risk-scoring system that combined up to 17 categories of previously siloed government data on citizens — applied only in so-labelled “problem districts”, with the court finding a risk of links drawn on lower socio-economic status or an immigration background — halted by the District Court of The Hague in 2020 as a violation of Article 8 ECHR; also reported as operationally ineffective

## IN BRIEF

The Dutch System Risk Indication (SyRI) combined up to 17 categories of previously siloed government data — tax, benefits, housing, education — to produce risk scores intended to flag potential welfare fraud. Deployed from 2014, it was applied only in so-labelled “problem districts”. On February 5, 2020 the District Court of The Hague halted the program as a violation of Article 8 ECHR (right to private life) — one of the first times a court anywhere stopped a welfare-AI system on human-rights grounds. Investigative reporting also found that, on its own terms, the system did not work: none of the algorithmic investigations had detected new fraud. The case is the paired contrast to the Open University (Case 81): SyRI’s governance objection was correct, not dissolvable by design — the system was both rights-violating and ineffective. The court left open that a more transparent system could pass, which makes the case a governance-objection-diagnostic teaching point rather than a verdict against all public-sector AI: design can correct some objections; some objections are correct, and no design dissolves them.

## BACKGROUND

SyRI — the System Risk Indication — was a Dutch government analytics program intended to surface potential welfare fraud by combining up to seventeen categories of previously siloed government data on citizens: tax, benefits, housing, education, employment. Deployed from 2014, it operated as a closed system: citizens were not informed they were being scored; the model and its inputs were not disclosed; the program was targeted at only so-labelled “problem districts”, the State confirmed at the hearing.<sup>o</sup>

## WHAT HAPPENED

The governance objection was raised by a coalition of civil society organizations and individual citizens who sued the Dutch state, arguing that SyRI violated Article 8 of the European

Convention on Human Rights — the right to private life. The case was heard by the District Court of The Hague. On February 5, 2020, the court halted the program: the system was found insufficiently transparent and verifiable to strike the fair balance Article 8(2) requires. The court accepted the State's case that combating fraud answers a pressing social need; what it held was that the SyRI legislation is "insufficiently transparent and verifiable to conclude that the interference with the right to respect for private life which the use of SyRI may entail is necessary, proportional and proportionate in relation to the aims the legislation pursues." The ruling is one of the first instances anywhere of a court stopping a welfare-AI system on human-rights grounds.<sup>1</sup>

#### THE INVESTIGATION

Investigative reporting alongside the court case surfaced a second finding: on its own operational terms, SyRI was ineffective. Reporting documented that none of the algorithmic investigations the program had launched had detected new fraud — across its first five years, five municipalities asked for a neighbourhood analysis and only two projects were actually run. The case is structurally important because it is both rights-violating **and** ineffective; the rights argument did not depend on the effectiveness gap, and the effectiveness gap did not depend on the rights argument. The two failures landed on the same system.<sup>2</sup>

#### THE CAPABILITY GAP

What the court explicitly left open is also pedagogically decisive. The ruling did not say no welfare-analytics system could pass Article 8; it said **this** system, in its specific opacity and disproportionality, could not. A more transparent system, with auditability and a credibly narrower scope, might pass. The case is therefore the governance-objection-diagnostic counter to the Open University (Case 81): both faced credible governance objections at the same era; OU's was about trust and accountability and was dissolved by a co-authored consent architecture; SyRI's was about rights and proportionality and could not be dissolved by design because the design was the rights violation.<sup>3</sup>

#### AFTERMATH & REFORM

The pair (Cases 81 + 189) is the case-grounded basis for the governance-objection diagnostic proposed in the v2 research backbone: distinguishing a governance objection that good design can dissolve from one that correctly signals the system should not deploy. The capability is to make the diagnostic call before deployment, not after — and the diagnostic itself is testable: a system whose operational claim has not been demonstrated after years of deployment, and whose data subjects have not been informed they are being scored, is not a case where design can fix the governance problem. The halt did not end Dutch welfare-fraud analytics: a broader successor statute — the WGS (Data Processing by Collaborative Networks Act), dubbed "Super-SyRI" by critics — passed the Senate in June 2024 and took effect in 2025, authorizing wider public-private data-sharing, and SyRI's era overlapped the toeslagenaffaire childcare-benefits scandal, in which roughly 26,000 families were wrongly profiled — a failure grave enough to collapse the Rutte government in January 2021 and draw a parliamentary inquiry in 2024, with compensation still unresolved years later.

*Some governance objections are correct. Design does not dissolve them — it is what the objection is to.*



EDITORS' SYNTHESIS OF THE SYRI RULING AND RACHOVITSA & JOHANN (2022).

REFERENCES

1.

District Court of The Hague (2020), judgment of 5 February 2020 (NJCM et al. v. State of the Netherlands), ECLI:NL:RBDHA:2020:865; the official English translation is ECLI:NL:RBDHA:2020:1878, from which the paragraph numbers below are taken. ¶16.93 records that SyRI had to date been applied only in so-called “problem districts” and that “there is in fact a risk that SyRI inadvertently creates links based on bias, such as a lower socio-economic status or an immigration background.”

2.

Rachovitsa & Johann (2022), “The Human Rights Implications of the Use of AI in the Digital Welfare State: Lessons Learned from the Dutch SyRI Case,” Human Rights Law Review 22(2), doi:10.1093/hrlr/ngac010; the transparency holding is at ¶16.86 of the judgment.

3.

AlgorithmWatch (2020), “How Dutch activists got an invasive fraud detection algorithm banned,” Automating Society Report 2020 — carrying the de Volkskrant reporting of late June 2019 that no SyRI investigation had detected new fraud, and that of five neighbourhood analyses requested by municipalities in five years only two were run.

4.

Library of Congress Global Legal Monitor (2020), report on the SyRI ruling — the ruling’s system-specific scope.

THE LEARNING ENGINEERING LENS

LE INSIGHT

SyRI is the canonical case in the corpus of a governance objection that was correct. The system was both rights-violating (Article 8 ECHR) and operationally ineffective (no new fraud detected). The court ruling is system-specific — it does not foreclose all public-sector AI — which is what makes the case a governance-objection diagnostic teaching point rather than a categorical verdict.

LENS APPROACH

SyRI is the negative Domain 5 / Problem Type 6 governance case (induced 5.1; LENS D5+D3/PT6) drafted as the contrast to the Open University (Case 81). LENS uses the pair in Domain 5 (Navigating Sociotechnical Constraints) for the governance-objection diagnostic — distinguishing dissolvable from correct objections — and in Domain 3 (Human-System Collaboration) for the delegation-with-revocation LEO: the court was the revocation channel because the system did not have one of its own. Adjacent to the AI-delegation typology (TREWS / Epic / SyRI / Watson) the v2 corpus builds.

| DOMAIN                                                                                                                      | FAILURE MODES  | LENS COMPETENCY                                                                   | PROBLEM TYPE | INDUCED | LEO   |
|-----------------------------------------------------------------------------------------------------------------------------|----------------|-----------------------------------------------------------------------------------|--------------|---------|-------|
| <div>GOVERNMENT/<br/>WELFARE<br/>(NETHERLANDS)</div> <div>ALGORITHMIC<br/>DECISION-MAKING</div> <div>PUBLIC-SECTOR AI</div> | <div>DGN</div> | <div>D1D2D3D4D5</div> <div>→ Navigating Sociotechnical<br/>Constraints + D3</div> | 6            | 5.1     | 5 · 3 |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Before deploying a public-sector analytics system, document the data subjects who will be scored, the inputs, the model behavior, and the proportionality argument; opacity is a governance liability that compounds over the life of the deployment.
- ▶ Specify in advance the operational evidence that would demonstrate the capability claim — for SyRI, fraud actually detected by algorithmic investigation — with a published reporting cadence and a pre-registered threshold for what would count as success or failure.

IN OPERATION / AFTER

- ▶ If a court or rights body halts the program, treat the ruling as governance information about the design specifically, not as a verdict about all related systems; redesign with the named deficiencies addressed.
- ▶ Publish the operational effectiveness evidence on the cadence specified at launch; six years without published evidence of the headline capability is itself a governance finding.
- ▶ Treat the rights-and-proportionality channel as a separate evaluation from the technical-accuracy channel; a system can be technically accurate and still fail the proportionality test, and SyRI failed both.

- ▶ Conduct the governance–objection diagnostic openly: is the objection about trust and accountability (potentially dissolvable by design) or about rights and proportionality (not dissolvable by design)? Different answers imply different deployment paths.

### REFLECTION QUESTIONS

1. Identify a public–sector analytics system in your jurisdiction. Are the data subjects informed they are being scored? Is the model and its inputs disclosed? What is the proportionality argument the deploying agency would make if challenged?
2. Specify the pre-registered operational evidence — what the system has to demonstrate, on what cadence — that would let a deployment be evaluated against its capability claim. SyRI ran for six years without published evidence of the headline claim.
3. The pair OU (Case 81) and SyRI (Case 189) teaches the governance–objection diagnostic. Construct a candidate diagnostic for your own domain: what features of an objection indicate that good design could dissolve it, and what features indicate that the objection is to the design itself?

**WHO BUILDS THIS** systems & human–factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

## CASE 190

AUTONOMOUS VEHICLES

REGULATORY GOVERNANCE

INCIDENT DISCLOSURE

2023 – 2024

# Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment

**G** Governance & Trust**K** Knowledge & Institutional Memory**N** Normalization of Deviance**EVIDENCE TIER** journalism-tier — source confidence flagged; future validation ongoing.

**IMPACT** On 24 October 2023 the California DMV suspended Cruise's driverless deployment and testing permits, citing the company's misrepresentation of safety-relevant information after a robotaxi dragged a pedestrian about 20 feet, reaching ~7 mph, following an initial stop; the order cites four regulatory provisions carrying two substantive findings — that the vehicles are not safe for the public's operation, and that the manufacturer misrepresented safety information

## IN BRIEF

On 2 October 2023 a pedestrian was struck by a human-driven vehicle in San Francisco and propelled into the path of a Cruise robotaxi. The robotaxi came to a stop, then performed a pullover maneuver that dragged the pedestrian about 20 feet, reaching about 7 mph, before stopping again. Cruise initially shared video of the initial stop with regulators and reporters but did not disclose the subsequent movement; the California DMV learned of the pullover from another agency and received the full video on 13 October. On 24 October the DMV suspended Cruise's driverless deployment and testing permits, citing misrepresentation of safety-relevant information. The case is the explicit contrast to Case 199 (Waymo): the same regulatory regime, the same delegation problem, the opposite governance choice. Partial disclosure was the failure mode, not the underlying collision sequence. The evidence-tier flag is journalism: the DMV's published Order of Suspension is investigation-grade, but the partial-disclosure mechanism and internal timeline are reconstructed from contemporaneous reporting at TechCrunch, NBC News, SF Standard, and Mission Local. The underlying Cruise internal post-mortem is public since 25 January 2024 as the Quinn Emanuel report. Future validation will continue on long-run regulatory consequences.

## BACKGROUND

The night of 2 October 2023, on a San Francisco street, a pedestrian was struck by a human-driven vehicle and propelled into the path of a Cruise robotaxi operating without a safety driver. The robotaxi detected the collision and came to a stop. The vehicle then executed what Cruise later described as a pullover maneuver, traveling about 20 feet and reaching a speed of 7 miles per hour before coming to a complete and final stop, with the pedestrian

pinned beneath it. The initiating collision was a hit-and-run driver's, not Cruise's; everything after it was Cruise's system acting on its own.<sup>0</sup>

#### WHAT HAPPENED

In the immediate days after the incident, Cruise never told the media about the dragging or showed them the full video. With regulators there are two readings: the DMV says the pullover was not shown to it on 3 October; Cruise employees say they played the full video and that connectivity problems hampered all but one of those meetings. The California DMV did not learn of the full sequence from Cruise; it learned of the pullover from another agency and obtained the full video on 13 October, eleven days after. The mechanism of failure shifted in the regulator's view from the collision to the company's disclosure posture.<sup>1</sup>

#### THE INVESTIGATION

On 24 October 2023 the DMV issued an Order of Suspension suspending Cruise's driverless deployment permit and its driverless testing permit. The order cites four regulatory provisions: that the vehicles "are not safe for the public's operation" (13 CCR §228.20(b)(6)); that the manufacturer "has misrepresented any information related to safety of the autonomous technology" (§228.20(b)(3)); an act or omission making testing on public roads an unreasonable risk to the public (§227.42(b)(5)); and the authority to suspend immediately where required for the safety of persons (§227.42(c)). Two substantive findings sit under them — unsafe performance and misrepresentation — and it is the second that makes this a disclosure case. The Order of Suspension is the investigation-grade artifact in this case; the reconstruction of the partial-disclosure mechanism and the internal timeline rests on contemporaneous journalism.<sup>2</sup>

#### THE CAPABILITY GAP

The case is the explicit foil to Case 199 (Waymo). Same regulatory regime, same delegation problem, opposite governance choice. Where Waymo answered an opacity objection by engineering a published safety-case framework and commissioning third-party audits, Cruise answered an incident-disclosure obligation by sharing partial video. The same DMV that permitted Waymo to continue suspended Cruise's permits. The journalism-tier flag is binding: the DMV's published order is investigation-grade, but the precise sequence of internal decision-making rests on TechCrunch, NBC News, SF Standard, and Mission Local reporting, superseded in part by the Quinn Emanuel report, which Cruise's board commissioned and published on 25 January 2024.<sup>3</sup>

#### AFTERMATH & REFORM

The LENS teaching point pairs directly with Waymo. The new subobjective **Delegation with revocation** (3.3) requires that the deploying organization design the disclosure architecture **before** the failure event — what will be reported, on what cadence, to which oversight body, with what verification. Partial disclosure under crisis is not a strategy; it is what happens when no architecture was designed. The permit suspension proved terminal rather than a pause: on 30 September 2024 Cruise took a \$1.5 million NHTSA consent order, and in November 2024 entered a deferred prosecution agreement with the DOJ, admitting a false report and paying \$500,000, and in December 2024 General Motors announced it would stop funding Cruise's robotaxi development and fold the unit into GM, cutting roughly half the remaining staff in early 2025. Future validation will continue on the long-run regulatory

consequences as the broader AV regulatory regime updates its disclosure requirements in light of the 2023 events.<sup>4</sup>

*The collision involved a human driver. The suspension was about what the company chose not to show.*

EDITORS' SYNTHESIS OF THE CALIFORNIA DMV ORDER OF SUSPENSION AND THE CONTEMPORANEOUS REPORTING.

REFERENCES

1. California DMV (24 October 2023), Order of Suspension, Cruise LLC — driverless deployment and testing permits: “The AV traveled about 20 feet and reached a speed of 7 miles per hour before coming to a complete and final stop.”

2. California DMV, Order of Suspension (24 October 2023) — the 3 October meeting at which the footage shown “ended with the AV’s initial stop”; “the department only learned of the AV’s subsequent movement via discussion with another government agency”; the full video provided on 13 October. Reconstruction of the internal decision-making rests on TechCrunch (October 2023).

3. California DMV, “DMV Statement on Cruise LLC Suspension” (24 October 2023) — the four regulatory provisions cited: 13 CCR §§228.20(b)(6), 228.20(b)(3), 227.42(b)(5) and 227.42(c).

4. Quinn Emanuel Urquhart & Sullivan, report to the Cruise board on the 2 October 2023 incident and the company’s regulatory disclosures, published 25 January 2024; NBC News, SF Standard and Mission Local reporting (October 2023) — the contemporaneous journalism-tier reconstruction.

5. NHTSA consent order with Cruise, 30 September 2024 (\$15 million civil penalty); U.S. Department of Justice deferred prosecution agreement with Cruise, November 2024 (admission of a false report; \$500,000); General Motors announcement of 10 December 2024 ending funding of Cruise robotaxi development.

6. Cruise public statements, October–November 2023, and the company’s blog post accepting the Quinn Emanuel findings (January 2024).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Cruise is the foil to Waymo: same regulatory regime, same delegation problem, opposite governance choice. The DMV’s Order of Suspension is investigation-grade; the partial-disclosure mechanism is reconstructed from journalism. The evidence-tier flag is binding — the internal timeline is journalism-tier and future validation continues as more of the company’s own post-mortem becomes public.

LENS APPROACH

Cruise is the AV partial-disclosure failure (induced 5.4; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the subobjective **Delegation with revocation** (3.3) — the case shows what revocation looks like when the disclosure architecture fails — and in Domain 3 (Human-System Collaboration) for the oversight obligations that follow when a system acts autonomously after a triggering event. Direct foil to Case 199 (Waymo); pairs with Case 200 (CPUC) on the regulator-side.

| DOMAIN                                                                                         | FAILURE MODES                 | LENS COMPETENCY                                                              | PROBLEM TYPE | INDUCED | LEO   |
|------------------------------------------------------------------------------------------------|-------------------------------|------------------------------------------------------------------------------|--------------|---------|-------|
| <div>AUTONOMOUS VEHICLES</div> <div>REGULATORY GOVERNANCE</div> <div>INCIDENT DISCLOSURE</div> | <div>G K N</div> <div>→</div> | <div>D1 D2 D3 D4 D5</div> <div>→ Navigating Sociotechnical Constraints</div> | 6            | 5.4     | 5 · 3 |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Design the incident-disclosure architecture before deployment, not under crisis — what will be reported, on what cadence, to which oversight body, with what verification by which agency.

IN OPERATION / AFTER

- Use the subobjective **Delegation with revocation** (3.3): revocation pathways must be designed and exercisable, and the deploying organization should

- ▶ Treat the disclosure obligation as decisive on its own — partial disclosure of a safety-relevant event is not “less bad” than non-disclosure; under the regulator’s framing it can be the failure mode that triggers revocation.
- ▶ Build the verification path the regulator will use into the architecture. The DMV learning the fuller sequence from another agency is the failure mode the architecture has to preclude.
- ▶ expect the regulator to exercise them when the disclosure architecture fails.
- ▶ Pair the case with Case 199 (Waymo) in any LENS Domain 4 module — the contrast between an engineered legitimacy artifact and partial crisis disclosure is the teaching point, not either case alone.
- ▶ Carry the journalism-tier flag honestly: the DMV Order is the investigation-grade primary; the internal-timeline reconstructions are journalistic; future validation will improve as Cruise’s own post-mortem and any subsequent litigation discovery enter the public record.

## REFLECTION QUESTIONS

1. Imagine you operate an autonomous system that has just been involved in a safety-relevant event. Design the disclosure decision: what is reported, to whom, on what cadence, with what verification — **before** you have the lawyer’s advice on what the disclosure obligation strictly requires. Where does your architecture leave you exposed?
2. Compare Cases 199 (Waymo) and 190 (Cruise) as a paired teaching unit. What is the smallest pre-incident artifact a deploying organization could publish that would make the post-incident disclosure architecture credible to a regulator?
3. The case rests partly on journalism-tier reconstruction of internal decisions. What evidence would you want to see — court discovery, the company’s full post-mortem, a multi-source corroboration — before treating any specific internal-timeline claim as decision-grade?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

# Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict

**D** Designed Out   **G** Governance & Trust   **N** Normalization of Deviance

**IMPACT** Income-averaging algorithm used by the Australian Department of Human Services and Centrelink raised approximately 470,000 wrongful debts against welfare recipients between 2016 and 2019; Royal Commission led by Catherine Holmes AC SC delivered its final report July 7 2023 finding the scheme unlawful and circumstantially attributing multiple deaths to its operation

## IN BRIEF

The Royal Commission into the Robodebt Scheme delivered its final report on July 7, 2023, under the leadership of Commissioner Catherine Holmes AC SC. The report concluded that the income-averaging algorithm operated by the Australian Department of Human Services and Centrelink between 2016 and 2019 had raised approximately 470,000 wrongful debts against welfare recipients and had operated outside the law. The algorithm averaged annual taxable income from the Australian Taxation Office across fortnightly Centrelink reporting periods and treated any resulting discrepancy as a debt the recipient had to disprove. The burden of proof was reversed onto the recipient. The Commission's attribution language on causation of deaths is careful — the attribution is circumstantial and not a direct legal finding of individual causation — but the finding that multiple deaths were associated with the scheme's operation is part of the adjudicated record. The case pairs with Case 189 (SyRI, the governance-objection-correct precedent), Case 48 (Johnson school surveillance, the algorithmic-public-administration parallel), and Case 5 (Epic Sepsis, the delegation-without-validation form).

## BACKGROUND

The Australian income-support architecture pairs Centrelink fortnightly income reporting with annual taxable income reporting to the Australian Taxation Office. The two reporting cadences exist because they measure different things — Centrelink measures earnings inside each fortnight so the income-test taper can be applied, and the ATO measures annual taxable income for the tax system. Between 2016 and 2019, the Department of Human Services deployed an automated debt-recovery system that took the ATO annual figure, divided it by 26, and compared the fortnightly average against the Centrelink reported figures. Any apparent shortfall was raised as a debt against the recipient, and the recipient was required to produce contemporaneous payslips to disprove it.<sup>o</sup>

## WHAT HAPPENED

The mathematical operation the algorithm performed cannot establish what it was being asked to establish. Annual income divided by 26 is not the income earned in any particular fortnight; a recipient who worked irregularly, or whose hours varied, would generate large arithmetic discrepancies between the averaged figure and the actual fortnightly earnings without ever having been overpaid. The Royal Commission's final report documents that the agency's own legal advice flagged this seam before deployment and that the advice was set aside. Approximately 470,000 debts were raised across the operating window; many recipients paid debts they did not owe, and the Commission identified cases in which recipients took their lives in proximity to the scheme's operation. The attribution language on those deaths is careful — "circumstantial" rather than a direct legal finding of individual causation — and the careful language is itself part of the record the case carries.<sup>1</sup>

## THE INVESTIGATION

The structural critique the Commission delivered is the reversal of the burden of proof. In a properly administered welfare-recovery scheme, the agency must establish that a debt is owed before pursuing recovery. The income-averaging method reversed this: the algorithm asserted a debt on the strength of arithmetic that could not establish overpayment, and the recipient had to produce documentary evidence — often payslips from years earlier, often from employers no longer reachable — to disprove the assertion. The Federal Court's 2019 Amato judgment had already found the method unlawful; the Royal Commission's 2023 report adjudicated the governance question of how the scheme had been built, approved, and operated for three years across multiple ministerial portfolios. The Commission named individuals; some are subject to subsequent referrals to the National Anti-Corruption Commission and to professional bodies.<sup>2</sup> Those referrals resolved slowly and unevenly: the NACC first declined to investigate in mid-2024 — a decision its own Inspector found affected by a conflict of interest — before an independent reconsideration reopened the matter as Operation Myrtleford in 2025. In its March 2026 report the NACC found two of the referred officials had engaged in serious corrupt conduct and cleared the others, including former prime minister Scott Morrison, while the Australian Public Service Commission separately found a dozen officials had breached the code of conduct.

## THE CAPABILITY GAP

The case pairs with Case 189 (SyRI, the Dutch System Risk Indication ruling by the Hague District Court) as the governance-objection-correct precedent — SyRI was struck down before it produced a debt-scale harm record; Robodebt operated for three years and the harm record is what the Commission adjudicated. Pair with Case 48 (Johnson school surveillance) for the algorithmic-public-administration parallel at a different population and a smaller scale; the structural form — algorithm asserts, affected party disprove — recurs across the two cases. Pair with Case 5 (Epic Sepsis) for the delegation-without-validation form: in Epic Sepsis the delegated system asserts a clinical risk; in Robodebt the delegated system asserts a financial debt; in both, the asserting party did not validate the assertion against the ground truth before consequences were transmitted to the affected person.<sup>3</sup>



AFTERMATH & REFORM

The hedges the case carries are not optional. The Royal Commission’s attribution language on deaths is circumstantial; the Commission did not, and could not on the evidence available, make individual legal findings of causation in those deaths. The case teaches the structural pattern — algorithmic public administration with the burden of proof reversed, deployed without the legal-advice seam being honored, operated for three years across multiple ministers — and the structural pattern is what makes Robodebt the reference case for an entire class of contemporary algorithmic-administration failures. The Commission’s careful attribution language and the case’s careful editorial framing travel together; neither is smoothed in the casebook’s rendering.

*The income-averaging method could not establish what it was being asked to establish, the burden of proof was reversed onto the recipient, and the Commission’s attribution on associated deaths is circumstantial — and these three facts together are what Robodebt teaches.*

EDITORS’ SYNTHESIS OF THE ROYAL COMMISSION FINAL REPORT (HOLMES AC SC, JULY 2023).

REFERENCES

1. Australian National Audit Office, Centrelink’s Compliance Activities — Income Compliance Program, performance audit reports across 2017–2020.

2. Royal Commission into the Robodebt Scheme, Final Report, Commissioner Catherine Holmes AC SC, July 7, 2023 (Commonwealth of Australia).

3. Prygodicz v Commonwealth of Australia (No 2) [2021] FCA 634 — Federal Court class-action settlement following the 2019 unlawfulness finding.

4. Whiteford, P. (2021), “Debt by Design: The Anatomy of a Social Policy Fiasco,” Australian Journal of Public Administration 80(2):340–360 — academic synthesis of the policy and administrative history.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Robodebt is the reference case for algorithmic public administration deployed at population scale without the burden of proof being honored. The Royal Commission’s final report adjudicated approximately 470,000 wrongful debts and circumstantially attributed multiple deaths to the scheme’s operation; the careful attribution language and the structural finding travel together, and neither is smoothed in the casebook’s rendering.

LENS APPROACH

Robodebt is the burden-of-proof-reversal case at population scale (induced 5.2; LENS D5+D3/PT6; LEO-5 and LEO-3). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the agency-legal-advice-as-binding-gate discipline and in Domain 3 (Human-System Collaboration) for the human-in- the-loop-for-consequential-decisions anchor. Pair with Case 189 (SyRI governance-objection-correct precedent), Case 48 (Johnson school surveillance algorithmic-public-administration parallel), and Case 5 (Epic Sepsis delegation-without- validation form). The Commission’s circumstantial attribution on deaths is the hedge that must survive.

| DOMAIN                                                                                | FAILURE MODES                          | LENS COMPETENCY                                                       | PROBLEM TYPE                               | INDUCED | LEO   |
|---------------------------------------------------------------------------------------|----------------------------------------|-----------------------------------------------------------------------|--------------------------------------------|---------|-------|
| <div>GOVERNMENT</div> <div>SOCIAL WELFARE</div> <div>ALGORITHMIC ADMINISTRATION</div> | <div>D</div> <div>G</div> <div>N</div> | <div>D1</div> <div>D2</div> <div>D3</div> <div>D4</div> <div>D5</div> | 6                                          | 5.2     | 5 · 3 |
|                                                                                       | →                                      | →                                                                     | Navigating Sociotechnical Constraints + D3 |         |       |

## APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- ▶ Treat agency legal advice on the lawfulness of an algorithmic-administration method as a binding gate, not a negotiable input; the Commission's record on Robodebt is that the seam was flagged in advance and that the override of the advice is itself the governance failure.
- ▶ Maintain the burden of proof on the asserting party for any algorithmically generated debt or risk; arithmetic that cannot establish the asserted fact cannot be the basis for transmitting consequences to an affected person.
- ▶ Build the cross-portfolio review surface that a multi-year algorithmic-administration scheme requires; Robodebt operated across multiple ministers and the cross-portfolio handoff was where the governance check kept being deferred.

## IN OPERATION / AFTER

- ▶ Carry the Commission's careful attribution language on deaths into print without softening; the case's value depends on the circumstantial nature of the attribution being preserved alongside the structural finding.
- ▶ Pair in syllabi with Case 189 (SyRI) so the governance-objection-correct precedent and the governance-objection-overridden harm record are taught together; the two cases together teach what advance objection can prevent and what its absence can produce.
- ▶ Use the case to anchor the human-in-the-loop LEO at population scale; the form Robodebt makes legible is what consequential-decision delegation looks like when the loop is removed and the asserting party operates on arithmetic that cannot establish its assertion.

## REFLECTION QUESTIONS

1. Identify an algorithmic-administration scheme in your domain whose arithmetic asserts a fact against an affected person. What is the asserting party's burden of proof, and what would the affected person have to produce to disprove the assertion?
2. Specify the legal-advice gate you would treat as binding in advance of deployment. What is the documented escalation path when the advice flags an unlawfulness seam, and who has authority to override it?
3. The Royal Commission's attribution on deaths is circumstantial — careful, not adjudicated as individual legal findings. Pick a setting where harm attribution to an algorithmic system is contested, and ask: what language would honor both the structural pattern and the limits of what the evidence can establish?

**WHO BUILDS THIS** systems & human-factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

## CASE 192

FINANCIAL SERVICES

CONSUMER CREDIT

ALGORITHMIC DECISION-MAKING

2019 – 2021

# Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model

**D** Designed Out   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** New York State Department of Financial Services investigation March 2021 did not produce evidence of deliberate or disparate-impact discrimination under New York and federal fair lending law in Apple Card credit-limit decisions following David Heinemeier Hansson's November 2019 viral allegation that his wife received a credit limit approximately 20 times lower despite shared assets; DFS documented "lack of transparency" as the structural problem, and the Bank and Apple had since taken steps to remedy the customer-service and transparency deficiencies

## IN BRIEF

On November 7, 2019, software developer David Heinemeier Hansson posted on Twitter that his Apple Card credit limit had been set approximately 20 times higher than his wife Jamie Heinemeier Hansson's, though the couple file jointly and, he said, she had the higher credit score. The thread went viral, with Apple co-founder Steve Wozniak reporting a similar disparity with his wife. The New York State Department of Financial Services opened an investigation. Its March 2021 report found no evidence that the credit-decisioning algorithm operated by Goldman Sachs (the issuing bank for Apple Card) discriminated — neither intentional disparate treatment nor a disparate-impact violation of New York and federal fair lending law. DFS documented "lack of transparency" as the structural problem: the Bank explained its decisions to DFS for every complainant, and to the complainants themselves it did not. The Bank and Apple had by then taken steps to remedy the deficiencies. The case pairs with Case 186 (Bartlett mortgage), Case 196 (Coots), and Case 86 (Gándara). The DFS finding of "no violation but lack of transparency" is the decisive nuance.

## BACKGROUND

The Apple Card launched in August 2019 as a co-branded consumer credit product issued by Goldman Sachs, with underwriting and credit-line decisions made by Goldman's consumer-banking unit. The product was Goldman's first retail consumer credit product at meaningful scale, and the decisioning architecture was built de novo against contemporary algorithmic-underwriting practice. On November 7, 2019, David Heinemeier Hansson posted that his Apple Card credit limit was approximately 20 times his wife's, though the couple file jointly and, he said, Jamie had the higher credit score. The thread surfaced similar reports

— Steve Wozniak named the same pattern with his wife — and rapidly moved from social media to regulatory attention.<sup>0</sup>

#### WHAT HAPPENED

The structural seam the case opens is that Goldman Sachs did not explain individual credit decisions to the applicants who challenged them. When Jamie Heinemeier Hansson asked Goldman to explain why her credit limit was lower, the customer-service response was that the algorithm had set the limit and that the decision could not be explained at the individual level. The escalation moved through standard customer-service channels, then to social-media-amplified public pressure, then to regulatory investigation. Asked by DFS, the Bank identified the factors behind each complainant's decision — credit score, indebtedness, income, credit utilization, missed payments — and DFS found each consistent with the Bank's credit policy. The explainability gap was operational rather than algorithmic in the narrow sense: even if the underlying model was defensible at the population level, the bank had not built the infrastructure to defend individual decisions to applicants.<sup>1</sup>

#### THE INVESTIGATION

The New York State Department of Financial Services opened its investigation on 9 November 2019 and released its findings in March 2021. The headline finding was that DFS did not find evidence of a violation of New York and federal fair lending law — neither intentional disparate treatment nor a disparate-impact violation under the applicable standards; women and men with equivalent credit characteristics, the report states, had similar Apple Card application outcomes. The investigation ran a regression analysis over underwriting data for nearly 400,000 New York applicants. The hedge is binding: DFS's "no violation" finding is specific to the statutory standard the agency applied, not a general finding that the decisioning was fair or non-discriminatory across all criteria. What DFS did find was "lack of transparency" as the structural problem. The report imposed no penalty and ordered no corrective action, closing instead with the judgment that credit scoring and the laws barring discrimination in lending "are in need of strengthening and modernization." The Bank had already dropped the six-month wait to appeal credit terms, and with Apple launched Path to Apple Card in June 2020 to help declined applicants build toward approval.<sup>2</sup>

#### THE CAPABILITY GAP

The case pairs with Case 186 (Bartlett mortgage) for the consumer-credit-fairness thread at adjacent scale and regulatory regime. Pair with Case 196 (Coots) for the competing-fairness-definitions thread; the DFS standard is one of several available standards, and the case teaches that "no violation under a specific statutory standard" is not "fair." Pair with Case 86 (Gándara) for the explainability-of-individual-predictions thread at a different population and scale. The case is unusual in the casebook for the speed of regulatory response — DFS opened the investigation two days after the viral thread — and for the structural conclusion that the failure was explainability rather than the decisioning algorithm itself.<sup>3</sup>

#### AFTERMATH & REFORM

The load-bearing hedge is the precise DFS finding. The case does not teach that Apple Card was unfair, and it does not teach that DFS found Apple Card was fair. It teaches that under the specific statutory standard the agency applied, no violation was found, and that the structural failure the agency named was "lack of transparency" — Goldman Sachs's inability to

explain individual credit decisions to applicants who challenged them. The human-in-the-loop LEO is anchored by the case at the appeal-and-explanation seam: a consequential-decision system that cannot explain its individual decisions to affected applicants does not have a functioning human-in-the-loop appeal mechanism, and the absence of the mechanism is the governance failure whether or not the decisioning is statistically defensible.

*DFS did not find evidence of a violation of New York and federal fair lending law; DFS did find lack of transparency as the structural problem, and the Bank and Apple took steps to remedy it without being ordered to – the case teaches that “no violation under a statutory standard” is not “fair.”*

EDITORS' SYNTHESIS OF THE NY DEPARTMENT OF FINANCIAL SERVICES REPORT ON THE APPLE CARD INVESTIGATION (MARCH 2021).

REFERENCES

1.

Heinemeier Hansson, D. (2019), Twitter thread of November 7, 2019, archived in DFS investigation record and contemporaneous press coverage (Bloomberg, New York Times, Wall Street Journal, November 2019).

2.

Goldman Sachs Bank USA, public response and credit-line-review process documentation submitted to DFS during the investigation (2019 – 2021).

3.

New York State Department of Financial Services, Report on Apple Card Investigation, March 23, 2021.

4.

Vigdor, N. (2019), “Apple Card Investigated After Gender Discrimination Complaints,” The New York Times, November 10, 2019.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Apple Card is the consumer-credit explainability case at deployment scale. DFS found no evidence of a violation of New York and federal fair lending law under the applicable standards, and DFS also found lack of transparency as the structural problem; the Bank dropped the six-month wait to appeal credit terms and, with Apple, launched Path to Apple Card for declined applicants. The hedge that matters is the precision of the DFS finding — neither “fair” nor “unfair,” but “no violation under this standard, transparency gap as the structural problem.”

LENS APPROACH

Apple Card is the explainability-as-governance case at consumer-credit scale (induced 5.2; LENS D5/PT6; LEO-5 and LEO-3). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the appeal-and-explanation-infrastructure discipline and in Domain 3 (Human-System Collaboration) for the human-in-the-loop-for-consequential-decisions anchor at the appeal seam. Pair with Case 186 (Bartlett mortgage), Case 196 (Coots competing fairness definitions), and Case 86 (Gándara explainability of individual predictions). The precise DFS finding — “no violation but lack of transparency” — is the hedge that matters.

| DOMAIN                      | FAILURE MODES | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO   |
|-----------------------------|---------------|---------------------------------------|--------------|---------|-------|
| FINANCIAL SERVICES          | D K N         | D1 D2 D3 D4 D5                        | 6            | 5.2     | 5 · 3 |
| CONSUMER CREDIT             | →             | →                                     |              |         |       |
| ALGORITHMIC DECISION-MAKING |               | Navigating Sociotechnical Constraints |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

IN OPERATION / AFTER

- ▶ Build the individual-applicant explanation infrastructure as part of the deployment, not as a customer-service afterthought; the Apple Card case demonstrates that a defensible population-level model paired with no individual-explanation infrastructure produces a regulatory finding of lack of transparency.
- ▶ Specify the appeal mechanism before the first decision is made; the human-in-the-loop LEO at consumer-credit scale is the appeal-and-explanation seam, and its absence is the governance failure the case names.
- ▶ Treat the customer-service escalation path as a deployment surface, not a support function; the case's escalation went from customer service to social media to regulation in days, and the deployment surface that mattered was the first one.
- ▶ Carry the precise DFS finding into print without softening; "no violation under the applicable statutory standard, but lack of transparency as the structural problem" is the nuance that matters and the case's pedagogical value depends on the nuance being preserved.
- ▶ Pair in syllabi with Case 186 (Bartlett) so the consumer-credit-fairness regulatory architecture is taught at both the mortgage and credit-card scales.
- ▶ Use the case as the anchor for the explainability-as-governance frame at consumer-credit scale; the curricular target is the appeal-and-explanation infrastructure that converts an algorithmic decision into a contestable one.

### REFLECTION QUESTIONS

1. Identify a consequential-decision system in your domain whose individual-applicant explanation infrastructure has not been built. What is the customer-service escalation path when an affected person challenges a decision, and what would the path look like with an appeal-and-explanation seam built into the deployment?
2. Specify the precise statutory or regulatory standard against which your deployment is being evaluated. What does "no violation under this standard" leave open about fairness across other standards, and how would the trade-off be disclosed?
3. The Apple Card escalation moved from customer service to social media to regulatory investigation within days. Pick a deployment in your domain and ask: what is the first deployment surface an affected person encounters when challenging a decision, and what would have to be true for that surface to resolve the challenge before it moves further?

**WHO BUILDS THIS** systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

# Bernard Madoff / SEC Failures

**G** Governance & Trust   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** ~\$65B Ponzi scheme — the largest in history; SEC repeatedly investigated and cleared Madoff; foundational regulator–capability case

## IN BRIEF

Bernard Madoff ran the largest Ponzi scheme in history — roughly \$65 billion in fictitious account value — for years while the SEC repeatedly investigated and cleared him. Financial analyst Harry Markopolos delivered the agency a detailed memo in 2005, “The World’s Largest Hedge Fund is a Fraud,” showing that Madoff’s steady returns were mathematically impossible. The SEC opened an investigation and concluded no action was warranted. Madoff operated until his sons turned him in during the 2008 crisis, when redemptions became impossible to honor. The SEC Inspector General later found the staff assigned relatively inexperienced, unable to appreciate the analysis in Markopolos’s complaint, and skeptical of it almost at once. The capability gap was at the regulator’s technical–evaluation pipeline, not at the evidence — which was specific and checkable.

## BACKGROUND

Bernard Madoff was a former NASDAQ chairman and a respected Wall Street figure whose investment arm reported remarkably steady returns for years. Those returns were entirely fictitious: client funds were never invested, and earlier investors were paid with later investors’ money — a Ponzi scheme that eventually represented some \$65 billion in fabricated account value. The very steadiness that reassured investors was the tell: real markets do not deliver an almost unbroken line of gains, and the absence of the normal volatility was itself evidence the numbers were manufactured.<sup>0</sup>

## WHAT HAPPENED

Financial analyst Harry Markopolos concluded by analysis that Madoff’s returns were mathematically impossible and delivered the SEC a detailed memorandum in 2005 titled “The World’s Largest Hedge Fund is a Fraud.” The SEC’s New York office opened a formal investigation in January 2006 and closed it in November 2007, the closing recommendation concluding that the violations it had found were “not so serious as to warrant an enforcement action” and never returning to the Ponzi allegation the case-opening report had raised. The warning was not a vague suspicion but a quantitative case any competent reviewer could in principle retrace, which is what makes the dismissal so telling. Madoff continued operating until December 2008, when the financial crisis made redemptions impossible and his sons reported him.<sup>1</sup>

## THE INVESTIGATION


The SEC Office of Inspector General's 2009 report found the agency had received credible, specific complaints across more than a decade and, "despite numerous credible and detailed complaints," never took "the necessary, but basic, steps to determine if Madoff was operating a Ponzi scheme." The staff assigned to Madoff lacked the expertise to evaluate Markopolos's technical arguments, and Madoff worked on that inexperience directly — emphasizing his standing in the industry, dropping the names of senior SEC officials, growing angry when examiners pressed for documents. The report's own account of the failure is plainer than deference: the staff took Madoff's verbal answers where an independent check was available and never made it. A letter to the NASD that would have shown the trades did not exist was drafted and never sent.<sup>2</sup>

## THE CAPABILITY GAP

The complaint was specific; the math was checkable; the institution simply did not have the people to check it. The capability gap was at the regulator's technical-evaluation pipeline, not at the evidence. A regulator whose technical depth lags the entities it oversees cannot act on even a correct and well-documented warning — and because the gap is in the evaluator rather than the tip, more tips would not have helped; the agency could not have used them.<sup>3</sup>

## AFTERMATH & REFORM

The collapse wiped out tens of thousands of investors and prompted SEC reforms to its handling of tips and referrals, its examination procedures, and its recruitment of staff with quantitative and trading expertise, including the creation of an Office of Market Intelligence. Each reform addressed a different part of the same pipeline — getting the tip in, getting it to someone who could read it, and having someone who could — so that a future Markopolos memo would meet an evaluator able to test it. Madoff is a case in which a regulator lacked the technical capability to challenge the evidence in front of it.<sup>4</sup> The aftermath complicates the "total loss" image the \$65-billion figure conjures: the fabricated account value was never real money, and trustee Irving Picard has recovered or agreed to recover more than \$15.3 billion for the customer fund — \$5 billion of it from the Picower estate, whose \$7.2 billion settlement sent a further \$2.2 billion to the government and seeded the Department of Justice's Madoff Victim Fund, whose claimants had recovered 93.71 percent of their losses when it closed at the end of 2024, a recovery without precedent for a Ponzi of this scale. Madoff himself died in federal prison in April 2021.





MADOFF — THE REGULATOR RECEIVED THE EVIDENCE; THE REGULATOR LACKED THE STAFF TO EVALUATE IT

*Despite numerous credible and detailed complaints, the SEC never properly examined or investigated Madoff's trading.*

SEC OFFICE OF INSPECTOR GENERAL, REPORT OF INVESTIGATION, CASE NO. OIG-509 (31 AUGUST 2009)

REFERENCES

1. SEC Office of Inspector General, Investigation of Failure of the SEC to Uncover Bernard Madoff's Ponzi Scheme, Report of Investigation Case No. OIG-509 (31 August 2009) — the quoted finding.

2. Markopolos, H. (2010), No One Would Listen — the 2005 memo and its dismissal.

3. United States v. Madoff (2009) — guilty plea and the \$65B figure.

4. U.S. Senate Committee on Banking, Housing, and Urban Affairs, The Madoff Investment Securities Fraud: Regulatory and Oversight Concerns and the Need for Reform, hearing, 10 September 2009 — SEC and OIG testimony on the examination and expertise failures.

5. U.S. Securities and Exchange Commission, The Securities and Exchange Commission Post-Madoff Reforms (2009–2010) — revamped tip and referral handling, the Office of Market Intelligence, and specialized enforcement units.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Madoff case is the canonical example of a regulator without the technical capability to evaluate the evidence it received. The complaint was specific. The math was checkable. The institution did not have the people to check it. The capability gap was at the regulator's expertise pipeline, not at the evidence.

LENS APPROACH

LENS uses Madoff in LEN 6 as an operator-to-institution feedback-channel failure: Harry Markopolos sent the SEC a specific, checkable memo — the channel existed and the signal was correct — and the institution had no one able to act on what came through it. The teaching point is the receiving capability a feedback channel requires to be real: a tip line without the technical expertise to evaluate the tip is not a channel. It pairs with the Mark 14 torpedo (Case 133), where the Bureau of Ordnance dismissed submariners' field reports for the same structural reason — the channel carried the truth and the institution could not receive it. The Ponzi narrative is not the lesson; the channel is.

| DOMAIN     | FAILURE MODES | LENS COMPETENCY                       | PROBLEM TYPE | INDUCED | LEO |
|------------|---------------|---------------------------------------|--------------|---------|-----|
| GOVERNMENT | → G K N       | → D1 D2 D3 D4 D5                      | 3            | 6.2     | 2   |
| TECHNOLOGY |               | Navigating Sociotechnical Constraints |              |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Staff the evaluation pipeline with quantitative and trading expertise matched to the entities being overseen, so a checkable claim meets someone able to check it.
- ▶ Build tip-and-referral triage that routes a specific, technical complaint to a qualified evaluator rather than letting the subject's reputation decide its fate.
- ▶ Require warnings to be assessed on their merits first, structurally separating the analysis of the math from any weighing of who the subject is.

IN OPERATION / AFTER

- ▶ Audit closed investigations for cases dismissed despite specific, checkable evidence, surfacing where the evaluator — not the evidence — was the failure point.
- ▶ Monitor for returns or results too smooth to be real as a leading indicator, treating the absence of expected volatility as a signal to investigate.
- ▶ Sustain the regulator's technical depth against the regulated industry's growing sophistication, since a depth gap quietly reopens the same blind spot over time.

## REFLECTION QUESTIONS

1. Identify a regulator in your domain whose technical evaluation capability has not kept pace with the entities it regulates. What is the resulting blind spot?
2. Design the regulator-side technical-capability deliverable that should have allowed the SEC to evaluate the Markopolos memo on its merits.
3. The SEC weighed Madoff's stature over Markopolos's math. Design a triage process that forces a tip to be evaluated on its technical merits before the subject's reputation is allowed to enter the decision.

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

Chapter 14

# Algorithms, Governance & Public Systems — What Works — and the Frontier

*When governance is designed as an artifact, not litigated after the fact.*

---

*The objection dissolved when the evidence was engineered to answer it.*

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

CASE 194

E-GOVERNMENT

DISTRIBUTED DATA EXCHANGE

NATIONAL IT INFRASTRUCTURE

2001 – PRESENT

# Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox)

**G** Governance & Trust   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**IMPACT** By December 2024 Estonia reported effectively 100% digitalization of government services across the X-Road data-exchange layer, with the tax authority reporting around 99% of tax declarations filed electronically and a three-minute online filing; the self-critique is that the country has now created its own legacy system — the very thing the program set out to avoid

## IN BRIEF

Estonia launched X-Road in 2001 as the answer to fragmented government databases: a distributed data-exchange layer that lets services interoperate without forcing a central monolith. By December 2024 the country reported effectively 100% digitalization of government services across X-Road, with the tax authority reporting around 99% of declarations filed electronically and a three-minute online filing. The case sits inside C7 not because of the headline outcomes but because of the decisive self-critique surfaced in the published academic analysis: by committing the country to X-Road as the data-exchange layer, Estonia has effectively created its own legacy system — the very thing its founders set out to avoid. The “no-legacy paradox” is the C7 teaching the corpus does not get from any other case: success-as-aging is the failure mode, and the modernization regime must contemplate its own future obsolescence as part of its current design discipline. The decisive self-critique — the no-legacy paradox — is written text in Kattel and Mergel’s own analysis (the 2019 Oxford chapter and the 2018 IIPP working paper behind it), not practitioner reflection, and the program is well-documented in academic and program-report sources alike. What remains open is whether the deliberate generational-replacement plans carry through. Cross-listed with Gap 5 (non-US/UK/EU coverage, Estonia).

## BACKGROUND

Estonia’s pre-2001 problem was the standard one for a small state with a recent administrative inheritance: government databases were fragmented, agency-by-agency, with no common exchange layer. X-Road was the architectural answer — a distributed data-exchange protocol that lets each agency keep its own system but interoperate via a shared, authenticated channel. The design choice avoided the trap of a central monolith and made cross-agency services possible without requiring agencies to surrender their data.<sup>o</sup>

## THE INTERVENTION

The deployment trajectory across two decades has been the defining national IT case study in the EU region. By December 2024 the country reported effectively 100% digitalization of government services, with the tax authority reporting around 99% of tax declarations filed electronically and a three-minute online filing. The Mission Mystique and the Hiding Hand chapter (Oxford 2019) and the ICEGOV 2021 historical analysis are the academic treatments; the X-Road Global program documentation extends the case to the institutions that have since adopted the protocol elsewhere.<sup>1</sup>

## HOW IT WORKED

What makes this a C7 case and not just a digital-government success story is the self-critique. The analysis says it outright. Kattel and Mergel write that with the decision to focus its digital development on X-Road only, Estonia has effectively created its own legacy system — a move that the initial thinkers wanted to avoid. The paradox is named against Estonia's own design rule: no legacy, meaning public digital infrastructure should not run on solutions older than thirteen years. The protocol is now twenty-plus years old; the agency systems that hang off it have accreted their own logic; the generational replacement of the platform itself is now the modernization problem the founders thought they had designed out. Success-as-aging is the failure mode.<sup>2</sup>

## THE EVIDENCE

The sourcing is worth stating precisely, because the self-critique is the decisive part. The academic sources — an edited-volume chapter (whose open-access working-paper version is where the paradox sentence is checkable) and a refereed conference paper, not journal articles — cover the program's design, its trajectory, and the self-critical framing including the no-legacy paradox; the program-report and practitioner material supplies the day-to-day operational picture around it. The country's generational-replacement plans for X-Road exist and have begun to materialize: X-Road version 7 ("Unicorn") is in production and version 8 ("Spaceship") is planned for 2026, re-scoping the platform toward a federated "dataspace" model and aligning its trust framework with the Gaia-X trust framework — the modernization regime acting on its own paradox rather than leaving it as theory. Future validation will continue as the next-generation data-exchange architecture is deployed.<sup>3</sup>

## WHAT TRANSFERRED

What the case adds at the LENS layer is the C7 teaching that no other corpus case supplies: a modernization regime must contemplate its own future obsolescence as part of its current design discipline. The Domain-1 frame applies: systems analysis of the data-exchange layer as an artifact whose own aging trajectory is part of the system's behavior. The Domain-5 frame applies because the sociotechnical commitment of a state to a single data-exchange protocol is itself the governance choice that determines what the future modernization will have to undo. The case is also a Gap-5 non-US/UK case from a small EU state, which the corpus needs.

*Success creates its own aging problem. The modernization regime must contemplate its own future obsolescence as part of its current design.*

EDITORS' SYNTHESIS OF THE NO-LEGACY PARADOX IN THE ESTONIA X-ROAD ANALYSIS.

REFERENCES

1. Jackson, Dreyling, & Pappel (2021), "A Historical Analysis on Interoperability in Estonian Data Exchange Architecture," ICEGOV 2021 proceedings, doi:10.1145/3494193.3494209.

2. Republic of Estonia, e-Estonia briefing materials and Year-in-Review documentation (2024) — program-report sourcing, including the e-Tax figures (around 99% of declarations filed electronically; three minutes to file) and the December 2024 100%-digital milestone.

3. Kattel, R., & Mergel, I. (2018), "Estonia's digital transformation: Mission mystique and the hiding hand," UCL Institute for Innovation and Public Purpose Working Paper 2018-09 — the open-access version carrying the no-legacy-paradox sentence verbatim and the thirteen-year no-legacy design principle.

4. X-Road Global / Nordic Institute for Interoperability Solutions (NIIS) — program documentation and deployment-partner case studies.

5. Kattel, R., & Mergel, I. (2019), "Estonia's Digital Transformation: Mission Mystique and the Hiding Hand," in Compton, M., & 't Hart, P. (eds.), Great Policy Successes (Oxford University Press, 2019) — the analytical chapter; an edited-volume chapter, not a peer-reviewed journal article.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Estonia X-Road is the C7 case the corpus needed for the success-as-aging failure mode: a 100%-digitalization program whose own success has now created the legacy system its founders set out to avoid. Evidence is mixed — academic analysis (an edited-volume chapter and a conference paper) plus program-report and practitioner reflection; the generational-replacement trajectory is the open empirical question. Non-US/UK case, Gap-5 cross-listed.

LENS APPROACH

X-Road is the C7 big-tier transition case with a small-tier self-critique (induced 7.1 and 5.3 alternate; LENS D1/PT1). LENS uses it in Domain 1 (Systems Analysis) for LEO-1 — the data-exchange layer is an artifact whose own aging is part of system behavior — and in Domain 5 (Navigating Sociotechnical Constraints) for LEO-5, since the state's commitment to the protocol is itself the governance choice determining future modernization cost. Pair with Case 116 (Eurocat) for the small-tier transition thread and with the big-tier C7 failures the corpus already documents — Ariane 5 (Case 101), Knight Capital (Case 143), TSB (Case 152), CrowdStrike (Case 172), Y2K (Case 174).

| DOMAIN                     | MODES ADDRESSED | LENS COMPETENCY    | PROBLEM TYPE | INDUCED | LEO   |
|----------------------------|-----------------|--------------------|--------------|---------|-------|
| E-GOVERNMENT               | G K N           | D1 D2 D3 D4 D5     | 1            | 7.1     | 1 · 5 |
| DISTRIBUTED DATA EXCHANGE  | →               | → Systems Analysis |              |         |       |
| NATIONAL IT INFRASTRUCTURE |                 |                    |              |         |       |

APPROACHES TO CONSIDER

- DURING DEVELOPMENT

▶ Design the data-exchange or platform layer so that its own future obsolescence is contemplated in the current design — the no-legacy paradox is the warning the case carries.

▶ Treat the generational-replacement plan as part of the deployment commitment, not as a future problem; the Estonia case is open precisely because the replacement is open.

▶ Document the sociotechnical commitment of the state or institution to the platform explicitly, so the cost of future replacement is in the record from the start.
- IN OPERATION / AFTER

▶ Track the platform's aging trajectory as a system variable — accreted agency logic, protocol drift, dependency depth — not only its operational uptime.

▶ Carry the self-critical framing honestly; the teaching is that success creates its own aging problem, not that the program failed.

▶ Pair with other C7 cases — Patriot/Dhahran (Case 129), Ariane 5 (Case 101), Knight Capital (Case 143), CrowdStrike (Case 172), TSB (Case 152) — for the failure-and-modernization contrast, and with Case 116 (Eurocat ATM modernization) for the small-tier transition thread.

## REFLECTION QUESTIONS

1. Identify a platform or data-exchange layer in your context whose success has created a generational-replacement problem the original design did not contemplate. What would the no-legacy-paradox-aware redesign look like?
2. Specify the aging-trajectory variables you would track on the platform (accreted agency logic, protocol drift, dependency depth) so the future-obsolescence question is in the operational record, not a future surprise.
3. The case is peer-reviewed analysis plus program-report plus practitioner reflection. What additional evidence would you require before committing your own institution to a platform that may itself become the legacy system in twenty years?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

# Tesla Autopilot — Recurring Fatalities

**T** Training Gap   **N** Normalization of Deviance   **G** Governance & Trust   **H** Human-System Interface

**IMPACT** 14 deaths across the 467 Autopilot crashes NHTSA analyzed in EA22-002; first U.S. cases of Level-2 automation contributing to fatal injury

## IN BRIEF

Tesla's Autopilot and Full Self-Driving Beta are Level-2 driver-assistance systems: the human driver remains legally and operationally responsible at all times. NHTSA's engineering analysis of Autopilot found 13 fatal crashes, and 14 deaths, among the 467 it examined, and the record runs back to the 2016 death of Joshua Brown in Florida. The pattern is consistent — the system performs capably for long stretches, the driver's monitoring attention attenuates, and an edge case (a stationary object, a faded lane line, a crossing vehicle) produces a collision the inattentive driver fails to catch. NHTSA's investigation has found Tesla's driver-engagement design inadequate to sustain the attention safe operation requires. At consumer scale, Autopilot is the live test of whether passive monitoring of good-enough automation is a role a human can actually perform.

## THE SHIFT

Partial driving automation has moved from research vehicles into millions of consumer cars. Level-2 systems like Tesla's Autopilot can steer, accelerate, and brake within their operational design domain, but they require the human driver to monitor continuously and take over instantly — a fundamentally new and demanding role assigned to ordinary, untrained consumers. Where a research program could screen, brief, and instrument its safety drivers, a consumer product reaches everyone who buys the car, with no curriculum and no qualification gate standing between purchase and the monitoring task itself.<sup>o</sup>

## WHAT IS EMERGING

Since the first fatal Autopilot crash — Joshua Brown, Florida, 2016 — NHTSA has documented 13 fatal crashes among the 467 Autopilot crashes it analyzed through August 2023. The pattern is consistent: long periods of capable operation, attenuating driver attention, and then an edge case — a stationary fire truck, a faded lane marking, a perpendicular crossing — that the disengaged driver fails to catch in time. The very reliability that makes the system attractive is what erodes the vigilance it depends on, so each uneventful mile quietly raises the odds that the next intervention will come too late.<sup>1</sup>

## THE CAPABILITY QUESTION

The proximate cause in each case is the driver, who holds the legal responsibility for the driving task — an assignment a Florida federal jury qualified for the first time in 2025 (below). But the deeper question is whether sustained vigilant monitoring of an automation that works well most of the time is a role a human can perform at all. Naming the feature



“Autopilot” and designing weak engagement checks shaped the very inattention the system then blamed on the operator — so the architecture both invited the disengagement and reserved the liability for the person least positioned to resist it.<sup>2</sup>

#### EARLY EVIDENCE

NHTSA’s engineering analysis (ODI EA22-002, closed April 2024 and superseded by the December 2023 recall of roughly two million vehicles) identified Tesla’s driver-engagement design as inadequate to maintain the attention safe operation requires, and the recurring fatality pattern across NTSB reports suggests passive monitoring is not a sustainable role as currently engineered. Decades of automation-complacency research point the same way — the finding is not that any one driver failed but that the role asks a human to stay alert to a system precisely because it almost never needs them, a demand the evidence keeps showing is not reliably met.<sup>3</sup>

#### OPEN PROBLEMS

Tesla Autopilot is the consumer-scale version of the Uber ATG problem (Case 183): a passive-monitoring role deployed without the capability infrastructure — training, engagement design, attention measurement — to make it performable. The open problem is what driver-engagement architecture, if any, could make Level-2 monitoring sustainable for an average driver over years of use, and whether the answer is a better attention check or a concession that the role itself has to be redesigned out of the human’s hands.<sup>4</sup> The question did not stay settled at the recall: in April 2024 NHTSA opened a new investigation into whether the December 2023 remedy was even adequate, and in October 2024 opened a separate probe of Full Self-Driving after four reduced-visibility crashes, one of them a pedestrian fatality. Then in June 2025 Tesla launched a limited driverless robotaxi service in Austin — moving the company past Level-2 monitoring into a service it calls autonomous, which drew regulator contact within days — though a Tesla safety monitor rode in the front passenger seat, so the human was relocated, not removed. The liability assignment moved too: in August 2025 a federal jury in Miami found Tesla 33% responsible for a 2019 Key Largo Autopilot crash and returned a \$243 million verdict — the first federal jury verdict in a fatal Autopilot case — which the trial judge declined to set aside in February 2026, with an appeal expected. The driver’s legal responsibility is no longer the whole of it.<sup>5</sup>

## L2

DRIVER RETAINED · ATTENTION NOT ENGINEERED

*the system works well most of the time — and then it does not*

TESLA AUTOPILOT — LEVEL-2 MONITORING AS A SUSTAINABLE ROLE

*Contributing to the car driver's overreliance on the vehicle automation was its operational design, which permitted his prolonged disengagement from the driving task and his use of the automation in ways inconsistent with guidance and warnings from the manufacturer.*

NTSB HIGHWAY ACCIDENT REPORT HAR-17/02 (WILLISTON, FLORIDA CRASH), 2017

REFERENCES

1. NTSB, Highway Accident Report HAR-17/02 (Williston, FL, 2016) — the quoted disengagement finding.

2. NHTSA Standing General Order 2021-01 reports — documented Autopilot fatal crashes.

3. NTSB, Highway Accident Report HAR-20/01 (Mountain View, CA, 2018) — Autopilot crash analysis.

4. NHTSA Office of Defects Investigation EA22-002 (opened 2022; closed April 2024, superseded by recall 23V-838) — driver-engagement adequacy.

5. Parasuraman, R. & Manzey, D. (2010) — automation complacency and monitoring.

6. Benavides v. Tesla (S.D. Fla.) — August 2025 federal jury verdict apportioning Tesla 33% of responsibility for the April 2019 Key Largo Autopilot crash, \$243 million; post-trial motions denied February 2026.

7. NHTSA recall-query RQ24009 (April 2024) into the adequacy of the December 2023 Autopilot remedy, and investigation PE24031 (October 2024) into Full Self-Driving; Tesla's June 2025 Austin robotaxi launch and the ensuing regulator scrutiny.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Tesla Autopilot at consumer scale is the largest live test of Level-2 monitoring as a sustainable role. The early evidence is that it is not. The case is the strongest currently available test of whether consumer-side training and engagement design can produce sustained attention to automation. The recurring fatality pattern suggests the answer.

LENS APPROACH

LENS uses Tesla Autopilot in LEN 2 as the live consumer-scale test of monitoring as a sustainable role and in LEN 7 for the governance dynamics of a Level-2 system marketed at the boundary of higher autonomy. Studio projects examine the driver-engagement design that would make the role performable.

| DOMAIN                     | MODES AT STAKE |   |   |   | LENS COMPETENCY |   |    |    |    | PROBLEM TYPE | INDUCED | LEO |     |   |
|----------------------------|----------------|---|---|---|-----------------|---|----|----|----|--------------|---------|-----|-----|---|
| AUTONOMOUS SYSTEMS         | →              | T | N | G | H               | → | D1 | D2 | D3 | D4           | D5      | 6   | 3.2 | 3 |
| Human-System Collaboration |                |   |   |   |                 |   |    |    |    |              |         |     |     |   |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Engineer driver-engagement monitoring to the demonstrated limits of human vigilance — verify attention against operational evidence before fielding, not against an assumption that warnings suffice.
- ▶ Name and present the feature so its capability boundary is unmistakable to an untrained consumer, rather than implying autonomy the system does not deliver.
- ▶ Constrain operation to the design domain the system can actually handle, so the human is not silently relied on as the backstop for edge cases.

IN OPERATION / AFTER

- ▶ Monitor the standing-order crash data for the disengagement pattern and treat a recurring signature as evidence the role, not the driver, needs redesign.
- ▶ Track attention and takeover performance across years of ownership, since vigilance erodes with the very reliability that accumulates over time.
- ▶ Hold the engagement design accountable to an independent regulator with authority to require changes when in-use evidence shows it is inadequate.

## REFLECTION QUESTIONS

1. Identify a passive-monitoring role in your domain. What evidence would tell you whether attention is sustainable over years of operation?
2. Design the driver-engagement architecture that would make Level-2 monitoring sustainable for an average consumer.
3. Autopilot assigns the operator legal responsibility for the driving task while engineering the conditions that erode their attention. Where in your domain does liability rest with the person an automated system has made least able to intervene — and how would you realign the two?

**WHO BUILDS THIS** learning science & instructional design · safety science & organizational analysis · governance, ethics & measurement · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards · governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs

# Fintech Lending Fairness Audit — When Including Race Reduces Disparity

**D** Designed Out **G** Governance & Trust **N** Normalization of Deviance

**EVIDENCE TIER** preprint-tier — source confidence flagged; future validation ongoing.

**IMPACT** A fintech-lending fairness audit finds a consumer-credit model miscalibrated by group when the protected attribute is withheld, and shows that using the attribute to correct the calibration reduces the disparity — surfacing the next teaching point: competing fairness definitions disagree, and the choice is a judgment

## IN BRIEF

The Coots et al. (2025) fintech lending fairness audit, posted as a preprint, examines a deployed consumer-lending model and finds it miscalibrated by group when the protected attribute is withheld from inputs — underestimating default risk for some borrowers and overestimating it for others. It then shows that using the protected attribute in a controlled fashion to correct the calibration reduces the resulting disparity — a small-scale echo of the pattern Bartlett documented at the mortgage-market scale. The case is a frontier candidate: it sharpens the teaching point that “fairness through unawareness” is not the conservative or safe choice it is often assumed to be, and it surfaces the next layer — competing fairness definitions disagree about what the same model is doing, and the practitioner has to decide **which definition** the deployment is being held to before the measured disparity can even be interpreted. The evidence-tier flag is preprint: the finding has not yet completed peer review, the metric choices and mitigation specifics may move, and future validation will continue. The teaching point survives those caveats, which is why the case is included with the flag rather than held.

## THE SHIFT

The lending pair (Cases 186 and 196) takes the practitioner past the first equity intuition — **just don’t use the protected attribute** — and into the harder territory the equity literature now operates in. Bartlett showed that omission preserves the disparity through correlated features. The Coots audit shows the next thing: the model is miscalibrated by group when the attribute is withheld, and using it to correct the calibration lowers the resulting disparity. Both findings are about the same family of models; they disagree about what makes a model fair.<sup>o</sup>

## WHAT IS EMERGING

The audit is a profit-and-calibration analysis of a deployed fintech-lending model. Under unawareness — the protected attribute withheld from inputs — the model is miscalibrated by group, underestimating default risk for some borrowers and overestimating it for others. Using the protected attribute to correct that calibration reduces the resulting disparity — small scale, but the pattern matches the mortgage-finance finding (Case 186) and the broader fair-ML literature that omission does not equal fairness.<sup>1</sup>

## THE CAPABILITY QUESTION

The case is a frontier candidate because it surfaces the layer Bartlett alone does not reach: competing fairness definitions disagree about the same model. A model that is more equitable on group calibration may be less equitable on equalized odds, and vice versa; the impossibility results of the fair-ML literature (Chouldechova; Kleinberg, Mullainathan, & Raghavan) make this explicit. The Coots audit is concrete enough to ground the choice: the practitioner cannot postpone the question of **which definition** the deployment is being held to.<sup>2</sup>

## EARLY EVIDENCE

The evidence-tier flag matters here and is rendered under the case title. The Coots audit is a preprint; the metric choices, the mitigation specifics, the dataset window, and the reported magnitude may move in peer review. The structural pattern — that using the protected attribute to correct group miscalibration can reduce disparity relative to omission — is consistent with the broader fair-ML literature and with Bartlett's mortgage-finance finding, but the specific magnitudes in the audit should be treated as the strongest current preprint claim, not as a settled fact. Future validation will continue, and this case will be updated when peer review lands.<sup>3</sup>

## OPEN PROBLEMS

What the pair (Cases 186 + 196) teaches together is the form of the equity capability deliverable: the practitioner must specify, in advance, the fairness definition the deployment will be evaluated against, audit on outputs rather than reasoning about inputs, and decide on judgment that the trade-offs across competing definitions are acceptable for the deployment context. This is the case-grounded basis for the subobjective **Fairness beyond omission** (4.6) and the subobjective **Judgment under inadequate evidence** (4.4) — the audit is itself a worked example of deciding under irreducible disagreement.<sup>4</sup>

*Once you accept that omission is not the answer, the next question — which fairness definition — is the one the deployment cannot avoid.*

EDITORS' SYNTHESIS OF COOTS ET AL. (2025) AND MITCHELL ET AL. (2021).

## REFERENCES

1. Coots, Bartlett, Nyarko, & Goel (2025), "Algorithmic Bias in Lending: Evidence from a Fintech Audit," arXiv:2512.20753 — audit of 80,000 personal loans showing model miscalibration disparities and that controlled inclusion of protected attributes could correct them; the evidence-tier flag is binding until peer review completes.
2. Kleinberg, Mullainathan, & Raghavan (2017), "Inherent Trade-Offs in the Fair Determination of Risk Scores," ITCS — competing fairness definitions are not jointly achievable.
3. Mitchell, Potash, Barocas, D'Amour, & Lum (2021), "Algorithmic Fairness: Choices, Assumptions, and Definitions," Annual Review of Statistics and Its Application 8:141–163 — practitioner-facing survey.

2.

Bartlett, Morse, Stanton, & Wallace (2022), “Consumer-lending discrimination in the FinTech era,” *Journal of Financial Economics* 143(1):30–56 — the paired big-tier case (186).

3.

Chouldechova (2017), “Fair Prediction with Disparate Impact,” *Big Data* 5(2):153–163 — the impossibility result for calibration and error-rate parity.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Coots audit is the small-tier frontier instance of the finding that using a protected attribute to correct group miscalibration can reduce disparity relative to omission. With Bartlett (Case 186) it forms the canonical lending pair: omission does not fix the harm; competing fairness definitions disagree about what fix is. Evidence is preprint-tier; future validation will continue.

LENS APPROACH

Coots is the small-tier frontier counter-case to Bartlett. LENS uses the pair in Domain 4 (Test and Evaluation) for the subobjective **Fairness beyond omission** (4.6); in Domain 4 again for the subobjective **Judgment under inadequate evidence** (4.4) (the pair is itself a decision under irreducible disagreement); and in Domain 3 (Human-System Collaboration) for delegation of consequential consumer-finance decisions to a model. The preprint-tier flag is binding until peer review completes.

| DOMAIN        | MODES AT STAKE | LENS COMPETENCY     | PROBLEM TYPE | INDUCED | LEO   |
|---------------|----------------|---------------------|--------------|---------|-------|
| FINANCE       | D G N          | D1 D2 D3 D4 D5      | 6            | 8.1     | 4 · 3 |
| ALGO FAIRNESS |                | Test and Evaluation |              |         |       |
| FINTECH       |                |                     |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify, before model selection, which fairness definition the deployment is being held to — group calibration, equalized odds, demographic parity — and the metric values that will count as acceptable.
- If the protected attribute is included under a controlled-mitigation regime, document the inclusion path (training-time, post-processing) and the auditing pipeline that verifies the mitigation does what it claims.
- Distinguish what is preprint-tier evidence from what is settled in the literature you are drawing on, so a deployment decision does not ride on the strongest unreviewed claim.

IN OPERATION / AFTER

- Audit the deployed model on outputs stratified by protected attribute, with the pre-registered fairness definition; publish the metric values and the cases that disconfirm them.
- Track the preprint to publication; when peer review lands, update the auditing pipeline to reflect the settled metric choices and any revised magnitudes.
- When competing fairness definitions disagree, name the trade-off explicitly in the deployment documentation; do not present the chosen definition as the technical optimum.

REFLECTION QUESTIONS

- Identify a fairness audit in your domain conducted at preprint or unpublished stage. What part of its claim survives if peer review modifies the metric choices? What part is contingent on the specific magnitudes?
- Specify which fairness definition your deployment is being held to before the audit is run. What trade-off — across calibration, equalized odds, demographic parity — does that choice accept?
- Coots’ finding is consistent with Bartlett (Case 186) and with the broader fair-ML literature, but the specific magnitudes are preprint-tier. What is the minimum follow-up evidence you would require before allowing this case to drive an operational decision in your context?

**WHO BUILDS THIS** systems & human-factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

# Predictive Policing — PredPol

**G** Governance & Trust   **H** Human-System Interface   **D** Designed Out

**IMPACT** Predictive-policing tools deployed across scores of U.S. police jurisdictions; Santa Cruz banned them outright in 2020 and other cities have since dropped or restricted their programs

## IN BRIEF

PredPol and similar predictive-policing tools use historical crime data to forecast where future crime is likely, directing patrols to those locations. A randomized field trial by the tool's own co-founder (Brantingham et al. 2018) found no significant arrest-rate disparity by racial-ethnic group between algorithm-directed and standard patrol; the structural objection is not to the arrest counts but to the data. Multiple analyses (Lum & Isaac 2016; Richardson, Schultz & Crawford 2019) found that because the historical data records where police have enforced, not where crime has occurred, the algorithm tends to reinforce existing enforcement patterns rather than predict underlying crime — a feedback loop that concentrates policing on already-over-policed neighborhoods. Cities including Santa Cruz, New Orleans, and Los Angeles have since suspended or abandoned predictive-policing deployments — Santa Cruz on the bias case, Los Angeles on pandemic budget cuts, New Orleans after the secrecy of its Palantir contract became public. The capability gap is at the construct definition: “where crime occurs” and “where crime is recorded” are different variables, and the tools treated them as one. It is the canonical algorithmic-governance case in U.S. policing.

## THE SHIFT

Police departments have adopted data-driven tools that forecast where crime will occur and direct patrol resources accordingly. PredPol and its peers promised to make policing more objective by replacing officer intuition with statistical prediction — relocating a discretionary judgment into an algorithm trained on historical records, and lending the output a veneer of neutrality that the human judgments embedded in those records did not actually possess.<sup>o</sup>

## WHAT IS EMERGING

Researchers examining these systems found a structural flaw. PredPol's place-based forecasts ran on reported crime incidents — type, location, date and time — not on arrest records. The feedback concern survives the correction and changes shape: reporting is itself uneven across neighbourhoods, and directing patrols to forecast cells generates more recorded incidents there, so the model can learn the distribution of police attention alongside the distribution of crime. Lum & Isaac (2016) showed the result is a feedback loop: patrols are sent where police already went, generating more recorded incidents that confirm the prediction — a loop that grows more confident the longer it runs, because its own output becomes the next cycle's evidence. The strongest counter-evidence is a randomized controlled trial across LAPD divisions (Brantingham, Valasik & Mohler 2018),



which found no statistically significant difference in the racial-ethnic proportion of arrests between algorithm-directed and standard-patrol conditions, and arrest totals that fell or held steady; the lead author co-founded PredPol, which is why the finding belongs here with its affiliation attached rather than either suppressed or treated as settling the question.<sup>1</sup>

#### THE CAPABILITY QUESTION

The capability gap is at the construct definition. “Where crime occurs” and “where crime is recorded” are not the same variable, yet predictive-policing tools have treated them as interchangeable. The question is whether a model trained on a record of institutional behavior can ever predict the underlying phenomenon, or only amplify the behavior that produced its data — a question no amount of modeling accuracy can answer, because the gap is in what the data measures, not in how well the model fits it.<sup>2</sup>

#### EARLY EVIDENCE

Richardson, Schultz & Crawford (2019) documented “dirty data” — records produced during periods of biased or unlawful policing — feeding directly into predictive systems. Santa Cruz banned predictive policing outright in June 2020, the first U.S. city to do so; Los Angeles had dropped PredPol that April citing pandemic budget cuts, a year after its Inspector General warned against drawing strong conclusions from the available statistics; New Orleans let a Palantir contract lapse in 2018 once the arrangement became public. Abandonment came only after the tools were already in service, the construct problem surfaced by external review rather than caught before the systems shaped where officers were sent.<sup>3</sup> The product itself did not outlast the critique: PredPol rebranded as Geolitica in 2021, and after a 2023 investigation by The Markup found its predictions for one New Jersey department accurate less than 0.5% of the time, SoundThinking (formerly ShotSpotter) hired its engineering team and moved to acquire some of its intellectual property, and Geolitica ceased operations at the end of 2023 — the construct-validity failure ending in commercial wind-down rather than correction.<sup>4</sup>

#### OPEN PROBLEMS

Predictive policing is the canonical algorithmic-governance case in U.S. policing and pairs with COMPAS (Case 187) and educational algorithmic bias (Case 46). The open problem is a construct-validity audit — a way to establish, before deployment, whether a predictive system’s training data is a record of ground truth or merely of institutional behavior — implemented in some jurisdictions and absent in most, so the same construct error remains available to the next department that mistakes a record of enforcement for a map of crime.<sup>5</sup>



*It is predicting future policing, not future crime.*

LUM, K. & ISAAC, W., "TO PREDICT AND SERVE?", SIGNIFICANCE 13(5), 2016, P. 17

REFERENCES

1. Lum, K. & Isaac, W. (2016), "To Predict and Serve?," Significance — the enforcement-vs-crime feedback loop (paraphrased).

2. Brantingham, P. J., Valasik, M., & Mohler, G. O. (2018), "Does Predictive Policing Lead to Biased Arrests? Results From a Randomized Controlled Trial," Statistics and Public Policy 5(1):1–6 — no statistically significant difference in arrests by racial-ethnic group between treatment and control; the lead author co-founded PredPol.

3. Richardson, Schultz & Crawford (2019), "Dirty Data, Bad Predictions" — biased records feeding predictive systems.

4. Municipal records on suspension and abandonment of predictive policing (Santa Cruz, New Orleans, Los Angeles).

5. Sankin, A. & Mattu, S. (2023), "Predictive Policing Software Terrible at Predicting Crimes," The Markup, 2 October 2023 (copublished with Wired) — the Plainfield, NJ accuracy analysis; and reporting on PredPol/Geolitica's 2021 rebrand and 2023 wind-down into SoundThinking.

6. Brayne, S. (2017), "Big Data Surveillance: The Case of Policing," American Sociological Review 82(5):977–1008 — ethnographic account of the LAPD's adoption of predictive and dragnet analytics.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Predictive policing is the canonical case for the difference between training data and ground truth at law-enforcement scale. The institution being predicted (where crime occurs) is not the institution producing the training data (where police make arrests). The capability gap is at the construct.

LENS APPROACH

LENS uses predictive policing in LEN 7 as a foundational AI- governance case in policing and in LEN 9 for the technical construct-validity analysis. The case is paired with COMPAS as criminal-justice algorithmic cases of different kinds.

| DOMAIN     | MODES AT STAKE      |   |   | LENS COMPETENCY |   |    |    |    | PROBLEM TYPE | INDUCED | LEO |     |   |
|------------|---------------------|---|---|-----------------|---|----|----|----|--------------|---------|-----|-----|---|
| GOVERNMENT | →                   | G | H | D               | → | D1 | D2 | D3 | D4           | D5      | 5   | 8.1 | 4 |
| TECHNOLOGY | Test and Evaluation |   |   |                 |   |    |    |    |              |         |     |     |   |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Establish a construct-validity audit before deployment that tests whether the training data records the phenomenon to be predicted or merely the institution's own behavior.
- ▶ Define the prediction target explicitly as "where crime occurs," not "where arrests are recorded," and reject data that cannot speak to the former.
- ▶ Engineer against the feedback loop — for example, decoupling patrol allocation from the data the model then re-ingests — before the system can amplify its own output.

IN OPERATION / AFTER

- ▶ Monitor for the enforcement feedback loop in deployment, watching whether predicted areas simply accumulate more recorded incidents that confirm the prediction.
- ▶ Require periodic equity review of in-use outcomes, since the construct flaw in these systems has surfaced through external analysis rather than internal metrics.
- ▶ Track the provenance of incoming training data and quarantine records produced during periods of biased or unlawful enforcement.

## REFLECTION QUESTIONS

1. Identify a predictive system in your domain whose training data is itself a record of institutional behavior rather than ground truth. What is the construct gap?
2. Design the construct–validity audit that should precede deployment of a predictive system in any institutional setting.
3. Predictive policing’s feedback loop grows more confident the longer it runs, because its own output becomes the next cycle’s training data. Where in your domain does a deployed model shape the data it later learns from — and how would you break the loop?

**WHO BUILDS THIS** governance, ethics & measurement · human factors & interaction design · systems & human–factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs · task analysis · design prototyping

# Launching the BRAIN Initiative — Governance of a Grand Challenge

**G** Governance & Trust

**K** Knowledge & Institutional Memory

**N** Normalization of Deviance

**DISCLOSURE** Institutional overlap: an editor works in the connectomics field adjacent to this case and shares institutional ties to the BRAIN research community; no editor authored or contributed to this work. Anchored to widely cited primary position papers and independent journalism, not the editor's own work; the critical retrospective is the deliberate safeguard.

**IMPACT** A multi-billion-dollar national research endeavor launched via a position-paper-to-policy iteration sequence, with governance contestation on the public record and a 2021 critical retrospective documenting that the unified-understanding framing exceeded what the science delivered

## IN BRIEF

The BRAIN Initiative is one of the few large-program launches in the corpus whose governance trail is publicly documented end-to-end: a 2011 Kavli symposium produced a six-author Neuron position paper (Alivisatos et al., 2012); the proposal was shepherded to the White House OSTP, became a Presidential initiative in 2013, and was operationalized by an NIH working group whose BRAIN 2025: A Scientific Vision report (2014; reviewed in peer-reviewed form by the group as Jorgenson et al., 2015) set milestones and cost estimates. What makes it teachable rather than triumphalist is the governance contestation visible in the same record. Yuste worried the advisory panel was “packing the committee with users, rather than tool builders” — diluting the original focus; Bargmann, who later co-chaired the working group, had written in February 2013 that “there is great concern in the neuroscience community that this sounds like a big central planning project that will take resources away from creative work.” A 2021 MIT Technology Review retrospective assessed that the big-science brain projects, BRAIN included, did not deliver the unified understanding their framing promised. The case pairs with the EU Human Brain Project (Case 181) as a contrast in governance models — distributed working-group versus top-down single-PI — at the same era, the same ambition, and opposite trajectories.

## BACKGROUND

In 2011 a small group of neuroscientists convened at a Kavli Foundation symposium argued that the field had a tractable grand-challenge target: recording the activity of large numbers of neurons across whole circuits. The argument was published in Neuron in 2012 as the Brain Activity Map proposal (Alivisatos, Chun, Church, Greenspan, Roukes, and Yuste) and was shepherded through Kavli's Miyoung Chun to the White House Office of Science and Technology Policy. The position paper named a capability the field could be organized

around and a sequence of tools that would have to be built; it was a launch artifact rather than a results report.<sup>9</sup>

### THE INTERVENTION

In April 2013 President Obama announced the BRAIN Initiative. The operational governance was an NIH working group, co-chaired by Cori Bargmann and Bill Newsome, charged with translating the vision into a milestone-and-cost-bearing plan. The group released the BRAIN 2025: A Scientific Vision report in June 2014, reviewed the following year in peer-reviewed form by members of the group (Jorgenson et al., *Phil. Trans. R. Soc. B*, 2015); it named seven priority areas, set timelines, and laid out funding ranges. Among the few large research-program launches in the corpus, this is one whose position-paper-to-policy-to-implementation sequence is openly auditable — every step has a published artifact attached.<sup>1</sup>

### HOW IT WORKED

What makes the case teachable rather than triumphalist is the governance contestation visible in the same record. Yuste, one of the original six authors, voiced concern that the advisory panel was being expanded with users of the tools the program was meant to build, rather than the tool-builders the original proposal had centered — a documented dilution of scope. Bargmann, before her appointment, had relayed in a February 2013 email the neuroscience community's concern that a big central planning project would draw resources away from creative work. The governance choices — who leads, tool-builders versus users, central plan versus distributed creativity — were not made in private and then defended; they were litigated in the public record while the program ran.<sup>2</sup>

### THE EVIDENCE

The 2021 MIT Technology Review retrospective took the ten-year view: big-science brain projects, BRAIN among them, did not deliver the unified understanding their founding framing had promised. The honest assessment is that the initiative produced tools, atlases, and a coordinated funding stream — meaningful capability — while drifting from the grand-challenge framing the position paper had used to mobilize political support. The case is the instructive form of the “enthusiasm-evidence gap” at field scale: the framing carried the politics, and the science delivered something different and more diffuse.<sup>3</sup>

### WHAT TRANSFERRED

What the case teaches is that large-program governance is the capability deliverable — not the framing, not the early tools, not the eventual results. The position-paper-to-policy sequence, the working-group composition decision, and the public airing of scope drift are the artifacts a future capability-development program can study. Paired with the EU Human Brain Project (Case 181), the case sets a distributed working-group model against a top-down single-PI one. No published assessment attributes the two trajectories to that difference; the record shows the HBP board replacing Markram's three-member executive committee with a 22-member governing board in 2015, after an open letter from hundreds of European scientists, while BRAIN ran on with funding through 2026. The framing, ambition, and era were comparable; the governance was not.<sup>4</sup>

*Governance contestation in the public record is not program failure. It is what makes the case teachable.*

EDITORS' SYNTHESIS OF THE BRAIN INITIATIVE RECORD (2012 – 2021).

REFERENCES

1. Alivisatos, Chun, Church, Greenspan, Roukes, & Yuste (2012), "The Brain Activity Map Project and the Challenge of Functional Connectomics," *Neuron* 74(6):970–974, doi:10.1016/j.neuron.2012.06.006.

2. Jorgenson et al. (2015), "The BRAIN Initiative: developing technology to catalyze neuroscience discovery," *Phil. Trans. R. Soc. B* 370(1668):20140164, doi:10.1098/rstb.2014.0164 — the BRAIN 2025 plan.

3. Yuste & Bargmann (2017), "Toward a Global BRAIN Initiative," *Cell* 168(6):956–959, doi:10.1016/j.cell.2017.02.023.

4. Underwood, E. (2013), "As White House Embraces BRAIN Initiative, Questions Linger," *Science / ScienceInsider* (April 3, 2013) — source of both quoted spans verbatim: Yuste on "packing the committee with users, rather than tool builders," and Bargmann's February 2013 email to Science on "a big central planning project that will take resources away from creative work."

5. Mullin, E. (2021), "How big science failed to unlock the mysteries of the human brain," *MIT Technology Review* (25 August 2021) — the critical ten-year assessment.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The BRAIN Initiative is one of the few large-program launches in the corpus whose position-paper-to-policy-to-implementation sequence is fully auditable. Governance choices were litigated in the public record, scope drifted from the founding framing, and a critical ten-year retrospective is part of the case materials. Its teaching value is the contested record, not a clean success or scandal.

LENS APPROACH

BRAIN is the field-scale stakeholder-and-governance case in the v2 corpus (induced 5.1; LENS D5+D1/PT4). LENS uses it in Domain 1 (Systems Analysis) for the position-paper-to-policy iteration sequence; in Domain 5 (Navigating Sociotechnical Constraints) for the public-record governance contestation; and in Domain 4 (Test and Evaluation) for the enthusiasm-evidence gap as the framing exceeded delivered science. Direct pair with Case 181 (EU Human Brain Project), same era and ambition, opposite governance model. The COI disclosure under the title is binding: the editor's research adjacency is what motivated the critical retrospective being included as a deliberate counterweight.

| DOMAIN         | MODES ADDRESSED                                | LENS COMPETENCY                                                                   | PROBLEM TYPE | INDUCED | LEO   |
|----------------|------------------------------------------------|-----------------------------------------------------------------------------------|--------------|---------|-------|
| NEUROSCIENCE   | → <span>G</span> <span>K</span> <span>N</span> | → <span>D1</span> <span>D2</span> <span>D3</span> <span>D4</span> <span>D5</span> | 4            | 5.1     | 1 · 5 |
| SCIENCE POLICY |                                                | Navigating Sociotechnical Constraints + D1                                        |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Make the position-paper-to-policy iteration auditable by attaching a published artifact to every step — proposal, OSTP transmittal, working-group composition, milestone plan.
- ▶ Treat working-group composition as a governance act, not a staffing decision: tool-builders vs. users vs. integrators is a framing-shaping choice, and naming the trade-off openly is part of the work.
- ▶ Carry the founding scope through the operational record: name where the program is keeping faith with the founding framing and where it is drifting, while the program is running rather than only in retrospect.

IN OPERATION / AFTER

- ▶ Commission and publish a long-window retrospective that compares the founding framing to delivered capability honestly — not as a closure ritual but as governance evidence the next program can use.
- ▶ When scope drift is documented, decide explicitly whether to reframe the program publicly or to re-baseline against the original framing; the worst case is leaving the gap unaddressed in the record.
- ▶ Treat governance contestation in the public record as program documentation, not as program failure; the contested record is what makes the case teachable.

## REFLECTION QUESTIONS

1. Identify a large research or capability-development program in your domain whose launch artifacts (position papers, working-group reports, milestone plans) are publicly auditable. What governance choices were made openly and which were made in private?
2. Specify the founding-framing-vs-delivered-capability comparison you would publish at year five and year ten of a program you are designing. What is the evidence form, who commissions it, and where does it live in the record?
3. BRAIN survived and adapted while the EU Human Brain Project (Case 181) unraveled. What is the minimum governance documentation that would let a future program-designer learn the difference, rather than reconstruct it from contested press accounts?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

CASE 199

AUTONOMOUS VEHICLES

TRANSPORTATION SAFETY

PUBLIC-SECTOR GOVERNANCE

2023 – 2026

# Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact

**G** Governance & Trust   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** After a California court let Waymo withhold trade-secret-laden safety data from a DMV public-disclosure request, the company answered the governance objection with a published, structured safety case framework — and in November 2025 commissioned the first independent third-party audits of both the safety case and the remote-assistance program

## IN BRIEF

In 2022 a California court permitted Waymo to withhold trade-secret-laden safety details from a public DMV disclosure, leaving regulators and the riding public with a credibility gap Waymo could not close by sharing the contested data. The company's response was to publish, in 2023, a structured **safety case framework**: a top-down argument with explicit claims, sub-claims, and the evidence types each rests on, accompanied by published operating-domain performance figures. In November 2025 Waymo released the first independent third-party audits of both the safety case and the remote-assistance program — the audits themselves disclosed, rather than the underlying trade-secret data the original objection targeted. The pattern is the OU-Analyse / inBloom move (governance objection dissolved by better design) transposed from education into physical-safety AV: the response to an opacity objection was a falsifiable argument structure auditable by third parties, not a defense of opacity. The evidence-tier flag rendered under the title is not a formality — the analysis rests on the practitioner whitepaper plus the 2025 audit summaries, not on a peer-reviewed safety-engineering evaluation. Future validation will continue as the audit cadence and post-deployment failure record accumulate.

## BACKGROUND

The precipitating event was not a crash. In 2022 a California court ruled that Waymo could withhold trade-secret-laden safety details from a DMV public-records process. The company had a legal answer to the disclosure request and no legitimacy answer to the public-trust gap that ruling created. The governance objection — “you are asking us to trust



an opaque system whose failure modes we cannot inspect” — could not be answered by disclosing the contested data without giving up the trade-secret position the court had just protected.<sup>9</sup>

#### THE INTERVENTION

Waymo’s 2023 response was to publish a structured **safety case framework**: a top-down argument with claims and sub-claims, the evidence categories each rests on (operational performance data, simulation and testing, hazard analysis, third-party assessment), and the operational-domain figures available at the time of publication. The artifact’s design move is that the **structure** of the safety argument is public even where individual evidence items remain proprietary — outside auditors can interrogate the chain of reasoning without seeing the trade secrets.<sup>1</sup>

#### HOW IT WORKED

In November 2025 Waymo commissioned and released the results of independent third-party audits of the safety case and of the remote-assistance program. The audits themselves — not the underlying data — were the disclosure artifact. The pattern is OU-Analyse / inBloom in the AV domain: a governance objection answered by a designed legitimacy artifact rather than by disclosure of the contested data. Where opacity could not be defended, a structured falsifiable argument plus audited assurance took its place.<sup>2</sup> The post-deployment failure record the framework anticipated has since begun to accumulate under real regulator action: a December 2025 recall over vehicles not stopping for school buses, a January 2026 NHTSA/NTSB investigation after a Waymo struck a child near a Santa Monica school, and a June 2026 recall of roughly 3,900 vehicles for entering freeway construction zones — the safety-case regime now being stress-tested by the revocation-and-recall machinery it was built to invite rather than resist.<sup>3</sup>

#### THE EVIDENCE

The evidence-tier flag rendered under the case title is doing real work. The framework and audit summaries are practitioner-authored or auditor-authored — not peer-reviewed safety-engineering analyses. Some audit-tier elements push toward investigation-grade, but the synthesis as a whole rests on Waymo and Montreal AI Ethics Institute documents. The honest framing in print is that the source confidence is flagged and that future validation — particularly post-deployment failure-record analysis and continued auditor independence — is ongoing.<sup>4</sup>

#### WHAT TRANSFERRED

The teaching point for LENS is that delegation of consequential decisions to an automated system creates a governance debt that the deploying organization owes the public. The subobjective **Delegation with revocation** (3.3) is the capability the case exercises: the safety case framework is the artifact that makes revocation possible — regulators or auditors can identify which sub-claim has failed, on what evidence, and require the deploying organization to act on the gap. Pair with Case 190 (Cruise) as the foil: the same regulatory regime, the opposite governance choice, the opposite outcome. Pair with Case 200 (CPUC permit framework) as the regulator-side counterpart of the deployer-side artifact.<sup>5</sup>

*Where opacity could not be defended, a structured falsifiable argument plus audited assurance took its place.*

EDITORS' SYNTHESIS OF THE WAYMO SAFETY CASE FRAMEWORK AND THE NOVEMBER 2025 THIRD-PARTY AUDITS.

REFERENCES

1. Cruise / California DMV Order of Suspension (October 2023) — paired Case 190 as the foil.

2. Waymo (2023), “A Blueprint for AV Safety: Waymo’s Toolkit For Building a Credible Safety Case,” whitepaper.

3. Waymo (November 2025), “Independent Audits of Waymo’s Safety Case and Remote Assistance Programs,” summary re-release.

4. NHTSA recall and investigation record (December 2025 school-bus recall; January 2026 NHTSA/NTSB investigation of a Santa Monica child-strike; June 2026 recall of 3,900 vehicles over freeway construction zones) — the post-deployment failure record under regulator action.

5. Montreal AI Ethics Institute (2023), summary and analysis of the Waymo safety case framework.

6. California Public Utilities Commission, AV passenger-service permit framework documents — paired Case 200 for the regulator-side artifact.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Waymo’s safety case framework is the AV-domain instance of the OU-Analyse / inBloom move: a governance objection dissolved by a designed artifact, not by disclosure of the contested data. Evidence-tier flag is practice-synthesis; the artifact pattern is teachable and the third-party audit posture pushes some elements toward investigation-grade, but the synthesis as a whole is practitioner-authored and future validation is ongoing.

LENS APPROACH

Waymo is the AV-safety governance case (induced 5.1; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the subobjective **Delegation with revocation** (3.3) — the safety case is what makes revocation possible — and in Domain 3 (Emerging Systems and Human-System Collaboration) for the deployer-side artifact that permits oversight of a system whose internals are trade secret. Pair with Case 190 (Cruise) as the foil and Case 200 (CPUC) as the regulator-side complement.

| DOMAIN                   | Modes Addressed | LENS Competency                         | Problem Type | Induced | LEO   |
|--------------------------|-----------------|-----------------------------------------|--------------|---------|-------|
| AUTONOMOUS VEHICLES      | G K N           | D1 D2 D3 D4 D5                          | 6            | 5.1     | 5 · 3 |
| TRANSPORTATION SAFETY    | →               | → Navigating Sociotechnical Constraints |              |         |       |
| PUBLIC-SECTOR GOVERNANCE |                 |                                         |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the disclosure objection as a design problem: what falsifiable artifact can you publish that addresses the legitimacy gap without requiring you to surrender trade-secret evidence?
- ▶ Structure the safety case as a top-down argument with explicit claims, sub-claims, and evidence types so an outside auditor can interrogate the **chain of reasoning** rather than only the contested data points.
- ▶ Commission and publish third-party audits of the argument structure and of the operational programs

IN OPERATION / AFTER

- ▶ Treat the safety case framework as a living document — update the claims and evidence as post-deployment failure data accumulates, and publish the updates so the legitimacy artifact does not calcify.
- ▶ Use the subobjective **Delegation with revocation** (3.3): design the framework so a regulator or auditor can identify which sub-claim has failed and trigger a revocation pathway, not only a “trust us, we will fix it” assurance.

(e.g. remote assistance) that the safety case rests on — the audits are the disclosure artifact when the data cannot be.

- Carry the practice-synthesis evidence-tier flag honestly in any program documentation citing the framework — the artifact pattern is teachable, but the magnitude of its public-trust effect is still being measured.

## REFLECTION QUESTIONS

1. Identify an automated system in your context that faces a public-trust objection it cannot answer by full disclosure. What falsifiable argument structure could you publish that would make the system's reasoning auditable without requiring disclosure of the contested data?
2. Specify how a regulator or independent auditor would **revoke** the delegation in your system if a sub-claim of the safety case failed. The subobjective **Delegation with revocation** (3.3) requires this pathway to exist before deployment, not only after a public-facing failure.

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

# CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver

**G** Governance & Trust   **K** Knowledge & Institutional Memory   **D** Designed Out

**EVIDENCE TIER** practice-synthesis-tier — source confidence flagged; future validation ongoing.

**IMPACT** California's Public Utilities Commission built an AV passenger-service permit framework whose Commission-authored conditions — chief among them a mandatory Passenger Safety Plan whose safety measures must be accessible to and apply to all passengers, including those with limited mobility, vision impairments or other disabilities — operationalize the accessibility objection, while the operating limits on hours, geography, speed and weather ride on the DMV-approved operational design domain and the Commission declined San Francisco's request to cap fleet, hours and service area

## IN BRIEF

The California Public Utilities Commission established an AV passenger-service permit framework whose conditions are explicitly designed to address common governance objections — safety, equity, fleet scale — by writing them into the permit rather than treating them as binary deploy / don't-deploy questions. The Commission's own conditions — a passenger notice-and-consent plan, an operational-design-domain attestation, and, decisively, a required Passenger Safety Plan whose safety measures must be accessible to and apply to all passengers, including those with limited mobility, vision impairments or other disabilities — sit on top of a permit whose limits on hours, geography, speed and weather arrive with the DMV-approved operational design domain. Draft DMV rules circulated for comment in 2024 and adopted in April 2026 gave the DMV authority to impose targeted operational restrictions for safety, and the case shows the regime in operation: the Cruise suspension (Case 190) shows the regime can revoke; the Waymo continuation (Case 199) shows it can permit conditionally. The structural complement to the Waymo deployer-side safety case framework is exactly this regulator-side artifact — the permit conditions make the safety case operative as a regulatory instrument. The evidence-tier flag rendered under the title is practice-synthesis: the program is documented in CPUC decisions and program guidance, but no peer-reviewed evaluation of the program's equity-of-service outcomes for disabled and disadvantaged-community riders yet exists. Future validation will continue as ridership and incident data accumulate.

## BACKGROUND

The governance question California faced from 2018 onward was how to regulate AV passenger service in a regime where the California DMV regulates the vehicle and the CPUC regulates the passenger service the vehicle provides. Treating each AV deployment as a binary deploy / don't-deploy decision risked either blocking deployment outright on objections the operator could not satisfy, or permitting deployment without a governance handle on the objections the public and the disability-rights community brought to the proceedings.<sup>0</sup>

## THE INTERVENTION

The CPUC's structural answer was to attach conditions of its own to a permit that otherwise carries the DMV-approved operational design domain. The AV passenger-service permits inherit the DMV's limits on hours, geography, speed and weather, and add a passenger notice-and-consent plan, an attestation and map of that domain, and — the condition that carries the equity objection — a mandatory Passenger Safety Plan, whose measures must be "accessible to and apply to all passengers, including those with limited mobility, vision impairments, or other disabilities." When San Francisco asked the Commission to cap fleet size and restrict hours and service area, Resolution TL-19144 declined: the Commission would "neither modify the DMV-approved ODD ... nor set limits on fleet size." The conditions are not policy aspirations — they are permit terms whose violation triggers enforcement.<sup>1</sup>

## HOW IT WORKED

In August 2024 the DMV released draft regulations for public comment that would strengthen its authority to impose targeted operational restrictions for safety reasons during the deployment lifecycle, not only at initial permit issuance. That authority was finalized: the DMV adopted new AV regulations on 28 April 2026 that codify targeted operational restrictions (fleet size, location, speed, weather), authorize heavy-duty AV trucks and medium-duty transit AVs, add law-enforcement notices of noncompliance, and require AVs to clear emergency zones within two minutes; the CPUC separately opened a successor passenger-service rulemaking in August 2025 covering airport and shared-ride operations, personally-owned AVs, and unaccompanied minors. The regulatory regime can therefore tighten the conditions in response to operational evidence. The pattern: rather than treat the governance objection as binary, the CPUC made the objection itself a design parameter the deployer must satisfy and that the regulator can re-tune.<sup>2</sup>

## THE EVIDENCE

The case is the structural complement to Case 199 (Waymo) and the inverse-outcome companion of Case 190 (Cruise). The Waymo safety case framework is the deployer-side artifact; the CPUC permit conditions are the regulator-side artifact that makes the safety case operative as a regulatory instrument. The Cruise suspension shows the regime can revoke when the disclosure architecture fails; the Waymo conditional continuation shows it can permit when the conditions are satisfied. Together the trio teaches the structural form of regulator-deployer interaction at the AV physical-safety C5 layer.<sup>3</sup>

## WHAT TRANSFERRED

The evidence-tier flag rendered under the title is practice-synthesis. The CPUC decisions, the program guidance, and the permit conditions themselves are public; no peer-reviewed

evaluation of the equity-of-service goals (disabled-rider access, disadvantaged-community access) has yet measured whether the Passenger Safety Plan conditions translate into measured ridership outcomes. The condition-as-objection-dissolver pattern is teachable and the regulatory architecture is explicit, but the outcome evidence is not yet decision-grade. Future validation will continue as the program ages and as the equity outcomes are independently measured.

*The governance objection was not refused. It was made a permit condition.*

EDITORS' SYNTHESIS OF THE CPUC AV PASSENGER-SERVICE PERMIT FRAMEWORK.

REFERENCES

1. California Public Utilities Commission, "Autonomous Vehicle Passenger Service Programs" — program page for the Drivered and Driverless pilot and deployment programs, including the Passenger Safety Plan requirement.

2. California Public Utilities Commission, Resolution TL-19144 (Waymo), issued 11 August 2023 (Commission vote 10 August) — approves driverless deployment in all of San Francisco 24 hours a day per the DMV-approved ODD, and declines to set limits on operating hours, geography or fleet size; the span quoted in the text is verbatim ("the Commission will neither modify the DMV-approved ODD submitted by Waymo ... nor set limits on fleet size").

3. California Department of Motor Vehicles, "New Autonomous Vehicle Regulations Strengthen Oversight and Enforcement, Authorize Trucks and Transit" (28 April 2026) — adopted authority to impose targeted operational restrictions on fleet size, location, speed and weather; draft language circulated for comment 30 August 2024.

4. CPUC permit decisions for Cruise and Waymo, 2020–2024.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The CPUC permit framework is the regulator-side counterpart to the Waymo safety case (Case 199): conditions operationalize objections rather than blocking deployment. Evidence-tier flag is practice-synthesis; the regulatory architecture is documented in CPUC decisions, but no peer-reviewed evaluation of the equity-of-service goals yet exists, and future validation is ongoing.

LENS APPROACH

CPUC is the AV regulator-side governance case (induced 5.1; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for cross-regime governance seams (DMV regulates the vehicle; CPUC regulates the service; both can act) and for the condition-as-design-parameter pattern that makes the deployer-side safety case (Case 199) operative as a regulatory instrument. Pairs with Case 190 as the revocation event under the regime.

| DOMAIN                                                                                                | MODES ADDRESSED                                 | LENS COMPETENCY                                                              | PROBLEM TYPE | INDUCED | LEO   |
|-------------------------------------------------------------------------------------------------------|-------------------------------------------------|------------------------------------------------------------------------------|--------------|---------|-------|
| <div><div>AUTONOMOUS VEHICLES</div><div>PUBLIC-UTILITY GOVERNANCE</div><div>ACCESSIBILITY</div></div> | <div><div>G</div><div>K</div><div>D</div></div> | <div><div>D1</div><div>D2</div><div>D3</div><div>D4</div><div>D5</div></div> | 6            | 5.1     | 5 · 3 |
| <div>→ Navigating Sociotechnical Constraints</div>                                                    |                                                 |                                                                              |              |         |       |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Treat governance objections as design parameters rather than as binary blockers — write them into the permit conditions where compliance is verifiable and violation triggers enforcement.

IN OPERATION / AFTER

- Pair the case with Case 199 (Waymo) as the regulator-deployer artifact dyad — the safety case framework and the permit conditions are co-constitutive, not independent moves.

- ▶ Require a documented Passenger Safety Plan (or equivalent equity-of-service artifact) at initial permit issuance so the equity objection has an operational answer the regulator can audit.
- ▶ Build the regulatory authority to re-tune the conditions during the deployment lifecycle into the regime itself; the Cruise suspension (Case 190) is what re-tuning under a failure event looks like.
- ▶ Use the case in LENS Domain 5 modules on cross-regime governance seams (CPUC and DMV) — the regime structure is itself part of the capability question.
- ▶ Carry the practice-synthesis flag honestly: the condition-design pattern is teachable, but the equity-of-service outcome data are not yet peer-reviewed; future validation requires independent evaluation of the Passenger Safety Plan deliverables.

### REFLECTION QUESTIONS

1. Identify a deployment regime in your context where governance objections are currently treated as binary deploy / don't-deploy questions. How could the objections be rewritten as permit-style conditions whose compliance is verifiable and whose violation triggers enforcement?
2. Specify the equity-of-service artifact your regime would require at permit issuance — the analogue of the Passenger Safety Plan — and the auditable evidence the deployer must provide that the artifact is operational.
3. The case is practice-synthesis tier. What independent outcome evidence — disabled-rider ridership figures, disadvantaged-community access measurements, comparative incident rates by permit condition — would you require before treating the condition-as-objection-dissolver pattern as a validated regulatory architecture?

**WHO BUILDS THIS** governance, ethics & measurement · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · task analysis · design prototyping

# Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India

**G** Governance & Trust   **N** Normalization of Deviance   **H** Human-System Interface

**EVIDENCE TIER** journalism-tier — source confidence flagged; future validation ongoing.

**IMPACT** India's Aadhaar biometric-identity system, with more than 1.2 billion numbers issued and some 800 million people drawing subsidised food under it, produced operational exclusion when authentication failed at the ration shop or pension window; the 2018 Puttaswamy II majority upheld the Section 7 welfare-linkage while Justice Chandrachud's dissent held that exclusion from technological error — through no fault of the individual — violates constitutional dignity, naming a principle the executive had already reached for in UIDAI's October 2017 exception-handling circular directing that no beneficiary be denied a benefit for want of Aadhaar

## IN BRIEF

Aadhaar — India's biometric digital-identity system, with more than 1.2 billion enrolled — was designed to streamline welfare delivery and reduce identity fraud. Across more than a decade of deployment the failure mode became operational exclusion: when biometric authentication failed for an individual (worn fingerprints, missing iris match, connectivity loss at the point of service), the welfare to which the person was entitled became inaccessible. The constitutional reckoning came in the 2018 Puttaswamy II Aadhaar judgment: the 4–1 majority upheld Section 7 of the Aadhaar Act, which conditions subsidies and benefits on authentication, while Justice D. Y. Chandrachud's lone dissent named the decisive principle — that “exclusion based on technological errors, with no fault of the individual, is a violation of dignity.” The empirical predicate is the documented exclusion record: the 2017 Jharkhand starvation deaths (Santoshi Kumari, 28 September 2017, one of 57 starvation deaths the Right to Food Campaign investigated across nine states from 2015, of which at least nineteen were directly linked to Aadhaar) and the Drèze–Khalid–Khera–Somanchi field study finding over forty percent of non-transacting households citing Aadhaar-authentication problems. The correction that followed was executive and doctrinal rather than a majority holding — circulars barring denial of rations for want of Aadhaar, and, in the distinct 2025 Pragya Prasun ruling, a right to accessible digital access read into Article 21 in the e-KYC context. The case is the non-US automated-welfare-delegation case the corpus needs alongside SyRI and the UK Post Office Horizon thread. The evidence tier is mixed: the judgments are investigation-grade; the lived-exclusion sourcing rests on journalism and advocacy reporting, with the field data peer-reviewed. Future validation continues on whether the fallback requirement is honored at the operator interface. Gap-5 non-US/UK/EU case (India).



## BACKGROUND

Aadhaar was designed as a population-scale solution to two problems the Indian welfare state had identified by the late 2000s: fragmented identification across welfare schemes and identity fraud in benefits delivery. The design choice was biometric authentication at the point of service — fingerprints and iris — keyed to a unique twelve-digit identifier issued by the UIDAI. Enrollment passed 1.2 billion. The capability the system promised was streamlined welfare access; the authority it required was delegation of identity verification to the biometric.<sup>0</sup>

## THE INTERVENTION

The failure mode emerged across deployment: operational exclusion. When biometric authentication failed at the ration-shop card-reader or the pension window — worn fingerprints from manual labor, missing iris matches, connectivity loss, device malfunction — the welfare to which the person was entitled became inaccessible. The failure was not the absence of identity; the person had the Aadhaar number. The failure was the brittleness of the authentication system at the operator interface. The people most dependent on the welfare were the people whose biometrics were most likely to fail.<sup>1</sup>

## HOW IT WORKED

The constitutional reckoning came in stages. In the 2018 Puttaswamy II judgment the Supreme Court's 4–1 majority upheld Aadhaar and Section 7 of the Aadhaar Act — the provision conditioning subsidies and benefits on authentication — reasoning that targeted delivery and leakage-reduction were legitimate, proportionate state aims. Justice D. Y. Chandrachud, dissenting alone, named the structural risk the majority did not: that Section 7 was over-broad and that “exclusion based on technological errors, with no fault of the individual, is a violation of dignity.” Across the following years the documented exclusion record accumulated, and the correction arrived not as a reversal of the majority but as executive circulars barring denial of benefits for want of Aadhaar — UIDAI's October 2017 exception-handling circular, issued to states and ministries weeks after Santoshi Kumari's death, is the first of them — and, in the distinct 2025 Pragya Prasun ruling (30 April 2025), a right to accessible digital access read into Article 21 in the e-KYC context for persons with disabilities.<sup>2</sup>

## THE EVIDENCE

The evidence-tier flag is split and decisive. The court judgments themselves are investigation-grade — the published opinions are the record. The empirical body rests on two tiers: the lived-exclusion account is journalism and advocacy reporting (the Right to Food Campaign's compilation of the Jharkhand starvation deaths; the Indian press), while the exclusion-rate evidence is peer-reviewed field work — Drèze, Khalid, Khera and Somanchi's Jharkhand survey (“Pain without Gain?”, EPW 2017) found over forty percent of non-transacting households citing Aadhaar-authentication problems, and Ranchi district's own point-of-sale records put about twenty percent of cardholders unable to transact in an average month. The journalism-tier flag travels with the lived-exclusion material; the field data and the judgments carry their own, higher, tiers. Whether the fallback requirement is honored at the operator interface is the open empirical matter the future will validate.<sup>3</sup>

## WHAT TRANSFERRED

What the case adds at the LENS layer is the non-US automated-welfare-delegation thread the corpus needs. The Domain-5 frame applies — a sociotechnical constraint (population-scale biometric delegation) tested against constitutional dignity — and the LEO around fairness beyond omission applies directly: the harm was not the absence of the system but the operational exclusion produced by it. The LEO on delegation with revocation applies too: the correction did not revoke the delegation but sought to bound it with a fallback requirement — and the gap between the Chandrachud dissent's principle and the majority's holding is precisely why the revocation remains contested rather than settled. The case pairs with SyRI (Dutch welfare-fraud system halted on rights grounds) and the UK Post Office Horizon thread as the global lineage of judicial reckoning with automated welfare delegation.<sup>4</sup>

*Exclusion based on technological errors, with no fault of the individual, is a violation of dignity.*

JUSTICE D. Y. CHANDRACHUD, DISSENTING, PUTTASWAMY II (2018)

## REFERENCES

1. Supreme Court of India (2018), Justice K. S. Puttaswamy (Retd.) v. Union of India, (2019) 1 SCC 1 — the Aadhaar judgment; 4–1 majority upholding Section 7, with Justice D. Y. Chandrachud's dissent on technological exclusion and dignity. The quoted span is verbatim from ¶262 of the dissent (Part H) — do not re-quote it from the judgment summaries, which paraphrase it.
2. Drèze, J., Khalid, N., Khera, R., & Somanchi, A. (2017), "Aadhaar and Food Security in Jharkhand: Pain without Gain?," *Economic & Political Weekly* 52(50), 16 December 2017 — field study of biometric-authentication exclusion in the PDS; more than 40% of non-transacting households cited ABBA-related problems in the month before the survey, and Ranchi district point-of-sale records put about 20% of cardholders unable to transact in an average month over January–July 2017.
3. Right to Food Campaign compilation and Indian journalism (The Hindu, The Wire, Scroll.in, 2017–2019) on the Jharkhand starvation deaths (Santoshi Kumari, 28 September 2017) — lived-exclusion sourcing with journalism-tier flag. The compilation runs to 57 deaths across nine states from 2015; 27 are reported as Aadhaar-related and at least 19 as directly linked to Aadhaar exclusion. The case carries the narrower 19 deliberately.
4. Supreme Court of India (2025), Pragma Prasun & Ors. v. Union of India, 2025 INSC 599 (30 April 2025) — the distinct, later ruling reading a right to accessible digital access into Article 21 in the e-KYC context for persons with disabilities.
5. Dixon, P. (2017), "A Failure to 'Do No Harm' — India's Aadhaar biometric ID program and its inability to protect privacy in relation to measures in Europe and the U.S.," *Health and Technology* 7(4):539–567 — comparative privacy and data-protection analysis, reporting biometric-authentication failure rates alongside it.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

Aadhaar exclusion is the non-US automated-welfare-delegation case the corpus needed: an operational-exclusion failure mode whose constitutional statement is Justice Chandrachud's 2018 dissent — exclusion from technological error, no fault of the individual, as a dignity violation — set against a majority that upheld the Section 7 linkage. The correction that followed was executive and doctrinal rather than a majority holding. Evidence tier is split: judgments investigation-grade, field data peer-reviewed, lived-exclusion sourcing journalism plus advocacy. Future validation ongoing on whether the fallback requirement is honored.

### LENS APPROACH

The designed teaching point is operational failure-mode attribution (induced 5.2, with 8.1 as the secondary construct reading; LENS D5/PT5). At the ration-shop card-reader and the pension window, a biometric authentication that did not match was logged as user error — the claimant's worn fingerprints, the elderly applicant's failed iris read — when it was in fact a designed exclusion mode of the system: manual laborers

and the elderly are precisely the populations whose biometrics the design could not reliably read, so the failure belongs to the authentication design, not to the excluded person. The capability discipline is to attribute an authentication failure to the system that produced it rather than to the user it shut out, and to name who the design predictably excludes before deployment. Justice Chandrachud’s dissent performed exactly this attribution — naming technological exclusion, through no fault of the individual, as the system’s dignity violation — even as the majority upheld the linkage; the correction has since been pursued through executive fallback mandates rather than a settled revocation. Pair with SyRI (Case 189) and the UK Post Office Horizon thread.

| DOMAIN                   | MODES ADDRESSED | LENS COMPETENCY | PROBLEM TYPE                          | INDUCED | LEO |
|--------------------------|-----------------|-----------------|---------------------------------------|---------|-----|
| DIGITAL IDENTITY         | G N H           | D1 D2 D3 D4 D5  | 5                                     | 5.2     | 1-5 |
| WELFARE                  | →               | →               | Navigating Sociotechnical Constraints |         |     |
| BIOMETRIC AUTHENTICATION |                 |                 |                                       |         |     |
| INDIA                    |                 |                 |                                       |         |     |

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the operator-interface brittleness of any biometric system as a design variable from the start: who is most likely to be excluded, and what is the alternative at the point of service.
- ▶ Specify the fairness-beyond-omission test the system must pass before deployment: not “is the system available,” but “does it work for the people most dependent on the welfare it gates.”
- ▶ Document the delegation explicitly — what authority is delegated to the biometric, under what conditions it can fail, and what the operator’s authority to override is — so the revocation question is answerable later.

IN OPERATION / AFTER

- ▶ Track operational exclusion as a primary outcome metric, not a downstream complication; the case teaches that the exclusion is the failure mode the design ignored.
- ▶ Carry the journalism-tier flag on the lived-exclusion material and the investigation-grade flag on the court judgments separately; the case is teachable because the two tiers are explicit.
- ▶ Pair with SyRI and the UK Post Office Horizon thread when teaching automated-welfare delegation; the cross-jurisdictional pattern is the lineage of judicial correction of these systems.

REFLECTION QUESTIONS

1. Identify a delegated authentication or eligibility-check system in your context whose failure mode is operational exclusion of the people most dependent on the service it gates. What is the alternative at the point of service, and is it actually available?
2. Specify the fairness-beyond-omission test such a system would have to pass before deployment. The Aadhaar pattern teaches that the harm is not the absence of the system but the exclusion the system produces — what would the equivalent test be in your domain?
3. The case has split evidence tiers: court judgments investigation-grade, lived-exclusion sourcing journalism plus advocacy plus comparative analysis. How would you carry the tier split into a teaching artifact without smoothing either tier away?

**WHO BUILDS THIS** governance, ethics & measurement · safety science & organizational analysis · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards · usability testing · cognitive walkthroughs

# NYC Local Law 144 – Bias Audits as Governance Artifact

**D** Designed Out   **G** Governance & Trust   **K** Knowledge & Institutional Memory

**IMPACT** New York City Local Law 144 of 2021, implemented through Department of Consumer and Worker Protection rules effective July 5 2023, requires employers using "automated employment decision tools" (AEDTs) for hiring or promotion decisions in NYC to conduct annual independent bias audits, publish audit summaries, and provide candidate notice; the first national municipal regulation of algorithmic hiring tools at this scope

## IN BRIEF

New York City Local Law 144 of 2021 became operationally effective on July 5, 2023, after the New York City Department of Consumer and Worker Protection finalized implementing rules. The law requires employers using "automated employment decision tools" (AEDTs) for hiring or promotion decisions in NYC to conduct annual independent bias audits, to publish a summary of the most recent audit on the employer's website, and to provide candidate notice that an AEDT will be used. The audit must compute selection-rate and impact-ratio metrics by sex, race/ethnicity, and intersectional categories. The law was the first municipal regulation of algorithmic hiring tools at this scope in the United States and has been an influential reference for subsequent state-level proposals. Independent academic critiques have surfaced two limitations that matter: bias audits without bias data — employers often lack the protected-attribute data the audit metrics require — and null compliance — 18 of 391 employers posted an audit at all, and non-compliance cannot be established. The case pairs with Case 85 (OU Analyse — governance-objection dissolved by design), Case 86 (Gándara community-college predictive equity), and Case 182 (Amazon hiring AI). The intervention is the audit-as- governance-artifact discipline; whether the audits reduce actual disparate impact at the hiring level remains under-evidenced.

## BACKGROUND

Local Law 144 was passed by the New York City Council on 10 November 2021 and became law without the Mayor's signature on 11 December 2021, the Mayor returning it unsigned on 13 December, with the operational rules to be specified by the Department of Consumer and Worker Protection. The rulemaking process extended through 2022 and into 2023, with two rounds of public comment that surfaced substantial industry and civil-society engagement. The final rules became effective on July 5, 2023, and the law moved from statute to operational regime on that date. The scope is municipal — employers using an AEDT for a hiring or promotion decision affecting a position in New York City — but the practical reach is

broad because many employers operating nationally use AEDTs that touch New York City positions.<sup>9</sup>

#### THE INTERVENTION

The operational requirements have three components. First, an annual independent bias audit by a person or entity that has not used or developed the AEDT, computing selection-rate and impact-ratio metrics by sex, race/ ethnicity, and the intersectional categories the rules specify. Second, publication of a summary of the most recent audit on the employer's website, including the date the audit was conducted, the date the AEDT was first used, and the audit's selection-rate and impact-ratio findings. Third, candidate notice — applicants must be told that an AEDT will be used and given a process to request alternative selection or accommodations. The audit- and-notice structure is the artifact the law produces; the law does not prohibit AEDTs or set bias thresholds for use, and the regulatory theory is disclosure-and-audit rather than substantive standards.<sup>1</sup>

#### HOW IT WORKED

The independent academic critiques have surfaced two limitations that matter. Andrus, Spitzer, Brown, and Xiang's 2021 interviews with 38 algorithmic-fairness practitioners found that they routinely lack access to the demographic data bias detection requires; an employer deploying an AEDT inherits that constraint, and the audits that proceed are either limited to attributes the employer happens to have or rely on imputation methods whose accuracy is itself under-evidenced. Wright et al.'s 2024 study put 155 investigators onto 391 employers and found 18 had posted an audit report and 13 a transparency notice; the posted audits were data tables with minimal explanation, none of them sufficient to inform a job seeker's decision to opt out. The paper names the deeper problem **null compliance** — employers hold enough discretion over whether the law applies that non-compliance cannot be established at all.

#### THE EVIDENCE

The case pairs with Case 85 (OU Analyse) for the governance-objection-dissolved-by-design thread: OU Analyse's deployment surfaced an equity question that the design process resolved structurally; Local Law 144's audit regime surfaces equity questions structurally through disclosure rather than through a design change. Pair with Case 86 (Gándara community-college predictive equity) for the predictive-equity thread at higher- education scale. Pair with Case 182 (Amazon hiring AI) for the same domain — the audit regime is the governance instrument that, had it been in place, would have applied to an internal recruiting tool of Amazon's described character had it been deployed against NYC candidates. The case is the rare example in the corpus of an intervention at the regulatory scale; whether the audit regime reduces actual disparate impact at the hiring level remains under-evidenced, and the open evaluation question is part of the case. A December 2025 New York State Comptroller audit sharpened that question from the enforcement side: it found the Department of Consumer and Worker Protection's enforcement of the law "ineffective" — flagging only one likely non-compliant disclosure among the 32 it reviewed, where the Comptroller identified at least 17, and finding that roughly 75% of test complaint calls to 311 never reached the agency — moving the open issue from whether disclosure reduces disparate impact to whether the regime is being enforced at all.<sup>2</sup>

## WHAT TRANSFERRED

The hedges are explicit. The bias-audit-as- governance-artifact intervention is an audit-and-disclosure regime, not a substantive-standards regime; the law does not require employers to achieve any specific impact ratio or to refrain from deploying an AEDT that performs poorly on the audit. Whether the disclosure-and-audit structure reduces actual disparate impact at the hiring level is an empirical question the published evidence does not yet resolve. The 5% posting rate and the unprovability of non-compliance are themselves findings the case carries — a governance artifact almost nobody publishes, and whose absence proves nothing, is a weaker instrument than the regulatory theory assumes. The intervention is real and its limits are real; the change-control-and-disclosure-as- governance-artifacts LEO is anchored by the case at the municipal-regulatory scale, and the evaluation arc the audit regime opens is at the start of its evidence development.

*The audit-and-notice regime is a disclosure-and-audit instrument, not a substantive-standards instrument; whether it reduces actual disparate impact at the hiring level is an empirical question the published evidence does not yet resolve.*

EDITORS' SYNTHESIS OF THE LOCAL LAW 144 RULE TEXT AND THE ANDRUS ET AL. AND WRIGHT ET AL. ACADEMIC CRITIQUES.

## REFERENCES

1. New York City Department of Consumer and Worker Protection, Rules Implementing Local Law 144 of 2021 (Automated Employment Decision Tools), effective July 5, 2023.
2. Wright, L., Muenster, R. M., Vecchione, B., Qu, T., Cai, S., Smith, A., Metcalf, J., & Matias, J. N. (2024), "Null Compliance: NYC Local Law 144 and the Challenges of Algorithm Accountability," in Proceedings of FAccT 2024, doi:10.1145/3630106.3658998.
3. Office of the New York State Comptroller (2025), audit of the Department of Consumer and Worker Protection's enforcement of Local Law 144 — finding enforcement "ineffective" (December 2025).
4. Engler, A. (2023), "The EU and U.S. diverge on AI regulation: A transatlantic comparison and steps to alignment," Brookings Institution commentary — regulatory-comparative frame for the municipal intervention.

## THE LEARNING ENGINEERING LENS

### LE INSIGHT

NYC Local Law 144 is the bias-audit-as-governance-artifact intervention at municipal-regulatory scale. The audit-and- notice regime is the first national municipal regulation of algorithmic hiring tools at this scope; the Andrus et al. and Wright et al. critiques name the data-availability and audit-quality limitations the regulatory theory does not provide for. The intervention is real; whether it reduces actual disparate impact at the hiring level is under- evidenced and is the open evaluation question the case carries.

### LENS APPROACH

NYC Local Law 144 is the change-control-and-disclosure-as- governance-artifacts case at municipal-regulatory scale (induced 5.4; LENS D5/PT5; LEO-4 and LEO-5). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the audit-as-governance-instrument discipline. Pair with Case 85 (OU Analyse governance-objection-dissolved-by-design), Case 86 (Gándara community-college predictive equity), and Case 182 (Amazon hiring AI). The intervention is real and its limits are real; the disclosure-and-audit structure is not a substantive-standards structure, and the reduction- of-actual-disparate-impact evaluation question is at the start of its evidence development.

| DOMAIN                                                                             | MODES ADDRESSED  | LENS COMPETENCY                                                            | PROBLEM TYPE | INDUCED | LEO   |
|------------------------------------------------------------------------------------|------------------|----------------------------------------------------------------------------|--------------|---------|-------|
| <div>GOVERNMENT</div> <div>HIRING ALGORITHMS</div> <div>AUDIT AND DISCLOSURE</div> | <div>D G K</div> | <div>D1 D2 D3 D4 D5</div> <div>Navigating Sociotechnical Constraints</div> | 5            | 5.4     | 4 · 5 |

### APPROACHES TO CONSIDER

#### DURING DEVELOPMENT

- Specify the protected-attribute data the audit metrics will require before the audit is commissioned; the Andrus et al. critique names data availability as the precondition the regulatory theory does not provide for, and the data infrastructure has to be built in advance of the audit.
- Choose an independent auditor whose methodology will produce a documentation-detailed audit rather than a single-paragraph summary; the thinness of the posted audits the Wright et al. study found is itself a deployment choice, and the choice of auditor is where it surfaces.
- Build the candidate-notice and alternative-selection process as part of the deployment, not as a compliance afterthought; the candidate-side governance interaction is the seam at which the disclosure-and-audit structure becomes contestable for the affected person.

#### IN OPERATION / AFTER

- Carry the hedges — disclosure-and-audit regime not substantive-standards regime; reduction of actual disparate impact under-evidenced; a 5% audit-posting rate, and non-compliance that cannot be established — into print without softening; the case's pedagogical value depends on the intervention's limits being preserved.
- Pair in syllabi with Case 85 (OU Analyse) so the governance-objection-dissolved-by-design and governance-objection-surfaced-by-disclosure threads are taught together as complementary intervention forms.
- Use the case as the change-control-and-disclosure-as-governance-artifacts anchor at the municipal-regulatory scale; the curricular target is the discipline of building the data infrastructure and the audit-quality criteria the disclosure regime presupposes.

### REFLECTION QUESTIONS

1. Identify a regulated decision domain in your setting in which a disclosure-and-audit regime has been proposed or adopted. What is the protected-attribute data infrastructure the audit metrics will require, and is the infrastructure in place before the regime's effective date?
2. Specify the audit-quality criteria you would apply when commissioning an independent audit. What is the format that distinguishes a documentation-detailed audit from a single-paragraph compliance summary, and what is the decision rule for accepting an auditor's methodology?
3. The Local Law 144 regime is at the start of its evidence development on whether disclosure-and-audit reduces actual disparate impact. Pick a regulatory intervention in your domain and ask: what is the empirical-evaluation arc that would surface the intervention's effect, and what evidence-development infrastructure would the arc require?

**WHO BUILDS THIS** systems & human-factors engineering · governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · knowledge bases · data pipelines



# The Discipline We Build Next

**T** Training Gap   **K** Knowledge & Institutional Memory   **N** Normalization of Deviance   **G** Governance & Trust  
**H** Human-System Interface   **D** Designed Out

**IMPACT** Open: the case that has not yet been written. The discipline that this casebook documents is the one practitioners now in training must build.

## IN BRIEF

The cases that precede this one are closed: the investigations are complete, the failures diagnosed, the interventions evaluated. This one is open. It belongs to the practitioners now in training — the students of the LDT program and the LENS concentration, the graduates who will take positions in healthcare, defense, education at scale, and the human-AI frontier, and the institutions they will build. The strongest evidence for the discipline is the cumulative record of what its absence has cost; the strongest argument for its possibility is the record of what its presence has produced. Capability engineering in 2026 sits roughly where INPO sat in 1979 or CRM in 1981 — evidence accumulating, practitioners in training, institutional architecture not yet complete. What the reader does next is what fills in this case.

## THE SHIFT

Every preceding case in this book is closed: the investigation is complete, the failure diagnosed or the intervention evaluated. This one is open. It belongs to the practitioners now in training — the students of the LDT program and the LENS concentration, and the institutions they will build across healthcare, defense, education, and the human-AI frontier. What distinguishes this case is not a verdict already reached but a verdict still to be earned, written by the choices its readers make once they leave the page.<sup>9</sup>

## WHAT IS EMERGING

What is emerging is a discipline. The argument of this volume is that capability is engineerable — that training, interfaces, evidence systems, and institutions can be designed rather than left to chance — and that the work of designing them is a teachable practice with its own methods and growing body of knowledge. If that argument holds, then the gaps the preceding cases document are not the price of progress but the absence of work that could have been done.<sup>1</sup>

## THE CAPABILITY QUESTION

The capability question this case poses is the largest one: can the discipline be built fast enough, and at enough scale, to close the gaps the preceding cases document? The strongest evidence for the discipline is the cumulative record of what its absence has cost; the strongest argument for its possibility is the record of what its presence has produced



— and the case stays open precisely because that question is answered not by argument but by what the next generation of practitioners actually builds.<sup>2</sup>

#### EARLY EVIDENCE

The early evidence is the success cases themselves. Capability engineering in 2026 sits roughly where INPO sat in 1979 or CRM in 1981 — a discipline whose evidence is accumulating, whose practitioners are being trained, and whose institutional architecture is not yet complete. Those precedents became mature institutions within a generation, which is the most that can yet be claimed here: a comparable trajectory is plausible, but it is not given, and the comparison is an invitation rather than a guarantee.<sup>3</sup>

#### OPEN PROBLEMS

No specific incident sits behind this case because the work is not done. The institutions of capability engineering that exist today are insufficient to the volume of capability work the world now requires. What the practitioners reading this volume do next — a failure prevented, a success engineered, an institution built — is what fills in this case, and the page is left deliberately unfinished so that the record of the discipline's possibility can still be added to.<sup>4</sup>

## Open

### THE UNWRITTEN CASE

*what the practitioners reading this volume do next is what fills it in*

THE CASE THAT HAS NOT YET BEEN WRITTEN

*Capability is engineerable. The institution to engineer it has not yet been built. That is the work.*

LDT / LENS PROGRAM STATEMENT, JOHNS HOPKINS SCHOOL OF EDUCATION

#### REFERENCES

1. LDT / LENS program statement, Johns Hopkins School of Education — the program and its commitments.
2. Goodell & Kolodner (2022), Learning Engineering Toolkit — the discipline's methods.
3. The preceding cases of this volume — the cumulative record of cost and of possibility.
4. INPO (Case 175) and CRM/CAST (Case 117) — institution-building precedents (1979, 1981).
5. IEEE ICICLE Learning Engineering Body of Knowledge (LEBoK); the introduction to this volume.

#### THE LEARNING ENGINEERING LENS

#### LE INSIGHT

Every case in this book is closed. This one is open. The discipline LENS represents is the discipline that closes cases that are now open. The institutions of capability engineering that exist in 2026 are insufficient to the volume of capability work the contemporary world requires. The graduates of this program are the people who will build the institutions that are insufficient now.

## LENS APPROACH

LENS is organized to produce the practitioners who can do this work. The curriculum's commitments — capability as a system parameter, evidence as a deliverable, ethics as design, implementation as a thread — are the operational form of the argument the preceding cases evidence. Every studio project (LEN 10) is a small attempt at this case.

## DOMAIN

## MODES AT STAKE

EDUCATION AT SCALE

HEALTHCARE

DEFENSE



T K N G H D

## APPROACHES TO CONSIDER

## DURING DEVELOPMENT

- Treat capability as a system parameter to be specified and engineered from the outset of any deployment, not left to chance and diagnosed only after failure.
- Build the evidence system as a deliverable alongside the capability, so the discipline can show what its presence produces rather than only argue for it.
- Design the institution that will carry the practice forward, taking INPO and CRM as precedents for how a discipline matures within a generation.

## IN OPERATION / AFTER

- Monitor whether the discipline is in fact closing the gaps the preceding cases document, treating the open case as a standing measure of progress.
- Sustain the institutional architecture once built, since the precedents show maturity is reached only by carrying the work across a generation.
- Keep the cumulative record of cost and of possibility current, so each failure prevented or success engineered is added to the evidence the discipline rests on.

## REFLECTION QUESTIONS

1. What is the case you are positioned to add to this dataset in the next ten years — a failure prevented, a success engineered, an institution built?
2. If this book is read by the practitioners who will write its next case, what should that case look like?
3. Capability engineering in 2026 is compared here to where INPO stood in 1979 — a plausible trajectory, not a guaranteed one. What would have to be true a generation from now for that comparison to hold, and what is the first institution you would build to make it so?

**WHO BUILDS THIS** learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis · governance, ethics & measurement · human factors & interaction design · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards · governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs · task analysis · design prototyping

DOMAIN INDEX

The cases by domain.

The matrix at the front lists every case in number order. This index reorganizes the cases by primary domain — a reader following healthcare, defense, education, or any other thread can find the relevant cases together. Each case appears once, under its primary domain. Where a section opens with a callout, it names the dataset’s standout failure and, where one exists in that domain, its standout success.

DEFENSE

STANDOUT FAILURE

Case 128

Operation Eagle Claw (1980). Cross-organizational capability that existed on no service’s deliverable list; the mission was improvised across the Army, Navy, Marines, Air Force, and CIA. Produced the Holloway Commission, USSOCOM, and Goldwater-Nichols.

STANDOUT SUCCESS

Case 137

U.S. Nuclear Navy / Rickover Training Model (1954–). The longest continuous capability-engineering program in any high-consequence domain. Sixty-plus years of zero reactor accidents through training treated as a system parameter.

|     |                                                   |                |       |
|-----|---------------------------------------------------|----------------|-------|
| 124 | USS Fitzgerald & USS John S. McCain               | 2017           | T·K·N |
| 126 | F-35 Sustainment & the Maintainer Capability Gap  | ongoing        | T·K·D |
| 127 | Military Fratricide — Desert Storm to Afghanistan | 1991 – 2014    | T·H·K |
| 128 | Operation Eagle Claw                              | 1980           | T·K   |
| 129 | Patriot Missile / Dhahran                         | 1991           | D·H·K |
| 130 | USS Vincennes & Iran Air Flight 655               | 1988           | H·T   |
| 131 | Stanislav Petrov / 1983 False Alert               | 1983           | H·T   |
| 133 | Mark 14 Torpedo Failures                          | 1941 – 1943    | T·K·G |
| 134 | V-22 Osprey                                       | 1991 – present | T·H·N |
| 135 | 9/11 Intelligence Sharing Failures                | 1996 – 2001    | G·K   |

|     |                                                                                                     |                |       |
|-----|-----------------------------------------------------------------------------------------------------|----------------|-------|
| 137 | U.S. Nuclear Navy / Rickover Training Model                                                         | 1954 – present | T·K·N |
| 138 | MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard<br>revision (series since 1968) | 2020           | G·K   |
| 139 | GIFT and the Adoption Gap                                                                           | 2012 – present | K·G·N |
| 140 | Navy Surface Warfare Readiness Reform                                                               | 2018 – present | T·K·N |
| 141 | DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks<br>– 2014                     | 2009           | H·K·D |

## AVIATION

### STANDOUT FAILURE

#### Case 97

**Boeing 737 MAX / MCAS (2018–19).** A single-string angle-of-attack sensor, a flight-control law that assumed pilot mental models from prior models, and a certification chain that did not knit the pieces together. 346 lives.

### STANDOUT SUCCESS

#### Case 117

**Crew Resource Management & CAST (1981–).** Authority gradients named and engineered into the cockpit, then a closed-loop industry data architecture sixteen years later. The 83% fatality-risk reduction CAST reports is a portfolio result; CRM's separate contribution is not isolated in any published estimate.

|     |                                  |                |       |
|-----|----------------------------------|----------------|-------|
| 96  | AeroPerú Flight 603              | 1996           | H·T   |
| 97  | Boeing 737 MAX / MCAS            | 2018 – 2019    | D·T·H |
| 102 | Air France Flight 447            | 2009           | T·H   |
| 103 | Kegworth / British Midland 92    | 1989           | T·H·K |
| 104 | Asiana Airlines Flight 214       | 2013           | T·H   |
| 105 | Helios Airways Flight 522        | 2005           | T·H   |
| 106 | TransAsia Airways Flight 235     | 2015           | T·H   |
| 107 | Eastern Air Lines Flight 401     | 1972           | H·T   |
| 109 | Colgan Air Flight 3407           | 2009           | T     |
| 110 | Atlas Air Flight 3591            | 2019           | T·K   |
| 112 | NASA Saturn V Documentation      | 1972 – present | K     |
| 117 | Crew Resource Management & CAST  | 1981 – present | T·H·N |
| 118 | Korean Air Safety Transformation | 2000 – present | T·N   |

|     |                                                                                        |                 |       |
|-----|----------------------------------------------------------------------------------------|-----------------|-------|
| 119 | Aviation Safety Reporting System (ASRS)                                                | 1976 – present  | T·K·N |
| 120 | EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation                        | 1996 – 2002     | H·K·G |
| 121 | TCAS — Coordinated Collision Avoidance and the Überlingen Lesson (TCAS II Version 7.1) | 1981 – 2008     | H·K·G |
| 122 | Singapore Airlines Safety Transformation                                               | 1980s – present | T·N   |
| 132 | F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap                           | 2008 – 2012     | H·K·N |
| 185 | Air Canada Chatbot Liability — Delegation Without Revocation                           | 2022 – 2024     | D·K·N |

## HEALTHCARE

### STANDOUT FAILURE

#### Case 2

**EHR / CPOE Implementation (2005–).** Roughly \$30B federal investment in systems designed for billing and administration, deployed against clinical workflow. Usability is now itself a patient-safety variable at scale.

### STANDOUT SUCCESS

#### Case 23

**WHO Surgical Safety Checklist (2008–).** One page, nineteen items, paired with a required pause. Halved surgical mortality in the multi-site pilot; the smallest effective capability intervention in the dataset.

|    |                                                                                   |                |           |
|----|-----------------------------------------------------------------------------------|----------------|-----------|
| 1  | Therac-25                                                                         | 1985 – 1987    | H·D·G     |
| 2  | EHR / CPOE Implementation                                                         | 2005 – present | H·D·G     |
| 4  | Mid Staffordshire NHS Foundation Trust                                            | 2005 – 2009    | G·N·K     |
| 5  | Epic Sepsis Model                                                                 | 2017 – 2021    | D·G·N     |
| 8  | Medical Errors as Systemic Failure                                                | 1999 – present | T·H·N·K·G |
| 9  | Vioxx Withdrawal                                                                  | 1999 – 2004    | G·D       |
| 11 | California Nurse Staffing Ratios — Legislating a Capability Requirement           | 2004 – 2010    | G·K       |
| 12 | ACGME 80-Hour Resident Duty-Hour Reform                                           | 2003–2017      | T·K·N     |
| 13 | Implementation Science in Healthcare — The 17-Year Gap                            | ongoing        | K·G·N     |
| 19 | Keystone ICU / Pronovost Checklist                                                | 2004 – present | T·N       |
| 20 | TREWS / Sepsis Watch                                                              | 2018 – 2022    | H·K·G     |
| 21 | Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Hand-off | 2024 – 2025    | H·K·N     |

|     |                                                                               |                |       |
|-----|-------------------------------------------------------------------------------|----------------|-------|
| 22  | ChatGPT in Healthcare — Hallucination Cases                                   | 2023 – present | H·D   |
| 23  | WHO Surgical Safety Checklist                                                 | 2008 – present | T·N   |
| 24  | Bristol Heart Babies Reform                                                   | 1984 – present | G·N   |
| 32  | Annual-Screening UI Redesign at University of Missouri Health Care            | 2020           | H·N   |
| 33  | Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal | 2019 – 2024    | T·D·N |
| 34  | COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case                | 2022 – 2024    | T·K·D |
| 35  | Radiology AI Miscalibration                                                   | 2018 – present | H·K·D |
| 36  | AlphaFold — Protein Structure Prediction                                      | 2020 – present | K·G   |
| 37  | DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback      | 2020           | K·G   |
| 39  | TeamSTEPPS                                                                    | 2006 – present | T·N   |
| 176 | Tylenol Recall                                                                | 1982           | G·N   |

## EDUCATION AT SCALE

### STANDOUT FAILURE

#### Case 54

**Summit Learning / Personalized Learning Rollout (2014–19).** A defensible pedagogical platform, well funded, deployed across 380 schools — and pulled by parent revolt because the governance architecture (consent, evidence, exit) was never engineered alongside it.

### STANDOUT SUCCESS

#### Case 80

**Georgia State University Predictive Analytics (2012–).** 800 risk factors per student, advising triggered by early warning signs, equity as a primary constraint. Six-year graduation 32% → 54%; equity gaps eliminated.

|    |                                                      |                 |       |
|----|------------------------------------------------------|-----------------|-------|
| 45 | Tennessee Voluntary Pre-K Study                      | 2009–2018       | G·D   |
| 46 | Algorithmic Bias in Educational Predictive Analytics | ongoing         | G·H·D |
| 51 | Atlanta Public Schools Cheating Scandal              | 2009 – 2015     | G·N   |
| 53 | inBloom                                              | 2014            | G     |
| 54 | Summit Learning / Personalized Learning Rollout      | 2014–2019       | G·T·K |
| 67 | Cognitive Tutor / Carnegie Learning                  | 1990s – present | T     |
| 80 | Georgia State University Predictive Analytics        | 2012 – present  | T·K   |

**142** xAPI / Total Learning Architecture — Interoperability Gap ongoing

K·G

**205** The Discipline We Build Next ongoing

T·K·N·G·H·D

## ENERGY

### STANDOUT FAILURE

#### Case 166

**Three Mile Island (1979).** Training matched the design-basis events the engineers had imagined; the accident was not one of them. The control-room interface confused command signal with valve state. Catalysed industry-wide reform.

### STANDOUT SUCCESS

#### Case 175

**INPO and the Nuclear Academy (1979–).** The industry built the institution TMI revealed it needed. No INES Level 4-or-above accident at a U.S. commercial reactor since.

**146** Northeast Blackout 2003

H·K

**160** Davis-Besse Reactor Head Corrosion 2002

N·K·G

**161** Texas City BP Refinery Explosion 2005

N·T·K·G

**162** Upper Big Branch Mine Explosion 2010

N·G·K

**163** Deepwater Horizon 2010

T·N·K

**165** Camp Fire / PG&E 2018

G·N·K

**166** Three Mile Island 1979

T·H

**169** Fukushima Daiichi 2011

N·G·K

**175** INPO and the Nuclear Academy 1979 – present

T·K·G

## GOVERNMENT

### STANDOUT FAILURE

#### Case 191

**Australia Robodebt (2016–2023).** An income-averaging algorithm that reversed the burden of proof onto welfare recipients, raising some 470,000 wrongful debts. The Royal Commission found the scheme unlawful and attributed deaths to its operation — circumstantially, not as a finding of individual causation.

### NOTE

#### Case 203

The reforms that follow these failures — Royal Commissions, statutory restructures — are documented inside the narratives rather than as separate cases. The government dataset's paired-intervention success is NYC Local Law 144 (Case 203), filed under its own primary domain.

|     |                               |                |       |
|-----|-------------------------------|----------------|-------|
| 7   | VA Wait-Time Scandal          | 2014           | G·K·N |
| 180 | Healthcare.gov Launch         | 2013           | K·T·G |
| 193 | Bernard Madoff / SEC Failures | 1992 – 2008    | G·K·N |
| 197 | Predictive Policing — PredPol | 2011 – present | G·H·D |

## INDUSTRIAL

STANDOUT  
FAILURE

## Case 159

**Bhopal (1984).** Training, knowledge, normalization of deviance, and governance all collapsed in the same plant. Thousands died within hours; total-death estimates run to 15,000–20,000, with roughly half a million exposed. The cautionary case for outsourced operational responsibility.

STANDOUT  
SUCCESS

## Case 155

**Toyota Production System / Andon Cord (1950s–).** Any operator can signal a defect at the source. Under the fixed-position stop system most pulls never halt the line — the smallest authority-architecture intervention in the dataset, and the most durably copied.

|     |                                       |                 |         |
|-----|---------------------------------------|-----------------|---------|
| 144 | Kodak Digital Camera Stagnation       | 1975–2012       | D·K·N   |
| 147 | Takata Airbag Inflators               | 2008 – 2026     | D·G     |
| 149 | Volkswagen Dieselgate                 | 2015            | D·G     |
| 151 | GM Ignition Switch                    | 2002 – 2014     | D·G     |
| 155 | Toyota Production System / Andon Cord | 1950s – present | N·G     |
| 159 | Bhopal                                | 1984            | T·K·N·G |
| 164 | Grenfell Tower                        | 2017            | G·T·K·N |



## TECHNOLOGY

STANDOUT  
FAILURE

## Case 184

**UK Post Office Horizon Scandal (1999–2015).** Institutional refusal to believe sub-postmasters' reports of computer error, sustained over fifteen years. More than 900 wrongful convictions, 236 imprisoned, and 13 suicides the inquiry could not rule out as Horizon-caused — the longest-running operator-feedback failure in the dataset.

## NOTE

## Case 36

The technology dataset's success story is filed under Healthcare: AlphaFold (Case 36), the contemporary worked example of computational capability-engineering done as a discipline.

|     |                                |                |       |
|-----|--------------------------------|----------------|-------|
| 143 | Knight Capital Trading Loss    | 2012           | D·K   |
| 148 | Wells Fargo Fake Accounts      | 2011 – 2016    | G·N   |
| 152 | TSB Bank IT Migration          | 2018           | H·G   |
| 153 | Equifax Data Breach            | 2017           | G·K   |
| 157 | AI-Augmented Coding Tools      | 2021 – present | T·H   |
| 172 | CrowdStrike Falcon Outage      | 2024           | D·K·G |
| 184 | UK Post Office Horizon Scandal | 1999 – 2015    | G·H·K |

## AI IN EDUCATION

|    |                                                                        |      |       |
|----|------------------------------------------------------------------------|------|-------|
| 88 | LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models | 2025 | T·K·N |
|----|------------------------------------------------------------------------|------|-------|

## K–12 EDUCATION

|    |                                                                                  |              |       |
|----|----------------------------------------------------------------------------------|--------------|-------|
| 48 | School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism | 2010s – 2022 | G·K·N |
| 60 | Houston EVAAS — Value-Added Teacher Evaluation on Trial                          | 2011–2017    | G     |
| 61 | Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap | 1997–2023    | K·T   |
| 62 | Gates Intensive Partnerships — Measurement Without a Coupled Lever               | 2009–2018    | G·K   |
| 63 | LAUSD iPad Program — The Common Core Technology Project                          | 2013–2015    | D·H·G |
| 66 | Growth-Mindset National Experiment — Heterogeneous Effects                       | 2019         | D·N·K |
| 75 | Chen et al. — Rural China AI Devices and the Equity-Direction Finding            | 2025         | T·K·D |

- 95** PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop 1997–2024 **T-G**

## K-12 MATHEMATICS

- 72** ASSISTments — National Replication and Long-Term Follow-Through 2014 – present **T-K-D**
- 84** Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive 2007 – 2014 **T-K-D**

## K-12 SCIENCE

- 74** Zhang/Scardamalia — Knowledge Building Across Three Cohorts 2009 **T-K-D**

## AIR TRAFFIC MANAGEMENT

- 116** Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change 2005 **G-K-H**

## ANESTHESIOLOGY

- 27** Anesthesia Monitoring Standards and the APSF — The Mortality Decline 1986 – present **H-K-G**

## AUTONOMOUS SYSTEMS

### STANDOUT FAILURE

### Case 183

**Uber ATG / Tempe Fatality (2018).** The safety driver was nominally responsible for what the system was designed to do without her attention — the classic vigilance contradiction at the heart of partial-autonomy operator roles.

### NOTE

### Case 205

This key's cases are all failures and frontiers; the paired-intervention successes — Waymo's safety case (199) and the CPUC permit regime (200) — are filed under autonomous vehicles. See Case 205 (**The Discipline We Build Next**) for the open question this volume closes on.

- 183** Uber ATG / Tempe Fatality 2018 **T-N-G-H**
- 195** Tesla Autopilot — Recurring Fatalities 2016 – present **T-N-G-H**

## AUTONOMOUS VEHICLES

- 190** Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment 2023 – 2024 **G-K-N**
- 199** Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact 2023 – 2026 **G-K-N**

- 200** CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver 2020 – 2026 **G-K-D**

## AVIATION INFRASTRUCTURE

- 115** FAA NextGen and the ADS-B Out Transition 2003 – 2020 **G-D-K**

## AVIATION SAFETY

- 114** FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule 1988 – 2010s **G-D-K**

- 123** FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike 1992 – 2000s **G-K-N**

## CLINICAL CARE

- 15** I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer 2014 **H-K-N**

- 29** Bar-Code Medication Administration — Cue at the Point of Care 2010 **H-K-D**

## CLINICAL MEDICINE

- 6** Racial Bias in Pain Assessment — The False-Belief Mechanism 2016 **T-K-N**

- 25** Removing the Race Coefficient from eGFR 2021 **D-G-N**

## CLINICAL/TRANSLATIONAL RESEARCH

- 40** Team Science Training for Clinical and Translational Scientists 2020 – 2022 **K-N**

## CORPORATE L&D

- 65** Training Transfer — The Gap Between Learning and Doing 2010 **K-N**

- 70** High-Impact Learning System — Engineering the Environment, Not Just the Event 2001 – present **K-N**

- 79** The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops 1959 – present **K-N**

- 83** Brinkerhoff Success Case Method — Tails as the Evaluation Instrument 2003 – present **K-N**

## CRIMINAL JUSTICE

- 187** COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate 2014 – 2018 **D-K-N**

## DIGITAL IDENTITY

- 201** Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India 2017 – 2025 **G-N-H**

## E-GOVERNMENT

- 
- |            |                                                                                                          |              |
|------------|----------------------------------------------------------------------------------------------------------|--------------|
| <b>194</b> | Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox) 2001 – present | <b>G-K-N</b> |
|------------|----------------------------------------------------------------------------------------------------------|--------------|
- 

## ED-TECH

- 
- |           |                                                    |              |
|-----------|----------------------------------------------------|--------------|
| <b>47</b> | Algorithmic Bias in Automated Exam Proctoring 2022 | <b>D-N-K</b> |
|-----------|----------------------------------------------------|--------------|
- 

## EDUCATION (NORWAY)

- 
- |           |                                                                       |              |
|-----------|-----------------------------------------------------------------------|--------------|
| <b>92</b> | Norway's National Expert Commission on Learning Analytics 2021 – 2023 | <b>G-K-N</b> |
|-----------|-----------------------------------------------------------------------|--------------|
- 

## EDUCATION AT SCALE

- 
- |           |                                                                         |              |
|-----------|-------------------------------------------------------------------------|--------------|
| <b>50</b> | Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction 2012 – 2023 | <b>D-K-N</b> |
|-----------|-------------------------------------------------------------------------|--------------|
- 

- |           |                                                                          |            |
|-----------|--------------------------------------------------------------------------|------------|
| <b>69</b> | Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale 2016 | <b>T-D</b> |
|-----------|--------------------------------------------------------------------------|------------|
- 

## EMERGENCY MANAGEMENT

- 
- |            |                                                        |            |
|------------|--------------------------------------------------------|------------|
| <b>167</b> | Hurricane Katrina — The FEMA/DHS Response Failure 2005 | <b>G-K</b> |
|------------|--------------------------------------------------------|------------|
- 

- |            |                                                                          |            |
|------------|--------------------------------------------------------------------------|------------|
| <b>171</b> | Hurricane Maria — Puerto Rico Response and Logistics Failure 2017 – 2018 | <b>G-N</b> |
|------------|--------------------------------------------------------------------------|------------|
- 

## FINANCE

- 
- |            |                                                                                            |              |
|------------|--------------------------------------------------------------------------------------------|--------------|
| <b>186</b> | Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity 2009 – 2022 | <b>D-G-N</b> |
|------------|--------------------------------------------------------------------------------------------|--------------|
- 

- |            |                                                                             |              |
|------------|-----------------------------------------------------------------------------|--------------|
| <b>196</b> | Fintech Lending Fairness Audit — When Including Race Reduces Disparity 2025 | <b>D-G-N</b> |
|------------|-----------------------------------------------------------------------------|--------------|
- 

## FINANCIAL SERVICES

- 
- |            |                                                                                       |              |
|------------|---------------------------------------------------------------------------------------|--------------|
| <b>192</b> | Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model 2019 – 2021 | <b>D-K-N</b> |
|------------|---------------------------------------------------------------------------------------|--------------|
- 

## GENERAL AVIATION

- 
- |            |                                                                              |              |
|------------|------------------------------------------------------------------------------|--------------|
| <b>100</b> | Glass-Cockpit Transition in Light Aircraft — Technology Outran Training 2010 | <b>H-N-K</b> |
|------------|------------------------------------------------------------------------------|--------------|
- 

## GLOBAL HEALTH

- 
- |           |                                                                                                         |              |
|-----------|---------------------------------------------------------------------------------------------------------|--------------|
| <b>18</b> | PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption 2023 | <b>K-N-D</b> |
|-----------|---------------------------------------------------------------------------------------------------------|--------------|
- 

- |           |                                                                                                 |            |
|-----------|-------------------------------------------------------------------------------------------------|------------|
| <b>43</b> | Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface 2013 – 2018 | <b>H-N</b> |
|-----------|-------------------------------------------------------------------------------------------------|------------|
- 

- |            |                                                                    |            |
|------------|--------------------------------------------------------------------|------------|
| <b>170</b> | West Africa Ebola — The Delayed International Response 2014 – 2016 | <b>G-N</b> |
|------------|--------------------------------------------------------------------|------------|
-

## GOVERNMENT

|     |                                                                                                    |       |
|-----|----------------------------------------------------------------------------------------------------|-------|
| 49  | UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days 2020 | D-K-N |
| 191 | Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict 2016 – 2023        | D-G-N |
| 203 | NYC Local Law 144 — Bias Audits as Governance Artifact 2023 – present                              | D-G-K |

## GOVERNMENT/WELFARE (NETHERLANDS)

|     |                                                                              |       |
|-----|------------------------------------------------------------------------------|-------|
| 189 | Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds 2014 – 2020 | D-G-N |
|-----|------------------------------------------------------------------------------|-------|

## HEALTH PROFESSIONS EDUCATION

|    |                                                              |     |
|----|--------------------------------------------------------------|-----|
| 28 | Interprofessional Education and the Evidence Gap 2013 – 2015 | K-N |
|----|--------------------------------------------------------------|-----|

## HEALTHCARE/ONCOLOGY

|   |                                                                           |       |
|---|---------------------------------------------------------------------------|-------|
| 3 | Watson for Oncology — Delegation Marketed Ahead of Capability 2013 – 2018 | D-K-N |
|---|---------------------------------------------------------------------------|-------|

## HIGHER EDUCATION

|    |                                                                                                                 |       |
|----|-----------------------------------------------------------------------------------------------------------------|-------|
| 55 | Enrollment-Algorithm Yield Optimization Across U.S. Higher Education 2010s – present (Brookings synthesis 2021) | T-K-N |
| 57 | GAO Online Program Manager Oversight Gap (GAO-22-104463) 2022                                                   | D-K-N |
| 58 | USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC) 2010s – 2024                        | D-K-N |
| 85 | OU Analyse — Predictive Learning Analytics and Teacher Use at Scale 2019 – 2023                                 | T-K-D |
| 86 | Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction 2024                  | D-K-N |
| 90 | Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions 2025                               | T-K-N |

## HIGHER EDUCATION (AFRICA)

|    |                                                                   |       |
|----|-------------------------------------------------------------------|-------|
| 89 | Data Privacy and Learning Analytics on the African Continent 2022 | K-G-N |
|----|-------------------------------------------------------------------|-------|

## HIGHER EDUCATION (BRAZIL)

|    |                                                                                    |     |
|----|------------------------------------------------------------------------------------|-----|
| 82 | MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil) 2024 | K-N |
|----|------------------------------------------------------------------------------------|-----|

## HIGHER EDUCATION (LATIN AMERICA)

- 
- |           |                                                                                            |            |
|-----------|--------------------------------------------------------------------------------------------|------------|
| <b>91</b> | LALA — Building Learning–Analytics Governance Capacity Across Latin America<br>2017 – 2020 | <b>K-N</b> |
|-----------|--------------------------------------------------------------------------------------------|------------|
- 

## HIGHER ED (UK)

- 
- |           |                                                              |              |
|-----------|--------------------------------------------------------------|--------------|
| <b>81</b> | Open University 'Ethical Use of Student Data' and OU Analyse | <b>G-K-N</b> |
|-----------|--------------------------------------------------------------|--------------|
- 

## HIGHER-ED ANALYTICS

- 
- |           |                                                               |              |
|-----------|---------------------------------------------------------------|--------------|
| <b>52</b> | Purdue Course Signals — The Reverse–Causality Retention Claim | <b>D-K-N</b> |
|-----------|---------------------------------------------------------------|--------------|
- 

## HOSPITAL PHARMACY

- 
- |           |                                                       |              |
|-----------|-------------------------------------------------------|--------------|
| <b>42</b> | Australian Hospital–Pharmacy Technician Role Redesign | <b>D-N-H</b> |
|-----------|-------------------------------------------------------|--------------|
- 

## HUMANITARIAN RELIEF

- 
- |            |                                                                                   |          |
|------------|-----------------------------------------------------------------------------------|----------|
| <b>178</b> | The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami | <b>G</b> |
|------------|-----------------------------------------------------------------------------------|----------|
- 

## IMPLEMENTATION SCIENCE

- 
- |           |                                                                           |            |
|-----------|---------------------------------------------------------------------------|------------|
| <b>41</b> | Translational–Workforce Training — Stated Goals Outrunning Named Practice | <b>K-N</b> |
|-----------|---------------------------------------------------------------------------|------------|
- 

## LEARNING ANALYTICS

- 
- |           |                                                                                     |              |
|-----------|-------------------------------------------------------------------------------------|--------------|
| <b>73</b> | The Doer Effect at Scale — Replication, AI–Generated Questions, Non–WEIRD Extension | <b>T-K-D</b> |
|-----------|-------------------------------------------------------------------------------------|--------------|
- 
- |           |                                                                         |              |
|-----------|-------------------------------------------------------------------------|--------------|
| <b>87</b> | Yu / Lee / Kizilcec — Protected Attributes in Learning–Analytics Models | <b>D-K-N</b> |
|-----------|-------------------------------------------------------------------------|--------------|
- 
- |           |                                                                                         |            |
|-----------|-----------------------------------------------------------------------------------------|------------|
| <b>94</b> | Learning Analytics on the African Continent — A Scoping Review as the Present–State Map | <b>K-N</b> |
|-----------|-----------------------------------------------------------------------------------------|------------|
- 

## MEDICAL DEVICES

- 
- |           |                                  |              |
|-----------|----------------------------------|--------------|
| <b>26</b> | Pulse Oximetry Across Skin Tones | <b>D-G-N</b> |
|-----------|----------------------------------|--------------|
- 

## MEDICAL EDUCATION

- 
- |           |                                                                                            |            |
|-----------|--------------------------------------------------------------------------------------------|------------|
| <b>14</b> | Barsuk SBML — Simulation–Based Mastery Learning Dissemination from North–western to the VA | <b>T-K</b> |
|-----------|--------------------------------------------------------------------------------------------|------------|
- 
- |           |                                           |              |
|-----------|-------------------------------------------|--------------|
| <b>17</b> | Spaced Education RCTs in Medical Training | <b>H-K-D</b> |
|-----------|-------------------------------------------|--------------|
-

## MEDICAL-DEVICE REGULATION

- 
- |           |                                                                                                   |            |
|-----------|---------------------------------------------------------------------------------------------------|------------|
| <b>44</b> | Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD 2014 – present | <b>G·N</b> |
|-----------|---------------------------------------------------------------------------------------------------|------------|
- 

## MULTIMODAL LEARNING ANALYTICS

- 
- |           |                                                                                         |              |
|-----------|-----------------------------------------------------------------------------------------|--------------|
| <b>76</b> | Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account 2023 | <b>T·K·N</b> |
|-----------|-----------------------------------------------------------------------------------------|--------------|
- 

## NAVAL ENGINEERING

- 
- |            |                                                                                        |              |
|------------|----------------------------------------------------------------------------------------|--------------|
| <b>173</b> | Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable 1963 – present | <b>G·K·H</b> |
|------------|----------------------------------------------------------------------------------------|--------------|
- 

## NEUROSCIENCE

- 
- |            |                                                                     |            |
|------------|---------------------------------------------------------------------|------------|
| <b>181</b> | EU Human Brain Project — Top-Down Vision That Unraveled 2013 – 2023 | <b>G·N</b> |
|------------|---------------------------------------------------------------------|------------|
- 
- |            |                                                                                        |              |
|------------|----------------------------------------------------------------------------------------|--------------|
| <b>198</b> | Launching the BRAIN Initiative — Governance of a Grand Challenge 2011 – 2015 – present | <b>G·K·N</b> |
|------------|----------------------------------------------------------------------------------------|--------------|
- 

## NEUROSCIENCE/CONNECTOMICS

- 
- |           |                                                                                 |              |
|-----------|---------------------------------------------------------------------------------|--------------|
| <b>78</b> | CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale 2017 – present | <b>H·K·N</b> |
|-----------|---------------------------------------------------------------------------------|--------------|
- 

## NUCLEAR ENGINEERING

- 
- |            |                                                                                                |              |
|------------|------------------------------------------------------------------------------------------------|--------------|
| <b>156</b> | INL / LWRs Control-Room Modernization — Sustainment Research for an Aging Fleet 2010 – present | <b>D·H·K</b> |
|------------|------------------------------------------------------------------------------------------------|--------------|
- 

## NUCLEAR POWER

- 
- |            |                                                                                           |              |
|------------|-------------------------------------------------------------------------------------------|--------------|
| <b>154</b> | INL Turbine-Control Upgrade — Verification and Validation as the Cutover Deliverable 2014 | <b>G·K·H</b> |
|------------|-------------------------------------------------------------------------------------------|--------------|
- 

## NURSING EDUCATION

- 
- |           |                                                                            |              |
|-----------|----------------------------------------------------------------------------|--------------|
| <b>16</b> | NCSBN National Simulation Study — Licensing the 50% Substitution Rule 2014 | <b>G·K·D</b> |
|-----------|----------------------------------------------------------------------------|--------------|
- 

## PHARMA SAFETY

- 
- |           |                                                                                        |              |
|-----------|----------------------------------------------------------------------------------------|--------------|
| <b>38</b> | iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms 2006 – 2011 | <b>G·D·N</b> |
|-----------|----------------------------------------------------------------------------------------|--------------|
- 

## PUBLIC HEALTH

- 
- |            |                                                                                              |            |
|------------|----------------------------------------------------------------------------------------------|------------|
| <b>179</b> | The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed 2020 – 2023 | <b>G·K</b> |
|------------|----------------------------------------------------------------------------------------------|------------|
-

## RAIL TRANSPORT

---

**177** CIRAS — Confidential Incident Reporting for UK Rail 1996 – present **G·K·N**

---

## SOFTWARE ENGINEERING EDUCATION

---

**71** Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development 2025 **K·N**

---

## SOFTWARE SUSTAINMENT

---

**174** Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed 1996 – 2000 **G·D·K**

---

## SPACE / AEROSPACE

STANDOUT  
FAILURE**Case 111**

**Challenger & Columbia (1986 / 2003).** Normalisation of deviance across seventeen years and two crews — the sociologist Diane Vaughan's term for a culture that accepted flying with flaws once problems were defined as normal and routine, and the culture the CAIB found still intact in 2003.

## NOTE

**Case 137**

No paired-intervention success in the space dataset — Part III's aviation half carries seven. See Defense (Rickover, Case 137) for the safety-engineering counterpart human spaceflight did not build.

---

**98** Mars Climate Orbiter — Unit Mismatch 1999 **D·K**

---

**101** Ariane 5 Flight 501 1996 **D·K·H**

---

**108** NASA EVA-23 — Water Intrusion Inside a Spacesuit 2013 **H·N·D**

---

**111** Challenger & Columbia 1986 / 2003 **N·K·G**

---

**113** Boeing Starliner 2019 – 2025 **K·D**

---

## SURGERY

---

**30** Surgical Skill Video Peer-Rating Predicts Complications 2013 **H·D·K**

---

**31** Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units 2006 – 2016 **T·K·H**

---

## TECHNOLOGY

---

**182** Amazon Hiring AI — Trained Bias, Deprecated 2017 2014 – 2018 **D·K·N**

---



TUTORING

|    |                                                         |      |       |
|----|---------------------------------------------------------|------|-------|
| 77 | Hybrid Human-AI Tutoring — Augmentation, Not Delegation | 2024 | T·K·D |
|----|---------------------------------------------------------|------|-------|

WORKFORCE DEV

|    |                                                                                |                |       |
|----|--------------------------------------------------------------------------------|----------------|-------|
| 68 | CIRCUIT — A Scalable, Equity-Centered Research Workforce Model<br>(six cycles) | 2017 – 2023    | T·K   |
| 93 | Singapore SkillsFuture — National Workforce Capability at Scale                | 2015 – present | G·K·D |

**About this index.** Each case is filed under its primary domain, so every case appears exactly once. Cases that span several domains (a healthcare-AI case is also a technology case, a defense-training case is also an education case) are cross-referenced in their narratives. The standout failures and successes are editorial picks — the case in that domain the editors return to most often in studio.

# The cases by LENS course.

Each case names the LEN courses it serves as a worked example for. This index inverts that mapping: for every course in the concentration, the cases that inform it, in case-number order. A case appears under every course it supports, so many appear more than once.

## LEN 1

36 cases

|     |                                                                                             |
|-----|---------------------------------------------------------------------------------------------|
| 6   | Racial Bias in Pain Assessment — The False-Belief Mechanism                                 |
| 8   | Medical Errors as Systemic Failure                                                          |
| 13  | Implementation Science in Healthcare — The 17-Year Gap                                      |
| 36  | AlphaFold — Protein Structure Prediction                                                    |
| 39  | TeamSTEPPS                                                                                  |
| 53  | inBloom                                                                                     |
| 63  | LAUSD iPad Program — The Common Core Technology Project                                     |
| 67  | Cognitive Tutor / Carnegie Learning                                                         |
| 68  | CIRCUIT — A Scalable, Equity-Centered Research Workforce Model                              |
| 80  | Georgia State University Predictive Analytics                                               |
| 97  | Boeing 737 MAX / MCAS                                                                       |
| 102 | Air France Flight 447                                                                       |
| 111 | Challenger & Columbia                                                                       |
| 112 | NASA Saturn V Documentation                                                                 |
| 114 | FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule                 |
| 115 | FAA NextGen and the ADS-B Out Transition                                                    |
| 116 | Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change |
| 117 | Crew Resource Management & CAST                                                             |

- 124 USS Fitzgerald & USS John S. McCain
- 138 MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard
- 139 GIFT and the Adoption Gap
- 141 DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks
- 154 INL Turbine-Control Upgrade — Verification and Validation as the Cutover Deliverable
- 156 INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet
- 166 Three Mile Island
- 167 Hurricane Katrina — The FEMA/DHS Response Failure
- 171 Hurricane Maria — Puerto Rico Response and Logistics Failure
- 173 Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable
- 174 Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed
- 175 INPO and the Nuclear Academy
- 178 The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami
- 180 Healthcare.gov Launch
- 181 EU Human Brain Project — Top-Down Vision That Unraveled
- 194 Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox)
- 198 Launching the BRAIN Initiative — Governance of a Grand Challenge
- 205 The Discipline We Build Next

## LEN 2

53 cases

- 1 Therac-25
- 2 EHR / CPOE Implementation
- 5 Epic Sepsis Model
- 14 Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA
- 15 I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer
- 16 NCSBN National Simulation Study — Licensing the 50% Substitution Rule
- 17 Spaced Education RCTs in Medical Training
- 18 PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption

- 
- 20 TREWS / Sepsis Watch

---

  - 22 ChatGPT in Healthcare — Hallucination Cases

---

  - 27 Anesthesia Monitoring Standards and the APSF — The Mortality Decline

---

  - 29 Bar-Code Medication Administration — Cue at the Point of Care

---

  - 30 Surgical Skill Video Peer-Rating Predicts Complications

---

  - 31 Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units

---

  - 52 Purdue Course Signals — The Reverse-Causality Retention Claim

---

  - 65 Training Transfer — The Gap Between Learning and Doing

---

  - 66 Growth-Mindset National Experiment — Heterogeneous Effects

---

  - 68 CIRCUIT — A Scalable, Equity-Centered Research Workforce Model

---

  - 69 Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale

---

  - 70 High-Impact Learning System — Engineering the Environment, Not Just the Event

---

  - 71 Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development

---

  - 72 ASSISTments — National Replication and Long-Term Follow-Through

---

  - 73 The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension

---

  - 74 Zhang/Scardamalia — Knowledge Building Across Three Cohorts

---

  - 76 Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account

---

  - 77 Hybrid Human-AI Tutoring — Augmentation, Not Delegation

---

  - 78 CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale

---

  - 84 Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive

---

  - 85 OU Analyse — Predictive Learning Analytics and Teacher Use at Scale

---

  - 93 Singapore SkillsFuture — National Workforce Capability at Scale

---

  - 95 PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop

---

  - 96 AeroPerú Flight 603

---

  - 97 Boeing 737 MAX / MCAS

---

  - 102 Air France Flight 447

---

  - 104 Asiana Airlines Flight 214

---

- 105 Helios Airways Flight 522
- 106 TransAsia Airways Flight 235
- 107 Eastern Air Lines Flight 401
- 116 Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change
- 117 Crew Resource Management & CAST
- 118 Korean Air Safety Transformation
- 127 Military Fratricide — Desert Storm to Afghanistan
- 130 USS Vincennes & Iran Air Flight 655
- 131 Stanislav Petrov / 1983 False Alert
- 141 DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks
- 144 Kodak Digital Camera Stagnation
- 146 Northeast Blackout
- 155 Toyota Production System / Andon Cord
- 157 AI-Augmented Coding Tools
- 172 CrowdStrike Falcon Outage
- 183 Uber ATG / Tempe Fatality
- 184 UK Post Office Horizon Scandal
- 195 Tesla Autopilot — Recurring Fatalities

### LEN 3

47 cases

- 4 Mid Staffordshire NHS Foundation Trust
- 17 Spaced Education RCTs in Medical Training
- 24 Bristol Heart Babies Reform
- 32 Annual-Screening UI Redesign at University of Missouri Health Care
- 33 Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal
- 34 COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case
- 37 DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback
- 39 TeamSTEPPS

- |     |                                                                                               |
|-----|-----------------------------------------------------------------------------------------------|
| 49  | UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days |
| 50  | Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction                                   |
| 55  | Enrollment-Algorithm Yield Optimization Across U.S. Higher Education                          |
| 57  | GAO Online Program Manager Oversight Gap (GAO-22-104463)                                      |
| 58  | USC x 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC)                   |
| 72  | ASSISTments — National Replication and Long-Term Follow-Through                               |
| 73  | The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension           |
| 74  | Zhang/Scardamalia — Knowledge Building Across Three Cohorts                                   |
| 75  | Chen et al. — Rural China AI Devices and the Equity-Direction Finding                         |
| 84  | Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive                         |
| 85  | OU Analyse — Predictive Learning Analytics and Teacher Use at Scale                           |
| 86  | Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction     |
| 87  | Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models                       |
| 88  | LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models                        |
| 90  | Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions                  |
| 108 | NASA EVA-23 — Water Intrusion Inside a Spacesuit                                              |
| 111 | Challenger & Columbia                                                                         |
| 117 | Crew Resource Management & CAST                                                               |
| 120 | EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation                               |
| 121 | TCAS — Coordinated Collision Avoidance and the Überlingen Lesson                              |
| 124 | USS Fitzgerald & USS John S. McCain                                                           |
| 126 | F-35 Sustainment & the Maintainer Capability Gap                                              |
| 132 | F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap                                  |
| 134 | V-22 Osprey                                                                                   |
| 135 | 9/11 Intelligence Sharing Failures                                                            |
| 137 | U.S. Nuclear Navy / Rickover Training Model                                                   |
| 144 | Kodak Digital Camera Stagnation                                                               |
| 154 | INL Turbine-Control Upgrade — Verification and Validation as the Cutover Deliverable          |

|       |                                                                                                    |          |
|-------|----------------------------------------------------------------------------------------------------|----------|
| 159   | Bhopal                                                                                             |          |
| 161   | Texas City BP Refinery Explosion                                                                   |          |
| 163   | Deepwater Horizon                                                                                  |          |
| 164   | Grenfell Tower                                                                                     |          |
| 169   | Fukushima Daiichi                                                                                  |          |
| 175   | INPO and the Nuclear Academy                                                                       |          |
| 182   | Amazon Hiring AI — Trained Bias, Deprecated 2017                                                   |          |
| 187   | COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate                                    |          |
| 192   | Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model                          |          |
| 203   | NYC Local Law 144 — Bias Audits as Governance Artifact                                             |          |
| 205   | The Discipline We Build Next                                                                       |          |
| LEN 4 |                                                                                                    | 68 cases |
| 4     | Mid Staffordshire NHS Foundation Trust                                                             |          |
| 5     | Epic Sepsis Model                                                                                  |          |
| 6     | Racial Bias in Pain Assessment — The False-Belief Mechanism                                        |          |
| 7     | VA Wait-Time Scandal                                                                               |          |
| 8     | Medical Errors as Systemic Failure                                                                 |          |
| 9     | Vioxx Withdrawal                                                                                   |          |
| 12    | ACGME 80-Hour Resident Duty-Hour Reform                                                            |          |
| 15    | I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer                                    |          |
| 18    | PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption |          |
| 19    | Keystone ICU / Pronovost Checklist                                                                 |          |
| 20    | TREWS / Sepsis Watch                                                                               |          |
| 21    | Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff                   |          |
| 23    | WHO Surgical Safety Checklist                                                                      |          |
| 24    | Bristol Heart Babies Reform                                                                        |          |
| 25    | Removing the Race Coefficient from eGFR                                                            |          |

- 
- 26 Pulse Oximetry Across Skin Tones

---

  - 28 Interprofessional Education and the Evidence Gap

---

  - 32 Annual-Screening UI Redesign at University of Missouri Health Care

---

  - 33 Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal

---

  - 35 Radiology AI Miscalibration

---

  - 40 Team Science Training for Clinical and Translational Scientists

---

  - 41 Translational-Workforce Training — Stated Goals Outrunning Named Practice

---

  - 42 Australian Hospital-Pharmacy Technician Role Redesign

---

  - 43 Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface

---

  - 45 Tennessee Voluntary Pre-K Study

---

  - 46 Algorithmic Bias in Educational Predictive Analytics

---

  - 48 School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism

---

  - 51 Atlanta Public Schools Cheating Scandal

---

  - 60 Houston EVAAS — Value-Added Teacher Evaluation on Trial

---

  - 61 Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap

---

  - 62 Gates Intensive Partnerships — Measurement Without a Coupled Lever

---

  - 65 Training Transfer — The Gap Between Learning and Doing

---

  - 67 Cognitive Tutor / Carnegie Learning

---

  - 70 High-Impact Learning System — Engineering the Environment, Not Just the Event

---

  - 79 The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops

---

  - 80 Georgia State University Predictive Analytics

---

  - 81 Open University 'Ethical Use of Student Data' and OU Analyse

---

  - 82 MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil)

---

  - 83 Brinkerhoff Success Case Method — Tails as the Evaluation Instrument

---

  - 89 Data Privacy and Learning Analytics on the African Continent

---

  - 91 LALA — Building Learning-Analytics Governance Capacity Across Latin America

---

  - 92 Norway's National Expert Commission on Learning Analytics

---

  - 93 Singapore SkillsFuture — National Workforce Capability at Scale

---



- 
- 94** Learning Analytics on the African Continent — A Scoping Review as the Present-State Map
  - 95** PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop
  - 109** Colgan Air Flight 3407
  - 110** Atlas Air Flight 3591
  - 117** Crew Resource Management & CAST
  - 119** Aviation Safety Reporting System (ASRS)
  - 123** FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike
  - 140** Navy Surface Warfare Readiness Reform
  - 142** xAPI / Total Learning Architecture — Interoperability Gap
  - 147** Takata Airbag Inflators
  - 148** Wells Fargo Fake Accounts
  - 149** Volkswagen Dieselgate
  - 151** GM Ignition Switch
  - 161** Texas City BP Refinery Explosion
  - 162** Upper Big Branch Mine Explosion
  - 177** CIRAS — Confidential Incident Reporting for UK Rail
  - 178** The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami
  - 179** The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed
  - 186** Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity
  - 189** Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds
  - 190** Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment
  - 193** Bernard Madoff / SEC Failures
  - 196** Fintech Lending Fairness Audit — When Including Race Reduces Disparity
  - 199** Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact
  - 200** CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver
- 

## LEN 5

84 cases

- 
- 1** Therac-25
-

- 
- 3 Watson for Oncology — Delegation Marketed Ahead of Capability

---

  - 11 California Nurse Staffing Ratios — Legislating a Capability Requirement

---

  - 12 ACGME 80-Hour Resident Duty-Hour Reform

---

  - 14 Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA

---

  - 16 NCSBN National Simulation Study — Licensing the 50% Substitution Rule

---

  - 17 Spaced Education RCTs in Medical Training

---

  - 19 Keystone ICU / Pronovost Checklist

---

  - 21 Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff

---

  - 27 Anesthesia Monitoring Standards and the APSF — The Mortality Decline

---

  - 29 Bar-Code Medication Administration — Cue at the Point of Care

---

  - 30 Surgical Skill Video Peer-Rating Predicts Complications

---

  - 31 Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units

---

  - 34 COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case

---

  - 37 DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback

---

  - 38 iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms

---

  - 42 Australian Hospital-Pharmacy Technician Role Redesign

---

  - 44 Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD

---

  - 47 Algorithmic Bias in Automated Exam Proctoring

---

  - 49 UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days

---

  - 50 Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction

---

  - 52 Purdue Course Signals — The Reverse-Causality Retention Claim

---

  - 55 Enrollment-Algorithm Yield Optimization Across U.S. Higher Education

---

  - 57 GAO Online Program Manager Oversight Gap (GAO-22-104463)

---

  - 58 USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC)

---

  - 62 Gates Intensive Partnerships — Measurement Without a Coupled Lever

---

  - 63 LAUSD iPad Program — The Common Core Technology Project

---

  - 66 Growth-Mindset National Experiment — Heterogeneous Effects

---

  - 68 CIRCUIT — A Scalable, Equity-Centered Research Workforce Model

---

- 
- 69 Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale
  - 76 Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account
  - 77 Hybrid Human-AI Tutoring — Augmentation, Not Delegation
  - 78 CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale
  - 88 LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models
  - 90 Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions
  - 95 PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop
  - 96 AeroPerú Flight 603
  - 97 Boeing 737 MAX / MCAS
  - 98 Mars Climate Orbiter — Unit Mismatch
  - 100 Glass-Cockpit Transition in Light Aircraft — Technology Outran Training
  - 101 Ariane 5 Flight 501
  - 102 Air France Flight 447
  - 103 Kegworth / British Midland 92
  - 104 Asiana Airlines Flight 214
  - 105 Helios Airways Flight 522
  - 106 TransAsia Airways Flight 235
  - 107 Eastern Air Lines Flight 401
  - 108 NASA EVA-23 — Water Intrusion Inside a Spacesuit
  - 109 Colgan Air Flight 3407
  - 113 Boeing Starliner
  - 120 EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation
  - 121 TCAS — Coordinated Collision Avoidance and the Überlingen Lesson
  - 124 USS Fitzgerald & USS John S. McCain
  - 126 F-35 Sustainment & the Maintainer Capability Gap
  - 127 Military Fratricide — Desert Storm to Afghanistan
  - 128 Operation Eagle Claw
  - 129 Patriot Missile / Dhahran
-

- 
- 130 USS Vincennes & Iran Air Flight 655
- 
- 132 F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap
- 
- 134 V-22 Osprey
- 
- 137 U.S. Nuclear Navy / Rickover Training Model
- 
- 138 MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard
- 
- 140 Navy Surface Warfare Readiness Reform
- 
- 141 DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks
- 
- 143 Knight Capital Trading Loss
- 
- 153 Equifax Data Breach
- 
- 159 Bhopal
- 
- 163 Deepwater Horizon
- 
- 166 Three Mile Island
- 
- 167 Hurricane Katrina — The FEMA/DHS Response Failure
- 
- 170 West Africa Ebola — The Delayed International Response
- 
- 171 Hurricane Maria — Puerto Rico Response and Logistics Failure
- 
- 172 CrowdStrike Falcon Outage
- 
- 173 Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable
- 
- 178 The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami
- 
- 180 Healthcare.gov Launch
- 
- 182 Amazon Hiring AI — Trained Bias, Deprecated 2017
- 
- 185 Air Canada Chatbot Liability — Delegation Without Revocation
- 
- 187 COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate
- 
- 191 Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict
- 
- 192 Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model
- 
- 194 Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox)
- 
- 201 Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India
- 
- 203 NYC Local Law 144 — Bias Audits as Governance Artifact
-

## LEN 6

18 cases

|     |                                                                                             |
|-----|---------------------------------------------------------------------------------------------|
| 8   | Medical Errors as Systemic Failure                                                          |
| 13  | Implementation Science in Healthcare — The 17-Year Gap                                      |
| 53  | inBloom                                                                                     |
| 61  | Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap            |
| 86  | Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction   |
| 87  | Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models                     |
| 113 | Boeing Starliner                                                                            |
| 116 | Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change |
| 129 | Patriot Missile / Dhahran                                                                   |
| 139 | GIFT and the Adoption Gap                                                                   |
| 142 | xAPI / Total Learning Architecture — Interoperability Gap                                   |
| 151 | GM Ignition Switch                                                                          |
| 154 | INL Turbine-Control Upgrade — Verification and Validation as the Cutover Deliverable        |
| 176 | Tylenol Recall                                                                              |
| 180 | Healthcare.gov Launch                                                                       |
| 193 | Bernard Madoff / SEC Failures                                                               |
| 194 | Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox)   |
| 195 | Tesla Autopilot — Recurring Fatalities                                                      |

## LEN 7

102 cases

|   |                                                               |
|---|---------------------------------------------------------------|
| 1 | Therac-25                                                     |
| 2 | EHR / CPOE Implementation                                     |
| 3 | Watson for Oncology — Delegation Marketed Ahead of Capability |
| 4 | Mid Staffordshire NHS Foundation Trust                        |
| 5 | Epic Sepsis Model                                             |
| 6 | Racial Bias in Pain Assessment — The False-Belief Mechanism   |
| 7 | VA Wait-Time Scandal                                          |

- 
- 9 Vioxx Withdrawal

---

  - 11 California Nurse Staffing Ratios — Legislating a Capability Requirement

---

  - 14 Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA

---

  - 15 I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer

---

  - 16 NCSBN National Simulation Study — Licensing the 50% Substitution Rule

---

  - 18 PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption

---

  - 20 TREWS / Sepsis Watch

---

  - 21 Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff

---

  - 22 ChatGPT in Healthcare — Hallucination Cases

---

  - 24 Bristol Heart Babies Reform

---

  - 25 Removing the Race Coefficient from eGFR

---

  - 26 Pulse Oximetry Across Skin Tones

---

  - 27 Anesthesia Monitoring Standards and the APSF — The Mortality Decline

---

  - 28 Interprofessional Education and the Evidence Gap

---

  - 29 Bar-Code Medication Administration — Cue at the Point of Care

---

  - 30 Surgical Skill Video Peer-Rating Predicts Complications

---

  - 31 Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units

---

  - 35 Radiology AI Miscalibration

---

  - 36 AlphaFold — Protein Structure Prediction

---

  - 38 iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms

---

  - 40 Team Science Training for Clinical and Translational Scientists

---

  - 41 Translational-Workforce Training — Stated Goals Outrunning Named Practice

---

  - 43 Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface

---

  - 44 Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD

---

  - 45 Tennessee Voluntary Pre-K Study

---

  - 46 Algorithmic Bias in Educational Predictive Analytics

---

  - 48 School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism

---

- 
- 51 Atlanta Public Schools Cheating Scandal

---

  - 53 inBloom

---

  - 54 Summit Learning / Personalized Learning Rollout

---

  - 60 Houston EVAAS — Value-Added Teacher Evaluation on Trial

---

  - 65 Training Transfer — The Gap Between Learning and Doing

---

  - 70 High-Impact Learning System — Engineering the Environment, Not Just the Event

---

  - 71 Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development

---

  - 72 ASSISTments — National Replication and Long-Term Follow-Through

---

  - 73 The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension

---

  - 74 Zhang/Scardamalia — Knowledge Building Across Three Cohorts

---

  - 75 Chen et al. — Rural China AI Devices and the Equity-Direction Finding

---

  - 76 Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account

---

  - 77 Hybrid Human-AI Tutoring — Augmentation, Not Delegation

---

  - 78 CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale

---

  - 79 The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops

---

  - 80 Georgia State University Predictive Analytics

---

  - 81 Open University 'Ethical Use of Student Data' and OU Analyse

---

  - 82 MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil)

---

  - 83 Brinkerhoff Success Case Method — Tails as the Evaluation Instrument

---

  - 84 Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive

---

  - 85 OU Analyse — Predictive Learning Analytics and Teacher Use at Scale

---

  - 88 LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models

---

  - 89 Data Privacy and Learning Analytics on the African Continent

---

  - 91 LALA — Building Learning-Analytics Governance Capacity Across Latin America

---

  - 92 Norway's National Expert Commission on Learning Analytics

---

  - 94 Learning Analytics on the African Continent — A Scoping Review as the Present-State Map

---

  - 100 Glass-Cockpit Transition in Light Aircraft — Technology Outtran Training

---

- 
- 108 NASA EVA-23 — Water Intrusion Inside a Spacesuit
  - 111 Challenger & Columbia
  - 114 FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule
  - 115 FAA NextGen and the ADS-B Out Transition
  - 120 EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation
  - 121 TCAS — Coordinated Collision Avoidance and the Überlingen Lesson
  - 123 FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike
  - 132 F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap
  - 133 Mark 14 Torpedo Failures
  - 135 9/11 Intelligence Sharing Failures
  - 138 MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard
  - 147 Takata Airbag Inflators
  - 148 Wells Fargo Fake Accounts
  - 149 Volkswagen Dieselgate
  - 151 GM Ignition Switch
  - 152 TSB Bank IT Migration
  - 153 Equifax Data Breach
  - 156 INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet
  - 159 Bhopal
  - 160 Davis-Besse Reactor Head Corrosion
  - 161 Texas City BP Refinery Explosion
  - 162 Upper Big Branch Mine Explosion
  - 164 Grenfell Tower
  - 165 Camp Fire / PG&E
  - 167 Hurricane Katrina — The FEMA/DHS Response Failure
  - 169 Fukushima Daiichi
  - 171 Hurricane Maria — Puerto Rico Response and Logistics Failure
  - 173 Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable
-



|     |                                                                                   |
|-----|-----------------------------------------------------------------------------------|
| 174 | Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed |
| 176 | Tylenol Recall                                                                    |
| 177 | CIRAS — Confidential Incident Reporting for UK Rail                               |
| 180 | Healthcare.gov Launch                                                             |
| 181 | EU Human Brain Project — Top-Down Vision That Unraveled                           |
| 184 | UK Post Office Horizon Scandal                                                    |
| 186 | Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity    |
| 189 | Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds                  |
| 193 | Bernard Madoff / SEC Failures                                                     |
| 195 | Tesla Autopilot — Recurring Fatalities                                            |
| 196 | Fintech Lending Fairness Audit — When Including Race Reduces Disparity            |
| 197 | Predictive Policing — PredPol                                                     |
| 198 | Launching the BRAIN Initiative — Governance of a Grand Challenge                  |

## LEN 8

90 cases

|    |                                                                                     |
|----|-------------------------------------------------------------------------------------|
| 7  | VA Wait-Time Scandal                                                                |
| 11 | California Nurse Staffing Ratios — Legislating a Capability Requirement             |
| 12 | ACGME 80-Hour Resident Duty-Hour Reform                                             |
| 13 | Implementation Science in Healthcare — The 17-Year Gap                              |
| 28 | Interprofessional Education and the Evidence Gap                                    |
| 32 | Annual-Screening UI Redesign at University of Missouri Health Care                  |
| 33 | Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal       |
| 39 | TeamSTEPs                                                                           |
| 40 | Team Science Training for Clinical and Translational Scientists                     |
| 41 | Translational-Workforce Training — Stated Goals Outrunning Named Practice           |
| 42 | Australian Hospital-Pharmacy Technician Role Redesign                               |
| 43 | Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface |
| 44 | Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD  |

- 
- 47 Algorithmic Bias in Automated Exam Proctoring

---

  - 48 School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism

---

  - 49 UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days

---

  - 50 Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction

---

  - 52 Purdue Course Signals — The Reverse-Causality Retention Claim

---

  - 54 Summit Learning / Personalized Learning Rollout

---

  - 55 Enrollment-Algorithm Yield Optimization Across U.S. Higher Education

---

  - 57 GAO Online Program Manager Oversight Gap (GAO-22-104463)

---

  - 58 USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC)

---

  - 61 Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap

---

  - 62 Gates Intensive Partnerships — Measurement Without a Coupled Lever

---

  - 63 LAUSD iPad Program — The Common Core Technology Project

---

  - 66 Growth-Mindset National Experiment — Heterogeneous Effects

---

  - 69 Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale

---

  - 71 Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development

---

  - 75 Chen et al. — Rural China AI Devices and the Equity-Direction Finding

---

  - 79 The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops

---

  - 81 Open University 'Ethical Use of Student Data' and OU Analyse

---

  - 82 MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil)

---

  - 83 Brinkerhoff Success Case Method — Tails as the Evaluation Instrument

---

  - 90 Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions

---

  - 91 LALA — Building Learning-Analytics Governance Capacity Across Latin America

---

  - 92 Norway's National Expert Commission on Learning Analytics

---

  - 93 Singapore SkillsFuture — National Workforce Capability at Scale

---

  - 94 Learning Analytics on the African Continent — A Scoping Review as the Present-State Map

---

  - 98 Mars Climate Orbiter — Unit Mismatch

---

  - 101 Ariane 5 Flight 501

---

- 
- 103 Kegworth / British Midland 92
- 
- 109 Colgan Air Flight 3407
- 
- 110 Atlas Air Flight 3591
- 
- 111 Challenger & Columbia
- 
- 112 NASA Saturn V Documentation
- 
- 113 Boeing Starliner
- 
- 114 FAA Aging–Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule
- 
- 115 FAA NextGen and the ADS–B Out Transition
- 
- 118 Korean Air Safety Transformation
- 
- 119 Aviation Safety Reporting System (ASRS)
- 
- 122 Singapore Airlines Safety Transformation
- 
- 123 FSF CFIT and ALAR Task Forces — Industry–Level Institution Building After a Spike
- 
- 124 USS Fitzgerald & USS John S. McCain
- 
- 126 F–35 Sustainment & the Maintainer Capability Gap
- 
- 127 Military Fratricide — Desert Storm to Afghanistan
- 
- 128 Operation Eagle Claw
- 
- 129 Patriot Missile / Dhahran
- 
- 133 Mark 14 Torpedo Failures
- 
- 134 V–22 Osprey
- 
- 135 9/11 Intelligence Sharing Failures
- 
- 137 U.S. Nuclear Navy / Rickover Training Model
- 
- 139 GIFT and the Adoption Gap
- 
- 142 xAPI / Total Learning Architecture — Interoperability Gap
- 
- 146 Northeast Blackout
- 
- 151 GM Ignition Switch
- 
- 152 TSB Bank IT Migration
- 
- 155 Toyota Production System / Andon Cord
- 
- 156 INL / LWRS Control–Room Modernization — Sustainment Research for an Aging Fleet
-

|     |                                                                                      |
|-----|--------------------------------------------------------------------------------------|
| 157 | AI-Augmented Coding Tools                                                            |
| 160 | Davis-Besse Reactor Head Corrosion                                                   |
| 163 | Deepwater Horizon                                                                    |
| 164 | Grenfell Tower                                                                       |
| 165 | Camp Fire / PG&E                                                                     |
| 169 | Fukushima Daiichi                                                                    |
| 174 | Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed    |
| 175 | INPO and the Nuclear Academy                                                         |
| 177 | CIRAS — Confidential Incident Reporting for UK Rail                                  |
| 181 | EU Human Brain Project — Top-Down Vision That Unraveled                              |
| 182 | Amazon Hiring AI — Trained Bias, Deprecated 2017                                     |
| 185 | Air Canada Chatbot Liability — Delegation Without Revocation                         |
| 187 | COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate                      |
| 190 | Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment              |
| 191 | Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict      |
| 192 | Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model            |
| 198 | Launching the BRAIN Initiative — Governance of a Grand Challenge                     |
| 199 | Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact  |
| 200 | CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver     |
| 201 | Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India |
| 203 | NYC Local Law 144 — Bias Audits as Governance Artifact                               |
| 205 | The Discipline We Build Next                                                         |

## LEN 9

30 cases

|    |                                                               |
|----|---------------------------------------------------------------|
| 2  | EHR / CPOE Implementation                                     |
| 3  | Watson for Oncology — Delegation Marketed Ahead of Capability |
| 25 | Removing the Race Coefficient from eGFR                       |
| 26 | Pulse Oximetry Across Skin Tones                              |

|     |                                                                                           |
|-----|-------------------------------------------------------------------------------------------|
| 34  | COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case                            |
| 35  | Radiology AI Miscalibration                                                               |
| 36  | AlphaFold — Protein Structure Prediction                                                  |
| 37  | DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback                  |
| 38  | iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms                |
| 46  | Algorithmic Bias in Educational Predictive Analytics                                      |
| 47  | Algorithmic Bias in Automated Exam Proctoring                                             |
| 67  | Cognitive Tutor / Carnegie Learning                                                       |
| 86  | Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction |
| 87  | Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models                   |
| 89  | Data Privacy and Learning Analytics on the African Continent                              |
| 100 | Glass-Cockpit Transition in Light Aircraft — Technology Outran Training                   |
| 129 | Patriot Missile / Dhahran                                                                 |
| 143 | Knight Capital Trading Loss                                                               |
| 170 | West Africa Ebola — The Delayed International Response                                    |
| 179 | The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed          |
| 185 | Air Canada Chatbot Liability — Delegation Without Revocation                              |
| 186 | Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity            |
| 189 | Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds                          |
| 190 | Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment                   |
| 191 | Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict           |
| 196 | Fintech Lending Fairness Audit — When Including Race Reduces Disparity                    |
| 197 | Predictive Policing — PredPol                                                             |
| 199 | Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact       |
| 200 | CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver          |
| 201 | Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India      |

## LEN 10

21 cases

|     |                                                                                  |
|-----|----------------------------------------------------------------------------------|
| 8   | Medical Errors as Systemic Failure                                               |
| 12  | ACGME 80-Hour Resident Duty-Hour Reform                                          |
| 13  | Implementation Science in Healthcare — The 17-Year Gap                           |
| 19  | Keystone ICU / Pronovost Checklist                                               |
| 22  | ChatGPT in Healthcare — Hallucination Cases                                      |
| 23  | WHO Surgical Safety Checklist                                                    |
| 39  | TeamSTEPPS                                                                       |
| 45  | Tennessee Voluntary Pre-K Study                                                  |
| 53  | inBloom                                                                          |
| 54  | Summit Learning / Personalized Learning Rollout                                  |
| 60  | Houston EVAAS — Value-Added Teacher Evaluation on Trial                          |
| 133 | Mark 14 Torpedo Failures                                                         |
| 139 | GIFT and the Adoption Gap                                                        |
| 140 | Navy Surface Warfare Readiness Reform                                            |
| 155 | Toyota Production System / Andon Cord                                            |
| 157 | AI-Augmented Coding Tools                                                        |
| 164 | Grenfell Tower                                                                   |
| 170 | West Africa Ebola — The Delayed International Response                           |
| 176 | Tylenol Recall                                                                   |
| 179 | The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed |
| 205 | The Discipline We Build Next                                                     |

## REFERENCES

# Works cited.

References 1–22 are cited by number in the Introduction; the program record that follows them supports the **About LENS** account at the end of this volume. Per-case sources are listed at the end of each case. The broader reading list in the last part of this section collects the core works the LENS curriculum draws on.

## CITED IN THE INTRODUCTION

1. Makary, M. A. & Daniel, M. (2016). Medical error — the third leading cause of death in the US. *BMJ* 353: i2139. doi:10.1136/bmj.i2139.
2. Morris, Z. S., Wooding, S. & Grant, J. (2011). The answer is 17 years, what is the question: understanding time lags in translational research. *Journal of the Royal Society of Medicine* 104(12): 510–520. doi:10.1258/jrsm.2011.110180.
3. Balas, E. A. & Boren, S. A. (2000). Managing clinical knowledge for health care improvement. *Yearbook of Medical Informatics*, 65–70.
4. U.S. Navy (2017). Comprehensive Review of Recent Surface Force Incidents; Strategic Readiness Review. Department of the Navy.
5. U.S. Government Accountability Office (2023). F-35 Sustainment: DOD Faces Several Uncertainties and Risks as it Defines Future Costs and Plans. GAO-23-105341. gao.gov.
6. Bulger, M., McCormick, P. & Pitcan, M. (2017). The Legacy of inBloom. Data & Society Research Institute. datasociety.net.
7. Office of Qualifications and Examinations Regulation, UK (2020). Awarding GCSE, AS, A level, advanced extension awards and extended project qualifications in summer 2020: interim report. Ofqual; see also Adam, D. (2020), MIT Technology Review.
8. Helmreich, R. L., Merritt, A. C. & Wilhelm, J. A. (1999). The evolution of Crew Resource Management training in commercial aviation. *International Journal of Aviation Psychology* 9(1): 19–32. doi:10.1207/s15327108ijap0901\_2.
9. Commercial Aviation Safety Team / Federal Aviation Administration (2016). Safety Enhancements and Outcomes. CAST/FAA.
10. Federal Aviation Administration (2017). Out Front on Airline Safety: Two Decades of Continuous Evolution. FAA Newsroom.
11. Rees, J. V. (1994). *Hostages of Each Other: The Transformation of Nuclear Safety since Three Mile Island*. University of Chicago Press. ISBN 978-0226706887.
12. Pronovost, P., Needham, D., Berenholtz, S. et al. (2006). An intervention to decrease catheter-related bloodstream infections in the ICU. *New England Journal of Medicine* 355(26): 2725–2732. doi:10.1056/NEJMoa061115.
13. Haynes, A. B., Weiser, T. G., Berry, W. R. et al. (2009). A surgical safety checklist to reduce morbidity and mortality in a global population. *New England Journal of Medicine* 360(5): 491–499. doi:10.1056/NEJMsa0810119.
14. Agency for Healthcare Research and Quality (2023). TeamSTEPPS 3.0 Curriculum. AHRQ Publication No. 23-0034.ahrq.gov.
15. Sawyer, R. K. (ed.) (2014). *The Cambridge Handbook of the Learning Sciences* (2nd ed.). Cambridge University Press. ISBN 978-1107033252.
16. Vicente, K. J. (1999). *Cognitive Work Analysis: Toward Safe, Productive, and Healthy Computer-Based Work*. Lawrence Erlbaum. ISBN 978-0805823974.
17. Wickens, C. D., Hollands, J. G., Banbury, S. & Parasuraman, R. (2021). *Engineering Psychology and Human Performance* (5th ed.). Routledge. doi:10.4324/9781003177616.
18. Fixsen, D. L., Naoom, S. F., Blasé, K. A., Friedman, R. M. & Wallace, F. (2005). *Implementation Research: A Synthesis of the Literature*. National Implementation Research Network. See also Aarons, G. A., Hurlburt, M. & Horwitz, S. M. (2011). Advancing a conceptual model of evidence-based practice implementation in public service sectors. *Administration and Policy in Mental Health* 38: 4–23. doi:10.1007/s10488-010-0327-7.
19. Reason, J. (1990). *Human Error*. Cambridge University Press, doi:10.1017/CBO9781139062367. Rasmussen, J. (1997). Risk management in a dynamic society: a modelling problem. *Safety Science* 27: 183–213, doi:10.1016/

- S0925-7535(97)00052-O. Leveson, N. (2011). *Engineering a Safer World: Systems Thinking Applied to Safety*. MIT Press. doi:10.7551/mitpress/8179.001.0001.
20. Weick, K. E. & Sutcliffe, K. M. (2015). *Managing the Unexpected: Sustained Performance in a Complex World* (3rd ed.). Wiley. See also Roberts, K. H. (1990). Some characteristics of one type of high reliability organization. *Organization Science* 1(2): 160–176.
  21. Saxberg, B. (2017). *Learning engineering: the art of applying learning science at scale*. In *Industry Connections Industry Consortium on Learning Engineering*. See also Goodell, J. & Kolodner, J. L. (eds.) (2022). *Learning Engineering Toolkit*. Routledge.
  22. IEEE ICICLE (2022). *Learning Engineering Body of Knowledge (LEBoK)*. lebok.wiki.

## THE LENS PROGRAM RECORD

Not cited by number in the Introduction. These are the sources behind the program account in About LENS at the end of this volume — the connectomics workforce work, its generalizations, and the incubator cohorts the Practice Flywheel was built from.

23. MICrONS Consortium (2025). Functional connectomics spanning multiple areas of mouse visual cortex. *Nature* 640: 167–179. doi:10.1038/s41586-025-08790-w. See also Hider, R. et al. (2022), BossDB, *Frontiers in Neuroinformatics*; Xenos, D. et al. (2022), NeuVue, bioRxiv 2022.07.18.500521; Gray-Roncal, W. (in press, 2026), “Calibrate the Scientists, Not Just the Microscopes,” *BRAIN CONNECTS* commentary.
24. Cervantes, M., Floryanzia, S., Sharp, J., Johnson, E. C. & Gray-Roncal, W. R. (2023). Empowering trailblazers toward scalable, systematized, research-based workforce development. 2023 ASEE Annual Conference & Exposition, Baltimore, MD. peer.asee.org. Documents the CIRCUIT model; Gray-Roncal is the senior author. See also Gray-Roncal, W. et al. (2025), MERIT: Mentoring Exceptional Researchers to Innovate and Thrive — A Structured, Scalable Research Training Framework for STEM Talent Development. JHU/APL whitepaper (the generalization of CIRCUIT). COMPASS — the layer of targeted mentor interventions that operates MERIT inside specific institutional settings, designed to help students uncover tacit and explicit knowledge gaps. COMPASS-Core: Curriculum Overview and Workshop Summaries (JHU/APL program documentation) is one well-documented instance, a workshop set grounded in the CCR framework and explicitly addressing the “hidden curriculum” of research practice.
25. TITAN — DARPA Artificial Intelligence Exploration program (JHU/APL pilot, c. 2020 – 2022). Program documentation on file with the LENS partner team at the JHU/APL AI Pathfinding Lab.
26. PISTA (2025). Perceptual Inference System for Tactical Agentic AI. JHU/APL AI Pathfinding Lab, Summer 2025 Research Incubator, sponsored by DTRA JSTO. Cohort poster and program documentation available from the LENS partner team.
27. Precision-medicine community-detection methods — preprint at medRxiv 2024.06.20.24309267.

## FOUNDATIONS · THE DISCIPLINE DRAWS ON

### Learning sciences and learning engineering

- Bransford, J. D., Brown, A. L. & Cocking, R. R. (eds.) (2000). *How People Learn: Brain, Mind, Experience, and School*. National Academy Press. doi:10.17226/9853.
- Sawyer, R. K. (ed.) (2014). *The Cambridge Handbook of the Learning Sciences* (2nd ed.).
- Koedinger, K. R., Corbett, A. T. & Perfetti, C. (2012). The Knowledge–Learning–Instruction framework. *Cognitive Science* 36(5): 757–798. doi:10.1111/j.1551-6709.2012.01245.x.
- Sweller, J., Ayres, P. & Kalyuga, S. (2011). *Cognitive Load Theory*. Springer. doi:10.1007/978-1-4419-8126-4.
- Soderstrom, N. C. & Bjork, R. A. (2015). Learning versus performance: an integrative review. *Perspectives on Psychological Science* 10(2): 176–199. doi:10.1177/1745691615569000.
- Goodell, J. & Kolodner, J. L. (eds.) (2022). *Learning Engineering Toolkit: Evidence–Based Practices from the Learning Sciences, Instructional Design, and Beyond*. Routledge. doi:10.4324/9781003276579.

### Human factors, cognitive engineering, and systems safety

- Vicente, K. J. (1999). *Cognitive Work Analysis*. Lawrence Erlbaum.
- Norman, D. A. (2013). *The Design of Everyday Things* (rev. ed.). Basic Books. ISBN 978-0465050659.
- Hollnagel, E. & Woods, D. D. (2005). *Joint Cognitive Systems: Foundations of Cognitive Systems Engineering*. CRC Press. doi:10.1201/9781420038194.
- Endsley, M. R. (1995). Toward a theory of situation awareness in dynamic systems. *Human Factors* 37(1): 32–64. doi:10.1518/001872095779049543.
- Perrow, C. (1984). *Normal Accidents: Living with High-Risk Technologies*. Basic Books. ISBN 978-0691004129 (1999 Princeton reprint).
- Leveson, N. (2011). *Engineering a Safer World: Systems Thinking Applied to Safety*. MIT Press. doi:10.7551/mitpress/8179.001.0001.
- Reason, J. (1997). *Managing the Risks of Organizational Accidents*. Ashgate. ISBN 978-1840141047.
- Rasmussen, J. (1997). Risk management in a dynamic society: a modelling problem. *Safety Science* 27: 183–213. doi:10.1016/S0925-7535(97)00052-O.



- Dekker, S. (2014). *The Field Guide to Understanding 'Human Error'* (3rd ed.). Ashgate. ISBN 978-1472439055.

### High-reliability organizing and team performance

- Weick, K. E. & Sutcliffe, K. M. (2015). *Managing the Unexpected* (3rd ed.). Wiley. ISBN 978-1118862414.
- Roberts, K. H. (1990). Some characteristics of one type of high reliability organization. *Organization Science* 1(2): 160–176. doi:10.1287/orsc.1.2.160.
- Edmondson, A. C. (2018). *The Fearless Organization: Creating Psychological Safety in the Workplace for Learning, Innovation, and Growth*. Wiley. ISBN 978-1119477242.
- Salas, E., Wilson, K. A., Burke, C. S. & Wightman, D. C. (2006). Does Crew Resource Management training work? An update, an extension, and some critical needs. *Human Factors* 48(2): 392–412. doi:10.1518/001872006777724444.
- Klein, G. (1998). *Sources of Power: How People Make Decisions*. MIT Press. ISBN 978-0262112277.
- Kanki, B. G., Helmreich, R. L. & Anca, J. (eds.) (2010). *Crew Resource Management* (2nd ed.). Academic Press. ISBN 978-0123749468.

### Implementation science and organizational learning

- Fixsen, D. L., Naoom, S. F., Blasé, K. A., Friedman, R. M. & Wallace, F. (2005). *Implementation Research: A Synthesis of the Literature*. NIRN.
- Aarons, G. A., Hurlburt, M. & Horwitz, S. M. (2011). Advancing a conceptual model of evidence-based practice implementation in public service sectors. *Administration and Policy in Mental Health* 38: 4–23. doi:10.1007/s10488-010-0327-7.
- Damschroder, L. J. et al. (2009). Fostering implementation of health services research findings into practice: a consolidated framework. *Implementation Science* 4: 50. doi:10.1186/1748-5908-4-50.
- Argyris, C. & Schön, D. A. (1996). *Organizational Learning II: Theory, Method, and Practice*. Addison-Wesley. ISBN 978-0201629835.
- Senge, P. M. (2006). *The Fifth Discipline: The Art and Practice of the Learning Organization* (rev. ed.). Doubleday. ISBN 978-0385517256.
- Vaughan, D. (1996). *The Challenger Launch Decision: Risky Technology, Culture, and Deviance at NASA*. University of Chicago Press. ISBN 978-0226851761.

### Ethics, governance, and equity in sociotechnical systems

- Eubanks, V. (2018). *Automating Inequality: How High-Tech Tools Profile, Police, and Punish the Poor*. St. Martin's. ISBN 978-1250074317.
- Friedman, B. & Nissenbaum, H. (1996). Bias in computer systems. *ACM Transactions on Information Systems* 14(3): 330–347. doi:10.1145/230538.230561.
- O'Neil, C. (2016). *Weapons of Math Destruction*. Crown. ISBN 978-0553418811.
- Selbst, A. D. et al. (2019). Fairness and abstraction in sociotechnical systems. In *FAT\*19*. doi:10.1145/3287560.3287598.
- Citron, D. K. (2008). Technological due process. *Washington University Law Review* 85: 1249–1313. SSRN.
- Selwyn, N. (2014). *Distrusting Educational Technology*. Routledge. ISBN 978-1138021068.

### Human-AI teaming and automation

- Parasuraman, R. & Manzey, D. H. (2010). Complacency and bias in human use of automation: an attentional integration. *Human Factors* 52(3): 381–410. doi:10.1177/0018720810376055.
- Bainbridge, L. (1983). Ironies of automation. *Automatica* 19(6): 775–779. doi:10.1016/0005-1098(83)90046-8.
- Sarter, N. B. & Woods, D. D. (1995). How in the world did we ever get into that mode? Mode error and awareness in supervisory control. *Human Factors* 37(1): 5–19. doi:10.1518/001872095779049516.
- Cummings, M. L. (2017). Automation bias in intelligent time-critical decision support systems. In *Decision Making in Aviation*.
- Casner, S. M. & Hutchins, E. (2019). What do we tell the drivers? Toward minimum driver-training standards for partially-automated cars. *Journal of Cognitive Engineering and Decision Making* 13(2): 55–66. doi:10.1177/15555343419830901.

### Patient safety and clinical improvement

- Kohn, L. T., Corrigan, J. M. & Donaldson, M. S. (eds.) (1999). *To Err Is Human: Building a Safer Health System*. Institute of Medicine. doi:10.17226/9728.
- Gawande, A. (2009). *The Checklist Manifesto: How to Get Things Right*. Metropolitan. ISBN 978-0805091748.
- Pronovost, P. & Vohr, E. (2010). *Safe Patients, Smart Hospitals*. Hudson Street Press. ISBN 978-1594630644.
- Wachter, R. (2015). *The Digital Doctor: Hope, Hype, and Harm at the Dawn of Medicine's Computer Age*. McGraw-Hill. ISBN 978-0071849463.
- Bates, D. W., Levine, D. M., Salmasian, H. et al. (2023). The safety of inpatient health care. *New England Journal of Medicine* 388(2): 142–153. doi:10.1056/NEJMSa2206117.

Each case spread in the body of the book includes its own primary sources. Sources cited only within case spreads are not repeated here; this list collects the works the discipline as a whole draws on, together with the references used in the introductory essay.

APPENDIX · LEOS AND COURSE COVERAGE

# The competencies, the LEOs, and the courses.

COMPETENCY → CONCENTRATION LEARNING OUTCOME

The five capabilities the cases turn on are named, in the curriculum, as five LENS Educational Objectives (LEOs) — the concentration-level objectives, held at the program level:

|                                       |                                                         |
|---------------------------------------|---------------------------------------------------------|
| Systems Analysis                      | LEO-1 · Systems Thinking and Analysis                   |
| Iterative Development                 | LEO-2 · Learning Engineering Design and Implementation  |
| Human-System Collaboration            | LEO-3 · Human-System Collaboration and Adaptive Systems |
| Test and Evaluation                   | LEO-4 · Data, Measurement, and Evaluation               |
| Navigating Sociotechnical Constraints | LEO-5 · Context and Domain Fluency                      |

THE LEN COURSES · CASES PER COURSE

Each case names the LEN courses it serves as a worked example for, and most cases serve several. The counts below are therefore not a partition of the 191 active cases — they sum to well more than 191, because a single case typically appears under three or four courses.

|       |                                    |    |
|-------|------------------------------------|----|
| LEN 1 | Principles of learning engineering | 36 |
| LEN 2 | Human-AI teaming                   | 53 |
| LEN 3 | Systems integration                | 47 |

|               |                          |     |
|---------------|--------------------------|-----|
| <b>LEN 4</b>  | Evidence and measurement | 68  |
| <b>LEN 5</b>  | Capability analysis      | 84  |
| <b>LEN 6</b>  | Applied problem-solving  | 18  |
| <b>LEN 7</b>  | Bias and governance      | 102 |
| <b>LEN 8</b>  | Knowledge transfer       | 90  |
| <b>LEN 9</b>  | Computational methods    | 30  |
| <b>LEN 10</b> | Capstone studio          | 21  |

#### WHAT THE DISTRIBUTION SHOWS

The distribution is itself diagnostic. The strongly covered topics are the ones the existing case literature most readily supports: bias and governance (LEN 7) and knowledge transfer (LEN 8) lead, with capability analysis (LEN 5) and evidence and measurement (LEN 4) close behind. Across the 191 active cases the corpus is saturated with examples of what goes wrong when ethics, governance, evidence, and institutional learning are not engineered — and, in the success cases, what it takes to engineer them.

The thinner coverage is equally informative. Computational methods (LEN 9) has the failure modes of algorithmic systems well represented but the computational tools for capability instrumentation under-represented. The capstone (LEN 10) maps to fewer cases than the number of studio projects each cohort undertakes. Systems integration (LEN 3) and applied problem-solving (LEN 6) required focused tagging to reach defensible coverage, surfacing that operational systems-engineering content is under-represented in the published case literature relative to its curricular centrality, and that stakeholder-analysis and problem-framing have not historically been packaged as case material by other disciplines. These are pedagogical gaps the LENS studio sequence is organized to close — not by adding more failure cases, but by having students produce the missing artifacts in real operational settings and entering them into the dataset for the next cohort.

# References, case by case.

Every reference cited in every case, organized by case in number order. The intent is that any reader can locate every source the book draws on. Each entry preserves the in-case citation form (author, year, venue) and adds a **Retrieved from:** rule beneath it — a worksheet line for the reviewer to enter the canonical access path they used: a publisher URL, a DOI, an agency landing page, or a stable archival locator. The rules print blank by design, because the per-case verification is being carried out against this appendix rather than reported by it; the running status of that work is in `casebook/verification-log.md`. Future editions will print the locators as the review closes them out.

**Style.** References are rendered in the form they appeared in the case — APA-style author/year/title/venue, with italicised book titles and journal names — supplemented by the explicit retrieval locator. Government and agency reports are cited by report number when one exists. Trade-book references that carry no DOI use ISBN as the stable identifier; that mapping is on the verification log.

## CASE 1 Therac-25

1985 – 1987

1. N. G. Leveson & C. S. Turner, "An Investigation of the Therac-25 Accidents," IEEE Computer 26(7): 18–41 (1993). doi:10.1109/MC.1993.274940 — the removed hardware interlocks and the adapted software. — Retrieved from: \_\_\_\_\_
2. Leveson & Turner (1993) — the two software defects (data-entry race condition and counter overflow), the uninformative "Malfunction 54", overdoses up to 100×, six accidents, and at least three deaths. — Retrieved from: \_\_\_\_\_
3. Leveson & Turner (1993); N. G. Leveson, Safeware: System Safety and Computers (Addison-Wesley, 1995) — the manufacturer's denial and the systemic findings. — Retrieved from: \_\_\_\_\_
4. N. G. Leveson, Engineering a Safer World: Systems Thinking Applied to Safety (MIT Press, 2011) — why removing a safeguard requires explicitly reassigning its safety function. — Retrieved from: \_\_\_\_\_
5. The case's role in medical-device software regulation and software-safety practice (FDA software guidance; IEC 62304 lineage); see also Ethics Unwrapped, UT Austin. — Retrieved from: \_\_\_\_\_

## CASE 2 EHR / CPOE Implementation

2005 – present

1. Health Information Technology for Economic and Clinical Health (HITECH) Act (2009) — the \$30B EHR incentive program. — Retrieved from: \_\_\_\_\_

2. Y. Han et al., "Unexpected Increased Mortality After Implementation of a Commercially Sold CPOE System," *Pediatrics* 116(6): 1506–1512 (2005) — the disputed pediatric-ICU mortality result. — Retrieved from: \_\_\_\_\_
3. M. Grissinger, "The Absence of a Drug–Disease Interaction Alert Leads to a Child's Death," *P&T* 43(2): 71–72 (2018); F. Schulte & E. Fry, "Death by 1,000 Clicks," *Fortune/KHN* (2019). — Retrieved from: \_\_\_\_\_
4. J. Howe, K. Adams, A. Z. Hettinger & R. Ratwani, "Electronic Health Record Usability Issues and Potential Contribution to Patient Harm," *JAMA* 319(12): 1276–1278 (2018) — 557 of 1.735M safety reports tied usability to possible harm. — Retrieved from: \_\_\_\_\_
5. D. Sittig & H. Singh, "Defining Health Information Technology–related Errors: New Developments Since To Err is Human," *Archives of Internal Medicine* 171(14): 1281–1284 (2011). — Retrieved from: \_\_\_\_\_
6. Pew / MedStar / AMA, *Ways to Improve Electronic Health Record Safety* (2018); R. Wachter, *The Digital Doctor* (2015). — Retrieved from: \_\_\_\_\_

### CASE 3 **Watson for Oncology — Delegation Marketed Ahead of Capability** 2013 – 2018

1. Somashekhar et al. and the WFO lung/colorectal early–experience abstracts, *J Clin Oncol* 35 (2017), 15\_suppl, presented at the ASCO Annual Meeting — the marketed concordance figures: 96% lung, 93% rectal, 81% colon. — Retrieved from: \_\_\_\_\_
2. Ross & Swetlitz, "IBM's Watson supercomputer recommended 'unsafe and incorrect' cancer treatments, internal documents show," *STAT* (25 July 2018), and "IBM pitched its Watson supercomputer as a revolution in cancer care. It's nowhere close," *STAT* (5 Sept 2017) — the primary sources; investigative journalism drawing on leaked IBM internal documents. — Retrieved from: \_\_\_\_\_
3. Strickland (2019), "How IBM Watson Overpromised and Underdelivered on AI Health Care," *IEEE Spectrum* — independent retrospective synthesis of the public record. — Retrieved from: \_\_\_\_\_
4. Topol (2019), *Deep Medicine*, Basic Books — secondary academic situating of Watson within the broader clinical–AI delegation discourse. — Retrieved from: \_\_\_\_\_
5. v2 paired cases: TREWS (Case 20), Epic Sepsis Model (Case 5), SyRI (Case 189) — the AI–delegation typology. — Retrieved from: \_\_\_\_\_

### CASE 4 **Mid Staffordshire NHS Foundation Trust** 2005 – 2009

1. R. Francis QC, *Report of the Mid Staffordshire NHS Foundation Trust Public Inquiry* (2013) — the staffing cuts and Foundation Trust targets. — Retrieved from: \_\_\_\_\_
2. Healthcare Commission, *Investigation into Mid Staffordshire NHS Foundation Trust* (2009) — ward conditions and excess mortality. — Retrieved from: \_\_\_\_\_
3. Francis QC (2013) — the three–volume, 1,782–page report and 290 recommendations. — Retrieved from: \_\_\_\_\_
4. Francis QC (2013) — "not just the Trust's Board but the system as a whole failed in its most essential duty — to protect patients from unacceptable risks of harm and from unacceptable, and in some cases inhumane, treatment" (quoted). — Retrieved from: \_\_\_\_\_
5. D. Berwick, *A Promise to Learn — A Commitment to Act* (National Advisory Group on the Safety of Patients in England, 2013). — Retrieved from: \_\_\_\_\_

6. K. Walshe & J. Higgins (2002) on NHS safety governance; cf. the VA wait-time scandal (Case 7). — Retrieved from: \_\_\_\_\_

## CASE 5 Epic Sepsis Model

2017 – 2021

1. Habib et al. (2021), commentary on Wong et al., JAMA Internal Medicine — on the implications for proprietary clinical AI. — Retrieved from: \_\_\_\_\_
2. Wong et al. (2021), “External Validation of a Widely Implemented Proprietary Sepsis Prediction Model in Hospitalized Patients,” JAMA Internal Medicine 181(8):1065–1070, doi:10.1001/jamainternmed.2021.2626. — Retrieved from: \_\_\_\_\_
3. FDA, Clinical Decision Support Software: Final Guidance (2022) — the post-Wong reframing of the EHR-embedded oversight question. — Retrieved from: \_\_\_\_\_
4. Ross, C. (2022), “Epic overhauls its sepsis algorithm,” STAT News (Oct 2022) — the ESM v2 redesign and the shift to locally-trained models. — Retrieved from: \_\_\_\_\_
5. Adams et al. (2022), Nature Medicine — the paired positive case (20). — Retrieved from: \_\_\_\_\_

## CASE 6 Racial Bias in Pain Assessment — The False-Belief Mechanism

2016

1. Anderson, Green, & Payne (2009), “Racial and ethnic disparities in pain: causes and consequences of unequal care,” Journal of Pain 10(12):1187–1204 — the population-level disparity. — Retrieved from: \_\_\_\_\_
2. Hoffman, Trawalter, Axt, & Oliver (2016), “Racial bias in pain assessment and treatment recommendations, and false beliefs about biological differences between blacks and whites,” PNAS 113(16):4296–4301, doi:10.1073/pnas.1516047113. — Retrieved from: \_\_\_\_\_
3. Sabin & Greenwald (2012), “The influence of implicit bias on treatment recommendations for 4 common pediatric conditions,” American Journal of Public Health — the diffuse-mechanism backdrop the Hoffman precision improves on. — Retrieved from: \_\_\_\_\_
4. Vyas, Eisenstein, & Jones (2020), NEJM — connecting race-in-clinical-algorithms to race-in-clinical-judgment. — Retrieved from: \_\_\_\_\_
5. Omiye, Lester, Spichak, Rotemberg, & Daneshjou (2023), “Large language models propagate race-based medicine,” npj Digital Medicine 6:195 — commercial LLMs reproduce the Hoffman false beliefs. — Retrieved from: \_\_\_\_\_

## CASE 7 VA Wait-Time Scandal

2014

1. D. Draper (GAO), House VA Committee testimony, GAO-19-687T (2019) — “given the high turnover among schedulers, it is important that VA remain vigilant about scheduler training” (quoted). — Retrieved from: \_\_\_\_\_
2. VA Office of Inspector General, Report 14-02603-267 (2014) — secret waiting lists and falsified appointment data. — Retrieved from: \_\_\_\_\_
3. VA OIG reports (2005, 2007, 2008) — prior, unactioned findings. — Retrieved from: \_\_\_\_\_
4. GAO Veterans Health Administration reports (2000–2019) — fifteen years of data-reliability warnings. — Retrieved from: \_\_\_\_\_

5. C. Argyris & D. Schön, *Organizational Learning: A Theory of Action Perspective* (1978); G. Bevan & C. Hood, "What's Measured Is What Matters," *Public Administration* 84(3): 517–538 (2006). — Retrieved from: \_\_\_\_\_

6. Veterans Access, Choice, and Accountability Act (2014) — the legislative response. — Retrieved from: \_\_\_\_\_

## CASE 8 Medical Errors as Systemic Failure

1999 – present

1. Institute of Medicine, *To Err Is Human: Building a Safer Health System* (1999); *Crossing the Quality Chasm* (2001); *Improving Diagnosis in Health Care* (2015) — the field-defining trilogy and the 44,000–98,000 estimate; the systems framing. — Retrieved from: \_\_\_\_\_
2. Makary, M. & Daniel, M. (2016), "Medical error — the third leading cause of death in the U.S.," *BMJ* 353:i2139 — the 250,000 estimate, the quoted framing, the extrapolation from four studies published since the IOM report (HealthGrades 2004; OIG 2010; Landrigan 2010; Classen 2011). — Retrieved from: \_\_\_\_\_
3. Shojania, K. & Dixon-Woods, M. (2017), "Estimating deaths due to medical error: the ongoing controversy and why it matters," *BMJ Quality & Safety* 26(5):423–428; with Bates, D. W. & Singh, H. (2018), *Health Affairs* 37(11):1736–1743 — methodological contestation of the Makary extrapolation and CDC-ranking comparison. — Retrieved from: \_\_\_\_\_
4. Makary & Daniel (2016), *BMJ* — death certificates do not capture medical error as a cause; ICD billing taxonomy as the structural reason. — Retrieved from: \_\_\_\_\_
5. Bates, D. W., Levine, D. M., Salmasian, H., et al. (2023), "The Safety of Inpatient Health Care," *NEJM* 388(2):142–153 — eleven-hospital Massachusetts cohort of 2,809 admissions; at least one adverse event in 23.6% of admissions, 25% of those preventable; persistence of harm at scale. — Retrieved from: \_\_\_\_\_
6. Agency for Healthcare Research and Quality, *National Healthcare Quality and Disparities Reports* (annual); CDC WONDER ICD-coded mortality data — institutional context for the missing national active-surveillance instrument. — Retrieved from: \_\_\_\_\_

## CASE 9 Vioxx Withdrawal

1999 – 2004

1. Graham, D. J. et al. (2005), "Risk of acute myocardial infarction and sudden cardiac death in patients treated with COX-2 selective and non-selective NSAIDs: nested case-control study," *Lancet* 365(9458):475–481 — Kaiser Permanente population-level cardiovascular risk analysis. — Retrieved from: \_\_\_\_\_
2. Bresalier, R. et al. (2005), "Cardiovascular events associated with rofecoxib in a colorectal adenoma chemoprevention trial," *APPROVe*, *NEJM* 352(11):1092–1102 — placebo-controlled confirmation of cardiovascular risk; early trial termination. — Retrieved from: \_\_\_\_\_
3. Ross, J. S., Hill, K. P., Egilman, D. S., Krumholz, H. M. (2008), "Guest authorship and ghostwriting in publications related to rofecoxib," *JAMA* 299(15):1800–1812 — guest-authorship and ghostwriting documentation from Merck litigation discovery. — Retrieved from: \_\_\_\_\_
4. U.S. Senate Committee on Finance, hearings on Vioxx and FDA's drug-safety system (November 18, 2004) — Graham testimony; "88,000 to 139,000 Americans" estimate of excess cardiovascular events; described FDA management pressure. — Retrieved from: \_\_\_\_\_
5. Bombardier, C. et al. (2000), "Comparison of upper gastrointestinal toxicity of rofecoxib and naproxen in patients with rheumatoid arthritis," *VIGOR trial*, *NEJM* 343(21):1520–1528 — the five-fold myocardial-infarction signal and the naproxen-protective hypothesis. — Retrieved from: \_\_\_\_\_

6. FDA Amendments Act of 2007 (P.L. 110–85) and FDA Sentinel Initiative documentation (2008–present) — REMS authority, post-market study requirements, and active distributed-data post-market surveillance. — Retrieved from: \_\_\_\_\_

## CASE 11 **California Nurse Staffing Ratios — Legislating a Capability Requirement** 2004 – 2010

1. Aiken, Clarke, Sloane, Sochalski, & Silber (2002), “Hospital nurse staffing and patient mortality, nurse burnout, and job dissatisfaction,” JAMA 288(16):1987–1993 — the workload-mortality relationship the 2010 modeling rests on. — Retrieved from: \_\_\_\_\_
2. California Department of Health Services (1999–2004), AB 394 regulatory implementation documentation. — Retrieved from: \_\_\_\_\_
3. Aiken, Sloane, Cimiotti, Clarke, Flynn, Seago, Spetz, & Smith (2010), “Implications of the California Nurse Staffing Mandate for Other States,” Health Services Research 45(4):904–921, doi:10.1111/j.1475-6773.2010.01114.x. — Retrieved from: \_\_\_\_\_
4. Spetz, Chapman, Herrera, Kaiser, Seago, & Dower (2009), “Assessing the impact of California’s nurse staffing ratios on hospitals and patient care,” California HealthCare Foundation — implementation-period analysis. — Retrieved from: \_\_\_\_\_

## CASE 12 **ACGME 80-Hour Resident Duty-Hour Reform** 2003–2017

1. B. H. Lerner, The Libby Zion Case and the Reform of Medical Education (2006); and the 1989 New York State (Bell Commission) duty-hour regulations — the origin of the reform. — Retrieved from: \_\_\_\_\_
2. Accreditation Council for Graduate Medical Education, Common Program Requirements (2003 and 2011 revisions) — the 80-hour weekly cap and the 16-hour first-year shift limit. — Retrieved from: \_\_\_\_\_
3. D. A. Asch, K. Y. Bilimoria & S. V. Desai, “Resident Duty Hours and Medical Education Policy — Raising the Evidence Bar,” NEJM 376: 1704–1706 (2017); Institute of Medicine (Ulmer, Wolman & Johns, eds.), Resident Duty Hours: Enhancing Sleep, Supervision, and Safety (2008) — hand-offs and continuity trade-offs. — Retrieved from: \_\_\_\_\_
4. Bilimoria, Chung, Hedges et al., “National Cluster-Randomized Trial of Duty-Hour Flexibility in Surgical Training” (FIRST), NEJM 374: 713–727 (2016). doi:10.1056/NEJMoa1515724. — Retrieved from: \_\_\_\_\_
5. Silber, Bellini, Shea et al., “Patient Safety Outcomes under Flexible and Standard Resident Duty-Hour Rules” (iCOMPARE), NEJM 380: 905–914 (2019). doi:10.1056/NEJMoa1810642. — Retrieved from: \_\_\_\_\_
6. ACGME, Common Program Requirements (2017 revision) — relaxation of the 16-hour first-year shift limit. — Retrieved from: \_\_\_\_\_
7. P. Pronovost et al., Keystone ICU intervention, NEJM 355: 2725–2732 (2006), doi:10.1056/NEJMoa061115 A. B. Haynes et al., surgical safety checklist, NEJM 360: 491–499 (2009), doi:10.1056/NEJMsa0810119 — integrated interventions that engineered the surrounding architecture. — Retrieved from: \_\_\_\_\_

## CASE 13 **Implementation Science in Healthcare — The 17-Year Gap** ongoing

1. E. A. Balas & S. A. Boren (2000), “Managing clinical knowledge for health care improvement,” Yearbook of Medical Informatics — the nine-service adoption analysis behind the seventeen-year figure, and the pipeline-attrition derivation behind the 14 percent. — Retrieved from: \_\_\_\_\_



2. Z. S. Morris, S. Wooding & J. Grant, "The answer is 17 years, what is the question: understanding time lags in translational research," *J. Royal Society of Medicine* 104(12):510–520 (2011) (quoted). — Retrieved from: \_\_\_\_\_
3. D. Fixsen et al., *Implementation Research: A Synthesis of the Literature* (2005) — the Active Implementation Frameworks. — Retrieved from: \_\_\_\_\_
4. G. Aarons et al. (2011), the EPIS framework; L. Damschroder et al. (2009), CFIR. — Retrieved from: \_\_\_\_\_
5. Cf. medical error as systemic failure (Case 8); Goodell & Kolodner, *Learning Engineering Toolkit* (2022). — Retrieved from: \_\_\_\_\_

## CASE 14 **Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA** 2009 – 2020s

1. Barsuk, McGaghie, Cohen, Balachandran, & Wayne (2009), "Use of simulation-based mastery learning to improve the quality of central venous catheter placement in a medical intensive care unit," *Journal of Hospital Medicine* 4(7):397–403. — Retrieved from: \_\_\_\_\_
2. Barsuk, Cohen, Feinglass, McGaghie, & Wayne (2009), "Use of simulation-based education to reduce catheter-related bloodstream infections," *Archives of Internal Medicine* 169(15):1420–1423, doi:10.1001/archinternmed.2009.215. — Retrieved from: \_\_\_\_\_
3. Cohen, Feinglass, Barsuk, Barnard, O'Donnell, McGaghie, & Wayne (2010), "Cost savings from reduced catheter-related bloodstream infection after simulation-based education for residents in a medical intensive care unit," *Simulation in Healthcare* 5(2):98–102, doi:10.1097/SH.0b013e3181bc8304. — Retrieved from: \_\_\_\_\_
4. Barsuk, Cohen, Nguyen, Mitra, O'Hara, Okuda, Feinglass, Cameron, McGaghie, & Wayne (2016), "Attending physician adherence to a 29-component central venous catheter bundle checklist during simulated procedures," *Critical Care Medicine* 44(10):1871–1881 — the published record of the train-the-trainer dissemination across 58 Veterans Affairs Medical Centers. — Retrieved from: \_\_\_\_\_
5. McGaghie, Issenberg, Cohen, Barsuk, & Wayne (2011), "Does simulation-based medical education with deliberate practice yield better results than traditional clinical medical education? A meta-analytic comparative review of the evidence," *Academic Medicine* 86(6):706–711. — Retrieved from: \_\_\_\_\_

## CASE 15 **I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer** 2014

1. Starmer, A. J., Spector, N. D., Srivastava, R., West, D. C., Rosenbluth, G., Allen, A. D., Noble, E. L., Tse, L. L., Dalal, A. K., Keohane, C. A., Lipsitz, S. R., Rothschild, J. M., Wien, M. F., Yoon, C. S., Zigmont, K. R., Wilson, K. M., O'Toole, J. K., Solan, L. G., Aylor, M., Bismilla, Z., Coffey, M., Mahant, S., Blankenburg, R. L., Destino, L. A., Everhart, J. L., Patel, S. J., Bale, J. F., Spackman, J. B., Stevenson, A. T., Calaman, S., Cole, F. S., Balmer, D. F., Hepps, J. H., Lopreiato, J. O., Yu, C. E., Sectish, T. C., & Landrigan, C. P. (2014). Changes in medical errors after implementation of a handoff program. *New England Journal of Medicine*, 371(19):1803–1812. doi:10.1056/NEJMSa1405556 — the case's primary evaluation. — Retrieved from: \_\_\_\_\_
2. Colligan, L., Brick, D., & Patterson, E. S. (2015). Changes in medical errors with a handoff program. *New England Journal of Medicine*, 372(5):490 — the correspondence cautioning against implementing the mnemonic alone; authors' reply at 372(5):490–491. — Retrieved from: \_\_\_\_\_
3. Sectish, T. C., Starmer, A. J., Landrigan, C. P., & Spector, N. D. (2010). Establishing a multisite education and research project requires leadership, expertise, collaboration, and an important aim. *Pediatrics*, 126(4):619–622 — the I-PASS Study Group methodology and design rationale. — Retrieved from: \_\_\_\_\_

4. Cohen, M. D., & Hilligoss, P. B. (2010). The published literature on handoffs in hospitals: deficiencies identified in an extensive review. *Quality and Safety in Health Care* — the broader handoff-failure-mode literature. — Retrieved from: \_\_\_\_\_

## CASE 16 **NCSBN National Simulation Study — Licensing the 50% Substitution Rule** 2014

1. Jeffries, P. R. (2012). *Simulation in Nursing Education: From Conceptualization to Evaluation* (2nd ed.). National League for Nursing — the simulation-quality framework the NCSBN study's quality conditions rest on. — Retrieved from: \_\_\_\_\_
2. INACSL Standards Committee (2016). *INACSL Standards of Best Practice: Simulation*. *Clinical Simulation in Nursing*, 12(S) — the simulation-practice standards downstream programs adopt as the analog of the NCSBN quality conditions. — Retrieved from: \_\_\_\_\_
3. Hayden, J. K., Smiley, R. A., Alexander, M., Kardong-Edgren, S., & Jeffries, P. R. (2014). The NCSBN National Simulation Study: A longitudinal, randomized, controlled study replacing clinical hours with simulation in prelicensure nursing education. *Journal of Nursing Regulation*, 5(2 Suppl):S1–S64. [https://www.journalofnursingregulation.com/article/s2155-8256\(15\)30062-4/fulltext](https://www.journalofnursingregulation.com/article/s2155-8256(15)30062-4/fulltext) — the case's primary study. — Retrieved from: \_\_\_\_\_
4. Smiley, R. A., & Martin, B. (2023). *Simulation in Nursing Education: Advancements in Regulation, 2014–2022*. *Journal of Nursing Regulation*. doi:10.1016/S2155-8256(23)00086-8 — the 2023 follow-up documenting the regulatory propagation. — Retrieved from: \_\_\_\_\_

## CASE 17 **Spaced Education RCTs in Medical Training** 2007 – 2009

1. Cepeda, N. J., Pashler, H., Vul, E., Wixted, J. T., & Rohrer, D. (2006). Distributed practice in verbal recall tasks: A review and quantitative synthesis. *Psychological Bulletin*, 132(3):354–380 — the basic-literature spacing-effect review the workplace-L&D RCT sits against. — Retrieved from: \_\_\_\_\_
2. Kerfoot, B. P., Baker, H. E., Koch, M. O., Connelly, D., Joseph, D. B., & Ritchey, M. L. (2007). Randomized, controlled trial of spaced education to urology residents in the United States and Canada. *Journal of Urology*, 177(4):1481–1487. doi:10.1016/j.juro.2006.11.074 — the case's primary RCT. — Retrieved from: \_\_\_\_\_
3. Kerfoot, B. P. (2009). Learning benefits of on-line spaced education persist for 2 years. *Journal of Urology*, 181(6):2671 — the two-year persistence follow-up. — Retrieved from: \_\_\_\_\_
4. Settles, B., & Meeder, B. (2016). A trainable spaced repetition model for language learning. *Proceedings of ACL 2016*, 1848–1858 — Duolingo half-life regression (Case 69), the operational deployment companion. — Retrieved from: \_\_\_\_\_

## CASE 18 **PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption** 2023

1. Kiguli-Malwadde, E., Forster, M., Eliaz, A., et al. (2023). "Comparing in-person, blended and virtual training interventions; a real-world evaluation of HIV capacity building programs in 16 countries in sub-Saharan Africa," *PLOS Global Public Health* 3(7):e0001654, doi:10.1371/journal.pgph.0001654 — the published version of medRxiv 2023.02.08.23285641. — Retrieved from: \_\_\_\_\_
2. The evaluated program is STRIPE HIV, led by the African Forum for Research and Education in Health (AFRE-health) and the University of California, San Francisco, and funded by HRSA through PEPFAR — named in the study's program description and funding statement. — Retrieved from: \_\_\_\_\_
3. Kirkpatrick & Kirkpatrick (2006), *Evaluating Training Programs* — the chain-of-evidence framework the L1–L2 limitation references (paired Case 79). — Retrieved from: \_\_\_\_\_

4. Blume et al. (2010), *Journal of Management* 36(4):1065–1105 — the meta-analytic transfer finding the L3 question references (paired Case 65). — Retrieved from: \_\_\_\_\_

### CASE 19 Keystone ICU / Pronovost Checklist

2004 – present

1. Pronovost & Vohr (2010), *Safe Patients, Smart Hospitals* — the checklist-plus-nurse-authority pairing (paraphrased). — Retrieved from: \_\_\_\_\_
2. Pronovost, P. et al. (2006), “An Intervention to Decrease Catheter-Related Bloodstream Infections in the ICU,” *NEJM* 355 — the trial and the near-zero result. — Retrieved from: \_\_\_\_\_
3. Bosk, C. et al. (2009), “Reality check for checklists,” *The Lancet* — the critique: the Keystone result cannot be attributed cleanly to the checklist, and the bundle’s active ingredient is not established. — Retrieved from: \_\_\_\_\_
4. Pronovost, P. et al. (2016), “Sustaining Reductions in CLABSI in Michigan ICUs: A 10-Year Analysis,” *Am J Med Qual* 31(3):197–202 — the effect sustained through 2013. — Retrieved from: \_\_\_\_\_
5. Agency for Healthcare Research and Quality, CUSP toolkit — dissemination across states. — Retrieved from: \_\_\_\_\_

### CASE 20 TREWS / Sepsis Watch

2018 – 2022

1. Sendak et al. (2020), “Real-world integration of a sepsis deep learning technology into routine clinical care: implementation study,” *JMIR Medical Informatics* 8(7):e15182 (Sepsis Watch implementation). — Retrieved from: \_\_\_\_\_
2. Adams et al. (2022), “Prospective, multi-site study of patient outcomes after implementation of the TREWS machine learning-based early warning system for sepsis,” *Nature Medicine* 28(7):1455–1460, doi:10.1038/s41591-022-01894-0. — Retrieved from: \_\_\_\_\_
3. Henry et al. (2022), “Factors driving provider adoption of the TREWS machine learning-based early warning system and its effects on sepsis treatment timing,” *Nature Medicine* 28(7):1447–1454, doi:10.1038/s41591-022-01895-z. — Retrieved from: \_\_\_\_\_
4. Wong et al. (2021), “External Validation of a Widely Implemented Proprietary Sepsis Prediction Model,” *JAMA Internal Medicine* 181(8):1065–1070 — the foil case (102). — Retrieved from: \_\_\_\_\_

### CASE 21 Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff

2024 – 2025

1. Federal Aviation Administration, 14 CFR § 121.542 (codified 1981) — origin of the aviation sterile-cockpit rule. — Retrieved from: \_\_\_\_\_
2. Starmer, A. J. et al. (2014), “Changes in medical errors after implementation of a handoff program,” *New England Journal of Medicine* 371(19):1803–1812 — I-PASS handoff intervention; structural cousin in the structured-information half of handoff capability. — Retrieved from: \_\_\_\_\_
3. Treloar, E., Herath, M., Altree, M., Potter, S., Penhall, M., Walsh, D., Kennedy, L., Bruening, M., Edwards, S., Ey, J. D., Bradshaw, E. L., & Maddern, G. J. (2025), “A Simple Solution for a Complex Problem: The ‘Sterile Cockpit’ to Improve Ward Rounds,” *World Journal of Surgery* 49(10):2769–2776, doi:10.1002/wjs.70074, PMID:40930848, PMCID:PMC12515027 — the cited adaptation study. — Retrieved from: \_\_\_\_\_
4. Broom, M. A., Capek, A. L., Carachi, P., Akeroyd, M. A., & Hilditch, G. (2011), “Critical phase distractions in anaesthesia and the sterile cockpit concept,” *Anaesthesia* 66(3):175–179 — prior anaesthesia-domain sterile-cockpit adaptation establishing the cross-domain transfer pattern. — Retrieved from: \_\_\_\_\_

**CASE 22 ChatGPT in Healthcare — Hallucination Cases**

2023 – present

1. Wachter, R. M., & Brynjolfsson, E. (2024). "Will Generative Artificial Intelligence Deliver on Its Promise in Health Care?" *JAMA* 331(1):65–69, doi:10.1001/jama.2023.25054 — genAI adoption and the productivity paradox in health care. — Retrieved from: \_\_\_\_\_
2. Morath, B., Chiriac, U., Jaszowski, E., et al. (2024). "Performance and risks of ChatGPT used in drug information: an exploratory real-world analysis," *European Journal of Hospital Pharmacy* 31(6):491–497 — only 13 of 50 answers correct and complete enough to act on without risk of harm. — Retrieved from: \_\_\_\_\_
3. Bhattacharyya, M., Miller, V. M., Bhattacharyya, D., & Miller, L. E. (2023). "High Rates of Fabricated and Inaccurate References in ChatGPT-Generated Medical Content," *Cureus* 15(5):e39238 — 47% of 115 generated references fabricated. — Retrieved from: \_\_\_\_\_
4. Triplett, S., Ness-Engle, G. L., & Behnen, E. M. (2025). "A comparison of drug information question responses by a drug information center and by ChatGPT," *American Journal of Health-System Pharmacy* 82(8):448–460, doi:10.1093/ajhp/zxae316 — pharmacist-rated overall accuracy of 50% across the ChatGPT responses. — Retrieved from: \_\_\_\_\_
5. World Health Organization (2024). Ethics and Governance of Artificial Intelligence for Health: Guidance on Large Multi-Modal Models, 18 January 2024 — names automation bias by health care professionals and the risk of false or inaccurate output. — Retrieved from: \_\_\_\_\_
6. Sallam, M. (2023). "ChatGPT Utility in Healthcare Education, Research, and Practice: Systematic Review on the Promising Perspectives and Valid Concerns," *Healthcare* 11(6):887 — 60 records reviewed; benefits cited in 85%, with the accuracy and citation concerns alongside. — Retrieved from: \_\_\_\_\_

**CASE 23 WHO Surgical Safety Checklist**

2008 – present

1. Haynes, A. et al. (2009), "A Surgical Safety Checklist to Reduce Morbidity and Mortality," *NEJM* — the 1.5%→0.8% result and major-complication reduction. — Retrieved from: \_\_\_\_\_
2. WHO Safe Surgery Saves Lives campaign documentation — the nineteen-item checklist and the three junctures. — Retrieved from: \_\_\_\_\_
3. Gawande, A. (2009), *The Checklist Manifesto* — the pause as the active mechanism. — Retrieved from: \_\_\_\_\_
4. Urbach, D. et al. (2014), "Introduction of Surgical Safety Checklists in Ontario, Canada," *NEJM* — null mortality result after mandate. — Retrieved from: \_\_\_\_\_
5. Bosk, C. et al. (2009), "Reality check for checklists," *The Lancet* — implementation fidelity. — Retrieved from: \_\_\_\_\_
6. Buissonniere, M. (2020), *Checking In on the Checklist*, Ariadne Labs and Lifebox — the decade review: facilities in 70% of the world's countries attest to using the checklist, with adoption far lower in low-income settings. — Retrieved from: \_\_\_\_\_

**CASE 24 Bristol Heart Babies Reform**

1984 – present

1. Kennedy, I. (2001), *Learning from Bristol: The Report of the Public Inquiry into Children's Heart Surgery at the Bristol Royal Infirmary 1984–1995* — the inquiry and its recommendations, paraphrased; the pull quote is verbatim from paragraph 18 of the Final Report Summary, and differs in wording from Recommendation 155, which carries the same requirement. — Retrieved from: \_\_\_\_\_
2. Society for Cardiothoracic Surgery in GB & Ireland, national outcomes reports — the registry and surgeon-level publication. — Retrieved from: \_\_\_\_\_

3. Berwick, D. (2013), A Promise to Learn — A Commitment to Act — the broader NHS–safety reform. — Retrieved from: \_\_\_\_\_
4. Sherlaw-Johnson, C., Gallivan, S., Treasure, T. & Nashef, S. A. M. (2004), “Computer tools to assist the monitoring of outcomes in surgery,” *European Journal of Cardio-Thoracic Surgery* 26(5):1032–1036 — risk-adjusted outcome monitoring (variable life-adjusted display) as it developed in UK cardiac surgery. — Retrieved from: \_\_\_\_\_
5. NHS national clinical audit and registry documentation — extension across specialties. — Retrieved from: \_\_\_\_\_

## CASE 25 Removing the Race Coefficient from eGFR

2021

1. Vyas, Eisenstein, & Jones (2020), “Hidden in Plain Sight — Reconsidering the Use of Race Correction in Clinical Algorithms,” *NEJM* 383(9):874–882 — broader race-in-algorithms survey. — Retrieved from: \_\_\_\_\_
2. Eneanya, Yang, & Reese (2019), “Reconsidering the Consequences of Using Race to Estimate Kidney Function,” *JAMA* 322(2):113–114 — the equity argument that motivated the revision. — Retrieved from: \_\_\_\_\_
3. Genzen, J. R., Souers, R. J., Pearson, L. N., et al. (2022), “Reported Awareness and Adoption of 2021 Estimated Glomerular Filtration Rate Equations Among US Clinical Laboratories, March 2022,” *JAMA* 328(20):2060–2062, doi:10.1001/jama.2022.15404 — the College of American Pathologists survey: 1,124 of 3,708 responding U.S. laboratories (30.3%) had adopted the 2021 CKD–EPI creatinine equation. — Retrieved from: \_\_\_\_\_
4. Delgado et al. (2021), “A Unifying Approach for GFR Estimation: Recommendations of the NKF–ASN Task Force on Reassessing the Inclusion of Race in Diagnosing Kidney Disease,” *American Journal of Kidney Diseases* 79(2):268–288 (published online 2021; in print vol. 79, 2022), doi:10.1053/j.ajkd.2021.08.003. Cited by online—first year, the year of the Task Force recommendation. — Retrieved from: \_\_\_\_\_
5. Inker et al. (2021), “New Creatinine- and Cystatin C–Based Equations to Estimate GFR without Race,” *New England Journal of Medicine* 385(19):1737–1749, doi:10.1056/NEJMoa2102953. — Retrieved from: \_\_\_\_\_
6. OPTN/UNOS (2022–2023), “Kidney waiting-time modifications for candidates affected by race-inclusive eGFR calculations” — the mandated backdating of Black candidates’ accrued waiting time (effective January 2023). — Retrieved from: \_\_\_\_\_

## CASE 26 Pulse Oximetry Across Skin Tones

1990 – 2020 – 2025

1. Jubran & Tobin (1990), “Reliability of Pulse Oximetry in Titrating Supplemental Oxygen Therapy in Ventilator-Dependent Patients,” *Chest* 97(6):1420–1425 — original finding, published thirty years earlier. — Retrieved from: \_\_\_\_\_
2. Fawzy et al. (2022), “Racial and Ethnic Discrepancy in Pulse Oximetry and Delayed Identification of Treatment Eligibility Among Patients With COVID-19,” *JAMA Internal Medicine* — downstream effect during the pandemic. — Retrieved from: \_\_\_\_\_
3. Sjoding, Dickson, Iwashyna, Gay, & Valley (2020), “Racial Bias in Pulse Oximetry Measurement,” *New England Journal of Medicine* 383(25):2477–2478, doi:10.1056/NEJMc2029240. — Retrieved from: \_\_\_\_\_
4. FDA (2025), “Pulse Oximeters for Medical Purposes — Non-Clinical and Clinical Performance Testing, Labeling, and Premarket Submission Recommendations: Draft Guidance for Industry and Food and Drug Administration Staff,” issued January 7 2025, Docket No. FDA-2023-N-4976; Federal Register notice 2024–31540,

published January 7 2025 — the regulatory corrective-action artifact, language may evolve in final. — Retrieved from: \_\_\_\_\_

## CASE 27 **Anesthesia Monitoring Standards and the APSF — The Mortality Decline** 1986 – present

1. Eichhorn, Cooper, Cullen, Maier, Philip, & Seeman (1986), "Standards for Patient Monitoring During Anesthesia at Harvard Medical School," JAMA 256(8):1017–1020. — Retrieved from: \_\_\_\_\_
2. American Society of Anesthesiologists, "Standards for Basic Anesthetic Monitoring," approved 21 October 1986 and amended since — the original standard and the amendment sequence that made pulse oximetry required (1989), expired-CO2 verification of tube placement required (early 1990s), and continual capnography through a general anesthetic required (effective July 1999). — Retrieved from: \_\_\_\_\_
3. Eichhorn (1989), "Prevention of intraoperative anesthesia accidents and related severe injury through safety monitoring," Anesthesiology 70(4):572–577. — Retrieved from: \_\_\_\_\_
4. Braz, L. G., Módolo, N. S. P., do Nascimento, P. Jr., et al. (2009), "Mortality in anesthesia: a systematic review," Clinics (São Paulo) 64(10):999–1006 — 33 studies, 1954–2007; anaesthesia-related mortality below one death per 10,000 anaesthetics in the final two decades. — Retrieved from: \_\_\_\_\_
5. Anesthesia Patient Safety Foundation (1985 – present), founding documents and the APSF Newsletter — institutional-history record of the broader change effort. — Retrieved from: \_\_\_\_\_
6. Sjoding et al. (2020), "Racial Bias in Pulse Oximetry Measurement," NEJM 383(25):2477–2478 — the racial-bias edge case the standard still carries (cross-reference Case 26). — Retrieved from: \_\_\_\_\_

## CASE 28 **Interprofessional Education and the Evidence Gap** 2013 – 2015

1. Reeves, Perrier, Goldman, Freeth, & Zwarenstein (2013), "Interprofessional education: effects on professional practice and healthcare outcomes (update)," Cochrane Database of Systematic Reviews, doi:10.1002/14651858.CD002213.pub3. — Retrieved from: \_\_\_\_\_
2. Institute of Medicine (2015), Measuring the Impact of Interprofessional Education on Collaborative Practice and Patient Outcomes, National Academies Press, doi:10.17226/21726, NCBI Bookshelf NBK338360. — Retrieved from: \_\_\_\_\_
3. WHO (2010), Framework for Action on Interprofessional Education and Collaborative Practice — the international policy backdrop. — Retrieved from: \_\_\_\_\_
4. v2 paired cases: Team-science training (Case 40), Implementation-science training (Case 41). — Retrieved from: \_\_\_\_\_

## CASE 29 **Bar-Code Medication Administration — Cue at the Point of Care** 2010

1. Poon, E. G., Keohane, C. A., Yoon, C. S., Ditmore, M., Bane, A., Levtzion-Korach, O., Moniz, T., Rothschild, J. M., Kachalia, A. B., Hayes, J., Churchill, W. W., Lipsitz, S., Whittemore, A. D., Bates, D. W., & Gandhi, T. K. (2010). Effect of bar-code technology on the safety of medication administration. New England Journal of Medicine, 362(18):1698–1707. doi:10.1056/NEJMs0907115 — the case's primary evaluation. — Retrieved from: \_\_\_\_\_
2. Bonkowski, J., Carnes, C., Melucci, J., Mirtallo, J., Prier, B., Reichert, E., Moffatt-Bruce, S., & Weber, R. J. (2013). Effect of barcode-assisted medication administration on emergency department medication errors. Academic Emergency Medicine, 20(8):801–806 — adjacent transfer evidence. — Retrieved from: \_\_\_\_\_



3. Thompson, K. M., Swanson, K. M., Cox, D. L., Kirchner, R. B., Russell, J. J., Wermers, R. A., Storlie, C. B., Johnson, M. G., & Naessens, J. M. (2018), "Implementation of Bar-Code Medication Administration to Reduce Patient Harm," *Mayo Clinic Proceedings: Innovations, Quality & Outcomes* 2(4):342–351, doi:10.1016/j.mayocpiqo.2018.09.001, PMID:30560236, PMCID:PMC6257885 — later single-site rollout reporting a 55.4% reduction in harmful medication errors (0.65 to 0.29 per 100,000 medications), by interrupted time series across fifty inpatient units. — Retrieved from: \_\_\_\_\_
4. Institute for Safe Medication Practices (2019), *Guidelines for Safe Electronic Communication of Medication Information — the display and communication standards that bar-code scanning systems and eMARs are built against*. — Retrieved from: \_\_\_\_\_

## CASE 30 Surgical Skill Video Peer-Rating Predicts Complications

2013

1. Birkmeyer, J. D., Finks, J. F., O'Reilly, A., Oerline, M., Carlin, A. M., Nunn, A. R., Dimick, J., Banerjee, M., & Birkmeyer, N. J. O., for the Michigan Bariatric Surgery Collaborative. (2013). Surgical skill and complication rates after bariatric surgery. *New England Journal of Medicine*, 369(15):1434–1442. doi:10.1056/NEJMs1300625 — the case's primary evaluation. — Retrieved from: \_\_\_\_\_
2. Birkmeyer, J. D., Stukel, T. A., Siewers, A. E., Goodney, P. P., Wennberg, D. E., & Lucas, F. L. (2003). Surgeon volume and operative mortality in the United States. *New England Journal of Medicine*, 349(22):2117–2127 — the volume-outcome literature the confound rests against. — Retrieved from: \_\_\_\_\_
3. Vassiliou, M. C., Feldman, L. S., Andrew, C. G., et al. (2005). A global assessment tool for evaluation of intraoperative laparoscopic skills. *American Journal of Surgery*, 190(1):107–113 — the laparoscopic skill-rating literature this video assessment sits in. — Retrieved from: \_\_\_\_\_
4. Jacobs, D. O. (2013). Cut well, sew well, do well? *New England Journal of Medicine*, 369(15):1466–1467. doi:10.1056/NEJMe1309785 — the accompanying editorial. — Retrieved from: \_\_\_\_\_

## CASE 31 Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units

2006 – 2016

1. Vedula, Malpani, Tao, Chen, Gao, Poddar, Ahmidi, Paxton, Vidal, Khudanpur, Hager, & Chen (2016), "Analysis of the Structure of Surgical Activity for a Suturing and Knot-Tying Task," *PLOS One* 11(3):e0149174, doi:10.1371/journal.pone.0149174. — Retrieved from: \_\_\_\_\_
2. Gao, Vedula, Reiley, Ahmidi, Varadarajan, Lin, Tao, Zappella, Bejar, Yuh, Chen, Vidal, Khudanpur, & Hager (2014), "JHU-ISI Gesture and Skill Assessment Working Set (JIGSAWS): A Surgical Activity Dataset for Human Motion Modeling," *MICCAI Workshop — JIGSAWS dataset release paper*. — Retrieved from: \_\_\_\_\_
3. Reiley, Lin, Yuh, & Hager (2011), "Review of methods for objective surgical skill evaluation," *Surgical Endoscopy* 25(2):356–366 — situates the decomposition within the broader skill-assessment literature. — Retrieved from: \_\_\_\_\_
4. Birkmeyer, Finks, O'Reilly, et al. (2013), *NEJM* — the paired skill-matters study (Case 30). — Retrieved from: \_\_\_\_\_

## CASE 32 Annual-Screening UI Redesign at University of Missouri Health Care

2020

1. Pierce RP, Eskridge BR, Rehard L, Ross B, Day MA, Belden JL (2020), "The Effect of Electronic Health Record Usability Redesign on Annual Screening Rates in an Ambulatory Setting," *Applied Clinical Informatics* 11(4):580–588, doi:10.1055/s-0040-1715828 — the source of record for every figure in this case. HIMSS carries a secondary write-up of the project; it describes the redesign as adding clinical decision support, which the peer-reviewed report does not, and this case follows the report. — Retrieved from: \_\_\_\_\_
2. Hussain MI, Reynolds TL, Zheng K (2019), "Medication safety alert fatigue may be reduced via interaction design and clinical role tailoring: a systematic review," *JAMIA* 26(10):1141–1149, doi:10.1093/jamia/ocz095 —

adjacent evidence that interaction–design changes move alert behavior. Note its outcome is prescriber acceptance of medication safety alerts, not screening completion; it is background, and no figure in this case rests on it. — Retrieved from: \_\_\_\_\_

3. Ratwani RM, Reider J, Singh H (2019), “A decade of health information technology usability challenges and the path forward,” *JAMA* 321(8):743–744, doi:10.1001/jama.2019.0161 — the call for lifecycle evaluation “from initial product design to local implementation followed by periodic reassessment,” and the framing for the site–configuration gap in federal certification requirements. — Retrieved from: \_\_\_\_\_
4. Middleton et al. (2013), “Enhancing patient safety and quality of care by improving the usability of electronic health record systems,” *JAMIA* 20(e1):e2–e8 — the framing peer–reviewed paper on EHR–usability–as–safety. — Retrieved from: \_\_\_\_\_

### CASE 33 **Alert–Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal** 2019 – 2024

1. Ancker et al. (2017), “Effects of workload, work complexity, and repeated alerts on alert fatigue in a clinical decision support system,” *BMC Medical Informatics and Decision Making* 17:36 — adjacent measurement evidence. — Retrieved from: \_\_\_\_\_
2. Hussain, M. I., Reynolds, T. L., & Zheng, K. (2019), “Medication safety alert fatigue may be reduced via interaction design and clinical role tailoring: a systematic review,” *JAMIA* 26(10):1141–1149, doi:10.1093/jamia/ocz095 — thirty–nine studies; interruptive modals accepted least often; role tailoring the only alternative that appeared to increase prescriber acceptance. — Retrieved from: \_\_\_\_\_
3. Fallon A, Haralambides K, Mazzillo J, Gleber C (2024), “Addressing Alert Fatigue by Replacing a Burden–some Interruptive Alert with Passive Clinical Decision Support,” *Applied Clinical Informatics* 15(1):101–110, doi:10.1055/a-2226-8144 — 80% cut in alert firings; COVID precautions ordered rose from 23% to 61%. — Retrieved from: \_\_\_\_\_
4. Van Dort, B. A., Zheng, W. Y., Sundar, V., & Baysari, M. T. (2021), “Optimizing clinical decision support alerts in electronic medical records: a systematic review of reported strategies adopted by hospitals,” *JAMIA* 28(1):177–183, doi:10.1093/jamia/ocaa279 — eight published optimization programs; six evaluated their changes and five of those reported a reduction in alert rate. — Retrieved from: \_\_\_\_\_
5. Office of the National Coordinator for Health IT, SAFER Guides on CDS — practitioner–tier guidance the redesigns instantiate. — Retrieved from: \_\_\_\_\_
6. Patterson E (2024), “Navigating Alert Fatigue: A Case Study in Electronic Health Record Alert Design Optimization,” *Studies in Health Technology and Informatics* 315:447–451, doi:10.3233/SHTI240188 — nursing–workflow QI redesign of four high–firing alerts; all four removed, 877 unactionable disruptive nursing hours released. — Retrieved from: \_\_\_\_\_

### CASE 34 **COMPOSER Sepsis Prediction — The Third Clinical–AI Sepsis Case** 2022 – 2024

1. Boussina, A., Shashikumar, S. P., Malhotra, A., Owens, R. L., El–Kareh, R., Longhurst, C. A., Quintero, K., et al. (2024), “Impact of a deep learning sepsis prediction model on quality of care and survival,” *NPJ Digital Medicine* 7:14, doi:10.1038/s41746-023-00986-6. — Retrieved from: \_\_\_\_\_
2. Shashikumar, S. P., Wardi, G., Malhotra, A., & Nemati, S. (2021), “Artificial intelligence sepsis prediction algorithm learns to say ‘I don’t know,’” *NPJ Digital Medicine* 4:134 — the methodological–foundation paper for the conformal–prediction abstention structure. — Retrieved from: \_\_\_\_\_
3. Wong, A., Otles, E., Donnelly, J. P., Krumm, A., McCullough, J., DeTroyer–Cooley, O., et al. (2021), “External Validation of a Widely Implemented Proprietary Sepsis Prediction Model in Hospitalized Patients,” *JAMA*



Internal Medicine 181(8):1065–1070 — the external-validation paper on Epic Sepsis (Case 5). — Retrieved from: \_\_\_\_\_

4. Adams, R., Henry, K. E., Sridharan, A., Soleimani, H., Zhan, A., Rawat, N., Johnson, L., et al. (2022), “Prospective, multi-site study of patient outcomes after implementation of the TREWS machine learning-based early warning system for sepsis,” *Nature Medicine* 28:1455–1460 — the prospective-positive TREWS deployment paper (Case 20). — Retrieved from: \_\_\_\_\_

## CASE 35 Radiology AI Miscalibration

2018 – present

1. Wu, E., Wu, K., Daneshjou, R., Ouyang, D., Ho, D. E., & Zou, J. (2021), “How medical AI devices are evaluated: limitations and recommendations from an analysis of FDA approvals,” *Nature Medicine* 27(4):582–584, doi:10.1038/s41591-021-01312-x — the 130-device audit (126 retrospective-only, none of the 54 high-risk devices prospective, 93 without publicly reported multi-site assessment, 17 reporting demographic subgroup performance) and the cross-site pneumothorax experiment behind this case’s diagram. — Retrieved from: \_\_\_\_\_
2. Seyyed-Kalantari, L., Zhang, H., McDermott, M. B. A., Chen, I. Y., & Ghassemi, M. (2021), “Underdiagnosis bias of artificial intelligence algorithms applied to chest radiographs in under-served patient populations,” *Nature Medicine* 27:2176–2182 — disparities across Black, Hispanic, female, and lower-socioeconomic subgroups; underdiagnosis highest for intersectional subpopulations such as Hispanic female patients. — Retrieved from: \_\_\_\_\_
3. Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019), “Dissecting racial bias in an algorithm used to manage the health of populations,” *Science* 366(6464):447–453 — proxy-label bias in a care-management algorithm; 200 million Americans/year are subject to tools of this class. — Retrieved from: \_\_\_\_\_
4. Daneshjou, R., Vodrahalli, K., Novoa, R. A., et al. (2022), “Disparities in dermatology AI performance on a diverse, curated clinical image set,” *Science Advances* 8(32):eabq6147, doi:10.1126/sciadv.abq6147 — the Diverse Dermatology Images set; models and labeling dermatologists both degrade on dark skin tones. — Retrieved from: \_\_\_\_\_
5. FDA, “Proposed Regulatory Framework for Modifications to AI/ML-Based Software as a Medical Device” (2019); FDA’s public “Artificial Intelligence-Enabled Medical Devices” authorization list; final Predetermined Change Control Plan guidance (December 2024) and the draft “Artificial Intelligence-Enabled Device Software Functions: Lifecycle Management and Marketing Submission Recommendations” (January 2025), with Good Machine Learning Practice principles — the regulatory trajectory; absence of mandatory demographic stratification at clearance and post-market monitoring of in-use performance. — Retrieved from: \_\_\_\_\_
6. Larrazabal, A. J., Nieto, N., Peterson, V., Milone, D. H., & Ferrante, E. (2020), “Gender imbalance in medical imaging datasets produces biased classifiers for computer-aided diagnosis,” *PNAS* 117(23):12592–12594 — AUC drops for under-represented genders on NIH ChestX-ray14 and CheXpert, across three network architectures. — Retrieved from: \_\_\_\_\_
7. Bernhardt, M., Jones, C., & Glocker, B. (2022), “Potential sources of dataset bias complicate investigation of underdiagnosis by machine learning algorithms,” *Nature Medicine* 28(6):1157–1158 — one of three Matters Arising contesting Seyyed-Kalantari et al.; the authors’ reply is at 28(6):1161–1162, doi:10.1038/s41591-022-01854-8. Cited here so the contested status travels with the finding. — Retrieved from: \_\_\_\_\_
8. Wachter, R. M. & Brynjolfsson, E. (2024), “Will Generative Artificial Intelligence Deliver on Its Promise in Health Care?” *JAMA* 331(1):65–69, doi:10.1001/jama.2023.25054 — the counterweight argument that health care is readier for generative AI than for its predecessors and the lag to benefit may be shorter. It is not a bias or evaluation paper; do not cite it for the proxy-and-population problem. — Retrieved from: \_\_\_\_\_

## CASE 36 AlphaFold — Protein Structure Prediction

2020 – present

1. Jumper, J., Evans, R., Pritzel, A., et al. (2021), “Highly accurate protein structure prediction with AlphaFold,” *Nature* 596(7873):583–589, doi:10.1038/s41586-021-03819-2 — the method, and the CASP14 validation. — Retrieved from: \_\_\_\_\_
2. Varadi, M., Bertoni, D., Magana, P., et al. (2024), “AlphaFold Protein Structure Database in 2024: providing structure coverage for over 214 million protein sequences,” *Nucleic Acids Research* 52(D1):D368–D375, doi:10.1093/nar/gkad1011 — the database at 214M, up from 300k in 2021, and its integration into PDB, UniProt, Ensembl and InterPro. — Retrieved from: \_\_\_\_\_
3. CASP14 press release, Protein Structure Prediction Center, 30 November 2020 — GDT above 90 on roughly two-thirds of targets. Note the Moulton quotation this case carries is **not** in the press release: it is from Google DeepMind’s post of the same date, “AlphaFold: a solution to a 50-year-old grand challenge in biology,” which also reports the 92.4 median GDT. — Retrieved from: \_\_\_\_\_
4. Tunyasuvunakool, K., Adler, J., Wu, Z., et al. (2021), “Highly accurate protein structure prediction for the human proteome,” *Nature* 596(7873):590–596, doi:10.1038/s41586-021-03828-1 — 98.5% of human proteins covered; 58% of residues confidently predicted, 36% of all residues at very high confidence, against 17% of residues covered by experimental structure beforehand. — Retrieved from: \_\_\_\_\_
5. AlphaFold 3 access chronology: publication with neither code nor weights (May 2024), the open letter signed by more than a thousand researchers, the *Nature* editorial “AlphaFold3 — why did *Nature* publish it without its code?” (2024, doi:10.1038/d41586-024-01463-0), and the 11 November 2024 GitHub release of the code under a non-commercial license with weights by application — the open-release governance decision as actually written. — Retrieved from: \_\_\_\_\_
6. Abramson, J., Adler, J., Dunger, J., et al. (2024), “Accurate structure prediction of biomolecular interactions with AlphaFold 3,” *Nature* 630(8016):493–500, doi:10.1038/s41586-024-07487-w; and the 2024 Nobel Prize in Chemistry, one half to David Baker for computational protein design and one half jointly to Demis Hassabis and John M. Jumper for protein structure prediction. — Retrieved from: \_\_\_\_\_

## CASE 37 DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback

2020

1. McKinney, S. M., Sieniek, M., Godbole, V., Godwin, J., Antropova, N., Ashrafian, H., Back, T., et al. (2020), “International evaluation of an AI system for breast cancer screening,” *Nature* 577:89–94, doi:10.1038/s41586-019-1799-6. — Retrieved from: \_\_\_\_\_
2. Haibe-Kains, B., Adam, G. A., Hosny, A., Khodakarami, F., MAQC Society Board of Directors, Waldron, L., Wang, B., et al. (2020), “Transparency and reproducibility in artificial intelligence,” *Nature* 586:E14–E16, doi:10.1038/s41586-020-2766-y. — Retrieved from: \_\_\_\_\_
3. McKinney, S. M., Karthikesalingam, A., Tse, D., Kelly, C. J., Liu, Y., Corrado, G. S., & Shetty, S. (2020), “Reply to: Transparency and reproducibility in artificial intelligence,” *Nature* 586:E17–E18, doi:10.1038/s41586-020-2767-x — the developers’ response, published alongside the addendum below. — Retrieved from: \_\_\_\_\_
4. McKinney et al. (2020), “Addendum: International evaluation of an AI system for breast cancer screening,” *Nature* 586:E19, doi:10.1038/s41586-020-2679-9 — published in the same issue as the Haibe-Kains comment; expands the supplementary methods with optimization hyperparameters and the augmentation strategy to assist replication. Code and trained models were not released. — Retrieved from: \_\_\_\_\_
5. Freeman, K., Geppert, J., Stinton, C., Todkill, D., Johnson, S., Clarke, A., & Taylor-Phillips, S. (2021), “Use of artificial intelligence for image analysis in breast cancer screening programmes: systematic review of test accuracy,” *BMJ* 374:n1872, doi:10.1136/bmj.n1872 — twelve studies, 131,822 women, McKinney included; no prospective

studies found, studies of poor methodological quality, 34 of 36 AI systems less accurate than a single radiologist and all less accurate than two in consensus. — Retrieved from: \_\_\_\_\_

6. MASAI trial (Lång et al.): clinical safety analysis, *Lancet Oncology* 2023;24(8):936–944, doi:10.1016/S1470-2045(23)00298-X; Hernström et al., screening performance, *Lancet Digital Health* 2025;7(3):e175–e183, doi:10.1016/S2589-7500(24)00267-X; Gommers et al., interval cancer, sensitivity and specificity, *The Lancet* 2026;407(10527):505–514, doi:10.1016/S0140-6736(25)02464-X. 105,934 women randomized; the AI system is Transpara v1.7.0 (ScreenPoint Medical), not the system evaluated by McKinney et al. — Retrieved from: \_\_\_\_\_

### CASE 38 **iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms** 2006 – 2011

1. FDA, iPLEDGE program documentation (2006 – present) — REMS architecture and enrollment requirements. — Retrieved from: \_\_\_\_\_
2. Sullivan et al. (2003), House Science Committee statement on SUBSAFE — the structural counterpoint (Case 173). — Retrieved from: \_\_\_\_\_
3. Tkachenko, E., Singer, S., Sharma, P., Barbieri, J., & Mostaghimi, A. (2019), “US Food and Drug Administration reports of pregnancy and pregnancy-related adverse events associated with isotretinoin,” *JAMA Dermatology*, 155(10):1175–1179 — source of the 218 to 310 annual reports after 2011. — Retrieved from: \_\_\_\_\_
4. Shin, J., Cheetham, T. C., Wong, L., Niu, F., Kass, E., Yoshinaga, M. A., Sorel, M., McCombs, J. S., & Sidney, S. (2011), “The impact of the iPLEDGE program on isotretinoin fetal exposure in an integrated health care system,” *Journal of the American Academy of Dermatology*, PMID:21565419. — Retrieved from: \_\_\_\_\_
5. Collins, M. K., Moreau, J. F., Opel, D., Swan, J., Prevost, N., Hastings, M., Schwarz, E. B., & Ferris, L. K. (2014), “Compliance with pregnancy prevention measures during isotretinoin therapy,” *Journal of the American Academy of Dermatology*, 70(1):55–59, PMID:24157382 — source of the 150 annual exposures figure. — Retrieved from: \_\_\_\_\_

### CASE 39 **TeamSTEPPS** 2006 – present

1. Salas, E., DiazGranados, D., Klein, C., Burke, C. S., Stagl, K. C., Goodwin, G. F. & Halpin, S. M. (2008), “Does Team Training Improve Team Performance? A Meta-Analysis,” *Human Factors* 50(6):903–933 — the cross-domain team-training evidence base. — Retrieved from: \_\_\_\_\_
2. Umscheid, C. & Haugstetter, M. (2023), “With TeamSTEPPS 3.0, AHRQ Refreshes a Landmark Patient Safety Training Curriculum,” *AHRQ Views* (12 September) — the DoD partnership and the first-eight-years adoption figures (1,500 hospitals, 5,000 master trainers, 300,000 professionals). — Retrieved from: \_\_\_\_\_
3. Weaver, S. J., Dy, S. M. & Rosen, M. A. (2014), “Team-training in healthcare: a narrative synthesis of the literature,” *BMJ Quality & Safety* 23(5):359–372 — bundled team-training interventions, with tools and organisational change to support transfer, show the largest effects. — Retrieved from: \_\_\_\_\_
4. Weld, L. R. et al. (2016), “TeamSTEPPS Improves Operating Room Efficiency and Patient Safety,” *American Journal of Medical Quality* 31(5):408–414 — the on-time first-start figure. — Retrieved from: \_\_\_\_\_
5. AHRQ, TeamSTEPPS 3.0 Curriculum (2023) — the framework and four competencies. — Retrieved from: \_\_\_\_\_

## CASE 40 Team Science Training for Clinical and Translational Scientists 2020 – 2022

1. Mook, A., Knerich, V., Komaie, G., Cicutto, L., & Cross, J. (2025), "Team science training for clinical and translational scientists: An assessment of effectiveness," *Journal of Clinical and Translational Science* 9(1):e158, doi:10.1017/cts.2025.10088. — Retrieved from: \_\_\_\_\_
2. Falk-Krzesinski et al. (2011), "Mapping a research agenda for the science of team science," *Research Evaluation* 20(2):143–156 — broader team-science literature backdrop. — Retrieved from: \_\_\_\_\_
3. National Research Council (2015), *Enhancing the Effectiveness of Team Science* (National Academies Press) — the National Academies team-science synthesis. — Retrieved from: \_\_\_\_\_
4. Paired cases in this volume: IPE evidence gap (Case 28), Implementation Science Training (Case 41) — the frontier/measurement companions. — Retrieved from: \_\_\_\_\_

## CASE 41 Translational-Workforce Training — Stated Goals Outrunning Named Practice 2020s

1. Morris, Wooding, & Grant (2011), "The answer is 17 years, what is the question: understanding time lags in translational research," *Journal of the Royal Society of Medicine* — the original 17-year-gap source for Case 13. — Retrieved from: \_\_\_\_\_
2. Sancheznieto, F., Sorkness, C. A., Attia, J., et al. (2021), "Clinical and translational science award T32/TL1 training programs: program goals and mentorship practices," *Journal of Clinical and Translational Science* 6(1):e13, doi:10.1017/cts.2021.884. — Retrieved from: \_\_\_\_\_
3. Brownson, Colditz, & Proctor (2018), *Dissemination and Implementation Research in Health* (2nd ed.) — the broader implementation-science synthesis. — Retrieved from: \_\_\_\_\_
4. Paired cases in this volume: Team-science training (Case 40), IPE evidence gap (Case 28). — Retrieved from: \_\_\_\_\_

## CASE 42 Australian Hospital-Pharmacy Technician Role Redesign 2016

1. Pharmacy Board of Australia and SHPA practice statements limiting technician work to activities not requiring professional judgement — the regulatory framing the project sits inside. — Retrieved from: \_\_\_\_\_
2. Society of Hospital Pharmacists of Australia (November 2016), "Exploring the role of hospital pharmacy technicians and assistants to enhance the delivery of patient centred care: A White Paper on the findings and outcomes of the Pharmacy Technician and Assistant Role Redesign within Australian Hospitals (Redesign Project)." — Retrieved from: \_\_\_\_\_
3. Anderson et al. (2021), "Perceptions of hospital pharmacists and pharmacy technicians towards expanding roles for hospital pharmacy technicians: a cross-sectional survey," *Journal of Pharmacy Practice and Research* 51(3):216–224, doi:10.1002/jppr.1697 — fielded September 2018; 61 of 122 staff (50%) responded. — Retrieved from: \_\_\_\_\_
4. Boughen, Sutton, Fenn, & Wright (2017), "Defining the Role of the Pharmacy Technician and Identifying Their Future Role in Medicines Optimisation," *Pharmacy* 5(3):40 — UK companion analysis. — Retrieved from: \_\_\_\_\_

### CASE 43 **Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface** 2013 – 2018

1. Rwanda Ministry of Health, community health program documentation and CHW scope-of-practice guidance, 2013–2018. — Retrieved from: \_\_\_\_\_
2. Ruton, H., Musabyimana, A., Gaju, E., Berhe, A., Grépin, K. A., Ngenzi, J., Nzabonimana, E., & Law, M. R. (2018), “The impact of an mHealth monitoring system on health care utilization by mothers and children: an evaluation using routine health information in Rwanda,” *Health Policy and Planning* 33(8):920–927, doi:10.1093/heapol/czy066 — interrupted time series on 461 health centres, 2012–2016; RapidSMS plus additional support raised utilization, RapidSMS alone did not. — Retrieved from: \_\_\_\_\_
3. Hategeka, C., Ruton, H., & Law, M. R. (2019), “Effect of a community health worker mHealth monitoring system on uptake of maternal and newborn health services in Rwanda,” *Global Health Research and Policy* 4:8, doi:10.1186/s41256-019-0098-y. — Retrieved from: \_\_\_\_\_
4. MIT News (2022), reporting on subsequent AI-augmented maternal-care work in Rwanda — journalism-tier companion to the peer-reviewed evaluation. — Retrieved from: \_\_\_\_\_
5. Cross-reference: Case 18 (PEPFAR HIV training-modality comparison) for the paired Sub-Saharan work-force-capability evidence. — Retrieved from: \_\_\_\_\_

### CASE 44 **Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD** 2014 – present

1. Kikuchi, et al. (2025), “Scoping Review of Regulatory Transparency in AI-based Radiology Software: Analysis of PMDA-approved SaMD Products,” *medRxiv* 2025.10.02.25336333 — preprint, not peer reviewed; 151 approved SaMD as of 31 December 2024, 40 using AI, 20 in radiology. — Retrieved from: \_\_\_\_\_
2. Aisu, N., Miyake, M., Takeshita, K., et al. (2022), “Regulatory-approved deep learning/machine learning-based medical devices in Japan as of 2020: A systematic review,” *PLOS Digital Health* 1(1):e0000001 (medRxiv preprint 2021). — Retrieved from: \_\_\_\_\_
3. Pharmaceuticals and Medical Devices Agency of Japan, PMD Act amendment (2019, enforced stepwise from September 2020) and DASH for SaMD — “DX Action Strategies in Healthcare for SaMD,” announced November 2020 — program documentation. — Retrieved from: \_\_\_\_\_
4. PMDA Science Board, Subcommittee on Software as a Medical Device Utilizing AI and Machine Learning (28 August 2023), Report on AI-based Software as a Medical Device (SaMD) — the regulator’s own statement that IDATEN “has yet to be fully used.” — Retrieved from: \_\_\_\_\_
5. Cross-reference: “A decade of review in global regulation and research of artificial intelligence medical devices (2015–2025),” *PMC12310608* — comparative regulatory context. — Retrieved from: \_\_\_\_\_

### CASE 45 **Tennessee Voluntary Pre-K Study** 2009–2018

1. National Institute for Early Education Research, State of Preschool yearbooks (2010–2022) — program scale and context. — Retrieved from: \_\_\_\_\_
2. K. Durkin, M. Lipsey, D. Farran & S. Wiesen, “Effects of a Statewide Pre-Kindergarten Program... Through Sixth Grade,” *Developmental Psychology* 58(3): 470–484 (2022) — the sixth-grade reversal (quoted). — Retrieved from: \_\_\_\_\_
3. Responses and counter-analyses to the TN-VPK findings (2018–2022) — the contested reception. — Retrieved from: \_\_\_\_\_

4. M. Lipsey, D. Farran & K. Durkin, "Effects of the Tennessee Pre-Kindergarten Program... Through Third Grade," *Early Childhood Research Quarterly* 45: 155–176 (2018) — the RCT and fade-out. — Retrieved from: \_\_\_\_\_
5. D. Phillips et al., *Puzzling It Out: The Current State of Scientific Knowledge on Pre-Kindergarten Effects* (2017); J. Heckman on durable early-childhood programs. — Retrieved from: \_\_\_\_\_

## CASE 46 Algorithmic Bias in Educational Predictive Analytics

ongoing

1. The Evolving State of Predictive Analytics, *The Chronicle of Higher Education* (2020), survey conducted by Maguire Associates, 589 institutions — 63 percent use predictive analytics for retention and graduation; only 20 percent institution-wide. — Retrieved from: \_\_\_\_\_
2. K. Bird, B. Castleman & Y. Song, "Are Algorithms Biased in Education? Exploring Racial Bias in Predicting Community College Student Success," *Journal of Policy Analysis and Management* 44(2) (2025), 379–402 — two Virginia Community College System models; bias over 5× higher when "at-risk" is the bottom decile than the bottom half for degree completion, and the reverse for course completion. — Retrieved from: \_\_\_\_\_
3. D. Gándara, H. Anahideh, M. Ison & L. Picchiarini, "Inside the Black Box: Detecting and Mitigating Algorithmic Bias across Racialized Groups in College Student-Success Prediction," *AERA Open* 10 (2024), doi:10.1177/23328584241258741 — higher false-negative rates for Black and Hispanic students in nationally representative ELS:2002 data; leading bias-mitigation techniques generally ineffective at eliminating the disparities. — Retrieved from: \_\_\_\_\_
4. R. Baker & A. Hawn, "Algorithmic Bias in Education," *International Journal of Artificial Intelligence in Education* 32 (2022), 1052–1092 — foundational review of algorithmic bias in education. — Retrieved from: \_\_\_\_\_
5. M. Ekowo & I. Palmer, *Predictive Analytics in Higher Education: Five Guiding Practices for Ethical Use*, New America (2017) — the deficit-versus-asset mindset and the handling of "at-risk," "low-risk," and "high-risk" labels from early-alert systems. — Retrieved from: \_\_\_\_\_
6. Cf. UK A-Level / Ofqual (Case 49); V. Eubanks, *Automating Inequality* (2018). — Retrieved from: \_\_\_\_\_

## CASE 47 Algorithmic Bias in Automated Exam Proctoring

2022

1. Yoder-Himes, D. R., Asif, A., Kinney, K., Brandt, T. J., Cecil, R. E., Himes, P. R., Cashion, C., Hopp, R. M. P., & Ross, E. (2022). Racial, skin tone, and sex disparities in automated proctoring software. *Frontiers in Education*, 7:881449. doi:10.3389/feduc.2022.881449 — the case's primary study. — Retrieved from: \_\_\_\_\_
2. Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. *Proceedings of Machine Learning Research*, 81:77–91 — the foundational intersectional-bias finding in face recognition. — Retrieved from: \_\_\_\_\_
3. Raji, I. D., & Buolamwini, J. (2019). Actionable auditing: Investigating the impact of publicly naming biased performance results of commercial AI products. *AAAI/ACM Conference on AI, Ethics, and Society* — the audit-and-disclosure mechanism the case calls for. — Retrieved from: \_\_\_\_\_
4. Sjoding, M. W., Dickson, R. P., Iwashyna, T. J., Gay, S. E., & Valley, T. S. (2020). Racial bias in pulse oximetry measurement. *New England Journal of Medicine*, 383(25):2477–2478 — the structural analog in the medical-device context (Case 26). — Retrieved from: \_\_\_\_\_



## CASE 48 **School Surveillance and Black Student Outcomes — Infra-structure as the Mechanism** 2010s – 2022

1. O. Johnson & J. Jabbari (2022), "Infrastructure of social control: A multi-level counterfactual analysis of surveillance and Black education," *Journal of Criminal Justice* 83:101983 — the peer-reviewed source for the case. — Retrieved from: \_\_\_\_\_
2. Hoffman, Trawalter, Axt, & Oliver (2016), "Racial bias in pain assessment and treatment recommendations, and false beliefs about biological differences between blacks and whites," *PNAS* 113(16):4296–4301 — race-construct trio at the cognitive-baseline layer (Case 6). — Retrieved from: \_\_\_\_\_
3. Inker, Eneanya, Coresh, et al. (2021), "New Creatinine- and Cystatin C-Based Equations to Estimate GFR without Race," *NEJM* — race-construct trio at the formula layer (Case 25). — Retrieved from: \_\_\_\_\_
4. Sjoding, Dickson, Iwashyna, Gay, & Valley (2020), "Racial Bias in Pulse Oximetry Measurement," *NEJM* 383:2477–2478 — race-construct trio at the device layer (Case 26). — Retrieved from: \_\_\_\_\_

## CASE 49 **UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days** 2020

1. Ofqual, Awarding GCSE, AS, A level, advanced extension awards and extended project qualifications in summer 2020: interim report, August 2020. — Retrieved from: \_\_\_\_\_
2. Royal Statistical Society, Submission to the Office for Statistics Regulation on the summer 2020 grading process, 2020 — independent statistical review of the standardisation methodology. — Retrieved from: \_\_\_\_\_
3. UK House of Commons Education Committee, Getting the grades they've earned: Covid-19: the cancellation of exams and 'calculated' grades, First Report of Session 2019–21, HC 617, 11 July 2020 — the pre-deployment warning; Office for Statistics Regulation, Ensuring statistical models command public confidence, 2 March 2021 — the post-withdrawal review. — Retrieved from: \_\_\_\_\_
4. Smith, H. (2020), "Algorithmic bias: should students pay the price?" *AI & Society* 35(4):1077–1078 — early academic commentary on the equity dimensions of the withdrawal. — Retrieved from: \_\_\_\_\_

## CASE 50 **Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction** 2012 – 2023

1. Perdomo, J. C., Britton, T., Hardt, M., & Abebe, R. (2025), "Difficult Lessons on Social Prediction from Wisconsin Public Schools," *Proceedings of FAccT 2025*, doi:10.1145/3715275.3732175 (also arXiv:2304.06205). — Retrieved from: \_\_\_\_\_
2. Wisconsin Department of Public Instruction (2021), internal equity analysis of the Dropout Early Warning System — unpublished; the presentation summarizing it, whose slide is headed "Is DEWS Fair?", was obtained by The Markup under public records. — Retrieved from: \_\_\_\_\_
3. Feathers, T. (2023), "False Alarm: How Wisconsin Uses Race and Income to Label Students 'High Risk,'" *The Markup*, April 27, 2023 — investigation documenting the disparate-impact finding and the agency response. — Retrieved from: \_\_\_\_\_
4. Wisconsin Department of Public Instruction, "WISEdash for Districts — Dropout Early Warning System (DEWS) Dashboards" — notice that early-warning dashboard data (DEWS/CCREWS) is no longer available as of 12 October 2023 and that DPI is evaluating the future of the systems. — Retrieved from: \_\_\_\_\_

5. Knowles, J. E. (2015), "Of Needles and Haystacks: Building an Accurate Statewide Dropout Early Warning System in Wisconsin," *Journal of Educational Data Mining* 7(3):18–67 — the original DEWS technical-methodology paper. — Retrieved from: \_\_\_\_\_

## CASE 51 Atlanta Public Schools Cheating Scandal

2009 – 2015

1. Office of the Governor, Special Investigators (Bowers, Wilson, Hyde), Special Investigation into Test Tampering in Atlanta's School System, Vol. 1 (June 30, 2011; released July 2011) — cheating in 44 of 56 schools examined, 178 educators named including 38 principals. The verified quoted span is "emphasized test results and public praise to the exclusion of integrity and ethics"; the subject clause is the reporting's, not the report's — do not re-quote it as the report's own sentence. — Retrieved from: \_\_\_\_\_
2. Atlanta Journal-Constitution investigative series (2008–2011) — the December 2008 analysis of statistically improbable CRCT scores and the subsequent erasure-rate reporting. — Retrieved from: \_\_\_\_\_
3. State of Georgia v. Hall et al. (2013–2015) — indictments and RICO convictions. — Retrieved from: \_\_\_\_\_
4. Campbell, D. (1976), "Assessing the Impact of Planned Social Change" — Campbell's Law. — Retrieved from: \_\_\_\_\_
5. Koretz, D. (2017), The Testing Charade — high-stakes-testing distortion. — Retrieved from: \_\_\_\_\_

## CASE 52 Purdue Course Signals — The Reverse-Causality Retention Claim

2012 – 2013

1. Arnold, K. E., & Pistilli, M. D. (2012). Course Signals at Purdue: Using learning analytics to increase student success. *Proceedings of LAK 2012*, 267–270. doi:10.1145/2330601.2330666 — the original study at the center of the critique. — Retrieved from: \_\_\_\_\_
2. Caulfield, M. (2013). "Why the Course Signals Math Does Not Add Up," Hapgood, September 26, 2013, and "Purdue Course Signals Data Issue Explainer," e-Literate, November 12, 2013 — the reverse-causality critique. — Retrieved from: \_\_\_\_\_
3. Straumsheim, C. (2013). "Mixed Signals," Inside Higher Ed, November 6, 2013 — documents Alfred Essa's placebo simulation, which substituted a randomly distributed piece of chocolate for a Signals class and reproduced the retention curve. — Retrieved from: \_\_\_\_\_
4. Feldstein, M. (2013). "Purdue's Non-Answer on Course Signals," e-Literate, November 2013 — the field-level critique of promoting research and then declining scrutiny. — Retrieved from: \_\_\_\_\_

## CASE 53 inBloom

2014

1. M. Bulger, P. McCormick & M. Pitcan, The Legacy of inBloom, Data & Society Research Institute (2017) — inBloom as a failure of technocratic reform. — Retrieved from: \_\_\_\_\_
2. Education Week and Hechinger Report coverage of the state withdrawals (2013–2014) — nine states exiting. — Retrieved from: \_\_\_\_\_
3. Bulger et al. (2017) — the diagnosis: low public tolerance for risk and uncertainty meeting a failure to communicate benefits and win stakeholder buy-in. — Retrieved from: \_\_\_\_\_
4. N. Selwyn, *Distrusting Educational Technology* (2014); d. boyd & K. Crawford, "Critical Questions for Big Data" (2012). — Retrieved from: \_\_\_\_\_



5. Parent Coalition for Student Privacy archives and the wave of state student-data-privacy legislation that followed inBloom. — Retrieved from: \_\_\_\_\_

## CASE 54 Summit Learning / Personalized Learning Rollout

2014–2019

1. N. Singer, “The Silicon Valley Billionaires Remaking America’s Schools,” New York Times (2017) — the CZI/ Summit rollout. — Retrieved from: \_\_\_\_\_
2. L. Genota, “Brooklyn Students Protest Use of Online Learning Platform Designed by Summit Learning,” Education Week (November 13, 2018) — the 380 schools and 72,000 students figures and the walkout of nearly 100 students at the Secondary School for Journalism. — Retrieved from: \_\_\_\_\_
3. M. Barnum, “Summit Learning declined to be studied, then cited collaboration with Harvard re-searchers anyway,” Chalkbeat (January 17, 2019) — the missing evaluation framework. — Retrieved from: \_\_\_\_\_
4. Chan Zuckerberg Initiative & Summit Learning public program documentation (2015–2019); cf. inBloom (Case 53). — Retrieved from: \_\_\_\_\_
5. M. Barnum, “Summit Learning, the Zuckerberg-backed platform, says 10% of schools quit using it each year. The real figure is higher,” Chalkbeat (May 23, 2019) — the 18 percent annual attrition figure from Summit’s own records. — Retrieved from: \_\_\_\_\_
6. N. Bowles, “Silicon Valley Came to Kansas Schools. That Started a Rebellion,” New York Times (2019) — the parent revolt. — Retrieved from: \_\_\_\_\_
7. S. Schwartz, “Two Districts Roll Back Summit Personalized Learning Program,” Education Week (December 22, 2017) — the Cheshire suspension and the Indiana Area rollback. — Retrieved from: \_\_\_\_\_

## CASE 55 Enrollment-Algorithm Yield Optimiza-tion Across U.S. Higher Education

2010s – present (Brookings synthesis 2021)

1. Engler, A. (2021), “Enrollment algorithms are contributing to the crises of higher education,” Brookings Institution, 14 September 2021 — the two-stage architecture, the vendor tally to at least 700 institutions, and the effect sizes. The phrase “algorithms excel at identifying a student’s exact willingness to pay” is Engler’s own, from the executive summary. — Retrieved from: \_\_\_\_\_
2. Franke, R. (2012), Towards the Education Nation: Revisiting the Impact of Financial Aid, College Experience, and Institutional Context on Baccalaureate Degree Attainment, UCLA doctoral dissertation — the aid-and-graduation effect sizes Engler cites. — Retrieved from: \_\_\_\_\_
3. Goldrick-Rab, S. (2016), Paying the Price — broader synthesis on net-price, unsubsidized loans, and low-income completion. — Retrieved from: \_\_\_\_\_
4. Vendor case studies cited in Engler (Othot, EAB, Ruffalo Noel Levitz, University of Washington) — vendor-reported and not externally audited; flagged at evidence-tier under the title. — Retrieved from: \_\_\_\_\_

## CASE 57 GAO Online Program Manager Oversight Gap (GAO-22-104463)

2022

1. U.S. Government Accountability Office (2022), “Higher Education: Education Needs to Strengthen Its Approach to Monitoring Colleges’ Arrangements with Online Program Managers,” GAO-22-104463. — Retrieved from: \_\_\_\_\_
2. U.S. Department of Education (2011), “Dear Colleague” guidance on incentive-compensation bundled-services exemption — the regulatory artifact the GAO audits. — Retrieved from: \_\_\_\_\_

3. Higher Education Amendments of 1992, Pub. L. No. 102-325, incentive-compensation prohibition, 20 U.S.C. § 1094(a)(20) — the statutory basis the 2011 guidance interpreted. — Retrieved from: \_\_\_\_\_
4. 2U Inc. (2024), Chapter 11 bankruptcy filing (July 25, 2024); U.S. Department of Education, “Dear Colleague” Letter (Jan. 14, 2025) reaffirming the 2011 bundled-services guidance (GEN-11-05) — the commercial-boundary closure alongside the federal guidance’s continued validity. — Retrieved from: \_\_\_\_\_

**CASE 58 USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC)** 2010s – 2024

1. Stephanie Luna v. University of Southern California, class action complaint, Los Angeles County Superior Court, filed May 4, 2023 (First Amended Complaint Sept. 2023; Second Amended May 2, 2024; Third Amended July 30, 2025). — Retrieved from: \_\_\_\_\_
2. Higher Ed Dive reporting on the Luna complaint, May 2023; classaction.org and topclassactions.com summaries. — Retrieved from: \_\_\_\_\_
3. Project on Predatory Student Lending statement on USC-2U partnership termination, November 2023. — Retrieved from: \_\_\_\_\_
4. 2U Inc. and USC public statements on partnership termination, November 2023; broader 2U commercial-collapse reporting referenced through Case 57. — Retrieved from: \_\_\_\_\_

**CASE 60 Houston EVAAS — Value-Added Teacher Evaluation on Trial** 2011–2017

1. American Statistical Association (2014), ASA Statement on Using Value-Added Models for Educational Assessment, April 8, 2014. — Retrieved from: \_\_\_\_\_
2. Amrein-Beardsley, A., & Collins, C. (2012), “The SAS Education Value-Added Assessment System (SAS® EVAAS®) in the Houston Independent School District (HISD): Intended and Unintended Consequences,” Education Policy Analysis Archives 20(12), doi:10.14507/epaa.v20n12.2012. — Retrieved from: \_\_\_\_\_
3. Education Week (October 2017), “Houston District Settles Lawsuit With Teachers’ Union Over Value-Added Scores”; Houston Public Media (October 10, 2017), “Federal Lawsuit Settled Between Houston’s Teacher Union and HISD” — settlement terms and the \$237,000 fee figure. — Retrieved from: \_\_\_\_\_
4. Hous. Fed’n of Teachers v. Hous. Indep. Sch. Dist., 251 F. Supp. 3d 1168 (S.D. Tex. 2017) — memorandum opinion denying summary judgment on the procedural due process claim. — Retrieved from: \_\_\_\_\_
5. Amrein-Beardsley, A., & Geiger, T. (2020), “Methodological Concerns About the Education Value-Added Assessment System (EVAAS): Validity, Reliability, and Bias,” SAGE Open 10(2), doi:10.1177/2158244020922224. — Retrieved from: \_\_\_\_\_
6. American Educational Research Association Council (2015), “AERA Statement on Use of Value-Added Models (VAM) for the Evaluation of Educators and Educator Preparation Programs,” Educational Researcher 44(8):448–452, doi:10.3102/OO13189X15618385. — Retrieved from: \_\_\_\_\_

**CASE 61 Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap** 1997–2023

1. National Reading Panel (2000), Teaching Children to Read: An Evidence-Based Assessment of the Scientific Research Literature on Reading and Its Implications for Reading Instruction, National Institute of Child Health and Human Development, NIH Pub. No. 00-4769. — Retrieved from: \_\_\_\_\_

2. Hanford, E. (2022 – 2024), “Sold a Story: How Teaching Kids to Read Went So Wrong,” APM Reports, podcast series launched October 20, 2022. — Retrieved from: \_\_\_\_\_
3. Hanford, E. (2018), “Hard Words: Why Aren’t Our Kids Being Taught to Read?” APM Reports, September 10, 2018. — Retrieved from: \_\_\_\_\_
4. Schwartz, S. (updated through 2024), “Which States Have Passed ‘Science of Reading’ Laws? What’s in Them?” Education Week — legislation tracker; 40 states plus DC by late 2024. — Retrieved from: \_\_\_\_\_
5. Castles, A., Rastle, K., & Nation, K. (2018), “Ending the Reading Wars: Reading Acquisition From Novice to Expert,” *Psychological Science in the Public Interest* 19(1):5 – 51. — Retrieved from: \_\_\_\_\_
6. Seidenberg, M. (2017), *Language at the Speed of Sight: How We Read, Why So Many Can’t, and What Can Be Done About It*, Basic Books. — Retrieved from: \_\_\_\_\_

## CASE 62 **Gates Intensive Partnerships — Measurement Without a Coupled Lever** 2009–2018

1. Stecher, B. M., Holtzman, D. J., Garet, M. S., Hamilton, L. S., et al. (2018), *Improving Teaching Effectiveness: Final Report: The Intensive Partnerships for Effective Teaching Through 2015–2016*, RAND Corporation, RR-2242–BMGF. — Retrieved from: \_\_\_\_\_
2. Kane, T. J., McCaffrey, D. F., Miller, T., & Staiger, D. O. (2013), *Have We Identified Effective Teachers? Validating Measures of Effective Teaching Using Random Assignment*, MET Project research paper, Bill & Melinda Gates Foundation, January 2013. — Retrieved from: \_\_\_\_\_
3. Sokol, M. (2015), “Sticker shock: How Hillsborough County’s Gates grant became a budget buster,” *Tampa Bay Times*, September 21, 2015 — with companion reporting, “Hillsborough schools to dismantle Gates-funded system that cost millions to develop” (fall 2015). — Retrieved from: \_\_\_\_\_
4. Stecher, B. M., Garet, M. S., Hamilton, L. S., Steiner, E. D., Robyn, A., Poirier, J., Holtzman, D. J., Fulbeck, E. S., Chambers, J., & Brodzia de los Reyes, I. (2016), *Improving Teaching Effectiveness: Implementation: The Intensive Partnerships for Effective Teaching Through 2013–2014*, RAND Corporation, RR-1295. — Retrieved from: \_\_\_\_\_
5. Will, M. (2018), “‘An Expensive Experiment’: Gates Teacher-Effectiveness Program Shows No Gains for Students,” *Education Week*, June 21, 2018. — Retrieved from: \_\_\_\_\_

## CASE 63 **LAUSD iPad Program — The Common Core Technology Project** 2013–2015

1. Gilbertson, A. (2014), “LA schools cancel iPad contracts after KPCC publishes internal emails,” KPCC 89.3 (Southern California Public Radio), August 25, 2014 — the public-records email investigation and the contract halt; see also KPCC’s “LAUSD iPads: Timeline of a troubled program.” — Retrieved from: \_\_\_\_\_
2. Margolin, J., Haynes, E., Heppen, J., Ruedel, K., Meakin, J., Hauser, A., Blum, J., Chavez, S., & Hubbard, A. (2014), *Evaluation of the Common Core Technology Project: Interim Report*, American Institutes for Research, September 2014 — 245 classroom observations across 19 schools (15 Phase 1 CCTP, four non-CCTP) visited in May 2014; the Pearson curriculum app observed in use in one, at a single elementary school. Followed by AIR’s *Evaluation of LAUSD’s Instructional Technology Initiative: Year 2 Report* (August 2015). — Retrieved from: \_\_\_\_\_
3. Herold, B. (2014), “Hard Lessons Learned in Ambitious L.A. iPad Initiative,” *Education Week*, September 10, 2014 — retrospective on the implementation failures; see also Herold (2014), “FBI Investigation Leaves L.A. iPad Initiative in Further Disarray,” *Education Week*, December 2014, and Herold (2015), “Pearson Settles With L.A. Over Curriculum Foul-Up,” *Education Week*, October 20, 2015 — the \$6.45 million settlement, paid by Pearson. — Retrieved from: \_\_\_\_\_

4. Blume, H., Los Angeles Times coverage of the CCTP (2013–2015) — the rollout, the security bypass, the contract suspension, the Deasy resignation, and the FBI document seizure. — Retrieved from: \_\_\_\_\_
5. Education Week Market Brief (2017), “Feds Drop Investigation Into Los Angeles District Over \$1 Billion iPad Purchase,” February 2017 — the reported closure of the federal inquiry without charges. — Retrieved from: \_\_\_\_\_

## CASE 65 Training Transfer — The Gap Between Learning and Doing

2010

1. Blume, Ford, Baldwin, & Huang (2010), “Transfer of Training: A Meta-Analytic Review,” *Journal of Management* 36(4):1065–1105, doi:10.1177/0149206309352880. — Retrieved from: \_\_\_\_\_
2. Baldwin & Ford (1988), “Transfer of Training: A Review and Directions for Future Research,” *Personnel Psychology* 41(1):63–105 — the foundational synthesis the 2010 meta-analysis updates. — Retrieved from: \_\_\_\_\_
3. Burke & Hutchins (2007), “Training Transfer: An Integrative Literature Review,” *Human Resource Development Review* 6(3):263–296 — the integrative-review companion synthesis. — Retrieved from: \_\_\_\_\_
4. Kirkpatrick & Kirkpatrick (2006), *Evaluating Training Programs* — the framework the meta-analysis informs (paired Case 79). — Retrieved from: \_\_\_\_\_

## CASE 66 Growth-Mindset National Experiment — Heterogeneous Effects

2019

1. Yeager, D. S., Hanselman, P., Walton, G. M., Murray, J. S., Crosnoe, R., Muller, C., Tipton, E., Schneider, B., Hulleman, C. S., Hinojosa, C. P., Paunesku, D., Romero, C., Flint, K., Roberts, A., Trott, J., Iachan, R., Buontempo, J., Yang, S. M., Carvalho, C. M., Hahn, P. R., Gopalan, M., Mhatre, P., Ferguson, R., Duckworth, A. L., & Dweck, C. S. (2019). A national experiment reveals where a growth mindset improves achievement. *Nature*, 573(7774):364–369. doi:10.1038/s41586-019-1466-y — the case’s primary trial. — Retrieved from: \_\_\_\_\_
2. National Study of Learning Mindsets, ICPSR 37353 — the trial dataset. — Retrieved from: \_\_\_\_\_
3. Dweck, C. S. (2006). *Mindset: The New Psychology of Success*. Random House — the broader theoretical framework the intervention rests on. — Retrieved from: \_\_\_\_\_
4. Sisk, V. F., Burgoyne, A. P., Sun, J., Butler, J. L., & Macnamara, B. N. (2018). To what extent and under which circumstances are growth mindsets important to academic achievement? Two meta-analyses. *Psychological Science*, 29(4):549–571 — the prior moderator-analysis literature the Yeager trial extends. — Retrieved from: \_\_\_\_\_
5. Macnamara, B. N., & Burgoyne, A. P. (2023). Do growth mindset interventions impact students’ academic achievement? A systematic review and meta-analysis. *Psychological Bulletin* 149(3–4):133–173 — near-null effects among the best-designed studies. — Retrieved from: \_\_\_\_\_
6. Education Endowment Foundation (Foliano, F., Rolfe, H., Buzzeo, J., Runge, J., & Wilkinson, D., 2019), *Changing Mindsets: Effectiveness Trial — Evaluation Report* — a 101-school, 5,018-pupil RCT with Year 6 pupils; point estimates of about two additional months’ progress in English and maths, neither statistically significant. — Retrieved from: \_\_\_\_\_

## CASE 67 Cognitive Tutor / Carnegie Learning

1990s – present

1. Anderson, J., Corbett, A., Koedinger, K. & Pelletier, R. (1995), “Cognitive Tutors: Lessons Learned,” *Journal of the Learning Sciences* — the ACT-R basis and design. — Retrieved from: \_\_\_\_\_

2. Koedinger, K. & Corbett, A. (2006), *Cambridge Handbook of the Learning Sciences — knowledge tracing and adaptive instruction*. — Retrieved from: \_\_\_\_\_
3. Pane, J. F., Griffin, B. A., McCaffrey, D. F. & Karam, R. (2014), "Effectiveness of Cognitive Tutor Algebra I at Scale," *Educational Evaluation and Policy Analysis* 36(2):127–144, doi:10.3102/0162373713507480 — no effect in year one; a statistically significant year-two effect for high schools, not for middle schools. — Retrieved from: \_\_\_\_\_
4. What Works Clearinghouse (2009), *Cognitive Tutor® Algebra I intervention report*, U.S. Department of Education, July 2009 — "more than 500,000 students in approximately 2,600 urban, rural, and suburban school districts" as of August 2008. — Retrieved from: \_\_\_\_\_
5. Lynch, C., Ashley, K., Pinkwart, N. & Aleven, V. (2009), "Concepts, Structures, and Goals: Redefining Ill-Defined-ness," *International Journal of Artificial Intelligence in Education* 19(3):253–266 — what makes a domain ill-defined and why the tutor pipeline has less to attach to there. — Retrieved from: \_\_\_\_\_
6. Ritter, S. & Fancsali, S. (2016), "MATHia X: The Next Generation Cognitive Tutor," *Proceedings of the 9th International Conference on Educational Data Mining (EDM 2016)*, 624–625 — the web re-platforming of the Cognitive Tutor software into MATHia; and Carnegie Learning's 2023 announcement of LiveHint AI (cf. Case 88). — Retrieved from: \_\_\_\_\_

**CASE 68**     **CIRCUIT — A Scalable, Equity-Centered Research Workforce Model**     2017 – 2023 (six cycles)

1. National Academies / NSF research on undergraduate research experience effects on STEM persistence (broader literature against which the case sits). — Retrieved from: \_\_\_\_\_
2. Cervantes, Floryanzia, Sharp, Gray-Roncal, & Johnson (2023), "Empowering Trailblazers toward Scalable, Systematized, Research-Based Workforce Development," *ASEE Annual Conference*, doi:10.18260/1-2-43271. — Retrieved from: \_\_\_\_\_
3. CIRCUIT program documentation (JHU Hub, 2017 – present) — institutional program description. — Retrieved from: \_\_\_\_\_
4. Paired case (78) — CIRCUIT proofreading + MICrONS — the deployed-capability companion. — Retrieved from: \_\_\_\_\_

**CASE 69**     **Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale**     2016

1. Ebbinghaus, H. (1885), *Über das Gedächtnis* — the empirical forgetting-curve foundation HLR formalizes. — Retrieved from: \_\_\_\_\_
2. Settles, B., & Meeder, B. (2016), "A Trainable Spaced Repetition Model for Language Learning," *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics*, pp. 1848–1858, doi:10.18653/v1/P16-1174. — Retrieved from: \_\_\_\_\_
3. Pimsleur, P. (1967), "A memory schedule," *Modern Language Journal* 51(2):73–75 — graduated-interval recall heuristic. — Retrieved from: \_\_\_\_\_
4. Leitner, S. (1972), *So lernt man lernen* — Leitner-box spacing heuristic. — Retrieved from: \_\_\_\_\_

## CASE 70 **High-Impact Learning System — Engineering the Environment, Not Just the Event** 2001 – present

1. Blume et al. (2010), *Journal of Management* 36(4):1065–1105 — the meta-analytic finding HILS operationalizes (paired Case 65). — Retrieved from: \_\_\_\_\_
2. Brinkerhoff, R. O., & Apking, A. M. (2001), *High Impact Learning: Strategies for Leveraging Performance and Business Results from Training Investments*, Perseus Publishing. — Retrieved from: \_\_\_\_\_
3. Brinkerhoff, R. O. (2006), “Increasing impact of training investments: an evaluation strategy for building organizational learning capability,” *Industrial and Commercial Training* 38(6):302–307, doi:10.1108/00197850610685824 — the author’s own statement that evaluation should address the whole performance-improvement process, not the training event. — Retrieved from: \_\_\_\_\_
4. Brinkerhoff, R. O. (2005), “The Success Case Method: A Strategic Evaluation Approach to Increasing the Value and Effect of Training,” *Advances in Developing Human Resources* 7(1):86–101 — SCM as the paired evaluation instrument (Case 83). — Retrieved from: \_\_\_\_\_

## CASE 71 **Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development** 2025

1. v2\_research/01\_LENS\_revised\_competencies\_and\_CLOs.md — the proposed LENS revisions, including the amended LEO-2 (first-person narration of design iteration) whose evaluation problem this case addresses. — Retrieved from: \_\_\_\_\_
2. Barr, M., Nabi, S. W., & Andrei, O. (2025), “The Development of Reflective Practice on a Work-Based Software Engineering Program: A Longitudinal Study,” *Proc. 37th IEEE Conference on Software Engineering Education and Training (CSEET)*, pp. 53–61, doi:10.1109/CSEET66350.2025.00012 (preprint: arXiv:2504.20956). — Retrieved from: \_\_\_\_\_
3. Boud, Keogh & Walker (eds.), *Reflection: Turning Experience into Learning* (1985) — reflection as a learning process, and the measurement problem it raises. — Retrieved from: \_\_\_\_\_
4. D. Schön, *The Reflective Practitioner* (1983) — the foundational account of reflection-in-action the genre rests on. — Retrieved from: \_\_\_\_\_

## CASE 72 **ASSISTments — National Replication and Long-Term Follow-Through** 2014 – present

1. Heffernan, N. T., & Heffernan, C. L. (2014), “The ASSISTments ecosystem,” *International Journal of AI in Education* 24:470–497 — platform design and research-loop description. — Retrieved from: \_\_\_\_\_
2. Roschelle, J., Feng, M., Murphy, R. F., & Mason, C. A. (2016), “Online Mathematics Homework Increases Student Achievement,” *AERA Open* 2(4):1–12, doi:10.1177/2332858416673968. — Retrieved from: \_\_\_\_\_
3. Feng, M., Huang, C., & Collins, K. (2023), *Technology-Based Support Shows Promising Long-Term Impact on Math Learning*, WestEd — the North Carolina replication, the equity heterogeneity finding, and the 8th-grade persistence result. — Retrieved from: \_\_\_\_\_
4. Arnold Ventures grant documentation of the virtual-training-adaptation arm — including the 8th-grade extension that WestEd and Arnold Ventures ended when the design departed from the planned RCT. — Retrieved from: \_\_\_\_\_

## CASE 73 The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension 2016 – 2025

1. Van Campenhout, R., Jerome, B., Dittel, J. S., & Johnson, B. G. (2023), "The Doer Effect at Scale: Investigating Correlation and Causation Across Seven Courses," LAK23, pp. 357–365, doi:10.1145/3576050.3576103. — Retrieved from: \_\_\_\_\_
2. Van Campenhout, R., Autry, K. S., Clark, M. W., & Johnson, B. G. (2025), "Scaling the Doer Effect: A Replication Analysis Using AI-Generated Questions," L@S '25, pp. 24–34, doi:10.1145/3698205.3729545 — two sections of a Cognitive Psychology course, following the established correlational doer-effect analysis. — Retrieved from: \_\_\_\_\_
3. Butler, D. et al. (2025), "Does the Doer Effect Generalize To Non-WEIRD Populations? Toward Analytics in Radio and Phone-Based Learning," LAK '25, doi:10.1145/3706468.3706505 (also arXiv 2412.20923). — Retrieved from: \_\_\_\_\_
4. Koedinger, K. R., McLaughlin, E. A., Jia, J. Z., & Bier, N. L. (2016), "Is the doer effect a causal relationship? How can we tell and why it's important," LAK '16, pp. 388–397, doi:10.1145/2883851.2883957 — the replication target the present case builds on. — Retrieved from: \_\_\_\_\_

## CASE 74 Zhang/Scardamalia — Knowledge Building Across Three Cohorts 2009

1. Zhang, J., Scardamalia, M., Reeve, R., & Messina, R. (2009), "Designs for Collective Cognitive Responsibility in Knowledge-Building Communities," Journal of the Learning Sciences 18(1):7–44, doi:10.1080/10508400802581676. — Retrieved from: \_\_\_\_\_
2. Scardamalia, M., & Bereiter, C. (2006), "Knowledge Building: Theory, pedagogy, and technology," in K. Sawyer (ed.), Cambridge Handbook of the Learning Sciences — programmatic backdrop. — Retrieved from: \_\_\_\_\_
3. Bereiter, C., & Scardamalia, M. (2014), "Knowledge Building and Knowledge Creation: One Concept, Two Hills to Climb," in S. C. Tan, H. J. So, & J. Yeo (eds.), Knowledge Creation in Education — extension into the post-2009 program. — Retrieved from: \_\_\_\_\_
4. Design-based research methodology references — Cobb, Confrey, diSessa, Lehrer, & Schauble (2003), "Design Experiments in Educational Research," Educational Researcher 32(1):9–13 — the methodological frame the 2009 study operates inside. — Retrieved from: \_\_\_\_\_

## CASE 75 Chen et al. — Rural China AI Devices and the Equity-Direction Finding 2025

1. Chen, R., Wu, Y., Chen, Z., & Zhou, P. (2025), "Advancing educational equity in rural China: the impact of AI devices on teaching quality and learning outcomes for sustainable development," Frontiers in Psychology 16:1588047, doi:10.3389/fpsyg.2025.1588047. — Retrieved from: \_\_\_\_\_
2. Paired Case 86 (Gándara et al., AERA Open) — bias in student-success prediction at the higher-education scale. — Retrieved from: \_\_\_\_\_
3. Paired Case 88 (LiveHint AI bias across foundation models, AIED 2025) — bias-surfacing in AI-supported instruction. — Retrieved from: \_\_\_\_\_
4. Paired Case 73 (Doer Effect non-WEIRD LAK 2025 radio-and-phone extension) — non-WEIRD methodological discipline at the heterogeneity-finding axis. — Retrieved from: \_\_\_\_\_



## CASE 76 **Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account** 2023

1. Blikstein, P., & Worsley, M. (2016), "Multimodal Learning Analytics and Education Data Mining: Using computational technologies to measure complex learning tasks," *Journal of Learning Analytics* 3(2):220–238 — broader MMLA literature backdrop. — Retrieved from: \_\_\_\_\_
2. Martinez-Maldonado, R., Echeverria, V., Fernandez-Nieto, G., et al. (2023), "Lessons Learnt from a Multimodal Learning Analytics Deployment In-the-wild," *ACM Transactions on Computer-Human Interaction* 31(1), Article 8, pp. 1–41 (November 2023), doi:10.1145/3622784; author preprint at arXiv:2303.09099. — Retrieved from: \_\_\_\_\_
3. v2\_research/03\_editor\_request\_memo.md, item B1 — the Practice Flywheel commissioning structure this case is offered as a published-first-person exemplar within. — Retrieved from: \_\_\_\_\_
4. Schön, D. (1983), *The Reflective Practitioner* — the genre's theoretical underpinning, referenced across the reflective-practice case tier. — Retrieved from: \_\_\_\_\_

## CASE 77 **Hybrid Human-AI Tutoring — Augmentation, Not Delegation** 2024

1. Thomas, D. R. et al. (2024), "Improving Student Learning with Hybrid Human-AI Tutoring: A Three-Study Quasi-Experimental Investigation," *LAK '24*, doi:10.1145/3636555.3636896. — Retrieved from: \_\_\_\_\_
2. Koedinger, K. R. et al. — Cognitive Tutor literature as the fully-automated track the augmentation track contrasts with. — Retrieved from: \_\_\_\_\_
3. Case 20 (TREWS) reference set — Henry et al. (2022), *Nature Medicine* — clinician-AI teaming analog. — Retrieved from: \_\_\_\_\_
4. Case 5 (Epic Sepsis) reference set — Wong et al. (2021), *JAMA Internal Medicine* — delegation-collapse analog. — Retrieved from: \_\_\_\_\_

## CASE 78 **CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale** 2017 – present

1. MICrONS program literature (IARPA) — connectomics method and automated segmentation evidence base. — Retrieved from: \_\_\_\_\_
2. APL BossDB documentation — petabyte-scale connectomics data infrastructure. — Retrieved from: \_\_\_\_\_
3. Cervantes, Floryanzia, Sharp, Gray-Roncal, & Johnson (2023), "Empowering Trailblazers toward Scalable, Systematized, Research-Based Workforce Development," *ASEE Annual Conference*, doi:10.18260/1-2-43271 — the paired peer-reviewed CIRCUIT case (Case 68). — Retrieved from: \_\_\_\_\_
4. CIRCUIT program documentation (JHU Hub, 2017 – present) — institutional/program description; the training-outcome evidence is at this tier and the evidence-tier flag is binding. — Retrieved from: \_\_\_\_\_
5. Wachter, R. M., & Brynjolfsson, E. (2024), "Will Generative Artificial Intelligence Deliver on Its Promise in Health Care?" *JAMA* 331(1):65–69 — the productivity-paradox framing: general purpose technologies pay off only after the complementary organizational work is done. — Retrieved from: \_\_\_\_\_



## CASE 79 **The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops** 1959 – present

1. Kirkpatrick (1959–1960), original four-level evaluation series in the *Journal of the American Society of Training Directors* (Nov 1959 – Feb 1960); updated as Kirkpatrick & Kirkpatrick, *Evaluating Training Programs* (3rd ed., 2006). — Retrieved from: \_\_\_\_\_
2. Devlin Peck, “The Kirkpatrick Model: 4 Levels, Examples, and Free Template” — synthesis of the stop-at-L2 pattern in corporate practice. — Retrieved from: \_\_\_\_\_
3. D2L, “Evaluating Training Programs: Models, Metrics and Business Impact,” practitioner guide documenting the same pattern. — Retrieved from: \_\_\_\_\_
4. Valamis, “Kirkpatrick model: how to evaluate whether training worked” — practitioner guide on evaluation-practice gaps. — Retrieved from: \_\_\_\_\_
5. Association for Talent Development (2016), *Evaluating Learning: Getting to Measurements That Matter* — survey of 199 talent-development professionals; source of the level-by-level figures ( 90% L1, 83% L2, 60% L3, 35% L4, 15% ROI). Association research, not peer-reviewed. — Retrieved from: \_\_\_\_\_
6. Blume, Ford, Baldwin, & Huang (2010), “Transfer of Training: A Meta-Analytic Review,” *Journal of Management* 36(4):1065–1105 — the paired peer-reviewed case (65); its opening \$125B figure is ASTD’s, via Paradise (2007). — Retrieved from: \_\_\_\_\_

## CASE 80 **Georgia State University Predictive Analytics** 2012 – present

1. Georgia State University institutional research and Strategic Plan reports — graduation-rate and equity data. — Retrieved from: \_\_\_\_\_
2. Fausset, R., “Georgia State, Leading U.S. in Black Graduates, Is Engine of Social Mobility,” *New York Times*, 15 May 2018 — the 800-factor advising model. — Retrieved from: \_\_\_\_\_
3. EDUCAUSE Review on GSU predictive advising — the human-loop architecture. — Retrieved from: \_\_\_\_\_
4. Renick, T. & Strom, A. (2020) on GSU’s advising transformation — the system design and outcomes. — Retrieved from: \_\_\_\_\_
5. Complete College America, *Game Changers* documentation — dissemination of the model. — Retrieved from: \_\_\_\_\_
6. Georgia State University National Institute for Student Success (founded 2021) — reported engagement with 130-plus campuses and 1.5 million students by 2024. — Retrieved from: \_\_\_\_\_

## CASE 81 **Open University 'Ethical Use of Student Data' and OU Analyse** 2014 – 2025

1. Slade & Prinsloo (2013), “Learning Analytics: Ethical Issues and Dilemmas,” *American Behavioral Scientist* 57(10):1510–1529, doi:10.1177/0002764213479366. — Retrieved from: \_\_\_\_\_
2. Open University (2014), “Ethical Use of Student Data for Learning Analytics” — first institutional policy of its kind in higher education. — Retrieved from: \_\_\_\_\_
3. Herodotou, Hlostá, Boroowa, Rienties, Zdrahal, & Mangafa (2019), “Empowering online teachers through predictive learning analytics,” *British Journal of Educational Technology* 50(6):3064–3079, doi:10.1111/bjet.12853 — OU Analyse evaluation across 559 teachers (189 with OUA access) and 14,000+ students in 15 undergraduate courses; average-use teachers benefited students most. — Retrieved from: \_\_\_\_\_

4. Rienties, Boroowa, Cross, Kubiak, Mayles, & Murphy (2016), "Analytics4Action Evaluation Framework: A Review of Evidence-Based Learning Analytics Interventions at the Open University UK," *Journal of Interactive Media in Education* 2016(1):2, doi:10.5334/jime.394. — Retrieved from: \_\_\_\_\_

## CASE 82 **MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil)** 2024

1. Freitas, Fonseca, Garcia, Pontual Falcão, Marques, Gašević, & Ferreira Mello (2024), "MMALA: Developing and Evaluating a Maturity Model for Adopting Learning Analytics," *Journal of Learning Analytics* 11(1):67–86, doi:10.18608/jla.2024.8099. The first version had 18 process areas across the same five categories; consolidation plus one added area ("result-based intervention") produced the 16 of record. — Retrieved from: \_\_\_\_\_
2. Open University Ethical Use of Student Data policy (2014) — institutional-policy companion (Case 81). — Retrieved from: \_\_\_\_\_
3. Hilliger et al. (2020), *Internet and Higher Education* — LALA participatory adoption companion (Case 91). — Retrieved from: \_\_\_\_\_
4. Norwegian Expert Commission on Learning Analytics, final NOU (2023) — national-scale companion (Case 92). — Retrieved from: \_\_\_\_\_

## CASE 83 **Brinkerhoff Success Case Method — Tails as the Evaluation Instrument** 2003 – present

1. Kirkpatrick & Kirkpatrick (2006), *Evaluating Training Programs* — the chain-of-evidence framework SCM operationalizes (paired Case 79). — Retrieved from: \_\_\_\_\_
2. Brinkerhoff, R. O. (2005), "The Success Case Method: A Strategic Evaluation Approach to Increasing the Value and Effect of Training," *Advances in Developing Human Resources* 7(1):86–101, doi:10.1177/1523422304272172 — the peer-reviewed statement of the method; the book-length treatment is Brinkerhoff (2003), *The Success Case Method: Find Out Quickly What's Working and What's Not* (Berrett-Koehler). — Retrieved from: \_\_\_\_\_
3. Blume, Ford, Baldwin, & Huang (2010), *Journal of Management* 36(4):1065–1105 — the meta-analytic environment-as-finding SCM samples around (paired Case 65). — Retrieved from: \_\_\_\_\_
4. Brinkerhoff Evaluation Institute deployment list — International Red Cross and Red Crescent, World Bank, Cargill, Merck, Ford Motor — practitioner channel. — Retrieved from: \_\_\_\_\_

## CASE 84 **Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive** 2007 – 2014

1. Evidence for ESSA, program record for Carnegie Learning High School Math Solution Algebra I — ESSA rating "Strong," weighted mean effect size +0.04 across three qualifying studies (Cabalo & Vu 2007; Campuzano et al. 2009; Pane et al. 2014). — Retrieved from: \_\_\_\_\_
2. Pane, J. F., Griffin, B. A., McCaffrey, D. F., & Karam, R. (2014), "Effectiveness of Cognitive Tutor Algebra I at Scale," *Educational Evaluation and Policy Analysis* 36(2):127–144, doi:10.3102/O162373713507480. — Retrieved from: \_\_\_\_\_
3. Pane, Griffin, McCaffrey, & Karam (2014), Addendum to Effectiveness of Cognitive Tutor Algebra I at Scale, RAND Working Paper WR-1050-DEIES — the addendum the WWC reviewed alongside the published paper. — Retrieved from: \_\_\_\_\_

4. Koedinger, K. R., Anderson, J. R., Hadley, W. H., & Mark, M. A. (1997), "Intelligent tutoring goes to school in the big city," IJAIED — the Case 67 system description Cognitive Tutor builds from. — Retrieved from: \_\_\_\_\_
5. What Works Clearinghouse (2016), Cognitive Tutor® intervention report, Secondary Mathematics, June 2016 — rates Cognitive Tutor Algebra I "mixed effects" on algebra; records the Pane et al. year-two difference as no longer significant after adjustment for multiple comparisons. — Retrieved from: \_\_\_\_\_
6. Escueta, Quan, Nickow, & Oreopoulos (2017), "Education Technology: An Evidence-Based Review," NBER Working Paper 23744 — reports that the Pane year-two gain tracked reduced use of the prescribed non-computer activities, not higher fidelity; published as Escueta, Nickow, Oreopoulos, & Quan (2020), "Upgrading Education with Technology: Insights from Experimental Research," Journal of Economic Literature 58(4):897–996. — Retrieved from: \_\_\_\_\_
7. Kraft, M. A. (2020), "Interpreting Effect Sizes of Education Interventions," Educational Researcher 49(4):241–253 — 1,942 effect sizes from 747 RCTs, median 0.10 SD, 0.03 SD for samples above 2,000 students, and the benchmark that treats 0.20 SD as large. — Retrieved from: \_\_\_\_\_

## CASE 85 **OU Analyse — Predictive Learning Analytics and Teacher Use at Scale** 2019 – 2023

1. Herodotou, C., Hlosta, M., Boroowa, A., Rienties, B., Zdrahal, Z., & Mangafa, C. (2019), "Empowering online teachers through predictive learning analytics," British Journal of Educational Technology 50(6):3064–3079, doi:10.1111/bjjet.12853. — Retrieved from: \_\_\_\_\_
2. Herodotou, C., Maguire, C., Hlosta, M., & Mulholland, P. (2023), "Predictive Learning Analytics and University Teachers: Usage and perceptions three years post implementation," LAK '23, pp. 68–78, doi:10.1145/3576050.3576061 — survey of 366 teachers across four faculties; 42% active users, 20% tried and stopped, 18% aware but unused, 19% unaware; UTAUT model of adoption. — Retrieved from: \_\_\_\_\_
3. Herodotou, Rienties, Boroowa, Zdrahal, & Hlosta (2019), "A large-scale implementation of predictive learning analytics in higher education: the teachers' role and perspective," Educational Technology Research and Development 67(5):1273–1306, doi:10.1007/s11423-019-09685-0 — the complementary teacher-perspective study (59 teachers, nine courses, 1,325 students). — Retrieved from: \_\_\_\_\_
4. Arnold, K. E., & Pistilli, M. D. (2012), "Course Signals at Purdue: Using Learning Analytics to Increase Student Success," LAK '12 — the structural precursor the corpus carries as a cautionary case at Case 52. — Retrieved from: \_\_\_\_\_

## CASE 86 **Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction** 2024

1. Gándara, Anahideh, Ison, & Picchiarini (2024), "Inside the Black Box: Detecting and Mitigating Algorithmic Bias across Racialized Groups in College Student-Success Prediction," AERA Open 10, doi:10.1177/23328584241258741 — primary case source on cross-group accuracy disparity, stratified evaluation, and bias mitigation in student-success prediction. — Retrieved from: \_\_\_\_\_
2. Barocas, Hardt, & Narayanan, Fairness and Machine Learning (fairmlbook.org) — methodological backdrop on construct definition and stratified evaluation. — Retrieved from: \_\_\_\_\_
3. Friedler, Scheidegger, & Venkatasubramanian (2021), "The (im)possibility of fairness: different value systems require different mechanisms for fair decision making," Communications of the ACM — the construct-definition argument at field level. — Retrieved from: \_\_\_\_\_

4. v2 cross-referenced cases: 25 (eGFR race correction), 26 (pulse oximetry across skin tones), 6 (Hoffman pain bias) — the race-construct trio at the construct-definition layer. — Retrieved from: \_\_\_\_\_

## CASE 87 **Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models** 2021 – 2022

1. Yu, R., Lee, H., & Kizilcec, R. F. (2021), "Should College Dropout Prediction Models Include Protected Attributes?" in Proceedings of the Eighth ACM Conference on Learning @ Scale (L@S '21), doi:10.1145/3430895.3460139 — primary paper on the include-or-exclude empirical analysis. — Retrieved from: \_\_\_\_\_
2. Kizilcec, R. F., & Lee, H. (2022), "Algorithmic Fairness in Education," in Holmes & Porayska-Pomsta (eds.), The Ethics of Artificial Intelligence in Education, Routledge, doi:10.4324/9780429329067-10 — broader synthesis of the fairness-in-learning-analytics frontier. — Retrieved from: \_\_\_\_\_
3. Dwork, Hardt, Pitassi, Reingold, & Zemel (2012), "Fairness through awareness," Proceedings of ITCS — foundational paper on the inadequacy of fairness-through-unawareness. — Retrieved from: \_\_\_\_\_
4. Barocas, Hardt, & Narayanan, Fairness and Machine Learning (fairmlbook.org) — technical-fairness backdrop. — Retrieved from: \_\_\_\_\_
5. v2 cross-referenced cases: 86 (Gándara), 25 (eGFR), 26 (pulse oximetry), 6 (Hoffman) — equity-construct five-case set. — Retrieved from: \_\_\_\_\_

## CASE 88 **LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models** 2025

1. Bommasani, R. et al. (2021), "On the Opportunities and Risks of Foundation Models," Stanford CRFM — the foundation-model framing the case builds on. — Retrieved from: \_\_\_\_\_
2. Vinodh, A., Harvey, E., Almoubayyed, H., Yu, R., Brooks, C., Koenecke, A., & Kizilcec, R. F. (2025), "Evaluating an AI Tutor for Bias Across Different Foundation Models," AIED 2025, LNAI 15882, pp. 341–348, doi:10.1007/978-3-031-98465-5\_43; author copy at renzheyu.com/papers/AIED2025\_Tutor.pdf. — Retrieved from: \_\_\_\_\_
3. Race-construct trio reference set: Hoffman et al. (2016), Sjoding et al. (2020) pulse oximetry, Inker et al. (2021) eGFR-without-race — paired with Cases 25, 26, 6. — Retrieved from: \_\_\_\_\_
4. Carnegie Learning LiveHint product documentation — the subject system; case framing is binding on LiveHint being in development, not deployment. — Retrieved from: \_\_\_\_\_

## CASE 89 **Data Privacy and Learning Analytics on the African Continent** 2022

1. Prinsloo, P., & Kaliisa, R. (2022), "Data privacy on the African continent: Opportunities, challenges and implications for learning analytics," British Journal of Educational Technology 53(4):894–913, doi:10.1111/bjet.13226. — Retrieved from: \_\_\_\_\_
2. African Union Convention on Cyber Security and Personal Data Protection (Malabo Convention, adopted 2014; in force 8 June 2023) — continental policy backdrop, ratified by a minority of AU member states. — Retrieved from: \_\_\_\_\_
3. Slade & Prinsloo (2013), American Behavioral Scientist — earlier framing on which the 2022 paper builds. — Retrieved from: \_\_\_\_\_
4. Open University Ethical Use of Student Data policy (2014) — single-regime governance-by-design companion. — Retrieved from: \_\_\_\_\_

## CASE 90 **Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions** 2025

1. Pyle, C., & Andalibi, N. (2025), "Algorithmic College Admissions in the U.S.: Distances Between Vendors' Claims and Applicants' Perceptions," *Proceedings of the ACM on Human-Computer Interaction* 9(7), CSCW369, doi:10.1145/3757550. — Retrieved from: \_\_\_\_\_
2. Engler (2021), Brookings — paired deployment-side mapping (Case 55). — Retrieved from: \_\_\_\_\_
3. GAO-22-104463 (2022) — paired regulator-side audit (Case 57). — Retrieved from: \_\_\_\_\_
4. Dunne, A., & Raby, F. (2013), *Speculative Everything* — methodological backdrop for speculative-design probes. — Retrieved from: \_\_\_\_\_

## CASE 91 **LALA — Building Learning-Analytics Governance Capacity Across Latin America** 2017 – 2020

1. Hilliger, Ortiz-Rojas, Pesántez-Cabrera, Scheihing, Tsai, Muñoz-Merino, Broos, Whitelock-Wainwright, & Pérez-Sanagustín (2020), "Identifying needs for learning analytics adoption in Latin American universities: A mixed-methods approach," *The Internet and Higher Education* 45:100726, doi:10.1016/j.iheduc.2020.100726. — Retrieved from: \_\_\_\_\_
2. LALA project — EU Erasmus+ grant 586120-EPP-1-2017-1-ES (2017–2020) — program documentation and the LALA CANVAS artifact. — Retrieved from: \_\_\_\_\_
3. Open University Ethical Use of Student Data policy (2014) and OU Analyse — UK companion governance-by-design case (Case 81). — Retrieved from: \_\_\_\_\_
4. Slade & Prinsloo (2013), "Learning Analytics: Ethical Issues and Dilemmas," *American Behavioral Scientist* 57(10):1510–1529 — the broader field-scale ethics framing. — Retrieved from: \_\_\_\_\_

## CASE 92 **Norway's National Expert Commission on Learning Analytics** 2021 – 2023

1. Ekspertgruppen for digital læringsanalyse (2022), *Læringsanalyse* — noen sentrale dilemmaer, interim report delivered to the Minister of Education and Research, June 2022. — Retrieved from: \_\_\_\_\_
2. NOU 2023: 19, *Læring, hvor ble det av deg i alt mylderet? Bruk av elev- og studentdata for å fremme læring* (June 2023) — the commission's final report. — Retrieved from: \_\_\_\_\_
3. Wasson, Giannakos, Blikstad-Balas, Uppstad, Langford & Bøhn (2024), "Implementing Learning Analytics in Norway: Four Central Dilemmas," *Journal of Learning Analytics* 11(2):268–280, doi:10.18608/jla.2024.8241 — written by commission members; the interim report and its four dilemmas (data, learning, governance, competence). — Retrieved from: \_\_\_\_\_
4. Hilliger et al. (2020), *Internet and Higher Education* — the LALA companion at multi-country participatory scale (Case 91). — Retrieved from: \_\_\_\_\_

## CASE 93 **Singapore SkillsFuture — National Workforce Capability at Scale** 2015 – present

1. SkillsFuture Singapore (SSG), *Year-in-Review 2024* — program metrics and outcome reporting. — Retrieved from: \_\_\_\_\_
2. Ministry of Education (MOE), Singapore, parliamentary replies on TRAQOM, 2020. — Retrieved from: \_\_\_\_\_

3. Ministry of Trade and Industry (MTI), Singapore, 2018 — WSQ wage-premium study. — Retrieved from: \_\_\_\_\_
4. “Future-Skilling the Workforce: SkillsFuture Movement in Singapore,” Springer, 2024 — peer-reviewed program analysis. — Retrieved from: \_\_\_\_\_
5. International Labour Organization (ILO), “Investigating an Upskilling Programme in Singapore” — international comparative analysis. — Retrieved from: \_\_\_\_\_

## CASE 94 **Learning Analytics on the African Continent — A Scoping Review as the Present-State Map** 2022

1. Prinsloo, P., & Kaliisa, R. (2022), “Learning Analytics on the African Continent: An Emerging Research Focus and Practice,” *Journal of Learning Analytics* 9(2):218–235, doi:10.18608/jla.2022.7539 — 15 studies (10 empirical, 5 conceptual); ten involve South Africa, two Nigeria, one each Kenya, Tanzania, Zambia. — Retrieved from: \_\_\_\_\_
2. Lemmens, J. C., & Henn, M. (2016), “Learning analytics: A South African higher education perspective,” in Botha & Muller (eds.), *Institutional Research in South African Higher Education*, 231–253, SUN PRESS — one of the fifteen included studies. — Retrieved from: \_\_\_\_\_
3. Janse van Vuuren, E. C. (2020), “Development of a contextualised data analytics framework in South African higher education: involvement of teacher (teaching) analytics as an indispensable component,” *South African Journal of Higher Education* 34(1) — one of the fifteen included studies. — Retrieved from: \_\_\_\_\_
4. Cross-reference: Case 89 (Prinsloo & Kaliisa’s companion African data-privacy paper, BJET 2022), for the construct-travel problem stated in governance terms. — Retrieved from: \_\_\_\_\_

## CASE 95 **PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop** 1997–2024

1. Bradshaw, Mitchell, & Leaf (2010), “Examining the Effects of Schoolwide Positive Behavioral Interventions and Supports on Student Outcomes: Results From a Randomized Controlled Effectiveness Trial in Elementary Schools,” *Journal of Positive Behavior Interventions*, 12(3), 133 – 148. — Retrieved from: \_\_\_\_\_
2. Bradshaw, Waasdorp, & Leaf (2012), “Effects of School-Wide Positive Behavioral Interventions and Supports on Child Behavior Problems,” *Pediatrics*, 130(5), e1136 – e1145. — Retrieved from: \_\_\_\_\_
3. Horner, Todd, Lewis-Palmer, Irvin, Sugai, & Boland (2004), “The School-Wide Evaluation Tool (SET): A Research Instrument for Assessing School-Wide Positive Behavior Support,” *Journal of Positive Behavior Interventions*, 6(1), 3 – 12. — Retrieved from: \_\_\_\_\_
4. Bradshaw, Pas, et al. (2012), “A State-Wide Partnership to Promote Safe and Supportive Schools: The PBIS Maryland Initiative,” *Administration and Policy in Mental Health*, 39(4), 225 – 237. — Retrieved from: \_\_\_\_\_
5. McIntosh, Mercer, Nese, Strickland-Cohen, & Hoselton (2016), “Predictors of Sustained Implementation of School-Wide Positive Behavioral Interventions and Supports,” *Journal of Positive Behavior Interventions*, 18(4), 209 – 218. — Retrieved from: \_\_\_\_\_
6. McIntosh, Massar, Algozzine, George, Horner, Lewis, & Swain-Bradway (2017), “Technical Adequacy of the SWPBIS Tiered Fidelity Inventory,” *Journal of Positive Behavior Interventions*, 19(1), 3 – 13. — Retrieved from: \_\_\_\_\_

**CASE 96 AeroPerú Flight 603**

1996

1. Peru Dirección General de Transporte Aéreo, Comisión de Investigación de Accidentes, final report on AeroPerú 603 (2 October 1997; with NTSB/FAA/Boeing participation) — principal cause given as error of the maintenance staff including the crew, and the instrument cascade (paraphrased). — Retrieved from: \_\_\_\_\_
2. Peru DGTA Comisión de Investigación de Accidentes report (2 October 1997) — the masking tape over the three left-side static ports and the contradictory warnings; the tape was not caught at release, at transfer to the apron, or in the pilot-in-command's walk-around. — Retrieved from: \_\_\_\_\_
3. Peru DGTA Comisión de Investigación de Accidentes report (2 October 1997) — the radio altimeter, independent of the air data computers, working normally and driving the GPWS terrain warnings; crew failure to follow the GPWS procedure and to read the radio altimeter named as contributing. CVR: "too low — terrain" 22 times in 45 seconds beginning about eight minutes before impact. — Retrieved from: \_\_\_\_\_
4. Leveson, N. (2011), Engineering a Safer World (STAMP) — common-cause failure. — Retrieved from: \_\_\_\_\_
5. Dismukes, Berman & Loukopoulos (2007), The Limits of Expertise — crew performance under instrument failure. — Retrieved from: \_\_\_\_\_

**CASE 97 Boeing 737 MAX / MCAS**

2018 – 2019

1. U.S. House Committee on Transportation and Infrastructure, The Boeing 737 MAX Aircraft: Costs, Consequences, and Lessons from Its Design, Development, and Certification (Final Committee Report, Sept. 2020) — the MCAS omission, the 2013 minutes, the 2016 survey, and the "diminished safety" conclusion. — Retrieved from: \_\_\_\_\_
2. U.S. House report (2020) and MCAS design summary — MCAS firing on a single angle-of-attack sensor; the Lion Air 610 (189 killed) and Ethiopian 302 (157 killed) accident sequences. — Retrieved from: \_\_\_\_\_
3. U.S. House Transportation Committee report (2020) — Boeing internal communications and the certification-and-training-impact reasoning (quoted). — Retrieved from: \_\_\_\_\_
4. U.S. Department of Transportation, Office of Inspector General, correspondence and review of FAA oversight of MCAS and the angle-of-attack systems (2019–2020) — the delegated-certification process. — Retrieved from: \_\_\_\_\_
5. J. Herkert, J. Borenstein & K. Miller, "The Boeing 737 MAX: Lessons for Engineering Ethics," Science and Engineering Ethics 26 (2020) — certification-by-omission as an engineering-ethics failure. — Retrieved from: \_\_\_\_\_
6. Federal Aviation Administration, Summary of the FAA's Review of the Boeing 737 MAX (Return-to-Service, 2020) — the MCAS redesign and the simulator-training requirement; and the Aircraft Certification, Safety, and Accountability Act (2020). — Retrieved from: \_\_\_\_\_

**CASE 98 Mars Climate Orbiter — Unit Mismatch**

1999

1. NASA, Mars Climate Orbiter Mishap Investigation Board: Phase I Report (Nov. 1999) — mission, the Lockheed/JPL split, and the program context. — Retrieved from: \_\_\_\_\_
2. NASA MCO MIB Report (1999) — the pound-force-second vs newton-second mismatch (factor 4.45) in the "Small Forces" ground-software file, the accumulated navigation error, the 57 km closest approach against a planned 226 km periapsis and an 80 km survival limit, and the loss on 23 Sept. 1999 (root-cause statement quoted). — Retrieved from: \_\_\_\_\_



3. NASA MCO MIB Report (1999) — the missing verified interface specification, the absent end-to-end validation, and the unresolved cruise-trajectory anomalies. — Retrieved from: \_\_\_\_\_
4. N. G. Leveson, *Engineering a Safer World* (MIT Press, 2011) — interfaces as engineering deliverables requiring an owner and a verification step. — Retrieved from: \_\_\_\_\_
5. Sauser, B. J., Reilly, R. R., & Shenhar, A. J. (2009), "Why projects fail? How contingency theory can provide new insights — A comparative analysis of NASA's Mars Climate Orbiter loss," *International Journal of Project Management* 27(7), 665–679. — Retrieved from: \_\_\_\_\_

## CASE 100 **Glass-Cockpit Transition in Light Aircraft — Technology Outran Training** 2010

1. Wiener, E. L., & Curry, R. E. (1980). Flight-deck automation: Promises and problems. *Ergonomics*, 23(10):995–1011 — the foundational literature on automation-induced failure modes that the glass-cockpit transition re-introduced at the GA scale. — Retrieved from: \_\_\_\_\_
2. National Transportation Safety Board (2010). Introduction of Glass Cockpit Avionics into Light Aircraft, Safety Study NTSB/SS-10/01. <https://www.nts.gov/safety/safety-studies/Documents/SS1001.pdf> — the case's primary investigation document. — Retrieved from: \_\_\_\_\_
3. NTSB Safety Recommendations A-10-36 through A-10-41 (2010), issued to the FAA — the engineering response the open verdict points to. — Retrieved from: \_\_\_\_\_
4. Sarter, N. B., Woods, D. D., & Billings, C. E. (1997). Automation surprises. In *Handbook of Human Factors and Ergonomics* (2nd ed.) — the mode-confusion / automation-surprise literature the NTSB findings cross-reference. — Retrieved from: \_\_\_\_\_

## CASE 101 **Ariane 5 Flight 501** 1996

1. J. L. Lions (chair), *Ariane 5 Flight 501 Failure Inquiry Board Report* (1996) — the data-conversion overflow (quoted). — Retrieved from: \_\_\_\_\_
2. Lions Report (1996) — SRI failure at HO+36.7 s (about 30 s after lift-off), disintegration under aerodynamic loads at about HO+39 s, and the automatic destruct. — Retrieved from: \_\_\_\_\_
3. Lions Report (1996) — the 64-bit-to-16-bit conversion overflow and the simultaneous shutdown of both inertial reference systems. — Retrieved from: \_\_\_\_\_
4. Lions Report (1996) — the reused code path neither removed nor re-verified for Ariane 5's envelope. — Retrieved from: \_\_\_\_\_
5. Cf. Patriot (Case 129) and Mars Climate Orbiter (Case 98). — Retrieved from: \_\_\_\_\_
6. N. Leveson, *Safeware* (1995) — software-reuse hazards; G. Le Lann, "An Analysis of the Ariane 5 Flight 501 Failure" (1997). — Retrieved from: \_\_\_\_\_
7. ESA Science & Technology, "Cluster — Summary" — the loss of the four Cluster spacecraft and the 214 MECU recovery mission approved by the SPC. — Retrieved from: \_\_\_\_\_

## CASE 102 **Air France Flight 447** 2009

1. Bureau d'Enquêtes et d'Analyses (BEA), *Final Report on the accident on 1 June 2009 to the Airbus A330-203 registered F-GZCP operated by Air France flight AF447 (July 2012)* — full report: flight, aircraft, and the known pitot-icing vulnerability with retrofit pending. — Retrieved from: \_\_\_\_\_



2. BEA, Final Report AF447 (2012) — accident sequence: autopilot/autothrust disconnection, degraded control law, sustained nose-up input, stall, and the stall-warning logic dropping out at high angle of attack. — Retrieved from: \_\_\_\_\_
3. BEA, Final Report AF447 (2012) — probable cause: airspeed inconsistency from pitot icing, inappropriate control inputs, and failure to recognize and recover from the stall. — Retrieved from: \_\_\_\_\_
4. BEA, Final Report AF447 (2012) — finding on warning-system ergonomics and stall-training conditions (quoted). — Retrieved from: \_\_\_\_\_
5. G. Palmer / E. Strickland, "Air France Flight 447 Crash Caused by a Combination of Factors," IEEE Spectrum (2014) — analysis of the divergence between the trained and operational envelopes. — Retrieved from: \_\_\_\_\_
6. BEA, Final Report AF447 (2012), safety recommendations — pitot-probe replacement, manual high-altitude flying, approach-to-stall and stall recovery, unreliable-airspeed procedures, and angle-of-attack indication. — Retrieved from: \_\_\_\_\_
7. European Union Aviation Safety Agency, Upset Prevention and Recovery Training (UPRT) requirements for airline pilots (phased in by 2019); ICAO, Manual on Aeroplane Upset Prevention and Recovery Training (Doc 10011). — Retrieved from: \_\_\_\_\_
8. Federal Aviation Administration, Qualification, Service, and Use of Crewmembers and Aircraft Dispatchers, final rule, 78 FR 67800 (2013, effective 2014) — Part 121 stall and upset recovery training, directed by the Airline Safety and FAA Extension Act of 2010; FAA Advisory Circular 120-109A, Stall Prevention and Recovery Training. — Retrieved from: \_\_\_\_\_

### CASE 103 **Kegworth / British Midland 92**

1989

1. Air Accidents Investigation Branch, Report on the accident to Boeing 737-400 G-OBME near Kegworth, Leicestershire, on 8 January 1989, AAIB Aircraft Accident Report 4/90 (1990) — aircraft, route, and the limited conversion training; see also U.S. FAA, Lessons Learned: G-OBME. — Retrieved from: \_\_\_\_\_
2. AAIB Report 4/90 (1990) — accident sequence: left-engine fan-blade failure, shutdown of the serviceable right engine, and total power loss on approach; 47 killed, 74 seriously injured. — Retrieved from: \_\_\_\_\_
3. AAIB Report 4/90 (1990) — the older-737 bleed-air mental model misapplied to the 737-400's changed configuration. — Retrieved from: \_\_\_\_\_
4. AAIB Report 4/90 (1990) — the new electronic engine instrumentation and the readability of the vibration indicator. — Retrieved from: \_\_\_\_\_
5. AAIB Report 4/90 (1990) — the cabin-to-flight-deck communication gap (crew resource management). — Retrieved from: \_\_\_\_\_
6. AAIB Report 4/90 (1990) — the 31 safety recommendations on difference training, instrument design, CRM, and crashworthiness. — Retrieved from: \_\_\_\_\_
7. J. Reason, Human Error (Cambridge Univ. Press, 1990) — Kegworth as a case in the misapplication of a correct mental model to a changed system. — Retrieved from: \_\_\_\_\_

### CASE 104 **Asiana Airlines Flight 214**

2013

1. NTSB, Aircraft Accident Report: Asiana Airlines Flight 214, NTSB/AAR-14/01 (2014) — probable cause and contributing factors (quoted). — Retrieved from: \_\_\_\_\_

2. NTSB/AAR-14/01 (2014) — the autothrottle HOLD reversion and airspeed decay to 103 kt. — Retrieved from: \_\_\_\_\_
3. NTSB/AAR-14/01 (2014) — the captain's mental model of the autothrottle. — Retrieved from: \_\_\_\_\_
4. Sarter, N. & Woods, D. (1995), "How in the world did we ever get into that mode?," Human Factors — automation surprise. — Retrieved from: \_\_\_\_\_
5. NTSB/AAR-14/01 (2014), safety recommendations A-14-37 to A-14-43 and A-14-55 on autoflight training, low-energy alerting, interface intuitiveness, and manual flight. — Retrieved from: \_\_\_\_\_

## CASE 105 Helios Airways Flight 522

2005

1. FAA Airworthiness Directive 2011-03-14, Amendment 39-16598, 76 FR 6529 (7 February 2011) — the mandated installation of separate CABIN ALTITUDE and TAKEOFF CONFIG warning lights on the 737 Classic fleet. — Retrieved from: \_\_\_\_\_
2. Hellenic Air Accident Investigation & Aviation Safety Board, Aircraft Accident Report 11/2006 — probable cause and the warning misinterpretation (paraphrased). — Retrieved from: \_\_\_\_\_
3. AAIASB Report 11/2006 — the manual pressurization selector and failure to pressurize. — Retrieved from: \_\_\_\_\_
4. AAIASB Report 11/2006 — the dual-meaning warning horn and the training emphasis. — Retrieved from: \_\_\_\_\_
5. Boeing 737 Flight Crew Operations Manual — configuration and cabin-altitude warnings. — Retrieved from: \_\_\_\_\_
6. Reason, J. (1990), Human Error — the cue-ambiguity / human-error framing the case draws on. — Retrieved from: \_\_\_\_\_

## CASE 106 TransAsia Airways Flight 235

2015

1. Aviation Safety Council (Taiwan), Aviation Occurrence Report: TransAsia Airways Flight GE235, final report (2016) — findings and the wrong-engine shutdown (paraphrased). — Retrieved from: \_\_\_\_\_
2. ASC final report (2016) — the autofeather event and the crash sequence. — Retrieved from: \_\_\_\_\_
3. ASC final report (2016) — non-use of the checklist and the captain's proficiency record. — Retrieved from: \_\_\_\_\_
4. AAIB (UK), British Midland Flight 92 (Kegworth) report (1990) — the antecedent wrong-engine case. — Retrieved from: \_\_\_\_\_
5. Dismukes, Berman & Loukopoulos (2007), The Limits of Expertise — memory-driven response under time pressure. — Retrieved from: \_\_\_\_\_

## CASE 107 Eastern Air Lines Flight 401

1972

1. NTSB, Aircraft Accident Report: Eastern Air Lines Flight 401, NTSB-AAR-73-14 (1973) — the night approach and full flight deck. — Retrieved from: \_\_\_\_\_
2. NTSB-AAR-73-14 (1973) — the landing-gear-bulb fixation, the inadvertent autopilot disengagement, and 101 deaths of 176 aboard. — Retrieved from: \_\_\_\_\_

3. NTSB-AAR-73-14 (1973) — the crew's preoccupation with the landing-gear indication that distracted them from maintaining altitude (paraphrased). — Retrieved from: \_\_\_\_\_
4. C. D. Wickens et al., *Engineering Psychology and Human Performance* — attention, monitoring, and alert design. — Retrieved from: \_\_\_\_\_
5. R. Helmreich & H. Foushee (1993), *aircrew-coordination research — the Crew Resource Management origins*. — Retrieved from: \_\_\_\_\_
6. Cooper, G. E., White, M. D. & Lauber, J. K. (1980), *Resource Management on the Flightdeck*, NASA CP-2120 — proceedings of the 1979 NASA/industry workshop that named Crew Resource Management. — Retrieved from: \_\_\_\_\_

### CASE 108 NASA EVA-23 — Water Intrusion Inside a Spacesuit

2013

1. NASA International Space Station Program (2013), "International Space Station EVA Suit Water Intrusion High Visibility Close Call" — Mishap Investigation Board final report, December 2013. — Retrieved from: \_\_\_\_\_
2. NASA Aerospace Safety Advisory Panel (2014), *Annual Report* — EVA-23 follow-up actions including helmet absorption pad and snorkel. — Retrieved from: \_\_\_\_\_
3. Parmitano, L. (2013), "EVA 23: exploring the frontier," ESA astronaut blog, 20 August 2013 — the crewmember's account of the water event and the return traverse. — Retrieved from: \_\_\_\_\_
4. Chappell, Norcross, Abercromby et al. (2017), "Risk of Injury and Compromised Performance Due to EVA Operations," NASA Human Research Program Evidence Report — broader EVA risk context. — Retrieved from: \_\_\_\_\_

### CASE 109 Colgan Air Flight 3407

2009

1. NTSB/AAR-10/01 (2010) — the captain's training history, and the first officer's fatigue and pay. — Retrieved from: \_\_\_\_\_
2. NTSB, *Aircraft Accident Report: Colgan Air Flight 3407*, NTSB/AAR-10/01 (2010) — probable cause (quoted). — Retrieved from: \_\_\_\_\_
3. NTSB/AAR-10/01 (2010), §2.7.3 — PRIA and FAA guidance did not require operators to obtain notices of disapproval; recommendation A-05-1 reiterated as "Open—Unacceptable Response". — Retrieved from: \_\_\_\_\_
4. DOT Office of Inspector General, *FAA Delays in Establishing a Pilot Records Database Limit Air Carriers' Access to Background Information*, AV-2015-079 (2015); the Families of Continental Flight 3407 campaign. — Retrieved from: \_\_\_\_\_
5. *Airline Safety and Federal Aviation Administration Extension Act of 2010*, Pub. L. 111-216 — the 1,500-hour rule and the Pilot Records Database. — Retrieved from: \_\_\_\_\_

### CASE 110 Atlas Air Flight 3591

2019

1. NTSB/AAR-20/02 (2020) — the first officer's training history and Atlas's record-access limits. — Retrieved from: \_\_\_\_\_
2. NTSB/AAR-20/02 (2020) — the inadvertent go-around activation and the dive into Trinity Bay. — Retrieved from: \_\_\_\_\_

3. NTSB, Aircraft Accident Report: Atlas Air Flight 3591, NTSB/AAR-20/02 (2020) — probable cause and contributing factors (paraphrased). — Retrieved from: \_\_\_\_\_
4. DOT Office of Inspector General, FAA Delays in Establishing a Pilot Records Database Limit Air Carriers' Access to Background Information, AV-2015-079 (2015) — the decade of implementation delay. — Retrieved from: \_\_\_\_\_
5. FAA Pilot Records Database final rule (Federal Register, 10 June 2021; subpart B/C compliance from June 2022, full historical coverage by September 2024); Pilot Records Improvement Act of 1996. — Retrieved from: \_\_\_\_\_

## CASE 111 Challenger & Columbia

1986 / 2003

1. W. Starbuck & M. Farjoun (eds.), Organization at the Limit: Lessons from the Columbia Disaster (2005) — independent academic re-analyses of the institutional-learning gap. — Retrieved from: \_\_\_\_\_
2. Columbia Accident Investigation Board, Report Vol. I (2003) — the foam strike; finding F3.2-5 (left bipod ramp shedding on 7 of the 72 missions for which launch or ET-separation imagery exists, including instances NASA was unaware of) and finding F3.2-7 (foam loss on more than 80 percent of the 79 missions with confirming imagery); the filing of strikes as “in-family,” the recurrence of the cultural pattern, and the call for an independent technical authority. — Retrieved from: \_\_\_\_\_
3. Rogers Commission, Report of the Presidential Commission on the Space Shuttle Challenger Accident (1986) — the O-ring failure, the Thiokol teleconference, and the cold-weather launch decision. — Retrieved from: \_\_\_\_\_
4. CAIB (2003) — “the NASA organizational culture had as much to do with this accident as the foam” (quoted); Rogers Commission (1986) and CAIB (2003) jointly on the suppressed upward flow of safety information across both accidents. — Retrieved from: \_\_\_\_\_
5. D. Vaughan, The Challenger Launch Decision: Risky Technology, Culture, and Deviance at NASA (Univ. of Chicago Press, 1996) — normalization of deviance; the working group's drift of decision rules. — Retrieved from: \_\_\_\_\_
6. NASA Engineering and Safety Center founding documents (2003 – present) — institutional response to the CAIB's call for independent technical authority; Shuttle retirement (STS-135, July 2011). — Retrieved from: \_\_\_\_\_

## CASE 112 NASA Saturn V Documentation

1972 – present

1. R. E. Bilstein, Stages to Saturn: A Technological History of the Apollo/Saturn Launch Vehicles (NASA SP-4206, 1980) — the program and its workforce. — Retrieved from: \_\_\_\_\_
2. NASA Constellation Program documentation and reviews (2005–2010) — the difficulty of reproducing Saturn V capability. — Retrieved from: \_\_\_\_\_
3. The documentation-vs-capability distinction (editors' synthesis of the Saturn V record). — Retrieved from: \_\_\_\_\_
4. M. Polanyi, The Tacit Dimension (1966); I. Nonaka & H. Takeuchi, The Knowledge-Creating Company (1995). — Retrieved from: \_\_\_\_\_
5. J. S. Brown & P. Duguid, The Social Life of Information (2000) — tacit and institutional knowledge. — Retrieved from: \_\_\_\_\_

**CASE 113 Boeing Starliner**

2019 – 2025

1. U.S. GAO, NASA Commercial Crew Program: Significant Work Remains to Begin Operational Missions to the Space Station, GAO-20-121 (Jan. 2020); Schedule Uncertainty Persists for Start of Operational Missions, GAO-19-504 (June 2019). — Retrieved from: \_\_\_\_\_
2. NASA and Boeing joint Independent Review Team, Orbital Flight Test findings (7 July 2020) — 80 recommendations across software requirements, testing and simulation, process, and hardware. — Retrieved from: \_\_\_\_\_
3. NASA OIG, NASA's Management of Its Commercial Crew Program, IG-26-011 (30 June 2026); NASA, Starliner Propulsion System Anomalies during the Crewed Flight Test (redacted, February 2026). — Retrieved from: \_\_\_\_\_
4. NASA Aerospace Safety Advisory Panel, annual reports (2020–2025) — Commercial Crew oversight and the Panel's criticism of NASA's mishap-classification requirements. — Retrieved from: \_\_\_\_\_
5. NASA, Starliner's uncrewed landing at White Sands Space Harbor (6 September 2024) and the Crew-9 return of Wilmore and Williams aboard a SpaceX Dragon (18 March 2025). — Retrieved from: \_\_\_\_\_
6. Cf. Saturn V (Case 112); N. Augustine, Augustine's Laws (1986). — Retrieved from: \_\_\_\_\_

**CASE 114 FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule**

1988 – 2010s

1. NTSB (1989), Aircraft Accident Report AAR-89/03, Aloha Airlines Flight 243, Boeing 737-200, N73711. — Retrieved from: \_\_\_\_\_
2. Swift, T. (1993), "Widespread Fatigue Damage Monitoring — Issues and Concerns," International Conference on Structural Airworthiness of New and Aging Aircraft, Hamburg, June 16–18, 1993 — technical synthesis of the WFD inspection regime. — Retrieved from: \_\_\_\_\_
3. Airworthiness Assurance Working Group (1999), Recommendations for Regulatory Action to Prevent Widespread Fatigue Damage in the Commercial Airplane Fleet, Final Report, March 11, 1999 (Rev. A, June 29, 1999) — the ARAC record behind the rule. — Retrieved from: \_\_\_\_\_
4. FAA, "Aging Airplane Program: Widespread Fatigue Damage," Final Rule, 75 FR 69746 (November 15, 2010, effective January 14, 2011) — the subpart C limit-of-validity provisions. — Retrieved from: \_\_\_\_\_
5. FAA, "Enhanced Airworthiness Program for Airplane Systems/Fuel Tank Safety (EAPAS/FTS)," Final Rule, 72 FR 63409 (November 8, 2007), Docket FAA-2004-18379 — the rule that created 14 CFR Part 26, "Continued Airworthiness and Safety Improvements for Transport Category Airplanes." — Retrieved from: \_\_\_\_\_

**CASE 115 FAA NextGen and the ADS-B Out Transition**

2003 – 2020

1. FAA, "Automatic Dependent Surveillance-Broadcast (ADS-B) Out Performance Requirements To Support Air Traffic Control (ATC) Service," Final Rule, 75 FR 30160 (May 28, 2010) — 14 CFR §§91.225 and 91.227. — Retrieved from: \_\_\_\_\_
2. Vision 100 — Century of Aviation Reauthorization Act, Public Law 108-176 (2003) — NextGen program statutory basis. — Retrieved from: \_\_\_\_\_
3. Government Accountability Office, "NextGen Air Transportation System" report series (2010s) — sustained external audit record on schedule slippage and benefit-realization gaps. — Retrieved from: \_\_\_\_\_

4. Department of Transportation Office of Inspector General (2021), NextGen Benefits Have Not Kept Pace With Initial Projections, Report No. AV2021023, March 30, 2021 — FAA's 2017 benefits projection \$113 billion below the JPDO estimate; \$6 billion in benefits realized 2010–2018. — Retrieved from: \_\_\_\_\_
5. FAA, NextGen Annual Reports (2010 – present) — program-self-report tier; read against the GAO and DOT IG reviews. — Retrieved from: \_\_\_\_\_

## CASE 116 Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change 2005

1. DelaPeyronnie, Newcomb, Morillo, Trimech, Nguyen, & Purtill (2010), "Modernization of the Eurocat Air Traffic Management System (EATMS)," in Information Systems Transformation: Architecture-Driven Modernization Case Studies (Elsevier / Morgan Kaufmann), Chapter 5. — Retrieved from: \_\_\_\_\_
2. Ulrich & Newcomb (eds., 2010), Information Systems Transformation — the host volume on architecture-driven modernization patterns. — Retrieved from: \_\_\_\_\_
3. Brodie & Stonebraker (1995), Migrating Legacy Systems — the framing reference on small-step legacy modernization. — Retrieved from: \_\_\_\_\_
4. Bisbal et al. (1999), "Legacy Information Systems: Issues and Directions," IEEE Software 16(5):103–111 — peer-reviewed framing companion. — Retrieved from: \_\_\_\_\_

## CASE 117 Crew Resource Management & CAST 1981 – present

1. Spanish CIAIAC / ALPA reports on the 1977 Tenerife collision — the overridden crew challenge. — Retrieved from: \_\_\_\_\_
2. Helmreich, Merritt & Wilhelm (1999), "The Evolution of Crew Resource Management Training in Commercial Aviation," International Journal of Aviation Psychology. — Retrieved from: \_\_\_\_\_
3. FAA Advisory Circular 120–51E, Crew Resource Management Training — CRM protocols and authority-gradient training. — Retrieved from: \_\_\_\_\_
4. CAST/FAA Safety Enhancement reports (2016, 2018) — the closed-loop data process and the 83% reduction. — Retrieved from: \_\_\_\_\_
5. Collier Trophy citation (2008); Kanki, Helmreich & Anca (2010), Crew Resource Management. — Retrieved from: \_\_\_\_\_

## CASE 118 Korean Air Safety Transformation 2000 – present

1. NTSB (2000), AAR–00/01, Controlled Flight Into Terrain, Korean Air Flight 801, Boeing 747–300, HL7468, Nimitz Hill, Guam, August 6, 1997 — probable cause at §3.2: the captain's failure to adequately brief and execute the nonprecision approach and the first officer's and flight engineer's failure to effectively monitor and cross-check it; contributing to those failures, the captain's fatigue and Korean Air's inadequate flight crew training; contributing to the accident, the FAA's intentional inhibition of the MSAW system at Guam and its failure to adequately manage it. — Retrieved from: \_\_\_\_\_
2. Helmreich, Wilhelm, Klinec & Merritt (2001), "Culture, Error, and Crew Resource Management," in Salas, Bowers & Edens (eds.), Improving Teamwork in Organizations (Erlbaum) — national culture and CRM adaptation. — Retrieved from: \_\_\_\_\_
3. Gladwell, M. (2008), Outliers — the Korean Air chapter on power distance and the English-language change. — Retrieved from: \_\_\_\_\_

4. Air Transport World Phoenix Award documentation (2006) — recognition of the turnaround. — Retrieved from: \_\_\_\_\_
5. Korean Air corporate safety reports — the post-2000 accident-free record. — Retrieved from: \_\_\_\_\_

### CASE 119 Aviation Safety Reporting System (ASRS)

1976 – present

1. Reason, J. (1997), Managing the Risks of Organizational Accidents — non-punitive reporting as a model (paraphrased). — Retrieved from: \_\_\_\_\_
2. NASA ASRS Program documentation and annual reports — system design, immunity provision, and report volume. — Retrieved from: \_\_\_\_\_
3. FAA Advisory Circular 00-46F, “Aviation Safety Reporting Program” — the program’s establishment in April 1975, the FAA–NASA Memorandum of Agreement of 15 August 1975, and the immunity provisions. — Retrieved from: \_\_\_\_\_
4. NASA ASRS technical reports (Connell et al.) — patterns first surfaced via ASRS data. — Retrieved from: \_\_\_\_\_
5. Dekker, S. (2012), Just Culture — the cultural commitment to non-punitive use. — Retrieved from: \_\_\_\_\_
6. CHIRP and patient-safety reporting-system documentation — cross-domain emulation. — Retrieved from: \_\_\_\_\_

### CASE 120 EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation

1996 – 2002

1. Royal Commission to Inquire into the Crash on Mount Erebus, Antarctica of a DC10 Aircraft Operated by Air New Zealand Limited (1981), final report (Mahon report). — Retrieved from: \_\_\_\_\_
2. Aeronáutica Civil of the Republic of Colombia (1996), Aircraft Accident Report: Controlled Flight Into Terrain, American Airlines Flight 965, Boeing 757-223, N651AA, near Cali, Colombia, December 20, 1995 (NTSB participating, DCA96RAO20). — Retrieved from: \_\_\_\_\_
3. Federal Aviation Administration (2000), 14 CFR §§ 91.223, 121.354, 135.154 — Terrain Awareness and Warning System (TAWS) equipage requirement. — Retrieved from: \_\_\_\_\_
4. Flight Safety Foundation (1998 – 2000), CFIT / ALAR Task Force reports — operational and outcome analyses motivating mandate. — Retrieved from: \_\_\_\_\_
5. Bateman, C. D. (1999), “The Introduction of Enhanced Ground Proximity Warning Systems (EGPWS) into Civil Aviation Operations Around the World,” in Proceedings of the 11th Annual European Aviation Safety Seminar, Flight Safety Foundation, pp. 259–274 — developer history. — Retrieved from: \_\_\_\_\_

### CASE 121 TCAS — Coordinated Collision Avoidance and the Überlingen Lesson

1981 – 2008 (TCAS II Version 7.1)

1. Federal Aviation Administration, 14 CFR § 121.356 — TCAS II equipage requirement. — Retrieved from: \_\_\_\_\_
2. RTCA (2008), DO-185B “Minimum Operational Performance Standards for Traffic Alert and Collision Avoidance System II (TCAS II)” — Version 7.1 with reversal logic. — Retrieved from: \_\_\_\_\_



3. Bundesstelle für Flugunfalluntersuchung (BFU) (2004), AX001-1-2/02 — Investigation Report on the mid-air collision on 1 July 2002 near Überlingen. — Retrieved from: \_\_\_\_\_
4. ICAO, PANS-OPS (Doc 8168), Amendment 12 — in force November 2003; the RA-takes-precedence rule and the requirement to report RAs to ATC. — Retrieved from: \_\_\_\_\_
5. Espindle, Griffith & Kuchar, "Safety Analysis of Upgrading to TCAS Version 7.1 Using the 2008 U.S. Correlated Encounter Model," MIT Lincoln Laboratory Project Report ATC-349 (2009) — the quantitative basis for the V7.1 change. — Retrieved from: \_\_\_\_\_
6. Commission Regulation (EU) No 1332/2011 — ACAS II version 7.1 required in EU airspace from 1 March 2012 for new aeroplanes and from 1 December 2015 for the existing fleet. — Retrieved from: \_\_\_\_\_

## CASE 122 Singapore Airlines Safety Transformation

1980s – present

1. Aviation Safety Council (Taiwan), Crashed on a Partially Closed Runway during Takeoff, Singapore Airlines Flight 006, Boeing 747-400, 9V-SPK, CKS Airport, Taoyuan, October 31, 2000 (2002) — the accident, the eight findings related to probable causes (§3.1: one weather, seven flight crew), and the Singapore MOT investigation team's comments appended at §7. — Retrieved from: \_\_\_\_\_
2. Henderson, J. C. (2003), "Communicating in a crisis: flight SQ 006," *Tourism Management* 24(3):279–287 — the analysis of Singapore Airlines' early crisis response. — Retrieved from: \_\_\_\_\_
3. Singapore Airlines, "Singapore Airlines keeps fleet young" (SilverKris/company account) and SIA annual reports — the young-fleet policy in the airline's own words; in-service average fleet age about 8.3 years as of 2026. — Retrieved from: \_\_\_\_\_
4. Helmreich, Wilhelm, Klinec & Merritt (2001), "Culture, error, and crew resource management," in Salas, Bowers & Edens (eds.), *Improving Teamwork in Organizations* — national culture and CRM adaptation. — Retrieved from: \_\_\_\_\_
5. Weick & Sutcliffe (2007), *Managing the Unexpected* — sustained high-reliability investment. — Retrieved from: \_\_\_\_\_

## CASE 123 FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike

1992 – 2000s

1. Flight Safety Foundation, "Killers in Aviation: FSF Task Force Presents Facts About Approach-and-landing and Controlled-flight-into-terrain Accidents," *Flight Safety Digest* (1998–1999). — Retrieved from: \_\_\_\_\_
2. Flight Safety Foundation, ALAR Tool Kit (2000; updated 2010) — distributed multi-element package built on the 34 ALAR Briefing Notes published in *Flight Safety Digest*, August–November 2000. — Retrieved from: \_\_\_\_\_
3. Khatwa & Helmreich, "Analysis of critical factors during approach and landing in accidents and normal flight," *Flight Safety Digest* (1998) — the analytical basis of the ALAR Task Force scope. — Retrieved from: \_\_\_\_\_
4. FAA, Terrain Awareness and Warning System (TAWS) Final Rule, 65 FR 16736 (29 March 2000), 14 CFR §§ 91.223, 121.354, 135.154 — new-build compliance 29 March 2002, retrofit by 29 March 2005. — Retrieved from: \_\_\_\_\_
5. ICAO Annex 6, ground-proximity warning with a forward-looking terrain-avoidance function — new certificates of airworthiness from 1 January 2004, the existing turbine fleet above 5,700 kg or over nine passengers by 1 January 2007; the international consolidation. — Retrieved from: \_\_\_\_\_



**CASE 124 USS Fitzgerald & USS John S. McCain**

2017

1. T. C. Miller, M. Faturechi & R. Rotella, "Years of Warnings, Then Death and Disaster," ProPublica (2019) — the 2003 SWOS-in-a-Box shift. — Retrieved from: \_\_\_\_\_
2. GAO, Navy Readiness, GAO-17-809T (Sept. 2017) — 37% of Japan-based warfare certifications expired by June 2017. — Retrieved from: \_\_\_\_\_
3. NTSB, Collision between USS Fitzgerald and MV ACX Crystal, NTSB/MAR-20/02 (2020), DCA17PM018. — Retrieved from: \_\_\_\_\_
4. NTSB, Collision between USS John S. McCain and Tanker Alnic MC, NTSB/MAR-19/01 (2019), DCA17PM029. — Retrieved from: \_\_\_\_\_
5. U.S. Navy, Comprehensive Review of Recent Surface Force Incidents (Nov. 2017) — both collisions judged avoidable. — Retrieved from: \_\_\_\_\_
6. NTSB/MAR-19/01 (2019) — probable cause: insufficient training and inadequate bridge procedures from a lack of Navy oversight. — Retrieved from: \_\_\_\_\_
7. R. Rotella et al., "The Navy Installed Touch-Screen Steering Systems to Save Money," ProPublica (2019). — Retrieved from: \_\_\_\_\_
8. U.S. Navy, Strategic Readiness Review (Dec. 2017) — risks that "accumulated over time and did so insidiously." — Retrieved from: \_\_\_\_\_
9. U.S. Navy corrective actions (2017); NTSB/MAR-19/01 recommendations; USNI News (2017–2020). — Retrieved from: \_\_\_\_\_
10. Surface Warfare Officers School Command, Basic Division Officer Course curriculum and the post-2017 return to in-person instruction; Naval Surface Group Western Pacific stand-up (2018) as the forward-readiness certification authority. — Retrieved from: \_\_\_\_\_
11. NTSB/MAR-19/01 (2019) — watchstander fatigue findings, including average sleep hours and the touch-screen helm misdiagnosis sequence. — Retrieved from: \_\_\_\_\_

**CASE 126 F-35 Sustainment & the Maintainer Capability Gap**

ongoing

1. U.S. Government Accountability Office, F-35 Aircraft: DOD and the Military Services Need to Reassess the Future Sustainment Strategy, GAO-23-105341 (Sept. 2023) — 2,500 planned aircraft and a lifecycle cost exceeding \$1.7 trillion, \$1.3 trillion of it in operating and support. — Retrieved from: \_\_\_\_\_
2. GAO-23-105341 (2023) — 55% fleet mission-capable rate (March 2023); over 10,000 components awaiting repair; 141-day average depot repair cycle time against a 60-to-90-day goal; depot stand-up behind schedule. — Retrieved from: \_\_\_\_\_
3. GAO-23-105341 (2023) — sustainment shortfall traced to insufficient technical data, training and support equipment for maintainers, depot stand-up delays, and contractor dependency. — Retrieved from: \_\_\_\_\_
4. GAO, F-35 Sustainment: DOD Faces Several Uncertainties..., GAO-22-105995 (2022), and the broader GAO F-35 series — the recurring, year-over-year diagnosis. — Retrieved from: \_\_\_\_\_
5. GAO-23-105341 (2023) — recommendation that DOD reassess the future sustainment strategy. — Retrieved from: \_\_\_\_\_

6. GAO, F-35 Sustainment: Costs Continue to Rise While Planned Use and Availability Have Decreased, GAO-24-106703 (2024) — costs rising while readiness stays below goal; \$1.58 trillion sustainment estimate and the above-\$2 trillion total. — Retrieved from: \_\_\_\_\_
7. GAO, F-35 Sustainment: Actions Needed to Ensure Updated Strategy Improves Persistent Readiness Challenges, GAO-26-108113 (11 June 2026) — mission-capable rate 67% (FY2021) to 44% (FY2025), full-mission-capable 38% to 25%, and the Global Support Solution Reset's \$13.7 billion of additional funding through FY2031. — Retrieved from: \_\_\_\_\_

## CASE 127 Military Fratricide — Desert Storm to Afghanistan

1991 – 2014

1. C. R. Shrader, *Amicide: The Problem of Friendly Fire in Modern War* (U.S. Army Combat Studies Institute, 1982) — the under-2% estimate from 269 incidents, a baseline later challenged as too low (OTA 1993: 15 to 20% may be the norm). — Retrieved from: \_\_\_\_\_
2. "Friendly Fire: Facts, Myths and Misperceptions," U.S. Naval Institute Proceedings (June 1994) — Desert Storm: 35 of 146 KIA (24%) and 72 of 467 wounded (15%) by friendly fire; critique of the 2% "historical norm." — Retrieved from: \_\_\_\_\_
3. Khafji and Warrior fratricide incidents (Feb. 1991) — see USNI Proceedings (1994) and Time, "They Didn't Have to Die". (Per-incident casualty figures vary across sources; see AUDIT.) — Retrieved from: \_\_\_\_\_
4. U.S. GAO, Operation Desert Storm fratricide investigations — Apache incident (OSI-93-4) and Army fratricide investigation (OSI-95-10) — causes and the Army's reviews. — Retrieved from: \_\_\_\_\_
5. U.S. Congress, Office of Technology Assessment, *Who Goes There: Friend or Foe?*, OTA-ISC-537 (June 1993) — the Desert Storm counts (615 battle casualties, 148 fatalities, 35 by friendly fire), the "15 to 20 percent may be the norm" reassessment, the absence of broad-based fratricide data, and the NTC observer distribution of causes (identification failure, command and control, planning, communication and navigation). — Retrieved from: \_\_\_\_\_
6. Later-conflict recurrences (2001 coordinate-reset strike; 2014 B-1 infrared-strobe strike) and S. Snook, *Friendly Fire* (Princeton Univ. Press, 2000) on the 1994 Black Hawk shootdown — the systems-integration archetype. — Retrieved from: \_\_\_\_\_
7. "IFF Update: Stalled Again," U.S. Naval Institute Proceedings (June 1994) — the combat-identification and identification-friend-or-foe programs pursued after Desert Storm, and how slowly they matured. — Retrieved from: \_\_\_\_\_

## CASE 128 Operation Eagle Claw

1980

1. Locher, J. (2002), *Victory on the Potomac* — the reform arc from Desert One to Goldwater-Nichols. — Retrieved from: \_\_\_\_\_
2. Holloway Commission (1980) — the Desert One sequence and the helicopter-C-130 collision. — Retrieved from: \_\_\_\_\_
3. Holloway Special Operations Review Group, *Rescue Mission Report* (1980) — 23 issues affecting the mission, of which command and control and the ad hoc nature of the organization and planning are the two the Review Group treated as fundamental; the helicopter and aircrew selection is among them. — Retrieved from: \_\_\_\_\_
4. Goldwater-Nichols Department of Defense Reorganization Act of 1986, Pub. L. 99-433. — Retrieved from: \_\_\_\_\_

5. Nunn-Cohen Amendment (1986), establishing U.S. Special Operations Command (1987). — Retrieved from: \_\_\_\_\_

### CASE 129 Patriot Missile / Dhahran

1991

1. U.S. General Accounting Office, Patriot Missile Defense: Software Problem Led to System Failure at Dhahran, Saudi Arabia, GAO/IMTEC-92-26 (1992) — the design-context mismatch and continuous-operation use. — Retrieved from: \_\_\_\_\_
2. R. Skeel, "Roundoff Error and the Patriot Missile," SIAM News 25(4): 11 (1992) — the 24-bit fixed-point time truncation and the 0.34-second range-gate drift after 100 hours. — Retrieved from: \_\_\_\_\_
3. GAO/IMTEC-92-26 (1992) — the prior Israeli warning, the software patch arriving in Dhahran on 26 February (the day after the strike), and the 28 killed. — Retrieved from: \_\_\_\_\_
4. GAO/IMTEC-92-26 (1992) — the vague "very long run time" advisory and the Army's presumption that batteries would not be run continuously for such periods (quoted). — Retrieved from: \_\_\_\_\_
5. M. Barr / Barr Group, "An Update on the Patriot Missile Software Problem" — an engineering post-mortem of the accumulated-truncation defect. — Retrieved from: \_\_\_\_\_
6. GAO/IMTEC-92-26 (1992) and Skeel (1992) — the design assumption that failed to travel across the change in operational context, and the case's teaching legacy. — Retrieved from: \_\_\_\_\_

### CASE 130 USS Vincennes & Iran Air Flight 655

1988

1. J. Barry & R. Charles, "Sea of Lies," Newsweek (July 13, 1992) — contemporaneous reinvestigation of the operational record, including the disputed account of Vincennes' position relative to Iranian territorial waters. — Retrieved from: \_\_\_\_\_
2. Rear Adm. W. Fogarty, Formal Investigation into the Circumstances Surrounding the Downing of Iran Air Flight 655 (U.S. Navy, August 1988) — the engagement, the IFF mode-confusion findings, and the 290 deaths. — Retrieved from: \_\_\_\_\_
3. Fogarty report (1988) — the Aegis SPY-1A radar functioned; the aircraft was ascending while the crew perceived a descent; "scenario fulfillment" as the psychological mechanism. — Retrieved from: \_\_\_\_\_
4. Fogarty report (1988) — "human error under extreme stress," confirmation bias, and "unconscious distortion of data" (quoted); shared-error finding across CIC operators. — Retrieved from: \_\_\_\_\_
5. LTC A. Tingle, USA, "The Human-Machine Team Failed Vincennes", U.S. Naval Institute Proceedings 144/7 (July 2018) — the human-AI-teaming reframing and the standing case-study role. — Retrieved from: \_\_\_\_\_
6. M. L. Cummings, "Automation Bias in Intelligent Time Critical Decision Support Systems," AIAA 1st Intelligent Systems Technical Conference (2004), doi:10.2514/6.2004-6313 — interface design and automation under time pressure as the engineering frame for the case. — Retrieved from: \_\_\_\_\_
7. G. Klein, Sources of Power (1998); M. Endsley, "Toward a Theory of Situation Awareness" (1995) — the naturalistic-decision-making and situation-awareness literatures that treat Vincennes as the canonical worked example. — Retrieved from: \_\_\_\_\_

### CASE 131 Stanislav Petrov / 1983 False Alert

1983

1. D. Hoffman, The Dead Hand: The Cold War Arms Race and Its Dangerous Legacy (2009) — the Oko system and Soviet launch-on-warning posture. — Retrieved from: \_\_\_\_\_

2. Accounts of the 26 Sept. 1983 incident (Hoffman; Petrov interviews) — the five-missile alert and Petrov's assessment. — Retrieved from: \_\_\_\_\_
3. Investigations of the Oko incident (incl. V. Aksenov, 2004) — sunlight-on-clouds satellite geometry as the false-positive cause. — Retrieved from: \_\_\_\_\_
4. The subsequent Oko algorithm modification and the incident's use in early-warning training. — Retrieved from: \_\_\_\_\_
5. S. Petrov, interview, The Washington Post (1999) — "a funny feeling in my gut" (quoted). — Retrieved from: \_\_\_\_\_
6. S. Sagan, *The Limits of Safety* (1993); M. L. Cummings (2017) — automation in critical decisions and human-in-the-loop. — Retrieved from: \_\_\_\_\_

### CASE 132 **F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap** 2008 – 2012

1. USAF Scientific Advisory Board, Report on Aircraft Oxygen Generation, SAB-TR-11-04 (1 February 2012) — the 2011 Quicklook Study on the F-22 hypoxia-like incidents. — Retrieved from: \_\_\_\_\_
2. NASA Engineering and Safety Center, F-22 Life Support System Independent Analysis Team (2012) — NESC support to the Air Force life-support investigation. — Retrieved from: \_\_\_\_\_
3. USAF Accident Investigation Board (2011), Capt. Jeffrey Haney F-22 accident report (Nov 2010). — Retrieved from: \_\_\_\_\_
4. U.S. House Armed Services Subcommittee on Tactical Air and Land Forces, F-22 Pilot Physiological Issues, H.A.S.C. No. 112-154 (13 September 2012) — congressional record. — Retrieved from: \_\_\_\_\_

### CASE 133 **Mark 14 Torpedo Failures** 1941 – 1943

1. Rowland, B. & Boyd, W. (1953), U.S. Navy Bureau of Ordnance in World War II — testing policy and the three defects. — Retrieved from: \_\_\_\_\_
2. Blair, C. (1975), *Silent Victory: The U.S. Submarine War Against Japan* — operator reports and the Bureau's resistance (paraphrased). — Retrieved from: \_\_\_\_\_
3. Blair (1975) — the USS Tinosa's July 1943 attack (eleven consecutive duds on a stopped target) and Lockwood's fleet-level testing. — Retrieved from: \_\_\_\_\_
4. Naval History and Heritage Command, Mark 14 torpedo files — depth error, Mark 6 magnetic exploder, contact-pin failure. — Retrieved from: \_\_\_\_\_
5. Edmondson, A. (2018), *The Fearless Organization* — suppression of field reports as an organizational-learning failure. — Retrieved from: \_\_\_\_\_

### CASE 134 **V-22 Osprey** 1991 – present

1. Compiled V-22 accident record — 16 hull losses and 62 fatalities since 1991, including the 2000 Marana test crash (19 Marines). (A compiled tally, not an official count; the four-crash, twenty-death figure for 2022–2025 is GAO's.) — Retrieved from: \_\_\_\_\_
2. U.S. Air Force Accident Investigation Board findings, via USNI News (Aug. 2024) — the 29 Nov. 2023 Yakushima CV-22B crash: transmission-gear cracks (X-53 inclusions) and continued flight despite warnings. — Retrieved from: \_\_\_\_\_

3. U.S. GAO, Osprey Aircraft: Additional Oversight and Information Sharing Would Improve Safety Efforts, GAO-26-107285 (Dec. 2025) — the joint program office's failure to assess and address risks. — Retrieved from: \_\_\_\_\_
4. GAO-26-107285 (2025) — serious-mishap rates exceeding comparable Navy/Air Force fleets (FY2015–FY2024); 28 unresolved system safety risks with a median age of about nine years, over half carried six to fourteen years. — Retrieved from: \_\_\_\_\_
5. NAVAIR independent review of the V-22 (Dec. 2025), reported in Defense One, “New GAO, Navy reports warn of serious V-22 Osprey safety risks” (Dec. 2025) — material and cross-service-coordination factors, unresolved catastrophic parts issues, the 5-of-12 mishap finding, and proprotor-gearbox fixes running to 2034. — Retrieved from: \_\_\_\_\_
6. USNI News, V-22 program coverage (2024–2025). — Retrieved from: \_\_\_\_\_

### CASE 135 9/11 Intelligence Sharing Failures

1996 – 2001

1. National Commission on Terrorist Attacks Upon the United States, The 9/11 Commission Report (2004) — the quoted “failure of imagination” and the specific sharing failures. — Retrieved from: \_\_\_\_\_
2. 9/11 Commission Report (2004) — the Kuala Lumpur tracking and the Phoenix/Minneapolis flagging. — Retrieved from: \_\_\_\_\_
3. Joint Inquiry into Intelligence Community Activities Before and After the Terrorist Attacks of September 11, 2001, S. Rept. 107-351 (2002) — cross-agency information-sharing failures. — Retrieved from: \_\_\_\_\_
4. Zegart, A. (2007), Spying Blind — structural-organizational analysis of the failures. — Retrieved from: \_\_\_\_\_
5. Intelligence Reform and Terrorism Prevention Act of 2004 — creation of the ODNI and NCTC. — Retrieved from: \_\_\_\_\_

### CASE 137 U.S. Nuclear Navy / Rickover Training Model

1954 – present

1. Polmar, N. & Allen, T. (2007), Rickover: Father of the Nuclear Navy — the program and Rickover's philosophy (paraphrased). — Retrieved from: \_\_\_\_\_
2. Naval Nuclear Propulsion Program documentation (NRC/DOE) — qualification standards and the accident record. — Retrieved from: \_\_\_\_\_
3. Admiral Hyman G. Rickover, “Doing a Job” (Columbia University, Egleston Medal address, 5 November 1982) — “people, not organizations... get things done” (quoted). — Retrieved from: \_\_\_\_\_
4. Duncan, F. (1990), Rickover and the Nuclear Navy — the qualification culture. — Retrieved from: \_\_\_\_\_
5. GAO-21-168 (2021), Navy Readiness: Actions Needed to Evaluate and Improve Surface Warfare Officer Career Path — the surface community's career path set against the submarine, aviation and EOD communities. — Retrieved from: \_\_\_\_\_

### CASE 138 MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard

2020 revision (series since 1968)

1. Fitts, P. M., & Jones, R. E. (1947), “Analysis of factors contributing to 460 ‘pilot error’ experiences in operating aircraft controls” — origin of designed-error analysis. — Retrieved from: \_\_\_\_\_

2. Department of Defense (2012), MIL-STD-1472G — the prior revision; revision history documents the 1968 origin and intermediate letters. — Retrieved from: \_\_\_\_\_
3. Chapanis, A. (1965), “Man-Machine Engineering” — foundational text for the discipline the standard codifies. — Retrieved from: \_\_\_\_\_
4. Department of Defense (2020), MIL-STD-1472H “Department of Defense Design Criteria Standard: Human Engineering,” 15 September 2020 — replaces MIL-STD-1472G (2012). — Retrieved from: \_\_\_\_\_

### CASE 139 GIFT and the Adoption Gap

2012 – present

1. K. VanLehn, “The Relative Effectiveness of Human Tutoring, Intelligent Tutoring Systems, and Other Tutoring Systems,” *Educational Psychologist* 46(4) (2011) — tutoring effectiveness. — Retrieved from: \_\_\_\_\_
2. GIFT Project, [gifttutoring.org](http://gifttutoring.org) — the framework, releases, and development under ARL / DEVCOM. — Retrieved from: \_\_\_\_\_
3. Editors’ synthesis of the GIFT adoption record — active development but limited ubiquitous fielding (quoted). — Retrieved from: \_\_\_\_\_
4. R. Sottolare, A. Graesser, X. Hu & H. Holden (eds.), *Design Recommendations for Intelligent Tutoring Systems* (U.S. Army Research Laboratory, 2013– ); IJAIED Special Issue on GIFT (2017). — Retrieved from: \_\_\_\_\_
5. J. Goodell & J. Kolodner, *Learning Engineering Toolkit* (2022) — adoption as an engineering problem. — Retrieved from: \_\_\_\_\_

### CASE 140 Navy Surface Warfare Readiness Reform

2018 – present

1. Navy and NTSB reports on the Fitzgerald and McCain collisions (2017–2019) — the training-degradation antecedent. — Retrieved from: \_\_\_\_\_
2. Readiness Reform Oversight Council, *One-Year Report* (2019) — restored training, assessments, and gates. — Retrieved from: \_\_\_\_\_
3. USNI News reform coverage (2020, 2022) — Ready-for-Sea Assessments (15 of 18 FDNF–Japan ships assessed ready for sea, the remaining three sidelined) and the surface community’s accounting of ten milestone assessments with four go/no-go gates. — Retrieved from: \_\_\_\_\_
4. GAO–20–154, *Navy Readiness: Actions Needed to Evaluate the Effectiveness of Changes to Surface Warfare Officer Training* (November 2019) — the absent evaluation processes and the planned threefold increase in ship-driving training hours. — Retrieved from: \_\_\_\_\_
5. SWOS Norfolk and San Diego Mariner Skills Training Center documentation — simulators and curriculum. — Retrieved from: \_\_\_\_\_

### CASE 141 DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks

2009 – 2014

1. Defense Advanced Research Projects Agency, *Digital Tutor program documentation* — program description and design rationale. — Retrieved from: \_\_\_\_\_
2. Fletcher, J. D. (2009). *From behaviorism to constructivism: a philosophical journey from drill and practice to situated learning*. — methodological grounding for the Digital Tutor’s tutorial discipline. — Retrieved from: \_\_\_\_\_

3. Morrison, J. E., & Fletcher, J. D. (September 2012). DARPA Digital Tutor: Assessment Data. IDA Document D-4686 (prepared for DARPA under contract DASW01-04-C-0003) — the sponsor-commissioned evaluation the case rests on. — Retrieved from: \_\_\_\_\_
4. Anderson, J. R., Corbett, A. T., Koedinger, K. R., & Pelletier, R. (1995). Cognitive tutors: Lessons learned. *Journal of the Learning Sciences*, 4(2):167–207. doi:10.1207/s15327809jls0402\_2 — the broader intelligent-tutoring evidence base the Digital Tutor program sits within. — Retrieved from: \_\_\_\_\_

## CASE 142 xAPI / Total Learning Architecture — Interoperability Gap

ongoing

1. Advanced Distributed Learning Initiative, Total Learning Architecture documentation — the cross-boundary vision. — Retrieved from: \_\_\_\_\_
2. IEEE Std 9274.1.1–2023, Standard for Learning Technology — JSON Data Model Format and RESTful Web Service for Learner Experience Data Tracking and Access (xAPI 2.0) — the technical standard of record; the IEEE LTSC project was authorized in May 2018, board-approved 30 March 2023, and published 6 October 2023. — Retrieved from: \_\_\_\_\_
3. B. Smith and L. Milham, Total Learning Architecture (TLA) Data Pillars and Their Applicability to Adaptive Instructional Systems (ADL Initiative, 2021; ERIC ED628093) — ADL’s own observation that while most LRS solutions offer dashboards, “they are commonly used to view xAPI statements that have been generated within a single learning activity,” against a design intent of data shared “vertically within an organization, horizontally across other organizations, or longitudinally throughout a career.” — Retrieved from: \_\_\_\_\_
4. ADL Initiative, Total Learning Architecture Standards: Digital Learning Acquisition Techniques (December 2023) — xAPI data that does not follow xAPI Profiles “will have interoperability issues outside of the implementing organization.” — Retrieved from: \_\_\_\_\_
5. Cf. inBloom (Case 53) — technology in advance of governance. — Retrieved from: \_\_\_\_\_

## CASE 143 Knight Capital Trading Loss

2012

1. U.S. Securities and Exchange Commission, Order Instituting Administrative and Cease-and-Desist Proceedings, In re Knight Capital Americas LLC (2013) — the firm, the Retail Liquidity Program launch, and the deployment. — Retrieved from: \_\_\_\_\_
2. SEC Order (2013) and Knight Capital Group 8-K filings of 2 and 6 August 2012 — the missed eighth server, the reactivation of the dormant “Power Peg” code, the 45-minute event, and the \$400 million rescue financing (preferred convertible at \$1.50 a share); the SEC puts the loss on the positions at over \$460 million, Knight’s own filing at a realized pre-tax loss of \$440 million. — Retrieved from: \_\_\_\_\_
3. SEC Order (2013), Rel. 34-70694, 16 October 2013 — the absence of a second-technician deployment review, the missing written code-development, deployment and unused-code procedures (§26), the capital thresholds unlinked to automated controls, and the \$12 million penalty (quoted). — Retrieved from: \_\_\_\_\_
4. B. Beyer, C. Jones, J. Petoff & N. R. Murphy (eds.), *Site Reliability Engineering* (O’Reilly, 2016) — deployment verification, dead-code removal, and automated safeguards as engineering deliverables. — Retrieved from: \_\_\_\_\_
5. SEC Market Access Rule (Rule 15c3-5) and subsequent automated-controls guidance — the regulatory response on pre-trade risk and market-access controls. — Retrieved from: \_\_\_\_\_
6. D. Seven, “Knightmare: A DevOps Cautionary Tale” — a widely cited engineering post-mortem on the deployment process, the orphaned eighth server, and the reused feature flag. — Retrieved from: \_\_\_\_\_



**CASE 144 Kodak Digital Camera Stagnation**

1975–2012

1. S. Sasson, "We Had No Idea," Kodak corporate blog, Kodak: Plugged In (2007) — first-hand account of the 1975 prototype, the demonstration, and the internal reception. — Retrieved from: \_\_\_\_\_
2. M. Spector & D. Mattioli, "Kodak Teeters on the Brink," The Wall Street Journal (5 January 2012); follow-up coverage of the patent-sale process in WSJ and Reuters (December 2012). — Retrieved from: \_\_\_\_\_
3. H. C. Lucas Jr. & J. M. Goh, "Disruptive technology: How Kodak missed the digital photography revolution," Journal of Strategic Information Systems 18(1):46–55 (March 2009). — Retrieved from: \_\_\_\_\_
4. Eastman Kodak Company, Voluntary Petition for Chapter 11 Reorganization, U.S. Bankruptcy Court, Southern District of New York, Case No. 12–10202 (19 January 2012). — Retrieved from: \_\_\_\_\_
5. A. R. Sorkin & M. J. de la Merced, "Eastman Kodak Files for Bankruptcy," The New York Times DealBook (19 January 2012). — Retrieved from: \_\_\_\_\_

**CASE 146 Northeast Blackout**

2003

1. U.S.–Canada Power System Outage Task Force, Final Report on the August 14, 2003 Blackout in the United States and Canada (2004) — the tree contact, the silent alarm, and the cascade. — Retrieved from: \_\_\_\_\_
2. Task Force (2004) — 50 million people affected across eight states and Ontario; the minute-by-minute sequence. — Retrieved from: \_\_\_\_\_
3. Task Force (2004), Causes Groups 1–4 — "inadequate situational awareness at FirstEnergy," FE's failure to manage tree growth, and the reliability organizations' failure to provide effective real-time diagnostic support (quoted). — Retrieved from: \_\_\_\_\_
4. FERC Order No. 693, Mandatory Reliability Standards for the Bulk-Power System, 118 FERC ¶ 61,218 (16 March 2007) — 83 of 107 proposed standards approved as mandatory and enforceable, effective 18 June 2007. — Retrieved from: \_\_\_\_\_
5. M. R. Endsley (1995), situation-awareness theory — the human-factors frame for silent-automation failure. — Retrieved from: \_\_\_\_\_
6. North American Electric Reliability Council reports (2004) and the creation of the Electric Reliability Organization. — Retrieved from: \_\_\_\_\_

**CASE 147 Takata Airbag Inflators**

2008 – 2026

1. U.S. National Highway Traffic Safety Administration, Takata air-bag inflator recall coordination materials — the ammonium-nitrate-without-desiccant design and the propellant-degradation rupture mechanism. — Retrieved from: \_\_\_\_\_
2. NHTSA recall record and investigative reporting (Reuters, New York Times) — roughly 67 million inflators recalled in the United States across 19 automakers (over 100 million worldwide); the largest automotive recall in history; 28 U.S. deaths confirmed by NHTSA and at least 400 U.S. injuries alleged. — Retrieved from: \_\_\_\_\_
3. U.S. DOJ (2017) and Takata internal documents released in litigation — the sustained misrepresentation of inflator test data to automakers and regulators. — Retrieved from: \_\_\_\_\_
4. U.S. Department of Justice, settlement and guilty plea with Takata Corporation (2017) — wire fraud, \$1B in fine and restitution, and Takata's subsequent bankruptcy. — Retrieved from: \_\_\_\_\_



5. NHTSA consumer alert on FCA's "Do Not Drive" warning for all vehicles with unrepaired Takata recalls (February 2026, 225,000 vehicles); CARFAX estimate of 4.8 million unrepaired Takata recalls as of July 2026; Honda's report that over 99% of the original "Alpha" inflators have been replaced or accounted for — the years-long completion of replacements and the continuing risk in unrepaired vehicles. — Retrieved from: \_\_\_\_\_
6. NHTSA Coordinated Remedy Program for the Takata recalls — the regulator's move to actively prioritize and manage replacement across nineteen automakers rather than leave pace to each manufacturer. — Retrieved from: \_\_\_\_\_

## CASE 148 Wells Fargo Fake Accounts

2011 – 2016

1. Consumer Financial Protection Bureau, Consent Order against Wells Fargo (2016) — the unauthorized accounts and penalties. — Retrieved from: \_\_\_\_\_
2. Office of the Comptroller of the Currency, Consent Order AA-EC-2016-66 (2016) — unsafe or unsound sales practices tied to the incentive-compensation structure (paraphrased). — Retrieved from: \_\_\_\_\_
3. Independent Directors of Wells Fargo, Sales Practices Investigation Report (2017) — how the sales-target architecture drove the conduct. — Retrieved from: \_\_\_\_\_
4. Enforcement record: 3.5 million accounts, \$3 billion in penalties, and the CEO's resignation. — Retrieved from: \_\_\_\_\_
5. A. C. Edmondson, *The Fearless Organization* (2018); incentive-design and corporate-governance literature. — Retrieved from: \_\_\_\_\_
6. Federal Reserve asset cap on Wells Fargo (imposed February 2018; lifted 3 June 2025) — the structural growth constraint and its termination after remediation. — Retrieved from: \_\_\_\_\_

## CASE 149 Volkswagen Dieselgate

2015

1. Volkswagen internal documents released through litigation — authorization of the defeat device within the engineering hierarchy. — Retrieved from: \_\_\_\_\_
2. U.S. EPA, Notice of Violation to Volkswagen (2015) — the defeat device and emissions exceedances. — Retrieved from: \_\_\_\_\_
3. West Virginia University / ICCT real-world diesel-emissions study (2014) — the discovery comparing road to lab. — Retrieved from: \_\_\_\_\_
4. U.S. Department of Justice Plea Agreement with Volkswagen AG (January 2017) — the institutional decision, USD 4.3 billion criminal-and-civil settlement, and convictions (quoted). — Retrieved from: \_\_\_\_\_
5. J. Ewing, *Faster, Higher, Farther: The Volkswagen Scandal* (2017); cf. Takata (Case 147). — Retrieved from: \_\_\_\_\_
6. EU Real Driving Emissions (RDE) testing regulation — the post-Dieselgate measurement reform. — Retrieved from: \_\_\_\_\_
7. Associated Press reporting on the German criminal proceedings — the Braunschweig regional court's conviction of four former VW managers on 26 May 2025 and its provisional termination of the Winterkorn trial in July 2025 on health grounds. — Retrieved from: \_\_\_\_\_

**CASE 151 GM Ignition Switch**

2002 – 2014

1. A. R. Valukas, Report to the Board of Directors of General Motors Company Regarding Ignition Switch Recalls (Jenner & Block, 2014) — the 2002 identification of the defect and the switch-torque mechanism. — Retrieved from: \_\_\_\_\_
2. U.S. NHTSA investigation of the GM ignition switch (2014) and the GM recall record — 2.6 million vehicles; 124 deaths via the GM compensation fund. — Retrieved from: \_\_\_\_\_
3. Valukas Report (2014) — the unchanged part number, the “GM nod” and “GM salute,” and the failure to use established escalation processes (quoted). — Retrieved from: \_\_\_\_\_
4. U.S. Department of Justice, deferred-prosecution agreement with General Motors (2015) — \$900 million forfeiture for concealing the defect; total penalties/settlements exceeding \$2.6 billion. — Retrieved from: \_\_\_\_\_
5. Valukas Report (2014) and A. C. Edmondson, *The Fearless Organization* (Wiley, 2018) — the part number as an organizational signal, and the suppression of upward safety information. — Retrieved from: \_\_\_\_\_
6. U.S. NHTSA consent order with General Motors (2014, \$35M civil penalty) and the GM ignition-switch victims’ compensation program (K. Feinberg, administrator) — the regulatory penalty and the basis for the 124-death figure. — Retrieved from: \_\_\_\_\_

**CASE 152 TSB Bank IT Migration**

2018

1. Slaughter and May, Independent Review of the TSB Migration (2019) — the single-weekend cutover and the testing failures. — Retrieved from: \_\_\_\_\_
2. Slaughter and May (2019) and FCA materials — 1.9 million customers locked out, £330M+ in costs, and the CEO’s resignation. — Retrieved from: \_\_\_\_\_
3. Slaughter and May (2019) — inadequate load testing and an unchallenged readiness certification. — Retrieved from: \_\_\_\_\_
4. Financial Conduct Authority, Final Notice on TSB Bank plc (20 December 2022) — the £29.75m penalty (with the PRA’s £18.9m) for planning, testing, risk-management and outsourcing failings. — Retrieved from: \_\_\_\_\_
5. House of Commons Treasury Committee, report on the TSB IT migration (2018). — Retrieved from: \_\_\_\_\_
6. Cf. Healthcare.gov (Case 180) and the migration-safety literature. — Retrieved from: \_\_\_\_\_

**CASE 153 Equifax Data Breach**

2017

1. U.S. Senate Permanent Subcommittee on Investigations, *How Equifax Neglected Cybersecurity and Suffered a Devastating Data Breach* (2019) — the unpatched Apache Struts vulnerability. — Retrieved from: \_\_\_\_\_
2. The breach record — 147 million affected, exploitation from May 2017, public disclosure in September 2017. — Retrieved from: \_\_\_\_\_
3. Senate PSI (2019) — “Equifax lacked a comprehensive IT asset inventory” (quoted). — Retrieved from: \_\_\_\_\_

4. U.S. FTC / CFPB / state settlement (at least \$575 million, up to \$700 million, 2019) and the CEO's resignation.  
— Retrieved from: \_\_\_\_\_
5. U.S. GAO, Actions Taken by Equifax and Federal Agencies in Response to the 2017 Breach, GAO-18-559 (2018).  
— Retrieved from: \_\_\_\_\_
6. Apache Struts CVE-2017-5638 advisory. — Retrieved from: \_\_\_\_\_

## CASE 154 **INL Turbine-Control Upgrade — Verification and Validation as the Cutover Deliverable** 2014

1. Ronald L. Boring (2014), "Human Factors Design, Verification, and Validation for Two Types of Control Room Upgrades at a Nuclear Power Plant," Proceedings of the Human Factors and Ergonomics Society Annual Meeting 58(1), pp. 2295–2299, doi:10.1177/1541931214581478 — the NUREG-0711 process for the paired plant-process-computer and turbine-control-system upgrade. Step 4 of the paper's implementation sequence is the source of the training claim verbatim: "A stand-alone training program will be developed in cooperation with the training organization. In addition, the scenarios will be reviewed by trainers and procedure writers to ensure that the operating procedures do not require modifications." The future tense is the paper's own.  
— Retrieved from: \_\_\_\_\_
2. Idaho National Laboratory, Light Water Reactor Sustainability Program reports on control-room modernization — series available via OSTI. — Retrieved from: \_\_\_\_\_
3. Nuclear Regulatory Commission (NUREG-0711), "Human Factors Engineering Program Review Model" — the regulatory framework the V&V deliverables are produced against. — Retrieved from: \_\_\_\_\_
4. O'Hara et al. (2008), "Human Factors Considerations with Respect to Emerging Technology in Nuclear Power Plants," NUREG/CR-6947 — peer-adjacent framing. — Retrieved from: \_\_\_\_\_

## CASE 155 **Toyota Production System / Andon Cord** 1950s – present

1. Liker, J. (2020), The Toyota Way (2nd ed.) — the andon cord, the cultural pairing, and the American imitation (paraphrased). — Retrieved from: \_\_\_\_\_
2. Shingo, S. (1989), A Study of the Toyota Production System — the technical mechanism. — Retrieved from: \_\_\_\_\_
3. Spear, S. & Bowen, H. (1999), "Decoding the DNA of the Toyota Production System," HBR — the response protocol and embedded learning. — Retrieved from: \_\_\_\_\_
4. Adler, P. (1993), "Time-and-Motion Regained," HBR 71(1) — NUMMI observed from the floor; the rehired workforce, the team structure, and the standardized-work practice that changed alongside the signalling system. — Retrieved from: \_\_\_\_\_
5. Womack & Jones (1996), Lean Thinking — diffusion and the limits of surface imitation. — Retrieved from: \_\_\_\_\_
6. Rother, M. (2009), Toyota Kata — the routines that sustain the practice. — Retrieved from: \_\_\_\_\_
7. NASA Engineering and Safety Center / NHTSA (2011), Technical Assessment of Toyota Electronic Throttle Control Systems — the unintended-acceleration review underlying the scope limit. — Retrieved from: \_\_\_\_\_

## CASE 156 **INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet** 2010 – present

1. O'Hara, Higgins, Brown, Fink, Persensky, Lewis, Kramer, & Szabo (2008), "Human Factors Considerations with Respect to Emerging Technology in Nuclear Power Plants," NUREG/CR-6947 — foundational human-factors backdrop. — Retrieved from: \_\_\_\_\_
2. Nuclear Regulatory Commission, Branch Technical Position 7-19, "Guidance for Evaluation of Diversity and Defense-in-Depth in Digital Computer-Based Instrumentation and Control Systems." — Retrieved from: \_\_\_\_\_
3. Nuclear Regulatory Commission, Regulatory Guide 1.180, "Guidelines for Evaluating Electromagnetic and Radio-Frequency Interference in Safety-Related Instrumentation and Control Systems." — Retrieved from: \_\_\_\_\_
4. Idaho National Laboratory, Light Water Reactor Sustainability (LWRS) program annual reports (2010 – present) — primary research-and-pilot record. — Retrieved from: \_\_\_\_\_
5. Electric Power Research Institute, control-room modernization technical reports — industry-side sustainment record. — Retrieved from: \_\_\_\_\_

## CASE 157 **AI-Augmented Coding Tools** 2021 – present

1. Peng et al. (2023), "The Impact of AI on Developer Productivity" — short-term productivity gains. — Retrieved from: \_\_\_\_\_
2. Pearce et al. (2022), "Asleep at the Keyboard? Assessing the Security of GitHub Copilot's Code Contributions." — Retrieved from: \_\_\_\_\_
3. L. Bainbridge, "Ironies of Automation," *Automatica* 19(6) (1983) — the classic deskilling and over-reliance problem, applied to AI-augmented work. — Retrieved from: \_\_\_\_\_
4. Dell'Acqua et al. (2023), "Navigating the Jagged Technological Frontier" (HBS / BCG) — professional LLM use. — Retrieved from: \_\_\_\_\_
5. METR (2025), "Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity," arXiv:2507.09089 — the randomized trial finding a 19% slowdown against a self-reported speedup. — Retrieved from: \_\_\_\_\_
6. Sandoval et al. (2023), "Lost at C," *USENIX Security* — AI assistance did not significantly increase critical security-bug rates. — Retrieved from: \_\_\_\_\_

## CASE 159 **Bhopal** 1984

1. Union Carbide and government investigation reports (1985) — MIC storage, the disabled safety systems, and plant understaffing. — Retrieved from: \_\_\_\_\_
2. Accounts of the 2–3 Dec. 1984 release — the contested toll (thousands of immediate deaths; 15,000–20,000 total estimates; 500,000 exposed). (Figures vary widely across sources; see AUDIT.) — Retrieved from: \_\_\_\_\_
3. New York Times investigation (1985) — "operating errors, design flaws, maintenance failures, training deficiencies and economy measures that endangered safety" (quoted). — Retrieved from: \_\_\_\_\_
4. N. Meshkati, "Human Factors in Large-Scale Technological Systems' Accidents," *Industrial Crisis Quarterly* (1991) — Bhopal alongside Three Mile Island on neglect of the human element. G. Joseph, M. Kaszniak & L.

Long, "Lessons after Bhopal: CSB a catalyst for change," *Journal of Loss Prevention in the Process Industries* 18(4–6):537–548 (2005) — ineffective employee training an underlying cause in nine of the CSB's first twenty-three completed investigations. — Retrieved from: \_\_\_\_\_

5. P. Shrivastava, *Bhopal: Anatomy of a Crisis* (1992); C. Perrow, *Normal Accidents* (1984). — Retrieved from: \_\_\_\_\_
6. The U.S. Chemical Safety and Hazard Investigation Board — authorized by the Clean Air Act Amendments of 1990, operational from January 1998 — and the post-Bhopal process-safety regime (EPCRA 1986; OSHA Process Safety Management, 1992). — Retrieved from: \_\_\_\_\_

## CASE 160 Davis-Besse Reactor Head Corrosion

2002

1. U.S. NRC Office of Inspector General, *NRC's Regulation of Davis-Besse Regarding Damage to the Reactor Vessel Head* (Case No. 02-03S, December 2002) — the post-TMI regulatory regime and oversight failures. — Retrieved from: \_\_\_\_\_
2. NRC event records and OIG (2002) — the 6-inch corrosion cavity, the remaining cladding, and the 2,200-psi coolant margin. — Retrieved from: \_\_\_\_\_
3. NRC Bulletin 2001-01, "Circumferential Cracking of Reactor Pressure Vessel Head Penetration Nozzles" (3 August 2001), and the compromise that let Davis-Besse run to a 16 February 2002 outage. — Retrieved from: \_\_\_\_\_
4. NRC OIG (2002), as quoted in U.S. GAO, GAO-04-415 (2004) — the finding that NRC's decision "was driven in large part by a desire to lessen the financial impact on FirstEnergy," and NRC's dispute of it. — Retrieved from: \_\_\_\_\_
5. D. Lochbaum / Union of Concerned Scientists analysis (2002); U.S. GAO, GAO-04-415 (2004). — Retrieved from: \_\_\_\_\_
6. J. V. Rees, *Hostages of Each Other* (1994) — INPO and the fragility of institutional safety capability. — Retrieved from: \_\_\_\_\_

## CASE 161 Texas City BP Refinery Explosion

2005

1. CSB (2007) — accumulated safety-culture decay, deferred maintenance, and the siting of occupied trailers beside a hazardous unit. — Retrieved from: \_\_\_\_\_
2. U.S. Chemical Safety and Hazard Investigation Board, *Refinery Explosion and Fire, Investigation Report 2005-04-I-TX* (2007) — the startup, malfunctioning instruments, and 15 killed / 180 injured. — Retrieved from: \_\_\_\_\_
3. A. Hopkins, *Failure to Learn: The BP Texas City Refinery Disaster* (CCH, 2008). — Retrieved from: \_\_\_\_\_
4. CSB (2007) — "indicators of personal safety are not indicators of process safety" (quoted). — Retrieved from: \_\_\_\_\_
5. OSHA Process Safety Management standard (29 CFR 1910.119, 1992) and the CCPS process-safety literature — the personal-vs-process-safety distinction. — Retrieved from: \_\_\_\_\_
6. Report of the BP U.S. Refineries Independent Safety Review Panel (the Baker Panel, 2007). — Retrieved from: \_\_\_\_\_

**CASE 162 Upper Big Branch Mine Explosion**

2010

1. U.S. MSHA, Internal Review of MSHA's Actions at Upper Big Branch and the accident investigation (2011–2012) — the dual records and the 29 deaths. — Retrieved from: \_\_\_\_\_
2. Governor's Independent Investigation Panel (J. Davitt McAteer, 2011) — mine conditions: methane, ventilation, and coal-dust inerting. — Retrieved from: \_\_\_\_\_
3. MSHA and U.S. DOJ findings — suppressed methane readings, disabled monitors, and falsified pre-shift examinations. — Retrieved from: \_\_\_\_\_
4. United States v. Blankenship (S.D.W. Va., 2015–2016) — the misdemeanor conviction and felony acquittals. — Retrieved from: \_\_\_\_\_
5. H. Berkes / NPR investigative reporting on Massey Energy. — Retrieved from: \_\_\_\_\_
6. A. Hopkins, Disastrous Decisions: The Human and Organisational Causes of the Gulf of Mexico Blowout (2012) — corporate decision-making and safety (comparative). — Retrieved from: \_\_\_\_\_

**CASE 163 Deepwater Horizon**

2010

1. Spill-volume estimates — the Flow Rate Technical Group's 4.9 million barrels against the 4.0 million barrels found by the district court in *In re Oil Spill by the Oil Rig "Deepwater Horizon"* (E.D. La., January 2015), less 810,000 barrels collected — and BP cost disclosures (>\$65 billion); the reorganization of the Minerals Management Service into BSEE/BOEM (Secretarial Order 3299, 2010); A. Lustgarten, *Run to Failure: BP and the Making of the Deepwater Horizon Disaster* (W.W. Norton, 2012) — long-arc account of accumulated procedural debt. — Retrieved from: \_\_\_\_\_
2. National Commission (2011) — the well-control sequence and mud-displacement decision. — Retrieved from: \_\_\_\_\_
3. National Commission (2011) — "systematic failures in risk management... place in doubt the safety culture of the entire industry" (quoted); U.S. Chemical Safety Board, *Drilling Rig Explosion and Fire at the Macondo Well* (final volumes, 2014–2016) — process-safety vs. personal-safety distinction and the improving personal-injury rates BP and Transocean reported as misleading indicators. — Retrieved from: \_\_\_\_\_
4. BOEMRE / U.S. Coast Guard Joint Investigation (2011) — blowout-preventer maintenance backlog and chain-of-command findings; National Academies, *Macondo Well Deepwater Horizon Blowout: Lessons for Improving Offshore Drilling Safety* (2012) — training gaps and the interlinked cascade of failed defenses; BP, *Deepwater Horizon Accident Investigation Report* (the Bly report, 8 September 2010) — BP's own summation, "a complex and interlinked series of mechanical failures, human judgments, engineering design, operational implementation and team interfaces" (quoted). — Retrieved from: \_\_\_\_\_
5. Deepwater Horizon Study Group (UC Berkeley, 2011) final report; N. Leveson, systems-theoretic analysis of Deepwater Horizon — the multi-layer drift and the limits of single-cause framings. — Retrieved from: \_\_\_\_\_
6. National Commission on the BP Deepwater Horizon Oil Spill, *Deep Water: The Gulf Oil Disaster and the Future of Offshore Drilling* (Report to the President, 2011) — systematic management failures as root cause; the misread negative-pressure test. — Retrieved from: \_\_\_\_\_

**CASE 164 Grenfell Tower**

2017

1. Grenfell Tower Inquiry, Phase 2 Report (2024) — decades of failure and the combustible-cladding decision. — Retrieved from: \_\_\_\_\_

2. Grenfell Tower Inquiry, Phase 1 Report (2019) — the fire’s spread up the cladding and the failure of “stay put.” — Retrieved from: \_\_\_\_\_
3. Phase 2 Report (2024) — cladding firms’ “systematic dishonesty” and the RBKC building control failures (ch. 62). — Retrieved from: \_\_\_\_\_
4. Phase 1 Report (2019) — London Fire Brigade unpreparedness; preparation and planning “was gravely inadequate” (quoted). — Retrieved from: \_\_\_\_\_
5. B. Hutter & M. Power (eds.), *Organizational Encounters with Risk* (2005) — distributed risk ownership. — Retrieved from: \_\_\_\_\_
6. UK Government response to the Grenfell Phase 2 report (2025) — building-safety and fire-service reform. — Retrieved from: \_\_\_\_\_

## CASE 165 Camp Fire / PG&E

2018

1. CalFire, Camp Fire Investigation Report (2019) — the worn transmission-line hardware as ignition source. — Retrieved from: \_\_\_\_\_
2. The Camp Fire record — 85 killed; PG&E’s guilty plea to 84 counts of involuntary manslaughter (2020). — Retrieved from: \_\_\_\_\_
3. Butte County District Attorney, The Camp Fire Public Report (2020) — PG&E’s knowledge and deferred maintenance (quoted). — Retrieved from: \_\_\_\_\_
4. California Public Utilities Commission, Order Instituting Investigation into PG&E (2019) — the regulatory dimension. — Retrieved from: \_\_\_\_\_
5. PG&E bankruptcy and California wildfire-mitigation reforms (2019– ). — Retrieved from: \_\_\_\_\_
6. Cf. the Northeast Blackout (Case 146); climate-and-infrastructure literature. — Retrieved from: \_\_\_\_\_

## CASE 166 Three Mile Island

1979

1. Kemeny Commission, Report of the President’s Commission on the Accident at Three Mile Island (Oct. 1979) — operator training oriented to large design-basis accidents. — Retrieved from: \_\_\_\_\_
2. M. Rogovin & G. T. Frampton, *Three Mile Island: A Report to the Commissioners and to the Public*, NUREG/CR-1250 (U.S. NRC, 1980) — the September 1977 Davis-Besse stuck-PORV precursor and the failure to disseminate its lessons. — Retrieved from: \_\_\_\_\_
3. U.S. Nuclear Regulatory Commission, Backgrounder on the Three Mile Island Accident — accident sequence, misleading PORV indication, throttled high-pressure injection, and the off-site dose — about 1 millirem average, under 100 millirem at the site boundary. — Retrieved from: \_\_\_\_\_
4. Kemeny Commission (1979) — central conclusion that the fundamental problems were people-related, not equipment-related (quoted). — Retrieved from: \_\_\_\_\_
5. J. V. Rees, *Hostages of Each Other: The Transformation of Nuclear Safety since Three Mile Island* (Univ. of Chicago Press, 1994) — TMI as the origin of systematic capability reform. — Retrieved from: \_\_\_\_\_
6. C. Perrow, *Normal Accidents* (1984); J. Reason, *Human Error* (1990) — why ambiguous cascades defeat design-basis training and command-not-state interfaces. — Retrieved from: \_\_\_\_\_

7. Institute of Nuclear Power Operations, "Our History" — INPO established December 1979 to set standards and force sharing of operating experience; see also NEI, "Lessons from the 1979 Accident at Three Mile Island." — Retrieved from: \_\_\_\_\_
8. U.S. Nuclear Regulatory Commission, Regulatory Guide 1.97, Rev. 2, Instrumentation for Light-Water-Cooled Nuclear Power Plants to Assess Plant and Environs Conditions During and Following an Accident (Dec. 1980), and NUREG-0700, Human-System Interface Design Review Guidelines — the control-room and post-accident-instrumentation standards that grew out of TMI; see also J. S. Walker, *Three Mile Island: A Nuclear Crisis in Historical Perspective* (Univ. of California Press, 2004) on the operator-training-and-licensing overhaul. — Retrieved from: \_\_\_\_\_

## CASE 167 Hurricane Katrina — The FEMA/DHS Response Failure

2005

1. Select Bipartisan Committee to Investigate the Preparation for and Response to Hurricane Katrina (2006), *A Failure of Initiative*, H. Rept. 109-377, U.S. House of Representatives, February 15, 2006 — including the dedicated Hurricane Pam findings chapter. — Retrieved from: \_\_\_\_\_
2. U.S. Senate Committee on Homeland Security and Governmental Affairs (2006), *Hurricane Katrina: A Nation Still Unprepared*, S. Rept. 109-322, May 2006 — 325+ interviews, 838,000+ pages of documents reviewed. — Retrieved from: \_\_\_\_\_
3. The White House (2006), *The Federal Response to Hurricane Katrina: Lessons Learned* (the Townsend report), February 23, 2006 — 17 lessons, 125 recommendations, including the situational-awareness finding. — Retrieved from: \_\_\_\_\_
4. Post-Katrina Emergency Management Reform Act of 2006, Title VI of Public Law 109-295, signed October 4, 2006 — rebuilt FEMA authorities, accelerated pre-request federal assistance, national preparedness goal and system. — Retrieved from: \_\_\_\_\_
5. U.S. Senate Committee on Homeland Security and Governmental Affairs (2006), *Preparing for a Catastrophe: The Hurricane Pam Exercise*, hearing, S. Hrg. 109-403, January 24, 2006 — testimony on the July 2004 exercise scope and unfinished follow-on work. — Retrieved from: \_\_\_\_\_

## CASE 169 Fukushima Daiichi

2011

1. National Diet of Japan Fukushima Nuclear Accident Independent Investigation Commission (NAIIC; K. Kurokawa, chair), Report (2012) — the internal tsunami evidence and the regulatory-capture finding. — Retrieved from: \_\_\_\_\_
2. NAIIC (2012) — the accident sequence: seawall overtopping, generator inundation, three core meltdowns, and "approximately 150,000 people were evacuated in response to the accident." The widely quoted 160,000 is the 2012 peak evacuee count in prefectural records, not NAIIC's figure; do not "correct" the text against it. Cleanup and compensation costs are from the Japanese government's December 2016 estimate of ¥21.5 trillion, not from NAIIC. — Retrieved from: \_\_\_\_\_
3. NAIIC (2012) — the "made in Japan" cultural and regulatory-capture conclusion (quoted). — Retrieved from: \_\_\_\_\_
4. Investigation Committee on the Accident (Hatamura government commission, 2012); IAEA Director General, *The Fukushima Daiichi Accident* (2015) — external-hazard under-estimation. — Retrieved from: \_\_\_\_\_
5. D. Lochbaum, E. Lyman & S. Stranahan, *Fukushima: The Story of a Nuclear Disaster* (2014). — Retrieved from: \_\_\_\_\_
6. Y. Funabashi & K. Kitazawa, *Fukushima in Review* (2012); cf. INPO (Case 175) and Davis-Besse (Case 160). — Retrieved from: \_\_\_\_\_



**CASE 170 West Africa Ebola — The Delayed International Response**

2014 – 2016

1. Médecins Sans Frontières (2015), Ebola: Pushed to the Limit and Beyond — MSF account of the international outbreak response, March 2015: the “out of control” declaration of 21 June 2014 and the reaction of WHO staff in Guinea and Sierra Leone. WHO declared the PHEIC on 8 August 2014, the Emergency Committee having met 6–7 August; as of 4 August, 1,711 cases and 932 deaths had been reported. — Retrieved from: \_\_\_\_\_
2. Ebola Interim Assessment Panel (2015), Report of the Ebola Interim Assessment Panel (the “Stocking Report,” chair Dame Barbara Stocking), World Health Organization, 7 July 2015 — the November 2014 IHR core-capacity figures (64 / 81 / 48), the “not acceptable” and “unacceptable” findings (quoted), and the Panel’s three review headings. — Retrieved from: \_\_\_\_\_
3. Moon, S., Sridhar, D., Pate, M. A., Jha, A. K., Clinton, C., et al. (2015), “Will Ebola change the game? Ten essential reforms before the next pandemic. The report of the Harvard-LSHTM Independent Panel on the Global Response to Ebola,” *The Lancet* 386(10009):2204–2221, 28 November 2015, doi:10.1016/S0140-6736(15)00946-0 — the “ample reason” finding (quoted) and the ten reform proposals. — Retrieved from: \_\_\_\_\_
4. World Health Organization (2016), Ebola Situation Reports — West Africa: approximately 28,600 cases and 11,300 deaths reported to 10 June 2016 (reported figures; undercount acknowledged). — Retrieved from: \_\_\_\_\_
5. High-level Panel on the Global Response to Health Crises (2016), Protecting Humanity from Future Health Crises (chair Jakaya M. Kikwete), UN doc. A/70/723, 9 February 2016 — twenty-seven recommendations, including political-leadership structures above the technical agencies. — Retrieved from: \_\_\_\_\_

**CASE 171 Hurricane Maria — Puerto Rico Response and Logistics Failure**

2017 – 2018

1. Houser, T., & Marsters, P. (2018), “The World’s Second Largest Blackout,” Rhodium Group, April 12, 2018 — approximately 3.4 billion lost customer-hours, “three times larger than the next biggest blackout in US history”; the largest blackout in U.S. history by that metric. — Retrieved from: \_\_\_\_\_
2. Federal Emergency Management Agency (2018), 2017 Hurricane Season FEMA After-Action Report, July 12, 2018 — the agency’s own self-assessment: force strength below target, Caribbean Distribution Center depletion, commodity-tracking shortfalls. — Retrieved from: \_\_\_\_\_
3. U.S. Government Accountability Office (2018), 2017 Hurricanes and Wildfires: Initial Observations on the Federal Response and Key Recovery Challenges, GAO-18-472, September 4, 2018 — “at the height of FEMA workforce deployments in October 2017, 54 percent of staff were serving in a capacity in which they did not hold the title of ‘Qualified’”; close to 14,000 federal employees deployed. — Retrieved from: \_\_\_\_\_
4. Kishore, N., Marqués, D., Mahmud, A., et al. (2018), “Mortality in Puerto Rico after Hurricane Maria,” *New England Journal of Medicine* 379(2):162 – 170, doi:10.1056/NEJMsa1803972 — household-survey estimate of 4,645 excess deaths (95% CI 793 – 8,498), September 20 – December 31, 2017. — Retrieved from: \_\_\_\_\_
5. Milken Institute School of Public Health, George Washington University (2018), Ascertainment of the Estimated Excess Mortality from Hurricane Maria in Puerto Rico, August 2018 — commissioned by the Government of Puerto Rico; 2,975 excess deaths, September 2017 – February 2018; basis for the official toll revision from 64. — Retrieved from: \_\_\_\_\_

**CASE 172 CrowdStrike Falcon Outage**

2024

1. CrowdStrike, Falcon Content Update Preliminary Post Incident Report (24 July 2024) — Rapid Response Content deploys as configuration data through Channel Files without the automated testing, manual validation and staged rollout that Sensor Content clears. — Retrieved from: \_\_\_\_\_
2. D. Weston, “Helping our customers through the CrowdStrike outage,” The Official Microsoft Blog (20 July 2024) — the 8.5 million affected Windows devices, “less than one percent of all Windows machines”; cf. the CrowdStrike PIR for the Channel File 291 out-of-bounds read and the kernel crash. — Retrieved from: \_\_\_\_\_
3. CrowdStrike PIR (24 July 2024) — the finding that the update did not clear the testing depth or staged rollout applied to sensor code; the staggered-deployment and canary commitment is quoted from this report. — Retrieved from: \_\_\_\_\_
4. U.S. Government Accountability Office, Cyber Resiliency: CrowdStrike Outage Highlights Challenges, GAO-24-107733 (23 September 2024); House Committee on Homeland Security, Subcommittee on Cybersecurity and Infrastructure Protection hearing, 24 September 2024 — supply-chain and concentration risk in endpoint security. — Retrieved from: \_\_\_\_\_
5. CrowdStrike PIR remediation commitments (2024); Microsoft, “Taking steps that drive resiliency and security for Windows customers,” Windows Experience Blog (12 September 2024), following the Windows Endpoint Security Ecosystem Summit of 10 September 2024 — new platform capabilities for security vendors rather than the elimination of kernel access. — Retrieved from: \_\_\_\_\_
6. Delta Air Lines v. CrowdStrike, Fulton County Superior Court, order on the motion to dismiss, 16 May 2025 (Judge Kelly Lee Ellerbe); the passenger class action against Delta, N.D. Ga., May 2025 (Judge Mark Cohen), allowed to proceed in part. — Retrieved from: \_\_\_\_\_
7. B. Beyer et al. (eds.), Site Reliability Engineering (2016) — staged rollout and canarying; cf. Knight Capital (Case 143). — Retrieved from: \_\_\_\_\_

**CASE 173 Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable**

1963 – present

1. NASA SMA (2006), “USS Thresher Lessons Learned” briefing — safety-message archive. — Retrieved from: \_\_\_\_\_
2. Rear Admiral Paul E. Sullivan, statement to House Science Committee (2003), NASA/Columbia archive — primary congressional record on SUBSAFE. — Retrieved from: \_\_\_\_\_
3. U.S. Navy NAVSEA, SUBSAFE program documentation — operating program publications. — Retrieved from: \_\_\_\_\_
4. Columbia Accident Investigation Board (2003), final report — Volume I, endorsement of SUBSAFE as a model for NASA. — Retrieved from: \_\_\_\_\_
5. MIT Press, “SUBSAFE: An Example of a Successful Safety Program” — book chapter (open access). — Retrieved from: \_\_\_\_\_

**CASE 174 Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed**

1996 – 2000

1. General Accounting Office, Year 2000 Computing Challenge report series (1997–2000), particularly GAO/AIMD-00-290, Lessons Learned Can Be Applied to Other Management Challenges (September 2000) — the

federal-program-management record and GAO's own retrospective; GAO designated Y2K a high-risk area in February 1997. — Retrieved from: \_\_\_\_\_

2. Office of Management and Budget, quarterly reports on federal Y2K remediation status (1997–1999) — program-self-report tier; OMB “directed the major departments and agencies to submit quarterly reports beginning May 15, 1997,” and reported 21 percent of their mission-critical systems compliant in May 1997 against 99.9 percent in December 1999 (figures as given in GAO/AIMD-00-290). — Retrieved from: \_\_\_\_\_
3. U.S. Department of Commerce, Economics and Statistics Administration, The Economics of Y2K and the Impact on the United States (November 1999) — total U.S. Y2K spending of about \$100 billion covering repairs started 1995 through 2001, of which the federal share was about \$8 billion. — Retrieved from: \_\_\_\_\_
4. President's Council on Year 2000 Conversion, The Journey to Y2K final report (2000) — institutional retrospective of the federal coordination effort; the Council was established by Executive Order 13073 in February 1998. — Retrieved from: \_\_\_\_\_
5. Manion & Evan (2000), “The Y2K Problem and Professional Responsibility: A Retrospective Analysis,” *Technology in Society* 22(3), 361–387, doi:10.1016/S0160-791X(00)00015-4 — retrospective literature on the counterfactual debate. — Retrieved from: \_\_\_\_\_

## CASE 175 INPO and the Nuclear Academy

1979 – present

1. Report of the President's Commission on the Accident at Three Mile Island (Kemeny Commission, 30 October 1979) — the pre-TMI culture, the disregarded September 1977 Davis-Besse precursor, and the recommendation that the industry “set and police its own standards of excellence.” — Retrieved from: \_\_\_\_\_
2. Willard, R. F. (2019), The Role of the Institute of Nuclear Power Operations in Supporting the United States Commercial Nuclear Power Industry's Focus on Nuclear Safety, testimony for the record, U.S. Senate Committee on Environment and Public Works, 13 November 2019 — INPO's design, funding, peer evaluations, and the formation of the National Academy for Nuclear Training in 1985. INPO's own history dates its formal establishment to December 1979, in response to the Commission's recommendations; the Kemeny report does not name INPO. — Retrieved from: \_\_\_\_\_
3. Rees, J. (1994), Hostages of Each Other: The Transformation of Nuclear Safety since Three Mile Island — INPO's design, the peer-evaluation mechanism, and the “hostages” premise (paraphrased). — Retrieved from: \_\_\_\_\_
4. World Nuclear Association — Three Mile Island Accident; INPO/WANO performance indicators. — Retrieved from: \_\_\_\_\_
5. Nuclear Energy Institute, “Lessons from the 1979 Accident at Three Mile Island”; National Academy for Nuclear Training — accreditation, and the founding of WANO after Chernobyl. — Retrieved from: \_\_\_\_\_

## CASE 176 Tylenol Recall

1982

1. Edmondson, A. (2018), The Fearless Organization — pre-committed values under stress. — Retrieved from: \_\_\_\_\_
2. Kaplan, T. (2014), The Tylenol Crisis — the recall and corporate response. — Retrieved from: \_\_\_\_\_
3. James Burke (J&J CEO), in Lasting Leadership (Wharton) — the Credo quote. — Retrieved from: \_\_\_\_\_

4. Greyser, S., Johnson & Johnson: The Tylenol Tragedy (HBS case, 1992) — market recovery and crisis management. — Retrieved from: \_\_\_\_\_

5. FDA Final Rule on Tamper-Resistant Packaging (1982) — the packaging reform. — Retrieved from: \_\_\_\_\_

### CASE 177 CIRAS — Confidential Incident Reporting for UK Rail

1996 – present

1. Davies, Wright, Courtney, & Reid (2000), "Confidential Incident Reporting on the UK Railways: The 'CIRAS' System," *Cognition, Technology & Work* 2, 117–125, doi:10.1007/PL00011494 — the interview-based method and the structural argument for a channel independent of the employer. — Retrieved from: \_\_\_\_\_
2. University of Strathclyde, CIRAS impact case study (REF) — the institutional history: the ScotRail pilot of 1995–97, the national mandate of 2000 after Ladbroke Grove, and "in 2008, CIRAS became an independently operated unit within the Rail Safety and Standards Board." Secondary accounts date the transfer to 2009, when the CIRAS Charitable Trust wound up its funding role; do not "correct" the 2008 date against them. — Retrieved from: \_\_\_\_\_
3. University of Strathclyde, CIRAS impact case study (REF) — the operating-program self-report: "between 2008 and 2012 CIRAS received 2228 reports; 45% of these resulted in tangible safety improvements and approximately 33% contained important information about safety that was new to the company concerned"; and the influence on the FRA's Confidential Close Call Reporting System (C3RS). — Retrieved from: \_\_\_\_\_
4. Rail Safety and Standards Board (RSSB), CIRAS program documentation 2008 – present — operating-program publications; the peer-reviewed literature describes the method and reports the program-supplied figures rather than independently validating them. — Retrieved from: \_\_\_\_\_
5. Ladbroke Grove Rail Inquiry (Cullen, 2001), final report — the public inquiry into the crash that preceded the national mandate. — Retrieved from: \_\_\_\_\_

### CASE 178 The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami

2005 – 2020

1. Adinolfi, Bassiouni, Lauritzen, & Williams (2005), Humanitarian Response Review, commissioned by the UN Emergency Relief Coordinator, New York/Geneva: OCHA, August 2005 — the diagnosis of uneven sectoral capacity and absent accountability across the international system. — Retrieved from: \_\_\_\_\_
2. Adinolfi et al. (2005), Humanitarian Response Review; Inter-Agency Standing Committee Principals' decision of 12 September 2005 — the cluster approach as one pillar of humanitarian reform: designated lead agencies, globally and per activated country response, accountable to the Humanitarian Coordinator and carrying provider-of-last-resort responsibility. Nine clusters were designated initially; the set later grew to eleven. — Retrieved from: \_\_\_\_\_
3. Street, A., & Parihar, G. (2007), "The UN cluster approach in the Pakistan earthquake response: an NGO perspective," *Humanitarian Exchange* 37, ODI Humanitarian Practice Network — the first field application: information-sharing and named sectoral accountability set against "cluster creep," meeting burdens and English-only coordination. — Retrieved from: \_\_\_\_\_
4. Stoddard, Harmer, Haver, Salomons, & Wheeler (2007), Cluster Approach Evaluation, Final Report, OCHA Evaluation and Studies Section, with NYU Center on International Cooperation, ODI Humanitarian Policy Group, and the Praxis Group — Phase 1: "stronger and more predictable leadership over sectors. However, there has been no observable increase in ultimate accountability." — Retrieved from: \_\_\_\_\_
5. Steets, Grünewald, Binder, de Geoffroy, Kauffmann, Krüger, Meier, & Sokpoh (2010), Cluster Approach Evaluation 2: Synthesis Report, GPPI / Groupe URD, commissioned by the IASC, Geneva: OCHA — Phase 2, six country

studies: the benefits “already slightly outweigh its costs and shortcomings,” against persistent marginalisation of national and local actors and weak inter-cluster coordination. — Retrieved from: \_\_\_\_\_

6. Inter-Agency Standing Committee (2011 – 2012), IASC Transformative Agenda — Principals’ decisions of December 2011 and associated protocols; Grünewald & Binder (2010), Inter-agency Real-Time Evaluation in Haiti: 3 Months After the Earthquake, Groupe URD / GPPI. — Retrieved from: \_\_\_\_\_

## CASE 179 **The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed** 2020 – 2023

1. Dong, Du, & Gardner (2020), “An interactive web-based dashboard to track COVID-19 in real time,” *The Lancet Infectious Diseases* 20(5), 533 – 534, doi:10.1016/S1473-3099(20)30120-1. — Retrieved from: \_\_\_\_\_
2. Dong et al. (2022), “The Johns Hopkins University Center for Systems Science and Engineering COVID-19 Dashboard: data collection process, challenges faced, and lessons learned,” *The Lancet Infectious Diseases* 22(12), e370 – e376, doi:10.1016/S1473-3099(22)00434-0. — Retrieved from: \_\_\_\_\_
3. Johns Hopkins University Hub (10 March 2023), “Johns Hopkins COVID-19 data hub ends after three years” — Coronavirus Resource Center closure announcement. — Retrieved from: \_\_\_\_\_
4. Lasker Foundation (2022), Lasker-Bloomberg Public Service Award citation — “The COVID-19 Dashboard” (Lauren Gardner); companion award essay in *JAMA*. — Retrieved from: \_\_\_\_\_
5. NPR (10 February 2023), “As the pandemic ebbs, an influential COVID tracker shuts down.” — Retrieved from: \_\_\_\_\_

## CASE 180 **Healthcare.gov Launch** 2013

1. HHS Office of Inspector General, *HealthCare.gov: Case Study of CMS Management of the Federal Marketplace*, OEI-06-14-00350 (February 2016) — the assembled team’s expertise, the unfilled technical positions, and the finding that CMS leadership “did not recognize the need to clearly outline responsibilities and delineate this role.” — Retrieved from: \_\_\_\_\_
2. U.S. GAO, *Healthcare.gov: CMS Has Taken Steps to Address Problems, but Needs to Further Implement Systems Development Best Practices*, GAO-15-238 (2015) — end-to-end testing did not occur prior to launch; CGI Federal, written testimony of Cheryl Campbell, House Committee on Energy and Commerce, 24 October 2013 — “CMS serves the important role of systems integrator or ‘quarterback’ on this project and is the ultimate responsible party for the end-to-end performance of the overall Federal Exchange.” The testimony names CMS, not HHS; do not “correct” it. — Retrieved from: \_\_\_\_\_
3. HHS OIG, OEI-06-14-00350 (2016), p. iii — “HealthCare.gov lacked clear project leadership to give direction and unity of purpose, responsiveness in execution, and a comprehensive view of progress.” — Retrieved from: \_\_\_\_\_
4. L. A. Schlesinger & P. D. Bhayani, “HealthCare.gov: The Crash and the Fix (A),” Harvard Business School case 315129 (2015); Mergel et al. (2018), digital-government literature — the rescue and what the code was not. — Retrieved from: \_\_\_\_\_
5. J. Pahlka, *Recoding America* (2023) — the founding of the U.S. Digital Service out of the rescue. — Retrieved from: \_\_\_\_\_

## CASE 181 **EU Human Brain Project — Top-Down Vision That Unraveled** 2013 – 2023

1. “How big science failed to unlock the mysteries of the human brain,” *MIT Technology Review*, 25 August 2021 — the principal critical assessment of HBP alongside BRAIN: the FET Flagship selection, the top-down governance under Markram, and the €1 billion headline. — Retrieved from: \_\_\_\_\_

2. Human Brain Project final reports (2023) — the project's own restructuring and concluding documentation: €607 million drawn across four funding periods, the re-scoping around EBRAINS, which entered the ESFRI Roadmap for Research Infrastructures in 2021, and the flagship's conclusion on 30 September 2023. — Retrieved from: \_\_\_\_\_
3. Alivisatos et al. (2012), Neuron — the BRAIN position paper that is the paired contrast (Case 198). — Retrieved from: \_\_\_\_\_
4. Science and Nature contemporaneous coverage of the July 2014 open letter (more than 800 signatories) and the mediation and restructuring (2014–2016); In Silico (2020), documentary by Noah Hutton — long-form follow of Markram and the HBP through the contestation period. This is the journalism-tier material the evidence flag is about. — Retrieved from: \_\_\_\_\_
5. MIT Technology Review (2021) — the governance-variable comparison across the two programmes. — Retrieved from: \_\_\_\_\_

## CASE 182 Amazon Hiring AI — Trained Bias, Deprecated 2017

2014 – 2018

1. Dastin, J. (2018), "Amazon scraps secret AI recruiting tool that showed bias against women," Reuters, October 10, 2018 — the primary investigation: the project initiated in 2014 to automate resume screening for technical roles. — Retrieved from: \_\_\_\_\_
2. Dastin (2018), Reuters — the model penalising resumes containing the word "women's" and downgrading graduates of two all-women's colleges; Amazon "edited the programs to make them neutral to these particular terms," but "that was no guarantee that the machines would not devise other ways of sorting candidates that could prove discriminatory." — Retrieved from: \_\_\_\_\_
3. Dastin (2018), Reuters — Amazon "disbanded the team by the start of last year because executives lost hope for the project," and recruiters "looked at the recommendations generated by the tool when searching for new hires, but never relied solely on those rankings." — Retrieved from: \_\_\_\_\_
4. Kim, P. T. (2017), "Data-Driven Discrimination at Work," William & Mary Law Review 58(3):857–936; Raghavan, M., Barocas, S., Kleinberg, J., & Levy, K. (2020), "Mitigating Bias in Algorithmic Hiring: Evaluating Claims and Practices," Proceedings of FAT\* 2020, pp. 469–481; Bogen, M., & Rieke, A. (2018), Help Wanted: An Examination of Hiring Algorithms, Equity, and Bias, Upturn report — the academic frame and the algorithmic-hiring landscape into which the case is positioned. — Retrieved from: \_\_\_\_\_

## CASE 183 Uber ATG / Tempe Fatality

2018

1. NTSB, Collision Between Vehicle Controlled by Developmental Automated Driving System and Pedestrian, Highway Accident Report HAR-19/03 (2019) — the Tempe collision and probable cause. — Retrieved from: \_\_\_\_\_
2. NTSB HAR-19/03 (2019) — the safety operator's distraction (watching a video) before impact. — Retrieved from: \_\_\_\_\_
3. NTSB HAR-19/03 (2019) — Uber "did not adequately recognize the risk of automation complacency"; training and policy failures (quoted in part). — Retrieved from: \_\_\_\_\_
4. NTSB HAR-19/03 (2019) — the absent functionality to anticipate a pedestrian crossing midblock; the one-second action-suppression period with no operator alert; and the design under which, where a collision could not be avoided within the allowed braking and jerk limits, the ADS sounded an auditory alert and slowed gradually rather than applying maximum braking. — Retrieved from: \_\_\_\_\_
5. NTSB HAR-19/03 (2019), Board Member Statement, p. 62 — Vice Chairman Bruce Landsberg: "Automation performs remarkably well most of the time and therein lies the problem." The report's own finding, at §2.2.2,

is that “when it comes to the human capacity to monitor an automation system for its failures, research findings are consistent—humans are very poor at this task.” — Retrieved from: \_\_\_\_\_

6. Uber’s sale of its Advanced Technologies Group to Aurora, announced 7 December 2020 and closed in January 2021; *State of Arizona v. Vasquez*, Maricopa County Superior Court (2023) — guilty plea to endangerment, three years’ supervised probation. — Retrieved from: \_\_\_\_\_

7. R. Parasuraman & D. Manzey (2010), complacency in automation — cited by the NTSB at §2.2.2; L. Bainbridge (1983), “Ironies of Automation.” — Retrieved from: \_\_\_\_\_

## CASE 184 UK Post Office Horizon Scandal

1999 – 2015

1. Post Office Horizon IT Inquiry hearings and exhibits (2020– ) — the system, the prosecutions, and the human toll. — Retrieved from: \_\_\_\_\_

2. *Hamilton & Others v. Post Office Limited* (Court of Appeal, 2021) — quashed convictions. — Retrieved from: \_\_\_\_\_

3. Internal Fujitsu and Post Office documents released through litigation — engineers’ knowledge of Horizon bugs. — Retrieved from: \_\_\_\_\_

4. Sir Wyn Williams, Post Office Horizon IT Inquiry, Volume 1 (July 2025) — senior employees “knew... that Legacy Horizon was capable of error” (quoted). — Retrieved from: \_\_\_\_\_

5. N. Wallis, *The Great Post Office Scandal* (2021). — Retrieved from: \_\_\_\_\_

6. The Post Office (Horizon System) Offences Act 2024 (mass exoneration) and the compensation schemes. — Retrieved from: \_\_\_\_\_

## CASE 185 Air Canada Chatbot Liability — Delegation Without Revocation

2022 – 2024

1. Air Canada bereavement-fare policy text (as in effect November 2022 and through the period covered by the ruling) — referenced in the tribunal decision as the divergence the chatbot’s representation produced. — Retrieved from: \_\_\_\_\_

2. Cecco, L. (2024), “Air Canada ordered to pay customer who was misled by airline’s chatbot,” *The Guardian*, February 16, 2024 — contemporaneous press coverage of the ruling. — Retrieved from: \_\_\_\_\_

3. *Moffatt v. Air Canada*, 2024 BCCRT 149 (British Columbia Civil Resolution Tribunal, February 14, 2024), Tribunal Member Christopher Rivers presiding. — Retrieved from: \_\_\_\_\_

4. Sookman, B. (McCarthy Tétrault LLP, 2024), “Moffatt v. Air Canada: A Misrepresentation by an AI Chatbot,” McCarthy Tétrault TechLex Blog, February 19, 2024 — practitioner-tier analysis of the tribunal’s negligent-misrepresentation holding, the rejection of the “separate legal entity” defense, and the duty-of-care framing for AI-mediated consumer interactions. mccarthy.ca. — Retrieved from: \_\_\_\_\_

## CASE 186 Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity

2009 – 2022

1. Consumer Financial Protection Bureau, Mortgage Market Activity and Trends (annual HMDA reports), supporting the population-level disparities backdrop of the fintech transition. — Retrieved from: \_\_\_\_\_

2. Bartlett, Morse, Stanton, & Wallace (2022), “Consumer-lending discrimination in the FinTech era,” *Journal of Financial Economics* 143(1):30–56, doi:10.1016/j.jfineco.2021.05.047 — the acceptance margin: “Face-to-



face lenders reject Latinx and African-American applications approximately 6% more often than they reject similarly situated non-minority applicants,” while FinTech lenders “do not discriminate at all in the decision to reject or accept”; the authors “note that these results are preliminary.” The 6% is the face-to-face figure specifically; the abstract’s shorter phrasing says “lenders,” and the case follows the body text. — Retrieved from: \_\_\_\_\_

3. Bartlett et al. (2022) — the pricing margin: 7.9 and 3.6 basis points on purchase and refinance mortgages, \$765M a year in aggregate extra interest, algorithms discriminating “40% less than face-to-face lenders,” and the rent-extraction reading (quoted). — Retrieved from: \_\_\_\_\_
4. Dwork et al. (2012), “Fairness Through Awareness,” ITCS 2012 — the foundational technical statement of the limits of unawareness. — Retrieved from: \_\_\_\_\_
5. Mitchell, Potash, Barocas, D’Amour, & Lum (2021), “Algorithmic Fairness: Choices, Assumptions, and Definitions,” Annual Review of Statistics and Its Application 8:141–163 — the competing-definitions framing. — Retrieved from: \_\_\_\_\_

### CASE 187 **COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate** 2014 – 2018

1. Angwin, J., Larson, J., Mattu, S., & Kirchner, L. (2016), “Machine Bias,” ProPublica, May 23, 2016 — the audit investigation of COMPAS scores in Broward County, Florida. — Retrieved from: \_\_\_\_\_
2. Dieterich, W., Mendoza, C., & Brennan, T. (2016), COMPAS Risk Scales: Demonstrating Accuracy Equity and Predictive Parity, Northpointe Inc. response document. — Retrieved from: \_\_\_\_\_
3. Chouldechova, A. (2017), “Fair Prediction with Disparate Impact: A Study of Bias in Recidivism Prediction Instruments,” Big Data 5(2):153–163, doi:10.1089/big.2016.0047. — Retrieved from: \_\_\_\_\_
4. Kleinberg, J., Mullainathan, S., & Raghavan, M. (2017), “Inherent Trade-Offs in the Fair Determination of Risk Scores,” Proceedings of ITCS 2017, doi:10.4230/LIPIcs.ITCS.2017.43. — Retrieved from: \_\_\_\_\_
5. Dressel, J., & Farid, H. (2018), “The accuracy, fairness, and limits of predicting recidivism,” Science Advances 4(1):eaao5580 — COMPAS no more accurate than untrained humans or a two-variable model. — Retrieved from: \_\_\_\_\_
6. Rudin, C., Wang, C., & Coker, B. (2020), “The age of secrecy and unfairness in recidivism prediction,” Harvard Data Science Review 2(1) — the age-nonlinearity reanalysis of ProPublica’s disparity claim. — Retrieved from: \_\_\_\_\_
7. Bao, M., et al. (2021), “It’s COMPASlicated: The Messy Relationship between RAI Datasets and Algorithmic Fairness Benchmarks,” NeurIPS Datasets & Benchmarks — the case against COMPAS as a fairness benchmark. — Retrieved from: \_\_\_\_\_

### CASE 189 **Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds** 2014 – 2020

1. District Court of The Hague (2020), judgment of 5 February 2020 (NJCM et al. v. State of the Netherlands), ECLI:NL:RBDHA:2020:865; the official English translation is ECLI:NL:RBDHA:2020:1878, from which the paragraph numbers below are taken. ¶6.93 records that SyRI had to date been applied only in so-called “problem districts” and that “there is in fact a risk that SyRI inadvertently creates links based on bias, such as a lower socio-economic status or an immigration background.” — Retrieved from: \_\_\_\_\_
2. Rachovitsa & Johann (2022), “The Human Rights Implications of the Use of AI in the Digital Welfare State: Lessons Learned from the Dutch SyRI Case,” Human Rights Law Review 22(2), doi:10.1093/hrlr/ngac010; the transparency holding is at ¶6.86 of the judgment. — Retrieved from: \_\_\_\_\_



3. AlgorithmWatch (2020), "How Dutch activists got an invasive fraud detection algorithm banned," Automating Society Report 2020 — carrying the de Volkskrant reporting of late June 2019 that no SyRI investigation had detected new fraud, and that of five neighbourhood analyses requested by municipalities in five years only two were run. — Retrieved from: \_\_\_\_\_
4. Library of Congress Global Legal Monitor (2020), report on the SyRI ruling — the ruling's system-specific scope. — Retrieved from: \_\_\_\_\_

## CASE 190 **Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment** 2023 – 2024

1. California DMV (24 October 2023), Order of Suspension, Cruise LLC — driverless deployment and testing permits: "The AV traveled about 20 feet and reached a speed of 7 miles per hour before coming to a complete and final stop." — Retrieved from: \_\_\_\_\_
2. California DMV, Order of Suspension (24 October 2023) — the 3 October meeting at which the footage shown "ended with the AV's initial stop"; "the department only learned of the AV's subsequent movement via discussion with another government agency"; the full video provided on 13 October. Reconstruction of the internal decision-making rests on TechCrunch (October 2023). — Retrieved from: \_\_\_\_\_
3. California DMV, "DMV Statement on Cruise LLC Suspension" (24 October 2023) — the four regulatory provisions cited: 13 CCR §§228.20(b)(6), 228.20(b)(3), 227.42(b)(5) and 227.42(c). — Retrieved from: \_\_\_\_\_
4. Quinn Emanuel Urquhart & Sullivan, report to the Cruise board on the 2 October 2023 incident and the company's regulatory disclosures, published 25 January 2024; NBC News, SF Standard and Mission Local reporting (October 2023) — the contemporaneous journalism-tier reconstruction. — Retrieved from: \_\_\_\_\_
5. NHTSA consent order with Cruise, 30 September 2024 (\$1.5 million civil penalty); U.S. Department of Justice deferred prosecution agreement with Cruise, November 2024 (admission of a false report; \$500,000); General Motors announcement of 10 December 2024 ending funding of Cruise robotaxi development. — Retrieved from: \_\_\_\_\_
6. Cruise public statements, October–November 2023, and the company's blog post accepting the Quinn Emanuel findings (January 2024). — Retrieved from: \_\_\_\_\_

## CASE 191 **Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict** 2016 – 2023

1. Australian National Audit Office, Centrelink's Compliance Activities — Income Compliance Program, performance audit reports across 2017–2020. — Retrieved from: \_\_\_\_\_
2. Royal Commission into the Robodebt Scheme, Final Report, Commissioner Catherine Holmes AC SC, July 7, 2023 (Commonwealth of Australia). — Retrieved from: \_\_\_\_\_
3. Prygodicz v Commonwealth of Australia (No 2) [2021] FCA 634 — Federal Court class-action settlement following the 2019 unlawfulness finding. — Retrieved from: \_\_\_\_\_
4. Whiteford, P. (2021), "Debt by Design: The Anatomy of a Social Policy Fiasco," Australian Journal of Public Administration 80(2):340–360 — academic synthesis of the policy and administrative history. — Retrieved from: \_\_\_\_\_

## CASE 192 **Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model** 2019 – 2021

1. Heinemeier Hansson, D. (2019), Twitter thread of November 7, 2019, archived in DFS investigation record and contemporaneous press coverage (Bloomberg, New York Times, Wall Street Journal, November 2019). — Retrieved from: \_\_\_\_\_
2. Goldman Sachs Bank USA, public response and credit-line-review process documentation submitted to DFS during the investigation (2019 – 2021). — Retrieved from: \_\_\_\_\_
3. New York State Department of Financial Services, Report on Apple Card Investigation, March 23, 2021. — Retrieved from: \_\_\_\_\_
4. Vigdor, N. (2019), “Apple Card Investigated After Gender Discrimination Complaints,” The New York Times, November 10, 2019. — Retrieved from: \_\_\_\_\_

## CASE 193 **Bernard Madoff / SEC Failures** 1992 – 2008

1. SEC Office of Inspector General, Investigation of Failure of the SEC to Uncover Bernard Madoff’s Ponzi Scheme, Report of Investigation Case No. OIG-509 (31 August 2009) — the quoted finding. — Retrieved from: \_\_\_\_\_
2. Markopolos, H. (2010), No One Would Listen — the 2005 memo and its dismissal. — Retrieved from: \_\_\_\_\_
3. United States v. Madoff (2009) — guilty plea and the \$65B figure. — Retrieved from: \_\_\_\_\_
4. U.S. Senate Committee on Banking, Housing, and Urban Affairs, The Madoff Investment Securities Fraud: Regulatory and Oversight Concerns and the Need for Reform, hearing, 10 September 2009 — SEC and OIG testimony on the examination and expertise failures. — Retrieved from: \_\_\_\_\_
5. U.S. Securities and Exchange Commission, The Securities and Exchange Commission Post-Madoff Reforms (2009–2010) — revamped tip and referral handling, the Office of Market Intelligence, and specialized enforcement units. — Retrieved from: \_\_\_\_\_

## CASE 194 **Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox)** 2001 – present

1. Jackson, Dreyling, & Pappel (2021), “A Historical Analysis on Interoperability in Estonian Data Exchange Architecture,” ICEGOV 2021 proceedings, doi:10.1145/3494193.3494209. — Retrieved from: \_\_\_\_\_
2. Republic of Estonia, e-Estonia briefing materials and Year-in-Review documentation (2024) — program-report sourcing, including the e-Tax figures (around 99% of declarations filed electronically; three minutes to file) and the December 2024 100%-digital milestone. — Retrieved from: \_\_\_\_\_
3. Kattel, R., & Mergel, I. (2018), “Estonia’s digital transformation: Mission mystique and the hiding hand,” UCL Institute for Innovation and Public Purpose Working Paper 2018-09 — the open-access version carrying the no-legacy-paradox sentence verbatim and the thirteen-year no-legacy design principle. — Retrieved from: \_\_\_\_\_
4. X-Road Global / Nordic Institute for Interoperability Solutions (NIIS) — program documentation and deployment-partner case studies. — Retrieved from: \_\_\_\_\_
5. Kattel, R., & Mergel, I. (2019), “Estonia’s Digital Transformation: Mission Mystique and the Hiding Hand,” in Compton, M., & ‘t Hart, P. (eds.), Great Policy Successes (Oxford University Press, 2019) — the analytical chapter; an edited-volume chapter, not a peer-reviewed journal article. — Retrieved from: \_\_\_\_\_

**CASE 195 Tesla Autopilot — Recurring Fatalities**

2016 – present

1. NTSB, Highway Accident Report HAR-17/02 (Williston, FL, 2016) — the quoted disengagement finding. — Retrieved from: \_\_\_\_\_
2. NHTSA Standing General Order 2021-01 reports — documented Autopilot fatal crashes. — Retrieved from: \_\_\_\_\_
3. NTSB, Highway Accident Report HAR-20/01 (Mountain View, CA, 2018) — Autopilot crash analysis. — Retrieved from: \_\_\_\_\_
4. NHTSA Office of Defects Investigation EA22-002 (opened 2022; closed April 2024, superseded by recall 23V-838) — driver-engagement adequacy. — Retrieved from: \_\_\_\_\_
5. Parasuraman, R. & Manzey, D. (2010) — automation complacency and monitoring. — Retrieved from: \_\_\_\_\_
6. Benavides v. Tesla (S.D. Fla.) — August 2025 federal jury verdict apportioning Tesla 33% of responsibility for the April 2019 Key Largo Autopilot crash, \$243 million; post-trial motions denied February 2026. — Retrieved from: \_\_\_\_\_
7. NHTSA recall-query RQ24009 (April 2024) into the adequacy of the December 2023 Autopilot remedy, and investigation PE24031 (October 2024) into Full Self-Driving; Tesla's June 2025 Austin robotaxi launch and the ensuing regulator scrutiny. — Retrieved from: \_\_\_\_\_

**CASE 196 Fintech Lending Fairness Audit — When Including Race Reduces Disparity**

2025

1. Coots, Bartlett, Nyarko, & Goel (2025), "Algorithmic Bias in Lending: Evidence from a Fintech Audit," arXiv:2512.20753 — audit of 80,000 personal loans showing model miscalibration disparities and that controlled inclusion of protected attributes could correct them; the evidence-tier flag is binding until peer review completes. — Retrieved from: \_\_\_\_\_
2. Bartlett, Morse, Stanton, & Wallace (2022), "Consumer-lending discrimination in the FinTech era," Journal of Financial Economics 143(1):30–56 — the paired big-tier case (186). — Retrieved from: \_\_\_\_\_
3. Chouldechova (2017), "Fair Prediction with Disparate Impact," Big Data 5(2):153–163 — the impossibility result for calibration and error-rate parity. — Retrieved from: \_\_\_\_\_
4. Kleinberg, Mullainathan, & Raghavan (2017), "Inherent Trade-Offs in the Fair Determination of Risk Scores," ITCS — competing fairness definitions are not jointly achievable. — Retrieved from: \_\_\_\_\_
5. Mitchell, Potash, Barocas, D'Amour, & Lum (2021), "Algorithmic Fairness: Choices, Assumptions, and Definitions," Annual Review of Statistics and Its Application 8:141–163 — practitioner-facing survey. — Retrieved from: \_\_\_\_\_

**CASE 197 Predictive Policing — PredPol**

2011 – present

1. Lum, K. & Isaac, W. (2016), "To Predict and Serve?," Significance — the enforcement-vs-crime feedback loop (paraphrased). — Retrieved from: \_\_\_\_\_
2. Brantingham, P. J., Valasik, M., & Mohler, G. O. (2018), "Does Predictive Policing Lead to Biased Arrests? Results From a Randomized Controlled Trial," Statistics and Public Policy 5(1):1–6 — no statistically significant difference in arrests by racial-ethnic group between treatment and control; the lead author co-founded PredPol. — Retrieved from: \_\_\_\_\_

3. Richardson, Schultz & Crawford (2019), "Dirty Data, Bad Predictions" — biased records feeding predictive systems. — Retrieved from: \_\_\_\_\_
4. Municipal records on suspension and abandonment of predictive policing (Santa Cruz, New Orleans, Los Angeles). — Retrieved from: \_\_\_\_\_
5. Sankin, A. & Mattu, S. (2023), "Predictive Policing Software Terrible at Predicting Crimes," The Markup, 2 October 2023 (copublished with Wired) — the Plainfield, NJ accuracy analysis; and reporting on PredPol/Geolitica's 2021 rebrand and 2023 wind-down into SoundThinking. — Retrieved from: \_\_\_\_\_
6. Brayne, S. (2017), "Big Data Surveillance: The Case of Policing," American Sociological Review 82(5):977–1008 — ethnographic account of the LAPD's adoption of predictive and dragnet analytics. — Retrieved from: \_\_\_\_\_

## CASE 198 **Launching the BRAIN Initiative — Governance of a Grand Challenge** 2011 – 2015 – present

1. Alivisatos, Chun, Church, Greenspan, Roukes, & Yuste (2012), "The Brain Activity Map Project and the Challenge of Functional Connectomics," Neuron 74(6):970–974, doi:10.1016/j.neuron.2012.06.006. — Retrieved from: \_\_\_\_\_
2. Jorgenson et al. (2015), "The BRAIN Initiative: developing technology to catalyse neuroscience discovery," Phil. Trans. R. Soc. B 370(1668):20140164, doi:10.1098/rstb.2014.0164 — the BRAIN 2025 plan. — Retrieved from: \_\_\_\_\_
3. Yuste & Bargmann (2017), "Toward a Global BRAIN Initiative," Cell 168(6):956–959, doi:10.1016/j.cell.2017.02.023. — Retrieved from: \_\_\_\_\_
4. Underwood, E. (2013), "As White House Embraces BRAIN Initiative, Questions Linger," Science / ScienceInsider (April 3, 2013) — source of both quoted spans verbatim: Yuste on "packing the committee with users, rather than tool builders," and Bargmann's February 2013 email to Science on "a big central planning project that will take resources away from creative work." — Retrieved from: \_\_\_\_\_
5. Mullin, E. (2021), "How big science failed to unlock the mysteries of the human brain," MIT Technology Review (25 August 2021) — the critical ten-year assessment. — Retrieved from: \_\_\_\_\_

## CASE 199 **Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact** 2023 – 2026

1. Cruise / California DMV Order of Suspension (October 2023) — paired Case 190 as the foil. — Retrieved from: \_\_\_\_\_
2. Waymo (2023), "A Blueprint for AV Safety: Waymo's Toolkit For Building a Credible Safety Case," whitepaper. — Retrieved from: \_\_\_\_\_
3. Waymo (November 2025), "Independent Audits of Waymo's Safety Case and Remote Assistance Programs," summary release. — Retrieved from: \_\_\_\_\_
4. NHTSA recall and investigation record (December 2025 school-bus recall; January 2026 NHTSA/NTSB investigation of a Santa Monica child-strike; June 2026 recall of 3,900 vehicles over freeway construction zones) — the post-deployment failure record under regulator action. — Retrieved from: \_\_\_\_\_
5. Montreal AI Ethics Institute (2023), summary and analysis of the Waymo safety case framework. — Retrieved from: \_\_\_\_\_

6. California Public Utilities Commission, AV passenger-service permit framework documents — paired Case 200 for the regulator-side artifact. — Retrieved from: \_\_\_\_\_

**CASE 200 CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver** 2020 – 2026

1. California Public Utilities Commission, “Autonomous Vehicle Passenger Service Programs” — program page for the Drivered and Driverless pilot and deployment programs, including the Passenger Safety Plan requirement. — Retrieved from: \_\_\_\_\_
2. California Public Utilities Commission, Resolution TL-19144 (Waymo), issued 11 August 2023 (Commission vote 10 August) — approves driverless deployment in all of San Francisco 24 hours a day per the DMV-approved ODD, and declines to set limits on operating hours, geography or fleet size; the span quoted in the text is verbatim (“the Commission will neither modify the DMV-approved ODD submitted by Waymo ... nor set limits on fleet size”). — Retrieved from: \_\_\_\_\_
3. California Department of Motor Vehicles, “New Autonomous Vehicle Regulations Strengthen Oversight and Enforcement, Authorize Trucks and Transit” (28 April 2026) — adopted authority to impose targeted operational restrictions on fleet size, location, speed and weather; draft language circulated for comment 30 August 2024. — Retrieved from: \_\_\_\_\_
4. CPUC permit decisions for Cruise and Waymo, 2020–2024. — Retrieved from: \_\_\_\_\_

**CASE 201 Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India** 2017 – 2025

1. Supreme Court of India (2018), Justice K. S. Puttaswamy (Retd.) v. Union of India, (2019) 1 SCC 1 — the Aadhaar judgment; 4–1 majority upholding Section 7, with Justice D. Y. Chandrachud’s dissent on technological exclusion and dignity. The quoted span is verbatim from ¶1262 of the dissent (Part H) — do not re-quote it from the judgment summaries, which paraphrase it. — Retrieved from: \_\_\_\_\_
2. Drèze, J., Khalid, N., Khera, R., & Somanchi, A. (2017), “Aadhaar and Food Security in Jharkhand: Pain without Gain?” Economic & Political Weekly 52(50), 16 December 2017 — field study of biometric-authentication exclusion in the PDS; more than 40% of non-transacting households cited ABBA-related problems in the month before the survey, and Ranchi district point-of-sale records put about 20% of cardholders unable to transact in an average month over January–July 2017. — Retrieved from: \_\_\_\_\_
3. Right to Food Campaign compilation and Indian journalism (The Hindu, The Wire, Scroll.in, 2017–2019) on the Jharkhand starvation deaths (Santoshi Kumari, 28 September 2017) — lived-exclusion sourcing with journalism-tier flag. The compilation runs to 57 deaths across nine states from 2015; 27 are reported as Aadhaar-related and at least 19 as directly linked to Aadhaar exclusion. The case carries the narrower 19 deliberately. — Retrieved from: \_\_\_\_\_
4. Supreme Court of India (2025), Pragya Prasun & Ors. v. Union of India, 2025 INSC 599 (30 April 2025) — the distinct, later ruling reading a right to accessible digital access into Article 21 in the e-KYC context for persons with disabilities. — Retrieved from: \_\_\_\_\_
5. Dixon, P. (2017), “A Failure to ‘Do No Harm’ — India’s Aadhaar biometric ID program and its inability to protect privacy in relation to measures in Europe and the U.S.,” Health and Technology 7(4):539–567 — comparative privacy and data-protection analysis, reporting biometric-authentication failure rates alongside it. — Retrieved from: \_\_\_\_\_

**CASE 203 NYC Local Law 144 — Bias Audits as Governance Artifact** 2023 – present

1. New York City Department of Consumer and Worker Protection, Rules Implementing Local Law 144 of 2021 (Automated Employment Decision Tools), effective July 5, 2023. — Retrieved from: \_\_\_\_\_

2. Wright, L., Muenster, R. M., Vecchione, B., Qu, T., Cai, S., Smith, A., Metcalf, J., & Matias, J. N. (2024), "Null Compliance: NYC Local Law 144 and the Challenges of Algorithm Accountability," in Proceedings of FAccT 2024, doi:10.1145/3630106.3658998. — Retrieved from: \_\_\_\_\_
3. Office of the New York State Comptroller (2025), audit of the Department of Consumer and Worker Protection's enforcement of Local Law 144 — finding enforcement "ineffective" (December 2025). — Retrieved from: \_\_\_\_\_
4. Engler, A. (2023), "The EU and U.S. diverge on AI regulation: A transatlantic comparison and steps to alignment," Brookings Institution commentary — regulatory-comparative frame for the municipal intervention. — Retrieved from: \_\_\_\_\_

## CASE 205 The Discipline We Build Next

ongoing

1. LDT / LENS program statement, Johns Hopkins School of Education — the program and its commitments. — Retrieved from: \_\_\_\_\_
2. Goodell & Kolodner (2022), Learning Engineering Toolkit — the discipline's methods. — Retrieved from: \_\_\_\_\_
3. The preceding cases of this volume — the cumulative record of cost and of possibility. — Retrieved from: \_\_\_\_\_
4. INPO (Case 175) and CRM/CAST (Case 117) — institution-building precedents (1979, 1981). — Retrieved from: \_\_\_\_\_
5. IEEE ICICLE Learning Engineering Body of Knowledge (LEBoK); the introduction to this volume. — Retrieved from: \_\_\_\_\_

## WHAT LENS OFFERS

## The program behind the casebook.

LENS — **Learning Engineering for Next-Generation Systems** — is a concentration within the Johns Hopkins University School of Education's **Learning Design & Technology** program. It operationalizes the casebook's central claim — that capability is a designable, measurable, engineerable property of a complex system — for high-consequence work in healthcare, defense, and education. It is housed in a School of Education because the binding constraint is rarely the engineering, which is mature; it is the learning sciences, educational measurement, implementation science, and equity-and-governance machinery that decide whether a capability intervention actually scales.

## THE FIVE PILLARS

LENS stands on five institutional pillars, load-bearing together rather than singly. **Mission Literacy** — reading what an operational setting actually requires and translating it into design and measurement decisions a serious reviewer would accept. **JHU Ecosystem** — the School of Education's centers, doctoral programs, and faculty in learning sciences, measurement, and design-based research, alongside university practice partners (the Armstrong Institute for Patient Safety and Quality; the Bloomberg School's implementation-science tradition; Whiting School engineering and cognitive science; the Berman Institute's governance and ethics work). **Intersectional Expertise** — the distinguishing pillar: holding the engineering, the learning, the measurement, the equity, and the operational context in one conversation. **Capability Focus** — keeping every artifact, measurement, and iteration accountable to what the system can actually do. **Flywheel Iteration** — the Practice Flywheel (**Identify → Activate → Prototype → Analyze → Transition**) that turns each cohort, case, and pilot into the starting condition of the next.

## THE RECORD AT JOHNS HOPKINS

The placement is not a default. The School of Education has studied the learner, designed the assessment, and engineered the intervention since before learning engineering had its current name. The Center for Talented Youth (1979) holds one of the longest continuous evidence bases in U.S. education on the identification and accelerated instruction of advanced learners; the Center for Research and Reform in Education has, since 2004, run large-scale evaluations against rigorous evidence standards (ESSA tiers, What Works Clearinghouse) and publishes the Best Evidence Encyclopedia. The Learning Design & Technology program inherits that environment: for two decades its M.Ed. and EdD tracks have graduated designers and evaluators of learning systems whose pedagogical core is design-based research. LENS formalizes that orientation as a concentration rather than departing from it.

## THE PILOTS


The claim that capability is engineerable rests on a documented record of learning-engineering pilots run at the Johns Hopkins Applied Physics Laboratory before LENS enrolled its first cohort. **CIRCUIT** was built as the workforce solution to the IARPA MICrONS connectomics program; the cohorts it trained contributed materially to the MICrONS reconstructions now published in *Nature*, with student development a deliberate byproduct of the mission result. The model is documented in the peer-reviewed engineering-education literature, and was generalized into the **MERIT** framework and the **COMPASS** mentor-intervention layer in JHU/APL program documentation. The **PISTA** incubator (DTRA, Summer 2025) compressed the same cycle into ten weeks: nine student fellows shipped a working CBRNE decision-support prototype to agency leadership through three full redesigns. No single pilot settles the proposition; the cumulative record shows the program is itself the institutional flywheel its case studies argue every high-consequence domain requires.



## ABOUT THE EDITORS

## The editors.

The casebook's central argument — that capability engineering is the applied form of the educational sciences, taught from the School of Education and informed by operational practice — is embodied in the two editors. Each has spent a career working in one half of that pairing and a substantial portion of it in the other. Each contributes to the five pillars on which LENS rests: Mission Literacy · JHU Ecosystem · Intersectional Expertise · Capability Focus · Flywheel Iteration.



William Gray-Roncal, Ph.D.

APPLIED RESEARCH · LEARNING ENGINEERING · STEM WORKFORCE

William Gray-Roncal, Ph.D., is Principal Research Engineer at the Johns Hopkins University Applied Physics Laboratory, where he applies systems engineering methods to challenges spanning defense, space, biomedical, and learning systems. His research — sponsored by DARPA, IARPA, NIH, and DoD — bridges computational neuroscience, artificial intelligence, and large-scale data infrastructure. He developed the College Prep and CIRCUIT workforce-development models, through which he has mentored nearly 500 students, and his current work develops quantitative measures of human cognitive performance in operational contexts. He leads the development of the LENS (Learning Engineering for Next-Generation Systems) concentration within the Johns Hopkins Master of Education in Learning Design and Technology, the program directed by James Diamond at the School of Education. Gray-Roncal holds a Ph.D. in Computer Science from Johns Hopkins, an M.S. in Electrical Engineering from USC, and a B.S. in Electrical Engineering from Vanderbilt.



James Diamond, Ph.D.

LEARNING SCIENCES · EDUCATIONAL DESIGN · EVALUATION

James Diamond, Ph.D., is Assistant Professor in the Johns Hopkins University School of Education and Program Director of the Master of Education in Learning Design and Technology, where he applies design-based research and learning-sciences methods to the design and evaluation of technology-mediated learning environments. His research — supported by the National Science Foundation, the Institute of Education Sciences, the Bill & Melinda Gates Foundation, the MacArthur Foundation's HASTAC Digital Media and Learning Initiative, the Robin Hood Learning + Technology Fund, the Wellcome Trust, and the UK Economic and Social Research Council — bridges digital games and simulations for learning, digital badges and micro-credentialing, classroom technology integration, and disciplinary literacy in history, social studies, civics, and STEM. Before joining Johns Hopkins, he was Senior Research Associate at the Education Development Center's Center for Children and Technology in New York City, where he led projects including Zoom In and the Who Built America? Teacher Mastery Badge System, contributed to Mission US and Possible Worlds, and taught at New York University as an adjunct. As Program Director, he oversees the LDT program within which the LENS (Learning Engineering for Next-Generation Systems) concentration is offered. Diamond holds a Ph.D. in Educational Communication and Technology from New York University, an Ed.M. in Educational Technology from Boston University, and a B.A. in History from Boston University.

These biographies reference careers documented in publications, program reports, faculty pages, and the curricular materials of the LDT program and the LENS concentration.