

Capability Matters

A CASEBOOK

Learning Engineering for Next-Generation Systems

Capability *Matters*

EDITED BY

William Gray-Roncal, PhD
James Diamond, PhD

JOHNS HOPKINS UNIVERSITY
LEARNING DESIGN AND TECHNOLOGY

EDITION · 15 JUNE 2026

An ongoing artifact of the LENS Practice Flywheel — Identify, Activate, Prototype, Analyze, Transition.

DEDICATION

*To everyone who believes in a dream and refuses to
give up.*

COLOPHON

Learning Engineering for Next-Generation Systems: Capability Matters

A Casebook.

First edition, 2026.

Edited by William Gray-Roncal, PhD and James Diamond, PhD.

Compiled for the LENS specialization within the Learning Design and Technology program at the Johns Hopkins University School of Education.

Set in Instrument Serif and DM Sans.

Printed via Lulu at 8 × 10 in, perfect-bound.

Copyright © 2026. All rights reserved.

HOW THIS VOLUME WAS MADE

This casebook was assembled using AI tools as part of an iterative learning-engineering process — the same Practice Flywheel (Identify → Activate → Prototype → Analyze → Transition) that the volume's case studies argue for. The methodology treats AI as the dual entity the curriculum names it: a creative partner that accelerated drafting, cross-referencing, layout, and citation lookup, and an epistemic risk that had to be hand-checked against the record. Every case in the book was reviewed by the editors and hand-checked by students for accuracy: confirming that the sources cited exist and are findable, that the quoted attributions are fairly represented, and that the case studies are accurate accounts of the incidents and the investigations they draw on. Items where the source could not be confirmed are marked "Paraphrasing..." so the attribution is honest about what is the author's reconstruction and what is verbatim from the record.

The volume is not finished. It is the first iteration of a reference dataset that LENS will continue to develop. We will revise the book based on reader feedback, reviewer findings, and our own ongoing practice of learning engineering — adding new cases as the operational record grows, correcting our record where it can be improved, and updating the curriculum crosswalk as the program itself iterates. Corrections, new cases, and reader feedback are welcomed.

SOURCES, ATTRIBUTION & LEGAL NOTICES

All cases reference publicly reported incidents and published investigations. Quoted material is attributed to source investigations, court documents, and academic publications. Reflection questions are original to this volume. The audit record of citations checked and changes made is maintained alongside the source repository.

References to convictions, settlements, and findings reflect the public record at the time of writing. Where individuals are named, they are named only as participants in investigations, court proceedings, or public inquiries already in the public domain. The casebook makes no legal characterization beyond what those records establish.

Quoted material is reproduced under fair use (17 U.S.C. § 107) for the purposes of criticism, comment, teaching, scholarship, and research. Trademarks and product names are used nominatively to identify the products and organizations under discussion; no affiliation or endorsement is claimed or implied.

Corrections, requests for clarification, and notice of any inaccuracy can be sent to LDT / LENS at the Johns Hopkins University School of Education. Substantive corrections will be incorporated in subsequent printings and noted in the casebook's public errata.

ABOUT THIS SUPPLEMENT

The argument, briefly

This volume is the digital supplement to **Capability Matters: A Casebook**. The printed main volume carries a selected set of cases — five to eleven per part — chosen to span the corpus’s domains, failure modes, evidence tiers, and signature lessons. This supplement carries every other case at full depth, in the same seven-part order, under the same editorial rules. Numbering is global across both books: a gap in the printed sequence is a case that lives here, and every case citation (by number or by slug) resolves identically in either volume.

The casebook’s claim is that **capability** — the interface between what a system requires of its human operators and the impact the system must deliver — is a designable, measurable, engineerable property of complex systems, and that treating it as an afterthought is the common mechanism behind failures that are usually filed under separate headings. The cases are organized topically in seven parts: Healthcare & Patient Safety; Education, Training & the Learning Workforce; Aviation & Aerospace; Defense & National Security; Industry, Energy & Enterprise Systems; Disaster Prevention & Recovery (read along the pre/post lifecycle of the event rather than a domain); and Algorithms, Governance & Public Systems. Each part splits into **What Fails** and **What Works — and the Frontier**.

Two analytic threads run across every part. The first is the six-mode failure taxonomy carried on each case header — T Training Gap · D Designed Out · N Normalization of Deviance · H Human-System Interface · G Governance & Trust · K Knowledge & Institutional Memory. The second is the **paired intervention**: every success case couples a mechanism to an authorization — a checklist to a nurse empowered to stop the line, a reporting form to immunity, a measurement instrument to a governance body that acts on it. Neither half alone moves the curve.

Every case carries three anchors (canonical LENS competency, induced sub-competency, and CLO), a scale tier, and an evidence-source tag. Weaker-evidence cases render an evidence-tier flag under the title with standing “future validation ongoing” language; cases touching the editors’ own institutions or work render a conflict-of-interest disclosure in the same position. Load-bearing causal hedges are never smoothed. For the full intellectual framing, the reading pathways, and the program context, see the main volume’s introduction; for the program documents of record, the LENS Companion; for the consolidated per-case references appendix with retrieval lines and the verification record, the Validation & Audit tracker.

The case matrix on the following pages covers the complete corpus; numbers marked ° are in this supplement, unmarked numbers are in the printed main volume.

THE CASE MATRIX

The cases at a glance

Gold numbers indicate failures and systemic conditions; teal numbers indicate paired-intervention successes and the open closing case. Numbers marked ° are carried in the digital supplement; unmarked cases appear in the printed main volume. Numbering is global across both.

1	Therac-25	1985 – 1987	H-D-G
2°	EHR / CPOE Implementation	2005 – present	H-D-G
3°	Watson for Oncology — Delegation Marketed Ahead of Capability	2013 – 2018	D-K-N
4°	Mid Staffordshire NHS Foundation Trust	2005 – 2009	G-N-K
5	Epic Sepsis Model	2017 – 2021	D-G-N
6	Racial Bias in Pain Assessment — The False-Belief Mechanism	2016	T-K-N
7°	VA Wait-Time Scandal	2014	G-K-N
8°	Medical Errors as Systemic Failure	1999 – present	T-H-N-K-G
9°	Vioxx Withdrawal	1999 – 2004	G-D
11°	California Nurse Staffing Ratios — Legislating a Capability Requirement	2004 – 2010	G-K
12°	ACGME 80-Hour Resident Duty-Hour Reform	2003–2017	T-K-N
13°	Implementation Science in Healthcare — The 17-Year Gap	ongoing	K-G-N
	Barsuk SBML — Simulation-Based		
14°	Mastery Learning Dissemination from Northwestern to the VA	2009 – 2020s	T-K
15°	I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer	2014	H-K-N
16°	NCSBN National Simulation Study — Licensing the 50% Substitution Rule	2014	G-K-D
17°	Spaced Education RCTs in Medical Training	2007 – 2009	H-K-D
18°	PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption	2023	K-N-D
19	Keystone ICU / Pronovost Checklist	2004 – present	T-N
20	TREWS / Sepsis Watch	2018 – 2022	H-K-G
21°	Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff	2024 – 2025	H-K-N
22°	ChatGPT in Healthcare — Hallucination Cases	2023 – present	H-D
23°	WHO Surgical Safety Checklist	2008 – present	T-N

24°	Bristol Heart Babies Reform	1984 – present	G·N
25°	Removing the Race Coefficient from eGFR	2021	D·G·N
26°	Pulse Oximetry Across Skin Tones	1990 – 2020 – 2025	D·G·N
27°	Anesthesia Monitoring Standards and the APSF — The Mortality Decline	1986 – present	H·K·G
28°	Interprofessional Education and the Evidence Gap	2013 – 2015	K·N
29°	Bar-Code Medication Administration — Cue at the Point of Care	2010	H·K·D
30°	Surgical Skill Video Peer-Rating Predicts Complications	2013	H·D·K
31°	Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units	2009 – 2016	T·K·H
32°	Annual-Screening UI Redesign + CDS at University of Missouri Health Care	2020	T·D·N
33°	Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal	2019 – 2024	T·D·N
34°	COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case	2022 – 2024	T·K·D
35°	Radiology AI Miscalibration	2018 – present	H·K·D
36°	AlphaFold — Protein Structure Prediction	2020 – present	T
37°	DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback	2020	T·K·D
38°	iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms	2006 – 2011	G·D·N
39°	TeamSTEPPS	2006 – present	T·N
40°	Team Science Training for Clinical and Translational Scientists	2020 – 2022	K·N
41°	Implementation-Science Training — Stated Goals Outrunning Operational Practice	2020s	K·N
42°	Australian Hospital-Pharmacy Technician Role Redesign	2016	D·N·H
43°	Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface	2013 – 2018	H·N
44°	Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD	2014 – present	G·N
45°	Tennessee Voluntary Pre-K Study	2009–2018	G·D
46°	Algorithmic Bias in Educational Predictive Analytics	ongoing	G·H·D
47°	Algorithmic Bias in Automated Exam Proctoring	2022	D·N·K
48°	School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism	2010s – 2022	G·K·N
49°	UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days	2020	D·K·N

Wisconsin DEWS — A Decade of		
50° Algorithmic Dropout Prediction	2012 – 2024	D-K-N
Atlanta Public Schools Cheating Scandal	2009 – 2015	G-N
Purdue Course Signals — The Reverse-Causality Retention Claim	2012 – 2013	D-K-N
53 inBloom	2014	G
Summit Learning / Personalized Learning Rollout	2014–2019	G-T-K
Enrollment-Algorithm Yield Optimization Across U.S. Higher Education	2010s – present (Brookings synthesis 2021)	T-K-N
GAO Online Program Manager Oversight Gap (GAO-22-104463)	2022	D-K-N
USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC)	2010s – 2024	D-K-N
60 Houston EVAAS — Value-Added Teacher Evaluation on Trial	2011–2017	G
Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap	1997–2023	K-T
62° Gates Intensive Partnerships — Measurement Without a Coupled Lever	2009–2018	G-K
63° LAUSD iPad Program — The Common Core Technology Project	2013–2015	D-H-G
65° Training Transfer — The Gap Between Learning and Doing	2010	K-N
66° Growth-Mindset National Experiment — Heterogeneous Effects	2019	D-N-K
67° Cognitive Tutor / Carnegie Learning	1990s – present	T
68 CIRCUIIT — A Scalable, Equity-Centered Research Workforce Model	2017 – 2023 (six cycles)	T-K
69° Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale	2016	T-D
70° High-Impact Learning System — Engineering the Environment, Not Just the Event	2001 – present	K-N
71° Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development	2025 (preprint)	K-N
72 ASSISTments — National Replication and Long-Term Follow-Through	2014 – present	T-K-D
73° The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension	2016 – 2025	T-K-D
74° Zhang/Scardamalia — Knowledge Building Across Three Cohorts	2009	T-K-D
75° Chen et al. — Rural China AI Devices and the Equity-Direction Finding	2025	T-K-D
76° Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account	2023	T-K-N
77 Hybrid Human-AI Tutoring — Augmentation, Not Delegation	2024	T-K-D

78°	CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale 2017 – present	H·K·N
79°	The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops 1959 – present	K·N
80°	Georgia State University Predictive Analytics 2012 – present	T·K
81°	Open University 'Ethical Use of Student Data' and OU Analyse 2014 – 2025	G·K·N
82°	MMALA — A Maturity Model for Responsible Learning–Analytics Adoption (Brazil) 2024	K·N
83°	Brinkerhoff Success Case Method — Tails as the Evaluation Instrument 2005 – present	K·N
84°	Cognitive Tutor Algebra I at Scale — Year–One Null, Year–Two Positive 2007 – 2014	T·K·D
85°	OU Analyse — Predictive Learning Analytics and Teacher Use at Scale 2019 – 2023	T·K·D
86°	Gándara — Detecting and Mitigating Algorithmic Bias in College Student–Success Prediction 2024	D·K·N
87°	Yu / Lee / Kizilcec — Protected Attributes in Learning–Analytics Models 2021 – 2024	D·K·N
88°	LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models 2025	T·K·N
89°	Data Privacy and Learning Analytics on the African Continent 2022	K·G·N
90°	Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions 2025	T·K·N
91°	LALA — Building Learning–Analytics Governance Capacity Across Latin America 2017 – 2020	K·N
92°	Norway's National Expert Commission on Learning Analytics 2021 – 2023	G·K·N
93°	Singapore SkillsFuture — National Workforce Capability at Scale 2015 – present	G·K·D
94°	Learning Analytics on the African Continent — A Scoping Review as the Present–State Map 2022	K·N
95°	PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop 1997–2024	T·G
96°	AeroPerú Flight 603 1996	H·T
97°	Boeing 737 MAX / MCAS 2018 – 2019	D·T·H
98°	Mars Climate Orbiter — Unit Mismatch 1999	D·K
100°	Glass–Cockpit Transition in Light Aircraft — Technology Outran Training 2010	H·N·K
101°	Ariane 5 Flight 501 1996	D·K·H
102°	Air France Flight 447 2009	T·H
103°	Kegworth / British Midland 92 1989	T·H·K
104°	Asiana Airlines Flight 214 2013	T·H
105°	Helios Airways Flight 522 2005	T·H

106°	TransAsia Airways Flight 235	2015	T-H	122°	Singapore Airlines Safety Transformation	1980s – present	T-N
107°	Eastern Air Lines Flight 401	1972	H-T		FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike	1992 – 2000s	G-K-N
108°	NASA EVA-23 — Water Intrusion Inside a Spacesuit	2013	H-N-D	123°	USS Fitzgerald & USS John S. McCain	2017	T-K-N
109°	Colgan Air Flight 3407	2009	T	124°	F-35 Sustainment & Maintainer Shortage	ongoing	T-K-D
110°	Atlas Air Flight 3591	2019	T-K		Military Fratricide — Desert Storm to Afghanistan	1991 – 2014	T-H-K
111°	Challenger & Columbia	1986 / 2003	N-K-G	125°	Operation Eagle Claw	1980	T-K
112°	NASA Saturn V Documentation	1972 – present	K	126°	Patriot Missile / Dhahran	1991	D-H-K
113°	Boeing Starliner	2019 – 2024	K-D	127°	USS Vincennes & Iran Air Flight 655	1988	H-T
114°	FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule	1988 – 2010s	G-D-K	128°	Stanislav Petrov / 1983 False Alert	1983	H-T
115°	FAA NextGen and the ADS-B Out Transition	2003 – 2020	G-D-K		F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap	2008 – 2012	H-K-N
116°	Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change	2005	G-K-H	129°	Mark 14 Torpedo Failures	1941 – 1943	T-K-G
117°	Crew Resource Management & CAST	1981 – present	T-H-N	130°	V-22 Osprey	1991 – present	T-H-N
118°	Korean Air Safety Transformation	2000 – present	T-N	131°	9/11 Intelligence Sharing Failures	1996 – 2001	G-K
119°	Aviation Safety Reporting System (ASRS)	1976 – present	T-K-N	132°	U.S. Nuclear Navy / Rickover Training Model	1954 – present	T-K-N
120°	EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation	1996 – 2002	H-K-G	133°	MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard	2020 revision (series since 1968)	G-K
121°	TCAS — Coordinated Collision Avoidance and the Überlingen Lesson	1981 – 2008 (TCAS II Version 7.1)	H-K-G	134°	GIFT and the Adoption Gap	2012 – present	K-G-N

140°	Navy Surface Warfare Readiness Reform 2018 – present	T-K-N	160°	Davis-Besse Reactor Head Corrosion 2002	N-K-G
141	DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks 2009 – 2014	H-K-D	161	Texas City BP Refinery Explosion 2005	N-T-K-G
142°	xAPI / Total Learning Architecture — Interoperability Gap ongoing	K-G	162°	Upper Big Branch Mine Explosion 2010	N-G-K
143°	Knight Capital Trading Loss 2012	D-K	163°	Deepwater Horizon 2010	T-N-K
144	Kodak Digital Camera Stagnation 1975–2012	D-K-N	164	Grenfell Tower 2017	G-T-K-N
146°	Northeast Blackout 2003	H-K	165°	Camp Fire / PG&E 2018	G-N-K
147°	Takata Airbag Inflators 2008 – 2023	D-G	166°	Three Mile Island 1979	T-H
148	Wells Fargo Fake Accounts 2011 – 2016	G-N	167	Hurricane Katrina — The FEMA/DHS Response Failure 2005	G-K
149°	Volkswagen Dieselgate 2015	D-G	169°	Fukushima Daiichi 2011	N-G-K
151°	GM Ignition Switch 2002 – 2014	D-G	170°	West Africa Ebola — The Delayed International Response 2014 – 2016	G-N
152	TSB Bank IT Migration 2018	H-G		Hurricane Maria — Puerto Rico Response and Logistics Failure 2017 – 2018	G-N
153°	Equifax Data Breach 2017	G-K	172°	CrowdStrike Falcon Outage 2024	D-K-G
154°	INL Turbine-Control Upgrade — Low-Burden Cutover as a Human-Factors Finding 2014	G-K-H	173	Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable 1963 – present	G-K-H
155	Toyota Production System / Andon Cord 1950s – present	N-G	174°	Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed 1996 – 2000	G-D-K
156°	INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet 2010 – present	D-H-K	175°	INPO and the Nuclear Academy 1979 – present	T-K-G
157	AI-Augmented Coding Tools 2021 – present	T-H	176	Tylenol Recall 1982	G-N
159°	Bhopal 1984	T-K-N-G	177°	CIRAS — Confidential Incident Reporting for UK Rail 1996 – present	G-K-N

178°	The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami 2005 – 2020	G
179	The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed 2020 – 2023	G-K
180°	Healthcare.gov Launch 2013	K-T-G
181°	EU Human Brain Project — Top-Down Vision That Unraveled 2013 – 2023	G-N
182°	Amazon Hiring AI — Trained Bias, Deprecated 2017 2014 – 2018	D-K-N
183°	Uber ATG / Tempe Fatality 2018	T-N-G-H
184	UK Post Office Horizon Scandal 1999 – 2015	G-H-K
185	Air Canada Chatbot Liability — Delegation Without Revocation 2022 – 2024	D-K-N
186°	Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity 2018 – 2022	D-G-N
187	COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate 2014 – 2018	D-K-N
189°	Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds 2014 – 2020	D-G-N
190°	Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment 2023	G-K-N
191	Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict 2016 – 2023	D-G-N
192°	Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model 2019 – 2021	D-K-N
193°	Bernard Madoff / SEC Failures 1992 – 2008	G-K-N
194°	Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox) 2001 – present	G-K-N
195°	Tesla Autopilot — Recurring Fatalities 2016 – present	T-N-G-H
196	Fintech Lending Fairness Audit — When Including Race Reduces Disparity 2025	D-G-N
197°	Predictive Policing — PredPol 2011 – present	G-H-D
198°	Launching the BRAIN Initiative — Governance of a Grand Challenge 2011 – 2015 – present	G-K-N
199	Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact 2023 / 2025	G-K-N
200°	CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver 2020 – 2024	G-K-D
201°	Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India 2018 – 2025	G-N-H
203	NYC Local Law 144 — Bias Audits as Governance Artifact 2023 – present	D-G-K
205	The Discipline We Build Next ongoing	T-K-N-G-H-D

Modes. T Training Gap · D Designed Out · N Normalization of Deviance · H Human–System Interface · G Governance & Trust · K Knowledge & Institutional Memory

Courses. The LDT / LENS courses each case serves as a worked example for are indexed separately in **The cases by LENS course** in the appendix.

Parts. The book is organized topically in seven parts — Healthcare & Patient Safety; Education, Training & the Learning Workforce; Aviation & Aerospace; Defense & National Security; Industry, Energy & Enterprise Systems; Disaster Prevention & Recovery; Algorithms, Governance & Public Systems. Each part splits into **What Fails** and **What Works — and the Frontier**; the open question closes the volume at Case 205.

Chapter 1

Healthcare & Patient Safety — What Fails

When care systems assume capability that was never specified, trained, or measured.

The harm arrived through the workflow, not around it.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

EHR / CPOE Implementation

H Human–System Interface **D** Designed Out **G** Governance & Trust

IMPACT ~\$30B federal investment under HITECH; documented increase in pediatric ICU mortality post–CPOE at one institution; ongoing usability harm at scale

IN BRIEF

The 2009 HITECH Act poured roughly \$30 billion into accelerating electronic health records, and adoption surged — but the systems had been built to billing and administrative specifications, not clinical workflow. New categories of harm followed. A disputed but canonical 2005 study found that deploying a commercial order–entry (CPOE) system at one pediatric ICU was associated with a more–than–doubling of mortality; later deployments showed mixed or improved results, yet the warning stuck. A 2023 survey across 200–plus hospitals found EHR usability is clinicians’ top complaint, and that usability scores track patient–safety outcomes. Specific harms recur — a missed abnormal result hidden behind a default, a fatal overdose an unactivated alert would have caught. There is still no regulatory framework monitoring EHR safety. It is the book’s live interface–mismatch case.

BACKGROUND

The 2009 HITECH Act authorized roughly \$30 billion in incentives to accelerate electronic health record adoption, and adoption surged. But the systems clinicians now had to use had been built to administrative and billing specifications, not to clinical workflow or human–factors specifications — a mismatch baked in before the first order was entered.⁰ Because the incentives rewarded adoption rather than usability, hospitals raced to install whatever cleared the bar, and the specification gap propagated to the bedside at national scale before anyone had to demonstrate that the tools supported the work of care rather than merely the work of billing.

WHAT HAPPENED

New categories of harm appeared alongside the new tools. A 2005 study reported that deploying a commercial computerized order–entry (CPOE) system at one pediatric ICU was associated with a more–than–doubling of mortality — a single–institution result that provoked debate and that later deployments elsewhere did not reproduce, but that became canonical as a warning that powerful tools without workflow integration can disrupt care at the moment of greatest acuity, when the seconds the interface adds to an order are the seconds a critical patient does not have.¹ Specific harms recur: a cancer treatment delayed for years because a default surfaced an old normal result instead of a recent abnormal one; a baby killed by an overdose an alert would have caught had it been switched on — each a case where the interface’s defaults and toggles, not the clinician’s competence, determined whether the right information reached the decision.²

THE INVESTIGATION


A 2023 KLAS Arch Collaborative survey across more than two hundred hospitals found EHR usability is consistently the top complaint of physicians, nurses, and pharmacists, and that end-user experience scores correlate with patient-safety outcomes — turning a stream of anecdotes into longitudinal evidence that usability is now itself a patient-safety variable. Spanning some two hundred hospitals, the survey converts what could be dismissed as one site's grievance into a reproducible signal, the kind of measurement that makes an interface problem legible as a safety problem rather than a satisfaction one.³

THE CAPABILITY GAP

EHR/CPOE is the canonical ongoing case of an interface designed for one set of requirements and deployed against another. The systems work for billing and administration because that was the design constraint; they fail for clinical workflow because clinical workflow was not. Decades and tens of billions of dollars later, usability remains among the largest contributors to in-system harm in U.S. healthcare — a cost that persists precisely because it was never a design requirement and so was never engineered out, only worked around by clinicians absorbing the friction shift after shift.⁴

AFTERMATH & REFORM

Usability and safety have risen on the agenda — AMA/Pew/MedStar usability work, KLAS benchmarking, safety-focused design guidance — but there is still no regulatory framework that monitors deployed EHR safety the way drugs or devices are monitored. The lesson is that buying the tool is not the same as engineering the capability: an interface deployed against the wrong specification keeps extracting a cost no one is formally counting, and without a monitoring regime the harm stays diffuse — spread across millions of encounters, attributable to no single failure, and therefore easy for the system to keep paying indefinitely.⁵



Reports of new types of errors directly related to EHR implementation — errors that can compromise quality of care and patient safety — have emerged.

PARAPHRASING SITTIG & SINGH, EHR-RELATED SAFETY RISKS, 2013

REFERENCES

1. Health Information Technology for Economic and Clinical Health (HITECH) Act (2009) — the \$30B EHR incentive program.

2. Y. Han et al., “Unexpected Increased Mortality After Implementation of a Commercially Sold CPOE System,” *Pediatrics* 116(6): 1506–1512 (2005) — the disputed pediatric-ICU mortality result.

3. KLAS Arch Collaborative, EHR usability and safety surveys (2023) — usability as the top clinician complaint, correlated with safety outcomes.

4. D. Sittig & H. Singh, “Defining health information technology-related errors,” and related EHR-safety work (2011–2013).

5. AHRQ Patient Safety Network, EHR-related harm cases — the default-result delay and the unactivated-alert overdose.

6. AMA / Pew / MedStar, EHR Usability and Patient Safety (2018); R. Wachter, *The Digital Doctor* (2015).

THE LEARNING ENGINEERING LENS

LE INSIGHT

EHR/CPOE is the canonical ongoing case of an interface designed for one set of requirements and deployed against another. The systems work for billing and administration. They fail for clinical workflow because clinical workflow was not the design constraint. Thirty billion dollars later, usability remains among the largest contributors to in-system harm in U.S. healthcare.

LENS APPROACH

LENS treats EHR/CPOE in LEN 7 as the live, ongoing example of Design-Out and Interface failure under governance opacity, in LEN 2 as a human-AI teaming problem (alert fatigue, default surfacing), and in LEN 9 as a computational systems problem (the alerting architecture itself is a learnable model). The Han 2005 pediatric ICU finding is the durable warning; the KLAS 2023 surveys across two hundred hospitals are the contemporary, longitudinal evidence that usability is now itself a patient-safety variable.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	→ H D G	→ D1 D2 D3 D4 D5	3	3.1	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Write clinical workflow and human-factors performance into the procurement specification, so a system that serves billing but not bedside care cannot clear the buy. ▶ Validate defaults, alerts, and result-surfacing against real clinical scenarios before deployment, since it was the default and the unactivated alert — not clinician skill — that determined the harm. ▶ Tie any adoption incentive to demonstrated usability, so the money rewards a tool that supports care rather than merely one that is installed.	▶ Stand up a post-deployment safety-monitoring regime for fielded EHRs comparable to how drugs and devices are surveilled, with thresholds and authority to force changes. ▶ Track end-user experience scores as a patient-safety variable, using the cross-institution evidence that usability correlates with outcomes. ▶ Audit recurring interface-driven harms — surfaced defaults, disabled alerts — and feed them back into vendor design requirements rather than absorbing them as clinician burden.

REFLECTION QUESTIONS

1. What is the equivalent system in your domain that was designed for one specification and deployed against another? How would you measure the harm?

2. Design the regulatory architecture that would surface EHR safety harms at scale. Be specific about signal, threshold, and authority.

3. The Han 2005 mortality result was disputed; the 2023 KLAS surveys are consistent across two hundred hospitals. What ongoing measurement architecture would have to exist for the equivalent emerging clinical-AI deployment in your domain to be evaluated honestly while in use?

WHO BUILDS THIS human factors & interaction design · systems & human-factors engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · task analysis · design prototyping · governance frameworks · bias & evidence audits

Watson for Oncology — Delegation Marketed Ahead of Capability

D Designed Out **K** Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER journalism-tier — source confidence flagged; future validation ongoing.

IMPACT IBM marketed Watson for Oncology with tumor-board concordance rates as high as 96%; STAT News investigation (Ross & Swetlitz, 2017–2018) and internal IBM documents documented unsafe and incorrect cancer-treatment recommendations and that the system was trained on a small number of synthetic cases rather than real patient outcomes; hospitals worldwide had purchased the tool on the marketed concordance claim

IN BRIEF

Watson for Oncology was IBM's heavily marketed clinical-decision- support system, sold to hospitals worldwide as a cancer- treatment recommendation engine whose concordance with expert tumor boards was advertised at rates as high as 96%. The capability the system was sold under was AI-grade treatment recommendation across cancer types, delivered with the institutional credibility of IBM and MD Anderson partnerships. The capability the system actually carried — documented in the investigative reporting by Casey Ross and Ike Swetlitz at STAT News through 2017 and 2018, drawing on leaked IBM internal documents — was substantially smaller. Watson was trained on a small number of synthetic cases curated by Memorial Sloan Kettering oncologists rather than on real patient outcomes, generated unsafe and incorrect recommendations in some documented cases, and the internal IBM record showed engineers raising concerns about the gap between marketed concordance claims and what the system actually delivered. Hospitals around the world had purchased Watson on the marketed capability; adoption stalled, deployments were discontinued (Manipal Hospitals in India ended its contract in 2018), and IBM ultimately wound down and sold off Watson Health. The case is the canonical instance in the AI-delegation typology of capability marketed ahead of capability validated. The evidence-tier flag rendered under the title is binding: STAT News journalism is the primary public source, and the academic record cites it secondarily. "Future validation ongoing" is the standing language; independent peer-reviewed evaluation of Watson for Oncology's operating accuracy at the level the marketing claimed was never produced.

BACKGROUND

Watson for Oncology was the most publicly visible commercial clinical-AI deployment of the mid-2010s. IBM had positioned Watson — originally the Jeopardy-winning question-answering system — as a healthcare engine after the 2011 Jeopardy performance, and oncology was the flagship domain. The product was sold as a cancer-treatment recommendation

system: a clinician would enter the patient's case, Watson would query its knowledge base against the training data provided by Memorial Sloan Kettering oncologists, and the system would return ranked treatment recommendations annotated with evidence citations. The marketed capability was treatment recommendation that concorded with expert tumor boards at rates as high as 96%. Hospitals worldwide purchased the system on that claim.⁰

WHAT HAPPENED

Casey Ross and Ike Swetlitz at STAT News produced the investigation of record across 2017 and 2018. Drawing on leaked IBM internal documents and interviews with engineers, oncologists, and former IBM staff, the reporting documented a substantially different operating picture. Watson had been trained not on real patient outcomes — the standard a treatment-recommendation engine would have to clear — but on a small number of synthetic cases curated by MSK oncologists to reflect what those oncologists would have done. The training corpus was a model of expert opinion at one institution, not an outcome-anchored learning record. In documented cases the system generated unsafe and incorrect treatment recommendations. Internal IBM documents showed engineers raising concerns about the gap between marketed capability and operating reality.¹

THE INVESTIGATION

The structural failure pattern the case carries is delegation of consequential authority — cancer-treatment recommendation — to a tool whose validated capability was substantially smaller than its marketed capability. The procurement mechanism was the carrier: hospitals purchased Watson through capital and IT procurement processes that evaluated the marketed concordance claims and the IBM/MSK institutional credentials, but did not independently validate operating accuracy on the receiving institution's patient population. The capability the procurement was buying — accurate, real-world treatment recommendation at the marketed concordance rate — was not the capability the system carried. Deployments were discontinued as the gap surfaced — Manipal Hospitals in India ended its contract, and IBM ultimately wound down and divested Watson Health.²

THE CAPABILITY GAP

The evidence-tier flag rendered under the title is binding and survives into the case verbatim. The STAT News investigation by Ross and Swetlitz is the primary public source, and the academic record cites it secondarily. No independent peer-reviewed evaluation of Watson for Oncology's operating accuracy at the level the marketing claimed was ever produced — partly because the system was proprietary and partly because the deployments were wound down before the kind of prospective study that would have established the operating record could be conducted. The case is teachable on the structural failure pattern at journalism-grade evidence; the standing "source confidence flagged; future validation ongoing" language is the honest tier the case rests on.³

AFTERMATH & REFORM

In the AI-delegation typology (Cases 20 TREWS, 102 Epic Sepsis Model, 111 SyRI, and this case), Watson is the marketing-ahead-of-capability failure mode: the procurement decision rode on marketed concordance rates that were not validated at the operating institutions, and the institutional credentials did not substitute for the validation. TREWS shows delegation with strong outcome evidence works; Epic Sepsis shows delegation

without validation fails by accuracy disconfirmation; SyRI shows delegation halted on rights grounds; Watson shows delegation driven by procurement and marketing rather than by validated capability. The four together teach the typology of when AI delegation is and is not legitimate. Watson is the case where the marketing finance — the institutional credibility and the headline concordance number — substituted for the evidence the procurement should have required, and the capability the hospitals thought they were buying was not the capability they got.

The capability the procurement was buying — accurate, real-world treatment recommendation at the marketed concordance rate — was not the capability the system carried.

EDITORS' SYNTHESIS OF ROSS & SWETLITZ (STAT NEWS 2017–2018).

REFERENCES

1. Ross & Swetlitz (2017–2018), “IBM Watson recommended unsafe and incorrect cancer treatments” and adjacent investigations, STAT News — the load-bearing primary source; investigative journalism drawing on leaked IBM internal documents.

2. Strickland (2019), “How IBM Watson Overpromised and Underdelivered on AI Health Care,” IEEE Spectrum — independent retrospective synthesis of the public record.

3. Topol (2019), Deep Medicine, Basic Books — secondary academic situating of Watson within the broader clinical-AI delegation discourse.

4. v2 paired cases: TREWS (Case 20), Epic Sepsis Model (Case 5), SyRI (Case 189) — the AI-delegation typology.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Watson for Oncology is the canonical instance of clinical-AI delegation marketed ahead of validated capability. The operating capability was substantially smaller than the marketed concordance claims; the procurement mechanism evaluated the marketing rather than the validation; the deployments were wound down as the gap surfaced. The evidence-tier flag is binding; journalism-grade reporting is the load-bearing source, and future validation ongoing remains the honest qualifier.

LENS APPROACH

Watson for Oncology is the marketing-ahead-of-capability AI-delegation failure (induced 1.1; LENS D3/PT6) — Domain 3 for **Delegation with revocation** (the discontinued deployments are the revocation half); Domain 4 for requirements-vs- marketing. Pair with Cases 20, 5, 189. Evidence-tier flag binds.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE/ ONCOLOGY	D K N	D1 D2 D3 D4 D5	6	2.4	1 · 4
COMMERCIAL AI		Human-System Collaboration			
PROCUREMENT					

APPROACHES TO CONSIDER

- DURING DEVELOPMENT

▶ Separate marketed capability from validated capability in any procurement decision involving consequential delegation; the marketed concordance number is the carrier of the procurement story, not the evidence the procurement should rest on.

▶ Require independent operating-accuracy validation on the receiving institution's patient population before
- IN OPERATION / AFTER

▶ Render the evidence tier honestly when the primary public source is journalism; the STAT investigation is the load-bearing source, and the academic record's secondary citation is the accurate description, not “peer-reviewed evidence.”

▶ Carry the AI-delegation typology — TREWS, Epic Sepsis, SyRI, Watson — as a unit in any curricular

delegation; the AI-delegation literature now treats this as the procurement floor, and Watson is one of the reasons.

- ▶ Anchor the procurement evidence to outcome-anchored training, not to expert-curated synthetic cases; expert opinion at one institution is not outcome data the procurement can rely on.

discussion of when delegation is legitimate; the four-case set teaches the typology more clearly than any single case.

- ▶ Treat the discontinuation of Watson for Oncology deployments as the case's own correction signal; the institutional decisions to discontinue are themselves evidence about what the procurement should have required up front.

REFLECTION QUESTIONS

1. Identify a procurement decision in your domain that delegated consequential authority to an automated system. Did the procurement evaluate marketed capability claims or validated capability on the receiving institution's data? Where is the gap most likely to have been smoothed by institutional credentials standing in for evidence?
2. Specify what operating-accuracy validation you would require before delegating: which prospective study, on which population, with which comparison and which outcome metric. The AI-delegation typology has established this as the procurement floor; the question is whether your domain's procurement processes meet it.
3. The evidence-tier flag under this case's title is binding because the primary public source is journalism. Identify a case in your domain where the strongest available evidence is journalism-grade. What is the responsible drafting discipline that surfaces the tier rather than implying peer-reviewed validation?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

Mid Staffordshire NHS Foundation Trust

G Governance & Trust **N** Normalization of Deviance **K** Knowledge & Institutional Memory

IMPACT Excess deaths at Stafford Hospital documented across years; the Francis Inquiry produced 290 recommendations and restructured UK patient-safety governance

IN BRIEF

Between roughly 2005 and 2009, Stafford Hospital, run by the Mid Staffordshire NHS Foundation Trust, subjected patients to appalling neglect — left without food, water, or basic care — while mortality ran above expected. The trust had been chasing financial targets that depended on staffing cuts. Robert Francis QC's public inquiry produced a 2,000-page report and 290 recommendations, identifying the structural problem as the gap between the trust's reported performance and patients' actual experience: every governance layer above the ward received reports that targets were being met, and none checked those reports against what was happening to patients. Mid Staffs is the dataset's strongest case for the harm done when measurement and reality diverge over years.

BACKGROUND

The Mid Staffordshire NHS Foundation Trust ran Stafford Hospital in England. Pursuing Foundation Trust status and the financial targets that came with it, the board cut staffing — and the cuts fell on the wards. The institution's reported performance, the numbers that travelled upward, stayed on target.⁰ The targets the board chased were financial and procedural, so cutting ward staff improved the very figures the trust was measured on even as it removed the people on whom patient care directly depended.

WHAT HAPPENED

On the wards the reality was appalling. Patients were left in their own excrement, denied food and water, given the wrong medication or none, for years; mortality ran substantially above expected. The harm was not a single incident but a sustained condition — visible to anyone on the ward and invisible in the reports that left it.¹ Because the suffering was a continuous state rather than a nameable event, it never generated the kind of discrete incident a reporting system is built to catch, and so it accumulated for years beneath numbers that stayed reassuringly on target.

THE INVESTIGATION

The Healthcare Commission investigated; the layers above did not. Robert Francis QC's public inquiry ran to some 2,000 pages and 290 recommendations.² Its structural finding was the divergence between reported performance and patient experience: every governance layer above the ward had received reports that the hospital was meeting its targets,


and none had checked them against what was happening to patients. “The system as a whole failed in its most essential duty — to protect patients from unacceptable risks of harm.”³ The phrase “the system as a whole” located the failure deliberately above any single ward or manager: no one layer was solely at fault, because each had trusted the layer below to be reporting reality rather than targets.

THE CAPABILITY GAP

Mid Staffs is the British analog of the VA wait-time scandal (Case 7): measurement and reality diverged over years while every layer of governance acted on the measurement. The capability that was missing was not clinical skill on the ward but the institutional habit of testing whether the numbers corresponded to the patients — a check no layer above the ward performed.⁴ Each layer reasonably assumed the check belonged to someone else, so the verification that the report matched the patient fell into the gap between layers — exactly the place a reporting chain built only to pass numbers upward is structurally unequipped to look.

AFTERMATH & REFORM

The Francis Inquiry restructured how the NHS treats patient safety, and the Berwick review that followed pressed for a culture of learning over targets. The lesson is the measurement one in its starkest form: a reporting chain can run clean from ward to Department of Health while, underneath it, the thing being reported on quietly fails.⁵ Berwick’s framing — learning over targets — named the deeper correction: as long as the target is the thing the institution rewards, the report will describe the target, and only a culture that prizes finding the gap will keep checking the report against the patient.



The system as a whole failed in its most essential duty — to protect patients from unacceptable risks of harm.

ROBERT FRANCIS QC, REPORT OF THE MID STAFFORDSHIRE NHS FOUNDATION TRUST PUBLIC INQUIRY, 2013

REFERENCES

1. R. Francis QC, Report of the Mid Staffordshire NHS Foundation Trust Public Inquiry (2013) — the staffing cuts and Foundation Trust targets.
4. Francis QC (2013) — “the system as a whole failed in its most essential duty — to protect patients from unacceptable risks of harm” (quoted).

2. Healthcare Commission, Investigation into Mid Staffordshire NHS Foundation Trust (2009) — ward conditions and excess mortality.
3. Francis QC (2013) — the 2,000-page report and 290 recommendations.
5. D. Berwick, A Promise to Learn — A Commitment to Act (National Advisory Group on the Safety of Patients in England, 2013).
6. K. Walshe & J. Higgins (2002) on NHS safety governance; cf. the VA wait-time scandal (Case 7).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Mid Staffordshire is the British analog of the VA wait-time scandal (Case 7). Measurement and reality diverged over years; every layer of governance above the operating environment was acting on the measurement; patients paid the cost of the divergence. The Francis Inquiry recommendations changed how the NHS thinks about patient safety as an institutional capability.

LENS APPROACH

LENS uses Mid Staffs in LEN 4 for the divergence-of-measurement problem and in LEN 7 for the governance failure across multiple layers. Studio projects examine the Francis recommendations as a template for institutional reform deliverables.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	→ G N K	→ D1 D2 D3 D4 D5	5	2.1	4
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design targets so that gaming them (e.g., cutting ward staff) cannot improve the metric while degrading the outcome it stands for.
- ▶ Build a direct patient-experience signal — independent of the financial and procedural targets — into the reporting chain from the start.
- ▶ Assign explicit ownership of the report-versus-reality check at each layer, so verification cannot fall into the gap between them.

IN OPERATION / AFTER

- ▶ Audit sustained conditions, not just discrete incidents, since continuous harm never trips an event-based reporting system.
- ▶ Cross-check upward reports against ground truth at the ward periodically, treating a clean report as a hypothesis to be tested.
- ▶ Reward a culture of learning over target attainment, so the institution keeps looking for the gap rather than describing the target.

REFLECTION QUESTIONS

1. Identify a multi-layer reporting chain in your domain. What would it take for the top layer to know whether the reports correspond to reality?
2. The Francis Inquiry produced 290 recommendations. Pick five that you think were most load-bearing and explain why.
3. The verification fell into the gap between layers because each assumed it belonged to someone else. In your domain, who explicitly owns the check that a report matches reality — and how would you know if no one does?

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards · knowledge bases · data pipelines

CASE 7 GOVERNMENT

2014

VA Wait-Time Scandal

G Governance & Trust **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Veterans died waiting for care; ~120,000 waiting 90+ days or never got onto lists; staff falsified records

IN BRIEF

In 2014 the VA Inspector General found that staff at the Phoenix VA — and then nationwide — had created secret waiting lists and falsified appointment data to hide that veterans were waiting weeks or months for care; some died waiting while the system reported success. A 14-day access target, unrealistic given staffing, was met by hiding reality rather than surfacing it. The warning signs ran back fifteen years: GAO had flagged data-reliability problems since 2000, the IG since 2005, with no systemic fix — and schedulers, the staff operating the measurement, are among the VA's highest-turnover roles, so the institution kept losing the knowledge even to see it was lying to itself. The VA case is the book's canonical example of normalization of deviance applied to measurement itself.

BACKGROUND

The Veterans Health Administration measured access to care against a 14-day appointment target — a target that, given staffing, was often unrealistic to meet honestly. Schedulers, the staff who operate that measurement, are among the VA's highest-turnover positions.⁰ The target functioned as the headline number leadership watched, so the pressure to show 14-day compliance bore down hardest on the very front-line role least equipped, through constant churn, to record the access data accurately or to question what the number was leaving out.

WHAT HAPPENED

In 2014 the VA Office of Inspector General found that staff at the Phoenix VA — and then across the system — had created secret waiting lists and falsified appointment data to hide that veterans were waiting weeks or months for care. Some veterans died while waiting — though the Inspector General could not conclusively tie the deaths to the scheduling delays — inside a system that, by its own metrics, was succeeding.¹ The secret lists were the mechanism by which an impossible target was reconciled with reality: official records showed appointments inside 14 days while the true wait accumulated off the books, so the more the metric was gamed, the more confidently the reporting line above it declared the access problem solved.

THE INVESTIGATION

The warning signs went back fifteen years: GAO had flagged data-reliability concerns since 2000, and the VA IG had identified problems in 2005, 2007, and 2008, with the incoming


administration warned in 2008 — none of which produced systemic change.² GAO named scheduler training as a root cause: with schedulers among the top-ten highest-turnover roles, the institution perpetually lost the knowledge required even to run the measurement honestly, and five years on still reported data-reliability concerns.³ Each warning had named a real defect in how access was recorded, yet because the schedulers who held the practical knowledge of the system kept turning over, every wave of findings landed on a workforce that had to relearn the instrument from scratch, so the same defect resurfaced report after report.

THE CAPABILITY GAP

This is normalization of deviance applied to measurement itself. When the measurement system cannot capture reality — and the people operating it churn before the gaming can be unlearned — the gap between reported and actual performance becomes invisible. Veterans died inside a system whose numbers said it was fine, which is the lethal form of the evidence problem this book treats as a design failure, not a reporting one.⁴ An institution that cannot retain the staff who run its measurement loses not only accuracy but the memory that the number was ever wrong, so the gaming stops registering as deviance at all and hardens into the ordinary way the work is done.

AFTERMATH & REFORM

The scandal forced resignations, the Veterans Access, Choice, and Accountability Act (2014), and a long, uneven effort to rebuild scheduling and data integrity.⁵ Its lesson is that decision-grade evidence is a design requirement: an institution must be able to surface its own failures without relying on the very people incentivized — and too transient — to hide them. The Choice Act bought access by routing care outside the VA, but the underlying capability — a measurement the institution could trust even as its schedulers churned — was the harder thing to rebuild, and the data-reliability concerns that persisted for years afterward show why.



Schedulers are among the top ten highest-turnover positions in the VA.

DEBRA DRAPER (GAO DIRECTOR OF HEALTH CARE), HOUSE VA COMMITTEE TESTIMONY, GAO-19-687T, JULY 2019

REFERENCES

1. VA Office of Inspector General, Report 14-02603-267 (2014) — secret waiting lists and falsified appointment data.
2. GAO Veterans Health Administration reports (2000–2019) — fifteen years of data-reliability warnings.
3. VA OIG reports (2005, 2007, 2008) — prior, unactioned findings.
4. D. Draper (GAO), House VA Committee testimony, GAO-19-687T (2019) — “schedulers are among the top ten highest-turnover positions in the VA” (quoted).
5. Veterans Access, Choice, and Accountability Act (2014) — the legislative response.
6. C. Argyris & D. Schön, Organizational Learning (1978); measurement-gaming literature.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The VA case is the canonical example of measurement failure as a capability failure. The system that should have surfaced the gap instead generated reports that hid it. The turnover among schedulers — the human capability operating the measurement instrument — meant the institution lost the knowledge to even notice it was lying to itself. The case stands as the strongest argument in this book for treating decision-grade evidence as a design requirement, not a reporting requirement.

LENS APPROACH

LENS treats the VA case in LEN 4 as the canonical evidence-gap case (the measurement system itself was the source of harm), in LEN 7 as a governance failure (multiple warnings unactioned over fifteen years), and in LEN 8 as a knowledge-loss case via turnover.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
GOVERNMENT	→ G K N	→ D1 D2 D3 D4 D5	5	2.2	4
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Set access targets against actual staffing capacity, so the metric cannot be met only by falsifying it.
- ▶ Design the scheduling instrument so correct data entry is the path of least resistance, not a discipline that churning staff must be retrained into.
- ▶ Build an independent read on real wait times — separate from the staff incentivized to report them — into the measurement architecture from the start.

IN OPERATION / AFTER

- ▶ Audit reported access against an out-of-band signal (direct veteran survey, third-party booking records) with authority to act when the two diverge.
- ▶ Treat scheduler turnover as a measurement risk: monitor it, and protect the knowledge of how to run the instrument honestly against constant churn.
- ▶ Track whether old data-reliability findings keep recurring — a repeat finding is the signal that the institution is relearning the same gap rather than closing it.

REFLECTION QUESTIONS

1. Identify a measurement system in your domain that is also operated by a high-turnover role. What is the institutional risk that the system stops measuring reality?
2. Design the evidence pipeline that would have surfaced the Phoenix VA gap without relying on the people who were gaming the metrics.
3. The 14-day target was unrealistic given staffing, so it was met by hiding reality. Identify a target in your organization that the people measured against it cannot honestly meet — and design the correction that surfaces the gap rather than burying it.

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

Medical Errors as Systemic Failure

T Training Gap **H** Human-System Interface **N** Normalization of Deviance
K Knowledge & Institutional Memory **G** Governance & Trust

DISCLOSURE Institutional overlap: an editor shares an institution (Johns Hopkins) with an author of this work; no editor was personally involved. The disputed mortality estimate is contested in the case itself.

IMPACT IOM 1999 estimate of 44,000–98,000 US deaths/year from medical error; Makary & Daniel (2016) estimate of ~250,000 deaths/year — substantively contested on methodological grounds; 2023 NEJM inpatient-harm study confirms persistence; 25-year reform arc with bounded successes and an unmoved population count

IN BRIEF

Medical error in the United States is not a single incident but a systemic condition that the system's own measurement instruments cannot see. The Institute of Medicine's *To Err Is Human* (1999) raised the alarm with a 44,000–98,000 annual-deaths estimate and reframed harm as a systems problem rather than an individual one; *Crossing the Quality Chasm* (2001) and the 2015 *Improving Diagnosis in Health Care* sequel extended the program. In 2016, Makary and Daniel of Johns Hopkins published a BMJ analysis estimating 250,000 deaths a year — which would make medical error the third leading cause of death behind heart disease and cancer. Death certificates do not record medical error as a cause, so the problem is structurally invisible to the systems meant to track it. The 250,000 figure has been substantively contested on methodological grounds (Shojania & Dixon-Woods, BMJ Quality & Safety, 2017; Carr in Health Affairs), a dispute that is itself a worked example of the gap-attribution problem. The field's 25-year arc shows bounded interventions that work — computerized order entry, handoff protocols, the WHO surgical checklist (Case 23), the Keystone ICU project (Case 19), TeamSTEPPS (Case 39) — alongside a population-scale mortality count that has resisted both intervention and precise estimation.

BACKGROUND

The Institute of Medicine's *To Err Is Human* (1999) was the field-defining moment: it estimated 44,000–98,000 deaths annually from medical error in US hospitals — at the lower bound, more Americans than die in motor-vehicle accidents — and made the explicit case that the problem was a systems problem, not an individual-clinician problem. The 2001 sequel *Crossing the Quality Chasm* set six aims for the redesign of care, and the 2015 sequel *Improving Diagnosis in Health Care* extended the framing to diagnostic error. The framing of harm as systemic mattered: it created the cultural permission for non-punitive incident reporting, root-cause analysis, and the broader patient-safety movement that followed.^o Yet modern medicine's core mortality-measurement instrument — the death certificate —

records a proximate physiological cause and has no field for medical error. Because the certificate is keyed to an ICD billing taxonomy built for disease, a death set in motion by a care-process breakdown is recorded under whatever organ ultimately failed, and the causal role of the system disappears into the physiology.

WHAT HAPPENED

In 2016, surgeon Martin Makary and Michael Daniel published an analysis in the *BMJ* estimating that medical errors cause more than 250,000 deaths a year in the United States — which would rank third behind heart disease and cancer. Their core claim was that “people don’t just die from heart attacks and bacteria, they die from system-wide failings and poorly coordinated care.” By relocating the cause from the individual clinician to the coordination of care, the framing recast a ledger of isolated mistakes as a single population-scale failure mode the existing statistics were never built to count. The number itself was computed by extrapolating from four prior studies — the IOM 1999 estimate, the 2010 OIG Medicare adverse-events study, the Landrigan 2010 *NEJM* study, and the Classen 2011 Global Trigger Tool study — to the contemporary inpatient population, an extrapolation the authors acknowledged as approximate.¹

THE INVESTIGATION


The 250,000 estimate was substantively contested almost immediately. Shojania and Dixon-Woods, writing in *BMJ Quality & Safety* in 2017, challenged both the extrapolation and the attribution method, arguing that counting deaths “due to” error is far harder than a single headline number implies. A *Health Affairs* commentary by Carr and a series of letters in the *BMJ* pressed the same point: the source studies use different definitions of “preventable death,” the extrapolation rests on prior estimates already known to be unstable, and the ranking against CDC cause-of-death categories is methodologically incompatible (CDC counts use ICD codes; the Makary estimate does not).² The objection turned on counterfactual attribution — how confidently one can say a frail, already-dying patient would have survived but for the error — a judgment that resists the clean tallying a headline number demands. The dispute is itself a worked example of the gap-attribution problem: how much of a counted death is the learning system, how much the system design, how much the underlying disease.³

THE CAPABILITY GAP

The deeper failure is one of measurement: a system that cannot see its own failure modes cannot manage them. With no field on the certificate, no reliable count, and no agreed attribution method, every safety program competes for resources against a harm that the official record renders invisible, so even effective interventions struggle to prove their worth at population scale. The bounded interventions that have worked — computerized physician order entry, structured handoff protocols (I-PASS), the WHO Surgical Checklist (Case 23), the Keystone ICU central-line bundle (Case 19), TeamSTEPPS (Case 39) — each move a specific harm in a specific setting, and each can prove it. What none can prove is movement on the population count, because the population count does not exist in a form precise enough to be moved. The missing capability is an instrument that captures medical error as a tracked cause of harm — and an attribution method robust enough that the resulting number can guide intervention rather than fuel a methodological stalemate.⁴

AFTERMATH & REFORM

The 25-year arc since *To Err Is Human* is the case's most important teaching artifact. The IOM framing catalyzed a patient-safety movement, an Agency for Healthcare Research and Quality patient-safety program, and the bounded interventions noted above. Yet later work, including the 2023 NEJM study by Bates et al. of inpatient harm across eleven Massachusetts hospitals, confirms that the problem persists at scale: about a quarter of admissions involved an adverse event, and roughly a quarter of those events were preventable. The system has not built the instrument the original report implied it would need — no national active-surveillance system for inpatient harm, no death-certificate field for care-process failure, no agreed attribution method for “deaths due to error.” The interventions exist; the measurement and the implementation still lag, and the headline mortality kept escaping measurement even as bounded harms fell. The case is the canonical worked example of the gap-attribution problem at population scale, and a standing reminder that an instrument the system cannot see through is one it cannot govern.⁵



People don't just die from heart attacks and bacteria, they die from system-wide failings and poorly coordinated care.

MARTIN MAKARY, PRESS STATEMENTS ACCOMPANYING MAKARY & DANIEL, BMJ (2016)

REFERENCES

1. Institute of Medicine, *To Err Is Human: Building a Safer Health System* (1999); *Crossing the Quality Chasm* (2001); *Improving Diagnosis in Health Care* (2015) — the field-defining trilogy and the 44,000–98,000 estimate; the systems framing.
2. Makary, M. & Daniel, M. (2016), “Medical error — the third leading cause of death in the US,” *BMJ* 353:i2139 — the 250,000 estimate, the quoted framing, the extrapolation from four prior studies (IOM 1999; OIG 2010; Landrigan 2010; Classen 2011).
3. Shojania, K. & Dixon-Woods, M. (2017), “Estimating deaths due to medical error: the ongoing controversy and why it matters,” *BMJ Quality & Safety* 26(5):423–428; with companion commentaries including Carr in *Health Affairs* — methodological contestation of the Makary extrapolation and CDC-ranking comparison.
4. Makary & Daniel (2016), *BMJ* — death certificates do not capture medical error as a cause; ICD billing taxonomy as the structural reason.
5. Bates, D. W., Levine, D. M., Salmasian, H., et al. (2023), “The Safety of Inpatient Health Care,” *NEJM* 388(2):142–153 — eleven-hospital Massachusetts cohort; adverse events in 25% of admissions, 25% of those preventable; persistence of harm at scale.
6. Agency for Healthcare Research and Quality, *National Healthcare Quality and Disparities Reports* (annual); CDC WONDER ICD-coded mortality data — institutional context for the missing national active-surveillance instrument.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Makary data is the anchor evidence for the LENS argument because of its scale, its provenance (Johns Hopkins), and its framing — system failure rather than individual error. The seventeen years between **To Err Is Human** and Makary's reassessment is also the implementation gap of Case 13 in another guise — and the cost of leaving it open is measured in lives at population scale.

LENS APPROACH

Medical error is the central evidence anchor of the curriculum (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation; CLO-4) for the foundational measurement problem: the death certificate cannot record the failure mode it is supposed to govern, and the contested mortality estimates show what happens when an instrument's gap is filled by extrapolation. LENS uses it in Domain 1 (Systems Analysis) as the foundational problem statement of the program and in Domain 5 (Navigating Sociotechnical Constraints) for industry-level institution building. The 25-year arc pairs the case with the WHO Surgical Checklist (Case 23), Keystone ICU (Case 19), and TeamSTEPPS (Case 39) as the bounded interventions that worked; pair with Vioxx (Case 9) at the population-scale-surveillance layer.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	→ T H N K G →	D1 D2 D3 D4 D5	5	2.1	4
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Give the mortality-recording instrument an explicit field for care-process failure, so a system-caused death is captured, not absorbed into the proximate cause.
- ▶ Specify the attribution method up front — the counterfactual test for “due to” error — so the count guides action rather than collapsing into dispute.
- ▶ Pair each safety intervention with the population measure it should move, gating reforms on demonstrated effect.

IN OPERATION / AFTER

- ▶ Audit reported mortality against an independent count of care-related harm.
- ▶ Use active surveillance of inpatient harm rather than waiting for the death certificate.
- ▶ Hold the count and the intervention to the same measurement discipline, closing the implementation gap.

REFLECTION QUESTIONS

1. Identify a measurement instrument in your domain that systematically fails to capture the failure modes it should be designed to surface. What would it cost to fix?
2. Two hundred fifty thousand deaths a year is the third leading cause of death in the U.S. Design the measurement and intervention regime that would shift the curve over a five-year horizon. Estimate the deliverable and the evidence.
3. Makary and Shojania disagreed not on whether error kills but on how to attribute a death to it. Specify an attribution method robust enough to survive that dispute — and name who would hold the authority to act on the number it produces.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design · safety science & organizational analysis · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs · incident analysis · safety dashboards · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

Vioxx Withdrawal

G Governance & Trust **D** Designed Out

IMPACT Tens of thousands of excess cardiovascular events estimated (FDA's Graham, 2004; Lancet 2005); Merck withdrew Vioxx in September 2004; ~\$4.85B settlement; FDA Amendments Act of 2007 and the Sentinel Initiative (2008) institutionalized active post-market surveillance

IN BRIEF

Merck's painkiller Vioxx (rofecoxib) was approved by the FDA in May 1999 and prescribed to an estimated 80 million people worldwide. The 2000 VIGOR trial (Bombardier et al., NEJM) found a roughly five-fold higher rate of myocardial infarction in the Vioxx arm than in the naproxen comparator — a signal Merck and many readers attributed to a hypothesized cardio-protective effect of naproxen rather than a cardiovascular risk of Vioxx, a reading that required the absent placebo arm to be true. Not until the placebo-controlled APPROVe trial (Bresalier et al., NEJM) was halted in September 2004 was the risk established and the drug withdrawn. For nearly four years the disclosure architecture between Merck's internal data, FDA reviewers, and prescribers failed to surface the magnitude of the danger to a decision point. FDA Office of Drug Safety scientist David Graham testified to the Senate Finance Committee in November 2004 that the risk had been visible well before withdrawal, and estimated 88,000–139,000 excess cardiovascular events attributable to Vioxx; the 2005 Graham et al. Lancet analysis put the figure in similar range. Merck litigation discovery produced internal Merck communications used in the Senate hearings to argue that publication-bias and authorship-by-Merck-employees patterns had suppressed the cardiovascular signal in the published record. The reforms that followed — Risk Evaluation and Mitigation Strategies (REMS), the FDA Amendments Act of 2007, and the Sentinel Initiative (2008) — built the active post-market surveillance architecture that had not previously existed.

BACKGROUND

Vioxx, a selective COX-2 inhibitor for arthritis pain, was approved by the FDA in May 1999 and became one of the most widely prescribed drugs of its era — an estimated 80 million people worldwide were exposed before withdrawal. Detecting rare or delayed harms in a drug at that scale depends on post-market surveillance — aggregating adverse-event data after approval — a function that, at the time, was thin relative to the size of the exposed population. The FDA's MedWatch system relied on voluntary submission from clinicians; no active query of the underlying claims and electronic-health-record data was operational. A risk too small to surface in a pre-approval trial becomes a large absolute toll once a drug reaches millions, which is precisely the regime that demands strong after-market monitoring — and precisely the regime in which voluntary reporting fails to detect it.^o

WHAT HAPPENED

The VIGOR trial (Bombardier et al., NEJM 2000) reported about a five-fold higher rate of myocardial infarction in 8,076 patients taking Vioxx than in those taking naproxen. Merck and many readers interpreted the gap as naproxen being cardio-protective rather than Vioxx being harmful. That reading was not absurd — it was the more comfortable of two explanations for the same numbers — but it required the absent placebo arm to be true, and the data could not adjudicate between the readings. The placebo-controlled APPROVE trial (Bresalier et al., NEJM 2005), originally designed to test Vioxx for colorectal-polyp prevention, was terminated early in September 2004 when its data safety monitoring board observed a doubling of cardiovascular events in the Vioxx arm relative to placebo. Merck withdrew the drug worldwide within days.¹

THE INVESTIGATION

Senate Finance Committee hearings in November 2004 and the subsequent FDA Office of Inspector General review found that signals of cardiovascular harm had been present in the trial record for years before withdrawal. FDA Office of Drug Safety scientist David Graham testified under oath that the cardiovascular risk had been visible to him by 2000, that he had been pressured by FDA management not to publish his estimate, and that he believed Vioxx had caused 88,000–139,000 excess heart attacks and strokes in the United States, of which 30–40% were probably fatal. The Graham et al. Lancet 2005 analysis, using Kaiser Permanente data, produced a population-level estimate in similar range. Merck litigation discovery, made public through New Jersey and federal court filings and reported in the NEJM editorial trail, included internal Merck communications and ghost-authorship patterns in published Vioxx papers; a 2008 JAMA analysis by Ross et al. documented the publication-bias and authorship patterns directly.² The harm was not hidden in some unmeasured corner — it sat in the trial record the whole time, waiting for an architecture that would carry it to a decision rather than leave it to interpretation.³


THE CAPABILITY GAP

The capability gap was in post-market surveillance: the disclosure architecture between manufacturer-held data, FDA reviewers, and prescribers did not aggregate the signal to a decision boundary for nearly four years. The FDA's MedWatch adverse-event reporting system relied on voluntary submissions structured to minimize signal, and was not adequate to the volume of the drug's distribution. A system that waits for a clinician to choose to file a report will always lag a harm spread thinly across millions of prescriptions, because no single prescriber sees enough of the pattern to recognize it. The Merck-side gap was structural in a different way: VIGOR was designed to test gastrointestinal safety (the COX-2 selling point), and the cardiovascular signal was a secondary finding that the trial was not powered to adjudicate against the naproxen-protective hypothesis. That trial-design choice compounded with a publication-bias pattern, an FDA adverse-event-reporting architecture insufficient for the population at risk, and a regulator-manufacturer disclosure protocol that had no decision boundary for ambiguous signals.⁴

AFTERMATH & REFORM

Merck eventually settled US litigation for about \$4.85 billion across approximately 27,000 plaintiffs. The case drove a coordinated regulatory response. The FDA Amendments Act of 2007 gave the FDA explicit authority to require Risk Evaluation and Mitigation Strategies

(REMS) and post-market study commitments as conditions of approval. In 2008 the FDA launched the Sentinel Initiative — an active post-market surveillance system that queries distributed health-data partners covering hundreds of millions of patient-years, rather than waiting for voluntary reports — a direct response to the Vioxx-era detection failure. The NEJM tightened conflict-of-interest disclosure for trial reports; the JAMA Ross et al. (2008) analysis became the reference point for publication-bias diagnosis in drug safety. By going out to the data instead of waiting for it to arrive, the reform inverted the logic that had let the signal sit unaggregated for years. The reform built the surveillance infrastructure that had not previously existed; the Vioxx record is what made the case for it.⁵



The cardiovascular risk was visible in Merck's internal data years before the drug was withdrawn.

PARAPHRASING DAVID GRAHAM (FDA OFFICE OF DRUG SAFETY), SENATE FINANCE COMMITTEE TESTIMONY, NOVEMBER 2004

REFERENCES

1. Bombardier, C. et al. (2000), "Comparison of upper gastrointestinal toxicity of rofecoxib and naproxen in patients with rheumatoid arthritis," VIGOR trial, NEJM 343(21):1520–1528 — the five-fold myocardial-infarction signal and the naproxen-protective hypothesis.
2. Bresalier, R. et al. (2005), "Cardiovascular events associated with rofecoxib in a colorectal adenoma chemoprevention trial," APPROVe, NEJM 352(11):1092–1102 — placebo-controlled confirmation of cardiovascular risk; early trial termination.
3. Graham, D. J. et al. (2005), "Risk of acute myocardial infarction and sudden cardiac death in patients treated with COX-2 selective and non-selective NSAIDs: nested case-control study," Lancet 365(9458):475–481 — Kaiser Permanente population-level cardiovascular risk analysis.
4. US Senate Committee on Finance, hearings on Vioxx and FDA's drug-safety system (November 18, 2004) — Graham testimony; "88,000 to 139,000 Americans" estimate of excess cardiovascular events; described FDA management pressure.
5. Ross, J. S., Hill, K. P., Egilman, D. S., Krumholz, H. M. (2008), "Guest authorship and ghostwriting in publications related to rofecoxib," JAMA 299(15):1800–1812 — publication-bias and ghost-authorship documentation from Merck litigation discovery.
6. FDA Amendments Act of 2007 (P.L. 110–85) and FDA Sentinel Initiative documentation (2008–present) — REMS authority, post-market study requirements, and active distributed-data post-market surveillance.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Vioxx is the canonical pharmaceutical case for post-market surveillance as a capability deliverable. The signal existed; the institutional architecture to aggregate it did not. The reform pattern — Sentinel — built the architecture. The case is the drug-industry analog of the EHR case (Case 2): a measurement architecture too thin for the system that depended on it.

LENS APPROACH

Vioxx is the canonical pharmaceutical post-deployment-surveillance case (induced 2.4; LENS D4+D3/PT5). LENS uses it in Domain 4 (Test and Evaluation; CLO-4) for post-market surveillance as a measurement deliverable — the signal existed; the institutional architecture to aggregate it did not. LENS uses it in Domain 3 (Human-System Collaboration; CLO-3) for the change-control and disclosure architecture between manufacturer, regulator, and prescriber that determines whether ambiguous safety signals reach a decision boundary. Pair with Radiology AI (Case 35) as the post-market-surveillance failure pattern at a new technological boundary; pair with EHR/CPOE (Case 2) at the measurement-architecture-too-thin layer; pair with Medical Errors and IOM (Case 8) at the population-scale-instrument layer.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	→ G D	→ D1 D2 D3 D4 D5	5	2.4	4 · 3
Test and Evaluation + D3					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Field an active post-market surveillance architecture sized to the exposed population before scale-up, so a harm spread thinly across millions can still aggregate to a decision boundary.
- ▶ Design trials and data collection to adjudicate between competing readings of a signal, rather than leaving an ambiguous result to the more comfortable interpretation.
- ▶ Treat signal-aggregation as a defined deliverable carrying evidence from manufacturer data through reviewers to prescribers, not a byproduct of voluntary reporting.

IN OPERATION / AFTER

- ▶ Audit the trial record and adverse-event data for signals already present but never carried to a decision, closing the lag between visibility and action.
- ▶ Monitor by querying health-data partners actively rather than waiting for voluntary reports that no single prescriber sees enough of the pattern to file.
- ▶ Sustain post-market surveillance funding against the commercial pressure to defer aggregation, since the cost of waiting compounds across the exposed population.

REFLECTION QUESTIONS

1. Identify a post-deployment surveillance architecture in your domain that is too thin for the scale of the system it monitors. What is the missing deliverable?
2. Design the Sentinel-equivalent post-market surveillance system for a new domain.
3. The VIGOR signal admitted two readings — harmful drug or protective comparator — and the data could not decide. Identify a measurement in your domain whose interpretation is underdetermined, and design the study or instrument that would force the question to a decision.

WHO BUILDS THIS governance, ethics & measurement · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · task analysis · design prototyping

Chapter 2

Healthcare & Patient Safety — What Works — and the Frontier

When measurement, teamwork, and design close the loop on patient harm.

Every durable safety gain shipped with an owner and an instrument.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

California Nurse Staffing Ratios — Legislating a Capability Requirement

G Governance & Trust **K** Knowledge & Institutional Memory

IMPACT California's mandated minimum nurse-to-patient ratios reduced nurse workload by 1–2 patients per nurse and were modeled to imply 10–14% fewer surgical patient deaths in comparison states if matched — observational, cross-sectional, no California baseline

IN BRIEF

California in 2004 became the first US state to mandate minimum nurse-to-patient ratios in acute-care hospitals — unit-level minimums written into law and enforced through inspection. The Aiken et al. (Health Services Research, 2010) study surveyed 22,336 nurses across California, Pennsylvania, and New Jersey (the latter two with no mandate), and found California nurses cared for 1–2 fewer patients each. Modeling implied that if the two comparison states had matched California's medical-surgical ratios, New Jersey would have seen 13.9% and Pennsylvania 10.6% fewer surgical patient deaths. The study is observational and cross-sectional, and the authors are explicit that the **absence of California baseline measures** fueled a contested stakeholder debate over whether the ratios themselves caused the improvement. The case is the canonical recent instance of a capability requirement (adequate staffing) being converted from a stated aspiration into an engineered, enforced specification by law — and the canonical instance of the methodological hedge such a conversion carries.

BACKGROUND

Nurse staffing is the load-bearing variable in the capability-of-care literature: a substantial body of evidence links nurse workload to medical errors, mortality, and rescue from complications. The professional consensus on the direction of the effect has been durable since the late 1990s. What had been contested is whether a mandated minimum, written into law and enforced through inspection, would convert the stated requirement (adequate staffing) into an engineered one (a specific ratio actually present at the bedside).^o

THE INTERVENTION

California in 2004 became the first US state to do it. The law specifies unit-by-unit minimum nurse-to-patient ratios — 1:5 in medical-surgical, 1:4 in step-down, 1:2 in ICU, and so on — enforceable through state inspection. The political process took five years from statutory enactment (1999) to regulatory implementation, with industry argument that the ratios would close hospitals and worsen access. The 2004 implementation went ahead substantially as written.¹

HOW IT WORKED

Aiken et al. (Health Services Research, 2010) is the strongest published evaluation. The study surveyed 22,336 nurses across California (mandated), Pennsylvania, and New Jersey (no mandate) on workload, work environment, and patient-outcome measures. The headline finding is that California nurses cared for 1–2 fewer patients each across comparable unit types. The modeled mortality implication — computed by applying the workload-mortality relationship from the broader Aiken et al. corpus to the workload difference — was that if Pennsylvania and New Jersey had matched California’s medical-surgical ratios, surgical mortality would have been 10.6% and 13.9% lower respectively.²

THE EVIDENCE

The hedge survives into the case verbatim. The Aiken et al. study is observational and cross-sectional, not a controlled experiment. There was no California baseline measurement to provide a pre/post within-state comparison. The modeled mortality estimates rest on the workload-mortality relationship from prior studies, and a contested stakeholder debate followed publication over whether the ratios themselves caused the observed improvement, or whether California hospitals differed in other ways the cross-section did not capture. The authors do not claim closure; they say the evidence is the strongest available and consistent with the broader literature.³

WHAT TRANSFERRED

What the case teaches is the form of converting a stated capability requirement into an engineered, enforced one by law, and the methodological hedge that conversion has to carry. The political path is part of the deliverable: five years from enactment to implementation, sustained industry opposition, contested evidence, and a study design constrained by the absence of a pre-implementation California baseline. In pair with SUBSAFE (Case 173), the case shows that engineered capability requirements at scale require both the requirements-discipline mechanism and an evidence architecture that survives the political process around the requirement.

The headline finding is direct; the modeled mortality estimate is the strongest available evidence and not closed proof.

EDITORS’ SYNTHESIS OF AIKEN ET AL. (2010).

REFERENCES

1. Aiken, Sloane, Cimiotti, Clarke, Flynn, Seago, Spetz, & Smith (2010), “Implications of the California Nurse Staffing Mandate for Other States,” *Health Services Research* 45(4):904–921, doi:10.1111/j.1475-6773.2010.01114.x.
2. Aiken, Clarke, Sloane, Sochalski, & Silber (2002), “Hospital nurse staffing and patient mortality, nurse burnout, and job dissatisfaction,” *JAMA* 288(16):1987–1993 — the workload-mortality relationship the 2010 modeling rests on.
3. California Department of Health Services (1999–2004), AB 394 regulatory implementation documentation.
4. Spetz, Chapman, Herrera, Kaiser, Seago, & Dower (2009), “Assessing the impact of California’s nurse staffing ratios on hospitals and patient care,” California HealthCare Foundation — implementation-period analysis.

THE LEARNING ENGINEERING LENS

LE INSIGHT

California’s nurse–ratio mandate is the canonical recent instance of converting a stated capability requirement (adequate staffing) into an engineered, enforced one by law. The workload effect is direct; the modeled mortality estimate is observational and cross–sectional with no California baseline. The hedge is the case.

LENS APPROACH

Nurse ratios is the legislated–requirement case (induced 1.1; LENS D1/PT3). LENS uses it in Domain 1 (Systems Analysis) for the conversion of stated requirement to engineered specification by law, and in Domain 4 (Test and Evaluation) for the CLO **Judgment under inadequate evidence** — the evaluation evidence is observational cross–sectional, the strongest available and not closed proof. Pair with Case 173 SUBSAFE at the requirements–as–deliverable layer.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	→ G K	→ D1 D2 D3 D4 D5	3	1.1	1 · 4
WORKFORCE POLICY		Systems Analysis			
PATIENT SAFETY					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ When converting a stated capability requirement into an engineered specification by law, design the evidence architecture before implementation: pre–registered measurement, a baseline window, and a comparison strategy that survives the political process. ▶ Specify the unit–level requirement at the granularity the work actually has — for nurse staffing, by unit type and acuity — rather than at an aggregate level that lets the requirement be averaged away. ▶ Anticipate sustained industry argument as part of the implementation cost; a five–year path from enactment to implementation is the political price of the engineered requirement.	▶ Report the workload effect (1–2 fewer patients per nurse) separately from the modeled mortality effect (10.6% and 13.9%); the workload finding is direct, the mortality estimate is modeled. ▶ Preserve the no–baseline hedge in any communication of the evidence; the strongest available evidence is not closed proof of causation. ▶ Treat the contested stakeholder debate as program documentation, not noise; the debate is part of why the case is teachable.

REFLECTION QUESTIONS

1. Identify a capability requirement in your domain that is currently stated but not engineered. What would a legislated minimum look like at unit–level granularity, and what evidence architecture would you put in place **before** implementation to make the conversion auditable?
2. Specify the political–process cost you would budget for: California’s path was five years from enactment to implementation under sustained industry argument. What is the analog in your context, and is the engineered requirement worth that cost?
3. The Aiken evidence is observational cross–sectional with no California baseline. What pre/post or controlled comparison would you build into the next state’s implementation to give the next round of evaluation a stronger foundation?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines

ACGME 80-Hour Resident Duty-Hour Reform

T Training Gap **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT ACGME capped U.S. resident physician work hours at 80/week to reduce fatigue-related errors; subsequent RCTs (FIRST, iCOMPARE) found flexible schedules non-inferior on patient outcomes, and the promised safety gain did not appear

IN BRIEF

After the 1984 death of Libby Zion focused decades of concern on resident-physician fatigue, the ACGME capped resident work at 80 hours a week in 2003 and limited first-year shifts to 16 hours in 2011. The logic was intuitive: tired doctors err, so cut the hours. But hours were one input to a many-variable system; cutting them multiplied error-prone hand-offs and cost residents continuity and procedural repetitions. Two large randomized trials — FIRST (2016) and iCOMPARE (2019) — found flexible schedules non-inferior to the strict caps on patient outcomes, so the promised safety gain never appeared, and in 2017 the ACGME relaxed the intern cap. The case is the book's clearest example of a single-variable intervention into a capability system — and of evidence catching up with a well-meant policy.

BACKGROUND

After the 1984 death of Libby Zion — blamed on overworked, under-supervised residents — resident fatigue became a decades-long argument. New York's Bell Commission produced the first hours limits in 1989, and pressure for a national standard built from there — the reform's intuition, that exhausted physicians make errors, was strong enough that the simplest lever, capping the hours, became the obvious answer long before anyone tested what the long shift was producing.^o

THE INTERVENTION

In 2003 the Accreditation Council for Graduate Medical Education capped resident work at 80 hours a week; in 2011 it limited first-year residents to 16-hour shifts. The logic was clean and intuitive — fatigue causes error, so reduce the hours. It was, in capability terms, a single-variable intervention: change one input and expect the outcome to move — a model that holds only if the rest of the system stays fixed, which in a teaching hospital it never does, since the hours removed had to be absorbed somewhere else.¹

HOW IT WORKED

But hours were one input to a system with many. Cutting them redistributed the work and multiplied patient hand-offs — themselves a documented site of error — while reducing residents' continuity and the procedural repetitions that build skill; many reported feeling

less prepared, not better rested. The long shift had quietly been doing work no one accounted for — sustaining a patient's care through one set of hands and accumulating the repetitions that turn a trainee into a clinician — and nothing was put in its place when it was cut.²

THE EVIDENCE

Two large randomized trials tested the policy. FIRST (Bilimoria et al., NEJM 2016), in surgery, found flexible duty hours non-inferior to the strict caps on patient outcomes and no worse for resident well-being — putting a randomized result against an intuition that had driven policy for years.³ iCOMPARE (Silber et al., NEJM 2019), in internal medicine, reached a parallel result in a second specialty, making the finding harder to dismiss as an artifact of surgery.⁴ Neither found the safety gain the cap had promised, and in 2017 the ACGME relaxed the 16-hour intern limit. The trials did not show fatigue is harmless — only that cutting one input, without rebuilding supervision and hand-offs, did not produce a safer system.⁵

WHAT TRANSFERRED

Duty-hour reform is the book's clearest case of a single-variable intervention into a multi-variable system. Read it against the successes here — the Keystone ICU project, crew resource management, the surgical safety checklist — which worked because they engineered supervision, hand-offs, and measurement **together with** the behavioral change, redesigning the surrounding architecture rather than pulling a single lever and hoping the rest would hold.⁶ The lesson is not that fatigue does not matter; it is that capability is a property of the whole system, and a reform that moves one variable while leaving the others untouched will be judged, in the end, by what it actually produced rather than by the plausibility of its intuition.



Flexible, less-restrictive duty hour policies for first-year residents were associated with non-inferior patient outcomes and no significant difference in residents' satisfaction with overall well-being and education quality.

PARAPHRASING ICOMPARE TRIAL (SILBER ET AL., NEJM 2019)

REFERENCES

1. B. H. Lerner, The Libby Zion Case and the Reform of Medical Education (2006); and the 1989 New York State (Bell Commission) duty-hour regulations — the origin of the reform.
5. Silber, Bellini, Shea et al., "Patient Safety Outcomes under Flexible and Standard Resident Duty-Hour Rules" (iCOMPARE), NEJM 380: 905–914 (2019). doi:10.1056/NEJMoa1810641.

2. Accreditation Council for Graduate Medical Education, Common Program Requirements (2003 and 2011 revisions) — the 80-hour weekly cap and the 16-hour first-year shift limit.
3. D. A. Asch et al., "Resident Duty Hours and Medical Education Policy," NEJM 370: 1671–1673 (2014); Institute of Medicine (Ulmer, Wolman & Johns, eds.), Resident Duty Hours: Enhancing Sleep, Supervision, and Safety (2008) — hand-offs and continuity trade-offs.
4. Bilimoria, Chung, Hedges et al., "National Cluster-Randomized Trial of Duty-Hour Flexibility in Surgical Training" (FIRST), NEJM 374: 713–727 (2016). doi:10.1056/NEJMoa1515724.
6. ACGME, Common Program Requirements (2017 revision) — relaxation of the 16-hour first-year shift limit.
7. P. Pronovost et al., Keystone ICU intervention, NEJM 355: 2725–2732 (2006), doi:10.1056/NEJMoa061115 A. B. Haynes et al., surgical safety checklist, NEJM 360: 491–499 (2009), doi:10.1056/NEJMsa0810119 — integrated interventions that engineered the surrounding architecture.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The clearest healthcare case in the dataset of a single-variable intervention into a multi-variable capability system. Pairs with Case 127 (fratricide) and Case 134 (V-22). The success cases — Keystone (14), CRM (12), Korean Air (23) — engineered supervisory, hand-off, and measurement architecture **together with** the behavioral change.

LENS APPROACH

LENS uses duty-hour reform in LEN 5 as the foundational capability-system case (what was the long-hour regime **producing** that was lost when the input was capped?), in LEN 4 to discuss measurement architecture (what FIRST and iCOMPARE measured, and what they did not), and in LEN 10 as a studio prompt for the integrated resident-training redesign that the reforms did not deliver.

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
HEALTHCARE	→ T K N	→ D1 D2 D3 D4 D5	3	2.3	2
Iterative Development					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Map the full set of variables the targeted input is coupled to — continuity, supervision, hand-offs, procedural reps — before changing one of them.
- ▶ Design the substitute for whatever the changed input was producing (e.g., structured hand-offs and supervision) into the same reform, not as a follow-on.
- ▶ Build the measurement that will test the reform's promised gain into the rollout, so the policy is falsifiable against the operational record.

IN OPERATION / AFTER

- ▶ Audit the intervention against patient and trainee outcomes with a controlled comparison, as FIRST and iCOMPARE did, rather than trusting the intuition.
- ▶ Monitor the variables that absorbed the change (hand-off frequency, procedural exposure) for the harms a single-lever fix can displace.
- ▶ Sustain a willingness to revise the policy when evidence catches up, as the 2017 relaxation did, rather than defending the original lever.

REFLECTION QUESTIONS

1. What capability is the long-hours / heavy-workload regime in your domain currently producing — supervisory exposure, continuity, procedural reps, tacit-knowledge transfer — that a simple cap would lose?
2. Design the integrated redesign — supervision, hand-off, measurement, exposure — that would substitute for the capability the input cap removes.
3. The reform was intuitive enough to set national policy years before FIRST and iCOMPARE tested it. Design the randomized or quasi-experimental check you would build into a future single-variable reform so its promised gain is measured before it is mandated, not after.

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

Implementation Science in Healthcare — The 17-Year Gap

K Knowledge & Institutional Memory **G** Governance & Trust **N** Normalization of Deviance

IMPACT Average time from research finding to clinical practice: 17 years; only ~14% of research findings ever reach practice

IN BRIEF

Implementation science has a canonical finding: it takes an average of about seventeen years for research evidence to reach clinical practice, and only roughly 14% of research findings ever make it at all. This is not a single incident but a systemic condition — effective interventions exist; the system to adopt, sustain, adapt, and measure them at scale does not. Frameworks like the Active Implementation Frameworks (Fixsen et al., 2005) and EPIS (Aarons et al., 2011) were built specifically to attack this gap, and LENS threads implementation science throughout its curriculum in direct response. The seventeen-year gap is the meta-case for the whole book: every success case is a closure of this gap in one domain, every failure case the gap left open. The discipline exists to make seventeen years shorter.

THE SHIFT

Medicine generates more validated knowledge than it can absorb. Implementation science arose to study a stubborn fact: knowing what works and having it practiced are different problems, separated by years.^o Treating the two as one problem is the error the field formed to correct: a validated finding is not yet a changed practice, and the distance between them is itself a phenomenon to be studied, measured, and engineered rather than waited out.

WHAT IS EMERGING

The canonical figures are stark: it takes an average of about seventeen years for research evidence to be integrated into clinical practice, and only roughly 14% of research findings ever make it at all.¹ Read together, the two figures describe a pipeline that is both slow and leaky: most of what is learned never reaches the bedside at all, and the fraction that does arrives long after the patients who first needed it, so the delay is compounded by sheer attrition.

THE CAPABILITY QUESTION

This is not a single case but a systemic condition — the same structural problem the medical-error data (Case 8) describes from the outcome side. Effective interventions exist; the institutional system to adopt, sustain, adapt, and measure them at scale does not.² Where the medical-error data counts the harm at the far end of the pipeline, the translation


figures name the mechanism that produces it: the same missing adoption-and-measurement system shows up as a delay from one vantage and as a body count from the other.

EARLY EVIDENCE

Frameworks built to attack the gap — the Active Implementation Frameworks (Fixsen et al., 2005) and the EPIS framework (Aarons et al., 2011) — show that implementation can be engineered rather than left to chance, and they inform LENS's choice to thread implementation science through every course rather than isolate it in a module.³ Threading the discipline through every course rather than confining it to a single module is itself a claim these frameworks support: if implementation is an engineerable property of any intervention, it cannot be quarantined as a specialty and must inform how every design is taught.

OPEN PROBLEMS

The seventeen-year gap is the meta-case for this book. Every success case in the chapters ahead is a closure of the gap in one domain; every failure case is the gap left open.⁴ The open problem is general and unglamorous: building, funding, and owning the adoption-and-measurement pathway that turns a proven intervention into routine practice — which is, in one sentence, what the LENS discipline exists to do. Because the problem is general rather than domain-specific, no single clinical result closes it; what closes it is the repeatable, owned, and funded pathway that any proven finding can be run through, again and again.



The answer is 17 years. What is the question?

MORRIS, WOODING & GRANT, J ROYAL SOC MED, 2011

REFERENCES

1. E. A. Balas & S. A. Boren (2000), Yearbook of Medical Informatics — the 17-year / 14% translation figures.
2. Z. Morris, S. Wooding & J. Grant, "The answer is 17 years, what is the question," J. Royal Society of Medicine (2011) (quoted).
3. D. Fixsen et al., Implementation Research: A Synthesis of the Literature (2005) — the Active Implementation Frameworks.
4. G. Aarons et al. (2011), the EPIS framework; L. Damschroder et al. (2009), CFIR.
5. Cf. medical error as systemic failure (Case 8); Goodell & Kolodner, Learning Engineering Toolkit (2022).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The 17-year gap is the structural problem that LENS exists to address. It is the difference between knowing what works and having it deployed. Every case in this book is a sample from a distribution governed by that gap. The success cases shorten it; the failure cases let it run.

LENS APPROACH

LENS uses this case in LEN 1 as the foundational problem statement of the discipline, in LEN 10 as a studio constraint (designs must consider implementation, not just efficacy), and in LEN 8 as the central knowledge-transfer challenge. The case is referenced at least once in every required course.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	→ K G N	→ D1 D2 D3 D4 D5	4	1.4	2
Iterative Development					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design every intervention with its adoption-and-measurement pathway attached, so implementation is engineered in rather than left to chance after the evidence is published.
- ▶ Use an implementation framework (Active Implementation, EPIS) from the outset to plan adoption, sustainment, and adaptation as deliverables of the project.
- ▶ Name an owner and a funding line for the pathway, since a proven finding with no one accountable for fielding it is exactly what the gap is made of.

IN OPERATION / AFTER

- ▶ Measure both reach and speed — what fraction of practice has adopted the intervention and how long it took — to see the slow-and-leaky pipeline rather than assume publication equals uptake.
- ▶ Sustain and adapt fielded interventions, treating drift back to old practice as a measurable failure mode, not a one-time rollout that holds itself.
- ▶ Track the gap as a standing metric across the institution, so closing it in one domain becomes a repeatable pathway rather than a one-off success.

REFLECTION QUESTIONS

1. Pick an evidence-based intervention in your domain. Estimate the gap between when the evidence became conclusive and when the intervention reached majority of practice. What did the gap cost?
2. Design the deliverable that would shorten that gap by half in your domain. Be specific about who funds it, who owns it, and what evidence demonstrates the reduction.
3. The translation pipeline is both slow and leaky — most findings never reach practice, and those that do arrive late. For your domain, is the binding problem the delay or the attrition, and what would you measure to tell which one to attack first?

WHO BUILDS THIS knowledge management & data engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS knowledge bases · data pipelines · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA

T Training Gap **K** Knowledge & Institutional Memory

IMPACT Barsuk et al. (Northwestern/Feinberg) demonstrated that simulation-based mastery learning for central venous catheter (CVC) insertion reduced catheter-related bloodstream infection rates and procedural complications, with cost-effectiveness shown in a single-hospital evaluation; the program was subsequently disseminated nationally to Department of Veterans Affairs medical centers

IN BRIEF

Central venous catheter (CVC) insertion is a high-volume, high-consequence procedure whose complications — pneumothorax, arterial puncture, catheter-related bloodstream infection (CRBSI) — are well-characterized and partly attributable to procedural skill at the bedside. Barsuk and colleagues at Northwestern/Feinberg published a series of studies through the late 2000s and 2010s establishing simulation-based mastery learning (SBML) — practice on a simulator to a defined performance standard, not to a clock — as an effective approach for trainee CVC insertion. The single-center evidence reported reduced CRBSI rates, reduced procedural complications, and favorable cost-effectiveness. The program was subsequently disseminated to Department of Veterans Affairs medical centers as a national rollout — the deliberate cross-institutional transfer that distinguishes the case from a successful single-site intervention. The case is the canonical small-tier instance of cross-domain adaptation of a proven intervention (C6.4), and pairs with the multidisciplinary-translation trio (Cases 40 team science, 122 IPE, 123 implementation-science training) as cross-domain workforce evidence. The honest hedge: the dissemination outcome literature is thinner than the original single-center evidence; the case is teachable on the SBML method and the documented dissemination effort, with measurement at the multi-site level a live evidence frontier.

BACKGROUND

Central venous catheter insertion is a procedure trainees in medicine, surgery, and critical care perform routinely; the complication profile — pneumothorax, arterial puncture, catheter-related bloodstream infection — is well-characterized and attributable in part to procedural skill at the bedside. The historical training model relied on graduated exposure on live patients, with supervision but without a defined performance standard before unsu-

pervised practice. The Barsuk group at Northwestern's Feinberg School of Medicine, working across critical care and medical education, set out to replace exposure-based training with simulation-based mastery learning.⁹

THE INTERVENTION

Simulation-based mastery learning is the specific structural form: trainees practice the procedure on a simulator to a defined performance standard, not to a clock. Practice continues until the trainee demonstrates competence; the standard, not the calendar, governs progression. Barsuk et al. (Archives of Internal Medicine, 2009; Critical Care Medicine, Journal of Hospital Medicine, and a series of subsequent papers) reported that the SBML cohort, compared with historical and concurrent non-SBML cohorts at the same hospital, performed CVC insertions with fewer needle passes, fewer arterial punctures, and lower CRBSI rates on the patients they subsequently catheterized.¹

HOW IT WORKED

The single-center economic evidence was the second pillar. Cohen et al. (Simulation in Healthcare, 2010) reported that the SBML program was cost-saving at the single hospital, driven principally by the reduction in CRBSI cases — each bloodstream infection averted is expensive enough that even a modest reduction recovers the simulation-training investment. The combination of the procedural-outcome evidence and the cost-effectiveness evidence made the program transferable beyond the original site.²

THE EVIDENCE

The dissemination step is what distinguishes the case from a successful single-site intervention. The SBML CVC program was subsequently disseminated to Department of Veterans Affairs medical centers as a national rollout, supported by the VA's investment in simulation infrastructure and clinical training architecture. The deliberate cross-institutional transfer — from an academic medical center to a federated health system — is the structural feature the C6.4 cross-domain-adaptation competency exists to name. The dissemination outcome literature is thinner than the single-center evidence; multi-site CRBSI tracking is harder to attribute, and the published evidence at the national-VA scale is at the program-report and small-evaluation tier rather than the controlled-comparison tier the single-center papers established.³

WHAT TRANSFERRED

In pair with the multidisciplinary-translation trio (Cases 178 team science, 122 IPE, 123 implementation-science training), the Barsuk SBML case completes the workforce-evidence picture: an intervention with strong single-center controlled-comparison evidence (this case) sits alongside a structured team-science training program with validated measurement (121), the field-scale enthusiasm-evidence gap in interprofessional education (122), and the operational-practice gap inside implementation-science training programs (123). The four cases together stage the cross-domain workforce-evidence pattern: the mechanism that works at single-program scale is the demonstrator; the field-scale measurement and the multi-site dissemination evidence are what the discipline is still building.⁴

The standard, not the calendar, governs progression. The trainee practices until they demonstrate competence; the simulator absorbs the cost of the practice.

EDITORS' SYNTHESIS OF BARSUK ET AL. (2009).

REFERENCES

1. Barsuk, Cohen, Feinglass, McGaghie, & Wayne (2009), "Use of simulation-based education to reduce catheter-related bloodstream infections," *Archives of Internal Medicine* 169(15):1420–1423, doi:10.1001/archinternmed.2009.215.
2. Cohen, Feinglass, Barsuk, Barnard, O'Donnell, McGaghie, & Wayne (2010), "Cost savings from reduced catheter-related bloodstream infection after simulation-based education for residents in a medical intensive care unit," *Simulation in Health-care* 5(2):98–102, doi:10.1097/SIH.0b013e3181bc8304.
3. Barsuk, McGaghie, Cohen, Balachandran, & Wayne (2009), "Use of simulation-based mastery learning to improve the quality of central venous catheter placement in a medical intensive care unit," *Journal of Hospital Medicine* 4(7):397–403.
4. McGaghie, Issenberg, Cohen, Barsuk, & Wayne (2011), "Does simulation-based medical education with deliberate practice yield better results than traditional clinical education? A meta-analytic comparative review," *Academic Medicine* 86(6):706–711.
5. Department of Veterans Affairs SimLEARN documentation — the operating-program record of the national dissemination effort.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Barsuk SBML for CVC insertion is the canonical small-tier intervention case for cross-institutional dissemination of a proven mechanism. The single-center evidence is controlled-comparison and cost-effective; the VA national dissemination is the cross-institutional transfer step. The multi-site outcome evidence is thinner than the single-center evidence — the live frontier of the dissemination case.

LENS APPROACH

Barsuk SBML is the cross-institutional dissemination case (induced 6.4; LENS D2/PT4) — Domain 2 for mastery-learning replacing exposure-based progression; Domain 4 for the procedural-outcome + cost-effectiveness pairing. Pair with the translation trio (Cases 40, 28, 41).

DOMAIN	Modes Addressed	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
MEDICAL EDUCATION	T K	D1 D2 D3 D4 D5	4	6.4	2 · 4
PATIENT SAFETY	→	Iterative Development			
WORKFORCE DISSEMINATION					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Replace exposure-based training with a defined performance standard on a simulator; the SBML deliverable is competence demonstrated, not hours accumulated.
- ▶ Design the simulator and the standard around the specific complication modes the procedure produces — pneumothorax, arterial puncture, CRBSI for CVC — so the training closes the gaps the outcome data identifies.
- ▶ Pair the procedural-outcome evidence with cost-effectiveness evidence at single-site scale; the combination is what makes the program disseminable beyond the original institution.

IN OPERATION / AFTER

- ▶ Plan the dissemination as a deliverable of the original program, not as a separate downstream activity; the VA rollout is what makes Barsuk SBML a cross-institutional intervention rather than a successful single-site study.
- ▶ Build multi-site outcome measurement into the dissemination from the start; the published evidence at national-VA scale is thinner than the single-center evidence, and the gap is a live frontier the case names.
- ▶ Carry the cross-domain pairing with team-science training (Case 40), IPE (Case 28), and implementation-science training (Case 41) into the curricular framing — the four cases together stage what cross-domain workforce dissemination looks like with measurement and what it looks like without.

REFLECTION QUESTIONS

1. Identify a high-volume, high-consequence procedural skill in your domain where training is currently exposure-based rather than performance-standard-based. What is the analog of SBML, and what is the simulator that would absorb the cost of the practice?
2. Specify the procedural-outcome and cost-effectiveness pairing you would design at single-site scale; both halves are what made the Barsuk program disseminable, and the pairing is what distinguishes a teachable single-site case from a national-rollout candidate.
3. The VA dissemination evidence is thinner than the single-center evidence. Design the multi-site outcome-tracking architecture you would build into the next dissemination — what instrument, what comparison, what cadence — so the multi-site evidence catches up to the demonstration evidence.

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines

I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer

H Human-System Interface **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Across nine pediatric residency programs, implementation of the I-PASS handoff bundle (mnemonic + training + faculty development + sustainability campaign) was associated with a 23% relative reduction in medical errors and a 30% reduction in preventable adverse events, without negatively affecting resident workflow — the study design 'precludes definitively establishing a causal link'

IN BRIEF

The shift-change handoff is the moment in inpatient care where patient state has to be transferred accurately across a human boundary under time pressure. Loss of safety-critical information at handoff is a documented failure mode, and the cognitive demand on the outgoing resident — synthesize, prioritize, and convey — exceeds what unaided improvisation can reliably deliver. I-PASS is a structured handoff bundle: a mnemonic (Illness severity, Patient summary, Action list, Situation awareness and contingency planning, Synthesis by receiver), paired with formal trainee education, faculty development, and a sustainability campaign. Starmer et al. (NEJM, 2014) studied its implementation at nine pediatric residency programs and reported a 23% relative reduction in medical errors and a 30% reduction in preventable adverse events, with no negative effect on resident workflow. The hedge survives verbatim: the authors state plainly that “our study design precludes definitively establishing a causal link.” Published correspondence cautions against implementing the mnemonic alone without the full bundle. The case is the structured-transfer companion to Case 177 (CIRAS) at the cultural-half-of-capability layer, and the small-tier intervention spine for state-transparency under stress.

BACKGROUND

The shift-change handoff is a structural risk point in inpatient care. The outgoing resident — who has the most detailed mental model of each patient's state and trajectory — has to synthesize, prioritize, and convey enough of that model to the incoming team that the patient's care is not disrupted by the boundary. Loss of safety-critical information at handoff is a documented failure mode in the patient-safety literature; preventable adverse events that trace to communication breakdown are a substantial fraction of inpatient harm. The cognitive demand exceeds what unaided improvisation can reliably deliver, and the institutional tolerance for that improvisation has been falling.^o

THE INTERVENTION

I-PASS is an explicitly bundled intervention. The mnemonic — Illness severity, Patient summary, Action list, Situation awareness and contingency planning, Synthesis by receiver — structures the handoff conversation around the information classes the receiver needs in order to take over safely. The mnemonic alone is the visible piece; the bundle around it is what makes the mnemonic land. The published implementation pairs the mnemonic with formal trainee education on its use, faculty development so that supervisors model and reinforce it, and a sustainability campaign so that the discipline does not erode after the rollout. The published correspondence cautions explicitly against implementing the mnemonic alone without the full bundle.¹

HOW IT WORKED

Starmer et al. (New England Journal of Medicine, 2014) studied the I-PASS bundle's implementation at nine pediatric residency programs in the United States and Canada. The design was a pre/post evaluation across the participating sites, with the bundle deployed sequentially and outcomes tracked through chart review and observation. The headline finding was a 23% relative reduction in medical errors and a 30% reduction in preventable adverse events from the pre-implementation period to the post-implementation period, with no observed negative effect on resident workflow — the bundle did not impose net friction that displaced other clinical work.²

THE EVIDENCE

The hedge is load-bearing and survives into the case verbatim. The authors state in the published paper that “our study design precludes definitively establishing a causal link” — pre/post evaluation across multiple sites is the strongest practical design in this setting, and it cannot rule out secular trends in patient-safety culture, simultaneous quality-improvement work, or the sites' own selection into the trial. The 23% reduction is the strongest available evidence; it is not closed proof that the bundle alone produced the effect. Treating it as either is a misreading of what the study design supports.³

WHAT TRANSFERRED

What the case carries for the corpus is the state-transparency-under-stress pattern at the human-human boundary (induced 3.3, LENS D2/PT5). The handoff is the moment of cognitive transfer at the team boundary; the bundle is the workflow artifact that makes the patient state legible across that boundary. Paired with Case 177 (CIRAS) at the cultural-half-of-capability layer, the case shows that the mnemonic alone is not the intervention — the institutional commitment to the bundle (trainee education, faculty development, sustainability) is what makes the mnemonic real in operation. The pattern is partly borrowed from aviation discipline, so the case also seeds the cross-domain-adaptation conversation that v1 Cases 117 (CRM) and 89 (ASRS) anchor in aviation evidence.

Our study design precludes definitively establishing a causal link.

STARMER ET AL., NEJM 2014.

REFERENCES

1. Starmer, A. J., Spector, N. D., Srivastava, R., West, D. C., Rosenbluth, G., Allen, A. D., Noble, E. L., Tse, L. L., Dalal, A. K., Keohane, C. A., Lipsitz, S. R., Rothschild, J. M., Wien, M. F., Yoon, C. S., Zigmont, K. R., Wilson, K. M., O'Toole, J. K., Solan, L. G., Aylor, M., Bismilla, Z., Coffey, M., Mahant, S., Blankenburg, R. L., Destino, L. A., Everhart, J. L., Patel, S. J., Bale, J. F., Spackman, J. B., Stevenson, A. T., Calaman, S., Cole, F. S., Balmer, D. F., Hepps, J. H., Lopreiato, J. O., Yu, C. E., Sectish, T. C., & Landrigan, C. P. (2014). Changes in medical errors after implementation of a handoff program. *New England Journal of Medicine*, 371(19):1803–1812. doi:10.1056/NEJMsa1405556 — the case's primary evaluation.
2. Starmer, A. J., et al. (2014). Implementation correspondence and follow-up. *New England Journal of Medicine* — the published correspondence cautioning against implementing the mnemonic alone.
3. Sectish, T. C., et al. (2010). Establishing a multisite education and research project: The I-PASS Study Group. *Academic Medicine* — the I-PASS Study Group methodology and design rationale.
4. Cohen, M. D., & Hilligoss, P. B. (2010). The published literature on handoffs in hospitals: deficiencies identified in an extensive review. *Quality and Safety in Health Care* — the broader hand-off-failure-mode literature.

THE LEARNING ENGINEERING LENS

LE INSIGHT

I-PASS is the structured-handoff intervention the patient-safety literature anchors on — a 23% relative reduction in preventable adverse events across nine residency programs, with no negative effect on resident workflow. The hedge that survives verbatim is the authors' own: the study design “precludes definitively establishing a causal link,” and the bundle, not the mnemonic alone, is what the evidence is about.

LENS APPROACH

I-PASS is the state-transparency-under-stress case at the human-human boundary (induced 3.3; LENS D2/PT5) — Domain 2 for bundle-as-intervention; Domain 5 for the institutional faculty-development commitment. Pair with Case 177 (CIRAS) and v1 Cases 117 (CRM) and 89 (ASRS).

DOMAIN	Modes Addressed	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
CLINICAL CARE	H K N	D1 D2 D3 D4 D5	5	3.3	2 · 5
PATIENT SAFETY	→	Iterative Development			
TEAM COMMUNICATION					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the information classes the receiver needs in order to take over safely (illness severity, patient summary, action list, situation awareness, synthesis by receiver) and structure the handoff conversation around them, rather than around the outgoing operator's narrative preference.
- Treat the mnemonic, the trainee education, the faculty development, and the sustainability campaign as one bundle; the published correspondence is explicit that the mnemonic alone does not carry the result.
- Design the evaluation to track preventable adverse events and resident workflow together, so that the intervention's effect on patient harm and on the work it adds are visible in the same evidence record.

IN OPERATION / AFTER

- Report the 23% relative reduction together with the verbatim hedge that “our study design precludes definitively establishing a causal link”; the strongest available evidence is not closed proof.
- Carry the bundle-not-mnemonic warning into any adaptation: implementations that drop the faculty-development and sustainability components are not implementations of the intervention the published evidence is about.
- Build the cross-domain adaptation conversation deliberately — aviation CRM and ASRS evidence (v1 Cases 117, 119) is part of the I-PASS lineage and CIRAS (Case 177) is the non-aviation companion at the cultural-commitment layer.

REFLECTION QUESTIONS

1. Identify a moment in your domain's workflow where state has to be transferred accurately across a human boundary under time pressure. What are the information classes (the analog of I-PASS's five) the receiver needs in order to take over safely, and what bundle around the structuring artifact would make it land?
2. The 23% relative reduction is the strongest available evidence and the authors are explicit it is not closed proof. What additional evidence — a cluster-randomized rollout, a residual-error analysis at the handoff point — would you require before treating the figure as closure rather than strong signal?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · incident analysis · safety dashboards

NCSBN National Simulation Study — Licensing the 50% Substitution Rule

G Governance & Trust **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT A longitudinal RCT randomized students across multiple US nursing programs to control, 25%, or 50% simulation substitution for traditional clinical hours; 660+ took the NCLEX with no statistically significant differences in clinical competency, nursing knowledge, or NCLEX pass rates — the number of nursing regulatory boards permitting up to 50% simulation substitution increased more than 20-fold from 2014 to 2022

IN BRIEF

Pre-licensure nursing education had long rested on a regulatory requirement for traditional clinical hours — time spent caring for patients in a real clinical setting under the supervision of a clinical faculty member. As clinical-placement capacity came under pressure, the question facing nursing regulators was whether high-quality simulation could substitute for some fraction of those hours without degrading the clinical capability of new nurses. The National Council of State Boards of Nursing (NCSBN) ran the study that the field then built on: a longitudinal RCT randomized students across multiple US nursing programs to control, 25%, or 50% simulation substitution. Hayden et al. (Journal of Nursing Regulation, 2014) reported no statistically significant differences in clinical competency, nursing knowledge, or NCLEX pass rates across the three groups; 660+ participants took the NCLEX. The follow-up (2023) documents the institutional transfer: the number of nursing regulatory boards permitting up to 50% substitution increased more than 20-fold from 2014 to 2022 — an unusually clean case of a single evidence base propagating across an entire regulatory field. The hedge survives: the result holds only “under conditions comparable to those described in the study” (high-quality simulation, trained faculty). Pair with Case 40 (Colorado CTSA team-science training) at the cross-domain workforce-intervention layer.

BACKGROUND

Pre-licensure nursing education in the United States has rested on a regulatory requirement, set by state nursing boards, that students complete a defined number of traditional clinical hours — time spent caring for patients in a real clinical setting under the supervision of a clinical faculty member. As nursing programs grew and clinical-placement capacity in hospitals came under pressure, the field was asking whether high-quality simulation — manikin-based or standardized-patient scenarios run in a controlled environment — could substitute for some fraction of traditional clinical hours without degrading the clinical capability of new nurses.^o

THE INTERVENTION

The National Council of State Boards of Nursing — the umbrella body for the state nursing-regulatory boards that together control nursing licensure — commissioned the definitive study. Hayden et al. (Journal of Nursing Regulation, 2014) reported the NCSBN National Simulation Study: a longitudinal randomized controlled trial that assigned pre-licensure nursing students at multiple US nursing programs to a control group, a 25%-substitution group, or a 50%-substitution group. The substitution condition replaced an equivalent fraction of traditional clinical hours with simulation hours run to specified quality standards (validated scenarios, trained simulation faculty, structured debriefing).¹

HOW IT WORKED

The headline finding was a null result on the outcome the regulators most cared about: no statistically significant differences in clinical competency, nursing knowledge, or NCLEX pass rates across the three groups. Six hundred and sixty-plus participants took the NCLEX licensure examination, providing the comparison the regulatory decision rested on. The longitudinal design followed cohorts through their first nursing positions, with comparable outcome measures at multiple time points. The study's qualifying language is precise: the result holds only "under conditions comparable to those described in the study," and those conditions are named — high-quality simulation, trained simulation faculty, structured debriefing.²

THE EVIDENCE

What followed is the unusual part. Nursing regulation in the US is genuinely decentralized: each state board sets its own pre-licensure requirements, and there is no federal forcing mechanism. The 2023 follow-up study documents that the number of nursing regulatory boards permitting up to 50% simulation substitution increased more than 20-fold from 2014 to 2022. A single evidence base — one study, run by the regulators' own collaborative — propagated across an entire regulatory field, on the strength of the published design and the null result on the licensure examination the boards control.³

WHAT TRANSFERRED

The case carries the cross-organization knowledge-transfer pattern at the regulatory-institution layer (induced 6.1, LENS D2/PT4). The mechanism that made the propagation work was the credibility of the study's design (longitudinal RCT, multi-site, blinded outcomes), the alignment of its outcome measure with the boards' own licensure mechanism (NCLEX), and the institutional ownership of the evidence (the regulators commissioned it and own it). Pair with Case 40 (Colorado CTSA team-science training) at the cross-domain workforce-intervention layer — both are small-tier interventions with measurable workforce-capability claims, and both depend on the legitimacy of the assessment instrument the institution then has to defend.

The result holds under conditions comparable to those described in the study — high-quality simulation, trained faculty, structured debriefing.

EDITORS' SYNTHESIS OF HAYDEN ET AL. (2014).

REFERENCES

1.

Hayden, J. K., Smiley, R. A., Alexander, M., Kardong-Edgren, S., & Jeffries, P. R. (2014). The NCSBN National Simulation Study: A longitudinal, randomized, controlled study replacing clinical hours with simulation in prelicensure nursing education. *Journal of Nursing Regulation*, 5(2 Suppl):S1–S64. [https://www.journalofnursingregulation.com/article/s2155-8256\(15\)30062-4/fulltext](https://www.journalofnursingregulation.com/article/s2155-8256(15)30062-4/fulltext) — the case’s primary study.

2.

Smiley, R. A., & Martin, B. (2023). Simulation in Nursing Education: Advancements in Regulation, 2014–2022. *Journal of Nursing Regulation*. doi:10.1016/S2155-8256(23)00086-8 — the 2023 follow-up documenting the regulatory propagation.

3.

Jeffries, P. R. (2012). *Simulation in Nursing Education: From Conceptualization to Evaluation* (2nd ed.). National League for Nursing — the simulation–quality framework the NCSBN study’s quality conditions rest on.

4.

INACSL Standards Committee (2016). *INACSL Standards of Best Practice: Simulation. Clinical Simulation in Nursing*, 12(S) — the simulation–practice standards downstream programs adopt as the analog of the NCSBN quality conditions.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The NCSBN National Simulation Study is the unusual case of a single, well-designed, regulator–commissioned study propagating a substantial workforce–capability change across an entire decentralized regulatory field. The null result on NCLEX is conditional on the quality conditions the study names — high-quality simulation, trained faculty, structured debriefing — and the qualifying language has to travel with the result.

LENS APPROACH

NCSBN is the regulator–owned cross–organization knowledge–transfer case (induced 6.1; LENS D2/PT4) — Domain 2 for the RCT–design discipline; Domain 5 for the institutional–ownership move that propagated evidence across decentralized regulators. Pair with Case 40 (Colorado CTSA).

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
NURSING EDUCATION	G K D	D1 D2 D3 D4 D5	4	6.1	2 · 4
CLINICAL SIMULATION	→	→ Iterative Development			
REGULATION					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Commission the study from inside the regulatory body that will then have to act on it; the institutional ownership of the evidence is what makes the propagation across decentralized regulators possible.
- ▶ Specify the quality conditions the result is conditional on (high–quality simulation, trained faculty, structured debriefing) at design time, so the qualifying language travels with the result and is not stripped at the citation stage.
- ▶ Align the study’s outcome measure with the licensure mechanism the boards control (NCLEX pass rate); the credibility of the propagation depends on the outcome being the one the regulators already use.

IN OPERATION / AFTER

- ▶ Carry the qualifying language (“under conditions comparable to those described in the study”) into every downstream adoption; the study’s null result is conditional on quality standards, not unconditional.
- ▶ Document the propagation as a transferable institutional pattern: a decentralized regulatory field can move on the strength of a single, well–designed, body–owned study when the outcome measure is the one the regulators already use.
- ▶ Treat the 20–fold expansion of board adoption from 2014 to 2022 as evidence of the mechanism of regulatory transfer, not as evidence that 50% substitution is safe outside the study’s quality conditions.

REFLECTION QUESTIONS

1. Identify a decentralized regulatory field in your domain where a substitution or adoption decision rests on whether the new modality preserves capability. What would a NCSBN-class RCT, commissioned from inside the regulatory body, look like — and what would the licensure-aligned outcome measure be?
2. The result is conditional on quality conditions (high-quality simulation, trained faculty, structured debriefing). What is the analog conditional in your context, and what mechanism would carry the qualifying language into the regulatory adoptions that follow?
3. The 20-fold expansion of board adoption is evidence of regulatory transfer, not of safety outside the study's quality conditions. What would you do operationally to ensure the adoptions in your field actually meet the conditions the propagation rests on?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · task analysis · design prototyping

Spaced Education RCTs in Medical Training

H Human–System Interface **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT A multi-institution RCT of urology residents across the US and Canada randomized participants to bolus versus spaced-pattern email delivery of validated study questions; spaced education improved acquisition and retention of medical knowledge, and a follow-up showed the learning benefit persisting for two years

IN BRIEF

Spacing — distributing study sessions over time rather than massing them — is one of the most robust findings in basic learning–science research, with effects across age, population, and content domain. Whether the basic finding transfers into the workplace–L&D context of practicing medical trainees is a separate empirical question. Kerfoot et al. (Journal of Urology, 2007) ran the test that closed the loop. A multi-institution RCT randomized urology residents across US and Canadian programs to receive validated study questions in either a bolus pattern (concentrated delivery) or a spaced pattern (distributed delivery), both via email. The spaced-education condition improved acquisition and retention of medical knowledge against the bolus comparison. The 2009 follow-up (Kerfoot, Journal of Urology, 181:2671) documented that the learning benefit persisted for two years. The case is a methodologically clean small-tier intervention with replication, with standard RCT scope: the result speaks to the delivery–pattern manipulation in the email–delivered validated–question context, not to spacing in general or to other modalities. The strongest randomized strength in the small-tier batch. Pair with Case 69 (Duolingo half-life regression) at the spacing–mechanism–in–deployment layer.

BACKGROUND

Spacing — distributing study sessions over time rather than massing them in a single concentrated block — is among the most robust findings in basic learning–science research. The spacing effect appears across age groups, populations, content domains, and outcome measures, and the broader retrieval–practice literature complements it. What is less obvious in the basic literature is whether the spacing effect transfers to the workplace–L&D context, where the learner is a practicing trainee, the content is the trainee’s own clinical specialty, and the delivery mechanism is one the trainee will tolerate as part of their work week.^o

THE INTERVENTION

Kerfoot et al. (Journal of Urology, 2007) ran the test that closed the loop. The study was a multi-institution randomized controlled trial of urology residents at multiple US and Canadian residency programs. Participants were randomized to receive validated study questions drawn from the urology in–service examination question pool, delivered by email, in either a bolus pattern (concentrated delivery over a short window) or a spaced pattern

(distributed delivery across a longer window). Outcomes were knowledge acquisition (immediate post-test) and retention (delayed post-test), measured against the same validated question pool the delivery drew from.¹

HOW IT WORKED

The result was direct. The spaced-education condition improved acquisition and retention of medical knowledge against the bolus comparison. The intervention and the measurement formed a tight closed loop: the validated questions delivered are the same content the post-test drew from, the email delivery is the manipulation, and the knowledge-outcome measure is the institution's own in-service question pool. The 2009 follow-up (Kerfoot, Journal of Urology, 181:2671) reported a separate analysis with the same residents and documented that the learning benefit persisted for two years — a duration that is rare in the workplace-L&D RCT literature.²

THE EVIDENCE

The case's scope is what makes it teachable rather than over-extended. The result speaks to the delivery-pattern manipulation (bolus vs. spaced) in the email-delivered validated-question context, with urology residents at multiple programs. It does not speak to spacing in general, to other modalities (mobile, in-clinic, simulation), or to content classes substantially different from in-service-examination knowledge. The strength of the design — multi-institution RCT, two-year persistence — is the strongest randomized strength in the small-tier batch, and the scope discipline is the qualifying language that makes the strength interpretable. The case is the methodologically clean small-tier closed-loop intervention the LENS Iterative Development domain anchors on.³

WHAT TRANSFERRED

Pair with Case 69 (Duolingo half-life regression) at the spacing-mechanism-in-deployment layer. Duolingo is the large-scale operational deployment of a spacing-and- retrieval-practice system in a consumer language-learning product; Kerfoot is the small-scale randomized evidence in the workplace-L&D medical-training context. Together the two cases anchor the spacing-mechanism conversation across laboratory, deployed-product, and randomized-workplace evidence. The Iterative Development discipline LENS teaches depends on having both the controlled-randomized closed loop (Kerfoot) and the operational-mechanism evidence (Duolingo) available in the corpus.

The learning benefit persisted for two years.

KERFOOT, JOURNAL OF UROLOGY 2009 FOLLOW-UP.

REFERENCES

1. Kerfoot, B. P., Baker, H. E., Koch, M. O., Connelly, D., Joseph, D. B., & Ritchey, M. L. (2007). Randomized, controlled trial of spaced education to urology residents in the United States and Canada. *Journal of Urology*, 177(4):1481–1487. doi:10.1016/j.juro.2006.11.074 — the case's primary RCT.
2. Kerfoot, B. P. (2009). Learning benefits of on-line spaced education persist for 2 years. *Journal of Urology*, 181(6):2671 — the two-year persistence follow-up.
3. Cepeda, N. J., Pashler, H., Vul, E., Wixted, J. T., & Rohrer, D. (2006). Distributed practice in verbal recall tasks: A review and quantitative synthesis. *Psychological Bulletin*, 132(3):354–380 — the basic-literature spacing-effect review the workplace-L&D RCT sits against.
4. Settles, B., & Meeder, B. (2016). A trainable spaced repetition model for language learning. *Proceedings of ACL 2016*, 1848–1858 — Duolingo half-life regression (Case 69), the operational deployment companion.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Kerfoot et al. is the methodologically clean small-tier closed-loop spaced-education RCT in workplace medical training, with two-year persistence in the follow-up. It is the strongest randomized strength in the small-tier batch. The scope discipline is what makes it interpretable: the result speaks to the email-delivered validated-question delivery pattern, not to spacing in general.

LENS APPROACH

Kerfoot is the canonical small-tier closed-loop spacing-in-workplace-L&D case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the intervention-and-measurement-as-single-design discipline, and in Domain 4 (Test and Evaluation) for the persistence- follow-up that makes the result a workplace-L&D headline. Pair with Case 69 (Duolingo half-life regression) at the spacing-mechanism-in-deployment layer — controlled- randomized evidence with deployed-product evidence together.

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
MEDICAL EDUCATION	H K D	D1 D2 D3 D4 D5	4	2.3	2 · 4
RETENTION		Iterative Development			
WORKFORCE L&D					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
- Specify the closed loop in advance: the intervention manipulation, the outcome measure, and the validated-question pool together, so the intervention and the measurement form a single design rather than two separately published artifacts. - Choose the delivery modality (email, in this case) that the trainees will tolerate as part of their work week; the workplace-L&D context is unforgiving of modalities that compete with clinical work. - Pre-register the persistence follow-up at the design stage; the two-year retention figure in Kerfoot 2009 is what gives the intervention its workplace-L&D strength, and follow-up has to be designed into the original RCT.	- Report the result in the scope it speaks to: email-delivered validated-question delivery pattern with urology residents; the spacing-in-general claim sits in the basic literature, and the case's contribution is the workplace-L&D closed-loop evidence. - Treat the two-year persistence figure as the workplace-L&D headline; persistence over years is rare in the workplace-L&D RCT literature, and the figure is what makes the result actionable as a training design. - Pair the result with Duolingo (Case 69) at the spacing-mechanism layer in any communication of the broader conversation, so the controlled-randomized evidence and the deployed-product evidence are read together.

REFLECTION QUESTIONS

1. Identify a workplace-L&D context in your domain where the basic-literature evidence (spacing, retrieval practice, deliberate practice) is robust and the workplace-transfer evidence is thinner. What would a Kerfoot-class closed-loop RCT look like — the intervention manipulation, the outcome measure, the validated-content pool together?
2. The two-year persistence figure is the workplace-L&D headline. What follow-up cadence would you design into your study at the start, so that persistence over years is part of the original RCT rather than a separately commissioned downstream check?
3. The result's scope is the email-delivered validated-question delivery pattern. What is the analog scope discipline in your context — the qualifying language that has to travel with the headline so the result is read as evidence about a specific mechanism rather than as evidence about a general principle?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. TOOLS usability testing · cognitive walkthroughs · knowledge bases · data pipelines · task analysis · design prototyping

CASE 18

GLOBAL HEALTH

HIV CARE

TRAINING-MODALITY COMPARISON

SUB-SAHARAN AFRICA

2023

PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption

K Knowledge & Institutional Memory**N** Normalization of Deviance**D** Designed Out**EVIDENCE TIER** preprint-tier — source confidence flagged; future validation ongoing.

IMPACT Across 16 PEPFAR-supported Sub-Saharan African countries, pre/post knowledge and self-confidence assessments showed mean increases across all training modalities — in-person (pre-COVID), virtual, and blended — when pandemic disruption forced a delivery-mode switch; the L1–L2 limitation is explicit and the L3/L4 question remains open

IN BRIEF

Across 16 PEPFAR-supported Sub-Saharan African countries, a 2023 real-world program evaluation compared in-person (pre-COVID), virtual, and blended HIV-care training delivery. Across all assessed knowledge domains and self-perceived confidence dimensions, pre/post assessments showed mean increases regardless of modality. The honest framing the case carries into print is that the outcome metric is knowledge and self-rated confidence — Kirkpatrick L1–L2 territory (Case 79) — not L3 on-job behavior change or L4 patient outcomes. Even so, the study is one of the few real-world cross-country modality comparisons at meaningful scale, and it is the L&D evaluation pattern playing out in global health at multi-country scope. The case is cross-listed with the corporate / workforce L&D cluster (Cases 79, 65, 83, 70) and with the non-US/UK/EU geographic-coverage gap (Cases 93, and the cases in the later supplemental batches). Evidence-tier flag is preprint-tier: the medRxiv version is preprint and the PMC version is journal-published — the editor's citation choice should be carried explicitly. Future validation on whether L1–L2 knowledge gains translate to L3 behavior change or L4 patient outcomes remains ongoing.

BACKGROUND

PEPFAR (the US President's Emergency Plan for AIDS Relief) has been one of the largest sustained global health programs of the past two decades. Workforce capability — clinicians and frontline workers trained on current HIV diagnosis, treatment, and prevention protocols — is one of its load-bearing deliverables. COVID disruption forced a delivery-mode switch

in 2020 across countries, from in-person training to virtual and blended modalities, before any of the evidence base on modality equivalence had been built for this workforce.⁰

THE INTERVENTION

The 2023 study compares pre/post-assessment outcomes across 16 PEPFAR-supported Sub-Saharan African countries and across the three modalities. The headline finding is that across all knowledge domains and self-perceived confidence dimensions assessed, mean increases were observed regardless of modality. But the modalities were not equivalent: the in-person format produced significantly greater knowledge gains (13.6% increase vs. 6.0% for the virtual workshop and 7.6% for the online course) and greater confidence gains, regardless of training level or professional cadre. The countervailing finding is equity of access — the virtual and online formats reached a significantly higher proportion of women — so the policy question under future disruption is a trade-off between L1–L2 learning gains and reach, not a claim of modality equivalence.¹

HOW IT WORKED

The honest framing the case must carry into print is the Kirkpatrick limitation. The study's outcomes are knowledge and self-rated confidence — Level 1 (reaction and confidence) and Level 2 (learning) in the Kirkpatrick framework (Case 79). It is not Level 3 (behavior change on the job — whether the clinicians actually changed practice) and it is not Level 4 (results — whether patient outcomes improved). The structural limitation Case 79 names is exactly the limitation this study sits inside.²

THE EVIDENCE

Even so, the study is one of the few real-world cross-country modality comparisons at meaningful scale, and the multi-country scope means the comparison is not trivially confounded by single-country labor-market or health-system factors. It is the L&D evaluation pattern playing out in global health, and the modality-comparison finding is policy-relevant for PEPFAR program design across the deployment region.³

WHAT TRANSFERRED

The evidence-tier flag is preprint-tier: the medRxiv version is the preprint; the PMC version is the journal-published article. The editor's citation choice between the two should be carried explicitly in the printed case. The LENS teaching point is the cross-listed pair with the corporate / workforce L&D cluster (Cases 79, 65, 83, 70) and the non-US/UK/EU cluster (Cases 93). The CLO **Judgment under inadequate evidence** is exercised: the study is the strongest evidence the field has for modality equivalence in this context, and it does not establish L3 / L4 outcomes. Future validation is ongoing.

The outcomes are knowledge and self-rated confidence. The L3 and L4 questions — did practice change, did patient outcomes improve — remain open.

EDITORS' SYNTHESIS OF THE PEPFAR 16-COUNTRY MODALITY COMPARISON.

REFERENCES

1. "Comparing in-person, blended and virtual training interventions; a real-world evaluation of HIV capacity building programs in 16 countries in sub-Saharan Africa," medRxiv 2023.02.08.23285641 (preprint) → PMC10365303 (published).
3. Kirkpatrick & Kirkpatrick (2006), Evaluating Training Programs — the chain-of-evidence framework the L1–L2 limitation references (paired Case 79).

2. PEPFAR program documentation — workforce-capability training as a load-bearing deliverable across Sub-Saharan African deployment countries.
4. Blume et al. (2010), Journal of Management 36(4):1065–1105 — the meta-analytic transfer finding the L3 question references (paired Case 65).

THE LEARNING ENGINEERING LENS

LE INSIGHT

PEPFAR’s 16-country modality comparison is the L&D evaluation pattern in global health: mean L1–L2 gains across all modalities, with the Kirkpatrick limitation (Case 79) explicit. Evidence-tier flag is preprint-tier — both the medRxiv preprint and the PMC published version are citable — and the L3 / L4 questions remain open. Future validation is ongoing.

LENS APPROACH

PEPFAR is the global-health workforce-capability case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the modality-design iteration COVID forced and in Domain 5 (Navigating Sociotechnical Constraints) for the multi-country program scope. The CLO **Judgment under inadequate evidence** is exercised because the study is the strongest available at L1–L2 and does not settle L3 / L4. Pairs with Case 93 (SkillsFuture) as the workforce L&D cross-listing.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
GLOBAL HEALTH	KND	D1D2D3D4	4	2.3	2 · 4 · 5
HIV CARE		D5			
TRAINING-MODALITY COMPARISON	→	→ Iterative Development			
SUB-SAHARAN AFRICA					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the modality comparison around what the assessment instrument can actually establish — knowledge and confidence at L1–L2 are tractable; L3 and L4 require longitudinal data and patient-outcome linkage the study did not have.
- ▶ Build the multi-country sampling so the comparison is not trivially confounded by single-country factors — the 16-country scope is part of the study’s evidential weight.
- ▶ Carry the Kirkpatrick limitation into the program documentation honestly — Case 79’s stop-at-L2 pattern is exactly the limitation this study sits inside, and naming it preserves the case’s teaching value.

IN OPERATION / AFTER

- ▶ Pair with Case 93 (SkillsFuture) as the workforce-capability counterparts at national- and multi-country scale; together they teach what L&D measurement looks like in non-US/UK/EU settings with the evidence-tier honesty intact.
- ▶ Use the CLO **Judgment under inadequate evidence**: the study is the strongest evidence the field has for modality equivalence in this context, and the practitioner has to decide modality policy on L1–L2 evidence while L3/L4 evidence develops.
- ▶ Carry the preprint-tier flag honestly: the medRxiv preprint and the PMC published article are both citable; future validation requires confirmatory replication and L3/L4 outcome measurement.

REFLECTION QUESTIONS

1. Identify a training program in your context that was forced to switch modality under disruption (COVID, budget constraint, geographic dispersion). What evidence would you have needed at the time to make the modality decision well, and at which Kirkpatrick level was the available evidence?
2. Specify the L3 / L4 measurement architecture you would build to extend the PEPFAR study into behavior-change and patient-outcome evidence — what data sources, what cadence, what linkage between training cohort and clinical outcomes.
3. The case is preprint-tier: both the medRxiv preprint and the PMC published version are citable. What additional confirmatory evidence — replication in different country sets, L3 / L4 follow-up studies, peer-reviewed meta-analysis — would you require before treating modality equivalence as a settled question for program-design decisions?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards · task analysis · design prototyping

Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff

H Human-System Interface **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT A clinical adaptation of the aviation 'sterile cockpit' rule — no non-essential communication or interruptions during high-workload phases — applied to ward rounds and handoffs; the published evaluation reports reductions in interruption frequency and improvements in information-transfer measures during the protected window, on a single-unit single-study evidence base whose generalization the authors are explicit about

IN BRIEF

The “sterile cockpit” rule in commercial aviation (FAR 121.542, codified 1981) prohibits non-essential communication among flight crew during taxi, takeoff, landing, and any flight phase below 10,000 feet — the high-workload, high-consequence windows when distraction is most likely to produce error. The principle has been adapted across other high-consequence domains; recent work in clinical care has extended the principle to ward rounds and clinical handoff, where interruption-driven information loss is a documented contributor to patient-safety events. Treloar et al. (2025), in *World Journal of Surgery*, describe a structured intervention on a single surgical ward: a defined window during the handoff or round during which non-essential pages, conversations, and interruptions are prohibited and information transfer is the protected workflow. The evaluation reports reductions in interruption frequency and improvements in information-transfer measures during the protected window. The hedges that survive into the case verbatim: this is a single-unit single-study evidence base, the interruption-reduction effect is direct while the patient-outcome inference rests on the established link between handoff quality and downstream events rather than on direct outcome measurement in this study, and the sustainability of the protected-window discipline beyond the study period is not yet established. The case is the cross-domain adaptation case: an aviation safety-culture rule, adapted with explicit attribution and adapted again for a clinical context, with the evidence at the tier the adaptation has reached.

BACKGROUND

The aviation sterile-cockpit rule (FAR 121.542) was codified in 1981 in response to a documented pattern of accidents in which flight-crew distraction during high-workload phases — non-operational conversation, cabin coordination, administrative communication — contributed to the error chain. The rule is operationally simple: below 10,000 feet, in taxi, takeoff, approach, and landing, only operationally necessary communication is permitted in the cockpit. The principle has been extended in safety-culture practice to other high-workload phases and operations.^o

THE INTERVENTION

Clinical care has long carried an analogous problem at the handoff and ward-round boundary. Interruption-driven information loss during handoff is a documented contributor to patient-safety events; I-PASS and SIGNOUT? and similar structured-handoff interventions address the information content, but the workflow context — the interruptions that fragment the handoff — has historically been treated as ambient noise rather than as a design variable. The structural argument the clinical adaptation makes is that the workflow context is a design variable, and that the sterile-cockpit principle provides a worked template for engineering it.¹

HOW IT WORKED

Treloar et al. (2025) describe a structured intervention on a surgical ward. A defined window during the handoff or round is designated as protected; non-essential pages are deferred, side conversations are prohibited, and the staff conducting the handoff are made unavailable for non-emergent interruption during the window. The intervention includes the workflow design (who is responsible for triaging pages, how emergent interruptions are preserved, how the window is signaled and ended) as well as the cultural commitment to honor it. The evaluation measured interruption frequency during the protected window and information-transfer quality against baseline.²

THE EVIDENCE

The published outcomes report reductions in interruption frequency during the protected window and improvements in information-transfer measures during that window. The structural learning the case carries is the adaptation discipline: an aviation safety- culture rule, with its specific operational scoping (below 10,000 feet, defined permitted communication), was carried into a clinical context with the operational scoping translated (a defined window, defined permitted interruption classes, defined triage responsibility) rather than the principle imported without translation. Pair the case with I-PASS (a structured-handoff intervention; Case 15 when drafted) at the handoff-as-capability layer, and with CRM (Case 117) at the aviation-safety-culture-to-clinical-care translation layer.³

WHAT TRANSFERRED

The hedges have to survive verbatim. This is a single-unit single-study evidence base; replication across units and settings is the natural next step the authors identify, and the case should not be generalized beyond a strong structural argument until that replication is in the literature. The interruption-reduction effect is direct; the inference that fewer interruptions and better information transfer reduce downstream patient-safety events depends on the established link in the handoff-quality literature rather than on direct patient-outcome measurement in this study. The sustainability of the protected-window discipline beyond the study period is not yet established — the operational cost of maintaining the workflow design over the long term is the open question. The case is teachable at the adaptation-discipline level today; the multi-site multi-cycle evidence base is the next deliverable.

The scoping is the adaptation. Sterile cockpit is “below 10,000 feet”; the clinical analog is a defined window with defined permitted interruption classes.

EDITORS' SYNTHESIS OF PMC12515027 AND THE STERILE-COCKPIT ADAPTATION LITERATURE.

REFERENCES

1. Treloar, E., Herath, M., Altree, M., Potter, S., Penhall, M., Walsh, D., Kennedy, L., Bruening, M., Edwards, S., Ey, J. D., Bradshaw, E. L., & Maddern, G. J. (2025), "A Simple Solution for a Complex Problem: The 'Sterile Cockpit' to Improve Ward Rounds," *World Journal of Surgery* 49(10):2769–2776, doi:10.1002/wjs.70074, PMID:40930848, PMCID:PMC12515027 — the cited adaptation study.

2. Federal Aviation Administration, 14 CFR § 121.542 (codified 1981) — origin of the aviation sterile-cockpit rule.

3. Starmer, A. J. et al. (2014), "Changes in medical errors after implementation of a handoff program," *New England Journal of Medicine* 371(19):1803–1812 — I-PASS handoff intervention; structural cousin in the structured-information half of handoff capability.

4. Broom, M. A., Capek, A. L., Carachi, P., Akeroyd, M. A., & Hilditch, G. (2011), "Critical phase distractions in anaesthesia and the sterile cockpit concept," *Anaesthesia* 66(3):175–179 — prior anesthesia-domain sterile-cockpit adaptation establishing the cross-domain transfer pattern.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The sterile-cockpit ward-rounds case is the worked example of cross-domain adaptation discipline at small scale: an aviation safety-culture rule, carried into clinical care with its operational scoping translated and its cultural half preserved. The single-unit evidence is direct on interruption frequency and information transfer; the patient-outcome inference rests on the established handoff-quality link, and the multi-site replication is the explicit next deliverable.

LENS APPROACH

Sterile-cockpit ward rounds is the cross-domain adaptation case (induced 3.2; LENS D3/PT5) — Domain 3 for workflow- context-as-design-variable; Domain 5 for the cultural- commitment half. Pair with I-PASS (Case 15), CRM (Case 117), and Case 27.

DOMAIN	MODES ADDRESSED			LENS COMPETENCY		PROBLEM TYPE	INDUCED	CLO					
HEALTHCARE	➔	H	K	N	➔	D1	D2	D3	D4	D5	5	3.2	4 · 5
PATIENT SAFETY		Human-System Collaboration											
HUMAN FACTORS													

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Adapt the principle with its operational scoping translated, not imported. Sterile cockpit is “below 10,000 feet, defined permitted communication”; the clinical analog is a defined window, defined permitted interruption classes, defined triage responsibility — the scoping is the adaptation.
- ▶ Design the triage workflow that handles emergent interruptions during the protected window; the rule cannot be “no interruptions” without a path for true emergencies, and the design of that path is part of the deliverable.
- ▶ Build the cultural-commitment half alongside the workflow half; the rule depends on the unit’s willingness to honor it under workload pressure, and the cultural commitment is what makes the rule operative.

IN OPERATION / AFTER

- ▶ Carry the single-unit single-study hedge into communication; the case is teachable at the adaptation-discipline level today, with the multi-site replication as the explicit next deliverable.
- ▶ Distinguish the direct measure (interruption-frequency and information-transfer-quality reductions during the protected window) from the inferred outcome (downstream patient-safety events), and report them at their respective evidence tiers.
- ▶ Plan the long-run sustainability measurement: does the protected-window discipline survive six months, twelve months, leadership turnover? The operational-cost half of the intervention is what the multi-cycle evidence base will eventually answer.

REFLECTION QUESTIONS

1. Identify a high-workload, high-consequence window in your domain that is currently treated as ambient workflow rather than as a protected period. What would the operational scoping of a sterile-window adaptation look like — permitted communication classes, triage responsibility, signaling?
2. Specify the cultural-commitment half of the rule. The workflow design is necessary but not sufficient; honor under workload pressure is what makes the rule operative, and that depends on leadership and unit culture more than on policy.
3. The case's evidence is single-unit single-study. Design the multi-site replication you would want to see before treating the adaptation as established, and the long-run sustainability measurement you would use to know whether the protected-window discipline survives leadership turnover.

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · incident analysis · safety dashboards

ChatGPT in Healthcare — Hallucination Cases

H Human–System Interface **D** Designed Out

IMPACT Documented cases of clinicians using LLMs that produce hallucinated citations, fabricated dosages, and fictitious clinical guidelines

IN BRIEF

Since ChatGPT's public release in late 2022, the clinical and peer-reviewed literatures have documented a recurring pattern: clinicians use large language models to draft patient education, summaries, or treatment guidance, and the output contains fabricated citations, hallucinated drug dosages, or fictitious clinical guidelines. The failures range from cosmetic — invented references in academic submissions — to potentially clinical, such as unsafe dosing or fabricated contraindications. The capability gap is at the human–verification interface: the model presents hallucinated content with the same fluent confidence as accurate content, and early reports suggest clinicians accept LLM output less critically than a colleague's. The case is the live, foundational case for LLM integration into clinical workflow — the discipline must specify what verification looks like at the moment of use, while deployment is already happening.

THE SHIFT

Large language models became broadly available with ChatGPT's release in late 2022, and clinicians began using them almost immediately — to draft patient-education materials, summarize records, and look up guidance. For the first time, a tool that produces fluent, authoritative-sounding medical text is in routine, informal use at the point of care — adopted ahead of any guideline, credential, or institutional sign-off, so the practice spread faster than any structure to govern it could be put in place.⁰

WHAT IS EMERGING

A recurring failure pattern has been documented across the clinical and peer-reviewed literatures: LLM output containing fabricated citations, hallucinated drug dosages, or fictitious clinical guidelines. The failures range from cosmetic — invented references in academic submissions — to potentially clinical, such as unsafe medication doses or invented contraindications — a span that matters because the same tool, used the same way, can produce a harmless error and a dangerous one with no change in how confident it sounds.¹

THE CAPABILITY QUESTION

The capability gap is at the human–verification interface. The model presents hallucinated content with exactly the same fluent confidence as accurate content; the interface does not distinguish the two. The question is whether clinicians can — and will — develop the routine verification practice that the tool's fluency actively discourages, since the very smoothness


that makes the output easy to accept is what removes the friction a reader would normally use as a warning.²

EARLY EVIDENCE

Early case reports suggest that clinicians who would carefully check a colleague's recommendation accept LLM output less critically, precisely because it reads so authoritatively. JAMA editorials and reviews through 2023–2024 have repeatedly flagged the absence of an established verification practice as the central risk of clinical LLM use — the concern is not that the model errs but that the practice for catching its errors at the moment of use has not yet been defined or taught.³

OPEN PROBLEMS

This is the live frontier case for human–AI teaming when the AI is fluent across both accurate and hallucinated content. The capability that does not yet exist is a routine clinical verification practice — an analog to the bibliographic discipline of academic writing — specified at the moment of use rather than after harm. The discipline is being asked to define what good looks like while deployment is already underway, so the verification standard has to be built around a tool already in millions of hands rather than gated in front of it.⁴



LLMs produce hallucinations indistinguishable in tone from accurate information, and clinicians have not yet developed the practice of routine verification.

EDITORS' SYNTHESIS OF JAMA EDITORIALS ON LLM CLINICAL USE (2023–2024)

REFERENCES

1. JAMA editorials on LLM integration into clinical practice (2023–2024) — the hallucination/verification problem (synthesized).
2. Sallam (2023), “ChatGPT Utility in Healthcare Education, Research, and Practice” — documented benefits and risks.
3. Case reports of LLM hallucinated citations and dosages in clinical and academic use (2023–2024).
4. WHO ethical guidance on AI for health (2024) — verification and oversight requirements.
5. Wachter & Brynjolfsson (2023), JAMA — generative AI in health care.

THE LEARNING ENGINEERING LENS

LE INSIGHT

LLM use in clinical practice is the live frontier case for human-AI teaming when the AI’s output is fluent across both accurate and hallucinated content. The capability that does not yet exist is the routine verification practice — a clinical analog to the bibliographic discipline of academic writing.

LENS APPROACH

The learning engineer’s deliverable here is a specific artifact, not a posture: a verification-at-point-of-use guardrail that counters cognitive offloading. Concretely, a confirm-before-act gate wrapped around the LLM output — every dosage, contraindication, or guideline citation the clinician would rely on must clear a forced source-attribution step that surfaces the underlying reference (or its absence) and requires an explicit confirmation against it before the output can be acted on. The gate is designed so the model’s fluency cannot satisfy it: a smooth, confident answer with no attributable source fails the check and is held back from patient-facing use. LENS builds and critiques exactly this guardrail in LEN 2 (the human-AI verification interface), stresses it in LEN 7 (governing the gate as deployment policy), and in LEN 10 (capstone) asks the student to design the confirm-before-act artifact for a clinical workflow already in use.

DOMAIN	MODES AT STAKE		LENS COMPETENCY					PROBLEM TYPE	INDUCED	CLO	
HEALTHCARE	→	H D	→	D1	D2	D3	D4	D5	6	5.2	4 · 3
TECHNOLOGY	Human-System Collaboration										

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Specify a routine verification practice at the moment of use — what a clinician must independently confirm before acting on LLM output — before the tool enters clinical workflow.
- ▶ Design the interface to mark provenance and uncertainty, so fluency alone cannot stand in for warranty of accuracy.
- ▶ Restrict the sanctioned uses to those where verification is feasible and cheap, keeping unverifiable high-stakes outputs out of patient-facing work.

IN OPERATION / AFTER

- ▶ Monitor the documented failure pattern — fabricated citations, dosages, guidelines — across in-use cases to keep the risk profile current as models change.
- ▶ Track whether the verification practice is actually being performed, since the tool’s fluency discourages exactly the checking it requires.
- ▶ Govern adoption against emerging guidance so the standard for verification keeps pace with a tool already in widespread informal use.

REFLECTION QUESTIONS

1. Identify a workflow in your domain currently being augmented by LLMs. What is the verification practice — and does it exist at the moment of use, or only after?
2. Design the verification deliverable that should accompany every clinician’s adoption of an LLM tool for patient-facing work.
3. The same fluency that makes LLM output easy to accept is what removes the cues a reader normally uses to doubt it. What interface signal would restore that friction at the moment of use without making the tool unusable?

WHO BUILDS THIS human factors & interaction design · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. TOOLS usability testing · cognitive walkthroughs · task analysis · design prototyping

WHO Surgical Safety Checklist

T Training Gap **N** Normalization of Deviance

IMPACT Death rate 1.5% → 0.8% in eight-site pilot; complications down >33%; adopted by the majority of surgical providers worldwide; Ontario population study (2014) found no significant mortality benefit after a mandated rollout

IN BRIEF

In 2008, Atul Gawande's team and the WHO introduced a single-page, nineteen-item surgical checklist applied at three junctures — before anesthesia, before incision, and before the patient leaves the operating room. Piloted across eight hospitals in eight countries, from Toronto to Tanzania, it nearly halved the surgical death rate (1.5% to 0.8%) and cut serious complications by more than a third — results published in the NEJM in 2009. The artifact was the checklist; the intervention was the system of forced pauses it created, three moments when a moving team had to stop and confirm shared state. A later Ontario study found no mortality benefit after a mandated rollout, surfacing the fidelity question: the artifact works only when the institution authorizes its honest use.

BACKGROUND

Surgical harm — wrong-site operations, retained instruments, missed allergies, post-operative infection — was widespread across health systems rich and poor, and much of it stemmed not from a lack of skill but from teams that never paused to confirm shared understanding before acting. The knowledge to prevent these events existed; the reliable practice did not. A surgical team in motion rarely stopped to verify that everyone held the same picture of the patient and procedure, so a mismatch any member could have caught passed silently into the operation. The cost was global and patterned — the WHO estimated that, of the roughly 234 million operations performed each year, at least seven million produced major complications and about a million ended in death — and the bulk of that harm sat in failure modes that a brief verbal cross-check would have closed.⁹

THE INTERVENTION

In 2008, the WHO and a Harvard team led by Atul Gawande introduced a single-page, nineteen-item Surgical Safety Checklist applied at three critical junctures: before anesthesia, before skin incision, and before the patient leaves the operating room. The team piloted it across eight hospitals in eight countries spanning very different economic and operational conditions — Toronto, London, Seattle, Auckland, Amman, New Delhi, Manila, and Ifakara — selecting sites deliberately to span the global income gradient. Testing the same artifact from Toronto to Tanzania was deliberate — it had to work in settings with vastly different resources to demonstrate that the gain came from the forced pause itself, not from any one system's wealth. The artifact was intentionally minimal: items the team had tested

against the surgical-safety literature for the smallest set that still spanned the highest-risk junctures.¹

HOW IT WORKED


The checklist was the artifact; the intervention was the system of pauses it imposed. At each juncture a team that would otherwise keep moving had to stop, look at one another, and confirm names, the procedure, allergies, and equipment aloud. The pauses were the load-bearing element — the requirement to halt and speak mattered more than the specific list of items. Saying the confirmations aloud to one another, rather than each member assuming them privately, is what surfaced the mismatches a moving team would otherwise have carried into the incision. The mechanism turned the operating room into a momentarily flat-hierarchy space: a circulating nurse who noticed a missing antibiotic dose at the time-out had institutional cover to say so before the incision, a transaction the pre-checklist culture had not reliably authorized. The artifact's authority — that the team had to stop — was what made the speech act safe.²

THE EVIDENCE

The 2009 NEJM study reported the surgical death rate falling from 1.5% to 0.8% and major complications dropping by more than a third across all eight sites — confirming Gawande's framing that "gaps in teamwork and safety practices in surgery are substantial in countries both rich and poor." The checklist was subsequently adopted by the majority of surgical providers worldwide. That the death rate roughly halved across all eight sites, despite their differences in wealth and practice, is what made the result so persuasive — the effect tracked the intervention rather than the setting. The published mortality reduction was contested almost from the moment of publication: observational, pre/post, no concurrent control, susceptible to secular trends and Hawthorne-style attention effects. The pilot was always best read as the formation of a strong hypothesis under pragmatic conditions, not as a closed proof of population-scale mortality benefit.³

WHAT TRANSFERRED

Urbach and colleagues, writing in NEJM in 2014, examined more than 200,000 procedures across 101 Ontario hospitals before and after a province-wide mandated rollout and found no statistically significant reduction in operative mortality, in-hospital complications, length of stay, or readmission. The null result did not refute the checklist so much as illustrate its dependence on implementation fidelity: where a mandate replaced genuine authorization of the pause, the measured effect attenuated, making the checklist a paired lesson in both minimal-artifact design and the limits of mere compliance. The contrast between the pilot and the mandated rollout is the lasting teaching point: when the pause was genuinely used it worked, and when it became a box to check before proceeding the same paper produced nothing. Subsequent multi-country replications and the WHO's own follow-up work have been mixed — the institutional uptake outran the closed evidence of population-scale mortality benefit, and the contested mortality reduction is the part of the case that does not resolve.⁴



Gaps in teamwork and safety practices in surgery are substantial in countries both rich and poor.

ATUL GAWANDE, HARVARD GAZETTE, JANUARY 2009

REFERENCES

1.

Haynes, A. et al. (2009), “A Surgical Safety Checklist to Reduce Morbidity and Mortality,” NEJM — the 1.5%→0.8% result and major-complication reduction.

2.

WHO Safe Surgery Saves Lives campaign documentation — the nineteen-item checklist and the three junctures.

3.

Gawande, A. (2009), The Checklist Manifesto — the pause as the active mechanism.

4.

Urbach, D. et al. (2014), “Introduction of Surgical Safety Checklists in Ontario, Canada,” NEJM — null mortality result after mandate.

5.

Bosk, C. et al. (2009), “Reality check for checklists,” The Lancet — implementation fidelity.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Surgical Safety Checklist is the canonical evidence that a tiny artifact — one page, nineteen items — can produce population-scale effects when paired with the structural change of **requiring a pause**. The checklist alone is paper. The pause alone is anxiety. Together they constitute the smallest effective capability intervention in the dataset. The Ontario follow-up underscores the secondary requirement: the artifact carries the effect only when the institution actually enforces the pause. Where mandate replaced authorization, the measured effect attenuated.

LENS APPROACH

LENS uses the WHO checklist as the canonical pre-committed institutional response paired with closed-loop evidence (induced 4.4; LENS D4/PT5). It runs in LEN 10 (capstone) as the end-to-end design exemplar — a problem identified, a minimal artifact prototyped, a multi-site pilot, a measurement regime, and global scale-out — and in LEN 4 for the measurement architecture that made the pilot effect provable and the Ontario null interpretable. Adjacent to SUBSAFE (Case 173) at the mandatory-mechanism layer, where the artifact’s authority is the intervention.

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
HEALTHCARE	→ T N	→ D1 D2 D3 D4 D5	5	2.3	4 · 5
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

IN OPERATION / AFTER

- ▶ Design the minimal artifact — a single-page checklist — around a small number of forced pauses at the highest-risk junctures, keeping the list short enough to be used every time.
- ▶ Make the confirmations spoken aloud and shared across the team, so the pause surfaces mismatches rather than letting each member assume them privately.
- ▶ Pilot across deliberately varied sites so the measured gain can be shown to track the intervention rather than any one setting's resources.
- ▶ Measure outcomes (death and complication rates) directly so the effect can be confirmed and the artifact is not assumed to work merely because it is in use.
- ▶ Guard implementation fidelity as the scale grows — a mandate that turns the pause into a box to check reproduces the artifact without the effect.
- ▶ Build a governance signal that distinguishes genuine, authorized use from compliance check-off, so the fidelity gap is visible before a rollout is judged a failure.

REFLECTION QUESTIONS

1. What is the smallest possible capability artifact in your domain that, paired with a required pause, would shift outcomes?
2. The WHO checklist was studied across eight countries. Design the multi-site evaluation that would establish whether your candidate intervention generalizes.
3. The Ontario mandated rollout produced no measurable mortality reduction. What governance signal would have surfaced the fidelity gap between authorized use and compliance check-off before the trial was declared a failure?

WHO BUILDS THIS learning science & instructional design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards

Bristol Heart Babies Reform

G Governance & Trust **N** Normalization of Deviance

IMPACT Foundational UK case in clinical outcomes transparency; produced specialty-level performance reporting in UK cardiac surgery

IN BRIEF

Through whistleblowing and a public inquiry, the Bristol Royal Infirmary was found to have had substantially worse pediatric cardiac-surgery outcomes than other UK centers for years. The Kennedy Inquiry — one of the most influential UK healthcare inquiries in modern times — located the capability gap in the absence of routine specialty-level outcomes reporting: surgeons did not know how their results compared with peers, patients did not know, and referrals did not respond to actual outcome data. The reform built national specialty-level outcomes registries, first in cardiac surgery and then across other specialties, making the UK one of the few countries that routinely publishes surgeon-level results — a paired intervention of technical registry plus institutional commitment to publish that ended outcomes data as the private property of clinicians.

BACKGROUND

In pediatric cardiac surgery, small differences in a unit's performance can mean the difference between a child living and dying — yet in the UK of the 1980s and early 1990s, no system routinely compared outcomes between centers. A surgeon, a hospital, or a referring physician had no reliable way to know whether a given unit's results were normal or dangerously poor. Without comparison across centers, a dangerously poor unit and an ordinary one looked alike from inside, so the very gap that cost children's lives was the one no one was positioned to see.⁹

THE INTERVENTION

After whistleblowing and a public inquiry into deaths at the Bristol Royal Infirmary between 1984 and 1995, the Kennedy Inquiry recommended routine, risk-adjusted, specialty-level outcomes reporting. The reform built national registries — beginning with cardiac surgery and extending to other specialties — together with a commitment to publish results, including at the level of individual surgeons. Starting in cardiac surgery and then extending outward was deliberate sequencing — the specialty where the harm had been exposed proved the model, and the registry then spread to fields that had not had their own Bristol.¹

HOW IT WORKED

The intervention was explicitly paired. The technical half — registry design, statistical risk adjustment so that surgeons taking hard cases are not penalized, and a publication architecture — was necessary but not sufficient. The cultural half — surgeons accepting that their comparative results would be visible — was equally load-bearing, and was the harder of

the two to secure. The risk adjustment was what made the cultural half securable: without it, surgeons who took the sickest patients would have been punished by the raw numbers, giving them every reason to resist publication or avoid hard cases.²

THE EVIDENCE

The UK became one of the few countries that routinely publishes surgeon- and unit-level outcomes, and the cardiac-surgery registry is among the most mature specialty-outcomes regimes in any country. Outcomes data ceased to be the private property of clinicians and became a shared resource for patients, referrers, and surgeons themselves. That surgeons themselves became users of the data, not just subjects of it, is part of why the regime endured — comparison that had once felt like exposure became a tool the profession relied on to know where it stood.³

WHAT TRANSFERRED

Bristol is the foundational UK case for outcomes transparency as a paired-intervention deliverable, and its registry model has been extended across NHS specialties and influenced later national-quality reforms. It pairs with Keystone ICU (Case 19) as healthcare-outcomes interventions operating at different layers — the bedside protocol and the system-level measurement regime. The two layers are complementary rather than competing: Keystone changes what happens at the point of care, while Bristol changes what the system can see about results across centers, and a mature regime needs both.⁴



Outcomes data must cease to be the private property of clinicians and become a shared institutional resource.

PARAPHRASING THE KENNEDY INQUIRY FINAL REPORT (LEARNING FROM BRISTOL), 2001

REFERENCES

1. Kennedy, I. (2001), Learning from Bristol: The Report of the Public Inquiry into Children's Heart Surgery at the Bristol Royal Infirmary 1984–1995 — the inquiry and recommendations (paraphrased).
2. Society for Cardiothoracic Surgery in GB & Ireland, national outcomes reports — the registry and surgeon-level publication.
3. Berwick, D. (2013), A Promise to Learn — A Commitment to Act — the broader NHS-safety reform.
4. Sherlaw-Johnson et al. — risk-adjusted outcome methodology.
5. NHS national clinical audit and registry documentation — extension across specialties.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Bristol is the canonical UK case for outcomes-transparency as a paired intervention. The technical capability — registry design, statistical risk adjustment, publication architecture — was necessary. The cultural capability — surgeons accepting that their results would be visible and comparable — was equally necessary. The pair has produced one of the most mature specialty-outcomes regimes in any country.

LENS APPROACH

LENS uses Bristol in LEN 4 for outcomes-transparency as a paired- intervention deliverable and in LEN 7 for the institutional reform that made surgeon-level publication acceptable. The case pairs with Keystone ICU (Case 19) as healthcare-outcomes interventions at different layers.

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
HEALTHCARE	→ G N	→ D1 D2 D3 D4 D5	5	2.1	4
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pair the technical instrument — a registry with statistical risk adjustment and a publication architecture — with the cultural change of practitioners accepting that comparative results will be visible.
- ▶ Build risk adjustment in from the start so those who take the hardest cases are not punished by raw numbers, which is what makes the cultural half securable.
- ▶ Sequence the rollout to begin where the harm was exposed, proving the model in one specialty before extending it to fields that have not had their own crisis.

IN OPERATION / AFTER

- ▶ Make practitioners users of the data, not just subjects of it, so comparison becomes a tool the profession relies on rather than an exposure it resents — which is what sustains the regime.
- ▶ Extend the registry model across specialties over time, turning a single reform into a system-wide measurement regime.
- ▶ Pair the system-level transparency layer with point-of-care interventions (as with Keystone), since a mature outcomes regime needs both what the system can see and what happens at the bedside.

REFLECTION QUESTIONS

1. What is the equivalent of surgeon-level outcomes transparency in your domain? What cultural change would have to accompany the technical instrument?
2. Design the registry and publication architecture for a specialty in your domain currently operating without outcomes transparency.
3. Bristol's risk adjustment was what let surgeons accept publication, by ensuring those who took the hardest cases were not punished by raw numbers. Identify a transparency measure in your domain that practitioners resist, and design the adjustment that would make the comparison fair enough to accept.

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

Removing the Race Coefficient from eGFR

D Designed Out **G** Governance & Trust **N** Normalization of Deviance

IMPACT A clinical estimating equation that for two decades raised estimated kidney function for Black patients was retired through a governance process; the documented effect on disparities remains unknown

IN BRIEF

For two decades the standard equation used to estimate glomerular filtration rate from serum creatinine — the kidney-function number on routine lab reports — included a coefficient that raised the estimate for patients reported as Black. The downstream effects were documented: later nephrology referral, later wait-listing for transplant. In 2021 the joint NKF-ASN Task Force, after reviewing over twenty candidate approaches with patient and clinician input, recommended immediate national implementation of the race-free 2021 CKD-EPI creatinine equation (Inker et al., NEJM 2021); clinical laboratories moved to adopt it. The case is a governance intervention — a change-control process that retired a construct after three decades of operational use. The honest hedge, preserved from the Task Force report and from follow-up commentaries, is that the new equation introduces a small bias for all groups, the disparities effect **remains unknown**, and the case is the construct-definition act, not a closed outcome proof. It pairs with pulse oximetry (Case 26) and pain assessment (Case 6) as the trio of “what counts as the same patient across race.”

BACKGROUND

Glomerular filtration rate (GFR) is the standard measure of kidney function and the threshold variable for nephrology referral, medication dosing, and transplant wait-listing. Because direct GFR measurement is impractical at clinic scale, equations estimate it from serum creatinine and demographic inputs. The dominant equations from the late 1990s onward included a Black-race coefficient that raised the estimate for patients reported as Black — making them appear to have better kidney function than the same creatinine value implied for White patients.⁹

THE INTERVENTION

The downstream effects of that coefficient are documented across the nephrology literature: later referral to specialty care, later transplant wait-listing, and a population-level shift in who was counted as having advanced chronic kidney disease. The defenders of the coefficient pointed to differences in creatinine generation across self-identified populations; the critics pointed to the logical and ethical problems of embedding a population-level adjustment in an individual-patient decision tool, and to the construct-definition question — what is the equation supposed to be measuring, and across whom does it have to mean the same thing?¹

HOW IT WORKED

The NKF–ASN Task Force, formed in 2020, ran the construct-revision process as a governance exercise: an external panel, patient–clinician input, more than twenty candidate alternatives, a sustained review window, and a published report (Delgado et al., *Am J Kidney Dis* 2021). The Task Force recommended immediate national adoption of the 2021 CKD–EPI creatinine equation, which eliminates the race coefficient. The replacement equation was published in parallel (Inker et al., *NEJM* 2021). Major laboratories and health systems moved to adopt the new equation within months.²

THE EVIDENCE

The hedge in both the Task Force report and the follow-up commentary is load-bearing and is preserved in the case. The new equation introduces a small bias for all groups relative to measured GFR; the **net effect on documented disparities in nephrology care and transplant access** is not yet measured at population scale, and the literature explicitly states that the disparities outcome **remains unknown**. The case is the construct-definition act — the right kind of governance intervention, run with the right kind of process — and the outcome evidence is the continuing work.³

WHAT TRANSFERRED

What the eGFR case teaches is that some equity capability deliverables are construct-definition acts: choosing what gets predicted and what counts as the same patient is the design decision, not a downstream remediation. It pairs with pulse oximetry (Case 26), where the failure was in the validation architecture rather than the equation; and with pain assessment (Case 6), where the failure was in clinician-held false beliefs rather than the instrument. The trio together is the case-grounded basis for the CLO **Gap attribution** — distinguishing capability gaps attributable to human development, system design, and organizational performance — and for the CLO **Fairness beyond omission**, of which eGFR is the construct-removal instance.

What counts as the same patient across race is a construct-definition decision, not a downstream remediation.

EDITORS' SYNTHESIS OF DELGADO ET AL. (2021) AND VYAS ET AL. (2020).

REFERENCES

1. Delgado et al. (2021), "A Unifying Approach for GFR Estimation: Recommendations of the NKF–ASN Task Force on Reassessing the Inclusion of Race in Diagnosing Kidney Disease," *American Journal of Kidney Diseases* 79(2):268–288 (published online 2021; in print vol. 79, 2022), doi:10.1053/j.ajkd.2021.08.003. Cited by online-first year, the year of the Task Force recommendation.
2. Inker et al. (2021), "New Creatinine- and Cystatin C-Based Equations to Estimate GFR without Race," *New England Journal of Medicine* 385(19):1737–1749, doi:10.1056/NEJMoa2102953.
3. Eneanya, Yang, & Reese (2019), "Reconsidering the Consequences of Using Race to Estimate Kidney Function," *JAMA* 322(2):113–114 — the equity argument that motivated the revision.
4. Vyas, Eisenstein, & Jones (2020), "Hidden in Plain Sight — Reconsidering the Use of Race Correction in Clinical Algorithms," *NEJM* 383(9):874–882 — broader race-in-algorithms survey.

THE LEARNING ENGINEERING LENS

LE INSIGHT

eGFR is the canonical recent instance of construct-definition as an equity capability deliverable. A race coefficient was embedded in a continuous estimating equation for two decades, documented downstream effects on referral and transplant access, and was retired through a governance process. The disparities effect of the change remains unknown; the case is the construct-revision act, not the closed outcome.

LENS APPROACH

eGFR is the construct-definition equity intervention in the race-construct trio (Cases 25, 26 and 6). LENS uses it in Domain 4 (Test and Evaluation) for the CLO **Gap attribution** — attributing the gap to construct definition, not to clinicians or patients — and in Domain 5 (Navigating Sociotechnical Constraints) for the governance process that revised the standard. Direct trio with Case 26 (pulse oximetry — the validation-architecture mechanism) and Case 6 (pain assessment — the human-development mechanism). Adjacent to the lending pair (Cases 186–113) at a different layer.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
CLINICAL MEDICINE	→ D G N	→ D1 D2 D3 D4 D5	5	8.1	4 · 5
HEALTH EQUITY	Test and Evaluation				

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Reconsider the construct: ask what the equation should predict and across whom it should mean the same thing, separately from optimizing the residual.
- ▶ Run construct-revision as a governance process: an external panel, patient and clinician input, candidate alternatives examined openly, and a published report.
- ▶ Carry the hedge — the small all-group bias introduced and the unmeasured disparities effect — into the deployment documentation, not the press release.

IN OPERATION / AFTER

- ▶ Instrument the downstream outcomes (nephrology referral, transplant wait-listing) demographically so the disparities effect of the construct revision is actually measurable at population scale.
- ▶ Treat construct revision as ongoing: a coefficient was removed; a new validation architecture is the longer work.
- ▶ Make the construct-definition decision auditable — publish the candidate alternatives that were considered and the basis on which the chosen alternative was selected.

REFLECTION QUESTIONS

1. Identify a clinical or algorithmic construct in your domain that embeds a population-level adjustment in an individual-patient decision. Was the construct revised, examined, or never questioned? What governance process would you run if it had to be revisited?
2. Specify the downstream outcomes you would instrument demographically to make the disparities effect of a construct revision measurable at population scale.
3. The new CKD-EPI equation introduces a small bias for all groups and the disparities effect **remains unknown**. What is the minimum follow-up evidence you would require before concluding the construct revision improved or worsened the equity outcome you care about?

WHO BUILDS THIS systems & human-factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

Anesthesia Monitoring Standards and the APSF — The Mortality Decline

H Human–System Interface **K** Knowledge & Institutional Memory **G** Governance & Trust

IMPACT Continuous pulse oximetry and capnography — encouraged in the 1986 Harvard/ASA minimum-monitoring standards and mandated in subsequent revisions (pulse oximetry ~1989–90; capnography ~1991) — converted silent hypoxemia and esophageal intubation from undetectable failures into monitored, recoverable ones; anesthesia-related mortality fell dramatically over subsequent decades — multifactorial decline

IN BRIEF

Through the early 1980s, anesthesia in the United States carried a documented patient–safety crisis: silent intraoperative hypoxemia and unrecognized esophageal intubation produced catastrophic outcomes — brain injury and death — that were structurally undetectable until the harm had occurred. A 1982 ABC news special on anesthesia catastrophes converged with a malpractice–insurance crisis to force institutional attention. In 1986 Eichhorn et al. (JAMA, 1986) published the Harvard Medical School minimum monitoring standards — which mandated an oxygen analyzer and disconnection alarm and **encouraged** the newer continuous pulse oximetry and capnography, the load-bearing additions designed to make hypoxemia and misplaced endotracheal tubes detectable early enough to recover; both were mandated in subsequent revisions (pulse oximetry 1989–90, capnography 1991). The ASA adopted essentially the same standards in 1986. The Anesthesia Patient Safety Foundation, founded in 1985, institutionalized the broader change–management effort. Anesthesia-related mortality fell dramatically over subsequent decades — one widely cited Brazilian series reports a fall toward zero — and malpractice premiums declined alongside. The field's own histories preserve the hedge: the decline has multiple co-varying causes (training, device design, pharmacology, team composition), and the device standards themselves still fail in documented edge cases.

BACKGROUND

Through the early 1980s, anesthesia in the United States carried a documented patient–safety crisis with two structurally similar failure modes at its center. Silent intraoperative hypoxemia — falling oxygenation in a paralyzed, ventilated patient — produced brain injury and death without the clinical signal a conscious patient would have given. And esophageal intubation — endotracheal tube misplaced in the esophagus instead of the trachea — produced no ventilation at all, often noticed only when the patient was already decompensating. Both failures were detectable in principle and in practice undetectable, because the cue did not reach the anesthetist in time.^o

THE INTERVENTION

The institutional forcing function arrived from outside the field. A 1982 ABC news special documented a string of catastrophic anesthesia outcomes, and the malpractice insurance environment for anesthesiology had reached a crisis level that made the status quo untenable. In 1985 the Anesthesia Patient Safety Foundation was founded — among the first specialty-organized patient-safety foundations in medicine. In 1986 Eichhorn et al. published in JAMA the Harvard Medical School minimum monitoring standards: a defined minimum set of continuous monitors for every anesthetized patient. Pulse oximetry and capnography — then just entering wide use — were **encouraged** rather than required in the 1986 text, and became the load-bearing additions when subsequent revisions mandated them (pulse oximetry 1989–90, capnography 1991).¹

HOW IT WORKED

Pulse oximetry made oxygenation continuously visible to the anesthetist; capnography made the end-tidal carbon-dioxide trace visible, which immediately distinguishes tracheal from esophageal intubation in the first breaths. Together they converted two structurally undetectable failures into monitored, recoverable ones. The American Society of Anesthesiologists adopted essentially the same minimum standards in 1986–87. Eichhorn's follow-up (Anesthesiology, 1989) reported declines in preventable mishaps in the Harvard hospitals consistent with the mechanism the standards targeted.²

THE EVIDENCE

Anesthesia-related mortality fell dramatically over the subsequent decades. A widely cited Brazilian series reported decline toward zero. Malpractice premiums for anesthesiology declined alongside — one of the few examples in American medicine of a specialty's malpractice cost falling as the specialty's safety record improved. The case is regularly cited as a canonical example of the cue/alert design as the capability deliverable: when the right signal is made visible at the right point in the workflow, an undetectable failure becomes recoverable, and the institutional outcome record moves.³

WHAT TRANSFERRED

The hedges the field's own histories preserve are load-bearing. The mortality decline has multiple co-varying causes — anesthetist training, device design improvements, pharmacological change (newer agents are inherently safer), team composition with broader monitoring presence, and the monitoring standards themselves — and attribution of the entire effect to the standards overstates what the evidence can support. The device standards themselves still fail in documented edge cases: pulse oximetry is unreliable across skin tones (Case 26 in this same v2 batch), and capnography can mislead in specific physiology. The case teaches the cue/alert-design form of capability engineering at its most durable, with the honest qualification that the institutional decline is consistent with the standards but not isolatable to them alone.⁴

The deliverable is not the device. It is the standard that makes the device's signal non-waiverable across the specialty.

EDITORS' SYNTHESIS OF EICHORN ET AL. (1986) AND APSF FOUNDING DOCUMENTS.

REFERENCES

1. Eichhorn, Cooper, Cullen, Maier, Philip, & Seeman (1986), “Standards for Patient Monitoring During Anesthesia at Harvard Medical School,” JAMA 256(8):1017–1020.

2. Eichhorn (1989), “Prevention of intraoperative anesthesia accidents and related severe injury through safety monitoring,” Anesthesiology 70(4):572–577.

3. Anesthesia Patient Safety Foundation (1985 – present), founding documents and the APSF Newsletter — institutional-history record of the broader change effort.

4. American Society of Anesthesiologists (1986), “Standards for Basic Anesthetic Monitoring” — original ASA standard.

5. Sjoding et al. (2020), NEJM — the racial-bias edge case the standard still carries (cross-reference Case 26).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Harvard / ASA / APSF anesthesia-monitoring intervention is the canonical cue/alert-design success: silent intraoperative hypoxemia and esophageal intubation were converted from undetectable failures into monitored, recoverable ones, and the specialty’s mortality and malpractice record moved over subsequent decades. The decline is multifactorial; the device standards still fail in documented edge cases (Case 26).

LENS APPROACH

Anesthesia monitoring is the canonical cue/alert intervention (induced 3.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the cue/alert design as deliverable; in Domain 5 (Navigating Sociotechnical Constraints) for the APSF / ASA institutional discipline that made the standard non-waiverable; and in Domain 3 (Human-System Collaboration) at the human-device boundary. Adjacent to Case 26 (pulse oximetry across skin tones) — same instrument, the canonical edge case the standard still carries. Pair with Case 23 (WHO Surgical Checklist) at the mandatory-standard layer.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ANESTHESIOLOGY	H K G	D1 D2 D3 D4 D5	5	3.1	4 · 3
PATIENT SAFETY		Test and Evaluation			
MEDICAL DEVICES					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the cue/alert to surface the failure mode that produces harm — silent hypoxemia, misplaced tube — at the point in the workflow where the operator can still recover, not in a post-hoc record.
- ▶ Pair the device standard with the institutional commitment (APSF, ASA) that makes the standard non-waiverable across the specialty, so adoption is a profession-level deliverable rather than a per-institution choice.
- ▶ Build the standard with edge-case acknowledgment: the device is the strongest available signal, not closed proof; pulse oximetry across skin tones (Case 26) is the canonical edge case the standard has to keep visible.

IN OPERATION / AFTER

- ▶ Track outcome metrics (mortality, malpractice premiums) alongside the standards, while attributing carefully — multifactorial declines are the norm, not the exception, in long-window safety interventions.
- ▶ Treat the malpractice-cost signal as institutional evidence the standard is recognized to have worked, not as the safety evidence itself; the safety evidence is mortality and harm.
- ▶ Maintain the edge-case surveillance the standard sets up: pulse oximetry across skin tones, capnography under specific physiology — each documented edge case is a place where the standard’s coverage is incomplete.

REFLECTION QUESTIONS

1. Identify an undetectable failure mode in your domain — one where the cue does not reach the operator in time to recover. What is the analog of pulse oximetry / capnography for that failure: the continuous signal that would convert it from undetectable to monitored?
2. Specify the institutional commitment (specialty foundation, regulatory standard) that would make the new signal non-waiverable. The device alone is not the deliverable; the standard is.
3. The anesthesia mortality decline is multifactorial. What is the minimum decomposition you would publish — training, device, pharmacology, team — to let downstream practitioners learn which components carried the most weight, rather than attributing the outcome to a single intervention?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

Interprofessional Education and the Evidence Gap

K Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Decades-long well-funded movement to educate health professionals together for collaborative care; Cochrane 2013 found only 15 studies between 1999 and 2011 met inclusion criteria; IOM 2015 made the gap the central finding — 'paucity of high-quality research' linking IPE to measurable changes in practice and patient outcomes

IN BRIEF

Interprofessional Education (IPE) is a decades-long, well-funded movement premised on the idea that doctors, nurses, pharmacists, and allied professionals should learn together so they can collaborate better in practice. The Cochrane review (Reeves et al., 2013) found only 15 studies between 1999 and 2011 met its inclusion criteria, and while those studies showed some positive outcomes, the evidence base for linking IPE to measurable changes in practice and patient outcomes was thin. The IOM report (2015) made the gap the central finding: there remains a "paucity of high-quality research" connecting IPE interventions to patient outcomes, and it proposed a conceptual model for doing the measurement properly. The case is the canonical instance in the corpus of a large, sincere, multidisciplinary translation effort whose core problem is that the field cannot yet measure whether the intervention works. It is the case-grounded basis for the enthusiasm-evidence gap as a curricular concept and motivates the Domain 3 sub-competency proposed in — the recurring pattern that a field instruments its enthusiasm faster than its outcomes.

THE SHIFT

The premise of Interprofessional Education is straightforward and credible: health-professions practice is collaborative, so health-professions training should be collaborative. Doctors, nurses, pharmacists, dietitians, social workers, and allied professionals should learn alongside each other, ideally in clinical and simulation contexts, so that the coordination capability is built during training rather than improvised on the wards. The movement is decades old, has institutional support across major accreditation bodies, and is well-funded by government and foundation sources.⁹

WHAT IS EMERGING

Reeves et al.'s 2013 Cochrane review applied the standard systematic-review machinery to the IPE outcome literature: what studies, between 1999 and 2011, met the inclusion criteria

for evaluating IPE interventions against collaborative-practice or patient outcomes? The answer was

15. Across that small set, the studies showed some positive outcomes — on clinician self-reported behavior, patient satisfaction, and a small number of clinical metrics — but the evidence base for linking IPE to measurable changes in practice and patient outcomes was thin. The reviewers' own conclusion was that the field had not produced the outcome evidence its scale of investment implied it should have.¹

THE CAPABILITY QUESTION

The IOM's 2015 report, *Measuring the Impact of Interprofessional Education on Collaborative Practice and Patient Outcomes*, made the evidence gap the central finding. There remains, the IOM concluded, a "paucity of high-quality research" connecting IPE interventions to patient outcomes. The report proposed a conceptual model for doing the measurement properly — a chain from IPE intervention through learner outcome, collaborative behavior, organizational practice, and patient outcome — and named the methodological work the field had to do to get from the current evidence state to a defensible causal claim. The report is itself the load-bearing artifact: a national-academy synthesis that the central problem of the field is the evidence architecture, not the intervention.²

EARLY EVIDENCE

What the case teaches at the LENS framing layer is the structural form of the enthusiasm-evidence gap at the multidisciplinary-translation scale. IPE is a sincere, well-funded, decades-long movement on which a great deal of curricular and operational investment has been made; the strongest synthesis of the outcome literature concludes that the field instrumented its enthusiasm faster than it instrumented outcomes. The case is the case-grounded basis for the Domain 3 sub-competency proposed in — naming the enthusiasm-evidence gap as a canonical anti-pattern — and for the CLO **Judgment under inadequate evidence**: practitioners and program designers in IPE have had to decide for decades on what to build, on evidence that does not establish the causal claim the field's premise rests on.³

OPEN PROBLEMS

In pair with the Colorado CTSA team-science training case (121) and implementation-science training (123), IPE is the frontier endpoint of the multidisciplinary-translation thread. Team science showed the measurement can be done at single-program scale with a validated instrument; IPE shows what happens when the field-scale evidence architecture has not been built. The pair-plus-trio teaches the Domain 3 sub-competency in both directions: collaboration as a unit of measurement is possible (Case 40), and the field-scale evidence is structurally absent (this case).

The field instrumented its enthusiasm faster than its outcomes. The strongest synthesis names the gap as the central finding, not the intervention as the failure.

EDITORS' SYNTHESIS OF REEVES ET AL. (2013) AND IOM (2015).

REFERENCES

1. Reeves, Perrier, Goldman, Freeth, & Zwarenstein (2013), "Inter-professional education: effects on professional practice and healthcare outcomes (update)," Cochrane Database of Systematic Reviews, doi:10.1002/14651858.CD002213.pub3.

2. Institute of Medicine (2015), Measuring the Impact of Inter-professional Education on Collaborative Practice and Patient Outcomes, National Academies Press, NCBI NBK338352.

3. WHO (2010), Framework for Action on Interprofessional Education and Collaborative Practice — the international policy backdrop.

4. v2 paired cases: Team-science training (121), Implementation-science training (123).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Interprofessional Education is the canonical enthusiasm-evidence-gap case in the corpus. A decades-long, well-funded movement; the strongest synthesis of the outcome literature concludes that the evidence base for the field’s central claim — that IPE produces measurable changes in practice and patient outcomes — is structurally thin. The IOM 2015 conceptual model is the proposed evidence architecture; the field’s task is to build to it.

LENS APPROACH

IPE is the field-scale enthusiasm-evidence-gap case (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) as the case-grounded basis for the enthusiasm-evidence-gap sub-competency and for the CLO **Judgment under inadequate evidence** — IPE is the pattern’s largest instance. Pair with Case 40 (team- science training, where measurement is possible at program scale) and Case 41 (imple- mentation-science training, where stated goals run ahead of operational practices).

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTH PROFESSIONS EDUCATION	K N	D1 D2 D3 D4 D5	5	2.1	4
INTERPROFESSIONAL COLLABORATION	→	→ Test and Evaluation			
PATIENT SAFETY					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the IPE intervention with an outcome-measurement chain in mind from the start: learner outcome → collaborative behavior → organizational practice → patient outcome. The IOM model is the published reference.
- ▶ Specify, before the intervention launches, which links in the chain the evaluation will measure and which it will leave as assumed; the field-scale gap is partly the result of every program leaving the same links unmeasured.
- ▶ Treat enthusiasm-as-evidence as a foreseeable failure mode in any field-scale capability movement; the IPE pattern recurs across other interdisciplinary translation efforts (see Case 41).

IN OPERATION / AFTER

- ▶ Commission field-level evidence architecture, not only program-level evaluation; the gap the IOM names is at the field level and the program-level studies cannot close it on their own.
- ▶ Publish negative or null findings as a contribution to the evidence base, not as a program failure; the field’s evidence gap is partly the result of selective publication on the positive side.
- ▶ Carry the IOM conceptual model into curriculum design conversations: a sub-competency that names the enthusiasm-evidence gap explicitly is the case-grounded curricular response to this pattern.

REFLECTION QUESTIONS

1. Identify a field-scale capability movement in your domain whose evidence architecture has not kept pace with its operational and curricular investment. What would the IOM-style conceptual model look like for that field's outcome chain?
2. Specify, for the next IPE-style program you would design or evaluate, which links in the outcome chain (learner / collaborative behavior / organizational practice / patient outcome) the evaluation will measure and which it will leave assumed. The field-scale gap is the accumulated result of leaving the same links assumed.

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

Bar-Code Medication Administration — Cue at the Point of Care

H Human-System Interface **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT A before-and-after study at an academic medical center associated bar-code electronic medication administration (bar-code eMAR) with a 41.4% reduction in nontiming administration errors and a 50.8% reduction in potential adverse drug events; transcription errors on order transcription were eliminated (6.1% to zero); a later single-site rollout reported a 55.4% reduction in actual patient-harm events

IN BRIEF

Wrong-drug and wrong-patient administration errors are a persistent failure mode in hospital pharmacy: the cue the human operator needs in order to catch the mismatch is structurally absent at the bedside, because the order, the dispensed medication, and the patient are connected only by paper documentation and clinical memory under time pressure. Bar-code medication administration (BCMA) supplies the cue in hardware: a bedside scan of the medication's bar code against the patient's wristband, gated by the electronic medication administration record. Poon et al. (NEJM, 2010) evaluated bar-code eMAR at an academic medical center using a before-and-after observational design and reported a 41.4% reduction in nontiming administration errors, a 50.8% reduction in potential adverse drug events, and the elimination of transcription errors on order transcription (6.1% to zero). A later single-site rollout (PMC6257885) reported a 55.4% reduction in actual patient-harm events. The study is explicit that the design is quasi-experimental — before-and-after / observational — not a randomized trial, and that significant workflow changes were required for the cue to land. The case is the canonical point-of-care cue/alert intervention, paired with Case 27 (anesthesia monitoring / APSF) at the cue-as-deliverable layer and with Case 23 (WHO Surgical Checklist) at the mandatory-mechanism layer.

BACKGROUND

The wrong-drug and wrong-patient administration error is a structural failure mode of the hospital medication chain. A physician's order, a pharmacist's dispense, a nurse's bedside administration, and a patient's wristband identification are connected only by paper documentation and the operator's clinical memory under time pressure. The cue that would let the bedside nurse catch a mismatch — "this medication's identity does not match this patient's record" — is not present in the workflow unaided. The conventional safeguards (the five rights, double-checks, naming protocols) are cognitive and procedural; the cue itself is structurally absent.^o

THE INTERVENTION

Bar-code medication administration supplies the cue in hardware. The medication unit-of-use carries a bar code, the patient wears a coded wristband, and the electronic medication administration record (eMAR) gates the administration on a bedside scan that matches medication-to-order-to-patient. The cue lands at the moment of administration, before the harm, and the system refuses to acknowledge the administration if any element does not match. The structural form is the same as the cue/alert interventions in Case 27 (continuous pulse oximetry, capnography under anesthesia): the cue at the right point in the workflow converts an undetectable failure into a monitored, recoverable one.¹

HOW IT WORKED

Poon et al. (New England Journal of Medicine, 2010) evaluated bar-code eMAR at a large academic medical center using a before-and-after observational design. The headline findings were a 41.4% reduction in nontiming administration errors (the failure mode the cue is designed to catch), a 50.8% reduction in potential adverse drug events (the harm the errors would have produced), and the elimination of transcription errors on order transcription (6.1% to zero) upstream of administration. A subsequent single-site rollout reported in PMC6257885 documented a 55.4% reduction in actual patient-harm events, supporting the transfer of the result across institutional contexts.²

THE EVIDENCE

The study's design is what the case's hedge rests on. The evaluation is quasi-experimental — before-and-after / observational — not a randomized trial, and the authors are explicit that significant workflow changes were required for the cue to land in operation. The cue is not the deliverable alone; the workflow is also the deliverable, and the attribution of the observed reduction to the cue rather than to the workflow change rests on plausibility and on the mechanism the cue is designed against. The 55.4% transfer number from the later rollout supports the mechanism, and that rollout figure also rests on observational data rather than randomized comparison.³

WHAT TRANSFERRED

What the case carries for the corpus is the cue-as-deliverable pattern at the point of care, with the explicit before-and-after-design hedge. Paired with Case 27 (anesthesia monitoring / APSF), bar-code eMAR shows the same structural form in a different specialty: the cue at the right point in the workflow converts an undetectable mismatch into a monitored one. Paired with Case 23 (WHO Surgical Checklist), the case sits at the mandatory-mechanism layer: the bar-code scan, like the checklist, is a workflow artifact the operator cannot route around, and the institutional commitment to that non-routability is part of why the cue lands.

The cue is supplied in hardware at the point of administration; the workflow change is part of the deliverable, and the design is observational rather than randomized.

EDITORS' SYNTHESIS OF POON ET AL. (2010) AND THE BCMA IMPLEMENTATION LITERATURE.

REFERENCES

1. Poon, E. G., Keohane, C. A., Yoon, C. S., Ditmore, M., Bane, A., Levitzion-Korach, O., Moniz, T., Rothschild, J. M., Kachalia, A. B.,
3. Thompson, K. M., Swanson, K. M., Cox, D. L., Kirchner, R. B., Russell, J. J., Wermers, R. A., Storlie, C. B., John-

Hayes, J., Churchill, W. W., Lipsitz, S., Whittemore, A. D., Bates, D. W., & Gandhi, T. K. (2010). Effect of bar-code technology on the safety of medication administration. *New England Journal of Medicine*, 362(18):1698–1707. doi:10.1056/NEJMsa0907115 — the case’s primary evaluation.

2. Bonkowski, J., Carnes, C., Melucci, J., Mirtallo, J., Prier, B., Reichert, E., Moffatt-Bruce, S., & Weber, R. J. (2013). Effect of barcode-assisted medication administration on emergency department medication errors. *Academic Emergency Medicine*, 20(8):801–806 — adjacent transfer evidence.

son, M. G., & Naessens, J. M. (2018), “Implementation of Bar-Code Medication Administration to Reduce Patient Harm,” *Mayo Clinic Proceedings: Innovations, Quality & Outcomes* 2(4):342–351, doi:10.1016/j.mayocpiqo.2018.09.001, PMID:30560236, PMCID:PMC6257885 — later single-site roll-out reporting a 55.4% reduction in actual patient-harm events.

4. Institute for Safe Medication Practices, Guidelines for Safe Electronic Communication of Medication Information — the institutional-commitment literature the workflow change rests on.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Bar-code medication administration is the canonical small-tier cue/alert intervention at the point of care. The cue lands at the bedside before the harm, the workflow change is part of the deliverable, and the headline results (41.4%, 50.8%, transcription errors eliminated, and a 55.4% transfer figure) rest on a before-and-after / observational design that the case preserves verbatim.

LENS APPROACH

BCMA is the small-tier point-of-care cue/alert intervention (induced 3.1; LENS D4/PT5) — Domain 4 for cue-as-deliverable; Domain 3 for the workflow-around-the- cue. Pair with Case 27 (anesthesia/APSf) and Case 23 (WHO Surgical Checklist).

DOMAIN	Modes Addressed	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
CLINICAL CARE	H K D	D1 D2 D3 D4 D5	5	3.1	4 · 3
PATIENT SAFETY		Test and Evaluation			
HEALTH IT					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Design the cue to land at the point in the workflow where the operator can still recover — for bar-code eMAR, the bedside administration scan, before the dose reaches the patient — not at a post-hoc reconciliation point.
- Treat the workflow change as part of the deliverable, not as friction to minimize; the cue’s effect is conditional on the workflow that surrounds it, and the evaluation has to acknowledge both.
- Pair the device standard with the institutional commitment that makes the scan non-routable around — the equivalent of the APSf / ASA monitoring standard in anesthesia, or the WHO checklist’s mandatory-mechanism status.

IN OPERATION / AFTER

- Report the headline reductions (41.4%, 50.8%, transcription errors eliminated, 55.4%) together with the design hedge — before-and-after / observational, significant workflow changes required — so the evidence the field cites is the evidence the field can actually use.
- Track the transfer figure (the 55.4% from the later single-site rollout) as evidence the mechanism survives a new institutional context, while preserving the same observational-design qualification on the new figure.
- Build the next round of evaluation against the residual error rate (the failures that survive the cue) so the institution learns where the workflow leaks remain, rather than treating the reduction as the closure of the failure mode.

REFLECTION QUESTIONS

1. Identify a point in your domain's workflow where the cue the operator needs in order to catch a mismatch is structurally absent. What would the hardware-supplied analog of bar-code eMAR look like — and at what point in the workflow does it have to land for the operator to still recover?
2. The headline results (41.4%, 50.8%, transcription errors eliminated, 55.4%) rest on a before-and-after observational design. What additional evidence — randomized comparison, segmented analysis by unit type, residual-error analysis — would you require before treating the figures as a closure rather than a strong signal?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · task analysis · design prototyping

Surgical Skill Video Peer-Rating Predicts Complications

H Human-System Interface **D** Designed Out **K** Knowledge & Institutional Memory

IMPACT Twenty bariatric surgeons each submitted a representative gastric-bypass video, rated blind by at least 10 peers; skill ratings ranged 2.6–4.8; the bottom skill quartile had a higher complication rate (14.5%) than the top quartile across a registry of 10,343 patients, and greater skill was associated with fewer reoperations, readmissions, and emergency department visits

IN BRIEF

Surgical complications are conventionally attributed to patient factors, hospital factors, and case-mix. Birkmeyer et al. (NEJM, 2013) asked a more direct question: can the surgeon's actual technical capability be measured well enough to predict patient outcomes? Twenty bariatric surgeons each submitted a representative video of a laparoscopic gastric bypass; the videos were rated blind by at least 10 peers on a structured skill scale. Skill ratings ranged from 2.6 to 4.8. Linked to a Michigan registry of 10,343 patients, the bottom skill quartile had a higher complication rate (14.5%) than the top quartile, and greater skill was associated with fewer reoperations, readmissions, and emergency department visits. The authors call the findings preliminary and name the skill-versus-volume confound explicitly: low-skill surgeons also did fewer cases and operated more slowly, so the contribution of skill itself versus the contribution of case volume is partly open. The hedge is part of the case. The proposed primary anchor is 2.1 (measuring the failure mode you care about) with C3 and C1 alternates; the editor may move it. Adjacent to JIGSAWS (Case 31) at the surgical-skill-measurement layer.

BACKGROUND

The conventional accounting for surgical complications attributes outcomes to patient factors (age, comorbidity, severity), hospital factors (volume, ICU support), and case-mix. The surgeon's actual technical capability — what the surgeon does with the instruments in this operation — is conventionally treated as too hard to measure at scale, and so as a hidden variable in the outcome equation. Birkmeyer et al. (New England Journal of Medicine, 2013) asked whether that variable could be brought into the open: can the surgeon's skill be measured well enough, at scale, to predict patient outcomes?^o

THE INTERVENTION

The study's design was deliberately simple. Twenty bariatric surgeons in a Michigan collaborative each submitted a representative video of a laparoscopic gastric bypass — a standardized, common procedure. The videos were rated blind by at least 10 peers on a

structured skill scale derived from the surgical-education literature: tissue handling, time and motion, instrument knowledge, exposure, completion of the operation. Skill ratings ranged from 2.6 to 4.8, with the inter-rater reliability adequate to support the rank-ordering the study then took into the outcomes analysis.¹

HOW IT WORKED

The outcomes evidence linked the peer-rated skill scores to a Michigan registry of 10,343 patients treated by the same twenty surgeons. The headline finding: the bottom skill quartile had a higher complication rate (14.5%) than the top quartile, and greater skill was associated with fewer reoperations, fewer readmissions, and fewer emergency department visits in the months after surgery. The skill score predicted the complication signal at the cohort scale. In the surgical-outcomes literature, this is the first large-registry evidence that operative skill, measured directly from operative video by peer rating, predicts the patient-outcome signal the institution actually cares about.²

THE EVIDENCE

The authors call the findings preliminary and name the load-bearing confound explicitly: low-skill surgeons also did fewer cases and operated more slowly. The skill-versus-volume confound is the central methodological hedge of the case. It is plausible — and consistent with the broader volume-outcome literature — that what the skill rating captured was partly the surgeon's accumulated operative experience, and that the rated skill is downstream of case volume rather than the other axis of the outcome. The published study does not separate the two; the editor may move the primary anchor away from 2.1 toward 1.1 (engineered vs. stated requirements) or 6.2 (operator-to-institution feedback channels) on that basis.³

WHAT TRANSFERRED

What the case carries for the corpus is the evidence-architecture move of measuring the operator's actual technical capability against the outcome the institution cares about, with the volume-confound hedge intact. The case is adjacent to JIGSAWS (Case 31) at the surgical-skill-measurement layer — JIGSAWS provides the controlled-task instrumented evidence and Birkmeyer the naturalistic operative-video evidence, and the two together anchor the small-tier C3 conversation about measuring skill in surgery. The preliminary-findings language is part of what the case teaches.

Greater skill was associated with fewer reoperations, readmissions, and emergency department visits; the low-skill surgeons also did fewer cases and operated more slowly.

EDITORS' SYNTHESIS OF BIRKMEYER ET AL. (2013) AND THE ACCOMPANYING EDITORIAL COMMENT.

REFERENCES

1. Birkmeyer, J. D., Finks, J. F., O'Reilly, A., Oerline, M., Carlin, A. M., Nunn, A. R., Dimick, J., Banerjee, M., & Birkmeyer, N. J. O., for the Michigan Bariatric Surgery Collaborative. (2013). Surgical skill and complication rates after bariatric surgery. *New England Journal of Medicine*, 369(15):1434–1442. doi:10.1056/NEJMsa1300625 — the case's primary evaluation.
2. Birkmeyer, J. D., Stukel, T. A., Siewers, A. E., Goodney, P. P., Wennberg, D. E., & Lucas, F. L. (2003). Surgeon volume and operative mortality in the United States. *New England Journal of Medicine*, 349(22):2117–2127 — the volume-outcome literature the confound rests against.
3. Vassiliou, M. C., et al. (2005). The MISTELS program — global objective assessment of laparoscopic skills. *Surgical Endoscopy* — the surgical-skill assessment literature the rating scale derives from.
4. Jacobs, D. O. (2013). Cut well, sew well, do well? *New England Journal of Medicine*, 369(15):1466–1467. doi:10.1056/NEJMe1309785 — the accompanying editorial.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Birkmeyer skill-outcomes study is the first large- registry evidence that operative skill, measured directly from operative video by blind peer rating, predicts the complication signal the institution cares about. The findings are preliminary; the skill-versus-volume confound is explicit; the multi-anchor (2.1 primary, 1.1 and 6.2 alternates) is the editor's call.

LENS APPROACH

The surgical-skill peer-rating study is the small-tier measure-the-failure-mode-you-care-about case (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the measure-the-capability-against-the-outcome design move, and in Domain 3 (Human-System Collaboration) for the peer-rating workflow that makes the assessment scalable. Adjacent to JIGSAWS (Case 31) at the surgical-skill-measurement layer — JIGSAWS the controlled-task instrumented evidence, Birkmeyer the naturalistic operative-video evidence. The skill-versus-volume confound is the case.

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
SURGERY CLINICAL OUTCOMES ASSESSMENT	H D K	D1 D2 D3 D4 D5 Test and Evaluation	5	2.1	4 · 3

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the capability the assessment is built to measure (operative skill on a standard procedure) and the outcome the institution actually cares about (complication rate, reoperation, readmission), so the evidence architecture connects the measured variable to the institutional signal.
- Use blind peer rating on standardized video to make the assessment scalable and resistant to the gaming patterns of self-report or volume-weighted reputation.
- Pre-register the threats-to-validity table — volume confound, selection of submitted video, inter-rater reliability — so the published headline is reported alongside the qualifications the design carries.

IN OPERATION / AFTER

- Report the 14.5% bottom-quartile complication finding together with the volume-confound hedge; the preliminary-findings language is what makes the result interpretable for the next study.
- Design the follow-up analysis to separate skill from volume — for example, by matching on case volume within skill quartile, or by instrumenting the case-volume time-course — so the next round of evidence resolves the confound the present study leaves open.
- Carry the result into the surgical-education and credentialing conversation as evidence that peer-rated operative-video skill is a measurable, outcome-predictive variable, while preserving the qualifications the authors themselves attach.

REFLECTION QUESTIONS

- Identify a capability in your domain where the operator's actual technical capability is conventionally treated as too hard to measure at scale and therefore left as a hidden variable in the outcome equation. What would a Birkmeyer-style scalable, blind peer-rated measure look like?
- The skill-versus-volume confound is the case. Design the follow-up analysis that would separate skill from volume in your context. What matching, what instrumentation of the volume time-course, what controlled-task companion would be required?
- Birkmeyer is anchored at 2.1 with C3 and C1 alternates. Which anchor would you choose for your own use, and what does that choice say about which conversation in your domain the case is the load-bearing evidence for?

WHO BUILDS THIS human factors & interaction design · systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · task analysis · design prototyping · knowledge bases · data pipelines

Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units

T Training Gap **K** Knowledge & Institutional Memory **H** Human-System Interface

DISCLOSURE Institutional overlap: this case originates at an editor's institution (Johns Hopkins); no editor was personally involved. Anchored to the peer-reviewed PLOS One paper and the public JIGSAWS dataset, not institutional press.

IMPACT JHU's Language of Surgery project treated surgical motion as language — decomposing tasks into gesture and sub-gesture motion primitives — and released JIGSAWS, a public da Vinci kinematic/video/gesture/skill-rating dataset that became a standard benchmark; experts used fewer gestures (26.29 vs 31.30) and fewer gesture errors than novices for a knot-tying task

IN BRIEF

The Language of Surgery project at Johns Hopkins, led by Gregory Hager and a cross-departmental team of roughly twenty investigators across engineering, computer science, and surgery, treated surgical motion as language — decomposing tasks into gestures (the “surgemes”) and sub-gesture motion primitives (the “dexemes”) fine enough to distinguish expert from novice. The project released JIGSAWS (JHU-ISI Gesture and Skill Assessment Working Set), a publicly available da Vinci surgical-robot dataset with synchronized kinematic, video, gesture-annotation, and skill-rating tracks for suturing, knot-tying, and needle-passing tasks. JIGSAWS became a standard benchmark in surgical skill-assessment and gesture-recognition research. Vedula et al. (PLOS One, 2016) used the dataset to show that experts used fewer gestures (26.29 vs. 31.30 on a knot-tying task) and made fewer gesture errors than novices, with quantifiable sub-task-level differences. The case establishes that surgical skill is decomposable and machine-measurable; the honest open question — preserved here verbatim — is whether automated motion-level feedback accelerates trainee skill acquisition or improves patient outcomes. The dataset enables the question more than it answers it. The case pairs directly with Case 30 (Birkmeyer's video-rated bariatric-surgical-skill outcome study), which establishes that skill matters; together they form the skill-measurement spine the corpus needed.

BACKGROUND

Surgical training has carried, for decades, a structural measurement gap: skill at the trainee-versus-attending level is universally acknowledged to matter, and is universally measured by rater judgment — case logs, milestone evaluations, OSATS scores — that resolves the construct only at the summary level. What such measures cannot resolve is the unit of skill: the specific motion sequence, the gesture choice, the sub-gesture smoothness, the corrective re-engagement after a near-error. The Language of Surgery project began from

the position that skill could be made measurable at the gesture level if surgical activity were decomposed the way language is decomposed into phonemes and words.⁰

THE INTERVENTION

The intervention is dataset-and-method, jointly. The team instrumented the da Vinci surgical robot to capture synchronized kinematic traces of the arms, stereo video of the surgical field, and frame-level human annotation of which gesture was being performed. Onto that base they layered expert skill ratings task by task. The technical decomposition was structured as a hierarchy: a task (e.g., knot-tying) consists of a sequence of gestures or surgemes (e.g., “reach for needle”, “position needle”, “drive needle”), each of which is composed of motion primitives or dexemes at the kinematic-trajectory level. The team released the resulting corpus as JIGSAWS in open form so the broader research community could test methods against a common benchmark, rather than each lab building its own private dataset.¹

HOW IT WORKED

Vedula et al. (PLOS One, 2016) used the dataset to analyze the structure of surgical activity for a suturing and knot-tying task and reported that experts used fewer gestures (26.29 on a knot-tying task) than novices (31.30), and that novices made more gesture errors. The differences were not at the gross-task-outcome level — both groups completed the task — but at the gesture-composition level the decomposition made visible. Subsequent work on JIGSAWS by other groups developed automated gesture-recognition methods, automated skill-rating classifiers, and motion-primitive analyses; the dataset’s role was to make these results comparable across labs.²

THE EVIDENCE

The honest open question survives into the case verbatim. The project demonstrates that surgical skill is decomposable and machine-measurable at the gesture level. What remains open is whether automated motion-level feedback delivered to trainees actually accelerates skill acquisition, and whether gesture-level skill differences translate into patient-outcome differences for the trainees once they reach the operating room. JIGSAWS enables both questions to be asked rigorously; it does not, by itself, answer either. The case’s learning-engineering content is the construction of the evidence architecture — the measurable unit, the open dataset, the cross-lab benchmark — that makes downstream measurement possible. The home-institution disclosure under the title is the standing safeguard against boosterism.³

WHAT TRANSFERRED

In pair with Case 30 (Birkmeyer et al.’s video-rated bariatric surgical-skill study), the case completes a skill-measurement pair the corpus had needed: Birkmeyer shows that rated skill predicts patient outcome at scale (skill matters), and this case shows that skill is decomposable into machine-measurable units (skill is engineerable). Drafted together they operationalize the new framing the v2 induced framework proposes — that capability is engineerable when the unit of capability is named and the measurement instrument follows. JIGSAWS’s continued use as a benchmark, more than a decade after release, is the field-scale evidence that the decomposition was the right resolution for the question.

Surgical motion is language. The unit of skill is the gesture; the corpus is the dataset; the benchmark is the field’s continued use of it more than a decade after release.

EDITORS’ SYNTHESIS OF VEDULA ET AL. (2016) AND THE JIGSAWS RELEASE.

REFERENCES

1. Vedula, Malpani, Tao, Chen, Gao, Poddar, Ahmidi, Paxton, Vidal, Khudanpur, Hager, & Chen (2016), “Analysis of the Structure of Surgical Activity for a Suturing and Knot-Tying Task,” PLOS One 11(3):e0149174, doi:10.1371/journal.pone.0149174.

2. Gao, Vedula, Reiley, Ahmidi, Varadarajan, Lin, Tao, Zappella, Bejar, Yuh, Chen, Vidal, Khudanpur, & Hager (2014), “JHU-ISI Gesture and Skill Assessment Working Set (JIGSAWS): A Surgical Activity Dataset for Human Motion Modeling,” MICCAI Workshop — JIGSAWS dataset release paper.

3. Reiley, Lin, Yuh, & Hager (2011), “Review of methods for objective surgical skill evaluation,” Surgical Endoscopy 25(2):356–366 — situates the decomposition within the broader skill-assessment literature.

4. Birkmeyer, Finks, O’Reilly, et al. (2013), NEJM — the paired skill-matters study (Case 30).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Language of Surgery / JIGSAWS is the corpus’s worked example of decomposing a tacit capability — surgical skill — into machine-measurable units, and releasing the measurement infrastructure openly so the field can build on it. The decomposition is established; the downstream question — whether motion-level feedback accelerates skill or improves patient outcomes — is open. The case enables the question rather than answering it.

LENS APPROACH

Language of Surgery is the skill-decomposition case (induced 2.1; LENS D4/PT5) — Domain 4 for construct-resolution; Domain 1 for task-decomposition. Pair with Case 30 (Birkmeyer) and Case 27. Home-institution disclosure under the title.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
SURGERY	T K H	D1 D2 D3 D4 D5	5	2.1	2 • 4
SKILL ASSESSMENT	→	→	Test and Evaluation		
HUMAN-MOTION ANALYSIS					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Decompose the capability into units fine enough to distinguish expert from novice but coarse enough to be reliably annotated; the surgeme / dexteme hierarchy is the worked example of the trade-off.
- ▶ Release the dataset openly with synchronized signal tracks (kinematic, video, annotation, rating) so that downstream methods can be compared on a common benchmark rather than each lab evaluating itself.
- ▶ Pair the dataset with a baseline analysis paper that demonstrates the decomposition resolves real expert/novice differences; the Vedula et al. analysis is what makes JIGSAWS more than a data dump.

IN OPERATION / AFTER

- ▶ Carry the open question explicitly — whether automated motion-level feedback accelerates skill acquisition is not what the dataset establishes — so downstream researchers and trainees see the gap as a research target, not a settled claim.
- ▶ Preserve the home-institution affiliation disclosure in any communication; the standing language anchors the case to the peer-reviewed paper and the public dataset, not to institutional press.
- ▶ Treat the cross-lab benchmark adoption as the case’s strongest evidence that the decomposition was correct at the resolution chosen; the field’s continued use of JIGSAWS more than a decade after release is itself the test.

REFLECTION QUESTIONS

1. Identify a tacit capability in your domain that is currently rated at the summary level. What is the unit of the capability — the gesture-equivalent — at which the decomposition would resolve expert/novice differences, and what signal tracks would the dataset need to synchronize?
2. Specify the open question your decomposition would not answer on its own. JIGSAWS does not establish that motion-level feedback accelerates skill; it enables the question. What is the analog in your context — the question the dataset enables but does not close?
3. The case's evidence of correctness is field-scale adoption of the benchmark a decade after release. What is the publication, release, and open-license strategy that would let your decomposition be tested by labs that have no stake in the original design?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · usability testing · cognitive walkthroughs

Annual-Screening UI Redesign + CDS at University of Missouri Health Care

T Training Gap **D** Designed Out **N** Normalization of Deviance

IMPACT A multidisciplinary EHR redesign of ambulatory annual-screening prompts (advance directives, depression, falls risk, alcohol/drug misuse), paired with embedded CDS, reported improvements in task time, error rates, System Usability Scale scores, and the downstream screening-rate outcomes the project was scoped to move

IN BRIEF

A multidisciplinary team at University of Missouri Health Care redesigned the EHR interface clinicians use to prompt and perform annual screening — advance directives, depression, falls risk, alcohol and drug misuse — and embedded clinical decision support inside the redesigned workflow. The team reported gains on the usability metrics (task time, error rate, System Usability Scale) and on the downstream process outcome the project was scoped to move: the actual rate at which guideline-recommended screening was completed. It is a small-tier intervention case for cue-and-alert design as a capability deliverable, with both human-factors and clinical-process outcomes in the same report. The corpus has long needed a small-tier C3 positive example to set against the interface failures already documented at the big tier (Therac-25, CPOE/EHR adoption, the Helios pattern). The evidence base is a single-institution QI study published peer-reviewed in *Applied Clinical Informatics* (2020), with a HIMSS chapter case-study write-up. Future validation will continue as the downstream clinical-outcome literature on screening-rate gains matures.

BACKGROUND

Ambulatory annual screening — advance directives, depression, falls risk, alcohol and drug misuse — is the kind of guideline-recommended care that is easy to declare and hard to land. The cue lives in the EHR; the action lives in a time-pressured encounter; and the gap between prompt and completion is where most screening programs lose their numbers. The University of Missouri Health Care project framed the problem squarely as cue-and-alert design: if the prompt cannot be acted on inside the workflow without friction, the screening will not happen.⁹

THE INTERVENTION

The redesign was multidisciplinary by construction — clinicians, informaticists, and usability specialists working against the existing screening interface. The team rebuilt the prompt presentation, added embedded clinical decision support that surfaced the next action at

the point of decision, and tightened the path between recognizing a positive screen and placing the appropriate order. The design move is the one the induced framework flags as the C3 deliverable: change the interface so that the desired action is the path of least resistance, not a separate sub-task layered on top of the visit.¹

HOW IT WORKED

The reported outcomes cross two layers. At the usability layer the team reported reductions in task time and error rate and a gain on the System Usability Scale. At the process layer they reported an increase in the screening-rate metric the project was scoped to move — the clinical process the cue exists to drive. That second layer is what makes the case a C3 small-tier intervention rather than a usability study: the interface gain translated into the downstream behavior, at least over the reported observation window.²

THE EVIDENCE

The case is a single-institution quality-improvement study, published peer-reviewed in Applied Clinical Informatics (2020), with a HIMSS chapter case-study write-up. The magnitudes reported sit inside a single-institution quality-improvement project with its own outcome metric; replication at other institutions and durability across EHR upgrades and personnel rotation have not been independently audited. Future validation will continue as the institution and others publish follow-on screening-rate data and as the relationship between screening completion and downstream patient outcomes is tracked.³

WHAT TRANSFERRED

What the case teaches at the LENS layer is that the C3 interface-failure pattern documented at the big tier is not a cosmic constraint — it is redressable by deliberate design at the small tier, when usability specialists, clinicians, and CDS authors are seated together. The capability deliverable is the redesigned cue plus the embedded decision support plus the measurement of the downstream process the cue exists to drive. The case is the missing positive example for induced 3.1 at the small tier and a paired teaching companion for the failures already in the corpus.

The C3 interface failure mode is not a cosmic constraint. It is redressable by deliberate design — when usability, clinicians, and CDS authors sit on the same redesign.

EDITORS' SYNTHESIS OF THE UMMC HIMSS CASE STUDY.

REFERENCES

1. Pierce RP, Eskridge BR, Rehard L, Ross B, Day MA, Belden JL (2020), "The Effect of Electronic Health Record Usability Redesign on Annual Screening Rates in an Ambulatory Setting," Applied Clinical Informatics 11(4):580–588, doi:10.1055/s-0040-1715828. HIMSS (Greater Kansas City chapter) case study is the secondary write-up.
2. Co et al. (2019), "Medication safety alert fatigue may be reduced via interaction design and clinical role tailoring: a systematic review," JAMIA 26(10):1141, doi:10.1093/jamia/ocz095 — adjacent systematic-review evidence on interaction-design redesign.
3. Office of the National Coordinator for Health IT, SAFER Guides on CDS design — practitioner-tier guidance the redesign instantiates.
4. Middleton et al. (2013), "Enhancing patient safety and quality of care by improving the usability of electronic health record systems," JAMIA 20(e1):e2–e8 — the framing peer-reviewed paper on EHR-usability-as-safety.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The UMHC redesign is the small-tier C3 positive example the corpus needed: cue-and-alert design as a capability deliverable, with both usability and downstream screening-rate gains in the same project. The evidence is a single-institution QI study, peer-reviewed in Applied Clinical Informatics (2020), with a HIMSS case-study write-up; magnitudes await independent replication and durability tracking. Future validation ongoing.

LENS APPROACH

UMHC is the C3 small-tier intervention case (induced 3.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the CLO-4 deliverable that cue redesign must show its downstream process effect, and in Domain 1 (Systems Analysis) for CLO-1 — the multidisciplinary team did the analysis of the screening workflow as the precondition for the redesign. The case is the paired positive example for the big-tier C3 failures (Therac-25, CPOE/EHR adoption, Helios) the corpus already documents.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	T D N	D1 D2 D3 D4 D5	5	3.1	4 · 1
EHR USABILITY	→	→			
CLINICAL DECISION SUPPORT		Test and Evaluation			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Frame the screening problem as cue-and-alert design from the start; do not separate usability from the clinical-process outcome the cue exists to drive.
- ▶ Seat usability specialists, clinicians, and CDS authors together on the redesign team; the C3 deliverable is the integrated artifact, not a hand-off.
- ▶ Specify the downstream process metric (screening-rate completion) before the redesign ships, so the usability-layer and process-layer outcomes are measured against the same scope.

IN OPERATION / AFTER

- ▶ Treat the QI report honestly: a single-institution peer-reviewed study is a small-tier intervention, not a settled magnitude. Carry the single-institution caveat into any downstream citation.
- ▶ Track durability — across EHR upgrades, personnel rotation, and downstream clinical outcomes — as a separate post-deployment commitment, not a footnote to the implementation report.
- ▶ Pair with the big-tier C3 failures already in the corpus when teaching; the failure-and-intervention pair is the teaching artifact, not either alone.

REFLECTION QUESTIONS

1. Identify a screening or recommended-care prompt in your EHR. What proportion of prompts result in the completed action? Where does the cue-action gap sit, and what is the redesign that would close it?
2. Specify the downstream process metric you would commit to before any usability redesign ships. The UMHC case is teachable because it reported the screening-rate outcome, not only the SUS gain. What would be the equivalent in your context?
3. The case is a single-institution peer-reviewed QI study. What is the minimum independent replication evidence you would require before treating the reported magnitudes as a basis for an institutional investment in your own setting?

WHO BUILDS THIS learning science & instructional design · systems & human-factors engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · task analysis · design prototyping · incident analysis · safety dashboards

CASE 33

HEALTHCARE

CLINICAL DECISION SUPPORT

EHR DESIGN

NURSING WORKFLOW

2019 – 2024

Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal

T Training Gap **D** Designed Out **N** Normalization of Deviance

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT Structured EHR alert redesign — fewer alerts, severity reclassification, interruptive-to-passive conversion, role-tailoring — reduced alert burden in published systematic-review and quality-improvement evidence; the 2024 case studies report alert-rate reduction with the underlying safety signal preserved

IN BRIEF

Alert fatigue is the structural failure mode the C3 thread names at the small tier: an EHR that fires so many alerts the actionable ones are lost in the noise. The 2019 JAMIA systematic review by Hussain and colleagues aggregates the evidence that structured redesign — interaction design changes and clinical-role tailoring — can reduce alert burden; most optimization studies in the review reported alert-rate reduction after intervention. Two 2024 quality-improvement publications extend the pattern with named maneuvers: a nursing-workflow redesign of four high-firing, low-action alerts using quantitative alert-firing analysis, empathy mapping, and iterative user feedback; and the replacement of an interruptive COVID-precautions alert with passive clinical decision support, targeting both alert burden and the timeliness of precautions orders. The case is the small-tier intervention companion to a C3.2 failure thread that v1 left almost entirely populated by failures (Uber ATG, Robodebt, Northeast Blackout, UK Post Office, Tesla Autopilot). The evidence tier is mixed: the systematic review is peer-reviewed; the per-site QI projects are practice-tier publications. Future validation will continue on whether the redesigns survive EHR upgrades and personnel rotation.

BACKGROUND

Monitoring under alert burden is the unsupportable role the C3.2 induced sub-competency names — the operator asked to keep attention on a stream of low-signal alerts and to spot the consequential one in real time. In the EHR setting the burden becomes structural: high-firing, low-action alerts train the clinician to dismiss alerts as the default, which is where the actionable alert is missed. Alert fatigue is the failure mode; the question for C3 is whether deliberate redesign can reduce the burden without cutting the safety signal.^o

THE INTERVENTION

The 2019 JAMIA systematic review by Hussain and colleagues aggregates the published evidence on EHR alert optimization. The headline finding: interaction–design changes and clinical–role tailoring reduce alert burden, and most evaluated optimization studies in the review reported alert–rate reduction post–intervention. The mechanism is not a single intervention but a family of moves — severity reclassification, conversion of interruptive alerts to passive decision support, role–based tailoring so the alert reaches the clinician who can act on it, and removal of alerts whose firing–to–action ratio shows the alert is no longer working.¹

HOW IT WORKED

The 2024 quality–improvement publications instantiate the moves. One project redesigned four high–firing, low–action alerts in the nursing workflow using mixed methods — quantitative analysis of firing data, empathy mapping of the nursing experience, and iterative user feedback as the redesign was refined. A second project replaced an interruptive COVID–precautions alert with passive CDS, with dual outcomes: reduce alert burden and improve the timeliness of precautions orders. Together the projects show the redesign–pattern is operable at the per–alert level and that the evaluation can report both halves of the trade–off the redesign exists to manage.²

THE EVIDENCE

The evidence–tier flag matters. The 2019 systematic review is peer–reviewed, and the 2024 Applied Clinical Informatics paper sits in a peer–reviewed informatics journal. The practical maneuver of treating the per–site QI projects as generalizable, though, rests on practice–synthesis logic across the body of work rather than on a single multi–site randomized evaluation. Magnitudes vary by site, EHR vendor, and alert category; the redesigns must be re–verified after EHR upgrades and personnel rotation. Future validation will continue as the optimization–study literature consolidates.³

WHAT TRANSFERRED

The teaching point pairs with Case 32 and with the v1 C3.2 failure thread. C3.2 in v1 is entirely failures; this case is the small–tier intervention that demonstrates the failure mode is redressable by design. The capability deliverable is the redesigned alert architecture itself — severity tiers, interruptive–vs–passive decisions, role–tailoring rules, and an ongoing measurement loop on firing–to–action ratios — not a one–time clean–up. The new CLO around delegation with revocation applies here: when automated decision support oversight is delegated to the bedside, the redesign discipline is part of the delegation.⁴

Alert fatigue is the failure mode the alert architecture trains. The redesign discipline is part of the delegation.

EDITORS' SYNTHESIS OF HUSSAIN ET AL. (2019) AND THE 2024 QI LITERATURE.

REFERENCES

1. Hussain et al. (2019), "Medication safety alert fatigue may be reduced via interaction design and clinical role tailoring: a systematic review," JAMIA 26(10):1141–1149, doi:10.1093/jamia/ocz095.
2. Patterson E. (2024), "Navigating Alert Fatigue: A Case Study in Electronic Health Record Alert Design Optimization," Studies
4. Office of the National Coordinator for Health IT, SAFER Guides on CDS — practitioner–tier guidance the redesigns instantiate.
5. Ancker et al. (2017), "Effects of workload, work complexity, and repeated alerts on alert fatigue in a clinical decision support system," BMC Medical Informatics and Decision Making 17:36 — adjacent measurement evidence.

- in Health Technology and Informatics, PubMed 39049299 — nursing-workflow QI redesign of four high-firing alerts.
3. Authors (2024), “Addressing Alert Fatigue by Replacing a Burdensome Interruptive Alert with Passive Clinical Decision Support,” Applied Clinical Informatics — interruptive-to-passive conversion with dual outcomes.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The 2019 JAMIA review plus the 2024 QI projects are the small-tier C3.2 intervention companion the corpus needed — the failure thread (Uber ATG, Robodebt, UK Post Office, Tesla) is redressable by deliberate alert redesign. The systematic review is peer-reviewed; the per-site QI publications are practice-tier; magnitudes and durability open. Future validation ongoing.

LENS APPROACH

Alert-fatigue redesign is the C3.2 small-tier intervention case (induced 3.1 and 3.2; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for CLO-4 — the redesign must report the safety signal alongside the alert-rate cut — and in Domain 3 (Human-System Collaboration) for CLO-3 oversight of automated decision support delegated to the bedside, with the CLO on delegation with revocation explicit. Pair with Case 32 for the small-tier C3 thread.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	T D N	D1 D2 D3 D4 D5	5	3.1	4 · 3
CLINICAL DECISION SUPPORT	→	→ Test and Evaluation			
EHR DESIGN					
NURSING WORKFLOW					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Instrument every consequential alert with its firing-to-action ratio; the metric that names the failure mode must be reported alongside the redesign.
- ▶ Decide the interruptive-vs-passive call deliberately for each alert; the conversion is the C3 move that reduces burden when the signal does not require the interruption to be actionable.
- ▶ Use mixed methods (firing data + clinician experience + iterative feedback) to design the per-alert redesign; quantitative alone misses why the alert is being dismissed.

IN OPERATION / AFTER

- ▶ Track the safety signal the alert exists to protect as a separate post-redesign outcome; an alert-rate cut that lost the signal is a failure, not a win.
- ▶ Re-audit the redesigned alerts after EHR upgrades and personnel rotation; the redesign is not a one-time clean-up and the v1 thread shows the failure mode returns.
- ▶ Treat the per-site QI publications as practice-synthesis evidence when generalizing; the systematic review supports the pattern, the magnitudes require local replication.

REFLECTION QUESTIONS

1. Identify an EHR alert in your context with a high firing-to-action ratio. Which of the redesign moves (severity reclassification, interruptive-to-passive conversion, role-tailoring, removal) would you apply, and what would the measured outcome be on both alert burden and the safety signal?
2. Specify the mixed-methods design (firing data + clinician experience + iterative feedback) you would use to make a per-alert redesign decision. What does the empathy-mapping step add beyond quantitative firing analysis alone?
3. The per-site QI publications are practice-synthesis tier. What is the minimum cross-site or randomized evidence you would require before generalizing the magnitudes from the 2024 reports to your own setting?

WHO BUILDS THIS learning science & instructional design · systems & human-factors engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · task analysis · design prototyping · incident analysis · safety dashboards

COMPOSER Sepsis Prediction – The Third Clinical-AI Sepsis Case

T Training Gap **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT Boussina et al. NPJ Digital Medicine 2024 prospective multi-site implementation of the COMPOSER (CONformal Multidimensional Prediction Of SEpsis Risk) deep-learning model at UC San Diego Health; reported a 1.9 percentage-point absolute decrease (17% relative) in in-hospital sepsis mortality and a 5.0 percentage-point absolute increase in sepsis bundle compliance during the deployment period, evaluated against the pre-deployment baseline period within the same two-emergency-department health system

IN BRIEF

Aaron Boussina, Shamim Nemati, and colleagues at UC San Diego Health published the COMPOSER (CONformal Multidimensional Prediction Of SEpsis Risk) prospective implementation study in NPJ Digital Medicine in early 2024. The model is a deep-learning sepsis-risk prediction system that uses conformal prediction to handle uncertainty — alerting only when the model's calibrated confidence threshold is met and abstaining when the input is sufficiently out-of-distribution. The deployment evaluation reported a 1.9 percentage-point absolute decrease in in-hospital sepsis mortality and a 5.0 percentage-point absolute increase in sepsis bundle compliance during the deployment period at two UC San Diego emergency departments, evaluated against the pre-deployment baseline period within the same health system. The case completes the AI-delegation typology in sepsis at three deployments — Case 20 (TREWS at Johns Hopkins), Case 5 (Epic Sepsis Score across multiple health systems), and COMPOSER at UCSD. The honest hedges from the source are binding: the deployment is prospective but not RCT-grade, the mortality reduction is multifactorial (the COMPOSER deployment ran alongside other process improvements at UCSD), and the author team is explicit about the structural attribution question. Pair with Case 20, Case 5, and Case 35 (Radiology AI Miscalibration).

BACKGROUND

The contemporary clinical-AI sepsis-prediction literature has converged on a small number of well-evidenced deployments. Case 20 documents TREWS — the Targeted Real-Time Early Warning System at Johns Hopkins — with its prospective evaluation showing mortality reduction associated with prompt clinician response to alerts. Case 5 documents the Epic Sepsis Score deployed across multiple U.S. health systems, with external validation by Wong and colleagues finding substantially lower sensitivity and higher false-positive rates than the vendor's validation data implied. COMPOSER is the third deployment in the typology and the

second prospective–positive case. The deployment is at UC San Diego Health across two hospital sites, integrated into the electronic health record and into emergency–department and inpatient workflows.⁹

THE INTERVENTION

The methodological feature that distinguishes COMPOSER is the conformal–prediction framework. Conformal prediction is a statistical method for producing calibrated uncertainty intervals around individual predictions: rather than producing a single risk score, the model produces a prediction together with a measure of how confident the model is in that prediction given the input. The operational implication is that the model can abstain from alerting when the input is sufficiently out-of-distribution — when the model’s calibrated confidence does not support a clinically actionable alert. The abstention structure is what handles one of the load-bearing failure modes of Case 5 (Epic Sepsis), where the model alerted at high rates against patient populations it had not been adequately calibrated against. COMPOSER’s abstention structure is one mechanism for refusing to alert in those circumstances.¹

HOW IT WORKED

The deployment outcomes Boussina and colleagues report are the load-bearing intervention–evidence. Across the deployment period at the two UC San Diego hospital sites, in-hospital sepsis mortality decreased by 1.9 percentage points in absolute terms, and sepsis bundle compliance increased by 5.0 percentage points, both evaluated against the pre-deployment baseline period within the same health system. The evaluation is prospective implementation, not RCT — the comparison is against the same hospital sites’ historical baseline rather than against a contemporaneous randomized control arm — and the authors are explicit about the limitation. The deployment ran alongside other process improvements at UCSD’s sepsis–care pathway, and the authors are honest that the mortality reduction is multifactorial: COMPOSER’s contribution cannot be cleanly separated from the contribution of the surrounding process changes.²

THE EVIDENCE

The case pairs with Case 20 (TREWS) for the prospective– positive sepsis–prediction thread: both deployments report mortality benefits, both are integrated into specific health-system workflows, and the comparison between TREWS and COMPOSER turns on the alerting architecture and on the role of the surrounding process. Pair with Case 5 (Epic Sepsis Score) for the external-validation–revealed– gaps thread: COMPOSER’s abstention structure is one response to the calibration–gap failure mode that the Epic Sepsis evidence base named. Pair with Case 35 (Radiology AI Miscalibration) for the broader medical-AI miscalibration–in–deployment thread. The three sepsis cases together — Case 20, Case 5, COMPOSER — define the AI-delegation typology in sepsis: prospective–positive with workflow integration (TREWS), externally validated and found wanting against vendor claims (Epic Sepsis), and prospective–positive with conformal–prediction abstention (COMPOSER).³

WHAT TRANSFERRED

The hedges the case carries are load-bearing and preserved in the prose. The deployment is prospective implementation, not an RCT; the mortality reduction and bundle–compliance gain are evaluated against the same health system’s historical baseline, and the comparison

cannot rule out concurrent improvements in sepsis care that would have produced similar reductions in the absence of COMPOSER. The mortality reduction is multifactorial — COMPOSER was deployed alongside other process improvements at UCSD, and the authors are explicit that the deployment evaluation cannot cleanly separate COMPOSER's contribution from the contribution of the surrounding process. The conformal-prediction abstention structure is the methodological contribution the case anchors, and the cue-and-alert design CLO is anchored at the deployment seam where abstention is a clinical-workflow capability: the model that can refuse to alert when its calibrated confidence does not support an alert is the deployment artifact the case names as the load-bearing one.

The deployment is prospective implementation, not RCT; the mortality reduction is multifactorial — COMPOSER was deployed alongside other process improvements at UCSD, and the contribution cannot be cleanly separated; the conformal-prediction abstention is the methodological contribution the case anchors.

EDITORS' SYNTHESIS OF BOUSSINA ET AL. (2024) NPJ DIGITAL MEDICINE AND THE SURROUNDING SEPSIS-PREDICTION DEPLOYMENT LITERATURE.

REFERENCES

1. Boussina, A., Shashikumar, S. P., Malhotra, A., Owens, R. L., El-Kareh, R., Longhurst, C. A., Quintero, K., et al. (2024), "Impact of a deep learning sepsis prediction model on quality of care and survival," NPJ Digital Medicine 7:14, doi:10.1038/s41746-023-00986-6.
2. Shashikumar, S. P., Wardi, G., Malhotra, A., & Nemati, S. (2021), "Artificial intelligence sepsis prediction algorithm learns to say 'I don't know,'" NPJ Digital Medicine 4:134 — the methodological-foundation paper for the conformal-prediction abstention structure.
3. Wong, A., Otles, E., Donnelly, J. P., Krumm, A., McCullough, J., De-Troyer-Cooley, O., et al. (2021), "External Validation of a Widely Implemented Proprietary Sepsis Prediction Model in Hospitalized Patients," JAMA Internal Medicine 181(8):1065–1070 — the load-bearing external-validation paper on Epic Sepsis (Case 5).
4. Adams, R., Henry, K. E., Sridharan, A., Soleimani, H., Zhan, A., Rawat, N., Johnson, L., et al. (2022), "Prospective, multi-site study of patient outcomes after implementation of the TREWS machine learning-based early warning system for sepsis," Nature Medicine 28:1455–1460 — the prospective-positive TREWS deployment paper (Case 20).

THE LEARNING ENGINEERING LENS

LE INSIGHT

COMPOSER is the third clinical-AI sepsis case in the delegation typology — prospective-positive with conformal-prediction abstention. The deployment reduced in-hospital sepsis mortality by 1.9 percentage points and raised sepsis bundle compliance by 5.0 percentage points; the load-bearing hedges are that the evaluation is prospective not RCT and that the mortality reduction is multifactorial. The abstention structure is the methodological contribution the case anchors.

LENS APPROACH

COMPOSER is the cue-and-alert-design-as-capability-deliverable case at clinical-deployment scale (introduced 3.1; LENS D4+D3/PT6; CLO-4 and CLO-3). LENS uses it in Domain 4 (Test and Evaluation) for the prospective-implementation-evaluation discipline and in Domain 3 (Human-System Collaboration) for the abstention-as-clinical-workflow-capability anchor. Pair with Case 20 (TREWS), Case 5 (Epic Sepsis), and Case 35 (Radiology AI Miscalibration). The three sepsis cases together — Case 20, Case 5, Case 34 — define the AI-delegation typology in sepsis at the deployment-evidence level.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	T K D	D1 D2 D3 D4 D5	6	3.1	4 · 3
CLINICAL DECISION SUPPORT	→	→			
MACHINE LEARNING		Test and Evaluation + D3			

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Build the abstention mechanism as part of the deployment, not as a post-hoc filter; the conformal-prediction framework’s clinical value is that the model can refuse to alert when its calibrated confidence does not support an alert, and the abstention is a deployment capability rather than an analysis artifact. ▶ Pre-specify the multifactorial-attribution question before evaluation; the case demonstrates that prospective implementation alongside concurrent process improvement requires the evaluation to acknowledge the attribution limit, and pre-specifying the acknowledgement is the discipline. ▶ Integrate the model into specific workflow surfaces — ED admission, inpatient deterioration — rather than as a hospital-wide alerting layer; the deployment’s specificity is part of what makes the conformal-prediction abstention structure operationally meaningful.	▶ Carry the multifactorial-mortality hedge into print without softening; the case’s pedagogical value depends on the attribution question being explicit alongside the deployment-outcome reporting. ▶ Pair in syllabi with Case 20 (TREWS) and Case 5 (Epic Sepsis) so the AI-delegation typology in sepsis is taught as a three-case set at the deployment-evidence level. ▶ Use the case to anchor the cue-and-alert design CLO; the curricular target is the discipline of building abstention into the alerting architecture as a capability deliverable, and of evaluating the abstention against the calibration-gap failure mode the prior cases named.

REFLECTION QUESTIONS

1. Identify a clinical-AI deployment in your setting whose alerting architecture does not include an abstention mechanism. What is the calibration-gap failure mode the deployment would encounter, and what would the conformal-prediction abstention structure look like in your workflow?
2. Specify the multifactorial-attribution acknowledgement you would carry in a prospective-implementation evaluation. What concurrent process improvements would have to be enumerated, and what is the documented decision rule for what the deployment evaluation can and cannot establish?
3. The three sepsis cases — TREWS, Epic Sepsis, COMPOSER — together define an AI-delegation typology. Pick a clinical-AI category in your setting and ask: what is the analogous typology, and which deployments would have to be evidenced to define it at the same level of resolution?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

Radiology AI Miscalibration

H Human–System Interface **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT Recurring documented cases of FDA-cleared radiology AI tools performing worse in deployment than in validation, often along demographic lines; the canonical v1 anchor for clinical–AI deployment without surveillance, cross-referenced by the Epic Sepsis (Case 5) and pulse-oximetry (Case 26) deployment-evidence cases

IN BRIEF

FDA-cleared radiology AI tools — for chest X-ray classification, mammography, CT triage — have been repeatedly documented performing worse in deployment than in their validation studies, often with the degradation concentrated in patient groups under-represented in the training data. Larrazabal et al. (PNAS 2020) showed this structurally for chest-X-ray classifiers across gender imbalance; Seyyed-Kalantari et al. (Nature Medicine 2021) extended the finding to under-served racial and socioeconomic subgroups across three large public chest-X-ray datasets. Obermeyer et al. (Science 2019) showed that bias in the labels used to train clinical AI can under-allocate care to Black patients even when the model looks well-calibrated on its chosen proxy. Wachter and Brynjolfsson (JAMA 2023) raised the generative-AI extension. The FDA’s 510(k) clearance pathway — the route most cleared radiology AI tools have taken — does not routinely require demographic stratification of validation metrics or post-market monitoring of how a tool performs in the population using it; the De Novo pathway used for a small number of novel tools imposes more, but is rarely the chosen route. The 2025 FDA draft guidance on AI/ML-based Software as a Medical Device (SaMD), with its Predetermined Change Control Plan and Good Machine Learning Practice provisions, begins to address this gap; the institutional architecture for demographic post-market surveillance is still being built. The capability gap is in the regulatory architecture, not the model: clearance is not the same thing as clinically performable deployment. The case is the v1 anchor for the cross-references in the Epic Sepsis (Case 5) and pulse-oximetry (Case 26) deployment-evidence cases.

THE SHIFT

Machine-learning tools are moving rapidly into radiology and other diagnostic medicine, cleared for market and integrated into clinical workflows that affect real patients. The FDA has now authorized more than 1,000 AI/ML-enabled medical devices, the majority through the 510(k) clearance pathway as devices substantially equivalent to a predicate. Unlike a drug, a model can pass its validation study and still behave very differently once it meets a population that differs from its training data — the same model file that scored well on the clearance set can quietly carry a different error profile into every hospital whose patients do not resemble it. The 510(k) substantial-equivalence framing was built for an era of physical devices whose behavior was largely determined by their design; it was not built for a class of systems whose behavior depends on the distribution of the data they meet.^o

WHAT IS EMERGING

Multiple FDA-cleared radiology tools — chest-X-ray classifiers, mammography aids, CT triage systems — have been documented in the peer-reviewed literature performing worse in deployment than in validation, with the degradation often concentrated in under-represented patient groups. Larrazabal et al. (PNAS 2020) demonstrated structural sensitivity drops for groups under-represented in chest-X-ray training data, using the NIH ChestX-ray14 and CheXpert datasets to show that classifier sensitivity for a given group tracks that group's prevalence in the training set. Seyyed-Kalantari et al. (Nature Medicine 2021) extended the finding directly: across three large public chest-X-ray datasets, classifiers under-diagnosed Black, Hispanic, female, and lower-socioeconomic patients at higher rates, with the disparity present across model architectures — evidence that the shortfall is not a stray bug but a predictable consequence of which patients the training set did and did not contain.¹

THE CAPABILITY QUESTION

The problem is not confined to imaging. Obermeyer et al. (Science 2019) showed that a widely deployed care-management algorithm — one of a class of commercial risk-prediction tools applied to an estimated 200 million Americans a year — systematically under-allocated resources to Black patients because it was trained on healthcare cost as a proxy label for need: equally sick Black patients had lower historical costs because they had received less care, so the model rated them as lower-need. The question is how a regulator can certify a model as safe without checking how it behaves across the populations and labels it will actually meet — since a model can look well-calibrated on its chosen proxy while the proxy itself encodes the inequity it then propagates. Wachter and Brynjolfsson (JAMA 2023) raised the generative-AI extension of the same question: the proxy-and-population problem deepens under LLM-class tools whose outputs are harder to validate against any well-defined label at all.²

EARLY EVIDENCE

Similar deployment-performance degradation has been reported in mammography AI (validation sets often under-represent dense-breast tissue and Black women), sepsis-prediction tools (the Epic Sepsis Model — Case 5 — is the canonical example of an AI tool deployed at scale whose external validation found it substantially worse than its developer-reported numbers), pulse-oximetry behavior (the structural racial miscalibration documented in Sjoding et al. — Case 26 — sits under several downstream AI sepsis and triage tools), and skin-lesion classifiers (training-set under-representation of darker skin tones produces accuracy gaps in deployment). Yet the FDA's 510(k) clearance pathway does not routinely require demographic stratification of validation metrics, nor post-market monitoring of in-use performance — so the divergence between clearance and deployment is largely invisible while the tool is in use, and a shortfall concentrated in one patient group can persist unmeasured across the entire period the tool is influencing care.³

OPEN PROBLEMS

The FDA's 2025 draft guidance on AI/ML-Based Software as a Medical Device begins to address this gap. The proposed Predetermined Change Control Plan would let manufacturers pre-specify the model updates and validation procedures that can be applied without a new clearance; the companion Good Machine Learning Practice principles emphasize

representative training data, transparent performance metrics, and lifecycle monitoring. The De Novo pathway, used for a small number of novel tools, imposes more — but is rarely the chosen route. What the guidance does not yet require, as of its draft form, is mandatory demographic stratification of validation metrics at clearance and mandatory population-level post-market surveillance of in-use performance for the cleared model. Radiology AI thus remains the live, recurring case for the gap between regulatory clearance and clinically performable deployment. The capability gap is at the regulatory architecture, not the model: the institutional machinery to require demographic post-market surveillance has not yet been built. It is the medical-AI analog of the Vioxx post-market-surveillance failure (Case 9) at a new technological boundary — a case where the harm comes not from a hidden defect but from the absence of any standing system to watch the tool once it is in the population's hands.⁴



Performance metrics on a clearance dataset are not the same as performance metrics in deployment populations.

EDITORS' SYNTHESIS OF THE FDA AI/ML-BASED SaMD DISCUSSION PAPER (2019)

REFERENCES

1. Larrazabal, A. J., Nieto, N., Peterson, V., Milone, D. H., & Ferrante, E. (2020), "Gender imbalance in medical imaging datasets produces biased classifiers for computer-aided diagnosis," PNAS 117(23):12592–12594 — sensitivity drops for under-represented groups on NIH ChestX-ray14 and CheXpert.
2. Seyyed-Kalantari, L., Zhang, H., McDermott, M. B. A., Chen, I. Y., & Ghassemi, M. (2021), "Underdiagnosis bias of artificial intelligence algorithms applied to chest radiographs in under-served patient populations," Nature Medicine 27:2176–2182 — disparities across Black, Hispanic, female, and lower-socioeconomic subgroups; persistence across model architectures.
3. Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019), "Dissecting racial bias in an algorithm used to manage the health of populations," Science 366(6464):447–453 — proxy-label bias in a care-management algorithm; 200 million Americans/year are subject to tools of this class.
4. Wachter, R. M. & Brynjolfsson, E. (2023), "Will Generative Artificial Intelligence Deliver on Its Promise in Health Care?" JAMA 331(1):65–69 — generative-AI extension of the proxy-and-population problem.
5. FDA, "Proposed Regulatory Framework for Modifications to AI/ML-Based Software as a Medical Device" (2019); FDA draft guidance on AI/ML-Based SaMD (2025), with Predetermined Change Control Plan and Good Machine Learning Practice principles — the regulatory trajectory; absence of mandatory demographic stratification at clearance and post-market monitoring of in-use performance.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Radiology AI is the canonical contemporary case for the gap between regulatory clearance and clinical deployment performance in medical AI. The clearance dataset and the deployment population diverge. The institutional architecture to surface the divergence — demographic post-market surveillance — has not yet been built.

LENS APPROACH

Radiology AI is the canonical contemporary clinical-AI deployment-and-surveillance case (induced 7.2; LENS D4+D3/PT5). LENS uses it in Domain 4 (Test and Evaluation; CLO-4) for the clearance-vs-deployment measurement architecture and for demographic stratification of validation as a deliverable. LENS uses it in Domain 3 (Human-System Collaboration; CLO-3) for the human-AI deployment-without-surveillance pattern, and in Domain 5 (Navigating Sociotechnical Constraints) for the FDA AI/ML regulatory trajectory itself. The Obermeyer (2019) finding generalizes the diagnosis: bias enters through the labels and through the population, both of which the 510(k) process currently treats as outside its scope. Pair with Vioxx (Case 9) as the post-market-surveillance-failure pattern at a new technological boundary; cross-references the Epic Sepsis (Case 5) and pulse-oximetry (Case 26) cases at the clinical-AI-without-surveillance layer.

DOMAIN	MODES AT STAKE			LENS COMPETENCY					PROBLEM TYPE	INDUCED	CLO		
HEALTHCARE	→	H	K	D	→	D1	D2	D3	D4	D5	5	7.2	4 · 3
TECHNOLOGY	Test and Evaluation + D3												

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
- Require validation metrics stratified by demographic group at clearance, so a tool’s performance is established on the populations it will actually meet, not an aggregate. - Specify the intended deployment population and label definition explicitly, and treat divergence from the training distribution as a known risk to be measured, not assumed away. - Build the data pipeline for in-use performance capture before deployment, so post-market signals are collectable from the first patient rather than reconstructed after harm.	- Mandate demographic post-market surveillance of in-use performance, the institutional machinery the clearance pathway currently lacks. - Monitor for the clearance-to-deployment performance gap continuously, since a shortfall concentrated in one group can otherwise persist invisibly. - Tie continued authorization to demonstrated in-population performance, so a tool that degrades in deployment can be withdrawn before the divergence compounds.

REFLECTION QUESTIONS

1. Identify a model in your domain whose deployment population diverges from its training population. What is the institutional architecture to surface the divergence?
2. Design the demographic post-market surveillance deliverable that should accompany every FDA clearance of medical AI.
3. FDA 510(k) clearance does not currently require demographic stratification of validation metrics, nor does it require post-market monitoring of how a cleared tool actually performs on the population using it. What is the minimum reporting deliverable a regulator should require so the gap is visible while the tool is in use?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. TOOLS usability testing · cognitive walkthroughs · knowledge bases · data pipelines · task analysis · design prototyping

AlphaFold — Protein Structure Prediction

T Training Gap

IMPACT Substantially solved a 50-year protein-structure prediction problem; 200M+ structures publicly released; foundational positive AI case

IN BRIEF

DeepMind's AlphaFold (2020) and AlphaFold 2 (2021) substantially solved the half-century-old protein-structure prediction problem in computational biology. The publicly released AlphaFold Protein Structure Database now contains predicted structures for more than 200 million proteins — essentially the entire known protein universe — and has been integrated into structural-biology and drug-discovery workflows worldwide. AlphaFold is the strongest positive AI case in this book, and its lesson is in the conditions that made the benefit possible: a benchmark (CASP) the field had agreed on for decades, high-quality training data, an output biologists could verify against experimental structures, and an open release that let the global community adopt it fast. Each was a precondition for the success — and none of them was the model itself.

THE SHIFT

Predicting a protein's three-dimensional structure from its amino-acid sequence was one of biology's grand challenges for half a century — slow, expensive experimental work that bottlenecked drug discovery and basic research. Deep learning offered, for the first time, the prospect of solving it computationally at scale, turning a problem that had been a years-long experimental undertaking per protein into one that could be approached for the whole proteome at once.⁹

WHAT IS EMERGING

DeepMind's AlphaFold (2020) and AlphaFold 2 (2021) substantially solved the problem, predicting structures at accuracies rivaling experiment. The publicly released AlphaFold Protein Structure Database now holds predicted structures for more than 200 million proteins — close to the entire known protein universe — and has been folded into research workflows worldwide, so the benefit arrived not as a single laboratory's advantage but as a shared resource the wider community could build on immediately.¹

THE CAPABILITY QUESTION


What made AlphaFold a benefit rather than just a benchmark win? The case poses the question of which conditions allow an AI capability to be safely and widely useful — and the answer turns out to lie around the model, not in it: an agreed benchmark, trustworthy training data, a verifiable output, and a deliberate decision about release. Each of these is an institutional or evidentiary precondition, not an artifact of the architecture, which is why the case reads as a lesson about capability infrastructure rather than about a model.²

EARLY EVIDENCE

AlphaFold's success rested on four features, each a precondition rather than the model: the CASP benchmark the field had used for decades, high-quality experimental training data, output that biologists could check against known structures, and an open release that let the global community adopt the tool quickly. Where those conditions hold, AI amplifies capability; the technical model alone does not — and the same architecture dropped into a domain missing any one of those four would not have produced a comparable, trusted, widely adopted result.³

OPEN PROBLEMS

AlphaFold is the strongest positive AI case in the dataset for what supports beneficial deployment in a well-defined scientific domain. The open problem is generalization: most consequential problems lack an agreed benchmark, clean training data, or verifiable output. The frontier question is how much of the AlphaFold pattern can be reconstructed where those preconditions are not given for free — that is, whether a field can deliberately build the benchmark, the data, and the verification path that protein structure happened to have accumulated over decades.⁴



This will be one of the most important achievements in AI in the past decade.

PARAPHRASING JOHN MOULT (ORGANIZER OF THE CASP BENCHMARK) ON ALPHAFOLD2, 2020

REFERENCES

1. Jumper et al. (2021), "Highly accurate protein structure prediction with AlphaFold," Nature — the method and accuracy.
2. Varadi et al. (2024), "AlphaFold Protein Structure Database in 2024," Nucleic Acids Research — the 200M+ structures and open release.
3. Moulton, J. (CASP organizer) commentary on AlphaFold2 (2020) — the benchmark and the achievement (paraphrased).
4. CASP benchmark documentation — the decades-long agreed evaluation.
5. Hassabis (DeepMind) public commentary — the open-release governance decision.

THE LEARNING ENGINEERING LENS

LE INSIGHT

AlphaFold is the strongest positive AI case in the dataset. The technical achievement is real. The conditions that made the benefit possible — agreed benchmark, training data, verifiable output, open release — are the capability infrastructure around the model, not the model itself. The case is the strongest available evidence for what supports beneficial AI deployment.

LENS APPROACH

LENS uses AlphaFold in LEN 1 as a problem-framing case for what productive AI deployment looks like, in LEN 9 as a technical achievement, and in LEN 7 for the open-release governance decision that distributed the benefit globally.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	→ T	→ D1 D2 D3 D4 D5	2	2.1	4
TECHNOLOGY		Test and Evaluation + D5			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Establish an agreed benchmark and high-quality training data for the target problem before building the model, treating these as preconditions rather than afterthoughts.
- ▶ Engineer the output to be verifiable against an independent ground truth, so users can check predictions rather than having to trust them.
- ▶ Decide the release and access terms deliberately as a governance choice, since open release is what distributed AlphaFold’s benefit globally.

IN OPERATION / AFTER

- ▶ Monitor downstream use to confirm the verifiable-output property holds in practice, and that users are in fact checking predictions where stakes are high.
- ▶ Track whether the four preconditions still hold as the tool is applied to new protein families or adjacent problems beyond its validated domain.
- ▶ Sustain the open resource and benchmark over time, so the community-wide benefit does not erode as the field and the data move on.

REFLECTION QUESTIONS

1. Identify a domain in your work where the conditions that supported AlphaFold’s success (benchmark, training data, verifiable output) are present. What is the analogous opportunity?
2. The open release of AlphaFold’s predictions was a governance decision. Design the equivalent decision for an AI capability your institution might develop.
3. AlphaFold inherited a benchmark, clean data, and a verification path that protein structure had accumulated over decades. Pick a problem in your domain that lacks one of those preconditions, and lay out how a field would deliberately build it.

WHO BUILDS THIS learning science & instructional design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models

DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback

T Training Gap **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT McKinney et al. 2020 Nature paper reported a deep-learning mammography screening system outperforming radiologists on retrospective UK and U.S. screening datasets, with reductions in false-positives (5.7 percentage points in the U.S. set, 1.2 in the UK set) and false-negatives (9.4 and 2.7 percentage points respectively); Haibe-Kains et al. October 2020 Nature comment critiqued the paper for failing to release code and trained models, arguing that reproducibility could not be assessed and that screening-comparison claims required deployment-grade evidence

IN BRIEF

Scott Mayer McKinney and colleagues at Google Health and DeepMind published “International evaluation of an AI system for breast cancer screening” in Nature on January 1, 2020. The paper reported that a deep-learning system outperformed radiologists on retrospective UK and U.S. screening datasets, with reductions in false-positives of 5.7 percentage points (U.S.) and 1.2 percentage points (UK) and reductions in false-negatives of 9.4 and 2.7 percentage points respectively. The paper drew unusual press attention and rapidly entered the policy conversation on AI-assisted screening. Haibe-Kains and colleagues’ October 14, 2020 Nature comment titled “Transparency and reproducibility in artificial intelligence” critiqued the McKinney paper for failing to release code, trained models, or sufficient methodological detail to permit independent reproduction. The load-bearing hedge the Haibe-Kains comment delivers is that a large fraction of the methodology was not reproducible, and the screening-comparison framing the original paper offered has been refined by subsequent deployment evidence rather than confirmed at the deployment scale the headline implied. The case pairs with Case 35 (Radiology AI Miscalibration), Case 5 (Epic Sepsis), and Case 26 (Pulse oximetry).

THE SHIFT

The McKinney et al. paper was published on January 1, 2020, in Nature — a top-tier venue and an unusually high-profile publication for a deep-learning medical-imaging study. The work was a collaboration across Google Health, DeepMind, and clinical partners at Cancer Research UK Imperial Centre, Northwestern University, the Royal Surrey County Hospital, and the National Cancer Institute. The retrospective evaluation used UK and U.S. screening datasets and compared the AI system’s outputs against single-reader and double-reader radiologist performance. The headline framing was that the AI system reduced both false-

positives and false-negatives relative to radiologists, with the U.S. dataset showing larger absolute reductions than the UK dataset.⁰

WHAT IS EMERGING

The press response was substantial. Mainstream coverage framed the result as “AI outperforms radiologists at breast cancer screening,” and the framing entered the policy conversation on AI-assisted medical imaging quickly. The framing carried more weight than the underlying retrospective comparison was designed to support: a retrospective evaluation against historical reader performance is informative about model output, but prospective deployment against current radiologists operating in their current workflow involves variables — reader fatigue, screen presentation, integration with reading worklists, recall thresholds — that the retrospective study does not measure. The McKinney paper itself was careful in its claims; the gap between the paper’s careful claims and the press’s framing of the headline is part of the case.¹

THE CAPABILITY QUESTION

The October 14, 2020 Nature comment by Benjamin Haibe-Kains, George Adam, Ahmed Hosny, Farnoosh Khodakarami, Massive Analysis Quality Control (MAQC) Society Board of Directors, Levi Waldron, Bo Wang, Chris McIntosh, Anna Goldenberg, Anshul Kundaje, Casey S. Greene, Tamara Broderick, Michael M. Hoffman, Jeffrey T. Leek, Keegan Korthauer, Wolfgang Huber, Alvis Brazma, Joelle Pineau, Robert Tibshirani, Trevor Hastie, John P. A. Ioannidis, John Quackenbush, and Hugo J. W. L. Aerts is the load-bearing reproducibility critique. The comment argued that the McKinney paper had not released code, had not released trained models, and had not provided sufficient methodological detail to permit independent reproduction. The comment was specific: a large fraction of the methodology was not reproducible from the published paper, and the screening-comparison claim could not be independently validated. The comment did not allege error in the paper; it argued that reproducibility had not been established.²

EARLY EVIDENCE

The case pairs with Case 35 (Radiology AI Miscalibration) for the medical-imaging-AI-deployment-evidence thread: retrospective evaluation produces one class of evidence; prospective deployment produces another, and the two are not interchangeable. Pair with Case 5 (Epic Sepsis) for the high-profile-result-versus-deployment-evidence thread in healthcare AI; Epic Sepsis is the load-bearing case in the corpus for the gap between vendor or developer claims and external evaluation, and DeepMind Mammography sits in the same conceptual family at a different domain. Pair with Case 26 (Pulse oximetry) for the population-heterogeneity-in-medical-AI thread; the McKinney paper’s UK-versus-U.S. effect-size difference (5.7 vs 1.2 pp on false-positives) is itself evidence that the system’s performance varies across screening populations, and the variation has implications for deployment.³

OPEN PROBLEMS

The honest hedges the case carries are load-bearing. The Haibe-Kains comment is not a finding that the McKinney paper was wrong; it is a finding that the paper as published did not establish reproducibility. The subsequent five years of deployment evidence on AI-assisted breast cancer screening have refined the screening-comparison framing — prospective evaluations have shown benefits in some settings and not in others, and the

operational variables the retrospective comparison did not measure have proved load-bearing in deployment. The case teaches the verification-as-deployment-event pattern: a high-profile retrospective result is the starting point of a verification arc, not its endpoint, and the reproducibility infrastructure the Haibe-Kains comment named is the condition for the arc to be possible.

A large fraction of the methodology was not reproducible from the published paper, and the screening-comparison framing has been refined by subsequent deployment evidence rather than confirmed at the deployment scale the headline implied.

EDITORS' SYNTHESIS OF THE MCKINNEY ET AL. (2020) AND HAIBE-KAINS ET AL. (2020) NATURE EXCHANGE.

REFERENCES

1. McKinney, S. M., Sieniek, M., Godbole, V., Godwin, J., Antropova, N., Ashrafian, H., Back, T., et al. (2020), "International evaluation of an AI system for breast cancer screening," Nature 577:89–94, doi:10.1038/s41586-019-1799-6.

2. Haibe-Kains, B., Adam, G. A., Hosny, A., Khodakarami, F., MAQC Society Board of Directors, Waldron, L., Wang, B., et al. (2020), "Transparency and reproducibility in artificial intelligence," Nature 586:E14–E16, doi:10.1038/s41586-020-2766-y.

3. McKinney et al. (2020), reply to Haibe-Kains et al., Nature 586:E17–E18 — the developers' response on the reproducibility-infrastructure question.

4. Freeman, K., Geppert, J., Stinton, C., Todkill, D., Johnson, S., Clarke, A., & Taylor-Phillips, S. (2021), "Use of artificial intelligence for image analysis in breast cancer screening programmes: systematic review of test accuracy," BMJ 374:n1872 — independent systematic review of subsequent AI-screening-deployment evidence.

THE LEARNING ENGINEERING LENS

LE INSIGHT

DeepMind Mammography is the verification-as-deployment-event case at the high-profile-publication scale. The McKinney paper's retrospective result was the starting point of a verification arc, not its endpoint; the Haibe-Kains comment named the reproducibility infrastructure as the condition for the arc, and subsequent deployment evidence has refined the screening-comparison framing the original paper offered.

LENS APPROACH

DeepMind Mammography is the deployment-and-reuse-as- verification-events case at the high-profile-publication seam (induced 7.2; LENS D4+D3/PT6; CLO-4 and CLO-3). LENS uses it in Domain 4 (Test and Evaluation) for the reproducibility-infrastructure-as-verification-condition discipline and in Domain 3 (Human-System Collaboration) for the retrospective-versus-prospective-evidence distinction. Pair with Case 35 (Radiology AI Miscalibration), Case 5 (Epic Sepsis), and Case 26 (pulse oximetry population heterogeneity). The Haibe-Kains comment is a reproducibility finding, not a finding of error; the distinction is the load-bearing hedge.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	T K D	D1 D2 D3 D4 D5	6	7.2	4-3
MEDICAL IMAGING		Test and Evaluation + D3			
MACHINE LEARNING					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Release code, trained models, and sufficient methodological detail to permit independent reproduction as a condition of publishing a high-profile retrospective medical-AI result; the Haibe-

IN OPERATION / AFTER

- Carry the Haibe-Kains comment's specific framing into print without softening; the comment is a finding on reproducibility, not a finding of error, and the case's

Kains comment names the reproducibility infrastructure as the condition for the verification arc that the original paper opens.

- ▶ Specify in advance the deployment variables — reader fatigue, recall threshold, worklist integration, screening-population characteristics — that a retrospective comparison does not measure but that a deployment will encounter.
- ▶ Treat the gap between the paper's careful claims and the press's framing as a deployment surface, not a communications problem; the framing the field receives is the framing the deployment will operate under in the policy conversation.

pedagogical value depends on the distinction being preserved.

- ▶ Pair in syllabi with Case 35 and Case 5 so the high-profile-result-versus-deployment-evidence pattern is taught across the medical-imaging-AI deployment seam at multiple instances.
- ▶ Use the case to anchor the verification-as-deployment-event frame; the curricular target is the discipline of treating a high-profile retrospective result as the starting point of a multi-year verification arc rather than as a deployment-ready endpoint.

REFLECTION QUESTIONS

1. Identify a high-profile retrospective result in your domain whose reproducibility infrastructure — code, trained models, methodological detail — has not been released. What would the verification arc the result opens require to proceed, and what currently blocks it?
2. Specify the deployment variables a retrospective comparison in your setting does not measure but a deployment will encounter. What is the prospective evaluation design that would surface those variables before deployment scale?
3. The press framing of a result often carries deployment implications the paper's careful claims do not. Pick a result in your domain and ask: what is the gap between the careful claim and the framing the field receives, and what would have to be true for the gap to be narrowed in advance of deployment decisions?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms

G Governance & Trust **D** Designed Out **N** Normalization of Deviance

IMPACT An FDA-mandated risk-evaluation and mitigation program with pregnancy testing, two-method contraception requirements, and lockout authorization before dispensing — Kaiser Permanente cohort (n=8,344; 9,912 treatment courses) reported 29 fetal exposures and 'no evidence that iPLEDGE significantly decreased the risk of fetal exposure' relative to the prior SMART program

IN BRIEF

Isotretinoin is a highly effective acne medication and a known teratogen: fetal exposure causes severe birth defects. The FDA has required risk-management programs since the 1980s; in 2006 the agency replaced the prior SMART program with iPLEDGE — a formal Risk Evaluation and Mitigation Strategy (REMS) requiring pregnancy testing, two contraception methods (or documented abstinence), and pharmacy lockout authorization before each dispense. The case is the productive counterpoint to SUBSAFE (Case 173) and the WHO Surgical Checklist (Case 23): the structural move is the same — mandatory pre-authorization gating a consequential action — and the measured outcome is very different. The Shin et al. Kaiser Permanente cohort (J Am Acad Dermatol, 2011; n=8,344 patients across 9,912 treatment courses) found 29 fetal exposures and concluded “no evidence that iPLEDGE significantly decreased the risk of fetal exposure” compared with the prior program. The broader literature reports approximately 150 isotretinoin-exposed pregnancies continue annually in the US despite the program, with non-adherence — missed pills, inconsistent condom use — the documented driver. The teaching point is that an authorization mechanism without adherence support does not reliably deliver the capability it is built to enforce.

THE SHIFT

Isotretinoin (Accutane and successors) is a highly effective systemic treatment for severe acne, and a well-characterized teratogen. Fetal exposure during treatment causes a recognized pattern of severe birth defects. The FDA has therefore required some form of risk-management program around its prescribing since the late 1980s, evolving through pregnancy testing requirements, contraception counseling, and structured enrollment programs. iPLEDGE, introduced in 2006, was the first to use the formal Risk Evaluation and Mitigation Strategy (REMS) architecture and to require pharmacy lockout: no authorization, no dispense.⁹

WHAT IS EMERGING

The mechanism is what the case turns on. iPLEDGE requires the patient to register, document either two contraception methods or abstinence (with declared method), complete monthly pregnancy tests, and receive a per-cycle authorization code before the pharmacy can fill the prescription. Each prescriber, pharmacy, and wholesaler must also enroll in the program. The structural form is the same one SUBSAFE (Case 173) and the WHO Surgical Checklist (Case 23) both use successfully: mandatory pre-authorization gating a consequential action, with the authorization withheld until pre-conditions are verifiably met.¹

THE CAPABILITY QUESTION

The measured outcome diverges sharply. Shin et al. (Journal of the American Academy of Dermatology, 2011) studied 8,344 Kaiser Permanente patients across 9,912 isotretinoin treatment courses in the iPLEDGE era and identified 29 fetal exposures. The paper's conclusion is the load-bearing sentence: "no evidence that iPLEDGE significantly decreased the risk of fetal exposure" relative to the prior SMART program. The broader literature consistently reports that approximately 150 isotretinoin-exposed pregnancies continue annually in the United States despite the program operating as designed. The documented driver in both is non-adherence: missed pills, inconsistent condom use, the assumption that one cycle of missed contraception is unlikely to coincide with the teratogenic window.²

EARLY EVIDENCE

The teaching point is precise and load-bearing. iPLEDGE is not a program-design failure in the sense that SUBSAFE succeeded and iPLEDGE was sloppy. The mechanism is elaborate, the enrollment burden on prescribers and pharmacies is real, and the program runs as specified. What the evidence shows is that an authorization mechanism without adherence support does not reliably deliver the capability it is built to enforce — because the capability sits downstream of the authorization in patient behavior, and the program does not reach into that behavior. The structural form (mandatory pre-authorization) is the same as SUBSAFE's; the missing piece is the equivalent of SUBSAFE's lifecycle audit and recurring training that keep the certification real in operation.³

OPEN PROBLEMS

Drafted alongside SUBSAFE (Case 173) and the WHO Surgical Checklist (Case 23), iPLEDGE is the most analytically useful "mixed" case in the v2 sweep. It is not a failure of the form (the form has demonstrated successes); it is evidence that the form alone does not deliver capability when the capability depends on a behavior the program does not instrument. The case is the case-grounded basis for the new CLO **Judgment under inadequate evidence** — the practitioner designing a REMS-class program has to decide on the available evidence what additional adherence-support architecture the mechanism needs in order to deliver, and the iPLEDGE evidence is the worked example of why the question matters.

The authorization mechanism is operating as designed. The capability sits downstream of the authorization, and the program does not reach that far.

EDITORS' SYNTHESIS OF SHIN ET AL. (2011).

REFERENCES

1.

Shin, J., Cheetham, T. C., Wong, L., Niu, F., Kass, E., Yoshinaga, M. A., Sorel, M., McCombs, J. S., & Sidney, S. (2011), "The impact of the iPLEDGE program on isotretinoin fetal exposure in an integrated health care system," Journal of the American Academy of Dermatology, PMID:21565419.

2.

FDA, iPLEDGE program documentation (2006 – present) – REMS architecture and enrollment requirements.

3.

Collins, M. K., Moreau, J. F., Opel, D., Swan, J., Prevost, N., Hastings, M., Schwarz, E. B., & Ferris, L. K. (2014), "Compliance with pregnancy prevention measures during isotretinoin therapy," Journal of the American Academy of Dermatology, 70(1):55–59, PMID:24157382 — source of the 150 annual exposures figure.

4.

Sullivan et al. (2003), House Science Committee statement on SUBSAFE — the structural counterpoint (Case 173).

THE LEARNING ENGINEERING LENS

LE INSIGHT

iPLEDGE is the productive counterpoint to SUBSAFE and the WHO Surgical Checklist. The structural form — mandatory pre-authorization gating a consequential action — is the same; the measured outcome is very different because the capability the program exists to enforce depends on patient adherence the program does not instrument. The “no significant decrease” finding is load-bearing and survives into the case verbatim.

LENS APPROACH

iPLEDGE is the analytically useful “mixed” pre-authorization case (induced 4.4; LENS D3/PT5). LENS uses it in Domain 3 (Human–System Collaboration) for the CLO **Delegation with revocation** — what the authorization mechanism delegates and what it does not — and in Domain 4 (Test and Evaluation) for the CLO **Judgment under inadequate evidence** on whether the form will deliver in a specific context. Direct pair with SUBSAFE (Case 173) at the form-vs-context layer; pair with WHO Surgical Checklist (Case 23) at the mandatory-mechanism layer.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
PHARMA SAFETY	G D N	D1 D2 D3 D4 D5	5	4.4	4 · 3
REGULATORY PROGRAMS	→	→			
CLINICAL CARE		Navigating Sociotechnical Constraints			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

Specify in advance what the authorization mechanism reaches and what it does not — for iPLEDGE, the authorization controls dispensing but does not reach into patient adherence — so the gap is named at design time rather than discovered in outcome data.

Pair the authorization architecture with adherence support that instruments the behavior the capability depends on: in this case, contraception adherence, not just contraception declaration.

Design the outcome metric around the harm the program exists to prevent (fetal exposure), not around the mechanism’s own throughput (enrollment counts, authorization codes issued), so the evaluation is of the capability and not of the bureaucracy.

IN OPERATION / AFTER

Treat a “no significant decrease” finding as program evidence about the mechanism’s reach, not as an argument against authorization architectures in general; SUBSAFE and the WHO checklist demonstrate the form can deliver under the right conditions.

Re-engineer the adherence-support layer based on the documented driver of failure — for iPLEDGE, missed pills and inconsistent condom use — rather than tightening the authorization layer that is already working as specified.

Carry the “approximately 150 exposures annually” figure honestly in any communication about the program; iPLEDGE is operating as designed, and the design’s reach is what the evidence is about.

REFLECTION QUESTIONS

1. Identify a mandatory pre-authorization program in your domain. What does the authorization mechanism actually reach, and what does it not? Where in the patient or operator behavior does the capability the program exists to deliver sit downstream of the authorization?
2. Specify the adherence-support architecture you would add to convert an iPLEDGE-class authorization into a delivered capability. For isotretinoin the gap is documented (missed pills, inconsistent condom use); what is the analog in your context?

WHO BUILDS THIS governance, ethics & measurement · systems & human-factors engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · task analysis · design prototyping · incident analysis · safety dashboards

TeamSTEPPS

T Training Gap **N** Normalization of Deviance

IMPACT Improved teamwork, communication, and patient-safety culture across diverse settings; OR on-time first start +21%; adopted by thousands of healthcare organizations

IN BRIEF

TeamSTEPPS — developed jointly by the Agency for Healthcare Research and Quality and the Department of Defense and released in 2006 — is the healthcare analog of Crew Resource Management (Case 117): an evidence-based team-training framework distilled from decades of aviation, military, and nuclear research and adapted for clinical settings. It trains four core competencies: communication, leadership, situation monitoring, and mutual support. It is explicitly the translation pathway from high-reliability research into bedside practice — the cross-domain capability transfer LENS is built to teach. Because its implementation infrastructure was funded as part of the program, TeamSTEPPS moved from research to scaled deployment in years rather than decades, and has been adopted by thousands of healthcare organizations with measurable gains in teamwork and safety culture.

BACKGROUND

Decades of research in aviation, the military, and nuclear power had shown that teamwork — not just individual expertise — drives safety in high-consequence work, and the IOM's *To Err Is Human* had identified poor communication as a leading cause of medical harm. What clinical settings lacked was a structured, evidence-based way to build those team skills. The evidence existed and the diagnosis was clear; what was missing was the translation — a route from the high-reliability research base into a curriculum a hospital could actually run at the bedside.⁹

THE INTERVENTION

In 2006, the Agency for Healthcare Research and Quality and the Department of Defense jointly released TeamSTEPPS — Team Strategies and Tools to Enhance Performance and Patient Safety — the healthcare analog of Crew Resource Management. It trains four core competencies: communication, leadership, situation monitoring, and mutual support, with a structured curriculum and ready-made implementation materials. The joint AHRQ-DoD authorship mattered: it paired a healthcare research agency's clinical reach with the military's deep team-training experience, so the framework arrived already grounded in both the evidence base and the practicalities of running it.¹

HOW IT WORKED


TeamSTEPPS is explicitly the translation pathway from high-reliability-organization research into clinical practice — decades of cross-domain evidence adapted for the bedside. Crucially, its implementation infrastructure (master-trainer programs, toolkits, an institutional support center) was funded as part of the intervention, so adopting organizations had a route from training to sustained practice rather than a binder on a shelf. Funding the implementation alongside the curriculum is precisely what compressed the usual decades-long gap between a proven method and its scaled use, because the path from adoption to sustained practice was paid for in advance.²

THE EVIDENCE

Studies across diverse settings report improved teamwork, communication, and patient-safety culture, with concrete operational gains such as a 21% improvement in on-time first surgical starts. Thousands of healthcare organizations have adopted the framework, and AHRQ has continued to develop it, releasing TeamSTEPPS 3.0 in 2023. The continued development into a third version is itself evidence the transfer took hold — a framework that is still being maintained and revised nearly two decades on is one institutions kept using, not one that was issued and abandoned.³

WHAT TRANSFERRED

TeamSTEPPS is the canonical evidence that cross-domain capability transfer is engineerable — and that the long implementation gap can be dramatically shortened when the transfer is funded as part of the intervention rather than as an afterthought. Its four competencies map directly onto the argument that capability engineering is itself a teachable discipline. The deeper lesson is that the bottleneck in cross-domain transfer is rarely the knowledge, which already existed, but the funded path that carries it from research into routine practice — the part programs most often leave unbudgeted.⁴



TeamSTEPPS represents the translation pathway from high-reliability research into clinical practice.

EDITORS' SYNTHESIS DRAWING ON AHRQ TEAMSTEPPS 3.0 (2023) AND SALAS ET AL.

REFERENCES

1. AHRQ, TeamSTEPPS 3.0 Curriculum (2023) — the framework and four competencies.
4. Weaver, S., Dy, S. & Rosen, M. (2014) — patient-safety team-training implementation and outcomes.

2. DoD / AHRQ partnership documentation — the joint development and implementation infrastructure.
3. Salas, E., Rosen, M. et al. (2009) — cross-domain team-training evidence base.
5. American Hospital Association Team Training Center — adoption and on-time-start metrics.

THE LEARNING ENGINEERING LENS

LE INSIGHT

TeamSTEPPS is the canonical evidence that cross-domain capability transfer is engineerable and that it can shorten the implementation gap dramatically when the transfer is funded as part of the intervention rather than as an afterthought. The four competencies map directly to the LENS curriculum’s argument for what capability engineering competence looks like as a deliverable.

LENS APPROACH

LENS treats TeamSTEPPS in LEN 8 as the canonical cross-domain transfer case and in LEN 1 as evidence that the program’s core proposition — that learning engineering is a discipline — has institutional precedent. Studio projects (LEN 10) reference TeamSTEPPS as the worked example of cross-domain adaptation methodology.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HEALTHCARE	→ T N	→ D1 D2 D3 D4 D5	4	6.4	5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Distill the cross-domain evidence base into a structured curriculum of a few teachable competencies rather than leaving adopters to translate the research themselves.
- ▶ Pair authorship across the source and target domains (as AHRQ and DoD did) so the framework arrives grounded in both the evidence and the practicalities of running it.
- ▶ Fund the implementation infrastructure — master-trainer programs, toolkits, a support center — as part of the intervention, not as an afterthought.

IN OPERATION / AFTER

- ▶ Maintain and revise the framework over time (as with successive TeamSTEPPS versions) so it stays in use rather than being issued and abandoned.
- ▶ Track concrete operational gains (such as on-time first surgical starts) alongside culture measures, so the transfer’s effect is shown, not assumed.
- ▶ Identify and budget the specific path from adoption to sustained practice when scaling, since that unbudgeted step is the usual bottleneck in cross-domain transfer.

REFLECTION QUESTIONS

1. What is the cross-domain transfer your domain has not yet executed? What evidence base in another industry should it draw on?
2. TeamSTEPPS funded its implementation infrastructure. Design the implementation-infrastructure budget for an equivalent transfer in your domain.
3. TeamSTEPPS shortened the usual decades-long gap because the path from adoption to sustained practice was paid for in advance. Identify a proven method in your domain that has not scaled, and name the specific unbudgeted step that is blocking it.

WHO BUILDS THIS learning science & instructional design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · incident analysis · safety dashboards

CASE 40

CLINICAL/TRANSLATIONAL RESEARCH

TEAM SCIENCE

HEALTHCARE WORKFORCE

2020 – 2022

Team Science Training for Clinical and Translational Scientists

K Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Colorado CTSA team-science training (221 registrants; pre/post survey N=117/46) reported statistically significant improvement ($p < 0.05$) across all three TEAMS instrument competencies — Team Planning, Managing a Team, and Interpersonal Relations — with the largest gains in structured/tool-based domains and the smallest in Interpersonal Relations

IN BRIEF

The Colorado Clinical and Translational Sciences Institute built a structured team-science training program — explicit competency targets, structured exercises, mentor pairings — and evaluated it across three 2020–2022 cohorts (221 registered participants; 117 pre-program and 46 post-program survey respondents) using a pre/post design. The evaluation introduces and validates a measurement instrument (TEAMS) that resolves team-science capability into three components: Team Planning, Managing a Team, and Interpersonal Relations. Participants showed statistically significant improvement ($p < 0.05$) across all three components — with the largest gains in the structured / tool-based domains (Team Planning, Managing a Team) and the smallest gain in Interpersonal Relations, the component the authors explicitly note is the hardest. The case treats interdisciplinary collaboration as an engineerable, measurable capability — the amended Domain 3 sub-competency (collaboration as a unit of measurement, distinct from any individual operator) — and is honest about the interpersonal half moving least. The honest limit the authors name: outcomes are self-reported and the design lacks longitudinal tracking and integration metrics; this is perceived-competency-gain evidence, not yet demonstrated downstream research– collaboration impact.

BACKGROUND

Clinical and translational research is the structural example of a discipline whose capability sits in the collaboration itself rather than in any single team member. A protocol moves across investigators, study coordinators, biostatisticians, clinical staff, and regulatory specialists; the science depends on the team's coordination as much as on any individual's expertise. The capability question the case addresses is whether interdisciplinary collaboration can be trained — and whether the resulting capability can be measured.⁹

THE INTERVENTION

The Colorado Clinical and Translational Sciences Institute built a structured team–science training program with explicit competency targets, structured exercises, and mentor pairings. The program ran across three 2020–2022 cohorts with 221 registered participants; the pre/post design drew 117 pre-program and 46 post-program survey respondents. The evaluation paper’s central methodological contribution is the validation of a measurement instrument — TEAMS — that resolves team–science capability into three competency components: Team Planning (the structured front–end of collaboration), Managing a Team (the structured operational half), and Interpersonal Relations (the unstructured half that the literature describes as the hardest to train).¹

HOW IT WORKED

Pre/post comparison showed statistically significant improvement ($p < 0.05$) across all three components. The pattern of gains is what makes the case teachable: the largest gains were in the structured / tool-based domains — Team Planning and Managing a Team — and the smallest gain was in Interpersonal Relations. The authors do not soften this; the paper states it plainly, and the smaller-gain-on-the-hardest-half finding is the load-bearing honest result.²

THE EVIDENCE

The honest limit the authors also name carries into the case. The TEAMS instrument and the pre/post comparison measure self-reported perceived competency, not downstream collaboration outcome. The design did not track participants longitudinally into their next collaborative projects, did not measure integration metrics on the projects they ran after the training, and did not compare to a non-trained control. The evidence is the strongest available current evidence for a structured team–science training program; it is not closed proof of downstream collaboration impact, and the case carries the qualification rather than collapsing it.³

WHAT TRANSFERRED

In pair with IPE (Case 28) and Implementation Science Training (Case 41), the Colorado CTSA case is the small-tier intervention companion to two frontier cases. Team–science training is one of the few cases in the corpus that operationalizes “collaboration as the unit of measurement” — the CLO — with a validated instrument and structured evaluation. The case is a worked example of how a discipline can convert collaboration capability from an aspiration to a measurable target, while preserving the honest qualifications about what the measurement does and does not establish.

The largest gains were in the structured half. The hardest half — interpersonal relations — moved least, and the paper states it plainly.

EDITORS’ SYNTHESIS OF COLORADO CTSA TEAM–SCIENCE TRAINING EVALUATION.

REFERENCES

1. Colorado CTSA, “Team science training for clinical and translational scientists: An assessment of effectiveness,” *Journal of Clinical and Translational Science*, PMC12392353.
2. National Research Council (2015), *Enhancing the Effectiveness of Team Science* — the IOM/National Academies team–science synthesis.
3. v2 paired cases: IPE evidence gap (122), Implementation Science Training (123) — the frontier/measurement companions.

2. Falk-Krzesinski et al. (2010), "Mapping a research agenda for the science of team science," Research Evaluation — broader team-science literature backdrop.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Colorado CTSA team-science training program is one of the few cases in the corpus that operationalizes "collaboration as a unit of measurement." The TEAMS instrument resolves the capability into three components; structured training moved all three with $p < 0.05$ but moved interpersonal relations least. The honest limit — self-reported perceived competency, no longitudinal tracking, no control — survives into the case.

LENS APPROACH

Team-science training is the case-grounded basis for the collaboration-as-measurement CLO (induced 4.3; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the team-coordination redesign and in Domain 4 (Test and Evaluation) for the TEAMS instrument validation and the perceived-competency vs. downstream-impact discipline. Direct pair with Cases 117 (IPE) and 123 (implementation science training) — the intervention with measurement against two frontier cases where the measurement is the gap.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
CLINICAL/ TRANSLATIONAL RESEARCH TEAM SCIENCE HEALTHCARE WORKFORCE	K N	D1 D2 D3 D4 D5	4	4.3	4 · 5
→ Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Resolve collaboration capability into measurable components — Team Planning, Managing a Team, Interpersonal Relations — and structure the training so each component has tools, exercises, and an assessment instrument.
- ▶ Validate the measurement instrument before relying on it for evaluation; the TEAMS instrument's three-component structure is the load-bearing methodological contribution.
- ▶ Acknowledge in the program design that the interpersonal component is the hardest to train and the slowest to move; design the exercises and the mentor pairings around that asymmetry rather than against it.

IN OPERATION / AFTER

- ▶ Report the pattern of gains, not just the aggregate; the smaller-gain-on-interpersonal finding is the honest signal about what the structured training can and cannot do.
- ▶ Commission longitudinal follow-up that tracks participants into their next collaborative projects and measures integration outcomes — the case names the gap and the next study addresses it.
- ▶ Carry the "collaboration as the unit of measurement" sub-competency into the curriculum: the Domain 3 sub-competency is grounded in evidence here, and the team-science training program is one of the few cases in the corpus that operationalizes it.

REFLECTION QUESTIONS

1. Identify a collaboration-dependent capability in your domain. What three or four components would you resolve it into for measurement purposes, and what instrument would you validate to test the resolution?
2. Specify the longitudinal follow-up you would build into the next iteration of a team-science training program — what integration metric, on what cadence, against what comparison — to convert perceived-competency-gain evidence into downstream-collaboration-impact evidence.
3. The Colorado CTSA finding is that the structured half moves first and the interpersonal half moves least. What does that imply about the curriculum you would design — and about which components are realistic targets for a single-program training intervention vs. an institutional / cultural intervention?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

CASE 41

IMPLEMENTATION SCIENCE

TRANSLATIONAL WORKFORCE

EDUCATION AT SCALE

2020S

Implementation-Science Training — Stated Goals Outrunning Operational Practice

K Knowledge & Institutional Memory**N** Normalization of Deviance

IMPACT Survey of CTSA-funded TL1 training programs (N=50): most name collaboration, team science, and multi/inter/cross-disciplinary training as goals, but far fewer name the specific competency practices as distinguishing features — experiential learning in 24% of programs, a competency-based curriculum in just 6%

IN BRIEF

Implementation science — the discipline of moving validated evidence into operational practice — is one of the most consequential cross-domain-transfer competencies in contemporary medicine and adjacent fields, and the CTSA T32 / TL1 program survey (N=50 programs) examined whether it is being taught systematically. The headline finding is the same structural pattern as the IPE case (122) at smaller scale: most programs name the right goals — collaboration, team science, multi/inter/cross-disciplinary training — but far fewer name the specific competency practices — a competency-based curriculum in just 6% of programs, experiential learning in 24%, program evaluation in 38% — that those stated goals imply. The picture is of a field that has converged on what interdisciplinary translation training should aim for, but where the operational practices lag the stated goals — the same enthusiasm-ahead-of-evidence pattern as IPE. The case is the workforce-training counterpart to Case 13 (the “17-year gap” between research evidence and practice adoption); the v2 frame is the gap between what implementation-science training programs **say** they do and what they **operationally** do.

THE SHIFT

Implementation science exists because the gap between research evidence and operational practice in medicine — the so-called “17-year gap” the v1 corpus documents at v1 Case 13 — is large enough to constitute a discipline-level capability question. Moving validated evidence into clinical practice reliably is itself a competency that has to be designed, trained, and evaluated. The CTSA (Clinical and Translational Science Awards) program is the primary US mechanism for building that workforce, and the TL1 / T32 programs are its training arm.^o

WHAT IS EMERGING

The case’s evidence base is a 2021 survey of CTSA-funded TL1 training programs (N=50) examining what those programs actually do at the operational level. The survey covers the stated goals — what the programs say they aim to develop — and the operational practices

— what the programs do, in practice, to develop it. The stated-goals side is where the programs converge: most name collaboration, team science, and multi/inter/cross-disciplinary training as central objectives. The operational-practices side is where the pattern surfaces: far fewer programs name the specific competency practices as distinguishing features — a competency-based curriculum in just 6%, experiential learning in 24%, program evaluation in 38% — at the level the stated goals would imply.¹

THE CAPABILITY QUESTION

The structural form is the same pattern IPE shows at field scale (Case 28): the field has converged on what interdisciplinary translation training should aim for, but the operational practices lag the stated goals. At the implementation-science training scale, the lag has specific consequences: programs that intend to produce graduates capable of moving evidence into practice often do not systematically assess whether they have produced that capability, and the field's measurement of its own workforce-development effectiveness is correspondingly sparse.²

EARLY EVIDENCE

What the case teaches is the operational-practice gap as a designable target, not as a curricular failure. The low uptake of the specific competency practices is not an indictment of the programs that fall on the wrong side of it; it is a finding about the operational-infrastructure investment that competency-based training and assessment require, and that field-scale stated-goal convergence does not by itself produce. The pair with Case 13 (the 17-year gap) gives the case its frame: the implementation-science workforce is the recovery mechanism for the research-to-practice gap, and the operational capacity of that workforce is itself the capability the training programs are trying to build.³

OPEN PROBLEMS

In the multidisciplinary-translation trio (Cases 40 + 30 + 31), implementation-science training sits between team science (where the measurement is possible at program scale) and IPE (where the measurement is the field-scale gap). The trio teaches the enthusiasm-evidence-gap sub-competency from three angles — the program-scale success, the field-scale gap, and the operational-practice gap inside training programs themselves — and it teaches the CLO **Judgment under inadequate evidence** by example: practitioners designing implementation-science training programs have to decide what to build on incomplete evidence of what works, while the field is still building the evidence architecture that would let them decide better.

Stated goals converge on the right targets. Operational practices lag. The gap is designable, not a curricular failure.

EDITORS' SYNTHESIS OF THE CTSA TL1 PROGRAM-GOALS STUDY (2021).

REFERENCES

1. CTSA T32/TL1 program-goals study (2021), *Journal of Clinical and Translational Science*, PMC8826009.
2. Morris, Wooding, & Grant (2011), "The answer is 17 years, what is the question: understanding time lags in translational research," *Journal of the Royal Society of Medicine — the original 17-year-gap source for Case 13*.
3. Brownson, Colditz, & Proctor (2018), *Dissemination and Implementation Research in Health* (2nd ed.) — the broader implementation-science synthesis.
4. v2 paired cases: Team-science training (121), IPE evidence gap (122).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The CTSA TL1 program survey is the workforce-training instance of the enthusiasm-evidence gap pattern: most programs name competency-based training, assessment, and program evaluation as goals; the specific competency practices that would deliver on those goals are named by far fewer — a competency-based curriculum in just 6% of programs, experiential learning in 24%. The implementation-science workforce is the recovery mechanism for the 17-year research-to-practice gap; the case names the gap inside the recovery mechanism itself.

LENS APPROACH

Implementation-science training is the frontier workforce-training case (induced 6.4; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the cross-disciplinary translation discipline; in Domain 4 (Test and Evaluation) for the enthusiasm-evidence-gap sub-competency; and in Domain 2 (Iterative Development) for the CLO **Judgment under inadequate evidence** as practitioners design training programs while the field-scale evidence is still being built. Pair with Cases 40, 28 as the multidisciplinary-translation trio; workforce-training counterpart to Case 13.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
IMPLEMENTATION SCIENCE	K N	D1 D2 D3 D4 D5	4	6.4	4 · 5
TRANSLATIONAL WORKFORCE	→	→ Navigating Sociotechnical Constraints			
EDUCATION AT SCALE					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the competency-based assessment instrument before launching the training program; the gap the survey documents is partly the result of stated competencies that were never operationalized into measurement.
- Build experiential learning into the operational structure of the program, not as an add-on; the survey found experiential learning named by only 24% of programs — one of the operational dimensions that lags.
- Plan program evaluation as a deliverable of the training program itself, with the cadence, instrument, and reporting venue named at launch.

IN OPERATION / AFTER

- Report the operational-practice gap honestly: stated goals are the convergence point, operational practices are the lag, and the lag is the addressable target.
- Treat the multidisciplinary-translation trio (Cases 40, 28, 41) as a unit; the three-angle teaching of the enthusiasm-evidence pattern is more useful than any single case can be.
- Connect the case explicitly to Case 13 (the 17-year gap): the implementation-science workforce is the recovery mechanism for that gap, and the operational capacity of the workforce is the capability that has to be built — which is the case's pedagogical point.

REFLECTION QUESTIONS

1. Identify a training program in your domain whose stated goals include competency-based assessment, program evaluation, or experiential learning. What proportion of those stated goals are operationalized into specific instruments, cadences, and reporting structures — and which are at the goal-statement layer only?
2. Specify the competency-based assessment instrument you would build into the next iteration of an implementation-science training program. The survey found a competency-based curriculum named by just 6% of programs — the instrument is what is missing more than the intent; what is the instrument?
3. The implementation-science workforce is the recovery mechanism for the 17-year research-to-practice gap (Case 13). What is the analog in your domain — the workforce whose operational capacity is the recovery for a documented systemic gap — and what is the case for investing in that workforce's training architecture?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

Australian Hospital-Pharmacy Technician Role Redesign

D Designed Out **N** Normalization of Deviance **H** Human-System Interface

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT Expanded pharmacy-technician scope (final accuracy checking, drugs-of-addiction management, clinical support) reportedly cut average prescription turnaround from 18.5 to 12.3 minutes, increased throughput from 220 to 295 prescriptions per shift, and decreased dispensing errors from 2.1% to 1.2% — throughput and the safety metric moving in the same direction

IN BRIEF

Australian hospital pharmacies faced the structural problem most healthcare workforces meet at some point: the pharmacist's capacity was being absorbed by dispensing-accuracy checking, which crowded out the clinical work that requires pharmacist judgment. The 2016 Pharmacy Technician and Assistant Role Redesign project expanded pharmacy-technician scope to include final accuracy checking, drugs-of-addiction management, and clinical-support duties, with the design intent of freeing pharmacist capacity without degrading dispensing safety. The reported operational outcomes moved throughput and the safety metric in the same direction: average prescription turnaround fell from 18.5 to 12.3 minutes, prescriptions per shift rose from 220 to 295, and dispensing errors declined from 2.1% to 1.2%. A 2021 Journal of Pharmacy Practice and Research cross-sectional survey extended the evidence base into workforce-acceptance attitudes. The evidence-tier flag renders under the title: the operational figures come from a program-report rather than an independent multi-site audit. Future validation will continue on long-term safety durability and on whether the role-redesign pattern generalizes to other healthcare-workforce roles.

BACKGROUND

The structural problem the case sits inside is the C3 thread where monitoring or checking duties are loaded onto the role that needs to be doing the clinical work — and the checking work absorbs the capacity. Australian hospital pharmacies in the early 2010s were in this pattern: the pharmacist's day was consumed by dispensing-accuracy checking, and the clinical-support work that requires pharmacist judgment was crowded out. The induced 3.2 framing is the precise diagnosis: the unsupportable role is the pharmacist as primary accuracy-checker.^o

THE INTERVENTION

The redesign expanded pharmacy-technician scope to include final accuracy checking, drugs-of-addiction management, and clinical-support duties. The design move is not the addition of a checker but the redistribution of authority: the technician is given the role that the pharmacist was doing, with the training, certification, and supervisory structure to back it. The pharmacist is freed to do the clinical-judgment work that the system needed pharmacists for. The redesign is the C3 role-design intervention the induced framework calls for at the small tier.¹

HOW IT WORKED

The reported operational outcomes moved throughput and the safety metric in the same direction — an unusual combination in healthcare workforce changes. Average prescription turnaround fell from 18.5 to 12.3 minutes; prescriptions per shift rose from 220 to 295; dispensing errors fell from 2.1% to 1.2%. The redesign did not simply trade safety for throughput. The 2021 Journal of Pharmacy Practice and Research cross-sectional survey extended the evidence into the workforce-attitudes layer — what hospital pharmacists and technicians thought about the expanded scope and where acceptance was strongest.²

THE EVIDENCE

The evidence-tier flag is load-bearing. The headline operational figures come from a program-report and rest on the project's internal measurement; the redesign was not subjected to an independent multi-site audit, and the magnitudes have not been replicated by a peer-reviewed controlled evaluation. The cross-sectional survey is peer-reviewed but measures attitudes rather than outcomes. The case is included because the pattern is teachable and the practitioner literature is consistent; the magnitudes should travel with the tier flag intact. Future validation will continue.³

WHAT TRANSFERRED

What the case teaches at the LENS layer is the role-design move as a sociotechnical-constraints intervention. The pharmacist-as-checker pattern was not a property of any individual; it was a property of the role architecture, and the redesign treated the architecture as the design variable. The Domain-5 frame applies: the work was to redraw the authority gradient and team-coordination boundary so that the system's capability — fast, safe dispensing plus pharmacist clinical judgment — emerged from the team rather than depending on the pharmacist's individual heroism. The case is also a Gap-5 echo: a non-US small-tier role-redesign success with documented operational outcomes.

The pharmacist-as-checker pattern was a property of the role architecture, not the pharmacist. The redesign treated the architecture as the design variable.

EDITORS' SYNTHESIS OF THE 2016 PROJECT REPORT AND ANDERSON ET AL. (2021).

REFERENCES

1. SHPA / Australian hospital-pharmacy network (2016), "Pharmacy Technician and Assistant Role Redesign within Australian Hospitals Project," outcomes report.
2. Anderson et al. (2021), "Perceptions of hospital pharmacists and pharmacy technicians towards expanding roles for hospital pharmacy technicians: a cross-sectional survey," Journal of Pharmacy Practice and Research, doi:10.1002/jppr.1697.
3. Boughen, Sutton, Fenn, & Wright (2017), "Defining the Role of the Pharmacy Technician and Identifying Their Future Role in Medicines Optimisation," Pharmacy 5(3):40 — UK companion analysis.
4. Royal Pharmaceutical Society and SHPA practice statements on technician scope expansion — practitioner-tier framing the project sits inside.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Australian pharmacy-technician redesign is a small-tier C3 role-redesign intervention with both throughput and safety moving in the same direction (turnaround, throughput, error rate). Evidence base is program-report plus a peer-reviewed attitudes survey; the operational magnitudes rest on the project's internal measurement and have not been independently audited at multi-site scale. Future validation ongoing.

LENS APPROACH

Australian pharmacy-technician redesign is the C3 small-tier role-redesign case (induced 3.2 and 4.3; LENS D5/PT3). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for CLO-5 — the work was to redraw the authority gradient — and in Domain 1 (Systems Analysis) for CLO-1, since the role-architecture diagnosis was the precondition for the redesign. Pair with Cases 32–132 for the small-tier C3 thread; Gap-5 echo as a non-US case.

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
HOSPITAL PHARMACY HEALTHCARE WORKFORCE ROLE REDESIGN	D N H	D1 D2 D3 D4 D5 Navigating Sociotechnical Constraints	3	3.2	1·5

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Diagnose the role-architecture failure before redesigning: which capability does the system need from which role, and where is checking work absorbing capacity that the system needs elsewhere.
- ▶ Redistribute authority, not just tasks: the technician's expanded scope must come with training, certification, and supervisory structure, not just a new line in the procedure manual.
- ▶ Specify both throughput and safety metrics before the redesign ships, so the joint movement (or trade-off) of both is what the report has to defend.

IN OPERATION / AFTER

- ▶ Treat the program-report figures as practice-synthesis evidence; carry the tier flag into any generalization to other sites or other workforce roles.
- ▶ Track durability of the safety metric over years, not weeks; the dispensing-error reduction is the load-bearing result and must hold under personnel rotation and volume changes.
- ▶ Use the workforce-acceptance evidence (2021 survey) to identify where the redesign meets resistance, and stage the rollout against that resistance rather than ignoring it.

REFLECTION QUESTIONS

1. Identify a healthcare role in your setting where checking or monitoring work is absorbing the capacity the system needs for clinical judgment. What is the technician- or assistant-equivalent role redesign that would redistribute the authority, and what training and supervisory structure would it require?
2. Specify the throughput and safety metrics you would commit to before the redesign ships, so the joint movement of both is what the report has to defend. The Australian project teaches because both metrics moved in the same direction — what would the equivalent be in your context?
3. The operational figures are program-report tier. What independent or multi-site evidence would you require before treating the magnitudes as generalizable to other workforce roles (nursing, radiology, lab) or other countries?

WHO BUILDS THIS systems & human-factors engineering · safety science & organizational analysis · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · incident analysis · safety dashboards · usability testing · cognitive walkthroughs

Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface

H Human–System Interface **N** Normalization of Deviance

EVIDENCE TIER preprint-tier — source confidence flagged; future validation ongoing.

IMPACT A Rwanda mHealth monitoring system equipped community health workers with mobile decision support for maternal and newborn care; the published evaluations were mixed — uptake of maternal and newborn services rose only where the system was paired with added training, supervision, and equipment, while a controlled national-survey time-series found no significant change from the system alone — with the technology designed to extend the CHW's diagnostic-and-referral role rather than replacing it

IN BRIEF

Between 2013 and 2018 Rwanda's Ministry of Health, working with research partners, deployed an mHealth monitoring system that placed mobile-phone-mediated decision support into the hands of community health workers (CHWs) responsible for maternal and newborn care, including surgical-site infection screening after Cesarean. Peer-reviewed evaluations were mixed: uptake rose only where the system was paired with added support (Ruton et al., 2018), while a controlled national-survey time-series found no significant change attributable to the system alone (Hategeka et al., 2019). The case sits at the small (program/study) tier and teaches a precise pattern: capability delivered at the periphery of the formal health system, with the technology designed to extend the CHW's diagnostic-and-referral role rather than replace it. The evidence-tier flag is honest — one peer-reviewed evaluation, with broader impact claims still resting partly on practitioner reporting. As a non-US small-tier case it pairs naturally with the PEPFAR Sub-Saharan training-modality comparison as the African workforce-capability evidence the corpus has been thin on, and it carries the standing "future validation ongoing" language into print rather than overclaiming what one study can settle.

BACKGROUND

Rwanda's health system, post-genocide, was rebuilt around community-level delivery: every village has elected CHWs who handle a defined scope of maternal, newborn, and child-health work and refer upward when their scope is exceeded. The capability question for the mHealth program was specific — could a mobile decision-support tool, in CHW hands,

increase the uptake of maternal and newborn health services in a way that the prior paper-based workflow had not.⁰

THE INTERVENTION

The system delivered structured prompts, reminders, and a referral pathway through the CHW's mobile phone, with an added screening flow for surgical-site infection after Cesarean delivery — a recognized post-discharge failure mode where signs first appear in the village, not the clinic. The design move worth naming is that the technology was scoped to extend the CHW's existing diagnostic-and-referral role, not to substitute for clinical judgment further up the chain. The CHW remained the capability interface; the phone was the cue and the record.¹

HOW IT WORKED

Ruton et al. (2018) report increased uptake of maternal and newborn health services only where the mHealth system was paired with additional training, supervision, and equipment; RapidSMS alone was ineffective. A controlled time-series on national survey data (Hategeka et al., 2019) found no significant change attributable to the system, largely for ceiling-effect reasons. The published evaluation is the anchor finding; downstream reporting (MIT News 2022 and subsequent AI-augmented maternal-care work) describes the program's continuation and adjacent developments but is journalism-tier, not investigation-grade. The case carries the preprint-tier evidence flag honestly: one peer-reviewed evaluation does not close the question of long-term outcome durability or generalization to other low-resource settings.²

THE EVIDENCE

The teaching point is the location of capability. The mHealth tool did not centralize the work; it instrumented the periphery. The decisive variables — household trust in the CHW, the CHW's standing in the village, the referral pathway upward — sit in the sociotechnical layer the tool could support but not constitute. This is the inverse of capability deployments that try to relocate judgment to the center; Rwanda's system kept judgment where it already was and used the technology to make that judgment more reliable and more visible to the formal health system.³

WHAT TRANSFERRED

What survives the evidence-tier flag is the structural pattern: capability extension at the frontline, with mobile-mediated decision support designed around an existing role, in a low-resource setting where the formal health system cannot reach every birth. Future validation ongoing — both replication in other African and South Asian settings, and longer-term outcome data linking process uptake to maternal and newborn outcomes downstream. The case is included with the flag intact and pairs with the PEPFAR cross-country modality comparison as the Sub-Saharan workforce-capability evidence the v2 corpus needs.

The mHealth tool did not centralize the work. It instrumented the periphery — and kept judgment where it already was.

EDITORS' SYNTHESIS OF HATEGEKA ET AL. (2019), RUTON ET AL. (2018), AND THE RWANDA CHW PROGRAM DOCUMENTATION.

REFERENCES

1. Hategeka, C., Ruton, H., Law, M. R., et al. (2019), "Effect of a community health worker mHealth monitoring system on uptake of maternal and newborn health services in Rwanda," Global Health Research and Policy, PMC6429813.
2. Rwanda Ministry of Health, community health program documentation and CHW scope-of-practice guidance, 2013–2018.
3. MIT News (2022), reporting on subsequent AI-augmented maternal-care work in Rwanda — journalism-tier companion to the peer-reviewed evaluation.
4. Cross-reference: Case 201 (PEPFAR HIV training-modality comparison) for the paired Sub-Saharan workforce-capability evidence.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Rwanda mHealth is a small-tier capability-extension case at the frontline: technology designed around an existing CHW role, with peer-reviewed evaluation showing uptake gains only when the system was paired with added support, and a controlled national-survey series showing no significant change from the system alone. The preprint-tier flag is load-bearing — one study does not close the durability or generalization question, and the broader impact claims rest partly on practitioner reporting. Future validation ongoing.

LENS APPROACH

Rwanda mHealth is the small-tier non-US frontline-capability case (induced 6.4; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the periphery-extension pattern — the technology supports the CHW's diagnostic-and-referral role without relocating judgment to the center — and in Domain 2 (Iterative Development) for the scope-the-tool-to-the-role design discipline. Direct pair with Case 201 (PEPFAR Sub-Saharan training-modality comparison) as the African workforce-capability evidence the v2 corpus needs.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
GLOBAL HEALTH	H N	D1 D2 D3 D4 D5	4	6.4	2 · 5
MHEALTH	→	→			
MATERNAL AND NEWBORN CARE		Navigating Sociotechnical Constraints			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Scope the tool to extend an existing frontline role rather than substitute for it; the CHW's village standing is the load-bearing capability the technology can support but cannot create.
- ▶ Instrument the post-discharge failure modes that the formal health system cannot see — surgical-site infection after Cesarean is the named example — so the tool turns village-level observation into health-system-visible signal.
- ▶ Carry the evidence-tier flag through the deployment documentation: one peer-reviewed evaluation is the anchor finding, not the closure of the question.

IN OPERATION / AFTER

- ▶ Track uptake of the targeted services as the primary process measure, and link to longer-term maternal and newborn outcomes as the data matures; resist treating short-term process gains as outcome evidence.
- ▶ Report the CHW workload and tool-acceptance trajectory separately from the headline uptake figures; the periphery-extension pattern only holds if the frontline role remains sustainable.
- ▶ When asked whether the pattern generalizes, name the conditions Rwanda's system supplies (village-level CHW network, post-conflict institutional rebuild, single payer) before asserting transfer to other low-resource settings.

REFLECTION QUESTIONS

1. Identify a frontline capability role in your context (CHW, field technician, ward nurse, line supervisor). What would it mean to scope a mobile decision-support tool to extend that role rather than relocate judgment to the center? Which failure modes only the frontline can see should the tool surface upward?
2. The case rests on one peer-reviewed evaluation with broader claims supported by practitioner reporting. What is the minimum additional evidence — durability follow-up, replication in a second setting, linkage to maternal/newborn outcomes — you would require before treating the uptake finding as settled for a transfer decision?
3. Specify the village-level and health-system-level instrumentation you would put in place for a comparable deployment so that the frontline-extension pattern can be evaluated against an alternative deployment that relocates judgment to the clinic level.

WHO BUILDS THIS human factors & interaction design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · incident analysis · safety dashboards

Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD

G Governance & Trust **N** Normalization of Deviance

EVIDENCE TIER preprint-tier — source confidence flagged; future validation ongoing.

IMPACT Japan's 2019 PMD Act amendment (enforced September 2020) introduced IDATEN (Improvement Design within Approval for Timely Evaluation and Notice): a pre-agreed device-improvement plan approved at initial approval, streamlined review for in-scope changes. The 2025 scoping review of PMDA-approved AI radiology software documents transparency variability — hedge preserved

IN BRIEF

Japan's Pharmaceuticals and Medical Devices Agency (PMDA) operates one of the regulatory regimes most explicitly designed for the AI/Software-as-a-Medical-Device (SaMD) update problem. A 2019 amendment to the Pharmaceutical and Medical Device Act (enforced September 2020) introduced conditional early approval and IDATEN (Improvement Design within Approval for Timely Evaluation and Notice): a company can submit a proposed device-improvement plan at initial submission; once the plan is approved, subsequent changes within its scope receive streamlined review rather than full re-approval. The DASH for SaMD initiative supports faster reviews and earlier application. A 2025 scoping review on PMDA-approved AI radiology software documents transparency variability across approvals — a load-bearing hedge the case preserves. The teaching point is structural: the regulator designed change-control as a deliverable rather than defaulting to "approve once, then watch," which is the governance failure pattern that Epic-Sepsis-class deployments surface. Evidence-tier flag is preprint-tier for the most recent systematic analyses; future validation ongoing on outcome durability across approved devices.

BACKGROUND

Medical-device regulation evolved for hardware whose design is largely fixed at approval. AI/SaMD breaks that assumption: a learning model can be updated, a training-data distribution can shift, and the device approved on Tuesday can be a meaningfully different device by Friday. The governance question for any regulator is whether the approval architecture

handles the update or defaults to approve–once–then–watch. Japan’s PMDA chose the former and built the architecture explicitly.⁹

THE INTERVENTION

The 2014 Pharmaceutical and Medical Device Act laid the groundwork; the 2019 amendment introduced conditional early approval and the IDATEN system (Improvement Design within Approval for Timely Evaluation and Notice). Under IDATEN a manufacturer submits, alongside the initial approval package, a structured description of the modifications the manufacturer expects to make over the device’s life — with the test and validation plan that will accompany each modification. Once the improvement plan is approved, subsequent changes within its scope receive streamlined regulatory review rather than full re–approval.¹

HOW IT WORKED

The DASH for SaMD initiative complements IDATEN by supporting faster reviews and earlier application of SaMD products, anchoring the change–control architecture in a workflow that can keep pace with iteration. Together IDATEN and DASH are the regulator–side design for delegation–with–revocation: the deployer is delegated the modification authority within a pre–agreed scope, and the regulator retains the authority to review and revoke. The structure is the inverse of static–approval regimes that either block iteration or fail to track it.²

THE EVIDENCE

The 2025 medRxiv scoping review on PMDA–approved AI radiology software documents transparency variability across approvals — what is published about each device’s training data, intended use, and post–approval modification history is not uniform. The hedge is load–bearing: IDATEN is a structural improvement in regulatory architecture, but the per–approval transparency the framework enables varies and has been documented as a research finding rather than assumed as a feature. The 2021 medRxiv–then–published systematic review on PMDA AI/ML medical devices through 2020 supplies the prior baseline.³

WHAT TRANSFERRED

The case pairs structurally with the Epic–Sepsis governance gap earlier in the corpus and with the FDA’s evolving change–control SaMD policy. PMDA shows what designing change–control as a deliverable looks like at the regulator layer; Epic–Sepsis shows what happens when neither the vendor nor the deploying health system holds the change–control deliverable explicitly. The preprint–tier flag is honest: the regulatory framework is documented in program–report sources; the per–approval transparency analyses are preprint or recent. Future validation ongoing on outcome durability across approved devices.

The PMDA designed change–control as a deliverable rather than defaulting to approve once, then watch.

EDITORS’ SYNTHESIS OF THE PMD ACT AMENDMENT AND THE PMDA SAMD SCOPING REVIEWS.

REFERENCES

1. Kikuchi, et al. (2025), “Scoping Review of Regulatory Transparency in AI–based Radiology Software: Analysis of PMDA–approved SaMD Products,” medRxiv 2025.10.02.25336333.
2. Aoki, T., et al. (2021 → published), “Regulatory–approved Deep Learning/Machine Learning–Based Medical Devices in Japan as of 2020: A Systematic Review,” PMC9931274.
3. Pharmaceuticals and Medical Devices Agency of Japan, PMD Act amendment (2019) and DASH for SaMD program documentation.
4. Cross–reference: “A decade of review in global regulation and research of artificial intelligence medical devices (2015–2025),” PMC12310608 — comparative regulatory context.

THE LEARNING ENGINEERING LENS

LE INSIGHT

PMDA's IDATEN and DASH for SaMD are the non-US regulatory architecture for AI/SaMD change control: the modification scope is pre-agreed at initial approval, in-scope changes get streamlined review, and the delegation-with-revocation structure is explicit. The 2025 scoping review documents transparency variability across approvals — load-bearing hedge preserved. Preprint-tier flag for the recent systematic analyses; future validation ongoing.

LENS APPROACH

PMDA is the non-US regulator-designed change-control case (induced 5.4; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the delegation-with-revocation structure and in Domain 3 (Human-System Collaboration) for the CLO **Delegation with revocation** — the regulator delegates in-scope modification authority and retains revocation, rather than defaulting to static approval. Pairs with the Epic-Sepsis governance gap as the structural contrast case.

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
MEDICAL-DEVICE REGULATION AI/SaMD JAPAN	G N	D1 D2 D3 D4 D5 Navigating Sociotechnical Constraints	6	5.4	5 · 3

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat post-approval change control as a deliverable at the initial-approval submission, not a downstream addendum: the modification scope, test plan, and revocation conditions should be on the approval package from the start.
- ▶ Specify the delegation boundary explicitly — what the manufacturer can modify without re-approval and what triggers a full review — so the delegation-with-revocation structure is operationally clear to both parties.
- ▶ Carry the preprint-tier evidence flag through any policy decision the scoping review supports; the regulatory framework is documented, but the per-approval transparency analyses are recent.

IN OPERATION / AFTER

- ▶ Audit transparency across approvals — what is published per device about training data, intended use, modification history — as a separate program-evaluation question; the 2025 scoping review documents this variability as a research finding.
- ▶ Use IDATEN as the structural contrast with Epic-Sepsis-class deployments: the regulator-designed change-control deliverable is the inverse of the static-approval-plus-watch pattern.
- ▶ When importing the IDATEN pattern into another jurisdiction, name what the Japanese regulatory institution supplies (centralized agency capacity, established conditional-approval precedent) before asserting transfer.

REFLECTION QUESTIONS

1. Identify an AI/SaMD-adjacent product in your context whose post-deployment modification is anticipated. What would a IDATEN-style submission look like — the pre-agreed modification scope, the per-modification test plan, the revocation conditions — and which party currently holds each element?
2. The case rests on program-report sources for the framework and preprint analyses for the per-approval transparency findings. What is the minimum additional evidence you would require — outcome durability across approved devices, independent transparency audits — before treating IDATEN as a settled best-practice template?
3. Specify the delegation-with-revocation boundary you would write for a deploying organization adopting a SaMD product: which modifications can ship under the pre-agreed scope, which trigger re-review, and what evidence the deployer must publish at each modification event.

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

Chapter 3

Education, Training & the Learning Workforce — What Fails

When learning systems scale faster than their evidence.

The pilot results were real. The scale-up assumed they would travel on their own.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

Algorithmic Bias in Educational Predictive Analytics

G Governance & Trust **H** Human–System Interface **D** Designed Out

IMPACT Predictive “at-risk” models show racial calibration bias that can misdirect support away from Black and Latinx students; a large majority of U.S. public colleges now use predictive analytics

IN BRIEF

Most U.S. public colleges now use predictive analytics to flag “at-risk” students for early support. Research finds these models carry racial calibration bias: they miscalibrate by race in ways that can misclassify Black and Latinx students — and so misdirect the very support the prediction is meant to trigger. The magnitude depends heavily on how “at-risk” is defined, making this partly a construct-definition problem: the choice of what to predict is itself a capability decision with equity consequences. The models inherit historical patterns of discrimination from their training data, and a “flag” can confirm a biased instructor’s low expectations rather than prompt help. It is the educational analog of the UK A-level case — an algorithm built to do good that can allocate help away from the students who need it most.

BACKGROUND

Colleges increasingly use predictive analytics to identify students at risk of failing or dropping out, so advisors can intervene early — a genuinely well-meant aim, to reach struggling students before they vanish. A large majority of U.S. public colleges now use some form of these models, which makes any bias they carry a quiet, system-wide feature of how support is allocated rather than an isolated experiment.⁰

WHAT HAPPENED

The intervention is well-intentioned, but research finds the models carry racial calibration bias: they miscalibrate by race in ways that can misclassify Black and Latinx students relative to their actual outcomes — and so misdirect the support the flag is meant to trigger, turning a tool built to help into one that can steer help away. Crucially, the magnitude of the bias depends on how “at-risk” is defined, which makes it partly a construct-definition problem rather than a coding bug — a flaw in what the model was asked to predict, not in how it computes the prediction.¹

THE INVESTIGATION

Researchers studying equity in completion-prediction models have documented these calibration gaps and traced them to training data that encodes historical patterns of


discrimination — so the model learns the past’s inequities and projects them forward as predictions.² Baker and Hawn’s review documents how algorithmic bias in education poses significant threats to educational equity, potentially amplifying existing social and economic disparities — and the harm compounds when an instructor with deficit assumptions reads an “at-risk” flag as confirmation rather than a cue to help, letting the prediction become a self-fulfilling label.³

THE CAPABILITY GAP

The bias is not in the math; it is in the definition of “at-risk.” Define it one way and support flows to one population; define it another way and it flows to another, so the construct silently sets who the system decides to help. The choice of what the model predicts — the construct — is a capability-engineering decision with measurable equity consequences, and it is the part most often made implicitly, by whoever assembles the training labels, rather than governed deliberately by anyone accountable for where help is sent.⁴

AFTERMATH & REFORM

Unlike the discrete failures elsewhere in this chapter, this one is ongoing and quiet — embedded in advising dashboards at hundreds of institutions — which is what makes it dangerous: there is no single collapse to force a reckoning, only a steady misallocation no headline announces.⁵ Its lesson, pushed upstream, is the chapter’s: governing an algorithm’s fairness begins not at deployment but at construct definition and label choice, with an equity audit of what the model is asked to predict and for whom the prediction allocates help — catching the bias where it is introduced rather than where it surfaces.



Algorithmic bias in educational systems poses significant threats to educational equity, potentially amplifying existing social and economic disparities.

EDITORS’ PARAPHRASE OF BAKER & HAWN (2021)

REFERENCES

1. Surveys of predictive-analytics adoption in U.S. higher education — a large majority of public colleges now use some form.
2. K. Bird et al., “Are Algorithms Biased in Education? Exploring Racial Bias in Predicting Community College Student Success,” *Journal of Policy Analysis and Management* 44 (2025), 379–402 — racial calibration bias, 5× higher at the bottom decile depending on the “at-risk” construct.
3. D. Gándara, H. Anahideh, M. Ison & L. Picchiarini, “Inside the Black Box: Detecting and Mitigating Algorithmic Bias across
4. R. Baker & A. Hawn, “Algorithmic Bias in Education,” *International Journal of Artificial Intelligence in Education* (2021) — foundational review of algorithmic bias in education.
5. Analyses of deficit framing and the interpretation of “at-risk” flags by faculty.
6. Cf. UK A-Level / Ofqual (Case 49); V. Eubanks, *Automating Inequality* (2018).

Racialized Groups in College Student–Success Prediction,” AERA Open (2024) — bias traced to training data encoding historical discrimination.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Educational predictive analytics is the ongoing live case for algorithmic bias at the construct level. The bias is not in the sigmoid; it is in the definition of “at-risk.” Defining at-risk in one way allocates support to one population; defining it in another way allocates support to another. The choice of definition is a capability–engineering decision with measurable equity consequences.

LENS APPROACH

LENS treats this case as the positive counterpart to Georgia State (Case 80). LEN 4 examines construct definition as the load-bearing measurement choice. LEN 7 examines the governance architecture that determines whose construct gets adopted. LEN 9 covers the technical bias-mitigation methods.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
EDUCATION AT SCALE	→ G H D	→ D1 D2 D3 D4 D5	5	8.1	4
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Govern the construct definition deliberately: have an accountable owner decide and document what “at-risk” means and whom the prediction will route help toward. ▶ Audit training data for encoded historical discrimination before fitting, so the model does not learn past inequities as future predictions. ▶ Test calibration across racial groups during development and treat disparate misclassification as a design defect, not a tolerable residual.	▶ Monitor deployed dashboards for the quiet, distributed misallocation that has no single failure event to announce it. ▶ Train advisors and instructors to read an “at-risk” flag as a cue to help rather than confirmation of a deficit assumption. ▶ Sustain a recurring equity audit of what each model predicts and for whom, since a bias embedded across hundreds of institutions persists until someone is tasked to find it.

REFLECTION QUESTIONS

1. Pick a predictive analytic in your institution. Reconstruct the construct definition behind it. What is the equity consequence of that definition?
2. Design the governance review that a new predictive model should pass before it allocates resources to or away from a population.
3. This failure is ongoing and quiet, with no single collapse to force a reckoning. What would it take to make a slow, distributed misallocation visible enough that someone with authority had to act on it?

WHO BUILDS THIS governance, ethics & measurement · human factors & interaction design · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs · task analysis · design prototyping

Algorithmic Bias in Automated Exam Proctoring

D Designed Out **N** Normalization of Deviance **K** Knowledge & Institutional Memory

IMPACT The first quantitative study of facial-detection bias in automated exam proctoring software found that students with darker skin tones and Black students were significantly more likely to be flagged for instructor review for potential cheating; at the race–sex intersection, women with the darkest skin tones were far more likely to be flagged than other groups

IN BRIEF

The COVID-era expansion of remote learning produced a rapid deployment of automated exam-proctoring software across higher education: computer-vision systems that monitor the student via webcam during an exam and flag suspicious behavior for instructor review. Yoder-Himes et al. (Frontiers in Education, 2022) ran the first quantitative study of facial-detection bias in this class of software. Students with darker skin tones and Black students were significantly more likely to be flagged for instructor review for potential cheating than students with lighter skin tones; at the race–sex intersection, women with the darkest skin tones were far more likely to be flagged than other groups. The study examined one major proctoring product and concludes the product “may employ biased AI algorithms that unfairly disadvantage students.” It documents the disparity but not a remediation, so it is a failure / diagnosis case. The harm comes directly from validation that did not stratify across skin tone, surfacing a group-specific failure in a deployed system. Trio with Cases 25 (eGFR), 26 (pulse oximetry), and 6 (Hoffman pain bias) at the race–construct–and–validation layer.

BACKGROUND

Automated exam proctoring is a recent and rapidly deployed class of ed-tech: a webcam-based monitoring system that observes the student during a remote-administered exam and uses computer vision to flag suspicious behavior — looking off-camera, multiple faces in the frame, the face leaving the frame — for instructor review. The COVID-era pivot to remote instruction expanded the deployment of this software across US higher education at speed. The systems’ face detection and face tracking are the load-bearing computer-vision components: a failure in face detection translates directly into a false flag for instructor review, and systematic failures in face detection across demographic groups translate into systematic group-level harm.^o

WHAT HAPPENED

Yoder-Himes et al. (Frontiers in Education, 2022) ran the first published quantitative study of facial-detection bias in this class of software. The study examined one major proctoring product across a sample of student-submitted sessions, with skin-tone classification on a standard scale and demographic information available. The headline finding ran in the direction the broader face-recognition-bias literature had documented (Buolamwini and Gebru, Raji, and others) and extended that finding into the assessment context. Students with darker skin tones, and Black students specifically, were significantly more likely to be flagged for instructor review for potential cheating than students with lighter skin tones.¹

THE INVESTIGATION

The intersectional analysis sharpened the finding. Women with the darkest skin tones were far more likely to be flagged than other groups — the race-by-sex intersection produced the largest disparity, consistent with the broader face-recognition literature's intersectional findings. The harm class is not abstract: a flag for instructor review under an academic-integrity process produces real downstream consequence — the student has to defend themselves against a suspicion the software generated, and the institutional resolution mechanism is not designed for the case where the suspicion was generated by a software bias rather than by a student behavior.²

THE CAPABILITY GAP

The study's claim is calibrated and direct: the product "may employ biased AI algorithms that unfairly disadvantage students." The study documents the disparity in a real deployment of a real product and does not document a remediation, so it is a failure / diagnosis case rather than a failure-to-intervention arc. The harm comes directly from validation that did not stratify across skin tone: the computer-vision face-detection model behind the proctoring system was deployed without demographic validation of its detection rates, and the group-specific failure was therefore present in the deployed system from day one and only surfaced post-hoc by external researchers.³

AFTERMATH & REFORM

Drafted alongside the race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman pain bias), the proctoring case extends the validation-must-be- demographically-stratified lesson into the ed-tech and assessment context. The structural form is the same as pulse oximetry: a deployed system measured on an unrepresentative sample, producing aggregate accuracy that conceals a group-specific failure for years, until external researchers stratify the validation post-hoc. The eGFR cross-reference is the construct-definition counterpart; the Hoffman cross-reference is the human-judgment counterpart. All four sit in the small-and-big-tier conversation about validation discipline as an equity design commitment.

The product may employ biased AI algorithms that unfairly disadvantage students.

YODER-HIMES ET AL., FRONTIERS IN EDUCATION 2022.

REFERENCES

1. Yoder-Himes, D. R., Asif, A., Kinney, K., Brandt, T. J., Cecil, R. E., Himes, P. R., Cashon, C., Hopp, R. M. P., & Ross, E. (2022). Racial, skin tone, and sex disparities in automated proctor-
3. Raji, I. D., & Buolamwini, J. (2019). Actionable auditing: Investigating the impact of publicly naming biased performance results of commercial AI products. AAAI/ACM Conference on AI, Ethics,

ing software. *Frontiers in Education*, 7:881449. doi:10.3389/educ.2022.881449 — the case’s primary study.

2.

Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. *Proceedings of Machine Learning Research*, 81:77–91 — the foundational intersectional-bias finding in face recognition.

4.

Sjoding, M. W., Dickson, R. P., Iwashyna, T. J., Gay, S. E., & Valley, T. S. (2020). Racial bias in pulse oximetry measurement. *New England Journal of Medicine*, 383(25):2477–2478 — the structural analog in the medical-device context (Case 26).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Automated exam proctoring is the assessment-domain counterpart to pulse oximetry’s medical-device bias. Yoder–Himes et al. 2022 is the first quantitative study of the disparity in this class of software, with the intersectional finding the broader face-recognition literature predicts. The case documents the disparity in one product; the remediation is not yet documented.

LENS APPROACH

Proctoring bias is the small-tier ed-tech validation-stratification failure (induced 8.2; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the validation-must-be-demographically-stratified discipline, and in Domain 5 (Navigating Sociotechnical Constraints) for the institutional resolution mechanism that has to handle the case where the flag was generated by software bias. Trio with Cases 25 (eGFR), 26 (pulse oximetry), and 6 (Hoffman pain bias) at the race-construct-and-validation layer.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ED-TECH	D N K	D1 D2 D3 D4 D5	5	8.2	4 · 5
ALGO FAIRNESS		Test and Evaluation			
ASSESSMENT					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Stratify validation across skin tone (and across the race–sex intersection) before deployment, not after, so the group-specific failure surfaces in the engineering record before it surfaces in the harm record.
- ▶ Specify the harm class the false-positive flag produces — academic-integrity defense, downstream consequence to the student — and design the institutional resolution mechanism to handle the case where the flag was generated by software bias.
- ▶ Require the proctoring vendor to publish demographic stratification of their detection rates, on a standard skin-tone scale, as a deployment condition rather than as a post-hoc disclosure.

IN OPERATION / AFTER

- ▶ Treat the Yoder–Himes finding as the diagnosis it is: the disparity is documented in one product; the remediation is not.
- ▶ Carry the intersectional reading explicitly (women with the darkest skin tones far more likely to be flagged) in any communication about the case; the broader race-by-sex intersectional finding is consistent with the face-recognition-bias literature and should not be smoothed into a single-dimension finding.
- ▶ Build the cross-case reading with eGFR (Case 25), pulse oximetry (Case 26), and Hoffman pain bias (Case 6): the validation-must-be-demographically-stratified lesson runs across clinical, device, and assessment domains, and the proctoring case is the assessment-domain anchor.

REFLECTION QUESTIONS

1. Identify a deployed system in your domain whose validation rests on an aggregate accuracy figure rather than a demographically stratified one. What would the stratified validation actually require, and who would have to commission it before deployment rather than after?
2. The Yoder–Himes finding is intersectional: women with the darkest skin tones were far more likely to be flagged. What is the analog intersectional structure of the harm in your domain, and is it visible in the engineering record before the harm record?
3. The case documents the disparity in one product; the remediation is not yet documented. What would a remediation look like — vendor disclosure, regulatory disclosure requirement, institutional resolution–mechanism redesign — and which of those is in your scope?

WHO BUILDS THIS systems & human–factors engineering · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · incident analysis · safety dashboards · knowledge bases · data pipelines

CASE 48

K-12 EDUCATION

SCHOOL SAFETY INFRASTRUCTURE

RACIAL DISPARITIES

2010S – 2022

School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism

G Governance & Trust**K** Knowledge & Institutional Memory**N** Normalization of Deviance

DISCLOSURE Institutional overlap: an editor shares an institution (Johns Hopkins School of Engineering) with an author of this work; no editor was personally involved. Included on the published peer-reviewed evidence (Journal of Criminal Justice, 2022); editorial framing keeps critical distance.

IMPACT Peer-reviewed study (Journal of Criminal Justice, 2022) examining school-surveillance infrastructure and Black student outcomes; the mechanism is the surveillance infrastructure (cameras, metal detectors, school resource officers), not the students, and the infrastructure drove a measured share of the outcome gap — one of the motivating cases for the CLO Gap attribution

IN BRIEF

Johnson and colleagues, writing in the Journal of Criminal Justice in 2022, analyzed the relationship between school-surveillance infrastructure — cameras, metal detectors, school resource officers, ID-check protocols — and outcomes for Black students across US schools. The study's learning-engineering content is in where it locates the mechanism. The disparity in outcomes between Black students and white students under surveillance-heavy schooling cannot be explained, the analysis shows, by student behavior alone; a measured share of the outcome gap is attributable to the surveillance infrastructure itself — to the institutional posture that schools serving predominantly Black student populations more often adopted. The mechanism, in other words, is the infrastructure, not the students. The case is one of the motivating cases in the v2 sweep for the CLO **Gap attribution** — the discipline of asking, when a disparity in outcomes is observed, what share of the disparity is attributable to the institutional or technical infrastructure rather than to the population the infrastructure is operating on. The standing COI rendered under the title is binding: an editor of this volume shares an institution with the study's authors, and the case is anchored to the peer-reviewed Journal of Criminal Justice evidence rather than to institutional press. The case extends the race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman) into the K-12 education domain at the infrastructure layer.

BACKGROUND

US public schools have, over the past two decades, deployed a progressively more elaborate surveillance infrastructure: cameras throughout common areas; metal detectors at entry points; school resource officers (sworn law enforcement stationed in schools); ID-check protocols and visitor management systems; behavioral-tracking software. The distribution

of that infrastructure across schools has not been uniform. Schools serving predominantly Black student populations have, on average, carried more of it. The institutional rationale has typically been safety — the surveillance infrastructure is described as keeping students safe — and the question the Johnson study confronts is whether the surveillance is operating as safety infrastructure or as something else, measured by the outcomes the infrastructure produces.⁰

WHAT HAPPENED

Johnson and colleagues (Journal of Criminal Justice, 2022) analyzed the relationship between school-surveillance infrastructure and outcomes for Black students across the US public-school sector. The study's design separates student-level variables from school-level infrastructure-deployment variables, so that the share of the outcome gap attributable to each can be estimated. The headline analytic result the case rests on: a measured share of the disparity in outcomes between Black students and white students under surveillance-heavy schooling is attributable to the surveillance infrastructure itself, not to differences in student behavior. The infrastructure is acting as a driver of the outcome gap, not only as a response to it.¹

THE INVESTIGATION

The mechanism is what makes the case the canonical gap-attribution case. When a disparity in outcomes is observed, the analytical default in many institutional settings is to attribute the disparity to the population the outcomes are measured on — to differences in preparation, behavior, family context. The Johnson study shows that, in the school-surveillance case, that default is wrong in a measurable sense: the institutional infrastructure is itself a driver of the disparity, and attributing the outcome gap to the students rather than to the infrastructure mis-locates the mechanism in a way the data does not support. The case extends the race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman) into K-12 education at the infrastructure layer: in eGFR and pulse oximetry the mechanism was in the device or the formula; in Hoffman the mechanism was in the practitioner's cognitive baseline; here the mechanism is in the institutional architecture itself.²

THE CAPABILITY GAP

The standing COI rendered under the title is binding. An editor of this volume shares an institution (Johns Hopkins School of Engineering) with the study's authors. The case is anchored to the published peer-reviewed evidence in the Journal of Criminal Justice, not to institutional press or to commentary by the editor or the authors outside the paper. The editorial framing has been written to maintain critical distance from the home-institution affiliation; the case's claim is what the published evidence supports, and the disclosure is the safeguard that makes the institutional adjacency visible at the point of reading.³

AFTERMATH & REFORM

The case is one of the motivating cases in the v2 sweep for the CLO **Gap attribution** — the discipline of asking, when a disparity in outcomes is observed, what share of the disparity is attributable to the institutional or technical infrastructure rather than to the population the infrastructure is operating on. The race-construct trio established the pattern at the device/formula/cognitive- baseline layers; this case carries the pattern at the institutional-infrastructure layer. The four-case set (Cases 25, 26, 6, 48) is the case-grounded basis for

Gap attribution as a designed competency: practitioners and program designers have to be trained to look for the mechanism in the infrastructure they built, not only in the population they are serving, and the evidence the case carries is what makes the proposed CLO defensible.

The mechanism is the infrastructure, not the students. Attributing the outcome gap to the population mis-locates the mechanism in a measurable sense the data does not support.

EDITORS' SYNTHESIS OF JOHNSON ET AL. (JOURNAL OF CRIMINAL JUSTICE, 2022).

REFERENCES

1. Johnson and colleagues (2022), school-surveillance infrastructure and Black student outcomes, Journal of Criminal Justice — the load-bearing peer-reviewed source for the case.
2. Hoffman, Trawalter, Axt, & Oliver (2016), "Racial bias in pain assessment and treatment recommendations, and false beliefs about biological differences between blacks and whites," PNAS 113(16):4296–4301 — race-construct trio at the cognitive-baseline layer (Case 6).
3. Inker, Eneanya, Coresh, et al. (2021), "New Creatinine- and Cystatin C–Based Equations to Estimate GFR without Race," NEJM — race-construct trio at the formula layer (Case 25).
4. Sjoding, Dickson, Iwashyna, Gay, & Valley (2020), "Racial Bias in Pulse Oximetry Measurement," NEJM 383:2477–2478 — race-construct trio at the device layer (Case 26).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Johnson school-surveillance study locates the mechanism of an outcome disparity in the institutional infrastructure, not in the population the infrastructure operates on. The case extends the race-construct trio (eGFR, pulse oximetry, Hoffman) into K-12 education at the infrastructure layer and is one of the motivating cases for the CLO Gap attribution. COI under the title — shared institution — is binding and rendered openly.

LENS APPROACH

Johnson school surveillance is the infrastructure-as-mechanism gap-attribution case (induced 8.1; LENS D4/PT5) — Domain 4 as the case-grounded basis for **Gap attribution**; Domain 5 for the institutional-architecture-as-mechanism framing. Pair with Cases 25, 26, 6. COI binds.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
K-12 EDUCATION	G K N	D1 D2 D3 D4 D5	5	8.1	4
SCHOOL SAFETY INFRASTRUCTURE	→	→			
RACIAL DISPARITIES		Test and Evaluation			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the analysis to separate population-level variables from infrastructure-level variables so the share of the disparity attributable to each can be estimated; the Johnson study's design is the worked example.
- ▶ Treat the institutional infrastructure as a candidate mechanism, not as a neutral background; the default attribution-to-population analysis cannot test the infrastructure if the infrastructure is not a variable in the model.

IN OPERATION / AFTER

- ▶ Carry the Gap-attribution discipline into curriculum design: practitioners have to be trained to look for the mechanism in the infrastructure they built, not only in the population they are serving.
- ▶ Pair the case with the race-construct trio (Cases 25, 26, 6) so the gap-attribution pattern is taught at multiple layers — device, formula, cognitive baseline, institutional architecture.
- ▶ Preserve the COI render — shared institution, anchored to peer-reviewed evidence, critical editorial

- ▶ Anchor the editorial framing to the peer-reviewed evidence rather than to institutional press; the COI is rendered openly under the title and the framing maintains critical distance.

distance — as the standing language for home-institution-shared cases across the corpus.

REFLECTION QUESTIONS

1. Identify a disparity in outcomes in your domain that is currently attributed primarily to the population the outcomes are measured on. What infrastructure-level variables would you have to add to the analysis to test whether the institutional architecture is itself a mechanism of the disparity?
2. Specify the design pattern the Johnson study uses: separating student-level from school-level variables so the share of the disparity attributable to each can be estimated. What is the analog in your context, and where is the default analysis most at risk of mis-locating the mechanism?
3. The case is one of the motivating cases for the CLO Gap attribution. What instance from your work — a device, a formula, a cognitive baseline, an institutional architecture — would you carry as the case-grounded basis for training practitioners in your context to look for the mechanism in the infrastructure they built?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction

D Designed Out **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Wisconsin Department of Public Instruction's Dropout Early Warning System (DEWS) deployed since 2012, producing risk scores for approximately 200,000 sixth- through ninth-grade students per year; Perdomo, Britton, Hardt, & Abebe FAccT 2025 regression-discontinuity analysis across ~10 years of data found the analysis cannot rule out zero treatment effect on graduation; DPI's own 2021 internal equity audit and The Markup's 2023 investigation found the system was less accurate for Black and Hispanic students

IN BRIEF

The Wisconsin Department of Public Instruction has operated the Dropout Early Warning System (DEWS) since 2012, producing dropout-risk scores for approximately 200,000 sixth- through ninth-grade students across the state each year. Two evidence streams converge on the deployment. Perdomo, Britton, Hardt, and Abebe's 2025 FAccT paper used a regression-discontinuity design on approximately ten years of DEWS-and-graduation data and found that the analysis cannot rule out zero treatment effect on graduation outcomes for students above versus below the DEWS risk threshold — the prediction did not appear to trigger interventions that affected the predicted outcome. The Wisconsin DPI's own 2021 internal equity audit, titled "Is DEWS Fair?", found that the system was less accurate for Black and Hispanic students, and the agency continued to operate it unchanged. The Markup's 2023 investigation by Todd Feathers documented the disparate-impact finding and the agency's response. Both findings are load-bearing. The case pairs with Case 48 (Johnson school surveillance), Case 86 (Gándara community-college predictive equity), and Case 52 (Purdue Course Signals — reverse causality). This case carries both the peer-reviewed and journalism-investigation evidence streams in one entry rather than parallel ones.

BACKGROUND

The Wisconsin Dropout Early Warning System was developed by the Wisconsin Department of Public Instruction in collaboration with researchers from the University of Wisconsin — Madison and deployed in 2012. The system ingests student-level administrative data — attendance, course performance, behavioral incidents, mobility, and demographics — and produces a dropout-risk score for each student in grades six through nine. The score is surfaced to school administrators, counselors, and teachers, with the operational theory that the early warning enables school-level interventions — additional counseling, tutoring,

mentoring, or family contact — that improve the student’s graduation trajectory. The system has operated at scale for more than a decade and produces risk scores for approximately 200,000 students each year.⁹

WHAT HAPPENED

Perdomo, Britton, Hardt, and Abebe’s 2025 paper at FAccT (also available as arXiv 2304.06205) is the load-bearing causal-inference evidence on the deployment. The researchers used a regression-discontinuity design, exploiting the DEWS risk-threshold cutoff to compare students just above the threshold (classified as high-risk and surfaced to the school’s intervention surface) against students just below it (not so surfaced), controlling for the continuous risk score itself. The analysis covered approximately ten years of DEWS data and the corresponding graduation outcomes. The headline finding is that the analysis cannot rule out zero treatment effect of being above the DEWS threshold on subsequent graduation. The confidence interval on the treatment effect includes zero, and the point estimate is small enough that the prediction’s operational role — triggering school-level interventions that would change the predicted outcome — is not evidenced in the ten-year record. The hedge the paper preserves is binding: the RDD analysis cannot establish a negative finding either (the interval includes small positive effects), but it can and does establish that the prediction-triggers-intervention-that-changes-outcome theory is not supported in the deployment record.¹

THE INVESTIGATION

The Wisconsin DPI’s own 2021 internal equity audit, titled “Is DEWS Fair?”, is the load-bearing disparate-impact evidence on the deployment. The audit found that DEWS was less accurate for Black and Hispanic students than for white students — that the false-positive and false-negative rates diverged across student-population subgroups in ways the audit characterized as inconsistent with fairness. The agency’s own response to its own audit was to continue operating DEWS unchanged. The Markup’s 2023 investigation by Todd Feathers documented both the finding and the agency’s response, and the journalism-investigation evidence stream entered the public record in 2023. The case carries both evidence streams — the peer-reviewed causal-inference null and the agency-and-investigation disparate-impact finding — in one entry, rather than as parallel cases, because the deployment is one deployment and the two streams describe complementary structural problems with it.²

THE CAPABILITY GAP

The case pairs with Case 48 (Johnson school surveillance) for the algorithmic-public-administration-in-education parallel; both cases involve administrative-data predictions deployed against student populations and both surface disparate-impact concerns at the deployment surface. Pair with Case 86 (Gándara community-college predictive equity) for the predictive-equity-in-education thread at adjacent population scale. Pair with Case 52 (Purdue Course Signals) for the reverse-causality and null-effect thread in education predictive analytics; the Purdue case named the same conceptual problem that the Perdomo et al. analysis evidences at population scale. DEWS is the rare deployment in the corpus that is evidenced from both the peer-reviewed causal-inference direction and the disparate-impact-investigation direction, and the editorial decision to carry both streams in one entry reflects the structural unity of the deployment.³

AFTERMATH & REFORM

The hedges the case carries are load-bearing and both streams are preserved. The Perdomo et al. RDD analysis cannot rule out zero treatment effect on graduation across approximately ten years of data — the prediction–triggers–intervention–that–changes–outcome theory is not supported in the deployment record. The DPI’s own equity audit found that DEWS is less accurate for Black and Hispanic students, and the agency continued operating the system unchanged after the audit’s release. Both findings are binding and travel together. The CLO on designing predictions to trigger support rather than gatekeeping is anchored by the case at the deployment–with–null–causal–effect–and–disparate–accuracy seam: a prediction system operating at population scale for more than a decade, without evidence that the prediction triggers outcome–changing intervention and with evidence that the prediction’s accuracy varies across protected–attribute subgroups, is the deployment form whose persistence the case asks the reader to account for.

The RDD analysis cannot rule out zero treatment effect on graduation across approximately ten years of data; the agency’s own equity audit found DEWS is less accurate for Black and Hispanic students; the agency continued operating the system unchanged.

EDITORS’ SYNTHESIS OF PERDOMO ET AL. (2025, FACCT), THE WISCONSIN DPI INTERNAL EQUITY AUDIT (2021), AND THE MARKUP INVESTIGATION (FEATHERS, 2023).

REFERENCES

1. Perdomo, J. C., Britton, T., Hardt, M., & Abebe, R. (2025), “Difficult Lessons on Social Prediction from Wisconsin Public Schools,” Proceedings of FAccT 2025, doi:10.1145/3715275.3732175 (also arXiv:2304.06205).
2. Wisconsin Department of Public Instruction (2021), Is DEWS Fair? — internal equity audit of the Dropout Early Warning System.
3. Feathers, T. (2023), “False Alarm: How Wisconsin Uses Race and Income to Label Students ‘High Risk,’” The Markup, April 27, 2023 — investigation documenting the disparate-impact finding and the agency response.
4. Knowles, J. E. (2015), “Of Needles and Haystacks: Building an Accurate Statewide Dropout Early Warning System in Wisconsin,” Journal of Educational Data Mining 7(3):18–67 — the original DEWS technical-methodology paper.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Wisconsin DEWS is the load-bearing case for a prediction system operating at population scale for more than a decade without evidence that the prediction triggers outcome–changing intervention and with evidence that the prediction’s accuracy varies across protected–attribute subgroups. Both the peer-reviewed null and the agency–audit disparate-impact finding travel together; the case carries both evidence streams in one entry rather than parallel ones.

LENS APPROACH

Wisconsin DEWS is the designing–predictions–to–trigger–support–not–gatekeeping case at population scale (induced 8.3; LENS D4+D5/PT6; CLO–4 and CLO–5). LENS uses it in Domain 4 (Test and Evaluation) for the causal-inference–on–multi–year–deployment discipline and in Domain 5 (Navigating Sociotechnical Constraints) for the equity–audit–as–binding–input anchor. Pair with Case 48 (Johnson school surveillance), Case 86 (Gándara community-college predictive equity), and Case 52 (Purdue Course Signals reverse causality). Both the peer-reviewed null and the journalism–and–agency–audit disparate-impact finding are load-bearing.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
EDUCATION AT SCALE	➔ D K N ➔	D1 D2 D3 D4 D5	6	8.3	4 · 5
PREDICTIVE ANALYTICS		Test and Evaluation + D5			
ALGORITHMIC EQUITY					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Pre-specify the causal-inference design that will evaluate whether the prediction triggers outcome-changing intervention; the Perdomo et al. RDD is the deployment-evidence form that the prediction-triggers-support theory requires for verification at multi-year horizons. ▶ Treat the agency's own equity audit as a binding governance input, not as an internal document; the DPI's "Is DEWS Fair?" audit was internally honest about the disparate-accuracy finding, and the case's load-bearing observation is the agency's continued operation of the system unchanged. ▶ Design predictions to trigger support, not gatekeeping; the case's central CLO anchor is the discipline of building the intervention surface that the prediction is meant to trigger, and of verifying — not assuming — that the prediction's operational role produces the outcome change the deployment theory requires.	▶ Carry both evidence streams into print as load-bearing and complementary; the case's pedagogical value depends on the peer-reviewed null and the agency-audit disparate-impact finding traveling together rather than being separated into parallel entries. ▶ Pair in syllabi with Case 48 (Johnson) and Case 52 (Purdue) so the algorithmic-public-administration-in-education and reverse-causality threads are taught alongside the multi-year-deployment-with-null-causal-effect finding. ▶ Use the case as the anchor for the designing-predictions-to-trigger-support CLO; the curricular target is the multi-year deployment record that demonstrates the gap between the prediction's operational theory and the evidence the deployment produces.

REFLECTION QUESTIONS

1. Identify a prediction system in your domain that has operated at scale for years without RDD-style causal evaluation of whether the prediction triggers outcome-changing intervention. What would the causal-evaluation design look like in your setting, and what would a null finding require of the deployment?
2. Specify the equity-audit-as-binding-input discipline you would apply when an internal audit finds disparate accuracy across protected-attribute subgroups. What is the documented decision rule for modifying or withdrawing the system, and what is the rule when the agency continues operation unchanged?
3. The case carries both the peer-reviewed null and the disparate-impact-investigation finding in one entry. Pick a deployment in your domain that has evidence streams in tension, and ask: what is the editorial discipline that carries the streams together rather than separating them, and what would separating them lose?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

Atlanta Public Schools Cheating Scandal

G Governance & Trust **N** Normalization of Deviance

IMPACT ~180 educators implicated; 35 indicted; 11 convicted under RICO statute; foundational U.S. education measurement-gaming case

IN BRIEF

Under a celebrated superintendent, Atlanta Public Schools was held up nationally as a high-performing urban district. The performance was substantially fabricated. A state special investigation, supported by the Georgia Bureau of Investigation, and reporting by the Atlanta Journal-Constitution found that for several years educators across dozens of schools had systematically erased and corrected students' answers on state standardized tests. The cheating was organized — principals pressured teachers, staff held "erasure parties" after testing days — and the incentive system rewarded the gaming with bonuses and promotions. Around 180 educators were implicated, 35 indicted, and 11 convicted under Georgia's racketeering statute. The capability gap lay in the measurement architecture: the institution being measured operated the instrument that measured it, with no independent audit.

BACKGROUND

During the 2000s, Atlanta Public Schools under superintendent Beverly Hall was celebrated nationally for rapid gains on Georgia's high-stakes standardized tests. Bonuses, public recognition, and job security were tied directly to those scores, and the district administered and reported the very tests on which it was being judged. With the reward attached to the number and the number produced in-house, the people under pressure to improve the score also controlled the answer sheets that determined it.⁹

WHAT HAPPENED

Over several years, educators at dozens of APS elementary and middle schools systematically changed students' answers after testing. The cheating was organized rather than incidental: principals pressured teachers to hit targets, and staff held after-hours "erasure parties" to correct answer sheets. Suspiciously high rates of wrong-to-right erasures flagged the pattern — a statistical fingerprint the gaming left behind precisely because correcting a wrong answer to a right one is far rarer, by chance, than the reverse.¹

THE INVESTIGATION

A 2009 Atlanta Journal-Constitution analysis of erasure data, followed by the state's Office of Student Achievement and a Governor-ordered special investigation supported by the Georgia Bureau of Investigation, documented the scheme in a 2011 report. Roughly 180 educators were implicated; 35 were indicted and 11 convicted under Georgia's RICO statute, the report finding the administration had emphasized results and praise "to the exclusion

of integrity and ethics.” That an investigation came from outside the district — a newspaper and state investigators, not the schools — is itself the mark of the missing independent check.²

THE CAPABILITY GAP

The capability gap was at the measurement architecture: the institution being measured also operated the instrument that measured it. High-stakes incentives without an independent audit are a textbook setup for Campbell’s Law — the more a measure is used for decision-making, the more it will be corrupted. The students were not learning; the system was rewarding the appearance of learning, so the score rose while the underlying capability it was meant to certify went unmeasured and, for the children, unmet.³

AFTERMATH & REFORM

The convictions made APS the most prominent U.S. case of high-stakes-testing fraud and fuelled a broader reassessment of accountability-testing regimes. The structural lesson — that a measurement used for consequential decisions needs an audit independent of the institution being measured — is direct, but is still rarely implemented at scale, because the independent audit costs money and political will that the high score, once reported, makes seem unnecessary.⁴

180

EDUCATORS IMPLICATED · DOZENS OF SCHOOLS

the institution being measured operated the instrument that measured it

ATLANTA PUBLIC SCHOOLS — MEASUREMENT GAMING UNDER HIGH-STAKES TESTING

Dr. Hall’s administration emphasized test results and public praise to the exclusion of integrity and ethics.

GEORGIA SPECIAL INVESTIGATORS’ REPORT ON ATLANTA PUBLIC SCHOOLS, 2011

REFERENCES

1. Special Investigators, Final Report on Atlanta Public Schools (2011) — the organized cheating and the quoted finding.
2. Atlanta Journal-Constitution investigative series (2009–2011) — the erasure-rate analysis.
3. State of Georgia v. Hall et al. (2013–2015) — indictments and RICO convictions.
4. Koretz, D. (2017), The Testing Charade — high-stakes-testing distortion.
5. Campbell, D. (1976), “Assessing the Impact of Planned Social Change” — Campbell’s Law.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The APS cheating scandal is the strongest available case for measurement-gaming under high-stakes incentives in education. The capability gap was at the measurement architecture: the institution being

measured operated the instrument that measured it. The reform pattern is direct — independent measurement audit — but rarely implemented at scale.

LENS APPROACH

APS is the protect-the-instrument-from-the-operator-who-controls-it case (induced 2.2): the educators measured on the test also administered and scored it, so the instrument was capturable by the very people it judged. LENS pairs it explicitly with VA wait-times (Case 7), a same-shape instrument-integrity failure where the unit being measured operated the measurement and quietly rewrote the number. The lesson is an instrument-integrity design choice — separate the measured from the measurer, and audit independently of the institution being judged — not the surface reading that people cheated under pressure.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
EDUCATION AT SCALE	→ G N	→ D1 D2 D3 D4 D5	3	2.2	4
GOVERNMENT		Test and Evaluation			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Separate test administration and scoring from the institution being judged, so the party with the incentive to inflate the number never controls the instrument that produces it.
- ▶ Engineer a statistical integrity check — such as wrong-to-right erasure-rate monitoring — into the measurement system before scores carry consequences.
- ▶ Anticipate Campbell's Law when attaching bonuses and job security to a metric, and design the audit as part of the incentive, not as an afterthought.

IN OPERATION / AFTER

- ▶ Run independent, ongoing audits of the high-stakes measure rather than relying on a newspaper or state investigators to surface fraud after years.
- ▶ Monitor for the appearance of learning diverging from independent evidence of learning, so the gap is caught while the score still means something.
- ▶ Sustain the audit's funding and authority against the complacency a high reported score creates, since the cost of checking looks unnecessary exactly when it is most needed.

REFLECTION QUESTIONS

1. Identify a high-stakes measurement in your domain whose audit is operated by the institution being measured. What would change with independent audit?
2. Apply Campbell's Law to a current high-stakes measurement system. What distortion is predicted?
3. The erasure-rate analysis caught APS only after the fact. Design a continuous integrity signal — a statistical fingerprint of gaming — that would flag the pattern as it emerged rather than years later.

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

CASE 52

HIGHER-ED ANALYTICS

EARLY-WARNING SYSTEMS

EVIDENCE ARCHITECTURE

2012 – 2013

Purdue Course Signals — The Reverse-Causality Retention Claim

D Designed Out **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Purdue's widely cited claim that students taking two or more Course Signals classes were 21% more likely to be retained was challenged by Mike Caulfield in 2013, who identified the dose-response curve as an artifact of selection: students persist and therefore take more Signals courses, not the reverse — Alfred Essa reproduced the apparent retention gain by substituting 'received a piece of chocolate' for 'took a Signals class' in a simulation

IN BRIEF

Purdue's Course Signals was one of the most-cited early-warning learning-analytics interventions of the early 2010s. The program's headline outcome claim — students who took two or more Signals courses were 21 percent more likely to be retained — was published by Arnold and Pistilli at LAK 2012 and circulated widely in the learning-analytics community and in vendor materials. In 2013 Mike Caulfield, blogging at e-Literate and in Inside Higher Ed, identified the dose-response curve as a reverse-causality artifact: students who persist in college therefore take more Signals courses, so the apparent retention gain reflects selection, not effect. Alfred Essa built a simulation that substituted "received a piece of chocolate" for "took a Signals class" and reproduced the apparent retention gain, demonstrating the methodological flaw. The original study was never peer-reviewed outside conference proceedings yet became one of the most-referenced learning-analytics studies of its era — which is itself the cautionary point about the field's evidence architecture. The case is the small-tier methodological companion to v1 Cases 46 (educational predictive-analytics bias) and 80 (Georgia State predictive analytics).

BACKGROUND

Course Signals at Purdue was a faculty-facing dashboard that classified students enrolled in a course as green, yellow, or red, based on a predictive model of academic risk built from LMS and grade-book signals. Faculty used the classification to send targeted communications to students flagged at risk. The design and the operational use were not the subject of the later critique; what became contested was the system's headline outcome claim, published by Arnold and Pistilli at Learning Analytics and Knowledge 2012: students who took two or more Signals courses were 21 percent more likely to be retained at the institution than students who took fewer.^o

WHAT HAPPENED

The figure circulated. It appeared in conference keynotes, vendor materials, accreditor presentations, and in widely shared accounts of what early-warning analytics could deliver. The claim's status in the literature outran its evidentiary base: the LAK 2012 paper was a conference paper, not a peer-reviewed journal article, and the institutional dataset behind the headline number was not made available for independent reanalysis. The field cited the result anyway, because it was the kind of result the field wanted to be true.¹

THE INVESTIGATION

In 2013 Mike Caulfield, writing at e-Literate and in Inside Higher Ed, asked a specific methodological question: was the dose-response curve — more Signals classes, more retention — what it appeared to be? Caulfield argued the relationship was a reverse-causality artifact. Students who persist at the institution have more semesters in which to take Signals courses; students who depart cannot. The “took two or more Signals courses” group was therefore an inadvertent selection on persistence — not a sample exposed to a different treatment intensity. Alfred Essa then built a simulation that substituted “received a piece of chocolate” for “took a Signals class,” with chocolate having no causal effect on anything, and reproduced the apparent retention gain. The reverse-causality reading survived the simulation replication.²

THE CAPABILITY GAP

The methodological point is precise: the published analysis did not isolate the treatment from the selection mechanism that determined treatment intensity, and the dose-response curve that looked like the effect was generated by the selection itself. The case is not an argument that Course Signals had no effect on retention. It is an argument that the published evidence could not distinguish effect from selection, and that the institution measured a number which **felt** like the failure mode it cared about (retention) using a design that could not actually answer the causal question. This is the textbook 2.1 failure: measuring the failure mode with a design the institution can deceive itself with.³

AFTERMATH & REFORM

Drafted as a deeper-evidence-of v1 Cases 46 and 80, the Purdue case carries a named methodological failure into the corpus's predictive-analytics conversation. The cautionary thread runs through three places at once: the original study's status (conference proceedings, never peer-reviewed outside that venue) outpaced its evidentiary base; the field's citation practice amplified the headline without probing the design; and the correction (Caulfield, Essa) was mounted from outside the original study's institutional network. The case teaches the evidence-architecture failure mode that the LENS Iterative Development domain and the Navigating Sociotechnical Constraints domain both have to protect against — and that v1 Cases 46 and 80 anchor at the bias and the institutional-deployment layers respectively.

Substituting “received a piece of chocolate” for “took a Signals class” in a simulation reproduces the apparent retention gain.

EDITORS' SYNTHESIS OF ESSA'S 2013 SIMULATION DEMONSTRATION.

REFERENCES

1.

Arnold, K. E., & Pistilli, M. D. (2012). Course Signals at Purdue: Using learning analytics to increase student success. *Proceedings of LAK 2012*, 267–270. doi:10.1145/2330601.2330666 — the original study at the center of the critique.

2.

Caulfield, M. (2013). A discussion of the Purdue Course Signals retention numbers. *e-Literate blog series and Inside Higher Ed commentary — the reverse-causality critique.*

3.

Essa, A. (2013). Simulation reproducing the Course Signals retention curve using a placebo treatment (“received a piece of chocolate”) — methodological demonstration of the artifact.

4.

Feldstein, M. (2013). What the Course Signals story says about learning-analytics evidence. *e-Literate — the broader field-level critique.*

THE LEARNING ENGINEERING LENS

LE INSIGHT

Course Signals is the named methodological failure in the predictive-analytics conversation: the institution measured a number that felt like the failure mode it cared about, using a design that could not isolate effect from selection. The original study was never peer-reviewed outside conference proceedings; the correction came from outside the original network. Both are part of the cautionary reading.

LENS APPROACH

Course Signals is the small-tier evidence-architecture failure (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for measurement design that cannot deceive the institution, and in Domain 5 (Navigating Sociotechnical Constraints) for the citation-practice failure mode in which a headline outcome outpaces the evidence it rests on. Deeper-evidence-of v1 Cases 46 (predictive-analytics bias) and 80 (Georgia State predictive analytics) — distinct because this is a named methodological-validity failure, not a bias finding.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER-ED ANALYTICS	D K N	D1 D2 D3 D4 D5	5	2.1	4 · 5
EARLY-WARNING SYSTEMS	→	→ Test and Evaluation			
EVIDENCE ARCHITECTURE					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the causal question the analysis is built to answer in advance, and design the comparison group so that the selection mechanism into “treatment intensity” is not itself the variable doing the work.
- Pre-register the dose-response analysis with the threats-to-validity table — selection, censoring, reverse causality — visible at design time, so the published headline cannot outrun the analysis it rests on.
- Choose the publication venue that fits the evidentiary claim: conference proceedings for an early result, peer-reviewed journal with independent reanalysis for the figure the field will cite.

IN OPERATION / AFTER

- Make the dataset available for independent reanalysis; the Caulfield / Essa correction succeeded because the published claim was simulatable on plausible data, not because the original dataset was inspected.
- Treat the simulation-replication finding as program evidence about the analysis design, not as a verdict on the intervention; Course Signals may have had an effect, and the published study could not detect or measure it.
- Carry the cautionary reading into the field’s citation practice: a headline outcome circulating widely without peer review outside conference proceedings is itself an evidence-architecture failure, separate from any specific methodological flaw.

REFLECTION QUESTIONS

1. Identify a predictive or early-warning analytics deployment in your domain whose published outcome claim circulates more widely than the peer-reviewed evidence supports. What threats to validity — selection, reverse causality, censoring — would a Caulfield-style external critic name first?
2. Specify the comparison-group design that would isolate effect from selection in your context. What pre-registered analysis plan, with simulated placebo treatment, would let the field check the claim before it circulates?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

Summit Learning / Personalized Learning Rollout

G Governance & Trust **T** Training Gap **K** Knowledge & Institutional Memory

IMPACT Personalized-learning platform deployed across ~380 U.S. schools; parent revolts in Brooklyn, Cheshire, McPherson; multiple districts withdrew within two years

IN BRIEF

Summit Learning, a personalized-learning platform from Summit Public Schools backed by the Chan Zuckerberg Initiative, was offered free to U.S. districts from 2015 and reached roughly 380 schools and 72,000 students by 2018. By 2019 prominent adopters were withdrawing under parent and student pressure — opt-out campaigns in Brooklyn, cancellations in Cheshire and McPherson — amid walkouts and complaints about screen time, disengagement, and data privacy. The pedagogy itself (competency-based progression, projects, mentoring) was defensible and often effective; what failed was deployment governance. There was no evaluation framework districts could read before adopting, no parent-facing data agreement, and no exit path independent of the vendor. Summit is a clean test of the book's claim that the governance architecture must be engineered alongside the tool.

BACKGROUND

Summit Learning was a personalized-learning platform developed by Summit Public Schools with technical and financial support from the Chan Zuckerberg Initiative, offered free to U.S. districts from 2015 — a price that lowered the bar to adoption while raising no governance questions at the door. Its pedagogy — competency-based progression, self-directed projects, mentor check-ins — was defensible and in many places effective, which is why the eventual revolt could not be blamed on the instructional design.⁰

WHAT HAPPENED

By 2018 the platform reached roughly 380 schools and an estimated 72,000 students, scaling fast on the strength of a free offer and a well-funded sponsor. By 2019 its most visible adopters were withdrawing under parent and student pressure: students at Brooklyn's Secondary School for Journalism walked out; districts in Cheshire and McPherson cancelled or scaled back after parent meetings where the unanswered governance questions surfaced all at once. Walkouts and complaints about screen time, eye strain, disengagement, and data privacy converged into a revolt that was not about the instructional design at all.¹

THE INVESTIGATION

Press coverage and later analyses located the failure in deployment governance, not pedagogy: there was no evaluation framework a district could read before adopting, no parent-facing data-handling agreement, and no exit pathway that did not depend on the vendor's goodwill — three absences that each became a grievance the moment families looked for them. The implementation never surfaced the governance questions parents would ask, so when those questions arrived they arrived as opposition rather than as design input, and the argument was lost before it started.²

THE CAPABILITY GAP

Summit is a clean test of the book's central claim: a technology that worked at the pedagogical level still failed because the governance architecture — consent, evidence, measurement, exit — had not been engineered alongside it. A working tool with no accountability contract is a liability waiting for the first organized objection. The pattern recurs across the ed-tech dataset (inBloom, Case 53): a well-intentioned tool, a well-funded rollout, and no institutional contract with the families and teachers operating inside it — the same omission producing the same collapse in a second case.³

AFTERMATH & REFORM

Several districts withdrew or rebranded their use, CZI and Summit revised their outreach, and the episode became a standard caution in ed-tech adoption — a reputational cost paid for governance work that would have been cheaper to do first.⁴ Its lesson for the field is concrete: an adoption decision should have to produce a public evidence summary at parent reading level, a data-handling agreement at the same resolution, and a documented exit path — governance artifacts that make a tool's deployment legitimate, not just its design sound, and that turn the questions parents will ask into inputs gathered before launch rather than weapons raised after it.

The tools were free. The accountability architecture had not been built.

EDITORS' SYNTHESIS OF SUMMIT LEARNING ROLLOUT COVERAGE (NEW YORK TIMES, WIRED, EDUCATION WEEK, 2018–2019)

REFERENCES

1. N. Bowles, "Silicon Valley Came to Kansas Schools. That Started a Rebellion," New York Times (2019) — the parent revolt.
2. N. Singer, "The Silicon Valley Billionaires Remaking America's Schools," New York Times (2017) — the CZI/Summit rollout.
3. B. Herold, Education Week coverage of Summit / CZI implementation (2019) — district adoptions and withdrawals.
4. A. Watters, "The Stories We Tell About Personalized Learning," Hack Education (2019) — the governance critique.
5. Chan Zuckerberg Initiative & Summit Learning public program documentation (2015–2019); cf. inBloom (Case 53).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Summit Learning is a clean test of the book's central claim: technology that worked at the pedagogical level still failed because the **governance** architecture (consent, evidence, measurement, exit) had not been engineered alongside it. The pattern — well-intentioned tool, well-funded rollout, no institutional contract with the families and teachers operating inside it — recurs across the educational-technology dataset (Cases 53, 139, 67) and is the educator's-side analog of the governance failures in Cases 49 and 191.

LENS APPROACH

LENS uses Summit Learning in LEN 7 as the foundational consent-and-evidence case for educational technology, and in LEN 10 as a studio prompt for the governance artifacts that any educational-technology adoption decision should produce: a public evidence summary at parent reading level, a data-handling agreement at the same resolution, and a documented exit pathway that does not depend on the vendor's goodwill.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
EDUCATION AT SCALE	→ G T K →	D1 D2 D3 D4 D5	4	5.4	5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Engineer the governance architecture — consent, evidence, measurement, exit — in lockstep with the pedagogy, so a sound tool ships with a sound accountability contract. ▶ Produce, before adoption, a public evidence summary and data-handling agreement at parent reading level that answer the questions families will raise. ▶ Build a documented exit pathway that does not depend on the vendor's goodwill, so a district can leave without being captured.	▶ Treat organized parent and student objections as governance signal about a missing contract, not as resistance to the instructional design. ▶ Monitor adopting districts for the early grievances — screen time, privacy, disengagement — that precede a withdrawal, and route them to a decision-maker. ▶ Maintain a re-adoption pathway so a withdrawn district can return only after completing the governance work that the first rollout skipped.

REFLECTION QUESTIONS

1.

What is the equivalent of the “free tool, free of governance” pattern in your domain — the offer that bypasses the accountability architecture because it does not yet exist?
2.

Design the parent-reading-level governance artifact that a district should require before adopting an educational-technology platform.
3.

Summit’s withdrawals in Brooklyn, Cheshire, and McPherson were led by parents, not regulators. What is the equivalent local constituency in your domain that institutional accountability has not yet accommodated, and how would they be heard before deployment rather than after?

WHO BUILDS THIS governance, ethics & measurement · learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · scenario simulation · competency models · knowledge bases · data pipelines

HIGHER EDUCATION

CASE 55

PREDICTIVE ANALYTICS

2010S – PRESENT (BROOKINGS SYNTHESIS 2021)

ACCESS PRICING

Enrollment-Algorithm Yield Optimization Across U.S. Higher Education

T Training Gap **K** Knowledge & Institutional Memory **N** Normalization of Deviance

DISCLOSURE Institutional overlap: John Katzman collaborates with an editor's institution (Johns Hopkins) but was not involved in this work. He is not associated with the specific vendors named here.

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT Vendor case studies report 23% yield increases (Washington), 33% net tuition increases with a 6-point cut to discount rate (EAB), and 173 additional freshmen without aid-budget increases (Othot); algorithms across at least seven major vendors price aid offers against each accepted applicant's modeled willingness to pay

IN BRIEF

Engler's 2021 Brookings paper documents the two-stage architecture of contemporary enrollment-management algorithms: predict each accepted applicant's probability of enrollment, then optimize the financial-aid offer to maximize either net tuition revenue or yield. He names the vendor landscape — Ruffalo Noel Levitz, EAB, Rapid Insight, Capture Higher Ed, Othot, Whiteboard, Civitas Learning — touching hundreds of U.S. institutions. Vendor-reported case studies cite large gains in yield or tuition revenue per matriculant. The structural critique is the inversion of Case 80 (Georgia State, where prediction triggered support): here, the algorithm identifies "willingness to pay" so the institution can offer the minimum scholarship that will still yield enrollment, reducing aid per low-income student. The honest hedges Engler preserves are binding: critical details are obscured by vendors and colleges; algorithmic optimization is hard to separate from manual leveraging; without auditing specific college data and models, fairness impacts on protected classes cannot be confirmed. The evidence-tier flag under the title carries the standing language; future validation ongoing.

BACKGROUND

The structural seam the case opens is that financial-aid distribution is itself a deployed prediction system, operating across hundreds of U.S. institutions, with measurable consequences for who attends college and how much they pay. Engler's Brookings synthesis is the most thorough public mapping of the deployment surface, drawing on vendor case studies

and the higher-education enrollment-management literature. The first stage of the two-stage architecture is prediction: estimate the probability that an accepted applicant will enroll, given observable attributes from the application, the FAFSA, and third-party data. The second stage is optimization: choose the financial-aid offer that maximizes a chosen objective — net tuition revenue per matriculant, total tuition revenue, or class size at a target discount rate.⁰

WHAT HAPPENED

The vendor landscape Engler names is large and concentrated. Ruffalo Noel Levitz works with roughly 300 institutions; EAB with 150; Rapid Insight with 150; Capture Higher Ed with 100; Othot with around 30; Whiteboard Higher Education and Civitas Learning round out the named tier. Vendor marketing reports characteristic effect sizes: the University of Washington's 23 percent yield gain; EAB's 33 percent net-tuition gain paired with a 6-point cut to the discount rate; Othot's 173 additional freshmen recruited without an increase to the aid budget. The vendor effect sizes are vendor-reported and are not auditable from outside the institution — a hedge Engler is explicit about, and one the evidence-tier flag preserves into the case.¹

THE INVESTIGATION

The structural critique runs through what the algorithm is optimized for. The chosen construct — willingness to pay — is the operational target the prediction system was built around, and Engler's argument is that this construct turns the financial-aid award into a pricing instrument rather than a need-or-merit one. The downstream evidence he marshals is the literature linking institutional aid to graduation outcomes: roughly a 0.9-point gain in graduation odds per additional \$1,000 in merit aid; a more than 5-point cut in low-income completion when unsubsidized loans substitute for grant aid. If the algorithm reduces aid per student to find the minimum that still yields enrollment, the downstream cost is the completion gap that grant aid was buying down.²

THE CAPABILITY GAP

The case sits as the structural inverse of Case 80 (Georgia State's predictive advising) and pairs with Case 186 (mortgage-lending fairness) and Case 86 (community-college predictive equity). Georgia State used prediction to trigger support; enrollment-management algorithms use prediction to reduce the help allocated. Bartlett's lending analysis names the same construct-substitution pattern at the pricing layer across a different deployed prediction system; Gándara's community-college work names it at the access layer. The anchor the three cases share is the inversion of the gatekeeping-vs-support frame: the prediction is used to gatekeep the help, and the gatekeeping is invisible at the applicant's end of the transaction.³

AFTERMATH & REFORM

The hedges Engler names are binding on the case's framing, and the evidence-tier flag's standing language — future validation ongoing — applies in a precise sense. Critical details of the optimization are obscured by vendor non-disclosure and college contracting confidentiality; algorithmic optimization is hard to distinguish from manual "leveraging" of the same logic by human enrollment officers; without auditing specific college data and models, the protected-class fairness questions cannot be answered with the precision Bartlett-class

auditing would require. The Brookings synthesis is the strongest public evidence available; it is policy-tier analysis built on vendor case studies and Engler's read of the operational record, and the case rests at that tier, not at the audited-deployment tier the corpus would prefer.

The algorithms excel at identifying a student's exact willingness to pay. The construct is the load-bearing fairness decision, and the institution that does not name it has nonetheless made it.

EDITORS' SYNTHESIS OF ENGLER (2021), BROOKINGS INSTITUTION.

REFERENCES

1. Engler, A. (2021), "Enrollment algorithms are contributing to the crises of higher education," Brookings Institution, 14 Sept 2021.
2. Bettinger, E. et al. (2019), "Merit aid and college completion," literature reviewed in Engler — graduation-odds effect sizes for \$1,000 of additional aid.
3. Goldrick-Rab, S. (2016), Paying the Price — broader synthesis on net-price, unsubsidized loans, and low-income completion.
4. Vendor case studies cited in Engler (Othot, EAB, Ruffalo Noel Levitz, University of Washington) — vendor-reported and not externally audited; flagged at evidence-tier under the title.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Enrollment-algorithm yield optimization is the construct-choice case at the pricing layer of higher-education access: prediction is used to reduce aid per applicant, not to trigger support, and the operational target is "willingness to pay." The evidence-tier flag is binding — vendor case studies are not auditable, the algorithmic-vs-manual distinction is bounded, and the protected-class fairness question requires audit access the public synthesis cannot supply. Future validation ongoing.

LENS APPROACH

Engler / enrollment-management is the construct-choice case at population scale (induced 8.3; LENS D5/PT5). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the construct-substitution anchor and the disclosure architecture the deployment lacks, and in Domain 4 (Test and Evaluation) for the evidence-tier discipline — practice-synthesis is the strongest available tier, and the case says so. Pair with Case 80 (Georgia State support-trigger inversion), Case 186 (Bartlett lending fairness), and Case 86 (Gándara community-college equity). coi-light render under the title is binding.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER EDUCATION	T K N	D1 D2 D3 D4 D5	5	8.3	4 · 5
PREDICTIVE ANALYTICS	→	Navigating Sociotechnical Constraints			
ACCESS PRICING					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Name the optimization construct explicitly. "Willingness to pay" is not "need" and is not "merit"; the choice of construct is the load-bearing fairness decision, and the institution that does not name it has nonetheless made it.
- ▶ Require vendor disclosure of the model's inputs, training data, and objective function as a condition of contracting; the case's evidence-tier limit is partly the

IN OPERATION / AFTER

- ▶ Commission an external audit of the deployed enrollment-management model against protected-class outcomes; the audit Engler says cannot be done from outside the institution is the audit institutions can choose to commission from inside.
- ▶ Publish the discount-rate-and-completion link as a paired metric; institutions that report only net-tuition

result of contractual non-disclosure that institutions could refuse to sign.

- ▶ Tie the deployed algorithm's outputs back to downstream completion data; the literature linking aid to graduation odds is the evidence base the optimization should be tested against, not abstracted from.

gains and not the completion consequences are publishing a partial scorecard.

- ▶ Treat the gatekeeping-vs-support inversion as a curricular pattern: pair this case in syllabi with Case 80 so the design choice — which direction the prediction points — is taught as the design choice, not as an institutional default.

REFLECTION QUESTIONS

1. Identify a deployed prediction system in your domain whose optimization target is named on the institutional side and obscure on the applicant or user side. What is the construct the optimization is built around — and what would change if the construct were named at the point of transaction?
2. Specify a vendor-disclosure clause you would require as a condition of contracting an enrollment-management or analogous optimization system. What inputs, training data, and objective function would the institution insist on auditing, and which would the vendor be willing to disclose under contract?
3. The case sits as the inversion of Case 80 (Georgia State, prediction to trigger support). Pick a prediction system in your domain and ask: in which direction does the prediction point — toward more help or less — and where is that design choice documented?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

GAO Online Program Manager Oversight Gap (GAO-22-104463)

D Designed Out **K** Knowledge & Institutional Memory **N** Normalization of Deviance

DISCLOSURE Institutional overlap: John Katzman collaborates with an editor's institution (Johns Hopkins) but was not involved in this work. Katzman founded 2U — central to the OPM debate, though not named in this GAO report.

IMPACT GAO found at least 550 colleges contracted with OPMs to support at least 2,900 academic programs as of July 2021; revenue-share arrangements typically gave the OPM 40–60% of tuition revenue per recruit, some up to 80%; ED lacked instructions for auditors to detect incentive-compensation violations and was not collecting the information it needed to oversee the arrangements

IN BRIEF

The Government Accountability Office's April 2022 audit (GAO-22-104463) names a structural oversight gap in the federal regime governing online program managers. As of July 2021, at least 550 colleges contracted with OPMs to support at least 2,900 academic programs; revenue-share arrangements typically gave the OPM 40 to 60 percent of tuition revenue per recruit, with some arrangements reaching 80 percent. The 1992 Higher Education Act amendments prohibited incentive compensation for student recruiters as a fraud-prevention measure; the Department of Education's 2011 "bundled-services" guidance exempted OPMs from the ban when recruiting was bundled with other services. The GAO found that colleges and their auditors lacked clear instructions to detect violations, and the Department was not collecting the information it needed to oversee the arrangements. 2U's July 2024 Chapter 11 filing closed one boundary of the policy debate at the commercial level; at the federal level the Department opened a review of the guidance but did not rescind it — the January 2025 Dear Colleague Letter reaffirmed the 2011 bundled-services guidance. The underlying delegation- and oversight problem persists for successor OPMs and revenue-share structures.

BACKGROUND

The federal regime against incentive compensation in higher-education recruiting dates to 1992 and was built to address a specific fraud pattern: recruiters paid by enrollment will enroll students the program cannot serve, because the recruiter's compensation is tied to the enrollment rather than to the student's outcome. The 1992 Higher Education Act amendments banned the practice for university-employee recruiters. The 2011 "bundled-services" guidance the Department of Education issued exempted online program managers

from the ban when student recruiting was bundled with other services such as program design, platform delivery, and student support — a construct that allowed the revenue-share contracting model to grow rapidly across the next decade.⁹

WHAT HAPPENED

The GAO's audit documents the scale the regime grew to. By July 2021, at least 550 colleges had contracted with OPMs to support at least 2,900 academic programs. Revenue-share arrangements typically transferred 40 to 60 percent of tuition revenue per recruit to the OPM; some arrangements ran to 80 percent. The OPM operated under the university brand — recruiting, marketing, and program operations conducted by OPM employees identifying as the university — while receiving compensation tied directly to the enrollments those operations generated. The structural seam the regime had created was an exemption that allowed the prohibited compensation structure as long as the structure was administered by a contracted vendor.¹

THE INVESTIGATION

The oversight gap the GAO documents is operational. Colleges and their auditors lacked clear instructions to detect incentive-compensation violations within the bundled-services exemption; the Department was not collecting the information — contract structures, revenue-share rates, recruiter-compensation arrangements — that would have let it monitor whether the exemption's boundaries were holding. The audit's central finding is not that the OPM regime was designed to fail; it is that the oversight architecture required to police the exemption's boundaries was never built, and the contracting structure that the exemption allowed grew faster than the monitoring capacity that would have surfaced violations.²

THE CAPABILITY GAP

2U's July 2024 Chapter 11 filing closed one boundary of the policy debate at the commercial level; at the federal level the Department opened a review of the bundled-services guidance in 2023–2024 but did not rescind it, and its January 2025 Dear Colleague Letter reaffirmed the 2011 guidance. Neither closed the underlying delegation-and-oversight problem. Successor OPMs and revenue-share structures continue to operate; the universities that delegated student recruitment under the pre-2024 regime retain the operational dependencies and the brand-and-program commitments built during the decade of growth. The pair with Case 58 (USC × 2U Luna class action) shows the consumer-side litigation half of the same delegation; the pair with Case 55 (Engler / enrollment algorithms) shows the pricing-side optimization half.³

AFTERMATH & REFORM

The case is investigation-grade — a GAO audit is the strongest tier of evidence the corpus carries for regulatory-oversight failure — and it is the structural delegation-with-revocation case at population scale. The university delegated student recruitment to a contracted vendor under a regulatory exemption that did not include the monitoring architecture the exemption's boundaries required. The delegation was reversible in principle: the contracts could be terminated, the exemption could be withdrawn, the operations could be brought back inside the university. In practice, the delegation accumulated operational and contractual lock-in across the decade the GAO audit covers, and the revocation when it came in

2024 was forced by commercial collapse rather than by the oversight architecture the audit recommended.

Colleges and their auditors lacked clear instructions to detect violations, and the Department was not collecting the information it needed to oversee the arrangements.

U.S. GOVERNMENT ACCOUNTABILITY OFFICE, GAO-22-104463 (2022).

REFERENCES

1. U.S. Government Accountability Office (2022), “Higher Education: Education Needs to Strengthen Its Approach to Monitoring Colleges’ Arrangements with Online Program Managers,” GAO-22-104463.

2. U.S. Department of Education (2011), “Dear Colleague” guidance on incentive-compensation bundled-services exemption — the regulatory artifact the GAO audits.

3. Higher Education Act of 1992, incentive-compensation prohibition — the statutory basis the 2011 guidance interpreted.

4. 2U Inc. (2024), Chapter 11 bankruptcy filing (July 25, 2024); U.S. Department of Education, “Dear Colleague” Letter (Jan. 14, 2025) reaffirming the 2011 bundled-services guidance (GEN-11-05) — the commercial-boundary closure alongside the federal guidance’s continued validity.

THE LEARNING ENGINEERING LENS

LE INSIGHT

GAO-22-104463 is the investigation-grade delegation-with- revocation case at population scale. A regulatory exemption permitted a contracting structure across 550+ colleges and 2,900+ programs; the monitoring architecture required to police the exemption’s boundaries was never built. 2U’s July 2024 Chapter 11 closed the commercial boundary; ED opened a review but did not rescind the guidance (reaffirmed in the January 2025 DCL); the underlying delegation problem persists for successor OPMs.

LENS APPROACH

GAO OPM oversight gap is the regulatory-seam case (induced 5.3; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the cross-regime delegation pattern and in Domain 3 (Human-System Collaboration) for the delegation-with-revocation frame — the contract is reversible in principle, locked-in in practice. Pair with Case 58 (USC × 2U Luna, the litigation half) and Case 55 (Engler / enrollment algorithms, the pricing-optimization half). coi-light render under the title is binding.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER EDUCATION	D K N	D1 D2 D3 D4 D5	6	5.3	5 · 3
REGULATORY OVERSIGHT	→	→ Navigating Sociotechnical Constraints			
ONLINE EDUCATION					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Build the monitoring architecture as a condition of any regulatory exemption; the bundled-services exemption created a foreseeable contracting structure, and the oversight infrastructure to police its boundaries should have been built with it.
- ▶ Require contract-disclosure as a federal reporting obligation, not as an institutional discretion; the GAO’s data-collection finding is operationally addressable by mandatory reporting of revenue-share rates and recruiter-compensation arrangements.
- ▶ Treat the delegation as reversible from the start: contract terms should preserve termination rights and

IN OPERATION / AFTER

- ▶ Carry the investigation-grade audit into the curriculum without softening: the GAO’s central finding is that the oversight architecture was not built, and that is the load-bearing teaching point.
- ▶ Pair the case with Case 58 (USC × 2U) so the regulator-side audit and the consumer-side litigation are taught together; one half names what the regulator missed, the other names what the delegated marketing actually did.
- ▶ Track post-review and post-2U-bankruptcy successor structures as a continuation of the case; the underlying delegation problem persists, and the case’s

operational-knowledge transfer; the lock-in the universities experienced was partly contractual and partly operational, and both halves are designable.

frame is the regime-level oversight gap, not the specific 2U arrangement.

REFLECTION QUESTIONS

1. Identify a regulatory exemption in your domain that permits a contracting structure without specifying the monitoring architecture required to police its boundaries. What would the audit-detectable artifact be — and is anyone currently collecting it?
2. Specify the contract terms you would require to preserve revocability of a delegated capability. Termination rights, operational-knowledge transfer, data portability — which of these are typically written into your domain's standard contracts, and which are left to negotiation?
3. GAO-22-104463 documents what happened when oversight architecture lagged contracting growth across a decade. What is the analog in your domain — a contracting pattern that grew under an exemption whose monitoring infrastructure was not built — and where would the equivalent investigation-grade audit have to come from?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

CASE 58

HIGHER EDUCATION

ONLINE PROGRAM MANAGEMENT

PROFESSIONAL LICENSURE

2010S – 2024

USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC)

D Designed Out **K** Knowledge & Institutional Memory **N** Normalization of Deviance

DISCLOSURE Institutional overlap: John Katzman collaborates with an editor's institution (Johns Hopkins) but was not involved in this work. Katzman founded 2U but had departed before the USC MSW expansion documented [here](#).

EVIDENCE TIER journalism-tier — source confidence flagged; future validation ongoing.

IMPACT USC's online MSW grew from ~300 students per cohort pre-2010 to >3,000 per cohort through the 2U partnership; 2023 class action alleges USC marketed the online program as the 'same' as the residential program while outsourcing recruiting, advising, and clinical placement to 2U employees; partnership wind-down announced November 2023

IN BRIEF

USC's online Master of Social Work program grew from about 300 students per cohort before 2010 to more than 3,000 students per cohort, almost entirely through the 2U partnership. The May 2023 Luna v. USC class-action complaint alleges that USC marketed the online program as the "same" as the residential program while outsourcing recruiting, advising, and clinical placement to 2U employees who carried usc.edu email addresses. Plaintiffs further allege that 2U marketers targeted students of color and veterans with aggressive, deceptive tactics, and that the online program's tuition (over \$100,000) reflected on-campus pricing while the delivered student experience did not. USC and 2U announced the partnership's termination on the MSW and other programs in November 2023. The case is in litigation, and the predatory-targeting reconstruction rests on the complaint and on contemporaneous reporting rather than on a fact-finder's ruling — the evidence-tier flag is binding, and the case carries the journalism-tier framing with the standing language: future validation ongoing.

BACKGROUND

The structural pattern the case names is the inversion of a stable institutional practice. USC's residential MSW program had been an established, modestly sized clinical training program with a known student experience. The 2U partnership took the program's name and credential and built a parallel online operation at roughly ten times the cohort scale, with student-facing operations — recruiting, enrollment counseling, clinical-placement coordination — run by 2U employees who identified to applicants and enrollees with usc.edu

email addresses. The university's delegation extended beyond program operations to the student-facing surface that defines what the credential means to the student paying for it.^o

WHAT HAPPENED

The Luna class-action complaint, filed in Los Angeles County Superior Court in May 2023, structures its case around two central allegations. The first is misrepresentation: USC marketed the online program as the "same" as the residential program — same faculty, same standards, same credential — while the operational substance of the program had been outsourced. The tuition price (over \$100,000) tracked the residential pricing structure while the delivered student experience tracked the OPM-operated structure. The second is targeting: 2U marketers, the complaint alleges, focused aggressive and deceptive recruitment tactics on students of color and veterans — populations for whom the credential cost-and-mobility calculation is structurally different and for whom the misrepresentation has differential consequences.¹

THE INVESTIGATION

The downstream consequences the complaint identifies extend into the licensure question. The MSW credential is a professional-licensure prerequisite, and clinical placement is the operational core of the training the licensure depends on. The complaint alleges that the delegated clinical-placement coordination — handled by 2U employees under usc.edu cover — did not consistently produce placements at the quality the residential program's reputation implied. The licensure-board half of the consequence chain — what graduates were actually able to do in practice, given clinical-placement quality — has not been independently studied, and the case carries that gap honestly rather than collapsing it.²

THE CAPABILITY GAP

USC and 2U announced the partnership's termination on the MSW and other programs in November 2023. The termination closed the operational arrangement at the boundary the litigation and the broader 2U commercial collapse forced; it did not adjudicate the legal claim, did not refund tuition paid under the prior arrangement, and did not produce an independent record of what the delegated operations actually delivered. The case sits at the consumer-side counterpart to Case 57 (GAO OPM oversight gap): one half is the regulator-side audit of the regime, the other is the litigation that names what happened to specific applicants and enrollees under the regime. Pair also with Case 55 (Engler / enrollment algorithms) for the pricing- optimization half.³

AFTERMATH & REFORM

The journalism-tier evidence flag is binding on the case's framing. The complaint and the partnership-termination record are investigation-grade artifacts. The predatory-targeting reconstruction relies on the complaint's allegations and on contemporaneous reporting — Higher Ed Dive, classaction.org summaries, the Project on Predatory Student Lending's statement — rather than on a fact-finder's ruling or an independent audit of 2U's marketing operations. Future validation ongoing on the litigation's outcome, on the licensure-board half of the consequence chain, and on whether the post-2024 successor-OPM contracts incorporate the disclosure architecture the case names as missing.

The tuition tracked residential pricing. The delivered student experience tracked the OPM-operated structure. The gap between what was promised and what was delivered is the case.

EDITORS' SYNTHESIS OF THE LUNA V. USC COMPLAINT AND CONTEMPORANEOUS REPORTING.

REFERENCES

1. Stephanie Luna v. University of Southern California, class action complaint, Los Angeles County Superior Court, May 2023.
2. Higher Ed Dive reporting on the Luna complaint, May 2023; classaction.org and topclassactions.com summaries.
3. Project on Predatory Student Lending statement on USC-2U partnership termination, November 2023.
4. 2U Inc. and USC public statements on partnership termination, November 2023; broader 2U commercial-collapse reporting referenced through Case 57.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Luna v. USC is the consumer-side journalism-tier counterpart to GAO-22-104463 (Case 57). The complaint alleges USC marketed the online MSW as the “same” as the residential program while outsourcing the student-facing operations to 2U; the licensure-board half of the consequence chain is independently unstudied. Journalism-tier flag is binding — the predatory-targeting reconstruction is allegation-tier; future validation ongoing on litigation outcome and licensure consequences.

LENS APPROACH

USC × 2U Luna is the disclosure-as-deliverable case at OPM-delegation scale (induced 5.4; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the promised-vs-delivered disclosure gap and in Domain 3 (Machine Teaming and Adaptation) for the delegation-with-revocation pattern — the partnership wind-down was announced in November 2023, but the consequences for students who enrolled under the prior arrangement persist. Pair with Case 57 (GAO regulator-side audit), Case 55 (Engler pricing optimization), and Case 5 (Epic Sepsis governance gap). coi-light render under the title is binding.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER EDUCATION	D K N	D1 D2 D3 D4 D5	6	5.4	5 · 3
ONLINE PROGRAM MANAGEMENT	→	→ Navigating Sociotechnical Constraints			
PROFESSIONAL LICENSURE					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the student-facing surface — recruiting communications, advising emails, clinical-placement coordination — as part of the disclosable program substance, not as operational discretion to be delegated under brand cover.
- ▶ Require that delegated personnel identify their actual employer in student-facing communications; the usc.edu email cover the complaint names is operationally addressable by branding policy.
- ▶ Build a paired disclosure that ties the delivered student experience back to the marketed program description; the gap the complaint alleges is between what was promised and what was delivered, and the gap is reportable.

IN OPERATION / AFTER

- ▶ Carry the journalism-tier framing into print without softening; the litigation is ongoing, the targeting reconstruction is allegation-tier, and the case's pedagogical power rests on naming the evidence tier honestly.
- ▶ Pair the case in syllabi with Case 57 so the regulator-side and consumer-side halves of the OPM regime are taught together; one half names what the audit missed, the other names what the affected students alleged happened.
- ▶ Track the licensure-board half over time; an independent study of clinical-placement quality and post-graduation practice capacity is the audit the case names as the missing evidence.

REFLECTION QUESTIONS

1. Identify a program in your domain where the student-facing surface — recruiting, advising, placement — has been delegated to a contracted vendor operating under institutional brand. What is the disclosure your institution makes to applicants about that delegation, and at what point in the transaction?
2. Specify the paired disclosure you would build to tie delivered student experience back to the marketed program description. What data — placement outcomes, advising-load ratios, vendor-employee proportion of student-facing communications — would you commit to publishing alongside enrollment marketing?
3. The journalism-tier flag is binding because the targeting reconstruction is allegation-tier. What is the investigation-grade study that would convert the case from allegation-tier to audit-tier — and who could commission it?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

CASE 62

K-12 EDUCATION

TEACHER EFFECTIVENESS REFORM

PROGRAM EVALUATION

2009–2018

Gates Intensive Partnerships — Measurement Without a Coupled Lever

G Governance & Trust **K** Knowledge & Institutional Memory

IMPACT ~\$575M Intensive Partnerships for Effective Teaching initiative (Gates Foundation ~\$212M; the remainder primarily site funds) across Hillsborough County FL, Memphis/Shelby County TN, Pittsburgh PA, and four California charter management organizations, 2009 – 2016; RAND/AIR six-year evaluation (Stecher et al. 2018, RR-2242) found the initiative did not achieve its goals for student achievement or graduation and produced no evidence of improved outcomes for low-income minority students relative to comparison sites; Hillsborough drew down reserves and dismantled its system in 2015

IN BRIEF

The Bill & Melinda Gates Foundation's Intensive Partnerships for Effective Teaching initiative (2009 – 2016) spent approximately \$575 million — roughly \$212 million from the foundation, the remainder primarily district and site funds — across three districts (Hillsborough County FL, Memphis/Shelby County TN, Pittsburgh PA) and four California charter management organizations to redesign teacher evaluation, feedback, and personnel decisions around multi-measure effectiveness ratings: value-added scores, structured classroom observations, and student surveys, informed by the foundation's parallel Measures of Effective Teaching (MET) project (2009 – 2013). The sites built the evaluation architecture. RAND and AIR's six-year evaluation — Stecher et al., *Improving Teaching Effectiveness: Final Report* (2018, RR-2242) — found the initiative did not achieve its goals for student achievement or graduation, and found no evidence of improvement in low-income minority students' outcomes or access to effective teaching relative to comparison sites. Hillsborough, the flagship implementation, drew down reserves reportedly by some \$200 million, absorbed costs well beyond the foundation's contribution, and dismantled the system in 2015. The measurement layer was real; the levers it was meant to drive — retention, placement, development — did not move the outcome, and the districts could not carry the system the theory of change required. The RAND evaluation itself is the exemplar of decision-grade evidence in the case; the intervention is the failure.

BACKGROUND

In 2009 the Gates Foundation launched the Intensive Partnerships for Effective Teaching, selecting three large districts — Hillsborough County Public Schools in Florida, Memphis City Schools in Tennessee (merged into Shelby County Schools in 2013), and Pittsburgh Public

Schools — and four California charter management organizations (Alliance College-Ready, Aspire, Green Dot, and Partnership to Uplift Communities). Each site agreed to redesign teacher evaluation around a multi-measure effectiveness rating combining growth in student achievement, structured classroom-observation rubrics, and student surveys, and then to couple the rating to the personnel system: recruitment, hiring, placement, tenure, dismissal, individualized professional development, career pathways, and performance-linked compensation. The theory of change was explicit — better measurement of teaching effectiveness would drive personnel and development decisions that would increase students' access to effective teaching, and that access would improve achievement and graduation, particularly for low-income minority students. Total spending across the sites reached approximately \$575 million, of which the foundation contributed roughly \$212 million; the remainder was primarily site funds.⁹

WHAT HAPPENED

The measurement layer was built with unusual rigor. The foundation's parallel Measures of Effective Teaching (MET) project (2009 – 2013) studied roughly three thousand teachers across six districts to validate the multi-measure composite, culminating in Kane, McCaffrey, Miller, and Staiger's January 2013 final report, *Have We Identified Effective Teachers?*, which used random assignment of approximately 800 teachers in grades four through eight to test whether composite effectiveness measures predicted student achievement under reassignment. The reported answer was broadly yes — the composite of value-added, observations, and student surveys carried predictive signal — with a load-bearing caveat the project itself documented: substantial noncompliance with the random assignment weakened the experimental protocol, and critics noted the validation spoke to prediction within the studied settings, not to what a district's personnel system would do with the score. The Intensive Partnerships proceeded on the measurement evidence. What the MET project did not and could not establish was the second half of the theory of change: that a district could operate the rating at scale, couple it to retention, placement, and development decisions, and sustain the machinery long enough for the coupling to move student outcomes.¹

THE INVESTIGATION

RAND and AIR's six-year evaluation — Stecher et al., *Improving Teaching Effectiveness: Final Report: The Intensive Partnerships for Effective Teaching Through 2015 – 2016* (2018, RR-2242) — is the load-bearing evidence on the deployment. The evaluation found that the sites did implement the measurement architecture: new multi-measure evaluation systems were built and operated, and the effectiveness ratings were generally regarded as an improvement over the prior binary evaluations. But by 2014 – 2015, student achievement, graduation rates, and low-income minority students' access to effective teaching were not detectably better than in comparison sites that did not participate. The evaluation's site-level implementation record explains the decoupling: the personnel levers the ratings were meant to drive moved weakly or not at all — dismissal and non-renewal touched few teachers, placement incentives did not shift effective teachers toward high-need schools at meaningful scale, individualized professional development linked to the ratings was the least-implemented component, and political and labor resistance narrowed the coupling further over time. The evaluation itself — pre-specified outcomes, comparison-site design, six years of implementation tracking, published in full at 500-plus pages — is the model

of decision-grade evidence the initiative's own theory of change was never subjected to in advance.²

THE CAPABILITY GAP

Hillsborough County was the flagship implementation and is the sustainability record. The district's Empowering Effective Teachers program, launched in 2009 with a foundation pledge of \$100 million, built a corps of more than 260 peer evaluators and mentors and a performance-linked pay system. Tampa Bay Times reporting (Sokol, 2015) documented the fiscal arc: the district ultimately spent more than \$180 million, reportedly \$80 million of the pledged \$100 million was paid by the foundation, and the district's reserves fell by roughly \$200 million in the two years after the new pay system took effect. In October 2015 superintendent Jeff Eakins announced the dismantling of the Gates-funded system, and the district wound down the peer-evaluator corps and the evaluation machinery it could no longer afford. The sustainment requirement — what the district would have to carry after the grant, at what recurring cost, against what revenue — was never engineered as a first-class requirement of the design, and when the foundation's share ended and costs ran past projections, the system's political and fiscal support collapsed together.³

AFTERMATH & REFORM

The hedges the case carries are load-bearing. The RAND evaluation is comparison-based, not randomized at the site level: it cannot rule out small effects, and implementation varied enough across sites that no single component is isolated as the point of failure. The MET validation evidenced the measurement, not the coupling, and its own random-assignment protocol was weakened by noncompliance. Hillsborough's dollar figures are drawn from contemporary journalism and district statements, not an audited public accounting. What the record does establish is structural: a rigorous, well-funded reform built a genuinely strong measurement layer and failed at the coupling between measurement and organizational capability — attribution (what the effectiveness rating actually measures and what decisions it can bear) and sustainability (what the district can carry) were both underengineered relative to the measurement itself. The case pairs with Case 51 (Atlanta Public Schools) for the inverse failure — measurement corrupted rather than uncoupled — and with Case 50 (Wisconsin DEWS) for the same prediction-without-evidenced-intervention structure in a student-facing system.

The sites implemented new multi-measure teacher-evaluation systems, but the initiative did not achieve its goals for student achievement or graduation, and low-income minority students' outcomes and access to effective teaching showed no evidence of improvement relative to comparison sites; Hillsborough drew down reserves and dismantled its system in 2015.

EDITORS' SYNTHESIS OF THE RAND/AIR FINAL EVALUATION (STECHEER ET AL., 2018, RR-2242) AND TAMPA BAY TIMES REPORTING ON THE HILLSBOROUGH IMPLEMENTATION (SOKOL, 2015).

REFERENCES

1. Stecher, B. M., Holtzman, D. J., Garet, M. S., Hamilton, L. S., et al. (2018). Improving Teaching Effectiveness: Final Report: The Intensive Partnerships for Effective Teaching Through 2015–2016, RAND Corporation, RR-2242-BMGF.
2. Kane, T. J., McCaffrey, D. F., Miller, T., & Staiger, D. O. (2013). Have We Identified Effective Teachers? Validating Measures
4. Stecher, B. M., et al. (2016). Improving Teaching Effectiveness: Implementation: The Intensive Partnerships for Effective Teaching Through 2013–2014, RAND Corporation, RR-1295.
5. Sawchuk, S. (2018). "An Expensive Experiment": Gates Teacher-Effectiveness Program Shows No Gains for Students," Education Week, June 21, 2018.

- of Effective Teaching Using Random Assignment, MET Project research paper, Bill & Melinda Gates Foundation, January 2013.
3. Sokol, M. (2015), "Sticker shock: How Hillsborough County's Gates grant became a budget buster," Tampa Bay Times, October 2015 — with companion reporting, "Hillsborough schools to dismantle Gates-funded system that cost millions to develop."

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Intensive Partnerships are the load-bearing case for a reform that engineered its measurement and did not engineer its coupling. A validated multi-measure effectiveness rating was built and deployed at scale, and the levers it was meant to drive — retention, placement, development — did not move the outcome, while the sustainment load was never specified as a requirement the district could be verified against. The intervention is the failure; the RAND evaluation that established it is the exemplar of decision-grade evidence.

LENS APPROACH

The Intensive Partnerships case is the requirements-as-a-sustainment-deliverable anchor at philanthropic-reform scale (induced 1.4; LENS D1/PT1; CLO-1). LENS uses it in Domain 1 (Systems Analysis) for the discipline of engineering the full theory-of-change chain — measurement, coupling, capacity, sustainment — as explicit requirements rather than letting a strong measurement layer stand in for the system. Pair with Case 51 (Atlanta Public Schools) for measurement corrupted rather than uncoupled, and Case 50 (Wisconsin DEWS) for prediction-without-evidenced-intervention. The RAND/AIR final report is taught alongside the failure as the model of the evaluation the design phase never had.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
K-12 EDUCATION	G K	D1 D2 D3 D4 D5	1	1.4	1
TEACHER EFFECTIVENESS REFORM	→	→ Systems Analysis			
PROGRAM EVALUATION					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Engineer the full theory of change as a requirements chain, not a measurement project; every link — rating to personnel decision, personnel decision to access to effective teaching, access to outcome — is a stated requirement with its own evidence plan, and a strong first link does not discharge the rest.
- ▶ Treat sustainment as a first-class design requirement; specify what the district must carry after the grant, at what recurring cost, against what revenue and what political tolerance, before the architecture is built — Hillsborough's collapse was fiscal and political before it was technical.
- ▶ Pre-commit the evaluation design and the decision rule; the RAND/AIR comparison-site evaluation with pre-specified student outcomes is the form the initiative's own go/no-go and mid-course decisions required, and it arrived as a retrospective rather than a control.

IN OPERATION / AFTER

- ▶ Publish the null in full; the 500-plus-page RAND final report, commissioned by the funder whose initiative it evaluates, is the standing model of an institution buying decision-grade evidence against itself and releasing it — carry that discipline forward as the norm for funded reforms.
- ▶ Attribute the failure to the coupling, not the measurement; the post-initiative lesson is not that multi-measure teacher evaluation is worthless but that measurement uncoupled from movable organizational levers cannot produce the outcome, and reforms that inherit the instruments must re-engineer the levers.
- ▶ Recover the sustainment accounting; reconstruct what each site actually carried — recurring cost, staffing, political capital — so the next reform's requirements include the load the last one could not bear.

REFLECTION QUESTIONS

1. Take a reform or system in your domain whose measurement layer is strong and write out the full theory-of-change chain from measurement to outcome. Which links are engineered requirements with evidence plans, and which are assumptions a strong first link is being allowed to stand in for?
2. Specify the sustainment requirement for a granted or pilot-funded capability you know: what must the institution carry after the external funds end, at what recurring cost, against what revenue and political tolerance? What would a Hillsborough-style drawdown look like in your setting, and when would you detect it?
3. The funder commissioned and published the evaluation that established its own initiative's failure. What is the standing mechanism in your organization for buying decision-grade evidence against your own program — and what is the pre-committed decision rule if the finding is null?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines

LAUSD iPad Program — The Common Core Technology Project

D Designed Out **H** Human-System Interface **G** Governance & Trust

EVIDENCE TIER journalism-tier — source confidence flagged; future validation ongoing.

IMPACT Los Angeles Unified's plan to put an iPad loaded with Pearson Common Core curriculum in every student's hands — publicly discussed at roughly \$1.3 billion (approximately \$500 million for devices and curriculum plus approximately \$800 million for network infrastructure, drawn from school-construction bonds) — collapsed within two years: students bypassed device security within about a week at initial schools, the district's own evaluator observed the Pearson curriculum in use in one of 245 classrooms, the Apple contract was halted in August 2014, Superintendent John Deasy resigned in October 2014, the FBI seized twenty boxes of documents in December 2014, the SEC examined the bond-funding disclosures, and the district recovered a settlement of approximately \$6.4 million in 2015

IN BRIEF

In June 2013 the Los Angeles Unified School District board approved the first phase of the Common Core Technology Project: an iPad, pre-loaded with Pearson's Common Core curriculum, for every student in the nation's second-largest district, publicly discussed at a full-build scale of roughly \$1.3 billion — approximately \$500 million for devices and curriculum and approximately \$800 million for the network infrastructure to support them, funded substantially from long-term school-construction bonds. The fall 2013 rollout at 47 initial schools outran every layer the deployment depended on: teachers were not prepared, the Pearson curriculum was incomplete, students at Roosevelt High and other campuses bypassed the device security filters within about a week, keyboards had not been budgeted, and home-use liability was unresolved. The procurement governance then unraveled in public: emails published by KPCC in August 2014 showed close contact between Superintendent John Deasy, his deputy, and Apple and Pearson executives predating the competitive bid; the district halted the Apple contract days later; Deasy resigned in October 2014; the FBI seized twenty boxes of documents in December 2014 at a federal grand jury's behest; and the SEC examined whether the bond funding was properly disclosed to investors. The district recovered a settlement of approximately \$6.4 million in 2015, ultimately borne by Pearson. The evidence base is largely journalism — KPCC, the Los Angeles Times, Education Week — plus the district's own commissioned evaluation by the American Institutes for Research, so the evidence-tier flag is rendered under the title; future validation will continue.

BACKGROUND

The Common Core Technology Project was the signature initiative of Los Angeles Unified Superintendent John Deasy: a one-to-one tablet program intended to put an iPad, pre-loaded with Pearson's digital Common Core curriculum, in the hands of every student and teacher in a district of more than 600,000 students. The board approved the initial Apple contract in June 2013, and the project's full build-out was publicly discussed at roughly \$1.3 billion — approximately \$500 million for devices and curriculum and approximately \$800 million set aside for upgrading school networks — funded substantially from voter-approved school-construction bonds, an instrument whose multi-decade repayment horizon critics noted would long outlive the devices it bought. The stated theory was equity and readiness at once: every student would have the device the new Common Core online assessments assumed, and the curriculum would arrive with it.⁰

WHAT HAPPENED

The fall 2013 rollout at 47 initial schools outran the operator layer the deployment's theory required. Within about a week of receiving the tablets, roughly 300 students at Roosevelt High School deleted their device profiles and bypassed the security filters — with the same breach reported at Westchester High and the Valley Academy of Arts and Sciences — and the district suspended home use of the devices while it worked out what it had deployed. Keyboards, which the Common Core online assessments would require, had not been budgeted in the original purchase; home-use liability for a district-owned device in a student's backpack was unresolved; and teacher preparation lagged the hardware. The district's own commissioned evaluator, the American Institutes for Research, observed 245 classrooms across CCTP schools in spring 2014 and reported that the Pearson curriculum — the pedagogical core of the purchase — was in use in just one of them; school staff reported that the content promised during initial training was incomplete or unavailable during the school year.¹

THE INVESTIGATION

The governance layer failed on a separate and compounding track. In August 2014, KPCC published internal emails, obtained by public-records request, showing that Deasy and his deputy had been in close contact with Apple and Pearson executives about the project well before the competitive bid opened, and that the bid requirements tracked curriculum specifications Pearson had suggested in those private exchanges. Three days after publication, the district announced it was halting the Apple contract and reopening the procurement. Deasy resigned in October

2014. In December 2014, at the behest of a federal grand

jury, FBI agents removed twenty boxes of documents from district offices; the SEC separately examined whether the district's use of construction-bond proceeds for tablets was properly disclosed to bondholders. The district sought a refund from Apple and Pearson and in 2015 reached a settlement of approximately \$6.4 million, ultimately borne by Pearson as the curriculum subcontractor. The federal investigation was later reported closed without charges.²

THE CAPABILITY GAP

The evidence-tier flag rendered under the case title is load-bearing. The public record on the CCTP is substantially journalism — KPCC's email investigation, sustained Los Angeles

Times coverage, and Education Week’s multi-year reporting and retrospectives — joined by the district’s own commissioned AIR evaluation and its board-committee oversight reports. No court adjudicated the procurement question; the federal inquiry closed without charges, so the pre-bid-contact finding rests on the published emails and the district’s own responsive conduct rather than on an adjudicated record. The dollar figures are publicly reported approximations — “roughly \$1.3 billion” is the scale the district and press discussed, not an audited final cost, and the settlement figure is reported at approximately \$6.4 million. The structural pattern is consistent across all sources; the specific magnitudes carry the hedges. Future validation will continue.³

AFTERMATH & REFORM

What the case teaches is capability designed out at \$1-billion scale. The deployment’s implicit theory was that device provision equals learning capability: buy the tablet, license the curriculum, and the classroom capability follows. Every layer that theory silently assumed — trained teachers, complete curriculum, working security, keyboards for the assessments the program existed to serve, a home-use policy — was unbuilt at rollout, and the district discovered each absence in production, at 47 schools simultaneously, rather than in an instrumented pilot honest about readiness. The governance layer was unbuilt in parallel: a procurement whose specifications tracked a vendor’s private suggestions could not survive disclosure, and when it surfaced, the program lost its superintendent, its contract, and its public legitimacy in under three months. The two failures are one failure. A deployment that skips the iteration cycle on the human and curricular layers is the same deployment that skips the contestability of its own procurement: both substitute the founding vision for the evidence architecture the vision required.⁴

The Pearson curriculum, touted as central to the effort to prepare students for the new Common Core standards, was in use in just one of the 245 classrooms the district’s own evaluators observed.

EDITORS’ SYNTHESIS OF THE AMERICAN INSTITUTES FOR RESEARCH PHASE 1 EVALUATION FINDING, AS REPORTED BY EDUCATION WEEK (SEPTEMBER 2014).

REFERENCES

1. Gilbertson, A. (2014), “LA schools cancel iPad contracts after KPCC publishes internal emails,” KPCC 89.3 (Southern California Public Radio), August 25, 2014 — the public-records email investigation and the contract halt; see also KPCC’s “LAUSD iPads: Timeline of a troubled program.”
2. American Institutes for Research (2014), interim evaluation of the LAUSD Common Core Technology Project, Phase 1 — 245 classroom observations across CCTP schools; Pearson curriculum observed in use in one.
3. Herold, B. (2014), “Hard Lessons Learned in Ambitious L.A. iPad Initiative,” Education Week, September 2014 — retrospective on the implementation failures; see also Herold (2014), “FBI Investigation Leaves L.A. iPad Initiative in Further Disarray,” Education Week, December 2014.
4. Blume, H., Los Angeles Times coverage of the CCTP (2013–2015) — the rollout, the security bypass, the contract suspension, the Deasy resignation, and the FBI document seizure.
5. Education Week Market Brief (2017), “Feds Drop Investigation Into Los Angeles District Over \$1 Billion iPad Purchase,” February 2017 — the reported closure of the federal inquiry without charges.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The LAUSD Common Core Technology Project is the \$1-billion-scale case for technology deployed without capability coupling: the deployment assumed device provision equals learning capability, and every layer that assumption depended on — teacher preparation, curriculum completeness, device security, keyboards, home-use policy — was unbuilt at rollout, while the procurement’s governance layer was unbuilt in parallel. The district’s own evaluator found the purchased curriculum in use in one of 245 observed classrooms; the published pre-bid emails cost the program its superintendent, its contract, and its legitimacy. Evidence base is journalism-tier plus the district’s own commissioned evaluation; the flag is rendered under the title.

LENS APPROACH

LAUSD CCTP is the ed-tech procurement-governance failure case at district scale (induced 5.1; LENS D5/PT6; CLO-5). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the stakeholder-trust-as-precondition anchor — a procurement that cannot survive disclosure forfeits the deployment’s legitimacy regardless of the technology’s merits — and as a Design-mode teaching case for the device-provision-equals-capability fallacy at the human-interface layer. The evidence-tier flag is binding until the AIR reports and the LA Times citations are re-confirmed against primary sources; the structural pattern survives the flag, the dollar magnitudes do not.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
K-12 EDUCATION	D H G	D1 D2 D3 D4 D5	6	5.1	5
ED-TECH PROCUREMENT	→	→			
DEVICE ROLLOUT		Navigating Sociotechnical Constraints			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pilot at capability scope, not device scope: an instrumented pilot honest about readiness would have surfaced the security bypass, the incomplete curriculum, the missing keyboards, and the teacher-preparation gap at one school’s cost instead of forty-seven — the pilot’s job is to falsify the device-provision-equals-capability theory before the bond money moves.
- ▶ Treat procurement contestability as a deliverable of the program, not an obstacle to it; a specification that tracks one vendor’s private suggestions is a governance breach whether or not it is ever adjudicated, and the disclosure record — not the intent — is what the deployment’s legitimacy will be judged on.
- ▶ Couple the financing instrument to the asset’s life: multi-decade construction bonds spent on three-to-four-year devices is a mismatch the district’s own investors, regulators, and public will eventually price, and the SEC’s disclosure inquiry is the form that pricing takes.

IN OPERATION / AFTER

- ▶ When the evidence base is journalism plus the institution’s own commissioned evaluation, flag it under the title and carry the standing “future validation ongoing” language into print; the structural pattern survives the flag, the dollar magnitudes do not.
- ▶ Preserve the district’s own evaluation as the load-bearing internal-evidence stream: the AIR finding that the purchased curriculum was in use in one of 245 observed classrooms is the deployment’s own instrument reporting the gap between the theory and the classroom.
- ▶ Teach the two failure tracks — the unbuilt operator layer and the unbuilt governance layer — as one failure with one root: the substitution of founding vision for evidence architecture, at the human layer and at the procurement layer simultaneously.

REFLECTION QUESTIONS

1. The CCTP's implicit theory was that device provision equals learning capability. Identify a deployment in your domain whose budget line is the artifact and whose unbudgeted lines are the operator layer — training, workflow, policy, support. What would an instrumented pilot honest about readiness have to measure to falsify the provision-equals-capability theory before full-scale rollout?
2. The procurement failed on disclosure, not adjudication: the federal inquiry closed without charges, but the published pre-bid emails ended the contract, the superintendent's tenure, and the program's legitimacy anyway. What is the disclosure test you would apply to a procurement you are running — and what specification, if published tomorrow, would your program not survive?
3. The case carries journalism-tier evidence plus the institution's own commissioned evaluation. When the deployment's own instrument (the AIR observation of one classroom in 245) agrees with the external journalism, what does that convergence license you to conclude — and which of the case's specific magnitudes still cannot bear weight?

WHO BUILDS THIS systems & human-factors engineering · human factors & interaction design · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · usability testing · cognitive walkthroughs · governance frameworks · bias & evidence audits

Chapter 4

Education, Training & the Learning Workforce — What Works — and the Frontier

When instruction, analytics, and workforce design run the full evidence loop.

The effect survived because somebody kept measuring after the launch.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

Training Transfer — The Gap Between Learning and Doing

K Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT A meta-analysis of 89 empirical studies finds training transfer is positively related to cognitive ability, conscientiousness, motivation — and decisively to a supportive work environment; the literature carries explicit hedges about inconsistent measurement and significant variability

IN BRIEF

Blume, Ford, Baldwin, and Huang (Journal of Management 2010) synthesized 89 empirical studies on training transfer — the extent to which what is learned in training produces meaningful change in on-the-job performance. The headline finding is that transfer is positively related to four categories of variable: cognitive ability, conscientiousness, motivation, and the work environment. Of these, the work environment — particularly supervisor and peer support — is among the strongest predictors, and the most decisive at the **system** layer rather than the individual layer. The authors are explicit, and the load-bearing hedge survives into the case: the literature is characterized by “inconsistent measurement of transfer and significant variability in findings,” and downstream practitioner summaries note that organizations frequently see limited return because learning fails to transfer to the workplace. The case is the paired peer-reviewed half of the corporate-L&D pair with Kirkpatrick (Case 79): together they close the gap the v1 corpus had open at the workforce-L&D layer, and they motivate the LENS framing that the human is the biggest variable at the interface — here, the return-to-work interface.

THE SHIFT

Training transfer is the structural question corporate L&D sits inside: does what learners did in training produce observable change on the job, and what predicts whether it does? The Kirkpatrick chain of evidence (Case 79) frames the question; the Blume et al. meta-analysis is the strongest single peer-reviewed answer the field has consolidated.⁰

WHAT IS EMERGING

Blume, Ford, Baldwin, and Huang synthesized 89 empirical studies, with several thousand learners in aggregate, and decomposed transfer predictors into four categories: cognitive ability, conscientiousness, motivation, and the work environment (particularly supervisor and peer support). All four categories were positively related to transfer; of them, the work environment was among the strongest predictors, and was the only one that is decisively a **system** variable rather than a learner variable.¹

THE CAPABILITY QUESTION

The load-bearing hedge survives. The authors are explicit that the literature is “characterized by inconsistent measurement of transfer and significant variability in findings.” Downstream practitioner summaries note that organizations frequently see limited training return because learning fails to transfer to the workplace — not because the training did not work in the classroom but because the return-to-work environment did not support application. The meta-analysis is the strongest current synthesis, and it also names what the field has not yet measured well.²

EARLY EVIDENCE

The teaching point at the LENS-framing layer is precise. The capability deliverable is not the training event; it is the transfer. And the decisive variable at the transfer layer is the work environment — supervisor support, peer support, the opportunity to practice on actual tasks, the absence of countervailing pressures that punish trying the new behavior. The human is the biggest variable at the interface, and the interface is the return-to-work boundary. Training that ignores this boundary is training that delivers Level 2 evidence (Kirkpatrick, Case 79) and does not deliver Level 3 capability.³

OPEN PROBLEMS

In pair with Kirkpatrick, the case closes the corporate-L&D gap the v1 corpus had open and is a worked example of the revised “decision-grade evidence” framing in the v2 research backbone: the strongest synthesis the field has is a meta-analysis with explicit hedges about measurement inconsistency. The practitioner must decide on what to build — training design, supervisor support, practice opportunity — knowing the evidence is the strongest available but not closed. The CLO **Judgment under inadequate evidence** is exactly the capability this meta-analysis pattern asks practitioners to bring.

The training event is not the capability deliverable. The transfer is. And the decisive variable at transfer is the work environment.

EDITORS' SYNTHESIS OF BLUME ET AL. (2010).

REFERENCES

1. Blume, Ford, Baldwin, & Huang (2010), “Transfer of Training: A Meta-Analytic Review,” *Journal of Management* 36(4):1065–1105, doi:10.1177/0149206309352880.
2. Baldwin & Ford (1988), “Transfer of Training: A Review and Directions for Future Research,” *Personnel Psychology* 41(1):63–105 — the foundational synthesis the 2010 meta-analysis updates.
3. Burke & Hutchins (2007), “Training Transfer: An Integrative Literature Review,” *Human Resource Development Review* 6(3):263–296 — the integrative-review companion synthesis.
4. Kirkpatrick & Kirkpatrick (2006), *Evaluating Training Programs* — the framework the meta-analysis informs (paired Case 79).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Blume et al. is the strongest current peer-reviewed synthesis on training transfer: cognitive ability, conscientiousness, motivation, and decisively the work environment predict whether training produces on-job behavior change. The literature carries explicit load-bearing hedges about inconsistent measurement and significant variability; the case is included with the hedges intact.

LENS APPROACH

Blume is the corporate-L&D transfer case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the design-iteration sub-competency (the work environment is the design variable) and in Domain 4 (Test and Evaluation) for the CLO **Judgment under inadequate evidence** — the meta-analysis is the strongest synthesis the field has, and it explicitly names what it cannot settle. Direct pair with Case 79 (Kirkpatrick); together they close the corporate-L&D gap the v1 corpus had open.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
CORPORATE L&D	K N	D1 D2 D3 D4 D5	4	2.3	2 · 4
TRAINING TRANSFER	→	Iterative Development			
ORGANIZATIONAL BEHAVIOR					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Design the training intervention with the work environment as a design variable — supervisor briefing, peer-support structures, practice opportunity — not only the classroom content. ▶ Specify the on-job behavior the training is supposed to enable and the conditions under which it will be observed; the meta-analysis identifies environment as decisive, so the environment must be instrumented. ▶ Carry the meta-analytic hedge into the deployment documentation: the literature is the strongest current synthesis, and it explicitly notes inconsistent measurement and significant variability.	▶ Track the work-environment variables — supervisor support, peer-support structures, practice opportunity — separately from the training event, and report transfer outcomes conditional on environment. ▶ Use the meta-analytic finding to argue for the work-environment investment in the program design, not to declare the question settled; future replications and better-measured studies will move the magnitudes. ▶ When transfer is low, attribute the gap with the meta-analysis in hand: was the trainee under-prepared, under-motivated, or returning to an environment that did not support application? The remediation depends on the attribution.

REFLECTION QUESTIONS

1. Identify a training program in your organization that produced strong Level 2 evidence (learning) but unclear Level 3 evidence (behavior change). What does the meta-analysis predict the decisive variable was, and how would you re-design the program with the work environment as a first-class design variable?
2. Specify the work-environment instrumentation you would put in place for a new training deployment — supervisor briefing artifacts, peer-support structures, on-job practice opportunity — so transfer can be measured conditional on environment, not just on training fidelity.
3. The meta-analysis is the strongest current synthesis and carries explicit hedges about measurement inconsistency. What is the minimum additional evidence you would require before treating any specific magnitude from this literature as settled in your decision-making?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS knowledge bases · data pipelines · incident analysis · safety dashboards

Growth-Mindset National Experiment — Heterogeneous Effects

D Designed Out **N** Normalization of Deviance **K** Knowledge & Institutional Memory

IMPACT A nationally representative RCT of US 9th-graders found a less-than-1-hour online growth-mindset intervention improved grades among lower-achieving students and increased advanced-math enrollment, but the effect was conditional on peer norms — the intervention changed grades only where peer norms aligned with the intervention's message

IN BRIEF

Growth-mindset interventions — short, scalable psychological interventions that teach students that intellectual ability is malleable rather than fixed — accumulated a substantial laboratory and small-school evidence base across the 2000s and 2010s. The open question, framed in the broader scalable-interventions literature, was whether the effects survived at population scale and what the structural moderators were. Yeager et al. (Nature, 2019) ran the test that became the field's reference point. A nationally representative RCT of US 9th-graders received a less-than-1-hour online growth-mindset intervention; the trial documented improved grades among lower-achieving students and increased advanced-math enrollment relative to control. The headline finding for the case is the conditional: the effect was conditional on peer norms. The intervention changed grades only where peer norms aligned with the intervention's message, and the study is notable for treating treatment-effect heterogeneity as the finding rather than as a nuisance. The intervention is not universally effective; naming where it does and does not work is the contribution. The "conditional on peer norms" language survives verbatim into the case. Pair with Case 86 (Gándara) at the scalable-equity-intervention layer.

THE SHIFT

Growth-mindset interventions teach students that intellectual ability is malleable — improvable through effort, strategy, and help-seeking — rather than fixed. The pedagogical claim, developed across two decades of research (Dweck and colleagues), is that students who hold the malleable view respond more constructively to academic challenge: they treat difficulty as informative rather than as a signal of fixed ability, and they persist on problems that the fixed-view student would interpret as evidence to stop. The laboratory and small-school evidence base accumulated substantially across the 2000s and 2010s; the open question, by the late 2010s, was whether the effects survived at population scale and what the structural moderators were.⁹

WHAT IS EMERGING

Yeager et al. (Nature, 2019) ran the trial that became the field's reference point. The design was a nationally representative RCT of US 9th-graders — drawn from a sample stratified to reflect the US ninth-grade population — with the intervention delivered as a less-than-1-hour online module. Outcomes included grades, course-taking, and contextual measures of the classroom and peer environment. The trial pre-registered the moderator analysis the case now anchors on: the effect of the intervention was hypothesized in advance to depend on the peer-norm environment the student returned to after the module.¹

THE CAPABILITY QUESTION

The headline outcome was a positive finding with structure. The growth-mindset intervention improved grades among lower-achieving students and increased advanced-math enrollment, relative to the active-control condition. The structural finding — the one the case is built around — is the conditional: the effect was conditional on peer norms. The intervention changed grades only where peer norms aligned with the intervention's message — where the peers treated the malleable view as legitimate and the help-seeking behavior the intervention encouraged as socially acceptable. In peer environments where the fixed view was the local norm, the intervention's effect on grades was substantially smaller or absent.²

EARLY EVIDENCE

What makes the case methodologically important is the authors' explicit treatment of treatment-effect heterogeneity as the finding rather than as a nuisance to be averaged away. The trial's structural answer is that the intervention is not universally effective, and naming where it does and does not work is the scientific contribution. The "conditional on peer norms" language survives verbatim into the case because it is the load-bearing hedge: a headline-only reading ("growth mindset works at population scale") misses the substance, and an opposite-headline reading ("growth mindset doesn't work") misses it equally. The contribution is the conditional, and the design — a pre-registered moderator analysis with a nationally representative sample — is what makes the conditional defensible.³

OPEN PROBLEMS

Drafted alongside Case 86 (Gándara) at the scalable-equity-intervention layer, the case carries the design-predictions-to-trigger-support pattern (induced 8.3, LENS D2/PT5). The intervention targets the students who benefit (lower-achieving students in peer environments where the norm permits the change), rather than applying a blanket treatment that the average effect would credit and the heterogeneity would conceal. The equity-relevant design commitment is to surface the heterogeneity as part of the deliverable, not to launder it through an average-effect headline. The case is the methodologically clean model of how a population-scale RCT can earn the heterogeneity-as-finding stance, and the "conditional on peer norms" qualification is what makes the result useful for the next adaptation.

The effect was conditional on peer norms.

YEAGER ET AL., NATURE 2019.

REFERENCES

1. Yeager, D. S., Hanselman, P., Walton, G. M., Murray, J. S., Crosnoe, R., Muller, C., Tipton, E., Schneider, B., Hulleman, C. S., Hinojosa, C. P., Paunesku, D., Romero, C., Flint, K., Roberts, A., Trott, J., Iachan, R., Buontempo, J., Yang, S. M., Carvalho, C. M., Hahn, P. R., Gopalan, M., Mhatre, P., Ferguson, R., Duckworth, A. L., & Dweck, C. S. (2019). A national experiment reveals where a growth mindset improves achievement. *Nature*, 573(7774):364–369. doi:10.1038/s41586-019-1466-y — the case's primary trial.
2. National Study of Learning Mindsets, ICPSR 37353 — the trial dataset.
3. Dweck, C. S. (2006). *Mindset: The New Psychology of Success*. Random House — the broader theoretical framework the intervention rests on.
4. Sisk, V. F., Burgoyne, A. P., Sun, J., Butler, J. L., & Macnamara, B. N. (2018). To what extent and under which circumstances are growth mindsets important to academic achievement? Two meta-analyses. *Psychological Science*, 29(4):549–571 — the prior moderator-analysis literature the Yeager trial extends.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The growth-mindset national experiment is the methodologically clean model of how a population-scale RCT can earn the heterogeneity-as-finding stance. The intervention improved grades among lower-achieving students and increased advanced-math enrollment — and the effect was conditional on peer norms. The qualifying language is the load-bearing hedge; headline-only readings in either direction miss the substance.

LENS APPROACH

Yeager et al. 2019 is the designing-predictions-to-trigger- support case (induced 8.3; LENS D2/PT5) — Domain 2 for pre-registered moderator analysis; Domain 5 for the equity-relevant commitment to target support to those who benefit. Pair with Case 86 (Gándara).

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
K-12 EDUCATION	D N K	D1 D2 D3 D4 D5	5	8.3	2-5
SCALABLE INTERVENTIONS	→	Iterative Development			
EQUITY					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pre-register the moderator analysis at the design stage; the heterogeneity-as-finding stance depends on the moderator being a planned analysis rather than a post-hoc inspection.
- ▶ Sample so that the moderator can be estimated — a nationally representative sample of US 9th-graders, in this case — so the heterogeneity is observed across the population the headline claim would otherwise speak to.
- ▶ Treat the intervention as scalable on the headline finding and conditional on the moderator at the same time; designing predictions to trigger support means targeting the support to those who benefit, in the environments where the support can land.

IN OPERATION / AFTER

- ▶ Carry the conditional language (“effect was conditional on peer norms”) verbatim into every downstream communication; a headline-only reading and an opposite-headline reading both miss the substance.
- ▶ Treat the treatment-effect heterogeneity as program documentation, not as a nuisance to be laundered through an average effect; the equity-relevant design commitment is to surface the heterogeneity, not to conceal it.
- ▶ Build the next adaptation against the moderator: where would the intervention work, what peer-norm context would it land in, and what is the institutional path to creating that context where it does not yet exist?

REFLECTION QUESTIONS

1. Identify a scalable intervention in your domain whose published evidence runs at the average-effect level. What pre-registered moderator analysis — and what sampling design — would let you treat treatment-effect heterogeneity as the finding rather than as a nuisance?
2. The Yeager finding is “effect was conditional on peer norms.” What is the analog conditional in your context — the structural moderator the intervention’s effect depends on — and how would you instrument the moderator at scale?

WHO BUILDS THIS systems & human-factors engineering · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · incident analysis · safety dashboards · knowledge bases · data pipelines

Cognitive Tutor / Carnegie Learning

T Training Gap

IMPACT Randomized controlled trials showed learning gains; RAND RCT found a significant Algebra I gain in year two (~0.2 SD, high school; no year-one effect); adopted across 3,000+ schools

IN BRIEF

Carnegie Learning's Cognitive Tutor, built from John Anderson's ACT-R cognitive architecture at Carnegie Mellon, is the most rigorously evaluated intelligent tutoring system in education. It uses Bayesian knowledge tracing to model each student's mastery and adapts instruction accordingly, and a RAND Corporation randomized trial found a statistically significant positive effect on Algebra I achievement — about 0.21 SD for high school students in the second year of use, with no significant effect in year one and a middle-school estimate that did not reach significance. The system is a learning-engineering success in the discipline's own terms — learning science to engineered software to randomized-trial evidence to deployment across 3,000-plus schools. Its limitations are instructive: it works best in well-defined domains like algebra and less well in ill-structured ones, making it the canonical evidence that the pipeline delivers for problems that fit it — leaving open whether the same discipline can deliver where problems do not.

BACKGROUND

For decades, intelligent tutoring systems promised individualized instruction at scale, but few were grounded in a validated theory of how people learn or rigorously tested for effect. The opportunity was to build a tutor from a real cognitive model and prove its impact with the methods of experimental science. The field's recurring weakness was that promising systems rested on intuition about learning rather than a validated theory, and were rarely subjected to the kind of controlled trial that could separate genuine effect from novelty.⁰

THE INTERVENTION

Carnegie Learning's Cognitive Tutor, developed from John Anderson's ACT-R cognitive architecture at Carnegie Mellon, models the specific skills underlying a subject and tracks each student's mastery using Bayesian knowledge tracing. It adapts problem selection to the individual learner and provides step-level feedback, embodying a full learning-science theory in software. Grounding the tutor in ACT-R rather than designer intuition is what made the system testable: a theory that decomposes a subject into specific skills can be turned into a measurement of mastery and an instrument that responds to it.¹

HOW IT WORKED

The tutor's effectiveness rests on the chain from theory to instrument: a decomposable skill model, a measurement method (knowledge tracing) that estimates mastery from student


actions, and an instrumentable interface that adapts in response. Instruction is targeted where the model detects weakness, so practice concentrates on skills not yet mastered rather than on a fixed sequence. Each link in that chain has to hold for the next to work — the skill model makes mastery measurable, the measurement makes adaptation possible, and the adaptation is what concentrates practice where it pays off.²

THE EVIDENCE

The RAND Corporation's multi-site evaluation found statistically significant positive effects on Algebra I achievement, and the program scaled to more than 3,000 schools. The case demonstrates the learning-engineering process working end to end: learning science, to engineered software, to randomized-controlled-trial evidence, to scaled implementation. A multi-site randomized evaluation is the strong form of the claim — it shows the effect survived contact with many real classrooms rather than a single favorable setting, which is what the field's earlier, untested systems had lacked.³

WHAT TRANSFERRED

The limitations are as instructive as the success. Cognitive Tutor performs best in well-defined domains like algebra and less well in ill-structured ones, making it the canonical evidence that the learning-engineering pipeline works for problems that fit the pipeline. The frontier question — whether the same discipline can deliver in operational, ill-structured domains where capability matters most — remains open. The dependence on a decomposable skill model is the boundary condition: where a subject cannot be cleanly broken into trackable skills, the very chain that made algebra tractable has nothing to attach to.⁴



Cognitive Tutors demonstrate the LE process working: theory → engineering → evidence → scale.

EDITORS' SYNTHESIS OF ANDERSON ET AL. (1995) AND KOEDINGER & CORBETT (2006)

REFERENCES

1. Anderson, J., Corbett, A., Koedinger, K. & Pelletier, R. (1995), "Cognitive Tutors: Lessons Learned," Journal of the Learning Sciences — the ACT-R basis and design.
2. Koedinger, K. & Corbett, A. (2006), Cambridge Handbook of the Learning Sciences — knowledge tracing and adaptive instruction.
3. RAND Corporation Algebra I evaluation — the statistically significant achievement effects.
4. Carnegie Learning program documentation — scale to 3,000-plus schools.
5. Aleven, V. et al. (2016), example-tracing tutors — the limits in ill-structured domains.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Cognitive Tutor is the canonical evidence that the learning–engineering process exists, works, and produces measurable benefits at scale — when the problem is tractable. It is also the canonical case for what tractability looks like: well-defined domain, decomposable skill model, instrumentable interface, rigorous evaluation. The frontier for the discipline is whether the same pipeline can deliver in the operational, ill-structured domains where capability matters most.

LENS APPROACH

LENS uses Cognitive Tutor in LEN 1 as the foundational LE-process exemplar, in LEN 4 as the canonical case for Bayesian knowledge tracing as a measurement instrument, and in LEN 9 as a technical case for model-based adaptive instruction.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
EDUCATION AT SCALE	→ T	→ D1 D2 D3 D4 D5	5	2.3	2
Iterative Development					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Ground the intervention in a validated theory that decomposes the subject into specific skills, rather than building on designer intuition about how people learn.
- ▶ Build the full chain — a skill model, a measurement method that estimates mastery from learner actions, and an interface that adapts in response — so each link supports the next.
- ▶ Target instruction where the model detects weakness, concentrating practice on unmastered skills rather than marching through a fixed sequence.

IN OPERATION / AFTER

- ▶ Prove the effect with a multi-site randomized evaluation, so the gain is shown to survive many real settings rather than one favorable classroom.
- ▶ Scale only where the problem fits the pipeline — well-defined, decomposable domains — and treat the decomposability of a skill model as the boundary condition for applicability.
- ▶ Keep the frontier question explicit when extending the method, testing whether the same theory-to-instrument chain can be built in less-structured domains before assuming it transfers.

REFLECTION QUESTIONS

1. Cognitive Tutor works best in well-defined domains. Identify a problem in your domain that is currently ill-structured. What would have to be true to make the LE pipeline applicable?
2. The case is the strongest evidence the LE process works. What is the equivalent piece of evidence required to demonstrate the same pipeline in a non-cognitive domain — for example, surgical skill, or operational watchstanding?
3. Cognitive Tutor's chain runs from a decomposable skill model through measurement to adaptation. Take a capability in your domain and attempt the decomposition — where it resists being broken into trackable skills, what does that tell you about whether the pipeline can apply?

WHO BUILDS THIS learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models

Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale

T Training Gap **D** Designed Out

IMPACT Settles & Meeder (ACL 2016) introduced Half-Life Regression (HLR), a trainable spaced-repetition model that learns per-item forgetting rates from large-scale learner data; HLR was deployed in Duolingo's production review-scheduling system and the published evaluation reports improvements over heuristic schedulers (Leitner, Pimsleur) on Duolingo's own predictive metric and on a daily (next-day) retention metric

IN BRIEF

Settles and Meeder (Association for Computational Linguistics, 2016) introduced Half-Life Regression (HLR), a trainable statistical model for spaced repetition in language learning. HLR combines the psychological theory of memory half-life (Ebbinghaus's forgetting curve and its descendants) with a learned regression that estimates each item's per-learner half-life from observed practice history and item features. The model was deployed in Duolingo's production review-scheduling system; the published evaluation compares HLR against heuristic schedulers (Leitner spacing, Pimsleur intervals) on Duolingo's predictive recall metric and on a daily (next-day) learner-retention metric, and reports improvements on both. The case is one of the few published instances of a spaced-repetition algorithm being deployed and evaluated against meaningful behavioral outcomes at consumer scale (Duolingo had tens of millions of active learners at the time of the study). The hedges that survive into the case verbatim: the evaluation is a single-vendor study, the "learning" outcome is measured by Duolingo's predictive metric and next-day retention rather than independent language-proficiency assessment, and the generalization to other content domains rests on the structural argument rather than a multi-vendor evidence base.

BACKGROUND

The psychological theory of memory has carried, since Ebbinghaus's 1885 forgetting-curve work, the finding that review timed near the point of forgetting produces stronger long-term retention than either massed practice or uniform-interval review. The applied descendants — Leitner boxes (1972), Pimsleur graduated-interval recall (1967), SuperMemo's SM-2 algorithm (1985, and successors) — are all attempts to operationalize the theory as a schedule. Each uses heuristics: items advance through fixed intervals on success and reset on failure. The heuristics work at the population level; they are coarse at the per-learner, per-item level where the actual half-life varies.^o

THE INTERVENTION

Settles and Meeder's contribution was to treat the half-life as a quantity that can be learned from data. Half-Life Regression (HLR) models the probability that a learner recalls an item at a given delay as an exponential decay whose half-life is a regression on item features (language, part of speech, difficulty proxies) and per-learner practice history. The model is fit on aggregated Duolingo practice data — roughly 13 million learning traces — and produces per-item half-life estimates that drive the scheduling: review when the estimated recall probability falls to a target threshold.¹

HOW IT WORKED

The published evaluation compares HLR against several baselines on two outcomes. The first is Duolingo's predictive metric — how well the scheduler's model of recall matches observed recall on held-out practice data. The second is a downstream behavioral outcome: daily (next-day) learner retention, which Duolingo treats as a proxy for engagement-driven learning continuation. HLR improves on both axes over Leitner, Pimsleur, and a simpler logistic baseline. The paper deploys HLR into the live review-scheduling system; the evaluation includes the production-deployment outcomes, not just offline validation.²

THE EVIDENCE

The case sits in the corpus because it is one of the few published, peer-reviewed deployments of a spaced-repetition algorithm at consumer scale with reported behavioral outcomes. Duolingo's scale at the time of publication — tens of millions of active learners — meant the study had statistical resolution that the academic spaced-repetition literature rarely achieves. The contribution to LENS is the worked example of iterative design in the Domain 2 sense: a learning-engineering model deployed against a measurable behavioral outcome, with the outcome instrumented in the production system rather than inferred from offline validation.³

WHAT TRANSFERRED

The hedges have to survive verbatim. This is a single-vendor study published by Duolingo researchers about a Duolingo product; independent replication at comparable scale is not in the literature. The behavioral outcome is Duolingo's daily next-day retention, which is a sensible proxy for engagement but is not a language-proficiency assessment — the case does not say HLR makes learners more proficient than Leitner does; it says HLR does better on the predictive metric and the return rate. Generalization to other content domains (clinical knowledge, technical skill, vocabulary outside language-learning) rests on the structural argument that per-item half-life estimation should outperform heuristic scheduling, not on a multi-domain replication base. The case teaches the trainable-scheduling form, with the qualification that the strongest available evidence for the form is single-vendor.

HLR does better on the predictive metric and on next-day (daily) retention. The case does not say it makes learners more proficient than Leitner does.

EDITORS' SYNTHESIS OF SETTLES & MEEDER (2016).

REFERENCES

1. Settles, B., & Meeder, B. (2016), "A Trainable Spaced Repetition Model for Language Learning," Proceedings of the 54th Annual
3. Pimsleur, P. (1967), "A memory schedule," Modern Language Journal 51(2):73–75 — graduated-interval recall heuristic.

Meeting of the Association for Computational Linguistics, pp. 1848–1858, doi:10.18653/v1/P16-1174.

2.

Ebbinghaus, H. (1885), Über das Gedächtnis — the empirical forgetting–curve foundation HLR formalizes.

4.

Leitner, S. (1972), So lernt man lernen — Leitner-box spacing heuristic.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Half-Life Regression is one of the few peer-reviewed spaced-repetition deployments at consumer scale with reported behavioral outcomes. The structural argument — learn the half-life rather than fix it by heuristic — is strong; the evidence is single-vendor, the outcome is an engagement proxy not a proficiency assessment, and the cross-domain generalization rests on the structural argument rather than replication.

LENS APPROACH

Duolingo HLR is the trainable-scheduling consumer-scale case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the production-deployment evaluation form, and in Domain 4 (Test and Evaluation) for the engagement-proxy-vs-proficiency-assessment distinction. Pair with the spaced-education clinical RCTs (Case 17) for the cross-domain replication base, and with the machine-teaming AI-delegation typology where the delegation is to a scheduler optimizing on a proxy metric.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
EDUCATION AT SCALE	T D	D1 D2 D3 D4 D5	4	2.3	2 · 4
MACHINE LEARNING		Iterative Development			
MEMORY AND SKILL					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the scheduling parameter (item half-life) as a quantity to learn from data rather than to fix by heuristic. The cost of learning it is the data infrastructure the production system already has; the heuristic was a substitute for missing data, not a principled choice.
- ▶ Instrument the production system to record the practice events that the model needs (item, learner, delay, outcome) so the model can be fit and re-fit on observed behavior at consumer scale.
- ▶ Choose the deployment outcome metric carefully. Duolingo’s choice — predictive recall + next-day retention — is defensible for an engagement-driven product; whether it is the right metric for a proficiency-driven application is a different decision.

IN OPERATION / AFTER

- ▶ Carry the single-vendor hedge into communication about the case. HLR is the strongest available evidence for trainable scheduling at consumer scale; it is not multi-vendor evidence and it is not a language-proficiency study.
- ▶ Distinguish the structural argument (per-item half-life estimation should outperform heuristic scheduling) from the domain transfer claim (HLR specifically generalizes to other learning content), which the published evidence does not yet support.
- ▶ Pair the case with spaced-education clinical RCTs (Case 17) when those are drafted, so the corpus has both a consumer-scale single-vendor deployment and a clinical-domain replication base for the same underlying mechanism.

REFLECTION QUESTIONS

1. Identify a scheduling, dosing, or pacing parameter in your domain that is currently set by heuristic. What data does your production system already record that would let the parameter be learned per-unit rather than fixed?
2. Specify the outcome metric you would use for evaluating a learned scheduler. Duolingo chose engagement proxy + predictive recall; a proficiency-driven application would choose differently, and the choice is part of the evaluation design.
3. The single-vendor hedge is structural to the case. What would a multi-vendor replication base look like in your domain, and what minimum independent evidence would you require before treating a single-vendor production result as the basis for a curriculum-wide claim?

WHO BUILDS THIS learning science & instructional design · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · task analysis · design prototyping

CASE 70

CORPORATE L&D

PERFORMANCE CONSULTING

LEARNING TRANSFER

2001 – PRESENT

High-Impact Learning System — Engineering the Environment, Not Just the Event

K Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT Brinkerhoff & Apking's HILS reframes corporate L&D as a system spanning pre-training (line-manager alignment, work-context preparation), the event itself, and post-training (supervisor support, on-the-job practice) — translating Blume's meta-analytic finding that the work environment dominates transfer into a deployable program model

IN BRIEF

The High-Impact Learning System (HILS), introduced by Brinkerhoff and Apking in 2001, reframes L&D as a system spanning pre-training, the event itself, and post-training. The design principle is that the training event alone explains a small fraction of transfer variance — Blume's 2010 meta-analysis (Case 65) identifies the work environment as the decisive variable — and so the program has to engineer the environment alongside the event. HILS deployments include pre-training line-manager alignment and work-context preparation; the event itself; and post-training supervisor support and on-the-job practice opportunities. Corporate deployments report transfer rates rising from the 10–20% baseline cited in the L&D literature to substantially higher figures, but the per-firm numbers live in vendor whitepapers and conference talks rather than peer-reviewed audits. The case is the deployed-program counterpart to Case 83 (SCM as evaluation instrument): SCM measures whether the program worked at the tails; HILS designs the program so that the conditions for transfer are engineered. Evidence-tier flag is practice-synthesis: the model is documented in practitioner publications and in Watershed and L-TEN summaries, the corporate effect sizes are self-reported, and future validation will continue.

BACKGROUND

Blume et al.'s 2010 meta-analysis (Case 65) names the work environment — supervisor support, peer support, practice opportunity — as the decisive transfer variable at the system layer rather than the learner layer. The L&D practitioner question that follows is what to **deploy** in response to the finding: if the training event is not the capability deliverable, what is the surrounding architecture that has to be built so the transfer actually happens?^o

THE INTERVENTION

Brinkerhoff and Apking's **High Impact Learning** (2001) is one of the answers in practitioner circulation. HILS reframes the L&D pipeline around three phases. **Before the event:** line-manager alignment, work-context preparation, learner readiness. **The event itself:** content delivery designed against the on-job task it is intended to enable. **After the event:** supervisor support, peer-support structures, deliberately engineered practice opportunities on the actual job tasks. The training event is one component, not the whole program.¹

HOW IT WORKED

The corporate adoption story — across firms identified in Brinkerhoff materials and in Watershed LRS and L-TEN summaries — reports transfer rates rising from the 10–20% baseline cited in the L&D literature to substantially higher figures. The per-firm numbers live in vendor whitepapers and conference talks, not in peer-reviewed evaluations. The pattern is consistent across the practitioner sources; the magnitudes are self-reported and have not been independently audited.²

THE EVIDENCE

HILS is the deployed-program counterpart to Case 83 (SCM as evaluation instrument). Where SCM measures whether the program worked at the tails of the outcome distribution, HILS designs the program so that the surrounding conditions for transfer are engineered. The pair is the operational answer to the Kirkpatrick chain-of-evidence problem (Case 79): design the environment so transfer can happen, and measure the tails to confirm it did. Neither component on its own crosses the Level-2 / Level-3 seam — together they structure the cross.³

WHAT TRANSFERRED

The LENS teaching point is the framework-level claim that capability lives at the operator-system interface, not in the individual. HILS is a documented practice pattern that names the environment-as-intervention move; it makes Blume's meta-analytic finding (Case 65) operational. The evidence-tier flag is practice-synthesis — the model and the deployment pattern are durable in practitioner literature, the per-firm effect magnitudes are self-reported, future validation will continue. The CLO **Judgment under inadequate evidence** is exactly the capability HILS asks the practitioner to bring: act on the strongest synthesis the field has produced while naming what is and is not independently audited.⁴

If the work environment is the decisive transfer variable, the work environment has to be a design variable. HILS makes it one.

EDITORS' SYNTHESIS OF BRINKERHOFF & APKING (2001) AND THE BLUME META-ANALYTIC FINDING.

REFERENCES

1. Brinkerhoff, R. O., & Apking, A. M. (2001), High Impact Learning: Strategies for Leveraging Performance and Business Results from Training Investments, Basic Books.
2. Watershed LRS practitioner summaries of HILS deployment patterns.
3. L-TEN (Life Sciences Trainers and Educators Network) practitioner summaries of HILS in life-sciences L&D.
4. Blume et al. (2010), Journal of Management 36(4):1065–1105 — the meta-analytic finding HILS operationalizes (paired Case 65).
5. Brinkerhoff (2005), Advances in Developing Human Resources 7(1):86–101 — SCM as the paired evaluation instrument (Case 83).

THE LEARNING ENGINEERING LENS

LE INSIGHT

HILS is the deployed-program operationalization of Blume’s environment-as-decisive finding (Case 65) and the design-side counterpart of SCM (Case 83). Evidence is practice-synthesis: the model is durable in practitioner literature, per-firm effect sizes are self-reported, future validation continues. The CLO **Judgment under inadequate evidence** is the capability the case asks for.

LENS APPROACH

HILS is the L&D environment-as-design-variable case (induced 2.4; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development / Learning Engineering Design) for the amended sub-competency that narrates the design iteration explicitly across pre / event / post phases, and exercises the CLO **Judgment under inadequate evidence** because the practitioner must decide on practice-synthesis-tier evidence. Pairs with Case 83 (SCM) and Cases 79 / 112 as the corporate-L&D cluster.

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
CORPORATE L&D PERFORMANCE CONSULTING LEARNING TRANSFER	K N	D1 D2 D3 D4 D5 Iterative Development	4	2.4	2 · 4

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Plan the pre-training phase as a first-class design deliverable: line-manager briefings, work-context preparation, learner readiness assessment, instrumented before the event runs.
- Design the post-training phase before delivery: supervisor support artifacts, peer-support structures, scheduled on-job practice opportunities on the specific tasks the training targets.
- Treat the training event as one component of a three-phase architecture, not as the whole program — Blume’s finding (Case 65) is the load-bearing reason the architecture has to span the boundary.

IN OPERATION / AFTER

- Pair with Case 83 (SCM) to measure whether the deployed program worked at the tails — HILS designs the environment, SCM samples the outcomes; together they structure the chain-of-evidence cross (Case 79).
- Carry the practice-synthesis flag honestly: the model is documented and durable, the per-firm effect sizes are self-reported, and any specific magnitude cited from vendor whitepapers should be flagged as such in program documentation.
- Use the CLO **Judgment under inadequate evidence**: the synthesis is the strongest the field has, and the practitioner has to decide whether to deploy HILS at organizational scale on practitioner-tier evidence while naming the validation that remains open.

REFLECTION QUESTIONS

- Identify a training deployment in your context that currently invests heavily in the event and lightly in the surrounding environment. What pre-training and post-training artifacts would you build to convert the deployment from a single-component program into a three-phase HILS-style architecture?
- Specify the line-manager and supervisor briefing materials, the peer-support structures, and the on-job practice opportunities you would design — and identify who in your organization would be accountable for each. The accountability map is part of the design.
- The HILS corporate effect sizes are self-reported. What independent evidence — third-party audit, peer-reviewed evaluation, cross-firm comparison — would you require before treating any specific magnitude as decision-grade, and how would you act on practitioner-tier evidence in the meantime?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development

K Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER preprint-tier — source confidence flagged; future validation ongoing.

IMPACT A 2025 longitudinal preprint study of reflective-practice development across a multi-year work-based software-engineering program — one of the few published instruments aimed at measuring the development of reflective capability rather than only its presence; preprint-tier flag is load-bearing

IN BRIEF

A 2025 arXiv preprint (“The Development of Reflective Practice on a Work-Based Software Engineering Program: A Longitudinal Study,” arXiv:2504.20956) reports a longitudinal study of how reflective-practice capability itself develops across a multi-year work-based software-engineering program. The signal the v2 corpus needs from this case is precise: it is one of the few published instruments aimed at measuring the development of reflective capability over time, not merely its presence at a single point. That is the LENS-revised CLO-2’s evaluation problem in miniature — if the program asks the learner to narrate and defend the design iteration in first person, the program also has to be able to evidence that the capability to do so is developing. The preprint-tier flag is load-bearing: not yet peer-reviewed at the time of this drafting, and the case carries the standing “future validation ongoing” language into print. It is the v2 corpus’s reference instance of an instrument built to measure the **development** of reflective capability over time, not only its presence at a single point — the prior art the editor-commissioned first-person Practice Flywheel accounts will sit alongside.

THE SHIFT

The amended CLO-2 in v2 asks the learner to narrate and defend the design iteration in first person. The evaluation problem this creates is real: most existing reflective-practice instruments measure whether a learner is reflecting at a given moment, not whether the capability to reflect well is developing across a program. A multi-year work-based software-engineering program is a useful site for this question because the work itself provides successive design iterations the learner can reflect on.^o

WHAT IS EMERGING

The 2025 preprint reports a longitudinal study of reflective–practice development across such a program. The methodological move worth naming is the focus on development rather than snapshot: the instrument is designed to detect change in the depth and structure of reflection over time, and the study reports the trajectory across a cohort. That is the LENS–revised evaluation problem in miniature — and one of the few published instances of an instrument built for the purpose.¹

THE CAPABILITY QUESTION

The teaching point is the construct boundary. Measuring the **presence** of reflection is well-trodden ground (rubrics, codings, reflective–essay scoring schemes). Measuring the **development** of reflective capability is a different construct: the comparison is intra–learner across time, the instrument has to be sensitive to depth change rather than presence change, and the study design has to handle the cohort–level variability that arises when learners enter the program with different reflective baselines.²

EARLY EVIDENCE

Why the case sits in the v2 supplemental tier rather than in the verified primary list: it is a preprint. The arXiv version has not been through peer review at the time of this drafting, and the case is included with the preprint–tier flag rendered under the title and the standing “future validation ongoing” language preserved. The signal the corpus extracts is structural — the instrument–design move and the construct distinction — rather than the specific magnitudes the preprint reports.³

OPEN PROBLEMS

What the case supplies the editor–commissioned first–person Practice Flywheel accounts (CIRCUIT, ERKS–class) is the evaluation pathway: evidence that reflective–practice capability can be measured as it develops, not only observed once. That is the prior art a Flywheel account’s evaluation design can build on. Future validation ongoing — both peer–review pipeline and replication across other work–based engineering programs.

The construct is the development of reflective capability, not its presence. The instrument has to be sensitive to depth change.

EDITORS’ PARAPHRASE OF THE ARXIV PREPRINT’S CONSTRUCT DISTINCTION.

REFERENCES

1. “The Development of Reflective Practice on a Work–Based Software Engineering Program: A Longitudinal Study,” arXiv:2504.20956 (2025) — preprint.
2. D. Schön, *The Reflective Practitioner* (1983) — the foundational account of reflection–in–action the genre rests on.
3. Boud, Keogh & Walker (eds.), *Reflection: Turning Experience into Learning* (1985) — reflection as a learning process, and the measurement problem it raises.
4. the proposed revisions — the amended CLO–2 (first–person narration of design iteration) the case evaluates.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The 2025 preprint is one of the few published instruments aimed at measuring the development of reflective capability across a multi–year work–based program, not only its presence at a snapshot. Preprint–tier flag load–bearing — not yet peer–reviewed at time of drafting; the case is included on the structural contri–

bution (construct distinction, instrument–design move) rather than specific magnitudes. Future validation ongoing.

LENS APPROACH

The longitudinal SE–program reflective–practice study is the evaluation–pathway case for first–person practice accounts (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the amended CLO–2 and in Domain 4 (Test and Evaluation) for the CLO **Judgment under inadequate evidence** — the preprint supplies the construct distinction (development vs. presence) without supplying peer–reviewed magnitudes.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
SOFTWARE ENGINEERING EDUCATION WORK-BASED LEARNING REFLECTIVE PRACTICE	K N	D1 D2 D3 D4 D5	4	2.3	2 · 4
	→	→ Iterative Development			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ When designing a program–level evaluation of reflective–practice capability, distinguish the development construct from the presence construct at the instrument–design stage; the preprint is the worked example of why the distinction matters.
- ▶ Specify the temporal cadence of the instrument — when in the program reflections are collected, against what successive iterations of design work — so the development trajectory can be evidenced rather than inferred.
- ▶ Carry the preprint–tier flag through any pedagogical decision the case informs; the structural contribution is the construct distinction, and the specific magnitudes await peer review.

IN OPERATION / AFTER

- ▶ Track intra–learner depth change in reflective output across the program as the primary measurement target; cross–learner presence comparisons are a secondary construct and should be reported as such.
- ▶ Use the preprint as the prior–art reference for any first–person Flywheel evaluation pathway you propose for CIRCUIT, an ERKS–class effort, or a LENS–graduate program; the construct distinction is portable even where the specific instrument is not.
- ▶ When the preprint’s peer–reviewed version appears, update the tier flag and re–evaluate the magnitudes; the case is included on the structural contribution, not on the preprint’s specific figures.

REFLECTION QUESTIONS

1. Identify a multi–year program in your context where reflective–practice capability is supposed to develop. What instrument would you build to measure development (intra–learner depth change over time) rather than presence (snapshot)? At what temporal cadence?
2. The case is a preprint not yet peer–reviewed. What is the minimum additional evidence — peer–review pipeline outcome, replication across other work–based programs, comparison with snapshot–based instruments — you would require before treating any specific magnitude from this work as settled in your program design?
3. Specify the prior art you would assemble around a first–person Flywheel evaluation pathway for an engineering–practice account; which construct distinctions (development vs. presence) would you carry forward, and which would you supplement with locally produced evidence?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension

T Training Gap **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT Van Campenhout et al. (LAK 2023) replicated the doer-effect causal claim across seven courses with 15.2 million data points; L@S 2025 replication held with AI-generated practice questions; LAK 2025 non-WEIRD radio/phone extension found weaker effect for learners with higher prior educational attainment — the load-bearing heterogeneity finding

IN BRIEF

The original “doer effect” causal claim — Koedinger and colleagues at LAK 2016 — held that students who interact with embedded practice activities learn more than students who only read, even after controlling for prior achievement and engagement, and that the effect appears causal. Van Campenhout et al.’s LAK 2023 paper replicated the claim across seven courses with 15.2 million data points; the L@S 2025 follow-up reported the effect held with AI-generated practice questions; the Butler et al. LAK 2025 non-WEIRD extension tested the effect for learners receiving lecture content via community radio and practice via basic mobile phones, and reported that the doer-effect relationship was weaker for learners with higher prior educational attainment — the load-bearing heterogeneity finding the corpus most needs. The case sits with Case 69 (Duolingo half-life) and Case 17 (spaced education RCTs) as the replication-arc thread. The closed loop is closed not by a single trial but by replication; the effect generalizes but not uniformly, and the heterogeneity is itself the teachable result.

BACKGROUND

The doer-effect causal claim is one of the cleanest published claims in the online-learning literature. Koedinger and colleagues, at LAK 2016, drew on student-level data from large online courses and isolated the effect of doing — interacting with embedded practice activities — from the effect of reading, controlling for prior achievement and engagement. The conclusion: doing causally improves learning more than reading does, by an effect size that has stood up across the field’s subsequent replications. The case Van Campenhout et al. open at LAK 2023 is what happens when the claim is treated not as a single result but as a replication target — what does the doer effect look like at scale, across courses, content domains, and delivery modalities?^o

THE INTERVENTION

The LAK 2023 replication is the broadest, by data volume. Van Campenhout et al. assembled data from seven courses and 15.2 million practice and reading interactions, and reported

the doer–effect relationship held in the direction and approximate magnitude the original claim established. The 2025 L@S follow-up addressed a question that did not exist when the original claim was published: what happens to the doer effect when the practice questions are AI-generated rather than human-authored? The L@S 2025 replication reported the effect held with auto-generated content — a meaningful generalization given the rising deployment of LLM-generated practice across educational platforms.¹

HOW IT WORKED

The 2025 LAK non-WEIRD extension by Butler and collaborators is the case’s load-bearing methodological contribution. The extension tested the effect for learners receiving lecture content via community radio and practice via basic mobile phones — a delivery substrate structurally different from the platform-mediated online courses the original claim and the LAK 2023 replication operated on. The doer–effect relationship was weaker for learners with higher prior educational attainment. The heterogeneity is not noise the analysis should adjust away; it is the result. The original causal claim generalizes across delivery modalities, but the effect size is conditioned on prior attainment in a way the WEIRD–population evidence base did not surface.²

THE EVIDENCE

The case is closed-loop in a sense the corpus most often aspires to. The loop is closed not by a single trial but by replication: the original claim, the seven-course large-N replication, the AI-generated-content replication, and the non-WEIRD-modality extension are four converging pieces of evidence. The closed loop is also honest about what it closes and what it does not. Long-term retention across multi-year intervals is not yet in the replication record; transfer beyond the included content domains is not yet in the replication record; the non-WEIRD-modality heterogeneity is documented but not yet decomposed into the components — content familiarity, modality affordances, attention conditions — that would let designers act on it. The case carries those open questions rather than collapsing them.³

WHAT TRANSFERRED

The case anchors with Case 69 (Duolingo half-life) and Case 17 (spaced education RCTs) as the replication-arc thread. All three demonstrate the closed-loop discipline at field scale: a design principle established as a single result, replicated across contexts, and surfaced as conditional on population and modality. The CLO on judgment under inadequate evidence is operative in a productive sense: the original claim was adequate evidence for the WEIRD–platform-online-course context, and the non-WEIRD-modality evidence the LAK 2025 extension adds is what extends the principle’s actionable scope. The case completes the replication-arc thread the corpus needs to teach design-principle generalization honestly.

The doer effect generalizes across delivery modalities, but the effect size is conditioned on prior attainment in a way the WEIRD–population evidence base did not surface. The heterogeneity is the result.

EDITORS’ SYNTHESIS OF VAN CAMPENHOUT ET AL. (2023, 2025) AND BUTLER ET AL. (2025).

REFERENCES

1. Van Campenhout, R., Jerome, B., Dittel, J. S., & Johnson, B. G. (2023), “The Doer Effect at Scale: Investigating
3. Butler, D. et al. (2025), “Does the Doer Effect Generalize To Non-WEIRD Populations? Toward Analytics in Radio and Phone-

Correlation and Causation Across Seven Courses,” LAK23, doi:10.1145/3576050.3576103.

2. Van Campenhout et al. (2025), “Scaling the Doer Effect: A Replication Analysis Using AI-Generated Questions,” L@S ’25, doi:10.1145/3698205.3729545.

Based Learning,” LAK ’25, doi:10.1145/3706468.3706505 (also arXiv 2412.20923).

4. Koedinger, K. R. et al. (2016), original doer-effect causal-claim paper at LAK 2016 — the replication target the present case builds on.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The doer-effect replication arc is the closed-loop-via- replication case in the corpus. Original claim, seven-course large-N replication, AI-generated-content replication, non- WEIRD-modality extension — four converging pieces of evidence with the prior-attainment heterogeneity finding as the load- bearing result. The case completes the replication-arc thread alongside Cases 69 and 17.

LENS APPROACH

Doer-effect replication arc is the closed-loop-by- replication case (induced 2.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the replication- arc discipline and in Domain 4 (Test and Evaluation) for the cross-population generalization-with-heterogeneity structure. Pair with Case 69 (Duolingo half-life) and Case 17 (spaced education RCTs) — the replication-arc thread teaches the closed-loop discipline at field scale.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
LEARNING ANALYTICS	T K D	D1 D2 D3 D4 D5	4	2.3	2 · 4
ONLINE LEARNING	→	→			
MOBILE/DISTANCE LEARNING		Iterative Development			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the design principle as the replication target, not the original trial; the doer-effect case’s structure is built on four converging pieces of evidence rather than on the original claim alone.
- ▶ Pre-specify the cross-context replication conditions before launching the replication: course scope, content domain, delivery modality, and population characteristics are all conditions the original effect was estimated under and that the replication should vary.
- ▶ Test design-principle generalization at the modality boundary before claiming generalization across modalities; the LAK 2025 non-WEIRD extension is the methodological model for that boundary test.

IN OPERATION / AFTER

- ▶ Publish the heterogeneity as the result, not the adjusted-away noise; the non-WEIRD prior-attainment finding is what extends the principle’s actionable scope and is the load-bearing pedagogical content of the case.
- ▶ Build the long-term-retention and cross-content-transfer studies the replication arc names as the next open questions; the closed loop the case completes is honest about what it does and does not close.
- ▶ Carry the case in syllabi alongside Cases 69 and 17 so the replication-arc thread is taught as a thread, not as three independent examples; the closed-loop discipline at field scale is more visible across the three cases than within any one.

REFLECTION QUESTIONS

1. Identify a design principle in your domain that has been replicated within WEIRD platform conditions but not across delivery modalities. What is the modality-boundary test you would design, and what would the analog of the radio-lecture-plus-mobile-phone substrate look like in your domain?
2. Specify the pre-specified heterogeneity outcomes you would build into the next replication of a design principle in your domain. What population characteristic do you expect would condition the effect size, and what would constitute disconfirming evidence?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

Zhang/Scardamalia — Knowledge Building Across Three Cohorts

T Training Gap **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT Three-year design study in a single Grade 4 classroom (Institute of Child Study, Toronto) using Knowledge Forum across an optics curriculum; documented progression from fixed small-group to opportunistic collaboration across cohorts, with associated gains in the depth and distribution of scientific explanations across the class community

IN BRIEF

Zhang, Scardamalia, Reeve, and Messina's 2009 Journal of the Learning Sciences paper reports a three-year design study in a single Grade 4 classroom at the Institute of Child Study in Toronto, working through an optics curriculum supported by the Knowledge Forum platform. Across three successive cohorts, the classroom's collaborative structure progressed from a fixed small-group organization to an opportunistic one in which students convened around emerging questions and dispersed when those questions were addressed. Outcomes on the depth and distribution of scientific explanations across the class community improved in step with the structural progression. The hedges that travel with the case from the source are binding: single teacher, single school, design study not causal in the trial sense; transferability to non-ICS contexts and to teachers without knowledge-building expertise is open. The case is the classroom-scale longitudinal counterpart in the small-tier evidence base; the LE Lens uses it as the JLS- anchored small-tier complement to v1 Case 67.

BACKGROUND

Knowledge Building is the pedagogical program Marlene Scardamalia and Carl Bereiter developed across the 1980s and 1990s, organized around the idea that classrooms can operate as knowledge-creating communities rather than as receivers of established content. The Knowledge Forum platform is the supporting infrastructure: a shared database in which students post conjectures, build on one another's notes, mark rise-above syntheses, and trace the evolution of community ideas. The 2009 Journal of the Learning Sciences paper is the systematic three-year study of how the collaborative structure inside a single Grade 4 classroom developed across cohorts as the teacher and students worked the program through.^o

THE INTERVENTION

The setting is deliberately constrained. A single Grade 4 classroom at the Institute of Child Study in Toronto, a laboratory school with a Knowledge Building tradition; a single teacher

across the three years; an optics curriculum delivered to three successive cohorts; Knowledge Forum as the supporting platform. The design study's interventions across cohorts were not randomized treatment arms but iterated redesigns of how collaboration was organized in the classroom — what kind of grouping the students worked in, how questions were surfaced and addressed, how the classroom community decided what counted as a productive direction. The design-based research method treats this iterative redesign as the unit of inquiry.¹

HOW IT WORKED

The structural progression Zhang and colleagues document is from fixed small-group collaboration to opportunistic collaboration. In the first cohort, students worked in stable small groups that took on assigned subtopics within optics. By the third cohort, the classroom had moved to an opportunistic structure: students convened around emerging questions as those questions surfaced from the community's shared work in Knowledge Forum, and dispersed when the questions had been addressed. The depth of scientific explanations in the community's shared notes improved in step with the structural change, and the distribution of explanatory contribution across students became less concentrated — more students were authoring substantive contributions, and the substantive contributions were reaching further into the underlying physics.²

THE EVIDENCE

The hedges that travel with the case from the source are binding and load-bearing. The study covers a single teacher, a single school, and a single curriculum domain; the design-based research method does not produce a causal estimate in the trial sense, and the authors are explicit that what they report is a developmental account of how the classroom's collaborative practice changed across cohorts, not a controlled comparison. Transferability to non-ICS contexts is open. Transferability to teachers without Knowledge Building expertise is open. The case is not the evidence that the Knowledge Building program works at scale across heterogeneous classrooms; it is the evidence that in a specific well-supported setting, a classroom's collaborative structure can be iteratively redesigned across cohorts in a direction that improves the depth and distribution of scientific explanation.³

WHAT TRANSFERRED

The case sits as the classroom-scale longitudinal counterpart in the small-tier evidence base. Pair with Case 40 (Colorado CTSA team science) for the collaboration-measurement thread at a different scale and domain; with Case 77 (Hybrid Human-AI Tutoring) for the small-tier deployment-success counterpart; with Case 81 and Case 85 (OU Analyse) for the distance-higher-education governance frame at a different population. The LE Lens uses the case as the JLS-anchored small-tier complement to v1 Case 67 (Cognitive Tutor's single-site arc); the two cases together teach that the longitudinal classroom record is the substrate that learning-engineering iterations operate on, and that the iteration unit is the cohort, not the lesson.

The cohort is the iteration unit. The collaborative structure of a single Grade 4 classroom moved from fixed small-group to opportunistic across three years, and the depth and distribution of scientific explanation moved with it.

EDITORS' SYNTHESIS OF ZHANG, SCARDAMALIA, REEVE, & MESSINA (2009).

REFERENCES

1. Zhang, J., Scardamalia, M., Reeve, R., & Messina, R. (2009), "Designs for Collective Cognitive Responsibility in Knowledge-Building Communities," *Journal of the Learning Sciences* 18(1):7–44, doi:10.1080/10508400802581676.

2. Scardamalia, M., & Bereiter, C. (2006), "Knowledge Building: Theory, pedagogy, and technology," in K. Sawyer (ed.), *Cambridge Handbook of the Learning Sciences — programmatic backdrop*.

3. Bereiter, C., & Scardamalia, M. (2014), "Knowledge Building and Knowledge Creation: One Concept, Two Hills to Climb," in S. C. Tan, H. J. So, & J. Yeo (eds.), *Knowledge Creation in Education — extension into the post-2009 program*.

4. Design-based research methodology references — Cobb, Confrey, diSessa, Lehrer, & Schauble (2003), "Design Experiments in Educational Research," *Educational Researcher* 32(1):9–13 — the methodological frame the 2009 study operates inside.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Zhang and colleagues’ three-year design study is the classroom-scale longitudinal record of how a collaborative structure can be iteratively redesigned across cohorts in a direction that improves the depth and distribution of scientific explanation. The hedges are binding — single teacher, single school, design study not causal — and the case’s pedagogical value depends on the hedges being preserved. The cohort is the iteration unit.

LENS APPROACH

Zhang/Scardamalia is the cohort-as-iteration-unit case at classroom scale (induced 2.2; LENS D2/PT4; CLO-2 and CLO-4). LENS uses it in Domain 2 (Iterative Development) for the multi-cohort design-based-research discipline and in Domain 4 (Test and Evaluation) for the depth-and-distribution paired outcome measure. Pair with Case 40 (Colorado CTSA team science collaboration measurement), Case 77 (Hybrid Human-AI Tutoring small-tier deployment), and Cases 81 and 85 (OU Analyse — distance higher-education governance at a different scale). The LE Lens uses it as the small-tier complement to v1 Case 67 (Cognitive Tutor single-site arc).

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
K-12 SCIENCE	T K D	D1 D2 D3 D4 D5	4	2.2	2 · 4
COLLABORATIVE LEARNING	→	→ Iterative Development			
DESIGN-BASED RESEARCH					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the cohort as the iteration unit when the classroom collaborative structure is the design target; the case demonstrates that a year is the right horizon for a meaningful redesign cycle and a multi-year horizon is required to surface the progression.
- ▶ Design the infrastructure — Knowledge Forum’s shared note database — to make the community’s evolving ideas inspectable; the redesign across cohorts depended on the teacher and students having a record of how prior cohorts’ ideas had moved.
- ▶ Track depth and distribution of contribution as paired outcome measures; reporting depth alone hides the equity-relevant distribution change, and reporting

IN OPERATION / AFTER

- ▶ Carry the binding hedges into print without softening; the case’s value to the corpus rests on its specificity, and the transferability questions are open and disclosable as open.
- ▶ Pair with Case 40 (collaboration measurement at team-science scale) so the collaboration-as-design-target thread is taught at both the classroom and team scales.
- ▶ Use the case as the JLS-anchored small-tier complement to v1 Case 67 (Cognitive Tutor’s single-site arc); the cohort-as-iteration-unit lesson is the bridge between the classroom-design literature and the at-scale evaluation literature.

distribution alone hides the substantive-progress signal.

REFLECTION QUESTIONS

1. Identify a classroom or team in your domain whose collaborative structure has not been redesigned across cohorts. What would a cohort-as-iteration-unit redesign cycle look like, and what shared-infrastructure artifact would make the prior cohort's evolution inspectable for the next?
2. Specify the paired outcome measures you would track when collaborative structure is the design target. Depth and distribution travel together in the Zhang case; what is the analogous pair in your setting, and what would reporting one without the other obscure?
3. The case's transferability hedges are open — non-ICS contexts, non-Knowledge-Building-expert teachers. Pick a feature of the ICS setting that the case depends on, and ask: what would have to be true in your setting for the redesign discipline to travel, and what wouldn't?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

CASE 75

K-12 EDUCATION

AI-SUPPORTED INSTRUCTION

EDUCATIONAL EQUITY

2025

Chen et al. — Rural China AI Devices and the Equity-Direction Finding

T Training Gap **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT Quasi-experimental study across 12 schools (4 urban, 8 rural) and 268 teachers, September to November 2024; rural experimental classes gained 17.93% on mathematics and 13.46% on history, while urban experimental classes gained 10.96% on mathematics and 9.55% on history — the rural gain exceeded the urban gain across both subjects

IN BRIEF

Chen, Wu, Chen, and Zhou's 2025 paper in *Frontiers in Psychology* reports a three-month quasi-experimental study of AI-supported instructional devices across 12 schools — 4 urban and 8 rural — in China, involving 268 teachers from September to November 2024. The headline result is the equity-direction finding: rural experimental classes gained 17.93% on mathematics and 13.46% on history, while urban experimental classes gained 10.96% on mathematics and 9.55% on history. The rural gain exceeded the urban gain across both subjects. The honest hedges that travel with the case from the source are load-bearing and preserved in the prose: the horizon is three months; the study covers 12 schools; assignment is non-randomized; the authors acknowledge self-report and observation biases in the measurement instruments. All four authors are at Chinese institutions — Hangzhou City University, Lingnan University, and Huazhong University of Science and Technology — and the case stands as the first-author-from-deployment-country evidence the corpus's non-WEIRD thread most needs.

BACKGROUND

The corpus's non-WEIRD deployment record is structurally thin. The published learning-engineering evidence base is built largely on WEIRD-population studies — Western, educated, industrialized, rich, democratic — and the deployments that operate at meaningful scale outside that population are rarely written up in peer-reviewed journals by authors based in the deployment country. Chen and colleagues' 2025 *Frontiers in Psychology* paper is one of the strongest documented exceptions: a deployment of AI-supported instructional devices across rural and urban Chinese schools, evaluated quantitatively, reported by a four-author team based at Chinese institutions, and published in a peer-reviewed journal with the load-bearing methodological hedges named.^o

THE INTERVENTION

The deployment covers 12 schools — 4 urban and 8 rural — and 268 teachers across the September to November 2024 window. The intervention is AI-supported instructional devices integrated into mathematics and history classroom instruction. The comparison structure is between experimental and control classes within each school setting across both subjects. The headline outcome the paper reports is the equity-direction finding: rural experimental classes gained 17.93% on mathematics and 13.46% on history, while urban experimental classes gained 10.96% on mathematics and 9.55% on history. The rural-over-urban gain pattern is the result the case teaches; the equity direction is unusual in the published EdTech literature, where deployments more often surface as widening the gap between higher- and lower-resourced settings rather than narrowing it.¹

HOW IT WORKED

The honest hedges the case carries are load-bearing and explicit in the source. The horizon is three months — a September to November 2024 window — and the case's pattern is conditioned on that horizon, not on a school-year or multi-year follow-through. The sample is 12 schools, which is right for the within-subject within-school comparison structure the study runs but is not right for prevalence claims about Chinese rural schooling at population scale. Assignment to experimental and control classes is non-randomized; the authors are explicit that they cannot rule out teacher- or school-level selection effects driving the heterogeneity. Self-report and observation biases in the measurement instruments are acknowledged by the authors. All four hedges travel into the prose without softening; the case is published in a peer-reviewed mid-tier journal, and the evidence-tier flag rendering is intentionally not set — the hedges are carried in the case's own argument rather than in the title-bar render.²

THE EVIDENCE

The case's structural placement in the corpus is the equity-direction-finding case at the non-WEIRD deployment seam. Pair with Case 86 (Gándara / community-college predictive equity in AERA Open) for the equity-direction thread at the higher-education scale; with Case 88 (LiveHint AI bias across foundation models) for the bias-surfacing thread in AI-supported instruction; with Case 73 (Doer Effect non-WEIRD radio-and-phone extension) for the non-WEIRD methodological discipline at the heterogeneity-finding axis. The first-author-from-deployment-country structure of Chen et al. is itself the methodological contribution the corpus's non-WEIRD thread needs: the strongest reading of a deployment is the one written by authors with access to the deployment context and published in the journal architecture the field can read.³

WHAT TRANSFERRED

The pedagogical seam the case opens is the equity-direction reading of an AI-supported instructional intervention. The conventional reading — that deployments widen the gap between higher- and lower-resourced settings — is not what Chen and colleagues report on the three-month horizon they studied. The rural gain exceeded the urban gain across both subjects, and the result is the load-bearing teaching point the case anchors. The case does not claim that the pattern generalizes past three months; it does not claim that the pattern holds outside the 12 schools the study covers; it does not claim causal identification past the bounds the non-randomized assignment supports. What it reports is the equity-direction

finding in the deployment window the study describes, and the case uses that finding to anchor the CLO on fairness beyond omission — the finding is the rare published example where an AI deployment in education narrowed rather than widened a between-setting gap on its evaluation horizon, and the methodological structure of the study supports treating the result as the case’s load-bearing observation rather than as a discounted outlier.

Rural experimental classes gained 17.93 percent on mathematics and 13.46 percent on history; urban experimental classes gained 10.96 percent on mathematics and 9.55 percent on history. The rural gain exceeded the urban gain across both subjects — and the load-bearing hedges are three-month horizon, twelve schools, non-randomized assignment, self-report and observation bias acknowledged by the authors.

EDITORS’ SYNTHESIS OF CHEN, WU, CHEN, & ZHOU (2025).

REFERENCES

1.

Chen, R., Wu, Y., Chen, Z., & Zhou, P. (2025), “Advancing educational equity in rural China: the impact of AI devices on teaching quality and learning outcomes for sustainable development,” *Frontiers in Psychology* 16:1588047, doi:10.3389/fpsyg.2025.1588047.

2.

Paired Case 86 (Gándara et al., AERA Open) — community-college predictive equity at the higher-education scale.

3.

Paired Case 88 (LiveHint AI bias across foundation models, AIED 2025) — bias-surfacing in AI-supported instruction.

4.

Paired Case 73 (Doer Effect non-WEIRD LAK 2025 radio-and-phone extension) — non-WEIRD methodological discipline at the heterogeneity-finding axis.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Chen and colleagues’ 12-school three-month deployment of AI- supported instructional devices across rural and urban Chinese classrooms reported an equity-direction finding — rural gain exceeded urban gain across both mathematics and history. The load-bearing hedges are explicit and carried in the prose: three-month horizon, twelve schools, non-randomized assignment, self-report and observation bias acknowledged by the authors. The case anchors the CLO on fairness beyond omission with a rare published equity-narrowing result.

LENS APPROACH

Chen et al. is the equity-direction-finding case at the non-WEIRD deployment seam (induced 8.3; LENS D2/PT4; CLO-4 and fairness beyond omission). LENS uses it in Domain 2 (Iterative Development) for the deployment-on-a-defined- horizon discipline and in Domain 5 (Navigating Sociotechnical Constraints) for the equity-direction reading that anchors the fairness-beyond-omission CLO. Pair with Case 86 (Gándara community-college predictive equity), Case 88 (LiveHint AI bias across foundation models), and Case 73 (Doer Effect non-WEIRD radio-and-phone extension). The case’s pedagogical value depends on the four binding hedges traveling with the result into print.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
K-12 EDUCATION	T K D	D1 D2 D3 D4 D5	4	8.3	4
AI-SUPPORTED INSTRUCTION	→	Iterative Development			
EDUCATIONAL EQUITY					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pre-specify the equity-direction outcome as the load-bearing reading; the case demonstrates that a deployment study can credibly report a rural-over-urban gain on a defined horizon when the methodological hedges are named.
- ▶ Carry the four binding hedges — horizon, sample, assignment structure, measurement bias — in the deployment write-up itself; the case's value to the corpus depends on the hedges being internal to the argument rather than buried in a limitations footnote.
- ▶ Treat first-author-from-deployment-country authorship as a methodological feature; the non-WEIRD evidence base is structurally improved by deployments written up by teams with access to the deployment context.

IN OPERATION / AFTER

- ▶ Commission the longer-horizon and larger-sample replication the three-month 12-school study cannot perform; the equity-direction finding's persistence past three months and across more schools is the validation question the case names.
- ▶ Pair the case in syllabi with Case 86 (Gándara community-college equity) and Case 73 (Doer Effect non-WEIRD extension) so the equity-direction reading is taught across deployment scales and population settings.
- ▶ Use the case to anchor the CLO on fairness beyond omission with a rare published equity-direction finding; the corpus needs equity-narrowing results as visible as equity-widening ones, and the methodological discipline that lets the result be reported credibly is the curricular target.

REFLECTION QUESTIONS

1. Identify a deployment in your domain whose equity-direction finding — narrowing or widening a between-setting gap — has not been pre-specified as a primary outcome. What would the pre-specification require, and what data infrastructure would support it?
2. Specify the four load-bearing hedges you would carry in the prose of a non-WEIRD deployment write-up — horizon, sample, assignment structure, measurement bias. Which of the four is hardest to name credibly in your domain, and what would naming it cost?
3. The case's structural feature is first-author-from-deployment-country authorship. Pick a deployment in your domain that has been written up primarily by external authors, and ask: what would change in the methodological reading if the write-up were led by authors with access to the deployment context?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

CASE 76

MULTIMODAL LEARNING ANALYTICS

CLASSROOM DEPLOYMENT

DESIGN-BASED RESEARCH

2023

Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account

T Training Gap **K** Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER preprint-tier — source confidence flagged; future validation ongoing.

IMPACT First-person practitioner reflection on multiple in-the-wild multimodal learning analytics (MMLA) deployments — eye-tracking, audio capture, spatial positioning in classroom contexts; documents what worked, what failed, what the team would have done differently

IN BRIEF

Martinez-Maldonado et al.'s 2023 arXiv paper, "Lessons Learnt from a Multimodal Learning Analytics Deployment In-the-Wild," is structured as a first-person practitioner reflection on lessons from multiple in-the-wild MMLA deployments — eye-tracking, audio capture, spatial positioning in classroom contexts. The paper documents what worked, what failed, and what the team would have done differently. The case is offered not as a deployment- results case (the deployment outcomes live in adjacent peer-reviewed papers) but as a published-first-person Practice Flywheel exemplar — the genre the front-matter "unpacking is the method" reframing calls for. The case pairs structurally with the reflective-practice cases elsewhere in the v2 supplemental tier and grounds the practitioner-reflection-as-evidence-tier discipline at the LE-specific layer. Preprint-tier evidence-flag is binding on the framing — the arXiv version is preprint, with sections published in adjacent peer-reviewed work, and the standing "future validation ongoing" language applies to both the peer-review pipeline for this version and the broader question of whether the genre takes hold across the LE community.

THE SHIFT

Multimodal learning analytics is the strand of the learning-analytics field where the data substrate moves beyond click-stream and assessment-response logs to include eye-tracking, audio capture, video, and spatial positioning. The lab-deployment record for MMLA is substantial; the in-the-wild classroom-deployment record is structurally different and structurally sparse. What happens when the sensor stack moves from the controlled lab environment into the classroom — where lighting varies, students move, audio overlaps, and the consent architecture has to accommodate the school's operational norms — is a question the published deployment-results papers can answer only partially.^o

WHAT IS EMERGING

Martinez-Maldonado et al.'s 2023 arXiv paper is structured as a first-person practitioner reflection on lessons from multiple in-the-wild MMLA deployments. The paper's content is what the deployment-results papers do not contain: what worked, what failed, what the team would have done differently. The structure is reflective rather than hypothesis-testing — the team narrates the deployments, identifies the recurring patterns, and surfaces the operational knowledge that does not fit into a methods section of a results paper. The honest framing the paper preserves is that the reflective genre is the right vehicle for the kind of operational knowledge the case contains, and that the peer-review structures of the LE field have not consistently supported the genre.¹

THE CAPABILITY QUESTION

The case is offered in the corpus not as a deployment- results case but as a published-first-person Practice Flywheel exemplar. The genre the editor's memo (B1) anticipates — first-person practitioner accounts of deployment iterations, intended to be paired with the front-matter "unpacking is the method" reframing — has structural analogs in adjacent fields (Lutz on reflective journaling, CBE-LSE on reflective-practice primers, SE work-based reflective-practice longitudinals) but has historically been under-published in LE. The Martinez-Maldonado paper is the LE-specific instance of the genre at deployment scale; the case carries it on that basis.²

EARLY EVIDENCE

The structural anchor the case grounds is the sustaining- tacit-capability-across-generations anchor. Practitioner knowledge — what to do when the eye-tracker calibration drifts mid-session, how to design the consent architecture for an audio capture in a classroom of twenty-five students, what the spatial-positioning sensor placement looks like when the classroom layout shifts — walks out the door if it is not narrated. The paper's first-person genre is the narration vehicle, and the case grounds the curriculum's response to the question of how operational knowledge accumulates and propagates across practitioner cohorts. The pair with the CIRCUIT cases (78, 68) is the workforce-and-capability layer; the pair with Case 77 (hybrid human-AI tutoring) is the design-iteration layer.³

OPEN PROBLEMS

The preprint-tier evidence-flag is binding on the framing. The arXiv version is preprint; sections have been published in adjacent peer-reviewed work, but the consolidated lessons-learned synthesis the case carries sits at the preprint tier. The standing "future validation ongoing" language applies along two dimensions. The peer- review pipeline for this specific consolidated synthesis is one dimension. The broader question — whether the first-person practitioner-reflection genre takes hold in the LE community at sufficient scale to function as the Practice Flywheel exemplar the framework names — is the other dimension. The case is included not despite the preprint-tier framing but with it; the framing is part of what the case teaches.

Practitioner knowledge walks out the door if it is not narrated. The first-person reflective genre is the narration vehicle, and the field has not consistently supported it.

EDITORS' SYNTHESIS OF MARTINEZ-MALDONADO ET AL. (2023).

REFERENCES

1. Martinez-Maldonado, R. et al. (2023), "Lessons Learnt from a Multimodal Learning Analytics Deployment In-the-Wild," arXiv:2303.09099 — preprint, sections published in adjacent LAK and IEEE TLT outlets.

2. Worsley, M., & Blikstein, P. (2018), "A multimodal multisensor framework for examining how students engage in design," Journal of Learning Analytics — broader MMLA literature backdrop.

3. Schon, D. (1983), The Reflective Practitioner — the genre's theoretical underpinning, referenced across the reflective-practice case tier.

4. Editors' memo (B1) — Practice Flywheel commissioning structure that the case is offered as a published-first-person exemplar within.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Martinez-Maldonado et al. is the LE-specific published- first-person Practice Flywheel exemplar at MMLA in-the-wild deployment scale. The case is offered not as a deployment- results case but as a genre exemplar — the reflective- practice account at the field's preprint tier. Preprint-tier flag binding; future validation ongoing on peer-review pipeline and on whether the genre takes hold across the LE community.

LENS APPROACH

MMLA in-the-wild is the sustaining-tacit-capability case in the LE-conferences tier (induced 6.3; LENS D2/PT4). LENS uses it in Domain 2 (Iterative Development) for the reflective-narration-of-design-iteration discipline and in Domain 4 (Test and Evaluation) for the evidence-tier discipline binding the preprint-tier framing to the genre's pedagogical role. Pair with Cases 78 and 68 (CIRCUIT workforce-and-capability layer) and Case 77 (hybrid human- AI tutoring design-iteration layer). Preprint-tier flag binding under the title.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
MULTIMODAL LEARNING ANALYTICS	T K N	D1 D2 D3 D4 D5	4	2.2	2-4
CLASSROOM DEPLOYMENT	→ Iterative Development				
DESIGN-BASED RESEARCH					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Narrate the deployment in first person while it is still operating; the operational knowledge the case names is contemporaneous, and the post-hoc reconstruction loses the texture the first-person genre preserves.
- ▶ Treat the reflective paper as a deliverable on par with the results paper; the case demonstrates that the operational knowledge has the same evidentiary status as the methods-section content of a results paper, and the publication structure should support it.
- ▶ Build the consent and ethics architecture around the in-the-wild deployment substrate from the start; the lessons-learned content includes consent-architecture failures that the lab-deployment record does not surface.

IN OPERATION / AFTER

- ▶ Move the consolidated synthesis through the peer-review pipeline; the preprint-tier evidence-flag is binding now, and converting the synthesis to peer-reviewed publication is the validation step the standing language anticipates.
- ▶ Carry the Practice Flywheel exemplar designation into the curriculum's first-person-account commissioning structure; the case is the LE-specific anchor for the genre the editor's memo (B1) calls for, and the curriculum is the vehicle that institutionalizes the genre.
- ▶ Pair the case with the front-matter "unpacking is the method" reframing so the genre's role in the casebook's pedagogical architecture is visible; the case is offered as an exemplar of the genre, and the genre is offered as the curriculum's response to the sustaining-tacit-capability question.

REFLECTION QUESTIONS

1. Identify a deployment in your domain whose operational knowledge — what worked, what failed, what the team would have done differently — has not been narrated outside the team. What would the first-person reflective account look like, and what publication venue would carry it?
2. Specify the consent and ethics architecture you would build into an in-the-wild deployment in your domain that the lab-deployment record would not have surfaced. The case's lessons-learned content includes consent-architecture failures; what would the deployment-substrate-specific architecture look like?
3. The case is offered as a published-first-person Practice Flywheel exemplar. Identify a practitioner in your domain whose operational knowledge you would commission a first-person account from. What would the commissioning structure look like, and what would the account contain that the published-results papers do not?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

CASE 78

NEUROSCIENCE/CONNECTOMICS

HUMAN-AI TEAMING

WORKFORCE TRAINING

2017 – PRESENT

CIRCUIT / MICrONS – The Human Correction Layer at Petabyte Scale

H Human-System Interface**K** Knowledge & Institutional Memory**N** Normalization of Deviance

DISCLOSURE Contribution: an editor has research adjacency to connectomics programs here, and the case originates at an editor's institution (Johns Hopkins / APL). The method is anchored to peer-reviewed work; program training-outcome evidence is institutional documentation, not independent evaluation — the gap rendered as the evidence-tier flag above.

EVIDENCE TIER internal-pm — source confidence flagged; future validation ongoing.

IMPACT Under IARPA's MICrONS program, automated electron-microscopy segmentation produces brain maps too large and too error-prone to deploy without human correction; CIRCUIT trains undergraduate cohorts as the proofreading workforce that is the recovery mechanism for automation failure at petabyte scale

IN BRIEF

Under IARPA's Machine Intelligence from Cortical Networks (MICrONS) program, automated electron-microscopy segmentation produces petabyte-scale brain maps that are too large and too error-prone to be deployed without human verification. APL's CIRCUIT program trains cohorts of undergraduates to proofread and correct these maps; APL's BossDB stores and serves the petabyte-scale data. The learning-engineering content the case carries is the human-in-the-loop correction layer: where automation fails at the petabyte boundary, a trained human capability is the recovery mechanism that makes the data scientifically usable. The honest evidentiary state — preserved as the evidence-tier flag rendered under the title — is that the connectomics method and infrastructure are documented in the peer-reviewed literature, but CIRCUIT's program training outcomes (did the cohorts reliably produce proofreading capability, at what error rates, with what retention, with what transfer to other tasks) sit in institutional documentation rather than peer-reviewed program evaluation. The case is included as a frontier — the structural pattern is the increasingly central question of how to design a human correction layer for generative and automated systems at scale, and future validation will continue as the program-evaluation literature builds out. The COI render under the title is binding: home institution, research adjacency, evidence-tier flag, all visible at the point of reading.

THE SHIFT

Connectomics — the reconstruction of detailed wiring diagrams of brain tissue from electron-microscopy imaging — is the modern incarnation of a problem with the same structural

shape as the early radiology-AI cases (Case 35): automation produces a primary output at scale that no human team could produce manually, and the output is too important and too error-prone to be deployed without verification. The MICrONS program, funded by IARPA, set out to produce reconstructions of cortical microcircuits at the volume scale required to make computational neuroscience comparisons possible.⁰

WHAT IS EMERGING

The automation produces petabytes of data. Automated segmentation labels each voxel with the neuron it belongs to, but at petabyte scale the per-voxel error rate, however low, compounds into a reconstruction that contains many incorrect neuron boundaries — exactly the boundaries the downstream science depends on. The structural form is familiar: the model is good enough to make the project tractable, and not good enough to deploy without a verification layer. The verification layer in this case is human: trained proofreaders work through the segmentation and correct it, neuron by neuron, with tools designed for the task.¹

THE CAPABILITY QUESTION

The CIRCUIT program at APL trained cohorts of undergraduates to be that proofreading workforce. The program is documented across institutional news, program description, and the MICrONS program literature; CIRCUIT trainees and APL infrastructure (BossDB) together produced the human and computational substrate of the verification layer. The learning-engineering content is the human capability as the recovery mechanism for automation failure at the petabyte boundary: where the model fails, the trained human is the design decision that makes the system's output usable.²

EARLY EVIDENCE

The honest evidentiary state is what the evidence-tier flag under the title encodes. The connectomics method, the automated-segmentation literature, and the BossDB infrastructure are documented in peer-reviewed connectomics publications. CIRCUIT's program **training outcomes** — did the cohorts reliably produce proofreading capability at the error rates and retention required, did the trainees transfer to other tasks, what is the program's measured effect on the downstream science — are documented in institutional news and program description rather than in peer-reviewed program evaluation. The case is teachable on the structural pattern; the operating outcomes are at the institutional-program-management tier and should be read at that tier. The COI under the title — research adjacency, home institution — makes the institutional-tier evidence claim auditable.³

OPEN PROBLEMS

The frontier note the case carries is the most forward-looking idea from the pass-4 sweep. The human correction layer for generative and automated systems at scale is going to be a defining capability across an expanding number of domains — automated transcription, AI-generated code, model-extracted structured data, document segmentation, scientific imaging reconstruction. CIRCUIT is one instance of the pattern in neuroscience; the structural question — how to design, staff, train, and govern a human correction layer for the gap between what models produce and what is operationally usable — recurs and is not well-named in the existing curriculum. The frontier flag on this case is the case-grounded basis

for proposing a sub-competency in this area, pairing with the broader v2 AI-delegation typology and the Domain 3 **Delegation with revocation** CLO.⁴

The recovery mechanism for automation failure is a designed human capability, not an emergent workforce.

EDITORS' SYNTHESIS OF THE CIRCUIT / MICRONS PROGRAM RECORD.

REFERENCES

1. MICrONS program literature (IARPA) — connectomics method and automated segmentation evidence base.

2. APL BossDB documentation — petabyte-scale connectomics data infrastructure.

3. CIRCUIT program documentation (JHU Hub, 2017 – present) — institutional/program description; the training-outcome evidence is at this tier and the evidence-tier flag is binding.

4. Cervantes et al. (2023), ASEE Annual Conference — the paired peer-reviewed CIRCUIT case (Case 68).

5. Wachter & Brynjolfsson (2023), JAMA — generative AI verification framing, applicable across domains.

THE LEARNING ENGINEERING LENS

LE INSIGHT

CIRCUIT / MICrONS is the frontier instance of designing the human correction layer for automated systems at scale. The connectomics method is peer-reviewed; the program training outcomes are institutional documentation, and the gap is rendered as the evidence-tier flag under the title. The forward-looking question the case names — how to design the human correction layer for generative and automated systems at scale — recurs across domains and is not well-named in the existing curriculum.

LENS APPROACH

CIRCUIT proofreading is the human-correction-layer frontier case (induced 3.4; LENS D3/PT6) — Domain 3 for **Delegation with revocation**; Domain 4 for the evidence- tier split. Pair with Case 68 and Cases 20, 5, 189, 3. COI binds.

DOMAIN	Modes At Stake	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
NEUROSCIENCE/ CONNECTOMICS	H K N	D1 D2 D3 D4 D5	6	3.4	4 · 3
HUMAN-AI TEAMING	→	→ Human-System Collaboration			
WORKFORCE TRAINING					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the human correction layer as a designed capability deliverable, not an emergent workforce: define the error rate to achieve, the verification protocol, the trainee qualification, and the tooling.
- ▶ Identify in advance which errors the automation will make at scale and design the correction tooling around those error classes — automation failure is structured, and the correction layer should be too.
- ▶ Carry the evidence-tier honesty: peer-reviewed sourcing for the method and infrastructure, institutional-program documentation for the program

IN OPERATION / AFTER

- ▶ Commission peer-reviewed program-evaluation work on the training-outcome side so the institutional record can converge toward independent evidence over time, while the program is operating.
- ▶ Track the human correction layer across automation upgrades: as the model improves, what changes about the kind of correction the trainees do, the error classes, the qualification requirements?
- ▶ Carry the frontier note into the curriculum conversation: the design of the human correction layer for generative and automated systems at scale is a sub-competency the existing framework does not yet

training outcomes, with the gap visible in the case rather than smoothed.

name, and CIRCUIT is one of several emerging instances.

REFLECTION QUESTIONS

1. Identify a domain in your work where automation produces a primary output at a scale that exceeds manual review. What is the human correction layer's design — error classes targeted, verification protocol, trainee qualification, tooling — and which of these decisions are designed vs. emergent?
2. Specify the peer-reviewed vs. institutional-evidence split for a human-correction-layer program you would propose. The connectomics method can be cited from the literature; the program training outcomes will not be. What evidence tier is honest for each layer?
3. The case's frontier note — designing the human correction layer for generative/automated systems at scale — is a sub-competency the existing curriculum does not yet name. What instance from your domain (automated transcription, AI code, structured-data extraction, document segmentation) would be a paired case to anchor this sub-competency against?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · incident analysis · safety dashboards

The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops

K Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT Across a US corporate-training market sized in the >\\$125B/year range, the dominant evaluation framework structurally collapses: most teams stop at Levels 1–2 (reaction + learning) and never measure Level 3 (behavior change) or Level 4 (business results) — measuring the variable that flatters the program

IN BRIEF

Donald Kirkpatrick's four-level model — Reaction, Learning, Behavior, Results — has been the dominant corporate-training evaluation framework for sixty years, increasingly described as a "chain of evidence" from satisfaction through workplace behavior change to business outcomes. The documented systemic pattern is that most organizations stop at Levels 1 and 2. Level 3 (behavior change on the job) and Level 4 (business results) require data that the training organization typically cannot access: longitudinal performance traces, business-unit outcome metrics, line-manager observation. Satisfaction scores do not predict on-job performance, and knowledge retention does not guarantee workplace application. The case is the corporate-scale instance of the enthusiasm-evidence gap and a direct illustration of the revised "decision-grade evidence" framing in the v2 research backbone: the evidence most L&D decisions ride on is structurally sub-decision-grade. The evidence base is practice-synthesis: Devlin Peck, D2L, Valamis, and related evaluation-practice guides documenting the stop-at-L2 pattern, plus the Blume meta-analysis (Case 65) for the transfer half of the chain. The evidence-tier flag is rendered under the title; future validation will continue as the corporate-L&D evaluation literature consolidates.

THE SHIFT

In 1959 Donald Kirkpatrick proposed four levels at which a training program could be evaluated: Reaction (did learners like it), Learning (did they acquire the content), Behavior (did they apply it on the job), and Results (did business outcomes improve). The four levels became the dominant corporate-training evaluation framework and were later described as a "chain of evidence" — each level meant to provide the evidence for the next.^o

WHAT IS EMERGING

The documented systemic pattern in corporate L&D evaluation practice is that most organizations stop at Levels 1 and 2. Practice-literature synthesis across major evaluation-

guidance sources (Devlin Peck, D2L, Valamis, and corporate L&D benchmarking studies) reports that Levels 3 and 4 are attempted by a minority of programs and reported reliably by fewer. The structural reason is that Level 3 (behavior change) requires longitudinal observation that the training organization cannot conduct, and Level 4 (results) requires business outcome data that often sits outside L&D's reach.¹

THE CAPABILITY QUESTION

The consequence is a field that spends heavily on training — US organizations reported >\$125B/year on workplace training, per the ASTD figure cited in Blume et al. (Case 65) — while measuring mostly the variable that flatters the program. Satisfaction scores do not predict on-job performance; knowledge retention measured immediately after training does not guarantee workplace application. The chain of evidence is cited as the framework; in practice the chain is broken at the link between Level 2 and Level 3, and the decisions made on the available data are not decisions about whether training is producing capability change.²

EARLY EVIDENCE

The evidence-tier flag rendered under the case title is load-bearing. The case is documented through evaluation-practice guides synthesizing the stop-at-L2 pattern across many organizations rather than through a single peer-reviewed study of the phenomenon. The pattern is consistent across the sources and is the practitioner consensus, but the magnitudes vary by sector and the sectoral breakdown is in flux as the field consolidates its evaluation evidence. Future validation will continue as the empirical synthesis improves; the case is included with the flag because the pattern itself is teachable and the practitioner consensus is durable.³

OPEN PROBLEMS

What the case teaches in pair with Blume (Case 65) is the structural form of the enthusiasm-evidence gap at corporate scale, and it is the case-grounded basis for the revised “decision-grade evidence” framing proposed in the v2 research backbone. The capability deliverable is an evaluation architecture that crosses the Level 2 / Level 3 seam — by partnering with line management for behavior observation, by instrumenting the workplace tasks the training targets, and by reporting honestly what evidence is and is not available — rather than declaring the chain satisfied at the Level the training organization can control. Until that crossing happens, most corporate L&D decisions are made on structurally sub-decision-grade evidence, and the CLO **Judgment under inadequate evidence** is exactly the capability the practitioner needs.⁴

The chain of evidence is cited as a framework. In practice it breaks at the link between Level 2 and Level 3, and the decisions made on the available data are not decisions about capability.

EDITORS' SYNTHESIS OF THE KIRKPATRICK MODEL AND CORPORATE L&D EVALUATION PRACTICE.

REFERENCES

1. Kirkpatrick (1959–1960), original four-level evaluation series in *Journal of the ASTD*; updated as Kirkpatrick & Kirkpatrick, *Evaluating Training Programs* (3rd ed., 2006).
2. Devlin Peck, “Kirkpatrick Model: A Guide to the Four Levels of Training Evaluation” — synthesis of the stop-at-L2 pattern in corporate practice.
3. Valamis, “Kirkpatrick Model” practitioner guide on evaluation-practice gaps.
4. Blume, Ford, Baldwin, & Huang (2010), “Transfer of Training: A Meta-Analytic Review,” *Journal of Management* 36(4):1065–1105 — the paired peer-reviewed case (65).

3. D2L, “Kirkpatrick’s 4 Levels of Training Evaluation,” practitioner guide documenting the same pattern.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Kirkpatrick stop-at-L2 pattern is the corporate-scale instance of the enthusiasm-evidence gap and the direct illustration of the revised “decision-grade evidence” point: the evidence most L&D decisions ride on is structurally sub-decision-grade. Evidence base is practice-synthesis tier; the pattern is consistent across sources, the magnitudes still consolidating; future validation will continue.

LENS APPROACH

Kirkpatrick is the corporate-L&D evaluation case (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the CLO **Judgment under inadequate evidence** — the practitioner must decide well on structurally sub-decision-grade evidence — and in Domain 5 (Navigating Sociotechnical Constraints) for the Level-2/Level-3 seam that the training organization cannot cross alone. Direct pair with Case 65 (Blume transfer meta-analysis) for the workplace-environment half of the chain.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
CORPORATE L&D	→ K N	→ D1 D2 D3 D4 D5	5	2.1	4
TRAINING EVALUATION		Test and Evaluation			
WORKFORCE DEV					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the evaluation across the Level-2/Level-3 seam before the training is delivered: partner with line management for behavior observation, identify the workplace tasks the training targets, and instrument them.
- ▶ Specify in advance the on-job behavior and business outcome that would count as evidence the training worked, with a reporting cadence long enough for transfer effects to surface.
- ▶ Treat the practice-synthesis evidence base honestly: name the pattern, flag the tier, and do not present a stop-at-L2 outcome as if it were a Level-3 or Level-4 finding.

IN OPERATION / AFTER

- ▶ When Level 3 and Level 4 data are not available, say so plainly in the evaluation report rather than substituting Level 2 metrics; the CLO **Judgment under inadequate evidence** is what the practitioner needs to do well here.
- ▶ Track the work environment as a separate variable — Blume’s meta-analysis (Case 65) names it as decisive — and report the training outcome conditional on environment, not as a property of the training alone.
- ▶ Use the structural pattern to argue for the evaluation architecture investment, not to abandon evaluation. The gap is a capability gap, not an argument against measurement.

REFLECTION QUESTIONS

1. Identify a recent corporate training program in your organization. At which Kirkpatrick level did evaluation stop? What would the Level 3 and Level 4 measurement have required, and who would have had to provide the data?
2. Specify the evaluation architecture you would build to cross the Level-2/Level-3 seam: which on-job behavior, which business outcome, on what cadence, in partnership with which line-management role.
3. The case is practice-synthesis tier. What is the minimum peer-reviewed or program-evaluation evidence you would require before relying on the stop-at-L2 pattern to justify an evaluation-architecture investment in your context?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

Open University 'Ethical Use of Student Data' and OU Analyse

G Governance & Trust **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT The first institutional 'Ethical Use of Student Data' policy in higher education (2014); an eight-principle consent architecture co-designed with students that unblocked predictive analytics on hundreds of thousands of learners and supported documented intervention work

IN BRIEF

Predictive learning analytics at the Open University faced a real governance objection: large-scale processing of student data for intervention, with pre-GDPR scrutiny and credible surveillance concerns. The OU's response in 2014 was to author the first institutional "Ethical Use of Student Data for Learning Analytics" policy in higher education — eight principles built through wide stakeholder consultation including students, framing students as participants rather than data subjects. The consent-and- transparency architecture was the enabling engineering, not a compliance afterthought; the deployment followed because trust was established first. OU Analyse — a weekly machine-learning at-risk prediction system — operated on top of that architecture, and the Analytics4Action framework (Rienties et al., JIME 2016) paired predictions with tutor judgment and documented interventions on modules of 3,000+ students. A 2019 evaluation (Herodotou et al., BJET) across 559 teachers and 14,000+ students examined how degree of system usage related to outcomes. The honest open question, raised by the OU's own researchers, is whether predictive analytics genuinely serves students versus surveils them — a tension that remains contested, and the policy has since been superseded by a broader Data Ethics Policy. Governance as a living artifact, not a solved problem.

BACKGROUND

The Open University runs higher education at distance scale — cohorts of tens of thousands across a single module, with no classroom signal that a student is falling behind. Predictive learning analytics on the engagement and assessment trace can, in principle, surface that signal in time for a tutor to act. The governance objection in the early 2010s was not abstract: the data volumes were large, the scrutiny pre-GDPR but tightening, and the surveillance concern — that institutional analytics treat students as objects rather than partners — was credible enough that several universities had paused similar programs.^o

THE INTERVENTION

The OU's intervention in 2014 was to author the "Ethical Use of Student Data for Learning Analytics" policy, the first such policy in higher education. The development process was the teaching point as much as the document was: wide stakeholder consultation including

the Students Union, eight principles that framed students as participants in the analytics rather than its subjects, and a published artifact the institution could be held to. The policy was the **enabling engineering** for analytics deployment, not a compliance overlay applied to a deployment that was happening anyway.¹

HOW IT WORKED

OU Analyse — a weekly machine-learning at-risk prediction system — operated on top of that architecture. The Analytics4Action framework (Rienties et al., JIME 2016) paired the prediction with tutor judgment rather than treating the model output as a decision: tutors reviewed flagged students, made the intervention call, and documented what action they took. The framework was evaluated on modules of more than 3,000 students. A 2019 evaluation (Herodotou et al., BJET) across 559 teachers and 14,000+ students examined how degree of teacher engagement with the system related to outcomes; the pattern was that engagement, not the raw prediction, was what tracked with intervention success.²

THE EVIDENCE

The honest open question, raised by the OU's own researchers, is whether predictive learning analytics genuinely serves students versus surveils them. The OU's policy did not foreclose that question; it made the deployment legible enough to argue about, and it superseded the 2014 policy with a broader Data Ethics Policy as practice and the scrutiny regime evolved. Governance as a living artifact is the case. The capability deliverable is not a final answer to the surveillance question — there isn't one — but a consent-and-transparency architecture that lets the institution decide and the students participate in the deciding.³

WHAT TRANSFERRED

What OU teaches in the pair (Cases 81 + 80) is the governance-objection diagnostic: when the objection is about trust and accountability — as it was at the OU — good design can dissolve it, and the deployment can proceed under a credibly co-authored consent architecture. The pair's contrast case is the Dutch SyRI (Case 189), where the governance objection was correct: the system was both rights-violating and operationally ineffective, and the District Court of The Hague stopped it on Article 8 ECHR grounds. The diagnostic capability is to tell those two situations apart before deployment, not after.

The consent architecture is the enabling engineering. The deployment is what follows once trust is established.

EDITORS' SYNTHESIS OF THE OU 2014 POLICY AND SLADE & PRINSLOO (2013).

REFERENCES

1. Slade & Prinsloo (2013), "Learning Analytics: Ethical Issues and Dilemmas," *American Behavioral Scientist* 57(10):1510–1529, doi:10.1177/0002764213479366.
2. Open University (2014), "Ethical Use of Student Data for Learning Analytics" — first institutional policy of its kind in higher education.
3. Rienties, Boroowa, Cross, Kubiak, Mayles, & Murphy (2016), "Analytics4Action Evaluation Framework: A Review of Evidence-Based Learning Analytics Interventions at the Open University UK," *Journal of Interactive Media in Education* 2016(1):2, doi:10.5334/jime.394.
4. Herodotou, Hlosta, Boroowa, Rienties, Zdrahal, & Mangafa (2019), "Empowering online teachers through predictive learning analytics," *British Journal of Educational Technology* 50(6):3064–3079, doi:10.1111/bjet.12853 — OU Analyse evaluation across 559 teachers (189 with OUA access) and 14,000+ students in 15 undergraduate courses; average-use teachers benefited students most.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Open University case is the cleanest instance in the sweep of a governance objection dissolved by design. The 2014 ethical-use policy was the enabling engineering; OU Analyse and Analytics4Action operated on top of it; the 2019 evaluation showed teacher engagement tracked with intervention success. The open question — serve vs. surveil — remains contested, and the policy is a living artifact, not a solved problem.

LENS APPROACH

OU is the positive Domain 4 / Problem Type 5 stakeholder governance case (induced 5.1; LENS D4/PT5). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the governance-objection diagnostic — the dissolvable-objection side — and in Domain 4 (Test and Evaluation) for the engagement-tracked-outcome evidence. Direct pair with Case 189 (Dutch SyRI), where the governance objection was correct and design could not have dissolved it. Adjacent to Case 5 (Epic Sepsis Model) as the inverse pattern: OU built consent before deployment; Epic deployed without validation, and the objection that should have been raised wasn't.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER ED (UK)	G K N	D1 D2 D3 D4 D5	5	5.1	4 · 5
LEARNING ANALYTICS	→	→			
DATA GOVERNANCE		Test and Evaluation			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Build the consent-and-transparency architecture as enabling engineering, not as compliance afterthought; the document is the artifact the deployment can be held to.
- ▶ Co-author the policy with the people the analytics will be applied to — students or operators — including their voice in the principles before the system is built.
- ▶ Pair predictions with human judgment by design: the prediction surfaces a candidate; the tutor or operator makes the decision and documents what they did.

IN OPERATION / AFTER

- ▶ Evaluate the system on engagement and intervention quality, not adoption alone; engagement is what tracks with student outcome at the OU.
- ▶ Treat the governance policy as a living artifact: when the scrutiny regime or practice evolves, supersede the policy openly rather than letting it fall into disuse.
- ▶ Keep the open question — does this genuinely serve students or surveil them? — visible in the institutional record; the legitimacy of the program depends on the question staying askable.

REFLECTION QUESTIONS

1. Identify a deployment in your domain that faced a governance objection about trust and accountability. Was the objection dissolved by design — a consent or transparency architecture co-authored with the affected parties — or was it managed as compliance overlay? What did the difference cost?
2. Specify the engagement metric you would track, separately from adoption, to know whether your deployment is producing intervention quality. At the OU it was teacher engagement; what is the analog in your context?
3. The OU's open question — does predictive analytics genuinely serve students or surveil them — remains contested, and the policy has been superseded. What is the institutional discipline that keeps the question visible and the policy a living artifact rather than a one-time document?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

CASE 82

HIGHER EDUCATION (BRAZIL)

LEARNING ANALYTICS

INSTITUTIONAL GOVERNANCE

2024

MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil)

K Knowledge & Institutional Memory**N** Normalization of Deviance

IMPACT MMALA is a maturity model spanning five categories — data management, data analysis, pedagogy, ethics, and privacy (16 process areas, maturity levels 0–4); expert evaluation rated it comprehensive and suitable; three-institution validation exercise in Brazilian universities found it could outline essential practices and support self-assessment for scaling — instrument for responsible adoption, downstream learning outcome open

IN BRIEF

Freitas, Fonseca, Garcia, Pontual Falcão, Marques, Gasevic, and Mello (Journal of Learning Analytics, 2024) developed and validated MMALA — a Maturity Model for Adopting Learning Analytics — designed to let an institution self-assess its readiness across the five categories responsible adoption actually depends on: data management, data analysis, pedagogy, ethics (governance and consent architecture), and privacy — resolved into 16 process areas with maturity levels 0 to 4. Experts evaluated the model as comprehensive and suitable; a three-institution validation exercise at Brazilian universities found that MMALA could outline essential practices and support self-assessment for scaling learning analytics responsibly. The case is one of the corpus's clearest worked examples of governance-as-instrument: a structured artifact an institution can use to convert the abstract goal "we should adopt learning analytics responsibly" into specific assessments of where its current capability sits and what it has to build next. The honest limit preserved verbatim: the instrument is validated by expert opinion and a three-institution exercise, not by longitudinal outcomes of institutions that used it to adopt LA — it is an instrument for responsible adoption, with the downstream effect on student learning still to be measured. The case pairs with the OU policy (Case 81) and the LALA CANVAS (Case 91) as the institutional-instrument layer of the non-US LA governance set.

BACKGROUND

Institutional adoption of learning analytics typically collapses into one of two patterns: a top-down deployment that meets resistance because the institution is not yet ready (governance is thin, staff are not analytically trained, the pedagogical integration was not designed), or a stalled aspiration where leadership wants responsible adoption but does not have a structured way to know what "ready" means. MMALA was built to address the second pattern: to give an institution a structured self-assessment instrument across the

dimensions readiness actually has, so the adoption decision becomes specific rather than aspirational.⁰

THE INTERVENTION

The instrument is dimensional. Infrastructure asks whether the technical capability — data pipelines, secure storage, analytical platforms — is in place at the level the planned adoption requires. Human resources asks whether the institution has the analytical and pedagogical staffing the tools will need. Ethics asks whether the governance and consent architecture has been built. Pedagogy asks whether the analytics integrate with the institution's learning design, or whether they would be bolted on as a separate track. Each dimension is resolved into maturity levels so the self-assessment yields a structured picture of which dimensions are ready and which are not, rather than a single overall readiness score.¹

HOW IT WORKED

The Freitas et al. validation has two halves. The first is expert evaluation: subject-matter experts in learning analytics, institutional research, and educational technology evaluated the instrument as comprehensive (covering the dimensions adoption actually depends on) and suitable (resolved at the right level for institutional self-assessment). The second is a three-institution validation exercise at Brazilian universities, which found that MMALA could outline essential practices and support self-assessment for scaling. The instrument was usable in practice and produced actionable structure for the institutions that piloted it.²

THE EVIDENCE

The honest limit survives verbatim. MMALA is validated by expert opinion and a three-institution exercise, not by longitudinal outcomes of institutions that used it to adopt LA. The instrument's claim is to be a suitable structure for responsible adoption — what the published evidence supports — and the downstream effect on student learning is the next study, not this one. The case is teachable on the instrument and the validation method; it is honest that the causal chain from "institution adopted via MMALA" to "students learned more" is not yet closed by evidence.³

WHAT TRANSFERRED

In the non-US LA governance set — OU (Case 81, policy and operating system), LALA (Case 91, participatory adoption framework), African data privacy (Case 89, cross-regime seam), Norway (Case 92, national commission), and MMALA (this case, institutional maturity model) — the cases together teach the level structure of governance for the same underlying capability. MMALA's specific contribution is the institutional-instrument layer: when a single institution is the unit of adoption and the governance artifact has to be usable by the institution itself, the maturity model is the form. The pentad demonstrates that governance is producible at multiple levels for the same capability, and the choice of level is itself a governance decision.

The instrument converts "we should adopt responsibly" into a structured per-dimension self-assessment. The instrument is validated; the downstream learning outcome is the next study.

EDITORS' SYNTHESIS OF FREITAS ET AL. (2024).

REFERENCES

1. Freitas, Fonseca, Garcia, Pontual Falcão, Marques, Gasevic, & Mello (2024), “MMALA: Developing and Evaluating a Maturity Model for Adopting Learning Analytics,” Journal of Learning Analytics 11(1):67–86.

2. Open University Ethical Use of Student Data policy (2014) — institutional-policy companion (Case 81).

3. Hilliger et al. (2020), Internet and Higher Education — LALA participatory adoption companion (Case 91).

4. Norwegian Expert Commission on Learning Analytics, final NOU (2023) — national-scale companion (Case 92).


THE LEARNING ENGINEERING LENS

LE INSIGHT

MMALA is the institutional-instrument instance of governance for learning-analytics adoption: a structured maturity model across five categories — data management, data analysis, pedagogy, ethics, and privacy. The validation evidence — expert evaluation and three-institution pilot — is what the case claims; the downstream effect on student learning is the next study, and the case is honest about that gap.

LENS APPROACH

MMALA is the institutional-maturity-model case (induced 5.4; LENS D4/PT4). LENS uses it in Domain 4 (Test and Evaluation) for the structured self-assessment instrument and the two-layer validation discipline; in Domain 5 (Navigating Sociotechnical Constraints) for the governance-as-instrument framing; and on the CLO **Judgment under inadequate evidence** for the instrument-validation-vs-outcome-validation tier distinction. Pair with Cases 81, 91, 89, 92 as the non-US LA governance pentad; MMALA is the institutional-instrument layer.



APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Resolve readiness into the categories adoption actually depends on — data management, data analysis, pedagogy, ethics, privacy — rather than a single overall score; the per-category structure is what makes the instrument actionable.
- ▶ Build maturity levels within each dimension so the self-assessment yields a structured next-step picture rather than a binary ready / not-ready.
- ▶ Validate the instrument in two layers: expert evaluation for comprehensiveness and suitability, and institutional pilot for usability — both are necessary, and the case carries both.

IN OPERATION / AFTER

- ▶ Report the evidence at its tier: expert opinion and three-institution validation establish the instrument; longitudinal outcome evidence on institutions that used it is the next study, and the case names the gap rather than glossing it.
- ▶ Commission follow-up studies on the institutions that pilot MMALA to convert the instrument-validation evidence into adoption-outcome evidence over time.
- ▶ Carry the level-of-governance lesson into the broader v2 governance set: institutional-instrument, multi-country participatory, national commission, and judicial review are different governance modes for the same capability; MMALA is the institutional-instrument instance.

REFLECTION QUESTIONS

1. Identify a capability your institution is considering adopting where readiness is currently discussed as a single overall question. What are the dimensions readiness actually has, and what maturity levels within each dimension would yield a structured next-step picture rather than a binary judgment?
2. Specify the two-layer validation you would build for the instrument: expert evaluation for comprehensiveness and suitability, and institutional pilot for usability. Where would each layer be at risk of being skipped, and what would that skip cost the instrument's credibility?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

Brinkerhoff Success Case Method — Tails as the Evaluation Instrument

K Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT When ROI-style evaluation of corporate training is intractable, Brinkerhoff's Success Case Method samples the tails of the outcome distribution — the highest- and lowest-impact participants — and reconstructs the system conditions that made the program work for some and fail for others; deployed at Cargill, Ford, Merck, World Bank, ICRC

IN BRIEF

The Success Case Method (SCM), introduced by Robert Brinkerhoff in his 2003 book and formalized in a 2005 peer-reviewed paper, deliberately samples the tails of a training program's outcome distribution. Instead of attempting to derive an average effect that flatters most programs and gives L&D nothing actionable, SCM identifies the highest- and lowest-impact participants, studies them in detail, and reconstructs the system conditions that made the program work for some and fail for others. The method's argument — and the one that places it inside the LENS framework — is that extreme cases reveal whether the program ever produces meaningful work-performance change and what the surrounding system has to provide for transfer to happen. The method is peer-reviewed in **Advances in Developing Human Resources**; the corporate deployments at Cargill, Ford, Merck, World Bank, and the International Committee of the Red Cross are documented in practitioner channels — case-study writeups, conference talks, vendor whitepapers — rather than in peer-reviewed evaluation literature. SCM is the operational answer to the chain-of-evidence problem named in Case 79 (Kirkpatrick): the practitioner instrument that crosses the Level-2 / Level-3 seam by sampling where the evidence is most informative. Evidence-tier flag is practice-synthesis; future validation will continue as more firms publish their SCM outcome data.

BACKGROUND

Corporate L&D evaluation faces a structural problem the Kirkpatrick chain-of-evidence pattern names (Case 79): Level 3 (behavior change on the job) and Level 4 (business results) require longitudinal data the training organization typically cannot access. Average-effect statistics across whole programs flatter most programs, because the distribution of training outcomes is heavily skewed — a minority of participants produce most of the work-performance change, and the average buries that signal under the participants for whom the program produced nothing.^o

THE INTERVENTION

Robert Brinkerhoff's 2005 **Advances in Developing Human Resources** paper proposes a different sampling logic: do not estimate the average; identify the participants in the tails of the outcome distribution and study them in detail. The Success Case Method identifies the highest-impact participants (where did the program work, and what conditions made the work-performance change possible), the lowest-impact participants (where did the program fail, and what conditions were missing), and reconstructs the system around each.¹

HOW IT WORKED

The method's argument is that the tails carry the decision-grade information. A program that produces meaningful work-performance change for any participants proves it **can** work; the conditions that distinguish the success cases from the failure cases name what the surrounding system has to provide for transfer. The argument is the practical complement of the Blume meta-analytic finding (Case 65) that work environment is the decisive transfer variable: SCM operationalizes the finding by sampling where the variable's effects are most visible.²

THE EVIDENCE

The Brinkerhoff Evaluation Institute lists deployments at Cargill, Ford, Merck, the World Bank, and the International Committee of the Red Cross. These deployments are documented in case-study writeups, conference talks, and vendor whitepapers, not in peer-reviewed program evaluations. The evidence-tier flag is therefore practice-synthesis: the method itself is peer-reviewed, but the per-firm impact data on which the corporate adoption story rests live in practitioner channels. The honest framing in print is that the methodological pattern is teachable and durable, but the per-firm effect sizes are not independently audited.³

WHAT TRANSFERRED

The LENS teaching point pairs SCM with the chain-of-evidence problem from Case 79 and the meta-analytic transfer finding from Case 65. SCM is the operational instrument that crosses the Level-2 / Level-3 seam by sampling where the evidence is most informative — and it exercises the CLO **Judgment under inadequate evidence** directly, because the method asks practitioners to act on detailed case reconstructions rather than wait for population-scale causal estimates that corporate L&D often cannot produce. Future validation will continue as more firms publish their SCM outcome data.

The tails carry the decision-grade information. The average flatters the program.

EDITORS' SYNTHESIS OF BRINKERHOFF (2005) AND THE CORPORATE SCM DEPLOYMENTS.

REFERENCES

1. Brinkerhoff, R. O. (2005), "The Success Case Method: A Strategic Evaluation Approach to Increasing the Value and Effect of Training," *Advances in Developing Human Resources* 7(1):86–101, doi:10.1177/1523422304272172.
2. Brinkerhoff Evaluation Institute deployment list — Cargill, Ford, Merck, World Bank, International Committee of the Red Cross — practitioner channel.
3. Kirkpatrick & Kirkpatrick (2006), *Evaluating Training Programs* — the chain-of-evidence framework SCM operationalizes (paired Case 79).
4. Blume, Ford, Baldwin, & Huang (2010), *Journal of Management* 36(4):1065–1105 — the meta-analytic environment-as-decisive finding SCM samples around (paired Case 65).

THE LEARNING ENGINEERING LENS

LE INSIGHT

SCM is the practitioner instrument that operationalizes Blume’s environment-as-decisive finding (Case 65) by sampling the tails of the outcome distribution. The method is peer-reviewed; the per-firm impact data at Cargill, Ford, Merck, World Bank, ICRC live in practitioner channels. Evidence-tier flag is practice-synthesis; future validation will continue as more firms publish.

LENS APPROACH

SCM is the corporate-L&D tail-sampling case (induced 2.1; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the CLO **Judgment under inadequate evidence** — the method asks practitioners to act on case-reconstruction evidence rather than population estimates — and in Domain 2 (Iterative Development) by way of the evaluation-feedback loop SCM enables. Pairs with Case 70 (HILS) for the redesign of the surrounding work environment.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
CORPORATE L&D	K N	D1 D2 D3 D4 D5	5	2.1	4 · 2
TRAINING EVALUATION	→	→ Test and Evaluation			
WORKFORCE DEV					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Before the program is delivered, plan the SCM sampling: define the outcome metric (specific work-performance change), identify how the success and failure cases will be surfaced (manager nominations, performance data, structured interviews), and pre-commit to the sampling logic.
- ▶ Structure the success-case and failure-case interviews around the system conditions Blume’s meta-analysis (Case 65) names as decisive — supervisor support, peer support, practice opportunity, environment — so the reconstruction is theory-grounded, not anecdote-grounded.
- ▶ Report the tails honestly: the success cases are existence proofs the program **can** work; the failure cases are evidence of what the surrounding system did not provide. Neither is a population effect size.

IN OPERATION / AFTER

- ▶ Translate the tail evidence into a redesign of the surrounding system — supervisor briefings, peer-support structures, on-job practice opportunity — rather than into a redesign of the training event alone. Pair with Case 70 (HILS).
- ▶ Use the CLO **Judgment under inadequate evidence**: act on the case-reconstruction evidence the method produces while being explicit that it is not a population causal estimate; document the uncertainty.
- ▶ Carry the practice-synthesis flag honestly in any program documentation citing SCM corporate deployments — the method is peer-reviewed, but the per-firm effect sizes are not.

REFLECTION QUESTIONS

1. Identify a recent training program in your organization. Define how you would identify the highest- and lowest-impact participants and what structured interview you would conduct with each to reconstruct the system conditions around them.
2. Map the SCM tails to the Blume meta-analytic variables (Case 65): supervisor support, peer support, practice opportunity, environment. What pattern would the tails have to show for you to redesign the surrounding system rather than the training event itself?
3. The corporate SCM deployments at Cargill, Ford, Merck, World Bank, ICRC are documented in practitioner channels. What additional evidence — independent program audit, peer-reviewed evaluation, cross-firm comparison — would you require before treating any specific per-firm effect size as decision-grade?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive

T Training Gap **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT Cluster-randomized 147 schools (middle and high schools) across seven states; year-one posttest scores showed no significant difference between Cognitive Tutor and control conditions; year-two posttest scores showed CTAI high schools significantly outperforming controls (the middle-school effect was not statistically significant); a one-year evaluation would have published the wrong answer

IN BRIEF

Pane, Griffin, McCaffrey, and Karam's 2014 paper in *Educational Evaluation and Policy Analysis* reports the central at-scale evaluation of Cognitive Tutor Algebra I (CTAI), the canonical learning-sciences-to-classroom translation. The RAND team cluster-randomized 147 middle and high schools across seven states to continue their current Algebra I curriculum or to adopt CTAI for two years. Year-one posttest scores: no significant difference between conditions. Year-two posttest scores: CTAI high schools significantly outperformed controls (the middle-school effect was not statistically significant). The honest reading the authors press is that a one-year evaluation would have published a null result and the two-year evaluation surfaced a real effect — and both findings were in the same trial. The case is the deeper-evidence-of update on v1 Case 67 (Cognitive Tutor), translating that case from a system-description case into a methodological-discipline case about evaluation horizons. The timeline of the evaluation is itself a falsifiable design choice; the case grounds the CLO on judgment under inadequate evidence where the inadequacy is the evaluation horizon, not the sample size.

BACKGROUND

Cognitive Tutor is the case the learning-sciences-to-classroom translation thread cannot avoid teaching at v1 Case 67. The system instantiates a model of student cognition, drives a problem-by-problem adaptive curriculum, and is the canonical published instance of an intelligent tutoring system in K-12 mathematics. The question the v1 case left open — and the question the Pane et al. 2014 paper answers — is what happens when the system is deployed at scale, in the classroom and curricular conditions that classroom adoption actually generates, rather than in the developer-supported conditions of an early efficacy trial.^o

THE INTERVENTION

The RAND team built the evaluation to support that question. 147 middle and high schools across seven states were cluster-randomized to either continue their current Algebra I curriculum or to adopt CTAI for two years. The sample was deliberately wide — multiple states, multiple district contexts, classroom adoption with the support structures the publisher could realistically supply. The outcome instrument was an Algebra I posttest applied at end of school year. The design supports an effect-size estimate at end of year one and a separate estimate at end of year two, with the same students in the same treatment condition for both years.¹

HOW IT WORKED

The year-one result was a null. Posttest scores in CTAI schools did not differ significantly from posttest scores in control schools, across the cluster-randomized sample. The year-two result was different. Posttest scores in CTAI high schools were significantly higher than in control schools (the middle-school effect was not significant), with an effect size that an Algebra I curriculum evaluation would treat as meaningful. The two findings were generated by the same trial, run in the same schools, with the same instrument. The methodological discipline the case teaches turns on what the field would have learned if the evaluation had been designed to a one-year horizon: it would have published a null result against an intervention that the same trial, on its planned horizon, established as positive.²

THE EVIDENCE

The authors press the honest reading without softening it. The timeline of the evaluation is itself a falsifiable design choice, not a methodological default. The reasons year one returned a null are recoverable from the trial record — teacher fluency with the system grew across the year, district-level scheduling and pacing adjusted across the year, student familiarity with the adaptive workflow stabilized across the year. None of these are noise; all are part of the deployment substrate the intervention depends on, and all required the second year to stabilize. The two-year horizon was the right horizon because the intervention's effect is conditioned on a deployment-substrate stabilization that takes more than one year.³

WHAT TRANSFERRED

The case is the deeper-evidence-of update on v1 Case 67. v1 Case 67 carries Cognitive Tutor as the system description; this case carries the at-scale evaluation and the timeline-of-evidence lesson. The CLO on judgment under inadequate evidence is directly motivated: the one-year evaluation would have been inadequate evidence, and a field that publishes the one-year null without noting the planned second-year horizon publishes a wrong answer on the available data. Pair with Case 72 (ASSISTments) for the multi-year follow-through paired case, and with Case 5 (Epic Sepsis) for the evaluation-horizon discipline in clinical AI. The closed loop the case completes is the two-year-horizon-was-the-right-horizon record that lets a field design the next at-scale evaluation honestly.

A one-year evaluation would have published a null. The two-year evaluation surfaced a real effect. Both findings were in the same trial.

EDITORS' SYNTHESIS OF PANE ET AL. (2014).

REFERENCES

1. Pane, J. F., Griffin, B. A., McCaffrey, D. F., & Karam, R. (2014), "Effectiveness of Cognitive Tutor Algebra I at Scale," Educational Evaluation and Policy Analysis 36(2):127–144, doi:10.3102/0162373713507480.

2. RAND Working Paper WR-1050 — addendum to the Pane et al. evaluation.

3. Koedinger, K. R., Anderson, J. R., Hadley, W. H., & Mark, M. A. (1997), "Intelligent tutoring goes to school in the big city," IJAIED — the v1 Case 67 system description Cognitive Tutor builds from.

4. What Works Clearinghouse — Cognitive Tutor evidence-base summary applying federal-grade evidence standards to the Pane et al. trial.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Pane et al. is the at-scale evaluation case for Cognitive Tutor and the worked example of evaluation horizon as a falsifiable design choice. Year one: null. Year two: significantly positive. The case is the deeper-evidence-of update on v1 Case 67 and the curriculum’s primary anchor for the CLO on judgment under inadequate evidence where the inadequacy is the horizon, not the sample size.

LENS APPROACH

Pane / CTAI at scale is the closed-loop evaluation-horizon case (induced 2.3; LENS D4/PT4). LENS uses it in Domain 4 (Test and Evaluation) for the horizon-as-design-choice discipline and in Domain 2 (Iterative Development) for the deployment-substrate-stabilization frame — teacher fluency, scheduling, workflow familiarity are designable parts of the intervention substrate, not noise. Pair with Case 72 (ASSISTments follow-through), v1 Case 67 (Cognitive Tutor system), and Case 5 (Epic Sepsis horizon discipline).

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
K-12 MATHEMATICS	T K D	D1 D2 D3 D4 D5	4	2.3	2 · 4
INTELLIGENT TUTORING	→	→	Test and Evaluation		
EDUCATION POLICY					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Plan the evaluation horizon as a design decision, not a budget default; the case demonstrates that the horizon choice is causally consequential and that a shorter horizon can produce a wrong-direction finding.
- Document the deployment-substrate stabilization that the intervention depends on: teacher fluency, scheduling adjustment, student familiarity with the workflow are all year-over-year stabilization mechanisms the trial should track.
- Design the trial so the year-one and year-two estimates are both interpretable; the case’s pedagogical value depends on having both estimates from the same trial, not on having only the longer-horizon estimate.

IN OPERATION / AFTER

- Publish the year-by-year result pattern, not the aggregate effect; the case’s teaching power depends on the year-one null being on the same page as the year-two positive.
- Carry the timeline-of-evidence lesson into the curriculum’s evaluation-design teaching; the case grounds the CLO on judgment under inadequate evidence with a rare worked example where the inadequacy is the horizon, not the sample size.
- Pair the case with Case 72 (ASSISTments multi-year follow-through) so the field-level discipline is taught with two converging examples; the methodological lesson is more useful as a pattern than as a single instance.

REFLECTION QUESTIONS

1. Identify an intervention in your domain whose effect depends on a deployment-substrate stabilization that takes more than a year (teacher fluency, workflow familiarity, scheduling alignment). What evaluation horizon is honest for that intervention, and what would the year-one estimate, taken alone, falsely conclude?
2. Specify the year-by-year result pattern your next evaluation will report. The case's pedagogical value depends on having both the year-one and year-two estimates from the same trial; what reporting structure would you commit to that supports the pattern view, not only the aggregate?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

OU Analyse — Predictive Learning Analytics and Teacher Use at Scale

T Training Gap **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT Across 15 courses and 559 teachers (189 with OU Analyse access), teachers' engagement with predictive learning analytics was associated with measurable improvements in student performance for >14,000 students; three-year post-implementation follow-up extends the evidence into sustained adoption and perceptions

IN BRIEF

Herodotou et al.'s 2019 British Journal of Educational Technology paper reports the evaluation of OU Analyse, the Open University UK's predictive-learning-analytics dashboard for online tutors, across 15 courses and 559 teachers (189 of whom had access to OU Analyse) with more than 14,000 students. Teachers' engagement with the predictive learning analytics was associated with measurable improvements in student performance. The LAK 2023 three-year-post-implementation follow-up extends the picture into questions of sustained adoption — how teachers' use of the predictions stabilized, what fraction continued to act on them, how perceptions evolved. The case is distinct from the OU consent-and-ethical- use frame Case 81 covers; this case carries the post- deployment teacher-use evaluation at multi-cohort scale. The authors' hedges are binding: causal attribution to OU Analyse use specifically — versus teacher selection effects — is bounded; the 2019 study uses propensity-style matching rather than RCT randomization. Future validation ongoing on multi-institution transfer as the system is licensed beyond the Open University.

BACKGROUND

Predictive learning analytics dashboards have been deployed across higher-education institutions for more than a decade. The deployment record is largely a collection of pilots and single-cohort studies, with Course Signals (the discontinued Purdue deployment) as the structural cautionary case. OU Analyse is structurally different: deployed across an entire distance-learning institution, used by hundreds of tutors across multiple cohorts and multiple courses, and evaluated at the scale the institutional deployment supports. The 2019 Herodotou et al. paper is the central peer-reviewed evaluation of the teacher-use side of the deployment.⁹

THE INTERVENTION

The evaluation design covers 15 courses, 559 teachers, and more than 14,000 students. 189 of the 559 teachers had access to OU Analyse; the comparison structure is between these

teachers and the OU Analyse-naïve teachers in the same courses. The outcome is student performance on course assessments. Teachers' engagement with the predictive learning analytics — operationalized as dashboard-usage and acted-upon-prediction proxies — was associated with measurable improvements in student performance for the cohorts those teachers taught. The effect size is meaningful at the scale the institutional deployment supports.¹

HOW IT WORKED

The 2023 LAK three-year-post-implementation follow-up extends the picture into a question the single-cohort pilots cannot address: what happens to teacher use of the predictions across multiple cohorts and across years? The follow-up paper documents how teachers' use stabilized, what fraction continued to act on the predictions, and how perceptions evolved as the institutional norm around the dashboard solidified. Across the three years, the pattern surfaces as one of sustained adoption with stratification: a fraction of teachers used the predictions actively, a fraction used them as background reference, and a fraction did not engage. The delegation-with-revocation structure is operative: each teacher chose, in each cohort, whether to act on the prediction, and the choice itself is the load-bearing capability the dashboard supports.²

THE EVIDENCE

The case is distinct from the OU consent-and-ethical-use frame Case 81 covers. Case 81 carries the governance story — the Policy on Ethical Use of Student Data, the institutional review of analytic deployments, the consent architecture. This case carries the post-deployment teacher-use evaluation at multi-cohort scale. The editor's decision (memo A6) anchors LENS primary, induced secondary, CLO carried; the two cases together teach the deployment across both governance and post-deployment-use frames without collapsing one into the other. The frame the present case adds to the corpus is the rare successful learning-engineering intervention with both deployment scale and longitudinal teacher-use evidence at the journal-graded evidence tier the corpus needs.³

WHAT TRANSFERRED

The authors' hedges are binding. The causal attribution to OU Analyse use specifically — as distinct from teacher selection effects — is bounded. The 2019 study uses propensity-style matching but does not randomize at the teacher level; the 2023 follow-up extends the evidence base on adoption and perception but does not close the teacher-selection question. Future validation ongoing on multi-institution transfer — OU Analyse has been licensed beyond the Open University, and the transfer outcomes are not yet in the peer-reviewed literature. Pair with Case 81 (OU consent governance), Case 84 (Cognitive Tutor at-scale evaluation), and the Purdue Course Signals cautionary case from v1 — the OU Analyse evaluation is the methodologically more careful descendant of the Course Signals lineage.

The delegation-with-revocation structure is operative: each teacher chose, in each cohort, whether to act on the prediction. The choice itself is the load-bearing capability the dashboard supports.

EDITORS' SYNTHESIS OF HERODOTOU ET AL. (2019, 2023).

REFERENCES

1. Herodotou, C., Hlosta, M., Boroowa, A., Rienties, B., Zdrahal, Z., & Mangafa, C. (2019), "Empowering online teachers through predictive learning analytics," *British Journal of Educational Technology* 50(6):3064–3079, doi:10.1111/bjet.12853.

2. Herodotou, C. et al. (2023), "Predictive Learning Analytics and University Teachers: Usage and perceptions three years post implementation," *LAK '23*, doi:10.1145/3576050.3576061.

3. Herodotou et al. (2019), "A large-scale implementation of predictive learning analytics in higher education: the teachers' role and perspective," *Educational Technology Research and Development*, ERIC EJ1227972 — complementary teacher-perspective paper.

4. Arnold, K. E., & Pistilli, M. D. (2012), "Course Signals at Purdue," *LAK '12* — the structural precursor v1 carries as a cautionary case.

THE LEARNING ENGINEERING LENS

LE INSIGHT

OU Analyse is the rare successful learning-engineering intervention with both deployment scale and longitudinal teacher-use evidence at journal tier. The 2019 BJET paper establishes the effect on student performance at the institutional deployment; the 2023 LAK follow-up extends the picture into three-year sustained adoption. The delegation-with-revocation structure is operative and teachable. Hedges binding on causal attribution and on multi-institution transfer.

LENS APPROACH

OU Analyse teacher-use is the human-in-the-loop predictive- analytics case at institutional scale (induced 5.2; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the multi-cohort longitudinal-follow-through design and in Domain 3 (Human-System Collaboration) for the delegation-with-revocation structure — the teacher chooses, each time, whether to act on the prediction. Pair with Case 81 (OU consent governance), Case 84 (Cognitive Tutor at-scale), and Case 72 (ASSISTments multi-year).

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER EDUCATION	T K D	D1 D2 D3 D4 D5	5	5.2	4 · 3
LEARNING ANALYTICS		Test and Evaluation			
DISTANCE LEARNING					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat teacher use of the prediction as the load-bearing capability the dashboard supports; the delegation-with-revocation structure — each teacher chooses, in each cohort, whether to act — is the design point, not an emergent behavior.
- ▶ Build the deployment at multi-cohort and multi-course scale from the start; the case's evidence-tier strength depends on the deployment having been institutional, not pilot, and the institutional deployment supports the longitudinal-follow-through evidence the case provides.
- ▶ Design the dashboard's prediction surface around the action the teacher would take — a flagged-student list with an intervention prompt, not a probability score the teacher has to interpret unaided.

IN OPERATION / AFTER

- ▶ Commission the next-tier evaluation that addresses the teacher-selection-effect question; an RCT at the teacher level, or an instrumental-variable design exploiting institutional dashboard-rollout variation, would close the residual causal-attribution gap.
- ▶ Publish the multi-institution-transfer evaluations as the system is licensed beyond the OU; the case's value as a generalizable instance depends on the transfer evidence the OU-internal evaluation cannot provide.
- ▶ Pair the case in the curriculum with Case 81 (OU consent governance) so the deployment is taught across both governance and post-deployment-use frames; the two cases together teach the institutional anchor without collapsing one frame into the other.

REFLECTION QUESTIONS

1. Identify a predictive-analytics dashboard in your domain whose operator-use stratification — active, background reference, non-engagement — has not been longitudinally tracked. What would the multi-cohort follow-up look like, and what data infrastructure would it require?
2. Specify the design choice for the prediction surface in your domain: a flagged-list-with-intervention-prompt vs. a probability-score-the-operator-interprets-unaided. Which design is in place currently, and what would change if the design switched?
3. The case's hedge on teacher selection effects is the open methodological question. What randomized or quasi-experimental design would close the residual attribution gap in your domain's analog deployment — and what institutional support would the design require?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction

D Designed Out **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Gándara, Anahideh, Ison, and Picchiarini (AERA Open, 2024) audited predictive models of college student success and showed that models which look acceptable on overall accuracy are systematically less accurate for Black and Hispanic students and overestimate success for white and Asian students — small-tier frontier evidence that the choice of construct and the stratification used in evaluation, not only model-bias mitigation, determine whether an equity-oriented prediction is fair to the groups the equity commitment is meant to protect

IN BRIEF

Predictive models of college student success — models that score students on predicted graduation, retention, or course completion to drive advising, outreach, and support decisions — have become a routine tool across community colleges and four-year institutions. Gándara, Anahideh, Ison, and Picchiarini, publishing in AERA Open (2024), audited such models across racialized groups and found that a model which looks acceptable on overall accuracy is systematically less accurate for Black and Hispanic students — making more prediction errors for them — while overestimating success for white and Asian students. The apparent fairness of the system depends materially on upstream choices: the construct the model is built to predict (predicted graduation vs. predicted benefit vs. predicted need), the stratification used in evaluation (overall accuracy vs. accuracy by income, race/ethnicity, first-generation status), and the decision context the prediction is consumed in (whether end users are trained to contextualize a flagged prediction or treat it as a verdict). The paper's contribution is the frontier-shaped finding that fairness in equity-oriented prediction is a construct-definition and stratified-evaluation problem before it is a model-bias problem. The case pairs explicitly with the v2 race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman) and with the broader equity-construct competency C8.2: demographic stratification of validation and outcomes as a design commitment.

THE SHIFT

Predictive modeling of student success in higher education is now routine. Community-college and four-year institutions operate models that score students on predicted graduation

likelihood, predicted retention, predicted benefit-from-intervention, or predicted financial need. The scores feed downstream decisions: which students get outreach emails, which get advising appointments, which get need-based aid or enrollment incentives. The structural capability question — whose interest does the model maximize, by what construct, and against what stratification — is the layer where the case operates.⁰

WHAT IS EMERGING

Gándara, Anahideh, Ison, and Picchiarini, publishing in AERA Open, audited predictive models of college student success across racialized groups. The central finding is that models which look acceptable on overall accuracy are systematically less accurate for Black and Hispanic students — making more prediction errors for them — while overestimating success for white and Asian students. The methodological move is to show that the choice of construct the model is built to predict, and the way the model is evaluated, have substantial consequences for which students the model treats accurately. A model optimized and validated on aggregate accuracy is not interchangeable with one whose validation is stratified by racialized group; the institution's choice of construct and evaluation regime is a normative decision the model architecture cannot make on its own.¹

THE CAPABILITY QUESTION

The second methodological move concerns evaluation stratification. Overall accuracy, the standard summary metric, can mask substantial disparity by income, race/ ethnicity, and first-generation status — populations whose base rates and observed outcomes differ from the modal student the training data over-represents. The paper's stratified evaluation shows the canonical equity-construct finding: a model that looks fair under overall accuracy can be substantially less accurate, or substantially biased in its prediction direction, for the subgroups the equity commitment is supposed to protect. The induced framework's C8.2 sub-competency — demographic stratification of validation and outcomes as a design commitment — is the analytic anchor.²

EARLY EVIDENCE

The third move concerns the decision context the prediction is consumed in and the mitigations available. A model is not deployed in isolation; its predictions are read by advisors and administrators who decide what to do with a flagged student. Gándara and colleagues stress training end users on the potential for algorithmic bias — so a prediction is contextualized for the individual student rather than treated as a verdict — and evaluate bias-mitigation techniques that reduce, but do not eliminate, the cross-group accuracy gap. The system's apparent fairness therefore depends not only on the model but on the evaluation regime and the decision practice the institution actually deploys.³

OPEN PROBLEMS

What the case teaches at the LENS framing layer is the frontier-shaped finding: fairness in equity-oriented prediction is a construct-definition problem before it is a model-bias problem. The induced framework's C8.2 sub-competency and the equity-construct CLOs proposed find their case-grounded basis here. The case explicitly cross-references the v2 race-construct trio — eGFR (Case 25), pulse oximetry (Case 26), and Hoffman pain bias (Case 6) — at the construct-definition layer: in those cases the construct (race correction in eGFR, single-sensor calibration in pulse oximetry, the pain-perception assumption in

Hoffman) was the design decision that produced the disparate outcome; in this case the construct (predicted enrollment vs. predicted need vs. predicted benefit) is the design decision the algorithmic– targeting system has to make explicitly.

Fairness in equity-oriented prediction is a construct-definition problem before it is a model-bias problem. Which student outcome the model is maximizing is the upstream decision.

EDITORS' SYNTHESIS OF GÁNDARA, ANAHIDEH, ISON, & PICCHIARINI, _AERA OPEN_ (2024).

REFERENCES

1. Gándara, Anahideh, Ison, & Picchiarini (2024), "Inside the Black Box: Detecting and Mitigating Algorithmic Bias across Racialized Groups in College Student-Success Prediction," AERA Open 10, doi:10.1177/23328584241258741 — primary case source on cross-group accuracy disparity, stratified evaluation, and bias mitigation in student-success prediction.

2. Barocas, Hardt, & Narayanan, Fairness and Machine Learning (fairmlbook.org) — methodological backdrop on construct definition and stratified evaluation.

3. Friedler, Scheidegger, & Venkatasubramanian (2021), "The (im)possibility of fairness: different value systems require different mechanisms for fair decision making," Communications of the ACM — the construct-definition argument at field level.

4. v2 cross-referenced cases: 25 (eGFR race correction), 26 (pulse oximetry across skin tones), 6 (Hoffman pain bias) — the race-construct trio at the construct-definition layer.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Gándara’s student-success-prediction audit is the small-tier frontier instance of fairness-as-construct-definition. Which construct the model maximizes (predicted graduation vs. benefit vs. need), which stratification is used in validation, and the decision context the prediction is consumed in each determine the fairness properties of the deployed system; the audit found models systematically less accurate for Black and Hispanic students. Cross-references the v2 race-construct trio at the construct-definition layer.

LENS APPROACH

Gándara student-success-prediction fairness is the equity-construct-definition case (induced 8.2; LENS D4/PT5) — Domain 4 for stratified evaluation; Domain 5 for the decision-context and mitigation question. Cross-reference Cases 25, 26, 6 — case-grounded basis for the equity- construct CLOs.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER EDUCATION	D K N	D1 D2 D3 D4 D5	5	8.2	4 · 5
ALGO FAIRNESS		Test and Evaluation			
POLICY ANALYSIS					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Make the construct choice explicit at design time: predicted enrollment, predicted benefit, predicted need are different constructs with different policy implications, and the institution’s normative position on which to optimize for has to be on the record.
- ▶ Build stratified evaluation by income, race/ethnicity, first-generation status, and other equity-relevant axes into the model’s validation; the C8.2 sub-competency operationalized.
- ▶ Pair the targeting model with the specific intervention it will feed; the fairness properties of the system depend on the pairing, and evaluating the model in

IN OPERATION / AFTER

- ▶ Report which construct the model maximizes, on which stratifications it was evaluated, and which intervention it is paired with — three pieces of information any deployed targeting system should publish together.
- ▶ Treat the fairness conversation as a construct-definition conversation first, and as a model-bias conversation second; the latter is operationally important, the former determines what the model is fair **about**.
- ▶ Carry the race-construct trio (Cases 25, 26, 6) as the cross-domain pair in the curriculum: in each, the

isolation underestimates the variation a deployed system actually exhibits.

construct definition is the upstream design decision; the framework's C8.2 sub-competency and the equity-construct CLOs are anchored here.

REFLECTION QUESTIONS

1. Identify a predictive-targeting model in your domain. Which construct does it maximize — observed outcome, predicted benefit, predicted need — and is the choice on the record as a normative institutional decision, or absorbed into the model's training objective without disclosure?
2. Specify the stratifications you would build into the model's validation: which equity-relevant axes (income, race/ethnicity, first-generation, geography, disability) and what comparison structure (per-axis accuracy, per-axis calibration, per-axis intervention effectiveness)?
3. The case cross-references the v2 race-construct trio (Cases 25, 26, 6). In each, the construct definition is the upstream design decision that produced the disparate outcome. Identify a construct decision in your domain that is currently absorbed into the design rather than on the record — and what would it take to make the choice explicit?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models

D Designed Out **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Yu, Lee, and Kizilcec, publishing in the LAK/EDM literature, examined whether and how protected attributes (race/ethnicity, gender, socioeconomic status) should be included in learning-analytics predictive models, and showed that whether including or excluding the attribute produces fairer outcomes depends on the construct, the model class, the downstream intervention, and the population — small-tier frontier evidence that the include-or-exclude question is the wrong framing

IN BRIEF

A long-running practical question in learning analytics is whether protected attributes — race/ethnicity, gender, socioeconomic status, first-generation status — should be included as features in predictive models. The intuitive “fairness through unawareness” answer is to exclude them. The technical-fairness literature has shown the unawareness answer is incomplete: omitted protected attributes are typically reconstructable from correlated features (zip code, course history, prior achievement), so excluding the attribute does not exclude its predictive footprint and can make discrimination harder to detect and audit. Yu, Lee, and Kizilcec, publishing in the LAK/EDM literature, examined the include-or-exclude question empirically across multiple learning-analytics prediction tasks. The headline finding is the frontier-shaped one: whether inclusion or exclusion produces fairer outcomes depends on the construct being predicted, the model class, the intervention the prediction feeds, and the population. The case is the small-tier frontier instance of “surfacing bias through governance, not just technique” (C8.4 in the induced framework). It cross-references the v2 race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman) at the model-fairness layer: in those cases the construct was the design decision; in this case the question is whether the attribute is allowed into the model that operationalizes the construct.

THE SHIFT

A practical question that recurs across learning-analytics deployments is whether protected attributes — race/ethnicity, gender, socioeconomic status, first-generation status, and similar — should be included as features in predictive models. The intuitive policy answer (“fairness through unawareness”) is to exclude them: a model that does not see race cannot discriminate on race. The intuition is operationally appealing and statutorily aligned in some domains where regulators have read the omission requirement strictly.^o

WHAT IS EMERGING

The technical-fairness literature has shown the unawareness answer is incomplete. Protected attributes that are omitted from a model are typically reconstructable from correlated features the model does see — zip code, school assignment, prior achievement, course history, language at home, family income proxies. Omitting the attribute removes the label, not the predictive footprint, and can in some cases make discrimination harder to detect and audit precisely because the auditable record no longer carries the attribute that would let an evaluator stratify the model's output.¹

THE CAPABILITY QUESTION

Yu, Lee, and Kizilcec, publishing in the LAK/EDM (Learning Analytics and Knowledge / Educational Data Mining) literature, examined the include-or-exclude question empirically across multiple learning-analytics prediction tasks — course completion, performance prediction, dropout risk — with multiple model classes (regression, tree-based, neural) and multiple downstream intervention contexts. The headline finding is the frontier-shaped one: whether including or excluding a protected attribute produces fairer outcomes depends on the construct being predicted, the model class, the intervention the prediction feeds, and the population. There is no general answer; the include-or-exclude question is the wrong framing.²

EARLY EVIDENCE

The right framing the paper develops is governance and audit. The decision to include or exclude a protected attribute is one of several decisions a learning-analytics deployment makes that determine its fairness properties; the decision has to be made deliberately, recorded with reasoning, and paired with stratified evaluation and audit cadence that catches the consequences of the decision in operation. The case is the induced framework's C8.4 instance at small scale: "surfacing bias through governance, not just technique." The technical-fairness machinery alone does not answer the question; the governance architecture around the model is the carrier of the answer in any specific deployment.³

OPEN PROBLEMS

In pair with Case 86 (Gándara on community-college targeting) and with the v2 race-construct trio (Cases 25 eGFR, 26 pulse oximetry, 6 Hoffman), the case completes the v2 equity-construct frontier picture. Case 86 names the construct-definition layer; this case names the protected-attribute-in-the-model layer; the race-construct trio names the construct-encoded-in-the-instrument layer (race correction in eGFR, sensor calibration in pulse oximetry, pain perception in Hoffman). Together, the five cases stage the equity-construct competency across the construct-definition / attribute-handling / instrument-construct axes — the case-grounded basis for the equity-construct sub-competencies proposed in the v2 research backbone.⁴

The include-or-exclude question is the wrong framing. The right framing is governance: a deliberate decision, recorded with reasoning, paired with stratified evaluation and audit cadence.

EDITORS' SYNTHESIS OF YU, LEE, & KIZILCEC (LAK/EDM).

REFERENCES

1. Yu, R., Lee, H., & Kizilcec, R. F. (2021), "Should College Dropout Prediction Models Include Protected Attributes?" in Proceedings of the Eighth ACM Conference on Learning @ Scale (L@S '21), doi:10.1145/3430895.3460139 — primary paper on the include-or-exclude empirical analysis.
2. Kizilcec & Lee, "Algorithmic Fairness in Education," in Holmes & Porayska-Pomsta (eds.), Ethics in Artificial Intelligence in Education — broader synthesis of the fairness-in-learning-analytics frontier.
3. Dwork, Hardt, Pitassi, Reingold, & Zemel (2012), "Fairness through awareness," Proceedings of ITCS — foundational paper on the inadequacy of fairness-through-unawareness.
4. Barocas, Hardt, & Narayanan, Fairness and Machine Learning (fairmlbook.org) — technical-fairness backdrop.
5. v2 cross-referenced cases: 86 (Gándara), 25 (eGFR), 26 (pulse oximetry), 6 (Hoffman) — equity-construct five-case set.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Yu, Lee, and Kizilcec's protected-attributes work is the frontier instance of surfacing bias through governance, not just technique. Whether including or excluding a protected attribute produces fairer outcomes depends on the construct, the model class, the intervention, and the population. The governance architecture around the model is the carrier of the answer in any specific deployment.

LENS APPROACH

Yu/Lee/Kizilcec protected attributes is the model-fairness- governance case (induced 8.4; LENS D4/PT5) — Domain 4 for stratified-evaluation-by-attribute; Domain 5 for the decision-plus-reasoning-plus-audit architecture. Cross- reference Case 86 and the race-construct trio (25, 26, 6).

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
LEARNING ANALYTICS	D K N	D1 D2 D3 D4 D5	5	8.4	4 · 5
ALGO FAIRNESS		Test and Evaluation			
ED-TECH RESEARCH					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the include-or-exclude decision for each protected attribute as a deliberate design choice, recorded with reasoning, rather than absorbed into the data-engineering pipeline.
- ▶ Build stratified evaluation by the protected attribute regardless of whether the model itself uses it as a feature; auditing the model's output by attribute is independent of whether the attribute is an input.
- ▶ Pair the decision with the downstream intervention context; the include-or-exclude answer that produces fairer outcomes in one intervention context may produce less fair outcomes in another.

IN OPERATION / AFTER

- ▶ Publish the protected-attribute handling decisions for any deployed learning-analytics model — included, excluded, and with what reasoning — as part of the model's governance documentation.
- ▶ Operate stratified audit on a regular cadence; the include-or-exclude consequences in operation are what the audit catches, and audit absence makes the decision functionally invisible to the institution.
- ▶ Carry the five-case equity-construct set into the curriculum: Cases 25 (eGFR), 26 (pulse oximetry), 6 (Hoffman) on the instrument-construct layer; Case 86 (Gándara) on the construct-definition layer; this case on the protected-attribute-in-the-model layer. The set is the case-grounded basis for the equity-construct sub-competencies.

REFLECTION QUESTIONS

1. Identify a learning-analytics or analogous predictive model in your domain. Which protected attributes are inputs to the model, which are not, and is the include-or-exclude decision on the record with reasoning, or absorbed into the data-engineering pipeline?
2. Specify the stratified-audit cadence you would operate for the model regardless of whether protected attributes are inputs; the audit catches the consequences of the include-or-exclude decision in operation, and its absence makes the decision functionally invisible.

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models

T Training Gap **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Repeated probing of LiveHint AI (an LLM-based tutor under development at Carnegie Learning) with identity-marked student queries surfaced response differences in tone and level of detail across identities; choice of foundation model materially affected the level of differentiation

IN BRIEF

The AIED 2025 paper “Evaluating an AI Tutor for Bias Across Different Foundation Models” tests LiveHint AI, an LLM-based tutor under development at Carnegie Learning, against a structured probing protocol. Realistic student queries are modified to include explicit or implicit statements of identity — nationality, dialect markers, demographic cues — and the tutor’s responses are assessed for tone and level of detail. The choice of foundation model materially affects the level of differentiation in responses. The authors are explicit that this is not a deployment-bias audit (LiveHint is in development); it is a methods- development paper proposing how foundation-model-level fairness evaluation should be done before deployment. The case extends the race-construct trio (Cases 25, 26, 6, 48) into the LLM-tutoring layer where the structurally new variable is the foundation model. Open questions: whether lab-style probing matches deployed-conversation patterns; whether vendor selection across foundation models becomes a routine fairness deliverable.

THE SHIFT

The deployment surface for LLM-based tutoring expanded faster than the fairness-evaluation literature for it. Most pre-deployment fairness work on educational AI was built for classifiers — proctoring systems, recommendation engines, automated essay scoring — where the output is a score or a flag and the fairness question can be formulated as a disparate-impact comparison across groups. LLM tutoring shifts the output to a conversational response, and the fairness question shifts with it: it is no longer whether two equivalently-prepared students receive equivalent flags, but whether they receive equivalent pedagogical engagement.⁹

WHAT IS EMERGING

The AIED 2025 paper builds the evaluation method against that shifted target. LiveHint AI, an LLM-based tutor under development at Carnegie Learning, is the subject system. The probing protocol generates realistic student queries — calibrated to the kinds of mathematics questions a tutor would receive — and produces matched variants of each

query that differ only in identity-marker content: explicit statements of nationality, dialect markers, demographic cues. The tutor's responses to the matched pairs are then assessed on two dimensions, using lexical metrics: tone (is the response respectful and appropriately framed for a tutoring context) and level of detail (does the response give the same quality of pedagogical content). Whether the resulting differentiation reflects sound pedagogy or undesirable behavior is the interpretive question the paper leaves open.¹

THE CAPABILITY QUESTION

The headline finding is that the choice of foundation model materially affects the level of differentiation in responses. Different foundation models — the underlying LLMs the tutoring layer wraps — produce different patterns of identity-conditioned response variation. The variation is not uniform across foundation models, not uniform across the three response dimensions, and not uniform across identity-marker types. The vendor-selection decision — which foundation model the tutoring product is built on — is itself a fairness-relevant design choice, and the case names it as such. The structurally new variable in the case, by comparison with the race-construct trio in the corpus, is the foundation-model layer, which did not exist as a deployment-side decision when the proctoring-bias and pulse-oximetry cases were studied.²

EARLY EVIDENCE

The authors' framing is explicit and binding on the case. LiveHint AI is in development, not deployment; this is a methods-development paper proposing how foundation-model-level fairness evaluation should be done before deployment. The case is not the deployment-bias-audit case the corpus carries at Cases 6 (Hoffman pain assessment), 26 (pulse oximetry), 25 (eGFR), and 48 (Johnson school surveillance). It is the structurally new methods-development case at the layer above those deployments — the foundation-model layer — and it grounds the curriculum's demographic-stratification anchor at that layer.³

OPEN PROBLEMS

The open questions the authors preserve are the case's load-bearing hedges. The authors also report that they found no evidence of harmful biases — the differentiation they measured may reflect sound pedagogy or undesirable behavior, an interpretive question they leave open. Whether the differentiation patterns documented in lab-style probing match what students encounter in deployed conversations — where session length, follow-up turns, and student adaptation affect the response trajectory — is a question the present study cannot answer. Whether vendor selection across foundation models becomes a routine fairness deliverable for the educational-AI procurement pipeline is a market-evolution question the present study can name but cannot resolve. The case pairs with the race-construct trio for the stratified-validation discipline and with Case 77 (hybrid human-AI tutoring) for the deployment-side complement — the augmentation pattern Case 77 documents depends on the foundation-model-level evaluation Case 88 is methodologically grounding.

The choice of foundation model is itself a fairness-relevant design decision. The evaluation should be a pre-deployment deliverable, not a post-deployment audit.

EDITORS' SYNTHESIS OF THE AIED 2025 LIVEHINT AI BIAS EVALUATION.

REFERENCES

1. AIED 2025, “Evaluating an AI Tutor for Bias Across Different Foundation Models,” Springer/ACM proceedings, doi:10.1007/978-3-031-98465-5_43; preprint at renzheyu.com/papers/AIED2025_Tutor.pdf.

2. Bommasani, R. et al. (2021), “On the Opportunities and Risks of Foundation Models,” Stanford CRFM — the foundation-model framing the case builds on.

3. Race-construct trio reference set: Hoffman et al. (2016), Sjoding et al. (2020) pulse oximetry, Inker et al. (2021) eGFR-without-race — paired with Cases 25, 26, 6.

4. Carnegie Learning LiveHint product documentation — the subject system; case framing is binding on LiveHint being in development, not deployment.

THE LEARNING ENGINEERING LENS

LE INSIGHT

LiveHint AI is the methods-development case at the foundation- model layer of LLM-tutoring fairness evaluation. The matched- pair identity-probing protocol surfaces tone and level-of-detail differences across identities; the foundation-model choice materially affects the differentiation level. The case extends the race-construct trio into the LLM- tutoring layer with a structurally new variable — the foundation model — that did not exist at the deployment-audit layer.

LENS APPROACH

LiveHint AI is the demographic-stratified validation case at the foundation-model layer (induced 8.2; LENS D4/PT6). LENS uses it in Domain 4 (Test and Evaluation) for the matched- pair probing methodology and in Domain 3 (Human-System Collaboration) for the foundation-model-selection-as- fairness-decision frame. Pair with the race-construct trio (Cases 25, 26, 6, 48) at the deployment-audit layer and with Case 77 (hybrid human-AI tutoring) as the augmentation-pattern complement.



APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Build the probing protocol to vary identity markers within matched query pairs; the methods-development contribution depends on the matched-pair design that isolates the identity-conditioned response variation.
- ▶ Treat the choice of foundation model as a fairness-relevant design decision, not as an upstream procurement decision; the case demonstrates that the foundation-model layer materially affects the tutoring-layer's identity-conditioned response patterns.
- ▶ Conduct the fairness evaluation before deployment, not after; the methods-development framing is that the evaluation should be a pre-deployment deliverable, parallel to the race-construct-trio cases at the deployment-audit layer.

IN OPERATION / AFTER

- ▶ Commission the deployed-conversation evaluation that the lab-style probing cannot perform; the open question on whether probing patterns match deployment patterns is testable against deployment logs as the system moves toward release.
- ▶ Publish the foundation-model-level fairness findings as part of the procurement record; the case argues that vendor selection across foundation models is a fairness deliverable, and the publication discipline is what would make that deliverable operational.
- ▶ Pair the case in the curriculum with the race-construct trio (Cases 25, 26, 6, 48) so the demographic-stratification anchor is taught across both the deployed-system layer and the foundation-model layer above it.

REFLECTION QUESTIONS

1. Identify an LLM-based system in your domain whose foundation-model choice was treated as an upstream procurement decision rather than as a fairness-relevant design decision. What would the matched-pair probing protocol look like for that system, and what response dimensions would you assess?
2. Specify the pre-deployment fairness deliverable you would build into the procurement record for an LLM-based product in your domain. What would the published artifact contain — probe set, response-dimension assessment, foundation-model comparison — and what would the procurement decision turn on?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

Data Privacy and Learning Analytics on the African Continent

K Knowledge & Institutional Memory

G Governance & Trust

N Normalization of Deviance

IMPACT Mapped the legal and regulatory privacy landscape across African jurisdictions and surfaced the structural governance seam — African higher education frequently uses learning platforms and analytics hosted by external providers, creating cross-regime gaps where student data crosses jurisdictions with inconsistent protection — and recommended common cross-border data-sharing frameworks

IN BRIEF

Prinsloo and Kaliisa (British Journal of Educational Technology, 2022) mapped the legal and regulatory privacy landscape across African jurisdictions and found a growing trend toward comprehensive data-protection legislation, though few frameworks are yet enacted and cross-border data-transfer policies differ sharply between countries. The core governance seam the paper surfaces — and what makes the case a frontier rather than a settled instance — is what the v2 sweep names **extraterritorial platform governance**: African higher education frequently uses learning platforms, learning-management systems, and analytics services hosted by external (typically Global-North or Asian) providers, creating a structural gap where student data is generated under one regulatory regime, processed under another, and the institution's governance authority does not reach the operating regime. The paper recommends common cross-border data-sharing frameworks as the architectural response. The case is included as a frontier — it documents the governance architecture needed for responsible adoption rather than a completed success — and it is valuable because it surfaces a seam (extraterritorial hosting) the US-centric canon rarely confronts, and that is increasingly universal as institutions everywhere build on cloud and AI services they do not control.

THE SHIFT

The standard framing of learning-analytics governance assumes that the operating institution, the regulatory regime, and the hosting infrastructure are co-located — a US university under US law on US-hosted infrastructure, a UK university under UK law on UK or EU infrastructure. African higher education sits structurally outside this assumption. Most universities across the continent rely on learning-management systems, predictive analytics tools, and adjacent platforms hosted by providers based in the Global North or in Asia, and the regulatory regimes governing those providers are not the same as the regimes governing the institutions or the students.⁹

WHAT IS EMERGING

Prinsloo and Kaliisa's 2022 paper in the *British Journal of Educational Technology* maps the legal and regulatory privacy landscape across African jurisdictions. The picture is one of a continent in motion: a growing trend toward comprehensive data-protection legislation following the model the EU established with GDPR, but with few frameworks yet enacted and substantial variation across countries in scope, enforcement capacity, and cross-border data-transfer policy. Some countries have comprehensive frameworks; some have sectoral protections; some have constitutional privacy provisions without operational data-protection legislation; and the patchwork makes the cross-border picture itself uneven.¹

THE CAPABILITY QUESTION

The structural seam the paper surfaces is what makes the case teachable. When student data is generated under one jurisdiction's regulatory regime, processed under another's because the platform is hosted abroad, and the institution's governance authority does not reach the operating regime, the governance question becomes architectural rather than institutional. The authors recommend common cross-border data-sharing frameworks as the architectural response — the kind of inter-regime instrument that GDPR's adequacy decisions and successor cross-border-transfer mechanisms began to construct between the EU and other regimes, but that no comparable continent-wide African framework yet provides. The case names the gap and the architectural response; what it does not yet document is the build-out.²

EARLY EVIDENCE

The honest evidentiary state is what the frontier tag carries. The paper is peer-reviewed and the mapping is grounded in legal documentation across jurisdictions; what the case does not establish — because the architecture is not yet built — is that any specific cross-border framework has been adopted and that the seam has been closed in practice. The case is teachable on the structural pattern: the governance question is the architectural one across regimes, not the policy one inside any single regime. The outcome is open in the sense that no continent-scale resolution yet exists. Future validation will continue as African data-protection law consolidates.³

OPEN PROBLEMS

The frontier note the case carries names what the v2 sweep calls extraterritorial platform governance — capability deployed on platforms governed by a different regime than the one operating them. The pattern is not African-specific. It is increasingly universal as institutions everywhere build on cloud, AI, and platform services they do not control: a US university on a US cloud is the historical case; a US university running analytics through a model-hosted-elsewhere service, a Latin American university on a US cloud, an African university on a European platform — all share the structural seam, and the African case surfaces it most starkly because the regime asymmetry is largest there. The case is the case-grounded basis for proposing extraterritorial-platform-governance as a sub-competency the existing curriculum does not yet name.

The governance question is architectural across regimes, not policy inside any single regime. The seam is at the joints, and the joints have to be drawn before the architecture can be designed.

EDITORS' SYNTHESIS OF PRINSLOO & KALIISA (2022).

REFERENCES

1. Prinsloo, P., & Kaliisa, R. (2022), "Data privacy on the African continent: Opportunities, challenges and implications for learning analytics," *British Journal of Educational Technology* 53(4):894–913, doi:10.1111/bjet.13226.
2. African Union Convention on Cyber Security and Personal Data Protection (Malabo Convention, 2014) — continental policy backdrop, partial ratification.
3. Slade & Prinsloo (2013), *American Behavioral Scientist* — earlier framing on which the 2022 paper builds.
4. Open University Ethical Use of Student Data policy (2014) — single-regime governance-by-design companion.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The African data-privacy case names a governance seam the existing curriculum underweights: extraterritorial-platform governance, where the operating, regulating, and hosting regimes diverge. The case is a frontier — peer-reviewed mapping, architectural recommendation, no continent-scale resolution yet built. The pattern is increasingly universal, and the case-grounded basis is the case for naming a sub-competency the framework does not yet have.

LENS APPROACH

African data privacy is the extraterritorial-platform governance case (induced 5.3; LENS D5/PT6) — Domain 5 for the cross-regime seam; Domain 3 for **Judgment under inadequate evidence** on a frontier where the architectural response is recommended but not built. Pair with Case 81 and Case 91.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER EDUCATION (AFRICA)	K G N	D1 D2 D3 D4 D5	6	5.3	5 · 3
DATA PRIVACY LAW	→	→			
LEARNING ANALYTICS		Navigating Sociotechnical Constraints			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Map the regime structure explicitly: which jurisdiction generates the data, which processes it on the platform, which the institution operates under, which the student lives under. The seam is at the joints, and the joints have to be drawn before the architecture can be designed.
- ▶ Choose the platform stack with the regime structure in view; a procurement decision is also a governance decision when the platform is in a different jurisdiction from the operating institution.
- ▶ Treat the architectural response as common cross-border frameworks rather than institution-by-institution contract negotiation; the inter-regime instrument is what scales.

IN OPERATION / AFTER

- ▶ Surface the extraterritorial-platform-governance pattern in any v2 case that involves cloud/AI infrastructure; the African case is the starkest instance, not the only one.
- ▶ Carry the outcome-open honesty: the architecture is not yet built, and the case names the gap rather than a completed success.
- ▶ Build the case-grounded basis for proposing extraterritorial-platform-governance as a sub-competency the existing curriculum does not yet name; pair with the OU and LALA cases at the governance-by-design layer.

REFLECTION QUESTIONS

1. Identify a platform or service your institution depends on whose operating regime is different from your own. Which jurisdiction generates the data, which processes it, which regulates the institution, which regulates the user — and where is the seam between regimes?
2. Specify the architectural response your context needs: an institution-by-institution contractual approach, a sectoral framework, a cross-border data-sharing instrument. Which level resolves the seam, and which level cannot?
3. The case names extraterritorial platform governance as an under-named sub-competency. What instance from your domain would anchor it — a cloud-hosted clinical analytics tool, a model-as-a-service inference layer, a learning-platform vendor in a different jurisdiction — and how is the curriculum response different from a single-regime governance case?

WHO BUILDS THIS knowledge management & data engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

CASE 90

HIGHER EDUCATION

ALGORITHMIC DECISION-MAKING

HUMAN-COMPUTER INTERACTION

2025

Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions

T Training Gap **K** Knowledge & Institutional Memory **N** Normalization of Deviance

DISCLOSURE Institutional overlap: John Katzman collaborates with an editor's institution (Johns Hopkins) but was not involved in this work. Pyle and Andalibi's study is independent of Katzman; the case sits on the applicant side of enrollment management.

IMPACT Eighteen semi-structured interviews with recent U.S. university applicants, using speculative-design probes, surfaced systematic distances between vendor marketing claims (efficiency, fairness, enhanced fit) and applicants' own perceptions (opacity, distrust, anticipated discrimination)

IN BRIEF

Pyle and Andalibi (CSCW 2025) report an interview study with 18 recent U.S. university applicants, using speculative-design probes to surface how applicants perceive algorithmic systems operating in college admissions and enrollment. The study's central contribution is naming the systematic distance between vendor marketing claims — efficiency, fairness, enhanced fit — and applicants' own perceptions of how the systems would treat them: opacity about the algorithm's existence, distrust of its objectives, and anticipated discrimination across protected characteristics. The case is the peer-reviewed consent-side companion to Case 55 (Engler / enrollment algorithms) and Case 57 (GAO OPM oversight gap). The authors' own hedge is explicit: 18 interviews is the right sample for the speculative-design method but not for prevalence claims, and "future validation ongoing" applies to whether the perception patterns generalize across applicant demographics and institution types.

THE SHIFT

The deployment surface for algorithmic systems in college admissions and enrollment has been studied largely from the institutional side — what vendors promise, what colleges buy, what optimization outcomes the institution reports. The applicants whose lives the deployment most directly affects have been the structurally absent voice in the deployment record. Pyle and Andalibi's CSCW 2025 paper is the systematic peer-reviewed correction at the consent-side seam: an interview study with 18 recent U.S. university applicants, using speculative-design probes calibrated to surface applicants' perceptions of algorithmic systems operating on their applications.^o

WHAT IS EMERGING

The speculative–design method is the right instrument for the question. Applicants in most cases do not know which specific algorithmic systems acted on their applications; the institutional side does not disclose the vendor stack or the optimization objective at the point of application. Speculative–design probes — scenario sketches that concretize plausible algorithmic interventions and ask applicants to react — let the study elicit perceptions that are not contingent on the applicant having direct knowledge of the deployed system. The 18-interview sample is calibrated to the method’s depth rather than to prevalence–claim breadth, and the authors are explicit about that calibration as a methodological choice.¹

THE CAPABILITY QUESTION

The findings name systematic distance between vendor marketing and applicant perception across three axes. Vendor marketing pitches efficiency — the system processes more applications faster — while applicants name opacity: they do not know which decisions are being made algorithmically and they cannot interrogate the system’s logic. Vendor marketing pitches fairness — the system treats applicants consistently — while applicants name distrust of the objective the system is consistent with. Vendor marketing pitches enhanced fit — the system matches applicants to programs likely to serve them well — while applicants name anticipated discrimination across protected characteristics, drawing on the broader public record of algorithmic disparate-impact findings.²

EARLY EVIDENCE

The case sits as the consent-side counterpart to Case 55 (Engler / Brookings, the deployment-side mapping) and Case 57 (GAO–22–104463, the regulator-side audit). Engler documents the deployed algorithmic optimization; the GAO audit documents the regulatory oversight gap; Pyle and Andalibi document what the affected applicants understand about the deployment. The three cases together name the structural seam: a deployed system about which the institutional side, the regulator, and the affected population each hold partial and non-overlapping information. The case also pairs with Cases 186 (Bartlett lending fairness) and 86 (Gándara community-college equity) as the applicant–perception strand of the equity-in–deployed–prediction thread.³

OPEN PROBLEMS

The authors’ hedges are binding on the case’s framing. Eighteen semi-structured interviews is the right N for the speculative–design method, surfacing the dimensions of perception that matter, but it is not the right N for claims about prevalence — what fraction of applicants hold each perception, how the perceptions distribute across demographic groups, how they correlate with admissions outcomes. The case is the strongest peer-reviewed evidence currently available on the consent-side question; future validation ongoing on whether the perception patterns generalize across applicant demographics, institution types, and the rapidly evolving algorithmic–deployment landscape.

Vendors pitch efficiency, fairness, and fit. Applicants name opacity, distrust, and anticipated discrimination. The distance between the two is the case.

EDITORS’ SYNTHESIS OF PYLE & ANDALIBI (2025).

REFERENCES

1. Pyle, C., & Andalibi, N. (2025), "Algorithmic College Admissions in the U.S.: Distances Between Vendors' Claims and Applicants' Perceptions," Proceedings of the ACM on Human-Computer Interaction 9(7), CSCW369, doi:10.1145/3757550.

2. Engler (2021), Brookings — paired deployment-side mapping (Case 55).

3. GAO-22-104463 (2022) — paired regulator-side audit (Case 57).

4. Dunne, A., & Raby, F. (2013), Speculative Everything — methodological backdrop for speculative-design probes.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Pyle and Andalibi is the peer-reviewed consent-side case at the applicant end of the enrollment-management deployment. Eighteen interviews surface systematic distance between vendor marketing and applicant perception across three axes; the authors' methodological hedge is binding on prevalence claims. The case completes the partial-information triangle with Cases 55 (deployment-side) and 57 (regulator-side).

LENS APPROACH

Pyle & Andalibi is the governance-rather-than-technique case at the consent boundary (induced 8.4; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the affected-voice inclusion lesson and in Domain 3 (Machine Teaming and Adaptation) for the disclosure-architecture design at the applicant interface. Pair with Cases 55 (Engler deployment), 57 (GAO oversight), 186 (Bartlett), and 86 (Gándara). coi-light render under the title is binding for cluster placement.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER EDUCATION	T K N	D1 D2 D3 D4 D5	6	8.4	5 · 3
ALGORITHMIC DECISION-MAKING	→	→ Navigating Sociotechnical Constraints			
HUMAN-COMPUTER INTERACTION					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Include the affected-population voice in the deployment-decision record from the start; the systematic distance the study documents is partly the result of decision processes that did not include the affected voice.
- ▶ Use speculative-design probes when the affected population cannot be expected to know what is deployed; the method is the right instrument for the consent-side question and the institutional side can commission it.
- ▶ Treat the three perceived gaps — opacity, distrust, anticipated discrimination — as designable targets, not as misperceptions to correct; the perceptions are responses to the actual disclosure architecture.

IN OPERATION / AFTER

- ▶ Commission the prevalence study the speculative-design study cannot perform; the perception patterns the 18-interview study surfaces are testable at survey scale, and the deployment-side institutions are best positioned to commission the survey.
- ▶ Pair the consent-side study with the regulator-side audit (Case 57) and the deployment-side mapping (Case 55) in the curriculum; the three cases together name the partial-information structure of the deployment.
- ▶ Track the perception findings over time as the algorithmic-deployment landscape evolves; the case's value as a longitudinal baseline depends on the comparison studies that come next.

REFLECTION QUESTIONS

1. Identify a deployed prediction system in your domain whose affected population was not consulted in the deployment–decision record. What would a speculative–design study look like for surfacing the affected voice — and who is positioned to commission it?
2. Specify the three perceived gaps — opacity, distrust, anticipated discrimination — for the system you identified. Which of the three is a disclosure–architecture target, which is an objective–choice target, and which is an audit–evidence target?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

CASE 91

HIGHER EDUCATION (LATIN AMERICA)

LEARNING ANALYTICS

CROSS-NATIONAL GOVERNANCE

2017 – 2020

LALA — Building Learning-Analytics Governance Capacity Across Latin America

K Knowledge & Institutional Memory**N** Normalization of Deviance

IMPACT An EU-funded multi-country project (Chile, Ecuador) that explicitly rejected lifting Global-North learning-analytics tools wholesale; structured interviews with administrators and focus groups with students and teachers produced the LALA CANVAS participatory adoption framework, with stakeholders demanding ethical responsibility as a precondition for data-driven feedback

IN BRIEF

The LALA (Building Capacity to Use Learning Analytics to Improve Higher Education in Latin America) project, funded under EU grant 586120-EPP-1-2017-1-ES, ran from 2017 to 2020 across Chilean and Ecuadorian universities. The project began from a deliberate refusal: not to lift US and European learning-analytics tools wholesale into Latin American institutional contexts, on the grounds that adoption fails when the tools do not integrate with local learning design and institutional decision-making. Through structured interviews with administrators and focus groups with students and teachers, the Hilliger et al. team (Internet and Higher Education, 2020) surfaced what stakeholders actually needed for adoption to be locally legitimate, and built the LALA CANVAS — a participatory adoption framework that puts ethical responsibility on the front-end of the adoption decision, not as a compliance afterthought. The honest limit preserved verbatim: this is adoption-readiness and capacity-building evidence, not yet long-run outcome evidence that the deployed systems improved student retention or learning. The contribution is the participatory governance method that made adoption locally legitimate — and the case is the non-US companion to OU Analyse (Case 81), where governance-by-design unblocked deployment in a different cross-cultural context.

BACKGROUND

Learning analytics as a field is heavily Global-North-shaped. The reference tools (predictive at-risk classifiers, dashboard analytics for instructors, early-warning systems) were developed at US and European universities with US and European governance assumptions about data, consent, and institutional decision-making. The LALA project began from the documented observation that lifting those tools wholesale into Latin American institutional contexts had failed before — adoption stalled, instructors did not use the dashboards, students did not trust the predictions — and that the failure was structural, not technical: the tools had not been integrated with the local learning design or the institutional decision-making process they were meant to inform.^o

THE INTERVENTION

The project ran across Chilean and Ecuadorian universities under an EU Erasmus+ grant from 2017 to 2020. The method's first commitment was participatory: structured interviews with administrators surfacing what their decisions actually needed evidence about; focus groups with teachers and students surfacing what they would accept, what they would resist, and what they wanted the analytics to do. Hilliger et al. (Internet and Higher Education, 2020) is the peer-reviewed mixed-methods report that documents the method and the findings. The headline result is that stakeholders consistently named ethical responsibility as a precondition for data-driven feedback rather than as a compliance burden — they wanted the analytics, conditional on the governance being right first.¹

HOW IT WORKED

The deliverable is the LALA CANVAS — a participatory adoption framework that walks an institution through the decisions that have to be made before a learning-analytics system is deployed: which questions the system is for, which stakeholders' consent is required, what disclosure is owed, what the operating governance will look like once the system runs. The framework's contribution is procedural: it converts the governance question from a yes/no gate at deployment time into a structured set of decisions taken openly during the adoption process. The participatory method made the framework locally legitimate across two regimes whose own data governance is differently mature.²

THE EVIDENCE

The honest limit survives into the case. The published evidence is adoption-readiness and capacity-building evidence: the CANVAS was developed and validated through the participatory process, and the project produced trained local teams with the capacity to lead adoption in their own institutions. What the evidence does **not** yet establish is that the deployed systems improved long-run student retention or learning outcomes. The case is teachable on the governance method — the participatory route to local legitimacy — and the outcome evidence is the next study, not this one. Drafting that softens this hedge overstates the claim.³

WHAT TRANSFERRED

In pair with the Open University case (Case 81, the UK consent-by-design intervention), LALA shows that the governance-by-design pattern is transferable across regimes: OU built consent for a single-institution intervention under pre-GDPR UK scrutiny; LALA built participatory adoption for multi-country capacity-building under two different Latin American regulatory regimes. The pair plus SyRI (Case 189) teaches the non-US LA governance triple — design that unblocked deployment (OU), participatory governance that built adoption capacity (LALA), and rights-grounded halt (SyRI) — three honest results for the same structural question of when delegation to analytic infrastructure is legitimate.

Stakeholders did not want the analytics if the governance was wrong. They wanted both, in that order, with ethical responsibility as the precondition rather than the afterthought.

EDITORS' SYNTHESIS OF HILLIGER ET AL. (2020).

REFERENCES

1. Hilliger, Ortiz-Rojas, Pesántez-Cabrera, Scheihing, Tsai, Muñoz-Merino, Broos, Whitelock-Wainwright, & Pérez-Sanagustín (2020), "Identifying needs for learning analytics adoption in Latin American universities: A mixed-methods approach," *The Internet and Higher Education* 45:100726, doi:10.1016/j.iheduc.2020.100726.
2. LALA project — EU Erasmus+ grant 586120-EPP-1-2017-1-ES (2017–2020) — program documentation and the LALA CANVAS artifact.
3. Open University Ethical Use of Student Data policy (2014) and OU Analyse — UK companion governance-by-design case (Case 81).
4. Slade & Prinsloo (2013), "Learning Analytics: Ethical Issues and Dilemmas," *American Behavioral Scientist* 57(10):1510–1529 — the broader field-scale ethics framing.

THE LEARNING ENGINEERING LENS

LE INSIGHT

LALA converted learning-analytics adoption from a Global-North template-lift into a participatory process that built local legitimacy across two Latin American regimes. The CANVAS is the artifact; the participatory method is the deliverable. The honest limit — adoption-readiness evidence, not yet long-run outcome evidence — is what the case carries, and the outcome study is the next one.

LENS APPROACH

LALA is the non-US participatory-governance case (induced 5.1; LENS D5/PT4) — Domain 5 for the cross-regime participatory method; Domain 3 on **Judgment under inadequate evidence** (adoption-readiness, not closed outcome proof). Pair with Case 81 (OU) and Case 189 (SyRI).

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HIGHER EDUCATION (LATIN AMERICA)	K N	D1 D2 D3 D4 D5	4	5.1	5 · 3
LEARNING ANALYTICS	→	→ Navigating Sociotechnical Constraints			
CROSS-NATIONAL GOVERNANCE					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Begin from the documented failure mode — Global-North tools that did not transfer — so the adoption project is designed against the actual local barrier, not against an imagined universal one.
- ▶ Run the participatory method with administrators, teachers, and students separately; the decisions, the daily use, and the consent each sit with a different stakeholder.
- ▶ Convert the governance question from a yes/no deployment gate into a structured set of decisions surfaced during adoption; the CANVAS is the worked artifact of that conversion.

IN OPERATION / AFTER

- ▶ Report the evidence at its tier: adoption-readiness and capacity-building is what the published study establishes; long-run student-outcome evidence is the next study, and the case names that gap rather than glossing it.
- ▶ Commission longitudinal follow-up at the institutions that adopted via the CANVAS to convert capacity-building evidence into deployed-outcome evidence over time.
- ▶ Carry the cross-regime lesson into the broader v2 governance set: governance-by-design is transferable, but the local participatory process is the half that does not transfer; each regime has to do it for itself.

REFLECTION QUESTIONS

1. Identify a tool or framework in your domain that has been lifted from one regime to another without local adaptation. What participatory method would surface what the destination stakeholders actually need, and what would convert the governance question into a structured set of decisions during adoption rather than a gate at deployment?
2. Specify the evidence tier honestly for an adoption-readiness case in your context: capacity-building and process evidence is what you can publish now; long-run outcome evidence is the next study. Where in your communication is the tier most at risk of being smoothed away?

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

Norway's National Expert Commission on Learning Analytics

G Governance & Trust

K Knowledge & Institutional Memory

N Normalization of Deviance

IMPACT Rather than let learning analytics diffuse unregulated or block it, Norway's Ministry of Education convened a national Expert Commission to investigate the pedagogical, ethical, legal, and privacy issues and establish a regulatory foundation before sector-wide deployment; interim report June 2022, final report 2023 (NOU), with central dilemmas explicitly framed

IN BRIEF

Norway's Ministry of Education convened a national Expert Commission on Learning Analytics in 2021 — a national-government response to a capability deployment question at sector scale. Rather than let learning analytics diffuse unregulated across Norwegian education, or block it on precautionary grounds, the ministry chose to construct the governance architecture first. The commission delivered an interim report to the Minister in June 2022 and a final report in 2023 (the NOU, Norges offentlige utredninger, the canonical form of Norwegian public commission reports) identifying central dilemmas across the pedagogical, ethical, legal, and privacy dimensions. The case is governance-as-deliberate-artifact at national scale — a country treating change-control and disclosure as the precondition for adoption rather than the consequence of it. The honest limit preserved verbatim: the commission's recommendations were guidance to a ministry, and downstream sector outcomes — whether deployed Norwegian LA systems actually improved learning outcomes or preserved trust under operation — are not yet documented. This is process-level evidence (a national governance artifact exists, the dilemmas are named), not yet deployment-outcome evidence. The case pairs with the OU (Case 81) and SyRI (Case 189) cases as the national Nordic complement to the institutional-UK and judicial-Dutch governance modes.

BACKGROUND

Learning analytics at sector scale presents a national government with a structural choice. Let the technology diffuse across schools and universities under existing patchwork governance (Sweden's path through the 2010s), or block it on precautionary grounds until questions are settled (a path some jurisdictions have pursued for specific applications), or construct the governance architecture first and let deployment follow under it. Norway in 2021 chose the third path. The Ministry of Education convened a national Expert Commission on Learning Analytics to investigate the pedagogical, ethical, legal, and privacy issues across the whole education sector and to establish a regulatory foundation for what sector-scale deployment should look like.^o

THE INTERVENTION

The commission's mandate covered the full chain. Pedagogically, what kinds of learning-analytics-driven interventions are defensible at primary, secondary, and tertiary levels; what claims about outcome the evidence supports; what teacher-student relationship the analytics should and should not be allowed to alter. Ethically, what is owed to students whose data drives the analytics, what consent architecture is defensible across age groups, and what disclosure structure the analytics-driven decisions should carry. Legally, how the Norwegian data-protection regime under GDPR interacts with the educational context, and what additional sectoral instruments are needed. Privacy, where the line is between pedagogically useful data and surveillance overreach.¹

HOW IT WORKED

The commission delivered an interim report to the Minister of Education and Research in June 2022 and a final report in 2023 in the canonical NOU form. The reports name central dilemmas the field has to live with rather than resolve once: the tension between predictive support and predictive gatekeeping; the tension between transparency to the student and the technical complexity of the models; the tension between cross-institutional benchmarking and student-data protection; the tension between national pedagogical consistency and institutional autonomy in how analytics are used. The reports' framing is that governance for learning analytics is the kind of artifact that has to be revisited as the technology and the evidence base change, not a one-time document.²

THE EVIDENCE

The honest evidentiary state is process-level. The case establishes that a national government can produce a structured governance artifact, identify the dilemmas sector-scale deployment will face, and deliver the architecture as guidance to a ministry. The case does **not** yet establish that downstream Norwegian sector outcomes — improved learning, preserved trust, defensible interventions — have been delivered, because deployment under the new architecture is too recent for outcome evidence. This is a governance-process success, not yet a measured deployment-outcome success, and the case carries the qualification rather than collapsing it.³

WHAT TRANSFERRED

In pair with the Open University (Case 81, institutional governance-by-design, UK) and SyRI (Case 189, judicial rights-grounded halt, Netherlands), Norway's commission is the national-scale governance-architecture mode. The three cases together teach that learning-analytics governance can be produced at the institutional level (OU), constrained by the courts after the fact (SyRI), or constructed by national deliberation before sector deployment (Norway). The structural lesson is that governance is producible as a deliverable at whichever level matches the deployment scope, and the choice of level is itself a governance decision the case literature names.

The commission did not resolve learning-analytics governance for Norway. It named the dilemmas the sector will live with and produced the artifact deployment can be governed under.

EDITORS' SYNTHESIS OF THE NORWEGIAN EXPERT COMMISSION INTERIM AND FINAL REPORTS.

REFERENCES

1. Norwegian Expert Commission on Learning Analytics, interim report to the Minister of Education and Research (June 2022).

2. Norwegian Expert Commission on Learning Analytics, final NOU report (2023), Norges offentlige utredninger.

3. Misiejuk & Wasson (2023), "Learning analytics in Norway: A national perspective," Journal of Learning Analytics — secondary academic synthesis of the commission and its dilemmas.

4. Hilliger et al. (2020), Internet and Higher Education — the LALA companion at multi-country participatory scale (Case 91).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Norway's national Expert Commission is the governance–architecture–at-national-scale instance: a country constructing the regulatory and pedagogical foundation before sector-scale learning-analytics deployment, rather than after diffusion or via judicial halt. The artifact exists and the dilemmas are named; downstream sector outcomes are not yet documented. Process-level success; deployment–outcome evidence is the next decade's work.

LENS APPROACH

Norway is the national-scale governance–architecture case (induced 5.4; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the change- control-and-disclosure-as-artifact discipline and in Domain 4 (Test and Evaluation) for the CLO **Judgment under inadequate evidence** on the process-vs-outcome tier. Pair with Case 81 (OU, institutional), Case 189 (SyRI, judicial), Case 91 (LALA, multi-country participatory), and Case 89 (African data privacy, frontier) — the non-US LA governance pentad teaching the level-of-governance decision.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
EDUCATION (NORWAY) NATIONAL POLICY LEARNING ANALYTICS	G K N	D1 D2 D3 D4 D5	4	5.4	5 · 3
→ Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Match the governance artifact to the deployment scope: institutional governance for a single university, national commission for sector-scale deployment, judicial review when neither operates. The choice of level is itself a governance decision.
- ▶ Convene the commission with the dimensional breadth the deployment actually has: pedagogical, ethical, legal, privacy — not only the dimension the most-visible objection sits on.
- ▶ Frame the artifact as living: governance for a moving technology has to be revisited as the technology and the evidence base change, and the artifact should say so.

IN OPERATION / AFTER

- ▶ Report the evidence at its tier: governance–process success is what the commission and NOU establish; deployment–outcome success is the next decade's work, and the case is honest about the gap.
- ▶ Carry the central dilemmas into the implementation conversation rather than treating them as resolved by publication of the NOU; the dilemmas are what implementation will live with.
- ▶ Use the case as the national-scale instance in the OU / LALA / SyRI / Norway non-US LA governance set; the joint teaching point is that governance is producible at whichever level matches the deployment scope.

REFLECTION QUESTIONS

1. Identify a sector-scale capability deployment question in your domain. Is the right governance level institutional, sectoral, national, or judicial — and what determines the match between deployment scope and governance level?
2. Specify the dimensional breadth your commission or governance artifact would have to cover. Norway's mandate was pedagogical, ethical, legal, and privacy. What is the analog in your context, and which dimension is most at risk of being narrowed away under stakeholder pressure?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

Learning Analytics on the African Continent — A Scoping Review as the Present-State Map

K Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT A 2022 scoping review found only 15 learning-analytics studies on the entire African continent, overwhelmingly from South Africa (10 of 15 studies), with the remainder from Nigeria, Tanzania, Zambia, and Kenya; the structural finding — limited LMS access, limited institutional resourcing, limited African-scholar visibility at SoLAR — is itself the evidence the field requires before construct-travel claims can be made

IN BRIEF

Prinsloo and colleagues (2022) published a scoping review of learning-analytics research on the African continent for the Journal of Learning Analytics. The review found only 15 studies meeting inclusion criteria, with publication output overwhelmingly from South Africa (10 of 15) and the remainder from Nigeria, Tanzania, Zambia, and Kenya. The structural findings — limited LMS access in many African higher-education institutions, limited institutional resourcing for learning-analytics infrastructure, and limited African-scholar visibility at the Society for Learning Analytics Research (SoLAR) conferences — are the present-state map the field needs before importing US/UK/EU learning-analytics constructs into African contexts. The case is included as a frontier scoping case at the practice-synthesis tier: a review of an early-stage research base where the absence of dense primary studies is itself the finding. It pairs with the African data-privacy governance case earlier in the corpus to articulate the construct-travel problem in both research-base and governance terms. Future validation ongoing as the African learning-analytics literature matures.

THE SHIFT

Learning analytics as a field consolidated around 2011 with the formation of SoLAR and a research agenda anchored heavily in US, Australian, UK, and EU higher-education data infrastructures. The capability question for the African higher-education community is whether the analytic constructs travel — and who is positioned to validate them when they arrive. Prinsloo et al. set out to map what learning-analytics research had actually been published from African institutions at the time of writing.^o

WHAT IS EMERGING

The headline finding is sparse. The scoping review located only 15 learning-analytics studies on the African continent, with publication output overwhelmingly from South Africa (10 of 15), the rest from Nigeria, Tanzania, Zambia, and Kenya. Adjacent South African higher-education studies — Lemmens and Henn (2015) and a 2020 paper on the development of a contextualised learning-analytics framework — extend but do not change the magnitude of the gap. The corpus exists; it is small, geographically concentrated, and dwarfed by the parent field's publication output from US/UK/EU institutions.¹

THE CAPABILITY QUESTION

The structural findings underneath the count are the load-bearing teaching. African higher-education institutions face limited LMS access at scale (many institutions still rely on heterogeneous and intermittent digital infrastructure); limited institutional resourcing for the data-engineering, ethics-review, and analyst capacity that learning analytics presupposes; and limited African-scholar visibility at SoLAR and related venues. The construct-travel problem is not abstract — it is the mismatch between what the analytic frameworks assume about data availability and what the institutions can actually generate.²

EARLY EVIDENCE

The case is positioned as frontier rather than failure or intervention because it does not document a deployed capability outcome. It documents the state of the evidence base, which is itself the relevant capability question for anyone considering whether to deploy a learning-analytics program in an African higher-education institution. The review surfaces the construct-validation work that has not yet been done; the absence of dense primary studies is the finding, and the implication is that any deployer must either undertake the validation work themselves or import constructs with the validation gap acknowledged.³

OPEN PROBLEMS

The practice-synthesis-tier flag is load-bearing. A scoping review is the field's snapshot at a moment in time; the individual primary studies it draws on vary in rigor, and the review is the strongest single map of the present state but not the final account. Future validation ongoing as the literature matures, as institutional infrastructure improves, and as African scholars author more of the learning-analytics research that travels out of the continent rather than into it. The case carries the standing tier language into print.

The absence of dense primary studies is itself the finding.

EDITORS' SYNTHESIS OF PRINSLOO ET AL. (2022).

REFERENCES

1. Prinsloo, P., & Kaliisa, R. (2022), "Learning Analytics on the African Continent: An Emerging Research Focus and Practice," *Journal of Learning Analytics*; ResearchGate publication 361096718.
2. Lemmens, J.-C., & Henn, M. (2015), South African Association for Institutional Research (SAAIR) proceedings — adjacent SA higher-education learning-analytics work.
3. SciELO (2020), "Development of a contextualised learning-analytics framework for South African higher education."
4. Cross-reference: the African data-privacy governance case earlier in the corpus, for the construct-travel problem stated in governance terms.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The scoping review is the present-state map of learning-analytics research on the African continent: 15 studies, geographically concentrated, with the structural barriers (LMS access, institutional resourcing, African-scholar SoLAR visibility) underneath the count. Practice-synthesis-tier — a snapshot of an early-stage literature; future validation ongoing as the field matures.

LENS APPROACH

African learning-analytics scoping is the non-US frontier-evidence case (induced 8.4; LENS D5/PT4). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the construct-travel problem and in Domain 4 (Test and Evaluation) for the CLO **Judgment under inadequate evidence** — the deployer must decide on what to import and what to validate locally when the validation literature is thin. Pairs with the African data-privacy governance case for the construct-travel problem stated in governance terms.

DOMAIN	MODES AT STAKE		LENS COMPETENCY					PROBLEM TYPE	INDUCED	CLO	
LEARNING ANALYTICS	➔	K N	➔	D1	D2	D3	D4	D5	4	8.4	4 · 5
HIGHER EDUCATION		Navigating Sociotechnical Constraints									
AFRICA											

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ When importing learning-analytics constructs into an African higher-education context, treat the scoping-review map as the precondition: do not assume the construct's validation literature applies and identify which validation work has to be repeated locally.
- ▶ Specify the institutional preconditions the analytic framework presupposes — LMS coverage, data-engineering capacity, ethics-review infrastructure — and audit them honestly against the deploying institution's actual conditions.
- ▶ Carry the practice-synthesis-tier flag through any decision document the scoping review supports; the review is the strongest current map but not the final account.

IN OPERATION / AFTER

- ▶ Treat each deployed learning-analytics program in an African institution as a contribution to the construct-validation literature, not only as a service intervention; publish the validation work where SoLAR and related venues can absorb it.
- ▶ Track institutional infrastructure conditions as a separate variable from analytic-construct performance; a construct that travels under one set of conditions may not travel under another.
- ▶ When the literature matures, retire the scoping review as the present-state map and replace it with denser primary-study evidence; the case is included because the present-state matters now, not because the present-state should persist.

REFLECTION QUESTIONS

1. Identify a learning-analytics construct your program imports from US/UK/EU literature. What does the validation literature for that construct assume about LMS coverage, data-engineering capacity, and ethics review? Audit those assumptions against your institution's actual conditions.
2. The case rests on a scoping review of an early-stage literature. What would the minimum additional evidence — replication primary studies, multi-institution validation, longitudinal outcome data — look like before you would treat any specific African learning-analytics finding as settled?
3. Specify the institutional preconditions you would identify and the validation work you would commit to before deploying an analytic framework in a low-resource higher-education setting; treat your deployment as a contribution to the literature, not only as a service intervention.

WHO BUILDS THIS knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · incident analysis · safety dashboards

Chapter 5

Aviation & Aerospace — What Fails

When automation, transition, and assumption meet an unforgiving envelope.

The aircraft did what it was designed to do. So did the crew. The designs disagreed.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

AeroPerú Flight 603

H Human–System Interface **T** Training Gap

IMPACT 70 killed; Boeing 757 crashed into the Pacific after maintenance tape was left over static ports

IN BRIEF

While polishing the fuselage of a Boeing 757 in Lima in October 1996, maintenance staff taped over the static ports and failed to remove the tape before flight. With the static system blocked, the aircraft's altimeters, airspeed indicators, and altitude–alert and overspeed warnings all read inconsistently. The crew received a cascade of contradictory warnings — overspeed, stick shaker, ground proximity — and could not determine their true altitude or speed. Believing they were higher than they were, they flew into the Pacific, killing all 70 aboard. The investigation named the maintenance error as the primary cause but stressed that every cockpit instrument depended on the same blocked sensor: the apparent redundancy was an illusion, and the crew's training had no procedure for "everything you see is wrong."

BACKGROUND

AeroPerú 603 was a Boeing 757 night departure from Lima to Santiago. Before the flight, ground crew had covered the aircraft's static ports with adhesive tape during cleaning and polishing of the fuselage, and the tape was not removed. The static ports feed the instruments that tell the crew their altitude and airspeed — a single physical source upstream of nearly every primary cockpit display, so blocking it corrupted not one instrument but the whole air–data picture the crew would rely on in the dark.^o

WHAT HAPPENED

On 2 October 1996, the aircraft took off into darkness over the ocean with its static system blocked. Altimeters, airspeed indicators, and the altitude–alert and overspeed systems all produced false and contradictory readings. The crew received simultaneous overspeed and stall warnings and ground–proximity alerts they could not reconcile. Believing they were higher than they were, they descended and struck the Pacific, killing all 70 aboard — the contradictory warnings offering no way to tell which, if any, instrument to trust, because every one of them drew on the same blocked source and so failed in concert rather than disagreeing usefully.¹

THE INVESTIGATION

The Peruvian accident investigation identified the failure to remove the static–port tape as the primary cause, but emphasized that the cockpit had no independent reference by which to detect the inconsistency. Every instrument that should have flagged the failure drew on the same blocked source. The crew's training had no procedure for the case in which all


primary instruments are simultaneously wrong — that failure had been assumed away rather than planned for, leaving the crew to improvise against a problem the system design had quietly made possible and the curriculum never named.²

THE CAPABILITY GAP

AeroPerú 603 is the case for redundancy that is not redundancy. The cockpit's apparent instrument redundancy was an illusion at the source — a common-cause failure that defeated every downstream display at once. The missing capability was both a maintenance control that made the blocked port impossible to miss and a procedure for flying when the instruments themselves cannot be trusted — the first preventing the common-cause failure at its source, the second giving the crew somewhere to stand when the apparent redundancy in front of them collapsed all at once.³

AFTERMATH & REFORM

The accident reinforced industry maintenance practices for conspicuous covering and mandatory removal-verification of static ports and pitot tubes, and entered the human-factors literature as a worked example of common-cause failure and the limits of crew training under total instrument corruption. It is paired with other pitot-static events in discussions of air-data integrity — its lesson being that redundancy counted at the display tells you nothing if the redundant paths share a single upstream point that can fail them all together.⁴



The crew received contradictory indications they could neither reconcile nor override.

PARAPHRASING THE PERUVIAN DGAC / COMISIÓN INVESTIGADORA DE ACCIDENTES DE AVIACIÓN REPORT ON AEROPERÚ 603, 1996

REFERENCES

1. Peru Dirección General de Aeronáutica Civil, accident investigation board, final report on AeroPerú 603 (October 1997; with NTSB/FAA/Boeing participation) — primary cause and the instrument cascade (paraphrased).
2. Peru CIIA report (1997) — the static-port tape and the contradictory warnings.
3. Peru CIIA report (1997) — the absence of an independent cockpit reference.
4. Leveson, N. (2011), Engineering a Safer World (STAMP) — common-cause failure.
5. Dismukes, Berman & Loukopoulos (2007), The Limits of Expertise — crew performance under instrument failure.

THE LEARNING ENGINEERING LENS

LE INSIGHT

AeroPerú 603 is the case for redundancy that is not redundancy. Every cockpit indicator depended on the same physical sensor. The apparent redundancy in the cockpit was an illusion at the source. The training did not include the failure case because the failure case had been assumed not to occur.

LENS APPROACH

LENS uses AeroPerú in LEN 5 to teach the difference between apparent and actual redundancy and in LEN 2 for the human role when all interface data is unreliable. Studio projects examine the “trust nothing” procedure design.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ H T	→ D1 D2 D3 D4 D5	1	3.1	1
Systems Analysis					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Map redundant displays back to their upstream sources, and eliminate or instrument any common point whose failure would corrupt them all at once.
- ▶ Make critical sensor covers conspicuous and require verified removal, so a blocked static port cannot reach takeoff unnoticed.
- ▶ Provide an independent reference (e.g., a source not fed by the common component) so the crew can detect inconsistency at the source.

IN OPERATION / AFTER

- ▶ Audit maintenance controls for tasks that cover or block critical sensors, confirming each has a conspicuous marker and a removal-verification step.
- ▶ Train and rehearse a “trust nothing” procedure for total instrument corruption, the case the curriculum had assumed away.
- ▶ Monitor for common-cause failure modes whose probability was assumed negligible, and revisit that assumption when the consequence is catastrophic.

REFLECTION QUESTIONS

1. Identify a redundant interface in your domain whose redundancy depends on a common upstream component. What is the operator procedure when the upstream component fails?
2. Design the “trust nothing” procedure for the AeroPerú crew. What information remains reliable when all instruments are corrupted?
3. A taped-over static port survived to takeoff because nothing made it conspicuous or forced its removal to be verified. Design the maintenance control in your domain that would make a blocked or covered critical component impossible to overlook before operation.

WHO BUILDS THIS human factors & interaction design · learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · scenario simulation · competency models

Mars Climate Orbiter — Unit Mismatch

D Designed Out **K** Knowledge & Institutional Memory

IMPACT ~\$125M spacecraft lost on approach to Mars; ground software produced thrust output in pound-force while navigation expected newtons

IN BRIEF

The Mars Climate Orbiter, a NASA spacecraft built by Lockheed Martin and navigated by JPL, was lost on arrival at Mars in September 1999. Lockheed's ground software reported thruster impulse in pound-force seconds; JPL's navigation expected newton-seconds. The conversion — a factor of about 4.45 — was never applied at the boundary between the two teams, so every trajectory correction was mis-modeled and the error accumulated over the cruise. The orbiter arrived too low, into the atmosphere, and burned up; the orbiter and its \$125 million were lost to a unit mismatch. The investigation found no individual blunder — both contractors did their own work correctly. What failed was the interface between them, which had no owner, specification, or verification step. It is the canonical case of interface-as-requirement.

BACKGROUND

The Mars Climate Orbiter, launched in December 1998 to study the Martian atmosphere, was one of NASA's "faster, better, cheaper" missions, built on a compressed budget that trimmed margins across the program. Lockheed Martin built the spacecraft and its ground software; JPL navigated it from Earth. Across the months of interplanetary cruise the two teams exchanged data so JPL could command the small trajectory corrections that keep a spacecraft on course — an exchange that crossed a software interface between the two organizations.⁰

WHAT HAPPENED

Lockheed's ground software reported thruster impulse in pound-force seconds; JPL's navigation expected newton-seconds. The conversion — a factor of about 4.45 — was never applied at the boundary where the two systems met. Each firing was mis-modeled by that factor, and the error accumulated steadily over the long cruise until the predicted and actual trajectories had quietly diverged. When the orbiter reached Mars on 23 September 1999 it arrived far too low, deep into the atmosphere it was built to study from orbit, and was destroyed. The orbiter, and its roughly \$125 million, were lost to a unit conversion no one made.¹

THE INVESTIGATION

The Mishap Investigation Board put the proximate cause exactly there — the failed English-to-metric translation in ground software — but was careful to name the deeper one rather than stop at the bug. No individual blundered; both contractors did their own work correctly


within their own assumptions. What failed was the boundary between them: there was no specified, verified interface fixing the units and checking that both sides agreed, and no end-to-end validation of the data flowing across the seam. Navigators had even noticed odd trajectory behavior in cruise, but the concern was never run fully to ground before arrival, and the chance to catch it passed.²

THE CAPABILITY GAP

The missing capability was ownership of the interface itself. Where a system is split across two organizations, the place they meet is not a documentation footnote but an engineering deliverable in its own right — with an owner, a specification, and a verification step. Here it had none of the three. Each team treated its own side as complete and the boundary as someone else's concern, so the one assumption that had to be shared and checked — what units are we speaking in? — was precisely the one no one verified. The spacecraft did not fail; it performed as built; the seam between the two halves of the organization did.³

AFTERMATH & REFORM

NASA tightened interface management and end-to-end verification and treated the loss as a cautionary tale about how far “faster, better, cheaper” could be pushed before the corners being cut turned out to be load-bearing. The orbiter became the canonical systems-engineering example of interface-as-requirement — the civilian, software parallel to the Patriot's untraveled assumption (Case 129): two competent halves of a system, a boundary nobody owned or verified, and a small unspecified thing at that boundary that quietly destroyed the whole.⁴



The root cause was the failed translation of English units to metric units in a segment of ground-based, navigation-related mission software.

PARAPHRASING THE NASA MARS CLIMATE ORBITER MISHAP INVESTIGATION BOARD, 1999

REFERENCES

1. NASA, Mars Climate Orbiter Mishap Investigation Board: Phase I Report (Nov. 1999) — mission, the Lockheed/JPL split, and the program context.
2. NASA MCO MIB Report (1999) — the pound-force-second vs newton-second mismatch (factor 4.45), the accumulated navigation error, and atmospheric destruction on 23 Sept. 1999 (root-cause statement quoted).
4. N. G. Leveson, *Engineering a Safer World* (MIT Press, 2011) — interfaces as engineering deliverables requiring an owner and a verification step.
5. B. Sauser et al. (2009), retrospective analysis of the Mars Climate Orbiter and the “faster, better, cheaper” trade space.

- 3. NASA MCO MIB Report (1999) — the missing verified interface specification, the absent end-to-end validation, and the unresolved cruise-trajectory anomalies.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Mars Climate Orbiter is the textbook case for interface boundaries as engineering deliverables. The contractors did their work. The boundary between them did not have an owner. The capability that was missing was the interface-specification verification step.

LENS APPROACH

LENS uses Mars Climate Orbiter in LEN 5 to teach interface-as- requirement and in LEN 8 to discuss cross-contractor capability boundaries. The case is the foundational software-engineering parallel to Patriot (Case 129).

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
SPACE / AEROSPACE	→ D K	→ D1 D2 D3 D4 D5	1	1.3	1
Systems Analysis					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Make the interface between two organizations an owned deliverable with a written specification of units and formats, not a footnote each side assumes the other handles.
- ▶ Fix and verify the unit convention at the boundary explicitly, requiring both sides to confirm agreement before data crosses the seam.
- ▶ Run end-to-end validation across the combined system, since each half can be correct within its own assumptions while their composition is wrong.

IN OPERATION / AFTER

- ▶ Treat in-cruise anomalies, like the noticed trajectory behavior, as signals to run to ground before a point of no return rather than concerns to revisit after arrival.
- ▶ Monitor predicted-versus-actual trajectory continuously so an accumulating boundary error surfaces while there is still time to correct it.
- ▶ Audit cross-contractor seams for unowned assumptions whenever a program trims margins under “faster, better, cheaper,” confirming the corners cut are not load-bearing.

REFLECTION QUESTIONS

1. Identify a contractor-to-contractor interface in your domain whose specification ownership is unclear. What would the equivalent unit-mismatch look like there?
2. Design the interface-verification deliverable that would have caught the pound-force / newton boundary before launch.
3. Navigators noticed odd trajectory behavior in cruise but the concern was never run to ground before arrival. Specify the threshold and process by which an in-flight anomaly in your domain must be resolved before a point of no return, rather than carried past it.

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines

Glass-Cockpit Transition in Light Aircraft — Technology Outran Training

H Human-System Interface **N** Normalization of Deviance **K** Knowledge & Institutional Memory

IMPACT An NTSB safety study of ~8,000 piston aircraft manufactured 2002–2006, with accidents analyzed for 2002–2008, found that glass-cockpit aircraft had no better overall safety record — and a higher fatal accident rate — than comparable conventional-instrument aircraft over the period studied; the Board attributed this to inadequate equipment-specific training and issued recommendations A-10-36 through A-10-41

IN BRIEF

Glass cockpits — integrated digital primary flight displays and multifunction displays replacing the inherited six-pack of analog instruments — were introduced into light piston aircraft over the 2000s as a fleet-wide modernization. The NTSB safety study NTSB/SS-10/01 examined approximately 8,000 small piston aircraft manufactured 2002–2006 and reached a finding the technology's own advocates did not expect: glass-cockpit aircraft had no better overall safety record than comparable conventional-instrument aircraft over the period studied, and in fact had a higher fatal accident rate. The NTSB attributed this not to the technology but to inadequate equipment-specific training. The Board issued recommendations A-10-36 through A-10-41 on knowledge-testing standards, simulator availability, and training requirements. The study is explicit that advanced avionics "can increase the safety potential" of light aircraft but that the potential "was not yet realized in the period studied" — an open, not closed, verdict. The case is the canonical 7.1 failure of an inherited capability regime (pilot proficiency) not being re-verified against the new envelope the technology transition introduced. Pair with the aging-system transition cases (Cases 114–116, 154, 156, 174, and 194, drafted in parallel).

BACKGROUND

Glass-cockpit avionics — integrated digital primary flight displays and multifunction displays replacing the inherited six-pack of analog round-dial instruments — were introduced into light piston aircraft over the 2000s. The fleet conversion was largely industry-driven; new-build aircraft came with glass as standard, retrofits became increasingly common, and the FAA's certification framework treated the transition as a positive safety move on the strength of the capability the displays delivered to pilots: integrated situational awareness, moving-map navigation, terrain and traffic display, system-status integration. The operationally critical assumption was that pilots transitioning from analog to glass would carry their proficiency across the change.^o

WHAT HAPPENED

The NTSB safety study NTSB/SS-10/01 examined that assumption directly. The Board studied approximately 8,000 small piston airplanes registered in the United States over the period 2002–2006, comparing accident rates between conventional-instrument and glass-cockpit fleets matched on model and operational class. The headline finding ran against expectation. Glass-cockpit aircraft had no better overall safety record than comparable conventional-instrument aircraft over the period, and the fatal accident rate was higher for the glass fleet. The expected fleet-wide safety gain from the displays had not yet materialized in the accident record.¹

THE INVESTIGATION

The NTSB's attribution is precise. The technology was not the cause; inadequate equipment-specific training was. The transitioning pilot population was certificated and current under the inherited training regime that assumed analog instruments, and the glass displays — for all their situational-awareness advantages — introduced new failure modes (mode confusion, automation surprises, attention capture by integrated displays, degraded scan in degraded modes) that the inherited training did not address. The Board issued recommendations A-10-36 through A-10-41 to the FAA: equipment-specific knowledge-testing standards, simulator and training-device availability for transitioning pilots, and structured training requirements before operating glass-cockpit aircraft.²

THE CAPABILITY GAP

The study's language is open rather than closed. Advanced avionics “can increase the safety potential” of light aircraft, but that potential “was not yet realized in the period studied.” The verdict is not that glass is unsafe; it is that the inherited certification of pilot proficiency, built around analog instruments, did not transfer to the new envelope without re-engineered training. The case is the canonical 7.1 failure: a capability regime (pilot proficiency) was not re-verified when the cockpit's interface envelope changed, and the inherited training assumptions silently became inadequate to the new operational reality.³

AFTERMATH & REFORM

Drafted alongside the aging-system transitions in Cases 22–76, the glass-cockpit GA study carries the capability-under-system-change pattern at the consumer-aviation scale where the analog cases run at the defense and commercial-aviation scale. The structural lesson is the same: when a system's interface envelope changes — and even when the change is a capability-enhancing one — the inherited certification of operator proficiency has to be re-verified against the new envelope, or the transition outruns the training and the accident record moves in the wrong direction. The NTSB recommendations are the engineering response to that pattern, and their implementation is what closes (or fails to close) the open verdict the study left.

Advanced avionics and electronic displays can increase the safety potential of general aviation aircraft operations, but more effort is needed to ensure that pilots are prepared to realize that potential.

NTSB SAFETY STUDY NTSB/SS-10/01 (2010).

REFERENCES

1. National Transportation Safety Board (2010). Introduction of Glass Cockpit Avionics into Light Aircraft, Safety Study NTSB/SS-10/01. <https://www.nts.gov/safety/safety-studies/Documents/SS1001.pdf> — the case’s primary investigation document.

2. NTSB Safety Recommendations A-10-36 through A-10-41 (2010), issued to the FAA — the engineering response the open verdict points to.

3. Wiener, E. L., & Curry, R. E. (1980). Flight-deck automation: Promises and problems. *Ergonomics*, 23(10):995–1011 — the foundational literature on automation-induced failure modes that the glass-cockpit transition re-introduced at the GA scale.

4. Sarter, N. B., Woods, D. D., & Billings, C. E. (1997). Automation surprises. In *Handbook of Human Factors and Ergonomics* (2nd ed.) — the mode-confusion / automation-surprise literature the NTSB findings cross-reference.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The NTSB’s glass-cockpit GA study is the canonical capability-under-system-change failure at the consumer-aviation scale: a positive technology transition that nevertheless outran the inherited certification of operator proficiency. The Board’s attribution is to inadequate equipment-specific training, and the verdict is open — the safety potential is there, and the transition has not yet realized it.

LENS APPROACH

Glass-cockpit GA is the small-fleet capability-under- transition failure (induced 7.1; LENS D3/PT1). LENS uses it in Domain 1 (Systems Analysis) for the inherited certification of operator proficiency that has to be re-verified against the new envelope, and in Domain 3 (Human-System Collaboration) for the new automation-induced failure modes (mode confusion, automation surprises, attention capture) the transition introduced. Pair with the aging-system transition cases (Cases 114–116, 154, 156, 174, 194) at the cross-scale capability-under-change layer.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
GENERAL AVIATION	H N K	D1 D2 D3 D4 D5	1	7.1	1-2
AVIONICS		Systems Analysis + D3			
PILOT TRAINING					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Specify the inherited certification of operator proficiency in advance of a technology transition, and design the re-verification against the new envelope as part of the transition’s deliverable, not as a downstream training catch-up.
- Treat capability-enhancing transitions with the same re-verification discipline as capability-degrading ones; the glass-cockpit transition was a positive technology move that nevertheless required training re-verification the field skipped.
- Name the new failure modes the transition introduces (mode confusion, automation surprises, attention capture) at the design stage of the equipment-specific training, rather than waiting for the accident record to surface them.

IN OPERATION / AFTER

- Preserve the NTSB’s open-verdict language (“can increase the safety potential ... not yet realized in the period studied”) in any report on the transition’s outcome; the study did not say glass is unsafe, and the precise language is what makes the recommendation set actionable.
- Track implementation of the NTSB recommendation set (A-10-36 through A-10-41) as the engineering response to the open verdict; the verdict closes when the recommendations are implemented and the next round of evidence is collected.
- Carry the structural lesson into adjacent transitions — the aging-system transition cases (Cases 114–116, 154, 156, 174, 194) — as evidence that the inherited certification of operator proficiency has to be re-verified against the new envelope across consumer, commercial, and defense scales.

REFLECTION QUESTIONS

1. Identify a technology transition in your domain that is positive on its capability claim and that nevertheless puts the inherited certification of operator proficiency under question. What re-verification of operator proficiency would the transition require — and is it currently part of the deliverable, or assumed away?
2. The NTSB attribution is to inadequate equipment-specific training, not to the technology. What is the analog distinction in your context — between the engineered artifact and the operator-proficiency regime that has to track it?

WHO BUILDS THIS human factors & interaction design · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · incident analysis · safety dashboards · knowledge bases · data pipelines

Ariane 5 Flight 501

D Designed Out **K** Knowledge & Institutional Memory **H** Human-System Interface

IMPACT Maiden flight destroyed itself 37 seconds after launch; ~\$500M payloads lost; reused Ariane 4 code never re-verified for Ariane 5

IN BRIEF

On its 1996 maiden flight, the Ariane 5 rocket destroyed itself 37 seconds after launch, losing around half a billion dollars in payloads. The cause was reused software: the inertial reference system inherited code from Ariane 4, where a horizontal-velocity value never exceeded a 16-bit integer's range. Ariane 5's steeper, faster trajectory pushed it past that range; the integer overflow crashed both redundant reference systems almost simultaneously, the vehicle lost guidance and broke up. The inquiry found the offending code path had been irrelevant on Ariane 4 and was never removed or re-verified for Ariane 5 — software reuse had been treated as risk-free. Ariane 5 is the foundational safety-critical-software case for the hazard of reusing code without re-verifying it against the new system's operating envelope.

BACKGROUND

The Ariane 5 heavy-lift rocket reused proven components from its successful predecessor, Ariane 4 — including, in its inertial reference system, software that had flown reliably for years. Reuse of flight-proven code was regarded as a way to reduce risk and cost.⁹ The reasoning was seductive precisely because the code's flight record was genuine: software that had worked on Ariane 4 looked like the safest possible choice, and that confidence is what allowed it to cross onto a different vehicle without being re-examined against the new vehicle's conditions.

WHAT HAPPENED

On 4 June 1996, Ariane 5 Flight 501 — its maiden flight — veered off course and broke up under aerodynamic stress 37 seconds after launch; the range-safety system destroyed it, and roughly half a billion dollars of payloads were lost.¹ The vehicle was destroyed in well under a minute of flight, before it had done anything but climb — the whole half-billion-dollar loss flowing from a fault that triggered almost immediately at liftoff, when the new trajectory first pushed the inherited code outside the range it had been built to handle.

THE INVESTIGATION

The inquiry board traced the cause to the reused code. A horizontal-velocity value that never exceeded a 16-bit signed integer on Ariane 4's trajectory did exceed it on Ariane 5's faster one; the resulting overflow shut down both redundant inertial reference systems almost simultaneously, and the vehicle lost guidance.² Redundancy gave no protection here because both units ran the identical inherited code and met the identical out-of-range

value at the same instant, so the backup failed in lockstep with the primary — duplicating a system defends against a part breaking, not against a shared assumption being wrong. The offending code path had become irrelevant to Ariane 5's flight after liftoff, yet had neither been removed nor re-verified against the new vehicle's envelope — reuse had been treated as risk-free.³

THE CAPABILITY GAP

Ariane 5 is the canonical software-engineering case for the fallacy of risk-free reuse. Code is fit for the envelope it was written and verified against; reusing it in a different envelope is a new design decision and must be re-verified as one. The institutional capability that was missing was the requirement that reused safety-critical code be re-verified, not merely trusted because it had worked before.⁴ The flaw was never in the Ariane 4 code, which was correct for the trajectory it was written against; the flaw was the institutional assumption that a record of working in one envelope carried over to another, an assumption no test was required to challenge.

AFTERMATH & REFORM

The fix and the lesson reshaped safety-critical software practice: reuse became a verification event, with explicit checking of every inherited assumption against the new system's operating conditions.⁵ The case rhymes with the Patriot (Case 129) and the Mars Climate Orbiter (Case 98) — small, unexamined assumptions inherited across a boundary, fatal when the boundary's conditions changed. What unites the three is that each inherited a quantity sound in its original setting and lethal in the new one, so the reform is the same in every case: make the boundary crossing the moment the assumption is re-checked, rather than the moment it is silently trusted.

37s

AFTER LAUNCH · 16-BIT INTEGER OVERFLOW

code path disabled by the previous vehicle's profile; re-enabled by the new one

ARIANE 5 — THE FALLACY OF RISK-FREE CODE REUSE

The internal SRI software exception was caused during execution of a data conversion from 64-bit floating point to 16-bit signed integer value.

ARIANE 5 FLIGHT 501 FAILURE INQUIRY BOARD, 1996

REFERENCES

1. J. L. Lions (chair), Ariane 5 Flight 501 Failure Inquiry Board Report (1996) — the data-conversion overflow (quoted).
2. Lions Report (1996) — the 37-second breakup and the lost payloads.
3. Lions Report (1996) — the reused code path neither removed nor re-verified for Ariane 5's envelope.
4. N. Leveson, Safeware (1995) — software-reuse hazards; G. Le Lann, "An Analysis of the Ariane 5 Flight 501 Failure" (1997).
5. Cf. Patriot (Case 129) and Mars Climate Orbiter (Case 98).

3. Lions Report (1996) — the 64-bit-to-16-bit conversion over-flow and the simultaneous shutdown of both inertial reference systems.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Ariane 5 is the canonical software-engineering case for the fallacy of risk-free code reuse. Code is fit for its envelope of operation. Reusing it in a different envelope is a new design decision and must be re-verified as one. The institutional capability that was missing was the requirement to re-verify.

LENS APPROACH

LENS uses Ariane 5 in LEN 5 for software-engineering capability deliverables and in LEN 8 for the institutional discipline that treats reuse as a verification event rather than as a savings.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
SPACE / AEROSPACE	→ D K H	→ D1 D2 D3 D4 D5	1	7.2	1
Systems Analysis					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat every reuse of safety-critical code as a new design decision: re-verify each inherited assumption against the new system’s operating envelope before flight.

▶ Document the original operating envelope a component was verified against, so a later reuse can be checked against it rather than blindly trusted.

▶ Design redundancy to fail independently — diverse implementations or inputs — so a shared inherited assumption cannot take out primary and backup at once.

IN OPERATION / AFTER

- ▶ Audit reused components for code paths that are irrelevant in the new system but still active, and remove or re-verify them rather than leaving them dormant.

▶ Make boundary crossings — code moving onto a new vehicle, a new envelope — a mandatory re-verification event in the institution’s process, not a savings.

▶ Review past reuse decisions for the Ariane / Patriot / Mars Climate Orbiter pattern: a quantity sound in its original setting and unchecked in the new one.

REFLECTION QUESTIONS

1. Identify a piece of reused infrastructure in your domain whose original operating envelope is not documented. What is the silent assumption?
2. Design the verification deliverable that should accompany every reuse of safety-critical software.
3. Both redundant reference systems failed at the same instant because they shared the same inherited assumption. Identify a place in your domain where redundancy gives false comfort because the duplicated units share a common flaw rather than fail independently.

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · usability testing · cognitive walkthroughs

Kegworth / British Midland 92

T Training Gap **H** Human-System Interface **K** Knowledge & Institutional Memory

IMPACT 47 killed and 74 seriously injured when the crew shut down the wrong engine

IN BRIEF

On 8 January 1989 British Midland Flight 92, a nearly new Boeing 737-400, crashed on the M1 motorway embankment near Kegworth, killing 47 and seriously injuring 74. After a fan blade fractured in the left engine, the crew shut down the right engine — the one still working. Reasoning from older 737s, on which the right engine fed the cabin air, they misattributed the smoke; new, harder-to-read electronic engine displays did not correct them, and the cabin crew who saw flames never told the flight deck. The pilots' mental model was right for the airplane they had flown the week before, but the brief conversion course never overwrote it. The AAIB issued 31 recommendations; Kegworth became the textbook case of capability degrading under system change.

BACKGROUND

On 8 January 1989 British Midland Flight 92, a nearly new Boeing 737-400 (G-OBME), left Heathrow for Belfast with 126 aboard. The -400 was a recent variant — bigger engines, a partly redesigned cockpit with new electronic engine instruments. The crew were experienced 737 pilots, but their conversion onto the variant had been brief and largely self-directed, the differences communicated as reading rather than drilled retraining — so the experience that made them confident on the type also seeded the prior model that would mislead them.⁰

WHAT HAPPENED

Climbing through 28,300 feet, a fan blade fractured in the left engine, filling the cabin with vibration and smoke. Reasoning from older 737s — on which the air-conditioning bleed drew from the right engine — the crew concluded the right engine was the source and throttled it back. The symptoms eased (partly because the autothrottle had also cut power on the failing left engine), seeming to confirm it, and they shut the right engine down. On final approach the damaged left engine failed completely; with the good engine off, the aircraft struck the M1 embankment short of the runway — 47 killed, 74 seriously injured. The brief calm after the shutdown was the trap: it appeared to confirm a diagnosis that was wrong, removing the doubt that might have prompted a reassessment.¹

THE INVESTIGATION

The Air Accidents Investigation Branch traced the disaster to the shutdown of the serviceable engine, and to a mental model carried from earlier 737s and applied to a -400 whose bleed-air configuration Boeing had changed — a difference the conversion training never disturbed, so the crew reasoned correctly from a configuration that no longer applied to the


airframe they were flying.² The new electronic displays did not help: the vibration indicator that would have pointed at the failing engine was harder to read at a glance than the dials it replaced, so the one instrument that could have corrected the diagnosis was the least likely to be consulted under pressure.³ And though cabin crew and passengers could see flames from the left engine, that never reached the flight deck — a crew-resource-management gap aviation would later train hard against, because those with the decisive evidence had no path to those making the decision.⁴

THE CAPABILITY GAP

The crew were skilled, and their mental model was correct — for the airplane they had flown the week before. The -400 was treated, for training, as an incremental change to a familiar type; under emergency pressure it behaved as a categorical one. Reading a manual note about revised bleed air is not the same as drilling a new reflex that fires before the old one, and the cockpit redesign had quietly removed perceptual cues a startled crew relies on. The change was filed as incremental precisely because nothing tested it until an emergency did. Capability had degraded under system change with no one noticing, because nothing failed until the day it did.⁵

AFTERMATH & REFORM

The AAIB issued thirty-one safety recommendations spanning conversion and difference training, engine-instrument and vibration-display design, cabin-to-flight-deck communication, and crashworthiness.⁶ Kegworth became a standard human-factors teaching case and a reference for how a transition program should be built — so the differences between an old system and a new one are made hard to overlook, and the people closest to the evidence have a path to those making the decision. Each recommendation traced back to a failure the accident exposed — the unforced mental model, the unreadable instrument, the silent cabin — treated together rather than as isolated faults.



The crew's mental model of the older 737 was inappropriately applied to the 737-400 on which they were operating.

PARAPHRASING AAIB AIRCRAFT ACCIDENT REPORT 4/90 ON BRITISH MIDLAND 92, 1990

REFERENCES

1. Air Accidents Investigation Branch, Report on the accident to Boeing 737-400 G-OBME near Kegworth, Leicestershire, on 8
5. AAIB Report 4/90 (1990) — the cabin-to-flight-deck communication gap (crew resource management).

January 1989, AAIB Aircraft Accident Report 4/90 (1990) — aircraft, route, and the limited conversion training; see also U.S. FAA, Lessons Learned: G-OBME.

2. AAIB Report 4/90 (1990) — accident sequence: left-engine fan-blade failure, shutdown of the serviceable right engine, and total power loss on approach; 47 killed, 74 seriously injured.

3. AAIB Report 4/90 (1990) — the older-737 bleed-air mental model misapplied to the 737-400's changed configuration.

4. AAIB Report 4/90 (1990) — the new electronic engine instrumentation and the readability of the vibration indicator.

6. AAIB Report 4/90 (1990) — the 31 safety recommendations on difference training, instrument design, CRM, and crashworthiness.

7. J. Reason, Human Error (Cambridge Univ. Press, 1990) — Kegworth as a case in the misapplication of a correct mental model to a changed system.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Kegworth is the textbook example of Capability Degradation Under System Change. The pilots were excellent. Their mental model was excellent — for the airframe they had flown the previous week. The difference training had not been engineered to disturb the prior model with enough force to keep it from being misapplied. The new airframe was treated as an incremental change. The crew's response revealed it was a categorical one.

LENS APPROACH

LENS treats Kegworth in LEN 5 as the canonical **system-change** capability problem and in LEN 8 as the institutional version: how knowledge about what has changed travels from engineering to operations and what gets lost in transit. The case sits alongside Patriot (Case 129) in the canonical problem-type pair for system transition.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ T H K	→ D1 D2 D3 D4 D5	6	3.3	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Classify any change that alters operator mental models — even a “minor” variant — as categorical, and gate it on drilled retraining rather than reading.
- ▶ Test new interface designs against startled-crew performance before fielding, so a redesign cannot quietly remove the perceptual cues operators rely on.
- ▶ Engineer a communication path from those closest to the evidence (cabin crew, passengers, sensors) to those making the decision, so decisive observations reach the flight deck.

IN OPERATION / AFTER

- ▶ Audit transition programs for capability that degraded silently — where nothing failed until an emergency tested the change — using simulator scenarios that force the new reflex.
- ▶ Monitor whether difference training actually overwrites the prior model, treating un-disturbed mental models as a measurable latent risk.
- ▶ Track that interface and communication recommendations are implemented together, not as isolated fixes, so the integrated failure mode does not recur.

REFLECTION QUESTIONS

1. The Kegworth crew's mental model was right for the previous variant. What change in your domain currently risks an analogous misapplication?
2. Difference training is a generic deliverable in transition programs. Design the artifact that would make differences hard to overlook rather than easy.
3. The cabin crew and passengers saw flames the flight deck never learned of. Identify a path in your domain by which the people closest to the evidence cannot reach the people making the decision, and design the channel that would carry it under emergency conditions.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs · knowledge bases · data pipelines

Asiana Airlines Flight 214

T Training Gap **H** Human-System Interface

IMPACT 3 killed, 187 injured; Boeing 777 struck the seawall short of SFO runway 28L

IN BRIEF

On a clear afternoon at San Francisco in July 2013, the crew of Asiana 214 allowed their Boeing 777 to slow far below approach speed without recognizing that the autothrottle was not maintaining it. The aircraft struck the seawall short of runway 28L, killing three of the 307 aboard and injuring nearly two hundred. The NTSB found the crew had mismanaged the approach and inadequately monitored airspeed — but also that the complexity of Boeing’s autoflight systems contributed: after the pilots disconnected the autopilot and pulled thrust to idle, the autothrottle reverted to a HOLD mode in which it would not wake to hold the selected speed. The captain believed it would. Asiana 214 is the most prominent recent case of automation surprise on a Western wide-body airliner.

BACKGROUND

Asiana 214 was a scheduled Boeing 777 flight from Seoul to San Francisco. The captain was experienced on other types but new to the 777 and flying it under the supervision of a training captain; the approach was a visual one, with the airport’s instrument glideslope out of service for construction — so a pilot still learning the type’s automation was hand-flying an approach stripped of the vertical guidance that would normally have anchored it, exactly the conditions under which an unfamiliar mode is most likely to bite.⁰

WHAT HAPPENED

On 6 July 2013, on final approach to runway 28L, the aircraft descended below the proper path and slowed. The crew had set an autopilot mode that left the autothrottle in HOLD, where it would not advance thrust to hold the selected speed. Airspeed decayed to about 103 knots against a target of 137; the 777 struck the seawall short of the runway, killing three and injuring nearly two hundred of the 307 aboard. The thirty-four-knot shortfall built up silently while the crew believed the automation was holding the speed, so the first unambiguous signal of trouble was the airframe itself running out of energy on short final.¹

THE INVESTIGATION

The NTSB found the probable cause to be the crew’s mismanagement of the approach and inadequate monitoring of airspeed, complicated by the complexity of the autothrottle and autopilot flight-director systems and by the crew’s faulty mental model of the automation. The captain told investigators he had assumed the autothrottle would maintain speed; in that configuration, it does not — a gap between the system’s documented behavior and the mental model the training had left him with, the kind of mismatch that surfaces only when the automation is asked to do what the operator wrongly believes it is already doing.²

THE CAPABILITY GAP

Asiana 214 is the textbook automation surprise on a Western wide-body: the mode the autothrottle was actually in and the mode the crew believed it was in diverged silently. The missing capability was not raw pilot skill but transparency — a system that made its own state, and the fact that the crew now owned the airspeed, salient at the moment it mattered, rather than leaving the handoff of responsibility to be inferred from a silent mode change the interface did nothing to announce.³

AFTERMATH & REFORM

The NTSB issued recommendations on autoflight training, energy-state monitoring, and the design of mode annunciation, and emphasized training that preserves manual-flying proficiency to counter automation dependence. The accident became a standard reference in airline automation-policy reviews and in the push toward “automation airmanship” curricula — shifting the lesson from one crew’s error to a recurring design problem about how modes are annunciated and how monitoring is taught when the automation is doing most of the flying.⁴



Complexities of the autothrottle and autopilot flight director systems contributed to the accident.

NTSB AIRCRAFT ACCIDENT REPORT AAR-14/01, CONTRIBUTING FACTORS, 2014

REFERENCES

1. NTSB, Aircraft Accident Report: Asiana Airlines Flight 214, NTSB/AAR-14/01 (2014) — probable cause and contributing factors (quoted).
2. NTSB/AAR-14/01 (2014) — the autothrottle HOLD reversion and airspeed decay to 103 kt.
3. NTSB/AAR-14/01 (2014) — the captain’s mental model of the autothrottle.
4. Sarter, N. & Woods, D. (1995), “How in the world did we ever get into that mode?,” Human Factors — automation surprise.
5. NTSB/AAR-14/01 (2014), safety recommendations on autoflight training and mode annunciation.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Asiana 214 is the aviation case for the LENS Human-AI Teaming proposition that automation transparency is a capability deliverable. The mode the autothrottle was actually in and the mode the crew believed it was

in diverged silently — a textbook automation surprise. The training, the documentation, and the interface together did not surface the divergence.

LENS APPROACH

LENS treats Asiana 214 in LEN 2 as a worked case for automation transparency requirements. Students design the interface and training intervention that would have made the mode disconnect salient to the crew at the moment the aircraft slowed.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ T H	→ D1 D2 D3 D4 D5	6	3.3	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Make automation transparency a design requirement: the system must announce its active mode and any silent reversion, not leave it to be inferred.
- ▶ Annunciate transfers of responsibility (e.g., when the crew now owns airspeed) explicitly at the moment they occur, so no handoff is implicit.
- ▶ Design and validate mode logic against operator mental models during development, so configurations that “do not do what pilots assume” are caught before fielding.

IN OPERATION / AFTER

- ▶ Train and audit energy-state monitoring so a slow, silent airspeed decay is caught by the crew before the airframe signals it.
- ▶ Preserve manual-flying proficiency through recurrent practice to counter the automation dependence that hid the decay.
- ▶ Monitor operations for mode-confusion events and feed them into automation-policy and annunciation-design reviews.

REFLECTION QUESTIONS

1. Identify a mode in an automated system you work with where the operator’s mental model and the system’s actual behavior can silently diverge. How would the operator know?
2. Design the interface change that would have alerted the Asiana crew to the mode mismatch before the airspeed decayed.
3. The autothrottle’s HOLD reversion handed airspeed back to a crew that did not know it now owned it. Design the annunciation that would make a transfer of responsibility between automation and operator unmissable at the instant it occurs.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · usability testing · cognitive walkthroughs

Helios Airways Flight 522

T Training Gap **H** Human-System Interface

IMPACT 121 killed; cabin failed to pressurize after a maintenance setting was left on "manual" and the crew misread the warning

IN BRIEF

Shortly after takeoff from Larnaca in August 2005, the crew of Helios 522 misread a warning horn. The same intermittent horn that signals an incorrect takeoff configuration on the ground sounds, in flight, when cabin altitude climbs past about 10,000 feet — and the pressurization selector had been left in "manual" after a leak check the night before. The cabin never pressurized. The crew, troubleshooting the wrong fault, lost consciousness to hypoxia; the Boeing 737 flew on autopilot until it ran out of fuel and crashed into a hillside near Athens, killing all 121 aboard. The investigation found a single cue carried two meanings with no differentiation, and the airline's training had drilled only the ground meaning.

BACKGROUND

Helios 522 was a Boeing 737 departing Larnaca, Cyprus, for Athens. The night before, engineers had run a pressurization leak check and left the cabin-pressurization selector in the "manual" position. The pre-flight checks did not catch it, and the aircraft departed with its pressurization system not set to maintain cabin altitude automatically — a maintenance setting left in an abnormal position the following morning, the kind of latent configuration error that a pre-flight verification exists precisely to catch and here did not.⁰

WHAT HAPPENED

As the aircraft climbed on 14 August 2005, the cabin failed to pressurize and a warning horn sounded. On the 737 the same horn indicates an incorrect takeoff configuration on the ground and excessive cabin altitude in flight. The crew interpreted it as a configuration warning and spent minutes troubleshooting the wrong problem while hypoxia set in. By the time the aircraft neared Athens the flight crew was incapacitated; it circled on autopilot until fuel exhaustion and crashed, killing all 121 aboard. The crew spent their few useful minutes diagnosing a configuration fault that did not exist, the ambiguous horn having steered them away from the one problem — climbing cabin altitude — that was quietly disabling them.¹

THE INVESTIGATION

The Hellenic Air Accident Investigation Board found the crew's misidentification of the cabin-altitude warning as a takeoff-configuration warning to be a critical, irrecoverable error. The single horn carried two meanings with no differentiation between them, and the airline's training had emphasized only the ground meaning; the in-flight case was an edge condition the crew had not been prepared to recognize — so an ambiguous cue and a one-sided

training regime combined to make the wrong interpretation the natural one, with hypoxia closing the window before the error could be caught.²

THE CAPABILITY GAP

Helios 522 is the textbook example of a single cue carrying two meanings without differentiation — an interface that lied by ambiguity while the training filled in only one of the two meanings. The missing capability was a cue, or a training regime, that distinguished the deadly in-flight meaning from the routine ground one before the operators had to guess under hypoxia — a way to resolve the ambiguity at the interface or in training, rather than leaving it to be resolved by a crew already losing the cognition the decision required.³

AFTERMATH & REFORM

The investigation drove changes to Boeing's warning logic and to pressurization-related training and checklists across operators; later 737 variants distinguish the cabin-altitude warning from the takeoff-configuration warning. The accident is a standard case study in cue ambiguity, pre-flight verification, and the physiology of hypoxia in human-factors curricula — its remedies attacking the failure on both fronts at once, differentiating the warning at the interface and reinforcing the in-flight meaning in training so neither gap could be papered over by the other.⁴

ACTUAL	INTERPRETED
cabin alt.	T/O config.

one horn, two meanings · the training had drilled only one

HELIOS 522 — A CUE CARRYING TWO MEANINGS WITHOUT DIFFERENTIATION

The crew's interpretation of the cabin-altitude warning horn as a takeoff-configuration warning was a critical and irrecoverable error.

PARAPHRASING THE HELLENIC AAIB FINAL REPORT ON HELIOS 522, 2006

REFERENCES

1. Hellenic Air Accident Investigation & Aviation Safety Board, Aircraft Accident Report 11/2006 — probable cause and the warning misinterpretation (paraphrased).
2. AAIASB Report 11/2006 — the manual pressurization selector and failure to pressurize.
3. AAIASB Report 11/2006 — the dual-meaning warning horn and the training emphasis.
4. Boeing 737 Flight Crew Operations Manual — configuration and cabin-altitude warnings.
5. Reason, J. (1990), Human Error — the cue-ambiguity / human-error framing the case draws on.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Helios 522 is the textbook example of a single cue carrying two meanings without differentiation. The interface lied by ambiguity. The training filled in only one meaning. The combination produced a twenty-minute window in which the crew was solving the wrong problem.

LENS APPROACH

LENS uses Helios in LEN 5 as a case in cue ambiguity and in LEN 2 as a Human-AI Teaming case where the interface failed to communicate which of two possible system states it was warning about. Students redesign the cue.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ T H	→ D1 D2 D3 D4 D5	3	3.1	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Differentiate cues that carry distinct meanings in distinct contexts, so a single signal cannot be read as the wrong one under stress.
- Train both meanings of any dual-purpose cue, not just the routine one, so the rare-but-deadly interpretation is recognized first.
- Design pre-flight verification to catch abnormal maintenance configurations, so a setting left in “manual” cannot survive into operation.

IN OPERATION / AFTER

- Audit checklists and warning logic for cues whose meaning depends on phase or context, and resolve the ambiguity before an operator must guess.
- Monitor maintenance-to-operations handoffs for latent configuration errors that pre-flight checks have historically missed.
- Sustain training on the physiology that degrades the operator (e.g., hypoxia), so the window for self-correction is understood and guarded.

REFLECTION QUESTIONS

1. Identify a cue in your domain that carries two meanings in different operational contexts. How would the operator know which?
2. Redesign the Helios configuration warning to distinguish ground- and in-flight meaning without adding cognitive load.
3. A maintenance setting left in “manual” survived the pre-flight checks and reached the air. Design the pre-flight verification step in your domain that would make an abnormal maintenance configuration impossible to depart with unnoticed.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · usability testing · cognitive walkthroughs

TransAsia Airways Flight 235

T Training Gap **H** Human-System Interface

IMPACT 43 killed in Taiwan; crew shut down the only working engine after the other failed

IN BRIEF

About thirty-seven seconds after takeoff from Taipei Songshan in February 2015, the right engine of a TransAsia ATR 72 auto-feathered following a sensor fault. Working from memory under acute time pressure, the crew shut down the left engine — the one still producing thrust — leaving the aircraft with no power and too little altitude to recover. It clipped a viaduct, struck a taxi, and crashed into the Keelung River, killing 43 of the 58 aboard. The Aviation Safety Council found the captain had failed proficiency checks in the year before the accident and that the crew had not used the published engine-shutdown checklist, responding instead from a flawed memory of the procedure. The wrong-engine error mirrors Kegworth (Case 103) twenty-six years earlier.

BACKGROUND

TransAsia 235 was an ATR 72-600 twin-turboprop departing Taipei Songshan. The captain had failed simulator proficiency checks in the year before the accident and had been noted for weak handling of abnormal procedures. The published response to an engine failure was a checklist-driven procedure; under stress, crews tend to revert to memory — so a captain already flagged for weak abnormal-procedure handling was exactly the operator most likely to skip the checklist precisely when it mattered, a known weakness meeting a known failure mode.⁰

WHAT HAPPENED

On 4 February 2015, about thirty-seven seconds after takeoff, the right engine's autofeather system activated following a faulty sensor signal — the engine itself was capable of producing power. The crew, identifying an engine problem, shut down the left engine, which was operating normally. With neither engine now producing useful thrust and the aircraft low and slow, it stalled, clipped a viaduct and a taxi, and crashed into the Keelung River, killing 43 of the 58 aboard. A recoverable single-engine event became unrecoverable the moment the working engine was shut down, leaving an aircraft with no thrust and no altitude in the seconds just after takeoff when both are least forgiving.¹

THE INVESTIGATION

The Taiwan Aviation Safety Council found the crew failed to identify which engine had actually failed and shut down the wrong one. They did not follow the published engine-flameout and shutdown checklist, acting from memory under time pressure; the misidentification


took roughly fifteen seconds. The captain's documented proficiency deficiencies and the airline's training and crew-resource-management shortfalls were contributing factors — the individual lapse sitting atop an organizational one, the carrier having left the very weaknesses its own proficiency checks had recorded uncorrected before the flight.²

THE CAPABILITY GAP

TransAsia 235 is the modern recurrence of the wrong-engine pattern that destroyed British Midland 92 at Kegworth in 1989 (Case 103). Crews under acute time pressure default to memory; if the memory is incomplete, the action follows the flawed memory rather than the checklist. The missing capability is a procedure design and training regime that keeps the verification step in the loop precisely when stress pushes operators to skip it — engineering the confirm-which-engine step so it fires under startle, rather than trusting it to a memory that time pressure is actively eroding.³

AFTERMATH & REFORM

The ASC issued recommendations on TransAsia's pilot training, proficiency-check standards, and crew-resource management; the airline ceased operations in 2016, after a second fatal accident the year before. The case is used to argue that the persistence of the wrong-engine pattern across a quarter-century reflects an un-engineered intervention, not merely individual error — the same failure recurring across crews, airframes, and decades points to a design and training gap that no amount of blaming the individual pilot has ever closed.⁴



The flight crew was unable to identify which engine had failed and shut down the operating engine in error.

PARAPHRASING THE AVIATION SAFETY COUNCIL (TAIWAN) FINAL REPORT ON TRANSASIA 235, 2016

REFERENCES

1. Aviation Safety Council (Taiwan), Aviation Occurrence Report: TransAsia Airways Flight GE235, final report (2016) — findings and the wrong-engine shutdown (paraphrased).
2. ASC final report (2016) — the autofeather event and the crash sequence.
3. ASC final report (2016) — non-use of the checklist and the captain's proficiency record.
4. AAIB (UK), British Midland Flight 92 (Kegworth) report (1990) — the antecedent wrong-engine case.
5. Dismukes, Berman & Loukopoulos (2007), The Limits of Expertise — memory-driven response under time pressure.

THE LEARNING ENGINEERING LENS

LE INSIGHT

TransAsia 235 is the modern recurrence of the Kegworth pattern under stress. Crews under acute time pressure default to memory; if the memory is incomplete or wrong, the action follows the memory. The capability gap was the same as in 1989. The intervention pattern needed is also the same.

LENS APPROACH

LENS uses TransAsia 235 in LEN 5 as a recurrence case for well-known capability problems (wrong-engine shutdown) and in LEN 2 to teach the difference between checklist-driven response and memory-driven response under stress.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ T H	→ D1 D2 D3 D4 D5	3	3.1	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the abnormal-procedure response so the confirm-which-unit verification step is forced under startle, not left to memory.
- ▶ Build proficiency checks whose red signals have the authority to ground an operator before the predicted failure, not merely to document a weakness.
- ▶ Train the checklist-driven response under realistic time pressure, so the trained reflex survives the stress that pushes crews to memory.

IN OPERATION / AFTER

- ▶ Audit whether documented proficiency deficiencies are actually acted on, treating an uncorrected red check as an organizational, not individual, failure.
- ▶ Monitor for recurrence of well-known patterns (wrong-engine shutdown) across crews and airframes as evidence the intervention is un-engineered, not merely human error.
- ▶ Sustain crew-resource-management and checklist discipline so the verification step does not erode back to memory under operational pressure.

REFLECTION QUESTIONS

1. Identify a procedure in your domain that is supposed to be checklist-driven but is actually memory-driven under stress. What is the gap?
2. Why has the wrong-engine shutdown pattern persisted across 25 years of aviation reform? What intervention has not yet been engineered?
3. The captain had failed proficiency checks in the year before the accident, yet kept flying. Design the proficiency-check regime in your domain whose red signal would actually remove an operator from duty before, not after, the failure it predicts.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · usability testing · cognitive walkthroughs

Eastern Air Lines Flight 401

H Human–System Interface **T** Training Gap

IMPACT 101 killed in the Everglades; the entire flight crew fixated on a landing–gear indicator bulb while the autopilot silently disengaged

IN BRIEF

On 29 December 1972 an Eastern Air Lines L-1011 descended into the Florida Everglades on a night approach to Miami, killing 101 of the 176 aboard. The cause was attention, not mechanics: all three flight-deck crew became absorbed in a landing-gear indicator bulb that had failed to light, and while they worked it, the autopilot's altitude hold was inadvertently disengaged. The aircraft sank slowly; an altitude alert chimed but went unheeded in the cockpit clutter, and no one was monitoring the flight path. The NTSB findings launched decades of research into attention, monitoring, and cockpit-alert design, and the accident is a formative event behind Crew Resource Management. Eastern 401 is the canonical case of a low-priority task crowding out a life-critical one.

BACKGROUND

Eastern Air Lines Flight 401, a wide-body Lockheed L-1011, was on a night approach to Miami on 29 December 1972 with a full flight deck — captain, first officer, and flight engineer.⁰ Three crew was the era's safeguard against any one person being overloaded, the assumption being that more eyes meant more coverage; the night would show that without a designed division of those eyes, three people could converge on a single trivial problem as easily as one.

WHAT HAPPENED

The crew noticed that the landing-gear indicator light had not illuminated, and all three focused on the bulb. While they worked the problem, one of them inadvertently nudged the controls and disengaged the autopilot's altitude hold. The TriStar began a slow descent into the Everglades; an altitude-warning chime sounded but, in the auditory clutter of the cockpit, no one registered it. The aircraft hit the swamp, killing 101 of the 176 aboard.¹ The descent was gentle enough to go unfelt, and the single chime carried no more urgency than the routine sounds around it — so the one cue that could have broken the fixation was indistinguishable, by design, from the background it sounded against.

THE INVESTIGATION

The NTSB found the crew had become so engrossed in the landing-gear difficulty that they failed to monitor the flight path.² A burned-out indicator bulb had absorbed the attention of an entire qualified crew while a wide-body airliner descended unwatched — the disproportion between the trigger and the outcome is exactly what made the case so durable a


teaching example. The findings inaugurated decades of research into attention, monitoring, and the design of cockpit alerts — the accident is cited in nearly every introductory cognitive-engineering course as the example of how a low-priority task can crowd out a life-critical one when attention is undivided among the available channels.³

THE CAPABILITY GAP

The missing capability was a designed division of attention across the crew, and an alert that carried the priority its message deserved. Three competent people attended to a burned-out bulb while no one watched the altitude; the chime existed but did not cut through. The flight deck had the people and the information — it lacked the design that would have kept one channel of attention on the thing that could kill them. Attention had been treated as something the crew would naturally allocate well, rather than as a resource the cockpit had to be built to protect, and that unexamined assumption was the gap.⁴

AFTERMATH & REFORM

Eastern 401 was one of the formative events behind Crew Resource Management (Case 117) — the explicit allocation of monitoring and cross-checking roles — and behind modern standards for prioritized cockpit alerting.⁵ Both reforms encode the same correction: CRM assigns someone to keep flying the aircraft while others troubleshoot, and prioritized alerting ensures the most dangerous condition is the loudest one — turning the night's two failures into permanent structural defenses. Its lesson is that attention is a designable parameter, not a personal virtue: a system that lets all eyes converge on one task must guarantee the critical channel is still watched.



The crew became so engrossed in the landing-gear difficulty that they failed to maintain altitude.

PARAPHRASING NTSB AIRCRAFT ACCIDENT REPORT NTSB-AAR-73-14 ON EASTERN 401, 1973

REFERENCES

1. NTSB, Aircraft Accident Report: Eastern Air Lines Flight 401, NTSB-AAR-73-14 (1973) — the night approach and full flight deck.
2. NTSB-AAR-73-14 (1973) — the landing-gear-bulb fixation, the inadvertent autopilot disengagement, and 101 deaths of 176 aboard.
3. C. D. Wickens et al., Engineering Psychology and Human Performance — attention, monitoring, and alert design.
4. R. Helmreich & H. Foushee (1993), aircrew-coordination research — the Crew Resource Management origins.
5. Eastern 401 as a formative case behind CRM (Case 117) and prioritized cockpit alerting standards.

3. NTSB-AAR-73-14 (1973) — the crew’s preoccupation with the landing-gear indication that distracted them from maintaining altitude (paraphrased).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Eastern 401 is the canonical attention-failure case. The capability that was missing was a designed division of attention across the crew. CRM exists because Eastern 401 happened, and the discipline of cockpit alert design exists because the altitude warning chime did not carry the priority weight it needed to.

LENS APPROACH

LENS uses Eastern 401 in LEN 5 to teach attention as a designable parameter and in LEN 2 to introduce alert prioritization as a capability deliverable. The case anchors the CRM origin story.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ H T	→ D1 D2 D3 D4 D5	3	3.1	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Build an explicit division of attention into the operating procedure, so at least one role is always assigned to the life-critical channel while others troubleshoot. ▶ Design alerts with a priority hierarchy, ensuring the most dangerous condition is also the most salient rather than one indistinguishable chime among many. ▶ Treat attention as a resource the system must protect by design, not a virtue the operators are assumed to supply on their own.	▶ Drill crews on scenarios where a trivial problem competes with a critical one, and verify the monitoring role actually holds under that pressure. ▶ Audit whether alerts in service are heard and acted on, retiring or redesigning any that get lost in routine clutter. ▶ Track distraction-related near-misses so attention-displacement failures are visible before they cause a loss.

REFLECTION QUESTIONS

1. Identify a low-priority task in your domain that could plausibly absorb all of an operator’s attention. What is the life-critical task that would be displaced?
2. Redesign the altitude warning chime of an L-1011 so that it cuts through a focused troubleshooting conversation.
3. Eastern 401 had three qualified people and still left the flight path unwatched. How would you assign and verify a “someone is always watching the critical channel” role so it cannot collapse when the whole team is drawn to one problem?

WHO BUILDS THIS human factors & interaction design · learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · scenario simulation · competency models

NASA EVA-23 — Water Intrusion Inside a Spacesuit

H Human-System Interface **N** Normalization of Deviance **D** Designed Out

IMPACT During the second of two ISS spacewalks, ESA astronaut Luca Parmitano's helmet filled with up to 1.5 liters of water from the suit's cooling-and-ventilation loop, blinding him, fouling his communications, and threatening drowning in vacuum; the EVA was terminated and Parmitano recovered to the airlock guided by tether feel — the NASA Mishap Investigation Board called the outcome a near-fatality

IN BRIEF

On July 16, 2013, during EVA-23 outside the International Space Station, water began collecting inside the helmet of European Space Agency astronaut Luca Parmitano. The contamination grew until the water covered his eyes, ears, and nose, fouled his communications microphone, and would have drowned him in vacuum had the spacewalk continued. The crew terminated the EVA and Parmitano made his way back to the airlock partly by feel along the tether. The NASA Mishap Investigation Board (MIB) report issued later that year identified the proximate cause as blocked separator holes in the suit's Fan/Pump/Separator assembly that allowed cooling-loop water into the ventilation loop, and the deeper cause as a missed opportunity at the previous EVA-22 a week earlier when a similar (but smaller) water event had been mis-attributed to in-suit drink-bag leakage. The case is the canonical instance in modern human spaceflight of a failure mode the suit was not instrumented to detect, surfacing first as a near-fatal event because the earlier weak signal was rationalized away. The MIB's recommendations span hardware redesign, in-suit water-detection instrumentation, and the safety-culture response to anomalous EVA telemetry.

BACKGROUND

EVA-23 was the second of two scheduled spacewalks for ISS Expedition 36. Parmitano and NASA astronaut Chris Cassidy exited the Quest airlock to continue tasks begun during EVA-22 a week earlier. Roughly 44 minutes into the EVA, Parmitano reported water at the back of his head inside the helmet. The crew continued briefly while ground controllers assessed; the water continued accumulating, eventually reaching the front of the helmet where it pooled across Parmitano's eyes, blocked his nose, and fouled the communications microphone. The EVA was terminated and Parmitano made the return traverse to the airlock with degraded vision and communication, partly guided by his sense of the tether.⁹

WHAT HAPPENED

The proximate hardware fault, identified by the NASA Mishap Investigation Board, was in the suit's Fan/Pump/Separator assembly — the component that separates cooling water from the breathing-gas ventilation loop. Contamination had blocked drum-holes in the water separator. Once the separator could not perform its function, water from the cooling loop migrated into the ventilation loop and was pushed up into the helmet by the suit's fan. The hardware had no in-suit instrumentation to detect free water in the ventilation path; the crew detected the failure only when Parmitano felt liquid on his skin.¹

THE INVESTIGATION

The MIB's deeper finding was about EVA-22, the prior spacewalk one week earlier. Parmitano had reported a smaller volume of water in his helmet during EVA-22 as well. At the time, the event was attributed to a drink-bag leak — a plausible but unverified hypothesis. No suit teardown or anomaly investigation was completed before EVA-23 was approved. The MIB characterized this as a missed opportunity: the EVA-22 event was the weak signal of the same failure mode that nearly killed Parmitano a week later, and the institutional response treated an unverified benign explanation as adequate clearance to proceed.²

THE CAPABILITY GAP

The corrective actions span three layers. At the hardware layer the suit's water-separator and ventilation-loop hygiene were re-engineered, and a Helmet Absorption Pad and a snorkel were added so a future water event would not immediately threaten the airway. At the instrumentation layer in-suit water-detection capability was added so the failure mode is no longer detectable only by skin contact. At the safety-culture layer the MIB pressed the "anomaly is anomaly" discipline: an unexplained EVA event requires investigation rather than the most convenient explanation, particularly when the next planned EVA reuses the same hardware.³

AFTERMATH & REFORM

The hedge that survives into the case is about attribution depth. The MIB was clear that the in-suit water-detection and hardware fixes address the specific failure mode that produced EVA-23. The harder, lifecycle question — how a suit certified for decades of use developed a contamination pathway not represented in its anomaly catalog, and how future suits will avoid the analogous category of uninstrumented failure — is left as an open sustainment-engineering problem. The case is the canonical modern human-spaceflight instance of a recovery from a near-fatal event whose proximate cause was hardware and whose deeper cause was the institutional decision to accept the most convenient explanation for an anomalous prior event.

The previous occurrence is the strongest available signal of the next one; the institutional decision to accept the convenient explanation is what made the second event near-fatal.

EDITORS' SYNTHESIS OF THE NASA MIB EVA-23 REPORT.

REFERENCES

1. NASA International Space Station Program (2013), "International Space Station EVA Suit Water Intrusion High Visibility
3. ESA (2013), debrief and operational statement on the Parmitano EVA-23 event.

Close Call” — Mishap Investigation Board final report, December 2013.

2. NASA Aerospace Safety Advisory Panel (2014), Annual Report — EVA-23 follow-up actions including helmet absorption pad and snorkel.

4. Chappell, Norcross, Abercromby et al. (2017), “Risk of Injury and Compromised Performance Due to EVA Operations,” NASA Human Research Program Evidence Report — broader EVA risk context.

THE LEARNING ENGINEERING LENS

LE INSIGHT

EVA-23 is the canonical modern human-spaceflight case of a failure mode the suit was not instrumented to detect, surfacing first as a near-fatality because the prior weak signal was rationalized away. The hardware fix and the new in-suit instrumentation are necessary; the safety-culture half — anomaly is anomaly — is what generalizes.

LENS APPROACH

EVA-23 is the human-spaceflight uninstrumented-failure case (induced 3.4; LENS D3/PT5). LENS uses it in Domain 3 (Human-System Collaboration) for the operator’s detection-by-skin-contact at the edge of the engineered envelope, and in Domain 4 (Test and Evaluation) for the anomaly-investigation discipline the MIB pressed. Pair with Case 27 (anesthesia monitoring) at the cue/alert-as- capability layer and with Case 132 (F-22 OBOGS) as the aerospace companion in instrumentation-gap failure.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
SPACE / AEROSPACE	→ HND →	D1D2D3D4D5	5	3.4	4 · 3
HUMAN SPACEFLIGHT		Human-System Collaboration			
SAFETY ENGINEERING					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Instrument the failure modes the system was not designed for. EVA-23 was detectable only by skin contact because the ventilation loop had no free-water sensor; instrumentation is the precondition for recovery.
- ▶ Treat the previous occurrence of an anomaly as the strongest available signal of the next one. EVA-22’s smaller water event was the warning EVA-23 ignored; the institutional decision to accept the convenient explanation is what made the second event near-fatal.
- ▶ Engineer the recovery margin into the suit itself — the post-EVA-23 Helmet Absorption Pad and snorkel — so a contamination event has a survivable buffer before it becomes an airway emergency.

IN OPERATION / AFTER

- ▶ Require an anomaly investigation, not a hypothesis, before reusing the same hardware for a subsequent high-consequence operation; the cost of a teardown is small against the cost of the failure mode.
- ▶ Build the in-suit instrumentation expectation into the certification process for future suits so the lesson generalizes beyond the specific Fan/Pump/Separator failure.
- ▶ Carry the sustainment-engineering hedge into program communication: hardware certified for decades can still develop pathways not represented in its anomaly catalog, and the certification process has to keep pace.

REFLECTION QUESTIONS

1. Identify an operational system in your domain where a failure mode is detectable only by direct human sensing (skin contact, smell, the operator noticing). What instrumentation would convert that failure mode from sensed-late to sensed-early?
2. Specify the anomaly-investigation threshold you would set in your program: an unexplained prior event of the same class requires what level of teardown before the next operation is approved? EVA-22 to EVA-23 is the worked example of the cost of getting that threshold wrong.
3. The post-EVA-23 fixes addressed the specific failure mode. What is the lifecycle question your program owes for decades-certified hardware: how does the certification process keep pace with contamination pathways the original anomaly catalog did not represent?

WHO BUILDS THIS human factors & interaction design · safety science & organizational analysis · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · incident analysis · safety dashboards · task analysis · design prototyping

Colgan Air Flight 3407

T Training Gap

IMPACT 50 killed near Buffalo; precipitated the FAA's 1,500-hour rule and the Pilot Records Database

IN BRIEF

On approach to Buffalo in February 2009, the captain of Colgan Air 3407 responded to a stall warning by pulling back on the controls instead of lowering the nose; the Bombardier Q400 stalled and crashed into a house in Clarence Center, killing all 49 aboard and one person on the ground. The NTSB found the captain had a documented history of training failures, and the first officer was ill, fatigued, and paid about \$16,000 a year. The airline and the hiring pipeline knew the training history; the regulator did not. Victims' families mounted one of the most effective aviation-safety campaigns in a generation, producing the 2010 law that raised first-officer experience requirements to 1,500 hours and created the Pilot Records Database. The capability gap was visible to everyone except the system that licensed the pilots.

BACKGROUND

Colgan Air 3407, a Bombardier Q400 flying a regional route to Buffalo, was crewed by a captain with a documented history of training failures and a first officer who was ill, fatigued, and paid roughly \$16,000 a year. The airline and the hiring pipeline knew the captain's record; the regulator did not — the information that should have shaped who sat in that seat lived inside the carrier's files and never crossed into the licensing system that was supposed to vouch for the crew.⁰

WHAT HAPPENED

On 12 February 2009, on approach to Buffalo, the aircraft slowed and its stall-warning stick shaker activated. The captain responded by pulling back on the control column — the opposite of stall recovery — and the Q400 stalled and crashed into a house in Clarence Center, killing all 49 aboard and one person on the ground. The reflex he reached for was exactly wrong for the situation, the kind of error a pattern of failed training events should have predicted long before he was handed the controls.¹

THE INVESTIGATION

The NTSB found the probable cause to be the captain's inappropriate response to the stick shaker, against a backdrop of fatigue, weak airline training, and a hiring system blind to the captain's history. The data that should have flagged him existed — multiple failed training events — but did not flow to the decision that put him in the seat, so the hiring choice was made on a record that omitted the very facts that mattered most, an absence that looked like a clean slate.²

THE CAPABILITY GAP

Colgan is the canonical case for the gap between the information an institution holds about its operators and the information that reaches the decisions about those operators. The capability that was missing was not pilot skill in the abstract but a data flow: a way to make a pilot's documented training history actionable at the hiring and licensing decision, rather than buried in files no one consulted — a record can exist in full and still fail completely if it never reaches the moment and the person whose choice it was meant to inform.³

AFTERMATH & REFORM

The victims' families organized one of the most effective aviation-safety lobbying efforts in a generation, producing the Airline Safety and FAA Extension Act of 2010 — which raised the minimum experience for airline first officers to 1,500 hours and created the Pilot Records Database to make pilot history flow between carriers.⁴ The 1,500-hour rule remains debated; the records database fundamentally restructured how the information that was missing in 2009 now moves, converting a buried file into a record a hiring carrier is required to retrieve — the difference between a fact that exists and one that arrives where the decision is made.

50

KILLED, CLARENCE CENTER NY

the data was in the system; the data flow was not

COLGAN 3407 — AND THE REFORM THAT FOLLOWED

The captain's inappropriate response to the activation of the stick shaker, which led to an aerodynamic stall from which the airplane did not recover.

NTSB AIRCRAFT ACCIDENT REPORT AAR-10/01, PROBABLE CAUSE, 2010

REFERENCES

1. NTSB, Aircraft Accident Report: Colgan Air Flight 3407, NTSB/AAR-10/01 (2010) — probable cause (quoted).
2. NTSB/AAR-10/01 (2010) — the captain's training history, and the first officer's fatigue and pay.
3. NTSB/AAR-10/01 (2010) — the information-flow gap between known pilot history and the hiring decision.
4. Airline Safety and Federal Aviation Administration Extension Act of 2010, Pub. L. 111-216 — the 1,500-hour rule and the Pilot Records Database.
5. GAO-12-203, Pilot Records Database Implementation (2012); the Families of Continental Flight 3407 campaign.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Colgan is the canonical case for the gap between the information an institution holds about its operators and the information that reaches the decisions about whether those operators should be in that seat. The data was already in the system — multiple failed training events. The data flow that would have made that data actionable at the hiring decision did not exist.

LENS APPROACH

LENS uses Colgan in LEN 4 as a case for evidence–system design (the PRD as a designed information flow) and in LEN 8 as a case for advocacy–driven institutional change. Studio projects examine how families of victims became the load-bearing element of a reform the industry had resisted for years.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ T	→ D1 D2 D3 D4 D5	2	2.1	4
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Make a pilot’s documented training history a required, structured input to the hiring and licensing decision, not a record held passively in carrier files. ▶ Design the data flow so that adverse signals (failed training events) are pushed to the decision point rather than waiting to be requested. ▶ Define record retention and transfer obligations across carriers up front, so a clean-looking application cannot conceal a known history.	▶ Audit whether the records actually reach hiring decisions, treating a complete file that never arrives as a live failure, not a solved one. ▶ Monitor the pipeline for operators whose adverse history is not surfacing, and act before the gap is paid for in an accident. ▶ Sustain the records database against drift and gaps in coverage, since a data-flow remedy decays as fast as its weakest link.

REFLECTION QUESTIONS

1. What information about operators in your domain exists somewhere in the system but does not flow to the decisions that depend on it?
2. Design the data-flow architecture that would make a Colgan–equivalent visible **before** the accident rather than after.
3. The reform in 2009 was driven by victims’ families, not the industry that had resisted it. What load-bearing constituency in your domain could force a stalled data-flow fix into existence, and what evidence would mobilize them?

WHO BUILDS THIS learning science & instructional design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models

Atlas Air Flight 3591

T Training Gap **K** Knowledge & Institutional Memory

IMPACT 3 killed; Boeing 767 cargo flight crashed in Texas after the first officer mistakenly activated go-around mode

IN BRIEF

On approach to Houston in February 2019, the first officer of Atlas Air 3591 — an Amazon Prime Air cargo flight — inadvertently triggered the go-around mode during a turbulent descent. Experiencing a somatogravic illusion that the aircraft was pitching up, he pushed the nose down sharply and dove the Boeing 767 freighter into Trinity Bay, killing all three aboard. The NTSB found the first officer had a long, undisclosed history of training failures across multiple carriers — and that Atlas could not see it, because the Pilot Records Database mandated by Congress after Colgan Air 3407 (Case 109) was not yet operational. Atlas relied on the older PRIA system, which surfaced only five years of history. The case is the live recurrence of Colgan: the information existed; the data-flow system was not yet complete.

BACKGROUND

Atlas Air 3591 was a Boeing 767 freighter operating for Amazon Prime Air from Miami to Houston. The first officer had failed training events at several previous employers and had not disclosed them on his Atlas application. At the time of hiring, the Pilot Records Database directed by Congress after Colgan Air 3407 was not yet operational; Atlas used the older PRIA system, which required only five years of records — a window too short to surface a pattern of training failures spread across several prior employers, so the record Atlas could legally obtain was structurally incapable of showing what it most needed to see.⁹

WHAT HAPPENED

On 23 February 2019, during a turbulent descent toward Houston, the first officer inadvertently activated the aircraft's go-around mode, which commanded a pitch-up. Experiencing a somatogravic illusion that the nose was rising into a stall, he pushed forward hard. The 767 entered a steep dive from which the crew did not recover and crashed into Trinity Bay, killing all three aboard. The startle and spatial disorientation that drove the fatal push were the same kind of breakdown his prior training failures had repeatedly recorded, had that history been visible to the carrier that hired him.¹

THE INVESTIGATION

The NTSB found the probable cause to be the first officer's inappropriate response to the inadvertent go-around activation, driven by a startle and spatial-disorientation response, and identified his pattern of training deficiencies and Atlas's inability to access his full record as contributing factors. The records that would have informed the hiring decision existed but


did not reach it — the same information-flow failure named after Colgan a decade earlier, recurring here because the remedy built to fix it was authorized but not yet carrying the full history.²

THE CAPABILITY GAP

Atlas Air 3591 is the iteration test of the Colgan reform. The reform partially worked — the Pilot Records Database had been authorized — but it was not yet operational and its coverage was incomplete, so the same information-flow gap that killed fifty people in 2009 was still partly open in 2019. Partial implementation of a remedy leaves a measurable harm inside the aperture — the years between authorization and full coverage are not a neutral transition but a window in which the original failure can recur on the cases the incomplete remedy does not yet reach.³

AFTERMATH & REFORM

The PRD final rule was published in 2021 and became effective for Part 121 carriers in 2022, with full historical coverage phasing in through 2024 — closing the gap the case exposed. Atlas 3591 is cited in implementation-fidelity discussions as evidence that a remedy is only as good as its completed coverage — that the date a rule is authorized and the date it actually reaches every carrier are different dates, and the distance between them is where the avoidable harm accumulates.⁴



Atlas Air did not have access to portions of the first officer's training record that would have informed its hiring decision.

PARAPHRASING NTSB AIRCRAFT ACCIDENT REPORT AAR-20/02, 2020

REFERENCES

1. NTSB, Aircraft Accident Report: Atlas Air Flight 3591, NTSB/AAR-20/02 (2020) — probable cause and contributing factors (paraphrased).
2. NTSB/AAR-20/02 (2020) — the inadvertent go-around activation and the dive into Trinity Bay.
3. NTSB/AAR-20/02 (2020) — the first officer's training history and Atlas's record-access limits.
4. FAA Pilot Records Database final rule (Federal Register, 10 June 2021; subpart B/C compliance from June 2022, full historical coverage by September 2024); Pilot Records Improvement Act of 1996.
5. GAO reports on Pilot Records Database implementation (2014–2024).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Atlas Air 3591 is the iteration test of the Colgan reform. The reform partially worked: the PRD exists. The reform was incomplete: not all carriers were covered. The case is a cautionary worked example of why partial implementation of a regulatory remedy leaves the original capability gap partially open.

LENS APPROACH

LENS uses Atlas Air 3591 in LEN 4 as a case for measurement–system coverage (the PRD had been built but its mandatory coverage was incomplete), and in LEN 8 for the iteration cycle of reform: an intervention that leaves an aperture creates a measurable harm inside the aperture.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ T K	→ D1 D2 D3 D4 D5	2	2.4	4
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Design the records system to require a full operator history, not a fixed-length window that can hide a multi-employer pattern. ▶ Plan the coverage phase-in explicitly, with interim controls protecting the cases the remedy does not yet reach. ▶ Specify the data the hiring decision needs and make its delivery mandatory before the decision, not optional or self-reported.	▶ Audit coverage completeness, not just authorization, and treat the gap between the two as an active source of harm until closed. ▶ Monitor for adverse events occurring inside the aperture of a partially implemented remedy, and use them to accelerate full coverage. ▶ Sustain attention through the full phase-in, since a remedy authorized but not yet universal leaves the original gap measurably open.

REFLECTION QUESTIONS

1. Identify a regulatory remedy in your domain whose coverage is partial. What harm is occurring inside the aperture?
2. Design the deliverable that would close the PRD coverage gap in advance of a future Atlas Air 3591.
3. The PRD was authorized in 2010 but not fully effective until 2024, and Atlas 3591 fell inside that gap. For a remedy in your domain, design the interim control that would protect the cases the remedy does not yet reach while its coverage phases in.

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · knowledge bases · data pipelines

Challenger & Columbia

N Normalization of Deviance **K** Knowledge & Institutional Memory **G** Governance & Trust

IMPACT 14 astronauts killed (7 per accident); 17 years between identical organizational pathologies

IN BRIEF

NASA lost two Space Shuttle crews to the same organizational pathology seventeen years apart: Challenger in 1986, when O-ring seals failed in cold weather, and Columbia in 2003, when foam debris breached the wing's thermal protection — fourteen astronauts in all. Both flaws had been seen repeatedly and accepted as routine; foam shedding had been documented on at least sixteen prior missions before Columbia, and O-ring erosion in the field joints had been on the engineering record since the early flights. Sociologist Diane Vaughan named the mechanism “normalization of deviance” from the Challenger investigation; the Columbia Accident Investigation Board found the same culture intact seventeen years later and concluded NASA's organizational culture had as much to do with the accident as the foam. The pair is the strongest single-institution evidence that culture is an engineerable property of a system — and that a pathology diagnosed but left unrepaired will recur, at the same cost.

BACKGROUND

The Space Shuttle flew with two flaws its engineers knew about. The solid-rocket booster joints sealed with O-rings that stiffened in cold; the external tank shed foam insulation that struck the orbiter on ascent. Both had appeared on flight after flight without catastrophe, and both were progressively reclassified from hazards to accepted, “routine” features of flying.⁰ Each survived flight became evidence that the flaw was tolerable, so the very absence of disaster fed the reclassification — the safety margin redefined downward not by decision but by the accumulating habit of getting away with it.

WHAT HAPPENED

On 28 January 1986 Challenger launched on an unusually cold morning; an O-ring failed to seal, and the vehicle broke apart 73 seconds after liftoff, killing all seven aboard. Seventeen years later, on 1 February 2003, foam that had struck Columbia's wing on ascent had opened a breach in its thermal protection; the orbiter disintegrated on reentry, killing its seven.¹ Both flaws were the ones engineers had already flagged and the institution had already filed as routine, so each accident was less a surprise than the arrival of a bill the system had decided, repeatedly, not to pay.

THE INVESTIGATION

The Rogers Commission traced Challenger to the O-ring and to a launch decision that overrode engineers' cold-weather warnings — the Thiokol teleconference on the eve of


launch, at which the engineering position (“do not launch below 53°F”) was reversed under management pressure and the launch went ahead at 36°F.² Diane Vaughan’s *The Challenger Launch Decision* (1996) re-examined the same teleconference with a decade’s documentary access and reframed the decision: not deviant individuals overriding good engineering, but a working group whose rules of evidence had drifted, one accepted anomaly at a time, until a launch at 36°F was inside what the group’s own decision rules counted as acceptable risk. Seventeen years later the Columbia Accident Investigation Board found the same patterns intact — flying with flaws defined as routine, foam strikes filed under “in-family” across at least sixteen prior missions, and a structure that suppressed the upward flow of safety concerns — and concluded that “the NASA organizational culture had as much to do with this accident as the foam.”³ That the same board found the same structure seventeen years on is the case’s sharpest point: the first reform had fixed the hardware and the schedule pressure but not the mechanism by which a flaw becomes redefined as acceptable.

THE CAPABILITY GAP

Diane Vaughan’s “normalization of deviance,” developed from Challenger, named the mechanism: deviations from the safety baseline become acceptable through production pressure, weak communication, and habit, one small step at a time. Columbia validated the concept against its author’s intent — Vaughan herself was retained by the CAIB and the Board’s final report cites the mechanism by name. The pathology had been diagnosed in 1986 and never engineered away — which is the point: culture is a system property, and a diagnosis without a remediation decays.⁴ Naming the mechanism did not arrest it, because a name lives in a report while the production pressure and the suppressed warnings live in the daily flow of work, where they kept doing what they had always done. The institutional-learning gap is itself the load-bearing finding: between 1986 and 2003 NASA had reorganized twice, lost an administrator, and adopted multiple safety initiatives, and the structural pathway from a dissenting engineer to a launch decision was substantially unchanged.

AFTERMATH & REFORM

Each accident produced reform — the Rogers Commission’s redesign of the booster joint, the CAIB’s call to treat culture as a safety variable, the establishment of the NASA Engineering and Safety Center, and the Board’s instruction to rebuild independent technical authority — and the Shuttle was retired in 2011 after completing the International Space Station assembly. The pair stands as the book’s strongest evidence that organizational culture is engineerable, and that leaving it unengineered is a choice with a recurring, lethal cost.⁵ The CAIB’s insistence on independent technical authority conceded the deeper lesson the booster-joint redesign alone had missed in 1986: the upward path for a dissenting engineer is itself a piece of safety hardware, and one that has to be rebuilt and defended rather than assumed. The hedge that survives into the case: Vaughan’s “normalization of deviance” is the load-bearing analytic claim across both accidents and is never to be smoothed; the seventeen-year institutional-learning gap is the empirical claim, and the CAIB’s cross-referencing of Rogers Commission language (“the same decision-making structures”) is the documentary anchor.



These repeating patterns mean that flawed practices embedded in NASA's organizational system continued for 20 years and made substantial contributions to both accidents.

COLUMBIA ACCIDENT INVESTIGATION BOARD, 2003

REFERENCES

1. Rogers Commission, Report of the Presidential Commission on the Space Shuttle Challenger Accident (1986) — the O-ring failure, the Thiokol teleconference, and the cold-weather launch decision.
2. D. Vaughan, *The Challenger Launch Decision: Risky Technology, Culture, and Deviance at NASA* (Univ. of Chicago Press, 1996) — normalization of deviance; the working group's drift of decision rules.
3. Columbia Accident Investigation Board, Report Vol. I (2003) — the foam strike, the sixteen prior missions of foam-shedding filed as "in-family," the recurrence of the cultural pattern, and the call for an independent technical authority.
4. CAIB (2003) — "the NASA organizational culture had as much to do with this accident as the foam" (quoted); Rogers Commission (1986) and CAIB (2003) jointly on the suppressed upward flow of safety information across both accidents.
5. W. Starbuck & M. Farjoun (eds.), *Organization at the Limit: Lessons from the Columbia Disaster* (2005) — independent academic re-analyses of the institutional-learning gap.
6. NASA Engineering and Safety Center founding documents (2003 – present) — institutional response to the CAIB's call for independent technical authority; Shuttle retirement (STS-135, July 2011).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Challenger/Columbia is the strongest documentary evidence that organizational culture is an engineerable property of a system, and that diagnoses without engineered remediations decay. The same pathology, twice, seventeen years apart, in the same institution, with the same human cost — and with the diagnosis already on the record. Vaughan's "normalization of deviance" names the mechanism and is the load-bearing analytic claim across both accidents. Capability engineering treats culture as a deliverable.

LENS APPROACH

Challenger/Columbia anchors the multi-layer-drift-and-cumulative- inadequacy competency (induced 7.4; LENS D1/PT4): a cascade of marginal-but-tolerable conditions across decision rules, schedule pressure, and communication structure aligned twice, seventeen years apart. LENS uses the pair in Domain 1 (Systems Analysis) for normalization of deviance as a systems concept and the institutional-learning gap as a measurable property; in Domain 5 (Navigating Sociotechnical Constraints) for the governance failure that allowed a diagnosed pathology to persist and for the upward-channel design the CAIB called a piece of safety hardware. Pair with Deepwater Horizon (Case 163) on the multi-layer-drift form, and with Bhopal (Case 159) and Fukushima (Case 169) on cumulative inadequacy in catastrophic-system operations.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
SPACE / AEROSPACE	NKG	D1D2D3D4D5	4	7.4	1·5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
- Hold known flaws as open hazards with explicit owners, so a clean flight cannot silently reclassify a deviation as routine. - Build independent technical authority into the launch decision, giving dissenting engineers a path that production pressure cannot override. - Define the safety baseline quantitatively at design time, so any drift below it is a flagged change rather than an unremarked habit.	- Re-audit every diagnosed-but-unrepaired pathology on a fixed cadence, verifying remediation in practice rather than on paper. - Track the upward flow of safety concerns as a measured signal — count and resolve dissents — so suppression becomes visible. - Treat the cultural mechanism, not just the hardware fix, as the deliverable that must persist between major incidents.

REFLECTION QUESTIONS

1. What is the equivalent “diagnosed but not repaired” pathology in your domain? What evidence would close the loop?
2. The CAIB called culture engineerable. Sketch the engineering deliverable for the cultural intervention you would propose in your domain — including its measurement signal.
3. Seventeen years separated two accidents with the same root pathology. What mechanism in your organization would verify that a past diagnosis has actually been remediated rather than merely documented?
4. Vaughan reframed the Thiokol teleconference not as deviant override but as a working group whose decision rules had drifted. Identify a recurring decision in your domain whose rules of evidence may have drifted incrementally, and specify the audit that would surface the drift before the next high-consequence call.

WHO BUILDS THIS safety science & organizational analysis · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. TOOLS incident analysis · safety dashboards · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

Boeing Starliner

K Knowledge & Institutional Memory **D** Designed Out

IMPACT Multiple delays; the 2024 crewed flight left two NASA astronauts at the ISS for months; contemporary case for capability erosion at a legacy contractor

IN BRIEF

Boeing's Starliner, meant to be the second U.S. commercial crew vehicle alongside SpaceX's Crew Dragon, accumulated failures across a decade: software errors on its 2019 uncrewed flight, valve corrosion scrubbing a 2021 launch, and propulsion-system trouble on its 2024 crewed test that left two NASA astronauts on the ISS rather than returning them as contracted — NASA brought them home on SpaceX months later. GAO and NASA reviews found capability erosion across multiple engineering disciplines at a contractor whose human-spaceflight record had once been definitive. The erosion looks partly generational (Apollo- and Shuttle-era engineers retired) and partly institutional (cost and schedule pressure, and the thinning of NASA's in-house depth to challenge the contractor). Starliner is the contemporary case for capability erosion at a legacy contractor.

BACKGROUND

Boeing was awarded a NASA Commercial Crew contract to build Starliner as the second U.S. vehicle to carry astronauts to the ISS, alongside SpaceX's Crew Dragon — drawing on a human-spaceflight heritage that had once been definitive in American aerospace.⁰ The award rested in part on that heritage: a contractor with a definitive human-spaceflight record was a presumed safe choice, and that presumption is exactly what the program would go on to test, since reputation was standing in for a current measurement of capability.

WHAT HAPPENED

Instead Starliner stumbled across a decade: software errors marred its 2019 uncrewed test flight, valve corrosion scrubbed a 2021 launch, and propulsion-system problems on the 2024 crewed test flight left the two NASA astronauts aboard the ISS rather than returning on the contracted spacecraft — NASA brought them home on a SpaceX vehicle months later.¹ The three failures fell in different subsystems across three separate years, which is itself telling: a single bad part is bad luck, but software, valves, and propulsion failing in succession points to something broader than any one component — a decline running across the engineering organization rather than within one of its parts.

THE INVESTIGATION

GAO and NASA reviews identified the program as exhibiting capability erosion across multiple engineering disciplines — software, valves, propulsion, integration testing — at a contractor whose track record had previously been definitive.² The erosion looked partly generational, as Apollo- and Shuttle-era engineers retired, and partly institutional: cost


pressure, schedule-driven decisions, and the loss of NASA's own in-house engineering depth to challenge the contractor.³ The two causes compound: as the contractor's senior depth thinned through retirement, the buyer's own engineering depth had also eroded, so the very expertise NASA would have needed to catch the slipping contractor was the expertise it had let go of.

THE CAPABILITY GAP

Starliner is the contemporary case for capability erosion at scale at a legacy contractor. The decline happened over decades and was visible only in retrospect, because the institutional architecture for catching a supplier's slow capability decay — and updating the buyer's confidence to match — did not exist. Reputation outran reality, and no instrument was measuring the gap.⁴ Slow decay is the hard case to catch: no single year's results look alarming, the brand keeps the buyer's confidence high, and without an instrument that tracks the supplier's actual current capability the divergence is only legible once a crewed flight forces the reckoning.

AFTERMATH & REFORM

NASA leaned harder on independent reviews and on SpaceX as the reliable alternative, and the episode sharpened questions about how to sustain — and verify — capability at sole-source and legacy suppliers.⁵ It pairs with Saturn V (Case 112): where that case lost a capability to retirement, Starliner shows the same erosion in slow motion at a living institution still carrying the brand of the capability it had let thin. Having a second supplier to fall back on is what let NASA absorb the failure, which is itself the lesson: where a capability can erode unseen, a verified alternative is the difference between an embarrassment and a crew with no way home.



Starliner has demonstrated significant readiness shortfalls across multiple engineering disciplines.

EDITORS' SYNTHESIS OF GAO AND NASA REVIEWS OF THE COMMERCIAL CREW PROGRAM

REFERENCES

1. U.S. GAO, NASA Commercial Crew Program, GAO-20-121 (Jan. 2020) and GAO-19-504 (2019) — schedule slips and technical risks across the program.
4. NASA Aerospace Safety Advisory Panel reports (2020–2024) — the institutional and generational factors.

2. The Starliner test history — 2019 software errors, 2021 valve scrub, and the 2024 crewed flight and ISS stay.

3. NASA reviews and reporting on the Starliner test program — issues across software, valves, propulsion, and integration testing (editors' synthesis).
5. The 2024–2025 return of the Starliner crew aboard a SpaceX vehicle.

6. Cf. Saturn V (Case 112); N. Augustine, Augustine's Laws (1986).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Starliner is the case for sustained capability erosion at a legacy contractor whose track record had previously been definitive. The erosion happened over decades and was visible in retrospect; the institutional architecture for catching it did not exist.

LENS APPROACH

LENS uses Starliner in LEN 5 to teach one clean capability gap: the contractor–NASA interface capability that thinned over a decade. As NASA shifted to fixed-price commercial contracting, its own oversight and insight capacity eroded, so the problems an integrated review would have caught instead surfaced in flight (against the cross-subsystem troubles as the visible symptom). The teaching point is that the buyer's review capability is an engineered deliverable in its own right — let it thin and a slipping supplier becomes legible only when a crewed flight forces the reckoning.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
SPACE / AEROSPACE	→ KD	→ D1D2D3D4D5	2	6.3	1
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Verify a supplier's current capability against present evidence rather than awarding on reputation and a once-definitive track record.
- ▶ Preserve enough in-house engineering depth to genuinely challenge a contractor, since a buyer who has let its own expertise erode cannot catch a slipping supplier.
- ▶ Maintain a verified second supplier where the stakes are crewed, so a capability that erodes unseen does not become a single point of failure.

IN OPERATION / AFTER

- ▶ Instrument supplier capability so slow, multi-year decay is visible before a high-stakes flight forces the reckoning.
- ▶ Watch for cross-subsystem patterns — failures in software, valves, and propulsion in succession — as a signal of organization-wide erosion, not isolated bad luck.
- ▶ Update institutional confidence to match measured reality, so a contractor's brand cannot keep outrunning its current performance.

REFLECTION QUESTIONS

1. Identify a legacy supplier in your domain whose capability track record may have eroded faster than your institutional confidence in them has updated. What is the signal?
2. Design the contractor–capability review that would have caught the Starliner gaps before the 2024 crewed flight.
3. Starliner's failures spanned software, valves, and propulsion across three separate years — a pattern, not a single bad part. In your domain, what would distinguish a one-off component failure from organization-wide capability erosion, and how soon could you tell them apart?

WHO BUILDS THIS knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · task analysis · design prototyping

Chapter 6

Aviation & Aerospace — What Works — and the Frontier

When an industry engineers its own error rate down and keeps the books open.

Flying got safer when reporting became cheaper than concealment.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

CASE 114

AVIATION SAFETY

STRUCTURAL ENGINEERING

REGULATORY RULEMAKING

1988 – 2010S

FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule

G Governance & Trust **D** Designed Out **K** Knowledge & Institutional Memory

IMPACT After Aloha Airlines 243 (1988) exposed widespread fatigue cracking in an aging Boeing 737-200, the FAA's Aging Aircraft Program and the Airworthiness Assurance Working Group produced two decades of structural inspection programs culminating in 14 CFR Part 26 (2010), which requires manufacturers to establish a Limit of Validity for each model and embed widespread fatigue damage protection into the maintenance program

IN BRIEF

On April 28, 1988, Aloha Airlines Flight 243 lost an 18-foot section of upper fuselage in flight; the aircraft, a Boeing 737-200, had accumulated 89,680 flight cycles and was operating well past the design assumptions of the original certification. The accident — one flight attendant killed, 65 injured — exposed the structural-engineering category of widespread fatigue damage (WFD): multiple small cracks across many similar structural details, simultaneous enough that established single-crack inspection assumptions did not apply. The FAA stood up the Aging Aircraft Program almost immediately, and the Airworthiness Assurance Working Group (AAWG) operated through the 1990s and 2000s producing structural inspection programs across the transport-airplane fleet. The work culminated in 14 CFR Part 26 (Continued Airworthiness and Safety Improvements; WFD provisions effective 2011), which requires manufacturers to establish a Limit of Validity (LOV) for each type — the operational service goal below which the maintenance program protects against WFD — and to embed WFD prevention into the structural maintenance program itself. Part 26 is one of the most concrete recent examples of legacy assets aging past their original oversight regime being pulled back under structured airworthiness governance, and closes a long-standing zero-state in the induced framework's C7 competency.

BACKGROUND

Aloha Airlines Flight 243 on April 28, 1988, was the sentinel event. The Boeing 737-200 had accumulated 89,680 flight cycles — well above the design service goal — and was operating in a high-cycle short-haul environment that stressed the fuselage skin and lap joints at a rate the original certification analysis had not anticipated. In cruise at 24,000 feet, an 18-foot section of upper fuselage separated. The crew recovered the aircraft and landed it; one flight attendant was lost overboard and 65 occupants were injured. The NTSB investigation identified multiple small fatigue cracks across many lap-joint rivet holes, linking up

catastrophically — the textbook presentation of widespread fatigue damage rather than a single-crack failure.⁰

THE INTERVENTION

The structural-engineering category WFD names is specifically the regime where established inspection assumptions break down. A single-crack model assumes one crack initiates, propagates, and is caught by scheduled inspection before reaching critical length. WFD assumes many small cracks initiate at similar structural details across the fleet at similar times, and that link-up between adjacent cracks becomes the dominant failure mode. The inspection cadence, detection threshold, and replacement program a designer would build under the single-crack assumption do not adequately cover the WFD regime, and the airworthiness oversight regime Aloha exposed had not made the distinction.¹

HOW IT WORKED

The FAA's response was to stand up the Aging Aircraft Program and convene the Airworthiness Assurance Working Group, which operated through the 1990s and 2000s. The AAWG produced mandatory structural inspection documents and supplemental structural inspection programs for transport-category aircraft model-by-model, working through the older fleet systematically. The work was technically substantial — fleet- survey data, fatigue-test campaigns, inspection-program revisions for in-service aircraft — and institutionally sustained across more than two decades. The intermediate deliverables were per-model inspection programs; the capstone was rulemaking.²

THE EVIDENCE

14 CFR Part 26 — Continued Airworthiness and Safety Improvements — was finalized in 2007 with the WFD provisions published November 15, 2010 and effective January 14, 2011. Two structural elements are load-bearing. First, every transport-category type must have a Limit of Validity established by the manufacturer: the operational service goal, in flight cycles or flight hours, below which the maintenance program is shown to protect against WFD. Operation beyond the LOV requires either the manufacturer establishing an extended LOV with supporting analysis, or operator-specific evidence approved by the regulator. Second, WFD prevention is embedded in the structural maintenance program itself, not handled as a separate one-time inspection event. The maintenance program becomes the carrier of WFD protection across the type's service life.³

WHAT TRANSFERRED

What the case teaches at the LENS framing layer is the structural form of pulling legacy assets back under structured airworthiness governance. The original 737-200 certification did not anticipate the cycle count and the WFD regime that high-cycle short-haul operation produced; the regulatory response was not to retire the type but to engineer the oversight regime forward — fleet survey, AAWG, model-by-model inspection programs, and finally Part 26 codifying the LOV concept across the transport-category fleet. The case is one of the v2 sweep's clearest closes of the C7 (Capability under system change, transition, and aging assumptions) zero-state in the induced framework: a sustained two-decade regulatory program that turned an aging-fleet structural surprise into a governed sustainment discipline.⁴

The deliverable was not a single rule. It was the two-decade per-model technical record that made the rule a codification of established practice rather than top-down regulation.

EDITORS' SYNTHESIS OF THE FAA AGING AIRCRAFT PROGRAM AND THE PART 26 RULEMAKING RECORD.

REFERENCES

1. NTSB (1989), Aircraft Accident Report AAR-89/03, Aloha Airlines Flight 243, Boeing 737-200, N73711.

2. FAA, 14 CFR Part 26, "Continued Airworthiness and Safety Improvements for Transport Category Airplanes," Final Rule (2007); WFD/LOV Final Rule published Nov 15, 2010, effective Jan 14, 2011.

3. Airworthiness Assurance Working Group reports (1990s-2000s) — per-model structural inspection program record.

4. FAA Aging Aircraft Program documentation (1988 – present) — institutional record of the sustained regulatory work.

5. Swift, "Widespread Fatigue Damage Monitoring — Issues and Concerns," International Conference on Aging Aircraft — technical synthesis of the WFD inspection regime.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The FAA Aging Aircraft Program and Part 26 are one of the v2 sweep's clearest closes of the C7 zero-state. Aloha 243 exposed a regime the original certification did not cover; the AAWG operated for two decades; Part 26 codified Limit of Validity and embedded WFD prevention into the maintenance program itself. The sustained two-decade rulemaking is the deliverable, not the rule alone.

LENS APPROACH

FAA aging aircraft is the canonical aging-asset-governance case (induced 7.3; LENS D1/PT3) — Domain 1 for LOV-as- operational-service-goal; Domain 5 for the AAWG institutional discipline. Closes the C7 zero-state. Pair with Cases 115, 174, 156 as the v2 aging-system quartet.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION SAFETY	GDK	D1D2D3D4D5	3	7.3	1-5
STRUCTURAL ENGINEERING	→	→ Systems Analysis			
REGULATORY RULEMAKING					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Name the structural regime the original certification analysis did not cover — WFD vs. single-crack here — explicitly, so the inspection program can be designed around the actual failure mode rather than retrofitted to the original assumption.
- ▶ Operate the working group across a long horizon (AAWG: two decades): the per-model deliverables build the technical record that supports the eventual rulemaking, and short-cycle deliverables alone do not.
- ▶ Make the operational service goal (LOV) the artifact the regulation rests on; the goal is auditable per type, and operation beyond it requires additional analysis rather than implicit extension.

IN OPERATION / AFTER

- ▶ Embed the new protection in the maintenance program rather than as a one-time inspection: the maintenance program is the carrier across the service life, and one-time events are not.
- ▶ Treat the model-by-model technical work as the legitimacy basis for the rule; without the AAWG's fleet-survey and inspection-program record, Part 26 would have looked like top-down regulation rather than codification of established practice.
- ▶ Carry the case as the C7 instance the induced framework needed: legacy assets aging past their original oversight regime, pulled back under structured airworthiness governance by sustained two-decade rulemaking.

REFLECTION QUESTIONS

1. Identify a legacy asset class in your domain whose original certification or design analysis no longer covers the regime the assets are now operating in. What is the analog of the WFD distinction — the failure mode the original analysis did not anticipate?
2. Specify the analog of the LOV: an operational service goal, auditable per asset, beyond which additional analysis is required. What instrument would you use, what cadence, and what would “operation beyond LOV” trigger in your domain?
3. The AAWG operated for two decades before Part 26 codified the practice. What is the institutional discipline you would build for a similar two-decade horizon — the working-group cadence, the per-asset technical record, the eventual rulemaking — and is there a body in your domain that could plausibly carry that horizon?

WHO BUILDS THIS governance, ethics & measurement · systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · task analysis · design prototyping · knowledge bases · data pipelines

FAA NextGen and the ADS-B Out Transition

G Governance & Trust **D** Designed Out **K** Knowledge & Institutional Memory

IMPACT The FAA's Next Generation Air Transportation System (NextGen) shifted US air-traffic management from radar-based surveillance to a satellite-based architecture; the ADS-B Out mandate effective January 1, 2020 required equipage across the controlled-airspace fleet, achieving substantial compliance — with documented schedule slippage and benefit-realization gaps preserved as load-bearing hedges

IN BRIEF

The FAA's Next Generation Air Transportation System (NextGen) is the multi-decade transition of US air-traffic management from a radar-based surveillance architecture to a satellite-based architecture built on Automatic Dependent Surveillance – Broadcast (ADS-B). The ADS-B Out final rule, published in 2010, set the equipage mandate for January 1, 2020 — aircraft operating in most controlled airspace must broadcast their GPS-derived position and identity, replacing the radar-only secondary-surveillance model that defined the era prior. The mandate was substantially met at the deadline; equipage compliance across the affected fleet was high, and ADS-B-based surveillance is now the operational backbone in much of US airspace. The case is one of the largest planned aging- infrastructure transitions in the recent regulatory record — closing a long-standing C7 zero-state in the induced framework — and it carries the load-bearing hedges that GAO and DOT Inspector General reviews have repeatedly documented: significant schedule slippage across the program, benefit- realization gaps relative to original projections, and contested cost-benefit accounting. The transition happened; the transition was harder, slower, and more partial than the original NextGen plan implied.

BACKGROUND

The pre-NextGen US air-traffic management architecture rested on ground-based primary and secondary radar surveillance, with voice communication, paper or quasi-paper flight progress strips at many facilities, and traffic-flow management built around scheduled equipage upgrades that lagged the broader civil-aviation hardware base. By the early 2000s the structural problem was named clearly in the policy literature: a 20th-century surveillance architecture was being asked to handle 21st-century traffic, and incremental modernization inside the radar-paradigm was approaching its limits.^o

THE INTERVENTION

NextGen, formally launched in 2003 under the Vision 100 – Century of Aviation Reauthorization Act, was the FAA's multi-decade response. The technical core was the transition from radar-based surveillance to ADS-B: aircraft broadcast their GPS-derived position once

per second, ground stations and other aircraft receive the broadcast, and the resulting surveillance picture is more precise, lower-latency, and less expensive to operate at scale than secondary radar. The broader NextGen program included additional elements (Data Communications, NAS Voice System, System-Wide Information Management) but ADS-B Out was the load-bearing equipage-mandate piece.¹

HOW IT WORKED

The ADS-B Out final rule was published in 2010 with the January 1, 2020 compliance deadline — a decade of lead time for operators to equip. At the deadline, substantial compliance was reported across the affected fleet; the FAA's surveillance picture in controlled airspace moved substantially onto the ADS-B architecture, and the operational transition was completed in the sense the rule intended. The transition is one of the largest infrastructure-replacement programs in the recent FAA record, and it executed.²

THE EVIDENCE

The hedges have to survive. GAO and DOT Inspector General reports across the 2010s repeatedly documented significant schedule slippage across NextGen — multiple elements missing original delivery dates, the broader program's realized benefits running below original projections, and contested cost-benefit accounting between FAA program reporting and external review. The benefit categories NextGen's business case rested on — fuel-burn reduction, delay reduction, increased airspace capacity — have materialized in some respects and not in others, and the attribution to NextGen specifically (vs. broader operational and weather variation) is contested in the published audit literature. Equipage happened; benefit realization is mixed and ongoing.³

WHAT TRANSFERRED

What the case teaches at the LENS framing layer is the sustainment-engineering form of a planned aging- infrastructure transition at continental scale. The transition happened: the regulatory architecture, the decade of lead time, the equipage mandate at the deadline, and the operational migration to ADS-B-based surveillance are real and large. The transition was harder, slower, and more partial than the original NextGen plan implied; the schedule slippage and benefit-realization gaps are not anomalies in the case but part of what infrastructure transitions at this scale look like in practice. Together with the FAA Aging Aircraft program (Case 114), Y2K (Case 174), and LWRS (Case 156), NextGen is part of the v2 quartet that closes the C7 zero-state in the induced framework — and is the instance where the hedges are largest.⁴

The transition happened. The transition was harder, slower, and more partial than the original plan implied. Both are part of the case.

EDITORS' SYNTHESIS OF THE GAO NEXTGEN AUDIT SERIES AND FAA PROGRAM REPORTING.

REFERENCES

1. FAA, "Automatic Dependent Surveillance – Broadcast (ADS-B) Out Performance Requirements to Support Air Traffic Control (ATC) Service," Final Rule, 14 CFR Part 91 (2010).
2. Vision 100 — Century of Aviation Reauthorization Act, Public Law 108-176 (2003) — NextGen program statutory basis.
3. Department of Transportation Office of Inspector General, NextGen audit report series — independent program review documenting cost growth and partial benefit realization.
4. FAA, NextGen Annual Reports (2010 – present) — program-self-report tier; read against the GAO and DOT IG reviews.

3. Government Accountability Office, “NextGen Air Transportation System” report series (2010s) — sustained external audit record on schedule slippage and benefit–realization gaps.

THE LEARNING ENGINEERING LENS

LE INSIGHT

NextGen / ADS-B is one of the largest planned aging– infrastructure transitions in the recent regulatory record. The equipage mandate executed at the January 2020 deadline. The broader NextGen program has documented significant schedule slippage and benefit–realization gaps in sustained external audit. The case closes a C7 zero–state with the largest acknowledged hedges in the v2 aging–system quartet.

LENS APPROACH

NextGen / ADS-B is the planned infrastructure–transition case (induced 7.1; LENS D1/PT4). LENS uses it in Domain 1 (Systems Analysis) for the separation of the auditable equipage mandate from the multifactorial benefit– realization business case, and in Domain 5 (Navigating Sociotechnical Constraints) for the decade of lead time and the sustained external–audit discipline. Closes the C7 zero–state alongside Cases 114 (FAA aging aircraft), 174 (Y2K), and 156 (LWRS) as the v2 aging–system quartet.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION INFRASTRUCTURE	G D K	D1 D2 D3 D4 D5	4	7.1	1–5
AIR TRAFFIC MANAGEMENT	→	→ Systems Analysis			
REGULATORY TRANSITION					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Specify the equipage mandate with the lead time the transition actually needs — a decade for ADS-B Out — so the deadline arrives with realistic compliance pathways rather than as a forcing function operators cannot meet.

▶ Separate the equipage–mandate piece (auditable, has a deadline) from the broader benefit–realization business case (multifactorial, harder to attribute); the equipage piece will deliver, the benefit piece will deliver partially, and conflating them sets the program up to look like a failure where it succeeded.

▶ Plan for sustained external audit (GAO, DOT IG) as part of the program’s operating environment — the schedule slippage and benefit–realization gap reporting is institutional discipline, not a sign the program is unworkable.

IN OPERATION / AFTER

- ▶ Report the equipage transition as the deliverable that executed and the benefit realization as the deliverable that is ongoing and mixed; both are real, and the load–bearing hedge is that the original NextGen plan understated the difficulty of the latter.

▶ Sustain the audit–and–rebaseline cadence after the headline equipage mandate; the benefit–realization picture continues to develop, and the institutional record of slippage and partial realization is part of how transitions at this scale are documented honestly.

▶ Treat the NextGen case as the instance where the C7 transition delivered with the largest acknowledged hedge — useful precisely because the transition is real and the hedges are also real, and conflating either with the other misreads the lesson.

REFLECTION QUESTIONS

1. Identify an aging-infrastructure transition in your domain whose original business case rested on benefit categories that are multifactorial and contested in attribution. What is the equivalent of the auditable equipage mandate — the piece of the transition that delivers cleanly — and how would you separate its reporting from the broader benefit case?
2. Specify the lead time the transition actually needs from mandate to deadline; ADS-B Out used a decade. What is the analog in your context, and what does the operator-equipage trajectory look like across that horizon?

WHO BUILDS THIS governance, ethics & measurement · systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · task analysis · design prototyping · knowledge bases · data pipelines

Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change

G Governance & Trust **K** Knowledge & Institutional Memory **H** Human-System Interface

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT A 17,000-line pilot modernization of the Eurocat air-traffic-management system was scoped explicitly to generate safety evidence that an automated transformation was non-distortive of original functionality; that pilot-tier evidence was used to convince customers to accept a system-wide architecture-driven modernization with 100% automated code transformation

IN BRIEF

The Eurocat Air Traffic Management System was the kind of safety-critical legacy software whose customers cannot accept a big-bang rewrite: the operational system in production cannot tolerate the discontinuity. The 2005 Thales-led pilot modernization was scoped narrowly — 17,000 lines of code — with the deliberate goal of generating safety-equivalence evidence that an automated transformation did not distort the original functionality. The pilot succeeded in producing that evidence, and the evidence was used to convince customers to accept a full architecture-driven modernization with 100% automated code transformation. The teaching pattern is the small-as-gateway-to-big move: the small-tier verification artifact dissolved the customer objection to the large-tier change. The case is documented in a vendor-authored Elsevier technical chapter rather than an independent academic evaluation; the evidence-tier flag is rendered under the title. The case is the small-tier C7 transition-verification companion to the big-tier failures (Patriot/Dhahran, Ariane 5, Knight Capital, CrowdStrike, TSB) already in the corpus. Future validation will continue on the long-run operational behavior of the modernized system.

BACKGROUND

Air traffic management software is a paradigmatic capability-system-misalignment-at-transition problem: the legacy system carries decades of accreted operational logic, the safety case the regulator and the airspace user expect rests on the legacy's documented behavior, and a big-bang rewrite is unacceptable because the operational discontinuity is

itself the safety hazard. The Eurocat problem in 2005 was the standard one — modernize without breaking the safety case the legacy already carries.⁰

THE INTERVENTION

The design move was to scope a pilot narrowly enough that the verification could be exhaustive. Seventeen thousand lines of code is small as a fraction of the full system; it is large enough that a working automated transformation pipeline can be exercised end-to-end and the safety-equivalence evidence can be generated against a realistic subset. The pilot was not a feasibility check. It was a deliberately designed evidence artifact: prove, at small scale and on the actual legacy, that the transformation does not distort the function the customer cares about.¹

HOW IT WORKED

The pilot succeeded in producing the evidence. Customers were then willing to accept a system-wide architecture-driven modernization with 100% automated code transformation. The causal claim the case rests on is not “the modernization worked” — that is a long-run operational question — but “the customer objection to the large-tier change was dissolved by the small-tier verification artifact.” That is the C7 teaching point at the small tier the corpus has not had: the same governance-objection-dissolver move that Case 199 / Case 200 show in the AV regulatory regime, transposed into a legacy-software transition.²

THE EVIDENCE

The evidence-tier flag matters. The case is documented in a practitioner-authored chapter in an Elsevier book on information-systems transformation. The authors are from the vendor side; there is no independent academic evaluation of the pilot’s safety-equivalence claims or of the post-modernization operational behavior. The pattern the case teaches — verification at the small tier as the gateway to the large-tier change — is robust across the practitioner literature on legacy modernization, but the Eurocat-specific magnitudes and the long-run system outcomes have not been audited in a peer-reviewed source. Future validation will continue.³

WHAT TRANSFERRED

What the case adds at the LENS layer is the small-tier C7 intervention that the big-tier failures already in the corpus do not teach: when the inherited design has to move to a new envelope, the verification work itself can be the capability deliverable that earns the change its legitimacy. The Domain-1 frame applies: systems analysis of the legacy’s safety case as an artifact, and design of the transformation pipeline as another artifact, are the capability-engineering moves the program made. The case is the missing positive example for induced 7.1 and 7.2 at the small tier.

The verification at the small tier is the artifact that earns the large-tier change its legitimacy. The pilot is not a feasibility check. It is the evidence.

EDITORS’ SYNTHESIS OF THE EUROCAT MODERNIZATION CHAPTER.

REFERENCES

1. DelaPeyronnie, Newcomb, Morillo, Trimech, Nguyen, & Purtill (2010), “Modernization of the Eurocat Air Traffic Management System (EATMS),” in *Information Systems Transformation: Ar-*
3. Brodie & Stonebraker (1995), *Migrating Legacy Systems* — the framing reference on small-step legacy modernization.

chitecture-Driven Modernization Case Studies (Elsevier / Morgan Kaufmann), Chapter 5.

2. Ulrich & Newcomb (eds., 2010), Information Systems Transformation — the host volume on architecture-driven modernization patterns.

4. Bisbal et al. (1999), "Legacy Information Systems: Issues and Directions," IEEE Software 16(5):103–111 — peer-reviewed framing companion.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Eurocat is the small-tier C7 transition-verification intervention the corpus needed: a narrowly scoped pilot designed to produce the safety-equivalence evidence that dissolves customer objection to a system-wide modernization. Evidence is vendor-authored practitioner literature; the long-run operational record of the modernized system is the open question. Future validation ongoing.

LENS APPROACH

Eurocat is the C7 small-tier transition case (induced 7.1 and 7.2; LENS D1/PT1). LENS uses it in Domain 1 (Systems Analysis) for CLO-1 — the safety case and the transformation pipeline are both engineered artifacts — and in Domain 2 (Iterative Development) for CLO-2, since the pilot is the iteration that earns the rollout. Pair with the big-tier C7 failures (Patriot/Dhahran, Ariane 5, Knight Capital, CrowdStrike, TSB) for the failure-and- intervention contrast.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AIR TRAFFIC MANAGEMENT	G K H	D1 D2 D3 D4 D5	1	7.1	1·2
LEGACY MODERNIZATION	→	→ Systems Analysis			
SOFTWARE TRANSFORMATION					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Scope the pilot to be small enough that the verification can be exhaustive against the actual legacy, not just a feasibility check on a synthetic subset.
- ▶ Design the safety-equivalence evidence as a deliberate deliverable from the pilot, not a byproduct: specify in advance what the customer needs to see to accept the large-tier change.
- ▶ Document the transformation pipeline itself as a verifiable artifact, so the customer can audit the transformation, not just inspect the transformed code.

IN OPERATION / AFTER

- ▶ Carry the practitioner-tier flag into any downstream citation; the case is a vendor account and the long-run operational behavior of the modernized system is the open question.
- ▶ Treat the small-as-gateway-to-big pattern as the teaching point; the magnitudes and the system-wide rollout are open and depend on the long-run operational record.
- ▶ Pair with the big-tier C7 failures (Patriot/Dhahran, Ariane 5, Knight Capital, TSB) when teaching; the pair shows the verification-as-deliverable principle is what separates the failure and intervention threads.

REFLECTION QUESTIONS

1. Identify a legacy system in your context whose customer or operator will not accept a big-bang rewrite. What is the small-tier pilot whose evidence would dissolve the objection to the large-tier change, and what specifically would the customer need to see?
2. Specify the safety-equivalence evidence the pilot would have to generate as a deliberate deliverable. The Eurocat pattern teaches because the pilot was scoped to produce the artifact the customer needed, not just to demonstrate feasibility.
3. The case is vendor-authored practitioner literature. What independent evidence (academic evaluation, multi-customer replication, long-run operational behavior) would you require before treating the small-as-gateway-to-big pattern as a settled engineering practice in your domain?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · usability testing · cognitive walkthroughs

Korean Air Safety Transformation

T Training Gap **N** Normalization of Deviance

IMPACT From industry pariah (16 aircraft written off, 700+ lives lost, 1970–1999) to spotless passenger safety record since 1999

IN BRIEF

Between 1970 and 1999, Korean Air had one of the worst safety records in commercial aviation — a loss rate roughly seventeen times United's, more than 700 lives lost — and after the 1997 Guam crash of Flight 801 killed 228, the president of South Korea called it “an embarrassment to the nation.” The NTSB traced root causes to cockpit authority gradients: junior officers were culturally unable to challenge a captain's errors. In 2000, Korean Air brought in Delta's David Greenberg to rebuild flight operations — mandating English as the cockpit language, adapting CRM for a high-power-distance culture, and bringing in outside consulting from Boeing and Delta. The airline has had no fatal passenger accident since, winning the 2006 Phoenix Award. Cultural legacy is not destiny when it is deliberately redesigned.

BACKGROUND

From 1970 to 1999, Korean Air suffered repeated fatal crashes — a loss rate roughly seventeen times United Airlines', with over 700 lives lost. The NTSB's investigation of the 1997 Guam crash of Flight 801, which killed 228, identified steep cockpit authority gradients as a root cause: junior officers were culturally unable to challenge a captain's erroneous decisions in time. The pattern echoed Tenerife's lesson but ran deeper, rooted in a national hierarchy the crew carried into the cockpit, so the silence in the right seat was not a lapse of training but an artifact of how rank was expressed in speech.⁹

THE INTERVENTION

After foreign carriers suspended code-shares over its safety record, Korean Air in 2000 hired David Greenberg from Delta to overhaul flight operations; the FAA's August 2001 downgrade of South Korea to Category 2 came a year later. The interventions were deliberate and cultural: mandated English fluency for all pilots, CRM training adapted for a high-power-distance setting, external consulting from Boeing and Delta, and fleet modernization. The external commercial pressure — suspended code-shares and a downgraded rating — supplied the leverage to act, turning a cultural problem the airline had long lived with into one it could no longer afford to ignore.¹

HOW IT WORKED

The load-bearing move operated on language. Making English the cockpit language stripped out the Korean honorific hierarchy that had silenced first officers, because English has no honorifics to enforce rank. The CRM adaptation then gave crews an explicit, culturally

workable protocol for raising concerns — converting deference into communication. The choice was elegant precisely because it changed the medium rather than asking pilots to override their own culture: with no honorifics to encode, the language itself flattened the gradient that had absorbed the first officer's warnings.²

THE EVIDENCE

Korean Air has recorded no fatal passenger accident since the reforms took hold. Air Transport World recognized the turnaround with its 2006 Phoenix Award, and the carrier moved from international pariah to a safety record indistinguishable from the best global operators. The categorical nature of the shift — from a loss rate many times the industry's to none — is what makes the case persuasive: the same crews and the same national culture produced opposite outcomes once the cockpit's communication structure was redesigned.³

WHAT TRANSFERRED

Korean Air is the strongest aviation evidence that cultural legacy is not destiny: a specific cultural feature — high power distance in the cockpit — was the binding constraint, and once it was redesigned, the safety record changed categorically. The case also shows that the interface for cultural redesign can itself be engineered, in this instance through language. The transferable lesson is to locate the single cultural feature that actually binds the capability rather than attempting to remake a culture wholesale, then change the medium through which that feature operates.⁴



Korean Air's record before 2000 was changed by an external intervention into cockpit culture, not by criticism of it.

EDITORS' SYNTHESIS OF NTSB FINDINGS ON KAL 801 AND THE KOREAN AIR TRANSFORMATION

REFERENCES

1. NTSB, Aircraft Accident Report: Korean Air Flight 801, Guam (2000) — authority gradients as a root cause.
2. Air Transport World Phoenix Award documentation (2006) — recognition of the turnaround.
3. Gladwell, M. (2008), *Outliers* — the Korean Air chapter on power distance and the English-language change.
4. Helmreich, Wilhelm, Klinec & Merritt (2001) — national culture and CRM adaptation.
5. Korean Air corporate safety reports — the post-2000 accident-free record.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Korean Air is the strongest aviation evidence that cultural legacy is not destiny. A specific cultural feature — high power distance in the cockpit — was the binding capability constraint. Once it was redesigned, the safety record changed categorically. The intervention is also methodologically interesting because it operated on language: English became the cockpit language because English has no honorifics to enforce. The interface for cultural redesign was a linguistic one.

LENS APPROACH

LENS uses Korean Air in LEN 8 as the canonical organizational- learning case under cultural constraint and in LEN 2 as a CRM- extension case for high-power-distance contexts. The case is paired in this book with the Toyota Andon Cord (Case 155) as cultural intervention success stories.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ T N	→ D1 D2 D3 D4 D5	3	4.3	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Locate the single cultural feature that actually binds the capability — here, high power distance expressed through honorifics — rather than attempting to remake the whole culture. ▶ Change the medium through which that feature operates (making English the cockpit language) so the constraint is removed structurally rather than by asking people to override their own culture. ▶ Pair the structural change with an adapted protocol — CRM tuned for a high-power-distance setting — that gives crews a culturally workable way to raise concerns.	▶ Use external pressure (suspended code-shares, a downgraded rating) as durable leverage to hold the reform in place against any drift back to the prior norm. ▶ Sustain the change with continued outside benchmarking and consulting so the new communication structure is reinforced rather than quietly eroding. ▶ Track the safety record over years to confirm the categorical shift holds, treating a maintained accident-free record as the evidence the redesign took.

REFLECTION QUESTIONS

1. Identify a cultural feature in an institution you work with that constrains capability. Is it engineerable? What would the redesign look like?
2. Korean Air’s intervention operated on language. What surface of culture is engineerable in your domain that you have not yet considered?
3. Korean Air acted only once suspended code-shares and a downgraded rating made the status quo unaffordable. What external pressure exists in your domain that could supply the leverage to redesign a long-tolerated cultural constraint — and how would you use it?

WHO BUILDS THIS learning science & instructional design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · incident analysis · safety dashboards

TCAS — Coordinated Collision Avoidance and the Überlingen Lesson



Human-System Interface



Knowledge & Institutional Memory



Governance & Trust

IMPACT TCAS II — the Traffic Alert and Collision Avoidance System — provides cockpit traffic display and coordinated Resolution Advisories (RAs) between aircraft on conflicting trajectories; mandated on US air-carrier and on most international turbine aircraft, TCAS converted mid-air collision in commercial aviation from a recurring fatality category to a rare event; the 2002 Überlingen mid-air (71 dead) exposed a specific coordination failure mode and drove the 2008 release of TCAS II Version 7.1 with the 'level off, level off' reversal logic

IN BRIEF

TCAS — the Traffic Alert and Collision Avoidance System, standardized in RTCA DO-185 and successors — is the cockpit automation that monitors transponder returns from nearby aircraft, computes potential conflicts, and issues Traffic Advisories and Resolution Advisories (RAs) to the crew. Operational TCAS II was mandated on US air-carrier aircraft by FAA rule in the early 1990s and on most international turbine aircraft by ICAO. RAs are coordinated: when two TCAS-equipped aircraft are in conflict, one is instructed to climb and the other to descend by negotiated inversion of the data-link. The intervention converted mid-air collision in commercial aviation from a recurring fatality category to a rare event. The case's load-bearing edge case is the 2002 Überlingen mid-air collision (71 dead), in which one crew followed the TCAS RA and the other followed an ATC instruction in the opposite direction. The BFU investigation identified the human-TCAS-ATC coordination failure mode and drove the 2008 release of TCAS II Version 7.1, which added the "level off, level off" reversal logic and clarified the precedence of TCAS RAs over ATC instructions. The hedge survives into the case: TCAS is among the strongest aviation automation interventions in the outcome record, and the Überlingen failure mode and its V7.1 correction are part of the case rather than smoothed away.

BACKGROUND

Mid-air collision in commercial aviation has been a recurring fatality category since the 1950s. The 1956 Grand Canyon mid-air (128 dead) prompted the modern US air-traffic-control system, but ATC alone cannot always prevent collision when traffic densities or coordination errors exceed the controller's reach. The FAA and RTCA developed the Traffic Alert and Collision Avoidance System (TCAS) through the 1980s as a cockpit collision-avoidance automation independent of ATC. TCAS II — the operational version with Resolution

Advisories — was mandated on US air-carrier aircraft beginning in the early 1990s, and on most international turbine aircraft by ICAO over the following decade.⁹

THE INTERVENTION

The TCAS II architecture is what the case rests on. Each equipped aircraft interrogates the transponders of nearby aircraft and computes a closest-point-of-approach projection from range, altitude, and rate data. If the projection enters the conflict envelope, TCAS issues a Traffic Advisory (TA) — a cue to the crew to acquire the other aircraft visually if possible. If the conflict persists, TCAS issues a Resolution Advisory (RA): a specific vertical-rate command (“Climb, climb” or “Descend, descend”). When two TCAS-equipped aircraft are in conflict, the two RAs are coordinated via the Mode S data link so the aircraft are instructed to diverge — one climbing, one descending — rather than both maneuvering in the same direction.¹

HOW IT WORKED

The deployed outcome record across the 1990s and 2000s was strong. Mid-air collision in TCAS-equipped fleets became rare. The case nevertheless turns on Überlingen. On July 1, 2002, a Russian Tu-154 and a DHL Boeing 757 approached on conflicting trajectories at FL360 over southern Germany. Both aircraft received coordinated TCAS RAs — the 757 to descend, the Tu-154 to climb. The Tu-154 crew received the climb RA and, almost simultaneously, an ATC instruction to descend from a Skyguide controller working alone on degraded console configuration. The Tu-154 crew followed the ATC instruction; the 757 followed the TCAS RA; both aircraft descended into each other. 71 people died, including 52 children on a school trip.²

THE EVIDENCE

The BFU investigation identified a specific failure mode at the human-automation-controller boundary: TCAS RA precedence over ATC instructions was insufficiently clear in the crew procedures, the data-link coordination between the two TCAS units had performed as designed but could not enforce the result on the crew, and the ATC single-controller / degraded-console context was a systemic failure. The aviation response was operational and technical: ICAO and national regulators clarified that TCAS RAs take precedence over conflicting ATC instructions; ATC procedures were tightened around coordination of conflict-resolution between controllers and TCAS; and RTCA / Eurocontrol developed TCAS II Version 7.1 (released 2008), which added the “level off, level off” reversal logic for the specific scenario where one aircraft does not follow its RA, and clarified RA wording and behavior at the boundary.³

WHAT TRANSFERRED

The hedge survives into the case. TCAS is among the strongest aviation automation interventions in the outcome record; the Überlingen failure mode is not a refutation of the system but a documented coordination limit that drove the V7.1 correction. The case teaches the coordinated-automation form at its most durable, with the discipline that the human-automation-human- operator triangle (crew, TCAS, ATC) has to be designed coherently. A single-controller / degraded-console situation and an unclear precedence rule converted a working automation into a fatal outcome; both were addressed in the post-Überlingen response, and both remain part of the case rather than smoothed away.⁴

The intervention is not less effective for having a documented coordination edge case; it is more credible because the edge case is named and addressed.

EDITORS' SYNTHESIS OF THE BFU ÜBERLINGEN REPORT AND THE V7.1 DEVELOPMENT RECORD.

REFERENCES

1. RTCA (2008), DO-185B “Minimum Operational Performance Standards for Traffic Alert and Collision Avoidance System II (TCAS II)” — Version 7.1 with reversal logic.

2. Bundesstelle für Flugunfalluntersuchung (BFU) (2004), AX001-1-2/02 — Investigation Report on the mid-air collision on 1 July 2002 near Überlingen.

3. Eurocontrol (2008 – 2011), ACAS II Bulletin and Programme for the Harmonised Implementation of Satellite Navigation — V7.1 mandate timing in Europe.

4. Federal Aviation Administration, 14 CFR § 121.356 — TCAS II equipage requirement.

5. NTSB / FAA TCAS historical safety studies (1990 – 2010) — outcome record context.

THE LEARNING ENGINEERING LENS

LE INSIGHT

TCAS is among the strongest aviation automation interventions in the outcome record; the Überlingen coordination failure mode is part of the case rather than smoothed away. The human–automation–ATC triangle has to be designed coherently; V7.1’s reversal logic and the clarified precedence rule are the iterative–design response to the documented edge case.

LENS APPROACH

TCAS is the coordinated-cockpit-automation case (induced 3.1; LENS D4/PT5) — Domain 4 for the Überlingen–driven V7.1 iteration; Domain 3 for the crew–TCAS–ATC precedence rule. Pair with EGPWS (Case 120) and Case 27.

DOMAIN	Modes Addressed	LENS Competency	Problem Type	Induced	CLO
AVIATION	H K G	D1 D2 D3 D4 D5	5	3.1	4 · 3
AUTOMATION		Test and Evaluation			
HUMAN FACTORS					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the coordinated–automation logic for the case where one of the coordinated agents does not comply — the Überlingen failure mode — not just for the case where both comply. The V7.1 reversal logic is the worked example of that design discipline.
- ▶ Specify precedence at the human–automation boundary unambiguously and in advance: TCAS RAs take precedence over conflicting ATC instructions. Leaving precedence to crew judgment under time pressure is the design choice that produced Überlingen.
- ▶ Treat the ATC context (single–controller, degraded console) as part of the system the automation operates in, not as a precondition the automation can assume away. Systemic failures at the boundary determine the boundary behavior.

IN OPERATION / AFTER

- ▶ Carry the Überlingen failure mode into the case rather than smoothing it away. The intervention is not less effective for having a documented coordination edge case; it is more credible because the edge case is named and addressed.
- ▶ Track the post–V7.1 outcome record as evidence about the correction, not just about the original intervention; the lesson is the iterative–design discipline that the human–automation triangle requires after a failure mode is exposed.
- ▶ Use the case as the canonical pair to EGPWS (Case 120): two cockpit automations, two outcome categories closed, one with a coordination edge case that drove a version revision and one without.

REFLECTION QUESTIONS

1. Identify a coordinated-automation system in your domain where two agents must comply for the resolution to work. What is the analog of the V7.1 reversal logic — the design for the case where one agent does not comply?
2. Specify the precedence rule at the human-automation boundary in your system. Überlingen turned on an ambiguous precedence rule; the post-2002 clarification is the worked example of why precedence has to be unambiguous in advance.
3. The systemic context (Skyguide single-controller / degraded console) was part of the failure. What contextual preconditions does your automation assume that, if they fail, would convert a working automation into a failure mode?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

Singapore Airlines Safety Transformation

T Training Gap **N** Normalization of Deviance

IMPACT Sustained safety record over decades despite challenging operating conditions; among the most safety-invested carriers in commercial aviation

IN BRIEF

Singapore Airlines has invested in safety capability across decades in a way that sets it apart from carriers operating under comparable conditions — early adoption of CRM, heavy simulator investment, a young-fleet policy, and a strong reporting culture, sustained even through rapid growth. The 2000 crash of Flight SQ006, which attempted takeoff from a closed, partly-constructed runway at Taipei during a typhoon and killed 83, prompted a sustained institutional self-examination — operational changes, training updates, and public transparency about what had happened — that became a model in the literature on post-accident learning. The airline is the operational successor to Korean Air (Case 118): an Asian carrier that engineered its safety capability deliberately and sustained the investment as a competitive differentiator, not only in response to crisis.

BACKGROUND

Commercial aviation runs on thin margins, and safety investment — simulators, training hours, fleet renewal — is a cost that competitive pressure constantly pushes downward. The question for any carrier is whether to treat capability as a floor set by regulation or as a deliberate, sustained investment ahead of the minimum. The pressure is structural rather than occasional — every budget cycle invites trimming the margin between regulatory minimum and actual capability, so sustaining the investment requires deciding the question deliberately rather than by default.⁰

THE INTERVENTION

Singapore Airlines chose sustained investment. From the 1980s it was an early adopter of Crew Resource Management and CRM-style culture work tuned to its operating context, and it committed to heavy simulator investment, a deliberately young fleet, and a strong internal reporting culture — maintaining these even during periods of rapid expansion. Holding the investment through rapid growth is the hard test — expansion is precisely when the temptation to let capability lag the fleet is strongest, and maintaining it then is what separates a sustained commitment from a fair-weather one.¹

HOW IT WORKED

The carrier treated safety capability as a competitive differentiator rather than a regulatory burden, pairing technical investment — training systems, modern aircraft — with a culture of transparency and reporting. Investing ahead of regulatory minimums made the capability a managed system parameter, not a residual of cost-cutting decisions made elsewhere.

Framing safety as a differentiator rather than a burden is what made the investment defensible against cost pressure — it tied capability to the brand and the premium the airline charged, giving the spend a commercial rationale, not just a safety one.²

THE EVIDENCE

The 2000 crash of Flight SQ006 — an attempted takeoff from a closed, partly-constructed runway at Taipei during Typhoon Xangsane, killing 83 — tested the institution. Its response, documented in the Taiwan investigation, became a reference case in post-accident institutional learning: operational changes, training updates, and public transparency about what had happened and why, rather than defensiveness. The test of a safety culture is how it behaves after its own accident, and choosing transparency over defensiveness is what turned SQ006 from a refutation of the airline's reputation into evidence the reporting culture extended to its own failures.³

WHAT TRANSFERRED

Singapore Airlines is the case for sustained capability investment under competitive pressure, and the operational successor to Korean Air (Case 118): where Korean Air is a transformation forced by crisis, Singapore Airlines is deliberate investment sustained without one. Together they show two routes — crisis-driven and voluntary — to the same engineered safety capability. The voluntary route is the harder one to hold, because it has no catastrophe to point to as justification, which is why framing the investment as a competitive differentiator matters: it supplies the rationale that a crisis would otherwise provide.⁴



Singapore Airlines has consistently invested in safety capability ahead of regulatory minimums.

EDITORS' SYNTHESIS ON SINGAPORE AIRLINES' SUSTAINED SAFETY INVESTMENT

REFERENCES

1. Aviation Safety Council (Taiwan), SQ006 Accident Investigation Final Report (2002) — the accident and the airline's response.
2. Singapore Airlines safety reports (multiple years) — training, fleet, and reporting—culture investment.
3. IATA Operational Safety Audit (IOSA) documentation — investment ahead of regulatory minimums (paraphrased).
4. Helmreich, Wilhelm, Klinec & Merritt (2001) — national culture and CRM adaptation.
5. Weick & Sutcliffe (2007), Managing the Unexpected — sustained high-reliability investment.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Singapore Airlines is the case for sustained capability investment in a competitive industry. The carrier has chosen safety capability investment as a primary differentiator. The result over decades is a safety record that distinguishes it from peers operating under comparable conditions.

LENS APPROACH

LENS uses Singapore Airlines in LEN 8 for sustained institutional capability investment under competitive pressure. The case pairs with Korean Air (Case 118) as Asian-carrier capability stories of different shapes — one transformation under crisis, the other sustained investment without crisis.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION	→ T N	→ D1 D2 D3 D4 D5	4	1.4	5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Decide deliberately to invest ahead of regulatory minimums — simulators, a young fleet, training hours — rather than letting capability settle at the floor by default under cost pressure.
- ▶ Pair the technical investment with a transparency and reporting culture, so capability is a managed system parameter rather than a residual of cost-cutting decisions elsewhere.
- ▶ Frame safety capability as a competitive differentiator tied to the brand, giving the spend a commercial rationale that can survive scrutiny, not just a safety one.

IN OPERATION / AFTER

- ▶ Hold the investment through periods of rapid growth — the moment the temptation to let capability lag the fleet is strongest — since maintaining it then is what makes the commitment sustained rather than fair-weather.
- ▶ Respond to the institution's own accidents with transparency over defensiveness, demonstrating that the reporting culture extends to its own failures and converting a setback into evidence the culture is real.
- ▶ Anchor the investment to the brand and premium so it survives leadership and budget cycles, supplying the durable justification a voluntary commitment lacks without a crisis to point to.

REFLECTION QUESTIONS

1. Identify an institution in your domain that has chosen capability investment as a primary differentiator. What pattern has it sustained?
2. Design the institutional architecture that makes sustained capability investment defensible against competitive cost pressure.
3. Singapore Airlines sustained its investment voluntarily, without a crisis to point to, by framing capability as a competitive differentiator. Identify a safety or capability investment in your domain that lacks a catastrophe to justify it, and construct the commercial rationale that would defend it against the next budget cut.

WHO BUILDS THIS learning science & instructional design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards

FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike

G Governance & Trust**K** Knowledge & Institutional Memory**N** Normalization of Deviance

IMPACT After Controlled Flight Into Terrain emerged as the leading cause of commercial-jet fatalities through the late 1980s, the Flight Safety Foundation convened industry-wide task forces that produced the CFIT Checklist, the ALAR Tool Kit, and the institutional momentum behind Terrain Awareness and Warning System (TAWS) mandates; CFIT and ALAR accident rates fell sharply over the subsequent decade

IN BRIEF

Through the late 1980s and into the early 1990s, Controlled Flight Into Terrain — a serviceable aircraft flown under control into the ground, water, or an obstacle — was the leading cause of commercial-jet fatalities worldwide. The Flight Safety Foundation (FSF), an independent industry body, convened the CFIT Task Force in 1992 and, in parallel with the broader ICAO response, produced the CFIT Checklist — a structured tool for operators to assess their own exposure. The Approach-and-Landing Accident Reduction (ALAR) Task Force followed in 1996, producing the ALAR Tool Kit (released 1998) covering the approach phase where roughly half of fatal accidents then occurred. The institutional momentum from those task forces sat behind the eventual Terrain Awareness and Warning System (TAWS) mandates in the US (2002) and ICAO (2007). CFIT and approach-and-landing accident rates fell sharply through the subsequent decade. The case teaches industry-level institution building after a catastrophe-class spike: the deliverable is the cross-operator tool, the diagnostic structure, and the coordinated path to mandate. The hedge survives — the rate decline is multifactorial (TAWS hardware, training, procedural change, fleet turnover) — and the institutional form is what the case is teachable on.

BACKGROUND

Through the late 1980s and into the early 1990s, Controlled Flight Into Terrain dominated commercial-jet fatality statistics. A serviceable aircraft, flown under control, was finding terrain or water — the crew's mental model of position and trajectory diverged from reality at the worst possible moment, and the existing Ground Proximity Warning Systems (GPWS) generation reached the crew too late to recover in many terrain encounters. The pattern was recognized across operators and regulators, and the response that emerged was industry coordination led by an independent body — the Flight Safety Foundation — rather than regulator-first action alone.^o

THE INTERVENTION

The CFIT Task Force convened in 1992 and operated as a working group across operators, manufacturers, regulators, and pilot organizations. Its central deliverable was the CFIT Checklist — a structured self-assessment instrument that let an operator score its own CFIT-risk exposure across route, equipment, training, and procedure dimensions, and identify where the gaps sat. The Checklist was distributed without restriction. The institutional theory was straightforward: a cross-operator diagnostic, owned by no single competitor, would let the industry move together on a problem that no single operator's own data could fully characterize.¹

HOW IT WORKED

The Approach-and-Landing Accident Reduction Task Force followed in 1996, scoped to the approach phase, which accounted for roughly half of fatal accidents in the analysis window. The ALAR Tool Kit, released in 1998, was the most substantial deliverable of the entire industry-coordination effort: a multi-element package of briefing notes, training aids, video content, and risk-assessment instruments covering stabilized approach criteria, runway excursion, monitored approach practice, and crew procedural design. The Tool Kit was adopted across operators of all sizes precisely because the FSF had no competitive stake in any one airline's adoption.²

THE EVIDENCE

The eventual regulatory action — the FAA's TAWS mandate (2002, retrofitting earlier GPWS-equipped aircraft with the newer terrain-database-driven warning system) and the ICAO TAWS requirement (effective 2007) — sat downstream of the task-force momentum, not upstream of it. CFIT accident rates fell sharply through the subsequent decade across both commercial transport and corporate aviation; approach-and-landing accidents declined alongside. The before/after pattern is robust across multiple independent operator and regulator datasets, and the case is treated in the FSF's own retrospective literature as one of the discipline's strongest examples of coordinated industry-led intervention.³

WHAT TRANSFERRED

The load-bearing hedge survives into the case. The accident-rate decline is multifactorial: TAWS hardware in the cockpit, training changes the task forces catalyzed, stabilized-approach criteria adopted at the operator level, and steady fleet turnover into airframes with more capable equipment all moved together. Attributing the entire decline to the FSF task forces alone overstates the evidence; what the evidence supports is that the industry-coordination form — independent convening body, cross-operator diagnostic tools released without competitive restriction, momentum sustained to a regulatory mandate — was the institutional mechanism that organized the response, and the response worked. The case is the canonical C6.1 instance of industry-level institution building after a catastrophe-class spike, paired with v1 ASRS (Case 119) and CAST (Case 117) at the industry-coordination layer.⁴

The deliverable is the cross-operator diagnostic, owned by no single competitor, released without restriction. The mandate consolidates what the field already does.

EDITORS' SYNTHESIS OF THE FSF CFIT AND ALAR TASK-FORCE LITERATURE.

REFERENCES

1. Flight Safety Foundation, "Killers in Aviation: FSF Task Force Presents Facts About Approach-and-landing and Controlled-flight-into-terrain Accidents," Flight Safety Digest (1998–1999).

2. Flight Safety Foundation, ALAR Tool Kit (1998) — distributed multi-element package of briefing notes, training aids, and risk-assessment instruments.

3. Khatwa & Helmreich, "Analysis of critical factors during approach and landing in accidents and normal flight," Flight Safety Digest (1998) — the analytical basis of the ALAR Task Force scope.

4. FAA, Terrain Awareness and Warning System (TAWS) Final Rule, 14 CFR Part 121 (2000, effective 2002).

5. ICAO Annex 6, TAWS requirement (effective 2007) — the international consolidation.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The FSF CFIT and ALAR task forces are the canonical case of industry-level institution building after a catastrophe-class spike. An independent foundation convened the cross-operator working bodies, released structured diagnostic tools without competitive restriction, and sustained momentum to a regulatory mandate. The accident-rate decline is multifactorial; the institutional form is what the case is teachable on.

LENS APPROACH

FSF CFIT/ALAR is the canonical industry-coordination case (induced 6.1; LENS D5/PT4) — Domain 5 for the independent- convening-body form; Domain 2 for the structured diagnostic instruments as iterating deliverables. Pair with Cases 117, 119, and 13.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION SAFETY	G K N	D1 D2 D3 D4 D5	4	6.1	5
INDUSTRY COORDINATION	→	→			
TRAINING		Navigating Sociotechnical Constraints			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Convene the cross-operator working body before the regulator does — an independent industry foundation, no competitive stake in any one operator’s adoption — so the diagnostic tool can be released without restriction and adopted across the field.
- ▶ Design the diagnostic as a structured self-assessment, not a benchmark league table; operators will use what they can apply privately, and the field-level signal aggregates from voluntary adoption.
- ▶ Build the tool kit around the phase of flight that accident analysis says dominates the harm — approach-and-landing in this case — rather than around evenly distributed coverage that no single operator can adopt fully.

IN OPERATION / AFTER

- ▶ Sustain the institutional momentum into the regulatory mandate; the task forces did the technical and political work that made the TAWS rule adoptable, and the rule consolidates what the field already does.
- ▶ Report the accident-rate decline with the multifactorial hedge intact; TAWS hardware, training, procedure, and fleet turnover all moved together, and isolating the task-force contribution overstates the available evidence.
- ▶ Treat the institutional form as the transferable artifact: an independent convening body, cross-operator diagnostic tools without competitive restriction, coordinated path to regulatory mandate. The form pairs with CAST (Case 117) and ASRS (Case 119) at the industry-coordination layer.

REFLECTION QUESTIONS

1. Identify a catastrophe-class failure pattern in your domain whose response has been operator-by-operator rather than industry-coordinated. What would the analog of an independent convening body look like, and which body could plausibly play that role without competitive stake?
2. Specify the cross-operator diagnostic — checklist, tool kit, structured self-assessment — that you would design as the first deliverable of an FSF-style task force. The deliverable has to be applicable privately by each operator, and aggregate into field-level signal.

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

Chapter 7

Defense & National Security — What Fails

When the capability the mission requires goes unspecified at the point of contact.

The system was fielded. The operator's task was not.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

F-35 Sustainment & Maintainer Shortage

T Training Gap **K** Knowledge & Institutional Memory **D** Designed Out

IMPACT Fleet mission-capable rate about 55% (March 2023), far short of program goals; lifecycle cost exceeds \$1.7T, with ~\$1.3T in operating and support; maintainer, technical-data, and depot shortfalls are the binding readiness constraint (GAO-23-105341)

IN BRIEF

The F-35 is the most expensive weapons program in history — roughly 2,500 jets planned and a lifecycle cost above \$1.7 trillion, about \$1.3 trillion of it in operating and support. The hard part was never the airplane but keeping a global fleet ready, and that half has lagged from the start. As of March 2023 the fleet's mission-capable rate was about 55%, far short of goal, with more than 10,000 components in the repair queue and depots behind schedule. GAO traced the shortfall to maintainer shortages, the military's lack of access to technical data, and contractor dependency, and urged a full reassessment of the sustainment strategy. It is the book's cleanest case of a platform fielded faster than the capability infrastructure to sustain it.

BACKGROUND

The F-35 is the most expensive weapons program in history: the Pentagon plans to buy nearly 2,500, at a lifecycle cost exceeding \$1.7 trillion — roughly \$1.3 trillion of it not the aircraft but the decades of operating and sustaining them.^o The flyaway jet is the finite part; keeping a global fleet ready — maintainers, technical data, depots — is the open-ended part, and the part that lagged, because the cost that dominates the program is not buying the aircraft but the decades of sustaining them, the very work that received the least attention as the jets rolled off the line.

WHAT HAPPENED

The failure is a standing condition. As of March 2023 the fleet's mission-capable rate was about 55%, far short of goal; more than 10,000 components waited in the repair queue, and the depots averaged about 72 days per repair while still behind schedule in standing up the capacity to do the work at all — a backlog and a turnaround time that compound, since parts stuck in the queue keep jets grounded and the under-built depots cannot clear the queue fast enough to recover the mission-capable rate.¹

THE INVESTIGATION

GAO's September 2023 review was bluntly titled: the Department and services **need to reassess the future sustainment strategy**. It traced the shortfall to maintainer shortages, the military's lack of access to the technical data needed to do its own repairs, and the


resulting dependence on the prime contractor.² None of it was new — GAO has repeated the same diagnosis year after year, through a troubled logistics–software backbone and slow progress, even as procurement continued and readiness stayed flat, so the program kept buying more jets it could not fully sustain while the same three shortfalls were named in review after review without being closed.³

THE CAPABILITY GAP

The F-35 is the cleanest modern case of a platform fielded faster than the capability infrastructure to sustain it. Aircraft arrived on schedule; the maintainers, technical data, and depots were treated as follow-on costs rather than deliverables that had to field with the jets. The hardest part is the data: much of what is needed to repair the aircraft stayed controlled by the contractor, so the services cannot freely write procedures, qualify depots, or compete the work, which locks the fleet into a single source for the knowledge needed to keep it flying — merely expensive in peacetime, dangerous in war, when a contractor-dependent sustainment chain is exactly the kind of bottleneck an adversary would seek to exploit.⁴

AFTERMATH & REFORM

GAO urged a full reassessment of the sustainment strategy — government access to technical data, depot capacity, and a maintainer pipeline — rather than more patching.⁵ Those recommendations remain a work in progress, with later reviews showing costs still rising and readiness still below goal. The F-35 sits in this book as the live argument for treating capability infrastructure — people, data, the means to sustain them — as a fielding gate, not an afterthought. The bill for skipping that gate does not disappear; it compounds, arriving as grounded jets, a swelling repair queue, and a sustainment chain the services cannot control rather than as a line item caught early enough to fix cheaply.



Organizational-level maintenance has been affected by a number of issues, including a lack of technical data and training.

PARAPHRASING GAO-23-105341, F-35 AIRCRAFT, 2023

REFERENCES

1. U.S. Government Accountability Office, F-35 Aircraft: DOD and the Military Services Need to Reassess the Future Sustainment Strategy, GAO-23-105341 (Sept. 2023) — 2,500 planned air-
4. GAO, F-35 Sustainment: DOD Faces Several Uncertainties..., GAO-22-105995 (2022), and the broader GAO F-35 series — the recurring, year-over-year diagnosis.

- craft and a lifecycle cost exceeding \$1.7 trillion, \$1.3 trillion of it in operating and support.

2. GAO-23-105341 (2023) — 55% fleet mission-capable rate (March 2023); over 10,000 components awaiting repair; 72-day average depot turnaround; depot stand-up behind schedule.

3. GAO-23-105341 (2023) — sustainment shortfall traced to maintainer shortages, lack of military access to technical data, and contractor dependency.
5. GAO-23-105341 (2023) — recommendation that DOD reassess the future sustainment strategy.

6. GAO, F-35 Sustainment: Costs Continue to Rise While Planned Use and Availability Have Decreased, GAO-24-106703 (2024) — costs rising while readiness stays below goal.

THE LEARNING ENGINEERING LENS

LE INSIGHT

F-35 is the live, current example of fielding a platform faster than its capability infrastructure can be built. The aircraft is the easy part; the maintainers are the hard part. A decade of program schedules treated maintainer training as a follow-on cost, not a fielding gate. The fleet is now operating at half of its design readiness, and the program of record is more than a trillion dollars over its 2018 estimate.

LENS APPROACH

LENS treats the F-35 in LEN 5 as the canonical case of **Capability- System Misalignment at Transition**, the program’s first canonical problem type. Students design the capability infrastructure that should have accompanied each lot of aircraft delivered, including the training pipeline, technical-data deliverables, and depot establishment.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
DEFENSE →	T K D →	D1 D2 D3 D4 D5	1	1.4	1
Systems Analysis					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Make sustainment infrastructure — maintainers, technical data, depot capacity — a contractual deliverable that fields with each lot of aircraft, not a follow-on cost.
- Secure government rights to the technical data needed to write procedures, qualify depots, and compete the work, so the fleet is not locked to a single source.
- Gate fielding on demonstrated sustainment capacity, not just aircraft delivery, so readiness is engineered alongside the platform.

IN OPERATION / AFTER

- Audit the mission-capable rate against repair-queue depth and depot turnaround, and act when the backlog signals a structural, not transient, shortfall.
- Track recurring GAO-style diagnoses to closure, refusing to let procurement outpace the sustainment fixes named year after year.
- Monitor contractor dependency as a wartime risk, and sustain a government maintainer pipeline and depot capacity that can hold under pressure.

REFLECTION QUESTIONS

- Pick a current technology platform in your domain. Estimate the capability infrastructure that must field with it. What happens if half of that infrastructure is years behind the hardware?
- The F-35 program treated maintainer training as follow-on cost. Design a fielding gate that would prevent that decision being available to a future program manager.
- Much of the data needed to repair the F-35 stayed controlled by the contractor, foreclosing the government’s ability to compete or qualify the work. For a platform in your domain, what data rights would you make a contractual deliverable up front, and how would you verify they were actually delivered?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · task analysis · design prototyping

Military Fratricide — Desert Storm to Afghanistan

T Training Gap **H** Human-System Interface **K** Knowledge & Institutional Memory

IMPACT 24% of U.S. KIA in Desert Storm from friendly fire (35 of 146) — well above the historical baseline

IN BRIEF

Friendly fire killed an unusual share of coalition forces in the 1991 Gulf War: 35 of 146 U.S. combat deaths (24%) and 72 of 467 wounded (15%). (The often-cited “2% historical baseline” from Shrader’s 1982 study is contested — later estimates run nearer 10–15%, and Shrader stepped back from it.) Post-war reviews blamed the chaos of combat, weak situational awareness and fire-control discipline, and combat-identification failures — and noted the military lacked a shared record to even study its own pattern. Fratricide is the failure of several systems to integrate; despite a generation of combat-ID investment, recurrences in Afghanistan and after show the rate never fell to a confidently low level, because integration is the hardest property to engineer by program.

BACKGROUND

Friendly fire is as old as war, but its true rate is hard to pin down. The most-cited estimate, from Charles Shrader’s 1982 study *Amicicide*, put it under 2% of battle casualties — a figure later analysts challenged as far too low (nearer 10–15%), and one Shrader himself stepped back from. The disputed baseline mattered because it became the yardstick against which a modern war would be measured — and a yardstick set too low makes any later rate look like a catastrophe.⁰

WHAT HAPPENED

The 1991 Gulf War made the question grim and concrete: of 146 U.S. service members killed in action, 35 — about 24% — died by friendly fire, and 72 of 467 wounded (15%) were hit by their own side.¹ An A-10 strike on U.S. LAV-25s near Khafji killed seven Marines; an A-10 attack on British Warrior vehicles killed nine — each an aircraft firing on friendly vehicles it had failed to identify, the recurring shape of the problem. Whatever the true baseline, a quarter of American combat deaths from one’s own forces could not be waved off as the ordinary friction of battle.²

THE INVESTIGATION

Post-war reviews converged on familiar causes: the chaos of combat, inadequate situational awareness, weak adherence to fire-control measures, and combat-identification failures.³ One finding cut deeper — the military lacked a comprehensive, shared record of fratricide


incidents, so it could not study its own pattern, separate training failures from doctrine or equipment, or tell whether a fix was working. Without a common database, every incident was investigated in isolation and the aggregate signal that might have driven reform never formed. The capability to **learn** was itself missing, a second-order gap beneath the first.⁴

THE CAPABILITY GAP

Fratricide is not one problem but the failure of several systems to integrate — situational awareness, fire-control discipline, combat identification, joint coordination — each the subject of dedicated programs, yet the rate never fell to anything negligible. In Afghanistan a satellite-guided strike went wrong after a controller's device reset its coordinates, and in 2014 a B-1's targeting pod could not detect the infrared strobes marking U.S. troops, killing five. Each contributor was worked on; the integration **across** them — which is what actually keeps friendly forces from killing each other — is the hardest thing to engineer by program, because no single office owns it and no single procurement can deliver it — and it is where the capability kept failing.⁵

AFTERMATH & REFORM

The response was a generation of combat-identification investment: better IFF systems, blue-force-tracking networks, and changes to fire-control doctrine and training. The improvements were real, yet fratricide never dropped to a confidently low, stable rate, and even measuring it remained contested.⁶ That is the lesson: where capability is an emergent property of many systems working together, no single program closes the gap — improving each contributor in isolation still leaves the integration between them unaddressed — and progress has to be measured against an honest baseline rather than a convenient one.



The lack of a comprehensive and accessible automated database prevented thorough examination of the problem.

PARAPHRASING POST-DESERT STORM FRATRICIDE REVIEWS, C. 1993

REFERENCES

1. C. R. Shrader, *Amicide: The Problem of Friendly Fire in Modern War* (U.S. Army Combat Studies Institute, 1982) — the under-2% estimate from 269 incidents, a baseline later challenged as too low (with estimates nearer 10–15%).
2. "Friendly Fire: Facts, Myths and Misperceptions," U.S. Naval Institute Proceedings (June 1994) — Desert Storm: 35 of 146 KIA
3. [Unlabeled reference]
4. [Unlabeled reference]
5. Post-war combat-identification reviews — leading causes (situational awareness, fire-control measures, combat ID) and the absence of a comprehensive incident database. (Synthesized from secondary analyses; see AUDIT.)
6. Later-conflict recurrences (2001 coordinate-reset strike; 2014 B-1 infrared-strobe strike) and S. Snook, *Friendly Fire* (Princeton

- (24%) and 72 of 467 wounded (15%) by friendly fire; critique of the 2% “historical norm.”

3. Khafji and Warrior fratricide incidents (Feb. 1991) — see USNI Proceedings (1994) and Time, “They Didn’t Have to Die”. (Per-incident casualty figures vary across sources; see AUDIT.)

4. U.S. GAO, Operation Desert Storm fratricide investigations — Apache incident (OSI-93-4) and Army fratricide investigation (OSI-95-10) — causes and the Army’s reviews.
- Univ. Press, 2000) on the 1994 Black Hawk shootdown — the systems-integration archetype.

7. “IFF Update: Stalled Again,” U.S. Naval Institute Proceedings (June 1994) — the combat-identification and identification-friend-or-foe programs pursued after Desert Storm, and how slowly they matured.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Fratricide is a multi-decade capability problem that resists single- intervention solutions. Each of the contributing causes — situational awareness, fire-control discipline, combat identification, joint coordination — has been the subject of dedicated programs. The rate persists because the integration across the contributors is itself the capability, and integration is the hardest property to engineer by program.

LENS APPROACH

LENS uses fratricide in LEN 5 to teach systems-of-systems capability analysis and in LEN 2 to introduce combat identification as a Human-AI Teaming problem (operators rely on automated IFF systems whose limitations are not in their training). LEN 8 examines why decades of awareness have not produced a sustained reduction.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
DEFENSE	→ T H K	→ D1 D2 D3 D4 D5	1	1.3	1
Systems Analysis					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat combat identification as a systems-of-systems requirement with a single owner for the integration, not as separate IFF, situational-awareness, and fire-control programs.
- ▶ Build a comprehensive, shared incident database from the outset so the pattern across events can be studied and fixes tested against it.
- ▶ Design fire-control discipline and identification into joint coordination, so the integration that prevents fratricide is engineered rather than left emergent.

IN OPERATION / AFTER

- ▶ Audit the friendly-fire rate against an honest baseline, distinguishing a temporary dip from a structural reduction before crediting any program.
- ▶ Monitor each contributor and the integration across them, since improving components in isolation can leave the joint failure mode untouched.
- ▶ Sustain the learning channel so later-conflict recurrences feed back into doctrine and training rather than being investigated in isolation.

REFLECTION QUESTIONS

1. Identify a capability problem in your domain that has been the subject of repeated programs without sustained improvement. What is the integration gap that the programs do not address?
2. Design a measurement system that would distinguish a temporary improvement in fratricide rate from a structural one.
3. The military lacked a shared incident database, so it could not study its own pattern. Identify a recurring harm in your domain that is investigated case-by-case but never aggregated, and design the shared record that would let the pattern become visible.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · usability testing · cognitive walkthroughs · knowledge bases · data pipelines

Operation Eagle Claw

T Training Gap **K** Knowledge & Institutional Memory

IMPACT 8 servicemembers killed in Iran; mission to rescue 53 American hostages aborted; catalyst for the founding of U.S. Special Operations Command

IN BRIEF

The April 1980 mission to rescue fifty-three American hostages held at the U.S. embassy in Tehran was aborted at a desert staging point when three of eight helicopters were lost to mechanical and weather problems. During the withdrawal a helicopter collided with a refueling C-130, killing eight servicemembers; the hostages remained captive for another nine months. The Holloway Commission found the operation had been assembled ad hoc: each service contributed its own units, equipment, and command relationships, the aircrews had not trained together, and there was no standing joint special-operations command to own the mission. The capability had to be improvised, and the improvisation failed. Eagle Claw catalyzed the creation of JSOC, the Goldwater-Nichols reforms, and ultimately U.S. Special Operations Command.

BACKGROUND

In November 1979, Iranian revolutionaries seized the U.S. embassy in Tehran and took more than sixty Americans hostage; fifty-three remained by April 1980. After months of failed diplomacy, the Carter administration authorized a military rescue. No standing joint special-operations command existed to plan or execute it; the force had to be assembled from units drawn separately from the Army, Navy, Marines, and Air Force — each bringing its own equipment, procedures, and chain of command to a mission that demanded they act as one, with no institution whose job it was to make them cohere before the night of the raid.^o

WHAT HAPPENED

On 24 April 1980, the mission staged at a remote site code-named Desert One. Three of the eight RH-53 helicopters were disabled by a dust storm and mechanical failure, dropping the force below the minimum needed; the commander aborted. During the withdrawal a helicopter collided with a C-130 tanker, and the resulting fire killed eight servicemembers. The mission failed before reaching Tehran — undone not by enemy action but by attrition and a chaotic withdrawal among aircraft and crews that had never operated together, the predictable cost of integration improvised under pressure rather than built and rehearsed in advance.¹

THE INVESTIGATION

The Holloway Special Operations Review Group identified the operation's "ad hoc-ery" as central: each service contributed its own units, equipment, command relationships, and communications; the aircrews had not trained together as a unit; the RH-53D had been

selected partly for a minesweeping cover story rather than for fitness for a desert rescue. There was no standing organization to own the mission end to end — no single authority responsible for the force’s training, equipment fit, and command architecture as a whole, so each gap was someone’s problem in part and no one’s in full.²

THE CAPABILITY GAP

Each service was competent inside its own boundary. The capability that did not exist was the integration across them — a joint command, a common communications architecture, and a force that had trained together. That integration had to be improvised for a single high-stakes mission, and the improvisation could not hold — because cross-service cohesion is not summoned on demand but accrued through standing structure and repeated joint training, neither of which existed when the order came.³

AFTERMATH & REFORM

Eagle Claw produced the Holloway Commission, the creation of the Joint Special Operations Command in 1980, and — alongside the 1983 Beirut bombing and the Grenada invasion — the impetus for the Goldwater-Nichols Act of 1986 and the Nunn-Cohen Amendment establishing U.S. Special Operations Command in 1987. The reform built the institution the mission had needed and not had — converting a capability that had been improvised once and failed into a standing command with its own forces, training, and authority, so the next mission would inherit cohesion rather than assemble it from scratch.⁴



The mission was ad hoc — assembled from units, equipment, and command relationships that had never operated together.

PARAPHRASING THE HOLLOWAY SPECIAL OPERATIONS REVIEW GROUP, 1980

REFERENCES

1. Holloway Special Operations Review Group, Rescue Mission Report (1980) — the ad-hoc assembly and equipment choices (paraphrased).
2. Holloway Commission (1980) — the Desert One sequence and the helicopter-C-130 collision.
3. Goldwater-Nichols Department of Defense Reorganization Act of 1986, Pub. L. 99-433.
4. Nunn-Cohen Amendment (1986), establishing U.S. Special Operations Command (1987).
5. Locher, J. (2002), Victory on the Potomac — the reform arc from Desert One to Goldwater-Nichols.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Eagle Claw is the canonical case in U.S. defense for the absence of institutionalized cross-service capability. Each service was competent inside its own boundary; the integration across them did not exist as a deliverable — until the reform that followed built the institution the mission had needed and not had.

LENS APPROACH

LENS uses Eagle Claw in LEN 5 as a worked case for cross- organizational capability requirements and in LEN 8 for the institutional reform that followed; it pairs with INPO (Case 175) as the defense analog of failure-driven institution building. The seven-year arc to Goldwater-Nichols (1986) and USSOCOM (1987) is itself a measurement: the institutional response time when the fix requires statutory action.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
DEFENSE →	T K →	D1 D2 D3 D4 D5	1	1.3	1
Systems Analysis					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Stand up a single command that owns the cross-organizational mission end to end — its force, equipment fit, communications, and training — rather than assembling it per mission.
- ▶ Require the contributing units to train together as one force on a common communications architecture before they are committed.
- ▶ Select equipment for fitness to the actual mission, not for a convenient cover story or parent-service availability.

IN OPERATION / AFTER

- ▶ Audit whether a cross-organizational capability that exists on paper has ever actually operated as one force, and treat the absence of joint reps as an unfilled gap.
- ▶ Sustain the standing institution and its training pipeline so the next mission inherits cohesion rather than improvising it.
- ▶ Monitor for the recurrence of ad-hoc assembly, since the conditions that produced Desert One reappear whenever no single authority owns the joint mission.

REFLECTION QUESTIONS

1. Where in your domain does a cross-organizational capability exist on paper but not in practice? What would force its institutionalization?
2. Eagle Claw produced USSOCOM and Goldwater-Nichols six years later. Sketch the institutional design that an equivalent failure in your domain would force into existence.
3. The Holloway Commission named the mission’s ad-hoc-ery as the diagnosis. What standing capability — institution, command, training pipeline — does your domain currently lack that an Eagle-Claw-class failure would force into existence?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · knowledge bases · data pipelines

USS Vincennes & Iran Air Flight 655

H Human-System Interface **T** Training Gap

IMPACT 290 civilians killed — one of the deadliest shootdowns of a commercial airliner by a military force; precipitating case for Aegis CIC display and decision-aid doctrine reform, and a central reference in the human-AI teaming literature

IN BRIEF

On 3 July 1988, during a surface skirmish with Iranian gunboats in the Persian Gulf, the USS Vincennes shot down Iran Air Flight 655 — a civilian Airbus A300 climbing on a scheduled route — killing all 290 aboard, one of the deadliest airliner shootdowns by a military force. The Aegis combat system worked: its radar correctly showed the aircraft ascending on the published civilian-traffic corridor. Yet operators in the Combat Information Center, primed by a simultaneous surface fight to expect a hostile inbound, tagged the contact as a military F-14 — from an IFF Mode II squawk mis-correlated from an aircraft at Bandar Abbas, not the airliner, which squawked only civilian Mode III — and, misreading its closing range as a descent, fired. The Fogarty Report (DoD, 1988) attributed the tragedy to human error under extreme stress — confirmation bias, “scenario fulfillment,” and the unconscious distortion of data — not equipment failure. The 1992 Newsweek “Sea of Lies” investigation and a 2018 US Naval Institute Proceedings retrospective reopened both the operational record and the interface-design lessons. Vincennes is the book’s foundational human-AI-teaming case: the most advanced surface combatant afloat failed because its CIC interface and the operational framing made the correct reading possible in principle and unsustainable in practice under combat stress.

BACKGROUND

In the closing weeks of the Iran-Iraq War, US warships patrolled a tense Persian Gulf where civilian airliners and hostile military aircraft shared the same crowded sky. The USS Vincennes, a Ticonderoga-class cruiser with the Navy’s most advanced Aegis combat system, was in a surface fight with Iranian gunboats — itself a contested engagement that the later Newsweek “Sea of Lies” investigation argued had taken the ship into Iranian territorial waters — when an aircraft took off from Bandar Abbas, an airfield used by both civil and military traffic.⁹ The dual-use field meant a single track could be either threat or scheduled flight, so the crew’s reading of the contact carried the whole burden of identification — and the Combat Information Center was already managing divided attention against the surface engagement and against an interface whose track display did not directly show altitude trend, the kind of conditions under which a wrong call becomes far easier to make than a right one.

WHAT HAPPENED

The aircraft was Iran Air Flight 655, a civilian Airbus A300 climbing through 12,000 feet on its scheduled route to Dubai. The Vincennes' crew identified it as a descending, hostile F-14 and fired two SM-2 surface-to-air missiles; all 290 people aboard, including 66 children, were killed — one of the deadliest shootdowns of a commercial airliner by a military force.¹ An ascending airliner and a diving attack jet are opposite behaviors, yet the crew converged on the second while the Aegis radar reported the first. The engagement consumed only seven minutes from radar detection to launch, collapsing identification, IFF interrogation, and decision into a window too narrow for anyone to slow down and reconcile the contradiction.

THE INVESTIGATION

The Aegis SPY-1A radar had not malfunctioned: it correctly showed the aircraft ascending. The IFF interrogation system added a structural confusion of its own: the Vincennes' IFF returns reflected both the airliner's Mode III civilian transponder and, apparently, a Mode II military code from another aircraft on the ground at Bandar Abbas — a mode-confusion failure that left the crew with ambiguous identification at the worst possible moment.² The Fogarty Report (Rear Adm. William Fogarty, DoD, August 1988) attributed the tragedy to human error under extreme stress — confirmation bias and the “stress and unconscious distortion of data” — naming “scenario fulfillment” as the psychological mechanism by which operators read every indication through the lens of a presumed hostile attack. An inexperienced operator misread the range as a descent and others in the CIC accepted it rather than checking,³ a shared error that is more damning than a single misread because it shows the interface offered no cross-check strong enough to break a wrong reading once the team had settled into it. The framing arrived before the data did, so each new return was fitted to the expected threat rather than weighed against it.⁴


THE CAPABILITY GAP

The system did not lie, and the operators did not act in bad faith. The interface and the operational framing combined to make a particular misreading not merely possible but likely. Correct performance was possible in principle and unsustainable in practice — and the gap between those two is the engineering problem. An Aegis CIC display that surfaced altitude as a numeric readout but not as a trend, that gave equal weight to civilian and military IFF returns under mode confusion, and that offered no mandatory disconfirmation step before weapons release, had not been designed for the stress it would actually meet. The burden of overriding a presumed hostile attack was left entirely to the operator's discipline, exactly when combat had stripped that discipline of the time and calm it needed to work.⁵

AFTERMATH & REFORM

The Navy's post-incident response included revised tactical doctrine for civilian-traffic deconfliction in the Gulf, updated Aegis training scenarios that explicitly rehearsed the scenario-fulfillment failure mode, and changes to the CIC display sequence and crew procedures around IFF interpretation under mode confusion. The case has remained a standing reference in the human-AI-teaming literature: a 2018 US Naval Proceedings retrospective placed it at the heart of how operators should be teamed with automated decision aids under stress, and the situation-awareness, naturalistic-decision-making, and automation-trust literatures (Klein, Endsley, Cummings) treat it as the canonical worked example.⁶ Its lesson is that interface design is a capability deliverable, not an aesthetic one — and that

the most advanced system afloat is only as good as the human reading it under fire. The retrospective reframed the loss not as a one-off error but as a predictable outcome of teaming a person with a decision aid that displayed truth without defending it, a pattern that recurs wherever automation is fast and the human is the last check.



The shootdown of Flight 655... reveals lessons for technology adoption and its use in stressful situations.

PARAPHRASING U.S. NAVAL PROCEEDINGS RETROSPECTIVE ON THE VINCENNES INCIDENT, 2018

REFERENCES

1. Rear Adm. W. Fogarty, Formal Investigation into the Circumstances Surrounding the Downing of Iran Air Flight 655 (US Navy, August 1988) — the engagement, the IFF mode-confusion findings, and the 290 deaths.
2. Fogarty report (1988) — the Aegis SPY-1A radar functioned; the aircraft was ascending while the crew perceived a descent; “scenario fulfillment” as the psychological mechanism.
3. Fogarty report (1988) — “human error under extreme stress,” confirmation bias, and “unconscious distortion of data” (quoted); shared-error finding across CIC operators.
4. J. Barry & R. Charles, “Sea of Lies,” *Newsweek* (July 13, 1992) — contemporaneous reinvestigation of the operational record, including the disputed account of Vincennes’ position relative to Iranian territorial waters.
5. US Naval Institute Proceedings retrospective on the Vincennes incident (2018) — the human-AI-teaming reframing and the standing case-study role.
6. M. L. Cummings, “Human Supervisory Control of Weapon Systems” (MIT) — interface design and automation under time pressure as the engineering frame for the case.
7. G. Klein, *Sources of Power* (1998); M. Endsley, “Toward a Theory of Situation Awareness” (1995) — the naturalistic-decision-making and situation-awareness literatures that treat Vincennes as the canonical worked example.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Vincennes is the canonical case for human-AI teaming under stress. The system did not lie. The operators did not act in bad faith. The interface and the operational framing combined to make a particular misreading not just possible but likely. Correct performance was possible in principle and unsustainable in practice — and that gap is the engineering problem.

LENS APPROACH

Vincennes is the foundational mode-and-state-transparency case under combat stress (induced 3.3; LENS D2/PT6 human-AI teaming). LENS uses it in Domain 2 (Iterative Development; CLO-2) for interface-as-deliverable: a CIC display tested against the worst case it would actually meet — confirmation bias during a simultaneous surface engagement — not against the calm of acceptance testing. LENS uses it in Domain 3 (Human-System Collaboration Adaptation; CLO-3) for the human-AI teaming problem itself: a decision aid that displays truth without defending it predictably fails wherever automation is fast and the human is the last check. The case sits at the heart of the program’s argument that interface design is a

capability deliverable, not an aesthetic one. Pair with EHR/CPOE (Case 2) at the interface-built-to-wrong-specification layer.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
DEFENSE	→ H T	→ D1 D2 D3 D4 D5	6	3.3	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Design the contact display to make ascent-versus-descent unmistakable at a glance, so the cue that contradicts a presumed threat cannot be passed over under time pressure. ▶ Test the interface against the worst case it will actually meet — confirmation bias during a simultaneous surface engagement — not against the calm conditions of acceptance testing. ▶ Build a mandatory disconfirmation step into the engagement sequence so identifying a hostile track requires actively ruling out the civilian one.	▶ Audit live engagements and exercises for cases where the crew's reading diverged from sensor data, treating each as a near-miss that exposes an interface gap. ▶ Train and drill operators specifically on the stress regime — divided attention, presumed-hostile framing — rather than only on nominal track identification. ▶ Track whether decision aids are trusted past their evidence, and feed that signal back into interface and procedure revisions.

REFLECTION QUESTIONS

1. Identify a high-stress interface in your domain. What framing arrives before the data, and how does it shape what operators see?
2. Vincennes' operators acted under tunnel vision. Design the procedural intervention that would have forced one of the five to call out the contradiction.
3. The Aegis radar reported the truth and the crew still read its opposite. What would an interface have to make impossible — not merely visible — to keep a pre-formed expectation from overriding correct data under stress?

WHO BUILDS THIS human factors & interaction design · learning science & instructional design — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · scenario simulation · competency models

F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap

H Human-System Interface **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Between 2008 and 2012 USAF F-22 Raptor pilots reported a cluster of physiological events consistent with hypoxia; one fatal accident (Capt. Jeffrey Haney, Nov 2010) was attributed in part to operator response to a bleed-air shutoff; the fleet was stood down in 2011, and the USAF Scientific Advisory Board found no single root cause but a contributing oxygen-delivery issue rooted in the On-Board Oxygen Generation System (OBOGS) and associated equipment

IN BRIEF

Between 2008 and 2012 the USAF accumulated a cluster of F-22 Raptor “physiological events” — pilot reports consistent with hypoxia, often with persistent post-flight symptoms (chronic cough, disorientation). One fatal class-A accident in November 2010 (Capt. Jeffrey Haney) was attributed in part to operator response after a bleed-air shutoff. In May 2011 the USAF stood the fleet down for several months. The USAF Scientific Advisory Board’s investigation, with NASA contribution, found **no single root cause** — the load-bearing finding of the case. The set of contributing factors included the On-Board Oxygen Generation System (OBOGS) and its associated equipment (regulators, hoses, upper-pressure-garment vest), aircrew flight equipment fit, and the absence of in-cockpit physiological monitoring that would have permitted earlier detection. The case is the canonical recent instance of a sustainment-era failure whose diagnosis was protracted because the system as fielded had no instrumentation of the variable that mattered — pilot blood oxygenation — and the failure mode therefore had to be inferred from pilot self-report and incident clustering.

BACKGROUND

The F-22 Raptor entered operational service in 2005. Like most modern fighters it uses an On-Board Oxygen Generation System (OBOGS) rather than carrying liquid oxygen — a molecular sieve separates oxygen from engine bleed air to supply the pilot’s breathing mask. Beginning in 2008, USAF Raptor pilots reported a cluster of “physiological events”: symptoms consistent with hypoxia during flight, often with persistent post-flight effects — chronic cough, disorientation, exercise intolerance. The cluster did not conform to a single platform or sortie profile, and the reports accumulated faster than the diagnostic process could resolve them.^o

WHAT HAPPENED

The forcing event was the loss of Capt. Jeffrey Haney and his aircraft in November 2010 over Alaska. The Air Force Accident Investigation Board found that a bleed-air leak had triggered an automatic bleed-air shutoff, which in turn cut OBOGS supply, and that the accident sequence involved pilot response under conditions consistent with hypoxia. The board's findings were contested by the pilot's family and reviewers; what the accident crystallized was that the platform's response to an OBOGS cutoff placed an unreasonable burden on pilot recognition and reaction in a window the platform did not instrument. In May 2011 the USAF stood the entire F-22 fleet down for several months pending investigation.¹

THE INVESTIGATION

The USAF Scientific Advisory Board investigation, with participation from NASA's review of the hypoxia incidents, reported the load-bearing finding the case turns on: there was **no single root cause**. The contributing factors included the OBOGS system and its associated equipment (regulators, hose configurations, the upper-pressure garment vest that under some conditions inhibited normal breathing), aircrew flight-equipment fit issues, and the absence of in-cockpit physiological monitoring. The investigation made several corrective recommendations — including modifications to the upper-pressure garment, changes to OBOGS components, and the addition of physiological monitoring during flight — and the symptom reporting decreased substantially in the period after implementation.²

THE CAPABILITY GAP

What makes this a sustainment-era case rather than a design-era one is the instrumentation gap. The F-22 as fielded had no in-cockpit measurement of the pilot's oxygenation or of the actual gas composition reaching the mask. The variable that mattered most for the failure mode — was the pilot getting enough oxygen? — had to be inferred from pilot self-report after the fact. The diagnostic process took years partly because the population of events had to be characterized statistically from a noisy reporting channel, and partly because the multi-cause structure of the failure made any single intervention only partially testable against the population.³

AFTERMATH & REFORM

The hedge survives into the case verbatim. The USAF SAB explicitly declined to identify a single root cause, and the corrective actions were a bundle: garment modification, OBOGS component changes, training emphasis on hypoxia recognition, and the long-overdue addition of physiological monitoring. Symptom reports decreased after the bundle was implemented, but the attribution of the decline to any single component within the bundle is not possible from the available evidence. The case teaches the sustainment-instrumentation-gap form: a high-performance platform fielded without measurement of the operator-physiology variable that determines whether the mission can be completed, and a diagnostic process forced to work from self-report and clustering because the platform did not surface the data.

The variable that mattered most — was the pilot getting enough oxygen? — was the variable the platform did not measure.

EDITORS' SYNTHESIS OF THE USAF SAB AND NASA REVIEWS.

REFERENCES

1. USAF Scientific Advisory Board (2012), “Aircraft Oxygen Generation Study” — final report on F-22 physiological-event investigation.

2. NASA Engineering and Safety Center (2012), report contributing to the F-22 hypoxia review.

3. USAF Accident Investigation Board (2011), Capt. Jeffrey Haney F-22 accident report (Nov 2010).

4. GAO-12-789 (2012), “Actions Needed to Establish Effective Oversight of F-22 Pilot Physiological Issues” — congressional record.

THE LEARNING ENGINEERING LENS

LE INSIGHT

F-22 OBOGS is the canonical recent sustainment-instrumentation- gap case. A high-performance platform was fielded without in-cockpit measurement of pilot oxygenation, the diagnostic process took years against a noisy self-report channel, and the corrective action was an irreducibly multi-cause bundle. The “no single root cause” finding is the case.

LENS APPROACH

F-22 OBOGS is the aerospace instrumentation-gap case (induced 2.4; LENS D4/PT5). LENS uses it in Domain 4 (Test and Evaluation) for the no-single-root-cause finding and the bundle-of-corrections evaluation problem, and in Domain 3 (Human-System Collaboration) for the operator burden when an automatic action cuts a life-critical supply. Pair with Case 108 (EVA-23) as the human- spaceflight instrumentation-gap companion and with Case 27 (anesthesia monitoring) at the cue/alert-as- capability layer.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AVIATION DEFENSE	H K N	D1 D2 D3 D4 D5	5	2.4	4 · 3
HUMAN FACTORS		Test and Evaluation			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Instrument the operator-physiology variable that determines mission completion before fielding a high-performance platform. The F-22 OBOGS case is the worked example of the cost of fielding without it: years of self-report-based diagnosis and one fatality.
- ▶ When a failure-event cluster does not conform to a single platform or sortie profile, treat the absence of a single root cause as a finding, not as a failure of the investigation. The bundle of contributing factors is the diagnosis.
- ▶ Design the platform’s automatic responses (e.g., bleed-air shutoff) with the pilot’s recognition window in mind; an automatic action that cuts a life-critical supply has to be paired with cueing the pilot can act on in time.

IN OPERATION / AFTER

- ▶ Implement the corrective actions as a bundle — garment modification, OBOGS components, training, physiological monitoring — and report the population-level outcome decline as bundle evidence, not as evidence about any single component.
- ▶ Carry the “no single root cause” hedge into program communication; the multi-cause structure is the case’s teaching point and the basis for the sustainment-instrumentation argument.
- ▶ Treat the F-22 OBOGS experience as a fielded-platform sustainment lesson for subsequent platforms (e.g., F-35) — the analogous instrumentation question should be answered before the cluster appears, not after.

REFLECTION QUESTIONS

1. Identify a fielded high-performance platform in your domain where the variable that determines mission completion is not directly measured. What is the analog of in-cockpit physiological monitoring, and what would it cost to add at sustainment vs. having designed it in?
2. Specify what your program would treat as “no single root cause” — the irreducibly multi-factor finding — and how you would report bundle-of-corrections evidence rather than overstating attribution to any single change.
3. The F-22 cluster appeared in 2008 and stand-down came in 2011. What signal-detection threshold would your program use on a noisy self-report channel to act earlier on a population-level pattern?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** usability testing · cognitive walkthroughs · knowledge bases · data pipelines · incident analysis · safety dashboards

V-22 Osprey

T Training Gap **H** Human-System Interface **N** Normalization of Deviance

IMPACT ~65 killed across ~17 hull-loss accidents since 1991; serious-mishap rate above comparable fleets (GAO-26-108905, 2025); some fixes stretch to the 2030s

IN BRIEF

The V-22 Osprey — the tiltrotor flown by the Marines, Air Force, and Navy — has had about 17 hull-loss accidents and roughly 65 fatalities since 1991, including 19 Marines in a single 2000 test crash and 8 airmen off Yakushima, Japan, in 2023. The Yakushima crash traced to cracks in a transmission gear (a flaw in the X-53 steel alloy) and a pilot's decision to keep flying through warnings. In December 2025, GAO and a NAVAIR review found the V-22's joint program office had failed for years to assess and address mounting safety risks, that serious-mishap rates exceeded comparable fleets, and that some fixes stretch toward

2034. The Osprey is the steady-state form of normalization of deviance:

a documented, reviewed shortfall accepted as the cost of flying the airframe.

BACKGROUND

The V-22 Osprey tilts its rotors to take off like a helicopter and cruise like a turboprop — an ambitious capability shared awkwardly across three services and a joint program office. Since development began in the 1980s it has suffered about 17 hull-loss accidents and roughly 65 fatalities, including 19 Marines in a single crash during 2000 testing.⁹ The same configuration that makes the tiltrotor uniquely useful makes it uniquely demanding to sustain, and the joint arrangement diffused ownership of that burden across three services that each carried the airframe but none of which fully owned its safety trajectory.

WHAT HAPPENED

On 29 November 2023 a CV-22B broke up off Yakushima, Japan, killing all eight airmen aboard. The Air Force traced it to cracks in a transmission gear — rooted in metallic inclusions in the X-53 steel alloy used for the gears — and to the pilot's decision to keep flying despite warnings to land. The crash grounded Osprey fleets worldwide for months.¹ The materiel flaw and the human decision to press on through warnings were the same paired factors earlier reviews had named, so Yakushima read less as a novel failure than as one more draw from a pattern the program had already characterized but not closed.

THE INVESTIGATION

In December 2025 the Government Accountability Office and a NAVAIR review reported together that the joint program office had failed for years to adequately assess and address

mounting safety risks, even as service members died.² Serious-mishap rates generally exceeded those of comparable Navy and Air Force fixed- and rotary-wing fleets from FY2015 to FY2024 and spiked in 2023–2024; a gearbox flaw dating to 2006 was not evaluated until 2024, and full fixes for some issues are not expected until the 2030s.³ An eighteen-year gap between a gearbox flaw arising and its being evaluated is the timescale of normalization made literal: the deviation persisted long enough to become the airframe’s accepted background condition rather than an open defect demanding action.

THE CAPABILITY GAP

The V-22 is the steady-state form of normalization of deviance. The shortfall is not unknown — it has been documented across successive reviews — but the system built to remediate it has accepted its own incompleteness as the cost of operating the airframe. Each crash carries a familiar set of factors (materiel failure, human error, weak coordination across the three services), and each is followed by adjustments that do not converge.⁴ Convergence is the missing property: three services adjusting in parallel, none holding final authority over the whole, produce motion without resolution, so the mishap rate stays elevated while every individual response looks reasonable on its own terms.

AFTERMATH & REFORM

Groundings, gear-inspection regimes, and a redesign program have followed, with full fixes for some gearbox issues stretching toward 2034. The case is instructive precisely because the problem is recognized: a known, reviewed capability gap can persist for decades when each incident is treated as an isolated event rather than a sample from a distribution the program keeps drawing from.⁵ Fixes that stretch toward 2034 mean the airframe will keep flying for years inside a margin already judged worse than comparable fleets — the program in effect electing to operate the distribution it has been told it is drawing fatalities from.



Materiel failure and human-error factors were the most frequently cited causal factors in serious Osprey accidents.

PARAPHRASING THE NAVAIR INDEPENDENT REVIEW OF THE V-22, 2025

REFERENCES

1.

Compiled V-22 accident record — 17 hull losses and 65 fatalities since 1991, including the 2000 Marana test crash (19 Marines).
2.

U.S. Air Force Accident Investigation Board findings, via USNI News (Aug. 2024) — the 29 Nov. 2023 Yakushima CV-22B crash: transmission-gear cracks (X-53 inclusions) and continued flight despite warnings.
3.

U.S. GAO, Osprey Aircraft: Additional Oversight and Information Sharing Would Improve Safety Efforts, GAO-26-108905 (Dec. 2025) — the joint program office’s failure to assess and address risks.
4.

GAO-26-108905 (2025) — serious-mishap rates exceeding comparable Navy/Air Force fleets (FY2015–FY2024); the 2006 gearbox flaw evaluated only in 2024; fixes into the 2030s.
5.

NAVAIR independent review of the V-22 (Dec. 2025) — materiel and cross-service-coordination factors and unresolved catastrophic parts issues.
6.

USNI News, V-22 program coverage (2024–2025).

THE LEARNING ENGINEERING LENS

LE INSIGHT

V-22 demonstrates the steady-state version of normalization of deviance: a platform whose shortfall has been documented, reviewed, and acted on across multiple administrations without producing sustained improvement. The capability gap has itself been normalized. Each incident is treated as an event rather than as a sample from a distribution the program continues to draw from.

LENS APPROACH

LENS treats V-22 in LEN 5 as a multi-service capability-coordination problem and in LEN 8 as a study in long-cycle organizational learning failure. Students design the evidence system that would distinguish a true reduction in mishap rate from the natural variation of a chronically marginal platform.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
DEFENSE	→ T H N	→ D1 D2 D3 D4 D5	4	7.4	5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Assign one accountable owner for the airframe’s safety trajectory rather than diffusing it across three services with no final authority. ▶ Treat each known materiel flaw as an open defect on a clock, so a problem like the gearbox cannot quietly become accepted background condition. ▶ Define an absolute mishap-rate threshold against comparable fleets that triggers mandatory action, not merely review.	▶ Aggregate incidents as samples from one distribution, not isolated events, so the elevated rate itself becomes the thing being managed. ▶ Audit whether parallel service-level adjustments are actually converging on a lower rate, and escalate when they are not. ▶ Gate continued operation on remediation progress, so fixes stretching toward 2034 do not silently license years of degraded margin.

REFLECTION QUESTIONS

1.

Identify a platform or process in your domain that has been operating in a documented shortfall for years. What measurement would have to change for the shortfall to become unacceptable?
2.

The V-22’s three services do not converge on remediation. Design the governance structure that would force convergence.
3.

A gearbox flaw went eighteen years between arising and being evaluated. What mechanism in your domain converts a long-tolerated defect back into an open item that demands action?

WHO BUILDS THIS learning science & instructional design · human factors & interaction design · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · usability testing · cognitive walkthroughs · incident analysis · safety dashboards

9/11 Intelligence Sharing Failures

G Governance & Trust **K** Knowledge & Institutional Memory

IMPACT 2,977 killed; the 9/11 Commission found systemic intelligence-sharing failures across the U.S. government

IN BRIEF

The September 11, 2001 attacks, which killed 2,977 people, were enabled in part by intelligence the U.S. government already held but never integrated. The CIA knew, from a January 2000 meeting in Kuala Lumpur, that two future hijackers had entered the country; it did not tell the FBI. The FBI separately flagged suspicious flight-training activity in Phoenix and Minneapolis in 2001, but that information was never aggregated. Visa issuance, immigration tracking, and watch-listing were each run by a different agency, and the handoffs between them depended on individual initiative that no institution required. The 9/11 Commission called it a “failure of imagination” — a framing critics say understates the structural gap. The reform that followed built the cross-agency architecture, the ODNI and the NCTC, that had not existed.

BACKGROUND

By 2001, U.S. counterterrorism depended on many agencies — CIA, FBI, NSA, State, INS, FAA — each holding a piece of the picture. No institution was responsible for integrating those pieces; cross-agency information sharing depended on individuals choosing to pass information along rather than on any architecture that required it. Each agency’s incentives, classification rules, and turf reinforced the boundary, so the natural tendency was to hold information rather than to push it across a line no one was charged with bridging.^o

WHAT HAPPENED

On 11 September 2001, nineteen hijackers seized four aircraft and killed 2,977 people. In the months and years before, the warning signs had been distributed across the government: the CIA tracked two future hijackers from a January 2000 meeting in Kuala Lumpur but did not watch-list them or notify the FBI; FBI field offices flagged suspicious flight-training activity in Phoenix and Minneapolis; none of it was aggregated into a single picture. Any one fragment looked minor in isolation; only assembled would they have shown the shape of the plot, and assembly was exactly the function no one performed.¹

THE INVESTIGATION

The 9/11 Commission and the earlier Congressional Joint Inquiry documented specific failures of information sharing across the FBI, CIA, and NSA. The Commission famously concluded that “the most important failure was one of imagination” — a framing later criticized as understating the structural nature of the gap, which was less a lack of imagination than an absence of any institution responsible for integration. The distinction matters for the


remedy: a failure of imagination invites exhortation to think harder, while a structural gap demands an institution be built to close it.²

THE CAPABILITY GAP

The intelligence sharing did not happen because no institution owned it as a deliverable. Each agency was competent inside its own boundary; the integration across boundaries existed nowhere as a required function. The missing capability was an architecture — shared watch-lists, mandated handoffs, a body responsible for fusing the picture — rather than more raw collection. The government did not lack data; it lacked the connective tissue to turn data held in many places into a single picture anyone was accountable for assembling.³

AFTERMATH & REFORM

The Intelligence Reform and Terrorism Prevention Act of 2004 created the Office of the Director of National Intelligence and the National Counterterrorism Center, and fusion centers followed — the institutional architecture for cross-agency information sharing that had not previously existed. By creating bodies whose explicit mandate was integration, the reform converted information sharing from an act of individual initiative into a required function someone owned. The case is foundational in U.S. national-security policy for treating cross-agency capability as an engineerable institutional deliverable.⁴



The most important failure was one of imagination.

THE 9/11 COMMISSION REPORT, 2004

REFERENCES

1. National Commission on Terrorist Attacks Upon the United States, The 9/11 Commission Report (2004) — the quoted “failure of imagination” and the specific sharing failures.
2. 9/11 Commission Report (2004) — the Kuala Lumpur tracking and the Phoenix/Minneapolis flagging.
3. Joint Inquiry into Intelligence Community Activities Before and After September 11, 2001 (2002) — cross-agency information-sharing failures.
4. Zegart, A. (2007), Spying Blind — structural-organizational analysis of the failures.
5. Intelligence Reform and Terrorism Prevention Act of 2004 — creation of the ODNI and NCTC.

THE LEARNING ENGINEERING LENS

LE INSIGHT

9/11 is the foundational U.S. case for cross-agency information- sharing as an engineering deliverable. The architecture had not been built. The reform built it. The cost of the missing architecture was paid in 2001. The discipline LENS represents is the kind of work that, applied across the U.S. intelligence community in the 1990s, would have produced the architecture earlier.

LENS APPROACH

LENS uses 9/11 in LEN 8 as the foundational case for cross- organizational capability and in LEN 7 for the governance architecture of multi-agency systems. Studio projects compare 9/11 with Eagle Claw (Case 128) as institutional-architecture failures of different kinds.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
DEFENSE	→ GK	→ D1D2D3D4D5	1	5.3	5
GOVERNMENT		→ Navigating Sociotechnical Constraints			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Assign integration as an explicit deliverable owned by a named body, so the function does not depend on individuals choosing to pass information across a boundary.
- ▶ Engineer mandated handoffs and shared watch-lists into the architecture, making cross-agency sharing a required function rather than an act of initiative.
- ▶ Design for fusion of fragments that look minor in isolation, so the system is built to assemble the picture no single holder can see.

IN OPERATION / AFTER

- ▶ Audit cross-boundary information flows for reliance on individual initiative, and treat any flow that depends on goodwill as an unowned failure mode.
- ▶ Monitor whether the integrating body actually receives and fuses what the component agencies hold, rather than assuming the mandate guarantees the practice.
- ▶ Sustain the integration architecture against the agency incentives, classification rules, and turf that pull information back behind boundaries over time.

REFLECTION QUESTIONS

1. Identify a cross-organizational information flow in your domain that depends on individual initiative rather than institutional architecture. What is the foreseeable failure mode?
2. Design the institutional deliverable that would have produced ODNI-level information sharing across U.S. intelligence agencies in the 1990s.
3. The Commission called it a “failure of imagination”; critics called it a structural gap. Take a near-miss in your domain and argue which framing fits — and show how the remedy differs depending on which you choose.

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · knowledge bases · data pipelines

Chapter 8

Defense & National Security — What Works — and the Frontier

When training, standards, and sustainment are engineered as one deliverable.

The readiness that survived contact was the readiness that had been measured.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard

G Governance & Trust **K** Knowledge & Institutional Memory

IMPACT DoD Design Criteria Standard MIL-STD-1472H, the 2020 revision of a series dating to 1968, converts human-factors and human-engineering findings into binding design criteria across DoD acquisition — controls, displays, anthropometry, workspace, environment, hazards — invoked by acquisition programs as a mandatory or tailored design specification

IN BRIEF

MIL-STD-1472 is the US Department of Defense's design-criteria standard for human engineering — the discipline of making equipment usable, safe, and effective for the human operator. The series originated in 1968 and has been revised through versions A, B, C, D, E, F, G, and most recently H, released September 15, 2020. The standard codifies established human-factors and ergonomic findings into specific binding design criteria: control and display design, anthropometric accommodation ranges, workspace dimensions and access, labeling and signage, environmental limits (noise, vibration, illumination, thermal), and hazard mitigation. Acquisition programs invoke MIL-STD-1472H either as a mandatory standard or with tailored deviation, making human-factors evidence binding rather than advisory. The 2020 revision integrated accumulated findings since 2012 (the prior G revision), including updates to anthropometric data, environmental criteria, and human-system-integration practices. The case is the structural archetype of converting a body of human-factors research into engineered design requirements; it works at the requirement-discipline layer rather than the per-platform layer, and is the standard that programs referenced when they specify human-engineering deliverables.

BACKGROUND

Human factors engineering as a defense discipline emerged from World War II — the documented mismatch between cockpit control layouts and pilot perception (Fitts & Jones's "designed errors" work) showed that platforms could be made un-flyable by avoidable design choices. Through the 1950s and 60s the discipline accumulated a body of findings on control / display design, anthropometry, environmental tolerance, and workspace layout that lived across academic publications, military handbooks, and service-specific guidance. What was missing was a single binding design-criteria standard that an acquisition program could specify in a contract.^o

THE INTERVENTION

The DoD response was MIL-STD-1472, first issued in 1968, and revised through letter suffixes A (1969), B (1974), C (1981), D (1989), E (1996), F (1999), G (2012), and H (September

2020). Each revision incorporates accumulated findings since the prior version, updated anthropometric data (the surveyed populations evolve), and refinements to specific criteria informed by service experience. The standard is not a list of recommendations; it is a Design Criteria Standard, structured so that an acquisition program can invoke it as a mandatory specification, or tailor specific paragraphs with documented deviation.¹

HOW IT WORKED

The structural mechanism is the conversion of human-factors findings into engineered requirements. Where the research literature might find that controls of a given size and force are operable across a defined population percentile, MIL-STD-1472H carries that finding as a binding design criterion that the program's controls must meet. Where the literature finds that illumination below a certain level degrades reading performance for given character sizes, the standard carries that as an environmental limit. Anthropometric accommodation — the body-size ranges the equipment must fit — is specified to defined population percentiles. Hazard analysis is required as part of the design process, with mitigation criteria for identified hazards.²

THE EVIDENCE

The case's value to LENS is the requirement-discipline form. MIL-STD-1472H is the human-engineering analog of what SUBSAFE (Case 173) does for submarine watertight-integrity: a binding, recurring, auditable requirement set that the program-management chain cannot trade away without a documented deviation. The acquisition contract invokes the standard, the acquirer's human-factors discipline verifies compliance against it, and the engineered design carries the research findings as binding criteria rather than as advice the program might or might not adopt. The 2020 H revision is notable for the cadence — the eight-year gap from G to H — and for the integration of human-systems-integration practices reflecting the rise of complex software-intensive systems.³

WHAT TRANSFERRED

The honest framing the case requires is that the standard is necessary but not sufficient. MIL-STD-1472H is the requirements-discipline mechanism; whether the requirements are met in any specific program depends on the program's human-factors engineering capacity, the acquirer's verification rigor, and the tailoring decisions made up front. The standard does not by itself guarantee usable equipment; it makes usability a contractable, auditable deliverable. Cases of operator-platform mismatch in fielded DoD systems (e.g., the F-22 OBOGS instrumentation gap in Case 132) are not failures of the standard's existence but of the tailoring and verification around it. The case teaches the requirements-as-deliverable form at the human-engineering scale, with the qualification that the standard is the mechanism, not the outcome.

The criterion that survives into a binding standard is the criterion the program is held to.

EDITORS' SYNTHESIS OF MIL-STD-1472 REVISION HISTORY.

REFERENCES

1. Department of Defense (2020), MIL-STD-1472H "Department of Defense Design Criteria Standard: Human Engineering," 15 September 2020 — replaces MIL-STD-1472G (2012).
3. Chapanis, A. (1965), "Man-Machine Engineering" — foundational text for the discipline the standard codifies.

2. Department of Defense (2012), MIL-STD-1472G — the prior revision; revision history documents the 1968 origin and intermediate letters.

4. Fitts, P. M., & Jones, R. E. (1947), “Analysis of factors contributing to 460 ‘pilot error’ experiences in operating aircraft controls” — origin of designed-error analysis.


THE LEARNING ENGINEERING LENS

LE INSIGHT

MIL-STD-1472H is the structural archetype of converting a body of human-factors research into binding engineered requirements an acquisition contract can specify. The standard is the mechanism by which usability becomes a contractable, auditable deliverable rather than design advice. It is necessary but not sufficient; tailoring and verification determine whether the contract is honored.

LENS APPROACH

MIL-STD-1472H is the binding-standard requirements case (induced 1.1; LENS D1/PT3). LENS uses it in Domain 1 (Systems Analysis) for the conversion of research findings into engineered requirements, and in Domain 5 (Navigating Sociotechnical Constraints) for the tailoring and verification disciplines around the standard. Direct pair with Case 173 (SUBSAFE) at the requirements-as-deliverable layer and with Case 132 (F-22 OBOGS) at the standard-versus-tailoring layer.



APPROACHES TO CONSIDER

- DURING DEVELOPMENT

 - ▶ Treat the human-factors literature in your domain as a source of engineered requirements, not as design advice. The criterion that survives into a binding standard is the criterion the program is held to.
 - ▶ Specify anthropometric and environmental criteria to defined population percentiles, not to “typical users.” The percentile framing is what makes accommodation auditable.
 - ▶ Design the tailoring discipline around the standard so that any deviation from a binding criterion is documented with rationale, and the program retains a record of which criteria it chose not to meet and why.
- IN OPERATION / AFTER

 - ▶ Maintain the revision cadence the field requires — anthropometric data ages, environmental tolerances drift, human-system-integration practices evolve — so the standard does not become an obsolete reference.
 - ▶ Verify compliance as a discipline distinct from the standard itself; the standard is the requirement, the verification process is the assurance that the requirement is met.
 - ▶ Carry the necessary-but-not-sufficient framing: a binding human-engineering standard is the mechanism by which research findings become contractable; whether the contract is honored is the program’s responsibility, not the standard’s.

REFLECTION QUESTIONS

1. Identify a body of research findings in your domain that programs treat as advisory. What would it take to convert those findings into a binding design-criteria standard an acquirer could specify in a contract?
2. Specify the tailoring discipline you would put around a binding standard: which paragraphs may be tailored, with what documented rationale, and what is the cadence at which the standard itself is revised against accumulated experience?
3. The standard is necessary but not sufficient. What verification capacity does your program need to know that the binding criteria are actually met in the delivered system, and where is the gap between standard and verification visible today?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines

GIFT and the Adoption Gap

K Knowledge & Institutional Memory **G** Governance & Trust **N** Normalization of Deviance

IMPACT Active open-source ITS framework with demonstrated learning effectiveness; ubiquitous fielded adoption in routine military training has not been achieved

IN BRIEF

The Generalized Intelligent Framework for Tutoring (GIFT) is an open-source framework, originated at the U.S. Army Research Laboratory, for authoring and delivering intelligent tutoring systems. Computer-based tutoring has been shown to be about as effective as expert human tutors in well-defined domains, and GIFT exists to lower the barrier to building it; the framework is actively developed, with regular releases and a peer-reviewed annual symposium. The puzzle is not that GIFT failed — it did not — but that ubiquitous fielded adoption across routine military training remains limited despite decades of supporting research. That gap is the canonical learning-engineering problem: the science is settled, the platform exists, the studies are positive — and the institutional pathway to scaled adoption is still being built.

THE SHIFT

Five decades of research established that computer-based tutoring can be about as effective as an expert human tutor in well-defined domains, and significantly better than traditional classroom instruction. The open question shifted from “does adaptive tutoring work?” to “why isn’t it everywhere?”^o That shift matters because it moves the problem out of the laboratory: once the efficacy question is answered, every remaining obstacle to scaled use is institutional rather than scientific, and the field’s research strength stops being the binding constraint.

WHAT IS EMERGING

The Generalized Intelligent Framework for Tutoring (GIFT), originated under the U.S. Army Research Laboratory and now developed at DEVCOM, is an open-source framework for authoring, delivering, and evaluating intelligent tutoring systems — an effort to lower the authoring barrier that has historically made ITS expensive to build. It is actively maintained, with regular releases and a peer-reviewed annual symposium.¹ Open-sourcing the framework and sustaining a research community around it directly attacks the cost-to-build problem, since the expense of authoring a tutor from scratch had long been the practical reason adaptive tutoring stayed confined to well-funded demonstrations.

THE CAPABILITY QUESTION


So the puzzle is not failure — GIFT was not discontinued. It is that ubiquitous fielded adoption across routine military training remains limited despite a working framework and decades of positive research. The science is settled and the platform exists; what is missing is the institutional pathway to scaled use.² This is the diagnostic feature of the case: the gap is not between idea and prototype but between a maintained, evidence-backed framework and the routine training pipelines that would have to adopt it, and that latter distance is the one no amount of further research closes.

EARLY EVIDENCE

The effectiveness evidence is strong — the tutoring-effectiveness literature is among the more robust in education — and GIFT-based studies continue to demonstrate learning gains. The bottleneck is not efficacy.³ Because the supporting literature is among the more robust in the field, a decision-maker hesitating to field adaptive tutoring cannot honestly point to weak evidence as the reason; the hesitation traces instead to the missing pathway that would let the proven approach be bought, integrated, and made routine.

OPEN PROBLEMS

What has not been built is the adoption pathway: procurement that can buy adaptive tutoring, integration into existing training pipelines, instructor-workflow redesign, and the authority structure to make adaptive tutoring a default rather than an experiment.⁴ GIFT is the case in this book closest to the LENS discipline itself — proof that a working technology and settled science do not adopt themselves, and that the adoption pathway is an engineering deliverable in its own right. Each of those pieces — a contracting vehicle, a pipeline integration, a redesigned instructor workflow, an owner with authority — is a concrete artifact someone must build, and their absence, not any technical shortfall, is what keeps the framework experimental.



The technology works. The institutional pathway to ubiquitous fielded use does not yet.

EDITORS' SYNTHESIS OF THE GIFT ADOPTION RECORD

REFERENCES

1. K. VanLehn, "The Relative Effectiveness of Human Tutoring, Intelligent Tutoring Systems, and Other Tutoring Systems," *Educational Psychologist* 46(4) (2011) — tutoring effectiveness.
4. R. Sottolare, A. Graesser, X. Hu & H. Holden (eds.), *Design Recommendations for Intelligent Tutoring Systems (GIFTSym volumes)*; IJAIED Special Issue on GIFT (2017).

2. GIFT Project, gifttutoring.org — the framework, releases, and development under ARL / DEVCOM.

3. Editors’ synthesis of the GIFT adoption record — active development but limited ubiquitous fielding (quoted).

5. J. Goodell & J. Kolodner, Learning Engineering Toolkit (2022) — adoption as an engineering problem.

THE LEARNING ENGINEERING LENS

LE INSIGHT

GIFT is the most directly relevant case in this book to the learning-engineering discipline itself. The technology works, the learning science works, the framework is active and supported. What has not been built is the institutional adoption pathway — procurement, training-pipeline integration, instructor workflow redesign, the authority structure — that would make adaptive tutoring a default rather than an experiment.

LENS APPROACH

LENS treats GIFT in LEN 1 as the problem-framing case for the discipline, in LEN 8 as the foundational adoption-pathway case, and in LEN 10 (capstone) as a prompt for designing the institutional deliverables that would close an adoption gap of this shape.

DOMAIN	MODES AT STAKE			LENS COMPETENCY					PROBLEM TYPE	INDUCED	CLO		
DEFENSE	→	K	G	N	→	D1	D2	D3	D4	D5	4	1.4	2
EDUCATION AT SCALE	Iterative Development												

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the adoption pathway — procurement, pipeline integration, instructor-workflow redesign, authority — as a named deliverable of the program, not a follow-on hope.
- ▶ Build the tutoring capability against an existing training pipeline so integration is designed in, rather than fielding a framework and expecting pipelines to bend to it.
- ▶ Specify who holds the authority to make adaptive tutoring a default, so the decision to scale is owned rather than left to volunteers.

IN OPERATION / AFTER

- ▶ Measure fielded routine use, not just study-level learning gains, so the institution can see whether adoption is actually happening.
- ▶ Sustain the open-source framework and its community so the authoring-cost barrier it lowered does not quietly rise again as releases age.
- ▶ Audit each stalled pipeline to find which adoption artifact is missing — contracting vehicle, integration, workflow, owner — and treat that as the engineering gap to close.

REFLECTION QUESTIONS

1. GIFT exists, is supported, and works. Adoption at scale does not. What is the equivalent in your domain — an effective intervention whose adoption pathway has not been engineered?
2. Design the institutional adoption deliverable that would move adaptive tutoring from “available framework” to “default routine practice” in one operational training pipeline.
3. GIFT’s bottleneck is procurement, pipeline integration, instructor-workflow redesign, and authority — not efficacy. For an effective tool in your domain, which of those four is the binding constraint, and who would have to own it for adoption to become the default?

WHO BUILDS THIS knowledge management & data engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS knowledge bases · data pipelines · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

Navy Surface Warfare Readiness Reform

T Training Gap **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT Tripled ship-driving training hours; 10 pass-or-fail career assessments; Ready-for-Sea Assessments — 3 of 18 forward-deployed ships immediately sidelined

IN BRIEF

After the fatal 2017 Fitzgerald and McCain collisions (Case 124) exposed gutted seamanship training, the U.S. Navy overhauled how it builds surface-warfare competence. It restored the Surface Warfare Officers School from CD-ROM self-study to classroom and simulator instruction, stood up Mariner Skills Training Centers on both coasts, created ten pass-or-fail career assessments — three of them no-go gates — and adopted aviation-style debriefing. New Ready-for-Sea Assessments evaluated forward-deployed ships against a deliverable standard; three of the first eighteen were immediately sidelined. The structural change is real and the investment substantial. What is missing is the third half: the GAO has noted the Navy still lacks systematic evaluation of whether the reforms work — a live success in structure, with evidence of effect still outstanding.

BACKGROUND

In 2017, the destroyers USS Fitzgerald and USS John S. McCain suffered fatal collisions (Case 124) that investigations traced in part to seamanship and navigation training degraded to CD-ROM self-study. The Navy faced a clear capability gap: officers were going to sea without the hands-on ship-driving competence the job required. The diagnosis was unusually unambiguous — two avoidable collisions in one year pointing at the same eroded fundamentals — which is what converted a long-tolerated shortfall into a mandate for structural change.⁰

THE INTERVENTION

Beginning in 2018, the Navy restored the Surface Warfare Officers School from self-study to classroom-plus-simulator instruction, established Mariner Skills Training Centers on both coasts, roughly tripled ship-driving training hours, and created ten pass-or-fail assessments across an officer's career path — three of them no-go gates that can halt advancement. The no-go gates were the structural teeth: by tying advancement to demonstrated competence rather than time served, they made the qualification something the system would stop on, not merely something it recorded.¹

HOW IT WORKED

The reform paired technical investment — simulators, restored curricula, qualification gates — with a cultural change borrowed from aviation: structured debriefing and explicit gate ownership. New Ready-for-Sea Assessments evaluated forward-deployed ships against a

deliverable standard rather than a paper one; three of the first eighteen ships assessed were immediately sidelined as not ready. That a sixth of the first cohort failed against the new standard showed the assessment had teeth — it measured demonstrated readiness rather than accepting the paper certifications that had masked the pre-collision gap.²

THE EVIDENCE

Here the case turns instructive. The Government Accountability Office has noted that the Navy still lacks systematic evaluation of whether the reforms actually improve readiness. The structural intervention happened and the investment was large, but decision-grade evidence on outcomes is incomplete — the measurement to confirm the effect has lagged the change itself. Without an outcome time-series, the Navy cannot yet distinguish a real capability gain from activity that looks like one, which is the same blindness, in milder form, that let the pre-collision gap go unseen.³

WHAT TRANSFERRED

The reform is a live, in-progress success: structural change real, evidence of effect outstanding. As a teaching case it argues that mature capability engineering must build the measurement infrastructure from the start — the time-series that lets an institution know whether the capability it bought is materializing — rather than treating evaluation as an afterthought. The two halves the reform did deliver, technical investment and a cultural change, are necessary but not sufficient; the missing third half is the evidence regime that would let the Navy prove, not assume, the gap has closed.⁴



The Navy still lacks systematic evaluation of whether the reforms work.

PARAPHRASING GAO-21-168 ON NAVY READINESS REFORM, 2021

REFERENCES

1. GAO-21-168 (and GAO-20-154), Navy readiness reform — the lack of systematic outcome evaluation (paraphrased).
2. Readiness Reform Oversight Council, One-Year Report (2019) — restored training, assessments, and gates.
3. Navy and NTSB reports on the Fitzgerald and McCain collisions (2017–2019) — the training-degradation antecedent.
4. SWOS Norfolk and San Diego Mariner Skills Training Center documentation — simulators and curriculum.
5. USNI News reform coverage (2020, 2022) — Ready-for-Sea Assessments and sidelined ships.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Navy reform is a paired intervention in progress: technical (training restored, simulators procured, assessments created) plus cultural (debriefing, gate ownership). What it lacks is the third half: evidence that the intervention has worked. LENS treats this as a teachable case for what mature capability engineering should include from the start — the measurement infrastructure that lets the institution know whether its investment is producing the capability it bought.

LENS APPROACH

LENS uses Navy SWO reform in LEN 10 (capstone) as a worked exercise in capability intervention at scale, and in LEN 4 as a case where the measurement infrastructure has lagged the intervention. Students design the evidence regime that should accompany the reform.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
DEFENSE	→ T K N	→ D1 D2 D3 D4 D5	4	2.3	4
Iterative Development + D4					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pair the technical investment — restored curricula, simulators, training hours — with a cultural change such as structured debriefing and explicit gate ownership, rather than buying tools alone.
- ▶ Install no-go gates that tie advancement to demonstrated competence, so qualification is something the system will stop on, not merely something it records.
- ▶ Assess units against a deliverable standard from the first cohort, accepting that a real test will sideline some — that failure rate is the evidence the gate has teeth.

IN OPERATION / AFTER

- ▶ Build the measurement infrastructure from the start — the outcome time-series that lets the institution distinguish a real capability gain from activity that merely looks like one.
- ▶ Give the senior commander a readiness dashboard that reports demonstrated capability over time, so the same blindness that masked the pre-collision gap cannot recur.
- ▶ Treat evidence of effect as the third half of the intervention, not an afterthought, so the reform can be proven rather than assumed to have closed the gap.

REFLECTION QUESTIONS

1. Navy SWO reform happened without an evidence regime to confirm it worked. Identify a current reform in your domain and the evidence regime that should accompany it.
2. Design the dashboard the Chief of Naval Operations should have to know whether SWO capability is improving in time-series.
3. The Ready-for-Sea Assessment sidelined three of the first eighteen ships because it measured readiness against a deliverable standard. Identify a certification in your domain that ratifies rather than tests, and design the gate that would be willing to fail a sixth of its first cohort.

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards

xAPI / Total Learning Architecture — Interoperability Gap

K Knowledge & Institutional Memory **G** Governance & Trust

IMPACT Despite xAPI adoption, most implementations remain siloed; cross-organizational learning-data interoperability largely unrealized

IN BRIEF

The Experience API (xAPI) was created to track learning experiences across platforms, and the Advanced Distributed Learning Initiative's Total Learning Architecture envisioned learning records, competencies, and credentials flowing across organizational boundaries to support continuous capability development — the kind of evidence infrastructure LENS depends on. In practice, xAPI implementations remain largely siloed inside individual learning-management systems; the cross-organizational data sharing that matters most for high-consequence domains has not materialized at scale. The technical standard exists and reference implementations exist; what lags is governance — who owns the data, what consent applies, how quality is assured. xAPI mirrors inBloom (Case 53) in structure — technology ahead of governance — across the whole learning-technology ecosystem. It is the book's case for treating data governance as a capability deliverable.

THE SHIFT

Learning increasingly happens across many systems — courses, simulators, on-the-job tools — and the field recognized it needed a common way to record those experiences, so capability could be tracked over a career rather than a single course.^o The career-long view is the point: a single course's records say little about whether a person can do the job, whereas experiences stitched across courses, simulators, and on-the-job tools are what let an institution reason about real capability rather than completed seat-time.

WHAT IS EMERGING

The Experience API (xAPI) was built to do exactly that, and the Advanced Distributed Learning Initiative's Total Learning Architecture envisioned learning records, competency frameworks, and credentials flowing across organizational boundaries — the evidence infrastructure a discipline like LENS depends on.¹ The vision was explicitly cross-boundary: records, competencies, and credentials that move with the learner between organizations are precisely the evidence base on which a capability discipline must stand, which is why the standard's promise mattered well beyond any single training shop.

THE CAPABILITY QUESTION

In practice, xAPI implementations remain largely siloed inside individual LMS platforms. The cross-organizational data sharing — the capability most relevant to high-consequence operational domains — has not materialized at scale. The standard exists; the ecosystem it promised does not.² The siloing is the diagnostic detail: the very cross-boundary flow that made the architecture worth building is the part that did not arrive, so the standard delivers tracking within each LMS while the career-long, cross-organizational record it envisioned stays out of reach.

EARLY EVIDENCE

The technical pieces work: the data model is sound and reference implementations exist. What has lagged is the governance — who owns the data, what consent frameworks apply, and how data quality is assured across organizations.³ With the data model proven and reference implementations in hand, the remaining obstacles are not engineering questions an organization can solve alone but agreements between organizations — ownership, consent, and assured quality — that no technical specification can settle on their behalf.

OPEN PROBLEMS

xAPI mirrors inBloom (Case 53) in essential structure — technology in advance of governance — but across the whole learning-technology ecosystem rather than one initiative.⁴ It is the book's clearest case that an interoperability standard is necessary but not sufficient: without the governance to make organizations willing and able to share, the data stays in its silos, and the evidence infrastructure remains a diagram rather than a system. The parallel to inBloom is instructive because it shows the pattern is not particular to one failed project: a sound standard arriving ahead of the ownership, consent, and quality-assurance arrangements will stall the same way at any scale.



The standard exists. The governance does not.

EDITORS' SYNTHESIS OF ICICLE / ADL TLA DISCUSSION

REFERENCES

1. Advanced Distributed Learning Initiative, Total Learning Architecture documentation — the cross-boundary vision.
2. xAPI specification (xapi.com) — the technical standard and reference implementations.
3. B. Saxberg, learning-engineering infrastructure essays; IEEE ICICLE LEBOK chapters on data and analytics.
4. Cf. inBloom (Case 53) — technology in advance of governance.

3. IEEE ICICLE proceedings on TLA adoption challenges — the persistence of siloed implementations.

THE LEARNING ENGINEERING LENS

LE INSIGHT

xAPI/TLA is the technical-standard analog of the implementation gap — the discipline has the data model, the spec, and the reference implementations. What it does not have is the governance and institutional architecture to make cross-organizational learning data flow. This is the canonical case in this book for treating data governance as a capability-engineering deliverable.

LENS APPROACH

LENS treats xAPI/TLA in LEN 4 as a data-governance and interoperability case and in LEN 8 as an example of organizational- learning infrastructure that has not scaled. The case is the technical underlay to the larger argument about evidence systems that decision-makers can trust.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
EDUCATION AT SCALE DEFENSE	→ K G	→ D1 D2 D3 D4 D5 Navigating Sociotechnical Constraints	4	5.3	5

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
- Specify the governance — data ownership, consent, and cross-organization quality assurance — as a deliverable alongside the technical standard, not after it. - Design for the cross-boundary flow from the start, so the standard is not merely implemented inside each LMS but engineered to move records between organizations. - Secure the inter-organizational agreements that no specification can settle, treating willingness-to-share as a thing to be built rather than assumed.	- Audit whether learning data is actually crossing organizational boundaries, not just whether xAPI is nominally adopted within each silo. - Monitor data quality across organizations continuously, since shared records are only trustworthy evidence if their quality is assured at the seams. - Watch for the inBloom pattern — a sound standard outrunning its governance — and treat any stall as a missing-governance signal, not a technical defect.

REFLECTION QUESTIONS

1. Why has the xAPI standard not produced cross-organizational interoperability at scale? What governance condition is missing?
2. Design the minimum governance architecture under which xAPI data could flow across two organizations in your domain.
3. The xAPI data model is sound; what stalled was ownership, consent, and quality assurance across organizations. For a data-sharing effort in your domain, which of those three is the unresolved question — and who would have to agree for the data to actually flow?

WHO BUILDS THIS knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** knowledge bases · data pipelines · governance frameworks · bias & evidence audits

Chapter 9

Industry, Energy & Enterprise Systems — What Fails

When incentive, inertia, and migration risk go unengineered.

The signal arrived years early. The iteration never ran.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

EDITOR'S NOTE · THE ITERATION GAP

Iterative development is what learning engineering does most often. It is also the capability the casebook documents least directly. Several cases in this chapter — and their counterparts across the book's failure halves — share a shape: an organization received an unambiguous signal, held the engineering or product position to act on it, and failed to run the iteration cycle through product, business model, and operations until the opportunity had passed. The signal was sometimes a field report (Therac-25, the adoption and sustainment cases), sometimes an external launch (BlackBerry meeting the iPhone), sometimes an internal invention the organization could not build a business around (Kodak's 1975 digital camera). In each case the organization had the engineering capability and the institutional position. What it lacked was the discipline of running the iteration through.

The casebook's broader corpus shows what iteration produced — the interventions collected in every part's What Works chapter, the failures whose recovery iterations followed. It does not yet show iteration **being done** in detail: the design moves, the dead ends, the reframes. That gap is exactly what the v2.1 amendment to D2 names — “narrate and defend the design iteration in first person” — and what the editor-commissioned first-person Practice Flywheel accounts will fill in the next edition. This note names the gap honestly so that the reader knows what to expect from the volume, and what the discipline is still working to make legible.

Knight Capital Trading Loss

D Designed Out **K** Knowledge & Institutional Memory

IMPACT \$440M loss in 45 minutes; firm forced to seek emergency capital and was effectively acquired within months

IN BRIEF

On 1 August 2012 Knight Capital, a major U.S. market maker, deployed new order-routing software to seven of its eight servers and missed the eighth. The new code reused a flag that, on the eighth server's old software, reactivated long-dead "Power Peg" code never removed from the repository. At the opening bell it fired millions of unintended orders; in about 45 minutes Knight lost roughly \$440 million — more than the firm was worth — and was effectively acquired within months. The SEC found Knight had no procedure to verify the deployment across all servers and no controls to halt the runaway orders. The capability designed out was deployment verification; the dead code was technical debt that exercised its option at the worst possible moment.

BACKGROUND

Knight Capital was one of the largest market makers in U.S. equities, routing an enormous share of retail order flow through automated systems wired directly into the market. On 1 August 2012 the New York Stock Exchange was launching a new Retail Liquidity Program, and Knight had updated its order-routing software to participate in it from the opening bell. The update went out to production the way countless updates had gone before — a routine deployment, on an ordinary morning, handled as unremarkable.⁰

WHAT HAPPENED

A technician deployed the new code to seven of the eight routing servers and missed the eighth, leaving one node running stale software. The new code reused a configuration flag that, on the old software still on that eighth server, had once activated long-dormant "Power Peg" test code — retired years earlier but never removed from the repository. At the open the dead code woke and began firing millions of unintended orders into the market. In about forty-five minutes Knight amassed a vast unwanted position; losses passed \$170 million almost at once and reached about \$440 million — more than the firm itself was worth. It survived only on emergency capital and was effectively acquired within months.¹

THE INVESTIGATION

The Securities and Exchange Commission's September 2013 enforcement order — In re Knight Capital Americas LLC — found Knight had no written procedure requiring a second technician to verify the deployment across all hosts, and no automated check that all eight servers were running the same code — nor controls able to recognize and halt the flood

of erroneous orders once it began. The order catalogued specific violations of the SEC's Market Access Rule (Rule 15c3-5), and Knight settled for a \$12 million penalty.² The dead Power Peg code was the proximate trigger and the reused flag the match that lit it; but the underlying cause was the absence of the verification and risk controls that should surround any change to a system wired directly into the live market. The order is unusually explicit that the institutional gap was financial-engineering practice catching up to software-engineering practice — release management, configuration control, and pre-trade risk limits were not yet treated as first-class deliverables on the trading desk.³

THE CAPABILITY GAP

The capability designed out was deployment verification — confirming, every time, that what runs in production is exactly what was intended, on every node, before it touches live money. The dead code was technical debt in its most literal form: a retired function left in the repository is a standing option on a future failure, and reusing its flag exercised that option at the worst possible moment. As at the Mars Climate Orbiter's interface, the boundary that mattered — between “deployed” and “verified as deployed everywhere” — had no owner and no automated check, and a large institution kept walking across it routinely until the day the floor gave way.⁴

AFTERMATH & REFORM

Knight became the canonical case in modern software-operations practice for why deployment is itself an engineering deliverable, not a clerical step: automated verification that every host runs the intended build, disciplined removal of dead code from the repository, pre-trade risk limits, and kill switches that can stop a runaway process in seconds. Regulators sharpened their attention to automated market-access controls in response. The SEC's 2014 Regulation Systems Compliance and Integrity (Regulation SCI) extended formal system-development, testing, change-management, and incident-reporting obligations to registered exchanges, clearing agencies, and large alternative trading systems — the institutional response that pulled the deployment discipline Knight had lacked into the regulatory perimeter. The lesson rhymes with the orbiter's across a forty-year, civilian-to-financial gap: a small, unowned boundary inside a large automated system is precisely where the institution is most exposed, and least watching — and as algorithmic trading continued to grow after 2012, the same class of fragility kept producing new incidents that the Knight template explained.⁵

\$440M

LOSS IN 45 MINUTES · 7 OF 8 SERVERS UPDATED

dead code on the eighth server, repurposed flag, no deployment verification

KNIGHT CAPITAL — THE COST OF AN UNVERIFIED DEPLOYMENT

Knight did not have written procedures requiring a second technician to verify the deployment.

PARAPHRASING THE SEC ORDER AGAINST KNIGHT CAPITAL, 2013

REFERENCES

1.

U.S. Securities and Exchange Commission, Order Instituting Administrative and Cease-and-Desist Proceedings, In re Knight Capital Americas LLC (2013) — the firm, the Retail Liquidity Program launch, and the deployment.

2.

SEC Order (2013) and Knight Capital 8-K filing (2012) — the missed eighth server, the reactivation of the dormant “Power Peg” code, the \$440 million loss in 45 minutes, and the near-collapse.

3.

SEC Order (2013) — the absence of a second-technician deployment verification, the lack of an automated code-consistency check, inadequate order controls, and the \$12 million penalty (quoted).

4.

B. Beyer, C. Jones, J. Petoff & N. R. Murphy (eds.), Site Reliability Engineering (O’Reilly, 2016) — deployment verification, dead-code removal, and automated safeguards as engineering deliverables.

5.

SEC Market Access Rule (Rule 15c3-5) and subsequent automated-controls guidance — the regulatory response on pre-trade risk and market-access controls.

6.

D. Seven, “Knightmare: A DevOps Cautionary Tale” — a widely cited engineering post-mortem on the deployment process, the orphaned eighth server, and the reused feature flag.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Knight Capital is the financial-industry version of Mars Climate Orbiter (Case 98): a small, unspecified boundary inside a large system that took the institution down. The capability that was missing was deployment verification. The dead code was the proximate trigger; the absent procedure was the cause.

LENS APPROACH

Knight Capital is the canonical change-control-and-disclosure governance case (induced 5.4; LENS D1/PT3). LENS uses it in LEN 5 to teach deployment-as-capability — students design the deployment deliverable that would have caught the eighth server — and in LEN 9 for the technical-debt argument: every line of dead code carries an option on a future failure. Adjacent to Mars Climate Orbiter (Case 98) at the small-boundary-no-owner layer and to Regulation SCI as the institutional response that codified the missing controls.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
TECHNOLOGY	→ D K	→ D1 D2 D3 D4 D5	3	7.1	1 · 4
Systems Analysis					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Remove retired code and dormant flags from the repository as a matter of discipline, since every dead function left behind is a standing option on a future failure.
- ▶ Never reuse a configuration flag whose old meaning still exists somewhere in the fleet; treat flag semantics as a versioned, owned interface.
- ▶ Design pre-trade risk limits and a kill switch into any system wired to the live market, so a runaway process can be stopped in seconds rather than minutes.

IN OPERATION / AFTER

- ▶ Require automated verification that every host runs the intended build before a change touches live money, with a second technician confirming the deployment across all nodes.
- ▶ Monitor for the runaway condition — an abnormal order rate — and halt automatically, rather than relying on humans to notice and intervene mid-flood.
- ▶ Audit deployments for the orphaned node: confirm no host is left on stale code, closing the gap between “deployed” and “verified as deployed everywhere.”

REFLECTION QUESTIONS

1. Identify a deployment procedure in your domain whose verification step depends on convention rather than on a designed check. What is the eighth server?
2. Design the deployment deliverable that would prevent a Knight Capital-equivalent loss in your organization.
3. A retired “Power Peg” function left in the repository became the trigger years later when its flag was reused. Identify dead code or a dormant feature flag in a system you run that is still an option on a future failure, and specify the policy that should have removed it.

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines

Northeast Blackout

H Human–System Interface **K** Knowledge & Institutional Memory

IMPACT 50 million people without power across eight U.S. states and Ontario; \$6B+ economic loss; FERC Order 693 followed

IN BRIEF

On 14 August 2003 a high-voltage transmission line in Ohio sagged into a tree and tripped — an event the grid should have absorbed. But FirstEnergy's control-room alarm system had been silently failing for over an hour, so operators did not know the line was gone. Further lines tripped, and a cascade swept across the Eastern Interconnection; within minutes 50 million people across eight U.S. states and Ontario lost power, at a cost above \$6 billion. The U.S.–Canada task force found FirstEnergy lacked situational awareness, its alarm system had failed without notice, vegetation management was poor, and the regional coordinator could not intervene. The reforms made reliability standards mandatory and enforceable (FERC Order 693). The gap sat at the automation–operator boundary: a silent failure left operators blind.

BACKGROUND

The Eastern Interconnection is built to ride through the loss of a single transmission line, and control rooms watch the grid through software and alarms. FirstEnergy, the Ohio utility at the center of the story, ran a control room whose alarm system had — unknown to anyone — been silently failing for over an hour.⁰ The interconnection's single-line resilience assumes the operators can see which line is gone; a silent alarm broke that assumption at its root, leaving a grid designed to tolerate one fault blind to the fault it was tolerating.

WHAT HAPPENED

On 14 August 2003 a 345 kV line in Ohio sagged into a tree and tripped. With the alarms silent, operators did not know the line was gone and so did not take the corrective steps that would have contained it. Further lines tripped and a cascade swept across the interconnection; within minutes 50 million people across eight states and Ontario lost power, at a cost above \$6 billion.¹ The first trip was the routine single-line loss the grid was built to absorb; what turned it into a cascade was not the tree but the hour in which the operators acted on a picture of the grid that no longer matched the grid.

THE INVESTIGATION

The U.S.–Canada Power System Outage Task Force found FirstEnergy operators “did not have adequate situational awareness,” that the alarm system had failed without notice, that vegetation management was inadequate, and that the regional reliability coordinator lacked the authority and information to intervene.² The reforms produced FERC Order 693, which for the first time made compliance with reliability standards mandatory and enforceable


rather than voluntary.³ Making the standards mandatory addressed the deeper finding that a voluntary regime had let vegetation management and operator awareness drift: when compliance is optional, the practices that prevent a cascade are exactly the ones a cost-pressured utility lets slip.

THE CAPABILITY GAP

The gap sat at the boundary between automation and operator. When the alarm system stopped doing its job, it did so silently — and the operators had no independent signal that they were losing control of their grid segment. The missing capability was the meta-monitor: the watch on the watchman, the check that tells you when the thing that is supposed to tell you has itself stopped.⁴ Silence is the most dangerous failure mode because it is indistinguishable from a quiet, healthy grid; the operators were not ignoring a warning but trusting an absence of warning that the broken system could no longer guarantee.

AFTERMATH & REFORM

FERC Order 693 and the new mandatory reliability regime — enforced by an Electric Reliability Organization with audits and penalties — treated grid reliability as a deliverable rather than a best practice. The 2003 blackout endures as the canonical case of silent automation failure: a system that quits without announcing it, leaving the humans nominally in charge with nothing to act on.⁵ Backing the standards with audits and penalties conceded that reliability could not be left to good intentions: the practices that would have kept the operators sighted had to be enforceable obligations, not best practices a utility could quietly defer.



FirstEnergy... did not have adequate situational awareness of conditions on its system.

U.S.-CANADA POWER SYSTEM OUTAGE TASK FORCE, FINAL REPORT ON THE AUGUST 14 2003 BLACKOUT, APRIL 2004

REFERENCES

1. U.S.-Canada Power System Outage Task Force, Final Report on the August 14, 2003 Blackout in the United States and Canada (2004) — the tree contact, the silent alarm, and the cascade.
2. Task Force (2004) — 50 million people affected across eight states and Ontario; the minute-by-minute sequence.
3. Task Force (2004) — FirstEnergy “did not have adequate situational awareness,” plus the vegetation-management and reliability-coordinator findings (quoted).
4. FERC Order No. 693, Mandatory Reliability Standards for the Bulk-Power System (2007) — enforceable standards.
5. North American Electric Reliability Council reports (2004) and the creation of the Electric Reliability Organization.
6. M. R. Endsley (1995), situation-awareness theory — the human-factors frame for silent-automation failure.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The 2003 blackout is the canonical case for silent automation failure: a system that stops doing its job without telling its operators. The capability that was missing was the meta-monitor — the system that watches the monitor. The reform that followed treated grid-reliability compliance as a deliverable rather than as a best practice.

LENS APPROACH

LENS uses the 2003 blackout in LEN 2 as a Human-AI Teaming case for silent-automation-failure handling and in LEN 8 for the legislative-reform arc that produced enforceable reliability standards.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ENERGY	→ H K	→ D1 D2 D3 D4 D5	6	3.3	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design monitoring with a positive heartbeat so a silent system is distinguishable from a healthy one, never assuming no alarm means no problem.
- ▶ Build a meta-monitor that watches the alarm system itself and announces when the watchman has stopped.
- ▶ Give the regional reliability coordinator the authority and information to intervene across utility boundaries before a cascade forms.

IN OPERATION / AFTER

- ▶ Audit vegetation management and operator situational awareness against enforceable standards, since voluntary practice drifts under cost pressure.
- ▶ Test alarm and monitoring systems for detectable failure, verifying that an outage of the monitor itself raises an alert.
- ▶ Back reliability obligations with audits and penalties so the practices that keep operators sighted cannot be quietly deferred.

REFLECTION QUESTIONS

1. Identify an automated monitoring system in your domain whose silent failure would not be detected by its operators. How would they know?
2. Design the meta-monitor that should have been watching FirstEnergy’s alarm system in 2003.
3. Silence was indistinguishable from a healthy grid. What positive heartbeat signal in your domain would let operators tell a quiet system from a dead one?

WHO BUILDS THIS human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS usability testing · cognitive walkthroughs · knowledge bases · data pipelines

Takata Airbag Inflators

D Designed Out **G** Governance & Trust

IMPACT More than 30 deaths and hundreds of injuries linked to inflator ruptures; largest automotive recall in history

IN BRIEF

Takata built its airbag inflators around ammonium nitrate — cheap and energetic but unstable, and used by almost no one else without a moisture-absorbing desiccant. Takata's inflators omitted the desiccant. Over years of heat and humidity the propellant degraded and could rupture the metal housing on deployment, spraying shrapnel at the driver; more than two dozen deaths and hundreds of injuries followed. Takata's own tests had shown the ruptures, but it reported them as isolated anomalies and, in places, falsified data; in 2017 it pleaded guilty to wire fraud, paid roughly \$1 billion, and went bankrupt. The recall — over 100 million inflators across 19 automakers — is the largest in automotive history. Two capabilities were designed out: the desiccant, and the independent verification regulators never had.

BACKGROUND

An airbag inflates by burning a propellant inches from a person's face, so the propellant must be energetic **and** stable across the years and climates a car will see. To cut cost, Takata built its inflators around ammonium nitrate — cheap and energetic but notoriously unstable, shifting crystalline structure and absorbing moisture as temperature cycles day after day. The few competitors who used it at all added a desiccant to keep it dry over the life of the car; Takata's inflators, for years, shipped with none.⁹

WHAT HAPPENED

Over years of heat and humidity the propellant degraded, and a degraded charge can burn too fast, generating more pressure than the housing was built to contain. A firing inflator could then rupture its own metal housing and spray shrapnel into the cabin — turning the device meant to save a life into a fragmentation hazard aimed at the driver. More than two dozen deaths worldwide and hundreds of injuries followed. The recall grew to more than a hundred million inflators across some nineteen automakers — the largest in automotive history — and Takata went bankrupt; the toll keeps rising as unrepaired vehicles stay on the road and their inflators keep aging.¹

THE INVESTIGATION

Takata's own testing had shown the ruptures, but internal documents revealed engineers raising the alarm internally while the company reported the failures to automakers and regulators as isolated anomalies rather than a systemic propellant problem — and, in places, manipulated the test data outright.² In 2017 Takata pleaded guilty to wire fraud and paid

roughly a billion dollars in fine, restitution, and a victims' fund. The legal finding was pointed: not that a part had failed, but that the company had spent years misrepresenting what its own engineers knew.³

THE CAPABILITY GAP

Two capabilities were designed out, and the second matters as much as the first. The product capability was the desiccant — a stabilizing component competitors used, omitted to save cost, in plain view of anyone comparing designs side by side. The system capability was independent verification: the recall regime treated the manufacturer's representations about its own safety data as authoritative, with no independent pipeline to test inflator behavior across the fleet as it aged in the field — which let an obvious failure mode hide for years inside a process that kept receiving the evidence and reporting it back as noise.⁴

AFTERMATH & REFORM

The recall is still being worked off, vehicle by vehicle, long after the company that built the inflators ceased to exist — concrete proof that a designed-out capability can outlive the firm that removed it and become someone else's burden. The episode pushed regulators toward more aggressive, coordinated recall management that does not leave pace to each manufacturer. Its central lesson is the pairing: Takata is the modern Pinto in its product failure, and more in its system failure — the evidence pipeline a regulator relies on is itself a safety capability, and omitting it is as consequential as omitting the desiccant from the inflator.⁵

100M+

INFLATORS RECALLED · 19 AUTOMAKERS

the desiccant competitors used was the designed-out capability

TAKATA — THE LARGEST AUTOMOTIVE RECALL ON RECORD

Takata engaged in a sustained pattern of misrepresenting inflator safety data to its automaker customers and to regulators.

PARAPHRASING THE U.S. DEPARTMENT OF JUSTICE SETTLEMENT WITH TAKATA, 2017

REFERENCES

1. U.S. National Highway Traffic Safety Administration, Takata air-bag inflator recall coordination materials — the ammonium-nitrate-without-desiccant design and the propellant-degradation rupture mechanism.
2. NHTSA recall record and investigative reporting (Reuters, New York Times) — 100M+ inflators across 19 automakers; the largest automotive recall in history; deaths and injuries from ruptures.
3. U.S. Department of Justice, settlement and guilty plea with Takata Corporation (2017) — wire fraud, \$1B in fine and restitution, and Takata's subsequent bankruptcy.
4. U.S. DOJ (2017) and Takata internal documents released in litigation — the sustained misrepresentation of inflator test data to automakers and regulators.
5. NHTSA Takata recall status reporting — the years-long completion of replacements and the continuing risk in unrepaired vehicles.
6. NHTSA Coordinated Remedy Program for the Takata recalls — the regulator's move to actively prioritize and manage replacement across nineteen automakers rather than leave pace to each manufacturer.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Takata is the modern Pinto: an engineering choice to omit a stabilizing component, internal data showing the consequence, and a regulatory architecture that treated the manufacturer’s representation as authoritative rather than as one input. The capability gap was at the regulator’s evidence pipeline as much as at the manufacturer’s bench.

LENS APPROACH

The load-bearing lesson is post-deployment surveillance. Field rupture reports arrived over years but were never aggregated and interpreted as a single class-level safety signal, so an obvious failure mode stayed hidden inside a process that kept receiving the evidence and reporting it back as noise. The ammonium-nitrate-without-desiccant chemistry is the substrate of the hazard, but the capability that should have caught it is the surveillance-aggregation function that turns scattered field reports into one population-level signal with authority to act. LENS uses Takata in LEN 5 to build that aggregate-the-field-signal capability and in LEN 7 as an industrial-governance failure spanning manufacturer, regulator, and customer auditors.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
INDUSTRIAL	→ D G	→ D1 D2 D3 D4 D5	3	2.4	4
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Qualify a safety-critical material against the full service life and climate range it will see, not just its as-delivered state, so slow degradation cannot turn a safe part lethal years later.
- ▶ When choosing a cheaper, less stable material, require the stabilizing measure competitors use — here a desiccant — rather than omitting it to save cost.
- ▶ Design inflator testing to detect the rupture failure mode as a class across temperature and humidity cycling, so it cannot be dismissed as an isolated anomaly.

IN OPERATION / AFTER

- ▶ Stand up an independent verification pipeline that tests fielded inflators across the aging fleet, so a regulator does not depend on the manufacturer’s own representation of its safety data.
- ▶ Aggregate field rupture reports as a population signal rather than closing them one by one, with authority to trigger recall when a pattern appears.
- ▶ Manage the recall as a coordinated program that prioritizes replacement across automakers, since a designed-out hazard outlives the firm and keeps aging in unrepaired cars.

REFLECTION QUESTIONS

1. Where in your domain does a regulator receive manufacturer test data without an independent verification pipeline? What is the load-bearing trust assumption?
2. Design the verification regime that should have surrounded ammonium-nitrate inflator testing. Who funds it, who runs it, and what does it produce?
3. Takata’s propellant degraded slowly with heat and humidity, so a part safe at delivery became lethal years later. Identify a component in your domain whose qualification testing does not cover its full service life, and specify the aging test that would close the gap.

WHO BUILDS THIS systems & human-factors engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits

Volkswagen Dieselgate

D Designed Out **G** Governance & Trust

IMPACT ~11 million vehicles equipped with defeat-device software; USD 4.3B U.S. DOJ criminal-and-civil plea (Jan 2017); total global cost above USD 30B across penalties, vehicle buybacks, and settlements; multiple criminal convictions

IN BRIEF

Volkswagen engineered a “defeat device” into its diesel emissions software: code that detected when a car was on a regulatory test bench and switched on full emissions controls only then. On the road the controls were largely disabled, producing nitrogen-oxide emissions up to forty times the legal limit, across roughly 11 million vehicles for years. The deception was uncovered not by regulators but by a small West Virginia university team comparing real-world to lab measurements. Internal documents showed the defeat device was an institutional decision — a deliberate response to a standard VW’s engineers did not believe they could meet within cost — not a rogue act. VW pleaded guilty to a USD 4.3 billion criminal and civil settlement with the U.S. Department of Justice (January 2017); total Dieselgate cost across penalties, vehicle buybacks, and settlements has been estimated above USD 30 billion globally, with multiple criminal convictions. Dieselgate is the book’s case for organized evasion of a measurement system.

BACKGROUND

Volkswagen had staked its U.S. growth on “clean diesel,” promising cars that met strict nitrogen-oxide limits without sacrificing performance or cost — a marketing position that left no room to admit the engineering could not deliver all three at once. Its engineers did not believe they could actually meet the standard within the platform’s cost constraints, putting the company’s public promise on a collision course with its own technical reality.⁰

WHAT HAPPENED

So they cheated by design. VW’s emissions software detected when a vehicle was on a regulatory test bench — by its steering, speed, and duration patterns, the telltale signature of a lab rather than a road — and switched on full emissions controls only during the test. On the road the controls were largely disabled, producing emissions up to forty times the legal limit, across roughly 11 million vehicles for years, so the pollution the standard existed to prevent flowed freely everywhere except where it was measured.¹

THE INVESTIGATION

The deception was uncovered not by the regulator but by a small university research team in West Virginia comparing real-world emissions to lab results — the gap between road and bench being exactly what the regulator’s own test could never reveal.² Internal documents then showed the defeat device had been authorized inside VW’s engineering hierarchy as


a deliberate institutional response to a standard the team could not meet — not the work of a rogue engineer, but a decision made up the chain of command. VW pleaded guilty to a USD 4.3 billion U.S. DOJ criminal and civil settlement (January 2017) — total global Dieselpgate cost across penalties, vehicle buybacks, and settlements has been estimated above USD 30 billion — and multiple executives were convicted.³

THE CAPABILITY GAP

The exploitable gap was at the regulator's instrument: the emissions test ran in a regime the vehicle could detect, so a manufacturer determined to evade could tune its behavior to the test rather than to the road, optimizing for the measurement instead of the goal it stood for. Dieselpgate is the canonical case of institutionally engineered evasion of a measurement system — and a reminder that any test the measured party can recognize is a test it can defeat, because a detectable test invites the very gaming it is meant to catch.⁴

AFTERMATH & REFORM

The reform closed the gap at the instrument: the EU introduced real-world driving-emissions testing, moving the measurement off the predictable bench and onto the road the standard actually cared about, making it far harder to game.⁵ The pattern parallels Takata (Case 147) — a manufacturer's fraud meeting a regulator whose evidence pipeline trusted the manufacturer's test conditions, and a fix that had to upgrade the measurement itself, not just punish the cheat, because punishing the cheat leaves the exploitable instrument in place for the next one.



The defeat device was the product of a long-standing institutional decision to evade emissions standards.

PARAPHRASING THE U.S. DEPARTMENT OF JUSTICE PLEA AGREEMENT WITH VOLKSWAGEN AG, 2017

REFERENCES

1. U.S. EPA, Notice of Violation to Volkswagen (2015) — the defeat device and emissions exceedances.
2. West Virginia University / ICCT real-world diesel-emissions study (2014) — the discovery comparing road to lab.
3. U.S. Department of Justice Plea Agreement with Volkswagen AG (January 2017) — the institutional decision, USD 4.3 billion criminal-and-civil settlement, and convictions (quoted).
4. Volkswagen internal documents released through litigation — authorization of the defeat device within the engineering hierarchy.
5. EU Real Driving Emissions (RDE) testing regulation — the post-Dieselpgate measurement reform.
6. J. Ewing, Faster, Higher, Farther: The Volkswagen Scandal (2017); cf. Takata (Case 147).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Dieselgate is the canonical case for institutionally engineered evasion of a measurement system. The capability gap was at the regulator’s test protocol: it operated in a regime the vehicle could detect. The reform — real-world emissions testing — was a capability deliverable upgrade at the regulator’s instrument boundary.

LENS APPROACH

LENS reads Dieselgate at the independent-instrument and operator-inspector boundary (induced 2.2): the defeat device worked by detecting the test condition itself, so the measurement was one the operator could see coming and game. The lesson is designing an emissions-measurement and inspection regime the operator cannot detect and defeat — real-driving emissions, surprise and outside-the-loop testing — rather than the bare observation that fraud occurred. The reform pattern parallels Takata (Case 147).

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
INDUSTRIAL	→ D G	→ D1 D2 D3 D4 D5	3	2.2	4
GOVERNMENT	Test and Evaluation				

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design measurement protocols the measured party cannot detect and tune to, so behavior is optimized for the goal rather than the test.
- ▶ Surface the gap between a public commitment and the engineering reality early, so an unmeetable standard is renegotiated rather than secretly evaded.
- ▶ Build a governance check between an impossible target and the team asked to meet it, so the response is escalation, not a defeat device.

IN OPERATION / AFTER

- ▶ Audit real-world behavior against test-condition behavior, treating a divergence between road and bench as the signature of engineered evasion.
- ▶ Upgrade the measurement instrument itself after a discovered cheat, not just penalize the cheat, since the exploitable test invites the next one.
- ▶ Sustain independent, outside-the-loop verification — as the university team provided — to catch fraud the regulator’s own protocol cannot see.

REFLECTION QUESTIONS

1. Identify a measurement protocol in your domain that operates in a regime detectable by the entity being measured. What is the evasion potential?
2. Design the upgrade to a regulatory test protocol that makes Dieselgate-style evasion structurally infeasible.
3. VW’s defeat device was authorized up the engineering hierarchy, not the act of a rogue engineer. What allows an institution to sanction fraud as a deliberate response to an unmeetable standard — and what governance check should sit between an impossible target and the team asked to meet it?

WHO BUILDS THIS systems & human-factors engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits

GM Ignition Switch

D Designed Out **G** Governance & Trust

IMPACT 124 deaths attributed; ~2.6M vehicles recalled for the defective switch; \$900M federal forfeiture (DOJ, 2015); the fix existed for eight years before the recall

IN BRIEF

A faulty ignition switch in several GM compact cars (the Chevrolet Cobalt, Saturn Ion) could slip from “run” to “accessory” while driving, cutting power steering and brakes and — fatally — disarming the airbags. GM engineers identified it in 2002. In 2006 an engineer approved a redesigned switch but did not change its part number, so the fix propagated as neither a revision nor a recall, and defective cars kept selling. The recall came only in 2014 — about 2.6 million vehicles, with 124 deaths attributed through GM’s compensation fund. The Valukas report found a culture that absorbed safety concerns, and GM paid a \$900 million federal penalty. What was designed out was not a part but the pathway by which a known safety problem reaches the decision to recall.

BACKGROUND

The ignition switches in several GM compact cars — the Chevrolet Cobalt and Saturn Ion among them — had too little detent torque to hold their position, so a jostle, a bump in the road, or a heavy keychain could rotate the switch out of “run” while driving. That cut power steering and braking and, fatally, disarmed the airbags so they would not deploy in the very crash that loss of control often produced. GM engineers identified the problem in 2002, during development, before the cars ever reached customers.^o

WHAT HAPPENED

In 2006 a GM engineer approved a redesigned switch but did not change its part number — and in an engineering organization the part number is the very mechanism by which the system knows something changed. The fix propagated as neither a revision nor a recall; defective cars kept selling, and crashes with non-deploying airbags were investigated piecemeal over years, no one connecting them to a switch the records said had “never changed.” The recall came only in 2014 — about 2.6 million vehicles, with 124 deaths attributed through the compensation fund. What finally broke the silence was a wrongful-death lawsuit, in which a family’s expert found the part had been quietly changed under the same number — the smoking gun GM’s own records had been built to miss.¹

THE INVESTIGATION

GM’s commissioned investigation by Anton Valukas (2014) found not a single villain but a culture that absorbed safety concerns until they dissipated — the “GM nod,” in which a room


agrees a thing should happen and then no one acts, and the “GM salute,” arms crossed and each person pointing elsewhere — and a fundamental failure to use the escalation processes the company already had on the books.² In 2015 GM paid a \$900 million federal forfeiture for concealing the defect; total penalties and settlements ultimately exceeded \$2.6 billion.³

THE CAPABILITY GAP

What was designed out of GM was not a part but a pathway — the institutional route by which a known safety problem reaches the decision to recall. The fix existed for eight years; the path from fix to recall did not, so the knowledge sat inert. The mechanism is mundane and all the more instructive for it: the part-number convention was the company’s own way of seeing what had changed, and quietly breaking it blinded the organization to its own action. Suppressing a signal need not be a conspiracy; here it was one engineer taking the path of least resistance through a process no one was assigned to guard.⁴

AFTERMATH & REFORM

Mary Barra, who became CEO as the recall broke, used the Valukas report to restructure GM’s safety decision-making — a global-safety leadership role, consolidated escalation channels, and a “Speak Up for Safety” program to give concerns a route upward that did not depend on one person’s persistence to survive. The case became a standard teaching example in governance and psychological safety: an organization can hold the fix for the better part of a decade and still fail to act if the channel that carries bad news upward has been allowed to fail quietly, with no one accountable for keeping it open.⁵



There was a fundamental failure to use the formal escalation processes that GM had established.

PARAPHRASING THE VALUKAS REPORT TO THE GM BOARD OF DIRECTORS, 2014

REFERENCES

1. A. R. Valukas, Report to the Board of Directors of General Motors Company Regarding Ignition Switch Recalls (Jenner & Block, 2014) — the 2002 identification of the defect and the switch-torque mechanism.
2. U.S. NHTSA investigation of the GM ignition switch (2014) and the GM recall record — 2.6 million vehicles; 124 deaths via the GM compensation fund.
3. Valukas Report (2014) — the unchanged part number, the “GM nod” and “GM salute,” and the failure to use established escalation processes (quoted).
4. U.S. Department of Justice, deferred-prosecution agreement with General Motors (2015) — \$900 million forfeiture for concealing the defect; total penalties/settlements exceeding \$2.6 billion.
5. Valukas Report (2014) and A. C. Edmondson, The Fearless Organization (Wiley, 2018) — the part number as an organizational signal, and the suppression of upward safety information.
6. U.S. NHTSA consent order with General Motors (2014, \$35M civil penalty) and the GM ignition-switch victims’ compensation program (K. Feinberg, administrator) — the regulatory penalty and the basis for the 124-death figure.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The GM ignition switch case is the canonical example of a corporate organizational structure that suppressed the upward flow of safety information by procedural design. The fix existed for eight years before the recall. The institutional path between the fix and the recall did not.

LENS APPROACH

The load-bearing lesson is change control. Engineers changed the ignition switch but kept the same part number, defeating configuration traceability so the field failures could never be tied back to a design change and the records insisted the part had “never changed.” The capability to build is configuration management that forbids altering a part without renumbering it, so the system can always see what changed and connect a field failure to the revision that caused it. LENS uses GM in LEN 4 for that change-control and part-number-as-signal discipline, and secondarily in LEN 7 and LEN 8 for the corporate-cover-up and Valukas-style retrospective accountability framing.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
INDUSTRIAL	→ D G	→ D1 D2 D3 D4 D5	5	5.4	4
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Treat the part number as a safety-critical signal: make changing a part without changing its number impossible by process, so the system cannot be blinded to its own revision. ▶ Specify the ignition switch’s detent torque as a hold-position requirement and verify it, so a known under-torque defect cannot pass into production as a tolerable quirk. ▶ Build the route from “fix approved” to “recall decision” as an explicit deliverable, not a path that depends on one person choosing to escalate.	▶ Correlate field crashes with non-deploying airbags as a population, so piecemeal incidents are connected to a common cause rather than investigated and closed one at a time. ▶ Run an escalation channel that carries bad news upward independent of any single person’s persistence, with a named owner accountable for keeping it open. ▶ Audit whether approved fixes have actually propagated to a recall or revision, closing the gap in which a known defect keeps shipping under an unchanged record.

REFLECTION QUESTIONS

1. What information channel in your organization carries the same load that GM’s part-number system did? Could it be silently bypassed?
2. Design the escalation deliverable that would have moved the GM ignition switch fix to a recall in 2006.
3. A single engineer changed the switch without changing its part number, and no one was assigned to guard that convention. Identify a process in your organization whose integrity depends on an unenforced convention, and name who should own enforcing it.

WHO BUILDS THIS systems & human-factors engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. TOOLS task analysis · design prototyping · governance frameworks · bias & evidence audits

Equifax Data Breach

G Governance & Trust **K** Knowledge & Institutional Memory

IMPACT 147 million Americans' personal data exposed; CEO resigned; ~\$700M settlement; foundational U.S. data-breach case

IN BRIEF

In 2017 attackers exploited a known, two-month-old vulnerability in Apache Struts on an Equifax web portal — a patch had been available, and Equifax's own security team had told IT to apply it; no one did. Over the next 2.5 months the attackers exfiltrated the personal data of 147 million Americans, and Equifax did not disclose the breach until September. A Senate investigation found systematically inadequate patching, an incomplete asset inventory (so the company did not know which systems needed the fix), and an incident-response function treated for years as a cost center. The CEO resigned and Equifax settled for about \$700 million. No single failure caused the breach; cumulative inadequacy across routine cybersecurity work did.

BACKGROUND

Equifax, one of the three U.S. credit bureaus, held the most sensitive financial data on virtually every American adult — a concentration of identity information that made any breach catastrophic by definition. One of its public web portals ran a version of Apache Struts with a known critical vulnerability for which a patch had been available for two months — and which Equifax's own security team had flagged for IT to apply, so the warning and the fix both existed inside the company and simply went unacted-upon.⁰

WHAT HAPPENED

The patch was not applied. Attackers exploited the vulnerability beginning in May 2017 and quietly exfiltrated personally identifying information on 147 million Americans over the next two and a half months, a window long enough that the theft was less a break-in than a sustained occupation. Equifax did not publicly disclose the breach until September, so the people whose identities were taken learned of it only well after the fact.¹

THE INVESTIGATION


The U.S. Senate Permanent Subcommittee on Investigations found Equifax's patching practices systematically inadequate and the company lacking "a comprehensive IT asset inventory" — so it could not reliably know which systems needed the patch, leaving the security team's warning with no map to act on.² Monitoring was weak, response delayed, and the incident-response architecture had been treated for years as a cost center rather than a capability worth funding; the CEO resigned and Equifax settled for roughly \$700 million.³

THE CAPABILITY GAP

Equifax is the institutional-cybersecurity case for cumulative inadequacy in routine work: patching, asset inventory, monitoring, response. Each function was below standard; none alone produced the breach; the combination did, the marginal weaknesses compounding into a single open door. The capability gap was the management of unglamorous, universally-agreed-necessary maintenance — exactly the work easy to defer because deferring it usually costs nothing, until the one time it costs everything, and a function starved as a cost center has no slack left on that day.⁴

AFTERMATH & REFORM

The settlement funded consumer compensation and credit monitoring, and the breach pushed patching discipline, asset inventory, and breach-disclosure timelines up the corporate agenda — elevating, after the loss, the unglamorous work that had been deferred before it.⁵ It was a cumulative-inadequacy failure: no dramatic single cause, just several routine defenses each left marginally inadequate, failing together on the day a determined attacker arrived to test all of them at once.



Equifax lacked a comprehensive IT asset inventory.

U.S. SENATE PERMANENT SUBCOMMITTEE ON INVESTIGATIONS, HOW EQUIFAX NEGLECTED CYBERSECURITY, MARCH 2019

REFERENCES

1. U.S. Senate Permanent Subcommittee on Investigations, How Equifax Neglected Cybersecurity and Suffered a Devastating Data Breach (2019) — the unpatched Apache Struts vulnerability.
2. The breach record — 147 million affected, exploitation from May 2017, public disclosure in September 2017.
3. Senate PSI (2019) — “Equifax lacked a comprehensive IT asset inventory” (quoted).
4. U.S. FTC / CFPB / state settlement (\$700 million, 2019) and the CEO’s resignation.
5. U.S. GAO, Actions Taken by Equifax and Federal Agencies in Response to the 2017 Breach, GAO-18-559 (2018).
6. Apache Struts CVE-2017-5638 advisory.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Equifax is the canonical institutional-cybersecurity case for cumulative inadequacy in routine work: patching, inventory, monitoring, response. Each function was below industry standard. None alone produced the breach. The combination was the breach.

LENS APPROACH

LENS uses Equifax in LEN 6 as a feedback-channel failure: the security team identified the Apache Struts patch and told IT to apply it — the operator-to-institution escalation existed, was correct, and died before reaching anyone with the asset map and the authority to act. The teaching point is not “patch your systems” but the engineered escalation deliverable: a flagged, agreed-upon fix has to be wired to an owner who can map it to every affected asset and is funded to close it before the gap costs everything. LEN 7 carries the institutional half — routine defenses run as a cost center have no slack on the bad day.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
TECHNOLOGY GOVERNMENT	→ G K →	D1 D2 D3 D4 D5 Navigating Sociotechnical Constraints	3	6.2	2

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Fund routine security maintenance — patching, monitoring, response — as a capability, not a cost center, so the unglamorous defenses have slack on the bad day.
- ▶ Build and maintain a comprehensive asset inventory so a known vulnerability can be mapped to every system that runs it.
- ▶ Wire the security team’s warnings to an owner with the authority and resources to apply a flagged patch promptly.

IN OPERATION / AFTER

- ▶ Audit the routine defenses together — patching, inventory, monitoring, response — since each marginally inadequate layer raises the odds the combination fails.
- ▶ Monitor for sustained, low-noise exfiltration, treating a months-long quiet window as the failure mode to detect, not just a single break-in.
- ▶ Hold breach-disclosure timelines short, so affected people learn of a theft of their identity without months of delay.

REFLECTION QUESTIONS

1. Identify a piece of routine work in your domain that is chronically deferred. What is the cumulative-inadequacy threshold?
2. Equifax did not know which assets ran Struts. Design the asset-inventory deliverable that an organization the size of Equifax should be able to produce on demand.
3. Equifax’s security team had flagged the patch, but the warning had no asset map to act on and no funded function to carry it out. What turns a known, agreed-upon fix into deferred work in your organization — and what would force it to be done before the one time it costs everything?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · knowledge bases · data pipelines

Chapter 10

Industry, Energy & Enterprise Systems — What Works — and the Frontier

When operations make capability visible on the floor.

The cord was there to be pulled, and pulling it was rewarded.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

INL Turbine–Control Upgrade — Low–Burden Cutover as a Human–Factors Finding

G Governance & Trust**K** Knowledge & Institutional Memory**H** Human–System Interface**EVIDENCE TIER** practice–synthesis–tier — source confidence flagged; future validation ongoing.

IMPACT An INL–affiliated study reported that operators were able to use the new digital turbine–control system without extensive additional training or rewriting of operating procedures — i.e., the human–factors verification–and–validation evidence supported a low–burden cutover from the legacy analog control

IN BRIEF

Nuclear–plant control–room modernization is one of the canonical C7 problems — legacy analog systems that have to be replaced with digital equivalents inside a regulatory regime that demands the safety case survive the transition. The Idaho National Laboratory’s Light Water Reactor Sustainability (LWRS) program has produced a body of technical reports and conference papers documenting the verification–and–validation work behind specific upgrades. The case here is a specific finding: in a study of turbine–control–system upgrade work, operators were able to use the new digital system without extensive additional training or rewriting of operating procedures. The human–factors V&V evidence supported a low–burden cutover. That finding is the small–tier complement to the big–tier LWRS program case the corpus also tracks: the program produces specific design–V&V case studies, not only program–level claims. The evidence base is INL technical reporting and OSTI–hosted conference papers, not independent academic evaluation; the tier flag is rendered under the title. Future validation will continue on whether the low–retraining finding generalizes beyond the studied subsystem and holds at multi–plant scale.

BACKGROUND

The nuclear–plant control room is the C7 transition problem with the safety–case constraint at the front: the regulator and the licensee both need to know that the upgraded digital control behaves equivalently to the legacy analog under the operating envelopes the safety case covers, and that the operator does not lose capability in the cutover. The LWRS program at INL exists to produce that evidence across the US fleet’s modernization needs. The specific study here is one of the program’s small–tier deliverables: a human–factors verification of a turbine–control–system upgrade.^o

THE INTERVENTION

The relevant finding the case rests on is that operators were able to use the new digital turbine-control system without extensive additional training or rewriting of operating procedures. That is a substantive human-factors-V&V claim: the cutover did not impose a retraining burden on the operator workforce that the schedule and the operational logic could not absorb. The study is a per-subsystem evaluation rather than a whole-control-room finding, and the LWRS program treats these small-tier evaluations as the building blocks of the larger fleet-modernization safety case.¹

HOW IT WORKED

The case is the small-tier complement to the big-tier LWRS program case the corpus also tracks. The big-tier case states the program-level claim: that nuclear plants can be modernized into digital control rooms with the safety case maintained. The small-tier case is the per-subsystem evidence that the program-level claim has to rest on — the individual design-V&V deliverables that, in aggregate, give the regulator and the licensee a reason to accept the modernization. Without the small-tier evidence the program-level claim is unfounded; without the program-level coordination the small-tier evidence does not roll up.²

THE EVIDENCE

The evidence-tier flag is load-bearing. The study sits in INL technical reporting and OSTI-hosted conference papers, not in independent academic evaluation. The low-burden-cutover finding is conditional on the specific subsystem, plant, and operator population studied; whether the finding generalizes to other subsystems in the same plant, to other plants in the fleet, or to other digital control products is the open question that the LWRS program itself acknowledges and continues to study. Future validation will continue as the program publishes follow-on reports.³

WHAT TRANSFERRED

What the case adds at the LENS layer is the small-tier example of verification-as-deliverable inside a regulated transition. The Domain-1 frame applies: the human-factors analysis of the operator's task under the digital cutover is a systems-analysis artifact that the regulatory regime consumes as part of its acceptance decision. The Domain-3 frame applies because the V&V evidence is itself the measurement instrument the program rests on. The case is a paired small-tier companion to Case 116 (Eurocat ATM) and sits inside the C7 thread alongside the big-tier failures.

The program-level claim cannot rest on un-rolled-up small-tier evidence. The per-subsystem V&V is the building block of the modernization safety case.

EDITORS' SYNTHESIS OF INL LWRS REPORTING.

REFERENCES

1. Ronald L. Boring (2014), "Human Factors Design, Verification, and Validation for Two Types of Control Room Upgrades at a Nuclear Power Plant," Proceedings of the Human Factors and Ergonomics Society Annual Meeting (HFES 2014), pp. 2295–2299 (ResearchGate publication 271728006).
2. Idaho National Laboratory, Light Water Reactor Sustainability Program reports on control-room modernization — series available via OSTI.
3. Nuclear Regulatory Commission (NUREG-0711), "Human Factors Engineering Program Review Model" — the regulatory framework the V&V deliverables are produced against.
4. O'Hara et al. (2008), "Human Factors Considerations with Respect to Emerging Technology in Nuclear Power Plants," NUREG/CR-6947 — peer-adjacent framing.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The INL turbine-control finding is a small-tier C7 verification-as-deliverable case inside the LWRs nuclear-modernization program: the human-factors V&V evidence supports a low-burden cutover. Evidence is INL technical reporting and OSTI conference papers, not independent academic evaluation; the generalization beyond the studied subsystem is the open question. Future validation ongoing.

LENS APPROACH

INL turbine-control upgrade is the C7 small-tier transition case (induced 7.1 and 3.1; LENS D1/PT1). LENS uses it in Domain 1 (Systems Analysis) for CLO-1 — the operator-task analysis under the digital cutover is the systems-analysis artifact the regulator consumes — and in Domain 4 (Test and Evaluation) for CLO-4, since the V&V evidence is itself the measurement instrument the program rests on. Pair with Case 116 (Eurocat) for the small-tier C7 thread and with the LWRs big-tier program case already in the corpus.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
NUCLEAR POWER	G K H	D1 D2 D3 D4 D5	1	7.1	1-4
CONTROL-ROOM MODERNIZATION	→	→ Systems Analysis			
HUMAN FACTORS					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat each subsystem upgrade as a per-subsystem human-factors V&V deliverable, not as a sub-task of the larger control-room modernization claim.
- ▶ Specify the operator-task analysis the cutover has to satisfy before the upgrade ships, so the V&V evidence has a target the regulator can audit.
- ▶ Document the conditioning factors — plant, operator population, subsystem boundary — explicitly, so the generalization question is honestly framed in the report.

IN OPERATION / AFTER

- ▶ Roll the per-subsystem V&V studies into the program-level fleet-modernization safety case deliberately; the program-level claim cannot rest on un-rolled-up small-tier evidence.
- ▶ Carry the practitioner-tier flag into any downstream citation; the small-tier finding is conditional on the studied subsystem and population.
- ▶ Track durability of the low-retraining finding under personnel rotation, procedure updates, and subsequent upgrades; the cutover is not finished when the digital system is in service.

REFLECTION QUESTIONS

1. Identify a control-system upgrade in your context whose cutover burden on the operator is the decisive constraint. What is the per-subsystem V&V deliverable that would settle the burden question, and what is its scope?
2. Specify the rollup logic from per-subsystem V&V evidence to a program-level modernization claim. The INL pattern teaches because the small-tier studies are the building blocks of a fleet-level claim — what would the equivalent be in your context?
3. The case is INL technical reporting, not independent academic evaluation. What additional evidence (independent audit, multi-plant replication, durability tracking under personnel rotation) would you require before treating the low-retraining finding as generalizable?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · usability testing · cognitive walkthroughs

CASE 156

NUCLEAR ENGINEERING

CONTROL-ROOM HUMAN FACTORS

RESEARCH AND DEVELOPMENT

2010 – PRESENT

INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet

D Designed Out **H** Human-System Interface **K** Knowledge & Institutional Memory

IMPACT The US Department of Energy's Light Water Reactor Sustainability (LWRS) program, executed at Idaho National Laboratory in partnership with utilities, has produced a structured research-and-pilot record for modernizing analog control-room instrumentation in the existing US nuclear fleet — pilot-scale evidence covering hybrid digital/analog operator interfaces, human-factors validation, and qualification pathways for digital instrumentation

IN BRIEF

The US commercial nuclear fleet is dominated by reactors originally licensed in the 1970s and 1980s with analog instrumentation and control-room layouts of that vintage. Aging analog components, vendor obsolescence, and the operational case for digital instrumentation make modernization a foreseeable sustainment requirement; the regulatory environment (Nuclear Regulatory Commission Regulatory Guide 1.180, Branch Technical Position 7-19, and related guidance) makes the qualification path for digital instrumentation deliberately stringent. The US Department of Energy's Light Water Reactor Sustainability (LWRS) program, executed at Idaho National Laboratory in partnership with US utilities, runs the federally-funded research-and-pilot work that supports the fleet's modernization decisions. The program has produced a structured record covering hybrid digital/ analog operator interfaces, human-factors validation studies in full-scope simulators, and qualification-pathway research for safety-related digital instrumentation. The honest hedge survives clearly: LWRS observations are pilot-scale, the sample of plants that have implemented major modernizations is small, and the operational-fleet evidence at scale is forward-looking. The case is teachable on the structured sustainment-research form — a federally-funded research program operating across decades to support an industry's aging-fleet decisions — and closes a C7 zero-state in the induced framework, paired with Cases 114, 115, 174 as the v2 aging-system quartet.

BACKGROUND

The US commercial nuclear fleet operates roughly 90 reactors originally licensed in the 1970s and 1980s. The control rooms of that era were built around analog instrumentation — strip-chart recorders, hardwired indicators, control-panel meters — designed for the operational envelope and the human-factors assumptions of the period. Aging analog components, vendor obsolescence, and the operational case for digital instrumentation

(better diagnostics, lower maintenance burden, integration with plant computer systems) make modernization a foreseeable sustainment requirement across the fleet.⁹

THE INTERVENTION

The regulatory environment is deliberately stringent. Nuclear Regulatory Commission guidance — Regulatory Guide 1.180 on electromagnetic and radio-frequency compatibility, Branch Technical Position 7-19 on common-cause failure analysis for digital instrumentation and control, and the broader Standard Review Plan Chapter 7 — sets the qualification path for safety-related digital instrumentation. The stringency is intentional: digital systems can introduce failure modes (common-cause software faults, cyber exposure) that analog systems do not, and the regulatory regime is designed to keep those failure modes bounded as plants modernize.¹

HOW IT WORKED

The US Department of Energy's Light Water Reactor Sustainability (LWRS) program is the federally-funded research-and-pilot work that supports the fleet's modernization decisions. Executed at Idaho National Laboratory in partnership with US utilities and EPRI, LWRS operates across a multi-decade horizon with annual research-report deliverables. The control-room modernization research line — including the Human Systems Simulation Laboratory at INL with full-scope plant simulators — covers hybrid digital/analog operator-interface designs, human-factors validation studies, and qualification-pathway research that supports utility-level licensing submissions.²

THE EVIDENCE

The honest hedge has to be visible. LWRS observations are pilot-scale: full-scope simulator studies, single-plant pilot implementations, and structured human-factors experiments with operator participants from utility partners. The sample of US plants that have implemented major control-room modernizations is small relative to the fleet, and the operational-fleet evidence at scale — fleet-wide reliability, fleet-wide human-error rate, fleet-wide maintenance burden under hybrid digital/analog control rooms — is forward-looking rather than retrospective. The program's own reports characterize the research at this evidence tier, and the case carries the same honesty.³

WHAT TRANSFERRED

What the case teaches at the LENS framing layer is the structured sustainment-research form — a federally-funded research program operating across decades in partnership with industry and a regulator, producing the research-and-pilot record that supports licensing decisions on aging-fleet modernization. The form is the sustainment-engineering analog of the FAA Aging Aircraft Program (Case 114): a long-horizon institutional discipline that produces the technical record that aging-fleet decisions can rest on. With NextGen/ADS-B (Case 115) and Y2K (Case 174), LWRS completes the v2 aging-system quartet that closes the C7 zero-state in the induced framework. The LWRS instance is where the evidence is most explicitly pilot-scale, and the case carries that as the program's tier acknowledgement rather than as a weakness to smooth.⁴

The institutional form is what transfers: federal research program plus industry partnership plus regulator engagement, operating across decades to produce a technical record aging-fleet decisions can rest on.

EDITORS' SYNTHESIS OF THE LWRS ANNUAL REPORT SERIES.

REFERENCES

1. Idaho National Laboratory, Light Water Reactor Sustainability (LWRS) program annual reports (2010 – present) — primary research-and-pilot record.

2. Nuclear Regulatory Commission, Regulatory Guide 1.180, "Guidelines for Evaluating Electromagnetic and Radio-Frequency Interference in Safety-Related Instrumentation and Control Systems."

3. Nuclear Regulatory Commission, Branch Technical Position 7-19, "Guidance for Evaluation of Diversity and Defense-in-Depth in Digital Computer-Based Instrumentation and Control Systems."

4. O'Hara, Higgins, Brown, Fink, Persensky, Lewis, Kramer, & Szabo (2008), "Human Factors Considerations with Respect to Emerging Technology in Nuclear Power Plants," NUREG/CR-6947 — foundational human-factors backdrop.

5. Electric Power Research Institute, control-room modernization technical reports — industry-side sustainment record.

THE LEARNING ENGINEERING LENS

LE INSIGHT

LWRS is the structured sustainment-research case — federally- funded research at INL in partnership with utilities and the regulator, producing the research-and-pilot record that supports aging-fleet modernization decisions across a multi- decade horizon. The observations are pilot-scale; the operational-fleet evidence at scale is forward-looking. The hedge is part of the case.

LENS APPROACH

LWRS control-room modernization is the structured sustainment-research case (induced 7.1; LENS D3/PT4) — Domain 1 for the long-horizon partnership; Domain 5 for the federal-research + industry + regulator triple structure; Domain 3 for the hybrid digital/analog research line. Closes C7 with Cases 114, 115, 174.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
NUCLEAR ENGINEERING	D H K	D1 D2 D3 D4 D5	4	7.1	1-3
CONTROL-ROOM HUMAN FACTORS	→	→ Human-System Collaboration			
RESEARCH AND DEVELOPMENT					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Pair the research program with full-scope simulator infrastructure (the INL Human Systems Simulation Laboratory) so the human-factors validation work runs in the operational envelope rather than in abstracted laboratory conditions.
- ▶ Structure the partnership across federal research (DOE/INL), industry (utilities, EPRI), and regulator (NRC) so the research record is recognized by all three parties — the licensing submissions ultimately rest on that recognition.
- ▶ Carry the pilot-scale evidence tier honestly in the program's reports; LWRS characterizes its research at

IN OPERATION / AFTER

- ▶ Track fleet-scale outcomes as utility implementations proceed: fleet-wide reliability, human-error rate, maintenance burden under modernized control rooms. The forward-looking evidence will accumulate over the next two decades; the program is the institutional carrier of that accumulation.
- ▶ Carry the LWRS case in pair with FAA Aging Aircraft (Case 114) as the sustainment-engineering analog at multi-decade horizon, and with Y2K (Case 174) and NextGen (Case 115) as the broader aging-system quartet.

that tier, and downstream practitioners reading the work should read it the same way.

- Treat the structured-sustainment-research institutional form as the transferable artifact: federal research program plus industry partnership plus regulator engagement, operating across decades to produce the technical record aging-fleet decisions can rest on.

REFLECTION QUESTIONS

1. Identify an aging fleet of long-lived assets in your domain whose modernization decisions rest on research that does not yet exist at fleet scale. What is the analog of a federally-funded research program — institutional partner, industry partner, regulator partner — that could operate across the horizon the decisions actually need?
2. Specify the full-scope simulator (or equivalent in-envelope test infrastructure) the program would need so its human-factors validation work runs in operational conditions rather than abstracted laboratory ones. What is the cost of building it, and what does its absence mean for the evidence the research produces?
3. LWRS evidence is pilot-scale and the program reports it as such. Identify a program in your domain whose evidence tier should be acknowledged in its reporting but is not. What institutional discipline keeps the evidence-tier honesty visible — to the program, to industry partners, and to the regulator?

WHO BUILDS THIS systems & human-factors engineering · human factors & interaction design · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · usability testing · cognitive walkthroughs · knowledge bases · data pipelines

Chapter II

Disaster Prevention & Recovery — What Fails

Before the event, prevention regimes that thinned out; after it, response plans that met reality.

Every barrier had been inspected. What failed was the system deciding what counted as a barrier.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

EDITOR'S NOTE · BEFORE AND AFTER

This part reads the corpus along a different axis than the domain parts that precede it: the lifecycle of a disaster. Every case here has two clocks. Before the event, there is a prevention regime — a certification, an inspection cadence, a defense-in-depth argument — whose quiet thinning is the real story. After the event, there is a response and a recovery, and the institutions that are (or are not) built so the same failure cannot recur. The chapter is therefore ordered in two arcs, marked below: prevention failures first, then response-and-recovery failures; and in the What Works chapter, prevention engineered as a deliverable, then recovery run as a discipline. Cases whose primary lessons belong to a domain — an aviation accident, a clinical harm — remain in their domain parts; what qualifies a case for this part is that its load-bearing lesson is the pre/post lifecycle itself.

BEFORE THE EVENT — PREVENTION

Bhopal

T Training Gap **K** Knowledge & Institutional Memory **N** Normalization of Deviance **G** Governance & Trust

IMPACT ≈ 15,000–20,000 killed; ≈ 500,000 injured; worst industrial disaster in history

IN BRIEF

On the night of 2–3 December 1984, about forty tons of methyl isocyanate gas escaped from a Union Carbide pesticide plant in Bhopal, India — the worst industrial disaster in history. Thousands died within hours; estimates of total deaths run to 15,000–20,000, with roughly half a million exposed or injured. Safety systems had been off-line for months, the plant was understaffed, and workers were inadequately trained to recognize or handle the emergency. Investigators traced the catastrophe to operating errors, design flaws, maintenance failures, training deficiencies, and cost-cutting that endangered safety. Bhopal catalyzed the creation of the U.S. Chemical Safety Board and reshaped industrial-safety regulation for decades. It is the largest-magnitude capability-and-governance failure on record.

BACKGROUND

Union Carbide's pesticide plant in Bhopal, India, stored methyl isocyanate (MIC) — an extraordinarily toxic intermediate — in bulk, holding a lethal hazard in tanks beside a populated city. By 1984 the plant was running under heavy cost pressure: understaffed, with several key safety systems out of service for months, and workers inadequately trained to handle an MIC emergency or read its warning signs. Each economy was individually defensible on a ledger; together they thinned every layer of defense the process depended on.⁰

WHAT HAPPENED

On the night of 2–3 December 1984, water entered an MIC storage tank and triggered a runaway reaction; the safety systems that should have contained it were non-operational, so the one event the plant existed to prevent met no working barrier on its way out. About forty tons of gas vented over the sleeping city. Thousands died within hours; estimates of total deaths run to 15,000–20,000, and roughly half a million people were exposed or injured — the worst industrial disaster in history, its toll set by who happened to be downwind.¹

THE INVESTIGATION

Investigations found the catastrophe “resulted from operating errors, design flaws, maintenance failures, training deficiencies and economy measures that endangered safety” — a list with no single villain, which is precisely what made it hard to govern.² Human-factors analysis placed Bhopal alongside Three Mile Island in its neglect of the human element, and the U.S. Chemical Safety Board would later find ineffective employee training an underlying cause

in nine of its first twenty-three chemical-incident investigations — a pattern that traces to Bhopal and shows the same gap recurring long after the lesson was available.³

THE CAPABILITY GAP

Bhopal is the largest-magnitude capability-and-governance failure on record, and a multi-layer one: training, maintenance, design, staffing, and oversight had all degraded together, and no layer above the plant was accountable for the whole. Because the erosion was spread across layers, no single inspection or metric saw it, and each degraded layer made the next one matter more. The capability to operate an extraordinarily hazardous process safely had been hollowed out by cost-cutting, and the governance that should have caught the hollowing did not exist to ask whether the whole was still safe.⁴

AFTERMATH & REFORM

The disaster reshaped industrial safety worldwide and, in the United States, catalyzed the creation of the Chemical Safety Board — an INPO-equivalent for industrial chemistry — and the process-safety regime that followed, an oversight layer built to own exactly the whole that no one had owned at Bhopal.⁵ The book's recurring arc runs through Bhopal in its starkest form: a catastrophe forces into being the institution the industry should have built before it, at a price no institution-building exercise should ever have to cost.



Operating errors, design flaws, maintenance failures, training deficiencies and economy measures that endangered safety.

NEW YORK TIMES INVESTIGATION, 1985

REFERENCES

1. Union Carbide and government investigation reports (1985) — MIC storage, the disabled safety systems, and plant understaffing.
2. Accounts of the 2–3 Dec. 1984 release — the contested toll (thousands of immediate deaths; 15,000–20,000 total estimates; 500,000 exposed). (Figures vary widely across sources; see AUDIT.)
3. New York Times investigation (1985) — “operating errors, design flaws, maintenance failures, training deficiencies and economy measures that endangered safety” (quoted).
4. N. Meshkati, “Human Factors in Large-Scale Technological Systems’ Accidents,” *Industrial Crisis Quarterly* (1991); U.S. CSB training-cause pattern.
5. P. Shrivastava, *Bhopal: Anatomy of a Crisis* (1992); C. Perrow, *Normal Accidents* (1984).
6. The creation of the U.S. Chemical Safety Board and the post-Bhopal process-safety regime.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Bhopal is the largest-magnitude capability-and-governance failure on record. It is also the catalyst for the creation of the U.S. Chemical Safety Board, an INPO-equivalent for industrial chemistry. The pattern that follows nearly every case in this book — diagnose, then engineer remediation at scale — runs through Bhopal in canonical form.

LENS APPROACH

LENS names the specific capability gap, not “governance matters” in the abstract: at a transferred and cost-pressured plant, the maintenance and operating capability had been hollowed out — refrigeration and scrubber safety systems off-line, maintenance deferred, staff and training cut, deviance quietly normalized. Each eroded layer was individually tolerated, so the failure is a multi-layer drift (induced 7.4) that no single inspection or metric could see. The teaching point is the multi-layer-drift analysis and the treatment of maintenance capability as a deliverable that must be sustained, not a ledger line available for trimming.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
INDUSTRIAL	→ T K N G →	D1 D2 D3 D4 D5	4	7.4	1·5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat each safety system, staffing level, and training requirement as a non-negotiable layer of defense, not a cost line available for trimming under pressure.
- ▶ Assign a single accountable owner for the integrated hazard, so degradation spread across layers cannot fall between everyone’s responsibilities.
- ▶ Require that operators be trained to recognize and respond to the specific emergency the process can produce before the process is run with that hazard in bulk.

IN OPERATION / AFTER

- ▶ Audit the whole defense-in-depth posture together rather than layer by layer, since each degraded barrier silently raises the stakes on the next.
- ▶ Track how long any safety system stays off-line and gate continued operation on its restoration, treating extended downtime as an unacceptable condition.
- ▶ Sustain an oversight layer — an INPO- or CSB-equivalent — with authority and reach above the individual plant to own the integrated risk.

REFLECTION QUESTIONS

1. Bhopal produced the CSB. What is the institution your domain has not yet built that a comparable disaster would force into existence?
2. Identify a plant or facility in your domain that has had safety systems off-line for an extended period. What is the measurement gap that allowed it?
3. At Bhopal training, maintenance, design, staffing, and oversight all degraded together while each looked acceptable alone. Whose job in your organization is to judge whether the whole is still safe — and what would they have to see to declare that it is not?

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · safety science & organizational analysis · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · incident analysis · safety dashboards · governance frameworks · bias & evidence audits

Davis–Besse Reactor Head Corrosion

N Normalization of Deviance **K** Knowledge & Institutional Memory **G** Governance & Trust

IMPACT Football-sized cavity discovered in the reactor pressure-vessel head; near-miss; ~\$600M recovery and extended outage

IN BRIEF

During a 2002 refueling outage at the Davis–Besse nuclear plant in Ohio, workers found a football-sized cavity eaten through six inches of the carbon-steel reactor-vessel head by leaking boric acid, leaving only a thin layer of stainless cladding between the cavity and reactor coolant at 2,200 psi — a serious near-miss for a loss-of-coolant accident. The leakage had been visible for years and inspections deferred; FirstEnergy had won an NRC deferral of a mandatory inspection to fit its refueling schedule, and the NRC’s Inspector General later found the decision had weighted economics over safety. Davis–Besse is the canonical post-Three-Mile-Island near-miss — evidence that even an industry that built INPO to engineer safety can let regulator–operator dynamics erode it.

BACKGROUND

After Three Mile Island, U.S. commercial nuclear power had built — through INPO and a strengthened NRC — a reputation for engineering safety as a system. Davis–Besse, a pressurized-water reactor in Ohio, operated inside that regime, with a known industry-wide hazard: boric acid leaking from cracked nozzles could corrode the carbon-steel reactor-vessel head.⁰ The hazard was generic and understood across the fleet, which is what makes the case sharp: the danger was not a surprise mechanism but a recognized one the post-TMI regime was precisely supposed to keep inspected and contained.

WHAT HAPPENED

During a refueling outage in early 2002, workers found a cavity about the size of a football eaten clean through six inches of the carbon-steel head, leaving only a thin layer of stainless-steel cladding between the corrosion and reactor coolant at 2,200 psi. Had the cladding given way, the result would have been a medium-break loss-of-coolant accident with a badly degraded safety margin.¹ The thin remaining cladding was the entire distance between routine operation and the kind of accident the entire regulatory regime existed to prevent — a margin measured in a fraction of an inch of unintended material rather than in any engineered barrier.

THE INVESTIGATION

The boric-acid leakage had been observable for years, and inspections had been deferred. FirstEnergy had requested — and the NRC had granted — a deferral of a mandatory inspection (NRC Bulletin 2001–01) so it would align with the plant’s February 2002 outage.² The

NRC's Office of Inspector General later found the agency had inappropriately weighted the utility's economic arguments over safety — its oversight, the OIG concluded, had not been adequate to ensure that safety would not be compromised.³ The deferral was not a covert lapse but a documented decision both parties signed off on, which is the unsettling part: the erosion happened through the regulator's normal process, in the open, agreed to be safe to wait.

THE CAPABILITY GAP

Davis-Besse is a near-miss in the dimension above operations: the regulator-operator relationship. The same industry that had shown, through INPO, that it knew how to engineer safety let a known corrosion mechanism go uninspected because inspecting it was inconvenient and expensive. Regulatory capture — the regulator adopting the operator's economic frame — is a capability failure at the institutional layer above the plant.⁴ The plant's own engineering competence was never the missing piece; what failed was the independence of the layer meant to overrule a utility's schedule when safety required it, and that layer had quietly adopted the schedule as its own.

AFTERMATH & REFORM

INPO and the NRC restructured significant parts of their inspection regimes, and reactor-head inspection requirements were tightened across the fleet. The lesson pairs with the book's TMI / INPO arc: institutional capability is not built once. It erodes if it is not re-engineered — and the erosion is quietest where the regulator and the regulated agree it is safe to wait.⁵ Tightening the head-inspection requirement across the fleet conceded that a mandatory inspection had been a real barrier all along — one the deferral process had been allowed to treat as negotiable against a refueling schedule.

1/4"

OF STAINLESS CLADDING REMAINED

between a 6-inch corrosion cavity and reactor coolant at 2,200 psi

DAVIS-BESSE — THE POST-TMI NEAR-MISS

The NRC's actions were not adequate to ensure that safety would not be compromised.

PARAPHRASING THE NRC OFFICE OF INSPECTOR GENERAL DAVIS-BESSE REPORT, 2002

REFERENCES

1. U.S. NRC Office of Inspector General, NRC's Regulation of Davis-Besse Regarding Damage to the Reactor Vessel Head (Case No. 02-03S, December 2002) — the post-TMI regulatory regime and oversight failures.
2. NRC event records and OIG (2002) — the 6-inch corrosion cavity, the remaining cladding, and the 2,200-psi coolant margin.
3. NRC OIG (2002) — economic arguments weighted over safety, and the finding that NRC oversight had not been adequate to ensure safety (paraphrased).
4. D. Lochbaum / Union of Concerned Scientists analysis (2002); U.S. GAO, GAO-04-415 (2004).
5. J. V. Rees, Hostages of Each Other (1994) — INPO and the fragility of institutional safety capability.

3. NRC Bulletin 2001-01 and the FirstEnergy inspection-deferral decision aligned to the February 2002 outage.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Davis-Besse is the case for how regulator-operator dynamics can erode the capability of an industry that had previously demonstrated, through INPO, that it knew how to engineer safety. Regulatory capture is a capability failure at the institutional layer above operations.

LENS APPROACH

LENS uses Davis-Besse in LEN 7 to study regulatory-capture dynamics and in LEN 8 to compare with INPO's earlier success. The case is a reminder that institutional capability requires sustained re-engineering.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ENERGY	→ N K G	→ D1 D2 D3 D4 D5	5	5.4	5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Make safety-critical inspections of known generic hazards non-deferrable, so a mandatory check cannot be traded against a refueling schedule.
- Require any deferral decision to be argued on safety margin alone, with the utility's economic case explicitly excluded from the record.
- Preserve the independence of the oversight layer so it can overrule an operator's timeline rather than adopt it.

IN OPERATION / AFTER

- Audit granted deferrals for whether the regulator reasoned from safety or from the operator's economics, and flag drift toward the latter.
- Track observable degradation (like boric-acid leakage) against inspection currency, so a known mechanism cannot run uninspected for years.
- Re-engineer institutional safety capability on a cadence, treating post-incident regimes like INPO as maintained, not permanent.

REFLECTION QUESTIONS

- Identify a regulator-operator relationship in your domain in which the regulator may be at risk of accepting the operator's economic argument over its own safety judgment. What signal would surface it?
- Design the institutional control that would prevent a Davis-Besse-style deferral from being granted.
- The erosion here happened openly, through the regulator's normal process. What audit would distinguish a defensible deferral from one in which the regulator has quietly adopted the operator's schedule as its own?

WHO BUILDS THIS safety science & organizational analysis · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. TOOLS incident analysis · safety dashboards · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

Upper Big Branch Mine Explosion

N Normalization of Deviance **G** Governance & Trust **K** Knowledge & Institutional Memory

IMPACT 29 killed in West Virginia coal mine; MSHA found systematic falsification of safety records; first U.S. mining-industry CEO convicted of a federal mine-safety charge (misdemeanor)

IN BRIEF

On 5 April 2010 coal dust and methane ignited at Massey Energy's Upper Big Branch mine in West Virginia, killing 29 miners — the worst U.S. mine disaster in forty years. Investigators found Massey had kept two sets of records: an internal log of actual conditions and a separate, clean log for federal inspectors. Foremen were directed to suppress methane readings, disable monitors, and falsify pre-shift examinations — not as an aberration but as a stable routine across years and management layers. CEO Don Blankenship was convicted of a misdemeanor conspiracy to violate mine-safety standards — the first U.S. mining CEO criminally convicted on such a charge. Upper Big Branch is the dataset's clearest case of a measurement system engineered deliberately to defeat the regulator.

BACKGROUND

Massey Energy's Upper Big Branch mine in West Virginia operated under federal safety rules enforced through inspections and the records the mine kept. Massey kept two: an internal log of actual conditions, and a separate, clean log maintained for the inspectors.⁰ The enforcement regime depended entirely on the records reflecting reality, so a second, sanitized set of books did not merely break a rule — it disabled the very mechanism by which the regulator was supposed to see the mine at all.

WHAT HAPPENED

On 5 April 2010 accumulated coal dust and methane ignited and tore through the mine, killing twenty-nine miners — the worst U.S. coal disaster in forty years. The conditions that fed the blast — high methane, inadequate ventilation, dust not rendered inert — were the very ones the clean, inspector-facing records had been built to hide.¹ The records had done their intended work right up to the blast: they kept the conditions that killed twenty-nine men out of the regulator's view precisely while those conditions were building toward ignition.

THE INVESTIGATION

MSHA's investigation and a parallel U.S. Department of Justice probe found foremen instructed to suppress methane readings, disable monitors, and falsify pre-shift examinations.² Massey CEO Don Blankenship was eventually convicted of a misdemeanor count of conspiring to willfully violate mine-safety standards — the first U.S. mining-industry CEO criminally convicted on a mine-safety charge — though acquitted on the felony counts.³


That the instructions ran down to foremen and the conviction reached up to the CEO marks the practice as vertical, not local: the deception was operated at the working level and conceived above it, spanning the management layers in between.

THE CAPABILITY GAP

The dual-records architecture was not the work of a few bad foremen; it was a stable institutional practice sustained across years and multiple management layers. The capability gap was not in the miners but in the executive ranks that designed and operated a measurement system specifically to defeat the regulator — which makes it the book's clearest case of measurement engineered as deception.⁴ Where other cases show measurement drifting or pointed at the wrong dimension, here it was deliberately constructed to mislead, so no improvement in the regulator's reading of the official records could ever have helped — the records were the lie.

AFTERMATH & REFORM

The case anchors the argument that decision-grade evidence needs structural protection from the institution that produces it, when that institution has a stake in what the evidence says. Blankenship's conviction set a marker for corporate-officer accountability — and left open the regulator-side question: what architecture would have detected two sets of books before twenty-nine people died?⁵ Holding a CEO criminally accountable raised the cost of designing such a system, but accountability after the fact is not detection; the unanswered question is what independent signal could have exposed the divergence between the two logs while the mine was still running.



Massey kept two sets of books — one for inspectors, one for itself.

PARAPHRASING FEDERAL INVESTIGATORS, MSHA UPPER BIG BRANCH REVIEWS (2011 – 2012)

REFERENCES

1. U.S. MSHA, Internal Review of MSHA's Actions at Upper Big Branch and the accident investigation (2011–2012) — the dual records and the 29 deaths.
2. Governor's Independent Investigation Panel (J. Davitt McAteer, 2011) — mine conditions: methane, ventilation, and coal-dust inerting.
3. MSHA and U.S. DOJ findings — suppressed methane readings, disabled monitors, and falsified pre-shift examinations.
4. United States v. Blankenship (S.D.W. Va., 2015–2016) — the misdemeanor conviction and felony acquittals.
5. H. Berkes / NPR investigative reporting on Massey Energy.
6. A. Hopkins, Disastrous Decisions: The Human and Organizational Causes of the Gulf of Mexico Blowout (2012) — corporate decision-making and safety (comparative).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Upper Big Branch is the case for deliberately engineered measurement deception. The dual-records practice was not error. It was capability design — by management, against the regulator. The case anchors the LENS argument that decision-grade evidence requires structural protection against the institution that produces it having a stake in what it reports.

LENS APPROACH

LENS uses Upper Big Branch in LEN 4 as the most explicit case for measurement-system protection and in LEN 7 as a corporate-criminal accountability case. Studio projects examine what regulator-side architecture would have detected the dual-records system earlier.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ENERGY INDUSTRIAL	→ N G K	→ D1 D2 D3 D4 D5	2	2.2	4
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Source safety-critical data from monitors the operator cannot disable or edit, so a clean official log cannot be constructed by hand.
- ▶ Design measurement so the producing institution has no unilateral control over the record the regulator relies on.
- ▶ Build cross-checks that compare independent streams, making a single sanitized set of books detectably inconsistent.

IN OPERATION / AFTER

- ▶ Establish whistleblower and anomaly-detection channels that can surface a divergence between actual and reported conditions while operations continue.
- ▶ Audit for the structural conditions that make dual records possible — operator-controlled monitors, no independent verification — not just for violations.
- ▶ Pair corporate-officer accountability with detection architecture, since raising the cost of deception does not by itself reveal it in time.

REFLECTION QUESTIONS

1. Where in your domain could two sets of records plausibly be kept? What architectural change would make that impossible?
2. What does it mean for a measurement system to be “structurally protected” from the institution that produces it?
3. Accountability after the fact is not detection. What independent signal could have exposed the divergence between the two logs while the mine was still running?

WHO BUILDS THIS safety science & organizational analysis · governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS incident analysis · safety dashboards · governance frameworks · bias & evidence audits · knowledge bases · data pipelines

Deepwater Horizon

T Training Gap **N** Normalization of Deviance **K** Knowledge & Institutional Memory

IMPACT 11 killed; largest marine oil spill in U.S. history; \$65B+ in damages

IN BRIEF

On 20 April 2010 the Macondo well blew out beneath the Deepwater Horizon rig in the Gulf of Mexico, killing 11 workers and releasing roughly 4.9 million barrels of oil — the largest marine spill in U.S. history, and more than \$65 billion in eventual costs to BP. The crew had misread the negative-pressure test, the primary check of well integrity, and kept working a well that was no longer sealed. The National Commission traced the blowout to a series of identifiable mistakes by BP, Halliburton, and Transocean — systematic failures in risk management that cast doubt on the safety culture of the whole industry. Every defense — procedure, training, supervision, contractor coordination, equipment — had drifted independently until none caught the others. It is the book's canonical multi-layer normalization failure. Four major investigations — the National Commission, the joint Coast Guard / BOEMRE inquiry, the Chemical Safety Board's process-safety review, and the Deepwater Horizon Study Group at Berkeley — each surfaced a distinct facet of the same drift; the disagreements among them are themselves load-bearing for the case.

BACKGROUND

The Deepwater Horizon was drilling BP's Macondo exploratory well in the Gulf of Mexico when, on 20 April 2010, the well blew out. Eleven workers were killed, the rig burned and sank, and roughly 4.9 million barrels of oil flowed into the Gulf over 87 days — the largest marine oil spill in U.S. history, eventually costing BP more than \$65 billion.^o The well sat at the end of a long chain of contractors and schedules, each operating to its own tolerance, so that the 87 days of uncontrolled flow stand as a measure not of one mistake but of how far the accumulated drift had to be unwound before the Gulf was sealed again.

WHAT HAPPENED

The proximate trigger was a misread safety check. The crew ran a negative-pressure test — the primary test of whether the well was sealed — got anomalous readings, accepted an incorrect explanation for them, and proceeded to displace the heavy drilling mud. The well was not sealed; hydrocarbons surged up the riser to the rig and ignited.¹ The anomalous pressure was the well speaking plainly that it remained live, but the explanation the crew accepted let the displacement proceed on schedule, and once the heavy mud was gone nothing stood between the reservoir and the rig but a barrier that had already failed.

THE INVESTIGATION

Four major investigations reached the wreck and pulled on different threads. The National Commission concluded that a cascade of mistakes across BP, Halliburton, and Transocean,


not mechanical failure alone, was the root cause, and that the mistakes revealed “such systematic failures in risk management that they place in doubt the safety culture of the entire industry.”² The joint Coast Guard / BOEMRE inquiry traced the blowout-preventer failure to a buckled drill pipe and a maintenance backlog that the leasing-and-safety regulator had not caught. The Chemical Safety Board’s separately published process-safety review made the load-bearing distinction explicit: BP’s lost-time-injury rate had been industry-leading on the rig, yet personal-safety metrics had no purchase on the catastrophic process-safety risks that produced the blowout — the same inversion the CSB had already documented at Texas City. Government and academic reviews found training that had not covered the well-control situation the crew faced, an unclear chain of command, and a cascade of failed defenses that BP’s own accident investigation report summed up as “a complex and interlinked series of mechanical failures, human judgments, engineering design, operational implementation and team interfaces.”³ That the Commission reached past the rig to the whole industry’s safety culture marked the failure as structural rather than local: the same drift, it judged, was latent wherever the same pressures and the same defenses operated unexamined.

THE CAPABILITY GAP

Deepwater Horizon is the canonical multi-layer failure: procedure, training, supervision, contractor coordination, and equipment had each drifted independently until none caught the others. The negative-pressure test was the trigger, but the system had been carrying accumulated procedural debt for years, and the capability to recognize an unsafe well-state was simply not present at the moment it was needed.⁴ Because each defense had degraded inside its own tolerance, no single review of any one layer would have flagged a crisis; the danger lived in the correlation between layers, which is precisely the property no individual procedure was designed to watch. The cement job, the blowout preventer’s design and maintenance, the well-control culture before the incident, and the regulator’s split mandate were each, on their own, a partial story; the failure was the alignment among them, and no defense layer had been instrumented to read that alignment in time.

AFTERMATH & REFORM

The blowout drove a restructuring of offshore regulation — the Minerals Management Service was broken up and replaced by the Bureau of Safety and Environmental Enforcement (BSEE) and the Bureau of Ocean Energy Management (BOEM), separating the leasing function from the safety function so revenue pressure could no longer erode well-integrity enforcement — and drilling-safety and well-integrity rules were tightened, while BP paid tens of billions in penalties and settlements. The deeper lesson is the normalization one: no single layer failed catastrophically on its own; each had quietly drifted within tolerance until the day the tolerances lined up.⁵ Splitting the regulator conceded that an agency both leasing acreage and policing safety carried a built-in conflict, and tightening the well-integrity rules conceded that the negative-pressure test had been a real barrier all along — one the system had been quietly permitted to misread. The hedges survive into the case: the process-safety / personal-safety distinction the CSB names is load-bearing and easy to lose; the OSHA-vs-EPA enforcement gap on offshore facilities was structural, not incidental; and the four investigations did not converge on a single cause precisely because the catastrophe had several, all of which had been quietly tolerated.



The immediate causes of the Macondo well blowout can be traced to a series of identifiable mistakes ... that reveal such systematic failures in risk management that they place in doubt the safety culture of the entire industry.

NATIONAL COMMISSION, DEEP WATER (REPORT TO THE PRESIDENT), 2011

REFERENCES

1. National Commission on the BP Deepwater Horizon Oil Spill, Deep Water: The Gulf Oil Disaster and the Future of Offshore Drilling (Report to the President, 2011) — systematic management failures as root cause; the misread negative-pressure test.
2. National Commission (2011) — the well-control sequence and mud-displacement decision.
3. National Commission (2011) — “systematic failures in risk management... place in doubt the safety culture of the entire industry” (quoted); U.S. Chemical Safety Board, Drilling Rig Explosion and Fire at the Macondo Well (final volumes, 2014–2016) — process-safety vs. personal-safety distinction and BP’s industry-leading lost-time-injury rate as a misleading indicator.
4. BOEMRE / U.S. Coast Guard Joint Investigation (2011) — blowout-preventer maintenance backlog and chain-of-command findings; National Academies, Macondo Well Deepwater Horizon Blowout: Lessons for Improving Offshore Drilling Safety (2012) — training gaps and the interlinked cascade of failed defenses.
5. Deepwater Horizon Study Group (UC Berkeley, 2011) final report; N. Leveson, systems-theoretic analysis of Deepwater Horizon — the multi-layer drift and the limits of single-cause framings.
6. Spill-volume estimates (4.9 million barrels) and BP cost disclosures (>\$65 billion); the reorganization of the Minerals Management Service into BSEE/BOEM (Secretarial Order 3299, 2010); A. Lustgarten, Run to Failure: BP and the Making of the Deepwater Horizon Disaster (W.W. Norton, 2012) — long-arc account of accumulated procedural debt.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Deepwater Horizon is the canonical multi-layer failure: training, procedure, supervision, contractor-coordination, and equipment all drifted independently until none caught the others. The negative-pressure test was the proximate trigger, but the system as a whole had been operating with accumulated procedural debt for years. The capability to recognize an unsafe well-state simply was not present at the moment it was needed. The CSB’s process-safety / personal-safety distinction is load-bearing: BP’s industry-leading lost-time-injury record was a measurement of the wrong construct, and reading it as safety was itself a normalization.

LENS APPROACH

Deepwater Horizon anchors the cue-and-alert-design competency (induced 3.1; LENS D4/PT5): the negative-pressure test was the cue, and the cue’s ambiguity at the moment of decision was the capability gap. LENS uses it in Domain 4 (Test and Evaluation) for the cue-design failure and the wrong-construct measurement; in Domain 1 (Systems Analysis) for multi-layer drift across procedure, training, supervision, contractor coordination, and equipment; and in Domain 5 (Navigating Sociotechnical Constraints) for the OSHA-vs-EPA enforcement gap and the leasing-vs-safety regulator conflict the BSEE/BOEM split

conceded. Pair with Texas City (Case 161) on the process-safety-vs-personal-safety inversion, and with Challenger / Columbia (Case 111) on the multi-layer-drift form.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ENERGY	→ T N K	→ D1 D2 D3 D4 D5	5	3.1	4 · 5
Test and Evaluation					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Make the negative-pressure test a hard, instrumented gate: define the pass criteria so an anomalous reading stops work rather than inviting a rationalizing explanation. ▶ Specify a single, unambiguous well-control chain of command across operator and contractors before drilling begins, so no decision falls between organizations. ▶ Build well-control training around the ambiguous small-anomaly case the crew actually faced, not only the textbook blowout.	▶ Audit each defense layer for drift independently — and audit the correlation between them, since the danger lived where tolerances aligned. ▶ Track accumulated procedural debt as a standing metric so silent degradation surfaces before tolerances coincide. ▶ Structurally separate the regulator’s leasing and safety roles so revenue pressure cannot erode well-integrity enforcement.

REFLECTION QUESTIONS

1. Identify a multi-defense process in your domain. For each layer, what is the procedural debt that has accumulated since it was designed?
2. The negative-pressure test was the proximate trigger. Design a capability check that would have caught the misinterpretation in real time — and specify what makes the cue unambiguous enough that an anomalous reading stops work rather than inviting a rationalizing explanation.
3. The Commission judged the drift industry-wide, not rig-specific. What measure in your domain would reveal whether a single failure is a local lapse or a sample from a systemic distribution?
4. BP’s lost-time-injury rate had been industry-leading on the rig. Name the measurement in your domain that is most at risk of being the wrong construct — a personal-safety analog where the catastrophic process-safety risk lives elsewhere — and design the second instrument that would surface the harm the first one cannot see.

WHO BUILDS THIS learning science & instructional design · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · incident analysis · safety dashboards · knowledge bases · data pipelines

Camp Fire / PG&E

G Governance & Trust **N** Normalization of Deviance **K** Knowledge & Institutional Memory

IMPACT 85 killed in Paradise, California; deadliest U.S. wildfire in a century; PG&E pleaded guilty to 84 counts of involuntary manslaughter

IN BRIEF

On 8 November 2018 a worn hook on a nearly century-old PG&E transmission line failed in high winds and drought, igniting the Camp Fire, which swept into the town of Paradise faster than people could evacuate. Eighty-five died — the deadliest U.S. wildfire in a century — and PG&E later pleaded guilty to 84 counts of involuntary manslaughter. Investigators found PG&E had known for years that its transmission infrastructure across high-fire-risk areas was deteriorating, and had deferred maintenance to fund other priorities, under a regulatory regime that let the deferrals continue. Infrastructure built for one climate was operating in another. Camp Fire is the book's foundational climate-era case for utility capability under changed conditions, and it restructured how California regulates utility wildfire risk.

BACKGROUND

PG&E operated aging transmission lines across the wildfire-prone foothills of Northern California — some hardware approaching a century old — in a climate growing hotter and drier than the one the grid had been built for, so the operating environment had drifted away from the assumptions the infrastructure was designed against.^o

WHAT HAPPENED

On 8 November 2018 a worn C-hook on a PG&E transmission tower failed in high winds, dropping a live line and igniting a fire under severe drought conditions — a single piece of aged hardware setting off a catastrophe the dry, windy conditions stood ready to amplify. The fire moved into the town of Paradise faster than its evacuation routes could clear it; 85 people died — the deadliest U.S. wildfire in a century. PG&E later pleaded guilty to 84 counts of involuntary manslaughter, accepting criminal responsibility at a scale rare for a utility.¹

THE INVESTIGATION

CalFire's investigation and the Butte County District Attorney's report found that PG&E had known for years about the deteriorating condition of its transmission infrastructure in high-fire-risk areas, and had deferred the maintenance to fund other corporate priorities — so the hazard was not unknown but a recognized risk repeatedly postponed.² The gap was simultaneously at the utility's asset-maintenance decisions and at the regulatory architecture that had allowed the deferrals to continue, neither side holding a line that would have forced the work.³

THE CAPABILITY GAP

Camp Fire is the climate-era case for utility-capability failure under changing risk. The infrastructure had been designed and maintained for one set of conditions and was operating in another, more dangerous one, so the safety margins the original design assumed had quietly eroded as the climate shifted beneath them. The capability to update operations — inspection cadence, vegetation management, de-energization, replacement — to match the actual risk did not exist as an institutional deliverable, on either the utility's side or the regulator's, leaving no one tasked with closing the widening gap.⁴

AFTERMATH & REFORM

PG&E entered bankruptcy under tens of billions in wildfire liability, and California restructured how it regulates utility wildfire-risk planning — mandatory mitigation plans, inspections, and public-safety power shutoffs, turning the deferred work into requirements with teeth behind them.⁵ Paired with the Northeast Blackout (Case 146), Camp Fire shows utility capability failing in a second way: not a silent control-room failure but a slow, known erosion of physical infrastructure against a rising hazard the institution declined to fund against — a failure measured in years of deferral rather than seconds of cascade.

85

KILLED IN PARADISE · SINGLE TRANSMISSION HOOK

infrastructure designed for one climate, operating in another

CAMP FIRE — CAPABILITY MISMATCH UNDER CHANGED CONDITIONS

PG&E knew its equipment was failing in high-fire-risk areas and continued operating without remediation.

PARAPHRASING THE BUTTE COUNTY DISTRICT ATTORNEY'S CAMP FIRE REPORT, 2020

REFERENCES

1. CalFire, Camp Fire Investigation Report (2019) — the worn transmission-line hardware as ignition source.
2. The Camp Fire record — 85 killed; PG&E's guilty plea to 84 counts of involuntary manslaughter (2020).
3. Butte County District Attorney, The Camp Fire Public Report (2020) — PG&E's knowledge and deferred maintenance (quoted).
4. California Public Utilities Commission, Order Instituting Investigation into PG&E (2019) — the regulatory dimension.
5. PG&E bankruptcy and California wildfire-mitigation reforms (2019–).
6. Cf. the Northeast Blackout (Case 146); climate-and-infrastructure literature.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Camp Fire is the canonical climate-era case for utility- capability failure under changing risk conditions. The infrastructure was designed for one set of conditions and operated in another. The capability to update operations to match actual conditions did not exist as an institutional deliverable.

LENS APPROACH

LENS uses the Camp Fire in LEN 7 as a legacy-asset case (induced 7.3): infrastructure designed for one operating regime and run unchanged into another as the climate shifted the fire risk underneath it. The teaching point is the assumption-migration review as a recurring engineered obligation — an asset whose safety case rests on conditions that no longer hold needs a standing, funded re-verification that the original design assumptions still match the world, not a one-time sign-off. It is the single taught example of the legacy-asset-aged-past-its-regime pattern; the case pairs with the Northeast Blackout (Case 146) as utility-capability failures of different kinds.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ENERGY	→ G N K	→ D1 D2 D3 D4 D5	1	7.3	1-5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
- Re-derive the infrastructure’s safety margins against current conditions, not the historical climate it was designed for, and treat the gap as a hazard to close. - Make updating operations to match rising risk — inspection cadence, vegetation management, replacement — an explicit institutional deliverable with an accountable owner. - Build de-energization and equipment-replacement triggers tied to known high-risk hardware in high-fire-risk areas.	- Audit deferred maintenance against the hazard it guards against, so a recognized, deteriorating risk cannot be postponed year after year. - Hold the regulator’s line with mandatory mitigation plans, inspections, and penalties, so the cost of resilience is forced before the disaster rather than after. - Monitor aging hardware in the highest-risk corridors as a standing priority, treating a worn critical component as an active threat, not a backlog item.

REFLECTION QUESTIONS

1. Identify infrastructure in your domain that was designed for one set of conditions and is now operating in another. What is the capability deliverable to bridge the gap?
2. Design the regulatory architecture that would prevent a utility from deferring critical wildfire-risk maintenance.
3. PG&E knew its hardware was deteriorating in high-fire-risk areas yet deferred the work for years, and the regulator let it. What in your domain lets a recognized, rising hazard be postponed indefinitely — and who would have to hold the line that forces the spending before the catastrophe?

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · incident analysis · safety dashboards · knowledge bases · data pipelines

AFTER THE EVENT — RESPONSE AND RECOVERY

Three Mile Island

T Training Gap **H** Human-System Interface

IMPACT Partial meltdown of a Babcock & Wilcox PWR; minimal off-site dose; most serious accident in U.S. commercial nuclear history; catalyst for industry-wide reform

IN BRIEF

On 28 March 1979 a small cooling fault at Three Mile Island's Unit 2, a Babcock & Wilcox reactor near Harrisburg, escalated into a partial core meltdown — the most serious accident in U.S. commercial nuclear history. A relief valve stuck open while a control-room light reported it closed, and operators trained for dramatic design-basis ruptures misread the slow, ambiguous cascade and cut the cooling the core needed. A nearly identical near-miss at Davis-Besse eighteen months earlier had never been propagated to the fleet. The Kemeny Commission concluded the fundamental problems were people-related, not equipment. Off-site radiation was minimal, but the accident produced a system of reform — most enduringly INPO — making it the book's paired example of a failure that engineered lasting capability.

BACKGROUND

Three Mile Island's Unit 2, a Babcock & Wilcox reactor near Harrisburg, was run by operators trained for large, fast, design-basis ruptures — not the slow, ambiguous small-break sequence that actually came.⁰ It should not have surprised anyone: in September 1977 a nearly identical stuck-open relief valve had occurred at Davis-Besse, but neither the utility, the vendor, nor the NRC grasped its significance or pushed the lesson out to the fleet — so the precise sequence that would later threaten the core had already been demonstrated and then quietly filed away rather than turned into a drill or a warning.¹ The industry's posture rested on a probabilistic argument that catastrophic accidents in light-water reactors were vanishingly unlikely; that argument had become institutional doctrine, shaping what was trained for, what was instrumented, and what near-misses were treated as.

WHAT HAPPENED

On 28 March 1979 a minor secondary-cooling upset tripped the reactor; a power-operated relief valve (PORV) opened and then stuck open, draining coolant from the primary loop. The control-room light reported the **command** to close it, not its actual position, so the panel read "closed" while the valve stayed open. Misreading the rising pressurizer level, the operators throttled back the high-pressure injection the starving core depended on — taking the one action that turned a recoverable upset into a meltdown, precisely because the instrument they trusted was telling them a state that was not true. About half the core melted; off-site radiation, as it turned out, was minimal, though the public-communication

failure that accompanied the accident — the contradictory updates from utility, state, and federal officials — would shape every nuclear emergency-response standard built after.²

THE INVESTIGATION


The Kemeny Commission (October 1979) inverted the industry's assumptions: "the fundamental problems are people-related problems and not equipment problems." The equipment was good enough that, but for the human failures, the accident would have been minor. The criticism reached past the operators to management, the utility, and the NRC — an institution-wide belief that serious accidents were effectively impossible, a complacency that had shaped training priorities, staffing, and oversight long before the relief valve ever stuck, so that the operators inherited a posture they had not chosen.³ The companion Rogovin report for the NRC documented the Davis-Besse precursor and the institutional failure to disseminate it, and the decades-long literature that followed — Walker's historical synthesis, Rees's institutional study — treated TMI as the moment U.S. nuclear stopped being able to assume away its own human factors.⁴

THE CAPABILITY GAP

The gap was not intelligence but the right capability for the event that arrived. Operators drilled on design-basis ruptures had no model for an ambiguous cascade, and a control room reporting commands rather than states made correct diagnosis nearly impossible — the interface and the training were mismatched to the failure that actually arrived. Beneath that lay a second failure: the capacity to **learn** — to turn Davis-Besse into fleet knowledge — had itself broken down, so the industry kept making the same diagnosis blind, which is why the case resists any single-cause reading.⁵ The third layer was institutional: there was no organization whose job was to force operating experience from one utility into the training and procedures of every other, and no body the NRC trusted to do it on the industry's behalf.

AFTERMATH & REFORM

TMI produced not another accident but a system of reform. Within nine months the industry created the Institute of Nuclear Power Operations (INPO, December 1979) to set standards, run plant evaluations, and force the sharing of operating experience the Davis-Besse failure had lacked.⁶ The NRC overhauled licensing and inspection: it required plant-referenced simulators, placed resident inspectors at every commercial reactor, tied operator training and re-qualification to a national standard, and in Regulatory Guide 1.97 specified the post-accident instrumentation needed to give crews state — not command — information across the regimes accidents actually produce, with subsequent control-room human-factors standards (NUREG-0700) carrying the principle into every new and modernized plant.⁷ Each measure attacked one of the failures the accident had exposed: the training gap, the broken learning channel, the gulf between what indicators displayed and what crews inferred, and the institutional absence of a body responsible for cross-fleet learning. TMI is paired later with INPO (Case 175) as the book's strongest argument that failure can engineer durable industry-level reform.



The fundamental problems are people-related problems and not equipment problems.

KEMENY COMMISSION REPORT ON THREE MILE ISLAND, 1979

REFERENCES

1. Kemeny Commission, Report of the President's Commission on the Accident at Three Mile Island (Oct. 1979) — operator training oriented to large design-basis accidents.
2. M. Rogovin & G. T. Frampton, Three Mile Island: A Report to the Commissioners and to the Public, NUREG/CR-1250 (U.S. NRC, 1980) — the September 1977 Davis-Besse stuck-PORV precursor and the failure to disseminate its lessons.
3. U.S. Nuclear Regulatory Commission, Backgrounder on the Three Mile Island Accident — accident sequence, misleading PORV indication, throttled high-pressure injection, 50% core damage, minimal off-site dose.
4. Kemeny Commission (1979) — central conclusion that the fundamental problems were people-related, not equipment-related (quoted).
5. C. Perrow, Normal Accidents (1984); J. Reason, Human Error (1990) — why ambiguous cascades defeat design-basis training and command-not-state interfaces.
6. Institute of Nuclear Power Operations, "Our History" — INPO established December 1979 to set standards and force sharing of operating experience; see also NEI, "Lessons from the 1979 Accident at Three Mile Island."
7. J. V. Rees, Hostages of Each Other: The Transformation of Nuclear Safety since Three Mile Island (Univ. of Chicago Press, 1994) — TMI as the origin of systematic capability reform.
8. U.S. Nuclear Regulatory Commission, Regulatory Guide 1.97, Criteria for Accident Monitoring Instrumentation for Nuclear Power Plants (revised post-TMI), and NUREG-0700, Human-System Interface Design Review Guidelines — the control-room and post-accident-instrumentation standards that grew out of TMI; see also J. C. Walker, Three Mile Island: A Nuclear Crisis in Historical Perspective (Univ. of California Press, 2004) on the operator-training-and-licensing overhaul.

THE LEARNING ENGINEERING LENS

LE INSIGHT

TMI is the textbook moment when an industry discovered that its capability assumptions did not survive contact with operational reality. Training to design-basis events left the operators blind to the genuinely ambiguous failures that complex systems actually produce; a control room that reported commands rather than states made correct diagnosis structurally hard; and the channel that should have carried Davis-Besse out to every plant did not exist. The reform — INPO, plant-referenced simulators, resident inspectors, RG 1.97 instrumentation, NUREG-0700 human-factors standards — built the missing infrastructure together rather than one piece at a time, which is the case's load-bearing teaching.

LENS APPROACH

TMI is the worked example of induced sub-competency 6.1 (industry-level institution building after catastrophe) and the LENS D5/PT4 pairing — Navigating Sociotechnical Constraints applied to the institutional architecture a catastrophic-system industry needs to learn at scale. Students reconstruct the capability requirements that the design-basis training framework missed (LENS D1), examine the control-room interface and post-accident instrumentation as evidence-architecture problems (LENS D3), and design the cross-fleet learning channel whose absence at Davis-Besse let TMI happen. The case pairs with INPO (Case 175) as the institution-building counterpart to the failure that produced it, and with Fitzgerald/McCain (Case 124) as the contrasting failure where the missing learning channel was internal to one service rather

than across a civilian industry. CLO mapping: CLO-5 (Sociotechnical Constraints) primary for the INPO/NRC institutional architecture; CLO-1 (Systems Analysis) for the interface-and-training requirements analysis the accident exposed.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ENERGY	→ T H	→ D1 D2 D3 D4 D5	4	3.1	3 · 1
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Design control-room indicators to report actual system state, not the command issued, so operators diagnose from truth rather than intent. ▶ Build training around ambiguous, slow-onset cascades — not only dramatic design-basis ruptures — using simulators that mirror the operators’ own plant. ▶ Engineer a learning channel that propagates precursor events across the fleet as drills, so a near-miss at one site becomes practiced knowledge everywhere.	▶ Audit whether near-misses are actually reaching operators as changed procedure and training, and force-share operating experience through an independent standards body (the INPO model). ▶ Sustain plant-referenced simulators and resident oversight so diagnostic skill on ambiguous transients does not decay back to textbook drills. ▶ Monitor the institutional assumption that serious accidents are impossible, and treat its persistence as a measurable readiness risk.

REFLECTION QUESTIONS

1.

TMI operators were trained for worst-case scenarios but failed in an ambiguous one. What is the equivalent training gap in your domain between the trained case and the messy case?
2.

The Kemeny Commission called the human element the dominant risk. What evidence would you need to demonstrate the same conclusion in your own domain?
3.

The Davis-Besse precursor occurred eighteen months earlier but never reached the operators who needed it. Design the channel in your domain that would turn a near-miss at one site into a drill at every other before the second event arrives.

WHO BUILDS THIS learning science & instructional design · human factors & interaction design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · usability testing · cognitive walkthroughs

Fukushima Daiichi

N Normalization of Deviance **G** Governance & Trust **K** Knowledge & Institutional Memory

IMPACT Three reactor meltdowns following the Tōhoku earthquake and tsunami; ~160,000 people displaced; cleanup projected at \$200B+

IN BRIEF

On 11 March 2011 the Tōhoku earthquake and the tsunami it spawned overwhelmed TEPCO's Fukushima Daiichi plant: the wave topped the seawall, flooded the emergency diesel generators, and cut cooling to the reactors. Three of the six cores melted down and hydrogen explosions spread radioactive material, displacing some 160,000 people; cleanup is projected above \$200 billion. The independent Diet commission (NAIIC), chaired by Kiyoshi Kurokawa, called it a disaster "made in Japan" — the product of regulatory capture and a culture reluctant to challenge utility assumptions; evidence of large historical tsunamis had been discussed internally but never forced a change to the seawall. Other reviews stressed under-estimated external hazards. Fukushima is the post-TMI evidence that safety institutions like INPO must be deliberately built, not assumed.

BACKGROUND

Fukushima Daiichi sat on Japan's northeast coast, its reactors and their backup diesel generators protected by a seawall sized to a design-basis tsunami. Evidence of much larger historical waves — back to the ninth-century Jōgan event — had been discussed in TEPCO's own internal assessments, but no institutional path turned that evidence into a higher seawall.⁰ The gap was not in the data but in the conversion: the larger-wave evidence lived inside the utility's own assessments, where it could be discussed indefinitely without ever becoming a binding requirement to raise the wall it implicated.

WHAT HAPPENED

On 11 March 2011 the Tōhoku earthquake struck and the tsunami that followed topped the seawall. The plant lost off-site power; the diesel generators meant to keep cooling running were inundated. Cooling failed, three of the six reactor cores melted down, and hydrogen explosions spread radioactive material across the region. Some 160,000 people were displaced, and cleanup is projected above \$200 billion.¹ Siting the backup generators where a wave that overtopped the seawall would reach them tied the entire cooling chain to that single design assumption, so once the wall was topped the loss of cooling followed almost mechanically from the layout itself.

THE INVESTIGATION

The independent investigation chaired by Kiyoshi Kurokawa for the National Diet (NAIIC) called the accident "made in Japan" — the product of regulatory capture, deference to authority, and an institutional reluctance to challenge utility assumptions.² Other major

reviews — the Hatamura government commission (2012) and the IAEA Director General's report (2015) — emphasized the under-estimation of external hazards and defense-in-depth assumptions over the cultural critique; the Kurokawa framing is the most-cited but not the only consensus reading.³ That serious independent reviews diverged on emphasis — cultural capture versus under-estimated external hazards — is itself part of the record, and the book treats Kurokawa's as the most-cited reading rather than the settled one.

THE CAPABILITY GAP

Fukushima is the post-TMI case showing that the INPO pattern (Case 175) is not self-executing. The U.S. industry built INPO to force operating discipline and shared learning; the Japanese industry did not build an equivalent with the independence to override a utility's optimistic assumptions. The internal tsunami evidence existed; the institutional capability to act on it did not.⁴ Evidence without an independent body empowered to act on it is inert: the larger-wave assessments could be acknowledged and shelved indefinitely because no institution stood outside the utility with the standing to convert them into a mandated change.

AFTERMATH & REFORM

The accident drove a restructuring of Japanese nuclear regulation — a new, more independent Nuclear Regulation Authority — and a global re-examination of external-hazard margins and station-blackout protection. Paired with INPO and Davis-Besse, it triangulates the book's claim: sustained nuclear-safety capability is an institution that must be deliberately built and rebuilt, not a property a competent industry simply has.⁵ Creating a more independent regulator after the fact conceded the structural diagnosis directly: the missing piece had been an authority sitting outside the utility, and the reform was precisely to build the independence that the pre-2011 arrangement had lacked.

3 of 6

REACTORS MELTED DOWN

tsunami exceeded design basis the institutional evidence had already questioned

FUKUSHIMA DAIICHI — "MADE IN JAPAN," PER THE DIET INQUIRY

What must be admitted — very painfully — is that this was a disaster "Made in Japan."

NATIONAL DIET OF JAPAN FUKUSHIMA NUCLEAR ACCIDENT INDEPENDENT INVESTIGATION COMMISSION, 2012

REFERENCES

1. National Diet of Japan Fukushima Nuclear Accident Independent Investigation Commission (NAIIC; K. Kurokawa, chair), Report (2012) — the internal tsunami evidence and the regulatory-capture finding.
4. Investigation Committee on the Accident (Hatamura government commission, 2012); IAEA Director General, The Fukushima Daiichi Accident (2015) — external-hazard under-estimation.
5. D. Lochbaum, E. Lyman & S. Stranahan, Fukushima: The Story of a Nuclear Disaster (2014).

2. NAIIC (2012) — the accident sequence: seawall overtopping, generator inundation, and three core meltdowns.
3. NAIIC (2012) — the “made in Japan” cultural and regulatory-capture conclusion (quoted).
6. Y. Funabashi & K. Kitazawa, Fukushima in Review (2012); cf. INPO (Case 175) and Davis-Besse (Case 160).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Fukushima is the post-TMI case that establishes that the INPO pattern (Case 175) is not self-executing. The U.S. industry built INPO; the Japanese industry did not. The cost of the difference, paid in 2011, is the strongest available evidence that capability institutions must be deliberately built, not assumed.

LENS APPROACH

LENS uses Fukushima in LEN 8 as the cross-cultural comparison to INPO and in LEN 7 for the regulator-utility dynamics study. The case is paired with INPO and Davis-Besse to triangulate what sustained nuclear-safety capability requires.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ENERGY →	N G K	→ D1 D2 D3 D4 D5	4	7.4	5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Build an institutional path that converts an acknowledged hazard (like the larger-wave evidence) into a binding design requirement, not a discussable assessment.
- ▶ Site critical backups so a single overtopped barrier cannot disable the whole cooling chain at once.
- ▶ Establish an oversight body independent of the utility, with standing to override optimistic in-house assumptions about external hazards.

IN OPERATION / AFTER

- ▶ Audit whether internally held hazard evidence has actually driven design changes, treating shelved assessments as an open finding.
- ▶ Re-examine external-hazard margins on a cadence rather than freezing them at a design-basis figure set once.
- ▶ Verify that the regulator retains the independence to act, since institutional safety capability erodes if it is assumed rather than rebuilt.

REFLECTION QUESTIONS

1. Identify a regulatory regime in your domain whose effectiveness depends on a cultural willingness to challenge authority. What if the culture changes?
2. INPO is U.S.-specific. Design the structural features that would have to be in place for a comparable institution to function in a different national context.
3. The tsunami evidence was discussed but never converted into a binding requirement. What institutional path in your domain turns an acknowledged hazard into a mandated change rather than a shelved assessment?

WHO BUILDS THIS safety science & organizational analysis · governance, ethics & measurement · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** incident analysis · safety dashboards · governance frameworks · bias & evidence audits · knowledge bases · data pipelines

West Africa Ebola – The Delayed International Response

G Governance & Trust **N** Normalization of Deviance

IMPACT The 2014–2016 West Africa Ebola epidemic reached approximately 28,600 reported cases and 11,300 reported deaths (WHO situation-report figures, undercount acknowledged) across Guinea, Liberia, and Sierra Leone; first cases December 2013, outbreak confirmed March 2014, MSF called it "out of control" in June 2014 while WHO leadership disputed that framing, and WHO declared the Public Health Emergency of International Concern only on 8 August 2014 — months into exponential growth

IN BRIEF

The 2014–2016 West Africa Ebola epidemic is the load-bearing case for the gap between declared international readiness and deployable international capability. The first cases appeared in Guinea in December 2013; the outbreak was confirmed as Ebola in March 2014. Médecins Sans Frontières declared the epidemic "out of control" in June 2014, while WHO leadership publicly disputed that framing as alarmist. WHO declared the Public Health Emergency of International Concern only on 8 August 2014 — months into exponential growth, and by the third PHEIC ever issued. The response then scaled: UNMEER, the first UN emergency health mission, and US and UK military logistics; and community-led change — safe burials and community care-centre networks — bent the curve. The reported toll reached approximately 28,600 cases and 11,300 deaths; the undercount is acknowledged and the figures are reported figures, not a census. Three converging post-mortems — the WHO Ebola Interim Assessment Panel (Stocking Report, July 2015), the Harvard-LSHTM Independent Panel (Moon et al., *The Lancet*, November 2015), and the UN High-Level Panel (2016) — found the same structural failures: WHO's outbreak-response capacity had been cut after the 2011 budget reductions; the IHR (2005) core capacities existed on paper in member states but not in practice; and early-warning signals were discounted for political and economic reasons. The reform — the WHO Health Emergencies Programme (2016) — carries the panels' own caveat that post-crisis reform pledges have historically decayed. COVID-19 later tested it.

BACKGROUND

The epidemiological timeline is the spine of the case. The first cases trace to Guinea in December 2013; the outbreak was laboratory-confirmed as Ebola virus disease in March 2014. In April, Guinea's health ministry and WHO reporting characterized the outbreak as slowing toward containment. By June, the case count in Guinea, Liberia, and Sierra Leone was climbing, and Médecins Sans Frontières stated publicly

that the epidemic was “out of control” and beyond the capacity of the responding organizations — a framing that WHO leadership disputed at the time as alarmist. WHO convened its Emergency Committee on 6–7 August 2014 and declared a Public Health Emergency of International Concern on 8 August 2014, by which point roughly 1,700 cases and 930 deaths had been reported. The declaration came months into exponential growth, not ahead of it. It was only the third PHEIC WHO had ever issued.^o

WHAT HAPPENED

The analytical spine is declared capability versus engineered capability. Under the International Health Regulations (2005), member states were required to build and certify a set of core capacities for detecting, assessing, and responding to public-health events. On the certification record, those capacities existed. In the three most-affected states, the deployable capability did not: surveillance, laboratory, and response infrastructure that had been declared present could not in fact absorb the epidemic. The IHR core-capacity regime had produced paper compliance without real capability — the readiness the system declared was not the readiness the system could produce. The Stocking Report named the IHR implementation gap directly as one of its three review areas, alongside WHO’s own emergency-response capacity and its role within the wider humanitarian system. The gap itself — declared readiness minus deployable readiness — is the artifact the case asks the reader to own.¹

THE INVESTIGATION

The escalation machinery failed as governance, not merely as forecasting. The early signals were present and were discounted: warnings of uncontrolled spread were treated as alarmist, and calls for a larger response were slowed by political and economic considerations, including the fear that a PHEIC declaration would damage trade and travel in the affected economies. This is normalization of deviance in the escalation loop — each week that the exponential signal was reframed as manageable lowered the threshold for reframing it again the next week. The Harvard-LSHTM Independent Panel (Moon et al., *The Lancet*, November 2015) located the most serious failure in the delay in raising the alarm, and its ten reform proposals targeted the governance of outbreak escalation — the machinery by which a discounted signal is supposed to become a declared emergency — rather than the science of detection. WHO’s outbreak-response capacity had itself been reduced after the 2011 budget cuts, so the institution charged with escalating had been hollowed at the point of need.²

THE CAPABILITY GAP

The response, once it scaled, showed what deployable capability looks like. The UN established UNMEER, its first dedicated emergency health mission; the United States and United Kingdom supplied military logistics, engineering, and treatment-unit construction. But the curve did not bend on logistics alone. Community-led change carried decisive weight: safe and dignified burial practices that interrupted a major transmission route, and community care-centre networks that moved isolation and care closer to where transmission was occurring. The reported toll settled at approximately 28,600 cases and 11,300 deaths across the three countries to mid-2016. The figures are reported figures and the undercount is acknowledged in the source record — the true toll is higher by an amount that the surveil-

lance system, the same system whose inadequacy is the case's subject, could not measure. That hedge is load-bearing and travels with every use of the numbers.³

AFTERMATH & REFORM

The reform response and its own caveat close the case. WHO established the Health Emergencies Programme in 2016, restructuring its outbreak-and-emergency function; the UN High-Level Panel on the Global Response to Health Crises (2016) issued twenty-seven recommendations for national, regional, and international action, including political- leadership structures above the technical agencies; and all three panels pressed for genuine — not certified — IHR core- capacity building. Each panel attached the same caveat: that reform pledges made in the wake of a health crisis have historically decayed as the crisis recedes and attention and funding move on. This is the case's binding hedge on its own resolution — the reform is real but its durability was, at the time of writing, unestablished. COVID-19 later tested the reformed system, and the question the panels raised — whether the escalation machinery had been engineered rather than re-certified — is the question the reader is asked to carry forward, not the question the case closes.

The most egregious failure was by WHO in the delay in sounding the alarm.

HARVARD-LSHTM INDEPENDENT PANEL ON THE GLOBAL RESPONSE TO EBOLA (MOON ET AL., THE LANCET, 28 NOVEMBER 2015).

REFERENCES

1. Ebola Interim Assessment Panel (2015), Report of the Ebola Interim Assessment Panel (the "Stocking Report," chair Dame Barbara Stocking), World Health Organization, 7 July 2015.
2. Moon, S., Sridhar, D., Pate, M. A., Jha, A. K., Clinton, C., et al. (2015), "Will Ebola change the game? Ten essential reforms before the next pandemic. The report of the Harvard-LSHTM Independent Panel on the Global Response to Ebola," The Lancet 386(10009):2204–2221, 28 November 2015, doi:10.1016/S0140-6736(15)00946-0.
3. High-level Panel on the Global Response to Health Crises (2016), Protecting Humanity from Future Health Crises (chair Jakaya M. Kikwete), UN doc. A/70/723, 9 February 2016.
4. World Health Organization (2016), Ebola Situation Reports — West Africa: approximately 28,600 cases and 11,300 deaths reported to 10 June 2016 (reported figures; undercount acknowledged).
5. Médecins Sans Frontières (2015), Ebola: Pushed to the Limit and Beyond — MSF account of the international outbreak response, March 2015.

THE LEARNING ENGINEERING LENS

LE INSIGHT

West Africa Ebola is the load-bearing case for declared international readiness that was certified but not deployable. The IHR (2005) core capacities existed on paper in member states but not in practice; WHO's own outbreak-response capacity had been cut after 2011; and the escalation loop discounted an exponential signal for months before the PHEIC of 8 August 2014. The reported toll carries an acknowledged undercount, and the reform carries the panels' own caveat that post-crisis pledges decay.

LENS APPROACH

West Africa Ebola is the declared-versus-engineered- capability case at the scale of the international system (induced 1.1; LENS D5/PT6; CLO-5). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the IHR paper- compliance-versus-real-capability seam and for the governance failure in the escalation machinery — the discounting of the escalating signal as normalization of deviance in the declaration loop. The three converging panels (Stocking, Moon et al., UN High-Level) are carried as the investigation evidence base; the acknowledged undercount and the reform- decay caveat are load-bearing hedges that travel with the case.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
GLOBAL HEALTH	G N	D1 D2 D3 D4 D5	6	1.1	5
OUTBREAK RESPONSE	→	→			
INTERNATIONAL GOVERNANCE		Navigating Sociotechnical Constraints			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Treat the gap between certified and deployable capability as the primary operational artifact: IHR core capacities existed on the certification record but not in the field, and the analyst's job during escalation is to measure the readiness the system can produce, not the readiness it has declared.
- ▶ Build the escalation loop to resist discounting an exponential signal for political or economic reasons; the recurring failure was that each week's reframing of the signal as manageable lowered the threshold for reframing it again, and the declaration lagged the growth it existed to trigger.
- ▶ Do not hollow the institution charged with escalation ahead of the emergency; WHO's outbreak-response capacity had been reduced after the 2011 budget cuts, so the body responsible for sounding the alarm had been degraded at the point of need — a sustainment failure in the readiness deliverable itself.

IN OPERATION / AFTER

- ▶ Preserve the acknowledged-undercount hedge with the reported figures; the 28,600 cases and 11,300 deaths are reported figures produced by the same surveillance system whose inadequacy is the case's subject, and the true toll is higher by an unmeasured amount.
- ▶ Carry the panels' own reform-decay caveat into any teaching of the resolution; the WHO Health Emergencies Programme (2016) is real, but the panels flagged that post-crisis reform pledges historically decay, and the case's honesty depends on holding the reform and its caveat together.
- ▶ Use COVID-19 as the deferred test rather than the closing chapter; the case asks whether the escalation machinery was engineered or merely re-certified, and the later pandemic is the evidence event the reader is directed toward, not a resolution the case supplies.

REFLECTION QUESTIONS

1. Identify a capability in your domain that is certified present but may not be deployable — a readiness that exists on the compliance record but has never been exercised at full load. What would it take to measure the gap between the declared readiness and the readiness the system could actually produce, and who owns that gap today?
2. The escalation loop discounted an exponential signal for months because raising the alarm carried political and economic cost. Specify the decision rule you would build so that the threshold for reframing a growing signal as "manageable" does not fall each time it is reframed. What makes an escalation trigger resistant to the incentive not to pull it?
3. The reform response carries the panels' own caveat that post-crisis pledges historically decay. Pick a reform your institution adopted after a crisis, and ask: was the underlying capability re-engineered, or was it re-certified? What evidence — short of the next crisis — would tell you which?

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

Hurricane Maria — Puerto Rico Response and Logistics Failure

G Governance & Trust **N** Normalization of Deviance

IMPACT Hurricane Maria struck Puerto Rico September 20, 2017, as the third major hurricane of the season after Harvey and Irma; FEMA's own 2017 Hurricane Season After-Action Report (July 12, 2018) found the agency entered Maria with force strength below target, its Caribbean Distribution Center inventory largely depleted by the Irma response, and commodity-tracking visibility it could not sustain to the ports; GAO-18-472 found 54% of deployed staff serving outside their qualified positions at the October 2017 peak; the grid collapse became the longest blackout in U.S. history (~11 months to full restoration); the commissioned GWU Milken Institute study estimated ~2,975 excess deaths (Sept 2017 – Feb 2018), adopted as the official toll in place of the initial 64

IN BRIEF

Hurricane Maria made landfall on Puerto Rico on September 20, 2017, as a high-end Category 4 storm, the third major hurricane to strike U.S. territory in under a month. The load-bearing evidence on the federal response is the responding institution's own self-assessment: FEMA's 2017 Hurricane Season After-Action Report, released July 12, 2018, found that the agency entered Maria with its commodity stockpiles largely drained by Harvey and Irma — the Caribbean Distribution Center in Puerto Rico had distributed more than 80 percent of its inventory of some commodities for Irma and had not been replenished — with force strength below its staffing target, and with logistics visibility so poor it could not comprehensively track commodities moving to Puerto Rico and the U.S. Virgin Islands. GAO-18-472 documented that at the October 2017 deployment peak, 54 percent of deployed staff were serving in a capacity in which they did not hold the "Qualified" title. The island's grid collapse became the longest blackout in U.S. history, with full restoration taking approximately eleven months. Excess-mortality estimates diverge by method and are carried here as estimates: the commissioned George Washington University Milken Institute study estimated approximately 2,975 excess deaths (September 2017 – February 2018), adopted as the official toll in place of the initial count of 64; the Kishore et al. household survey in the New England Journal of Medicine estimated 4,645 with a 95 percent confidence interval of 793 to 8,498 (September 20 – December 31, 2017). The capability gap the case anchors is a logistics and workforce capability designed to a mainland-hurricane envelope and never re-verified for an island scenario with a destroyed grid and ports.

BACKGROUND

Hurricane Maria struck Puerto Rico on September 20, 2017, two weeks after Hurricane Irma had passed north of the island and less than a month after Hurricane Harvey had drawn down federal response capacity on the Gulf Coast. Maria destroyed the island's electrical grid outright: the resulting outage became the longest blackout in U.S. history, with power not restored to all customers until August 2018 — approximately eleven months — and an estimated 3.4 billion lost customer-hours of electricity service, roughly three times the next-largest blackout in the U.S. record. The operational context was categorically different from the mainland-hurricane scenario the federal response apparatus was sized against: an island 1,000 miles from the mainland, reachable only by sea and air, with ports damaged, the grid destroyed, and no adjacent states from which mutual-aid convoys could drive in.^o

WHAT HAPPENED

FEMA's own 2017 Hurricane Season After-Action Report, released July 12, 2018, is the load-bearing self-assessment and the reason the case carries the program-report evidence tag. The report found that the agency entered the season with force strength below its staffing target, and that the sequential demands of Harvey and Irma had largely drained the national commodity posture before Maria arrived: when Irma struck, FEMA opened the Caribbean Distribution Center — its warehouse of disaster supplies in Puerto Rico — and within days had distributed more than 80 percent of its inventory of some commodities; Maria struck before the supplies were replenished. The report further found that FEMA could not comprehensively track commodities moving across multiple modes of transportation to Puerto Rico and the U.S. Virgin Islands, due to staffing shortages and business-process shortfalls — the agency lost visibility of its own relief supplies at the ports. A 2011 exercise after-action report for Puerto Rico had anticipated that the territory would require extensive federal support in moving commodities from the mainland and to distribution points across the island; the anticipated requirement had not been converted into verified capability by 2017.¹

THE INVESTIGATION

GAO-18-472 (September 2018) documented the workforce side of the capability gap. At the height of the 2017 deployments in October — close to 14,000 federal employees deployed across Harvey, Irma, Maria, and the California wildfires — 54 percent of FEMA staff were serving in a capacity in which they did not hold the title of "Qualified" under FEMA's own qualification system standards. The workforce was overwhelmed by the concurrent demand, and the shortage of trained personnel with program expertise complicated the response, particularly after Maria, where the staffing shortfall landed on the disaster with the greatest logistics difficulty and the least pre-existing territorial response capacity. The GAO finding and the FEMA after-action "Force Strength" finding describe the same structural condition from outside and inside the agency: a workforce sized and qualified for the routine disaster-response envelope, deployed into a season three standard deviations outside it.²

THE CAPABILITY GAP

The mortality evidence is carried as estimates, with their methods, and the divergence is not smoothed. The government of Puerto Rico's initial official toll was 64. The Kishore et al. study, published in the *New England Journal of Medicine* in May 2018, surveyed households across the island and estimated 4,645 excess deaths from September 20 through December 31,

2017, with a 95 percent confidence interval of 793 to 8,498 — a 62 percent increase in the mortality rate over the same period in 2016, and an interval whose width is itself a finding about what a survey design can resolve. The commissioned George Washington University Milken Institute study, released in August 2018, used death-registry data and demographic modeling that accounted for post-storm population displacement, and estimated approximately 2,975 excess deaths from September 2017 through February 2018. The government of Puerto Rico adopted 2,975 as the official toll, revising the count from 64. The two studies used different methods, different windows, and different uncertainty structures; they agree that the initial official count understated the mortality by more than an order of magnitude, and they do not agree on a single number.³

AFTERMATH & REFORM

The hedges the case carries are load-bearing. The excess-mortality figures are model-dependent estimates, not counts; the GWU and NEJM designs are not directly comparable and neither supports an attribution chain from any individual logistics failure to any individual death. What the record does support is the capability-envelope finding, and it is the responding institution's own: the logistics and workforce capability had been specified and verified against a mainland-hurricane scenario — adjacent states, intact ports, a grid that fails partially — and was never re-verified for the island scenario that the 2011 exercise report had already described. The after-action report itself is the intervention-shaped element inside the failure: an institution measuring its own capability gap honestly, in public, with its own name on the findings. The systems-analysis discipline the case anchors is the one FEMA's report performs after the fact and the one the deployment record shows was not performed before it — deriving the requirement from the operational reality of the deployed context rather than from the envelope the capability was built and staffed to.

FEMA entered Maria with force strength below target and its Caribbean Distribution Center largely emptied by the Irma response; it could not comprehensively track commodities to the ports; the grid collapse became the longest blackout in U.S. history; the official toll was revised from 64 to 2,975 on the strength of the commissioned GWU excess-mortality study.

EDITORS' SYNTHESIS OF THE FEMA 2017 HURRICANE SEASON AFTER-ACTION REPORT (2018), GAO-18-472 (2018), THE GWU MILKEN INSTITUTE STUDY (2018), AND KISHORE ET AL. (2018, NEJM).

REFERENCES

1. Federal Emergency Management Agency (2018), 2017 Hurricane Season FEMA After-Action Report, July 12, 2018 — the agency's own self-assessment: force strength below target, Caribbean Distribution Center depletion, commodity-tracking shortfalls.
2. U.S. Government Accountability Office (2018), 2017 Hurricanes and Wildfires: Initial Observations on the Federal Response and Key Recovery Challenges, GAO-18-472, September 2018 — 54% of deployed staff serving outside qualified positions at the October 2017 peak.
3. Milken Institute School of Public Health, George Washington University (2018), Ascertainment of the Estimated Excess Mortality from Hurricane María in Puerto Rico, August 2018 — commissioned by the Government of Puerto Rico; 2,975 excess deaths, September 2017 – February 2018; basis for the official toll revision from 64.
4. Kishore, N., Marqués, D., Mahmud, A., et al. (2018), "Mortality in Puerto Rico after Hurricane Maria," *New England Journal of Medicine* 379(2):162 – 170, doi:10.1056/NEJMs1803972 — household-survey estimate of 4,645 excess deaths (95% CI 793 – 8,498), September 20 – December 31, 2017.
5. Houser, T., & Marsters, P. (2018), "The World's Second Largest Blackout," Rhodium Group, April 12, 2018 — approximately 3.4 billion lost customer-hours; the largest blackout in U.S. history by that metric.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Hurricane Maria is the load-bearing case for a capability specified and verified against one operational envelope — the mainland hurricane, with adjacent states, intact ports, and a partially failing grid — and deployed into another. The stockpile depletion, the workforce qualification gap, and the lost logistics visibility were not separate failures; they were one requirements failure surfacing in three subsystems. The after-action report is the case's second lesson: the institution measured its own capability gap honestly, and that honesty is what makes the gap analyzable.

LENS APPROACH

Hurricane Maria is the engineered-versus-stated-readiness case at national-response scale (induced 1.1; LENS D1/PT1; CLO-1). LENS uses it in Domain 1 (Systems Analysis) for the discipline of deriving requirements from the deployed operational context — island logistics, destroyed grid and ports, sequential-storm depletion — rather than from the envelope the capability was built to, and for the after-action report as the requirements artifact produced after the fact. Pair with Case 124 (USS Fitzgerald and McCain) for the stated-versus-real-readiness thread and Case 126 (F-35 sustainment) for capability requirements as a sustainment deliverable. The excess-mortality estimates are carried as estimates with their methods; the divergence between the GWU and NEJM figures is load-bearing and is not smoothed.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
EMERGENCY MANAGEMENT DISASTER LOGISTICS PUBLIC HEALTH	G N	D1 D2 D3 D4 D5 Systems Analysis	1	1.1	1

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Specify the operational envelope from the deployed context, not the prior platform: an island response with destroyed ports and grid is a different requirement than a mainland response with adjacent-state mutual aid, and the 2011 exercise report that anticipated the difference was the requirements artifact that needed an owner.
- ▶ Treat stockpile posture and force strength as live capability measurements under sequential-event depletion; the Caribbean Distribution Center's inventory state after Irma was a known quantity before Maria's landfall, and the replenishment gap was a decision surface, not a surprise.
- ▶ Instrument logistics visibility as a first-class capability, not a reporting afterthought; an agency that cannot track its own commodities at the ports cannot re-plan distribution under the conditions where re-planning matters most.

IN OPERATION / AFTER

- ▶ Write the after-action report the way FEMA wrote this one — the institution's own name on the capability-gap findings, in public; the report's honesty is what makes the case teachable and is the standard the casebook holds after-action self-assessment to.
- ▶ Carry the excess-mortality estimates with their methods and their divergence; the GWU registry-based 2,975 and the NEJM survey-based 4,645 (95% CI 793 – 8,498) answer differently designed questions, and revising the official toll from 64 to 2,975 is the load-bearing correction, not the reconciliation of the two studies.
- ▶ Convert the envelope finding into verification events: re-verify inherited response designs against the island scenario explicitly — pre-positioned stock floors that survive a sequential season, workforce qualification targets audited against deployment reality, and port-condition logistics exercised rather than assumed.

REFLECTION QUESTIONS

1. Identify a capability in your domain that was specified and verified against one operational envelope and is routinely deployed near or beyond it. What is the equivalent of the island scenario — the context that breaks the envelope's quiet assumptions — and what verification event would surface the gap before deployment does?
2. FEMA's after-action report named its own force-strength, stockpile, and visibility gaps in public. Specify what an equivalently honest self-assessment would look like for your institution: who signs it, what it is permitted to find, and what pre-committed response follows a finding that the capability was below requirement.
3. The two excess-mortality studies estimate different quantities with different methods and do not agree on a number, and both were necessary to displace the official count of 64. When evidence streams in your domain diverge, what is the discipline for carrying both — with methods attached — rather than picking one or averaging them?

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · incident analysis · safety dashboards

CrowdStrike Falcon Outage

D Designed Out **K** Knowledge & Institutional Memory **G** Governance & Trust

IMPACT 8.5 million Windows machines crashed; airlines, hospitals, broadcasters, and banks affected simultaneously; largest IT outage on record

IN BRIEF

On 19 July 2024 CrowdStrike pushed a content update to its Falcon endpoint sensor that contained a logic error in a small configuration file. On every Windows machine running the affected sensor, the file drove the kernel driver to crash, looping the blue screen of death and requiring hands-on recovery of each device. Roughly 8.5 million machines went down at once — across hospitals, airlines, banks, broadcasters, and governments worldwide — the largest IT outage on record. CrowdStrike's post-incident review found the content update had not been put through the same automated testing or staged rollout as code: the pipeline treated "content" as a lesser category than "code," though the operating system did not. CrowdStrike is the cybersecurity-vendor analog of Knight Capital, six orders of magnitude larger.

BACKGROUND

CrowdStrike's Falcon sensor runs deep inside Windows — in the kernel — to detect threats, and it updates constantly: not only its code, but "content," the rapid-response detection configuration pushed to customers continuously. The deployment pipeline treated that content as a lighter category than code, with less testing and no staged rollout.⁰ The distinction had a logic: content shipped fast and often, precisely so the sensor could keep pace with new threats, and slowing it down with full code-grade testing seemed to defeat its purpose. That speed was the very reason the safety gate was lowered on the artifact that could still crash a kernel.

WHAT HAPPENED

On 19 July 2024 a content update contained a logic error in a small configuration file. On every Windows machine running the affected sensor, the file drove the kernel driver to read out of bounds and crash, looping the blue screen of death and requiring manual recovery of each device. Roughly 8.5 million machines failed at once — hospitals, airlines, banks, broadcasters, and governments worldwide — the largest IT outage on record.¹ Because the file went to every affected sensor simultaneously with no staged rollout, there was no first wave to catch the fault and no blast radius short of the whole install base; the requirement for hands-on recovery of each device turned a single bad push into weeks of physical labor across the world.

THE INVESTIGATION

CrowdStrike's own post-incident review found the content update had not been put through the same depth of testing and staged rollout as its code releases. The fault was not exotic: a category boundary in the deployment pipeline — content treated as safer than code — that did not match the operational reality, in which a bad content file could crash the kernel exactly as bad code could. The boundary was an organizational convenience rather than a technical truth, and the kernel, which executes whatever reaches it, recognized no such distinction at all.²

THE CAPABILITY GAP

The missing capability was the recognition that, for deployment safety, content **is** code: anything that can crash the kernel must clear the same testing and staged-rollout gates. CrowdStrike's customers had trusted the vendor's deployment safety, and that trust turned out to be load-bearing for the operation of a large slice of the global economy on a single morning. Each customer had implicitly outsourced a safety gate to the vendor's pipeline, so the one missing gate inside that pipeline was multiplied across every institution that ran the sensor, with no independent check standing between a bad push and their kernels.³

AFTERMATH & REFORM

CrowdStrike moved content updates onto staged rollouts and stronger validation, Microsoft revisited kernel-level access for security vendors, and the episode prompted scrutiny of concentration risk in endpoint security.⁴ Each response targets a different layer of the same failure: staged rollout limits the blast radius of any one push, reconsidering kernel access limits how much a vendor fault can break, and the concentration-risk scrutiny acknowledges that a single vendor had become a shared point of failure for much of the economy. It is the cybersecurity-vendor analog of Knight Capital (Case 143) — an unverified deployment to a system wired into critical operations — at a scale six orders of magnitude larger.

8.5M

MACHINES · SINGLE CONFIGURATION FILE

content treated differently from code in the deployment pipeline

CROWDSTRIKE — THE COST OF A CATEGORY BOUNDARY IN A DEPLOYMENT

Our content configuration update process did not include the same depth of testing and staged rollout as our code releases.

PARAPHRASING CROWDSTRIKE PRELIMINARY POST-INCIDENT REVIEW, JULY 2024

REFERENCES

1.

CrowdStrike, Falcon Content Update: Preliminary Post-Incident Review (July 2024) — the content-vs-code testing and staged-rollout gap (paraphrased).

2.

CrowdStrike PIR (2024) — the configuration-file fault and the kernel crash; the 8.5 million affected Windows machines figure is Microsoft's (D. Weston, 20 July 2024).

3.

Microsoft resilient-engineering analyses and Windows kernel-access review (2024).

4.

U.S. GAO post-incident analysis and Senate Homeland Security hearings (2024) — concentration risk in endpoint security.

5.

B. Beyer et al. (eds.), Site Reliability Engineering (2016) — staged rollout and canarying; cf. Knight Capital (Case 143).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The CrowdStrike outage is the live case for what happens when a deployment safety architecture treats categories of artifact differently than the operational system actually does. Content looked operationally identical to code; it was treated as different in the deployment pipeline. The mismatch became the largest IT outage in history.

LENS APPROACH

LENS uses CrowdStrike in LEN 5 as a categories-and-boundaries capability case and in LEN 2 for the vendor-customer trust architecture: customers trusted CrowdStrike's deployment safety; that trust was load-bearing for the operation of the global economy on a single day.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
TECHNOLOGY	→ D K G	→ D1 D2 D3 D4 D5	4	5.4	5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Define deployment-safety gates by operational impact, not artifact category: anything that can crash the kernel clears the same testing and staged rollout as code.
- ▶ Make staged rollout mandatory for every update, so a faulty push is caught by a first wave instead of reaching the entire install base at once.
- ▶ Design for recovery, not just prevention — assume a bad update will ship and ensure it does not require hands-on intervention at every device.

IN OPERATION / AFTER

- ▶ Audit vendor deployment-safety practices your institution depends on, rather than treating the outsourced safety gate as trustworthy by default.
- ▶ Map and reduce concentration risk so a single vendor fault cannot take down a large slice of critical operations simultaneously.
- ▶ Run post-incident reviews that interrogate category boundaries in the pipeline and feed the findings back into validation gates.

REFLECTION QUESTIONS

1.

Identify a vendor relationship in your domain whose deployment-safety practice your institution does not audit. What would the audit reveal?
2.

Design the cross-vendor staged-rollout protocol that should be standard for endpoint security software.
3.

CrowdStrike's pipeline treated content as safer than code, but the kernel did not. What category boundary in one of your systems is an organizational convenience that the operational reality ignores — and what would it cost if it broke?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

Chapter 12

Disaster Prevention & Recovery — What Works — and the Frontier

Prevention as a deliverable before the event; recovery as a discipline after it.

The disaster that did not happen had a budget, an owner, and a deadline.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

BEFORE THE EVENT — PREVENTION

CASE 174

SOFTWARE SUSTAINMENT

CRITICAL INFRASTRUCTURE

FEDERAL PROGRAM MANAGEMENT

1996 – 2000

Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed

G Governance & Trust **D** Designed Out **K** Knowledge & Institutional Memory

IMPACT The US federal government and the broader public and private sectors invested an estimated 100 billion dollars (US) over four years remediating two-digit-year date handling in legacy systems; the January 1, 2000 rollover passed with minimal disruption to critical infrastructure — the success contributed to the durable counterfactual debate about whether the threat justified the spending

IN BRIEF

The “Year 2000 problem” — Y2K — was the structural consequence of decades of legacy software representing year fields as two digits, ambiguous between 1900 and 2000 at the rollover. The affected code ran banking systems, embedded controllers in industrial infrastructure, federal benefit-payment systems, air traffic and rail systems, and the broader public and private software base. From 1996 through December 1999 the US federal government, under sustained Office of Management and Budget reporting and GAO audit, drove an inventory-and-remediation program across mission-critical federal systems, while the private sector executed a parallel multi-year effort. Estimates of the total US investment range around \$100 billion. The January 1, 2000 rollover passed with minimal disruption to critical infrastructure. The case is the canonical instance of a believed-and-treated aging-system transition in the recent regulatory record: the threat was specific, the program was inventoried at line-item granularity, the deadline was immovable, and the institutional discipline was sustained over four years. The hedge survives in the durable counterfactual debate the success itself produced — would the rollover have passed similarly with less spending? — and the published GAO record characterizes the program as a major federal management success without claiming the counterfactual is closed.

BACKGROUND

The Y2K problem was an artifact of decades of software development in which year fields had been stored as two digits to save memory and database space when memory and database space were expensive. By the mid-1990s the code base had aged into a regime where the two-digit representation was a foreseeable failure: date arithmetic across the January 1, 2000 boundary could return wildly incorrect results — interest calculations,

scheduling logic, embedded controller timestamps, federal benefit eligibility checks. The structural form is the C7 case par excellence: legacy code aging past the assumptions of its original design, with a hard, foreseeable deadline.⁹

THE INTERVENTION

The federal-program-management response began in earnest in 1996 and accelerated through 1997–1998 with the creation of the President’s Council on Year 2000 Conversion (established February 1998) and the OMB quarterly reporting cycle. Federal mission-critical systems were inventoried at line-item granularity — the GAO report series tracked the proportion of federal mission-critical systems Y2K-compliant on a quarterly cadence, agency by agency. The discipline of the program rested on three institutional features: line-item inventory at the level of the actual systems, OMB-enforced quarterly status reporting, and GAO-sustained external audit that named agencies whose remediation lagged.¹

HOW IT WORKED

The parallel private-sector effort was as large and is harder to characterize precisely. Total US investment estimates cluster around the \$100 billion figure; global estimates are larger. Major financial institutions, utilities, telecommunications providers, and industrial operators ran their own multi-year inventory-and-remediation programs. Industry-coordination bodies (banking, electric utility, telecom) operated parallel structures to the federal program’s coordination role. The deadline forced parallel execution across the public and private sectors at a scale and pace that legacy-software remediation does not usually see.²

THE EVIDENCE

January 1, 2000 passed with minimal disruption to critical infrastructure. A small number of localized incidents occurred and were absorbed; the catastrophic-scenario possibilities the program had been built to prevent — major financial system failure, power-grid cascade, air-traffic disruption, federal benefit payment failure — did not materialize. The GAO published a sustained record of the federal program through and past the rollover; the broader retrospective literature characterizes the four-year effort as one of the major federal program-management successes of the era.³

WHAT TRANSFERRED

The hedge that survives the case is the counterfactual debate the success itself produced. The rollover passed; the counterfactual — whether the same outcome would have obtained with substantially less spending — is structurally unobservable, because the program ran. The post-2000 retrospective literature includes serious arguments on both sides: that the threat was real and the investment was the reason the rollover passed quietly, and that the investment was substantially overestimated relative to the actual latent failure population. The case is teachable on the institutional discipline — line-item inventory, OMB-enforced reporting, GAO audit, sustained four-year cadence, immovable deadline — and on the structural feature that the C7 transition succeeded because it was believed and treated. The counterfactual debate is preserved as part of the case rather than smoothed away.⁴

The C7 transition succeeded because it was believed and treated. The counterfactual — whether the threat was as large as the response implies — the program structurally cannot answer.

EDITORS' SYNTHESIS OF THE GAO Y2K REPORT SERIES AND THE POST-2000 RETROSPECTIVE LITERATURE.

REFERENCES

1. Government Accountability Office, Year 2000 Computing Challenge report series (1996–2000), particularly GAO/AIMD-99-225, GAO/T-AIMD-00-30, and GAO/AIMD-00-1 — line-item federal-program-management record.
2. Office of Management and Budget, quarterly reports on federal Y2K remediation status (1997–1999) — program-self-report tier.
3. President's Council on Year 2000 Conversion, The Journey to Y2K final report (2000) — institutional retrospective of the federal coordination effort.
4. Anson, "The Y2K Bug: A Historical and Retrospective Analysis," Computer (IEEE), retrospective literature on the counterfactual debate.
5. National Research Council, Continued Review of the Tax Systems Modernization of the Internal Revenue Service — Y2K-related sustainment-engineering record.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Y2K remediation is the canonical case of a believed-and- treated aging-system transition. Line-item inventory, OMB quarterly reporting, GAO sustained audit, an immovable deadline, and a sustained four-year program-management cadence converted a foreseeable legacy-software failure into a transition the rollover passed quietly. The counterfactual debate the success produced is part of the case.

LENS APPROACH

Y2K is the legacy-software-sustainment case (induced 7.1; LENS D1/PT4) — Domain 1 for line-item inventory + immovable deadline; Domain 5 for OMB-reporting + GAO-audit. Closes C7 with Cases 114, 115, 156. Counterfactual-debate hedge preserved.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
SOFTWARE SUSTAINMENT	G D K	D1 D2 D3 D4 D5	4	7.1	1-5
CRITICAL INFRASTRUCTURE	→	→ Systems Analysis			
FEDERAL PROGRAM MANAGEMENT					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Inventory the legacy at line-item granularity — actual systems, not categories — so the remediation status can be reported and audited against a denominator the program can defend.
- ▶ Build the immovable deadline into the program's operating discipline; the Y2K deadline could not be rebaselined, and the program's discipline came from the deadline's hardness rather than from management exhortation.
- ▶ Pair internal OMB reporting (the program's own status discipline) with external GAO audit (independent verification); the combination is what produced the institutional record the retrospective rests on.

IN OPERATION / AFTER

- ▶ Preserve the counterfactual hedge: a transition that succeeds because it was treated cannot prove the threat was as large as the response implies. The case is teachable on the institutional discipline, not on the closed answer to "was it worth it."
- ▶ Carry the public/private parallel-execution lesson: the federal program and the broader private-sector effort moved together because the deadline was external to both, and the coordination mechanisms operated alongside each other rather than depending on each other.
- ▶ Treat the Y2K case as part of the v2 aging-system quartet (Cases 114, 115, 156) — the instance where the transition was a software-sustainment problem with the largest counterfactual-debate hedge.

REFLECTION QUESTIONS

1. Identify a legacy–software or aging–system transition in your domain whose deadline is foreseeable but rebaselinable. What would the program look like if the deadline were treated as immovable — line-item inventory, quarterly reporting, sustained external audit — and what is the institutional cost of treating it that way?
2. Specify the public/private parallel–execution structure you would build if the transition reached beyond a single organization. The Y2K coordination did not depend on one body controlling the others; it depended on the external deadline being equally hard for everyone.
3. The Y2K success produces a counterfactual debate. Identify a similar success in your domain whose justification rests on the absence of the failure it was built to prevent. What institutional discipline keeps the historical record honest about the counterfactual without diminishing the discipline that produced the outcome?

WHO BUILDS THIS governance, ethics & measurement · systems & human–factors engineering · knowledge management & data engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · task analysis · design prototyping · knowledge bases · data pipelines

AFTER THE EVENT — RESPONSE AND RECOVERY

INPO and the Nuclear Academy

T Training Gap **K** Knowledge & Institutional Memory **G** Governance & Trust

IMPACT No INES-level event at U.S. commercial reactors post-INPO; sustained improvement in INPO/WANO performance indicators across the industry

IN BRIEF

Three Mile Island did not produce a reactor accident at the next plant over — it produced an institution. Within months of the 1979 Kemeny Commission report, the U.S. commercial nuclear industry founded the Institute of Nuclear Power Operations on a stark premise: an accident at any single plant would threaten every operator's license, and no utility could engineer its safety capability alone. Funded by the utilities it evaluated and operating without statutory authority, INPO set training and certification standards, accredited every plant's programs through the National Academy for Nuclear Training, and ran peer evaluations in which operators from one utility scrutinized another's control rooms and records. The pre-TMI culture of complacency gave way to mandated vigilance. No U.S. commercial reactor has had a significant INES-level event since.

BACKGROUND

The 1979 partial meltdown at Three Mile Island exposed not just a plant-level failure but an industry with no shared mechanism for learning. The Kemeny Commission traced the accident in part to a pervasive "mindset" of complacency, in which each utility operated alone and no institution carried lessons from one plant to the rest. The structural problem sat above any single control room: a lesson learned at one plant had no path to the others, so the same latent failure could surface repeatedly across an industry that never compared notes.⁹

THE INTERVENTION

Within months of the Kemeny report, the utilities founded the Institute of Nuclear Power Operations. Its premise was that an accident anywhere threatened everyone's license to operate. INPO set training and certification standards for operators and supervisors, and in 1985 the National Academy for Nuclear Training began accrediting each facility's programs. The shared-exposure premise was what gave a body with no statutory power its teeth — every utility had a direct stake in every other utility's competence, because one failure could end the whole industry's license to operate.¹

HOW IT WORKED

INPO's load-bearing mechanism was honest peer review: teams of operators from one utility examined another's procedures, control rooms, and incident records, reporting candidly because every utility was, in the title of one history, a hostage of the others. Funded by the utilities it evaluated and holding no statutory authority, INPO depended on shared catastrophic exposure to make its findings stick. Peer review by working operators rather

than distant regulators also meant the examiners knew what to look for and the findings carried the weight of professional judgment, not just rule compliance.²

THE EVIDENCE

The post-TMI culture shifted from smugness to mandated vigilance, and U.S. commercial reactors have recorded no significant INES-level event since INPO's founding. Industry performance indicators tracked by INPO and its international counterpart WANO improved steadily and broadly across the fleet. The broad, steady improvement across the whole fleet — not just the strongest plants — is the signature of a learning mechanism working as designed: the laggards were pulled up by the same peer-review architecture that held the leaders to standard.³

WHAT TRANSFERRED

INPO is the strongest evidence in any domain that capability engineering can be undertaken at the level of an entire industry rather than a single organization. Its enabling conditions — shared catastrophic exposure, regulatory legitimacy, and an honest peer-review architecture — recur wherever one operator's failure can damage every operator, and it informed the founding of WANO after Chernobyl. That the model crossed national borders to WANO is itself evidence that the design is portable: the enabling conditions, not the particular American institution, are what make the mechanism work.⁴



Every utility recognized that an accident at any single plant would affect every operator's license to operate.

PARAPHRASING THE INSTITUTIONAL ANALYSIS IN REES, HOSTAGES OF EACH OTHER, 1994

REFERENCES

1. Rees, J. (1994), Hostages of Each Other: The Transformation of Nuclear Safety since Three Mile Island — INPO's design and the "hostages" premise (paraphrased).
2. Report of the President's Commission on the Accident at Three Mile Island (Kemeny Commission, 1979) — the pre-TMI culture.
3. Nuclear Energy Institute, "Lessons from the 1979 Accident at Three Mile Island"; National Academy for Nuclear Training — accreditation and peer evaluation.
4. World Nuclear Association — Three Mile Island Accident; INPO/WANO performance indicators.
5. Marsh (2019), "INPO and the Transformation of Nuclear Safety Culture."

THE LEARNING ENGINEERING LENS

LE INSIGHT

INPO is the strongest evidence in any domain that capability engineering can be undertaken at the level of an **industry**, not just an organization. The conditions that made it possible — shared catastrophic exposure, regulatory legitimacy, an honest peer-review architecture — appear wherever a single failure can damage every operator.

LENS APPROACH

LENS uses INPO in LEN 8 as the canonical example of industry-level learning: students identify the structural conditions in their own domain that would permit (or block) an INPO-equivalent and design the peer-review architecture required. LEN 1 uses the founding moment — nine months after TMI — to discuss the **speed** a credible response to catastrophe demands.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
ENERGY	→ T K G	→ D1 D2 D3 D4 D5	4	6.1	5
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Found the body on shared catastrophic exposure — make explicit that one operator's failure threatens every operator's license — so members have a direct stake in each other's competence.
- ▶ Set common training, certification, and accreditation standards across the industry rather than leaving each organization to learn alone.
- ▶ Staff peer review with working operators, not distant regulators, so examiners know what to look for and findings carry professional weight beyond rule compliance.

IN OPERATION / AFTER

- ▶ Track fleet-wide performance indicators and confirm the laggards are being pulled up, not just the leaders held to standard — the signature of a learning mechanism working.
- ▶ Sustain candid peer review by keeping the body funded by and accountable to its members, so the honest examination that makes it effective does not erode into formality.
- ▶ Export the enabling conditions rather than the institution when scaling (as INPO informed WANO), adapting the shared-exposure-plus-peer-review design to each new context.

REFLECTION QUESTIONS

1. What is the equivalent of “an accident at any single plant affects every operator” in your domain? If the answer is “nothing,” what does that tell you?
2. INPO was stood up in nine months. Pick a current cross-organizational capability problem and write the nine-month deliverable that would constitute a credible response.
3. INPO held no statutory authority yet made its findings stick through shared catastrophic exposure and peer review. Design the non-statutory mechanism that could enforce a standard in your domain, and name the shared stake that would give it teeth.

WHO BUILDS THIS learning science & instructional design · knowledge management & data engineering · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · knowledge bases · data pipelines · governance frameworks · bias & evidence audits

CIRAS — Confidential Incident Reporting for UK Rail

G Governance & Trust

K Knowledge & Institutional Memory

N Normalization of Deviance

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT Between 2008 and 2012 the UK rail Confidential Incident Reporting and Analysis System received 2,228 reports — 45% led to tangible safety improvements and about 33% contained important safety information (program self-report); directly influenced a confidential reporting system in the US

IN BRIEF

CIRAS began as a 1995–1997 ScotRail pilot — a structured channel for rail workers to report hazards and near-misses confidentially, insulated from the employer’s disciplinary process. After the Ladbroke Grove crash in 1999, the program was mandated across UK mainline rail in 2000; in 2008 it became an independent unit within the Rail Safety and Standards Board. The published record states that between 2008 and 2012 CIRAS received 2,228 reports, of which the operating program reports 45% led to tangible safety improvements and about 33% contained important safety information. The program directly influenced the design of a confidential reporting system in the United States. The interview-based method surfaces motive and intent — the **why** behind a near-miss — that company databases miss precisely because those databases are tied to discipline. The honest hedge that survives into the case: the 45%-led-to-improvement figure is self-reported by the operating program, not independently audited. The case is the non-aviation companion to ASRS / CRM in v1 (Cases 117 + 119), keeping the non-punitive-reporting-with-credible-commitment competency from resting entirely on aviation evidence.

BACKGROUND

Confidential incident reporting as a safety-culture intervention rests on a specific structural argument: the richest information about a near-miss — the operator’s intent, the team’s local pressures, the workaround that almost worked — is exactly the information that an employer-controlled reporting channel cannot collect, because reporters with anything to lose will not put it there. The aviation experience (ASRS, CRM; v1 Cases 117 and 119) establishes the pattern at the canonical safety-culture scale. CIRAS is the same pattern, in a different industry, with the credible commitment supplied differently — and the corpus needs the non-aviation evidence.^o

THE INTERVENTION

CIRAS began in 1995–1997 as a ScotRail pilot. The structural design was the interview, not the form: trained interviewers took reports from rail workers, anonymized them, and surfaced patterns to the operating companies. The credible commitment that made the channel safe to use was not just confidentiality — it was institutional independence from the employer's disciplinary process. After the Ladbroke Grove crash in October 1999 (31 deaths, the accident that defined the decade's UK rail-safety reform), CIRAS was mandated across UK mainline rail in 2000. In 2008 it became an independent unit within the Rail Safety and Standards Board (RSSB), with the independence from employer discipline written into its operating structure.¹

HOW IT WORKED

The published record on outcomes is what the case rests on. Davies et al. (Cognition, Technology & Work) describe the method and its yield. The operating program reports that between 2008 and 2012, CIRAS received 2,228 reports, of which 45% led to tangible safety improvements and approximately 33% contained important safety information that fed back into operating practice. The program's design has directly influenced the construction of a confidential reporting system in the United States. The interview-based method surfaces motive and intent in a way that incident-database schemas tied to discipline cannot.²

THE EVIDENCE

The honest hedge has to survive. The 45%-led-to-improvement figure is self-reported by the operating program — CIRAS itself characterizes its outputs, and there is not yet an independent audit of that classification. The peer-reviewed literature on CIRAS describes the method and the institutional design, and reports the program-supplied figures rather than independently validating them. The case is teachable on the evidence of the institutional design and the published method; the operating outcome statistics are the strongest current claim, and the editor and downstream readers should treat them as program self-report rather than audited finding.³

WHAT TRANSFERRED

What CIRAS adds to the corpus is non-aviation depth in the non-punitive-reporting-with-credible-commitment competency. The mechanism (a structured confidential channel) paired with the credible commitment (independence from the employer's disciplinary process, written into operating structure) is the cultural half of capability that the aviation cases also teach. Drafted together, ASRS (Case 119), CRM/CAST (Case 117), and CIRAS (this case) show that the structural form is transferable across high-consequence operational industries — and that the credible commitment, not the channel alone, is what makes the reporting safe to use.

The 45%-led-to-improvement figure is the operating program's self-report. The institutional design is the audited finding.

EDITORS' SYNTHESIS OF DAVIES ET AL. AND THE STRATHCLYDE IMPACT CASE STUDY.

REFERENCES

1. Davies, Wright, Courtney, & Reid, "Confidential Incident Reporting on the UK Railways: The CIRAS System," Cognition, Technology & Work, doi:10.1007/PLOO011494.

2. Rail Safety and Standards Board (RSSB), CIRAS program documentation 2008 – present — operating-program publications.

3. University of Strathclyde, CIRAS impact case study — the operating-program-self-report on outcomes between 2008 and 2012.

4. Ladbroke Grove Rail Inquiry (Cullen, 2001), final report — the regulatory forcing event for national mandate.

THE LEARNING ENGINEERING LENS

LE INSIGHT

CIRAS is the non-aviation companion to ASRS and CRM in the corpus. The pattern — confidential channel paired with credible commitment (institutional independence from discipline) — works in rail as it does in aviation. The operating outcome statistics are program self-report and deserve their tier acknowledgement; the institutional design is the audited finding.

LENS APPROACH

CIRAS is the non-aviation pairing-mechanism case (induced 4.2; LENS D5/PT2). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the credible-commitment design — independence written into operating structure — and in Domain 4 (Test and Evaluation) for the CLO **Judgment under inadequate evidence** on the operating self-report tier. Pair with ASRS (Case 119) and CRM/CAST (Case 117) as the cross-industry C4 set, and with the WHO Surgical Checklist (Case 23) at the mandatory-mechanism layer.

DOMAIN		MODES ADDRESSED		LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
RAIL TRANSPORT	→	G K N	→	D1 D2 D3 D4 D5	2	4.2	5
SAFETY REPORTING				Navigating Sociotechnical Constraints			
SAFETY CULTURE							

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the credible commitment first: institutional independence from the employer's disciplinary process, written into operating structure, not relying on goodwill that can be withdrawn.
- ▶ Choose the reporting method (interview, not just form) for the information class you want to surface — motive and intent — that form-based incident databases tied to discipline cannot collect.
- ▶ Anchor the program's authority in a sector-wide forcing event (Ladbroke Grove for rail; in other domains, the analog) so the mandatory-adoption decision rides on the safety crisis the program is built to address.

IN OPERATION / AFTER

- ▶ Report operating outcome statistics with their evidence tier: program self-report vs. independently audited. Treating the 45% figure as audited overstates the evidence; treating it as untested under-states the program's institutional credibility.
- ▶ Commission an independent audit of the program's classification of its own reports at a sustainable cadence — five years is a reasonable interval — so the operating record builds toward audited evidence over time.
- ▶ Carry the structural lesson across domains: ASRS, CRM, CIRAS — the cultural half of capability is the credible commitment, not the channel alone.

REFLECTION QUESTIONS

1. Identify a high-consequence operational domain in your context where confidential incident reporting is absent or weak. What is the credible-commitment design that would make the channel safe to use — and is the institutional independence written into operating structure, or relying on goodwill?
2. Specify the information class your reporting program is built to surface: form-based incident counts, or interview-based motive and intent. The choice of method follows the information class, and they answer different questions.
3. The 45%-led-to-improvement figure is program self-report. Design the independent-audit cadence that would convert the operating record into audited evidence over time, without compromising the confidentiality the channel depends on.

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

The UN Cluster Approach – Engineering Humanitarian Coordination After the Tsunami

G Governance & Trust

IMPACT IASC cluster approach adopted after the 2005 Humanitarian Response Review diagnosed unowned sectors in the Indian Ocean tsunami and Darfur responses; first applied in the October 2005 Pakistan earthquake response; eleven global clusters, each with a designated lead agency accountable for coverage and carrying provider-of-last-resort responsibility; two formal system-wide evaluations (Stoddard et al. 2007; Steets et al. 2010) found improved gap identification and coverage at real transaction cost, with national and local actors often marginalized; revised under the 2011 IASC Transformative Agenda after the Haiti earthquake exposed scale limits

IN BRIEF

The response to the December 2004 Indian Ocean tsunami — hundreds of NGOs converging on the same coastlines, duplicated relief in accessible areas, unassigned gaps in hard ones — and the parallel critiques of the Darfur response exposed a structural fact about the international humanitarian system: no one was accountable for whole sectors of a response. The 2005 Humanitarian Response Review, commissioned by UN Emergency Relief Coordinator Jan Egeland, named the diagnosis, and the Inter-Agency Standing Committee's answer, rolled out from late 2005, was the cluster approach: each sector — logistics, shelter, health, water and sanitation, protection, camp management, and the rest — assigned a designated lead agency accountable for coverage and carrying provider-of-last-resort responsibility. Coordination was redesigned as a role architecture with named ownership rather than as goodwill. First applied in the October 2005 Pakistan earthquake response, the approach was evaluated twice at system scale: Phase 1 (Stoddard et al., 2007) found early gains in gap identification at real transaction cost; Phase 2 (Steets et al., GPPI/Groupe URD, 2010) found mostly positive coverage and gap effects with persistent weaknesses in local-actor inclusion and inter-cluster coordination. The 2010 Haiti earthquake exposed the mechanism's scale limits, and the IASC's 2011 Transformative Agenda revised it. The evaluations' own hedges are load-bearing: benefits were uneven across clusters and countries, the costs were real, and national and local actors were often marginalized.

BACKGROUND

The 2004 Indian Ocean tsunami produced the largest humanitarian mobilization to that date and, with it, the clearest exhibit of the system's chronic coordination failure. Hundreds of NGOs, UN agencies, and bilateral responders arrived with independent mandates, independent funding, and no mechanism that made any of them answerable for a sector as a whole. Relief duplicated where access was easy and cameras were present; gaps went unassigned where it was not; and when a sector was underserved, there was no one whose job it was to notice. The Darfur response drew the same critique in the same period, in a conflict setting: sectoral coverage depended on which agencies happened to show up, and the persistent gaps — protection above all — belonged to no one. The failure was not of effort or resources but of architecture: the system had participants, not owners.⁰

THE INTERVENTION

The 2005 Humanitarian Response Review, commissioned by Emergency Relief Coordinator Jan Egeland from an independent consultant team, made the structural diagnosis explicit: capacity was unevenly distributed across sectors, and accountability for sectoral outcomes was absent by design, because no design existed. The Inter-Agency Standing Committee's response was the cluster approach, agreed in September 2005 as one pillar of a broader humanitarian reform. Each sector of response — logistics, emergency shelter, health, water/sanitation/hygiene, protection, camp coordination and camp management, nutrition, education, and others, eventually eleven global clusters — was assigned a designated lead agency, globally and in each activated country response. The lead is accountable to the Humanitarian Coordinator for the adequacy of coverage in its sector and carries the provider-of-last-resort commitment: where critical gaps persist and no other actor steps in, the lead must do its utmost to fill them. The redesign is the case's analytic core — coordination converted from a voluntary good into a role architecture with named ownership.¹

HOW IT WORKED

The mechanism was paired with an authorization, and the pairing is what distinguishes the cluster approach from the coordination exhortations that preceded it. The lead agency's sectoral accountability was mandated by the IASC and the Emergency Relief Coordinator — the closest thing the voluntary humanitarian system has to an institutional spine — rather than left to inter-agency goodwill, and it was tested almost immediately: the October 2005 Pakistan earthquake, two months after the approach was agreed, became its first field application, with the initial set of clusters stood up mid-response. The Pakistan experience showed both faces at once — sectoral gap analysis that had never existed before, and meeting burdens heavy enough that NGO observers questioned whether the coordination was worth its transaction cost.²

THE EVIDENCE

The formal evaluations are the reason the case can be told with discipline. The Phase 1 evaluation (Stoddard, Harmer, Haver, Salomons, and Wheeler, 2007) found early gains — better identification of gaps, fewer duplications, clearer sectoral leadership — purchased at real transaction cost in meetings, reporting, and staff time, with benefits uneven across clusters and countries. The Phase 2 evaluation (Steets, Grünewald, Binder, and colleagues, GPPI/Groupe URD, 2010), commissioned by the IASC across six country studies, found mostly positive effects on coverage and gap identification and endorsed the approach as

the right one, while documenting persistent weaknesses that the evaluators declined to smooth: national and local actors were often marginalized or excluded by an architecture built around international agencies, inter-cluster coordination remained weak, and cluster performance varied widely by sector and country. Those hedges are carried here as findings, not caveats — the outcome evidence is genuinely mixed at the margins even where the architectural gain is real.³

WHAT TRANSFERRED

The 2010 Haiti earthquake stress-tested the mechanism at a scale it had not been engineered for — hundreds of organizations in a single cluster's meetings, a coordination apparatus that became process-driven enough to be accused of impeding the delivery it existed to enable — and the IASC Principals' December 2011 Transformative Agenda revised the approach in response, restating the clusters' original purpose around strategic gap analysis and adding system-wide activation protocols for the largest emergencies. That arc is why the case sits on the works side of this book: an institution diagnosed its own coordination failure after catastrophe, redesigned accountability as a property of the multi-organization system rather than of any single agency, evaluated the redesign formally and in public, and revised it when the next catastrophe found its limits. The cluster approach did not solve humanitarian coordination; it made sectoral coverage someone's named job, made the residual gaps visible, and left an evaluation record honest enough to build on.

Coordination redesigned as a role architecture with named ownership rather than as goodwill — and evaluated, in public, at system scale.

EDITORS' SYNTHESIS OF THE HUMANITARIAN RESPONSE REVIEW (2005) AND THE TWO CLUSTER EVALUATIONS (2007, 2010).

REFERENCES

1. Adinolfi, Bassiouni, Lauritzen, & Williams (2005), Humanitarian Response Review, commissioned by the UN Emergency Relief Coordinator, New York/Geneva: OCHA.
2. Stoddard, Harmer, Haver, Salomons, & Wheeler (2007), Cluster Approach Evaluation, Final Report, OCHA Evaluation and Studies Section, with NYU Center on International Cooperation, ODI Humanitarian Policy Group, and the Praxis Group.
3. Steets, Grünewald, Binder, de Geoffroy, Kauffmann, Krüger, Meier, & Sokpoh (2010), Cluster Approach Evaluation 2: Synthesis Report, GPPI / Groupe URD, commissioned by the IASC, Geneva: OCHA.
4. Inter-Agency Standing Committee (2011 – 2012), IASC Transformative Agenda — Principals' decisions of December 2011 and associated protocols.
5. Grünewald & Binder (2010), Inter-agency Real-Time Evaluation in Haiti: 3 Months After the Earthquake, Groupe URD / GPPI.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The cluster approach is the paired-intervention pattern at global scale: a mechanism — sector lead agencies with coverage accountability and provider-of-last-resort responsibility — paired with an authorization, the ERC/IASC mandate. Accountability became a designed property of a multi-organization system rather than an aspiration. The evaluation record keeps the case honest: gains in gap identification and coverage were real, and so were the transaction costs, the uneven benefits across clusters and countries, and the marginalization of national and local actors. The 2011 revision after Haiti is part of the intervention, not evidence against it.

LENS APPROACH

The cluster approach is the institution-building-after- catastrophe case at the largest organizational scale in the corpus (induced 6.1; LENS D5/PT6; with 4.4 alternate for the mechanism-plus-authorization pairing). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) as the anchor example of governance architecture for a multi-organization system — named sectoral ownership, provider of last resort, and a mandate that binds a voluntary system; in Domain 1 (Systems Analysis) for the Humanitarian Response Review’s structural diagnosis of unowned sectors; and in Domain 4 (Test and Evaluation) for the two formal system-wide evaluations whose hedged findings — uneven benefits, real costs, marginalized local actors — drove the 2011 Transformative Agenda revision. Pair with the industry-level institution-building cases in Part VI: the discipline is the same, the scale is the test.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
HUMANITARIAN RELIEF INTERNATIONAL COORDINATION DISASTER RESPONSE	G	D1 D2 D3 D4 D5	6	6.1	5
→ Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Assign every sector of a multi-organization response a named owner with explicit coverage accountability; a gap with no owner is a design decision, not an accident.
- ▶ Pair the ownership mechanism with an authorization that can bind — the lead-agency role works because the ERC/IASC mandate stands behind it, not because agencies agree to coordinate in principle.
- ▶ Make provider-of-last-resort responsibility explicit at designation time, so the question “who acts when no one volunteers?” has an answer before the gap appears.

IN OPERATION / AFTER

- ▶ Commission formal, published, system-wide evaluations of the coordination architecture itself — not only of individual operations — and let the findings, including the unflattering ones, drive revision.
- ▶ Treat local and national actors’ marginalization as an architectural defect to be engineered out, not a cultural footnote; an inclusion weakness documented in 2010 that persists is a standing design debt.
- ▶ Budget the transaction cost of coordination honestly and re-verify the mechanism at each new scale; Haiti showed that a role architecture sized for one magnitude of response can fail at the next.

REFLECTION QUESTIONS

1. Map a multi-organization response system in your domain — a hospital network’s surge plan, a district’s crisis-response stack, a consortium research program. For each sector of the work, who is the named owner accountable for coverage, and who is the provider of last resort? If the answer is “everyone” or “no one,” you have found the tsunami-response architecture.
2. The cluster evaluations found real gains purchased at real transaction cost, with benefits uneven across sectors and countries. Design the evaluation you would commission for a coordination mechanism in your domain: what would count as a gap-identification gain, how would you price the meeting and reporting burden against it, and who is independent enough to say the mechanism is not working?
3. The Phase 2 evaluation documented national and local actors marginalized by an architecture built around international agencies — a weakness named in 2010 and still central to humanitarian-reform debates a decade later. Identify the actors your coordination design treats as participants rather than owners. What would it take, structurally rather than rhetorically, to move them into the ownership architecture?

WHO BUILDS THIS governance, ethics & measurement — plus domain experts and a learning engineer to integrate.
TOOLS governance frameworks · bias & evidence audits

Chapter 13

Algorithms, Governance & Public Systems — What Fails

When delegation to systems outruns accountability for them.

No one could say who had decided, because no one had.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

Healthcare.gov Launch

K Knowledge & Institutional Memory **T** Training Gap **G** Governance & Trust

IMPACT ~27,000 federal–marketplace enrollments in the first month vs. a 7M first–year target; hundreds of millions in remediation

IN BRIEF

When Healthcare.gov launched on 1 October 2013 it collapsed under load it had never been tested for — about 27,000 enrollments through the federal marketplace in its first month against a seven-million first-year target. The people assembled to build it knew insurance markets and government programs but not technology product launches; key technical roles went unfilled, no office clearly owned the integration (CMS acted as its own systems integrator but lacked the capacity for the role, while CGI reported that no party was performing true integration), and no end-to-end test was run before launch. The fix-it team that rescued the site became the U.S. Digital Service. At root it was a capability mismatch at scale: the organization lacked the human capabilities the system required, and the governance chain surfaced no signal of the gap before launch. It is the rare governance failure that produced lasting institutional reform.

BACKGROUND

Healthcare.gov was the federal insurance marketplace at the center of the Affordable Care Act — a high-visibility online system that millions would hit on day one, with a political deadline that could not slip. The people assembled to build it understood insurance markets and large government programs, but not the launch of a consumer technology product at that scale, and key technical positions went unfilled — so the very expertise the launch most needed was the expertise the team most lacked.⁰

WHAT HAPPENED

There was no clear division of responsibility among the many government offices involved; CMS believed it was itself serving as the systems integrator — a role it lacked the capacity to perform — while CGI reported that no party was performing true systems integration, so the single most important role on the program was one no party actually held. No end-to-end test was run before launch, meaning the assembled pieces were never exercised together as a user would exercise them. The site went live on 1 October 2013 and immediately collapsed under load it had never been validated for: about 27,000 enrollments through the federal marketplace in its first month against a seven-million first-year target.¹

THE INVESTIGATION

Reviews by the GAO and the HHS Inspector General found that no single person had a clear understanding of the project's status, and that the governance chain had no signal that would surface the readiness gap before launch. With ownership diffused across offices and

no integration test to fail, the chain could report progress at every level while the whole remained untested — a structure in which bad news had nowhere to enter and no one positioned to act on it.²

THE CAPABILITY GAP

Healthcare.gov is a capability failure wearing a technology costume. The site was salvageable in weeks once the right people arrived, which is the clearest proof that the code was never the binding constraint; the original failure was that the wrong people had been assembled, and that the governance chain meant to catch the mismatch had no mechanism to see it. The missing capability was the matching of human capability to system requirement — and the institutional signal that would have flagged its absence before a deadline locked the launch in place.³

AFTERMATH & REFORM

The rescue effort pulled together the team that became the U.S. Digital Service — a permanent institution born from a failure visible on the news every night, its mandate built directly from what the launch had lacked.⁴ It is the rare case in this book of a governance failure that produced durable organizational reform, and a reminder that the technical narrative (“the website crashed”) often hides the real one: the wrong capability was assembled, unnoticed, by a chain that had no way to notice.



No single person had a clear understanding of the project's status.

PARAPHRASING THE HHS OFFICE OF INSPECTOR GENERAL REVIEW OF HEALTHCARE.GOV, 2016

REFERENCES

1. U.S. GAO, Healthcare.gov reports (2014–2016) — the launch, the capability gaps, and the absent end-to-end testing.
2. HHS Office of Inspector General, Case Study of CMS Management of the Federal Marketplace, OEI-06-14-00350 (2016) — unclear ownership and the CMS/CGI integration confusion.
3. HHS OIG (2016) — “no single person had a clear understanding of the project's status” (paraphrased).
4. J. Pahlka, Recoding America (2023) — the founding of the U.S. Digital Service out of the rescue.
5. Eaves & Goldenfein, “The Healthcare.gov Failure” (Harvard, 2014); Mergel et al. (2018), digital-government literature.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Healthcare.gov is a capability case wearing a technology costume. The technology was salvageable in weeks once the right people arrived. The original failure was that the wrong people had been assembled in the first place, and the governance chain that should have noticed the mismatch had no signal to surface it. The aftermath — USDS — is the rare case in this book of a governance failure that produced permanent institutional reform.

LENS APPROACH

LENS uses Healthcare.gov in LEN 1 as a problem-framing case (the technical narrative obscures the capability narrative), in LEN 5 to teach capability-requirements analysis for a large government program, and in LEN 8 to discuss the founding of USDS as organizational learning that survived.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
GOVERNMENT	→ [K][T][G]	→ [D1][D2][D3][D4][D5]	1	1.1	1
Systems Analysis					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
- Match assembled human capability to the system's real requirements at staffing time — fill the technical roles a consumer-scale launch demands before committing to the deadline. - Name a single accountable integration owner and confirm every party shares that understanding, so no critical role is one each office assumes someone else holds. - Gate launch on a full end-to-end test that exercises the assembled system the way users will, not on per-component sign-offs.	- Build a governance signal that can surface a readiness gap upward, with someone holding the authority to delay a politically fixed launch date. - Independently audit project status against demonstrated capability, since a diffuse chain can report green at every level while the whole remains untested. - Institutionalize the rescue capability — as USDS did — so the talent that fixes a launch outlives the crisis that summoned it.

REFLECTION QUESTIONS

- Healthcare.gov shipped without end-to-end testing. What is the equivalent missing deliverable in a current high-stakes deployment in your domain?
- USDS was born from the Healthcare.gov failure. What is the institutional capability your domain still lacks, and what failure would have to occur to produce it?
- No single office believed it owned integration, so the gap stayed invisible until launch day. In a project you know, who actually owns the seam between components — and what signal would let them, rather than the news, learn the system is not ready?

WHO BUILDS THIS knowledge management & data engineering · learning science & instructional design · governance, ethics & measurement — plus domain experts and a learning engineer to integrate. TOOLS knowledge bases · data pipelines · scenario simulation · competency models · governance frameworks · bias & evidence audits

EU Human Brain Project — Top-Down Vision That Unraveled

G Governance & Trust **N** Normalization of Deviance

EVIDENCE TIER journalism-tier — source confidence flagged; future validation ongoing.

IMPACT A €1B EU-flagship initiative to simulate a living brain governed top-down under a single PI; the project unraveled as the field disputed both feasibility and approach, restructured under protest, and concluded in 2023 with the founding framing abandoned

IN BRIEF

The EU Human Brain Project, launched in 2013 as one of the EU's Future and Emerging Technologies Flagship programs, set out to build a computer simulation of a living brain under a single principal investigator (Henry Markram). The governance model was top-down: a central vision, a small leadership group, a decade-long funding commitment of about €1 billion, and a research community asked to align around the simulation goal. The case is drafted as the paired contrast to the BRAIN Initiative (Case 198), not as a standalone study: the same era, the same field-scale ambition, opposite governance models, opposite trajectories. The evidence base is largely journalism — MIT Technology Review retrospective, In Silico documentary, science press, the project's own restructuring records — so the evidence-tier flag is rendered under the case title; future validation will continue. The teaching point survives the flag: the governance choice (top-down single-PI versus distributed working-group) was the variable that materially changed the trajectory of the two programs at field scale.

BACKGROUND

In 2013 the European Commission selected the Human Brain Project as one of two Future and Emerging Technologies Flagship programs, with a ten-year horizon and approximately €1 billion in committed funding. The founding vision, articulated and championed by Henry Markram, was a computer simulation of a living human brain. The governance was top-down: a small leadership group, a unified research framing, and a community of European neuroscience labs asked to organize around the simulation goal.^o

WHAT HAPPENED

The contestation surfaced quickly. Open letters from a sizable fraction of the European neuroscience community questioned both the feasibility of the simulation goal and the project's governance — the breadth of expertise on the leadership group, the relationship between the cognitive neuroscience the field practiced and the molecular-scale simulation

the project proposed, and the absence of a process for scope revision. The EU commissioned a mediation and restructuring exercise that broadened the leadership and re-scoped the research agenda around infrastructure platforms rather than a single simulation. The project ran to its scheduled 2023 conclusion; the founding framing was not what it delivered.¹

THE INVESTIGATION

The case pairs directly with the BRAIN Initiative (Case 198). The two programs launched within months of each other, at comparable field-scale ambition, in the same decade of neuroscience. The governance models were not comparable: BRAIN organized a working group that visibly contested its own composition and re-baselined its scope on the public record; HBP organized around a single PI and a unified framing, and the community had to mediate the project away from that framing. The trajectory difference is not a verdict on either field; it is evidence about the governance variable at the launch moment.²

THE CAPABILITY GAP

The evidence-tier flag rendered under the case title is load-bearing here. The detailed account of the early contestation, the open letters, the mediation, and the restructuring lives largely in MIT Technology Review's retrospective, in the In Silico documentary, and in science-press coverage from Science, Nature, and the European press; the project's own restructuring records add the formal layer. The case is teachable from this material — the structural pattern is consistent across sources — but the specific magnitudes and the exact language of the early contestation should be re-confirmed against primary sources before publication. Future validation will continue.³

AFTERMATH & REFORM

What the pair (Cases 198 + 181) teaches is that the governance model — distributed working-group with public-record contestation versus top-down single-PI with unified framing — is itself the variable that explains why one field-scale program survives and adapts while the other unravels and is mediated away from its founding goal. The capability deliverable at program launch is not the vision and not the tools; it is a governance architecture that can absorb contestation, document scope drift, and re-baseline without losing the program. BRAIN had one; HBP had to be retrofitted one.⁴

A program that cannot re-scope without external mediation is a program designed to fail open.

EDITORS' SYNTHESIS OF THE HUMAN BRAIN PROJECT RECORD.

REFERENCES

1. MIT Technology Review (2021), retrospective on big-science brain projects — the principal critical assessment of HBP alongside BRAIN.
2. In Silico (2020), documentary by Noah Hutton — long-form follow of Markram and the HBP through the contestation period.
3. Science and Nature contemporaneous coverage of the open letters and the mediation/restructuring (2014–2016).
4. Human Brain Project final reports (2023) — the project's own restructuring and concluding documentation.
5. Alivisatos et al. (2012), Neuron — the BRAIN position paper that is the paired contrast (Case 198).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The EU Human Brain Project is the paired contrast to the BRAIN Initiative (Case 198): same era, same field-scale ambition, opposite governance models, opposite trajectories. Top-down single-PI governance with a unified framing did not survive community contestation; the program was mediated away from its founding goal and ran to its scheduled conclusion in 2023. Evidence base is journalism-tier and the flag is rendered under the title; future validation will continue.

LENS APPROACH

HBP is the field-scale failure case in the v2 corpus (induced 5.1; LENS D1+D5/PT4) drafted as the contrast to BRAIN (Case 198). LENS uses the pair in Domain 1 (Systems Analysis) for the governance-variable comparison and in Domain 5 (Navigating Sociotechnical Constraints) for the contestation-absorption capability. The evidence-tier flag is binding until primary-source confirmation completes; the structural pattern survives the flag, the specific magnitudes do not.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
NEUROSCIENCE	→ G N	→ D1 D2 D3 D4 D5	4	5.1	1-5
SCIENCE POLICY (EU)		Systems Analysis + D5			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the scope-revision process before the program launches; a program that cannot re-scope without external mediation is a program designed to fail open.
- ▶ Stress-test the leadership group’s breadth before launch: does it span the disciplines the program will have to coordinate, or is it narrow enough to be captured by the founding framing?
- ▶ Treat early community contestation as governance information, not as friction; absorbing it programmatically is the capability the program needs to demonstrate.

IN OPERATION / AFTER

- ▶ When journalism-tier sourcing is what the record provides, flag it under the title and carry the standing “future validation ongoing” language into the printed case — the teaching value survives the flag, but the magnitudes do not.
- ▶ Commission a structured retrospective on the governance variable specifically: who held authority, how scope drift was named, how the program re-baselined. This is the documentation a future field-scale launch needs more than another success narrative.
- ▶ Cross-reference governance evidence across paired programs: a single program’s trajectory is not evidence about governance; a paired comparison at the same era and ambition is.

REFLECTION QUESTIONS

1. Identify a large program in your domain that ran a top-down single-PI governance model. What was the scope-revision process designed in at launch, and what was the process the program actually used when contestation arrived?
2. Specify the paired-program comparison you would commission to learn from a field-scale program’s trajectory, rather than from the program alone. What is the matched comparison case and what is the governance variable you are isolating?
3. The evidence base for this case is largely journalism. What primary-source confirmation would you require — interviews, restructuring documents, mediation reports — before relying on this case for a deployment decision?

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · incident analysis · safety dashboards

Amazon Hiring AI – Trained Bias, Deprecated 2017

D Designed Out **K** Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER journalism-tier — source confidence flagged; future validation ongoing.

IMPACT Amazon internal recruiting-algorithm project initiated 2014, team disbanded by early 2017 (Reuters reported Oct 2018) after engineers determined the model could not be debiased; trained on ten years of historical resume data in which men predominated in technical roles; the model downgraded resumes containing the word "women's" and resumes from all-women's colleges; Reuters single-source investigation Oct 10 2018

IN BRIEF

Jeffrey Dastin's Reuters investigation, published October 10, 2018, reported that Amazon had initiated an internal recruiting-algorithm project in 2014 with the goal of automating resume screening for technical roles, and had deprecated the project in 2017 – 2018 after engineers determined the model could not be debiased. The model was trained on ten years of historical resume data in which men predominated in technical roles. The training-data composition encoded an association between gender-correlated features and role suitability; engineers found that the model downgraded resumes containing the word "women's" (as in "women's chess club captain") and downgraded resumes from all-women's colleges. Attempts to remove the offending features surfaced additional features carrying the same signal — the trained bias could not be debiased through feature engineering inside the model. The case rests on Reuters single-source reporting; Amazon never published the technical detail. The journalism-tier evidence-flag under the title carries the standing language. The load-bearing teaching point — that trained bias cannot in general be debiased through downstream feature manipulation — is the case's curricular value even as the specific case details remain limited to journalism-tier sourcing.

BACKGROUND

The recruiting-algorithm project that Reuters reported on was an internal Amazon effort initiated in 2014, with the stated goal of automating the early-stage screening of resumes for technical roles. The architectural premise was to train a model on historical resume data — applicants who had been hired and applicants who had not, across the preceding decade — and to score new resumes against the learned pattern. The premise was operationally appealing to a company hiring at Amazon's scale and was consistent with contemporary

practice in algorithmic hiring across the technology sector. The seam the project's deprecation surfaced is structural to the premise itself.⁰

WHAT HAPPENED

The training data composition encoded the gender imbalance of historical Amazon technical hiring. Men predominated in technical roles across the ten-year window, and the resumes that had been hired carried gender-correlated features — wording, extracurricular activities, college affiliations, vocabulary choice — that the model learned to associate with hire-suitability. The Reuters investigation reports that engineers found the model downgrading resumes that contained the word “women’s” (in contexts like “women’s chess club captain”) and downgrading resumes from all-women’s colleges. Attempts to remove the offending features did not eliminate the pattern; the remaining features carried correlated signal that reproduced the same downgrade. The trained bias was structural to the training-data composition and could not be debiased through downstream feature engineering inside the model architecture.¹

THE INVESTIGATION

Amazon disbanded the project team by early 2017 and has said the model was never used by recruiters to evaluate candidates. Reuters reported, however, that recruiters “looked at the recommendations” the tool generated without relying solely on them — so “never deployed at production scale” is the company’s account rather than an independently verified fact. The withdrawal is the load-bearing decision in the case: the engineering team determined that the model could not be made fair, and the organization withdrew the project rather than deploying it. The case’s evidentiary structure rests on Reuters’ single-source reporting — the company never published the technical detail, and the specific mechanism by which engineers verified the irreducibility of the bias is not available in the public record at the level of a peer-reviewed study or a regulator’s audit. The evidence-flag under the title is binding: the case is journalism-tier, and future validation of the specific technical detail remains ongoing in the sense that the public record has not deepened beyond the 2018 investigation.²

THE CAPABILITY GAP

The case pairs with Case 186 (Bartlett mortgage) for the fairness-through-unawareness-fails thread: removing the gender feature from the training data does not eliminate the gender signal when the remaining features carry correlated signal. Pair with Case 48 (Johnson school surveillance) for the algorithmic-employment-and-surveillance-decisions parallel at a different population. Pair with Case 88 (LiveHint AI bias across foundation models) for the trained-bias-in-foundation-models thread at contemporary scale. The Amazon case is unusual in that the project was deprecated rather than deployed; most cases in the corpus document deployments that ran for years before withdrawal, and the case’s curricular value is partly that the engineering team’s verification of irreducibility led to the decision not to deploy.³

AFTERMATH & REFORM

The load-bearing teaching point is that trained bias cannot in general be debiased through downstream feature manipulation. When the training-data composition encodes a historical disparity in the outcome the model is being trained to predict, the model will learn to repro-

duce the disparity through whatever features remain in the input space, and the remediation is not feature engineering. The remediation is at the construct-validity layer — what is the model being asked to predict, and is the historical record from which the prediction is being learned a defensible target — or at the deployment-architecture layer — what is the role of the model in the decision process, and what human-in-the-loop infrastructure surrounds it. The case’s evidence-tier hedge is binding: the journalism-tier sourcing limits the specificity of the teaching point’s mechanism, but the structural form of the teaching point — choose the construct, then ask whether the historical record supports learning to predict it — is the case’s curricular anchor.

Trained bias cannot in general be debiased through downstream feature manipulation; when the training-data composition encodes a historical disparity in the outcome the model is being trained to predict, the disparity will be reproduced through whatever features remain in the input space.

EDITORS’ SYNTHESIS OF THE REUTERS INVESTIGATION (DASTIN, 2018) AND THE SURROUNDING ALGORITHM-MIC-HIRING LITERATURE.

REFERENCES

1. Dastin, J. (2018), “Amazon scraps secret AI recruiting tool that showed bias against women,” Reuters, October 10, 2018 — the primary investigation.

2. Subsequent commentary: Kim, P. T. (2017), “Data-Driven Discrimination at Work,” William & Mary Law Review 58(3):857–936 — academic frame for the structural pattern the Amazon case instantiates.

3. Raghavan, M., Barocas, S., Kleinberg, J., & Levy, K. (2020), “Mitigating Bias in Algorithmic Hiring: Evaluating Claims and Practices,” Proceedings of FAT* 2020, pp. 469–481 — survey of the algorithmic-hiring landscape into which the Amazon case is positioned.

4. Bogen, M., & Rieke, A. (2018), Help Wanted: An Examination of Hiring Algorithms, Equity, and Bias, Upturn report — contemporary practice survey of algorithmic hiring at the time of the Amazon deprecation.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Amazon Hiring AI is the trained-bias-cannot-be-debiased case at major-technology-company scale. The engineering team determined that the model could not be made fair through feature engineering and the organization deprecated the project rather than deploying it. The journalism-tier evidence-flag is binding for the specific technical detail; the structural teaching point — choose the construct, then ask whether the historical record supports learning to predict it — is the case’s curricular anchor.

LENS APPROACH

Amazon Hiring AI is the choose-the-construct case at hiring- algorithm scale (induced 8.1; LENS D2+D4/PT6; CLO-4 and CLO-5). LENS uses it in Domain 2 (Iterative Development) for the irreducibility-verification-as-deployment-gate discipline and in Domain 4 (Test and Evaluation) for the construct-validity-at-the-training-data-layer anchor. Pair with Case 186 (Bartlett mortgage), Case 48 (Johnson school surveillance), and Case 88 (LiveHint AI bias across foundation models). The journalism-tier evidence-flag is binding; the structural teaching point is the case’s curricular value.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
TECHNOLOGY	D K N	D1 D2 D3 D4 D5	6	8.1	4 · 5
HIRING ALGORITHMS		Iterative Development + D4			
MACHINE LEARNING					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Choose the construct the model is being asked to predict with construct validity in mind before training begins; the Amazon case demonstrates that when the historical record encodes a disparity in the outcome, the disparity will be learned regardless of feature engineering at the input layer.
- ▶ Verify the irreducibility of bias as a deployment gate; the engineering team's determination that the model could not be debiased is what enabled the deprecation decision, and the verification discipline is the curricular target.
- ▶ Treat the deprecation decision as the operational artifact that the verification supports; an organization able to deprecate a project at the verification finding is operating with a different decision architecture than one that defaults to deployment.

IN OPERATION / AFTER

- ▶ Carry the journalism-tier evidence-flag under the title without softening; the case rests on single-source reporting and the future-validation-ongoing language is binding for the specific technical detail.
- ▶ Pair in syllabi with Case 186 (Bartlett) so the fairness-through-unawareness-fails thread is taught at adjacent scales and in adjacent decision domains.
- ▶ Use the case to anchor the construct-choice CLO; the curricular target is the discipline of refusing to deploy when the historical record from which the model would be learned does not defensibly support the prediction the deployment requires.

REFLECTION QUESTIONS

1. Identify a deployment in your domain whose training-data composition encodes a historical disparity in the outcome the model is being trained to predict. What construct-validity question would have to be answered before training proceeds, and who has authority to refuse training when the answer is "no"?
2. Specify the irreducibility-verification protocol your team would run before deploying a model whose training data is known to carry historical bias. What is the operational threshold for "cannot be debiased," and what is the decision rule when the threshold is met?
3. Amazon deprecated the project rather than deploying it. Pick a deployment in your domain that proceeded despite known bias-irreducibility findings, and ask: what would have had to be different in the decision architecture for the verification finding to have led to deprecation instead of deployment?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

Uber ATG / Tempe Fatality

T Training Gap **N** Normalization of Deviance **G** Governance & Trust **H** Human-System Interface

IMPACT One pedestrian killed — the first fatality involving a self-driving vehicle striking a pedestrian

IN BRIEF

On the night of 18 March 2018 a modified Volvo running Uber's self-driving system struck and killed Elaine Herzberg as she crossed a road in Tempe, Arizona — the first pedestrian killed by an autonomous vehicle. The NTSB found the safety operator had been watching a video on her phone, but placed heavy blame on Uber: it had not recognized the risk of automation complacency, trained for it, or enforced its own no-phone policy. The system itself was programmed not to brake when a crash was deemed unavoidable, and could not classify a pedestrian who was not near a crosswalk. The human was kept in the loop as a passive monitor — a role the NTSB noted is chronically unperformable — with no infrastructure to make it work. Uber ATG is the book's defining human-AI-teaming case.

BACKGROUND

Uber's Advanced Technologies Group tested self-driving cars on public roads with a safety operator behind the wheel, present to take over if the automation failed. The role was passive surveillance: watch a system that drove itself well almost all of the time, and intervene in the rare moment it did not.^o That structure asks a person to stay vigilant for an event that almost never comes, the precise condition under which human attention is known to lapse — so the role was set up to demand exactly the kind of sustained monitoring that people are least able to deliver.

WHAT HAPPENED

On 18 March 2018 a modified Volvo SUV in autonomous mode struck and killed Elaine Herzberg as she crossed a road at night in Tempe, Arizona — the first pedestrian killed by a self-driving vehicle. The safety operator was looking down at a video on her phone in the seconds before impact.¹ The phone was not an aberration but the predictable filling of an attention vacuum: a role with nothing to do for hours invites exactly that drift, and nothing in the car's design or the monitoring around the seat pulled the operator's eyes back to the road when it finally mattered.

THE INVESTIGATION

The NTSB's probable cause centered on the operator's failure to monitor the road, but it placed heavy blame on Uber, which "did not adequately recognize the risk of automation complacency and develop effective countermeasures": training was inadequate and the no-phone policy unenforced.² Naming the company alongside the operator was the board's way


of locating the fault in the design of the role rather than the lapse of the person filling it — the policy existed on paper but had no mechanism behind it to make compliance the default. The system compounded the human gap. It was programmed not to apply emergency braking when a crash was judged unavoidable — removing the automated backstop — and it could not classify an object as a pedestrian unless it was near a crosswalk, so the very situation on the road that night fell into a blind spot the software was not built to see.³

THE CAPABILITY GAP

A human was retained not because the designers believed a person could meaningfully catch the failure, but because the regulatory and public posture required one present. The role of “monitor” was assigned without the interface, training, or authority to make it performable — a placeholder for safety rather than an instrument of it. As NTSB board members observed, humans tend to tune out when tasked with monitoring automated systems that work well most of the time. The design was safe only on the assumption that the failure case would not arrive — until it did, and the assumption that had quietly held the whole arrangement together was paid for with a life.⁴

AFTERMATH & REFORM

Uber suspended testing, later exited self-driving, and the case reshaped how the industry and regulators treat safety drivers — toward two-operator teams, driver-monitoring systems, and honest accounting of what a monitor can and cannot do.⁵ Each of those reforms is a concession that the single passive observer had been an unsupported role all along: a second operator shares the vigilance burden, and driver-monitoring closes the attention vacuum the original design left open. Its lasting contribution is the reframing of passive monitoring as a role that must be engineered to be performable — or not assigned at all.



Repeatedly, humans tend to tune out when tasked with monitoring automated systems that work well most of the time.

NTSB BOARD MEMBERS, UBER TEMPE HEARING, 2019

REFERENCES

1. NTSB, Collision Between Vehicle Controlled by Developmental Automated Driving System and Pedestrian, Highway Accident Report HAR-19/03 (2019) — the Tempe collision and probable cause.
2. NTSB HAR-19/03 (2019) — the safety operator’s distraction (watching a video) before impact.
3. NTSB HAR-19/03 (2019) — the suppressed emergency braking and the inability to classify a pedestrian away from a crosswalk.
4. NTSB board members, Tempe hearing (2019) — “humans tend to tune out...” (paraphrase reported by Consumer Reports).
5. R. Parasuraman & D. Manzey (2010), complacency in automation; L. Bainbridge (1983), “Ironies of Automation.”

- NTSB HAR-19/03 (2019) — Uber “did not adequately recognize the risk of automation complacency”; training and policy failures (quoted in part).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Uber ATG is the defining case for the LENS Human-AI Teaming competency. A human was retained in the system not because the designers believed a human could meaningfully act, but because the regulatory architecture required a human be present. The role of “monitor” was assigned without the capability infrastructure to support it. The system that resulted was performable only on the assumption that the failure case would not arrive — until it did.

LENS APPROACH

LENS treats this case in LEN 2 as the live exemplar of monitoring as an unsupportable role. Students reconstruct the capability requirements for the safety operator and design the interface, training, and authority structure that would have made the role performable — or made the case for not retaining the role at all.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AUTONOMOUS SYSTEMS	→ T N G H →	D1 D2 D3 D4 D5	6	3.2	3
Human-System Collaboration					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Engineer the monitoring role to be performable — give the operator an active task, a usable interface, and the authority to act — or do not assign a human backstop you do not expect to work.
- Keep the automated emergency-braking backstop active rather than suppressing it, so the system does not silently remove its own last line of defense.
- Validate object classification against the real operating environment, including pedestrians away from crosswalks, before fielding the system on public roads.

IN OPERATION / AFTER

- Enforce the no-phone and attention policies with driver-monitoring that detects and corrects drift in real time, not a written rule alone.
- Audit safety-operator attention data continuously, treating sustained lapses as a design failure of the role rather than a fault of the individual.
- Deploy two-operator teams or equivalent redundancy so the vigilance burden does not rest on a single person doing an unsupportable job.

REFLECTION QUESTIONS

- Identify a passive-monitor role in your domain. What evidence would tell you the role is or is not performable as designed?
- The Tempe vehicle was programmed not to brake when a crash was unavoidable. Reconstruct the design rationale and propose the deliverable that should have prevented that decision.
- Uber had a no-phone policy with nothing to enforce it. What is a rule in your domain that exists on paper but lacks the mechanism to make compliance the default — and how would you engineer that mechanism?

WHO BUILDS THIS learning science & instructional design · safety science & organizational analysis · governance, ethics & measurement · human factors & interaction design — plus domain experts and a learning engineer to integrate. **TOOLS** scenario simulation · competency models · incident analysis · safety dashboards · governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs

Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity

D Designed Out **G** Governance & Trust **N** Normalization of Deviance

IMPACT Algorithmic mortgage underwriting reduced face-to-face discrimination but preserved a measured pricing disparity even when race was excluded from the inputs — the variable's omission did not fix the harm

IN BRIEF

Bartlett, Morse, Stanton, and Wallace (Journal of Financial Economics, 2022) analyzed several million US mortgage applications across the fintech transition and documented two patterns that together unsettle a common intuition. Algorithmic underwriting was, on average, less discriminatory than face-to-face underwriting along the **acceptance** margin — fintech lenders accept-rejected Black and Latino applicants more equitably than traditional lenders. But the **pricing** margin persisted: minority borrowers were charged systematically higher rates for equivalent loans even when race was not in the model inputs. The excluded variable did not stay excluded; it re-entered through geography, credit-score history, and other features that correlate with race in the present US population. The case is the canonical instance of why “fairness through unawareness” is not, by itself, fairness. It pairs directly with the Coots et al. fintech fairness audit (Case 196, preprint-tier): the next teaching step is that competing fairness definitions can disagree about what counts as fair even when the inputs are held constant.

BACKGROUND

The US mortgage market is the most consequential consumer-credit market in which algorithmic underwriting now competes with face-to-face underwriting at large scale. Fintech entrants over the 2010s automated significant parts of the application-to-approval pipeline, and the public-policy question that followed was whether automation reduced or preserved the discrimination documented in traditional lending.⁰

WHAT HAPPENED

Bartlett et al. assembled a corpus of several million applications from a period spanning fintech adoption, and decomposed discrimination along two distinct margins: the **acceptance** decision (do you get a loan) and the **pricing** decision (what interest rate you pay if you do). On acceptance, algorithmic lenders did better: Black and Latino applicants with equivalent observable characteristics were accepted at rates closer to those of comparable White applicants than face-to-face lending produced. The result is consistent with the intuition that removing the loan officer removes some of the channel through which bias entered.¹

THE INVESTIGATION

The pricing margin tells the harder story. Even when race was not among the model inputs, minority borrowers were charged systematically higher rates than equivalent White borrowers — by about 8 basis points on purchase loans and 4 on refinances in the paper’s central estimate. The disparity did not vanish because the variable was omitted; the variable returned through features correlated with race in the present US population — geography, credit-history depth, and other proxies that the model is allowed to use and that carry the historical signal of where lending has and has not flowed.²

THE CAPABILITY GAP

What the case teaches is the structural form of “fairness through unawareness”: when protected attributes are excluded from a model that operates over a population in which other admissible features correlate with the protected attribute, the model can preserve the disparity it was meant to remove. Omission shifts the channel of discrimination but does not close it. The capability deliverable is a measurement architecture that surfaces disparate impact in the **outputs**, not assurance about the **inputs**. The acceptance/pricing asymmetry also reframes the policy question: a model can be more equitable on one decision margin and unchanged on another within the same transaction.³

AFTERMATH & REFORM

The case is the headline mortgage-finance instance of a pattern that now surfaces in clinical algorithms (eGFR, Cases 25, 26 and 6), in hiring and proctoring tools (Case 47 small-tier proctoring bias), and in welfare administration (SyRI). It pairs with the Coots et al. fintech fairness audit (Case 196) which shows the next layer: once the practitioner accepts that omission is not the answer, competing fairness definitions disagree about which adjustment is the right one — and the choice has to be made on judgment under irreducible uncertainty, not on a technical optimum.

The omitted variable does not stay omitted. It re-enters through every feature that carries the same signal.

EDITORS’ SYNTHESIS OF BARTLETT ET AL. (2022) AND DWORK ET AL. (2012).

REFERENCES

1. Bartlett, Morse, Stanton, & Wallace (2022), “Consumer-lending discrimination in the FinTech era,” *Journal of Financial Economics* 143(1):30–56, doi:10.1016/j.jfineco.2021.05.047.
2. Consumer Financial Protection Bureau, Mortgage Market Activity and Trends (annual HMDA reports), supporting the population-level disparities backdrop.
3. Dwork et al. (2012), “Fairness Through Awareness,” ITCS 2012 — the foundational technical statement of the limits of unawareness.
4. Mitchell, Potash, Barocas, D’Amour, & Lum (2021), “Algorithmic Fairness: Choices, Assumptions, and Definitions,” *Annual Review of Statistics and Its Application* 8:141–163 — the competing-definitions framing.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Bartlett et al. is the canonical mortgage-finance instance of why “fairness through unawareness” is not fairness. Algorithmic underwriting reduced the acceptance-margin disparity but the pricing-margin disparity persisted even with race excluded — because correlated admissible features carry the same signal. The capability deliverable is output-level disparate-impact measurement, not input-level assurance.

LENS APPROACH

Bartlett is the headline equity-and-construct case in consumer credit (induced 8.2; LENS D4+D3/PT6). LENS uses it in Domain 4 (Test and Evaluation) for the CLO **Fairness beyond omission**, in Domain 3 (Human-System Collaboration) for delegation to underwriting models, and in Domain 5 (Navigating Sociotechnical Constraints) for the fair-lending regulatory regime. Direct pair with Case 196 (Coots fintech fairness audit). Adjacent to the race-construct trio in clinical medicine (Cases 25, 26 and 6) — same structural lesson at the construct-definition layer rather than the pricing layer.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
FINANCE	→ D G N	→ D1 D2 D3 D4 D5	6	8.2	4 · 5 · 3
ALGO FAIRNESS	Test and Evaluation + D3				

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Specify, before deployment, the disparate-impact measurement on outputs (acceptance rate, pricing) stratified by protected attribute, with a pre-registered threshold for what would trigger remediation.
- ▶ Audit the model's correlated features for the channel through which an excluded attribute can re-enter — geography, credit-history depth, transaction venue — and decide whether each feature's predictive value justifies its disparate-impact contribution.
- ▶ Separate acceptance and pricing as distinct decision margins; do not assume parity on one implies parity on the other.

IN OPERATION / AFTER

- ▶ Publish stratified outcome metrics at intervals long enough for selection effects to settle; the central finding required a multi-year panel.
- ▶ Treat omission of the protected attribute as a baseline, not a remedy; the test of fairness is the output distribution, not the input set.
- ▶ When the measurement surfaces a disparity, name the fairness definition under which it is a problem — group calibration, equalized odds, demographic parity — and the trade-offs of the chosen remediation.

REFLECTION QUESTIONS

1. Identify a model in your domain where a protected attribute is excluded from the inputs. Which admissible features correlate with the excluded attribute in your population? What is the channel through which the excluded variable could re-enter?
2. Design the disparate-impact measurement you would publish at intervals after deployment. Specify the fairness definition, the decision margin (acceptance, pricing, escalation), and the threshold that would trigger remediation.
3. Bartlett's central finding is that the acceptance margin can be more equitable than the face-to-face baseline while the pricing margin remains unchanged. What policy or design intervention would address each margin without assuming parity on one implies parity on the other?

WHO BUILDS THIS systems & human-factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

CASE 189

GOVERNMENT/WELFARE (NETHERLANDS)

ALGORITHMIC DECISION-MAKING

PUBLIC-SECTOR AI

2014 – 2020

Dutch SyRI — Welfare–Fraud Risk Scoring Halted on Rights Grounds

D Designed Out **G** Governance & Trust **N** Normalization of Deviance

IMPACT An opaque risk-scoring system that combined up to 17 categories of previously siloed government data on citizens — disproportionately deployed in low-income and minority neighborhoods — halted by the District Court of The Hague in 2020 as a violation of Article 8 ECHR; also reported as operationally ineffective

IN BRIEF

The Dutch System Risk Indication (SyRI) combined up to 17 categories of previously siloed government data — tax, benefits, housing, education — to produce risk scores intended to flag potential welfare fraud. Deployed from 2014, it was targeted at low-income and minority neighborhoods. On February 5, 2020 the District Court of The Hague halted the program as a violation of Article 8 ECHR (right to private life) — one of the first times a court anywhere stopped a welfare-AI system on human-rights grounds. Investigative reporting also found that, on its own terms, the system did not work: none of the algorithmic investigations had detected new fraud. The case is the paired contrast to the Open University (Case 81): SyRI's governance objection was correct, not dissolvable by design — the system was both rights-violating and ineffective. The court left open that a more transparent system could pass, which makes the case a governance-objection-diagnostic teaching point rather than a verdict against all public-sector AI: design can correct some objections; some objections are correct, and no design dissolves them.

BACKGROUND

SyRI — the System Risk Indication — was a Dutch government analytics program intended to surface potential welfare fraud by combining up to seventeen categories of previously siloed government data on citizens: tax, benefits, housing, education, employment. Deployed from 2014, it operated as a closed system: citizens were not informed they were being scored; the model and its inputs were not disclosed; the program was targeted at specific low-income and minority neighborhoods.^o

WHAT HAPPENED

The governance objection was raised by a coalition of civil society organizations and individual citizens who sued the Dutch state, arguing that SyRI violated Article 8 of the European Convention on Human Rights — the right to private life. The case was heard by the District

Court of The Hague. On February 5, 2020, the court halted the program: the system was found insufficiently transparent and disproportionate to its stated aim. The ruling is one of the first instances anywhere of a court stopping a welfare-AI system on human-rights grounds.¹

THE INVESTIGATION

Investigative reporting alongside the court case surfaced a second finding: on its own operational terms, SyRI was ineffective. Reporting documented that none of the algorithmic investigations the program had launched had detected new fraud — the system had been deployed for six years and the headline capability claim had not been demonstrated. The case is structurally important because it is both rights-violating **and** ineffective; the rights argument did not depend on the effectiveness gap, and the effectiveness gap did not depend on the rights argument. The two failures landed on the same system.²

THE CAPABILITY GAP

What the court explicitly left open is also pedagogically load-bearing. The ruling did not say no welfare-analytics system could pass Article 8; it said **this** system, in its specific opacity and disproportionality, could not. A more transparent system, with auditability and a credibly narrower scope, might pass. The case is therefore the governance-objection-diagnostic counter to the Open University (Case 81): both faced credible governance objections at the same era; OU's was about trust and accountability and was dissolved by a co-authored consent architecture; SyRI's was about rights and proportionality and could not be dissolved by design because the design was the rights violation.³

AFTERMATH & REFORM

The pair (Cases 81 + 189) is the case-grounded basis for the governance-objection diagnostic proposed in the v2 research backbone: distinguishing a governance objection that good design can dissolve from one that correctly signals the system should not deploy. The capability is to make the diagnostic call before deployment, not after — and the diagnostic itself is testable: a system whose operational claim has not been demonstrated after years of deployment, and whose data subjects have not been informed they are being scored, is not a case where design can fix the governance problem.

Some governance objections are correct. Design does not dissolve them — it is what the objection is to.

EDITORS' SYNTHESIS OF THE SYRI RULING AND RACHOVITSA & JOHANN (2022).

REFERENCES

1. District Court of The Hague (2020), judgment of 5 February 2020 (NJCM et al. v. State of the Netherlands), ECLI:NL:RBDHA:2020:865.
2. Rachovitsa & Johann (2022), "The Human Rights Implications of the Use of AI in the Digital Welfare State," Human Rights Law Review 22(2), doi:10.1093/hrlr/ngac010.
3. Library of Congress Global Legal Monitor (2020), report on the SyRI ruling.
4. Algorithm Watch (2020), case study of the SyRI ruling and its implications for public-sector AI in Europe.

THE LEARNING ENGINEERING LENS

LE INSIGHT

SyRI is the canonical case in the corpus of a governance objection that was correct. The system was both rights-violating (Article 8 ECHR) and operationally ineffective (no new fraud detected). The court ruling is system-specific — it does not foreclose all public-sector AI — which is what makes the case a governance-objection diagnostic teaching point rather than a categorical verdict.

LENS APPROACH

SyRI is the negative Domain 5 / Problem Type 6 governance case (induced 5.1; LENS D5+D3/PT6) drafted as the contrast to the Open University (Case 81). LENS uses the pair in Domain 5 (Navigating Sociotechnical Constraints) for the governance-objection diagnostic — distinguishing dissolvable from correct objections — and in Domain 3 (Human-System Collaboration) for the delegation-with-revocation CLO: the court was the revocation channel because the system did not have one of its own. Adjacent to the AI-delegation typology (TREWS / Epic / SyRI / Watson) the v2 corpus builds.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
GOVERNMENT/ WELFARE (NETHERLANDS) ALGORITHMIC DECISION-MAKING PUBLIC-SECTOR AI	D G N	D1 D2 D3 D4 D5	6	5.1	5 · 3
→ Navigating Sociotechnical Constraints + D3					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Before deploying a public-sector analytics system, document the data subjects who will be scored, the inputs, the model behavior, and the proportionality argument; opacity is a governance liability that compounds over the life of the deployment.
- ▶ Specify in advance the operational evidence that would demonstrate the capability claim — for SyRI, fraud actually detected by algorithmic investigation — with a published reporting cadence and a pre-registered threshold for what would count as success or failure.
- ▶ Conduct the governance-objection diagnostic openly: is the objection about trust and accountability (potentially dissolvable by design) or about rights and proportionality (not dissolvable by design)? Different answers imply different deployment paths.

IN OPERATION / AFTER

- ▶ If a court or rights body halts the program, treat the ruling as governance information about the design specifically, not as a verdict about all related systems; redesign with the named deficiencies addressed.
- ▶ Publish the operational effectiveness evidence on the cadence specified at launch; six years without published evidence of the headline capability is itself a governance finding.
- ▶ Treat the rights-and-proportionality channel as a separate evaluation from the technical-accuracy channel; a system can be technically accurate and still fail the proportionality test, and SyRI failed both.

REFLECTION QUESTIONS

1. Identify a public-sector analytics system in your jurisdiction. Are the data subjects informed they are being scored? Is the model and its inputs disclosed? What is the proportionality argument the deploying agency would make if challenged?
2. Specify the pre-registered operational evidence — what the system has to demonstrate, on what cadence — that would let a deployment be evaluated against its capability claim. SyRI ran for six years without published evidence of the headline claim.
3. The pair OU (Case 81) and SyRI (Case 189) teaches the governance-objection diagnostic. Construct a candidate diagnostic for your own domain: what features of an objection indicate that good design could dissolve it, and what features indicate that the objection is to the design itself?

WHO BUILDS THIS systems & human-factors engineering · governance, ethics & measurement · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · governance frameworks · bias & evidence audits · incident analysis · safety dashboards

Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment

G Governance & Trust **K** Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER journalism-tier — source confidence flagged; future validation ongoing.

IMPACT On 24 October 2023 the California DMV suspended Cruise's driverless deployment and testing permits, citing the company's misrepresentation of safety-relevant information after a robotaxi dragged a pedestrian roughly 20 feet at ~7 mph following an initial stop — disclosure posture, not the underlying collision sequence, was the load-bearing failure

IN BRIEF

On 2 October 2023 a pedestrian was struck by a human-driven vehicle in San Francisco and propelled into the path of a Cruise robotaxi. The robotaxi came to a stop, then performed a pullover maneuver that dragged the pedestrian roughly 20 feet at about 7 mph before stopping again. Cruise initially shared video of the initial stop with regulators and reporters but did not disclose the subsequent movement; the California DMV learned of the pullover from another agency and obtained the fuller video weeks later. On 24 October the DMV suspended Cruise's driverless deployment and testing permits, citing misrepresentation of safety-relevant information. The case is the explicit contrast to Case 199 (Waymo): the same regulatory regime, the same delegation problem, the opposite governance choice. Partial disclosure was the load-bearing failure mode, not the underlying collision sequence. The evidence-tier flag is journalism: the DMV's published Order of Suspension is investigation-grade, but the partial-disclosure mechanism and internal timeline are reconstructed from contemporaneous reporting at TechCrunch, NBC News, SF Standard, and Mission Local. The underlying Cruise internal post-mortem is referenced in public statements but not fully public. Future validation will continue on long-run regulatory consequences.

BACKGROUND

The night of 2 October 2023, on a San Francisco street, a pedestrian was struck by a human-driven vehicle and propelled into the path of a Cruise robotaxi operating without a safety driver. The robotaxi detected the collision and came to a stop. The vehicle then executed what Cruise later described as a pullover maneuver, traveling roughly 20 feet at about 7 mph with the pedestrian pinned beneath it. The collision sequence itself involved a human driver, not Cruise's software, but the post-impact behavior was Cruise's system acting on its own.^o

WHAT HAPPENED

In the immediate days after the incident, Cruise shared video of the initial stop with reporters and regulators. The pullover maneuver and the dragging of the pedestrian were not included in those initial disclosures. The California DMV did not learn of the full sequence from Cruise; it learned of the pullover from another agency and obtained the fuller video weeks after the incident. The mechanism of failure shifted in the regulator's view from the collision to the company's disclosure posture.¹

THE INVESTIGATION

On 24 October 2023 the DMV issued an Order of Suspension revoking Cruise's driverless deployment permit and its driverless testing permit. The order cited misrepresentation of safety-relevant information — that the company had "misrepresented information related to the safety of its autonomous technology" — as a load-bearing reason for the revocation. The Order of Suspension is the investigation-grade artifact in this case; the reconstruction of the partial-disclosure mechanism and the internal timeline rests on contemporaneous journalism.²

THE CAPABILITY GAP

The case is the explicit foil to Case 199 (Waymo). Same regulatory regime, same delegation problem, opposite governance choice. Where Waymo answered an opacity objection by engineering a published safety-case framework and commissioning third-party audits, Cruise answered an incident-disclosure obligation by sharing partial video. The same DMV that permitted Waymo to continue revoked Cruise's permits. The journalism-tier flag is load-bearing: the DMV's published order is investigation-grade, but the precise sequence of internal decision-making rests on TechCrunch, NBC News, SF Standard, and Mission Local reporting that has not been independently corroborated by the company's full post-mortem.³

AFTERMATH & REFORM

The LENS teaching point pairs directly with Waymo. The new CLO **Delegation with revocation** requires that the deploying organization design the disclosure architecture **before** the failure event — what will be reported, on what cadence, to which oversight body, with what verification. Partial disclosure under crisis is not a strategy; it is what happens when no architecture was designed. Future validation will continue on the long-run regulatory consequences as Cruise pursues reinstatement and as the broader AV regulatory regime updates its disclosure requirements in light of the 2023 events.⁴

The collision involved a human driver. The suspension was about what the company chose not to show.

EDITORS' SYNTHESIS OF THE CALIFORNIA DMV ORDER OF SUSPENSION AND THE CONTEMPORANEOUS REPORTING.

REFERENCES

1. California DMV (24 October 2023), Order of Suspension, Cruise LLC — driverless deployment and testing permits.
2. NBC News (October 2023), reporting on the Cruise pedestrian incident and DMV suspension.
3. TechCrunch (October 2023), incident-disclosure reconstruction.
4. SF Standard (October 2023), San Francisco AV regulatory reporting.
5. Mission Local (October 2023), San Francisco-specific incident reporting.
6. Cruise public statements and partial post-mortem references, October–November 2023.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Cruise is the foil to Waymo: same regulatory regime, same delegation problem, opposite governance choice. The DMV’s Order of Suspension is investigation-grade; the partial-disclosure mechanism is reconstructed from journalism. The evidence-tier flag is load-bearing — the internal timeline is journalism-tier and future validation continues as more of the company’s own post-mortem becomes public.

LENS APPROACH

Cruise is the AV partial-disclosure failure (induced 5.4; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the CLO **Delegation with revocation** — the case shows what revocation looks like when the disclosure architecture fails — and in Domain 3 (Emerging Systems and Human-System Collaboration) for the oversight obligations that follow when a system acts autonomously after a triggering event. Direct foil to Case 199 (Waymo); pairs with Case 200 (CPUC) on the regulator-side.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AUTONOMOUS VEHICLES	G K N	D1 D2 D3 D4 D5	6	5.4	5 · 3
REGULATORY GOVERNANCE	→	→ Navigating Sociotechnical Constraints			
INCIDENT DISCLOSURE					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the incident-disclosure architecture before deployment, not under crisis — what will be reported, on what cadence, to which oversight body, with what verification by which agency.
- ▶ Treat the disclosure obligation as load-bearing on its own — partial disclosure of a safety-relevant event is not “less bad” than non-disclosure; under the regulator’s framing it can be the failure mode that triggers revocation.
- ▶ Build the verification path the regulator will use into the architecture. The DMV learning the fuller sequence from another agency is the failure mode the architecture has to preclude.

IN OPERATION / AFTER

- ▶ Use the CLO **Delegation with revocation**: revocation pathways must be designed and exercisable, and the deploying organization should expect the regulator to exercise them when the disclosure architecture fails.
- ▶ Pair the case with Case 199 (Waymo) in any LENS Domain 4 module — the contrast between an engineered legitimacy artifact and partial crisis disclosure is the teaching point, not either case alone.
- ▶ Carry the journalism-tier flag honestly: the DMV Order is the investigation-grade primary; the internal-timeline reconstructions are journalistic; future validation will improve as Cruise’s own post-mortem and any subsequent litigation discovery enter the public record.

REFLECTION QUESTIONS

1. Imagine you operate an autonomous system that has just been involved in a safety-relevant event. Design the disclosure decision: what is reported, to whom, on what cadence, with what verification — **before** you have the lawyer’s advice on what the disclosure obligation strictly requires. Where does your architecture leave you exposed?
2. Compare Cases 199 (Waymo) and 190 (Cruise) as a paired teaching unit. What is the smallest pre-incident artifact a deploying organization could publish that would make the post-incident disclosure architecture credible to a regulator?
3. The case rests partly on journalism-tier reconstruction of internal decisions. What evidence would you want to see — court discovery, the company’s full post-mortem, a multi-source corroboration — before treating any specific internal-timeline claim as decision-grade?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

CASE 192

FINANCIAL SERVICES

CONSUMER CREDIT

ALGORITHMIC DECISION-MAKING

2019 – 2021

Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model

D Designed Out **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT New York State Department of Financial Services investigation March 2021 found no violation of New York anti-discrimination law in Apple Card credit-limit decisions following David Heinemeier Hansson's November 2019 viral allegation that his wife received a credit limit approximately 20 times lower despite shared assets; DFS documented "lack of transparency" as the structural problem and required Goldman Sachs to overhaul its customer-service process

IN BRIEF

On November 7, 2019, software developer David Heinemeier Hansson posted on Twitter that his Apple Card credit limit had been set approximately 20 times higher than his wife Jamie Heinemeier Hansson's, despite the couple filing taxes jointly and Jamie having a higher credit score. The thread went viral, with Apple co-founder Steve Wozniak reporting a similar disparity with his wife. The New York State Department of Financial Services opened an investigation. Its March 2021 report concluded that the credit-decisioning algorithm operated by Goldman Sachs (the issuing bank for Apple Card) did not violate New York anti-discrimination law, finding no statutory finding of intent or disparate-impact violation. DFS documented "lack of transparency" as the structural problem: Goldman Sachs could not adequately explain individual credit decisions to applicants who challenged them. DFS required Goldman Sachs to overhaul its customer-service process. The case pairs with Case 186 (Bartlett mortgage), Case 196 (Coots), and Case 86 (Gándara). The DFS finding of "no violation but lack of transparency" is the load-bearing nuance.

BACKGROUND

The Apple Card launched in August 2019 as a co-branded consumer credit product issued by Goldman Sachs, with underwriting and credit-line decisions made by Goldman's consumer-banking unit. The product was Goldman's first retail consumer credit product at meaningful scale, and the decisioning architecture was built de novo against contemporary algorithmic-underwriting practice. On November 7, 2019, David Heinemeier Hansson posted that his Apple Card credit limit was approximately 20 times his wife's, despite jointly filed taxes and Jamie's higher credit score. The thread surfaced similar reports — Steve Wozniak

named the same pattern with his wife — and rapidly moved from social media to regulatory attention.⁹

WHAT HAPPENED

The structural seam the case opens is that Goldman Sachs could not adequately explain individual credit decisions to the applicants who challenged them. When Jamie Heinemeier Hansson asked Goldman to explain why her credit limit was lower, the customer-service response was that the algorithm had set the limit and that the decision could not be explained at the individual level. The escalation moved through standard customer-service channels, then to social-media-amplified public pressure, then to regulatory investigation, and at no point along the path did Goldman surface an account of why the decision had been made. The explainability gap was operational rather than algorithmic in the narrow sense: even if the underlying model was defensible at the population level, the bank had not built the infrastructure to defend individual decisions to applicants.¹

THE INVESTIGATION

The New York State Department of Financial Services opened its investigation in late November 2019 and released its findings in March 2021. The headline finding was that DFS did not find a violation of New York anti-discrimination law — neither intentional discrimination nor an actionable disparate-impact violation under the applicable standards. The investigation reviewed approximately 400,000 New York State credit-line decisions. The hedge is binding: DFS's "no violation" finding is specific to the statutory standard the agency applied, not a general finding that the decisioning was fair or non-discriminatory across all criteria. What DFS did find was "lack of transparency" as the structural problem and required Goldman Sachs to overhaul its customer-service process, build appeal mechanisms, and document its credit-decisioning explanations at the individual-applicant level.²

THE CAPABILITY GAP

The case pairs with Case 186 (Bartlett mortgage) for the consumer-credit-fairness thread at adjacent scale and regulatory regime. Pair with Case 196 (Coots) for the competing-fairness-definitions thread; the DFS standard is one of several available standards, and the case teaches that "no violation under a specific statutory standard" is not "fair." Pair with Case 86 (Gándara) for the explainability-of-individual-predictions thread at a different population and scale. The case is unusual in the casebook for the speed of regulatory response — DFS opened the investigation within weeks of the viral thread — and for the structural conclusion that the failure was explainability rather than the decisioning algorithm itself.³

AFTERMATH & REFORM

The load-bearing hedge is the precise DFS finding. The case does not teach that Apple Card was unfair, and it does not teach that DFS found Apple Card was fair. It teaches that under the specific statutory standard the agency applied, no violation was found, and that the structural failure the agency named was "lack of transparency" — Goldman Sachs's inability to explain individual credit decisions to applicants who challenged them. The human-in-the-loop CLO is anchored by the case at the appeal-and-explanation seam: a consequential-decision system that cannot explain its individual decisions to affected applicants does not have a functioning human-in-the-loop appeal mechanism, and the absence of the mechanism is the governance failure whether or not the decisioning is statistically defensible.

DFS did not find a violation of New York anti-discrimination law; DFS did find lack of transparency as the structural problem, and required Goldman Sachs to overhaul its customer-service process – the case teaches that “no violation under a statutory standard” is not “fair.”

EDITORS’ SYNTHESIS OF THE NY DEPARTMENT OF FINANCIAL SERVICES REPORT ON THE APPLE CARD INVESTIGATION (MARCH 2021).

REFERENCES

1. New York State Department of Financial Services, Report on Apple Card Investigation, March 23, 2021.

2. Heinemeier Hansson, D. (2019), Twitter thread of November 7, 2019, archived in DFS investigation record and contemporaneous press coverage (Bloomberg, New York Times, Wall Street Journal, November 2019).

3. Vigdor, N. (2019), “Apple Card Investigated After Gender Discrimination Complaints,” The New York Times, November 10, 2019.

4. Goldman Sachs Bank USA, public response and credit-line-review process documentation submitted to DFS during the investigation (2019 – 2021).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Apple Card is the consumer-credit explainability case at deployment scale. DFS found no violation of New York anti- discrimination law under the applicable statutory standard, and DFS also found lack of transparency as the structural problem; Goldman Sachs was required to overhaul its customer-service process and build individual-applicant appeal mechanisms. The load-bearing hedge is the precision of the DFS finding – neither “fair” nor “unfair,” but “no violation under this standard, transparency gap as the structural problem.”

LENS APPROACH

Apple Card is the explainability-as-governance case at consumer-credit scale (induced 5.2; LENS D5/PT6; CLO-5 and CLO-3). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for the appeal-and-explanation-infrastructure discipline and in Domain 3 (Human-System Collaboration) for the human-in-the-loop-for-consequential-decisions anchor at the appeal seam. Pair with Case 186 (Bartlett mortgage), Case 196 (Coots competing fairness definitions), and Case 86 (Gándara explainability of individual predictions). The precise DFS finding – “no violation but lack of transparency” – is the load-bearing hedge.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
FINANCIAL SERVICES	D K N	D1 D2 D3 D4 D5	6	5.2	5 · 3
CONSUMER CREDIT	→	→			
ALGORITHMIC DECISION-MAKING		Navigating Sociotechnical Constraints			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- Build the individual-applicant explanation infrastructure as part of the deployment, not as a customer-service afterthought; the Apple Card case demonstrates that a defensible population-level model paired with no individual-explanation infrastructure produces a regulatory finding of lack of transparency.
- Specify the appeal mechanism before the first decision is made; the human-in-the-loop CLO at consumer-credit scale is the appeal-and-explanation

IN OPERATION / AFTER

- Carry the precise DFS finding into print without softening; “no violation under the applicable statutory standard, but lack of transparency as the structural problem” is the load-bearing nuance and the case’s pedagogical value depends on the nuance being preserved.
- Pair in syllabi with Case 186 (Bartlett) so the consumer-credit-fairness regulatory architecture is taught at both the mortgage and credit-card scales.

- seam, and its absence is the governance failure the case names.
- ▶ Treat the customer-service escalation path as a deployment surface, not a support function; the case's escalation went from customer service to social media to regulation in days, and the deployment surface that mattered was the first one.
- ▶ Use the case as the anchor for the explainability-as-governance frame at consumer-credit scale; the curricular target is the appeal-and-explanation infrastructure that converts an algorithmic decision into a contestable one.

REFLECTION QUESTIONS

1. Identify a consequential-decision system in your domain whose individual-applicant explanation infrastructure has not been built. What is the customer-service escalation path when an affected person challenges a decision, and what would the path look like with an appeal-and-explanation seam built into the deployment?
2. Specify the precise statutory or regulatory standard against which your deployment is being evaluated. What does "no violation under this standard" leave open about fairness across other standards, and how would the trade-off be disclosed?
3. The Apple Card escalation moved from customer service to social media to regulatory investigation within days. Pick a deployment in your domain and ask: what is the first deployment surface an affected person encounters when challenging a decision, and what would have to be true for that surface to resolve the challenge before it moves further?

WHO BUILDS THIS systems & human-factors engineering · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** task analysis · design prototyping · knowledge bases · data pipelines · incident analysis · safety dashboards

Bernard Madoff / SEC Failures

G Governance & Trust **K** Knowledge & Institutional Memory **N** Normalization of Deviance

IMPACT ~\$65B Ponzi scheme — the largest in history; SEC repeatedly investigated and cleared Madoff; foundational regulator-capability case

IN BRIEF

Bernard Madoff ran the largest Ponzi scheme in history — roughly \$65 billion in fictitious account value — for years while the SEC repeatedly investigated and cleared him. Financial analyst Harry Markopolos delivered the agency a detailed memo in 2005, “The World’s Largest Hedge Fund is a Fraud,” showing that Madoff’s steady returns were mathematically impossible. The SEC opened an investigation and concluded no action was warranted. Madoff operated until his sons turned him in during the 2008 crisis, when redemptions became impossible to honor. The SEC Inspector General later found the staff assigned lacked the expertise to evaluate Markopolos’s arguments and had deferred to Madoff’s industry stature. The capability gap was at the regulator’s technical-evaluation pipeline, not at the evidence — which was specific and checkable.

BACKGROUND

Bernard Madoff was a former NASDAQ chairman and a respected Wall Street figure whose investment arm reported remarkably steady returns for years. Those returns were entirely fictitious: client funds were never invested, and earlier investors were paid with later investors’ money — a Ponzi scheme that eventually represented some \$65 billion in fabricated account value. The very steadiness that reassured investors was the tell: real markets do not deliver an almost unbroken line of gains, and the absence of the normal volatility was itself evidence the numbers were manufactured.⁰

WHAT HAPPENED

Financial analyst Harry Markopolos concluded by analysis that Madoff’s returns were mathematically impossible and delivered the SEC a detailed memorandum in 2005 titled “The World’s Largest Hedge Fund is a Fraud.” The SEC opened an investigation and concluded that no enforcement action was warranted. The warning was not a vague suspicion but a quantitative case any competent reviewer could in principle retrace, which is what makes the dismissal so telling. Madoff continued operating until December 2008, when the financial crisis made redemptions impossible and his sons reported him.¹

THE INVESTIGATION

The SEC Office of Inspector General’s 2009 report found the agency had received credible, specific complaints across more than a decade and “missed numerous opportunities” to uncover the fraud. The staff assigned to Madoff lacked the expertise to evaluate Markopolos’s technical arguments, and the institutional culture had defaulted to treating


Madoff as a respected industry figure rather than following the evidence. Stature stood in for analysis: the reviewers weighed who Madoff was instead of whether the math could be true, and the reputational halo did the work that scrutiny should have.²

THE CAPABILITY GAP

The complaint was specific; the math was checkable; the institution simply did not have the people to check it. The capability gap was at the regulator's technical-evaluation pipeline, not at the evidence. A regulator whose technical depth lags the entities it oversees cannot act on even a correct and well-documented warning — and because the gap is in the evaluator rather than the tip, more tips would not have helped; the agency could not have used them.³

AFTERMATH & REFORM

The collapse wiped out tens of thousands of investors and prompted SEC reforms to its handling of tips and referrals, its examination procedures, and its recruitment of staff with quantitative and trading expertise, including the creation of an Office of Market Intelligence. Each reform addressed a different part of the same pipeline — getting the tip in, getting it to someone who could read it, and having someone who could — so that a future Markopolos memo would meet an evaluator able to test it. Madoff is a case in which a regulator lacked the technical capability to challenge the evidence in front of it.⁴



The agency missed numerous opportunities to discover the fraud.

PARAPHRASING SEC OFFICE OF INSPECTOR GENERAL REPORT OIG-509 (2009)

REFERENCES

1. SEC Office of Inspector General, Investigation of Failure of the SEC to Uncover Madoff's Ponzi Scheme, Report OIG-509 (2009) — the quoted finding.
2. Markopolos, H. (2010), No One Would Listen — the 2005 memo and its dismissal.
3. United States v. Madoff (2009) — guilty plea and the \$65B figure.
4. SEC OIG (2009) — staff expertise gap and deference to Madoff's stature.
5. SEC post-Madoff reforms (2009–2010), including the Office of Market Intelligence for tip and referral triage.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The Madoff case is the canonical example of a regulator without the technical capability to evaluate the evidence it received. The complaint was specific. The math was checkable. The institution did not have the people to check it. The capability gap was at the regulator’s expertise pipeline, not at the evidence.

LENS APPROACH

LENS uses Madoff in LEN 6 as an operator-to-institution feedback-channel failure: Harry Markopolos sent the SEC a specific, checkable memo — the channel existed and the signal was correct — and the institution had no one able to act on what came through it. The teaching point is the receiving capability a feedback channel requires to be real: a tip line without the technical expertise to evaluate the tip is not a channel. It pairs with the Mark 14 torpedo (Case 133), where the Bureau of Ordnance dismissed submariners’ field reports for the same structural reason — the channel carried the truth and the institution could not receive it. The Ponzi narrative is not the lesson; the channel is.

DOMAIN	FAILURE MODES	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
GOVERNMENT	→ G K N →	D1 D2 D3 D4 D5	3	6.2	2
TECHNOLOGY		Navigating Sociotechnical Constraints			

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
- Staff the evaluation pipeline with quantitative and trading expertise matched to the entities being overseen, so a checkable claim meets someone able to check it. - Build tip-and-referral triage that routes a specific, technical complaint to a qualified evaluator rather than letting the subject’s reputation decide its fate. - Require warnings to be assessed on their merits first, structurally separating the analysis of the math from any weighing of who the subject is.	- Audit closed investigations for cases dismissed despite specific, checkable evidence, surfacing where the evaluator — not the evidence — was the failure point. - Monitor for returns or results too smooth to be real as a leading indicator, treating the absence of expected volatility as a signal to investigate. - Sustain the regulator’s technical depth against the regulated industry’s growing sophistication, since a depth gap quietly reopens the same blind spot over time.

REFLECTION QUESTIONS

1. Identify a regulator in your domain whose technical evaluation capability has not kept pace with the entities it regulates. What is the resulting blind spot?
2. Design the regulator-side technical-capability deliverable that should have allowed the SEC to evaluate the Markopolos memo on its merits.
3. The SEC weighed Madoff’s stature over Markopolos’s math. Design a triage process that forces a tip to be evaluated on its technical merits before the subject’s reputation is allowed to enter the decision.

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

Chapter 14

Algorithms, Governance & Public Systems — What Works — and the Frontier

When governance is designed as an artifact, not litigated after the fact.

The objection dissolved when the evidence was engineered to answer it.

AN OBSERVATION RECURRING ACROSS THE CHAPTER'S CASES.

CASE 194

E-GOVERNMENT

DISTRIBUTED DATA EXCHANGE

NATIONAL IT INFRASTRUCTURE

2001 – PRESENT

Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox)

G Governance & Trust **K** Knowledge & Institutional Memory **N** Normalization of Deviance

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT By December 2024 Estonia reported effectively 100% digitalization of government services across the X-Road data-exchange layer, with sub-five-minute tax filing and >95% of tax declarations filed electronically; the load-bearing self-critique is that the country has now created its own legacy system — the very thing the program set out to avoid

IN BRIEF

Estonia launched X-Road in 2001 as the answer to fragmented government databases: a distributed data-exchange layer that lets services interoperate without forcing a central monolith. By December 2024 the country reported effectively 100% digitalization of government services across X-Road, with sub-five-minute tax filing for more than 95% of the population. The case sits inside C7 not because of the headline outcomes but because of the load-bearing self-critique surfaced in the peer-reviewed analysis: by committing the country to X-Road as the data-exchange layer, Estonia has effectively created its own legacy system — the very thing its founders set out to avoid. The “no-legacy paradox” is the C7 teaching the corpus does not get from any other case: success-as-aging is the failure mode, and the modernization regime must contemplate its own future obsolescence as part of its current design discipline. The evidence-tier flag is rendered under the title: the program is well-documented in peer-reviewed and program-report sources; the self-critical framing rests partly on practitioner reflection that the corpus must carry honestly. Future validation will continue on whether the deliberate generational-replacement plans materialize. Cross-listed with Gap 5 (non-US/UK/EU coverage, Estonia).

BACKGROUND

Estonia’s pre-2001 problem was the standard one for a small state with a recent administrative inheritance: government databases were fragmented, agency-by-agency, with no common exchange layer. X-Road was the architectural answer — a distributed data-exchange protocol that lets each agency keep its own system but interoperate via a shared,

authenticated channel. The design choice avoided the trap of a central monolith and made cross-agency services possible without requiring agencies to surrender their data.⁰

THE INTERVENTION

The deployment trajectory across two decades has been the defining national IT case study in the EU region. By December 2024 the country reported effectively 100% digitalization of government services, with sub-five-minute tax filing and more than 95% of tax declarations filed electronically. The *Mission Mystique* and the *Hiding Hand* chapter (Oxford 2019) and the ICEGOV 2021 historical analysis are the peer-reviewed treatments; the X-Road Global program documentation extends the case to the institutions that have since adopted the protocol elsewhere.¹

HOW IT WORKED

The load-bearing self-critique is what makes this a C7 case and not just a digital-government success story. The analysis surfaces what practitioners and researchers have begun to call the no-legacy paradox: by committing the country to X-Road as the data-exchange layer, Estonia has effectively created its own legacy system. The protocol is now twenty-plus years old; the agency systems that hang off it have accreted their own logic; the generational replacement of the platform itself is now the modernization problem the founders thought they had designed out. Success-as-aging is the failure mode.²

THE EVIDENCE

The evidence-tier flag matters. The peer-reviewed sources cover the program's design and trajectory and the self-critical framing; the program-report and practitioner reflections supply the day-to-day operational picture and the no-legacy paradox in its sharpest form. The country's generational-replacement plans for X-Road exist; whether they materialize at the pace the analysis implies is necessary is the open empirical question. Future validation will continue as the next-generation data-exchange architecture is designed and deployed.³

WHAT TRANSFERRED

What the case adds at the LENS layer is the C7 teaching that no other corpus case supplies: a modernization regime must contemplate its own future obsolescence as part of its current design discipline. The Domain-1 frame applies: systems analysis of the data-exchange layer as an artifact whose own aging trajectory is part of the system's behavior. The Domain-5 frame applies because the sociotechnical commitment of a state to a single data-exchange protocol is itself the governance choice that determines what the future modernization will have to undo. The case is also a Gap-5 non-US/UK case from a small EU state, which the corpus needs.

Success creates its own aging problem. The modernization regime must contemplate its own future obsolescence as part of its current design.

EDITORS' SYNTHESIS OF THE NO-LEGACY PARADOX IN THE ESTONIA X-ROAD ANALYSIS.

REFERENCES

1. Kattel, R., & Mergel, I. (2019), "Estonia's Digital Transformation: Mission Mystique and the Hiding Hand," in Compton, M., & 't Hart, P. (eds.), *Great Policy Successes* (Oxford University Press, 2019) — peer-reviewed analytical chapter.
2. X-Road Global / Nordic Institute for Interoperability Solutions (NIIS) — program documentation and deployment-partner case studies.
3. Republic of Estonia, e-Estonia briefing materials and Year-in-Review documentation (2024) — program-report sourcing.

2. Jackson, Dreyling, & Pappel (2021), "A Historical Analysis on Interoperability in Estonian Data Exchange Architecture," ICEGOV 2021 proceedings, doi:10.1145/3494193.3494209.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Estonia X-Road is the C7 case the corpus needed for the success-as-aging failure mode: a 100%-digitalization program whose own success has now created the legacy system its founders set out to avoid. Evidence is mixed — peer-reviewed analytical chapters plus program-report and practitioner reflection; the generational-replacement trajectory is the open empirical question. Future validation ongoing. Non-US/UK case, Gap-5 cross-listed.

LENS APPROACH

X-Road is the C7 big-tier transition case with a small-tier self-critique (induced 7.1 and 5.3 alternate; LENS D1/PT1). LENS uses it in Domain 1 (Systems Analysis) for CLO-1 — the data-exchange layer is an artifact whose own aging is part of system behavior — and in Domain 5 (Navigating Sociotechnical Constraints) for CLO-5, since the state’s commitment to the protocol is itself the governance choice determining future modernization cost. Pair with Cases 25–26 for the C7 thread and with the big-tier C7 failures the corpus already documents.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
E-GOVERNMENT	G K N	D1 D2 D3 D4 D5	1	7.1	1-5
DISTRIBUTED DATA EXCHANGE	→	→ Systems Analysis			
NATIONAL IT INFRASTRUCTURE					

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Design the data-exchange or platform layer so that its own future obsolescence is contemplated in the current design — the no-legacy paradox is the warning the case carries.
- ▶ Treat the generational-replacement plan as part of the deployment commitment, not as a future problem; the Estonia case is open precisely because the replacement is open.
- ▶ Document the sociotechnical commitment of the state or institution to the platform explicitly, so the cost of future replacement is in the record from the start.

IN OPERATION / AFTER

- ▶ Track the platform’s aging trajectory as a system variable — accreted agency logic, protocol drift, dependency depth — not only its operational uptime.
- ▶ Carry the self-critical framing honestly; the load-bearing teaching is that success creates its own aging problem, not that the program failed.
- ▶ Pair with other C7 cases (Patriot/Dhahran, Ariane 5, Knight Capital, CrowdStrike, TSB) for the failure-and-modernization contrast, and with Cases 116–26 for the small-tier transition thread.

REFLECTION QUESTIONS

1. Identify a platform or data-exchange layer in your context whose success has created a generational-replacement problem the original design did not contemplate. What would the no-legacy-paradox-aware redesign look like?
2. Specify the aging-trajectory variables you would track on the platform (accreted agency logic, protocol drift, dependency depth) so the future-obsolescence question is in the operational record, not a future surprise.
3. The case is peer-reviewed analysis plus program-report plus practitioner reflection. What additional evidence would you require before committing your own institution to a platform that may itself become the legacy system in twenty years?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

Tesla Autopilot — Recurring Fatalities

T Training Gap **N** Normalization of Deviance **G** Governance & Trust **H** Human-System Interface

IMPACT Dozens of fatalities documented in NHTSA's Standing General Order data; first U.S. cases of Level-2 automation contributing to fatal injury

IN BRIEF

Tesla's Autopilot and Full Self-Driving Beta are Level-2 driver-assistance systems: the human driver remains legally and operationally responsible at all times. NHTSA's Standing General Order data has documented dozens of fatal crashes involving Autopilot since the 2016 death of Joshua Brown in Florida, and the pattern is consistent — the system performs capably for long stretches, the driver's monitoring attention attenuates, and an edge case (a stationary object, a faded lane line, a crossing vehicle) produces a collision the inattentive driver fails to catch. NHTSA's investigation has found Tesla's driver-engagement design inadequate to sustain the attention safe operation requires. At consumer scale, Autopilot is the live test of whether passive monitoring of good-enough automation is a role a human can actually perform.

THE SHIFT

Partial driving automation has moved from research vehicles into millions of consumer cars. Level-2 systems like Tesla's Autopilot can steer, accelerate, and brake within their operational design domain, but they require the human driver to monitor continuously and take over instantly — a fundamentally new and demanding role assigned to ordinary, untrained consumers. Where a research program could screen, brief, and instrument its safety drivers, a consumer product reaches everyone who buys the car, with no curriculum and no qualification gate standing between purchase and the monitoring task itself.^o

WHAT IS EMERGING

Since the first fatal Autopilot crash — Joshua Brown, Florida, 2016 — NHTSA's Standing General Order data has documented dozens of fatal crashes involving the system. The pattern is consistent: long periods of capable operation, attenuating driver attention, and then an edge case — a stationary fire truck, a faded lane marking, a perpendicular crossing — that the disengaged driver fails to catch in time. The very reliability that makes the system attractive is what erodes the vigilance it depends on, so each uneventful mile quietly raises the odds that the next intervention will come too late.¹

THE CAPABILITY QUESTION

The proximate cause in each case is the driver, who was legally responsible. But the deeper question is whether sustained vigilant monitoring of an automation that works well most of the time is a role a human can perform at all. Naming the feature "Autopilot" and designing weak engagement checks shaped the very inattention the system then blamed on the

operator — so the architecture both invited the disengagement and reserved the liability for the person least positioned to resist it.²

EARLY EVIDENCE

NHTSA's engineering analysis (ODI EA22-002, closed April 2024 and superseded by the December 2023 recall of roughly two million vehicles) identified Tesla's driver-engagement design as inadequate to maintain the attention safe operation requires, and the recurring fatality pattern across NTSB reports suggests passive monitoring is not a sustainable role as currently engineered. Decades of automation-complacency research point the same way — the finding is not that any one driver failed but that the role asks a human to stay alert to a system precisely because it almost never needs them, a demand the evidence keeps showing is not reliably met.³

OPEN PROBLEMS

Tesla Autopilot is the consumer-scale version of the Uber ATG problem (Case 183): a passive-monitoring role deployed without the capability infrastructure — training, engagement design, attention measurement — to make it performable. The open problem is what driver-engagement architecture, if any, could make Level-2 monitoring sustainable for an average driver over years of use, and whether the answer is a better attention check or a concession that the role itself has to be redesigned out of the human's hands.⁴

L2

DRIVER RETAINED · ATTENTION NOT ENGINEERED

the system works well most of the time — and then it does not

TESLA AUTOPILOT — LEVEL-2 MONITORING AS A SUSTAINABLE ROLE

The operational design ... permitted his prolonged disengagement from the driving task.

NTSB HIGHWAY ACCIDENT REPORT HAR-17/02 (WILLISTON, FLORIDA CRASH), 2017

REFERENCES

1. NTSB, Highway Accident Report HAR-17/02 (Williston, FL, 2016) — the quoted disengagement finding.
2. NTSB, Highway Accident Report HAR-20/01 (Mountain View, CA, 2018) — Autopilot crash analysis.
3. NHTSA Standing General Order 2021-01 reports — documented Autopilot fatal crashes.
4. NHTSA Office of Defects Investigation EA22-002 (opened 2022; closed April 2024, superseded by recall 23V-838) — driver-engagement adequacy.
5. Parasuraman, R. & Manzey, D. (2010) — automation complacency and monitoring.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Tesla Autopilot at consumer scale is the largest live test of Level-2 monitoring as a sustainable role. The early evidence is that it is not. The case is the strongest currently available test of whether consumer-

side training and engagement design can produce sustained attention to automation. The recurring fatality pattern suggests the answer.

LENS APPROACH

LENS uses Tesla Autopilot in LEN 2 as the live consumer-scale test of monitoring as a sustainable role and in LEN 7 for the governance dynamics of a Level-2 system marketed at the boundary of higher autonomy. Studio projects examine the driver-engagement design that would make the role performable.

DOMAIN	MODES AT STAKE				LENS COMPETENCY					PROBLEM TYPE	INDUCED	CLO		
AUTONOMOUS SYSTEMS	→	T	N	G	H	→	D1	D2	D3	D4	D5	6	3.2	3
Human-System Collaboration														

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Engineer driver-engagement monitoring to the demonstrated limits of human vigilance — verify attention against operational evidence before fielding, not against an assumption that warnings suffice. ▶ Name and present the feature so its capability boundary is unmistakable to an untrained consumer, rather than implying autonomy the system does not deliver. ▶ Constrain operation to the design domain the system can actually handle, so the human is not silently relied on as the backstop for edge cases.	▶ Monitor the standing-order crash data for the disengagement pattern and treat a recurring signature as evidence the role, not the driver, needs redesign. ▶ Track attention and takeover performance across years of ownership, since vigilance erodes with the very reliability that accumulates over time. ▶ Hold the engagement design accountable to an independent regulator with authority to require changes when in-use evidence shows it is inadequate.

REFLECTION QUESTIONS

1. Identify a passive-monitoring role in your domain. What evidence would tell you whether attention is sustainable over years of operation?
2. Design the driver-engagement architecture that would make Level-2 monitoring sustainable for an average consumer.
3. Autopilot assigns full legal responsibility to the operator while engineering the conditions that erode their attention. Where in your domain does liability rest with the person an automated system has made least able to intervene — and how would you realign the two?

WHO BUILDS THIS learning science & instructional design · safety science & organizational analysis · governance, ethics & measurement · human factors & interaction design — plus domain experts and a learning engineer to integrate. TOOLS scenario simulation · competency models · incident analysis · safety dashboards · governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs

Predictive Policing — PredPol

G Governance & Trust **H** Human–System Interface **D** Designed Out

IMPACT Predictive–policing tools deployed across scores of U.S. police jurisdictions; several cities have abandoned them after equity analyses

IN BRIEF

PredPol and similar predictive–policing tools use historical crime data to forecast where future crime is likely, directing patrols to those locations. Multiple analyses (Lum & Isaac 2016; Richardson, Schultz & Crawford 2019) found that because the historical data records where police have enforced, not where crime has occurred, the algorithm tends to reinforce existing enforcement patterns rather than predict underlying crime — a feedback loop that concentrates policing on already-over-policed neighborhoods. Cities including Santa Cruz, New Orleans, and Los Angeles have since suspended or abandoned predictive–policing deployments after equity reviews. The capability gap is at the construct definition: “where crime occurs” and “where crime is recorded” are different variables, and the tools treated them as one. It is the canonical algorithmic–governance case in U.S. policing.

THE SHIFT

Police departments have adopted data–driven tools that forecast where crime will occur and direct patrol resources accordingly. PredPol and its peers promised to make policing more objective by replacing officer intuition with statistical prediction — relocating a discretionary judgment into an algorithm trained on historical records, and lending the output a veneer of neutrality that the human judgments embedded in those records did not actually possess.^o

WHAT IS EMERGING

Researchers examining these systems found a structural flaw. Because the training data records where police have made arrests rather than where crime has actually occurred, the model learns enforcement patterns, not crime patterns. Lum & Isaac (2016) showed the result is a feedback loop: patrols are sent where police already went, generating more recorded incidents that confirm the prediction — a loop that grows more confident the longer it runs, because its own output becomes the next cycle’s evidence.¹

THE CAPABILITY QUESTION

The capability gap is at the construct definition. “Where crime occurs” and “where crime is recorded” are not the same variable, yet predictive–policing tools have treated them as interchangeable. The question is whether a model trained on a record of institutional behavior can ever predict the underlying phenomenon, or only amplify the behavior that produced its data — a question no amount of modeling accuracy can answer, because the gap is in what the data measures, not in how well the model fits it.²

EARLY EVIDENCE

Richardson, Schultz & Crawford (2019) documented “dirty data” — records produced during periods of biased or unlawful policing — feeding directly into predictive systems. After equity reviews, cities including Santa Cruz, New Orleans, and Los Angeles suspended or abandoned their predictive-policing deployments — abandonment that came only after the tools were already in service, the construct problem surfaced by external review rather than caught before the systems shaped where officers were sent.³

OPEN PROBLEMS

Predictive policing is the canonical algorithmic-governance case in U.S. policing and pairs with COMPAS (Case 187) and educational algorithmic bias (Case 46). The open problem is a construct-validity audit — a way to establish, before deployment, whether a predictive system’s training data is a record of ground truth or merely of institutional behavior — implemented in some jurisdictions and absent in most, so the same construct error remains available to the next department that mistakes a record of enforcement for a map of crime.⁴



Predictive policing systems learn from a record of past policing, not from a record of past crime.

PARAPHRASING LUM & ISAAC, SIGNIFICANCE, 2016

REFERENCES

1. Lum, K. & Isaac, W. (2016), “To Predict and Serve?” Significance — the enforcement-vs-crime feedback loop (paraphrased).
2. Richardson, Schultz & Crawford (2019), “Dirty Data, Bad Predictions” — biased records feeding predictive systems.
3. Brantingham et al. (2018) — predictive-policing field experiments.
4. Brayne, S. (2017), “Big Data Surveillance: The Case of Policing.”
5. Municipal records on suspension and abandonment of predictive policing (Santa Cruz, New Orleans, Los Angeles).

THE LEARNING ENGINEERING LENS

LE INSIGHT

Predictive policing is the canonical case for the difference between training data and ground truth at law-enforcement scale. The institution being predicted (where crime occurs) is not the institution producing the training data (where police make arrests). The capability gap is at the construct.

LENS APPROACH

LENS uses predictive policing in LEN 7 as a foundational AI- governance case in policing and in LEN 9 for the technical construct-validity analysis. The case is paired with COMPAS as criminal-justice algorithmic cases of different kinds.

DOMAIN	MODES AT STAKE	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
GOVERNMENT	→ G H D	→ D1 D2 D3 D4 D5	5	8.1	4
TECHNOLOGY		Test and Evaluation			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Establish a construct-validity audit before deployment that tests whether the training data records the phenomenon to be predicted or merely the institution's own behavior.
- ▶ Define the prediction target explicitly as "where crime occurs," not "where arrests are recorded," and reject data that cannot speak to the former.
- ▶ Engineer against the feedback loop — for example, decoupling patrol allocation from the data the model then re-ingests — before the system can amplify its own output.

IN OPERATION / AFTER

- ▶ Monitor for the enforcement feedback loop in deployment, watching whether predicted areas simply accumulate more recorded incidents that confirm the prediction.
- ▶ Require periodic equity review of in-use outcomes, since the construct flaw in these systems has surfaced through external analysis rather than internal metrics.
- ▶ Track the provenance of incoming training data and quarantine records produced during periods of biased or unlawful enforcement.

REFLECTION QUESTIONS

1. Identify a predictive system in your domain whose training data is itself a record of institutional behavior rather than ground truth. What is the construct gap?
2. Design the construct-validity audit that should precede deployment of a predictive system in any institutional setting.
3. Predictive policing's feedback loop grows more confident the longer it runs, because its own output becomes the next cycle's training data. Where in your domain does a deployed model shape the data it later learns from — and how would you break the loop?

WHO BUILDS THIS governance, ethics & measurement · human factors & interaction design · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · usability testing · cognitive walkthroughs · task analysis · design prototyping

Launching the BRAIN Initiative — Governance of a Grand Challenge

G Governance & Trust

K Knowledge & Institutional Memory

N Normalization of Deviance

DISCLOSURE Institutional overlap: an editor works in the connectomics field adjacent to this case and shares institutional ties to the BRAIN research community; no editor authored or contributed to this work. Anchored to widely cited primary position papers and independent journalism, not the editor's own work; the critical retrospective is the deliberate safeguard.

IMPACT A multi-billion-dollar national research endeavor launched via a position-paper-to-policy iteration sequence, with governance contestation on the public record and a 2021 critical retrospective documenting that the unified-understanding framing exceeded what the science delivered

IN BRIEF

The BRAIN Initiative is one of the few large-program launches in the corpus whose governance trail is publicly documented end-to-end: a 2011 Kavli symposium produced a six-author Neuron position paper (Alivisatos et al., 2012); the proposal was shepherded to the White House OSTP, became a Presidential initiative in 2013, and was operationalized by an NIH working group whose BRAIN 2025: A Scientific Vision report (2014; the group's plan later published as Jorgenson et al., 2015) set milestones and cost estimates. What makes it teachable rather than triumphalist is the governance contestation visible in the same record. Yuste worried the advisory panel was “packing the committee with users, rather than tool builders” — diluting the original focus; Bargmann, who later co-chaired the working group, had earlier expressed skepticism that the proposal “sounds like a big central planning project that will take resources away from creative work.” A 2021 MIT Technology Review retrospective assessed that the big-science brain projects, BRAIN included, did not deliver the unified understanding their framing promised. The case pairs with the EU Human Brain Project (Case 181) as a contrast in governance models — distributed working-group versus top-down single-PI — at the same era, the same ambition, and opposite trajectories.

BACKGROUND

In 2011 a small group of neuroscientists convened at a Kavli Foundation symposium argued that the field had a tractable grand-challenge target: recording the activity of large numbers of neurons across whole circuits. The argument was published in Neuron in 2012 as the Brain Activity Map proposal (Alivisatos, Chun, Church, Greenspan, Roukes, and Yuste) and was shepherded through Kavli's Miyoung Chun to the White House Office of Science and Technology Policy. The position paper named a capability the field could be organized

around and a sequence of tools that would have to be built; it was a launch artifact rather than a results report.⁰

THE INTERVENTION

In April 2013 President Obama announced the BRAIN Initiative. The operational governance was an NIH working group, co-chaired by Cori Bargmann and Bill Newsome, charged with translating the vision into a milestone-and-cost-bearing plan. The group released the BRAIN 2025: A Scientific Vision report in June 2014, later published in peer-reviewed form (Jorgenson et al., Phil. Trans. R. Soc. B, 2015); it named seven priority areas, set timelines, and laid out funding ranges. Among the few large research-program launches in the corpus, this is one whose position-paper-to-policy-to-implementation sequence is openly auditable — every step has a published artifact attached.¹

HOW IT WORKED

What makes the case teachable rather than triumphalist is the governance contestation visible in the same record. Yuste, one of the original six authors, voiced concern that the advisory panel was being expanded with users of the tools the program was meant to build, rather than the tool-builders the original proposal had centered — a documented dilution of scope. Bargmann, before her appointment, had publicly expressed skepticism that a big central planning project would draw resources away from creative independent research. The governance choices — who leads, tool-builders versus users, central plan versus distributed creativity — were not made in private and then defended; they were litigated in the public record while the program ran.²

THE EVIDENCE

The 2021 MIT Technology Review retrospective took the ten-year view: big-science brain projects, BRAIN among them, did not deliver the unified understanding their founding framing had promised. The honest assessment is that the initiative produced tools, atlases, and a coordinated funding stream — meaningful capability — while drifting from the grand-challenge framing the position paper had used to mobilize political support. The case is the instructive form of the “enthusiasm-evidence gap” at field scale: the framing carried the politics, and the science delivered something different and more diffuse.³

WHAT TRANSFERRED

What the case teaches is that large-program governance is the capability deliverable — not the framing, not the early tools, not the eventual results. The position-paper-to-policy sequence, the working-group composition decision, and the public airing of scope drift are the artifacts a future capability-development program can study. Paired with the EU Human Brain Project (Case 181), the case shows that the governance model — distributed working-group versus top-down single-PI — was the variable that explained why one program survived and adapted while the other unraveled. The framing, ambition, and era were comparable; the governance was not.⁴

Governance contestation in the public record is not program failure. It is what makes the case teachable.

EDITORS' SYNTHESIS OF THE BRAIN INITIATIVE RECORD (2012 – 2021).

REFERENCES

1. Alivisatos, Chun, Church, Greenspan, Roukes, & Yuste (2012), “The Brain Activity Map Project and the Challenge of Functional Connectomics,” *Neuron* 74(6):970–974, doi:10.1016/j.neuron.2012.06.006.

2. Jorgenson et al. (2015), “The BRAIN Initiative: developing technology to catalyse neuroscience discovery,” *Phil. Trans. R. Soc. B* 370(1668):20140164, doi:10.1098/rstb.2014.0164 — the BRAIN 2025 plan.

3. Yuste & Bargmann (2017), “Toward a Global BRAIN Initiative,” *Cell* 168(6):956–959, doi:10.1016/j.cell.2017.02.023.

4. Underwood (2013), “As White House Embraces BRAIN Initiative, Questions Linger,” *Science / ScienceInsider* (April 3, 2013) — source of the Yuste and Bargmann public-record contestation quotes.

5. MIT Technology Review (2021), retrospective on big-science brain projects — the critical ten-year assessment.

THE LEARNING ENGINEERING LENS

LE INSIGHT

The BRAIN Initiative is one of the few large-program launches in the corpus whose position-paper-to-policy-to-implementation sequence is fully auditable. Governance choices were litigated in the public record, scope drifted from the founding framing, and a critical ten-year retrospective is part of the case materials. Its teaching value is the contested record, not a clean success or scandal.

LENS APPROACH

BRAIN is the field-scale stakeholder-and-governance case in the v2 corpus (induced 5.1; LENS D5+D1/PT4). LENS uses it in Domain 1 (Systems Analysis) for the position-paper-to-policy iteration sequence; in Domain 5 (Navigating Sociotechnical Constraints) for the public-record governance contestation; and in Domain 4 (Test and Evaluation) for the enthusiasm-evidence gap as the framing exceeded delivered science. Direct pair with Case 181 (EU Human Brain Project), same era and ambition, opposite governance model. The COI disclosure under the title is binding: the editor’s research adjacency is what motivated the critical retrospective being included as a deliberate counterweight.

DOMAIN		MODES ADDRESSED		LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
NEUROSCIENCE	→	G K N	→	D1 D2 D3 D4 D5	4	5.1	1 · 5
SCIENCE POLICY				Navigating Sociotechnical Constraints + D1			

APPROACHES TO CONSIDER

DURING DEVELOPMENT

- ▶ Make the position-paper-to-policy iteration auditable by attaching a published artifact to every step — proposal, OSTP transmittal, working-group composition, milestone plan.
- ▶ Treat working-group composition as a governance act, not a staffing decision: tool-builders vs. users vs. integrators is a framing-shaping choice, and naming the trade-off openly is part of the work.
- ▶ Carry the founding scope through the operational record: name where the program is keeping faith with the founding framing and where it is drifting, while the program is running rather than only in retrospect.

IN OPERATION / AFTER

- ▶ Commission and publish a long-window retrospective that compares the founding framing to delivered capability honestly — not as a closure ritual but as governance evidence the next program can use.
- ▶ When scope drift is documented, decide explicitly whether to reframe the program publicly or to re-baseline against the original framing; the worst case is leaving the gap unaddressed in the record.
- ▶ Treat governance contestation in the public record as program documentation, not as program failure; the contested record is what makes the case teachable.

REFLECTION QUESTIONS

1. Identify a large research or capability-development program in your domain whose launch artifacts (position papers, working-group reports, milestone plans) are publicly auditable. What governance choices were made openly and which were made in private?
2. Specify the founding-framing-vs-delivered-capability comparison you would publish at year five and year ten of a program you are designing. What is the evidence form, who commissions it, and where does it live in the record?
3. BRAIN survived and adapted while the EU Human Brain Project (Case 181) unraveled. What is the minimum governance documentation that would let a future program-designer learn the difference, rather than reconstruct it from contested press accounts?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · safety science & organizational analysis — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · incident analysis · safety dashboards

CASE 200

AUTONOMOUS VEHICLES

PUBLIC-UTILITY GOVERNANCE

ACCESSIBILITY

2020 – 2024

CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver

G Governance & Trust **K** Knowledge & Institutional Memory **D** Designed Out

EVIDENCE TIER practice-synthesis-tier — source confidence flagged; future validation ongoing.

IMPACT California's Public Utilities Commission built an AV passenger-service permit framework whose conditions — time-of-day limits, weather restrictions, fleet caps, geographic carve-outs, mandatory Passenger Safety Plan for riders with disabilities — are explicitly designed to operationalize the governance objections that would otherwise block deployment outright

IN BRIEF

The California Public Utilities Commission established an AV passenger-service permit framework whose conditions are explicitly designed to address common governance objections — safety, equity, fleet scale — by writing them into the permit rather than treating them as binary deploy / don't-deploy questions. The framework includes time-of-day limits, weather restrictions, fleet caps, geographic carve-outs, and a required Passenger Safety Plan documenting how the driverless service handles riders with limited mobility, vision impairments, and other disabilities. A 2024 update strengthened the DMV's authority to impose targeted operational restrictions for safety, and the case shows the regime in operation: the Cruise suspension (Case 190) shows the regime can revoke; the Waymo continuation (Case 199) shows it can permit conditionally. The structural complement to the Waymo deployer-side safety case framework is exactly this regulator-side artifact — the permit conditions make the safety case operative as a regulatory instrument. The evidence-tier flag rendered under the title is practice-synthesis: the program is documented in CPUC decisions and program guidance, but no peer-reviewed evaluation of the program's equity-of-service outcomes for disabled and disadvantaged-community riders yet exists. Future validation will continue as ridership and incident data accumulate.

BACKGROUND

The governance question California faced from 2018 onward was how to regulate AV passenger service in a regime where the California DMV regulates the vehicle and the CPUC regulates the passenger service the vehicle provides. Treating each AV deployment as a binary deploy / don't-deploy decision risked either blocking deployment outright on objections the operator could not satisfy, or permitting deployment without a governance

handle on the objections the public and the disability-rights community brought to the proceedings.⁹

THE INTERVENTION

The CPUC's structural answer was to build the objections into the permit conditions. The AV passenger-service permits specify time-of-day limits, weather restrictions (fog, heavy rain), fleet caps, geographic carve-outs (initially excluding certain neighborhoods), and a mandatory Passenger Safety Plan documenting how the driverless service will accommodate riders with limited mobility, vision impairments, hearing impairments, and other disabilities. The conditions are not policy aspirations — they are permit terms whose violation triggers enforcement.¹

HOW IT WORKED

In August 2024 the DMV released draft regulations for public comment that would strengthen its authority to impose targeted operational restrictions for safety reasons during the deployment lifecycle, not only at initial permit issuance. The regulatory regime can therefore tighten the conditions in response to operational evidence. The pattern: rather than treat the governance objection as binary, the CPUC made the objection itself a design parameter the deployer must satisfy and that the regulator can re-tune.²

THE EVIDENCE

The case is the structural complement to Case 199 (Waymo) and the inverse-outcome companion of Case 190 (Cruise). The Waymo safety case framework is the deployer-side artifact; the CPUC permit conditions are the regulator-side artifact that makes the safety case operative as a regulatory instrument. The Cruise suspension shows the regime can revoke when the disclosure architecture fails; the Waymo conditional continuation shows it can permit when the conditions are satisfied. Together the trio teaches the structural form of regulator-deployer interaction at the AV physical-safety C5 layer.³

WHAT TRANSFERRED

The evidence-tier flag rendered under the title is practice-synthesis. The CPUC decisions, the program guidance, and the permit conditions themselves are public; no peer-reviewed evaluation of the equity-of-service goals (disabled-rider access, disadvantaged-community access) has yet measured whether the Passenger Safety Plan conditions translate into measured ridership outcomes. The condition-as-objection-dissolver pattern is teachable and the regulatory architecture is explicit, but the outcome evidence is not yet decision-grade. Future validation will continue as the program ages and as the equity outcomes are independently measured.

The governance objection was not refused. It was made a permit condition.

EDITORS' SYNTHESIS OF THE CPUC AV PASSENGER-SERVICE PERMIT FRAMEWORK.

REFERENCES

1. California Public Utilities Commission, "Autonomous Vehicle Passenger Service Programs," CPUC program page and August 2024 application guidance.
2. CPUC permit decisions for Cruise and Waymo, 2020–2024.
3. California Department of Motor Vehicles, AV regulatory program — strengthened safety-restriction authority, 2024.
4. Paired Cases 199 (Waymo deployer-side artifact) and 190 (Cruise revocation under regime).

THE LEARNING ENGINEERING LENS

LE INSIGHT

The CPUC permit framework is the regulator-side counterpart to the Waymo safety case (Case 199): conditions operationalize objections rather than blocking deployment. Evidence-tier flag is practice-synthesis; the regulatory architecture is documented in CPUC decisions, but no peer-reviewed evaluation of the equity-of-service goals yet exists, and future validation is ongoing.

LENS APPROACH

CPUC is the AV regulator-side governance case (induced 5.1; LENS D5/PT6). LENS uses it in Domain 5 (Navigating Sociotechnical Constraints) for cross-regime governance seams (DMV regulates the vehicle; CPUC regulates the service; both can act) and for the condition-as-design-parameter pattern that makes the deployer-side safety case (Case 199) operative as a regulatory instrument. Pairs with Case 190 as the revocation event under the regime.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
AUTONOMOUS VEHICLES PUBLIC-UTILITY GOVERNANCE ACCESSIBILITY	G K D	D1 D2 D3 D4 D5	6	5.1	5 · 3
Navigating Sociotechnical Constraints					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Treat governance objections as design parameters rather than as binary blockers — write them into the permit conditions where compliance is verifiable and violation triggers enforcement. ▶ Require a documented Passenger Safety Plan (or equivalent equity-of-service artifact) at initial permit issuance so the equity objection has an operational answer the regulator can audit. ▶ Build the regulatory authority to re-tune the conditions during the deployment lifecycle into the regime itself; the Cruise suspension (Case 190) is what re-tuning under a failure event looks like.	▶ Pair the case with Case 199 (Waymo) as the regulator-deployer artifact dyad — the safety case framework and the permit conditions are co-constitutive, not independent moves. ▶ Use the case in LENS Domain 5 modules on cross-regime governance seams (CPUC and DMV) — the regime structure is itself part of the capability question. ▶ Carry the practice-synthesis flag honestly: the condition-design pattern is teachable, but the equity-of-service outcome data are not yet peer-reviewed; future validation requires independent evaluation of the Passenger Safety Plan deliverables.

REFLECTION QUESTIONS

1. Identify a deployment regime in your context where governance objections are currently treated as binary deploy / don't-deploy questions. How could the objections be rewritten as permit-style conditions whose compliance is verifiable and whose violation triggers enforcement?
2. Specify the equity-of-service artifact your regime would require at permit issuance — the analogue of the Passenger Safety Plan — and the auditable evidence the deployer must provide that the artifact is operational.
3. The case is practice-synthesis tier. What independent outcome evidence — disabled-rider ridership figures, disadvantaged-community access measurements, comparative incident rates by permit condition — would you require before treating the condition-as-objection-dissolver pattern as a validated regulatory architecture?

WHO BUILDS THIS governance, ethics & measurement · knowledge management & data engineering · systems & human-factors engineering — plus domain experts and a learning engineer to integrate. **TOOLS** governance frameworks · bias & evidence audits · knowledge bases · data pipelines · task analysis · design prototyping

Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India

G Governance & Trust **N** Normalization of Deviance **H** Human-System Interface

EVIDENCE TIER journalism-tier — source confidence flagged; future validation ongoing.

IMPACT India's Aadhaar biometric-identity system, gating welfare for roughly one billion people, produced operational exclusion when authentication failed at the ration shop or pension window; the 2018 Puttaswamy II majority upheld the Section 7 welfare-linkage while Justice Chandrachud's dissent held that exclusion from technological error — through no fault of the individual — violates constitutional dignity, the reasoning behind executive fallback mandates issued after documented starvation deaths

IN BRIEF

Aadhaar — India's biometric digital-identity system, with roughly one billion enrolled — was designed to streamline welfare delivery and reduce identity fraud. Across more than a decade of deployment the load-bearing failure mode became operational exclusion: when biometric authentication failed for an individual (worn fingerprints, missing iris match, connectivity loss at the point of service), the welfare to which the person was entitled became inaccessible. The constitutional reckoning came in the 2018 Puttaswamy II Aadhaar judgment: the 4–1 majority upheld Section 7 of the Aadhaar Act, which conditions subsidies and benefits on authentication, while Justice D. Y. Chandrachud's lone dissent named the load-bearing principle — that “exclusion based on technological errors, with no fault of the individual, is a violation of dignity.” The empirical predicate is the documented exclusion record: the 2017 Jharkhand starvation deaths (Santoshi Kumari among at least nineteen hunger deaths the Right to Food Campaign tied to ration-card cancellation or Aadhaar-authentication failure) and the Drèze-Khera-Somanchi field study finding roughly a fifth of cardholders unable to transact in an average month. The correction that followed was executive and doctrinal rather than a majority holding — circulars barring denial of rations for want of Aadhaar, and, in the distinct 2025 Pragya Prasun ruling, a right to accessible digital access read into Article 21 in the e-KYC context. The case is the non-US automated-welfare-delegation case the corpus needs alongside SyRI and the UK Post Office Horizon thread. The evidence tier is mixed: the judgments are investigation-grade; the lived-exclusion sourcing rests on journalism and advocacy reporting, with the field data peer-reviewed. Future validation continues on whether the fallback requirement is honored at the operator interface. Gap-5 non-US/UK/EU case (India).

BACKGROUND

Aadhaar was designed as a population-scale solution to two problems the Indian welfare state had identified by the late 2000s: fragmented identification across welfare schemes and identity fraud in benefits delivery. The design choice was biometric authentication at the point of service — fingerprints and iris — keyed to a unique twelve-digit identifier issued by the UIDAI. Enrollment reached approximately one billion. The capability the system promised was streamlined welfare access; the authority it required was delegation of identity verification to the biometric.⁰

THE INTERVENTION

The load-bearing failure mode emerged across deployment: operational exclusion. When biometric authentication failed at the ration-shop card-reader or the pension window — worn fingerprints from manual labor, missing iris matches, connectivity loss, device malfunction — the welfare to which the person was entitled became inaccessible. The failure was not the absence of identity; the person had the Aadhaar number. The failure was the brittleness of the authentication system at the operator interface. The people most dependent on the welfare were the people whose biometrics were most likely to fail.¹

HOW IT WORKED

The constitutional reckoning came in stages. In the 2018 Puttaswamy II judgment the Supreme Court's 4–1 majority upheld Aadhaar and Section 7 of the Aadhaar Act — the provision conditioning subsidies and benefits on authentication — reasoning that targeted delivery and leakage-reduction were legitimate, proportionate state aims. Justice D. Y. Chandrachud, dissenting alone, named the structural risk the majority did not: that Section 7 was over-broad and that “exclusion based on technological errors, with no fault of the individual, is a violation of dignity.” Across the following years the documented exclusion record accumulated, and the correction arrived not as a reversal of the majority but as executive circulars barring denial of rations for want of Aadhaar and, in the distinct 2025 Pragya Prasun ruling, a right to accessible digital access read into Article 21 in the e-KYC context for persons with disabilities.²

THE EVIDENCE

The evidence-tier flag is split and load-bearing. The court judgments themselves are investigation-grade — the published opinions are the record. The empirical body rests on two tiers: the lived-exclusion account is journalism and advocacy reporting (the Right to Food Campaign's compilation of the Jharkhand starvation deaths; the Indian press), while the exclusion-rate evidence is peer-reviewed field work — Drèze, Khalid, Khera and Somanchi's Jharkhand survey (“Pain without Gain?”, EPW 2017) found roughly twenty percent of cardholders unable to transact in an average month, with over forty percent of non-transacting households citing Aadhaar-authentication problems. The journalism-tier flag travels with the lived-exclusion material; the field data and the judgments carry their own, higher, tiers. Whether the fallback requirement is honored at the operator interface is the open empirical matter the future will validate.³

WHAT TRANSFERRED

What the case adds at the LENS layer is the non-US automated-welfare-delegation thread the corpus needs. The Domain-5 frame applies — a sociotechnical constraint (population-

scale biometric delegation) tested against constitutional dignity — and the CLO around fairness beyond omission applies directly: the harm was not the absence of the system but the operational exclusion produced by it. The CLO on delegation with revocation applies too: the correction did not revoke the delegation but sought to bound it with a fallback requirement — and the gap between the Chandrachud dissent's principle and the majority's holding is precisely why the revocation remains contested rather than settled. The case pairs with SyRI (Dutch welfare-fraud system halted on rights grounds) and the UK Post Office Horizon thread as the global lineage of judicial reckoning with automated welfare delegation.⁴

Exclusion based on technological errors, with no fault of the individual, is a violation of dignity.

JUSTICE D. Y. CHANDRACHUD, DISSENTING, PUTTASWAMY II (2018)

REFERENCES

1. Supreme Court of India (2018), Justice K. S. Puttaswamy (Retd.) v. Union of India, (2019) 1 SCC 1 — the Aadhaar judgment; 4–1 majority upholding Section 7, with Justice D. Y. Chandrachud's dissent on technological exclusion and dignity.
2. Drèze, J., Khalid, N., Khera, R., & Somanchi, A. (2017), "Aadhaar and Food Security in Jharkhand: Pain without Gain?" Economic & Political Weekly 52(50) — field study of biometric-authentication exclusion in the PDS.
3. Right to Food Campaign compilation and Indian journalism (The Hindu, The Wire, Scroll.in, 2017–2019) on the Jharkhand starvation deaths (Santoshi Kumari, September 2017) — lived-exclusion sourcing with journalism-tier flag.
4. Supreme Court of India (2025), Pragya Prasun & Ors. v. Union of India, 2025 INSC 599 — the distinct, later ruling reading a right to accessible digital access into Article 21 in the e-KYC context for persons with disabilities.
5. "A Failure to Do No Harm," PMC5741784 — peer-reviewed comparative analysis of biometric-ID exclusion in welfare delivery.

THE LEARNING ENGINEERING LENS

LE INSIGHT

Aadhaar exclusion is the non-US automated-welfare-delegation case the corpus needed: an operational-exclusion failure mode whose constitutional statement is Justice Chandrachud's 2018 dissent — exclusion from technological error, no fault of the individual, as a dignity violation — set against a majority that upheld the Section 7 linkage. The correction that followed was executive and doctrinal rather than a majority holding. Evidence tier is split: judgments investigation-grade, field data peer-reviewed, lived-exclusion sourcing journalism plus advocacy. Future validation ongoing on whether the fallback requirement is honored.

LENS APPROACH

The designed teaching point is operational failure-mode attribution (induced 5.2; canonical competency 8.1; LENS D5/PT5). At the ration-shop card-reader and the pension window, a biometric authentication that did not match was logged as user error — the claimant's worn fingerprints, the elderly applicant's failed iris read — when it was in fact a designed exclusion mode of the system: manual laborers and the elderly are precisely the populations whose biometrics the design could not reliably read, so the failure belongs to the authentication design, not to the excluded person. The capability discipline is to attribute an authentication failure to the system that produced it rather than to the user it shut out, and to name who the design predictably excludes before deployment. Justice Chandrachud's dissent performed exactly this attribution — naming technological exclusion, through no fault of the individual, as the system's dignity violation — even as the majority upheld the linkage; the correction has since been pursued through executive fallback mandates rather than a settled revocation. Pair with SyRI (Case 189) and the UK Post Office Horizon thread.

DOMAIN	MODES ADDRESSED	LENS COMPETENCY	PROBLEM TYPE	INDUCED	CLO
DIGITAL IDENTITY	G N H	D1 D2 D3 D4 D5	5	5.2	1·5
WELFARE	→	→	Navigating Sociotechnical Constraints		
BIOMETRIC AUTHENTICATION					
INDIA					

APPROACHES TO CONSIDER

DURING DEVELOPMENT	IN OPERATION / AFTER
▶ Treat the operator–interface brittleness of any biometric system as a design variable from the start: who is most likely to be excluded, and what is the alternative at the point of service. ▶ Specify the fairness–beyond–omission test the system must pass before deployment: not “is the system available,” but “does it work for the people most dependent on the welfare it gates.” ▶ Document the delegation explicitly — what authority is delegated to the biometric, under what conditions it can fail, and what the operator’s authority to override is — so the revocation question is answerable later.	▶ Track operational exclusion as a primary outcome metric, not a downstream complication; the case teaches that the exclusion is the failure mode the design ignored. ▶ Carry the journalism–tier flag on the lived–exclusion material and the investigation–grade flag on the court judgments separately; the case is teachable because the two tiers are explicit. ▶ Pair with SyRI and the UK Post Office Horizon thread when teaching automated–welfare delegation; the cross–jurisdictional pattern is the lineage of judicial correction of these systems.

REFLECTION QUESTIONS

1. Identify a delegated authentication or eligibility–check system in your context whose failure mode is operational exclusion of the people most dependent on the service it gates. What is the alternative at the point of service, and is it actually available?
2. Specify the fairness–beyond–omission test such a system would have to pass before deployment. The Aadhaar pattern teaches that the harm is not the absence of the system but the exclusion the system produces — what would the equivalent test be in your domain?
3. The case has split evidence tiers: court judgments investigation–grade, lived–exclusion sourcing journalism plus advocacy plus comparative analysis. How would you carry the tier split into a teaching artifact without smoothing either tier away?

WHO BUILDS THIS governance, ethics & measurement · safety science & organizational analysis · human factors & interaction design — plus domain experts and a learning engineer to integrate. TOOLS governance frameworks · bias & evidence audits · incident analysis · safety dashboards · usability testing · cognitive walkthroughs

DOMAIN INDEX

The cases by domain.

The matrix at the front lists every case in number order. This index reorganises the cases by primary domain — a reader following healthcare, defense, education, or any other thread can find the relevant cases together. Each case appears once, under its primary domain. Where a section opens with a callout, it names the dataset’s standout failure and, where one exists in that domain, its standout success.

DEFENSE

STANDOUT FAILURE

Case 5

Operation Eagle Claw (1980). Cross-organizational capability that existed on no service’s deliverable list; the mission was improvised across the Army, Navy, Marines, Air Force, and CIA. Produced the Holloway Commission, USSOCOM, and Goldwater-Nichols.

STANDOUT SUCCESS

Case 18

U.S. Nuclear Navy / Rickover Training Model (1954–). The longest continuous capability-engineering program in any high-consequence domain. Sixty-plus years of zero reactor accidents through training treated as a system parameter.

124	USS Fitzgerald & USS John S. McCain	2017	T·K·N
126	F-35 Sustainment & Maintainer Shortage	ongoing	T·K·D
127	Military Fratricide — Desert Storm to Afghanistan	1991 – 2014	T·H·K
128	Operation Eagle Claw	1980	T·K
129	Patriot Missile / Dhahran	1991	D·H·K
130	USS Vincennes & Iran Air Flight 655	1988	H·T
131	Stanislav Petrov / 1983 False Alert	1983	H·T
133	Mark 14 Torpedo Failures	1941 – 1943	T·K·G
134	V-22 Osprey	1991 – present	T·H·N
135	9/11 Intelligence Sharing Failures	1996 – 2001	G·K

137	U.S. Nuclear Navy / Rickover Training Model	1954 – present	T-K-N
138	MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard revision (series since 1968)	2020	G-K
139	GIFT and the Adoption Gap	2012 – present	K-G-N
140	Navy Surface Warfare Readiness Reform	2018 – present	T-K-N
141	DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks – 2014	2009	H-K-D

AVIATION

STANDOUT FAILURE

Case 7

Boeing 737 MAX / MCAS (2018–19). A single-string angle-of-attack sensor, a flight-control law that assumed pilot mental models from prior models, and a certification chain that did not knit the pieces together. 346 lives.

STANDOUT SUCCESS

Case 70

Crew Resource Management & CAST (1981–). First-officer authority engineered into the cockpit, a continuous learning architecture across the industry, and the safety record that followed.

96	AeroPerú Flight 603	1996	H-T
97	Boeing 737 MAX / MCAS	2018 – 2019	D-T-H
102	Air France Flight 447	2009	T-H
103	Kegworth / British Midland 92	1989	T-H-K
104	Asiana Airlines Flight 214	2013	T-H
105	Helios Airways Flight 522	2005	T-H
106	TransAsia Airways Flight 235	2015	T-H
107	Eastern Air Lines Flight 401	1972	H-T
109	Colgan Air Flight 3407	2009	T
110	Atlas Air Flight 3591	2019	T-K
112	NASA Saturn V Documentation	1972 – present	K
117	Crew Resource Management & CAST	1981 – present	T-H-N
118	Korean Air Safety Transformation	2000 – present	T-N
119	Aviation Safety Reporting System (ASRS)	1976 – present	T-K-N

120	EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation	1996 – 2002	H·K·G
121	TCAS — Coordinated Collision Avoidance and the Überlingen Lesson (TCAS II Version 7.1)	1981 – 2008	H·K·G
122	Singapore Airlines Safety Transformation	1980s – present	T·N
132	F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap	2008 – 2012	H·K·N
185	Air Canada Chatbot Liability — Delegation Without Revocation	2022 – 2024	D·K·N

HEALTHCARE

STANDOUT FAILURE

Case 61

EHR / CPOE Implementation (2005–). Roughly \$30B federal investment in systems designed for billing and administration, deployed against clinical workflow. Usability is now itself a patient-safety variable at scale.

STANDOUT SUCCESS

Case 109

WHO Surgical Safety Checklist (2008–). One page, nineteen items, paired with a required pause. Halved surgical mortality in the multi-site pilot; the smallest effective capability intervention in the dataset.

1	Therac-25	1985 – 1987	H·D·G
2	EHR / CPOE Implementation	2005 – present	H·D·G
4	Mid Staffordshire NHS Foundation Trust	2005 – 2009	G·N·K
5	Epic Sepsis Model	2017 – 2021	D·G·N
8	Medical Errors as Systemic Failure	1999 – present	T·H·N·K·G
9	Vioxx Withdrawal	1999 – 2004	G·D
11	California Nurse Staffing Ratios — Legislating a Capability Requirement	2004 – 2010	G·K
12	ACGME 80-Hour Resident Duty-Hour Reform	2003–2017	T·K·N
13	Implementation Science in Healthcare — The 17-Year Gap	ongoing	K·G·N
19	Keystone ICU / Pronovost Checklist	2004 – present	T·N
20	TREWS / Sepsis Watch	2018 – 2022	H·K·G
21	Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Hand-off	2024 – 2025	H·K·N
22	ChatGPT in Healthcare — Hallucination Cases	2023 – present	H·D

23	WHO Surgical Safety Checklist	2008 – present	T·N
24	Bristol Heart Babies Reform	1984 – present	G·N
32	Annual-Screening UI Redesign + CDS at University of Missouri Health Care	2020	T·D·N
33	Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal	2019 – 2024	T·D·N
34	COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case	2022 – 2024	T·K·D
35	Radiology AI Miscalibration	2018 – present	H·K·D
36	AlphaFold — Protein Structure Prediction	2020 – present	T
37	DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback	2020	T·K·D
39	TeamSTEPPS	2006 – present	T·N
176	Tylenol Recall	1982	G·N

EDUCATION AT SCALE

STANDOUT
FAILURE

Case 149

Summit Learning / Personalized Learning Rollout (2014–19). A defensible pedagogical platform, well funded, deployed across 380 schools — and pulled by parent revolt because the governance architecture (consent, evidence, exit) was never engineered alongside it.

STANDOUT
SUCCESS

Case 110

Georgia State University Predictive Analytics (2012–). 800 risk factors per student, advising triggered by early warning signs, equity as a primary constraint. Six-year graduation 32% → 54%; equity gaps eliminated.

45	Tennessee Voluntary Pre-K Study	2009–2018	G·D
46	Algorithmic Bias in Educational Predictive Analytics	ongoing	G·H·D
51	Atlanta Public Schools Cheating Scandal	2009 – 2015	G·N
53	inBloom	2014	G
54	Summit Learning / Personalized Learning Rollout	2014–2019	G·T·K
67	Cognitive Tutor / Carnegie Learning	1990s – present	T
80	Georgia State University Predictive Analytics	2012 – present	T·K
142	xAPI / Total Learning Architecture — Interoperability Gap	ongoing	K·G

205 The Discipline We Build Next ongoing

T·K·N·G·H·D

ENERGY

STANDOUT
FAILURE**Case 69**

Three Mile Island (1979). Training matched the design-basis events the engineers had imagined; the accident was not one of them. The control-room interface confused command signal with valve state. Catalysed industry-wide reform.

STANDOUT
SUCCESS**Case 172**

INPO and the Nuclear Academy (1979–). The industry built the institution TMI revealed it needed. No INES-level event at U.S. commercial reactors since.

146 Northeast Blackout 2003

H·K

160 Davis-Besse Reactor Head Corrosion 2002

N·K·G

161 Texas City BP Refinery Explosion 2005

N·T·K·G

162 Upper Big Branch Mine Explosion 2010

N·G·K

163 Deepwater Horizon 2010

T·N·K

165 Camp Fire / PG&E 2018

G·N·K

166 Three Mile Island 1979

T·H

169 Fukushima Daiichi 2011

N·G·K

175 INPO and the Nuclear Academy 1979 – present

T·K·G

GOVERNMENT

STANDOUT
FAILURE**Case 162**

Australia Robodebt (2016–20). An income-averaging algorithm that reversed the burden of proof onto welfare recipients. The Royal Commission found “venality, incompetence and cowardice”; the scheme was linked to deaths, including suicide.

NOTE

Case 111

No paired-intervention success in the government dataset. The reforms that follow these cases — Royal Commissions, statutory restructures — are documented in the narratives but did not produce a LENS-style success case in time for this volume.

7 VA Wait-Time Scandal 2014

G·K·N

180 Healthcare.gov Launch 2013

K·T·G

193 Bernard Madoff / SEC Failures 1992 – 2008

G-K-N

197 Predictive Policing — PredPol 2011 – present

G-H-D

INDUSTRIAL

STANDOUT FAILURE

Case 147

Bhopal (1984). Training, knowledge, normalisation of deviance, and governance all collapsed in the same plant. Roughly 4,000 immediate deaths; chronic harms continue. The cautionary case for outsourced operational responsibility.

STANDOUT SUCCESS

Case 73

Toyota Production System / Andon Cord (1950s–). Any operator can stop the line. The smallest authority-architecture intervention in the dataset, and the most durably copied.

144 Kodak Digital Camera Stagnation 1975–2012

D-K-N

147 Takata Airbag Inflators 2008 – 2023

D-G

149 Volkswagen Dieseldiegate 2015

D-G

151 GM Ignition Switch 2002 – 2014

D-G

155 Toyota Production System / Andon Cord 1950s – present

N-G

159 Bhopal 1984

T-K-N-G

164 Grenfell Tower 2017

G-T-K-N

TECHNOLOGY

STANDOUT FAILURE

Case 66

UK Post Office Horizon Scandal (1999–2015). Institutional refusal to believe sub-postmasters' reports of computer error, sustained over fifteen years. Hundreds of wrongful convictions; ruined lives and at least four suicides — the longest-running operator-feed-back failure in the dataset.

NOTE

Case 132

The technology dataset's success story is filed under Healthcare: AlphaFold (Case 36), the contemporary worked example of computational capability-engineering done as a discipline.

143 Knight Capital Trading Loss 2012

D-K

148 Wells Fargo Fake Accounts 2011 – 2016

G-N

152 TSB Bank IT Migration 2018

H-G

153	Equifax Data Breach	2017	G·K
157	AI-Augmented Coding Tools	2021 – present	T·H
172	CrowdStrike Falcon Outage	2024	D·K·G
184	UK Post Office Horizon Scandal	1999 – 2015	G·H·K

AI IN EDUCATION

88	LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models	2025	T·K·N
----	--	------	-------

K–12 EDUCATION

48	School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism	2010s – 2022	G·K·N
60	Houston EVAAS — Value-Added Teacher Evaluation on Trial	2011–2017	G
61	Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap	1997–2023	K·T
62	Gates Intensive Partnerships — Measurement Without a Coupled Lever	2009–2018	G·K
63	LAUSD iPad Program — The Common Core Technology Project	2013–2015	D·H·G
66	Growth-Mindset National Experiment — Heterogeneous Effects	2019	D·N·K
75	Chen et al. — Rural China AI Devices and the Equity-Direction Finding	2025	T·K·D
95	PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop	1997–2024	T·G

K–12 MATHEMATICS

72	ASSISTments — National Replication and Long-Term Follow-Through	2014 – present	T·K·D
84	Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive	2007 – 2014	T·K·D

K–12 SCIENCE

74	Zhang/Scardamalia — Knowledge Building Across Three Cohorts	2009	T·K·D
----	---	------	-------

AIR TRAFFIC MANAGEMENT

116	Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change	2005	G·K·H
-----	---	------	-------

ANESTHESIOLOGY

- 27** Anesthesia Monitoring Standards and the APSF — The Mortality Decline 1986 – present **H·K·G**

AUTONOMOUS SYSTEMS

STANDOUT
FAILURE

Case 62

Uber ATG / Tempe Fatality (2018). The safety driver was nominally responsible for what the system was designed to do without her attention — the classic vigilance contradiction at the heart of partial-autonomy operator roles.

NOTE

Case 194

No paired-intervention success in the autonomous-systems dataset. The capability engineering this category needs is still being built — see Case 205 (**The Discipline We Build Next**) for the open question this volume closes on.

- 183** Uber ATG / Tempe Fatality 2018 **T·N·G·H**

- 195** Tesla Autopilot — Recurring Fatalities 2016 – present **T·N·G·H**

AUTONOMOUS VEHICLES

- 190** Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment 2023 **G·K·N**

- 199** Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact 2023 / 2025 **G·K·N**

- 200** CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver 2020 – 2024 **G·K·D**

AVIATION INFRASTRUCTURE

- 115** FAA NextGen and the ADS-B Out Transition 2003 – 2020 **G·D·K**

AVIATION SAFETY

- 114** FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule 1988 – 2010s **G·D·K**

- 123** FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike 1992 – 2000s **G·K·N**

CLINICAL CARE

- 15** I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer 2014 **H·K·N**

- 29** Bar-Code Medication Administration — Cue at the Point of Care 2010 **H·K·D**

CLINICAL MEDICINE

6 Racial Bias in Pain Assessment — The False-Belief Mechanism 2016 **T·K·N**

25 Removing the Race Coefficient from eGFR 2021 **D·G·N**

CLINICAL/TRANSLATIONAL RESEARCH

40 Team Science Training for Clinical and Translational Scientists 2020 – 2022 **K·N**

CORPORATE L&D

65 Training Transfer — The Gap Between Learning and Doing 2010 **K·N**

70 High-Impact Learning System — Engineering the Environment, Not Just the Event 2001 – present **K·N**

79 The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops 1959 – present **K·N**

83 Brinkerhoff Success Case Method — Tails as the Evaluation Instrument 2005 – present **K·N**

CRIMINAL JUSTICE

187 COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate 2014 – 2018 **D·K·N**

DIGITAL IDENTITY

201 Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India 2018 – 2025 **G·N·H**

E-GOVERNMENT

194 Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox) 2001 – present **G·K·N**

ED-TECH

47 Algorithmic Bias in Automated Exam Proctoring 2022 **D·N·K**

EDUCATION (NORWAY)

92 Norway's National Expert Commission on Learning Analytics 2021 – 2023 **G·K·N**

EDUCATION AT SCALE

50 Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction 2012 – 2024 **D·K·N**

69 Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale 2016 **T·D**

EMERGENCY MANAGEMENT

-
- 167** Hurricane Katrina — The FEMA/DHS Response Failure 2005 **G-K**
- 171** Hurricane Maria — Puerto Rico Response and Logistics Failure 2017 – 2018 **G-N**

FINANCE

-
- 186** Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity 2018 – 2022 **D-G-N**
- 196** Fintech Lending Fairness Audit — When Including Race Reduces Disparity 2025 **D-G-N**

FINANCIAL SERVICES

-
- 192** Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model 2019 – 2021 **D-K-N**

GENERAL AVIATION

-
- 100** Glass-Cockpit Transition in Light Aircraft — Technology Outran Training 2010 **H-N-K**

GLOBAL HEALTH

-
- 18** PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption 2023 **K-N-D**
- 43** Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface 2013 – 2018 **H-N**
- 170** West Africa Ebola — The Delayed International Response 2014 – 2016 **G-N**

GOVERNMENT

-
- 49** UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days 2020 **D-K-N**
- 191** Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict 2016 – 2023 **D-G-N**
- 203** NYC Local Law 144 — Bias Audits as Governance Artifact 2023 – present **D-G-K**

GOVERNMENT/WELFARE (NETHERLANDS)

-
- 189** Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds 2014 – 2020 **D-G-N**

HEALTH PROFESSIONS EDUCATION

-
- 28** Interprofessional Education and the Evidence Gap 2013 – 2015 **K-N**

HEALTHCARE/ONCOLOGY

-
- 3** Watson for Oncology — Delegation Marketed Ahead of Capability 2013 – 2018 **D·K·N**

HIGHER EDUCATION

-
- 55** Enrollment-Algorithm Yield Optimization Across U.S. Higher Education 2010s – present (Brookings synthesis 2021) **T·K·N**
-
- 57** GAO Online Program Manager Oversight Gap (GAO-22-104463) 2022 **D·K·N**
-
- 58** USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC) 2010s – 2024 **D·K·N**
-
- 85** OU Analyse — Predictive Learning Analytics and Teacher Use at Scale 2019 – 2023 **T·K·D**
-
- 86** Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction 2024 **D·K·N**
-
- 90** Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions 2025 **T·K·N**

HIGHER EDUCATION (AFRICA)

-
- 89** Data Privacy and Learning Analytics on the African Continent 2022 **K·G·N**

HIGHER EDUCATION (BRAZIL)

-
- 82** MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil) 2024 **K·N**

HIGHER EDUCATION (LATIN AMERICA)

-
- 91** LALA — Building Learning-Analytics Governance Capacity Across Latin America 2017 – 2020 **K·N**

HIGHER ED (UK)

-
- 81** Open University 'Ethical Use of Student Data' and OU Analyse 2014 – 2025 **G·K·N**

HIGHER-ED ANALYTICS

-
- 52** Purdue Course Signals — The Reverse-Causality Retention Claim 2012 – 2013 **D·K·N**

HOSPITAL PHARMACY

-
- 42** Australian Hospital-Pharmacy Technician Role Redesign 2016 **D·N·H**

HUMANITARIAN RELIEF

-
- | | | |
|------------|---|----------|
| 178 | The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami 2005 – 2020 | G |
|------------|---|----------|
-

IMPLEMENTATION SCIENCE

-
- | | | |
|-----------|--|------------|
| 41 | Implementation-Science Training — Stated Goals Outrunning Operational Practice 2020s | K·N |
|-----------|--|------------|
-

LEARNING ANALYTICS

-
- | | | |
|-----------|---|--------------|
| 73 | The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension 2016 – 2025 | T·K·D |
|-----------|---|--------------|
-
- | | | |
|-----------|---|--------------|
| 87 | Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models 2021 – 2024 | D·K·N |
|-----------|---|--------------|
-
- | | | |
|-----------|--|------------|
| 94 | Learning Analytics on the African Continent — A Scoping Review as the Present-State Map 2022 | K·N |
|-----------|--|------------|
-

MEDICAL DEVICES

-
- | | | |
|-----------|---|--------------|
| 26 | Pulse Oximetry Across Skin Tones 1990 – 2020 – 2025 | D·G·N |
|-----------|---|--------------|
-

MEDICAL EDUCATION

-
- | | | |
|-----------|--|------------|
| 14 | Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA 2009 – 2020s | T·K |
|-----------|--|------------|
-
- | | | |
|-----------|---|--------------|
| 17 | Spaced Education RCTs in Medical Training 2007 – 2009 | H·K·D |
|-----------|---|--------------|
-

MEDICAL-DEVICE REGULATION

-
- | | | |
|-----------|---|------------|
| 44 | Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD 2014 – present | G·N |
|-----------|---|------------|
-

MULTIMODAL LEARNING ANALYTICS

-
- | | | |
|-----------|---|--------------|
| 76 | Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account 2023 | T·K·N |
|-----------|---|--------------|
-

NAVAL ENGINEERING

-
- | | | |
|------------|--|--------------|
| 173 | Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable 1963 – present | G·K·H |
|------------|--|--------------|
-

NEUROSCIENCE

-
- | | | |
|------------|---|------------|
| 181 | EU Human Brain Project — Top-Down Vision That Unraveled 2013 – 2023 | G·N |
|------------|---|------------|
-

-
- 198** Launching the BRAIN Initiative — Governance of a Grand Challenge 2011 – 2015 – present **G·K·N**
-

NEUROSCIENCE/CONNECTOMICS

- 78** CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale 2017 – present **H·K·N**
-

NUCLEAR ENGINEERING

- 156** INL / LWRs Control-Room Modernization — Sustainment Research for an Aging Fleet 2010 – present **D·H·K**
-

NUCLEAR POWER

- 154** INL Turbine-Control Upgrade — Low-Burden Cutover as a Human-Factors Finding 2014 **G·K·H**
-

NURSING EDUCATION

- 16** NCSBN National Simulation Study — Licensing the 50% Substitution Rule 2014 **G·K·D**
-

PHARMA SAFETY

- 38** iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms 2006 – 2011 **G·D·N**
-

PUBLIC HEALTH

- 179** The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed 2020 – 2023 **G·K**
-

RAIL TRANSPORT

- 177** CIRAS — Confidential Incident Reporting for UK Rail 1996 – present **G·K·N**
-

SOFTWARE ENGINEERING EDUCATION

- 71** Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development 2025 (preprint) **K·N**
-

SOFTWARE SUSTAINMENT

- 174** Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed 1996 – 2000 **G·D·K**
-

SPACE / AEROSPACE

STANDOUT
FAILURE

Case 140

Challenger & Columbia (1986 / 2003). Normalisation of deviance across two decades and two crews, captured in the Rogers and CAIB reports as a culture that accepted flying with flaws when problems were defined as normal and routine.

NOTE

Case 18

No paired-intervention success in the aerospace dataset. See Defense (Rickover, Case 137) for the safety-engineering counterpart that aerospace did not build.

98	Mars Climate Orbiter — Unit Mismatch	1999	D-K
101	Ariane 5 Flight 501	1996	D-K-H
108	NASA EVA-23 — Water Intrusion Inside a Spacesuit	2013	H-N-D
111	Challenger & Columbia	1986 / 2003	N-K-G
113	Boeing Starliner	2019 – 2024	K-D

SURGERY

30	Surgical Skill Video Peer-Rating Predicts Complications	2013	H-D-K
31	Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units	2009 – 2016	T-K-H

TECHNOLOGY

182	Amazon Hiring AI — Trained Bias, Deprecated	2017 – 2018	D-K-N
-----	---	-------------	-------

TUTORING

77	Hybrid Human-AI Tutoring — Augmentation, Not Delegation	2024	T-K-D
----	---	------	-------

WORKFORCE DEV

68	CIRCUIT — A Scalable, Equity-Centered Research Workforce Model (six cycles)	2017 – 2023	T-K
93	Singapore SkillsFuture — National Workforce Capability at Scale	2015 – present	G-K-D

About this index. Each case is filed under its primary domain, so every case appears exactly once. Cases that span several domains (a healthcare-AI case is also a technology case, a defense-training case is also an

education case) are cross-referenced in their narratives. The standout failures and successes are editorial picks — the case in that domain the editors return to most often in studio.

The cases by LENS course.

Each case names the LEN courses it serves as a worked example for. This index inverts that mapping: for every course in the specialization, the cases that inform it, in case-number order. A case appears under every course it supports, so many appear more than once.

LEN 1

36 cases

6	Racial Bias in Pain Assessment — The False-Belief Mechanism
8	Medical Errors as Systemic Failure
13	Implementation Science in Healthcare — The 17-Year Gap
36	AlphaFold — Protein Structure Prediction
39	TeamSTEPPS
53	inBloom
63	LAUSD iPad Program — The Common Core Technology Project
67	Cognitive Tutor / Carnegie Learning
68	CIRCUIT — A Scalable, Equity-Centered Research Workforce Model
80	Georgia State University Predictive Analytics
97	Boeing 737 MAX / MCAS
102	Air France Flight 447
111	Challenger & Columbia
112	NASA Saturn V Documentation
114	FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule
115	FAA NextGen and the ADS-B Out Transition
116	Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change
117	Crew Resource Management & CAST

124	USS Fitzgerald & USS John S. McCain
138	MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard
139	GIFT and the Adoption Gap
141	DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks
154	INL Turbine-Control Upgrade — Low-Burden Cutover as a Human-Factors Finding
156	INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet
166	Three Mile Island
167	Hurricane Katrina — The FEMA/DHS Response Failure
171	Hurricane Maria — Puerto Rico Response and Logistics Failure
173	Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable
174	Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed
175	INPO and the Nuclear Academy
178	The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami
180	Healthcare.gov Launch
181	EU Human Brain Project — Top-Down Vision That Unraveled
194	Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox)
198	Launching the BRAIN Initiative — Governance of a Grand Challenge
205	The Discipline We Build Next

LEN 2

53 cases

1	Therac-25
2	EHR / CPOE Implementation
5	Epic Sepsis Model
14	Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA
15	I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer
16	NCSBN National Simulation Study — Licensing the 50% Substitution Rule
17	Spaced Education RCTs in Medical Training
18	PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption

-
- 20 TREWS / Sepsis Watch

 - 22 ChatGPT in Healthcare — Hallucination Cases

 - 27 Anesthesia Monitoring Standards and the APSF — The Mortality Decline

 - 29 Bar-Code Medication Administration — Cue at the Point of Care

 - 30 Surgical Skill Video Peer-Rating Predicts Complications

 - 31 Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units

 - 52 Purdue Course Signals — The Reverse-Causality Retention Claim

 - 65 Training Transfer — The Gap Between Learning and Doing

 - 66 Growth-Mindset National Experiment — Heterogeneous Effects

 - 68 CIRCUIT — A Scalable, Equity-Centered Research Workforce Model

 - 69 Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale

 - 70 High-Impact Learning System — Engineering the Environment, Not Just the Event

 - 71 Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development

 - 72 ASSISTments — National Replication and Long-Term Follow-Through

 - 73 The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension

 - 74 Zhang/Scardamalia — Knowledge Building Across Three Cohorts

 - 76 Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account

 - 77 Hybrid Human-AI Tutoring — Augmentation, Not Delegation

 - 78 CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale

 - 84 Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive

 - 85 OU Analyse — Predictive Learning Analytics and Teacher Use at Scale

 - 93 Singapore SkillsFuture — National Workforce Capability at Scale

 - 95 PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop

 - 96 AeroPerú Flight 603

 - 97 Boeing 737 MAX / MCAS

 - 102 Air France Flight 447

 - 104 Asiana Airlines Flight 214

- 105 Helios Airways Flight 522
- 106 TransAsia Airways Flight 235
- 107 Eastern Air Lines Flight 401
- 116 Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change
- 117 Crew Resource Management & CAST
- 118 Korean Air Safety Transformation
- 127 Military Fratricide — Desert Storm to Afghanistan
- 130 USS Vincennes & Iran Air Flight 655
- 131 Stanislav Petrov / 1983 False Alert
- 141 DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks
- 144 Kodak Digital Camera Stagnation
- 146 Northeast Blackout
- 155 Toyota Production System / Andon Cord
- 157 AI-Augmented Coding Tools
- 172 CrowdStrike Falcon Outage
- 183 Uber ATG / Tempe Fatality
- 184 UK Post Office Horizon Scandal
- 195 Tesla Autopilot — Recurring Fatalities

LEN 3

47 cases

- 4 Mid Staffordshire NHS Foundation Trust
- 17 Spaced Education RCTs in Medical Training
- 24 Bristol Heart Babies Reform
- 32 Annual-Screening UI Redesign + CDS at University of Missouri Health Care
- 33 Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal
- 34 COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case
- 37 DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback
- 39 TeamSTEPPS

- | | |
|-----|---|
| 49 | UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days |
| 50 | Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction |
| 55 | Enrollment-Algorithm Yield Optimization Across U.S. Higher Education |
| 57 | GAO Online Program Manager Oversight Gap (GAO-22-104463) |
| 58 | USC x 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC) |
| 72 | ASSISTments — National Replication and Long-Term Follow-Through |
| 73 | The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension |
| 74 | Zhang/Scardamalia — Knowledge Building Across Three Cohorts |
| 75 | Chen et al. — Rural China AI Devices and the Equity-Direction Finding |
| 84 | Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive |
| 85 | OU Analyse — Predictive Learning Analytics and Teacher Use at Scale |
| 86 | Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction |
| 87 | Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models |
| 88 | LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models |
| 90 | Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions |
| 108 | NASA EVA-23 — Water Intrusion Inside a Spacesuit |
| 111 | Challenger & Columbia |
| 117 | Crew Resource Management & CAST |
| 120 | EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation |
| 121 | TCAS — Coordinated Collision Avoidance and the Überlingen Lesson |
| 124 | USS Fitzgerald & USS John S. McCain |
| 126 | F-35 Sustainment & Maintainer Shortage |
| 132 | F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap |
| 134 | V-22 Osprey |
| 135 | 9/11 Intelligence Sharing Failures |
| 137 | U.S. Nuclear Navy / Rickover Training Model |
| 144 | Kodak Digital Camera Stagnation |
| 154 | INL Turbine-Control Upgrade — Low-Burden Cutover as a Human-Factors Finding |

159	Bhopal
161	Texas City BP Refinery Explosion
163	Deepwater Horizon
164	Grenfell Tower
169	Fukushima Daiichi
175	INPO and the Nuclear Academy
182	Amazon Hiring AI — Trained Bias, Deprecated 2017
187	COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate
192	Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model
203	NYC Local Law 144 — Bias Audits as Governance Artifact
205	The Discipline We Build Next
LEN 4	
68 cases	
4	Mid Staffordshire NHS Foundation Trust
5	Epic Sepsis Model
6	Racial Bias in Pain Assessment — The False-Belief Mechanism
7	VA Wait-Time Scandal
8	Medical Errors as Systemic Failure
9	Vioxx Withdrawal
12	ACGME 80-Hour Resident Duty-Hour Reform
15	I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer
18	PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption
19	Keystone ICU / Pronovost Checklist
20	TREWS / Sepsis Watch
21	Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff
23	WHO Surgical Safety Checklist
24	Bristol Heart Babies Reform
25	Removing the Race Coefficient from eGFR

-
- 26 Pulse Oximetry Across Skin Tones

 - 28 Interprofessional Education and the Evidence Gap

 - 32 Annual-Screening UI Redesign + CDS at University of Missouri Health Care

 - 33 Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal

 - 35 Radiology AI Miscalibration

 - 40 Team Science Training for Clinical and Translational Scientists

 - 41 Implementation-Science Training — Stated Goals Outrunning Operational Practice

 - 42 Australian Hospital-Pharmacy Technician Role Redesign

 - 43 Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface

 - 45 Tennessee Voluntary Pre-K Study

 - 46 Algorithmic Bias in Educational Predictive Analytics

 - 48 School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism

 - 51 Atlanta Public Schools Cheating Scandal

 - 60 Houston EVAAS — Value-Added Teacher Evaluation on Trial

 - 61 Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap

 - 62 Gates Intensive Partnerships — Measurement Without a Coupled Lever

 - 65 Training Transfer — The Gap Between Learning and Doing

 - 67 Cognitive Tutor / Carnegie Learning

 - 70 High-Impact Learning System — Engineering the Environment, Not Just the Event

 - 79 The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops

 - 80 Georgia State University Predictive Analytics

 - 81 Open University 'Ethical Use of Student Data' and OU Analyse

 - 82 MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil)

 - 83 Brinkerhoff Success Case Method — Tails as the Evaluation Instrument

 - 89 Data Privacy and Learning Analytics on the African Continent

 - 91 LALA — Building Learning-Analytics Governance Capacity Across Latin America

 - 92 Norway's National Expert Commission on Learning Analytics

 - 93 Singapore SkillsFuture — National Workforce Capability at Scale

-
- 94 Learning Analytics on the African Continent — A Scoping Review as the Present-State Map
 - 95 PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop
 - 109 Colgan Air Flight 3407
 - 110 Atlas Air Flight 3591
 - 117 Crew Resource Management & CAST
 - 119 Aviation Safety Reporting System (ASRS)
 - 123 FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike
 - 140 Navy Surface Warfare Readiness Reform
 - 142 xAPI / Total Learning Architecture — Interoperability Gap
 - 147 Takata Airbag Inflators
 - 148 Wells Fargo Fake Accounts
 - 149 Volkswagen Dieselgate
 - 151 GM Ignition Switch
 - 161 Texas City BP Refinery Explosion
 - 162 Upper Big Branch Mine Explosion
 - 177 CIRAS — Confidential Incident Reporting for UK Rail
 - 178 The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami
 - 179 The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed
 - 186 Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity
 - 189 Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds
 - 190 Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment
 - 193 Bernard Madoff / SEC Failures
 - 196 Fintech Lending Fairness Audit — When Including Race Reduces Disparity
 - 199 Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact
 - 200 CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver
-

LEN 5

84 cases

-
- 1 Therac-25
-

- | | |
|----|---|
| 3 | Watson for Oncology — Delegation Marketed Ahead of Capability |
| 11 | California Nurse Staffing Ratios — Legislating a Capability Requirement |
| 12 | ACGME 80-Hour Resident Duty-Hour Reform |
| 14 | Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA |
| 16 | NCSBN National Simulation Study — Licensing the 50% Substitution Rule |
| 17 | Spaced Education RCTs in Medical Training |
| 19 | Keystone ICU / Pronovost Checklist |
| 21 | Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff |
| 27 | Anesthesia Monitoring Standards and the APSF — The Mortality Decline |
| 29 | Bar-Code Medication Administration — Cue at the Point of Care |
| 30 | Surgical Skill Video Peer-Rating Predicts Complications |
| 31 | Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units |
| 34 | COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case |
| 37 | DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback |
| 38 | iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms |
| 42 | Australian Hospital-Pharmacy Technician Role Redesign |
| 44 | Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD |
| 47 | Algorithmic Bias in Automated Exam Proctoring |
| 49 | UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days |
| 50 | Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction |
| 52 | Purdue Course Signals — The Reverse-Causality Retention Claim |
| 55 | Enrollment-Algorithm Yield Optimization Across U.S. Higher Education |
| 57 | GAO Online Program Manager Oversight Gap (GAO-22-104463) |
| 58 | USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC) |
| 62 | Gates Intensive Partnerships — Measurement Without a Coupled Lever |
| 63 | LAUSD iPad Program — The Common Core Technology Project |
| 66 | Growth-Mindset National Experiment — Heterogeneous Effects |
| 68 | CIRCUIT — A Scalable, Equity-Centered Research Workforce Model |

-
- 69 Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale
-
- 76 Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account
-
- 77 Hybrid Human-AI Tutoring — Augmentation, Not Delegation
-
- 78 CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale
-
- 88 LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models
-
- 90 Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions
-
- 95 PBIS Implementation Fidelity — Why the Framework Travels Only with Its Measurement Loop
-
- 96 AeroPerú Flight 603
-
- 97 Boeing 737 MAX / MCAS
-
- 98 Mars Climate Orbiter — Unit Mismatch
-
- 100 Glass-Cockpit Transition in Light Aircraft — Technology Outran Training
-
- 101 Ariane 5 Flight 501
-
- 102 Air France Flight 447
-
- 103 Kegworth / British Midland 92
-
- 104 Asiana Airlines Flight 214
-
- 105 Helios Airways Flight 522
-
- 106 TransAsia Airways Flight 235
-
- 107 Eastern Air Lines Flight 401
-
- 108 NASA EVA-23 — Water Intrusion Inside a Spacesuit
-
- 109 Colgan Air Flight 3407
-
- 113 Boeing Starliner
-
- 120 EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation
-
- 121 TCAS — Coordinated Collision Avoidance and the Überlingen Lesson
-
- 124 USS Fitzgerald & USS John S. McCain
-
- 126 F-35 Sustainment & Maintainer Shortage
-
- 127 Military Fratricide — Desert Storm to Afghanistan
-
- 128 Operation Eagle Claw
-
- 129 Patriot Missile / Dhahran
-

-
- 130 USS Vincennes & Iran Air Flight 655
-
- 132 F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap
-
- 134 V-22 Osprey
-
- 137 U.S. Nuclear Navy / Rickover Training Model
-
- 138 MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard
-
- 140 Navy Surface Warfare Readiness Reform
-
- 141 DARPA Digital Tutor — Compressing Years of IT Expertise into 16 Weeks
-
- 143 Knight Capital Trading Loss
-
- 153 Equifax Data Breach
-
- 159 Bhopal
-
- 163 Deepwater Horizon
-
- 166 Three Mile Island
-
- 167 Hurricane Katrina — The FEMA/DHS Response Failure
-
- 170 West Africa Ebola — The Delayed International Response
-
- 171 Hurricane Maria — Puerto Rico Response and Logistics Failure
-
- 172 CrowdStrike Falcon Outage
-
- 173 Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable
-
- 178 The UN Cluster Approach — Engineering Humanitarian Coordination After the Tsunami
-
- 180 Healthcare.gov Launch
-
- 182 Amazon Hiring AI — Trained Bias, Deprecated 2017
-
- 185 Air Canada Chatbot Liability — Delegation Without Revocation
-
- 187 COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate
-
- 191 Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict
-
- 192 Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model
-
- 194 Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox)
-
- 201 Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India
-
- 203 NYC Local Law 144 — Bias Audits as Governance Artifact
-

LEN 6

17 cases

8	Medical Errors as Systemic Failure
13	Implementation Science in Healthcare — The 17-Year Gap
53	inBloom
61	Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap
86	Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction
87	Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models
113	Boeing Starliner
116	Eurocat ATM Pilot Modernization — Small-Tier Verification as the Gateway to Big-Tier Change
129	Patriot Missile / Dhahran
139	GIFT and the Adoption Gap
142	xAPI / Total Learning Architecture — Interoperability Gap
151	GM Ignition Switch
154	INL Turbine-Control Upgrade — Low-Burden Cutover as a Human-Factors Finding
176	Tylenol Recall
180	Healthcare.gov Launch
194	Estonia X-Road — Continuous Migration as a Governance Pattern (and the No-Legacy Paradox)
195	Tesla Autopilot — Recurring Fatalities

LEN 7

102 cases

1	Therac-25
2	EHR / CPOE Implementation
3	Watson for Oncology — Delegation Marketed Ahead of Capability
4	Mid Staffordshire NHS Foundation Trust
5	Epic Sepsis Model
6	Racial Bias in Pain Assessment — The False-Belief Mechanism
7	VA Wait-Time Scandal
9	Vioxx Withdrawal

-
- 11 California Nurse Staffing Ratios — Legislating a Capability Requirement

 - 14 Barsuk SBML — Simulation-Based Mastery Learning Dissemination from Northwestern to the VA

 - 15 I-PASS Handoff Bundle — Structuring the Human-to-Human Transfer

 - 16 NCSBN National Simulation Study — Licensing the 50% Substitution Rule

 - 18 PEPFAR HIV Training Across 16 Sub-Saharan African Countries — Modality Comparison Under Disruption

 - 20 TREWS / Sepsis Watch

 - 21 Sterile-Cockpit Ward Rounds — Adapting an Aviation Principle to Clinical Handoff

 - 22 ChatGPT in Healthcare — Hallucination Cases

 - 24 Bristol Heart Babies Reform

 - 25 Removing the Race Coefficient from eGFR

 - 26 Pulse Oximetry Across Skin Tones

 - 27 Anesthesia Monitoring Standards and the APSF — The Mortality Decline

 - 28 Interprofessional Education and the Evidence Gap

 - 29 Bar-Code Medication Administration — Cue at the Point of Care

 - 30 Surgical Skill Video Peer-Rating Predicts Complications

 - 31 Language of Surgery / JIGSAWS — Decomposing Skill into Measurable Units

 - 35 Radiology AI Miscalibration

 - 36 AlphaFold — Protein Structure Prediction

 - 38 iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms

 - 40 Team Science Training for Clinical and Translational Scientists

 - 41 Implementation-Science Training — Stated Goals Outrunning Operational Practice

 - 43 Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface

 - 44 Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD

 - 45 Tennessee Voluntary Pre-K Study

 - 46 Algorithmic Bias in Educational Predictive Analytics

 - 48 School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism

 - 51 Atlanta Public Schools Cheating Scandal

-
- 53 inBloom

 - 54 Summit Learning / Personalized Learning Rollout

 - 60 Houston EVAAS — Value-Added Teacher Evaluation on Trial

 - 65 Training Transfer — The Gap Between Learning and Doing

 - 70 High-Impact Learning System — Engineering the Environment, Not Just the Event

 - 71 Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development

 - 72 ASSISTments — National Replication and Long-Term Follow-Through

 - 73 The Doer Effect at Scale — Replication, AI-Generated Questions, Non-WEIRD Extension

 - 74 Zhang/Scardamalia — Knowledge Building Across Three Cohorts

 - 75 Chen et al. — Rural China AI Devices and the Equity-Direction Finding

 - 76 Multimodal Learning Analytics In-the-Wild — A First-Person Lessons-Learned Account

 - 77 Hybrid Human-AI Tutoring — Augmentation, Not Delegation

 - 78 CIRCUIT / MICrONS — The Human Correction Layer at Petabyte Scale

 - 79 The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops

 - 80 Georgia State University Predictive Analytics

 - 81 Open University 'Ethical Use of Student Data' and OU Analyse

 - 82 MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil)

 - 83 Brinkerhoff Success Case Method — Tails as the Evaluation Instrument

 - 84 Cognitive Tutor Algebra I at Scale — Year-One Null, Year-Two Positive

 - 85 OU Analyse — Predictive Learning Analytics and Teacher Use at Scale

 - 88 LiveHint AI — Evaluating an AI Tutor for Bias Across Foundation Models

 - 89 Data Privacy and Learning Analytics on the African Continent

 - 91 LALA — Building Learning-Analytics Governance Capacity Across Latin America

 - 92 Norway's National Expert Commission on Learning Analytics

 - 94 Learning Analytics on the African Continent — A Scoping Review as the Present-State Map

 - 100 Glass-Cockpit Transition in Light Aircraft — Technology Outran Training

 - 108 NASA EVA-23 — Water Intrusion Inside a Spacesuit

111 Challenger & Columbia

114 FAA Aging–Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule

115 FAA NextGen and the ADS-B Out Transition

120 EGPWS / TAWS — Closing the CFIT Category in Commercial Aviation

121 TCAS — Coordinated Collision Avoidance and the Überlingen Lesson

123 FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike

132 F-22 OBOGS Hypoxia — The Sustainment-Era Instrumentation Gap

133 Mark 14 Torpedo Failures

135 9/11 Intelligence Sharing Failures

138 MIL-STD-1472H — Human Engineering as a Binding Acquisition Standard

147 Takata Airbag Inflators

148 Wells Fargo Fake Accounts

149 Volkswagen Dieselgate

151 GM Ignition Switch

152 TSB Bank IT Migration

153 Equifax Data Breach

156 INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet

159 Bhopal

160 Davis-Besse Reactor Head Corrosion

161 Texas City BP Refinery Explosion

162 Upper Big Branch Mine Explosion

164 Grenfell Tower

165 Camp Fire / PG&E

167 Hurricane Katrina — The FEMA/DHS Response Failure

169 Fukushima Daiichi

171 Hurricane Maria — Puerto Rico Response and Logistics Failure

173 Navy SUBSAFE — Requirements as a Non-Negotiable Sustainment Deliverable

174 Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed

176	Tylenol Recall
177	CIRAS — Confidential Incident Reporting for UK Rail
180	Healthcare.gov Launch
181	EU Human Brain Project — Top-Down Vision That Unraveled
184	UK Post Office Horizon Scandal
186	Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity
189	Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds
193	Bernard Madoff / SEC Failures
195	Tesla Autopilot — Recurring Fatalities
196	Fintech Lending Fairness Audit — When Including Race Reduces Disparity
197	Predictive Policing — PredPol
198	Launching the BRAIN Initiative — Governance of a Grand Challenge

LEN 8

90 cases

7	VA Wait-Time Scandal
11	California Nurse Staffing Ratios — Legislating a Capability Requirement
12	ACGME 80-Hour Resident Duty-Hour Reform
13	Implementation Science in Healthcare — The 17-Year Gap
28	Interprofessional Education and the Evidence Gap
32	Annual-Screening UI Redesign + CDS at University of Missouri Health Care
33	Alert-Fatigue Redesign — Cutting EHR Alerts Without Cutting the Safety Signal
39	TeamSTEPPS
40	Team Science Training for Clinical and Translational Scientists
41	Implementation-Science Training — Stated Goals Outrunning Operational Practice
42	Australian Hospital-Pharmacy Technician Role Redesign
43	Rwanda mHealth Maternal Care — Community Health Workers as the Capability Interface
44	Japan PMDA — Pre-Approved Change Management as Regulatory Architecture for AI/SaMD
47	Algorithmic Bias in Automated Exam Proctoring

-
- 48 School Surveillance and Black Student Outcomes — Infrastructure as the Mechanism

 - 49 UK Ofqual A-Level Algorithm — National-Scale Grading Replaced by Algorithm, Withdrawn in Days

 - 50 Wisconsin DEWS — A Decade of Algorithmic Dropout Prediction

 - 52 Purdue Course Signals — The Reverse-Causality Retention Claim

 - 54 Summit Learning / Personalized Learning Rollout

 - 55 Enrollment-Algorithm Yield Optimization Across U.S. Higher Education

 - 57 GAO Online Program Manager Oversight Gap (GAO-22-104463)

 - 58 USC × 2U Online MSW — When the Delegation Becomes the Product (Luna v. USC)

 - 61 Sold a Story — The Science of Reading and the Decades-Long Research-Practice Gap

 - 62 Gates Intensive Partnerships — Measurement Without a Coupled Lever

 - 63 LAUSD iPad Program — The Common Core Technology Project

 - 66 Growth-Mindset National Experiment — Heterogeneous Effects

 - 69 Duolingo Half-Life Regression — Spaced Repetition at Consumer Scale

 - 71 Reflective Practice on a Work-Based Software Engineering Program — Longitudinal Capability Development

 - 75 Chen et al. — Rural China AI Devices and the Equity-Direction Finding

 - 79 The Kirkpatrick Chain of Evidence — Where Corporate L&D Stops

 - 81 Open University 'Ethical Use of Student Data' and OU Analyse

 - 82 MMALA — A Maturity Model for Responsible Learning-Analytics Adoption (Brazil)

 - 83 Brinkerhoff Success Case Method — Tails as the Evaluation Instrument

 - 90 Algorithmic College Admissions — Vendors' Claims vs. Applicants' Perceptions

 - 91 LALA — Building Learning-Analytics Governance Capacity Across Latin America

 - 92 Norway's National Expert Commission on Learning Analytics

 - 93 Singapore SkillsFuture — National Workforce Capability at Scale

 - 94 Learning Analytics on the African Continent — A Scoping Review as the Present-State Map

 - 98 Mars Climate Orbiter — Unit Mismatch

 - 101 Ariane 5 Flight 501

 - 103 Kegworth / British Midland 92

-
- 109 Colgan Air Flight 3407
 - 110 Atlas Air Flight 3591
 - 111 Challenger & Columbia
 - 112 NASA Saturn V Documentation
 - 113 Boeing Starliner
 - 114 FAA Aging-Aircraft Program — Widespread Fatigue Damage and the Part 26 Rule
 - 115 FAA NextGen and the ADS-B Out Transition
 - 118 Korean Air Safety Transformation
 - 119 Aviation Safety Reporting System (ASRS)
 - 122 Singapore Airlines Safety Transformation
 - 123 FSF CFIT and ALAR Task Forces — Industry-Level Institution Building After a Spike
 - 124 USS Fitzgerald & USS John S. McCain
 - 126 F-35 Sustainment & Maintainer Shortage
 - 127 Military Fratricide — Desert Storm to Afghanistan
 - 128 Operation Eagle Claw
 - 129 Patriot Missile / Dhahran
 - 133 Mark 14 Torpedo Failures
 - 134 V-22 Osprey
 - 135 9/11 Intelligence Sharing Failures
 - 137 U.S. Nuclear Navy / Rickover Training Model
 - 139 GIFT and the Adoption Gap
 - 142 xAPI / Total Learning Architecture — Interoperability Gap
 - 146 Northeast Blackout
 - 151 GM Ignition Switch
 - 152 TSB Bank IT Migration
 - 155 Toyota Production System / Andon Cord
 - 156 INL / LWRS Control-Room Modernization — Sustainment Research for an Aging Fleet
 - 157 AI-Augmented Coding Tools
-

- 160 Davis-Besse Reactor Head Corrosion
- 163 Deepwater Horizon
- 164 Grenfell Tower
- 165 Camp Fire / PG&E
- 169 Fukushima Daiichi
- 174 Y2K Remediation — The Aging-System Transition That Worked Because It Was Believed
- 175 INPO and the Nuclear Academy
- 177 CIRAS — Confidential Incident Reporting for UK Rail
- 181 EU Human Brain Project — Top-Down Vision That Unraveled
- 182 Amazon Hiring AI — Trained Bias, Deprecated 2017
- 185 Air Canada Chatbot Liability — Delegation Without Revocation
- 187 COMPAS Recidivism Prediction — Calibration vs. Equal Error Rate
- 190 Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment
- 191 Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict
- 192 Apple Card / Goldman Sachs — When the Lender Cannot Explain Its Own Model
- 198 Launching the BRAIN Initiative — Governance of a Grand Challenge
- 199 Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact
- 200 CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver
- 201 Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India
- 203 NYC Local Law 144 — Bias Audits as Governance Artifact
- 205 The Discipline We Build Next

LEN 9

30 cases

- 2 EHR / CPOE Implementation
- 3 Watson for Oncology — Delegation Marketed Ahead of Capability
- 25 Removing the Race Coefficient from eGFR
- 26 Pulse Oximetry Across Skin Tones
- 34 COMPOSER Sepsis Prediction — The Third Clinical-AI Sepsis Case

35	Radiology AI Miscalibration
36	AlphaFold — Protein Structure Prediction
37	DeepMind Mammography — High-Profile Nature Paper, Replicability Pushback
38	iPLEDGE Isotretinoin REMS — When the Authorization Mechanism Underperforms
46	Algorithmic Bias in Educational Predictive Analytics
47	Algorithmic Bias in Automated Exam Proctoring
67	Cognitive Tutor / Carnegie Learning
86	Gándara — Detecting and Mitigating Algorithmic Bias in College Student-Success Prediction
87	Yu / Lee / Kizilcec — Protected Attributes in Learning-Analytics Models
89	Data Privacy and Learning Analytics on the African Continent
100	Glass-Cockpit Transition in Light Aircraft — Technology Outran Training
129	Patriot Missile / Dhahran
143	Knight Capital Trading Loss
170	West Africa Ebola — The Delayed International Response
179	The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed
185	Air Canada Chatbot Liability — Delegation Without Revocation
186	Algorithmic Mortgage Lending — Omitting the Variable Did Not Fix the Disparity
189	Dutch SyRI — Welfare-Fraud Risk Scoring Halted on Rights Grounds
190	Cruise's Partial Disclosure — How Disclosure Posture Decides Deployment
191	Australia Robodebt — Algorithmic Debt-Recovery and the Royal Commission Verdict
196	Fintech Lending Fairness Audit — When Including Race Reduces Disparity
197	Predictive Policing — PredPol
199	Waymo's Safety Case Framework — Governance Objection Dissolved by Designed Artifact
200	CPUC AV Passenger-Service Permits — Conditions as a Designed Objection-Dissolver
201	Aadhaar Exclusion — Biometric Welfare Delegation and Its Judicial Reckoning in India

LEN 10

21 cases

8	Medical Errors as Systemic Failure
---	------------------------------------

-
- 12 ACGME 80-Hour Resident Duty-Hour Reform
-
- 13 Implementation Science in Healthcare — The 17-Year Gap
-
- 19 Keystone ICU / Pronovost Checklist
-
- 22 ChatGPT in Healthcare — Hallucination Cases
-
- 23 WHO Surgical Safety Checklist
-
- 39 TeamSTEPPS
-
- 45 Tennessee Voluntary Pre-K Study
-
- 53 inBloom
-
- 54 Summit Learning / Personalized Learning Rollout
-
- 60 Houston EVAAS — Value-Added Teacher Evaluation on Trial
-
- 133 Mark 14 Torpedo Failures
-
- 139 GIFT and the Adoption Gap
-
- 140 Navy Surface Warfare Readiness Reform
-
- 155 Toyota Production System / Andon Cord
-
- 157 AI-Augmented Coding Tools
-
- 164 Grenfell Tower
-
- 170 West Africa Ebola — The Delayed International Response
-
- 176 Tylenol Recall
-
- 179 The Johns Hopkins COVID-19 Dashboard — Evidence Infrastructure at Decision Speed
-
- 205 The Discipline We Build Next
-